

The Art of SEO

Mastering Search Engine Optimization

Eric Enge, Stephan Spencer,
and Jessie C. Stricchiola

3rd
Edition



Early Release
RAW & UNEDITED

In 2004, when I did my first professional work in the field of SEO, Google was estimated to receive ~200 million searches each day. In 2014, that number jumped to more than 6 billion, a 30X increase over 10 years. Since its inception, web search has been a powerful tool for people to find what they need and, as a result, it's also been a powerful channel for those seeking to earn attention, traffic, actions, and customers. But in 2004, discovering how search engines worked and how to drive that traffic was a daunting challenge.

Web forums, blogs, and a handful of industry news sites were the first to take on the problem. They built communities and published resources to help both industry insiders and those aspiring to learn Search Engine Optimization. That's where I first cut my teeth. But learning from those sources was hard work - and there were no alternatives. Schools didn't teach SEO (they still don't!), online courses were extremely rare (and many were low quality or straight-up misleading), and there were almost no books on the subject.

When I founded the SEOMoz blog (which became the company known today as "Moz"), my mission was to educate and to learn by doing so. Nothing has made me a better student of SEO and of marketing, than being forced to write about and to teach it to others.

That's why it was a great honor to have been one of the contributing writers for the original two editions of this book - The Art of SEO. When we finished that first full draft, I remember skimming through it and thinking "man... I wish they'd had this when I was learning SEO." That was a good feeling.

Search engines are complex. The web is massively complex. And the human beings doing the searching, sifting, clicking, and converting -- they're the most complex of all.

But all that complexity shouldn't create an impenetrable wall for those seeking knowledge. In the past, when it has, the reputation of SEO itself has suffered. CMOs and marketing managers and small business owners have hired or contracted professionals to perform SEO and been frustrated by the process, the requirements, and the results, often because they themselves didn't know enough about the practice to make a good choice or to create the right expectations.

By pulling back the veil on SEO, those of us tasked with teaching are enabling the right expectations to be set, the right people to be hired, and the right results from SEO investments. Arguably, no one is more qualified or has done more than the team of writers, editors, and practitioners behind this new edition of The Art of SEO. In this book, you'll find not only the wisdom of its authors, but the work of thousands around the world of search engine optimization who've contributed their insight directly (and are properly cited, of course) and indirectly (we thank you, too, unacknowledged teachers of us all) to transparency in a sometimes too-opaque field.

While I could not personally contribute as I would have liked to this edition, it was my honor and privilege to once again grace these pages (even if only in this paltry introduction) alongside such an excellent team. Welcome to the Art of SEO.

Sincerely,
Rand Fishkin
December 2014

Search: Reflecting Consciousness and Connecting Commerce

Search has become integrated into the fabric of our society. With more than 197.9 billion searches performed worldwide each month as of April 2014 (according to comScore, <http://www.comscore.com>), approximately 6.6 billion web searches are performed every day. This means that on average about 7,500 searches are performed every single second of every day. In addition, users have grown to expect that the responses to their search queries will be returned in less than one second.

If you have ever come across the various “SEO is dead” claims that make the rounds periodically in the online marketing world, rest assured: SEO is both alive and well - and has not, as many feared, been “killed” by social media marketing as a driver of web traffic. In fact, in November, 2014, Twitter (<https://twitter.com/>) announced that by renewing their focus on SEO, they increased the number of logged-out visitors coming to their site tenfold – from 7.5 million to 75 million visitors per month: <http://searchengineland.com/twitter-seo-more-visitors-208160>.

The high demand for search continues to exist because people can now obtain in mere seconds information that 20 years ago would have required a trip to the library, the use of a card catalog and the Dewey Decimal System, and a foot search through halls of printed volumes—a process that could easily have consumed two hours or more. Through the new channel of search, people can conduct many of their shopping, banking, and social transactions online—something that has changed the way our global population lives and interacts.

This dramatic shift in behavior represents what investors like to label a *disruptive event*—an event that has changed something in a fundamental way. Search engines have been at the center of this disruptive event, and having a business’s website content appear prominently in the search engines when people are looking for the service, product, or resource it provides is critical to the survival of that business. As is the case with most paths to success, obtaining such prime search result real estate is not a simple matter, but it is one that this book aims to deconstruct and demystify as we examine, explain, and explore the ever-changing art of search engine optimization (SEO).

The Mission of Search Engines

Since web searchers are free to use any of the many available search engines to find what they are seeking, the burden is on the search engines to develop a relevant, fast, and fresh search experience. For the most part, search engines accomplish this by being *perceived* as having the most relevant results and delivering them the fastest, as users will go to the search engine they think will get them the answers they want in the least amount of time.

As a result, search engines invest a tremendous amount of time, energy, and capital in improving their relevance. This includes performing extensive studies of user responses to their search results, comparing their results against those of other search engines, conducting eye-tracking and click-through rate studies (discussed later in this chapter), and constructing PR and marketing campaigns.

Search engines generate revenue primarily through paid advertising. The great majority of this revenue comes from a *pay-per-click* (or *cost-per-click*) model, in which the advertisers pay only for users who click on their ads. Because the search engines' success depends so greatly on the relevance of their search results, manipulations of search engine rankings that result in non-relevant results (generally referred to as *spam*) are dealt with very seriously. Each major search engine employs a team of people who focus solely on finding and eliminating spam from their search results (generally referred to as "web spam" teams). This matters to SEO professionals because they need to be careful that the tactics they employ will not be considered spam by the search engines, as this would carry the risk of resulting in penalties for the websites they work on.

The Market Share of Search Engines

Figure 1-1 shows the US market share for search engines in September 2014, according to comScore. As you can see, Google is the dominant search engine on the Web in the United States

comScore Explicit Core Search Share Report*			
September 2014 vs. August 2014			
Total U.S. – Home & Work Locations			
Source: comScore qSearch			
Core Search Entity	Explicit Core Search Share (%)		
	Aug-14	Sep-14	Point Change
<i>Total Explicit Core Search</i>	100.0%	100.0%	N/A
Google Sites	67.3%	67.3%	0.0
Microsoft Sites	19.4%	19.4%	0.0
Yahoo Sites	10.0%	10.0%	0.0
Ask Network	2.0%	2.0%	0.0
AOL, Inc.	1.3%	1.3%	0.0

***"Explicit Core Search" excludes contextually driven searches that do not reflect specific user intent to interact with the search results.*

[[figs/print/0101.png]]

Figure 1-1. Search engine market share (September 2014)

In many European countries, the disparity is even greater. However, in some markets Google is not dominant. In China, for instance, Baidu is the leading search engine. The fact remains, however, that in most world markets, a heavy focus on Google is a smart strategy for SEO.

The Human Goals of Searching

The basic goal of a human searcher is to obtain information relevant to a specific set of keywords and/or phrases entered into a search field, also known as a query. A searcher may formulate their query as a question, but the vast majority of searches are performed by users simply entering word combinations – leaving the search engines to do the work of determining “intent.” One of the most important elements to building an SEO strategy for a website is developing a thorough understanding of the psychology of your target audience, and how they use words and concepts to obtain information about the services and/or products you provide. Once you understand how the average searcher—and, more specifically, your target market—utilizes query-based search engines, you can more effectively reach and keep those users.

Search engine usage has evolved over the years, but the primary principles of conducting a search remain largely unchanged. Most search engine use is comprised of the following steps:

1. Users experience the need for information. The user may be looking for information on a specific website, and they will search for that website; (a navigational query); the user might want to buy something (a transactional query), or they might want to learn something (an informational query). We will discuss this in more detail in the following section.
2. Users formulate that need using a string of words and phrases (comprising the query). Most people formulate their queries in one to three words, though as users are becoming more web savvy, their queries can become longer to receive more specific results more quickly. Table 1-1 gives a more detailed look at the percentages of searches per query length as of August, 2011, the most recent comScore study on this aspect of search.
3. Users execute the query, check the results, and if they seek further additional information, then will try a refined query.

Table 1-1. Searches by query length (comScore, April 2014 data)

Words	Percent of searches
1	25.8%
2	22.8%
3	18.7%
4	13.2%
5+	19.5%

When this process results in the satisfactory completion of a task, a positive experience is created for the user, the search engine, and the site providing the information or result.

Who Searches and What Do They Search For?

comScore reported that the number of search queries performed worldwide on the Web was approximately 197.9 billion across all search engines in April 2014.

comScore data also shows over 1.7 billion people were using a search engine on a given day in that month. Search engine users in the US were slightly more likely to be women than men (50.8% versus 49.2%). According to comScore, about two-thirds of US Internet users had an income of \$40,000 or more (as shown in Table 1-2).

Table 1-2. Internet users by household income (April 2014)

US household income	Internet users
Less than \$15,000	15,649 (6.9%)
\$15,000–\$24,999	14,313 (6.3%)
\$25,000–\$39,999	26,318 (11.5%)

\$40,000–\$59,999	36,557 (16%)
\$60,000–\$74,999	25,763 (11.3%)
\$75,000–\$99,999	35,028 (15.3%)
\$100,000 or more	74,626 (32.7%)

You can find additional data from studies, surveys, and white papers on Search Engine Land's Stats & Behaviors page (<http://searchengineland.com/library/stats-search-behavior>).

All of this research data leads us to some important conclusions about web search and marketing through search engines. For example:

- Search is a major source of valuable, targeted traffic.
- Google is the dominant search player in most world markets.
- Users tend to use shorter search phrases, but these are gradually getting longer, especially for non-navigational queries.
- Search covers all types of markets.

Search is undoubtedly still one of the best and most important ways to reach consumers and build a business, regardless of that business's size, reach, or focus.

Determining Searcher Intent: A Challenge for Search Marketers and Search Engines

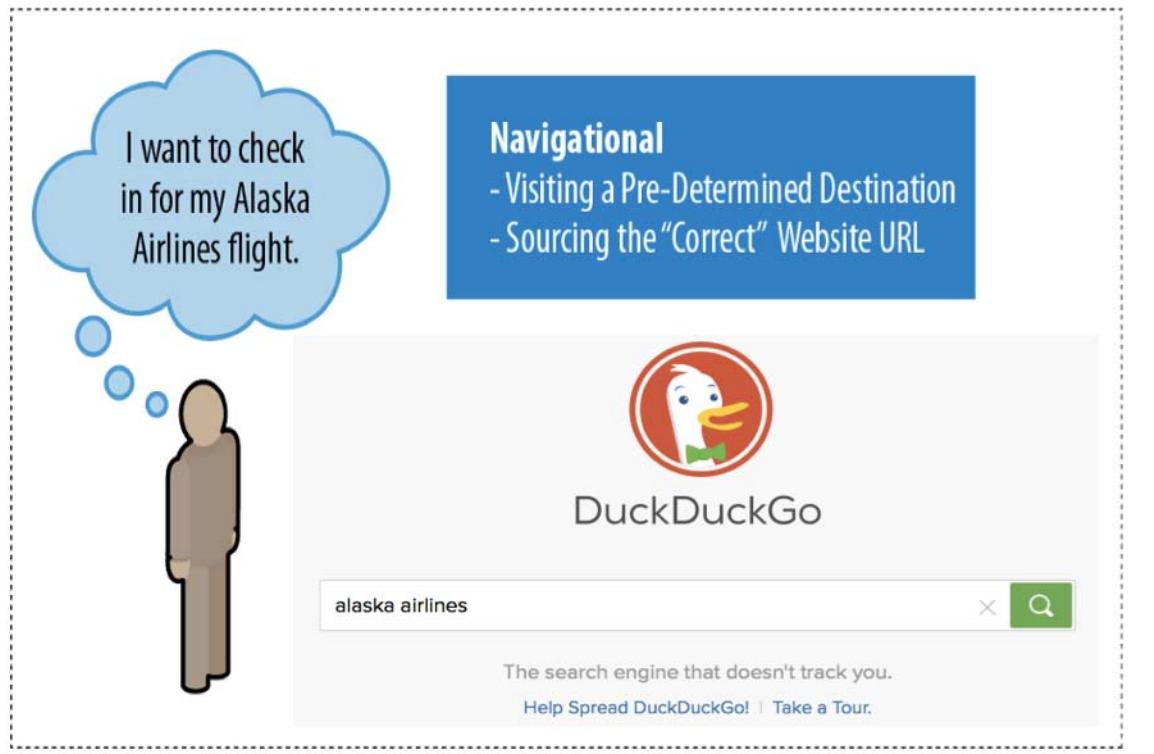
Good marketers are empathetic, and smart SEO practitioners - as well as the search engines - have a common goal of providing searchers with results that are relevant to their queries. Therefore, a crucial element to building an online marketing strategy around SEO and organic (sometimes referred to as "natural") search visibility is understanding your audience and how they think about, discuss, and search for your service, product, and brand.

Search engine marketers need to be aware that search engines are *tools*—resources driven by intent towards a content destination. Using the search box is fundamentally different from entering a URL into the browser's address bar, clicking on a bookmark, or clicking on a link to go to a website. Searches are performed with *intent* - the user wants to find specific information, rather than just land on it by happenstance; search is also different from "browsing," or clicking around links on a web page. Because of this specificity, search traffic is generally of higher value to marketers than these other types of web traffic.

What follows is an examination of the different types of search queries and their categories, characteristics, and processes.

Navigational Queries

Navigational searches are performed with the intent of going directly to a specific website. In some cases, the user may not know the exact URL, and the search engine serves as the "White Pages." Figure 1-2 shows an example of a navigational query.



[[figs/print/0102.png]]

Figure 1-2. Navigational query

Opportunities: Pull searcher away from destination; get ancillary or investigatory traffic.

Average traffic value: Very high when searches are for the publisher's own brand. These types of searches tend to lead to very high conversion rates. However, these searchers are already aware of the company brand, and some percentage of these queries may not represent new customers. For brands other than the one being searched on, the click-through rates will tend to be low, but this may represent an opportunity to take a customer away from a competitor.

Informational Queries

Informational searches involve an incredibly broad range of queries. Consider the many types of information people might look for: local weather, driving directions, a celebrity's recent interview, disease symptoms, self-help information, how to train for a specific type of career... the possibilities are as endless as the human capacity for thought. Informational searches are primarily non-transaction-oriented (although they can include researching information about a product or service); the information itself is the goal, and no interaction beyond clicking and reading is required for the searcher's query to be satisfied. Figure 1-3 shows an example of an informational query.



[[figs/print/0103.png]]

Figure 1-3. *Informational query*

Opportunities: Brand searchers with positive impressions of your site, information, company, and so on; attract inbound links; receive attention from journalists/researchers; potentially convert to sign up or purchase.

Average traffic value: The searcher may not be ready to buy anything just yet, or may not even have a long-term intent to buy anything, so the value tends to be “medium,” at best. However, many of these searchers will later enter in a more targeted search, and this represents an opportunity to capture mindshare with those potential customers. For example, informational queries that are focused on researching commercial products or services can have high value.

Transactional Queries

Transactional searches don’t necessarily have to involve a credit card or immediate financial transaction. Creating a Pinterest account, signing up for a free trial account at DomainTools (<http://www.domaintools.com>), or finding the best local Japanese restaurant for dinner tonight are all transactional queries. Figure 1-4 shows an example of a transactional query.



[[figs/print/0104.png]]

Figure 1-4. Transactional query

Opportunities: Achieve transaction (financial or other).

Average traffic value: Very high.

Research from Pennsylvania State University and the Queensland University of Technology (http://ist.psu.edu/faculty_pages/jjanssen/academic/pubs/jansen_user_intent.pdf) shows that more than 80% of searches are informational in nature, and only about 10% of searches are navigational or transactional.

The researchers went further and developed an algorithm to automatically classify searches by query type. When they tested the algorithm, they found that it was able to correctly classify queries 74% of the time. The difficulty in classifying the remaining queries was vague user intent—that is, the queries could have multiple meanings. Here is additional academic research on this topic:

- <http://blog.strategynode.com/keyword-search-intent/>

Adaptive Search

The search engines also look at sequences of search queries to determine intent. This was confirmed in co-author Eric Enge's interview with Jack Menzel, Product Management Director for Google Search (<http://www.stonetemple.com/how-google-does-personalization-with-jack-menzel/>). You can verify this by trying search sequences such as a search on *Rome* followed by a search on *hotels*.

Normally, a search on *hotels* would not include results for hotels in Rome, but when the preceding query was for *Rome*, some results for hotels in Rome will be included. Keeping track of users' previous search queries and taking them into account when determining which results to return for a new query—known as *adaptive search*—is intended to help the search engines get a better sense of a user's intent. The search engines need to do this with care: excessive changes to the results they return based on recent query history are likely to lead to problems, so usually these types of changes are fairly limited in scope. Nonetheless, it is useful to be aware of the types of sequences of searches that users go through in their quest for information.

Leveraging Searcher Intent

When you are building keyword research charts for clients or on your own sites, it can be incredibly valuable to determine the intent of each of your primary keywords. Table 1-3 shows some examples.

Table 1-3. Sample search queries and intent

Term	Queries	Intent	\$\$ value
<i>Beijing Airport</i>	980	Nav	Low
<i>Hotels in Xi'an</i>	2,644	Info	Mid
<i>7-Day China tour package</i>	127	Trans	High
<i>Sichuan jellyfish recipe</i>	53	Info	Low

This type of analysis can help to determine where to place ads and where to concentrate content and links.

Hopefully, this data can help you to think carefully about how to serve different kinds of searchers based on their individual intents, and how to concentrate your efforts in the best possible areas.

Although informational queries are less likely to immediately convert into sales, this does not mean you should forego pursuing rankings on these queries; getting your informative content in front of users seeking information can be incredibly valuable, and can turn users into potential customers. If you are able to build a relationship with users who find your site after an informational query, they may be more likely to come to you to make a related purchase at a later date, and they may also decide to share your information with others via their own website, or through social media engagement – an indirect, but potentially more valuable result than converting the single user into a paying customer.

One problem in search is that when most searchers formulate their search queries their input is limited to just a handful of (generally one to four) words. Since most people don't have a keen understanding of how search engines work, they often provide queries that are too general or that are presented in a way that does not provide the search engine (or the marketer) with what it needs to determine, with 100% accuracy 100% of the time, their specific intent. Additionally, search engine users may not have a specific intent for an individual search query beyond "curiosity" about a general trending topic, or subject matter. While this poses potential difficulty in delivering relevant results, it also poses great opportunity to capture the mind of someone who may not know what they are looking for, specifically, but who takes an interest in the subsequent variety of results the search engine (and search marketers) deliver in response.

These types of general queries are important to most businesses because they often get the brand and site on the searcher's radar, and this initiates the process of building trust with the user. Over time, the user will move on to more specific searches that are more transactional or navigational in nature.

If, for instance, companies buying pay-per-click (PPC) search ads bought only the high-converting navigational and transactional terms and left the informational ones to competitors, they would lose market share to those competitors. Over the course of several days, a searcher may start with *digital cameras*, hone in on *Olympus OMD*, and then ultimately buy from the store that showed up in his or her search for *digital cameras* and pointed him or her in the direction of the Olympus OMD model.

Given the general nature of how query sessions start, though, determining intent is quite difficult, and it can result in searches being performed where the user does not find what he wants - even after multiple tries. A July 2013 ForeSee Results (<http://www.foreseeresults.com>) study for the American Customer Satisfaction Index (ACSI) found that 77% of Google users and 76% of Bing users were satisfied with their experiences.

While 77% satisfaction is an amazing accomplishment given the complexity of building a search engine, this study still showed that more than 22% of users were not satisfied with their search results; additionally, these numbers represent the lowest levels of customer satisfaction with search engines since 2003, and could reflect users' dissatisfaction with the number of ads that are infiltrating search engine results pages (SERPs). Danny Sullivan of Search Engine Land provides a broad analysis of these issues: <http://searchengineland.com/google-sees-lowest-score-ever-in-customer-satisfaction-survey-167600>. It will

be interesting to see how these satisfaction numbers evolve over time. The important takeaway here is that in all instances, it is still difficult to determine exactly a searcher's intent; and when the searcher's intent is broad, there is ample opportunity to leverage your content assets with SEO.

As an SEO practitioner, you should be aware that some of the visitors that you succeed in attracting to your site may have arrived for the wrong reasons (i.e., they were really looking for something else), and these visitors are not likely to help your business goals. Part of your task as an SEO is to maintain a high level of relevance in the content placed on the pages you manage, to help minimize this level of waste – while still attempting to maximize SERP “mindshare” and gain brand exposure.

How People Search

Search engines invest significant resources into understanding how people use search, enabling them to produce better (i.e., faster, fresher, and more relevant) search engine results. For website publishers, the information regarding how people use search can be used to help improve the usability of a site as well as search engine compatibility.

User interactions with search engines can also be multistep processes. Witness the user search session documented by Microsoft and shown in Figure 1-5.

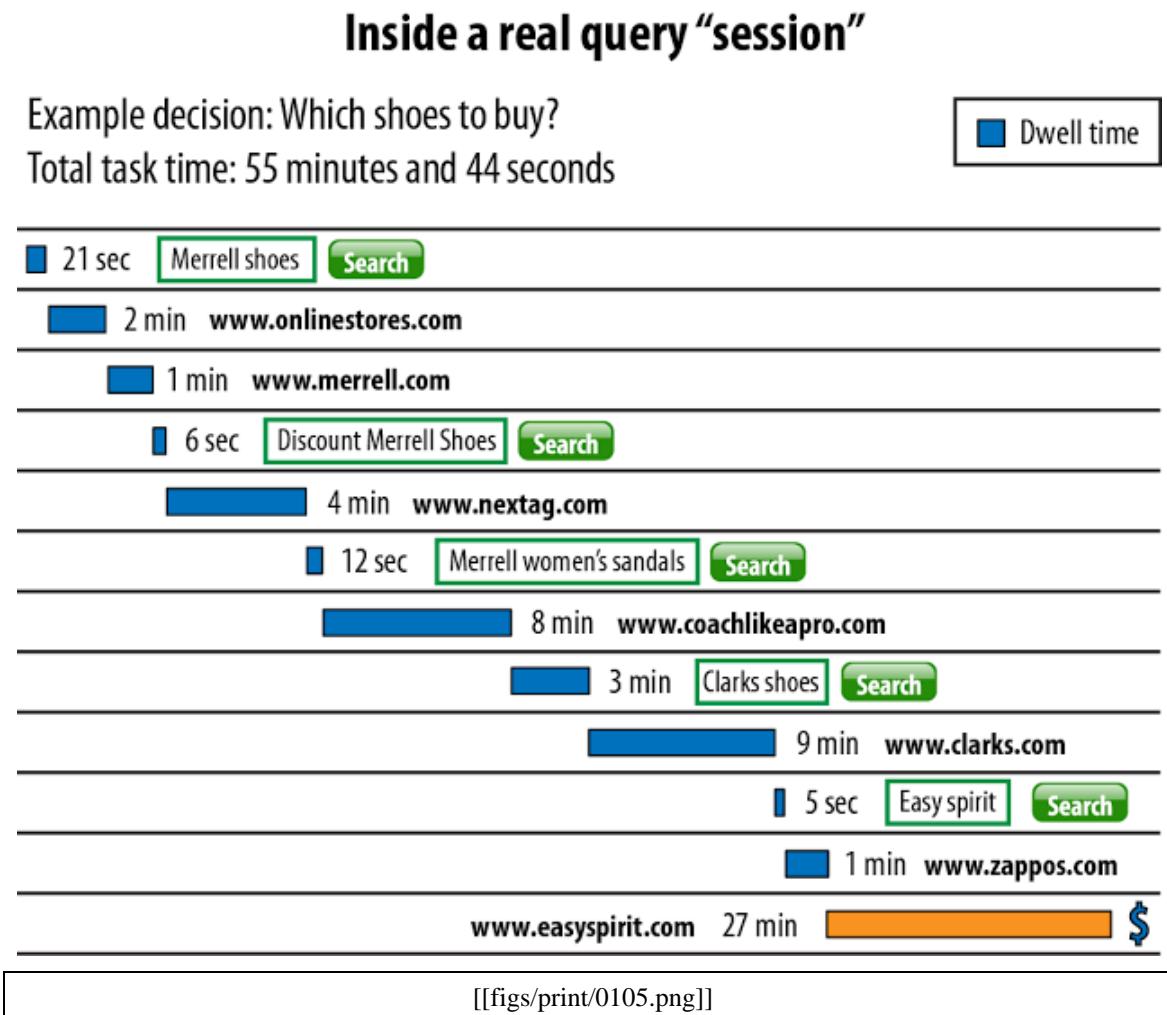


Figure 1-5. Merrell shoes user search session

In this sequence, the user performs five searches over a 55+ minute period before making a final selection. The user is clearly trying to solve a problem and works at it in a persistent fashion until the task is done.

Marin Software provided us with data on one consumer durable retailer (whose products represent high-cost, considered purchases) for whom 50% of the orders involved more than 10 clicks leading up to the conversion event.

For this particular retailer, when you look at the number of different ad groups that were clicked on in those 10 clicks, the clicks were mostly on the same keyword. In fact, for more than 75% of all conversions that came from multiple paid clicks, all the clicks were from the same ad group. Only 7% of conversions came from three different ad groups (and none from more than that).

Table 1-5 shows the average delay between the first click received by the site and the resulting purchase for this example retailer.

Table 1-4. Delay between first click and purchases

Delay between first click and purchases	Percentage of users
Same day	50%
2 to 7 days	9%
8 to 30 days	12%
31 to 90 days	26%
More than 90 days	3%

This behavior pattern indicates that people are thinking about their tasks in stages. As in our Merrell shoes example in Figure 1-5, people frequently begin with a general term and gradually get more specific as they get closer to their goal. They may also try different flavors of general terms. In Figure 1-5, it looks like the user did not find what she wanted when she searched on *Merrell shoes*, so she then tried *discount Merrell shoes*. You can then see her refine her search, until she finally settles on Easy Spirit as the type of shoe she wants.

This is just one example of a search sequence, and the variety is endless. Figure 1-6 shows another search session, once again provided courtesy of Microsoft.

Typical session: health

10 unique queries
7 partial re-queries and refinements
57 minutes

2:58	Google	nausea headache pregnant what's wrong
		pregnancy.families.com
2:58		nausea headache pregnant 3rd trimester
	2 min	answers.yahoo.com/question
3:00		Define malaise
	1 min	www.chefjournal.org/content
3:01		Nausea headache pregnant 3rd trimester Headache pregnant 3rd trimester
	5 min	answers.yahoo.com/question
3:06		gestational diabetes gestational diabetes personal story
	16 min	www.experienceproject.com
3:22		pregnancy complications stories
	4 min	www.diabetictmommy.com/articles.html
3:26		gestational diabetes itchy skin
	14 min	www.steadyhealth.com/Diabetes
3:40		first signs of gestational diabetes
	15 min	www.articlebase.com; www.man-health-fitness-solutions.com; www.healthspy.net; www.diabeteslife.com/diabetes; www.optimumdiabetics.com; www.everydayhealth.com; www.squidoo.com; mascan-files-anti.com/200092/scan
		[[figs/print/0106.png]]

Figure 1-6. Health user search session

In this search session, the user has a health concern. This particular user starts with a five-word search, which suggests that she may have some experience using search engines. At 3:01 her search on *headache pregnant 3rd trimester* leads her to Answers.yahoo.com (<http://answers.yahoo.com>). After visiting this site, her search suddenly gets more specific.

She begins to focus on gestational diabetes, perhaps because something she saw on Answers.yahoo.com (<http://answers.yahoo.com>) led her to believe she may have it. The session culminates in a search for *first signs of gestational diabetes*, which suggests that she has concluded that this is quite possibly the issue she is facing.

The session stops there. It may be that at this point the user feels she has learned what she can. Perhaps her next step is to go to her doctor with her concerns, prepared to ask a number of questions based on what she has learned.

Our next search session example begins with a navigational search, where the user simply wants to locate the travel website Orbitz.com (<http://www.orbitz.com>) (see Figure 1-7). The user's stay there is quite short, and she progresses to a search on *Cancun all inclusive vacation packages*. Following that she

searches on a few specific resorts and finally settles on *cancun riviera maya hotels*, after which it appears she may have booked her hotel—the final site visited on that search is Bookings.occidentalhotels.com (<http://bookings.occidentalhotels.com>), and the direction of her searches changes after that.

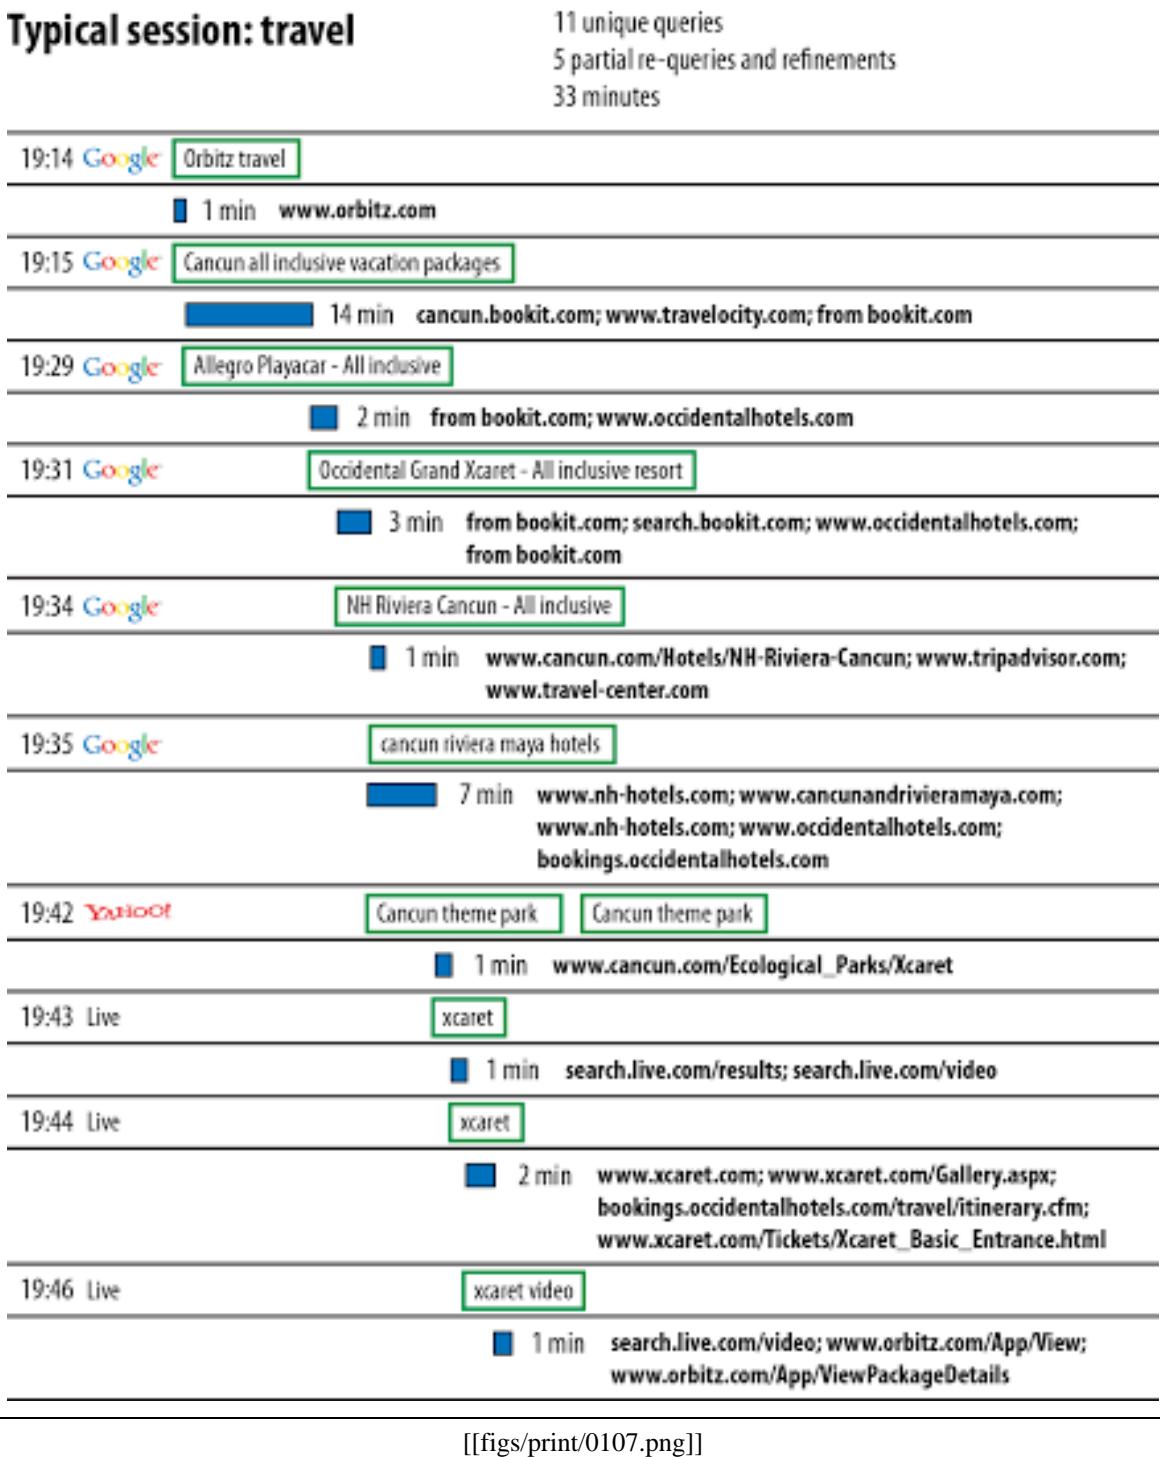


Figure 1-7. Travel user search session

At that point, the user begins to look for things to do while she is in Cancun. She conducts a search for *cancun theme park* and then begins to look for information on *xcaret*, a well-known eco park in the area.

Users traverse countless different scenarios when they are searching for something. These example search sessions represent traditional desktop interactions.

Recent data from mobile search shows different behavior for mobile searchers, who are more likely to be close to completing a transaction. Data from a May 2011 eMarketer study showed that 55% of people visited a business they found in the search results after searching for information on their smartphone devices. Search engines do a lot of modeling of these different types of scenarios to enable them to provide better results to users. The SEO practitioner can benefit from a basic understanding of searcher behavior as well. We will discuss this in more detail in Chapter 2.

How Search Engines Drive E-Commerce

People make use of search engines for a wide variety of purposes, with some of the most popular being to research, locate, and buy products. Ecommerce sales reported by the US Census Bureau were a healthy \$71.2 billion (http://www.census.gov/retail/mrts/www/data/pdf/ec_current.pdf) in the first quarter of 2014. Forrester Research forecasts that US ecommerce retail sales will reach \$370 billion by 2017, outpacing sales growth at brick and mortar stores:



[[figs/print/0108.png]]

Figure 1-8. Forrester Research Online Retail Forecast to 2017

It is important to note that search and offline behavior have a heavy degree of interaction, with search playing a growing role in driving offline sales. In April 2014, Google announced the launch of a pilot program within their AdWords system called “In-Store Attribution Transaction Reporting,” geared towards solving the online search/offline conversion attribution dilemma by teaming up with large data providers Axiom Corp. and DataLogix Holdings Inc. to combine cookie data with offline personal information databases. It will be interesting to see how this program performs, and to what extent it helps search marketers close the attribution gap between search and brick and mortar conversions.

The Mobile Shift

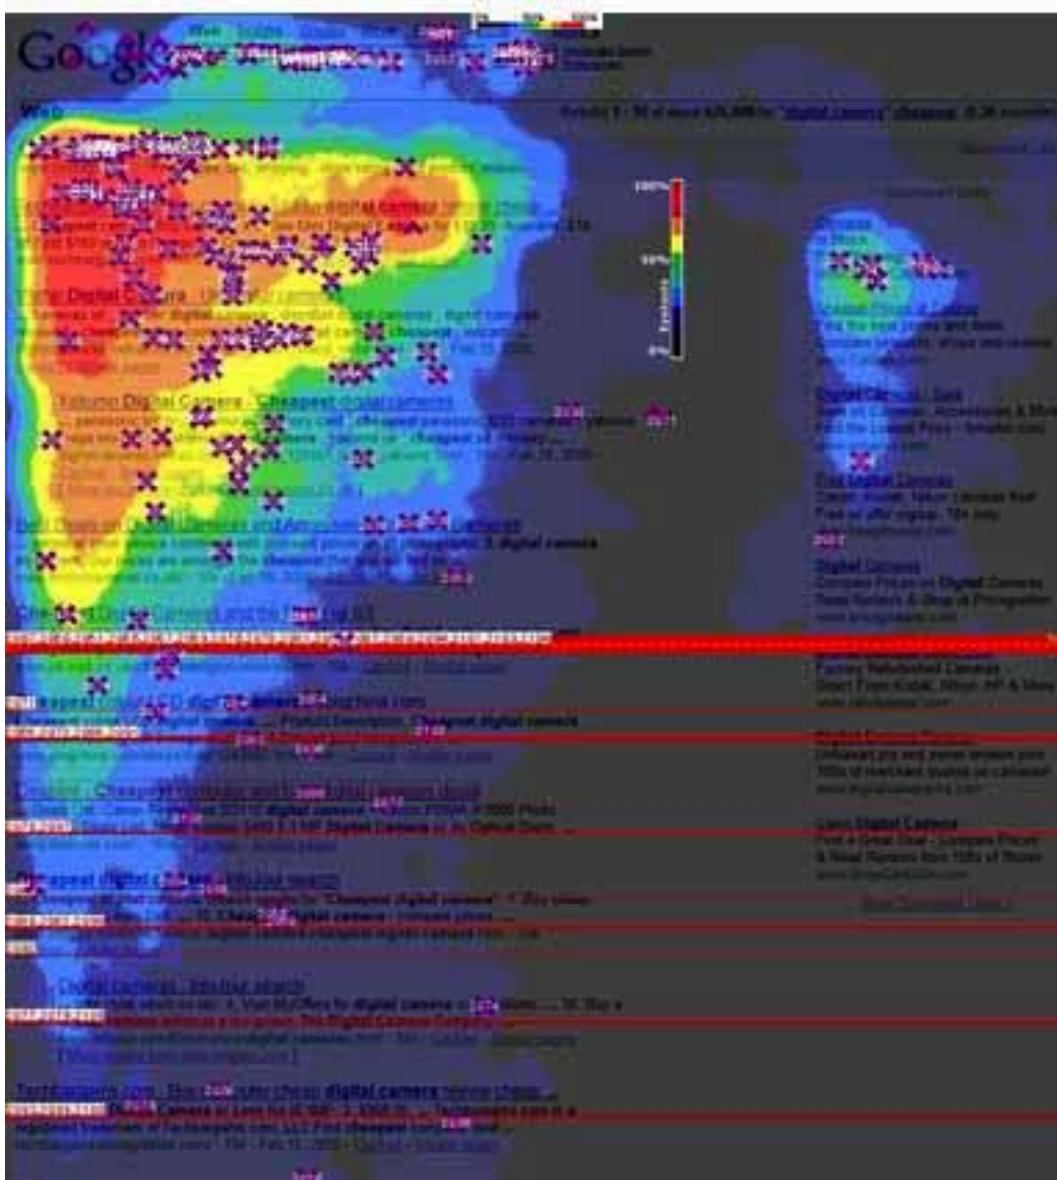
In March 2015, eMarketer published a study showing that mobile advertising spend continues to increase at the expense of desktop advertising, and projects that by 2019, mobile ad spending will rise to \$65.87 billion, or 72.2% of total US digital advertising spend (<http://www.emarketer.com/Article/Mobile-Will-Account-72-of-US-Digital-Ad-Spend-by-2019/1012258>). Figure 1-9 shows the steep upward trend projection of mobile ad spend through 2019, reflecting consumers’ increasing use of mobile devices and tablets for search – and search marketers are taking notice.

[[figs/print/0109.png]]

Figure 1-9. Forrester Research Online Retail Forecast to 2017

The History of Eye Tracking: How Users Scan Results Pages

Way back in 2006, research firm Enquiro (now called Mediactive) conducted heat-map testing with search engine users (<http://www.mediactive.com/research/eyetrackingreport.asp>) that produced fascinating results related to what users see and focus on when engaged in search activity. Figure 1-9 depicts a heat map showing a test performed on Google – the graphic indicates that users spent the most amount of time focusing their eyes in the top-left area, where shading is the darkest. This has historically been referred to in search marketing as the “Golden Triangle.”



[[figs/print/0110.png]]

Figure 1-10. Enquiero eye-tracking results

This particular study perfectly illustrated how little attention has traditionally been paid to results lower on the page versus those higher up, and how users' eyes are drawn to bold keywords, titles, and descriptions in the natural (generally referred to throughout this book as "organic") results versus the paid search listings, which receive comparatively little attention. It also showed that different physical positioning of on-screen search results resulted in different user eye-tracking patterns. When viewing a standard Google results page, users tended to create this "F-shaped" pattern with their eye movements, focusing first and longest on the upper-left corner of the screen, then moving down vertically through the first two or three results, across the page to the first paid page result, down another few vertical results, and then across again to the second paid result. (This study was done only on left-to-right language search results—results for Chinese, Hebrew, and other non-left-to-right-reading languages would be different.)

In May 2008, Google introduced the notion of Universal Search (discussed in more detail in Chapter 10), which was a move from simply showing the 10 most relevant web pages (referred to as "10 blue links") to

showing other types of media, such as videos, images, news results, and so on, as part of the results of a basic web search. The other search engines followed suit within a few months, and the industry now refers to this general concept as *Blended Search*.

Blended Search created more of a chunking effect of user eye attention, with the focus “chunks” hovering over various rich media objects such as images or video. Understandably, our eyes are drawn to the image first, then to look at the text beside it to see whether it corresponds to the image or video thumbnail. Based on an updated eye tracking study published by Enquiro in September 2007, Figure 1-10 shows what the eye-tracking pattern on a Blended Search page looked like:

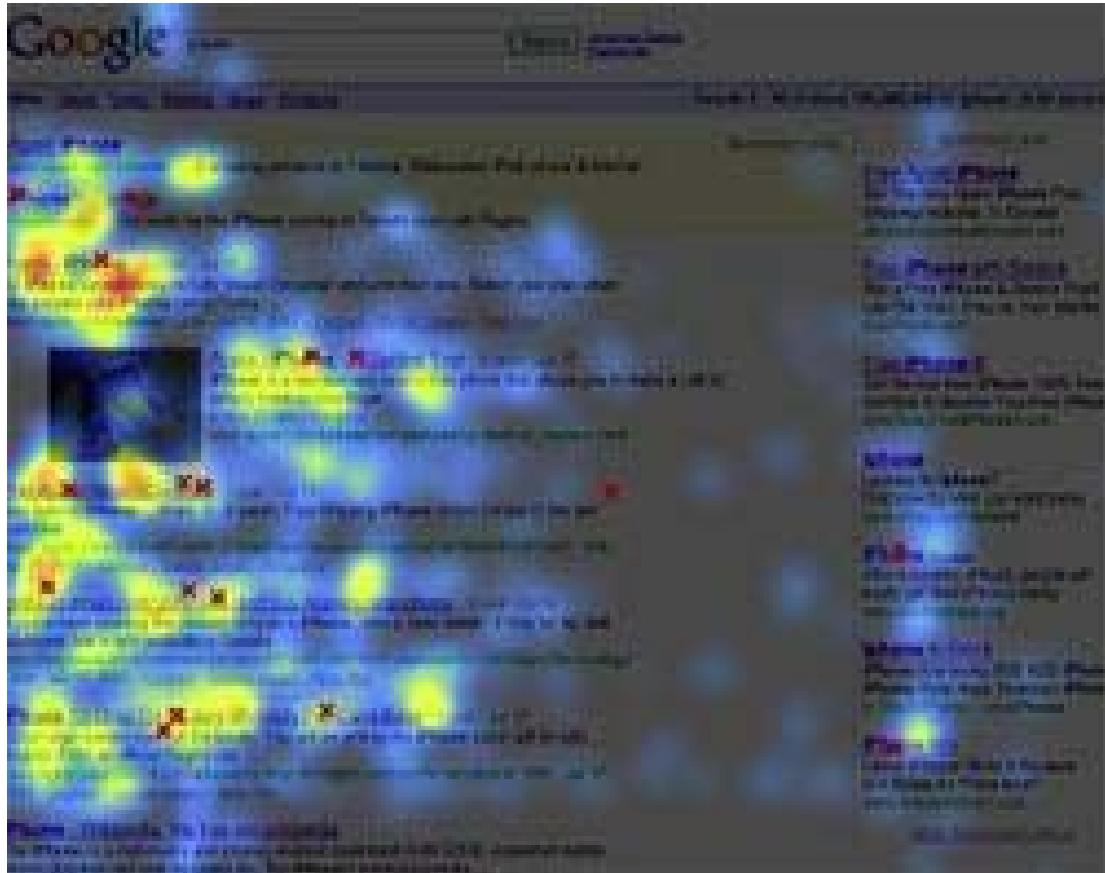
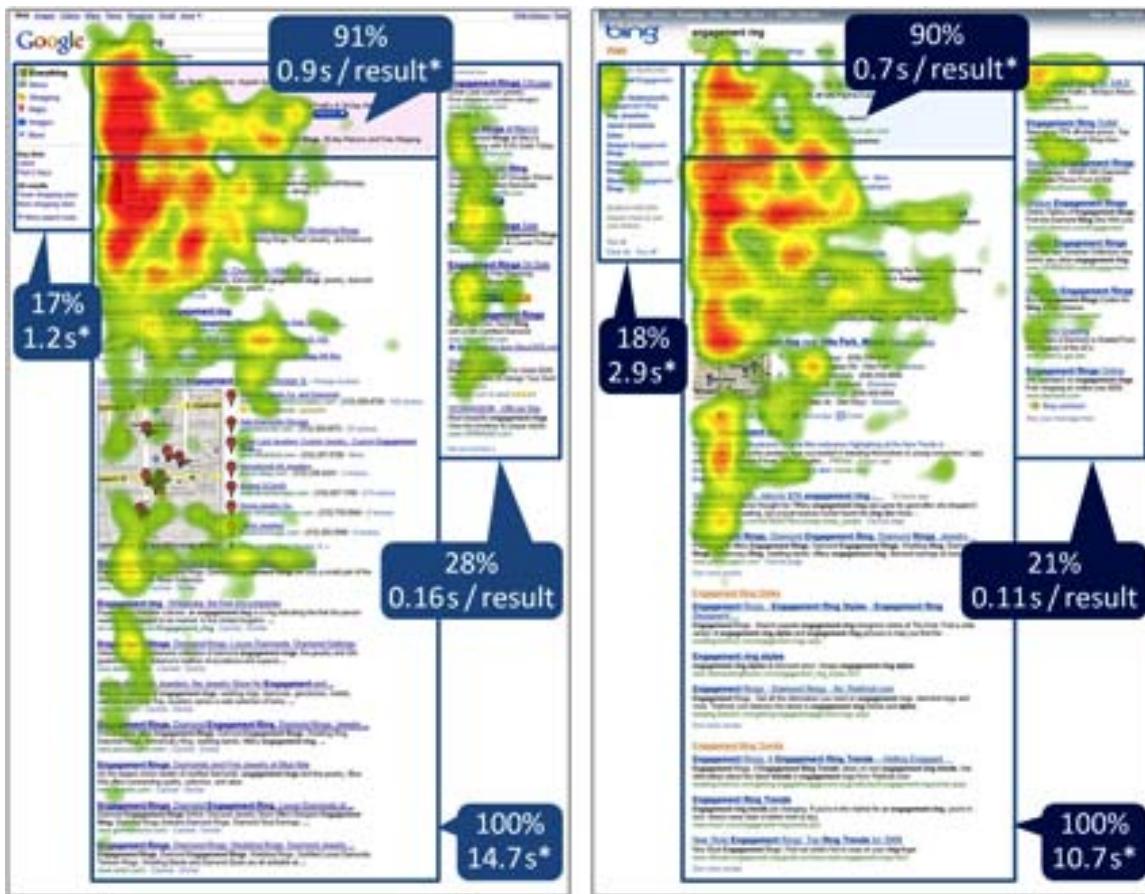


Figure 1-11. Enquiro eye-tracking results, Blended Search

A January 2011 study performed by User Centric (<http://www.usercentric.com/news/2011/01/26/eye-tracking-bing-vs-google-second-look>) showed similar results, as shown in Figure 1-11.



Heatmaps showing the aggregate gaze time of all 24 participants on Google (left) and Bing (right) for one of the transactional tasks. The red color indicates areas that received the most total gaze time (4.5 seconds and above). Each callout includes the percentage of participants who looked at the area and the time (in seconds) they spent looking there. The numerical data are an average across all four tasks. Asterisks indicate values that were significantly different between Google and Bing at $\alpha = .1$.

[[figs/print/0112.png]]

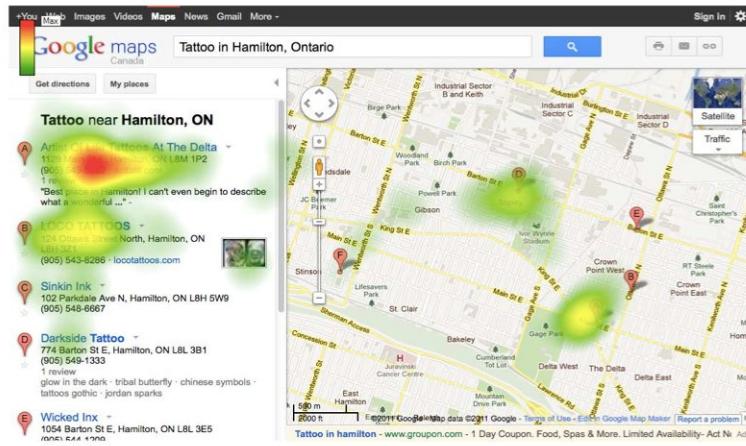
Figure 1-12. User Centric eye-tracking results

In 2010, Enquiero investigated the impact of Google Instant on search usage and attention (<http://ask.enquiero.com/2010/eye-tracking-google-instant/>), noting that for queries in their study:

- Percent of query typed decreased in 25% of the tasks, with no change in the others
- Query length increased in 17% of the tasks, with no change in the others
- Time to click decreased in 33% of the tasks and increased in 8% of the tasks

In more recent studies, we can see how the continuously changing search results are shifting how users focus on a results page.

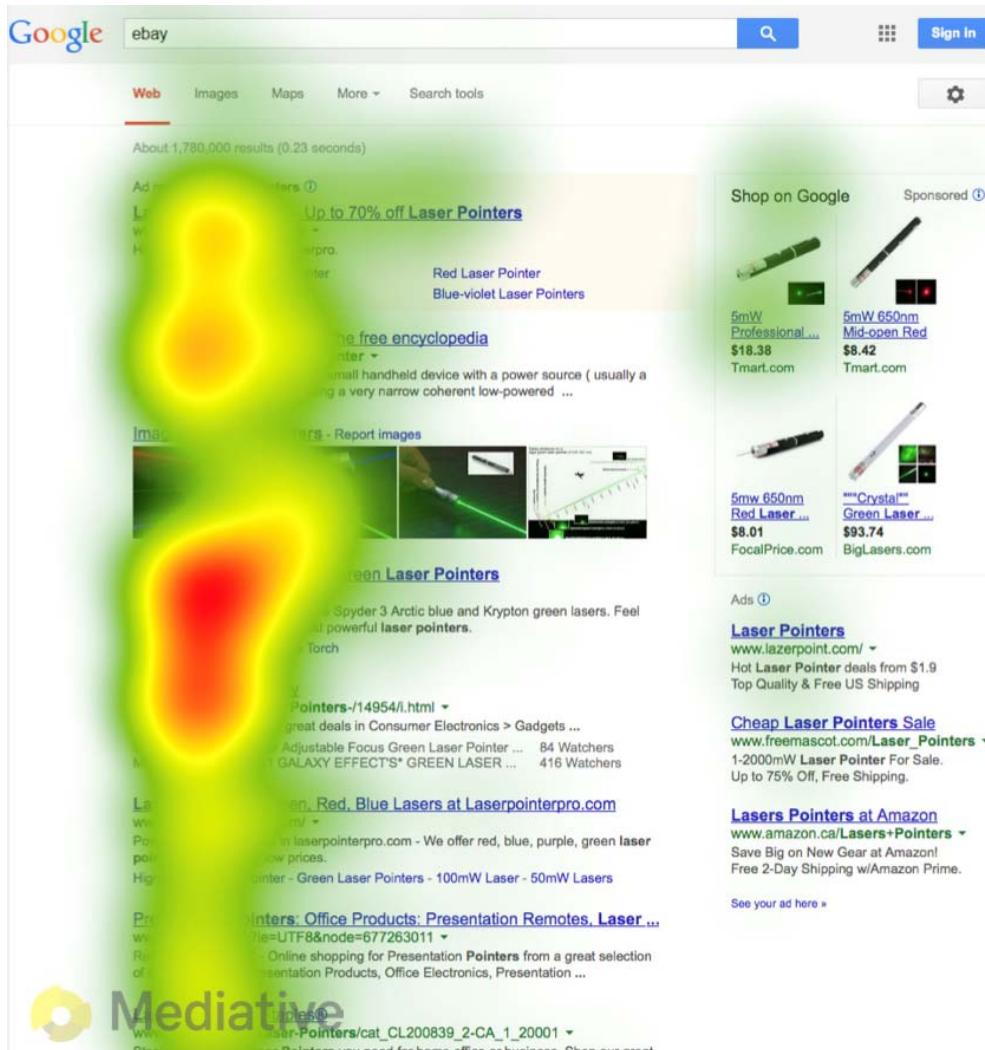
In October 2012, Mediative conducted a new eye tracking study analyzing the areas of attention on Google Maps results pages for locally-targeted searches, and, similar to a Google Web search results page, users eyes were most drawn to the top-most search result:



[[figs/print/0113.png]]

Figure 1-13. User Centric eye-tracking results

A more recent study by Mediactive, issued in 2014, called “The Evolution of Google Search Results Pages and Their Effect on User Behavior” (<http://pages.mediactive.com/SERP-Research>) shows us a fascinating progression away from the “Golden Triangle” as the search results – as well as search platforms (mobile!) – evolve. Figure 1-13 below is the heatmap from this study for an updated Google SERP showing how user attention is moving more vertically than horizontally. Some have theorized that this shift in attention towards a vertical, rather than horizontal scan, is due to the increased adoption of mobile search by mobile device users, while others think that the ever-changing SERP makeup that pushes the “blue text link” results lower on the page (such as the increase in the number prominence and content of top-of-page paid search ads, as well universal and local listings for various queries) is contributing to this change.



[[figs/print/0114.png]]

Figure 1-14. Google SERP - 2014 Eye Tracking Results, Vertical Scanning

Another notable trend in user eye attention is that because the #1 organic site's actual position on the SERP is moving lower, sites that are positioned lower on the results page are seeing more click activity than in years past – making ranking in organic at these lower positions more valuable than in the past as well. Figure 1-14 below highlights the increased attention being placed lower on the page:

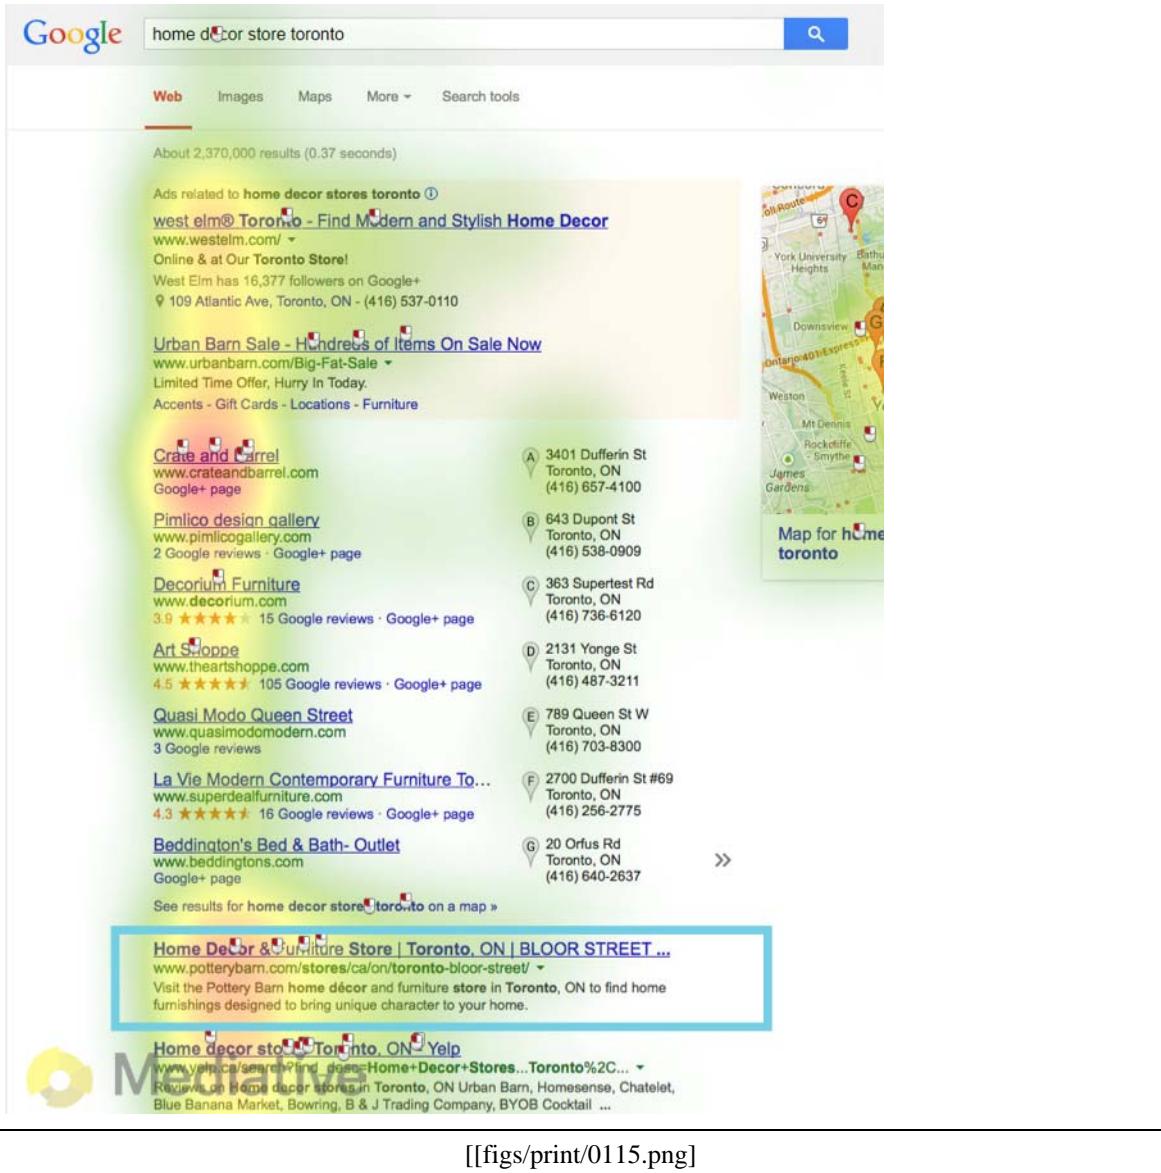


Figure 1-15. Google SERP - 2014 Eye Tracking Results, Lower SERPs Get More Attention

These types of studies are a vivid reminder of how important search engine results pages really are. And as the eye-tracking research demonstrates, as the search engines – and therefore the search results – continue to evolve, users' search and engagement patterns will continue to follow suit. There will be more items on the page for searchers to focus on, more ways for searchers to remember and access the search listings, and more interactive, location-based delivery methods and results layouts that will keep changing as other search environments and platforms, such as mobile, continue to evolve.

Figure 1-15 below shows a “Mega SERP” published by Moz (<http://moz.com>) – a display of the current universe of potential results within a Google SERP.

Google The Amazing World of Tacos

A Best tacos in Chicago

B Taco songs

C Tacos for New Years Eve

D Tacos for New Years Eve on Google

E 156 calories

F Tacos la Mex - Mexican

G All Day, Every Day - Orange, CA 92701

H Tacos

I How To Make A Taco - YouTube

J Santa Fe (Carmel), White Farms Market

K Texas Johnston - Austin

L Taco Doctor in Decatur, Atlanta, Georgia - Food Network

M Mexico (D.F.), Mexico

N El Tacos De Juan Carlos - Mexican

O Tacos la Mex - Mexican

P Mi Jardines, Av. 20, Mexico, Mexico - Mexican

Q In-depth articles

R Mexican Health News

S Search results for Tacos

T Shoutbox Mexico

U Taco

V Taco Bell

W Counterattack Mexico

X See results about

Y Map for pumping tacos

Google >   

Google Home | Adwords Program | Business Solutions | Privacy & Safety | Help | Google+ | Google News | Adwords Help | Business Help | Privacy & Safety | Help | Google News

Figure 1-16. User Centric eye-tracking results

Click Tracking: How Users Click on Results, Natural Versus Paid

By now, you should be convinced that you want your site content to be prominently displayed within SERPs. It never hurts to be #1 in the natural search results.

In contrast, data shows that you may *not* want to be #1 in the paid search results, because the resulting cost to gain the #1 position in a PPC campaign can reduce the total net margin on your campaign. A study released by AdGooroo in June 2008 (http://www.adgooroo.com/how_keyword_length_and_ad_posi.php) found that:

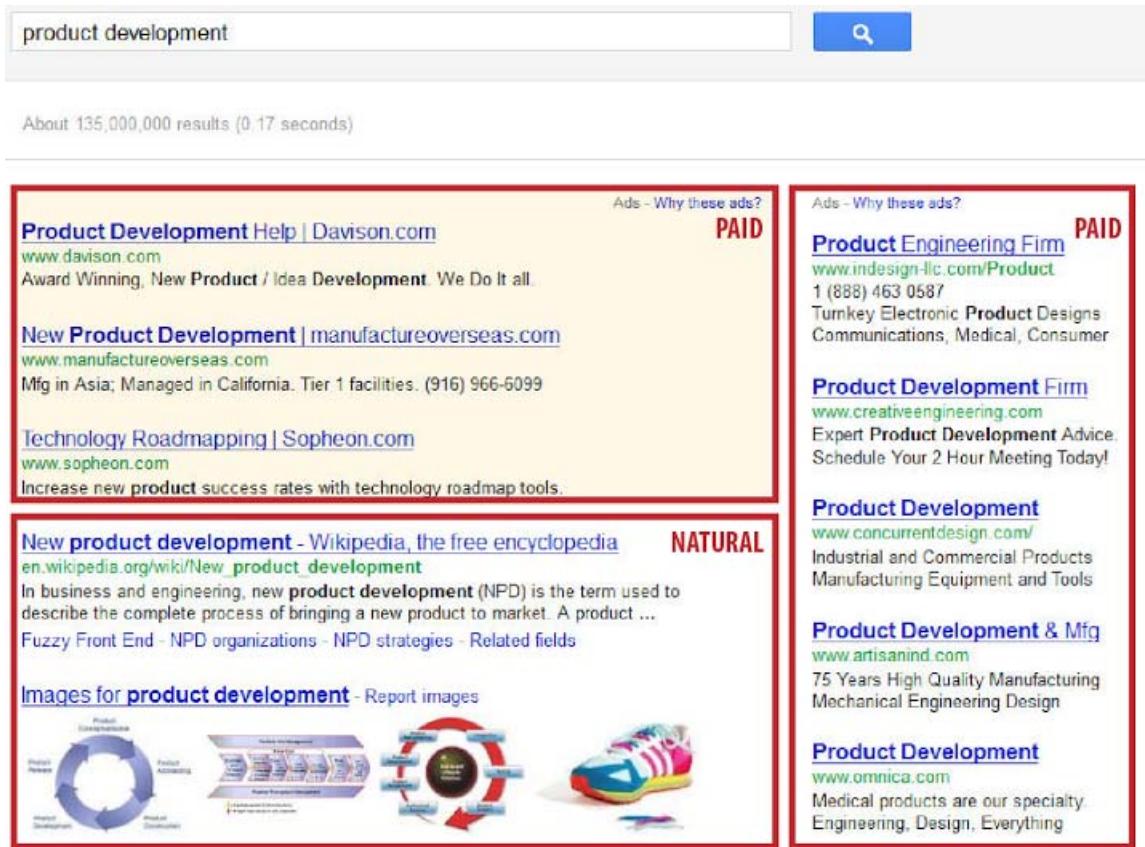
Bidding for top positions usually makes financial sense only for high-budget, brand-name advertisers. Most other advertisers will find the optimal position for the majority of their keywords to lie between positions 5–7.

Of course, many advertisers may seek the #1 position in paid search results, for a number of reasons. For example, if they have a really solid backend on their website and are able to make money when they are in the #1 position, they may well choose to pursue it. Nonetheless, the data from the survey suggests that there are many organizations for which being #1 in paid search does not make sense.

Even if your natural ranking is #1, you can still increase the ranking page's click rate by having a sponsored ad above it or in the right hand column. The AdGooroo survey showed that having a prominent paid ad on the same search results page makes your #1 natural ranking receive as much as 20% more clicks.

Distribution of Search Results and Traffic

To start breaking this down a bit, Figure 1-16 shows the screen real estate occupied by the two types of search results. This screenshot was taken prior to Google's January 2012 Search, plus Your World announcement, but is the type of screen layout related to studies that will help us understand which portions of the search results receive the most clicks.



[[figs/print/0117.png]]

Figure 1-17 Paid and natural search results

This example from Google shows how the paid results appear above and to the right of the natural search results. Note that Google often does not show paid results above the natural results, in which case the paid results show up only on the right.

Your position in the results has a huge impact on the traffic you will receive. Studies on the impact of SERP position have shown widely varying results, but do agree that the advantage of higher positions is significant.

In 2013, Catalyst Search Marketing published the most recent comprehensive study of Google CTR (click-through rate), which shows CTR data for both desktop vs. mobile searches, as well as for branded vs. unbranded searches (<http://www.catalystsearchmarketing.com/pubs/google-ctr-study/>). In this study, Catalyst found that the top 4 Google SERP positions received 83% of first page organic clicks, and that 48% of searches result in a page one organic click. Figure 1-17 below shows the average CTR (percentage of clicks received) by Google SERP position:

AVERAGE POSITION	AVERAGE CTR
1.00	17.16
2.00	9.94
3.00	7.64
4.00	5.31
5.00	3.50
6.00	1.63
7.00	1.09
8.00	1.04
9.00	0.44
10.00	0.51

[[figs/print/0118.png]]

Figure 1-18- Google CTR by Position

This data provides insight around platform CTR differentiators (desktop vs. mobile), searcher intent CTR differentiators (branded vs. non-branded), and query length CTR differentiators (one, two, three, and four+ word query data).

In Figure 1-18 below, we can see that the more refined a user's search query is – that is, the more words used in the query – the greater the likelihood that the user will click on search results farther down the page:

AVERAGE POSITION	ONE WORD	TWO WORDS	THREE WORDS	FOUR+ WORDS
1.00	17.68	16.32	16.28	18.76
2.00	4.41	8.30	10.58	12.41
3.00	3.56	6.86	7.93	8.64
4.00	1.50	4.10	5.12	6.38
5.00	1.86	2.64	3.61	4.11
6.00	0.37	1.10	1.60	2.18
7.00	0.00	1.08	1.12	1.36
8.00	0.00	1.13	1.45	0.84
9.00	0.08	0.52	0.36	0.53
10.00	0.44	0.42	0.54	0.62

[[figs/print/0119.png]]

Figure 1-19 – Google CTR by Position & Query Length

As the search engines, search engine results pages, and searcher behavior evolve, the practice of search engine optimization must evolve in tandem.

Conclusion

Search has penetrated the very fabric of global society. The way people learn, work, share, play, shop, research, socialize and interact has changed forever – and organizations, causes, brands, charities, individuals – almost all entities – need to have an internet presence of some kind, and they need the search engines – and search functionality within all platforms – to generate exposure and facilitate engagement. As our society moves ever closer to a professional consumer (“prosumer”) economy, the ways in which people create, publish, distribute, and ultimately find information and resources on the web will continue to be of great importance. This book will investigate further just how search, and therefore search engine optimization, is at the center of the web ecosystem – and is still your key to success in the ever-evolving web economy.

2

Search Engine Basics

In this chapter, we will begin to explore how search engines work. Building a strong foundation on this topic is essential to understanding the SEO practitioner's craft.

As we discussed in Chapter 1, people have become accustomed to receiving nearly instantaneous answers from search engines after they have submitted a search query. In Chapter 1 we also discussed the volume of queries (more than 7,500 per second), and that as early as 2008 Google knew about 1 trillion pages on the Web (<http://googleblog.blogspot.com/2008/07/we-knew-web-was-big.html>). At Pubcon Las Vegas in 2012, Matt Cutts was heard to say that Google knew about 100 trillions pages on the web. The scale of the Internet/Web (se sometimes call this the Interwebs) is growing fast!

Underlying this enormous problem of processing all these pages is the complex nature of the web itself. Web pages include text, video, images and more. It's easy for humans to understand these and to transition seamlessly between them, but software lacks the intelligence we take for granted. This limitation and others affect how search engines understand the web pages they come across. We'll discuss some of these limitations in this chapter.

Of course, this is an ever changing landscape. The search engines continuously invest in improving their ability to better process the content of web pages. For example, advances in image and video search have enabled search engines to inch closer to human-like understanding, a topic that will be explored more in the section "Vertical Search Engines".

Understanding Search Engine Results

In the search marketing field, the pages the engines return to fulfill a query are referred to as *search engine results pages* (SERPs). Each engine returns results in a slightly different format and these may include *vertical results* (these are results which can be derived from different data sources or presented on the results page in a different format, which we'll illustrate shortly).

Understanding the Layout of Search Results Pages

Figure 2-1 shows the SERPs in Google for the query *stuffed animals*.

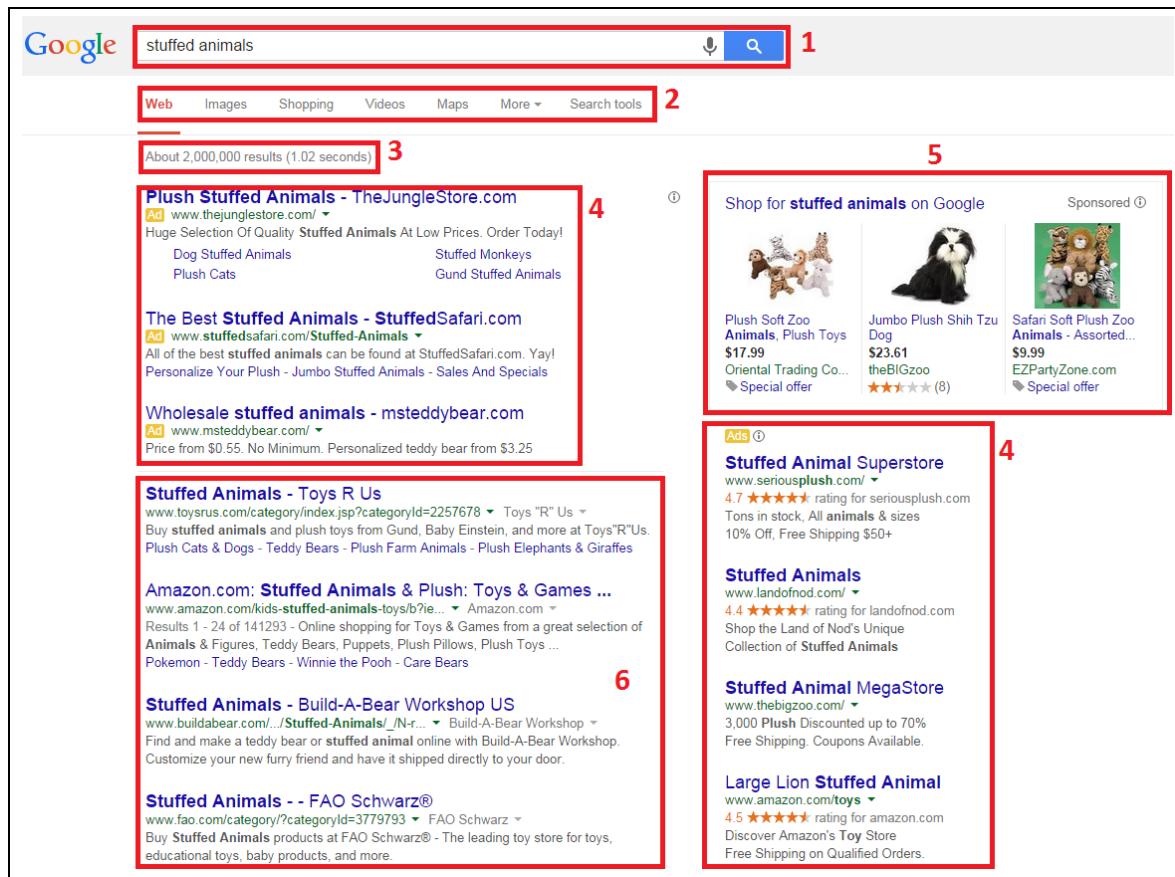


Figure 2-1. Layout of Google search results

The various sections outlined in the Google search results are as follows:

1. Search query box
2. Vertical navigation
3. Results information
4. PPC Advertising
5. Google Product Search Results
6. Natural/organic/algorithmic results

Even though Yahoo! no longer does its own crawl of the web, or provides its own search results information (it sources them from Bing), it does format the output differently. Figure 2-2 shows Yahoo!'s results for the same query.

1

Home Mail News Sports Finance Weather Games Groups Answers Screen Flickr Mobile More

YAHOO! stuffed animals

2

Ads related to stuffed animals

3

Web 3
Images
Video
News
Local
Shopping
Maps
More
Anytime
Past day
Past week
Past month

4

Stuffed Animals at JCP®
jcpenney.com/StuffedAnimals
Save on Stuffed Animals at JCP®. JCPenney Offers Free Ship Over \$99.
573 Donald Lynch Blvd, Marlborough, MA (508) 624-6070 Directions

Arts & Crafts Boys' Toys
Dolls and Dollhouses Educational Toys
Girls' Toys Remote Control Cars

Stuffed Animal Ball Pets™ - As Seen On Tv [Stuffed Animal Pets.](#)
www.BallPets.com/StuffedAnimalPets
Kids Favorite Plush Pets Order Now

Contact Us Ball Pets™ Reviews
Kids Love Ball Pets™ Special Offer \$12.99

Wholesale stuffed animals | msttedybear.com
www.msttedybear.com
Quality stuffed animal at low price Personalized teddy bear from \$3.25

Stuffed Animals - Highest Quality Stuffed Animals
www.TheJungleStore.com
Browse Our Huge Selection!

See more ads for: [stuffed animals](#), [stuffles](#) [stuffed animals](#)

5

StuffedAnimals.com™ - Stuffed Animals, Plush Teddy Bears ...
www.stuffedanimals.com Cached
You're a Valued Stuftomer at StuffedAnimals.com Whether you want to buy **stuffed animals**, teddy bears, toys, baby gifts, dolls, **plush** toys or bulk and wholesale ...

Amazon.com: Stuffed Animals & Plush Toys & Games: Animals
www.amazon.com/kids-stuffed-animals-toys/b?ie=UTF8&node=... Cached
Stuffed Animals from Amazon.com. Do you remember cuddling up with your favorite **stuffed animals**, listening to a goodnight story or the rain patter softly against the ...

Stuffed Animals - Image Results

6

Ads

Birthday Party Invitations **Plush Animal Assortment** **Personalized Plush Pink...** **Raja the Embroidered**
Tiny Prints \$20.40 Windy City \$11.80 Personalization \$39.95 Pier 1 Imports \$20.00

Bed Bug Stuffed Animal **Standard Business Cards** **Kids Plush Toys: Orange Cat...** **Bedtime Buddies**
Bed Bug Supply Vistaprint \$8.99 Land of Nod \$29.00 Bedtime Buddies \$19.99

[More Stuffed Animals on Yahoo Shopping](#)

stuffed animals **stuffles stuffed animals** **stuffed animals plush toys** **stretchkins stuffed animals** **stuffed animals amp plush toys** **gund stuffed animals**

4

Ads

Pure Alpaca Lux Sweaters
www.MtCaesarAlpacas.com
25% Off Biggest Holiday Sale. Handcrafted Artisan Alpaca Products

Stuffed Animals In Bulk

Figure 2-2. Layout of Yahoo! search results

The sections in the Yahoo! results are as follows:

1. Vertical navigation
2. Search query box
3. Horizontal navigation
4. PPC Advertising
5. Natural/organic/algorithmic results
6. Navigation to More Advertising

Figure 2-3 shows the layout of the results from Microsoft's Bing for *stuffed animals*.

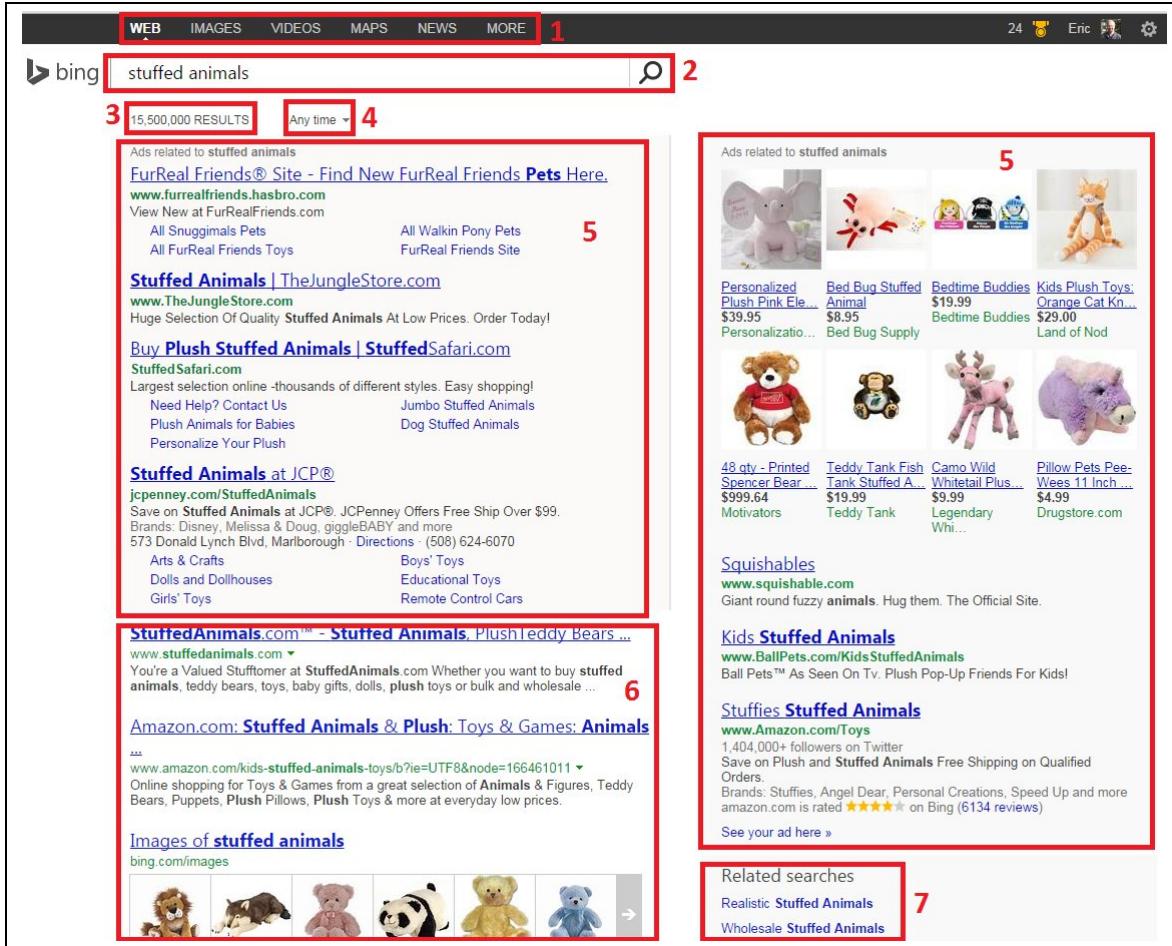


Figure 2-3. Layout of Bing search results

The sections in Bing's search results are as follows:

1. Vertical navigation
2. Search query box
3. Results information
4. Time based refinement options
5. PPC Advertising
6. Natural/organic/algorithmic results
7. Query Refinement Options

Each unique section represents a snippet of information provided by the engines. Here are the definitions of what each piece is meant to provide:

Vertical navigation

Each engine offers the option to search different verticals, such as images, news, video, or maps. Following these links will result in a query with a more limited index. In Figure 2-3, for example, you might be able to see news items about stuffed animals or videos featuring stuffed animals.

Horizontal navigation

All three engines used to have some form of horizontal navigation, but as of October 2014 only Yahoo! continues to do so.

Search query box

All of the engines show the query you've performed and allow you to edit or reenter a new query from the search results page. If you begin typing, you may notice that Google gives you a list of suggested searches below. This is the Google autocomplete suggestions feature, and it can be incredibly useful for targeting keywords. Next to the search query box, the engines also offer links to the advanced search page, the features of which we'll discuss later in the book. In addition, you will also see a microphone icon in the right of the search box that allows you to speak your query. In Google image search, this shows up as a camera icon that allows you to upload an image or get similar images back

Results information

This section provides a small amount of meta information about the results that you're viewing, including an estimate of the number of pages relevant to that particular query (these numbers can be, and frequently are, wildly inaccurate and should be used only as a rough comparative measure).

PPC (a.k.a. paid search) advertising

The text ads are purchased by companies that use either Google AdWords or Bing. The results are ordered by a variety of factors, including relevance (for which click-through rate, use of searched keywords in the ad, and relevance of the landing page are factors in Google) and bid amount (the ads require a maximum bid, which is then compared against other advertisers' bids).

Natural/organic/algorithmic results

These results are pulled from the search engines' primary indices of the Web and ranked in order of relevance and importance according to their complex algorithms. This area of the results is the primary focus of this section of the book.

Query refinement suggestions

Query refinements are offered by Google, Bing, and Yahoo!. The goal of these links is to let users search with a more specific and possibly more relevant query that will satisfy their intent.

In March 2009, Google enhanced the refinements by implementing Orion Technology, based on technology Google acquired in 2006. The goal of this enhancement is to provide a wider array of refinement choices. For example, a search on *principles of physics* may display refinements for the Big Bang, angular momentum, quantum physics, and special relativity.

Navigation to more advertising

Only Yahoo! shows this in the search results. Clicking on these links will bring you to additional paid search results related to the original query.

Be aware that the SERPs are always changing as the engines test new formats and layouts. Thus, the images in Figures Figure 2-1 through Figure 2-3 may be accurate for only a few weeks or months until Google, Yahoo!, and Bing shift to new formats.

How Vertical Results Fit into the SERPs

These "standard" results, however, are certainly not all that the engines have to offer. For many types of queries, search engines show *vertical* results, or *instant answers*, and include more than just links to other sites to help answer a user's questions. These types of results present many additional challenges and opportunities for the SEO practitioner.

Figure 2-4 shows an example of these types of results. The query in Figure 2-4 brings back a business listing showing an address and the option to get directions. This result attempts to provide the user with the answer he is seeking directly in the search results.

Google

Web Images Maps Shopping More Search tools

About 29,400 results (0.36 seconds)

Palo Alto - Il Fornaio Restaurants and Bakeries
www.ilfornaio.com/paloalto/ Il Fornaio
520 Cowper Street (at the Garden Court Hotel) Palo Alto, CA 94301. Phone: 650.853.3888. Email: paloalto@ilf.com - Steve Boyden, Managing Partner
At the Garden Court Hotel - Palo Alto - Dining Room

Il Fornaio Restaurants and Bakeries
www.ilfornaio.com/ Il Fornaio
The official website for Il Fornaio Authentic Italian Restaurant and Bakery, a full-service Italian restaurant offering on and off-site catering, private dining, Italian ...
3.6 ★★★★☆ 217 Google reviews · Write a review

 See photos

520 Cowper St, Palo Alto, CA 94301
(650) 853-3888
Menus - Restaurants - Online Reservations

Palo Alto - Il Fornaio Restaurants and Bakeries
www.ilfornaio.com/ Home > Restaurants > Palo Alto > Il Fornaio
Let us cater your next office meeting or special event. We offer a great variety for breakfast, lunch and afternoon snacks. Il Fornaio Palo Alto 520 Cowper Street

Il Fornaio - Palo Alto, CA | Yelp
www.yelp.com > Restaurants > Italian > Yelp, Inc.
★★★★☆ Rating: 3 - 368 reviews - Price range: \$\$
Not too many years ago Il Fornaio was one of the top restaurants in Palo Alto ... celebrating her birthday tonight and made a reservation at Il Fornaio Palo Alto.

Il Fornaio - Palo Alto Restaurant - Palo Alto, CA | OpenTable
www.opentable.com > ... > Palo Alto restaurants > OpenTable
★★★★☆ Rating: 3.9 - 190 reviews - Price range: \$30 and under
Book now at Il Fornaio - Palo Alto in Palo Alto, explore menu, see photos and read 190 reviews: "The night we dined the restaurant was exceptionally bu..."

 Map data ©2014 Google

IL Fornaio

Address: 520 Cowper St, Palo Alto, CA 94301
Phone: (650) 853-3888
Prices: \$\$\$
Hours: Open today - 7:00 am - 11:00 pm
Reservations: opentable.com
Menu: viewmenu.com

Reviews
3.6 ★★★★☆ 217 Google reviews
Upmarket chain for Italian dining, including bread from an in-house bakery. - Google
 "They always bring out bread and a platter of balsamic vinegar with olive oil."
 "Holy cow, they have awesome squash ravioli."

Figure 2-4. Local search result for a business

Figure 2-5 shows another example. The Google search in Figure 2-5 for weather plus a city name returns a direct answer. Once again, the user may not even need to click on a website if all she wanted to know was the temperature.

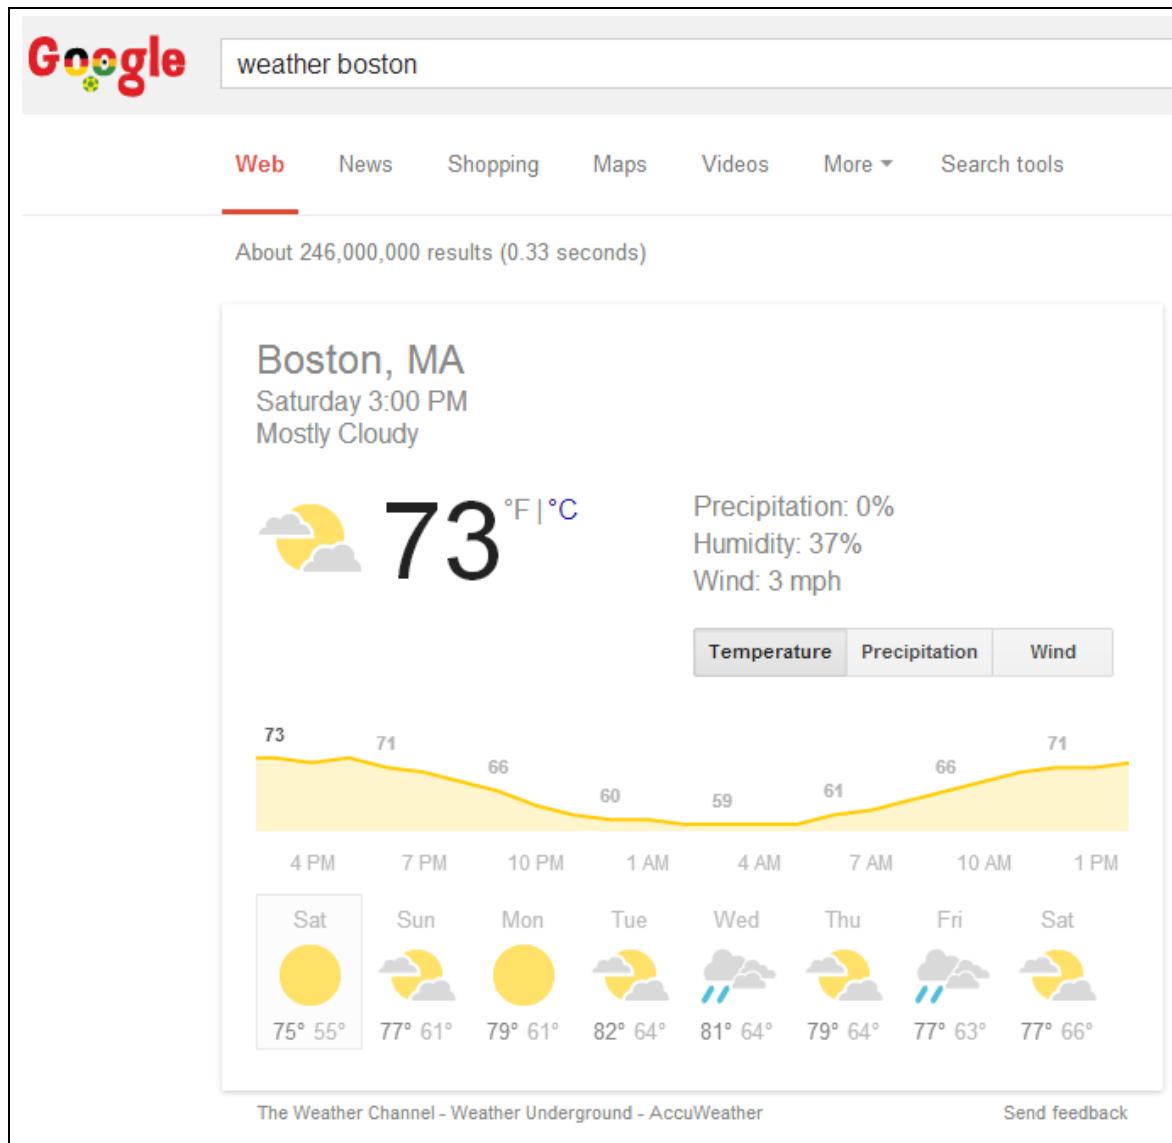


Figure 2-5. Weather search on Google

Figure 2-6 is an example of a search for a well-known painter. A Google search for the famous painter *Edward Hopper* returns image results of some of his most memorable works (shown in the bottom right of the screen shot). This example is a little different from the “instant answers” type of result shown in Figure 2-4 and Figure 2-5. If the user is interested in the first painting shown, he may well click on it to see the painting in a larger size or to get more information about it. For the SEO practitioner, getting placed in this vertical result could be a significant win.

Google Search

Web Images Books Videos Shopping More ▾ Search tools

About 2,000,000 results (0.27 seconds)

Edward Hopper - Wikipedia, the free encyclopedia
en.wikipedia.org/wiki/Edward_Hopper Wikipedia

Edward Hopper (July 22, 1882 – May 15, 1967) was a prominent American realist painter and printmaker. While he was most popularly known for his oil ...
Nighthawks - Early Sunday Morning - Automat (painting) - Office in a Small City

Edward Hopper - paintings, biography, quotes of Edward ...
www.edwardhopper.net/

Edward Hopper was born in 1882, in NY, into a middle class family, which encouraged the art work and career that he wanted to pursue. From 1900 to 1906 he ...
Nighthawks - Edward Hopper Paintings - Early Sunday Morning, 1930 - Automat

Edward Hopper (1882–1967) - Metropolitan Museum of Art
www.metmuseum.org/toah/.../hd_hopp.htm Metropolitan Museum of Art

Edward Hopper was born in Nyack, New York, a town located on the west side of the Hudson River, to a middle-class family that encouraged his artistic abilities.

News for edward hopper

 **Finalists announced in Art Everywhere US: Edward Hopper's ...**
 Dallas Morning News (blog) - by Michael Cranberry - 1 day ago
 Voting began in April, when people throughout the United States were given a list of 100 great American artworks and asked to whittle them ...

America, Online: Love for Hopper's *Nighthawks*
 New York Times (blog) - 1 day ago

Artworks chosen to be displayed across US
 Boston Herald - 6 hours ago

[More news for edward hopper](#)

 More images

Edward Hopper
 Painter

Edward Hopper was a prominent American realist painter and printmaker. While he was most popularly known for his oil paintings, he was equally proficient as a watercolorist and printmaker in etching. [Wikipedia](#)

Artwork View 10+ more

 **Nighthawks**
 1942

 **Automat**
 1927

 **Chop Suey**
 1929

 **Office at Night**
 1940

 **Office in a Small City**
 1953

Figure 2-6. Google search on an artist's name

Figure 2-7 shows an example from Yahoo!. A query on Yahoo! for *chicago restaurants* brings back a list of popular dining establishments from Yahoo!'s local portal. High placement in these results has likely been a good thing for Lou Malnati's Pizzeria.

YAHOO!

Web

[Images](#)

[Video](#)

[News](#)

[Local](#)

[Shopping](#)

[Maps](#)

[More](#)

[Anytime](#)

[Past day](#)

[Past week](#)

[Past month](#)

Ads related to chicago restaurants

[GrubHub - Food Delivery - See Local Restaurants that Deliver.](#)
 grubhub.com
View Menus & Order Online Here.

[LivingSocial Chicago - The Best Local Dining, Spas & More.](#)
www.LivingSocial.com
Up to 70% Off Deals. Join Free!

[Best Chicago Restaurants - See Our Top 10 Recommendations.](#)
ClippingBook.com/Chicago
Discover These Must See Hot Spots!

[Chicago Restaurants | Chicago TripAdvisor.com](#)
Chicago.TripAdvisor.com
 1,170 reviews for tripadvisor.com
Research Chicago Hotels, Chicago Deals, and Chicago Attractions!

See more ads for: [chicago restaurants](#), [north chicago restaurants](#)

Restaurants near Chicago

1 Lou Malnati's Pizzeria
 1568 reviews on Yelp
439 N Wells St, Chicago (312) 828-9800 Closed now

2 Giordano's
 239 reviews on Yelp
1040 W Belmont Ave, Chicago (773) 327-1200 Closed now

3 Hot Doug's
 3276 reviews on Yelp
3324 N California Ave, Chicago (773) 279-9550 Closed now

Hot Doug's
American Restaurant

3276 reviews on **yelp**

Is this your business? [Verify your listing](#)

3324 N California Ave, Chicago, IL 60618
Cross Streets: Between W Henderson St and W Roscoe St
Neighborhoods: Avondale
(773) 279-9550
[hotdougs.com](#) · more info

Today: 10:30 am ~ 4:00 pm Closed now
[See all hours](#)

[Directions](#) [Menu](#)

Reviews

Sandi S. 09/03/14
I've been warned about the Hot Doug's wait, so I wasn't surprised to see the line. However I never expected to be waiting almost 5 hours! I thought maybe 3, at most 4...I've waited 3 hours for ...[more](#)

Philip B. 09/04/14
A long wait, but amazing hot dogs. After reading reviews we decided to get there early to avoid waiting too long. We arrived at around 9:30 (an

Figure 2-7. Yahoo! search for Chicago restaurants

Figure 2-8 is an example of a celebrity search on Bing.

bing charlie chaplin

4,770,000 RESULTS Any time ▾

Charlie Chaplin - Wikipedia, the free encyclopedia
en.wikipedia.org/wiki/Charlie_Chaplin ▾

 Sir Charles Spencer "Charlie" Chaplin, KBE (16 April 1889 – 25 December 1977) was an English actor, comedian, and filmmaker, who rose to fame in ...
 Biography · Filmmaking · Legacy · Awards and recognition · Filmography

Charlie Chaplin - Official Site
www.charliechaplin.com ▾

Charlie's Mother: Hannah Chaplin. Hannah was a singer and character comedienne in the British music halls with the stage name of Lily Harley, and she did enjoy some ...

Images of charlie chaplin
bing.com/images



Charlie Chaplin Biography - Facts, Birthday, Life Story ...
www.biography.com/people/charlie-chaplin-9244327 ▾

 Iconic actor Charlie Chaplin was known for his lovable tramp figure from vintage film footage. Learn more at Biography.com.

Charles Chaplin - IMDb
www.imdb.com/name/nm0000122 ▾

 Charles Chaplin, Writer: *Modern Times*. Charlie Chaplin, considered to be one of the most pivotal stars of the early days of Hollywood, lived an ...

Videos of charlie chaplin
bing.com/videos



Charlie Chaplin - The Lion's Cage YouTube 3:25 Charlie Chaplin receiving an Hon... YouTube 5:09 Charlie Chaplin Google Doodle YouTube 2:07 Charlie Chaplin final speech in T... YouTube 4:54

Charlie Chaplin

 Sir Charles Spencer "Charlie" Chaplin, KBE was an English actor, comedian, and filmmaker, who rose to fame in the silent era. Chaplin became a worldwide icon through his screen persona 'the Tramp' and is considered one of the most im... +

en.wikipedia.org
fotosytumbas.blogs...

Lived: Apr 16, 1889 - Dec 25, 1977 (age 88)
 Height: 5' 5" (1.65 m)
 Spouse: [Dona O'Neill](#) (1943 - 1977) · [Paulette Goddard](#) (1936 - 1942) · [Lita Grey](#) (1924 - 1927) · [Mildred Harris](#) (1918 - 1920)
 Children: [Geraldine Chaplin](#) · [Christopher Chaplin](#) · [Charles Chaplin, Jr.](#) · [Michael Chaplin](#) · [Sydney Chaplin](#) · [Eugene Chaplin](#) · [Victoria Chaplin](#) +
 Parents: [Hannah Chaplin](#) · [Charles Chaplin, Sr.](#)
 Siblings: [Sydney Chaplin](#) · [Wheeler Dryden](#)

Movies



The Great Dictator Chaplin 1992 Modern Times The Kid City Lights

People also search for



Dona O'Neill Spouse Jackie Coogan Buster Keaton Douglas Fairbanks Edna Purviance

Data from: [wikipedia](#) · [imdb](#) · [firebase](#)

Figure 2-8. Bing result for Charlie Chaplin

The results in Figure 2-9 include a series of images of the famous actor Charlie Chaplin. As a last example, Figure 2-9 is a look at the Bing search results for videos with Megan Fox.

b bing megan fox videos 

9,820,000 RESULTS Any time ▾

Videos of megan fox
bing.com/videos

 2:42  1:19  3:10  1:26

Behind the scenes of *This is 40*... MTV Megan Fox - Esquire Magazine... worldstarhiphop.... Megan Fox in "the Tip" (Full Video) YouTube

 5:46  1:27  3:57  3:58

Megan Fox metatube FHM's Exclusive Megan Fox Video! YouTube Megan Fox YouTube Megan Fox Boob Trailer Jaw-dropper Fox News

Megan Fox: 5 Hottest Videos of the Sexy Star | HEAVY
heavy.com/action/girls/2013/06/megan-fox-hot-videos-5-sexiest... ▾

This is the shoot that started it all – back when Megan Fox's fame as a Transformers gal propelled her into the universe of celebrity sex symbols.

News about Megan Fox Videos
bing.com/news

 *Teenage Mutant Ninja Turtles'*: Watch Our On Set Visit With **Megan Fox** MTV - 12 days ago

we visited the set of the upcoming movie and spoke to stars Megan Fox and Will Arnett. And Fox, who plays news reporter and friend of the turtles April O'...

 **Check out these views from Megan Fox's L.A. home** HLN - 1 day ago

Exclusive: Brian Austin Green And Megan Fox List Eclectic Los Feliz Mansion For Nearly \$4 M Forbes - 1 day ago

FHM's Exclusive Megan Fox Video! - YouTube
www.youtube.com/watch?v=2UJOA4gJ2U ▾

By FHM - 1 min - 9,193,209 views - Added Jun 27, 2008 <http://www.fhm.com/girls/covergirls/m...> Megan Fox's exclusive lingerie shoot for FHM Magazine! She looks pretty good... we won't argue.

Megan Fox

 Megan Fox is an American actress and model. She began her acting career in 2001, with several minor television and film roles, and played a regular role on the *Hope & Faith* television show. In 2004, she made her film debut with a role in... + en.wikipedia.org

www.imdb.com

 [Facebook](#)

Born: May 16, 1986 (age 28) - Oak Ridge, Tennessee
Net worth: \$5 million USD (2014)
Spouse: Brian Austin Green (2010)
Romance: Cam Gigandet - Angus Mullen
Children: Noah Shannon Green - Bodhi Ransom Green
Upcoming movies: *Teenage Mutant Ninja Turtles*

Movies and TV shows

 Jennifer's Body 2009  This Is 40 2012  Transformers 2007  Passion 2010  Transformers: Revenge of the Fallen 2009

People also search for

 Brian Austin Green Spouse  Shia LaBeouf  Jennifer Lawrence  Scarlett Johansson  Rosie Huntington-Whiteley

Data from: wikipedia | imdb | freebase

Figure 2-9. Bing result for Megan Fox videos

At the top of the search results in Figure 2-9, you get a series of popular videos provided. Click on a video in the results and the video begins playing right there in the search results.

As you can see, the vast variety of vertical integration into search results means that for many popular queries you can expect to receive significant amounts of information in the SERPs themselves. Engines are competing by providing more relevant results and more targeted responses to queries that they feel are best answered by vertical results, rather than web results.

As a direct consequence, site owners and web marketers must take into account how this incorporation of vertical search results may impact their rankings and traffic. For many of the searches shown in the previous figures, a high ranking—even in position #1 or #2 in the algorithmic/organic results—may not produce much traffic because of the presentation of the vertical results above them.

The vertical results also signify an opportunity, as listings are available in services from images to local search to news and products. We will cover how to get included in these results in Chapter 11.

Google's Knowledge Graph

The search engines are actively building structured databases of information that allow them to show answers to questions that are not simply links to web pages. In Figure 2-6 the information on the upper right is an example of this. Google provides direct answers in the result, including Edward Hopper's birth date, place of birth, the date and place of his death, his spouse, and more. In Figure 2-9, Bing provides similar information for Megan Fox.

Not only is additional information shown, it is not just a data dump, it shows that the search engines are working to develop their own knowledge of the relationships between people and things. In the case of Figure 2-6, we can see that Google understands that:

- Edward Hopper is the name of a person.
- People have dates and place of birth
- People have dates and places of death
- People might have spouses

The search engines are actively mapping these types of relationships as part of their effort to offer more complete information directly in the search results themselves.

Algorithm-Based Ranking Systems: Crawling, Indexing, and Ranking

Understanding how crawling, indexing, and ranking works is helpful to SEO practitioners, as it helps them determine what actions to take to meet their goals. This section primarily covers the way Google and Bing operate, and does not necessarily apply to other search engines that are popular in other countries, such as Yandex (Russia), Baidu (China), Seznam (Czechoslovakia), and Naver (Korea).

The search engines must execute many things very well to provide relevant search results. Put simplistically, you can think of these as:

- Crawling and indexing trillions of documents (pages and files) on the Web (note that they ignore pages that they consider to be "insignificant", perhaps because they are perceived as adding no new value or are not referenced at all on the web).
- Responding to user queries by providing lists of relevant pages

In this section, we'll walk through the basics of these functions from a nontechnical perspective. This section will start by discussing how search engines find and discover content.

Crawling and Indexing

To offer the best possible results, search engines must attempt to discover all the public pages on the World Wide Web and then present the ones that best match up with the user's search query. The first step in this process is *crawling* the web. The search engines start with a seed set of sites that are known to be very high quality sites, and then visit the links on each page of those sites to discover other web pages.

The link structure of the Web serves to bind together all of the pages that were made public as a result of someone linking to them. Through links, search engines' automated robots, called *crawlers*, *spiders* can reach the many billions of interconnected documents.

In Figure 2-10 you can see the home page of <http://www.usa.gov>, the official US government web site. The links on the page are outlined in red. Crawling this page would start with loading the page analyzing the content, and then seeing what other pages that <http://www.usa.gov> links to.

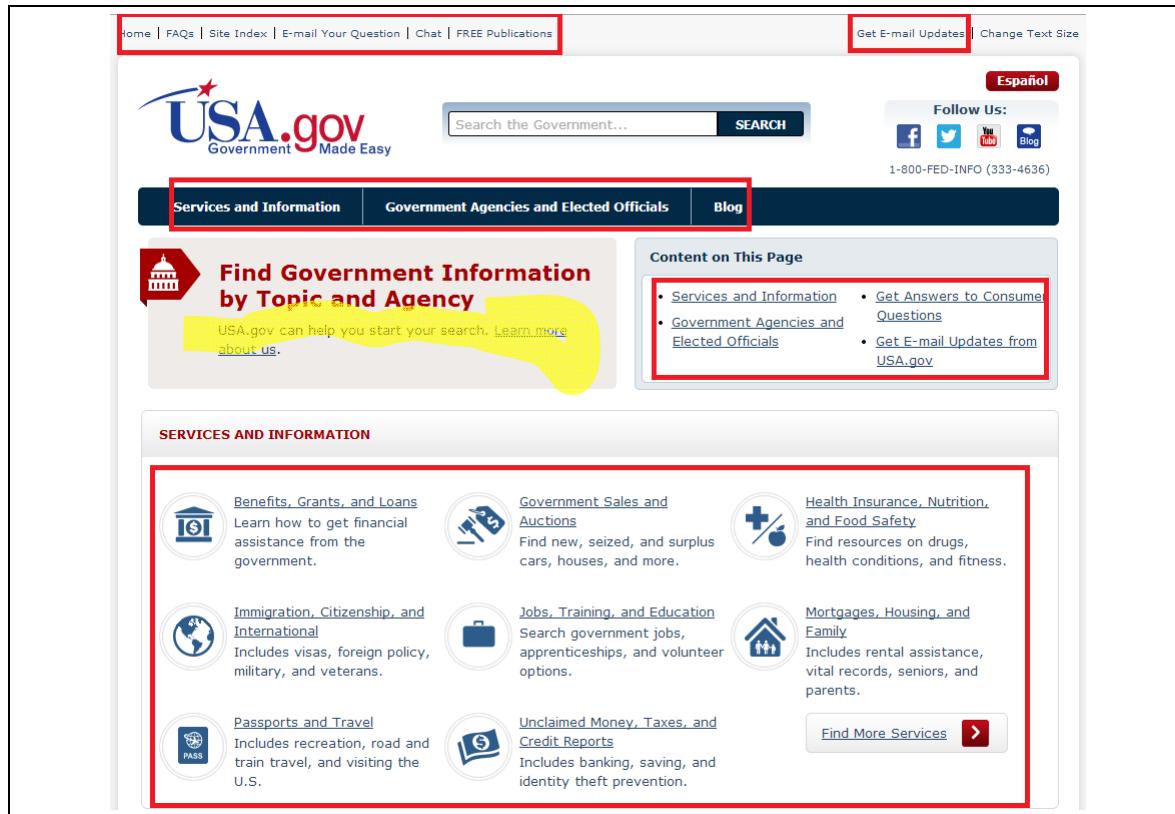


Figure 2-8. Crawling the US Government Website

The search engine will then load those other pages and analyze that content as well. This process repeats over and over again until the crawling process is complete. This process is an enormously complex one as the web is a large and complex place.

Search engines do not attempt to crawl the entire web every day. In fact they may become aware of pages that they choose not to crawl because they are not likely to be important enough to return in a search result. We will discuss the role of *importance* in the next section, "Retrieval and Rankings".

The first step in this process is to build an index of terms. This is a massive database that catalogs all the significant terms on each page crawled by the search engine.

A lot of other data is also recorded, such as a map of all the pages that each page links to, the clickable text of those links (we call this the anchor text), whether or not those links are considered ads, and more.

To accomplish the monumental task of holding data on hundreds of billions (or trillions) of pages that can be accessed in a fraction of a second, the search engines have constructed massive data centers to deal with all this data.

One key concept in building a search engine is deciding where to begin a crawl of the Web. Although you could theoretically start from many different places on the Web, you would ideally begin your crawl with a trusted seed set of websites.

Starting with a known trusted set of web sites enables search engines to measure how much they trust the other web sites that they find through the crawling process. We will discuss the role of trust in search algorithms in more detail in "How Links Influence Search Engine Rankings" in Chapter 7.

Retrieval and Ranking

For most searchers, the quest for an answer begins as shown in Figure 2-11.

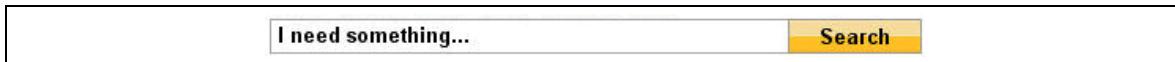


Figure 2-9. Start of a user's search quest

The next step in this quest occurs when the search engine returns a list of relevant pages on the Web in the order most it believes is most likely to satisfy the user. This process requires the search engines to scour their corpus of hundreds of billions of documents and do two things: first, return only the results that are related to the searcher's query; and second, rank the results in order of perceived importance (taking into account the trust and authority associated with the site). It is both relevance and importance that the process of SEO is meant to influence.

Relevance is the degree to which the content of the documents returned in a search matches the user's query intention and terms. The relevance of a document increases if the page contains terms relevant to the phrase queried by the user, or if links to the page come from relevant pages and use relevant anchor text.

You can think of relevance as the first step to being "in the game." If you are not relevant to a query, the search engine does not consider you for inclusion in the search results for that query. We will discuss how relevance is determined in more detail in "Determining Searcher Intent and Delivering Relevant, Fresh Content".

Importance refers to the relative importance, measured via *citation* (the act of one work referencing another, as often occurs in academic and business documents) of a given document that matches the user's query. The importance of a given document increases with every other document that references it. In today's online environment citations can come in the form of links to the document or references to it on social media sites. Determining how to weight these signals is known as *citation analysis*.

You can think of importance as a way to determine which page, from a group of equally relevant pages, shows up first in the search results, which is second, and so forth. The relative authority of the site, and the trust the search engine has in it, are significant parts of this determination. Of course, the equation is a bit more complex than this and not all pages are equally relevant. Ultimately, it is the combination of relevance and importance that determines the ranking order.

So, when you see a search results page such as the one shown in Figure 2-12, you can surmise that the search engine (in this case, Bing) believes the Superhero Stamps page on ebay.com (shop.ebay.com/i.html?_nkw=superhero+stamps) has the highest combined score for relevance and importance for the query *marvel superhero stamps*.

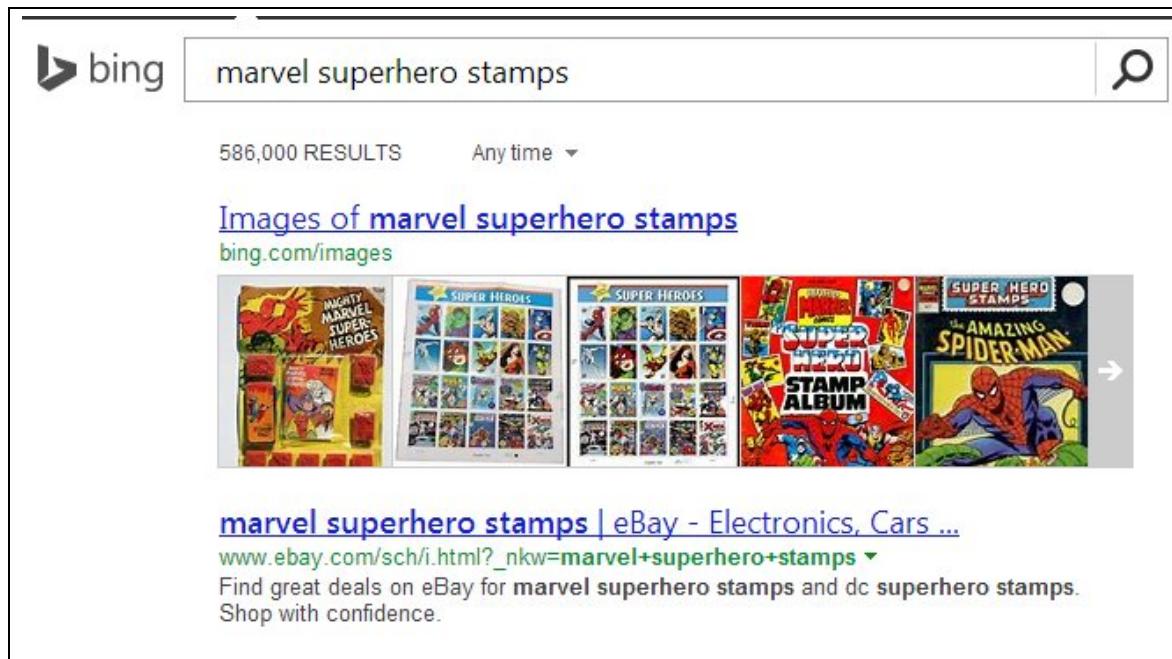


Figure 2-10. Sample search result for “marvel superhero stamps”

Importance and relevance aren’t determined manually (those trillions of man-hours would require Earth’s entire population as a workforce). Instead, the engines craft careful, mathematical equations—algorithms—to sort the wheat from the chaff and to then rank the wheat in order of quality. These algorithms often comprise hundreds of components. In the search marketing field, they are often referred to as *ranking factors* or *algorithmic ranking criteria*.

We discuss ranking factors or signals (*signals* is the term Google prefers) in more detail in “Analyzing Ranking Factors”.

Evaluating Content on a Web Page

Search engines place a lot of weight on the content of each web page. After all, it is this content that defines what a page is about, and the search engines do a detailed analysis of each web page they find during their crawl to help make that determination.

You can think of this as the search engine performing a detailed analysis of all the words and phrases that appear on a web page, and then building a map of that data for it to consider showing your page in the results when a user enters a related search query. This map, often referred to as a *semantic map*, seeks to define the relationships between those concepts so that the search engine can better understand how to match the right web pages with user search queries.

If there is no semantic match of the content of a web page to the query, the page has a much lower possibility of showing up. Therefore, the words you put on the page, and the “theme” of that page, play a huge role in ranking.

Figure 2-13 shows how a search engine will break up a page when it looks at it, using a page on the Forbes website.

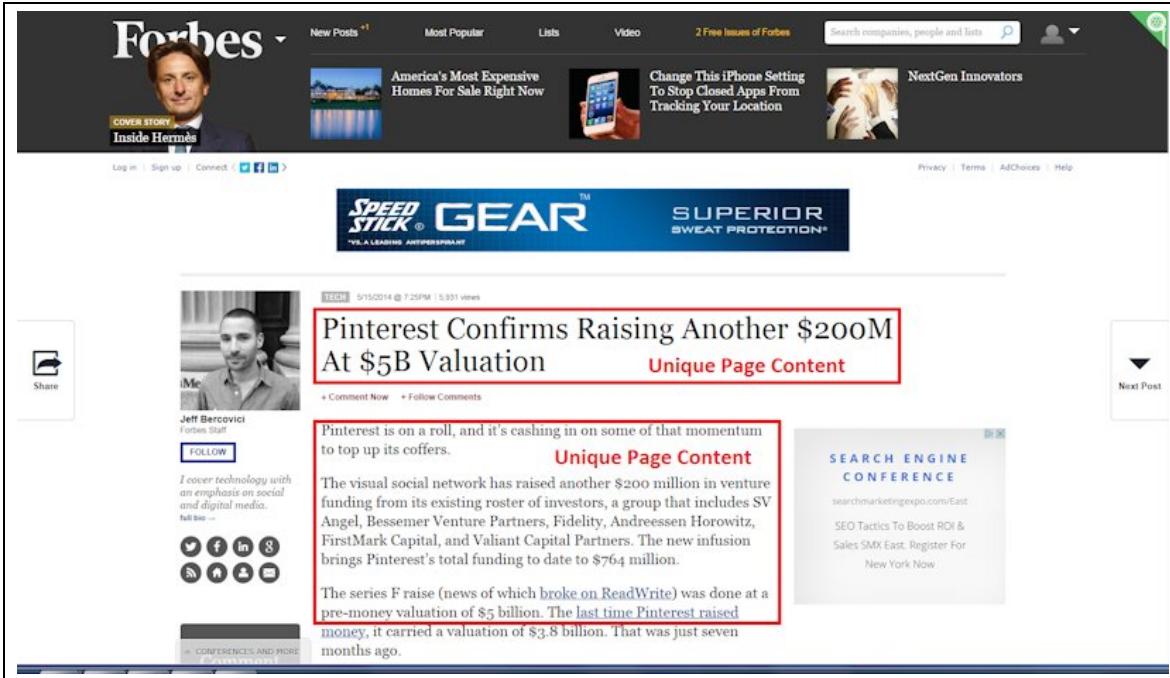


Figure 2-11. Breaking up a web page

The navigational elements of a web page are likely similar across the many pages of a site. These navigational elements are not ignored, and they do play an important role, but they do not help a search engine determine what the unique content is on a page. To do that, the search engine focuses on the part of Figure 2-13 that is labeled “Unique Page Content”.

Determining the unique content on a page is an important part of what the search engine does. It is this understanding of the unique content on a page that the search engine uses to determine the types of search queries for which the web page might be relevant. Since site navigation is generally not unique to a single web page, it does not help the search engine with that task.

This does not mean navigation links are not important, because they most certainly are—however, they simply do not count when trying to determine the unique content of a web page because those navigation links are shared among many web pages.

One task the search engines face is judging the value of content. Although evaluating how the community responds to a piece of content using link analysis is part of the process, the search engines can also draw some conclusions based on what they see on the page.

For example, is the exact same content available on another website? Is the unique content the search engine can see two sentences long or 500+ words long? Does the content repeat the same keywords excessively? These are a few examples of things that the search engine can look at when trying to determine the value of a piece of content.

What Content Can Search Engines “See” on a Web Page?

Search engine crawlers and indexing programs are basically software programs. These programs are extraordinarily powerful. They crawl trillions of web pages, analyze the content of all these pages, and analyze the way all these pages link to each other. Then they organize this into a series of databases that can respond to a user search query with a highly tuned set of results in a few tenths of a second.

This is an amazing accomplishment, but it has its limitations. Software is very mechanical, and it can understand only portions of most web pages. The search engine crawler analyzes the raw HTML form of a web page. If you want to see what this looks like, you can do so by using your browser to view the source.

Figure 2-14 shows how to do that in Chrome, and Figure 2-15 shows how to do that in Firefox (bottom). Typically you can access it most easily by right-clicking with your mouse on a web page to access a hidden menu.

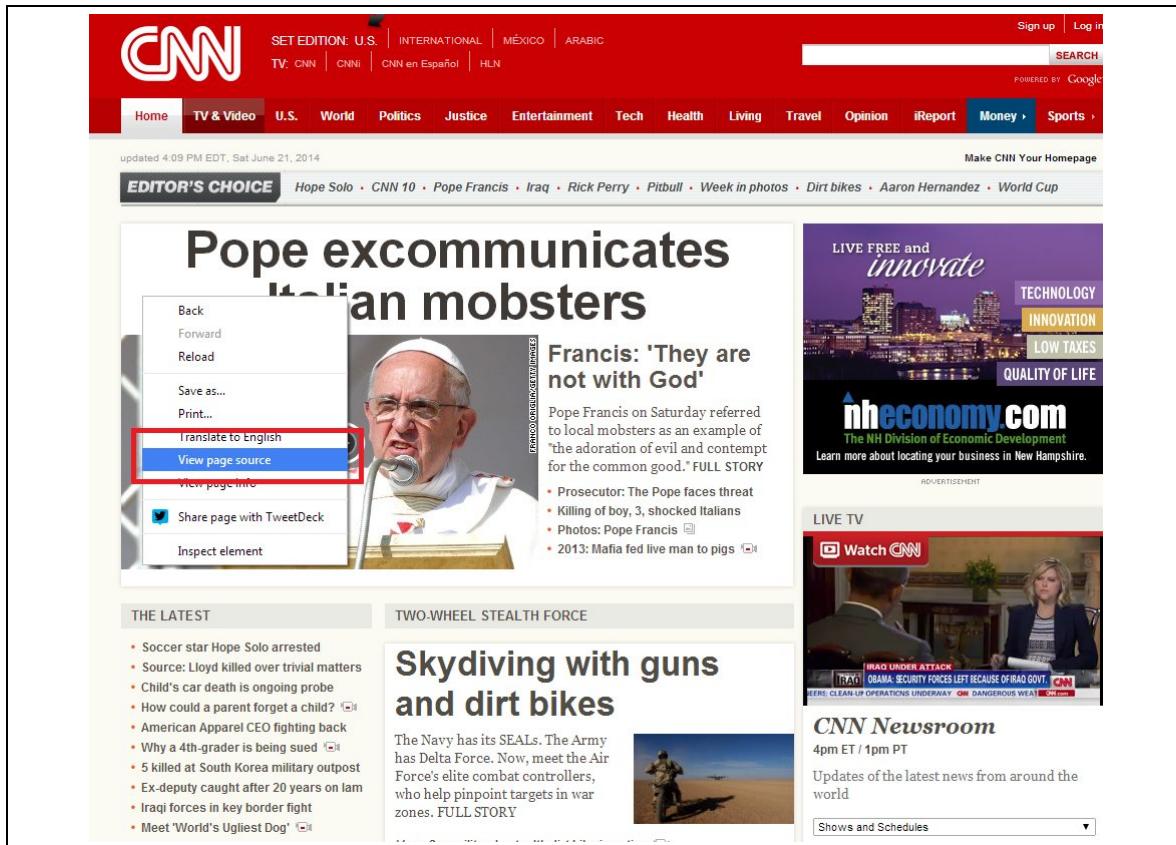


Figure 2-12. Viewing source in Chrome – right click on web page to access menu

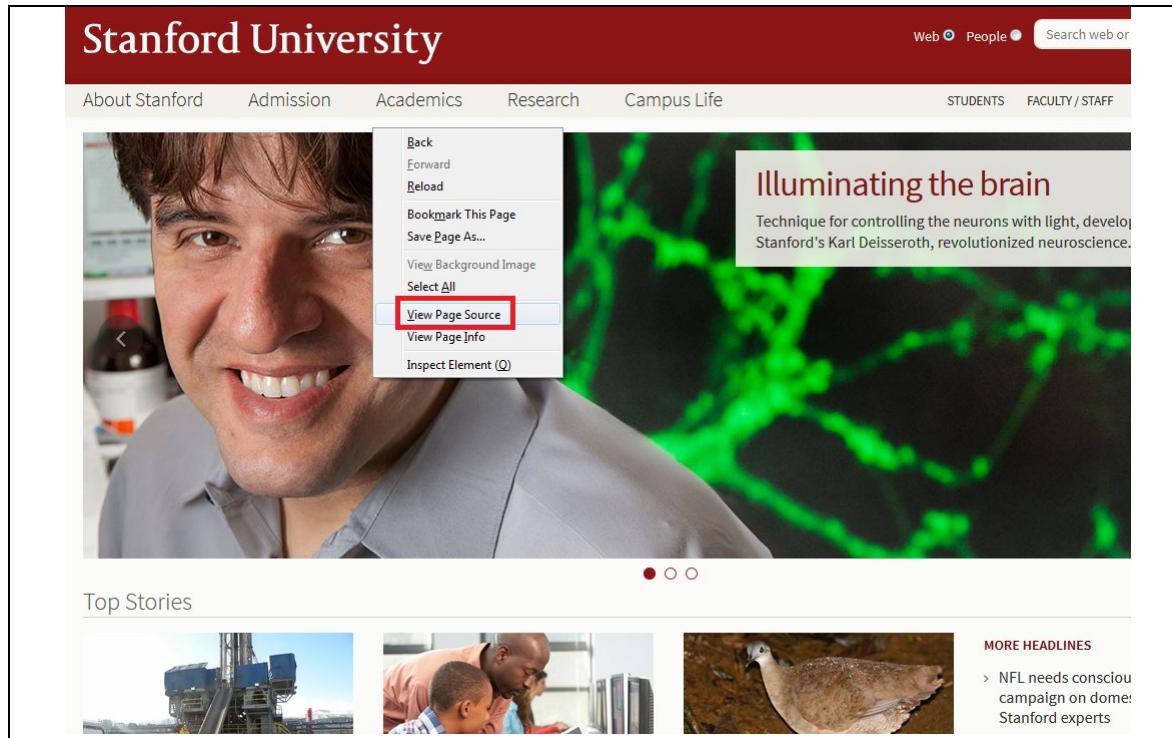


Figure 2-135. Viewing source in Firefox

Once you view the source, you will be presented with the exact code for the web page that the web server sent to your browser. This is most of what the search engine crawler sees (the search engine also sees the HTTP headers for the page, which are status codes it receives from the web server where the page is hosted). When trying to analyze the user visible content on a web page, search engines largely ignore code related to the navigation and display of the page such as that shown in Figure 2-16, as it has nothing to do with the content of the web page.

```

<body class="home home-index layout-default">
  <div id="root" class="no-overflow">
    <nav class="globalnav-container default">
      <ul class="globalnav">
        <li class="active">
          <a href="http://moz.com" onclick="_gaq.push(['_trackEvent', 'multiproduct_nav', 'button', 'moz_com']);">Moz.com</a>
        </li>
        <li class="">
          <a href="http://moz.com/pro" onclick="_gaq.push(['_trackEvent', 'multiproduct_nav', 'button', 'moz_pro']);">Moz Pro</a>
        </li>
        <li class="">
          <a href="https://moz.com/local" onclick="_gaq.push(['_trackEvent', 'multiproduct_nav', 'button', 'moz_local']);">Moz Local</a>
        </li>
      </ul>
    </nav><header class="background-blue-canvas masthead-container">
      <div class="container masthead">
        <div class="row">
          <div id="logo" class="span3 center">
            <a href="/">
              <img alt="Moz logo" width="120" height="35" data-bbox="188 578 308 698" />
            </a>
          </div>
        </div>
      </div>
    </header>
    <div class="content">
      <div>
        <h1>Illuminating the brain</h1>
        <img alt="Micrograph of neurons" data-bbox="478 154 828 354" />
        <div>
          <h2>Technique for controlling the neurons with light, developed by Stanford's Karl Deisseroth, revolutionized neuroscience.</h2>
        </div>
      </div>
    </div>
  </div>
</body>

```

Figure 2-14. Sample web page source code

The information the search engine crawler is most interested in is in the HTML text on the page. Figure 2-17 is an example of HTML text for a web page, using the moz.com (<http://www.moz.com>) home page.

```
        <h2 class="h3 top0 bottom1"><a href="http://moz.com/blog/tips-and-tactics-for-amplifying-your-content-whiteboard-friday" class="slate">Tips and Tactics for Amplifying Your Content - Whiteboard Friday</a></h2>
        <p class="bottom1 small">
            <time datetime="2014-06-20 00:12:00">June 20th, 2014</time>
        </p>
        - Posted by <a href="http://moz.com/community/users/17229">Ben Lloyd</a> to
            <a href="http://moz.com/blog/category/content">Content</a> and
            <a href="http://moz.com/blog/category/whiteboard-friday">Whiteboard Friday</a>
        </p>
    </header>
    <p>Content marketing should never be approached with a "set it and forget it" mentality. It needs to be structured and shared in the right ways, and in today's Whiteboard Friday, the folks from Add3 are here to show you what that means.</p>
</div>

<div class="pull-right">
<a href="http://moz.com/blog/tips-and-tactics-for-amplifying-your-content-whiteboard-friday">Read Full Entry</a>
|
<span class="comment-count">
<a href="http://moz.com/blog/tips-and-tactics-for-amplifying-your-content-whiteboard-friday#comments">
    <i class="icon icon-speech"></i>
    32 comments
</a>

```

Figure 2-15. Sample HTML text in the source code showing real content

Although Figure 2-17 still shows some HTML encoding, you can see the “regular” text clearly in the code. This is the unique content that the crawler is looking to find.

In addition, search engines read a few other elements. One of these is the page title. The page title is one of the most important factors in ranking a given web page. It is the text that shows in the browser's title bar (above the browser menu and the address bar).

Figure 2-18 shows the code that the crawler sees, using Trip Advisor (<http://www.tripadvisor.com>) as an example.

The first red circle in Figure 2-18 is for the title tag. The title tag is also often (but not always) used as the title of your listing in search engine results.

```
-----  
<meta http-equiv="imagetoolbar" content="no"/>  
<title>Reviews of Hotels, Flights and Vacation Rentals - TripAdvisor</title>  
<meta http-equiv="pragma" content="no-cache"/>  
<meta http-equiv="cache-control" content="no-cache,must-revalidate"/>  
<meta http-equiv="expires" content="0"/>  
<meta property="og:image" content="http://c1.tacdn.com/img2/postimg.jpg" height="150px"  
width="150px"/>  
<meta name="keywords" content="vacation, vacations, vacation packages, vacation package,  
travel package, travel packages, travel, planning, hotel, hotels, motel, bed and  
breakfast, inn, guidebook, review, reviews, popular, plan, airfare, cheap, discount, map,  
maps, golf, ski, articles, attractions, advice, restaurants"/>  
<meta name="description" content="TripAdvisor - Unbiased hotel reviews, photos and travel  
advice for hotels and vacations - Compare prices with just one click."/>  
<link rel="alternate" hreflang="en" href="http://www.tripadvisor.com/"/>
```

Figure 2-16. Meta tags in HTML source

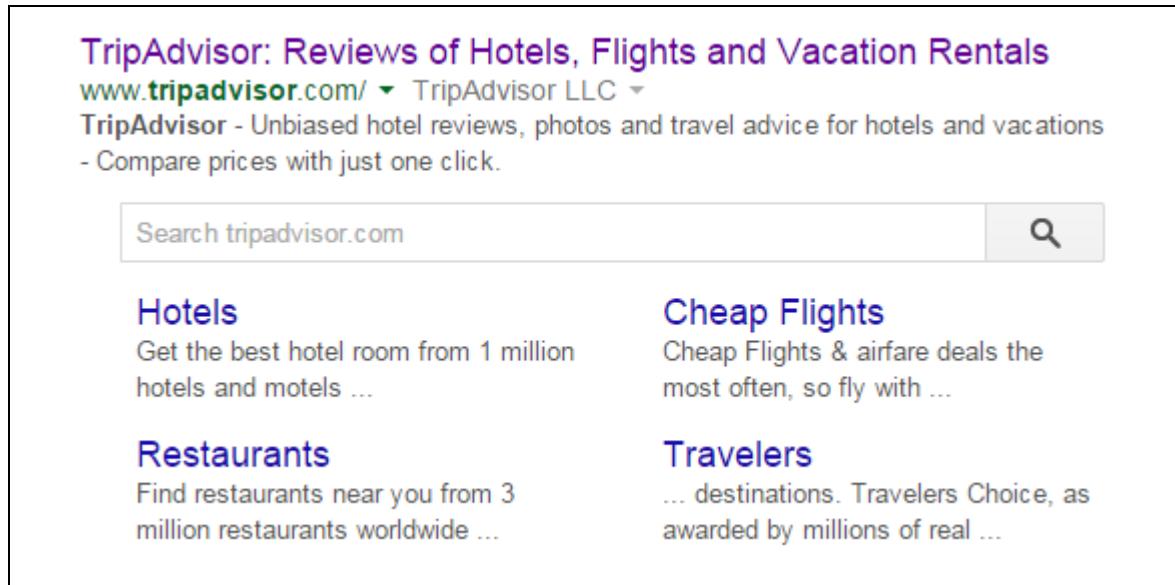


Figure 2-17. Search result showing title tag

In addition to page titles, the search engines also used to use the “meta keywords” tag. This is a list of keywords that you wish to have associated with the page. *Spammers* (people who attempt to manipulate search engine results in violation of the search engine guidelines) ruined the SEO value of this tag many years ago, so its value is now negligible, as search engines don’t use it anymore. Spending time on meta keywords is not recommended because of the lack of SEO benefit.

The second red circle in Figure 2-18 shows an example of a meta keywords tag.

Search engines also read the meta description tag (the third red circle in the HTML in Figure 2-18). However, the content of a meta description tag is not directly used by search engine in their ranking algorithms (<http://searchengineland.com/21-essential-seo-tips-techniques-11580>).

Nonetheless, the meta description tag plays a key role as search engines often use it as a part or all of the description for your page in search results. Therefore, a well-written meta description can have a significant influence on how many clicks you get on your search listing, and the click through rate on your search listing can impact your ranking. Time spent on meta descriptions is quite valuable as a result. Figure 2-20 uses a search on *trip advisor* to show an example of the meta description being used as a description in the search results.

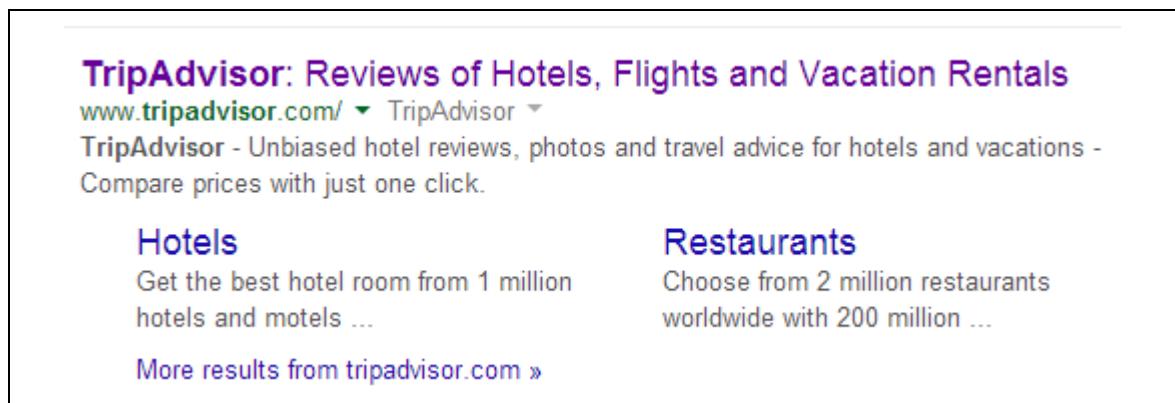


Figure 2-18. Meta description used in search results

The user’s keywords are typically shown in boldface when they appear in the search results (sometimes close synonyms are shown in boldface as well). As an example of this,

in Figure 2-20, *TripAdvisor* is in boldface at the beginning of the description. This is called KWIC (keywords in context).

A fourth element that search engines read is the `alt` attribute for images. The `alt` attribute was originally intended to allow something to be rendered when viewing of the image is not possible. There were two basic audiences for this:

- Vision-impaired people who do not have the option of viewing the images.
- People who turn off images for faster surfing. This is generally an issue only for those who do not have a broadband connection.

Support for the vision-impaired remains a major reason for using the `alt` attribute. You can read more about this by visiting the Web Accessibility Initiative page on the W3C website (<http://www.w3.org/WAI>).

Search engines also read the text contained in the `alt` attribute of an image tag. An image tag is an element that is used to tell a web page to display an image. Here is an example of an image tag from the Alchemist Media site:

```

|
```

The `src=` part of the tag is where the image to be displayed is located. The part that starts with `alt=` and is then followed with "Top Search Agencies" is considered the `alt` attribute.

The `alt` attribute is something that the search engines read. The search engines interpret it to help them determine what the image is about and to get more of a sense as to what the page is about.

Another element that search engines read is the `noscript` tag. Historically, the ability of search engines to read JavaScript was quite limited, but this has been changing over time and Google says they execute more Javascript today (<http://googlewebmastercentral.blogspot.com/2014/05/understanding-web-pages-better.html>). However, a small percentage of users do not allow JavaScript to run when they load a web page (the authors' experience is that it is about 2%). For those users, nothing would be shown to them where the JavaScript is on the web page, unless the page contains a `noscript` tag.

Here is a very simple JavaScript example that demonstrates this:

```
<script type="text/javascript">
  document.write("It is a Small World After All!")
</script>
<noscript>Your browser does not support JavaScript!</noscript>
```

The `noscript` portion of this is Your browser does not support JavaScript!. In this example, you could also choose to make the `noscript` tag contain the text "It is a Small World After All!". The `noscript` tag should be used only to represent the content of the JavaScript.

What search engines cannot see

It is also worthwhile to review the types of content that search engines cannot "see" in the human sense.

For instance, although search engines are able to detect that you are displaying an image, they have little idea what the image is a picture of, except for whatever information you provide them in the `alt` attribute, as discussed earlier. They can only recognize some very basic types of information within images, such as the presence of a face, or whether images have pornographic content by how much flesh tone there is in the image. A search engine cannot easily tell whether an image is a picture of Bart Simpson, a boat, a house, or a tornado. In addition, search engines typically don't recognize any text rendered in the image.

The reality is that the search engines have the technology to handle these types of tasks to some degree. For example, you can take a picture of the Taj Majal and drag it into the search box in Google image search and it will recognize it. However, because of the processing power to do such image recognition, the search

engines do not currently try to do that for all of the images they encounter on the hundreds of trillions of pages across the web.

The search engines are also experimenting with technologies to use optical character recognition (OCR) to extract text from images, but this technology is not yet in general use within search. The main problem with applying OCR and image processing technology is they are very computationally intensive, and it is not practical to apply them at the scale of the web.

In addition, conventional SEO wisdom has always held that the search engines cannot read Flash files, but this is a little overstated. Search engines have been extracting some information from Flash for years, as indicated by this Google announcement in 2008: <http://googlewebmastercentral.blogspot.com/2008/06/improved-flash-indexing.html>. However, the bottom line is that it's not easy for search engines to determine what is in Flash. One of the big issues is that even when search engines look inside Flash, they are still looking for textual content, but Flash is a pictorial medium and there is little incentive (other than the search engines) for a designer to implement text inside Flash. All the semantic clues that would be present in HTML text (such as heading tags, boldface text, etc.) are missing too, even when HTML is used in conjunction with Flash.

A third type of content that search engines cannot see is the pictorial aspects of anything contained in Flash, so this aspect of Flash behaves in the same way images do. For example, when text is converted into a vector-based outline (i.e., rendered graphically), the textual information that search engines can read is lost. Chapter 6 discusses methods for optimizing Flash.

Audio and video files are also not easy for search engines to read. As with images, the data is not easy to parse. There are a few exceptions where the search engines can extract some limited data, such as ID3 tags within MP3 files, or enhanced podcasts in AAC format with textual “show notes,” images, and chapter markers embedded. Ultimately, though, a video of a soccer game cannot be distinguished from a video of a forest fire.

Search engines also cannot read any content contained within a program. The search engine really needs to find text that is readable by human eyes looking at the source code of a web page, as outlined earlier. It does not help if you can see it when the browser loads a web page—it has to be visible and readable in the source code for that page.

One example of a technology that can present significant human-readable content that the search engines cannot see is AJAX. AJAX is a JavaScript-based method for dynamically rendering content on a web page after retrieving the data from a database, without having to refresh the entire page. This is often used in tools where a visitor to a site can provide some input and the AJAX tool then retrieves and renders the correct content.

The problem arises because the content is retrieved by a script running on the client computer (the user's machine) only after receiving some input from the user. This can result in many potentially different outputs. In addition, until that input is received the content is not present in the HTML of the page, so the search engines cannot see it. Google does offer specific tips on how to make AJAX applications crawlable, that you can see here: <http://code.google.com/web/ajaxcrawling/docs/getting-started.html>.

Similar problems arise with other forms of JavaScript that don't render the content in the HTML until a user action is taken. New forms of Javascript, such as Angular JavaScript (<https://angularjs.org/>) make this even more challenging for search engines (<http://stackoverflow.com/questions/13499040/how-do-search-engines-deal-with-angularjs-applications>).

As of HTML 5, a construct known as the embed tag (`<embed>`) was created to allow the incorporation of plug-ins into an HTML page. Plug-ins are programs located on the user's computer, not on the web server of your website. This tag is often used to incorporate movies or audio files into a web page. The `<embed>` tag tells the plug-in where it should look to find the data file to use. Content included through plug-ins may or may not be invisible to search engines.

Frames and iframes are methods for incorporating the content from another web page into your web page. Iframes are more commonly used than frames to incorporate content from another website. You can execute an iframe quite simply with code that looks like this:

```

<iframe src = "http://accounting.careerbuilder.com" width="100%" height="300">
  <p>Your browser does not support iframes.</p>
</iframe>

```

Frames are typically used to subdivide the content of a publisher's website, but they can be used to bring in content from other websites, as was done in Figure 2-21 with <http://accounting.careerbuilder.com> on the *Chicago Tribune* website.

The screenshot shows a framed page from careerbuilder.com. The main content area displays search results for 'Elgin, IL Jobs' with 591 results found. The search filters are set to 'Elgin, IL' and 'Last 30 Days'. The results are presented in a grid format. At the top of the page, there is a banner for email alerts for new Elgin, IL Jobs, with options for daily or weekly updates and a 'Save' button. The navigation bar at the top includes links for Home, My CareerBuilder, Find Jobs, Recommendations, Post Resumes, and Advice & Resources. The careerbuilder logo is at the top left, and an 'Employee' link is at the top right.

Figure 2-19. Framed page rendered in a browser

Figure 2-21 is an example of something that works well to pull in content (provided you have permission to do so) from another site and place it on your own. However, the search engines recognize an iframe or a frame used to pull in another site's content for what it is, and therefore may ignore the content inside the iframe or frame as it is content published by another publisher. In other words, they don't consider content pulled in from another site as part of the unique content of your web page.

Determining Searcher Intent and Delivering Relevant, Fresh Content

Modern commercial search engines rely on the science of information retrieval (IR). This science has existed since the middle of the twentieth century, when retrieval systems powered computers in libraries, research facilities, and government labs. Early in the development of search systems, IR scientists realized that two critical components comprised the majority of search functionality: relevance and importance (which we defined earlier in this chapter). To measure these factors, search engines perform document analysis (including semantic analysis of concepts across documents) and link (or citation) analysis.

Document Analysis and Semantic Connectivity

In *document analysis*, search engines look at whether they find the search terms in important areas of the document—the title, the metadata, the heading tags, and the body of the text. They also attempt to automatically measure the quality of the document based on document analysis, as well as many other factors.

Reliance on document analysis alone is not enough for today's search engines, so they also look at semantic connectivity. *Semantic connectivity* refers to words or phrases that are commonly associated with one another. For example, if you see the word *aloha* you associate it with Hawaii, not Florida. Search engines actively build their own thesaurus and dictionary to help them determine how certain terms and topics are related. By simply scanning their massive databases of content on the Web, they can use Fuzzy Set Theory and certain equations (described at <http://forums.searchenginewatch.com/showthread.php?threadid=48>) to connect terms and start to understand web pages/sites more like a human does.

The professional SEO practitioner does not necessarily need to use semantic connectivity measurement tools to optimize websites, but for those advanced practitioners who seek every advantage, semantic connectivity measurements can help in each of the following sectors:

- Measuring which keyword phrases to target
- Measuring which keyword phrases to include on a page about a certain topic
- Measuring the relationships of text on other high-ranking sites/pages
- Finding pages that provide “relevant” themed links

Although the source for this material is highly technical, SEO specialists need only know the principles to obtain valuable information. It is important to keep in mind that although the world of IR has hundreds of technical and often difficult-to-comprehend terms, these can be broken down and understood even by an SEO novice.

Table 2-1 explains some common types of searches in the IR field.

Table 2-1. Common types of searches

Proximity searches

A proximity search uses the order of the search phrase to find related documents. For example, when you search for “*sweet German mustard*” you are specifying only a precise proximity match. If the quotes are removed, the proximity of the search terms still matters to the search engine, but it will now show documents that don't exactly match the order of the search phrase, such as *Sweet Mustard—German*.

Fuzzy logic

Fuzzy logic technically refers to logic that is not categorically true or false. A common example is whether a day is sunny (is 50% cloud cover a sunny day, etc.). In search, fuzzy logic is often used for misspellings.

Boolean searches

These are searches that use Boolean terms such as AND, OR, and NOT. This type of logic is used to expand or restrict which documents are returned in a search.

Term weighting

Term weighting refers to the importance of a particular search term to the query. The idea is to weight particular terms more heavily than others to produce superior search results. For example, the

appearance of the word *the* in a query will receive very little weight in selecting the results because it appears in nearly all English language documents. There is nothing unique about it, and it does not help in document selection.

IR models (search engines) use Fuzzy Set Theory (an offshoot of fuzzy logic created by Dr. Lotfi Zadeh in 1969) to discover the semantic connectivity between two words. Rather than using a thesaurus or dictionary to try to reason whether two words are related to each other, an IR system can use its massive database of content to puzzle out the relationships.

Although this process may sound complicated, the foundations are simple. Search engines need to rely on machine logic (true/false, yes/no, etc.). Machine logic has some advantages over humans, but machine logic doesn't have a way of thinking like humans, and things that are intuitive to humans can be quite hard for a computer to understand. For example, both oranges and bananas are fruits, but both oranges and bananas are not round. To a human this is intuitive.

For a machine to understand this concept and pick up on others like it, semantic connectivity can be the key. The massive human knowledge on the Web can be captured in the system's index and analyzed to artificially create the relationships humans have made. Thus, a machine knows an orange is round and a banana is not by scanning thousands of occurrences of the words *banana* and *orange* in its index and noting that *round* and *banana* do not have great concurrence, while *orange* and *round* do.

This is how the use of fuzzy logic comes into play, and the use of Fuzzy Set Theory helps the computer to understand how terms are related simply by measuring how often and in what context they are used together.

For example, a search engine would recognize that *trips* to the *zoo* often include *viewing wildlife* and *animals*, possibly as part of a *tour*.

To see this in action, conduct a search on Google for *zoo trips*. Note that the boldface words that are returned match the terms that are italicized in the preceding paragraph. Google is setting “related” terms in boldface and recognizing which terms frequently occur concurrently (together, on the same page, or in close proximity) in their indexes.

Search companies have been investing in these types of technologies for many years. In September of 2013, Google quietly let the world know that it had rewritten its search engine and given it the name of Hummingbird (<http://searchengineland.com/google-birthday-event-172791>). This rewrite was in large part done to enable a whole new set of capabilities for recognizing the relationships between things.

For example, if you use Google's voice search (Click on the microphone icon at the right of the search box on Google.com) and ask it “Who is Tom Brady?”, it will answer that question for you with a search result, but then use audio to tell you that he is an “American football quarterback for the New England Patriots of the National Football League”.

This shows that Google understands many aspects of Tom Brady. For example:

- He has an occupation: quarterback, playing American football (as distinct from the way the term football is used outside of the US and Canada).
- That he plays on a team, the New England Patriots
- That the New England Patriots belong to a league, the NFL

This is far more sophisticated than search was in 2012. You can take this much further. For example, if you now use the voice search feature to ask “who is his wife?”, it will answer that question too (see Figure 2-22).

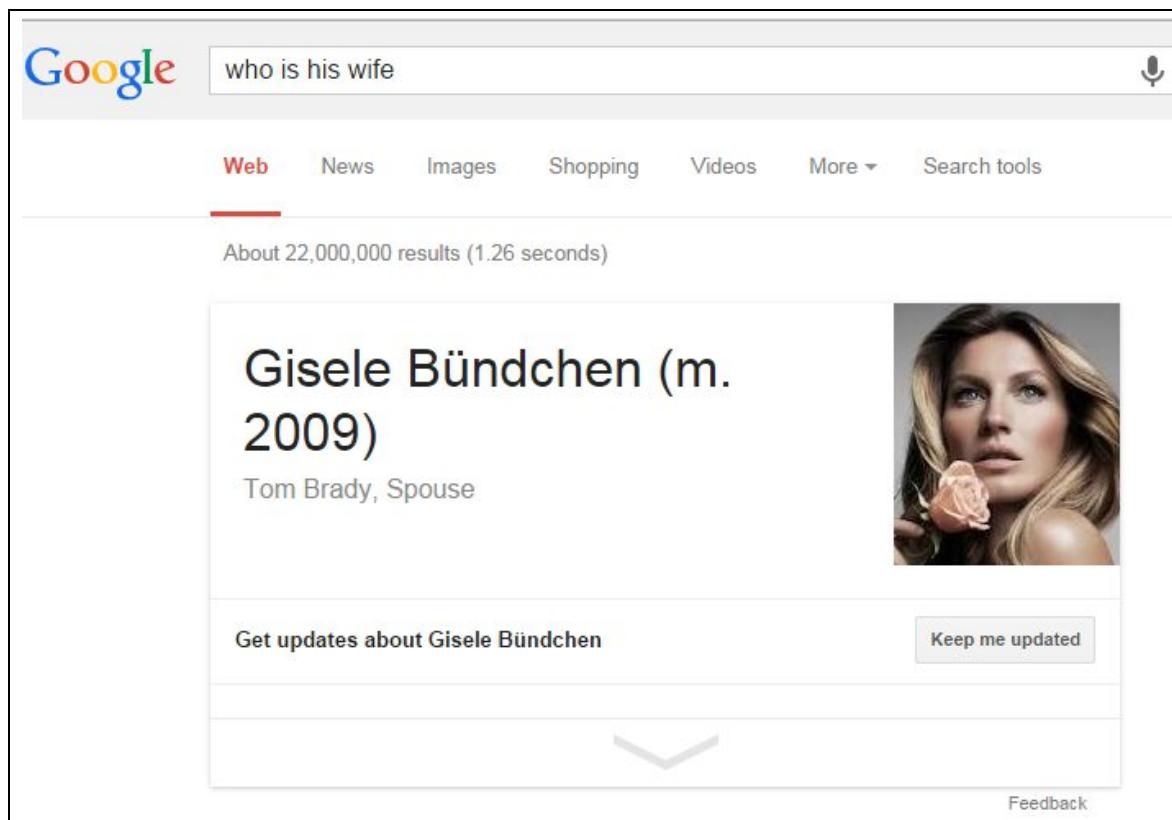


Figure 2-22. *Gisele Bündchen is Tom Brady's wife*

Notice too that in the above query that we did not repeat Tom Brady's name, and that Google remembered the context of the conversation, in that "his" refers to Tom Brady. You can continue with questions such as does he have children, and Google will answer that one as well.

For SEO purposes, this usage opens our eyes to realizing how search engines recognize the connections between words, phrases, and ideas on the Web. As semantic connectivity becomes a bigger part of search engine algorithms, you can expect greater emphasis on the theme of pages, sites, and links. It will be important going into the future to realize the search engines' ability to pick up on ideas and themes and recognize content, links, and pages that don't fit well into the scheme of a website.

Measuring Content Quality and User Engagement

The search engines also attempt to measure the quality and uniqueness of a web site's content. One method they may use for doing this is by evaluating the document itself. For example, if a web page has lots of spelling and grammatical errors, that can be taken as a sign that little editorial effort was put into that page (you can see more on this here: <http://searchengineland.com/google-pagerank-spelling-correlation-95821>).

They can also analyze the reading level of the document. One popular formula for doing this is the Flesch-Kincaid Grade Level Readability Formula which considers things like the average word length and the words per sentence to determine the level of education needed to be able to understand the sentence. Imagine a scenario where the product being sold on a page is children's toys and the reading level calculated suggests a grade level of a senior in college to read the page. This could be another indicator of a poor editorial effort.

The other method that search engines can use to evaluate the quality of a web page is to measure actual user interaction. For example, if a large number of users who visit the web page after clicking on a search result immediately return to the search engine and click on the next result that would be a strong indicator of poor quality.

Engagement with a web site began to publicly emerge as a ranking factor with the release of the Panda update by Google on February 23, 2011 (<http://searchengineland.com/google-forecloses-on-content-farms-with-farmer-algorithm-update-66071>). Google has access to a large number of data sources that it can use to measure how visitors interact with your web site. Just because Google has access to this data, it does not mean that they are definitely using it as a ranking factor. That noted, some of those sources include:

Interaction with Web Search Results

For example, if a user clicks through on a SERP listing and comes to your site, clicks the Back button, and then clicks on another result in the same set of search results, that could be seen as a negative ranking signal. Or if the results below you in the SERPs are getting clicked on more than you are, that could be seen as a negative ranking signal for you and a positive ranking signal for them. Whether search engines use this signal or not, or how much weight they might put on it, is not known.

Google Analytics

It is hard to get a firm handle on just what percentage of web sites run Google Analytics. A 2008 survey of web sites by immeria.net showed their share at 59% (<http://blog.immeria.net/2008/01/web-analytics-vendors-market-shares.html>), and the Metric Mail Blog checked the top 1 Million sites in Alexa and found that about 50% of those had Google Analytics (<http://metricmail.tumblr.com/post/904126172/google-analytics-market-share>). Suffice it to say that Google is able to collect detailed data about what is taking place on a large percentage of the world's web sites.

This provides Google with a rich array of data on that site, including:

- Bounce Rate: The percentage of visitors who visit only one page on your web site.
- Time on Site: The time spent by the user on the site. Note that Google Analytics only receives information when each page is loaded, so if you view only one page it does not know how much time you spent on that page. More precisely then, this metric tells you the average time between the loading of the first page and the loading of the last page, but does not take into account how long visitors spent on the last page loaded.
- Page Views per Visitor: The average number of pages viewed per visitor on your site.

Google Toolbar

It is not known how many users out there use the Google Toolbar, but the authors believe that it numbers in the millions. For these users, Google can track their entire web surfing behavior. Unlike Google Analytics, the Google Toolbar can measure the time from when a user first arrives on a site to the time when they load a page from a different web site. It can also get measurements of bounce rate and page views per visitor.

Google +1 Button

This enables users to vote for a page on the page itself. There is currently no evidence that Google uses this as a ranking factor, but in theory, it could. You can see a lot more about this in Chapter 8.

Chrome Personal Blocklist Extension

Google offers a Chrome add-on called the Personal Blocklist Extension (<http://googleblog.blogspot.com/2011/02/new-chrome-extension-block-sites-from.html>). This provides users of the Chrome browser a way to indicate a search result they don't like. This was first used by Google as a part of its Panda algorithm, which attempts to measure the quality of a piece of content. You can read more about this algorithm in Chapter 9.

Goo.gl

Google has its own URL shortener. This tools allows Google to see what content is being shared, and which content is being clicked on, even in closed environments where Google web crawlers are not allowed to go.

What matters most is how your site compares to that of your competition. If your site has better engagement metrics this is likely to be seen as an indication of quality and can potentially boost your rankings with respect to your competitors. Little has been made public about the way search engines use these types of signals, so the above comments are speculation by the authors on what Google may be doing in this area. One of the most interesting posts that Google put up on this topic can be found here: <http://googlewebmastercentral.blogspot.com/2011/05/more-guidance-on-building-high-quality.html>. Social and User Engagement ranking factors are discussed in more detail in Chapter 8 "How Social Media and User Data Play a Role in Search Results and Rankings"

Link Analysis

In link analysis, search engines measure who is linking to a site or page and what they are saying about that site/page. They also have a good grasp on who is affiliated with whom (through historical link data, the site's registration records, and other sources), who is worthy of being trusted based on the authority of sites linking to them, and contextual data about the site on which the page is hosted (who links to that site, what they say about the site, etc.).

Link analysis goes much deeper than counting the number of links a web page or website has, as links are not created equal (one link can be worth 10 million times more than another one). Links from a highly authoritative page on a highly authoritative site will count more than other links of lesser authority. A website or page can be determined to be authoritative by combining an analysis of the linking patterns and semantic analysis.

For example, perhaps you are interested in sites about dog grooming. Search engines can use semantic analysis to identify the collection of web pages that focus on the topic of dog grooming. The search engines can then determine which of these sites about dog grooming have the most links from the set of dog grooming sites. These sites are most likely more authoritative on the topic than the others.

The actual analysis is a bit more complicated than that. For example, imagine that there are five sites about dog grooming with a lot of links from pages across the Web on the topic, as follows:

- Site A has 213 topically related links.
- Site B has 192 topically related links.
- Site C has 203 topically related links.
- Site D has 113 topically related links.
- Site E has 122 topically related links.

Further, it may be that Site A, Site B, Site D, and Site E all link to each other, but none of them link to Site C. In fact, Site C appears to have the great majority of its relevant links from other pages that are topically relevant but have few links to them. In this scenario, Site C may not be considered authoritative because it is not linked to by the right sites.

The concept of grouping sites based on who links to them and who they link to, is referred to as grouping sites by *link neighborhood*. The neighborhood you are in says something about the subject matter of your site, and the number and quality of the links you get from sites in that neighborhood say something about how important your site is to that topic.

The degree to which search engines rely on evaluating link neighborhoods is not clear, and links from irrelevant pages can still help the rankings of the target pages. Nonetheless, the basic idea remains that a link from a relevant page or site should be more valuable than a link from a non-relevant page or site.

Another factor in determining the value of a link is the way the link is implemented and where it is placed. For example, the text used in the link itself (i.e., the actual text that will go to your web page when the user clicks on it) is also a strong signal to the search engines.

This is referred to as *anchor text*, and if that text is keyword-rich (with keywords relevant to your targeted search terms), it can potentially do more for your rankings in the search engines than if the link is not keyword-rich. For example, anchor text of "Dog Grooming Salon" may bring more value to a dog grooming salon's website than anchor text of "Click here". However, take care. If you get 10,000 links

using the anchor text “Dog Grooming Salon” and you have few other links to your site, this definitely does not look natural and could lead to a penalty.

The semantic analysis of a link’s value does go deeper than just the anchor text. For example, if you have that “Dog Grooming Salon” anchor text on a web page that is not really about dog grooming at all, the value of the link is less than if the page is about dog grooming. Search engines also look at the content on the page immediately surrounding the link, as well as the overall context and authority of the website that is providing the link.

All of these factors are components of link analysis, which we will discuss in greater detail in Chapter 7.

Evaluating Social Media Signals

Sites such as Facebook (<http://www.facebook.com>), Twitter (<http://www.twitter.com>), and Google+ (<http://plus.google.com>) have created whole new ways for users to share content or indicate that they value it. This has led many to speculate that search engines could be using these signals as a ranking factor. Fueling that speculation, in August of 2013 Moz released the data from their latest correlation study (<http://moz.com/search-ranking-factors>) and this showed a very strong correlation between +1s and ranking in Google.

Figure 2-23 shows the top 10 results in that data, and Google +1s had the second strongest correlation with rankings.

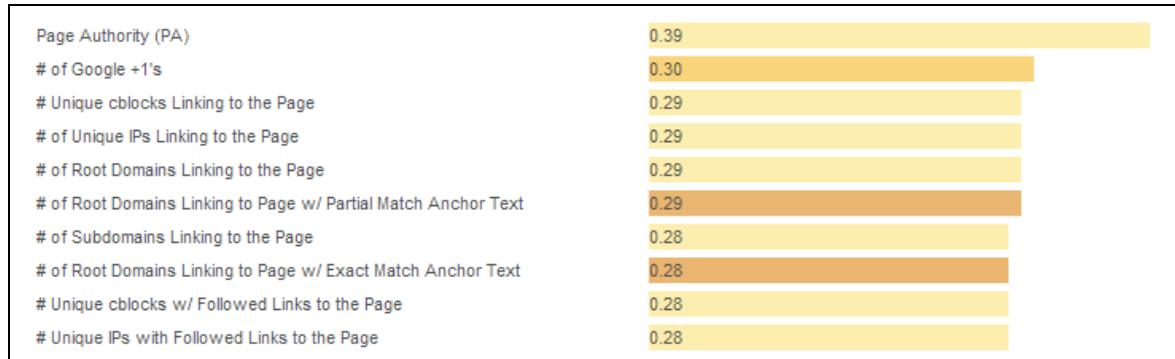


Figure 2-23. Top 10 Results in Moz’ 2013 Correlation Study

However, the fact that there is a correlation in no way means that +1s are used by Google as a ranking signal, or that they cause higher rankings. It can simply mean that good content that receives lots of links (which is known to be a signal that causes higher rankings) also happens to get many +1s.

In fact, Stone Temple Consulting did a different study targeted at measuring whether or not Google Plus activity was used as a ranking factor by Google (<http://www.stonetemple.com/measuring-google-plus-impact-on-search-rankings/>). This study showed that there was no material evidence that Google Plus +1s or shares had any material impact on rankings. The potential for social signals as a ranking factor is discussed in depth in Chapter 8.

Problem Words, Disambiguation, and Diversity

On the opposite side of the coin are words that present an ongoing challenge for the search engines. One of the greatest challenges comes in the form of disambiguation. For example, when someone types in *boxers*, does he mean the prize fighter, the breed of dog, or the type of underwear? Another example of this is *jaguar*, which is at once a jungle cat, a car, a football team, an operating system, and a guitar. Which does the user mean?

Search engines deal with these types of ambiguous queries all the time. The two examples offered here have inherent problems built into them, but the problem is much bigger than that. For example, if someone types in a query such as *cars*, does he:

- Want to read reviews?
- Want to go to a car show?
- Want to buy one?
- Want to read about new car technologies?

The query *cars* is so general that there is no real way to get to the bottom of the searcher's intent based on this one query alone. One way that search engines deal with this is by looking at prior queries by the same searcher may provide additional clues to their intent - we discuss this a bit more in "Adaptive Search" in Chapter 1).

Another solution they use is to offer diverse results. As an example, Figure 2-24 shows a generic search, this time using *GDP*.

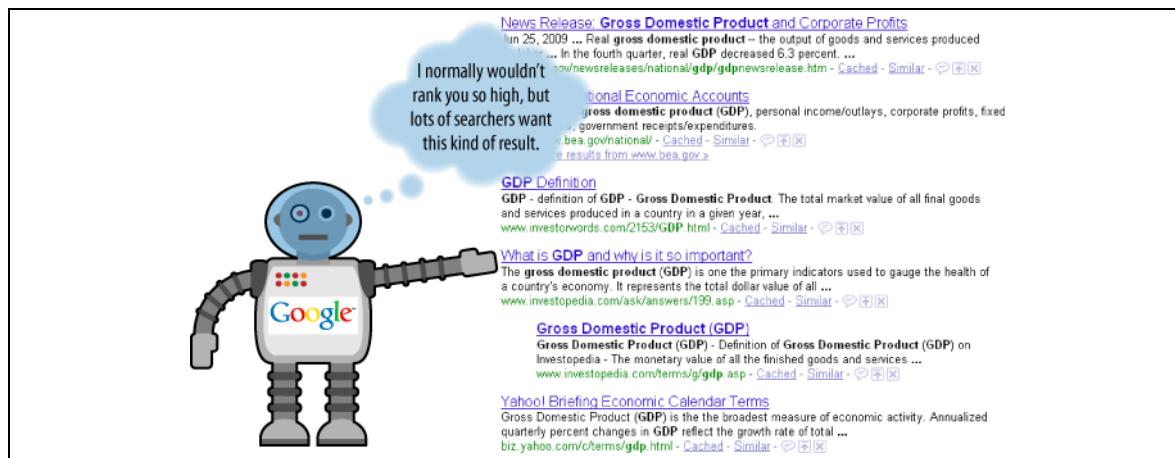


Figure 2-20. Diverse results example

This brings up an important ranking concept. It is possible that a strict analysis of the relevance and link-driven importance scores in Figure 2-24 would not have resulted by itself in the Investopedia.com page being on the first page, *but* the need for diversity caused the ranking of the page to be elevated. This concept of altering the results in this manner is sometimes referred to as *query deserves diversity* (QDD).

A strict relevance- and importance-based ranking system might have shown a variety of additional government pages discussing the GDP of the United States. However, a large percentage of users will likely be satisfied by the government pages already shown, but for those users who are not satisfied with those results, showing more of the same types of pages is not likely to raise the level of satisfaction with the results.

Introducing a bit of variety allows Google to also provide a satisfactory answer to those who are looking for something different from the government pages. Google's testing has shown that this diversity-based approach has resulted in a higher level of satisfaction among its users.

For example, the testing data for the non-diversified results may have shown lower click-through rates in the SERPs, greater numbers of query refinements, and even a high percentage of related searches performed subsequently.

The idea to deliberately introduce diversity into the result algorithm makes sense and can enhance searcher satisfaction for queries such as:

- Company names (where searchers might want to get positive and negative press, as well as official company domains)

- Product searches (where e-commerce-style results might ordinarily fill up the SERPs, but Google tries to provide some reviews and noncommercial, relevant content)
- News and political searches (where it might be prudent to display “all sides” of an issue, rather than just the left- or right-wing blogs that did the best job of obtaining links)

Search engines also personalize results for users based on their search history or past patterns of behavior. For example, if a searcher has a history of searching on card games, and then does a search for *dominion*, the search engine may choose to push some of the results related to the Dominion card game higher in the results, instead of emphasizing the power company.

Where freshness matters

Much of the time, it makes sense for the search engines to deliver results from older sources that have stood the test of time. However, other times the response should be from newer sources of information.

For example, when there is breaking news, such as an earthquake, the search engines begin to receive queries within seconds, and the first articles begin to appear on the Web within 15 minutes.

In these types of scenarios, there is a need to discover and index new information in near-real time. Google refers to this concept as *query deserves freshness* (QDF). According to the *New York Times* (<http://www.nytimes.com/2007/06/03/business/yourmoney/03google.html?pagewanted=3>), QDF takes several factors into account, such as:

- Search volume
- News coverage
- Blog coverage

QDF applies to up-to-the-minute news coverage, but also to other scenarios such as hot, new discount deals or new product releases that get strong search volume and media coverage. There has also been speculation that Google will apply QDF more to sites that have higher PageRank (<http://www.copyblogger.com/query-deserves-freshness/>).

A Few Reasons Why These Algorithms Sometimes Fail

As we’ve outlined in this chapter, the search engines do some amazing stuff. Nonetheless, there are times when the process does not work as well as you would like to think. Part of this is because users often type in search phrases that provide very little information about their intent (e.g., if they search on *car*, do they want to buy one, read reviews, learn how to drive one, learn how to design one, or something else?). Another reason is that some words have multiple meanings, such as the *jaguar* example we used previously in this section.

For more information on reasons why search algorithms sometimes fail, you can read the following Moz article, which was written by Hamlet Batista:

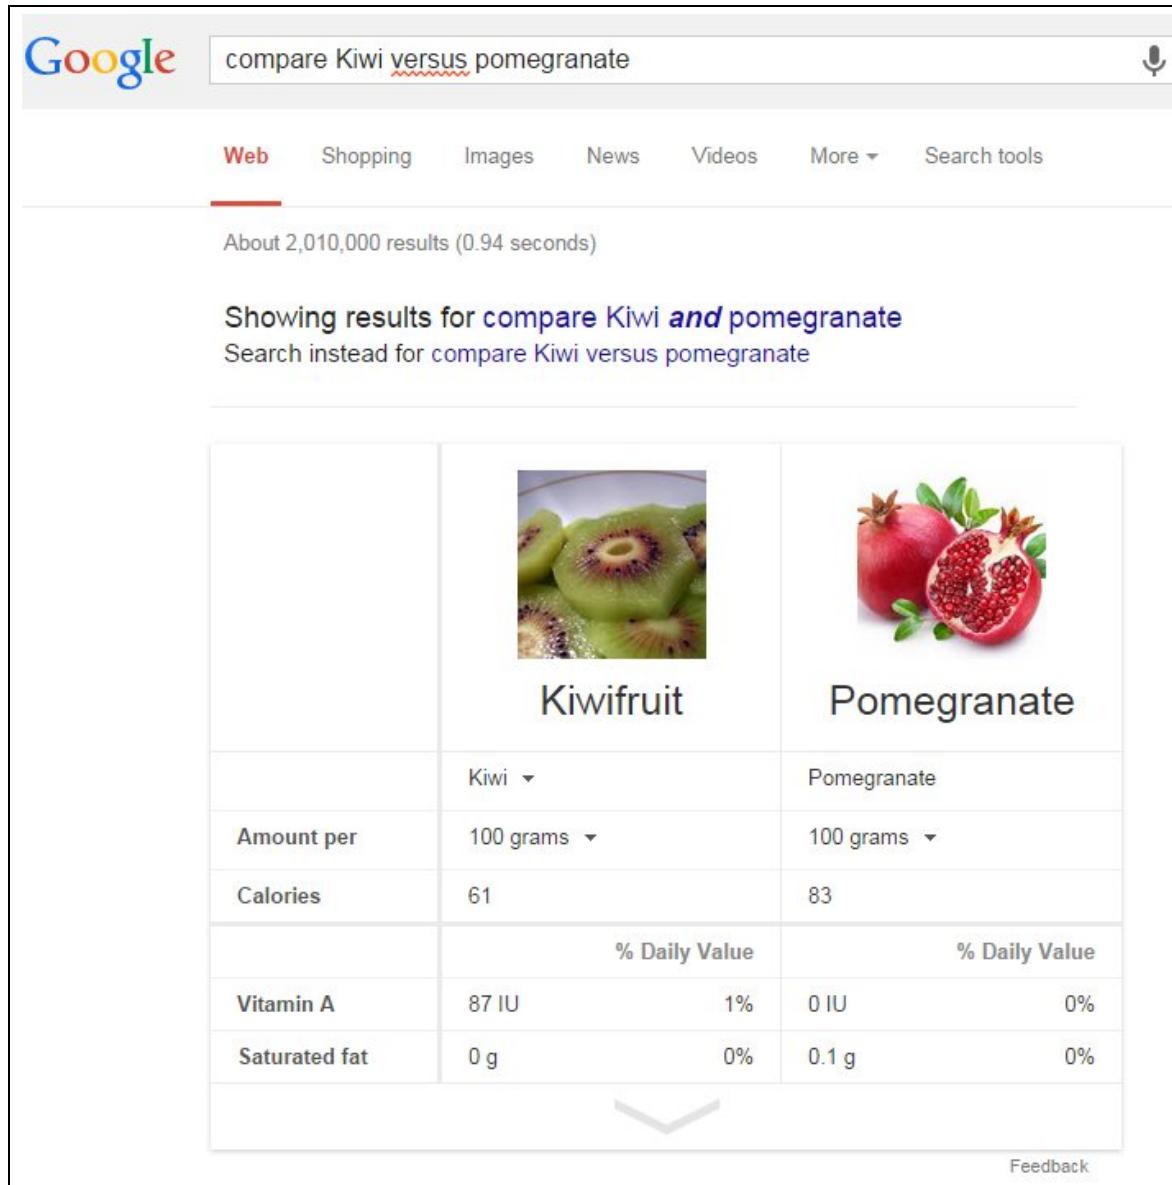
<http://www.moz.com/blog/7-reasons-why-search-engines-dont-return-relevant-results-100-of-the-time>

The Knowledge Graph

Traditional search results are derived by search engines crawling and analyzing web pages and then presenting that information in the search results. However, Google’s mission “is to organize the world’s information and make it universally accessible and useful” (<https://www.google.com/about/company/>). They are actively pursuing initiatives to build databases of information that go far beyond traditional web based search.

Note that earlier in this chapter we wrote about vertical search. Vertical search relates to breaking search into different categories, such as a search for images, videos, or local business information. The Knowledge Graph is more about providing richer answers directly in the search results, often answering the user’s question directly without their having to click through to a web site.

In May of 2012, Google announced the Knowledge Graph (<http://googleblog.blogspot.com/2012/05/introducing-knowledge-graph-things-not.html>). Initially, this was a set of structured databases of information that allows Google to access information without deriving it from the web. You can see an example of the type of data that Google might extract from its Knowledge Graph database in Figure 2-25.



The screenshot shows a Google search results page for the query "compare Kiwi versus pomegranate". The results are for the "Web" category. The search took 0.94 seconds and found about 2,010,000 results. The results are titled "Showing results for compare Kiwi and pomegranate" and suggest "Search instead for compare Kiwi versus pomegranate". Below the search bar, there are images of kiwi fruit and pomegranates. A comparison table follows, showing nutritional data for 100 grams of each fruit. The table includes columns for Kiwi and Pomegranate, with dropdown menus for amount per (100 grams) and rows for Calories, Vitamin A, and Saturated fat. A "Feedback" link is at the bottom right of the table.

	Kiwifruit	Pomegranate
Kiwi ▾		Pomegranate
Amount per	100 grams ▾	100 grams ▾
Calories	61	83
	% Daily Value	% Daily Value
Vitamin A	87 IU	1% 0 IU 0%
Saturated fat	0 g	0% 0.1 g 0%

Figure 2-25. Kiwi vs Pomegranate Search Result

Google initially built the Knowledge Graph using data from Freebase, Wikipedia, and the CIA Fact Book. This allowed them to answer many questions, but really only satisfied a very small number of search queries. For that reason, Google is constantly working on expanding the information in the Knowledge Graph.

In addition, Google is investing in ways to more reliably extract information from other sources, including web sites, to present as direct answers in search. Figure 2-26 shows the search result for buying a car.

Google search results for "how to buy a car". The results page shows the following:

- The Future of Car Buying - truecar.com**
Ad www.truecar.com/ ▾
4.3 ★★★★★ rating for truecar.com
Never overpay and enjoy a hassle-free car buying experience
customer support · certified dealers · guaranteed savings · hassle-free pricing
Get upfront pricing before even visiting a dealership – Nationwide
Search Used Cars See What Others Paid
2014 New Car Pricing Select Make/Model
- How To Buy A Car - Auto Buying Tips, Reviews, Pricing**
Ad www.edmunds.com/ ▾
4.5 ★★★★★ rating for edmunds.com
Get Local Dealer Invoice Prices.
Ratings: Site content 10/10 - Site design 9/10 - Community 9/10
Edmunds.com has 1,788,746 followers on Google+
\$15,000 to \$25,000 - \$55,000 to \$85,000 - \$45,000 to \$55,000
- CARFAX® Used Car Listings - CARFAX.com**
Ad www.carfax.com/Cars_For_Sale ▾ 4.6 ★★★★★ advertiser rating
Find The Right Car With The Right History w/ CARFAX Used Car Listings
- Buying a Car - Money Essentials, Lesson 17 - CNN Money**
money.cnn.com/magazines/moneymag/money101/lesson17/ CNNMoney ▾
Feedback

A red box highlights the CNN Money result, specifically the list of steps under the heading "Buying a Car - Money Essentials, Lesson 17 - CNN Money".

- Make sure you are getting the right vehicle. ...
- Assess the worth of your old car. ...
- Decide whether new or used is best for you. ...
- Consider whether leasing or buying makes more sense. ...
- Do your homework and set your target price. ...
- Shop for money before you shop for the car. ...
- Negotiating a lease. ...
- Negotiate a purchase.

2 more items

Figure 2-26. Step by Step Instructions for Buying a Car

In this result, Google provides a set of step by step instructions extracted from the CNN Money web site. Note that 2 steps are omitted, so to get the complete procedure, or additional details on each step, the user needs to click through to the CNN Money web site.

In some cases, Google does provide the complete instructions in the search results, but most of the time it does not. A study performed by Stone Temple Consulting examined 276 examples of step by step instructions, and found that 217 of these (79%) did not provide the complete instructions.

A related concept that people speak about is called **Semantic Search**, which overlaps to some degree with the Knowledge Graph, but also takes into account many other factors to personalize results for the searcher. You can see a depiction of some of these factors in Figure 2-27.

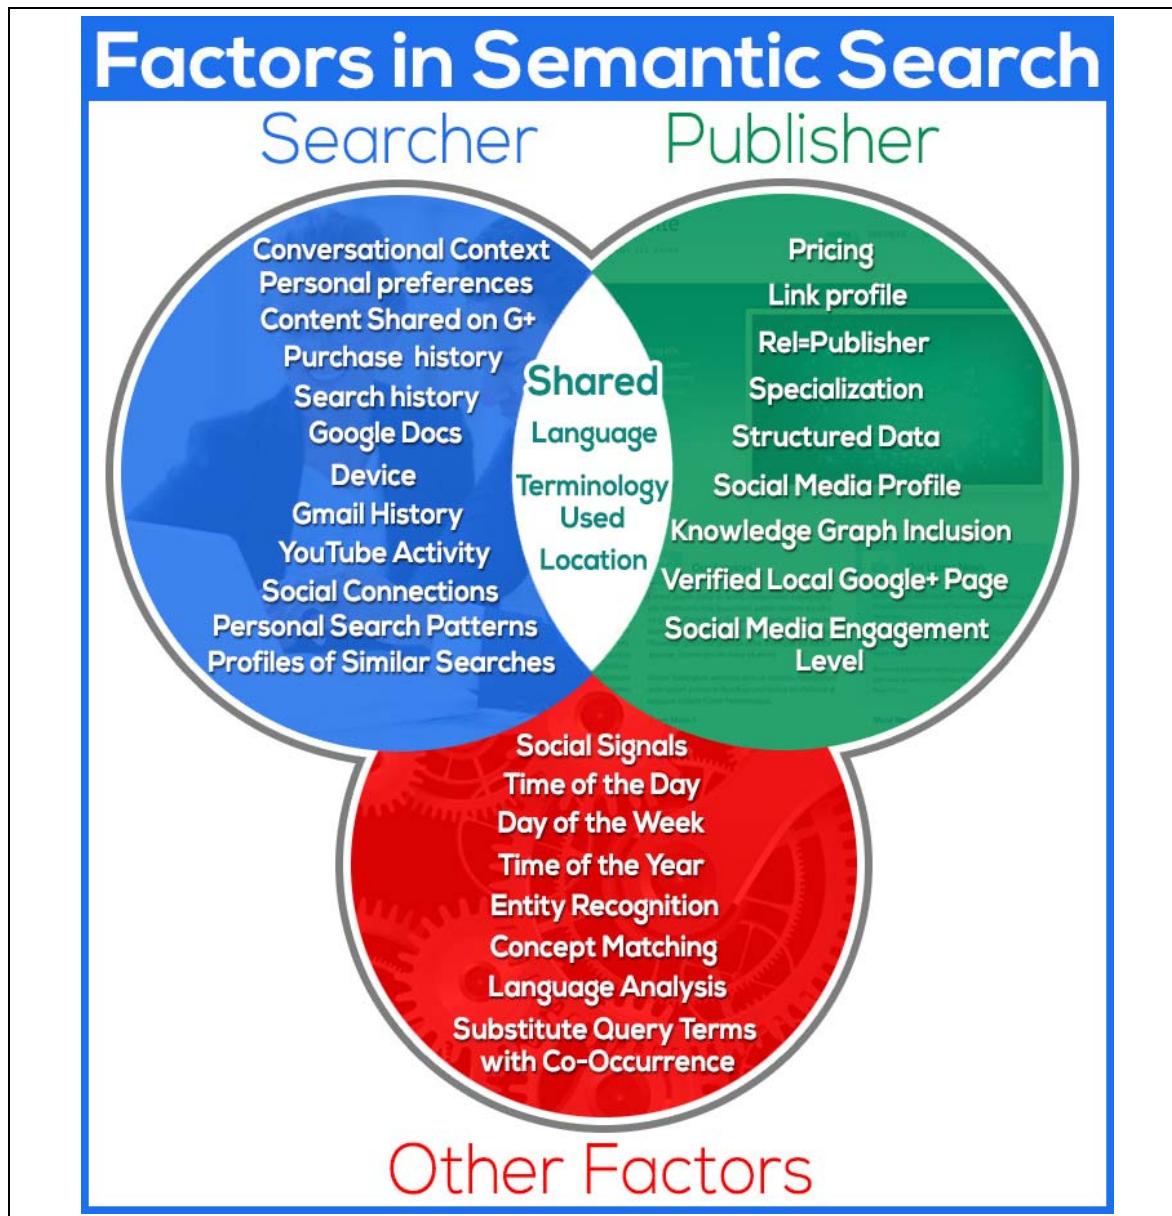


Figure 2-27. Factors Involved in Semantic Search

The Knowledge Graph and semantic search are discussed in more detail in Chapter 6.

Analyzing Ranking Factors

Moz periodically conducts surveys of leading SEOs to determine what they think are the most important ranking factors (<http://www.moz.com/article/search-ranking-factors>). Here is a high-level summary of the top nine results, in priority order (as suggested by the referenced study):

- Domain Level Link Authority Features
- Page Level Link Metrics
- Page Level Keyword & Content Based
- Page Level Keyword Agnostic Features
- Domain Level Brand Metrics
- Usage & Traffic/Query Data
- Page Level Social Metrics
- Domain Level Keyword Usage
- Domain Level Keyword Agnostic Features

Here is a brief look at each of these:

Domain Level Link Authority Features

Domain Level Link Authority is based on a cumulative link analysis of all the links to the domain. This includes factors such as the number of different domains linking to the site, the trust/authority of those domains, rate at which new inbound links are added, the relevance of the linking domains, and more.

Page Level Link Metrics

This refers to the links as related to the specific page, such as the number of links, the relevance of the links, and the trust and authority of the links received by the page.

Page Level Keywords & Content

This describes use of the keyword term/phrase in particular parts of the HTML code on the page (title element, H1s, alt attributes, etc.).

Page Level Features Other than Keywords

Factors included here are on page elements such as number of links on the page, number of internal links, number of followed links, number of “NoFollow” links, and other similar factors.

Domain Level Brand Metrics

This factor includes search volume on the web site's brand name, mentions, whether they have a presence in social media and other brand related metrics.

Page Level Traffic/Query Data

Elements of this factor are Click Through Rate (CTR) to the page in the search results, bounce rate of visitors to the page, and other similar measurements.

Page Level Social Metrics

Social metrics considered include mentions, links, shares, Likes and other social media site based metrics. It should be emphasized that many SEO practitioners believe that this is a ranking factor even though studies have since shown otherwise, and representatives from Google clearly state that social signals are not part of their algorithm.

Domain Level Keyword Usage

This refers to how keywords are used in the root or subdomain name, and how impactful that might be on search engine rankings.

Domain Level Keyword Agnostic Features

Major elements of this factor in the survey included number of hyphens in the domain name, number of characters in the domain name, and domain name length.

Negative Ranking Factors

It's also possible to have negative ranking factors. For example, if a site has a large number of low quality inbound links that appear to be the result of artificial efforts by the publisher to influence search rankings, the site's rankings can be lowered. This is, in fact, exactly what Google's *Penguin* algorithm does. This algorithm is discussed more in Chapter 9. Some other potential negative ranking factors include:

Malware being hosted on the site

The search engines will act rapidly to penalize sites that contain viruses or trojans.

Cloaking

Search engines want publishers to show the same content to the search engine as is shown to users.

Pages on the sites with links for sale

Google has a strong policy against paid links (<http://www.google.com/support/webmasters/bin/answer.py?answer=66736>), and sites that sell links may be penalized.

Content which advertises paid links on the site

As an extension of the prior negative ranking factor, promoting the sale of paid links may be a negative ranking factor.

Page Speed

Back in 2010, Google's Matt Cutts announced that Google was making Page Speed a ranking factor (<https://www.mattcutts.com/blog/site-speed/>). In general, it is believed that it is a negative factor for pages that are exceptionally slow.

Other Ranking Factors

The ranking factors we've discussed so far are really just the basics. Search engines potentially factor in many more signals. Some of these include:

Rate of acquisition of links

If, over time, your site has acquired an average of 5 links per day, and then the links suddenly start to come in at a rate of 10 per day, then that could be seen as a positive ranking signal. On the other hand, if the rate of new links drops to two per day, that could be a signal that your site has become less relevant.

However, it gets more complicated than that. If your site suddenly starts to get 300 new links per day, you have either become a lot more relevant or started to acquire links in a spammy way. The devil is in the details here, with one of the most important details being from where those new links are coming. The concept of considering temporal factors in link analysis is documented in a U.S. patent held by Google and found at the following URL:

<http://appft1.uspto.gov/netacgi/nph-Parser?Sect1=PTO1&Sect2=H1OFF&d=PG01&p=1&u=%2Fnetacgi%2FPTO%2Fsrchnum.html&r=1&f=G&l=50&s1=%2220050071741%22.PGNR.&OS=DN/20050071741&RS=DN/20050071741>

Alternatively, you can look this patent up by searching for patent number 20050071741.

User data

Personalization is one of the most talked about frontiers in search. There are a few ways personalization can take place. For one, a search engine can perform a geolocation lookup to figure out where you are approximately located. Based on this, a search engine can show results tailored to your current location. This is very helpful, for example, if you are looking for a local restaurant.

Another way a search engine can get some data on a user is if the user creates a profile with the search engine and voluntarily provides some information. A simple example would be a language preference. If the user indicates he prefers Portuguese, the search engine can tailor the results to that preference.

Search engines can also look at the search history for a given user. Basically, the search engine maintains a log of all the searches you have performed when you are logged in. Based on this, it can see that you have been checking out luxury cars recently, and can use that knowledge to tweak the results you see after you search on *jaguar*. This is sometimes referred to as adaptive search.

To reduce the level of personalization, users can log out of their Google account. However, this does not disable all personalization, as Google may still tie some history to the person's computer. Disabling all personalization can be done by using Google's Chrome browser in *Incognito* mode. This will allow the user to see Google results which are not personalized based on search history. However, the results will still be personalized to your location.

You can also de-personalize your search results by performing your search query, and then appending &pws=0 to the end of the search page URL and reloading the page. Note, this only works if you have turned off Google Instant (which is Google's feature of showing results instantly as you type). Or, you can choose the option "Disable customizations based on web history" under the "webhistory" under the Gear icon in the SERPS.

Using Advanced Search Techniques

One of the basic tools of the trade for an SEO practitioner is the search engines themselves. They provide a rich array of commands that can be used to perform advanced research, diagnosis, and competitive analysis. Some of the more basic operators are below (note that square brackets in the examples are not meant to be part of the query):

`[-keyword]`

Excludes the keyword from the search results. For example, `[loans -student]` shows results for all types of loans *except* student loans.

`["key phrase"]`

Shows search results for the exact phrase—for example, `["seo company"]`. You can also use “” to force the inclusion of specific word. This is particularly useful for including *stopwords* (keywords that are normally stripped from a search query because they usually do not add value, such as the word *the*) in a query, or if your keyword is getting converted into multiple keywords through automatic stemming. For example, if you mean to search for the TV show *The Office*, you would want the word *The* to be part of the query. As another example, if you are looking for Patrick Powers, who was from Ireland, you would search for “*patrick powers*” *Ireland* to avoid irrelevant results for Patrick Powers.

`[keyword1 OR keyword2]`

Shows results for *at least one* of the keywords—for example, `[google OR Yahoo!]`.

These are the basics, but for those who want more information, what follows is an outline of the more advanced search operators available from the search engines.

Advanced Google Search Operators

Google supports a number of advanced search operators (<http://www.google.com/support/websearch/bin/static.py?hl=en&page=guide.cs&guide=1221265&answer=136861&rd=1>) that you can use to help diagnose SEO issues. Table 2-2 gives a brief overview of the queries, how you can use them for SEO purposes, and examples of usage.

Table 2-2. Google advanced search operators

Operator	Short description	SEO application	Examples
(site:)—Domain-restricted search	Narrows a search to a specific domain(s)/directories	(a) Shows approximately how many URLs are indexed by Google From a directory Includes all subdomains Shows sites of a specific top-level domain (TLD)	site:www.google.com site:mit.edu/research/ site:google.com site:org
(inurl:)/(allinurl:)—URL keyword restricted search	Narrows the results to documents containing (a) search term(s) in the URLs	Find web pages having your keyword in a file path	inurl:seo inurl:company = allinurl:seo company
(intitle:)/(allintitle:)—Title keyword restricted search	Restricts the results to documents containing (a) search term(s) in a page title	Find web pages using your keyword in a page title	intitle:seo intitle:company = allintitle:seo company
(inanchor:)/(allinanchor:)—Anchor text keyword restricted search	Restricts the results to documents containing (a) search term(s) in the anchor text of backlinks pointing to a page	Find pages having the most backlinks/the most powerful backlinks with the keyword in the anchor text	inanchor:seo inanchor:company = allinanchor:seo company
(intext:)—Body text keyword restricted search	Restricts the results to documents containing (a) search term(s) in the body text of a page	Find pages containing the most relevant/most optimized body text	intext:seo
(ext:)/(filetype:)—File type restricted search	Narrows search results to the pages that end in a particular file extension	A few possible extensions/file types:	filetype:pdf ext:pdf

		PDF (Adobe Portable Document Format) HTML or .htm (Hypertext Markup Language) .xlsx (Microsoft Excel) .pptx (Microsoft PowerPoint) .docx (Microsoft Word)
(related:)—Similar URLs search	Shows <i>related pages</i> by finding pages linking to the site and looking at what else they tend to link to (i.e., “co-citation”); usually 25 to 31 results are shown.	Evaluate how relevant the site’s “neighbors” are
(info:)—Information about a URL search	Gives information about the given page	Learn whether the page has been indexed by Google; provides links for further URL information; this search can also alert you to possible site issues (duplicate content or possible DNS problems)
(cache:)—See what the page looked like when Google crawled it	Shows Google’s saved copy of the page	Google’s <i>text</i> version of the page works the same way as SEO Browser

When using the `site:` operator, some indexed URLs might not be displayed (even if you use the “repeat the search with omitted results included” link to see the full list). The `site:` query is notoriously inaccurate. You can obtain a more accurate count of the pages of your site indexed by Google by appending `&start=990&filter=0` to the URL of a Google result set for a `site:` command.

This tells Google to start with result 990, which is the last page Google will show you since it limits the results to 1,000. This must take place in two steps. First, enter a basic `site:yourdomain.com` command, and then get the results. Then go up to the address bar and append the `&start=990&filter=0` parameters to the end of the URL. Once done with this, you can look at the total pages returned to get a more accurate count. Note that this only works if Google Instant is turned off.

To see more results, you can also use the following search patterns:

- `[site:yourdomain.com/subdirectory1] + [site:yourdomain.com/subdirectory2] + etc.` (the “deeper” you dig, the more/more accurate results you get)
- `[site:yourdomain.com inurl:keyword1] + [site:yourdomain.com inurl:keyword2] + etc.` (for subdirectory-specific keywords)
- `[site:yourdomain.com intitle:keyword1] + [site:yourdomain.com intitle:keyword2] + etc.` (for pages using the keywords in the page title)

To learn more about Google advanced search operators, check out Stephan Spencer’s book “Google Power Search”.

Combined Google queries

To get more information from Google advanced search, it helps to learn how to effectively combine search operators. Table 2-3 illustrates which search patterns you can apply to make the most of some important SEO research tasks.

Table 2-3. Combined Google search options

What for	Description	Format	Example
Competitive analysis	Find recent mentions of your competitor not on their own site; use the date range option under the search tools in the SERPs; the following brand-specific search terms can be used: [domainname.com], [domain name], [domainname], [site owner name], etc.	[domainname.com site:domainname.com] Pick Search Tools, then Any time, then Past 24 hours to select one day	- moz -site:moz.com during past 24 hours
Keyword research	Evaluate the given keyword competition (sites that apply proper SEO to target the term) Find more keyword phrases	[inanchor:keyword intitle:keyword] [key * phrase]	inanchor:seo intitle:seo free * tools
SEO site auditing	Learn whether the site has canonicalization problems Find the site’s most powerful pages	[site:domain.com inurl:www] [www site:domain.com] [inurl:domain site:domain.com] [domain site:domain.com]	- site:stephanspence r.com -inurl:www www site:alchemistmedi a.com inurl:stonetemple site:stonetemple.c om alchemistmedia site:alchemistmedi a.com

	Find the site's most powerful page related to the keyword	[site:domain.com keyword]	site:scienceofseo.com seo
		[site:domain.com intitle:keyword]	site:stonetemple.com intitle:seo
		[site:domain inanchor:keyword]	site:moz.com inanchor:seo
Link building	Find authority sites offering a backlink opportunity	[site:org bookmarks/links/"favorite sites"] [site:gov bookmarks/links/"favorite sites"] [site:edu bookmarks/links/"favorite sites"]	site:org donors
	Search for relevant forums and discussion boards to participate in discussions and probably link back to your site	[inurl:forum OR inurl:forums keyword]	OR inurl:forum OR inurl:forums seo

Firefox plug-ins for quicker access to Google advanced search queries

You can use a number of plug-ins with Firefox to make accessing these advanced queries easier:

- Advanced Dork (<https://addons.mozilla.org/en-US/firefox/addon/advanced-dork/>), for quick access to (intitle:), (inurl:), (site:), and (ext:) operators for a highlighted word on a page, as shown in Figure 2-25
- SearchStatus (<http://www.quirk.biz/searchstatus/>), for quick access to a (site:) operator to explore a currently active domain, as shown in Figure 2-26

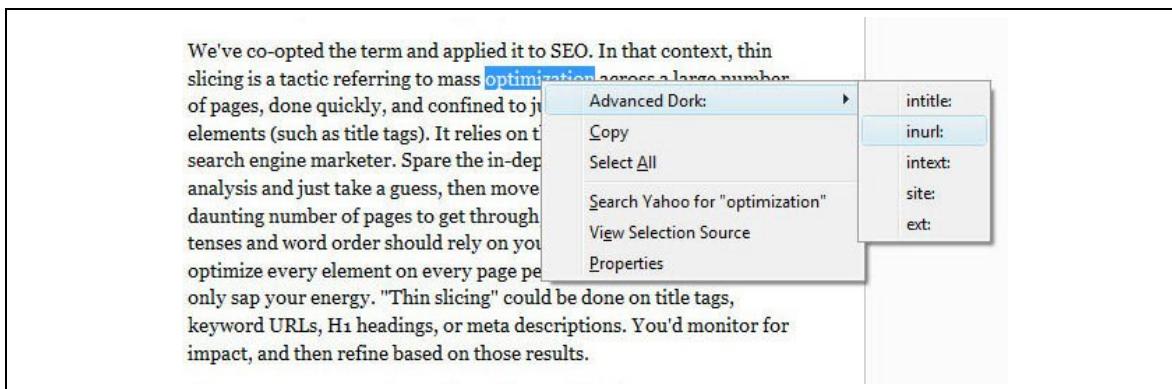


Figure 2-28. Advanced Dork plug-in for Firefox



Figure 2-29. *SearchStatus* plug-in for Firefox

Bing Advanced Search Operators

Bing also offers several unique search operators worth looking into, as shown in Table 2-5.

Table 2-5. *Bing advanced operators*

Operator	Short description	SEO application	Examples
(linkfromdomain:)—Domain outbound links restricted search	Finds all pages the given domain links out to	Find most relevant sites your competitor links out to	linkfromdomain:moz.com seo
(contains:)—File type links restricted search, filetype: is also supported	Narrows search results to linking to a document of the specified file type	Find pages linking to a specific document type containing relevant information	contains:wma seo
(ip:)—IP restricted search	Shows sites sharing one IP address	ip:xxx.xxx.xxx.xxx	ip:207.182.138.245
(inbody:)—Body keyword restricted search	Restricts the results to documents	Find pages containing the most relevant/best optimized body text	inbody:seo (equivalent to Google's intext:)

	containing query word(s) in the body text of a page	
(location:)/(loc:)— Location specific search	Narrows search results to a specified location (multiple location options can be found under Bing's advanced search)	Find geospecific seo loc:AU documents using your keyword
(feed:)—Feed keyword restricted search	Narrows search results to terms contained in RSS feeds	Find relevant feeds feed:seo
(hasfeed:)—Feed keyword restricted search	Narrows search results to pages linking to feeds that contain the specified keywords	Find pages linking to hasfeed:seo site:cnn.com

More Advanced Search Operator Techniques

You can also use more advanced SEO techniques to extract more information.

Keyword difficulty

When building a web page, it can be useful to know how competitive the keyword is that you are going after, yet this can be difficult to obtain. However, there are steps you can take to get some idea as to how difficult it is to rank for a keyword. For example, the `intitle:` operator shows pages that are more focused on your search term than the pages returned without that operator (e.g., `intitle:"dress boots"`).

You can use different ratios to give you a sense of how competitive a keyword market is (higher results mean that it is more competitive). For example:

`dress boots` (108,000,000) versus “`dress boots`” (2,020,000) versus `intitle:"dress boots"` (375,000)

Ratio: $108,000/375 = 290:1$

Exact phrase ratio: $2,020/37 = 5.4:1$

Another significant parameter you can look at is the `inanchor:` operator; for example, `inanchor:"dress boots"`. You can use this operator in the preceding equation instead of the `intitle:` operator.

Using number ranges

The number range operator can help restrict the results set to a set of model numbers, product numbers, price ranges, and so forth. For example:

```
site:stevespanglerscience.com "product/1700..1750"
```

Unfortunately, the number range combined with `inurl:` is not supported. So, the product number must be on the page. The number range operator is also great for copyright year searches (to find abandoned sites to acquire). Combine it with the `intext:` operator to improve the signal-to-noise ratio; for example, `intext:"copyright 1993..2011" -2014 blog`.

Advanced doc type search

The `filetype:` operator is useful for looking for needles in haystacks. Here are a couple of examples:

```
confidential business plan -template filetype:doc  
forrester research grapevine filetype:pdf
```

Determine listing age

You can label results with dates that give a quick sense of how old (and thus trusted) each listing is; for example, by appending the `&as_qdr=m199` parameter to the end of a Google SERP URL, you can restrict results to those within the past 199 months.

Uncover subscriber-only or deleted content

You can sometimes get to subscriber-only or deleted content from the Cached link in the listing in the SERPs (found under the down arrow after the URL in the search listing) or by using the `cache:` operator. Don't want to leave a footprint? Add `&strip=1` to the end of the Google cached URL. Images on the page won't load.

If no Cached link is available, use Google Translate to take your English document and translate it from Spanish to English (this will reveal the content even though no Cached link is available):

```
http://translate.google.com/translate?prev=&hl=en&u=URL-GOES-  
HERE&sl=es&tl=en
```

Identify neighborhoods

The `related:` operator will look at the sites linking (the "Linking Sites") to the specified site, and then see which other sites are commonly linked to by the Linking Sites. These are commonly referred to as *neighborhoods*, as there is clearly a strong relationship between sites that share similar link graphs.

Find Creative Commons (CC) licensed content

Use the `as_rights` parameter in the URL to find Creative Commons licensed content. Here are some example scenarios to find CC-licensed material on the Web:

Permit commercial use

```
http://google.com/search?as_rights=(cc_publicdomain|cc_attribute|cc_  
sharealike|cc_nonderived).-(cc_noncommercial)&q=KEYWORDS
```

Permit derivative works

```
http://google.com/search?as_rights=(cc_publicdomain|cc_attribute|cc_  
sharealike|cc_noncommercial).-(cc_nonderived)&q=KEYWORDS
```

Permit commercial and derivative use

[http://google.com/search?as_rights=\(cc_publicdomain|cc_attribute|cc_sharealike\).-\(cc_noncommercial|cc_nonderived\)&q=KEYWORDS](http://google.com/search?as_rights=(cc_publicdomain|cc_attribute|cc_sharealike).-(cc_noncommercial|cc_nonderived)&q=KEYWORDS)

Make sure you replace *KEYWORDS* with the keywords that will help you find content that is of relevance to your site. The value of this to SEO is an indirect one. Creative Commons content can potentially be a good source of content for a website. An easier way if you don't need this same freed in your Creative Commons searches is to use Google's Advanced Search page, where you can specify your Creative Commons license type.

Vertical Search Engines

Vertical search is the term people sometimes use for specialty or niche search engines that focus on a limited data set. Examples of vertical search solutions provided by the major search engines are image, video, news, and blog searches. These may be standard offerings from these vendors, but they are distinct from the engines' general web search functions.

Vertical search engines sometimes come in the form of specialty web sites, such as travel sites (such as TripAdvisor, <http://www.tripadvisor.com>), and local business listing sites (such as YellowPages.com, <http://www.yellowpages.com>). Any site, such as these, that focus on vertically oriented niche markets could be considered a vertical search engine.

Vertical search results can provide significant opportunities for the SEO practitioner. High placement in these vertical search results can equate to high placement in the web search results, often above the traditional 10 blue links presented by the search engines.

Vertical Search from the Major Search Engines

The big three search engines offer a wide variety of vertical search products. Here is a partial list:

Google

Google Maps, Google Images, Google Shopping, Google Blog Search, Google Video, Google News, Google Custom Search Engine, Google Book Search, etc.

Yahoo!

Yahoo! News, Yahoo! Local, Yahoo! Images, Yahoo! Video, Yahoo! Shopping, Yahoo! Autos, etc.

Bing

Bing Images, Bing Videos, Bing News, Bing Maps, etc.

Image search

All three search engines offer image search capability. Basically, image search engines limit the data that they crawl, search, and return in results to images. This means files that are in GIF, TIF, JPG, PNG, and other similar formats. Figure 2-27 shows the image search engine from Bing.

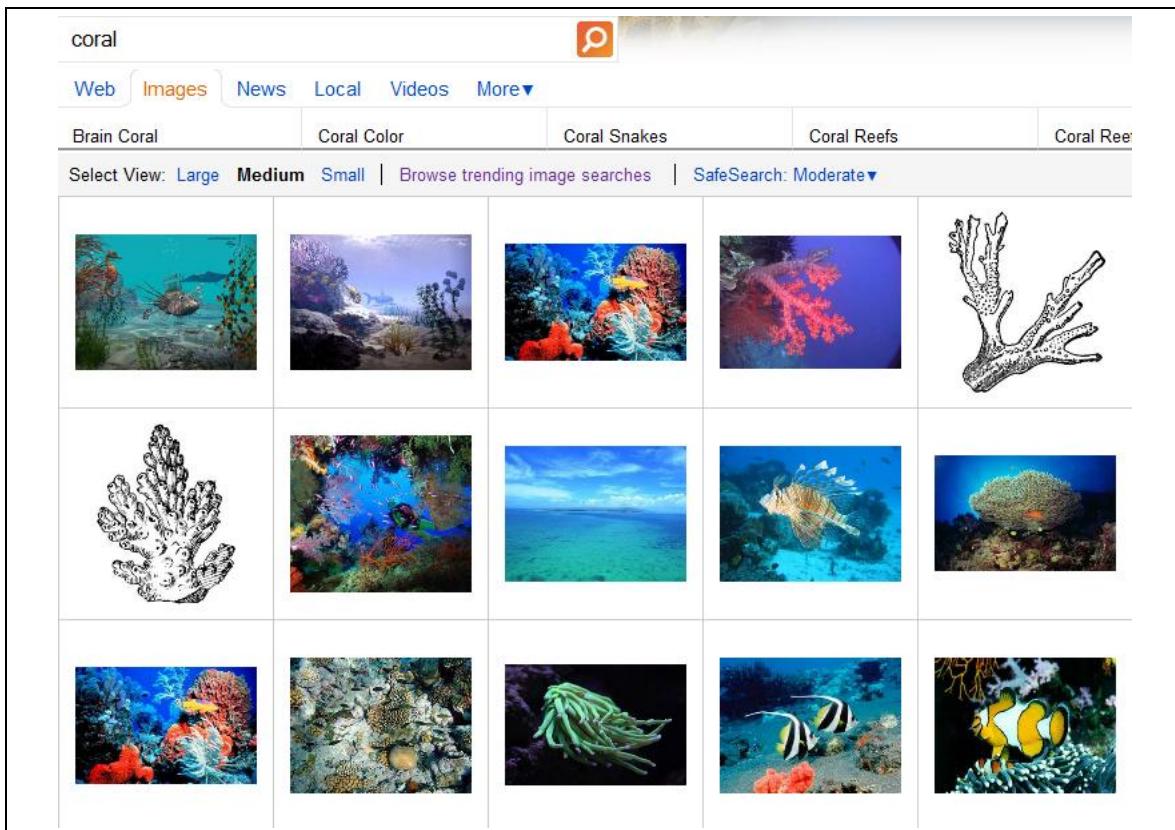


Figure 2-30. Image search results from Bing

Image search engines get a surprisingly large number of searches performed on them. Unfortunately, market data on these volumes is not often published, but according to comScore, more than 1 billion image searches were performed on Google Image Search (<http://images.google.com>) per month. However, it is likely that at least that many images related search queries occurred within Google web search during that same timeframe. However, since an image is a binary file, it cannot be readily interpreted by a search engine crawler.

Search engines have had to historically rely on text surrounding the image, the `alt` attribute within the `img` tag, and the image filename. However, Google now offers a search by image feature (<http://www.google.com/insidesearch/searchbyimage.html>) where you can drag an image file into the Google Image Search box and it will attempt to identify the subject matter of the image and show relevant results. Optimizing for image search is its own science, and we will discuss it in more detail in “Optimizing for Image Search” in Chapter 10.

Video search

As with image search, video search engines focus on searching specific types of files on the Web, in this case video files, such as MPEG, AVI, and others. Figure 2-31 shows a quick peek at video search results from YouTube.

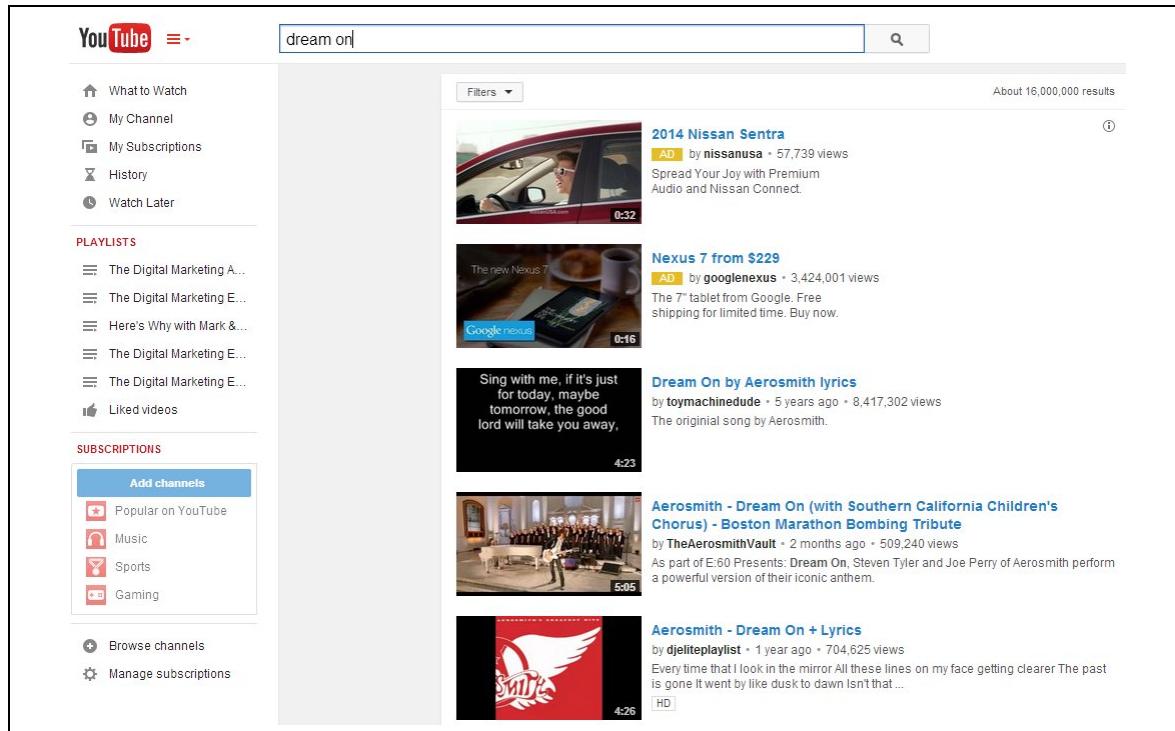


Figure 2-31. Video search results from YouTube

A very large number of searches are also performed in video search engines. YouTube (<http://www.youtube.com>) is the dominant video search engine, with over 3.8 billion searches performed in June 2011, representing more than 14% of all search queries performed on the web (need updated numbers). This makes YouTube the 3rd largest search engine on the web (Bing is larger when you consider the cumulative search volume of Bing + Yahoo!). As with image search, many video searches are also performed directly within Google web search.

There is significant traffic to be gained by optimizing for video search engines and participating in them. Once again, these are binary files and the search engine cannot easily tell what is inside them.

This means optimization is constrained to data in the header of the video and on the surrounding web page. We will discuss video search optimization in more detail in “Others: Mobile, Video/Multimedia Search” in Chapter 10.

However, each search engine is investing in technology to analyze images and videos to extract as much information as possible. For example, the search engines are experimenting with OCR technology to look for text within images, and transcription and other advanced technologies are being used to analyze video content. Flesh-tone analysis is also in use to detect porn or recognize facial features. The application of these technologies is in its infancy, and is likely to evolve rapidly over time.

News search

News search is also unique. News search results operate on a different time schedule. News search results have to be very, very timely. Few people want to read the baseball scores from a week ago when several other games have been played since then.

News search engines must be able to retrieve information in real time and provide near instantaneous responses. Modern consumers tend to want their news information now. Figure 2-32 is a quick look at the results from a visit to Yahoo! News.

The screenshot shows a Yahoo News search results page. The search bar at the top contains the query "world cup". Below the search bar, there are navigation links for "Web", "News", and "More". A sidebar on the left provides filtering options: "Filter by Source" (All sources, Yahoo! News, Yahoo! Sports, Yahoo! Finance, Los Angeles Times, More), "Filter by Time" (Any time, Past hour, Past day, Past week), and "Sort by" (Relevance, Time). The main content area displays news articles and ads. One news article is titled "World Cup: England Lose 1-2 to Uruguay" and discusses a controversial goal by Luis Suarez. Another article is titled "World Cup day 12: Suarez bit someone – did he expect applause?". To the right of the news, there are several ads for the World Cup, including one for McDonald's and one for Soccer Tickets on StubHub. A sidebar on the right lists "See more ads for: world cup, world cup 2014 schedule, fifa world cup 2014 schedule, world cup betting, fifa world cup, fifa world cup 2014".

Figure 2-32. News search results from Yahoo!

As with the other major verticals, there is a lot of search volume here as well. To have a chance of receiving this volume, you will need to become a news source. This means timely, topical news stories generated on a regular basis. There are other requirements as well, and we will discuss them further in “Optimizing for News, Blog, and Feed Search” in Chapter 10.

Local search/maps

Next up in our hit parade of major search verticals is local search (a.k.a. map search). Local search results are now heavily integrated into the traditional web search results, so a presence in local search can have a large impact on organizations that have one or more brick and mortar locations. Local search engines search through databases of locally oriented information, such as the name, phone number, and location of local businesses around the world, or just provide a service, such as offering directions from one location to another. Figure 2-33 shows Google Maps local search results.

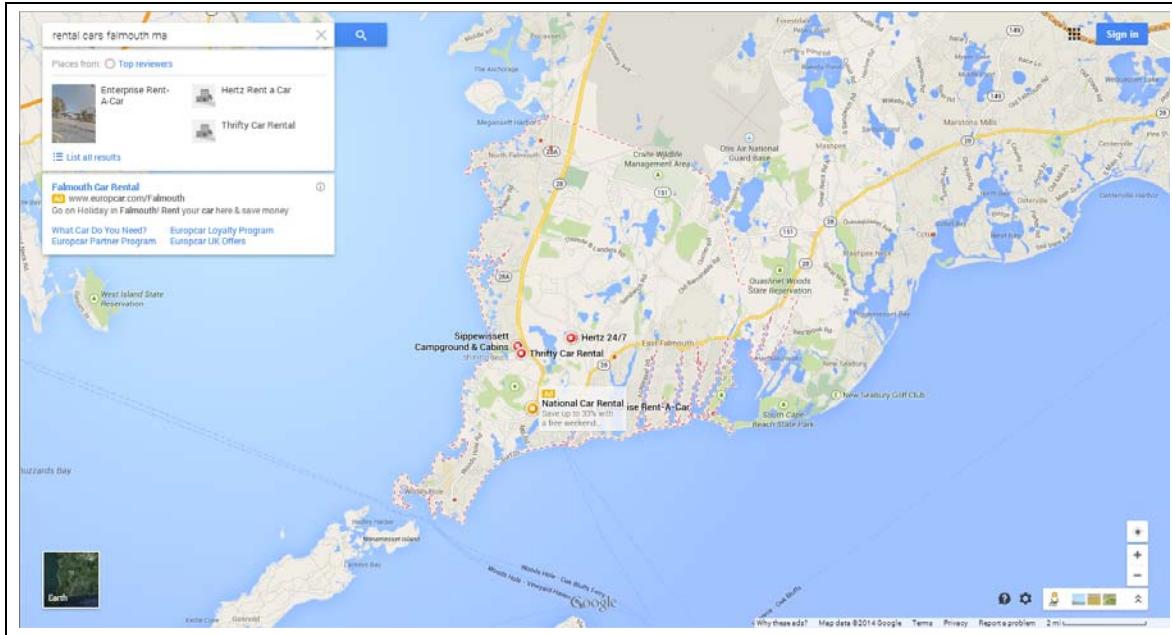


Figure 2-33. Local search results from Google Maps

The integration of local search results into regular web search results has dramatically increased the potential traffic that can be obtained through local search. We will cover local search optimization in detail in “Optimizing for Local Search” in Chapter 10.

Blog search

Google has implemented a search engine focused just on blog search called Google Blog Search (misnamed because it is an RSS feed engine and not a blog engine). This search engine will respond to queries, but only search blogs (more accurately, feeds) to determine the results. Figure 2-34 is an example search result for the search phrase *barack obama*.

The screenshot shows a Google search results page for the query "barack obama". The search bar at the top contains the text "barack obama". Below the search bar are navigation links for "Web", "Images", "News", "Videos", "More", and "Search tools". The search results section displays five entries, each with a title, a snippet of text, and a link to the source. The results are as follows:

- Blog homepages for barack obama**
Barack Obama
www.barackobama.com/ ▾
2 days ago - OFA works to ensure the voices of ordinary Americans are heard in Washington, while training the next generation of ...
- President Barack Obama | The White House**
www.whitehouse.gov ▾ The Administration ▾
Jan 20, 2009 - Barack H. Obama is the 44th President of the United States. His story is the American story — values from the heartland, ...
- Flickr: Barack Obama's Photostream**
https://flickr.com/photos/barackobamadotcom/ ▾
Jan 15, 2006 - Photos from the president organized by set and tags. Includes archives, favorites and profile.
- Barack Obama and the Three Hundred - The New Yorker**
www.newyorker.com/.../barack-obama-and-the-three-hundred-iraq-adviser... ▾
2 days ago - Three hundred American "advisers" are going to Iraq, President Obama said on Thursday. He could hardly have looked less enthusiastic.
More results from www.newyorker.com
- Kanye West Disses Barack Obama: 'You Can't Effect Change'**
hollywoodlife.com/2014/06/.../kanye-west-disses-barack-obama-intervie... ▾
by meganrosshi
2 days ago - The rapper's latest diss was pointed directly at President Barack Obama and his time in the White House. Kanye and the Prez have a shade-throwing history. One of their earliest beefs happened when President Obama called ...
More results from hollywoodlife.com

Figure 2-34. Results from Google Blog Search

We explore the subject of optimizing for Google Blog Search in “Optimizing for News, Blog, and Feed Search” in Chapter 10.

Book search

The major search engines also offer a number of specialized offerings. One highly vertical search engine is Google Book Search, which specifically searches only content found within books, as shown in Figure 2-35.

Google whales

Web Images Videos News Books More Search tools

About 2,830,000 results (0.30 seconds)

Whales
 books.google.com/books?id=FWVHAAAAYAAJ
Robert Hamilton - 1843 - Read - More editions

Whales
 books.google.com/books?isbn=1612111653
Ann Herriges - 2012 - Preview - More editions
Whales are giant sea creatures capable of performing amazing acrobatics! They groan and sing to talk to each other. This book lets children look closely at the life and behavior of whales.

Whales
 books.google.com/books?isbn=0307204049
Golden Books - 1999 - Snippet view - More editions
Introduces the food, growth, habitats, shapes, and sizes of whales such as the blue whale, the humpback whale, and the sperm whale

Whales
 books.google.com/books?isbn=0761413898
Daniel A. Greenberg - 2003 - Preview - More editions
Describes in detail the physical characteristics, behavior, and migration and life cycle of various kinds of whales, among the largest creatures ever known to have lived on Earth, and discusses the history of human interaction with these ...

Figure 2-35. Google Book Search

Product Search

Bing also has some unique vertical search features. One of the more interesting ones is their product search solution. Instead of having a separate shopping search engine, Bing has integrated the product results into the main body of their search results, as shown in the right side of Figure 2-36.

Figure 2-36. Bing product search

Universal Search/Blended Search

Google made a big splash in 2007 when it announced Universal Search. This was the notion of integrating images, videos, and results from other vertical search properties directly into the main web search results. Prior to this announcement, all the search engines showed their search results in separate vertical search engines. We have already shown an example of this in Figure 2-36 which shows Bing's way of integrating product search features directly into the main search results.

After Google's announcement both Bing and Yahoo quickly followed with their own implementations of blended search. Each type of result you see on a search results page offers different opportunities for obtaining traffic from search engines. People now refer to this general concept as Blended Search (since Universal Search is specifically associated with Google).

Figure 2-37 shows an example of blended search results from a Google search.

Google Search More

Web News Maps Images Videos More Search tools

About 79,100,000 results (0.57 seconds)

News for Iraq

 **Iran's Participation in Meeting to Aid Iraq is 'Not Appropriate' ...**
New York Times - 12 hours ago
France is the host of the Monday meeting, which is to coordinate aid to the new Iraqi government for its fight against the Islamic State in Iraq and

 **Strikes in Iraq, Syria Expected Under New War Plan**
ABC News - 4 hours ago

More news for Iraq

Iraq - Wikipedia, the free encyclopedia
en.wikipedia.org/wiki/Iraq - Wikipedia
It has a northern section of Palestine measuring 58 km (36 mi) on the northern Persian Gulf and its territory encompasses the Mesopotamian Alluvial Plain, the Iraq insurgency (2011–present) - Islamic State of Iraq - Iraq War - Nouri al-Maliki

Iraq News - Breaking World Iraq News - The New York Times
topics.nytimes.com/topics/news/international/Iraq - The New York Times -
World news about Iraq. Breaking news and archival information about its people, politics and economy from The New York Times.

Iraq | World news | The Guardian
www.theguardian.com | News | World news ▾ The Guardian -
Latest news and comment on Iraq from the Guardian.

Iraq | Reuters.com
www.reuters.com/places/iraq ▾ Reuters -
PARIS - Around 100 militants who fought with rebel groups in Syria and Iraq have returned to France, requiring "massive" resources for surveillance and other ...

Images for Iraq Report images

More images for Iraq

Iraq news, all the latest and breaking Iraq news - Telegraph
www.telegraph.co.uk ▾ Middle East ▾ The Daily Telegraph -
Beset by rising discontent at home, the unpopular French president seizes a chance to cast himself as an international statesman with Iraq visit. 12 Sep 2014 ...

Barack Obama and Islamic State: Back to Iraq | The Economist
www.economist.com/2161659-barack-obama-vows-st - The Economist -
14 hours ago - BARACK OBAMA's prime-time address of September 10th, bracing America for an open-ended campaign against Islamic State (IS) in Iraq and ...

Iraq News - Bloomberg
topics.bloomberg.com/iraq ▾ Bloomberg LP -
Breaking news about Iraq. Find the latest articles, videos, photos and blogs about Iraq.

Iraq Body Count
https://www.iraqbodycount.org/ ▾ Iraq Body Count project -
The project tries to establish an independent public database of civilian deaths in Iraq resulting directly from military actions by the USA and its allies in 2003.

In-depth articles

President Obama Talks to Thomas L. Friedman ...
The New York Times - Aug 2014
In a wide-ranging interview, the president explains American policy on Iraq, Putin and the ... President Obama said on Friday that Iraq's leaders must form an inclusive government in ... Explore: obama interview

The Iraq We Left Behind
Foreign Affairs
Nine years after U.S. troops toppled Saddam Hussein and just a few months after the last U.S. soldier left Iraq, the country has become something close to a failed state. Prime Minister Nouri al-Maliki ...

What We Left Behind - The New Yorker
The New Yorker - Apr 2014
When the last American soldiers left Iraq, at the end of 2011, the bloody civil war between the country's Sunni and Shiite sects had been stilled but not resolved. Now the sectarian ...

+ More in-depth articles

Searches related to Iraq

iraq oil iraq information
iraq currency iraq cities
history of iraq iraq religion
iraq facts iraq pronunciation

Google > 1 2 3 4 5 6 7 8 9 10 Next

● Marlborough, MA - From your Internet address - Use precise location - Learn more
Help Send feedback Privacy & Terms

Figure 2-37. Google Universal Search results

More specialized vertical search engines

Vertical search can also come from third parties. Here are some examples:

- Comparison shopping engines, such as PriceGrabber, Shopzilla, and Nextag
- Travel search engines, such as Expedia, Travelocity, and Kayak
- Real estate search engines, such as Trulia and Zillow
- Job search engines, such as Indeed, CareerBuilder, and SimplyHired
- Music search engines, such as iTunes Music Store
- B2B search engines, such as KnowledgeStorm and ThomasNet

There is an enormous array of different vertical search offerings from the major search engines, and from other companies as well. It is to be expected that this explosion of different vertical search properties will continue.

Effective search functionality on the Web is riddled with complexity and challenging problems. Being able to constrain the data types (to a specific type of file, a specific area of interest, a specific geography, or whatever) can significantly improve the quality of the results for users.

Country-Specific Search Engines

At this stage, search is truly global in its reach. Google is the dominant search engine in many countries, but not all of them. How you optimize your website depends heavily on the target market for that site, and the search engine(s) that is (are) the most important in that market.

According to comScore, Google receives 54.3% of all searches performed worldwide as of April 2014. In many countries, that market share is 80% or more.

Here is some data on countries where other search engines are major players:

China

China Internet Watch (<http://www.chinainternetwatch.com/category/search-engine/>) reported in September 2014 that Baidu had about 70% market share. This is significant since China has the largest Internet usage in the world with 618 million users in 2010 according to China Internet Network Information Center (<http://www.cnnic.net.cn/>).

Russia

According to figures announced by Yandex, the company's market share in Russia comprised about 62% of all searches in April 2014 (<http://searchengineland.com/yandex-reports-36-growth-q1-2014-now-owns-62-share-russian-search-market-189860>).

South Korea

Naver (<http://www.naver.com>) was estimated to have about 70% market share in South Korea in March 2014 (<http://www.korea-marketing.com/should-naiver-worry-about-google/>).

Czech Republic

The Startup Yard blog reported in January 2014 that Seznam (<http://seznam.cz>) was reported to have more than 60% market share in Czechoslovakia (<http://startupyard.com/meet-the-only-company-in-europe-that-is-beating-google-seznam-cz/>).

Optimizing for Specific Countries

One of the problems international businesses continuously need to address with search engines is identifying themselves as "local" in the eyes of the search engines. In other words, if a search engine user is located in France and wants to see where the wine shops are in Lyons, how does the search engine know which results to show?

Here are a few of the top factors that contribute to international ranking success:

- Owning the proper domain extension (e.g., .com.au, .uk, .fr, .de, .nl) for the country that your business is targeting
- Hosting your website in the country you are targeting (with a country-specific IP address)
- Registering with local search engines:
 - Google My Business: <https://www.google.com/business/>
 - Yahoo Small Business: <https://smallbusiness.yahoo.com/>
 - Bing Places: <https://www.bingplaces.com/>
- Having other sites from the same country link to you
- Using the native language on the site (an absolute requirement for usability)
- Help Google serve the correct language or regional URL in the search results by adding hreflang attribute (<https://support.google.com/webmasters/answer/189077?hl=en>).
- Placing your relevant local address data on major pages of the site
- Set your geographic target in Google Webmaster Tools (you can read more about this here: <https://support.google.com/webmasters/answer/62399>). Note that Google does not really need you to do this if your site is on a ccTLD such as “.de” or “.co.uk” as the preferred regional target is assumed.

All of these factors act as strong signals to the search engines regarding the country you are targeting, and will make them more likely to show your site for relevant local results.

The complexity increases when targeting multiple countries. We will discuss this in more depth in “Best Practices for Multilanguage/Country Targeting” in Chapter 6.

Conclusion

Understanding how search engines work is an important component of SEO. The search engines are constantly tuning their algorithms. For that reason, the successful SEO professional is constantly studying search engine behavior and learning how they work.

3

SEO Planning: Customizing Your Strategy

SEO, once a highly specialized, and often marginalized task relegated to the back rooms of website development teams, is now a mainstream online marketing activity. This dramatic rise can be attributed to various emerging trends, including:

- Search engines now drive a large volume of highly targeted traffic, from people intent on accomplishing research-based, purchase-oriented, and content consumption goals. Businesses can earn significant revenues by leveraging the quality and relevance of this traffic for direct sales, customer acquisition, and content and branding/awareness campaigns.
- Visibility in search engines creates an implied endorsement effect, whereby searchers associate quality, relevance, and trustworthiness with sites that rank highly for their queries.
- Dramatic growth in the interaction between offline and online marketing necessitates investment by organizations of all kinds in a successful search strategy. Consumers are increasingly turning to the internet before making purchases in verticals ranging from real estate, autos, and technology to finding local resources for health and wellness services. At this stage, businesses simply cannot afford to ignore their customers' needs as expressed through searches conducted on the major search engines.

Search engine optimization is a technical marketing function that is reliant upon both content creation and technological development, and success in organic search is greatly dependent upon various technical and analytical tools that provide access to valuable data necessary for ongoing SEO and website improvement, as well as overall user experience optimization. The intersection and interdependence between website technology and online marketing has become clear, and essential – and SEO was at the forefront of this trend. New conference events such as MarTech (produced by Third Door Media, the company behind SearchEngineLand.com and Search Marketing Expo) – are beginning to provide platforms for learning and growth within the powerful intersection of marketing and technology.

As with all marketing functions, specific goal-setting is required for success to be measured and achieved. While SEO can be viewed as a *project* (and there are certainly situations in which SEO “projects” exist) the best investment, in our opinion, is to treat it as more of a *process*—one that is iterative, ongoing, and requires steady commitment from the stakeholders of an organization.

An ongoing investment is needed to create a site with the proper content, architecture and overall foundation for SEO success. The results may not appear instantly, but a business that makes a patient and prudent commitment to SEO, invests in quality content development, constructs a crawlable site architecture, and earns legitimate linking relationships will be handsomely rewarded.

Strategic Goals SEO Practitioners Can Fulfill

Organic search is one of many viable online marketing channels as a source of highly targeted traffic. And while SEO is not a cure-all for businesses, it can fit into a company's overall business strategy in several critical ways.

Visibility (Branding)

Many consumers assume that top placement in the search engines is like a stamp of approval on a brand. Surely a company could not rank highly in the search results if it were not one of the best in its field, right?

If you are an experienced search engine user, you probably recognize that the preceding statement is not always true. However, the fact is that many consumers, and even sophisticated web searchers, interpret high search rankings as an implicit brand endorsement.

Therefore, for critical brand terms, an SEO strategy should be geared towards increasing organic search exposure for branded keywords.

In addition to targeting brand terms, you will want to rank well for non-branded search terms that are specific to your core business (a highly competitive, but highly rewarding branding effort). When searchers see you ranking highly on these types of search terms, they will associate your brand with the product or service and assume you are one of the best places to purchase from.

The list of situations where the brand can also limit the SEO strategy is quite long, and the opposite can happen too, where the nature of the brand makes a particular SEO strategy pretty compelling. Ultimately, your goal is to dovetail SEO efforts with branding as seamlessly as possible.

Website Traffic

Long gone are the days of a “build it and they will come” paradigm online. Today’s search environment is highly competitive, and you need great SEO to capture targeted, high-quality traffic to your site.

While a business that engages with many of its customers through offline channels can drive traffic by telling those customers to visit its website, SEO fills the different, more critical role of bringing new prospects to your website from an audience of people who might not otherwise have been interested in, or perhaps even aware of, your business at all.

Experienced SEO professionals understand that users search for products, services, and information using an extraordinarily wide variety of search queries and query types. Developing an SEO strategy involves performing extensive keyword research (which we will discuss in Chapter 5) to determine which search queries people actually use. For example, when interested in purchasing a fuel efficient car, a searcher might type in “hybrid efficient cars.” In this instance, the search marketing manager for a manufacturer of electric cars might be interested in search exposure for this traffic, even though technically, “electric cars” was not in the original search query.

Some users may not even know that a company specializing in a specific type of product exists until they perform that search. Or, if they have at one time learned about such a manufacturer, they might not remember enough about it to seek out the manufacturer’s website directly.

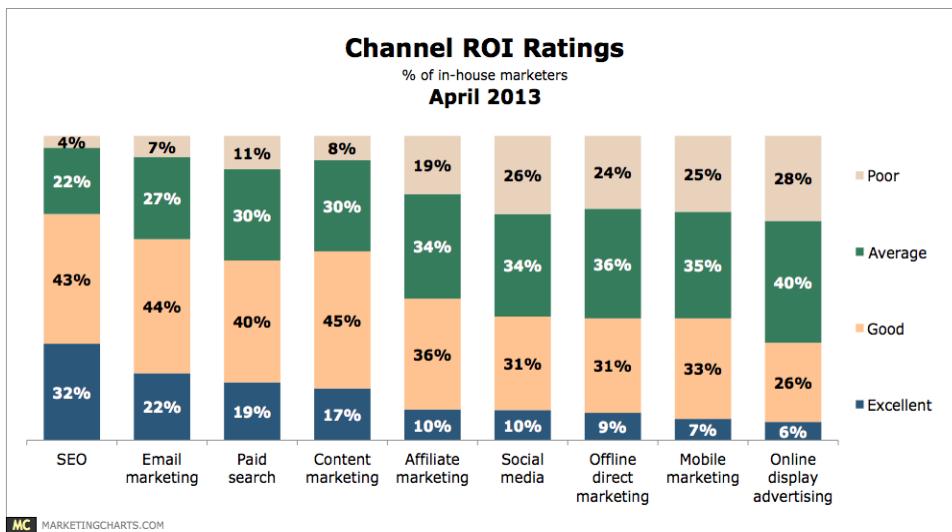
Capturing that traffic could provide the manufacturer with incremental sales of its electric vehicles that it probably would not have gotten otherwise. Knowing these factors, the SEO process involves developing a site architecture strategy (see Chapter 6) and a content development/editorial strategy (we cover this in Chapter 7) to help the site’s pages achieve competitive search engine exposure for a broad range of potentially relevant terms.

High ROI

Increasing online visibility and driving targeted traffic are the first steps to success with SEO – the next step is measuring the performance of that traffic to determine whether the overall website and business

objectives are being achieved. This is an ongoing, iterative process that will always accompany all marketing efforts, and SEO is no exception. For most organizations, goals are defined as generating sales, leads, or advertising revenue. For others, it may mean the promotion of a particular message, the consumption, sharing, or perhaps the creation of specific types of content, or the simple signup for a newsletter. An important component of SEO is to deliver not just traffic, but *targeted* traffic that has the possibility of converting into these actions. Whether you are selling products and services, advertising for branding value, or trying to promote specific content to the world, a well-designed SEO strategy can result in a very high return on investment when contrasted with other methods of marketing.

SEO generally brings a higher ROI when compared to TV, print, and radio advertising campaigns; and while traditional media is not in danger of being replaced by SEO, organic search can provide some high-margin returns that complement and enhance the use of offline media – especially when done properly over the long haul.



[[figs/print/0301.png]]

Figure 3-1. SEO, a high-ROI activity

In addition, a growing number of businesses operate purely online. Some examples of these are LinkedIn (<http://www.linkedin.com>), Zappos (<http://www.zappos.com>), Amazon (<http://www.amazon.com>), and eBay (<http://www.ebay.com>).

Every SEO Strategy Should be Customized

There is no such thing as a cookie-cutter SEO strategy (and if someone is trying to sell you one, be sure to dig deeper and get second opinions!). There is a difference between best practices, which are universal, and a business and website-specific SEO strategy. The ever-changing, dynamic nature of the search marketing industry requires constant diligence, and SEO professionals must maintain a research process for analyzing how the search landscape is changing.

The following factors must be taken into account when developing your SEO strategy:

- What the organization is trying to promote (service, product, content)
- Who the target market is (can be as simple as “women” or as detailed as personas)
- Brand (includes copy & messaging)

- Website structure (includes site architecture, navigational elements, and file/URL naming conventions)
- Current site content assets (includes images, videos, PDF files, white papers, case studies, articles)
- Ease with which the content and site structure can be modified (involves the CMS & web development teams)
- Editorial resources and calendar for content development (what content is developed, by whom, and on what timeline)
- Competitive landscape

Learning about the specific market a business is in is an obvious first step towards developing industry familiarity, trends, and vernacular - but it often makes sense for two businesses offering the same products in the market to use different SEO strategies.

For example, if Competitor 1 published its website four years ago, and Competitor 2 is just launching their site now, Competitor 2 may gain the most traction by focusing on specific vertical areas where Competitor 1's offering, or SEO implementation, is weak.

It is important to not underestimate the importance of a well thought out SEO strategy. Skipping over this process, or not treating it seriously, can hamper your prospects for growth.

Understanding Search Engine Traffic and Visitor Intent

As we discussed in “The Mission of Search Engines” in Chapter 1, searchers enter many different types of queries. These can generally be classified into three major categories:

Navigational query

This is a query with the intent to arrive at a specific website or page (e.g., the person types in your company domain name, www.companyname.com, or simply types in the word *facebook*).

Informational query

This is a search performed to receive an answer to a broad or direct question, or to research and explore information around a specific topic with no specific source in mind (e.g., *yoga poses*).

Transactional query

A person who types in *digital camera* may be looking to buy one now, but it is equally possible that she is researching digital cameras to learn about how they are different from film cameras. This is an example of an initial transactional query, which can evolve in stages. For example, below are some other types of transactional queries that occur at a later stage in the buying cycle:

- The user types in *best online digital camera store*. Although there is no information in the query about which one she wants to buy, the intent is clearer that the searcher is seeking a store, not simply information about types of digital cameras.
- The searcher types in *olympus OMD lowest price*. The chances are very high that this user is looking to buy that particular camera.

Part of an SEO strategy is to understand how the various types of searches relate to the content and architecture of your website.

Developing an SEO Plan Prior to Site Development

It is widely understood in the SEO industry that SEO should be built in, as early as possible, to the entire site development strategy - from choosing a content management system (CMS) and planning site architecture to creating, optimizing, and publishing site content. As you will see in Chapter 6, SEO is deeply intertwined with these areas.

SEO plans have many moving parts, and decisions made in consideration of SEO can and usually do have a significant impact on other departments such as web development, content development and editorial, other marketing groups (direct, offline, etc), and sales. Integrating these moving parts and aligning them with each other is essential to developing an SEO-friendly website and establishing a strong foundation for enduring organic visibility.

Business Factors That Impact Your SEO Strategy

Here are some examples of business issues that can impact SEO:

Revenue & business models

The effective SEO strategy takes into account the purpose of the site – whether it is to sell products, sell advertising, obtain leads, or gain membership signups. We will discuss this more in the later sections of this chapter.

Target customers

Who are you trying to reach? This could be an age group, a gender group, or as specific as people looking to buy a house within a specific neighborhood of San Francisco, CA.

Competitors

The competitive landscape is another big factor in your overall strategy. Competition may be strongly entrenched in one portion of the market online, and it may make sense to focus on a different segment. Or you may be the current leader in your market, and want to protect this position while continuing to build your customer base.

Branding goals

There may be search terms for which it is critical that you have top search exposure, for branding reasons.

Content development

An important part of SEO and general online success is the creation and optimization of high quality content for your users. For most businesses, your capacity to create quality content on an ongoing basis can significantly improve your SEO efforts, both from a content-availability perspective (more content in the search engines) and from a user engagement and link development perspective (great content breeds great links and social sharing, both of which can positively influence SEO).

How people search for products like yours

Understanding what customers do when they are searching for products or services like yours is one of the most basic functions of SEO, which we will discuss it in detail in Chapter 5. This involves mapping the actual search queries your target customers use when they go to a search engine.

Understanding Your Audience and Finding Your Niche

Determining your target audience should be something incorporated into the initial site design, and should inform your SEO and content development strategies. As you will see in this section, many factors enter into this including the competition, the particular strengths or weaknesses of your own company, and more.

Mapping Your Products and Services

Successful SEO requires a thorough understanding of your business. What products, services, and types of information and resources does you have to offer your site visitors and potential customers?

As we outlined in the preceding section, a critical SEO activity is understanding who is searching for what you are trying to promote, which requires thoroughly understanding all aspects of your offering. You will need to understand the broad and specific market categories that your products fall into, as each of these

categories might relate to sections of your website that you will likely need to create. By having content areas of the site for those categories, you create the opportunity to obtain organic search traffic related to those categories.

You also should consider business development and the company's expansion strategy at the outset of the SEO planning process. Consider Amazon, which began as a bookseller but has evolved into a general purpose e-tailer; sites that go through these types of changes may need to be substantially restructured, and such restructurings can be a source of major SEO headaches. Anticipating and planning for these types of changes in advance provides the opportunity to recommend appropriate architectural approaches to developing and optimizing the site.

Yes, Content Is King

And so it has gone since the late 1990's – the debate about whether content is king in SEO. Not only is content king for SEO purposes, but it is one of the most important elements of any online effort. Development of high quality, engaging content for your users increases the available real estate for organic search queries, provides your users reason to enjoy your site and compels them to become customers, and invites promotion and exposure through users sharing your content online, on social media platforms and via direct links from other websites. Determining your available content assets is the first step towards leveraging them within the context of SEO. You may have a deep library of "how to" content, great testimonial or product demonstration videos, a unique photo gallery, or an awesome tool that people are interested in using. All of these content types can be invaluable in building a world-class website that does well in the both search engines and the overall web ecosystem.

The content you have available to you will also affect your keyword research and your site architecture, as your content is the major source of information that search engines use to determine what your site is about. As we discussed in "Algorithm-Based Ranking Systems: Crawling, Indexing, and Ranking", you need relevant content to even be "in the game" in search (i.e., if someone searches for *hybrid efficient cars* and you don't have any content related to hybrid efficient cars, chances are good that you won't rank for that search query).

As we will discuss in Chapter 7, on-site content also affects content marketing and link development. Content marketing can be very similar to PR in that the success of your efforts is integrally related to what you are promoting (i.e., what type of content, and of what quality, are you hoping others will link to?).

Consider Site A, a site that has built a substantial, well-researched and well-written set of articles on a specific topic. However, 20 other sites out there have equally substantial sets of articles on the same topic, and many of these other sites have been in the major search engine indexes for much longer than Site A.

Site A has a content quality concern – namely, why would someone link to their articles over the articles from any of the other 20 websites? There is nothing new there. Chances are that Site A will succeed in getting some links to its articles; however, it will likely never be able to establish itself as a leader if it has nothing new, or unique, to offer.

To establish itself as a leader, Site A must create fresh, unique, engaging content that differentiates the site from its competitors. Perhaps it can offer a solution to a problem that no one else has been able to solve before, or perhaps it focuses on a specific vertical niche and establishes itself as a leader in that niche—for example, by being the first to release a high-quality video series on the topic it covers.

One of the most important decisions Site A's leadership needs to make is where and how they are going to establish themselves as an authoritative resource in their industry. If they plan to make their website a major player in capturing market-related search engine traffic, this is not an optional step.

When looking at content plans and hopefully establishing an editorial calendar, it is critical to consider not only what you already have, but also what you could develop. This, of course, relates to budget. A publisher with no budget to spend on content development has few choices that she can make in her SEO strategy, whereas another publisher who has a team of in-house content developers has a myriad of options.

Segmenting Your Site's Audience

Let's not forget the audience! It is important for the SEO practitioner to understand the target audience. For example, Site A may be a website that sells sneakers. As a result, the site's developers go out and implement a brilliant campaign to rank for the "sneaker" terms they consider relevant. Being young and energetic, they focus on the way their peers search for sneakers—but what if the target audience for the sneakers Site A sells are age 50 or older? This segmentation goes to the very core of deciding what pages to create on your site, and what content to put on them.

The target audience for Site A (the over-50 crowd) may use different search terms than the younger generation to search for gadgets, which means Site A may well be bringing in search traffic from people who are not interested in its products, and not bringing in traffic from those who might be! For example, Skechers, a California-based shoe company, made the business decision to pivot from being a youth brand to targeting an older demographic with their shoes – so an SEO strategy for selling Skechers sneakers at this stage, then, would be very different from an SEO strategy for the more youth-oriented Nike sneakers.

Similar things can happen with gender, as women and men may not search for their shoes the same way. Figure 3-2 lists the top shoe-related search terms from Wordtracker.com (<http://www.wordtracker.com>) for searches related to *womens shoes* and *mens shoes*.

Keyword	Volume
womens shoes	33,419
womens shoes online	4,465
womens boots shoes	2,912
womens high heel shoes	2,887
cheap womens shoes	2,740
womens boot shoes	2,680
salomon womens shoes sale	2,093
womens fashion shoes	2,084
clarks womens shoes	1,697
ebay womens shoes	1,679

Keyword	Volume
mens shoes	14,641
mens tennis shoes	2,035
mens dress shoes	1,440
mens shoes stores	1,387
buy mens shoes online	1,351
mens fashion shoes	1,217
mens clothing shoes	1,176
mens walking shoes	1,111
mens shoes online	1,049
mens sports shoes	1,007

[[figs/print/0302a.png]]

[[figs/print/0302b.png]]

Figure 3-2 – Keyword Search Term Variation Examples by Gender

As you can see in the figure above, search terms used can vary significantly by gender.

Another major element to consider might be location. Searchers in Austin, TX, may want a different version of your product than searchers in Chicago, IL. For that matter, because they want different products, and they are in different regions, they may be looking for something specific to their region, and they may use different vernacular – all factors that influence a user’s choice of search terms. Understanding these variables and identifying appropriate traffic targets requires extensive keyword research—yet another critical aspect of the SEO process.

Context: Market Competitiveness

A critical factor to understand before you launch your SEO efforts is the nature of the market in which you are competing. This tells you how competitive the environment is in general, and augmented with additional research, you can use this information to tell how competitive the SEO environment is. In some markets, natural search is intensively competitive.

Market competitiveness does not mean you should give up on competing, especially if it is already the focus of your business; however, you might choose to focus your SEO efforts on less competitive terms that can still bring you many qualified leads.

To get a very quick read on keyword competitiveness, use the Google AdWords Keyword Planner (<https://adwords.google.com/KeywordPlanner>) to see what your cost per click could be if you bid on your target phrase in a PPC campaign. Higher cost per click prices in Google AdWords often mean that the terms are more competitive in organic search as well.

Your SEO strategy can also be influenced by your competitors’ strategies, so understanding what they are doing is a critical part of the process for both SEO and business intelligence objectives. There are several scenarios you might encounter:

- The competitor discovers a unique, highly converting set of keywords.
- The competitor discovers a targeted, high-value link from an authoritative source.
- The competitor saturates a market segment, justifying your focus elsewhere.

- Weaknesses appear in the competitor's strategy, which provide opportunities for exploitation.

Understanding the strengths and weaknesses of your competition from an SEO perspective is a significant part of formulating your own SEO strategy.

SEO tools such as SEMRush (<http://semrush.com>) and SearchMetrics (<http://searchmetrics.com>) can provide insight into your competitors' SEO performance.

SEO for Raw Traffic

Optimizing for search engines and creating thematically-targeted content helps a site rank for key search queries, which typically leads to direct traffic, social sharing, and referral traffic from links as more and more people find, use, and enjoy the content you've produced. Thousands of sites on the Web leverage this traffic to serve advertising, directly monetizing the traffic sent from the engines. From banner ads to contextual services such as Google's AdSense, to affiliate and social media marketing - internet advertising spending has become a massive industry. In October 2014, the Internet Advertising Bureau (IAB) measured internet advertising revenues for Q3 2014 at \$12.4 billion, a 17% increase over Q4 2013 (http://www.iab.net/about_the_iab/recent_press_releases/press_release_archive/press_release/pr-121814).

Here are some factors to think about when considering SEO for raw traffic:

When to employ SEO for raw traffic:

Use it when you can monetize traffic without actions or financial transactions taking place on your site (usually through advertising).

Keyword targeting:

Keyword targeting in this scenario can be very broad. The goal here isn't typically to select specific keywords, but rather to create high-quality content that naturally targets interesting, searched-for terms. Instead of singular optimization on specific terms, the focus is on accessibility and best practices throughout the site to earn traffic through both high-volume and long-tail queries (for more on what the *long tail* is, see Chapter 5). Concentrate efforts on great content, and use keyword-based optimization as a subsequent application to confirm the titles, headlines, filenames, metadata, and other elements of the content you create.

Page and content creation/optimization

A shallow, highly crawlable link structure is critical to getting all of your content indexed—follow good information architecture practices (see “Creating an Optimal Information Architecture (IA)” in Chapter 6) and use intelligent, detailed category and subcategory structures to get the most benefit out of your work. You'll also need to employ good on-page optimization in title tags, headlines, internal links, etc. and make your articles easy to share and optimized for viral spreading (see “Root Domains, Subdomains, and Microsites” and “Optimization of Domain Names/URLs” in Chapter 6 for more on this).

SEO for Ecommerce Sales

One of the most direct monetization strategies for SEO is driving relevant traffic to an ecommerce shop to boost sales. Search traffic is among the highest quality traffic on the Web, primarily because a search user has expressed a specific interest through his or her query, and when this matches a service, product, or brand a website carries, conversion rates are often extremely high. Below are some factors to think about when considering SEO for ecommerce sales:

When to employ SEO for ecommerce sales

Use it when you have products/services that are directly for sale on your website.

Keyword targeting

Pay per click (PPC) advertising is an excellent way to test the efficacy and potential ROI of keyword targets. Find those that have reasonable traffic and convert well, and pursue them further. You'll often find that the more specific the query is—brand-inclusive, product-inclusive, and so on—the more likely the visitors are to make the purchase. The best use of this tactic is for generic terms that you will find harder to win on than brand/company named terms, so you can decide if they are worth the effort.

Quality content creation & optimization

Your site will need to provide interesting, unique, and accessible content for both users and search engines in order to begin gaining traction in organic search. The links and social sharing that come from such content are highly influential in increasing overall site traffic, as well as improving organic search performance. Creating link-worthy, deeply engaging content should be the primary focus of any website seeking search and referral traffic, and should be incorporated into both the content development and SEO strategies. Manual link building is always an option, but scalable strategies that leverage a community or customers can be equally, or even more, valuable. Remember: content that keeps users on the page, instead of clicking away rapidly is a signal of quality – and also serves to increase the value of traffic to the page (improving page monetization through advertising).

SEO for Mindshare & Branding

An equally powerful application of SEO is to use it for branding purposes. Bloggers, social media platforms, community websites, content producers, news outlets, and dozens of other web publishing outlets have found tremendous value in appearing atop the SERPs and using the resulting exposure to bolster their brand recognition and authority.

The process is fairly simple, much like the goal in traditional advertising of ad repetition to enter the buyer's consideration set. Online marketers have observed that being at the top of the search rankings around a particular subject has a positive impact on traffic, consideration, and perceived authority. A 2012 study by Conductor, Inc. highlights the significant branding value of a website appearing in organic search results: <http://www.conductor.com/blog/2012/07/search-reimagined-the-branding-value-of-page-one-study/>

Here are some factors to think about when considering SEO for mindshare and branding:

When to employ SEO for mindshare/branding

Use it when branding or communicating a message is your goal.

Keyword targeting

A keyword focus is less critical here—you'll likely have a few broad terms that receive the high traffic you want, but the long tail may be far more achievable and the better target. Focus on keywords that are going to bring you visitors who are likely to be interested in and remember your brand.

Page and content creation/optimization

Make your site content easily crawled by the search engines, intuitively navigable for users with intelligent linking structures, and implement SEO best practices.

SEO for Lead Generation and Direct Marketing

Although online lead generation is less direct than an ecommerce transaction, it is arguably just as valuable and important for building customers, revenue, and long-term value. Millions of search queries have commercial intents that can't be (or currently aren't) fulfilled directly online. These can include searches for services such as legal consulting, contract construction, commercial loan requests, alternative energy providers, or virtually any service or product people source via the internet.

Here are some factors to think about when considering SEO for lead generation and direct marketing:

When to employ SEO for lead generation and direct marketing

Useful for non-e-commerce products/services/goals that you want users to accomplish on your site, or for which you are hoping to attract inquiries and/or direct contact over the Web.

Keyword targeting

As with ecommerce, choose phrases that have measurable traffic and have previously converted well in PPC campaigns. In a “considered sale” or enterprise-scale business decision for users, include long-tail keywords that might pertain to customer research activities or pain points, to educate and persuade via content. Avoid arcane jargon terms that aren’t typically understood outside your company.

Page and content creation/optimization

Although you might think it would be easier to rank high in the SERPs for lead-generation programs than for ecommerce, it is often equally challenging. A solid combination of content development and on-site optimization are necessary to be competitive in the more challenging arenas.

SEO for Reputation Management

Since one’s own name—whether personal or corporate—is one’s identity, establishing and maintaining the reputation associated with that identity is generally of great interest.

Imagine that you search for your brand name in a search engine, and high up in the search results is a web page that is highly critical of your organization.

SEO for reputation management is a process for neutralizing negative mentions of your name in the SERPs. In this type of SEO project, you would strive to occupy additional spots in the top 10 results to push the critical listing lower, and hopefully off the first page.

SEO enables this process through both content creation and promotion via link development, as well as through optimization of content on third party platforms, such as Pinterest, Facebook, and LinkedIn. Although reputation management is among the most challenging of SEO tasks (primarily because you are optimizing many results for a query rather than one), demand for these types of services is rising as more and more companies become aware of the issue.

Here are some factors to think about when considering SEO for reputation management:

When to employ SEO for reputation management

If you’re trying to either protect your brand from negative results appearing on page 1 or push down already existing negative content.

Keyword targeting

Chances are this is very easy—the keyword you are targeting is a person’s name, your brand name, or some common variant (and you already know what it is). You might want to use keyword research tools just to see whether there are popular variants you’re missing.

Page and content creation/optimization

Unlike the other SEO tactics, reputation management involves optimizing pages on many different domains to demote negative listings. This involves using social media profiles and other third party platform pages, public relations, press releases, and links from networks of sites you might own or control, along with classic optimization of internal links and on-page elements. It is certainly among the most challenging of SEO practices, especially in Google, where the use of the *query deserves diversity* (QDD) algorithm can mean you have to work much harder because of how it favors diverse content.

On the topic of negative search results in Google, in May, 2014 the Luxembourg-based European Union Court of Justice (ECJ) ruled that people can ask Google to remove search results containing information about them, with the new, EU-specific “right to be forgotten”

(<http://www.reuters.com/article/2014/05/13/eu-google-dataprotection-idUSL6N0NZ23Q20140513>). While some view this as a victory for privacy laws, others view it as censorship – time will tell how this plays out internationally and on a case-by-case basis for individuals.

SEO for Ideological Influence

For those seeking to sway public (or private) opinion about a particular topic, SEO can be a powerful tool. By promoting ideas and content within the search results for queries likely to be made by those seeking information about a specific topic, you can influence the perception of even very large groups. Politicians and political groups and individuals are the most likely employers of this tactic, but it can certainly be applied to any subject, from the theological to the technical or civic.

Some factors to consider when using SEO for ideological influence:

When to employ

Use this tactic when you need to change minds or influence decisions/thinking around a subject—for example, a group of theoretical physicists attempting to get more of their peers to consider the possibility of alternative universes as a dark matter source.

Keyword targeting

Generally in these types of campaigns, you probably know the primary keywords you’re chasing and can use keyword research query expansion to find others. Consider inclusion of both “pro” and “con” phrases for maximum visibility in a debate. For example, “Obamacare” typically is considered a conservative/critical label while “Affordable care act” is the official label and is more commonly used by supporters.

Page and content creation/optimization

This is classic SEO, but with a twist. Since you’re engaging in ideological warfare in the SERPs, chances are you’ve got allies you can rally to the cause. Leverage your combined links and content to espouse your philosophical preferences.

Advanced Methods for Planning and Evaluation

There are many methodologies for business planning. One of the better-known ones is the SWOT (Strengths, Weaknesses, Opportunities, Threats) analysis. There are also methodologies for ensuring that the plan objectives are the right type of objectives, such as the SMART (Specific, Measurable, Achievable, Realistic, Timelined) plan. We will take a look at both of these in the context of SEO.

SWOT analysis

Sometimes you need to get back to the basics and carry out a simple evaluation of where you are in the marketplace, as well as where you would like to be – for this, a SWOT analysis is a great starting point. It creates a grid from which to work and is very simple to execute.

As you can see from the SWOT chart in Figure 3-3, strengths and weaknesses usually stem from internal (on-site, business operational, business resource) sources, whereas opportunities and threats are from external sources.

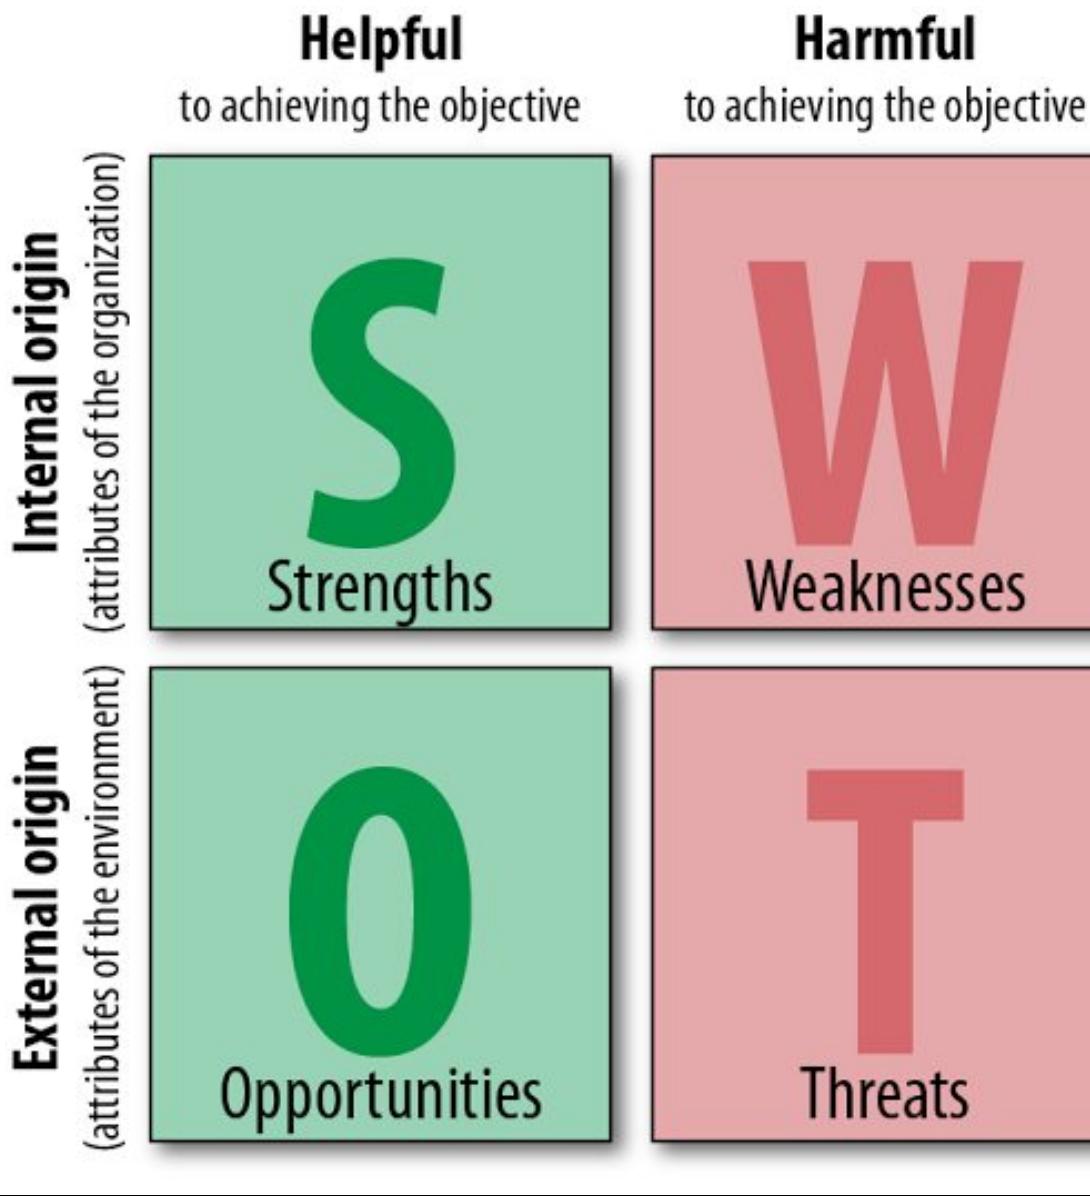


Figure 3-3. Example SWOT chart

Where does SEO fit in here? To explore this, it is helpful to use an example. Take Business X. It has a website that was built on WordPress, makes use of category tagging, adds at least one page of content every two days, and has excellent knowledge of its industry. Its domain name isn't ideal—Businessnameandkeyword.com—but it is decent.

Business X does not get much traffic from search engines, but its rival, Business Y, does because Business Y has had its website up for a long period of time and received some great links along the way. Business Y doesn't have any SEO plan and relies on its main page to bring in all of its search traffic. This is because Business Y has a keyword-rich domain name and people have used those keywords in their links to Business Y's website (giving it keyword-rich anchor text), and because of its longevity on the Web.

There aren't a lot of target search queries; in fact, there are fewer than 50,000 searches per month for the core set of keywords. Business X's site ranks on the second page of Google results, whereas Business Y is ranked #3, with Wikipedia and About.com taking up the top two positions.

Neither of the businesses is spending money on PPC (paid search) traffic, and the niche doesn't have much room for other entrants (there may be 10 to 15 competitors). Both sites have similar link authority in terms of strengths and numbers. The businesses deal in impulse purchases—the products evoke strong emotions.

Figure 3-4 shows what the SWOT for Business X might look like.

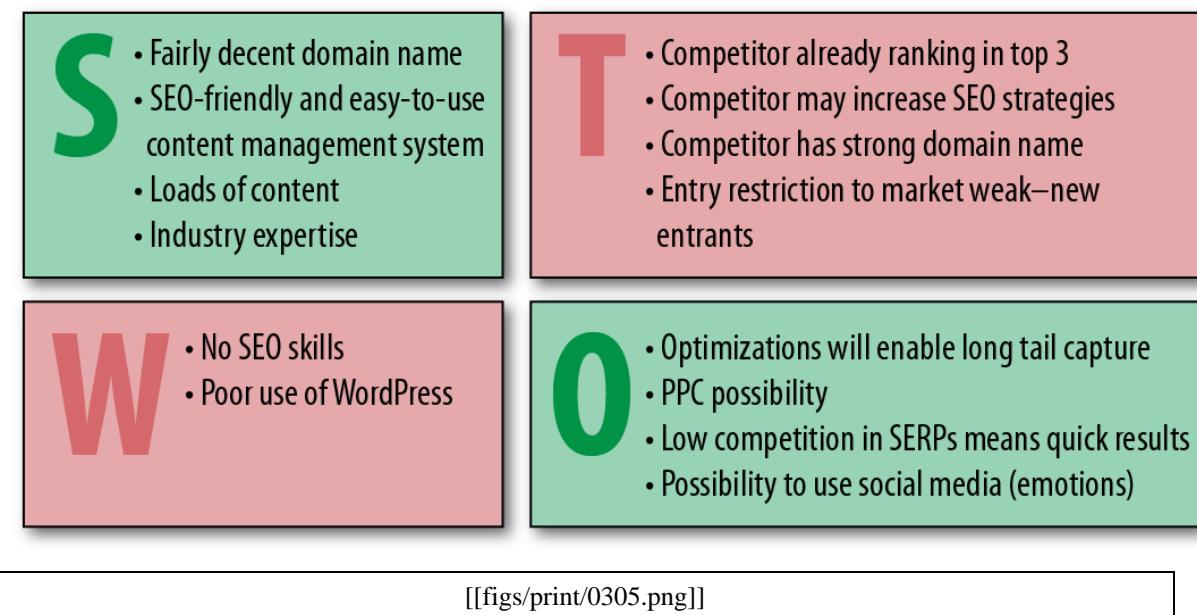


Figure 3-4. Sample SWOT chart data for Business X

The preceding analysis suggests where Business X can get some quick wins for its site, as well as where the priorities are. It also forms a great starting point for a long-term strategy and tactical maneuvers. This example is simplistic, but it illustrates how instructive a fleshed out SWOT can be. It does require you to have analyzed your site, your main competitor(s), your target keywords, and the search engine results pages (SERPs).

SWOT Guidelines

Identifying *strengths* is typically one of the easier objectives:

- What sources of traffic are working well (converting on established site goals) for your site/business?
- Which projects/properties/partnerships are driving positive momentum toward traffic/revenue goals?
- Which of your content sections/types produces high traffic, conversions, and ROI?
- What changes have you made historically that produced significant value?

Determining the *weaknesses* can be tougher (and requires emotional separation from the data):

- What content is currently driving low levels of search/visitor traffic?
- Which changes that were intended to produce positive results have shown little/no value?
- Which traffic sources are underperforming or under-delivering?
- What projects/properties/partnerships are being leveraged poorly, or not at all?

Parsing *opportunities* requires a combination of strength and weakness analysis. You want to find areas that are doing well but have room to expand, as well as those that have yet to be explored:

- What brainstormed but undeveloped or untested projects/ideas can have a significant, positive impact?
- What traffic sources currently sending good-quality traffic could be expanded to provide more value?

- What areas of weakness have direct paths to recovery?
- Which website changes have had positive results? Can these be applied more rigorously or to other areas for increased benefit?
- What new markets or new content areas are potentially viable/valuable for expansion?
- What sources of new content/new links have yet to be tapped?
- What third party platforms (social media, content curators, etc.) can be utilized to expand reach and increase engagement?

Determining *threats* can be the most challenging of the tasks. You'll need to combine creative thinking with an honest assessment of your weaknesses and your competitors' strengths, and consider the possibilities of macro-events that could shape your website/company's future:

- In your areas of weakness, which players in your market (or other, similar markets) are strong? How have they accomplished this?
- What shifts in human behavior, web usage, or market conditions could dramatically impact your business/site? (For example, consider the "what if people stopped searching and instead navigated the Web in different ways" perspective. It is a bit "pie in the sky," but we have already seen Craigslist make classifieds obsolete, and have witnessed Facebook start to take advertising market share from the search engines.)
- Which competitors have had the most success in your arena? How have they accomplished this? Where do they intersect with your business/customers?
- Are there any strategies implemented by start-ups in similar businesses that have had massive success in a particular arena that could be dangerous to your business if they were replicated in your market?

Conducting SWOT analysis from a web marketing and SEO perspective is certainly one of the most valuable first steps you can take as an organization poised to expend resources. If you haven't taken the time to analyze the landscape from these bird's-eye-view perspectives, you might end up like a great runner who's simply gone off the course—sure, you'll finish fast, but where will it take you?

Get SMART

Every company is unique, so naturally their challenges are unique. Even a second SEO initiative within the same company will not be the same as the first initiative. Your first SEO efforts will have changed things, creating new benchmarks, new expectations, and different objectives. Thus, each SEO effort is a new endeavor.

One way to start a new project is to set SMART objectives. Let's look at how to go about doing that in the world of SEO.

Specific objectives are important. It is easy to get caught up in the details of the plan and lose sight of the broader site objectives. You may think you want to rank #1 for this phrase or that, but in reality what you want is more granular than that – more leads, more page views, more customers. Perhaps you don't even need more customers from organic search, but you want higher sales volumes, so in fact having the same number of orders but with a higher average order value would meet your objectives better.

Measurable objectives are essential if one is to manage the performance in meeting them—you can't manage what you can't measure. SEO practitioners have to help their clients or organizations come to grips with analytics, and not just the analytics software, but the actual processes of how to gather the data, how to sort it, and, most importantly, how to use it to make informed decisions.

Achievable objectives are ones that can be accomplished with the available resources. You could decide to put a man on Mars next year, for example, but it is just too big an undertaking to be feasible. You can be ambitious, but it is important to pick goals that can be met. You cannot possibly sell to more people than exist in your market. There are limits to markets, and at a certain point the only growth can come from opening new markets, or developing new products for the existing market.

Aside from basic business achievability, there are also limits to what can rank at #1 for a given search query. The search engines want the #1 result to be the one that offers the most value for users, and unless you are close to having the website that offers the most value to users, it may be unreasonable to expect to get to that position, or to maintain it if you succeed in getting there.

Realistic objectives are about context and resources. It may be perfectly achievable to meet a certain objective, but only with greater resources than may be presently available. Even a top ranking on the most competitive terms around is achievable for a relevant product, but it is a realistic goal only if the resources required for such an effort are available.

Time-bound is the final part of a SMART objective. If there is no timeline, no project can ever fail, since it can't run out of time. SEO generally tends to take longer to implement and gather momentum than a paid search advertising campaign. It is important that milestones and deadlines be set so that expectations can be managed and course corrections made.

“We want to rank at #1 for loans” is not a SMART objective. It doesn’t identify the specific reason why the company thinks a #1 ranking will help it. It doesn’t have a timeline, so there is no way to fail. It doesn’t state an engine on which to be #1, so there’s a guaranteed argument if the intention is to rank well on both Google and Bing, but the result is only high rankings on Bing.

“We want to increase approved loan applications generated by organic search by 30% over six months” is a far better objective. There is a deadline, and the company can certainly gauge progress toward the specific objective. The company can look at its current market share and the resources committed to see whether this is an achievable and realistic goal.

Conclusion

To bring this all together successfully, your objectives, strategies, and tactics need to be aligned – and they need to take into account your market, your business, and the competition. Don’t spread yourself too thin. Remember to ask yourself the tough questions, such as:

- Does your company need direct sales, traffic, branding, or some combination of these?
- Are there specific influencers you’re trying to reach with a message?
- Is the organization/brand subject to potentially negative material that needs to be controlled/mitigated?
- Do you have products/services you sell, either directly over the Web or through leads established online?
- Do you have the resources to develop new, unique, and interesting content?

Getting the answers won’t be easy, but it will be worth the effort!

4

SEO Implementation: First Stages

SEO efforts require forethought and planning to obtain the best results, and a business's SEO strategy is ideally incorporated into the planning stages of a website development or redevelopment project, as well into ongoing web development efforts. Website architecture (including the selection of a content management system, or CMS), the overall marketing plan (including branding objectives), content development efforts and more are affected by SEO.

In this chapter, we will discuss several aspects of how SEO efforts for both desktop and mobile generally begin, including:

- Putting together an SEO strategy
- Performing a technical SEO audit of site versions
- Setting a baseline for measuring results and progress

The Importance of Planning

As discussed in Chapter 3, your SEO strategy should be incorporated into the site planning process long before your site goes live – and your process should be well outlined before you make even the most basic technology choices, such as deciding on your hosting platform and content management system. However, this is not always possible—and in fact, more often than not, SEO efforts in earnest often begin well after a site has been launched and in use for some time.

In all scenarios, there are major components to any SEO strategy that need to be addressed long before you craft your first HTML title tag.

Identifying the Site Development Process and Players

Before you start the SEO process, it is important to identify who your target audience is, what your message is, and how your message is relevant. There are no web design tools or programming languages that tell you these things. Your company's marketing, advertising, and PR teams have to set the objectives before you can implement them—successful SEO requires a team effort.

Your SEO team should be cross-functional and multidisciplinary, consisting of the team manager, the technical team, the creative team, the data and analytics team (if you have one), and the major stakeholders from marketing, advertising, and PR. In a smaller organization, you may have to wear all of those hats yourself (we never said SEO was easy!).

The SEO team leader wants to know who the website's target audience is. What does the marketing team know about them? How did we find them? What metrics will we use to track them? All of this is key information that should have an impact on various aspects of the project's technical implementation.

The PR team has to take your story to the media and entice them into writing and talking about it. What message do they want to deliver? You have to mirror that message in your content. If they say you're relevant to organic cotton clothes but your project plan says you're relevant to yoga attire, the whole project is in trouble. When you're creating visibility, the people who build up your brand have to see a clear, concise focus in what you do. If you provide them with anything less, their interest in your business will be brief, and they'll find someone else to talk about.

The technical and creative teams are responsible for implementing the majority of an SEO strategy. They take direction from marketing, advertising, and PR on what needs to be accomplished, but from there on out they have to put the pieces into place. As the SEO implementation process unfolds, marketing has to come back and say whether the target audience is being reached, advertising has to come back and say whether the message is clear, and PR has to come back and say whether the media like what they see.

Ongoing feedback is essential because the success of your strategy is determined solely by whether you're meeting your goals. A successful SEO team understands all of these interactions and is comfortable relying on each team member to do his part, and establishing good communication among team members is essential.

And even if you are a team of one, you still need to understand all of these steps, as addressing all aspects of the marketing problem (as relates to SEO) is a requirement for success.

Development Platform & Information Architecture

Whether you're working with an established website or not, you should plan to research the desired site architecture (from an SEO perspective) as a core element of your SEO strategy. This task can be divided into two major components: technology decisions and structural decisions.

Technology Decisions

SEO is a technical process, and as such, it impacts major technology choices. For example, a content management system (CMS) can facilitate—or, possibly, undermine—your SEO strategy. For example, some platforms do not allow you to write customized titles and meta descriptions that vary from one web page to the next, some create hundreds (or thousands) of pages of duplicate content (not good for SEO!). For a deep dive into the technical issues you'll need to familiarize yourself with in order to make the right technology decisions for your SEO needs, turn to Chapter 6, “Developing an SEO Friendly Website.” The technology choices you make at the outset of developing your site and publishing your content can make or break your SEO efforts – and it is best to make the right choices in the beginning to save yourself headaches down the road.

As we outlined previously in this chapter, your technology choices can have a major impact on your SEO results. The following is an outline of the most important issues to address at the outset:

Dynamic URLs

Dynamic URLs are URLs for dynamic web pages (which have content generated “on the fly” by user requests). These URLs are generated in real-time as the result of specific queries to a site’s database – for example, a search for “leather bag” on Etsy.com results in the dynamic search result URL <https://www.etsy.com/search?q=leather%20bag>. However, Etsy also has a static URL for a static page showing leather bags, here: https://www.etsy.com/market/leather_bag.

Although Google has stated for some time that dynamic URLs are not a problem for the search engine to crawl, it is wise to make sure your dynamic URLs are not “running wild” by making sure your CMS does not render your pages on URLs with too many convoluted parameters in them. In addition, be sure to make proper use of the `rel="canonical"` tag as outlined by Google: <https://support.google.com/webmasters/answer/139066?hl=en>.

Finally, while dynamic URLs are crawlable, don't overlook the value of static URLs for the purpose of controlling your URL structure for brevity, descriptiveness, user-friendliness, and ease of sharing.

Session IDs or user IDs in the URL

It used to be very common for CMSs to track individual users surfing a site by adding a tracking code to the end of the URL. Although this worked well for this purpose, it was not good for search engines, because they saw each URL as a different page rather than variants of the same page. Make sure your CMS does not ever serve up session IDs. If you are not able to do this, making sure you use `rel="canonical"` on your URLs (what this is, and how to use it, is explained in Chapter 6).

Superfluous parameters in the URL

Related to the preceding two items is the notion of extra characters being present in the URL. This may bother the search engines, and it interferes with the user experience for your site.

Links or content based in Flash

Search engines often cannot see links and content implemented using Flash technology. Make sure there is a plan is to expose your links and content in simple HTML text, and be aware of Flash's limitations.

Content behind forms (including pull-down lists)

Making content accessible only after the user has completed a form (such as a login) or made a selection from an improperly implemented pull-down list is a great way to hide content from the search engines. Do not use these techniques unless you want to hide your content!

Temporary (302) redirects

This is also a common problem in web server platforms and CMSs. The 302 redirect blocks a search engine from recognizing that you have permanently moved the content, and it can be very problematic for SEO as 302 redirects block the passing of PageRank. You need to make sure the default redirect your systems use is a 301, or understand how to configure it so that it becomes the default.

All of these are examples of basic technology choices that can adversely affect your chances for a successful SEO project. Do not be fooled into thinking that SEO issues are understood, let alone addressed, by all CMS vendors out there—unbelievably, some are still very far behind the SEO curve. It is also important to consider whether a “custom” CMS is truly needed when many CMS vendors are creating ever more SEO-friendly systems—often with much more flexibility for customization and a broader development base, and customizable, SEO-specific “modules” that can quickly and easily add SEO functionality. There are also advantages to selecting a widely used CMS, including portability in the event that you choose to hire different developers at some point.

Also, do not assume that all web developers understand the SEO implications of what they develop. Learning about SEO is not a requirement to get a software engineering degree or become a web developer (in fact, there are still very few college courses that adequately cover SEO). It is up to you, the SEO expert, to educate the other team members on this issue as early as possible in the development process.

Structural Decisions

One of the most basic decisions to make about a website concerns internal linking and navigational structures, which are generally mapped out in a site architecture document. What pages are linked to from the home page? What pages are used as top-level categories that then lead site visitors to other related pages? Do pages that are relevant to each other link to each other? There are many, many aspects to determining a linking structure for a site, and it is a major usability issue because visitors make use of the links to surf around your website. For search engines, the navigation structure helps their crawlers determine what pages you consider the most important on your site, and it helps them establish the relevance of the pages on your site to specific topics.

This section outlines a number of key factors that you need to consider before launching into developing or modifying a website. The first step will be to obtain a current site information architecture (IA) document for reference, or to build one out for a new site. From here, you can begin to understand how your content types, topics, and products will be organized.

Target keywords

Keyword research is a critical component of SEO. What search terms do people use when searching for products or services similar to yours? How do those terms match up with your site hierarchy? Ultimately, the logical structure of your pages should match up with the way users think about products and services like yours. Figure 4-2 shows how this is done on the Amazon site.

Figure 4-1. Example of a well thought out site hierarchy

Cross-link relevant content

Linking between articles that cover related material can be very powerful. It helps the search engine ascertain with greater confidence how relevant a web page is to a particular topic. This can be extremely difficult to do well if you have a massive ecommerce site, but Amazon solves the problem very well, as shown in Figure 4-3.

[[figs/print/0402.png]]

Figure 4-2. Product cross-linking on Amazon

The “Frequently Bought Together” and “What Do Customers Ultimately Buy After Viewing This Item?” sections are brilliant ways to group products into categories that establish the relevance of the page to certain topic areas, as well as to create links between relevant pages.

In the Amazon system, all of this is rendered on the page dynamically, so it requires little day-to-day effort on Amazon’s part. The “Customers Who Bought...” data is part of Amazon’s internal databases, and the “Tags Customers Associate...” data is provided directly by the users themselves.

Of course, your site may be quite different, but the lesson is the same. You want to plan on having a site architecture that will allow you to cross-link related items.

Use anchor text, intuitively

Anchor text has generally been one of the golden opportunities of internal linking, and exact-match keyword anchor text was generally the protocol for internal linking for many years. However, in these days of aggressive anchor text abuse (and crackdown by the search engines), while keyword-infused anchor text in internal links is still often the most intuitive and user-friendly, we generally advocate for a more broad-minded approach to crafting internal anchor text. Use descriptive text in your internal links and avoid using irrelevant text such as “More” or “Click here.” Try to be as specific and contextually relevant as possible and include phrases when appropriate within your link text. For example, as a crystal vendor, you might use “some of our finest quartz specimens” as anchor text for an internal link, vs. “quality quartz crystals.” Make sure that the technical, creative, and editorial teams understand this approach, as it will impact how content is created, published, and linked to within your site.

Use breadcrumb navigation

Breadcrumb navigation is a way to show the user where he is in the navigation hierarchy. Figure 4-4 shows an example from PetSmart.

The screenshot shows the PetSmart homepage with a navigation bar at the top. The main content area displays a breadcrumb navigation path: Home > CAT > FOOD & HEALTH > Food. This path is circled in red. Below the breadcrumb, there is a section titled "SHOP BY:" with categories like FOOD and HEALTH CONSIDERATION. The FOOD category is selected, showing "536 Items Available". There is a "COMPARE" section where users can add up to 3 items. The main content area displays three bags of Blue Buffalo cat food: "WILDERNESS" natural grain-free dry cat food (chicken flavor), "INDOOR" natural grain-free dry cat food (chicken flavor), and another bag of "WILDERNESS" natural grain-free dry cat food (chicken flavor). The bags feature images of a cat and the Blue Buffalo logo.

Figure 4-3. Breadcrumb bar on PetSmart.com (<http://www.petSMART.com>)

This page is currently four levels down from the home page. Also, note how the anchor text in the breadcrumb is keyword-rich, as is the menu navigation on the left. This is helpful to both users and search engines.

Minimize link depth

Search engines (and people) look to the site architecture for clues as to what pages are most important. A key measurement is how many clicks from the home page it takes a person, and a search engine crawler, to reach a page. A page that is only one click from the home page is clearly important, while a page that is five or six clicks away from the home page is not nearly as influential. In fact, the search engine spider may never even find such a page, depending in part on the site's link authority.

Standard SEO advice is to keep the site architecture as flat as possible, to minimize clicks from the home page to important content. The bottom line is that you need to plan out a site structure that is as flat as you can reasonably make it without compromising your user experience.

In this and the preceding sections, we outlined common structural decisions that you need to incorporate into your SEO strategy prior to implementation. There are other considerations, such as how you might be able to make your efforts scale across a very large site (thousands of pages or more). In such a situation, you cannot feasibly review every page one by one.

Mobile Sites & Mobile Apps

If you are building a website, you need to build a mobile version if you want to take full advantage of organic search through SEO – and depending on your business, you may benefit from developing a mobile app as well. The main consideration regarding your site's mobile version is whether to host the mobile version on same or separate URLs as your desktop version – and, if utilizing same URLs, whether to choose Responsive Design or Dynamic Serving (aka Adaptive Design), discussed in detail in Chapter 10.

Auditing an Existing Site to Identify SEO Problems

Auditing an existing site is one of the most important tasks that SEO professionals encounter. SEO is still a relatively new field, and many of the limitations of search engine crawlers are non-intuitive. In addition, many web developers, unfortunately, are still not well versed in SEO. Even more unfortunately, some stubbornly refuse to learn, or, worse still, have learned the wrong things about SEO! This includes those who have developed CMS platforms, so there is a lot of opportunity to find problems when conducting a site audit. While you may have to deal with some headaches in this department (trust us, we still deal with this on a regular basis) – your evangelism for SEO, and hopefully support from key stakeholders, will set the stage for an effective SEO strategy.

Elements of an Audit

Your website needs to be a strong foundation for the rest of your SEO efforts to succeed. An SEO site audit is often the first step in executing an SEO strategy. Both your desktop and mobile site versions need to be audited for SEO effectiveness. The following sections identify what you should look for when performing a site audit.

Page load time

Is the page load time excessive? Too long a load time may slow down crawling and indexing of the site – and can virtually eliminate your site from competitiveness in mobile search.

Mobile-friendliness

Your site should have a fast, mobile-friendly version that is served to mobile devices.

Usability

Although this may not be seen as a direct SEO issue, it is a very good place to start. Usability affects many factors, including conversion rate as well as the propensity of people to link to a site.

Accessibility/spider-ability

Make sure the site is friendly to search engine spiders (discussed in more detail in “Making Your Site Accessible to Search Engines” and “Creating an Optimal Information Architecture (IA)” in Chapter 6).

Search engine health check

Here are some quick health checks:

- Perform a `site:yourdomain.com` search in the search engines to check how many of your pages appear to be in the index. Compare this to the number of unique pages you believe you have on your site. Also, check indexation numbers in your Google and Bing Webmaster Tools accounts.
- Check the Google cache to make sure the cached versions of your pages look the same as the live versions of your pages.
- Check to ensure major search engine Webmaster Tools have been verified for the domain (and any subdomains, for mobile or other content areas). Google and Bing currently offer site owner validation to “peek” under the hood of how the engines view your site).
- Test a search on your brand terms to make sure you are ranking for them (if not, you may be suffering from a penalty; be sure to check your associated Webmaster Tools accounts to see if there are any identifiable penalties, or any other helpful information).

Keyword health checks

Are the right keywords being targeted? Does the site architecture logically flow from the way users search on related keywords? Does more than one page target the same exact keyword (a.k.a. *keyword cannibalization*)? We will discuss these items in “Keyword Targeting”.

Duplicate content checks

The first thing you should do is to make sure the non-*www* versions of your pages (i.e., `http://yourdomain.com`) 301-redirect to the *www* versions of your pages (i.e., `http://www.yourdomain.com`), or vice versa (this is often called the *canonical redirect*). While you are at it, check that you don’t have *https*: pages that are duplicates of your *http*: pages. You should check the rest of the content on the site as well.

The easiest way to do this is to take unique text sections from each of the major content pages on the site and search on them in Google. Make sure you enclose the string inside double quotes (e.g., “*a phrase from your website that you are using to check for duplicate content*”) so that Google will search for that exact string. If you see more than one link showing in the results, look closely at the URLs and pages to determine why it is happening.

If your site is monstrously large and this is too big a task, make sure you check the most important pages, and have a process for reviewing new content before it goes live on the site.

You can also use commands such as `inurl:` and `intitle:` (see Table 2-1) to check for duplicate content. For example, if you have URLs for pages that have distinctive components to them (e.g., “1968-mustang-blue” or “1097495”), you can search for these with the `inurl:` command and see whether they return more than one page.

Another duplicate content task to perform is to make sure each piece of content is accessible at only one URL. This probably trips up more big commercial sites than any other issue. The issue is that the same content is accessible in multiple ways and on multiple URLs, forcing the search engines (and visitors) to

choose which is the canonical version, which to link to, and which to disregard. No one wins when sites fight themselves - if you have to deliver the content in different ways, rely on cookies so that you don't confuse search engine spiders.

URL check

Make sure you have clean, short, descriptive URLs. *Descriptive* means keyword-rich but not keyword-stuffed (eg. site.com/outerwear/mens/hats is keyword rich; site.com/outerwear/mens/hat-hats-hats-for-men is keyword-stuffed!). You don't want parameters appended (have a minimal number if you must have any), and you want them to be simple and easy for users (and search engine spiders) to understand.

HTML title tag review

Make sure the title tag on each page of the site is unique and descriptive. If you want to include your company brand name in the title, consider putting it at the end of the title tag, not at the beginning, as placement of keywords at the front of a page title (generally referred to as prominence) brings ranking benefits due to the weighting added to words near the beginning of the title. Also check to ensure the title tag is fewer than 70 characters long, or 512 pixels wide.

Content review

Do the main pages of the site have enough text content to engage and satisfy a site visitor? Do these pages all make use of header tags? A subtler variation of this is making sure the number of pages on the site with little content is not too high compared to the total number of pages on the site.

Meta tag review

Check for a meta `robots` tag on the pages of the site. If you find one, you may have already spotted trouble. An unintentional `NoIndex` or `NoFollow` tag (we define these in "Content Delivery and Search Spider Control") could adversely affect your SEO efforts.

Also make sure every page has a unique meta description. If for some reason that is not possible, consider removing the meta description altogether. Although the meta description tags are generally not a direct factor in ranking, they may well be used in duplicate content calculations, and the search engines frequently use them as the description for your web page in the SERPs; therefore, they can affect click-through rate.

Sitemaps file and robots.txt file verification

Use the Google Webmaster Tools "Robots.txt fetch" to check your `robots.txt` file. Also verify that your `Sitemap` file is correctly identifying all of your site pages.

URL redirect checks

Check all redirects to make sure the right redirect is in place, and it is pointing to the correct destination URL. This also includes checking that the canonical redirect is properly implemented. Use a server header checker such as redirectcheck.com, redirect-checker.org, or when using Firefox, install the browser extension "Redirect Check Client" at <https://addons.mozilla.org/en-US/firefox/addon/redirectcheck-client/>.

Not all URL redirect types are created equal, and some can create problems for SEO. Be sure to research and understand proper use of URL redirects before implementing them, and minimize the number of redirects needed over time by updating your internal linking and navigation so that you have as few redirects as possible.

Internal linking checks

Look for pages that have excessive links. As discussed earlier, make sure the site makes intelligent use of anchor text in its internal links. This is a user-friendly opportunity to inform users and search engines what the various pages of your site are about. Don't abuse it, though. For example, if you have a link to your home page in your global navigation (which you should), call it "Home" instead of picking your juiciest keyword. The search engines can view that particular practice as spammy, and it does not engender a good user experience.

A brief aside about hoarding PageRank: many people have taken this to an extreme and built sites where they refused to link out to other quality websites, because they feared losing visitors and link authority. Ignore this idea! You should link out to quality websites. It is good for users, and it is likely to bring you ranking benefits (through building trust and relevance based on what sites you link to). Just think of your human users and deliver what they are likely to want. It is remarkable how far this will take you.

Avoidance of unnecessary subdomains

The engines may not apply the entirety of a domain's trust and link authority, weight to subdomains. This is largely due to the fact that a subdomain could be under the control of a different party, and therefore in the search engine's eyes it needs to be separately evaluated. In the great majority of cases, content that gets placed within its own subdomain can easily go in a subfolder – such as site.com/content, as opposed to the subdomain content.site.com.

Geolocation

If the domain is targeting a specific country, make sure the guidelines for country geotargeting outlined in “Best Practices for Multilanguage/Country Targeting” in Chapter 6 are being followed. If your concern is primarily about ranking for “*san francisco chiropractor*” because you own a chiropractic office in San Francisco, CA, make sure your address is on every page of your site. You should also claim and ensure the validity of your Google Places listings to ensure data consistency - this is discussed in detail in Chapter 10.

External linking

Check the inbound links to the site by performing a backlink analysis. Use a backlinking tool such as LinkResearchTools (<http://www.linkresearchtools.com/>), Open Site Explorer (<http://www.opensiteexplorer.org>), Majestic SEO (<http://www.majesticseo.com>), or Ahrefs Site Explorer & Backlink Checker (<https://ahrefs.com>) to collect data about your links. Look for bad patterns in the anchor text, such as 87% of the links having the critical keyword for the site in them. Unless the critical keyword happens to also be the name of the company, this is a sure sign of trouble. This type of distribution is quite likely the result of link purchasing or other manipulative behavior, and will (if it hasn't already) likely earn you a manual Google penalty or trigger Google's Penguin algorithm to lower your rankings.

On the flip side, make sure the site's critical topics and keywords are showing up sometimes.. A lack of the topically related anchor text is not entirely good, either. You need to find a balance, and err on the side of caution, intuitiveness, and usability.

Also look to see that there are links to pages other than the home page. These are often called *deep links* and they will help drive the ranking of key sections of your site. You should look at the links themselves, by visiting the linking pages and see whether the links appear to be paid for. They may be overtly labeled as sponsored, or their placement may be such that they are clearly not natural endorsements. Too many of these are another sure sign of trouble in the backlink profile.

Lastly, check how the backlink profile for the site compares to the backlink profiles of its major competitors. Make sure that there are enough external links to your site, and that there are enough high quality links in the mix.

Image alt attributes

Do all the images have relevant, keyword-rich image alt attribute text and filenames? Search engines can't easily tell what is inside an image, and the best way to provide them with some clues is with the alt attribute and the filename of the image. These can also reinforce the overall context of the page itself.

Code quality

Although W3C validation is not something the search engines require, checking the code itself is a good idea (you can check it with the W3C validator located at <http://validator.w3.org>). Poor coding can have some undesirable impacts. You can use a tool such as SEO Browser (<http://www.seo-browser.com>) to see how the search engines see the page.

The Importance of Keyword Reviews

Another critical component of an architecture audit is a keyword review. Basically, this involves the following steps.

Step 1: Keyword research

It is vital to examine your topic and keyword strategy as early as possible in any SEO effort. You can read about this in more detail in Chapter 5.

Step 2: Site architecture

Coming up with architecture for a website can be very tricky. At this stage, you need to look at your keyword research and the existing site (to make as few changes as possible). You can think of this in terms of your site map.

You need a hierarchy that leads site visitors to your *high-value pages* (i.e., the pages where conversions are most likely to occur). Obviously, a good site hierarchy allows the parents of your "money pages" to rank for relevant keywords, which are likely to be shorter tail.

Most products have an obvious hierarchy they fit into, but for products with descriptions, categories and concepts that can have multiple hierarchies, deciding on a site's information architecture can become very tricky. Some of the trickiest hierarchies, in our opinion, can occur when there is a location involved. In London alone there are London boroughs, metropolitan boroughs, Tube stations, and postcodes. London even has a city ("The City of London") within it.

In an ideal world, you will end up with a single hierarchy that is natural to your users and gives the closest mapping to your keywords - but whenever there are multiple ways in which people search for the same product, establishing a hierarchy becomes challenging.

Step 3: Keyword mapping

Once you have a list of keywords and a good sense of the overall architecture, start mapping the major relevant keywords to URLs. When you do this, it is very easy to spot pages that you were considering creating that aren't targeting a keyword (perhaps you might skip creating these), and, more importantly, keywords that don't have a page.

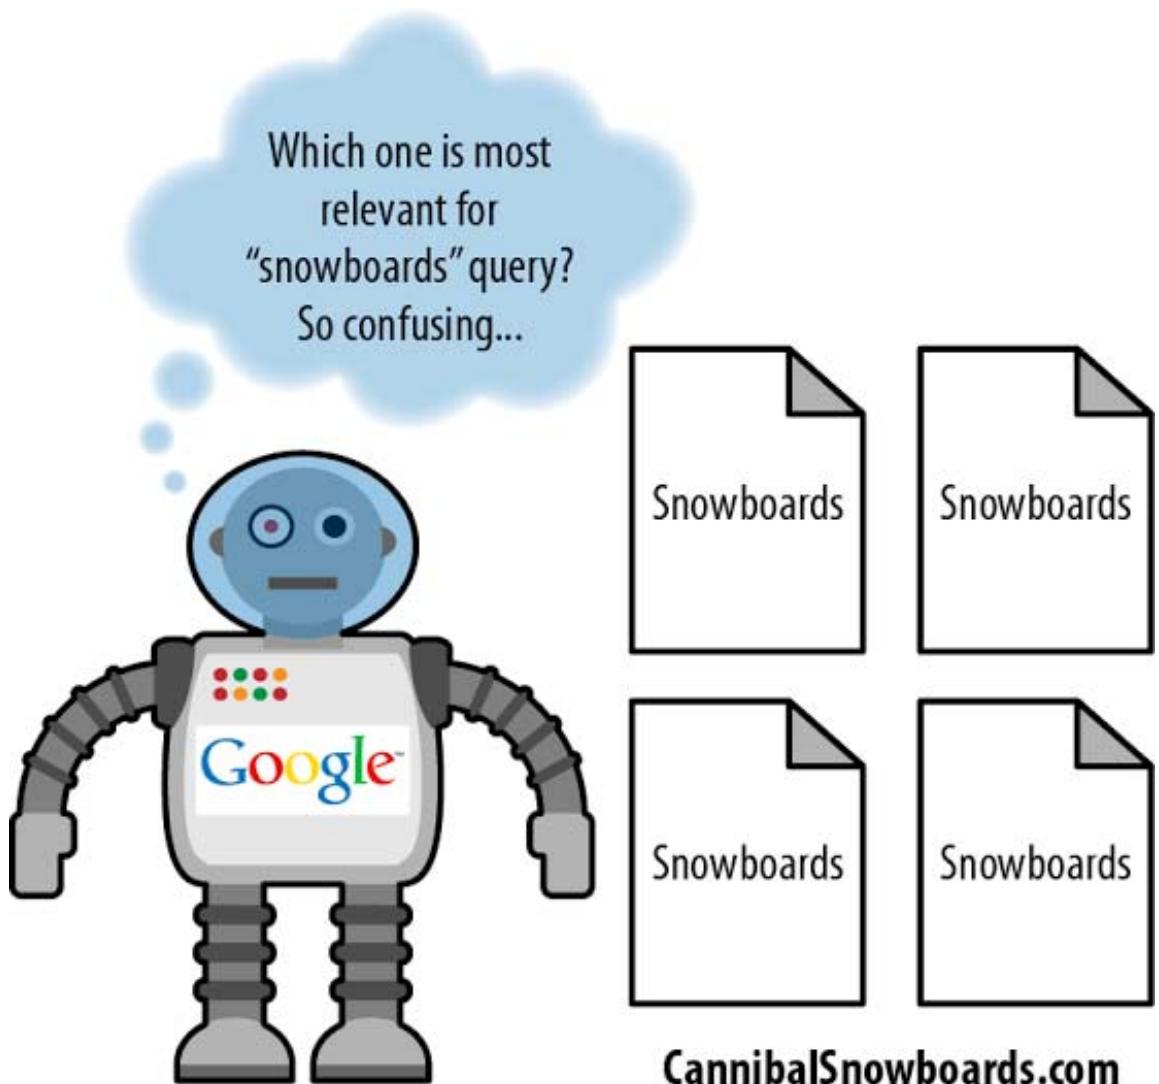
If this stage is causing you problems, revisit step 2. Your site architecture should lead naturally to a mapping that is easy to use and includes your keywords.

Step 4: Site review

Once you are armed with your keyword mapping, the rest of the site review will flow more easily. For example, when you are looking at crafting your title tags and headings, you can refer back to your keyword mapping and see not only whether there is appropriate use of tags (such as the heading tag), but also whether it includes the appropriate keyword targets.

Keyword Cannibalization

Keyword cannibalization typically starts when a website's information architecture calls for the targeting of a single term or phrase on multiple pages of the site. This is often done unintentionally, but it can result in several or even dozens of pages that have the same keyword target in the title and header tags. Figure 4-5 shows the problem.



[[figs/print/0404.png]]

Figure 4-4. Example of keyword cannibalization

Search engines will spider the pages on your site and see 4 (or 40) different pages, all seemingly relevant to one particular keyword (in the example in Figure 4-5, the keyword is *snowboards*). For clarity's sake, Google doesn't interpret this as meaning that your site as a whole is more relevant to snowboards or should rank higher than the competition. Instead, it forces Google to choose among the many versions of the page and pick the one it feels best fits the query. When this happens, you lose out on a number of rank-boosting features:

Internal anchor text

Since you're pointing to so many different pages with the same subject, you can't concentrate the value and weight of internal, thematically relevant anchor text on one target.

External links

If four sites link to one of your pages on snowboards, three sites link to another of your snowboard pages, and six sites link to yet another snowboard page, you've split up your external link value among three topically similar pages, rather than consolidating it into one.

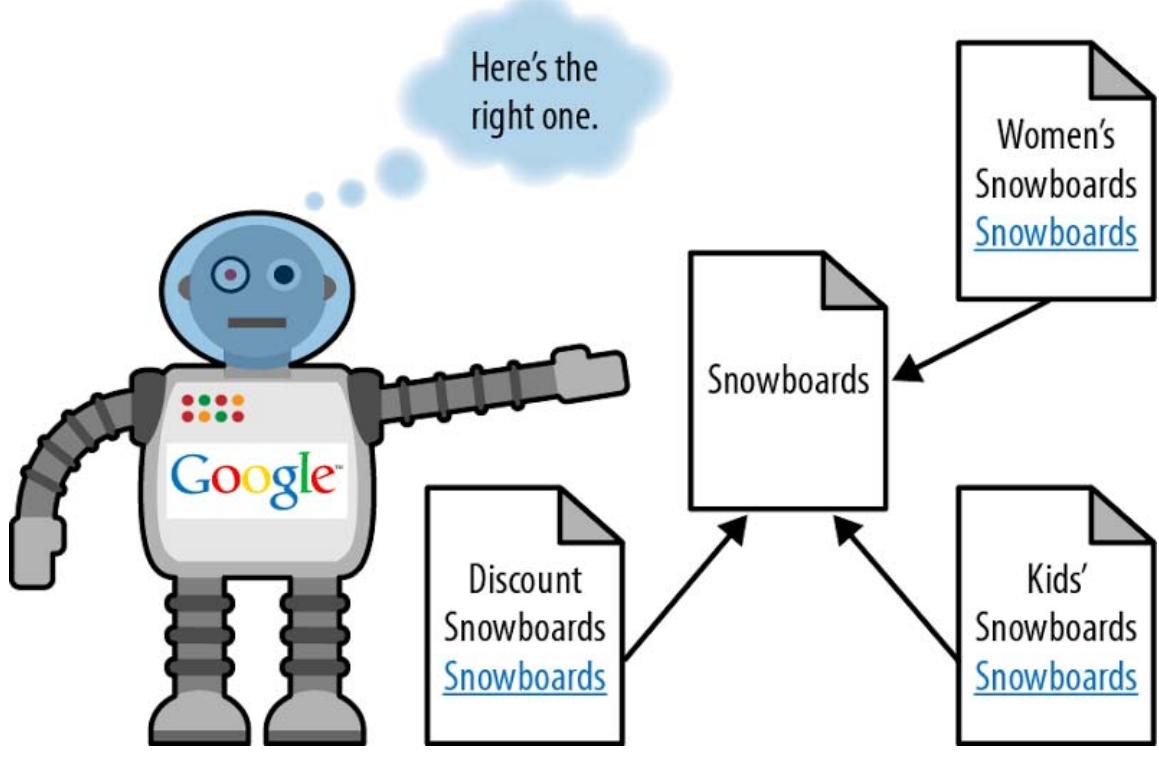
Content quality

After three or four pages of writing about the same primary topic, the value of your content is going to suffer. You want the best possible single page to attract links and referrals, not a dozen bland, repetitive pages.

Conversion rate

If one page is converting better than the others, it is a waste to have multiple lower-converting versions targeting the same traffic. If you want to do conversion tracking, use a multiple-delivery testing system (either A/B or multivariate), such as Optimizely (<http://www.optimizely.com/>).

So, what's the solution? Take a look at Figure 4-6.



[[figs/print/0405.png]]

Figure 4-5. Solution to keyword cannibalization

The difference in this example is that instead of every page targeting the single term *snowboards*, the pages are focused on unique, valuable variations and all of them link back to an original, canonical source for the single term. Google can now easily identify the most relevant page for each of these queries. This isn't just valuable to the search engines; it also represents a far better user experience and overall information architecture.

What should you do if you already have a case of keyword cannibalization? Employ 301s liberally to eliminate pages competing with each other within SERPs, or figure out how to differentiate them. Start by identifying all the pages in the architecture with this issue and determine the best page to point them to, and then use a 301 from each of the problem pages to the page you wish to retain. This ensures not only that visitors arrive at the right page, but also that the link equity and relevance built up over time are directing the engines to the most relevant and highest-ranking-potential page for the query. If 301's aren't practical or feasible, use canonical tags.

Example: Fixing an Internal Linking Problem

Enterprise sites range between 10,000 and 10 million pages in size. For many of these types of sites, an inaccurate distribution of internal link juice is a significant problem. Figure 4-7 shows how this can happen.



▲ Page/section rich in link juice
▼ Page/section poor in link juice

[[figs/print/0406.png]]

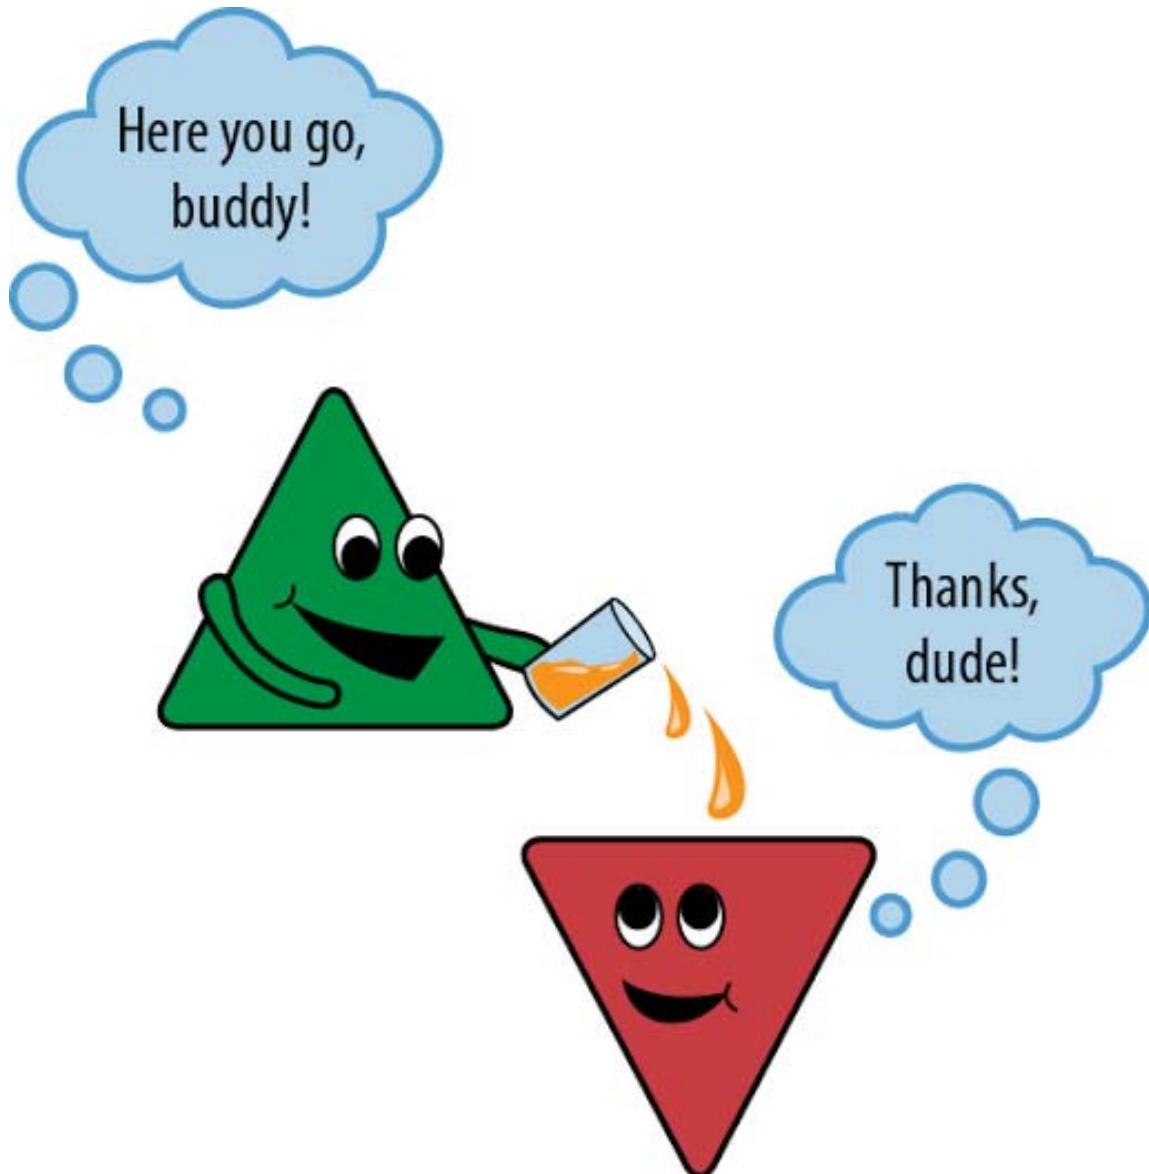
Figure 4-6. Link authority distribution on a very large site

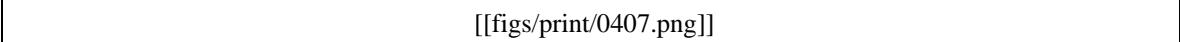
Figure 4-7 is an illustration of the link authority distribution issue. Imagine that each of the tiny pages represents between 5,000 and 100,000 pages in an enterprise site. Some areas, such as blogs, articles, tools, popular news stories, and so on, might be receiving more than their fair share of internal link attention. Other areas—often business-centric and sales-centric content—tend to fall by the wayside. How do you fix this problem? Take a look at Figure 4-8.

The solution is simple, at least in principle: have the link-rich pages spread the wealth to their link-bereft brethren. As easy as this may sound, in execution it can be incredibly complex. Inside the architecture of a site with several hundred thousand or a million pages, it can be nearly impossible to identify link-rich and link-poor pages, never mind adding code that helps to distribute link authority equitably.

The answer, sadly, is labor-intensive from a programming standpoint. Enterprise site owners need to develop systems to track inbound links and/or rankings and build bridges (or, to be more consistent with Figure 4-8, spouts) that funnel authority between the link-rich and link-poor.

An alternative is simply to build a very flat site architecture that relies on relevance or semantic analysis. This strategy is more in line with the search engines' guidelines (though slightly less perfect) and is certainly far less labor-intensive.





[[figs/print/0407.png]]

Figure 4-7. Using cross-links to push link authority where you want it

Server and Hosting Issues

Thankfully, only a handful of server or web hosting dilemmas affect the practice of search engine optimization. However, when overlooked, they can spiral into massive problems, and so are worthy of review. The following are some server and hosting issues that can negatively impact search engine rankings:

Server timeouts

If a search engine makes a page request that isn't served within the bot's time limit (or that produces a server timeout response), your pages may not make it into the index at all, and will almost certainly rank very poorly (as no index-able text content has been found).

Slow response times

Although this is not as damaging as server timeouts, it still presents a potential issue. Not only will crawlers be less likely to wait for your pages to load, but surfers and potential linkers may choose to visit and link to other resources because accessing your site is problematic. Again, Net Neutrality concerns are relevant here.

Shared IP addresses

Basic concerns include speed, the potential for having spammy or untrusted neighbors sharing your IP address, and potential concerns about receiving the full benefit of links to your IP address (discussed in more detail at <http://www.seroundtable.com/archives/002358.html>).

Blocked IP addresses

As search engines crawl the Web, they frequently find entire blocks of IP addresses filled with nothing but egregious web spam. Rather than blocking each individual site, engines do occasionally take the added measure of blocking an IP address or even an IP range. If you're concerned, search for your IP address at Bing using the `ip:address` query.

Bot detection and handling

Some sys admins will go a bit overboard with protection and restrict access to files to any single visitor making more than a certain number of requests in a given time frame. This can be disastrous for search engine traffic, as it will constantly limit the spiders' crawling ability.

Bandwidth and transfer limitations

Many servers have set limitations on the amount of traffic that can run through to the site. This can be potentially disastrous when content on your site becomes very popular and your host shuts off access. Not only are potential linkers prevented from seeing (and thus linking to) your work, but search engines are also cut off from spidering.

Server geography

While the search engines of old did use the location of the web server when determining where a site's content is, Google specifically makes it clear that in today's search environment, actual server location is, for the most part, irrelevant. According to Google, if a site is using a ccTLD or gTLD in conjunction with Webmaster Tools to set geo-location information for the site, then the location of the server itself becomes immaterial. There is one caveat, and that is the fact that content hosted closer to end users tends to be delivered more quickly, and speed of content delivery is taken into account by Google – impacting mobile search significantly.

Identifying Current Server Statistics Software and Gaining Access

In Chapter 11, we will discuss in detail the methods for tracking results and measuring success, and we will also delve into how to set a baseline of measurements for your SEO projects. But before we do that, and before you can accomplish these tasks, you need to have the right tracking and measurement systems in place.

Web Analytics

Analytics software can provide you with a rich array of valuable data about what is taking place on your site. It can answer questions such as:

- How many unique visitors did you receive yesterday?
- Is traffic trending up or down?
- What site content is attracting the most visitors from organic search?
- What are the best-converting pages on the site?

We strongly recommend that if your site does not currently have any measurement systems in place, you put something in place immediately. High-quality, free analytics tools are available, such as the powerful and robust Google Analytics (<http://www.google.com/analytics/>), as well as the open source platform Piwik (<http://www.piwik.com>). Of course, higher-end analytics solutions are also available, which we will discuss in greater detail in Chapter 11.

Logfile Tracking

Logfiles contain a detailed click-by-click history of all requests to your web server. Make sure you have access to the logfiles and some method for analyzing them. If you use a third-party hosting company for your site, chances are it provides some sort of free logfile analyzer, such as AWStats (<http://awstats.sourceforge.net>), Webalizer (<http://www.webalizer.com>), or something similar (the aforementioned Piwik also provides a log analytics tool, as well). Identify, and obtain access to, whatever tool is in use as soon as you can.

One essential function that these tools provide that JavaScript-based web analytics software cannot record is search engine spider activity on your site. Although spidering will typically vary greatly from day to day, you can still see longer-term trends of search engine crawling patterns, and whether crawling activity is trending up (good) or down (bad). Although this web crawler data is very valuable, do not rely on these free solutions provided by hosting companies for all of your analytics data, as there is a lot of value in what traditional analytics tools can capture.

Some web analytics software packages read log files as well, and therefore can report on crawling activity. We will discuss these in more detail in Chapter 11.

Google and Bing Webmaster Tools

As mentioned earlier, other valuable sources of data include Google Webmaster Tools (<http://www.google.com/webmasters/tools>) and Bing Webmaster Tools (<http://www.bing.com/webmaster>). We cover these extensively in “Using Search Engine–Supplied SEO Tools” in Chapter 13.

From a planning perspective, you will want to get these tools in place as soon as possible. Both tools provide valuable insight into how the search engines see your site, including things such as external link data, internal link data, crawl errors, high-volume search terms, and much, much more. In addition, these tools provide important functionality for interacting directly with the search engine via sitemap

submissions, URL and sitelink removals, disavowing bad links, setting parameter handling preferences, tagging structured data, and more – enabling you to have more control over how the search engines view, crawl, index, and output your page content within their index, and within search results.

Determining Top Competitors

Understanding the competition should be a key component of planning your SEO strategy. The first step is to understand who your competitors in the search results really are. It can often be small players who give you a run for your money. For example, consider the previously mentioned *credit card* search in Google (Figure 4-1), whereby all fail to reach the #1 position in the Google results. Instead, mostly banks and affiliate players dominate the SERPs for this query.

Two Spam Examples

Affiliates that cheat tend to come and go out of the top search results, as only sites that implement ethical tactics are likely to maintain their positions over time.

How do you know whether a top-ranking site is playing by the rules? Look for dubious links to the site using a backlink analysis tool such as Majestic SEO or Open Site Explorer (discussed earlier in this chapter). Since the number of links is one factor search engines use to determine search position, less ethical websites will obtain links from a multitude of irrelevant and low-quality sites.

This sort of sleuthing can reveal some surprises. For instance, here are examples of two devious link schemes:

- GiftCertificates.com's short-lived nemesis was FindGiftCards.com (<http://findgiftcards.com>), which came out of nowhere to command the top two spots in Google for the all-important search term *gift certificates*, thus relegating GiftCertificates.com (<http://gftcertificates.com>) to the third position. How did FindGiftCards.com (<http://findgiftcards.com>) do it? It operated a sister site, 123counters.com (<http://123counters.com>), with a free hit counter that propagated "link spam" across thousands of sites, all linking back to FindGiftCards.com (<http://findgiftcards.com>) and other sites in its network.

Sadly for FindGiftCards.com (<http://findgiftcards.com>), co-author Stephan Spencer discussed the company's tactics in an article he wrote for *Multichannel Merchant* (http://multichannelmerchant.com/catalogage/ar/marketing_casing_competition/), and Google became aware of the scam. The end result? Down to two pages in the Google index, as shown in Figure 4-9.

- CraigPadoa.com (<http://craigpadoa.com>) was a thorn in the side of SharperImage.com (<http://sharperimage.com>), outranking the latter for its most popular product, the Ionic Breeze, by frameset trickery and guestbook spamming (in other words, defacing vulnerable websites with fake guestbook entries that contained spammy links back to its own site). As soon as The Sharper Image realized what was happening, it jumped on the wayward affiliate. It also restricted such practices in its affiliate agreement and stepped up its monitoring for these spam practices.



[[figs/print/0408.png]]

Figure 4-8. Site with only two pages in the index

Seeking the Best

Look for competitors whose efforts you would like to emulate (or “embrace and extend,” as Bill Gates would put it)—usually a website that consistently dominates the upper half of the first page of search results in the search engines for a range of important keywords that are popular and relevant to your target audience.

Note that your “mentor” competitors shouldn’t just be good performers, they should also demonstrate that they know what they’re doing when it comes to SEO. To assess competitors’ competence at SEO, you need to answer the following questions:

- Are their websites fully indexed by Google and Bing? In other words, are all their web pages, including product pages, making it into the search engines’ databases? You can go to each search engine and type in `site:theirdomain.com` to find out. A competitor with only a small percentage of its site indexed in Google probably has a site that is unfriendly to search spiders.
- Do their product and category pages have keyword-rich page titles (title tags) that are unique to each page? You can easily review an entire site’s page titles within Google or Bing by searching for `site:www.yourcompetitor.com`.

Incidentally, this type of search can sometimes yield confidential information. A lot of webmasters do not realize that Google has discovered and indexed commercially sensitive content buried deep in their sites. For example, a Google search for `business plan confidential filetype:doc` can yield a lot of real business plans among the sample templates.

- Do their product and category pages have reasonably high PageRank scores?
- Is anchor text across the site, particularly in the navigation, descriptive but not keyword-stuffed?
- Are the websites getting penalized? You can overdo SEO. Too much keyword repetition or too many suspiciously well-optimized text links can yield a penalty for over-optimization. Sites can also be penalized for extensive amounts of duplicate content, for lack of high quality content (referred to as “thin” content), and for participating in guest blog/article writing schemes. You can learn more about how to identify search engine penalties in Chapter 9 (“Panda, Penguin, and Penalties”).

Uncovering Their Secrets

Let’s assume your investigation has led you to identify several competitors who are gaining excellent search placement using legitimate, intelligent tactics. Now it is time to identify their strategy and tactics:

- What keywords are they targeting? You can determine this by looking at the page titles (up in the bar above the address bar at the top of your web browser, which also appears in the search results listings) of each competitor’s home page and product category pages. You can also use various online tools to see what keywords they may be targeting with PPC (pay per click) advertising – while not always an indication that they are investing in SEO, you can still get a solid grasp on their overall keyword strategy.
- Who’s linking to their home page, or to their top-selling product pages and category pages? A link popularity checker can be quite helpful in analyzing this.
- If it is a database-driven site, what technology workarounds are they using to get search engine spiders such as Googlebot to cope with the site being dynamic? Nearly all the technology workarounds are tied to the ecommerce platforms the competitors are running. You can check to see whether they are using the same server software as you by using the “What’s that site running?” tool at http://toolbar.netcraft.com/site_report/ Figure 4-10 shows a screenshot of a segment of the results for HSN.com (<http://hsn.com>):

Server-Side Includes all the main technologies that Netcraft detects as running on the server such as PHP.		
Technology	Description	Popular sites using this technology
Atlas 	A set of ASP.NET extensions for implementing Ajax functionality	www.ighome.com , www.ldlc.com , www.geocaching.com
Using ASP.NET 		
Client-Side Includes all the main technologies that run on the browser (such as JavaScript and Adobe Flash).		
Technology	Description	Popular sites using this technology
Image Map	No description	www.ebay.it , www.amazon.com , www.leboncoin.fr
JavaScript 	Open source programming language commonly implemented as part of a web browser	www.imdb.com , mail.google.com
Client-Side Scripting Frameworks Frameworks or libraries allow for easier development of applications by providing an Application Program Interface (API) or a methodology to follow whilst developing.		
Technology	Description	Popular sites using this technology
Modernizr	No description	www.samsung.com , www.wunderground.com , www.gutefrage.net
Mobile Technologies Mobile technology is the technology used for hand held mobile devices.		
Technology	Description	Popular sites using this technology
Click to call 	Markup language syntax intended for devices that can place calls (e.g. phones, VoIP, etc.)	www.eprice.it , www.istockphoto.com , www.nic.ca
Character Encoding A character encoding system consists of a code that pairs each character from a given repertoire with something else such as a bit pattern, sequence of natural numbers, octets, or electrical pulses in order to facilitate the transmission of data (generally numbers or text) through telecommunication networks or for data storage.		
Technology	Description	Popular sites using this technology
UTF8 	UCS Transformation Format 8 bit	www.googleadservices.com
HTTP Compression HTTP compression is a capability that can be built into web servers and web clients to make better use of available bandwidth, and provide greater transmission speeds between both.		
Technology	Description	Popular sites using this technology
Gzip Content Encoding 	Gzip HTTP Compression protocol	www.lifattoquotidiano.it , www.programme-tv.net , www.zerohedge.com

[[figs/print/0409.png]]

Figure 4-09. Sample Netcraft output

While you are at it, look at “cached” (archived) versions of your competitors’ pages by clicking on the Cached link next to their search results in Google to see whether they’re doing anything too aggressive, such as *cloaking*, where they serve up a different version of the page to search engine spiders than to human visitors. The cached page will show you what the search engine actually saw, and you can see how it differs from the page you see when you go to the web page yourself.

- What effect will their future SEO initiatives have on their site traffic? Assess the success of their SEO not just by the lift in rankings. Periodically record key SEO metrics over time—the number of pages indexed, the PageRank score, the number of links—and watch the resulting effect on their site traffic. If you utilize one of the many web-based SEO tool platforms (read more about these in Chapter 11), you can set your competitors’ sites up as additional sites or campaigns to track.

You do not need access to competitors’ analytics data or server logs to get an idea of how much traffic they are getting. Simply go to Compete.com (<http://www.compete.com/>), Quantcast.com (<http://www.quantcast.com/>), Search Metrics (<http://searchmetrics.com/>), or SEMRush (<http://www.semrush.com/>) and search on the competitor’s domain. If you have the budget for higher-end competitive intelligence tools, you can use comScore.com (<http://comscore.com>) or Experian’s Hitwise.com (<http://www.experian.com/hitwise/>).

The data these tools can provide is limited in its accuracy, and can often be unavailable if the site in question receives too little traffic - but they are still very useful in giving you a general assessment of where your competitors are. The tools are most useful when making relative comparisons between sites in the same market space. To get an even better idea of where you stand, use their capabilities to compare the traffic of multiple sites, including yours, to get a pretty accurate idea as to how your traffic compares to theirs.

- How does the current state of their sites’ SEO compare with those of years past? You can reach back into history and access previous versions of your competitors’ home pages and view the HTML source to see which optimization tactics they were employing back then. The Wayback Machine (<https://archive.org/>) provides an amazingly extensive archive of web pages.

Assessing Historical Progress

Measuring the results of SEO changes can be challenging, partly because there are so many moving parts and partly because months can elapse between when changes are made to a site and when results are seen in organic search exposure and traffic. This difficulty only increases the importance of measuring progress and being accountable for results. This section will explore methods for measuring the results from your SEO efforts.

Maintain a Timeline of Site Changes

Keeping a log of changes to your site is absolutely recommended. If you're not keeping a timeline (which could be as simple as an online spreadsheet or as complex as a professional project management visual flowchart), you will have a harder time executing your SEO plan and managing the overall SEO process. Sure, without one you can still gauge the immediate effects of content additions/revisions, link acquisitions, and development changes, but visibility into how technical modifications to the website might have altered the course of search traffic, whether positively or negatively, is obscured.

One of the most useful ways to document this information is by utilizing Annotations in Google Analytics, whereby you can document specific events along your site timeline for future reference – and for ongoing visual association with your analytics data. Figure 4-10 below shows how annotations appear along your timeline graph in Google Analytics reports; note the text balloons along the date timeline (each text balloon can then be clicked on individually, to display the notes added at that time):

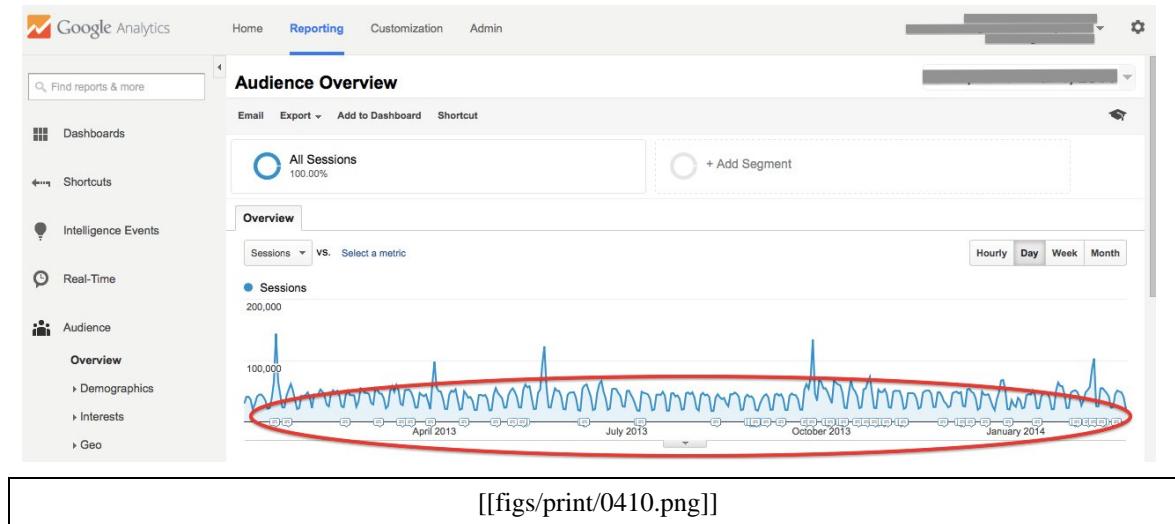


Figure 4-10. Annotations in Google Analytics

If you don't record and track site changes—both those intended to influence SEO, and those for which SEO wasn't even a consideration—you will be eventually optimizing blind and could miss powerful signals that could help dictate your strategy going forward. Ideally, you should track more than just site changes as well, as external factors that can have a big impact on your SEO results can be just as important - they include confirmed search engine algorithm updates, competitor news events (e.g., product or company launches), breaking news, and trending topics in social media. Factors within your own business can have an impact as well, such as major marketing or PR events, IPOs, or the release of earnings statements.

There are many scenarios in which you will want to try to establish cause and effect, such as:

If search traffic spikes or plummets

Sudden changes in organic traffic are obviously notable events. If traffic plummets, you will be facing lots of questions about why, and having a log of site changes will put you in a better position to assess whether any changes you recommended could have been the cause. Of course, if traffic spikes you will want to be able to see whether an SEO-related change was responsible as well.

When gradual traffic changes begin

Changes do not always come as sudden spikes or drop-offs. If you see the traffic beginning a gradual climb (or descent), you will want to be able to assess the likely reasons.

To track and report SEO progress

Accountability is a key component of SEO. Budget managers will want to know what return they are getting on their SEO investment. This will inevitably fall into two buckets: itemizing specific work items worked on, and analyzing benefits to the business. Keeping an ongoing change log makes tracking and reporting SEO progress much easier to accomplish.

Creation of new, noteworthy content onsite and/or tied to social media campaigns

Your social media efforts can directly impact your site's performance, both directly (traffic from third party social media platforms) to indirectly via links, social sharing, search query trending, and other factors. All of this can be tied back via annotations in Google Analytics for meaningful data analysis.

Types of Site Changes That Can Affect SEO

Your log should track all changes to the website, not just those that were made with SEO in mind. Organizations make many changes that they do not think will affect SEO, but that have a big impact on it. Here are some examples:

- Adding content areas/features/options to the site (this could be anything from a new blog to a new categorization system).
- Changing the domain of the site. This can have a significant impact, and you should document when the switchover was made.
- Modifying URL structures. Changes to URLs on your site will likely impact your rankings, so record any and all changes.
- Implementing a new CMS. This is a big one, with a very big impact. If you must change your CMS, make sure you do a thorough analysis of the SEO shortcomings of the new CMS versus the old one, and make sure you track the timing and the impact.
- Establishing new partnerships that either send links or require them (meaning your site is earning new links or linking out to new places).
- Acquiring new links to pages on the site other than the home page (referred to as "deep links").
- Making changes to navigation/menu systems (moving links around on pages, creating new link systems, etc.).
- Implementing redirects either to or from the site.
- New/updated sitemaps, canonical tag implementations, schema markup usage, etc.

When you track these items, you can create an accurate storyline to help correlate causes with effects. If, for example, you've observed a spike in traffic from Bing that started four to five days after you switched your menu links from the page footer to the header, it is likely that there is a correlation; further analysis would determine whether, and to what extent, there is causation.

Without such documentation it could be months before you noticed the impact—and there would be no way to trace it back to the responsible modification. Your design team might later choose to switch back to footer links, your traffic might fall, and no record would exist to help you understand why. Without the lessons of history, you cannot leverage the positive influencing factors.

Previous SEO Work

When you are brought on to handle the SEO for a particular website, one of the first things you need to find out is which SEO activities have previously been attempted. There may be valuable data there, such as a log of changes that you can match up with analytics data to gauge impact.

If no such log exists, you can always check the Wayback Machine (<https://archive.org/>) to see whether it has historical logs for your website. This offers snapshots of what the site looked like at various points in time.

Even if a log was not kept, spend some time building a timeline of when any of the types of changes that affect SEO (as discussed in the previous section) took place. In particular, see whether you can get copies of the exact recommendations the prior SEO consultant made, as this will help you with the timeline and the specifics of the changes made.

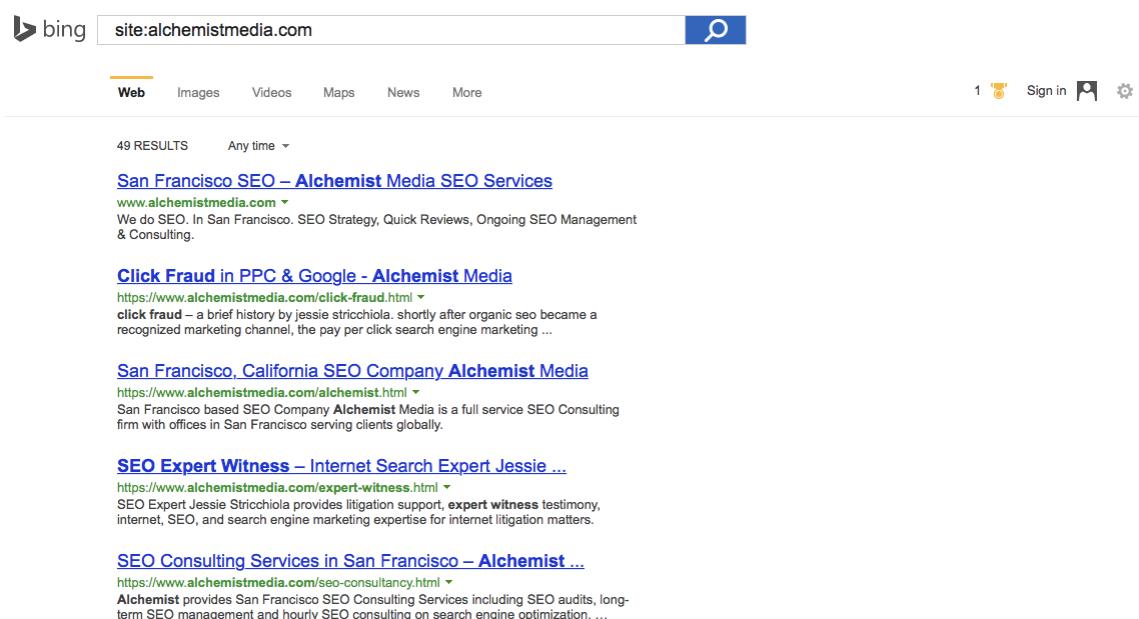
You should also pay particular attention to understanding the types of link-building activities that took place. Were shady practices used that carry a lot of risk? Was there a particular link-building tactic that worked quite well? Going through the history of the link-building efforts can yield tons of information that you can use to determine your next steps. Google undoubtedly keeps a 'rap sheet' on us webmasters/SEO practitioners.

Benchmarking Current Indexing Status

The search engines have an enormous task: that of indexing the world's online content—well, more or less. The reality is that they try hard to discover all of it, but they do not choose to include all of it in their indexes. There can be a variety of reasons for this, such as the page being inaccessible to the spider, being penalized, or not having enough link juice to merit inclusion.

When you launch a new site or add new sections to an existing site, or if you are dealing with a very large site, not every page will necessarily make it into the index. You will want to actively track the indexing level of your site in order to develop and maintain a site with the highest accessibility and crawl efficiency. If your site is not fully indexed, it could be a sign of a problem (not enough links, poor site structure, etc.).

Getting basic indexation data from search engines is pretty easy. The major search engines support the same basic syntax for that: `site:yourdomain.com`. Figure 4-12 shows a sample of the output from Bing.



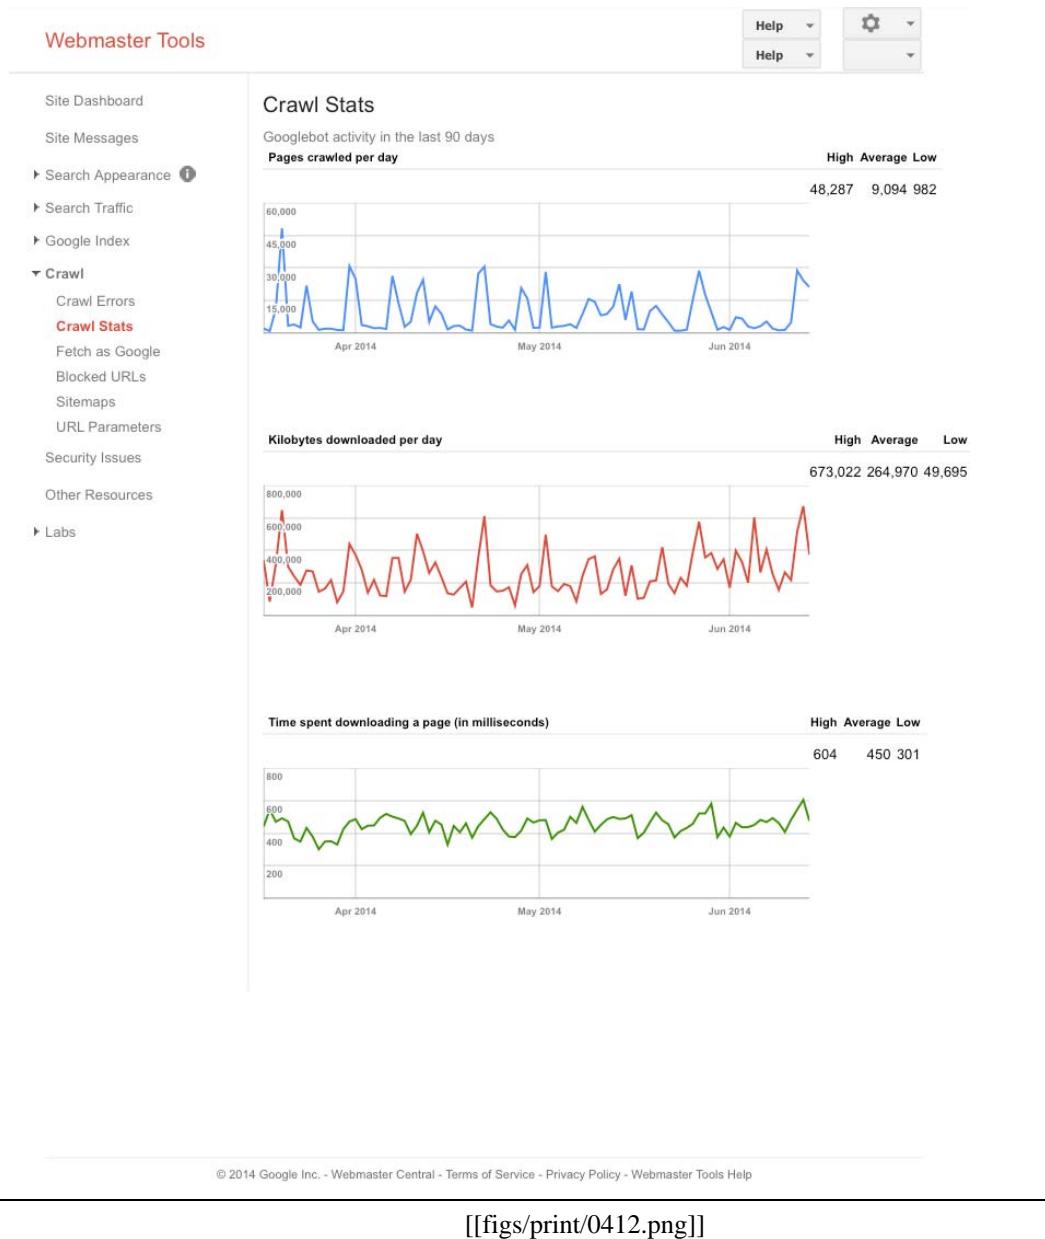
The screenshot shows a Bing search results page. The search bar at the top contains the query `site:alchemistmedia.com`. Below the search bar is a navigation bar with links for Web, Images, Videos, Maps, News, and More. The results section shows 49 results for the query. The first result is a link to [San Francisco SEO – Alchemist Media SEO Services](http://www.alchemistmedia.com), with a snippet: "www.alchemistmedia.com ▾ We do SEO. In San Francisco. SEO Strategy, Quick Reviews, Ongoing SEO Management & Consulting." The second result is a link to [Click Fraud in PPC & Google - Alchemist Media](http://www.alchemistmedia.com/click-fraud.html), with a snippet: "https://www.alchemistmedia.com/click-fraud.html ▾ click fraud – a brief history by jessie stricchiola. shortly after organic seo became a recognized marketing channel, the pay per click search engine marketing ...". The third result is a link to [San Francisco, California SEO Company Alchemist Media](http://www.alchemistmedia.com/alchemist.html), with a snippet: "https://www.alchemistmedia.com/alchemist.html ▾ San Francisco based SEO Company Alchemist Media is a full service SEO Consulting firm with offices in San Francisco serving clients globally." The fourth result is a link to [SEO Expert Witness – Internet Search Expert Jessie ...](http://www.alchemistmedia.com/expert-witness.html), with a snippet: "https://www.alchemistmedia.com/expert-witness.html ▾ SEO Expert Jessie Stricchiola provides litigation support, **expert witness** testimony, internet, SEO, and search engine marketing expertise for internet litigation matters." The fifth result is a link to [SEO Consulting Services in San Francisco – Alchemist ...](http://www.alchemistmedia.com/seo-consultancy.html), with a snippet: "https://www.alchemistmedia.com/seo-consultancy.html ▾ Alchemist provides San Francisco SEO Consulting Services including SEO audits, long-term SEO management and hourly SEO consulting on search engine optimization. ...".

[[figs/print/0411.png]]

Figure 4-11. Indexing data from Bing

Keeping a log of the level of indexation over time can help you understand how things are progressing, and this information can be tracked in a spreadsheet.

Related to indexation is the crawl rate of the site. Google and Bing provide this data in their respective Webmaster Tools. Figure 4-13 shows a screenshot representative of the crawl rate charts that are available in Google Webmaster Tools.



[[figs/print/0412.png]]

Figure 4-12. Crawl data from Google Webmaster Tools

Short-term spikes are not a cause for concern, nor are periodic drops in levels of crawling. What is important is the general trend. Bing Webmaster Tools provides similar data to webmasters, and for other search engines their crawl-related data can be revealed using logfile analyzers mentioned previously in this chapter – and then a similar timeline can be created and monitored.

Benchmarking Organic Rankings

People really love to check their search rankings. Many companies want to use this as a measurement of SEO progress over time, but it is a bit problematic, for a variety of reasons. Here is a summary of the major problems with rank checking:

- Google results are not consistent:
 - Different geographies (even in different cities within the United States) often give different results.
 - Results are personalized for logged-in users based on their search histories, and personalized for non-logged in users as well based on their navigational behaviors within the SERPs, browser history, location data, and more.
 - No rank checker can monitor and report all of these inconsistencies (at least, not without scraping Google hundreds of times from all over the world with every possible setting).
- Obsessing over rankings (rather than traffic) can result in poor strategic decisions:
 - When site owners obsess over rankings for particular keywords, the time and energy they expend on those few keywords and phrases often produces far less value than would have been produced if they had spent those resources on the site as a whole.
 - Long-tail traffic very often accounts for 70% to 80% of the demand curve, and it is much easier to rank in the long tail and get valuable traffic from there than it is to concentrate on the few rankings at the top of the demand curve.

So, indulge your desire to check rankings by going to the search engine and typing in a few queries or use a tool like AuthorityLabs (<http://authoritylabs.com/>) and automate it, but be sure to also keep an ongoing focus on your visitor and conversion statistics.

Benchmarking Current Traffic Sources and Volume

The most fundamental objective of any SEO project should be to drive the bottom line. For a business, this means delivering more revenue with favorable ROI. As a precursor to determining the level of ROI impact, the SEO pro must focus on increasing the volume of relevant traffic to the site. This is a more important objective than anything related to rankings or number of links obtained. More relevant traffic should mean more revenue for the business (or more conversions, for those whose websites are not specifically selling something).

Today's web analytics tools make the gathering of such data incredibly easy. Google Analytics offers a robust set of tools sufficient for many smaller to medium sites, while larger sites will probably need to use the paid version of Google Analytics (Google Analytics Premium, <http://www.google.com/analytics/premium/>) in lieu of or in conjunction with other paid solutions such as Omniture.com (<http://www.omniture.com>), IBM Experience One (<http://www-01.ibm.com/software/marketing-solutions/experienceone/>), or Webtrends.com (<http://webtrends.com>).

The number of data points you can examine in analytics is nearly endless. It is fair to say that there is too much data, and one of the key things that an SEO expert needs to learn is what data is worth looking at and what data is not. Generally speaking, the creation of custom Dashboards in Google Analytics can greatly reduce the steps needed to see top level data points for SEO; in Chapter 10, we will discuss in greater detail how to leverage the power of analytics.

Leveraging Business Assets for SEO

Chances are your company/organization has numerous digital media assets beyond the website that can be put to good use to improve the quality and quantity of traffic you receive through search engine optimization efforts. We discuss some of these things in the subsections that follow.

Other Domains You Own/Control

If you have multiple domains, the major items to think about are:

- Can you 301-redirect some of those domains back to your main domain or to a subfolder on the site for additional benefit? Be sure to check the domain health in Webmaster Tools before performing any 301 redirects, as you want to ensure the domain in question has no penalties before closely associating it with your main domain(s).
- If you're maintaining those domains as separate sites, are you linking between them intelligently?

If any of those avenues produce valuable strategies, pursue them—remember that it is often far easier to optimize what you're already doing than to develop entirely new strategies, content, and processes.

Relationships On and Off the Web

Business relationships and partnerships can be leveraged in similar ways, particularly on the link development front. If you have business partners that you supply or otherwise work with—or from whom you receive service—chances are good that you can safely create legitimate, user-friendly linking relationships between your sites. Although reciprocal linking carries a bit of a bad reputation, there is nothing wrong with building a “partners,” “clients,” “suppliers,” or “recommended” list on your site, or with requesting that your organizational brethren do likewise for you. Just do this in moderation and make sure you link only to highly relevant, trusted sites.

Content or Data You've Never Put Online

Chances are that you have content that you have never published on your website. This content can be immensely valuable to your SEO efforts. However, many companies are not savvy to the nuances of publishing that content in a manner that is friendly to search engines. Those hundreds of lengthy articles you published when you were shipping a print publication via the mail could be a great fit for your website archives. You should take all of your email newsletters and make them accessible on your site. If you have unique data sets or written material, you should apply it to relevant pages on your site (or consider building out if nothing yet exists).

Customers Who Have Had a Positive Experience

Customers are a terrific resource for earning links, but did you also know they can write? Customers and website visitors can contribute all kinds of content. Seriously, if you have user-generated content (UGC) options available to you and you see value in the content your users produce, by all means reach out to customers, visitors, and email list subscribers for content opportunities.

Followers, Friends, and Fans

For many businesses that operate offline or work in entertainment, hard goods, or any consumer services, chances are good that if your business is respected, there are people out there who've used your products or services and would love to share their experiences. Do you make video games? Hopefully you've established a presence on social media – and can leverage third party platforms (Facebook, Pinterest, Instagram, Twitter, Vine, and others) to market your content for link acquisition, content curation positive testimonials, and brand buzz.

Conclusion

The first steps of SEO are often the most challenging ones. It's tempting to launch into the effort just to get things moving – but spending some focused time properly planning your SEO strategy before implementation will pay big dividends in the long run. Establish a strong foundation, and you will have a firm launching pad for SEO success.

5

Keyword Research

The Theory Behind Keyword Research

Keyword research is one of the most important, valuable, and high-return activities in the search engine marketing field. Through the detective work of dissecting your market's keyword demand, you learn not only which terms and phrases to target with SEO, but also more about your customers as a whole.

With keyword research you can predict shifts in demand, respond to changing market conditions, and produce the products, services, and content that web searchers are already actively seeking. In the history of marketing, there has never been such a low barrier to entry in understanding the motivations of consumers in virtually every niche.

Every search phrase that's typed into an engine is recorded in one way or another, and keyword research tools such as the ones we discuss in this chapter allow you to retrieve this information. However, those tools cannot show you (directly) how valuable or important it might be to rank for and receive traffic from those searches. This chapter seeks to expose the details of this process and the tools that can best assist.

Start by Thinking Strategically

Keyword research tools provide valuable insight into the thinking of your potential customers. When users go to search engines and type out their search query they may use language that is entirely different from what you expect. Even if your product or service provides a solution they can use they may start with their problem. Someone with diabetes might simply type in *diabetes* in the search box. Then their next search might be *diabetes medication* or *relief for diabetes symptoms*.

As we laid out in Chapter 1, searchers often go through a progression where they try certain searches, check out some sites, refine their searches, and repeat this process until they finally find what they want. Taking the time to understand typical search sequences is one aspect that impacts your keyword strategy.

Other elements that influence search behavior include searcher demographics (male/female, age, income, etc.), geographical location, and time of year. Seasonal products such as Valentine's Day cards, for example, go through sharp peaks in volume during the time preceding February 14th, and then decline rapidly once the season is past.

The keyword research tools in this Chapter will provide you with methods to investigate all of these factors. Take the time to go beyond the surface and use the tools to learn how your customer thinks, get your thinking in alignment with theirs, and then build your web site strategy (and perhaps even your product strategy) around this powerful information.

Understanding the Long Tail of the Keyword Demand Curve

It is wonderful to deal with keywords that have 5,000 searches per day, or even 500 searches per day, but in reality these “popular” search terms may actually comprise less than 30% of the overall searches performed on the Web. The remaining 70% lie in what’s commonly called the “long tail” of search (as published at <http://moz.com/blog/rewriting-the-beginners-guide-part-v-keyword-research>); see Figure 5-1. The tail contains hundreds of millions of unique searches that might be conducted a few times in any given day, or even only once ever, but when assessed in aggregate they comprise the majority of the world’s demand for information through search engines.

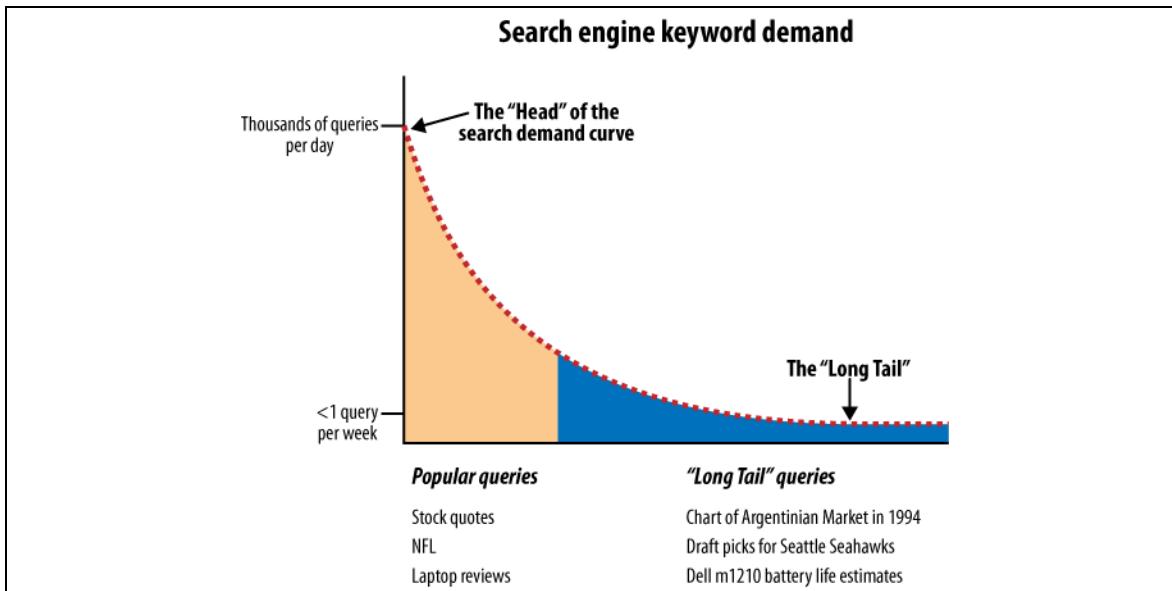


Figure 5-1. Long tail of search

Impact of Google Hummingbird

In September 2013, Google announced, to coincide with its 15th birthday, a major change to their search algorithm called “Hummingbird” (<http://insidesearch.blogspot.com/2013/09/fifteen-years-onand-were-just-getting.html>). Since then, although Hummingbird has been much discussed, Google has not released any more official documentation on this update. Hummingbird is a major change to the way Google interprets searches and in how searches relate to one another.

Hummingbird is intended to get at the heart of what the user wants, not just the exact keywords they search for. In large part, this is related to Google getting more prepared for mobile search. In mobile searches, users are less apt to type traditionally formatted queries, and in many cases, actually do voice searches resulting in conversational queries. When users use voice search, their queries may also be much more conversational in format, such as *please find me the closest gas station*. In addition, with Hummingbird, Google will use many other factors to determine the intent of the user, such as considering previous related searches by that user.

The Hummingbird algorithm attempts to determine the true meaning behind what a user is searching for, rather than simply returning results for the exact query they use. Indeed, in many cases, Google may simply relate the terms and consider them synonyms for the purposes of returning search results.

Let’s look at this hypothetical, yet fully-functional, example of a series of queries that a mobile user might ask Google:

“where is the empire state building?”

“who built it?”

“pictures”

“how tall is it?”

“restaurants”

“Italian restaurants”

“show me the second one”

“directions”

No longer is Google’s response to the query “where is the empire state building?” simply the web page that is most closely optimized to that search query. Google uses its knowledge of entities, which include notable buildings and monuments and their attributes and locations, nearby restaurants, average review ratings of those restaurants (partially from semantic markup), and so forth, to return meaningful responses to these queries, including actual Italian restaurant recommendations near the Empire State Building, rather than merely a match on the keywords searched.

It’s worth noting in the above example the presence of one-word queries that Google expands (interprets) into longer tail queries. This changes the game in terms of chasing after “long tail” keywords. Now like never before, we must be like detectives who anticipate the searcher’s intent and next move before they even arrive on our website and be ready to meet their needs based on that intent, without the keywords to guide us.

In terms of Hummingbird’s impact on keyword research, it’s also important to consider co-occurrence analysis. While there is some debate in the search engine optimization marketing discipline as to the impact of various types of co-occurrence there is strong evidence that the co-occurrence of keywords in on-page content can affect search engine rankings. When it comes to keyword research co-occurrence refers to an analysis of what words and phrases appear most commonly on a page, how often they appear, and how they relate to each other.

Co-author Eric Enge gives a summary of a study by Searchmetrics (<http://www.searchmetrics.com/>) in his post “How Content Quality Analysis Works With SEO” (<http://searchengineland.com/content-quality-analysis-seo-194314>). In the example he shows how on a sample page the terms: women, shoes, boots, shoe, etc. co-occur on an e-commerce page with words like “view”, “sale”, “cart”, “shipping”, and other e-commerce related terms. This sends a signal to the search engine that this should be a page where users can shop for and buy women’s footwear. The phrase “buy women’s shoes” may never actually appear on the page and yet due to the co-occurrence of the product type terms and e-commerce terms Google should understand that the “theme” or “purpose” of the page is to offer a place where users can buy women’s shoes and should potentially rank this page for that term (depending on the myriad of other quality signal surrounding the page). If an e-commerce page lacks those words and phrases which help to signal that the intent of the page is to be a place where products may be purchased it might lower the chances of that page ranking for purchase-oriented search phrases.

There is a related type of keyword co-occurrence which may affect rankings similarly, that of co-occurrence across various search engine queries during a session. In this type of co-occurrence the engines would look at how similar groups of keywords are searched across a users’ search session. Following the example above, if a large number of queries for “women’s shoes” co-occurred during users’ search sessions with a query for “buy women’s shoes online” Google should know that it is likely that users searching for “women’s shoes” are looking for where to buy women’s shoes online. In this instance it may be that sites or pages which are well optimized for the term “women’s shoes” would also rank well for “buy women’s shoes online” without that exact phrase having to appear in the on page copy or metadata for the page. In this case Google would extrapolate that the user’s actual intent was to find a place to buy women’s shoes online, even though the query may simply have been “women’s shoes”.

Ultimately, in both cases, the goal is for the search engines to understand the user’s intent and to return a set of results which satisfy that intent, regardless of whether the exact phrase searched for explicitly appears on the pages which are returned as results for a given query.

Hummingbird is neither about shortening or lengthening queries, it's a complete rewrite of the search platform. You can read about this in detail at <http://www.wordtracker.com/blog/a-hummingbird-is-not-just-for-christmas>. You can also watch this hour-long video hangout with Danny Sullivan on Hummingbird at <https://www.youtube.com/watch?v=MdzQKxcwtNk>.

Hummingbird was a platform rewrite, much like Caffeine. Caffeine rewrote the crawler and indexing parts of Google. Hummingbird rewrote the actual search engine piece (the part that receives user search queries and then pulls results out of the index).

Marcus Tober of Searchmetrics (<http://www.searchmetrics.com/>) presented some results of their study into Google Hummingbird at the SMX West search marketing show in the spring of 2014 (<http://www.marketing-mojo.com/blog/hummingbird-entity-search-revolution-smx-west-2014/>). Among his findings were that the context of the search (not just the exact query), and previous searches matter, that Hummingbird attempts to understand conversational search, and that Google is trying to understand entities (people, places, things, etc.) and how they relate to each other. Searchmetrics found that Google is now delivering search results for many related queries as if they were the same search query.

What this means from a search engine optimization perspective is that the exact query a user may be searching for is less important than the intent behind it. While keyword research is still crucial, creating pages highly optimized to a specific keyword is less important than creating extremely high quality unique content which answers the need or question which lies behind the keyword query.

When doing your keyword research be flexible and open to using tools which show you "related queries" to those you are specifically looking at. Also watch for "informational" queries where it is apparent that the user is attempting to solve a problem. Much of your keyword research should be centered around the true meaning behind a query and why a user searches on those specific terms.

From the perspective of how to get your content to rank for the queries you uncover through your keyword research, the focus should now be on creating the best site and the best content while still incorporating keywords into your content where it is natural to do so.

Keyword "not provided" and Co-Occurrence Analysis

Two major changes in recent years which affect the way we approach keyword research are Google's move to obscuring referring keywords (keyword "not provided") and a fundamental change to the way Google approaches looking at the entire body of keywords on a page, often referred to as *co-occurrence analysis*.

With "not provided" Google has blocked the vast majority of the data which lets you see which keywords lead to organic search visits to your site. While there are some workaround solutions to try to retrieve or re-create this data such as Searchmetrics (www.searchmetrics.com) other changes in the way Google views individual keywords make this data somewhat less vital than it has been in the past. Going forward it is going to become more important to not just think about how to replace lost keyword data but to understand your site's relationship to keywords and to better understand your users' needs and intent. Rather than simply thinking about optimizing a given page to a specific keyword think about how that page can more effectively answer your users' needs and provide a breadth of useful information to satisfy that need.

As referred to in the Hummingbird section above, Google's intention is to now be able to understand the meaning behind a query rather than just the specific words used in that query. They also are actively working towards understanding the meaning and theme of the content on your pages rather than just the specific keywords used within that content. Rather than simply ranking a given page for how well it is optimized for a specific query they will attempt to rank the content which most closely answers the true intent behind the user's query.

Related to this notion of intent is the idea of co-occurrence analysis. Co-occurrence analysis refers to an analysis of what words appear most commonly on a page and under what circumstances they appear together. Rather than simply looking at keywords in isolation on a page, Google appears to have begun

looking at what words commonly appear together under varying circumstances. We explore the notion of co-occurrence in depth in the “Co-occurrence Analysis” and the “Integrating Keyword Research, Co-Occurrence Analysis, and Knowledge of User Intent” sections below. For now it is useful to note that Google is not looking at just specifically the repetition of a single keyword phrase over and over again but rather at what kinds of words appear together to signal the true intent and theme of the page and how well the content answers the true intent behind the user’s query.

As you continue through this chapter consider that the keyword research that you will perform is no longer simply about finding the exact best keyword phrases to specifically optimize pages on your site for but rather to also help uncover what your potential users are actually needing to get from your site. Rather than highly optimizing each page on your site for one or two specific keyword phrases think about which phrases might belong together on a page and what the theme and intent of the page is. Consider each page more holistically and how you might improve the page to better answer your users’ needs. Rather than trying to over-optimize for the exact right number or combination of phrases, write naturally and for your users, using a wide variety of related terms and synonyms which you uncover during the keyword research process.

Traditional Approaches: Domain Expertise, Site Content Analysis

One of the smartest things you can do when initially conducting keyword research is to brainstorm original ideas with the participants in the business before getting keyword tools involved. This can be surprisingly effective for coming up with numerous critical keywords, and it can help you understand if your organization thinks about your offerings using different language than your customers - in which case you may want to adapt!

Start by generating a list of terms and phrases that are relevant to your (or your client’s) industry and pertain to what your site or business offers. The brainstorming phase should ideally result in a list of several dozen to several hundred or more keyword searches that will bring relevant, qualified visitors to your site. It can be a great idea to get sales, customer services, or whichever branch works most directly with clients to participate in the brainstorm, as they may have input into keywords or phrases the customer uses or expresses interest in that aren’t currently targeted.

One easy way to begin this process is to gather your team in a conference room and then follow these steps:

1. Produce a list of key one- to three-word phrases that describe your products/services.
2. Spend some time coming up with synonyms that your potential customers might use for those products and services. Use a thesaurus to help you with this process.
3. Create a taxonomy of all the areas of focus in your industry. It can be helpful to imagine creating a directory for all the people, projects, ideas, and companies connected to your site. You can also look at sites that are leaders in the industry and study their site hierarchy as a way to start your thinking about a taxonomy.
4. Broaden your list by thinking of higher-level terms, and topics, of which your products or services are a subset.
5. Review your existing site and extract what appear to be key phrases from your site.
6. Review industry association and/or media sites to see what phrases they use to discuss your topic area.
7. List all of your various brand terms.
8. List all of your products. If your site has a massive number of products, consider stepping back a level (or two) and listing the categories and subcategories.
9. Have your team imagine that they are a potential customer, and ask them what they would type into a search engine if they were looking for something similar to your product or service.

10. Supplement this by asking some people outside your business what they would search for, preferably people who are not directly associated with the company. Consider also the value of performing actual market research with a test group of consumers in your demographic, and ask them the same question.
11. Use various tools (such as Google Webmaster Tools) to see what terms people are already using to come to your site, or what terms they are using within your site search tool if you have one.

Gathering intelligence on how potential customers discuss related products and services is what a traditional marketer might have done prior to initiating a marketing campaign before the Web existed. And of course, if any of this data is available to you from other departments of the company, be sure to incorporate it into your research process.

Include Competitive Analysis

Your competitors face the same problem, and unless you are very lucky, they are also probably resourceful and creative. You can likely count on their having invested in learning how their customers think and the best ways to appeal to them. So, add these steps to the process:

1. Review your competitors' websites and see what keywords and phrases they use for their products and services that compete with yours. Look also for unique variations and synonyms they incorporate into their language, and see if these indicate shifting trends in vernacular in your industry.
2. Record what non-brand terms they use for their business.
3. Read any articles they have written that are published on sites other than their own.
4. Observe what the media may have had to say about them.

Add these ideas into the mix and you will have a wonderfully robust set of keywords to use as a starting point.

You may ask why you should go through all this trouble. Don't the keyword tools take care of all this for you? There are two reasons why the extra effort is critical:

- Your internal team has a rich array of knowledge that the keyword tools do not: they know where to start. Keyword tools require the initial input of information, and the quality of the data they provide is only as good as the quality of the "seeds" you give them.
- The upfront brainstorming helps your organization's stakeholders better understand the market and the opportunities. It also helps them "buy into" the eventual prioritization of keywords.

Once you have completed these steps you will have in hand a rich set of terms of interest. The next step is to expand those terms of interest using keyword research tools.

Integrating Keyword Research, Co-Occurrence Analysis, and Knowledge of User Intent

What changes like co-occurrence and Hummingbird mean for keyword research and strategy is that it is not as simple as it once was. No longer can you simply research what terms get the most traffic and build pages optimized to each and every one of the terms you would like to target. You still need to do keyword research to understand the universe of keywords your potential users might be looking for, and the relative scale of the opportunity in ranking well on each of those terms. You must also be able to uncover the intent your users have behind the searches they are performing and be able to build quality pages and content which meet those needs.

Not only must your pages be well written and researched, and meet the actual needs of your potential users, but each page must present a coherent theme so that the engines can truly understand what need it is attempting to satisfy, and thus what body of related keywords it should be considered to rank for. Searchmetrics (<http://www.searchmetrics.com/>) which is mentioned elsewhere in this Chapter has tools in their Content Optimization section which can help you understand the terms and phrases which are highly relevant and related to your optimal keyword targets.

It is also vital that you understand what your customer's needs are and whether they are able to meet those needs with the information they find on your pages. There are lots of ways to understand more about what your users want and whether they are getting it from your site from services like Usertesting.com (<http://www.usertesting.com/>), to simply spending time talking to your customers about how and why they use your site. It is becoming more and more important that you ensure that your content is both meeting these needs in as simple a way as possible, that it includes keywords relevant to those needs, and that it is integrated into well-designed, clearly themed pages.

Keyword Research Tools

A wide variety of options are available for performing keyword research, including tools provided by the search engines, tools developed by third parties, and tools for complex keyword analysis of terms culled during research. In this section, we will review each of these; but first we'll provide some perspective on how to use these tools.

Things to Keep in Mind

It is important to keep in mind that when you are using the various keyword research tools to brainstorm keywords, they are all based on relatively limited data. You should be aware that each tool will provide different search count numbers than the other tools. Rather than focusing on the exact search count of various terms, you should think of each tool as a good way to get a general comparison of two search terms, or of groups of terms, relative to one another.

For example, if you compare two terms and see that one term is more popular than the other because it returns a higher search count, you at least know that Term A may be more popular and searched for more often than Term B, but you can treat the search count values as only (rough) estimates.

If you are just starting out with keyword research, consider starting with the Google AdWords Keyword Planner (which is free with an AdWords account) and at least one of the paid tools, such as KeywordDiscovery.com. This will give you a rich, varied data set with which to begin your keyword research. Over time you can experiment with the other tools and adjust your process as you find tools that you prefer for one task or another.

In addition, as you look at these keywords, keep in mind the semantic aspects of search, and the probable intent of the person performing the search. You don't want to simply slap popular search terms on various pages of your site. You are using the keyword research information to provide you with insight into how your potential customers think, but you need to take into consideration many other inputs as well.

Keyword Research Data from the Engines

The search engines provide a number of tools that can help you with keyword research. Many of these are not designed specifically for that purpose, but they can be used to obtain interesting keyword research information if they are used in the right manner. The data in these tools show the number of pages that are related to a search phrase, not the number of searches on the phrase. This is still a useful indicator of the importance of a keyword phrase as more web pages tend to get built for more popular topics.

Related terms

Several of the engines offer "related" terms, including Google, Yahoo!, Bing, Ask, and Yippy (which shows related terms in clusters, as shown in Figure 5-3). This data can be invaluable if you're looking to find related terms that may not have come up through competitive analysis or brainstorming.

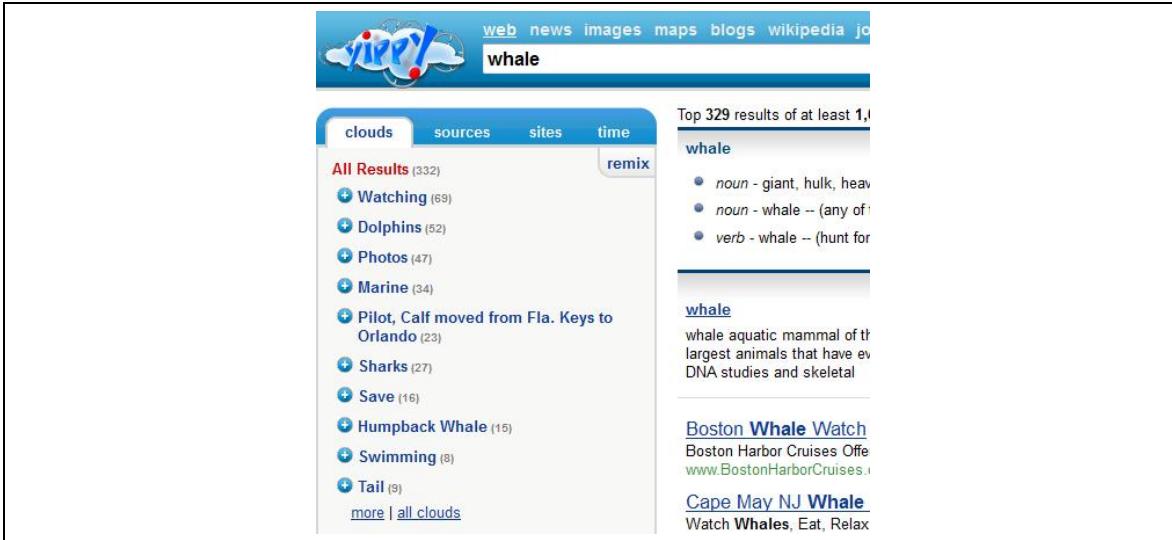


Figure 5-2. Yippy related terms clusters

Common usage and phrase combinations

Using a search with the * character can give you a good idea of what terms/phrases commonly precede or follow a given term/phrase. For example, using * *ringtones* can show you phrases that are commonly associated with the term *ringtones*, as shown in Figure 5-3.

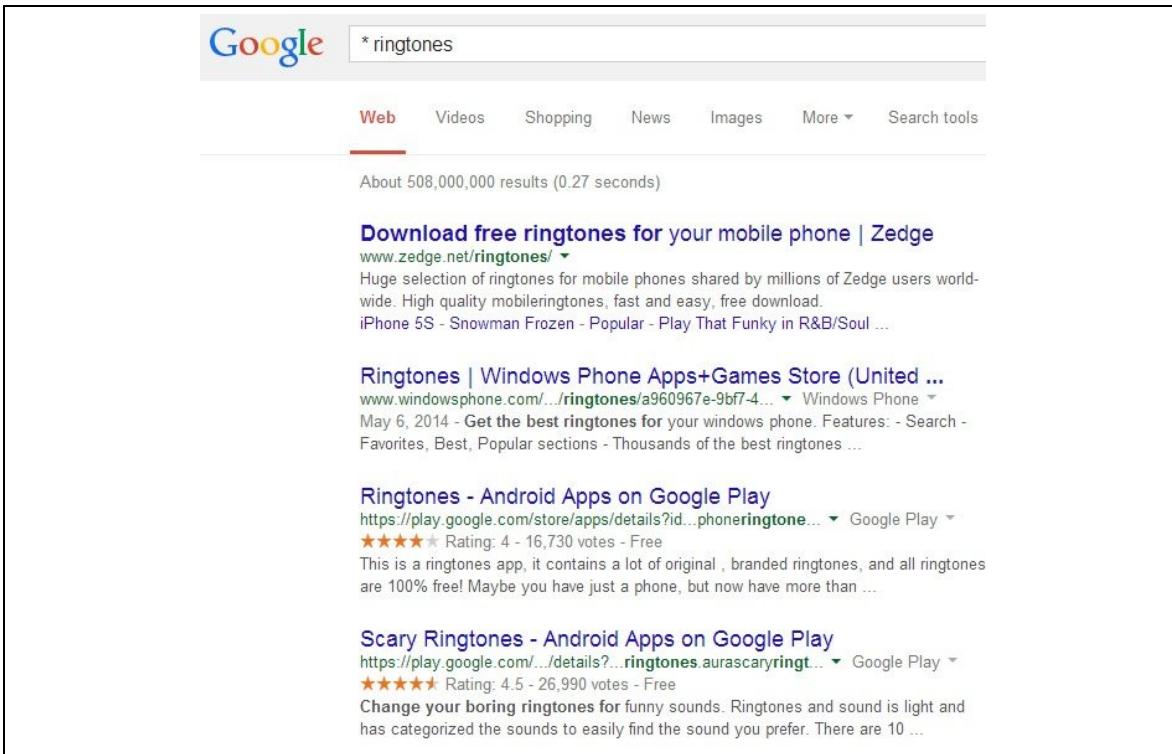


Figure 5-3. Finding common phrases

Frequency of recent usage

Using the very cool Google date range operator, shown in Figure 5-4, you can determine how many times in the past day, week, month, or year new content related to your term was added to the Google index. The easiest way to do this is to click on search tools on the top of the Google results. Once you do that, you can

pick from “Any Time” (which is the default), “Past hour,” “Past 24 hours,” “Past week,” “Past month,” “Past year”, and “Custom range.” This will limit you to the results that were added to the index during the referenced time frame.



Figure 5-4. Google pages indexed in past 24 hours

Picking a custom range provides you with a calendar method for picking the date range you want to focus the search on, so you can pick any time interval you want. For example, you might pick November 1, 2014 to December 24, 2014 if you wanted to see what happened during the 2014 holiday season. For additional flexibility, you can perform a normal search, get your result, and add a parameter to the end of the results page URL, using the operators shown in Table 5-1 (Note: Google Instant must be disabled in your search settings for this to work).

Table 5-1. Google date search operators

&as_qdr=d	Past 24 hours
&as_qdr=d4	Past four days
&as_qdr=w	Past week
&as_qdr=w5	Past five weeks
&as_qdr=m6	Past six months
&as_qdr=y2	Past two years

This can give you some seasonal data if you follow it closely, and it can also show you who is producing content in your arena. For example, try a search for *President Obama (past 24 hours)*. You can also get information on activity level from Yahoo! News, as shown in Figure 5-5.

The screenshot shows the Yahoo! News search interface. The search bar at the top contains the query 'president obama'. Below the search bar, there is a link to 'Get alerts about president obama by email' with a bell icon. The main content area displays news articles. The first article is titled 'Obama at West Point: US must lead by 'empowering partners'' and is from 'Christian Science Monitor via Yahoo! News' on May 28 02:16pm. The second article is titled 'Five Ways Obama Wants To Change Foreign Policy' and is from 'ABC News via Yahoo! News' on May 28 03:32pm. The third article is titled 'Obama avoids controversial questions in concussion speech' and is from 'Yahoo News via Yahoo! News' 2 hours, 10 minutes ago. The sidebar on the left provides filtering options: 'Web', 'News' (selected), 'More', 'Filter by Source' (with links to 'All sources', 'Yahoo! News', 'NPR', 'Washington Post', 'ABC News', and 'More'), and 'Filter by Time' (with links to 'Any time', 'Past hour', 'Past day', and 'Past week').

Figure 5-5. Yahoo! News activity level

Both Google News and Yahoo! News are great places to do a bit of digging into anyone who is publishing press releases or getting news coverage on the terms/phrases you might be researching. If there's a lot of activity in these arenas (and it is not all press release spam), you can bet the terms are going to be even more competitive. For example, the following URLs show SEO news in Google News (<http://news.google.com/news?q=seo>) and in Yahoo! News (<http://news.search.yahoo.com/news/search?p=seo>).

You can combine all of this data to form a very well-rounded view of a particular term or phrase, and although it is probably overkill for most keyword research projects, it is certainly a valuable exercise and something to monitor closely if you're basing a lot of your success off of a single search query (or just a handful of queries). Even if you're just trying to get a better sense of what's going on infrequently and informally, these pieces of the keyword puzzle can be remarkably valuable.

Keyword Research with Tools

It is great to get this data from search engine queries, and it can certainly help you get a sense of the importance of a given keyword. However, a large array of tools exists to give you direct insight into the volume of searches performed on specific keywords, and also to help you discover new keywords to consider.

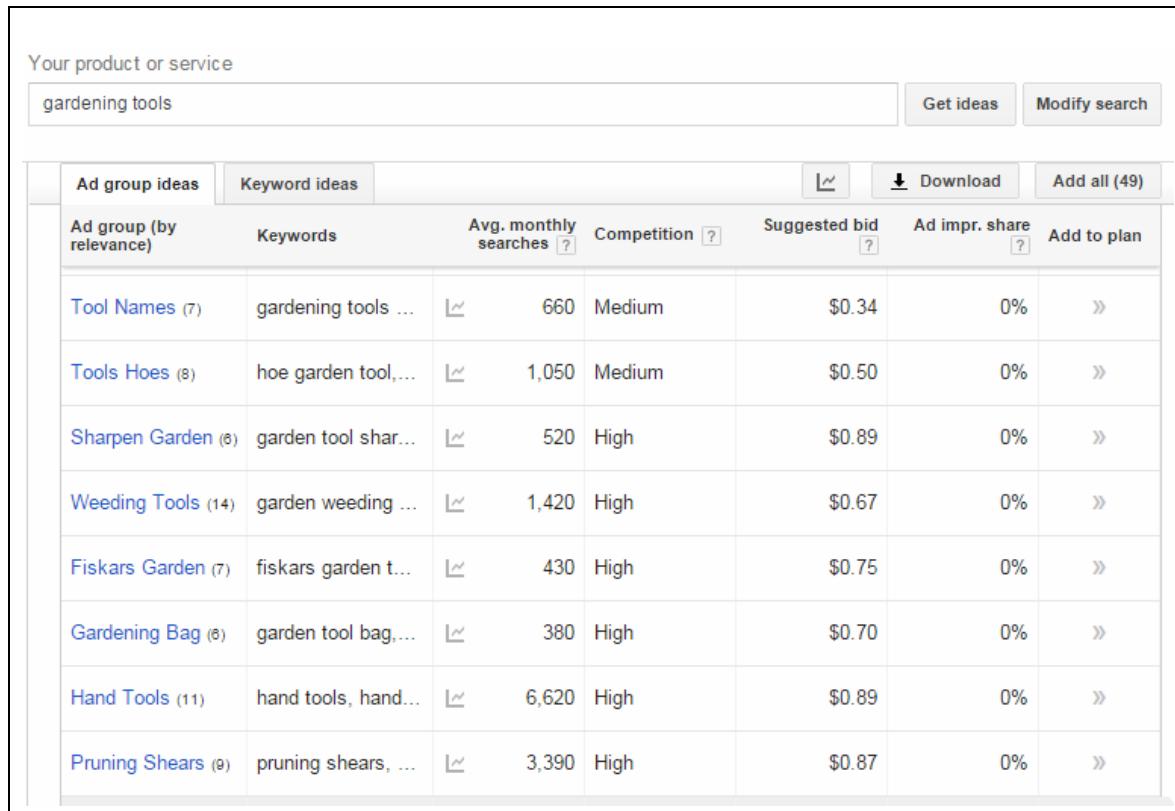
We review many of these leading tools on the pages that follow.

Google AdWords Keyword Planner

Google provides a couple of tools specifically designed for use in keyword research. Although they are primarily meant to help their paid search customers, they can also be used to obtain information for organic search. This tool is most effectively utilized when coupled with an active Adwords campaign.

What the Keyword Planner provides

Google's AdWords Keyword Planner (<https://adwords.google.com/KeywordPlanner>) provides related terms, search volume estimates, search trends, and ad cost estimates for any keyword or URL that you enter (see Figure 5-6).



The screenshot shows the Google AdWords Keyword Planner interface. At the top, there is a search bar labeled "Your product or service" containing "gardening tools". Below the search bar are buttons for "Get ideas" and "Modify search". The main area is a table with the following columns: "Ad group ideas" (highlighted in blue), "Keyword ideas" (highlighted in red), "Keywords", "Avg. monthly searches", "Competition", "Suggested bid", "Ad impr. share", and "Add to plan". The table lists several keyword suggestions:

Ad group ideas	Keyword ideas	Keywords	Avg. monthly searches	Competition	Suggested bid	Ad impr. share	Add to plan
Tool Names (7)	gardening tools ...	garden tools	660	Medium	\$0.34	0%	»
Tools Hoes (8)	hoe garden tool,...	hoe garden tools	1,050	Medium	\$0.50	0%	»
Sharpen Garden (6)	garden tool shar...	garden tool sharpener	520	High	\$0.89	0%	»
Weeding Tools (14)	garden weeding ...	garden weeding tools	1,420	High	\$0.67	0%	»
Fiskars Garden (7)	fishkars garden t...	fishkars garden tools	430	High	\$0.75	0%	»
Gardening Bag (6)	garden tool bag,...	garden tool bags	380	High	\$0.70	0%	»
Hand Tools (11)	hand tools, hand...	hand tools	6,620	High	\$0.89	0%	»
Pruning Shears (9)	pruning shears, ...	pruning shears	3,390	High	\$0.87	0%	»

Figure 5-6. Google AdWords Keyword Planner

The Keyword Planner provides multiple ways to search: based on words/phrases, based on websites/landing pages, and based on categories. In the "Your product or service" box, you can enter key words/phrases (one per line) and the Keyword Planner will return keyword ideas related to the term you entered (note that you have to click on the "Keyword ideas" tab; the "Ad group ideas" is what displays by default). The output of a search will show you:

Keyword (by relevance)

The related keywords in a list, including the phrase or phrases you entered.

Average Monthly Searches

Shows the search query volume for the keyword for the locations, languages, and Google properties you defined.

Monthly Searches bar chart

Shows the search volume for the keyword broken out by month for the last 12 months (requires that you hover over the small gray chart icon immediately to the left of the numbers in the Average Monthly Searches column).

Competition

Displays the relative competitiveness of the keyword (in paid search).

There are a number of settings you can use to tune your search. These include:

Include/Exclude

Provides a way to specify or exclude additional terms. For the include option, the Keyword Planner will only show you suggestions that included these terms (see Figure 5-7). The exclude option allows you to add a negative keyword for any keyword phrase that does not pertain to your business. This feature is not necessarily useful for researching keywords for organic search; rather, it is more valuable when planning your AdWords account bids.

Locations

Use this to set the country used for the "Average Monthly Searches" part of the results.

Language

Set the default language to use.

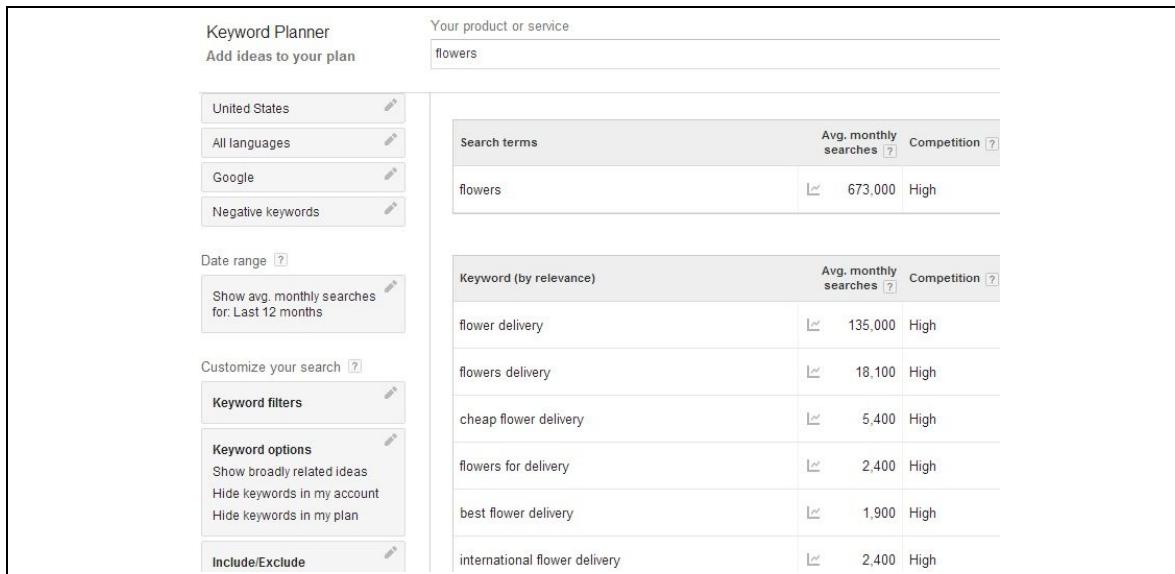
Search Networks

Allows you to specify whether you want traffic estimates returned for only searches done specifically on the Google search engine or whether you also wish to include Google search partners. Google search partners include other non-Google websites like AOL and other Google properties such as Maps, Video, and other Google sites.

Date Range

By default the tool shows search volumes for the past 12 months. You can also specify a custom range of months that you would like to see and compare volumes for 2 different time ranges.

Note: you must log in to your Google AdWords account to use the Keyword Planner, but fortunately, you don't have to spend money on AdWords to access this free tool.



The screenshot shows the Google Keyword Planner interface. On the left, there are several filter and search parameters. The main area displays search results for the term 'flowers'.

Search terms:

Search terms	Avg. monthly searches	Competition
flowers	673,000	High

Keyword (by relevance):

Keyword (by relevance)	Avg. monthly searches	Competition
flower delivery	135,000	High
flowers delivery	18,100	High
cheap flower delivery	5,400	High
flowers for delivery	2,400	High
best flower delivery	1,900	High
international flower delivery	2,400	High

Figure 5-7. Specifying a required term in the Keyword Planner.

By entering a webpage URL, you can get the Keyword Planner to show you keywords related to that website or specific page. (see Figure 5-8).

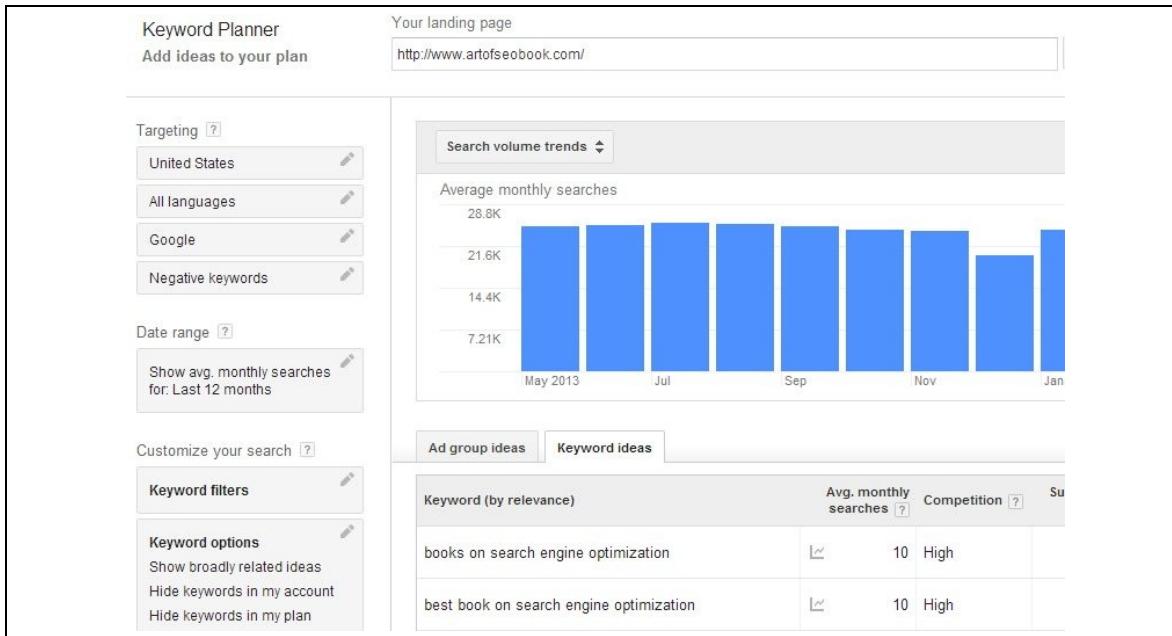


Figure 5-8. Keyword Planner Site-Related Keywords

What the Traffic Estimator provides

Within the Google AdWords Keyword Planner is a feature called the Traffic Estimator (see Figure 5-10) that allows you to get estimates of traffic on different keywords (i.e., the potential click-throughs you may see to your site, instead of just the number of impressions provided by most keyword research tools).

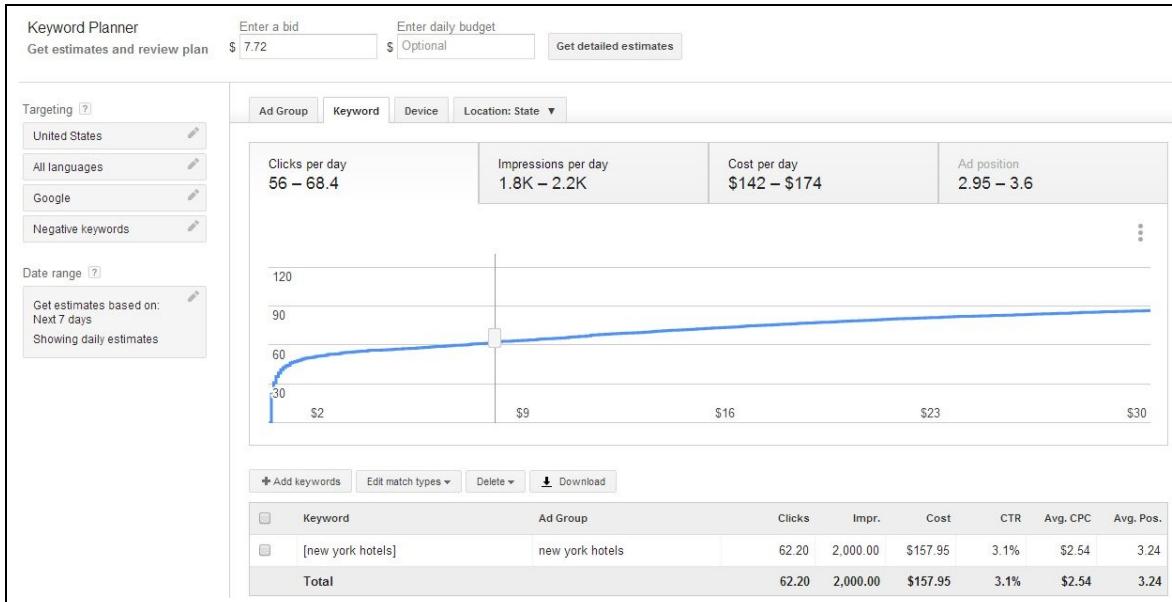


Figure 5-9. Google AdWords Keyword Planner's traffic estimates feature

When you specify one or more keywords to get traffic estimates on, the Keyword Planner will return the predicted number of clicks per day, impressions per day, cost per day for those clicks and average ad position, as well as additional breakdown on a per keyword basis, including clicks, impressions, cost, click through rate, average cost-per-click, and average ad position. The cost information can provide you with additional insight into how competitive a keyword is in organic search as well. You can get more exact estimates by selecting specific budgets or costs per click.

When using the Keyword Planner's traffic estimates feature for keyword research, it is best to enter your keywords as "exact match," meaning the keywords specified were exactly the ones the Google user searched on.

After you've entered your keywords, you can leave Daily Budget blank. Select your language, search network, and the location you're targeting (for U.S.-focused campaigns, use the default of "Countries and territories" and enter "United States").

You'll then see data for each keyword you entered. Useful data for keyword research purposes includes Estimated Clicks/Day and Cost/Day. You can compare each keyword's estimated clicks to see which term is more likely to be searched for and clicked on than others.

In the results shown in Figure 5-10, at the maximum possible bid, *internet marketing* is estimated to have 3 clicks per day, while *search engine marketing* has 2, *search engine optimization* has 11, and *seo* has 68. Based on this data, it is clear that *seo* is the most popular term of the four and is likely to be one of the more competitive terms. In this particular case, an additional factor enters into the equation, because *seo* is a "trophy term" on which people put an extra focus for branding reasons. Nonetheless, the value of the traffic data is considerable.

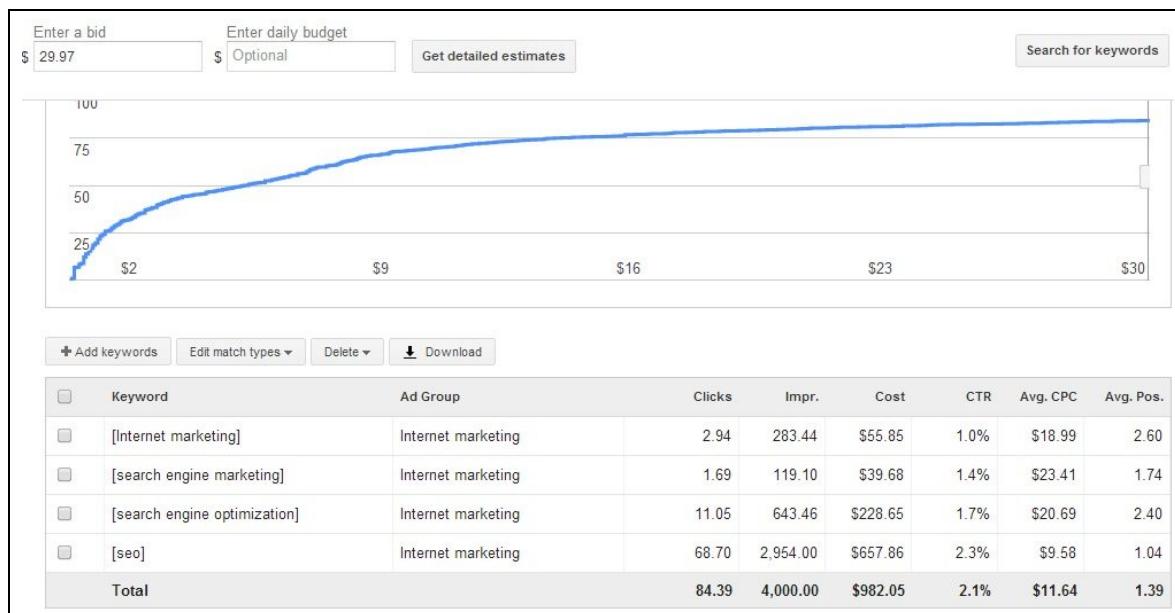


Figure 5-10. Google AdWords Keyword Planner traffic estimates

Where the tool gets its data

The Google AdWords Keyword Planner gets its data from Google's search query database.

How the tool is useful

The Keyword Planner offers some useful information about your keyword campaigns, such as suggestions for similar keywords, an estimate of the keyword's popularity, ad costs and positions, general search volume trend information, and keyword campaign suggestions for your site or your competitor's site. The tool is great for compiling a lot of general information about a keyword.

The traffic estimates feature provides a rough estimate of your keyword's click-through rate. Based on the estimated clicks per day, you can get a relative idea of which of your keywords are the most popular and can potentially bring you the most traffic. Practitioners should use other tools to cross-reference these figures, as these numbers can be inaccurate.

Cost

The Google AdWords Keyword Planner is free to use as long as you have a Google AdWords account.

Google Trends

Google Trends (<http://www.google.com/trends>) allows you to view the popularity trend of a keyword over time and by geography, as well as compare two or more search terms to each other to see relative popularity and seasonality/trending over time. If you enter the terms into the search bar and separate them with commas, you'll see the requested terms' trend history depicted in different colors on a graph spread over a certain time period. You can modify the results by changing the time period and/or region (see Figure 5-12).

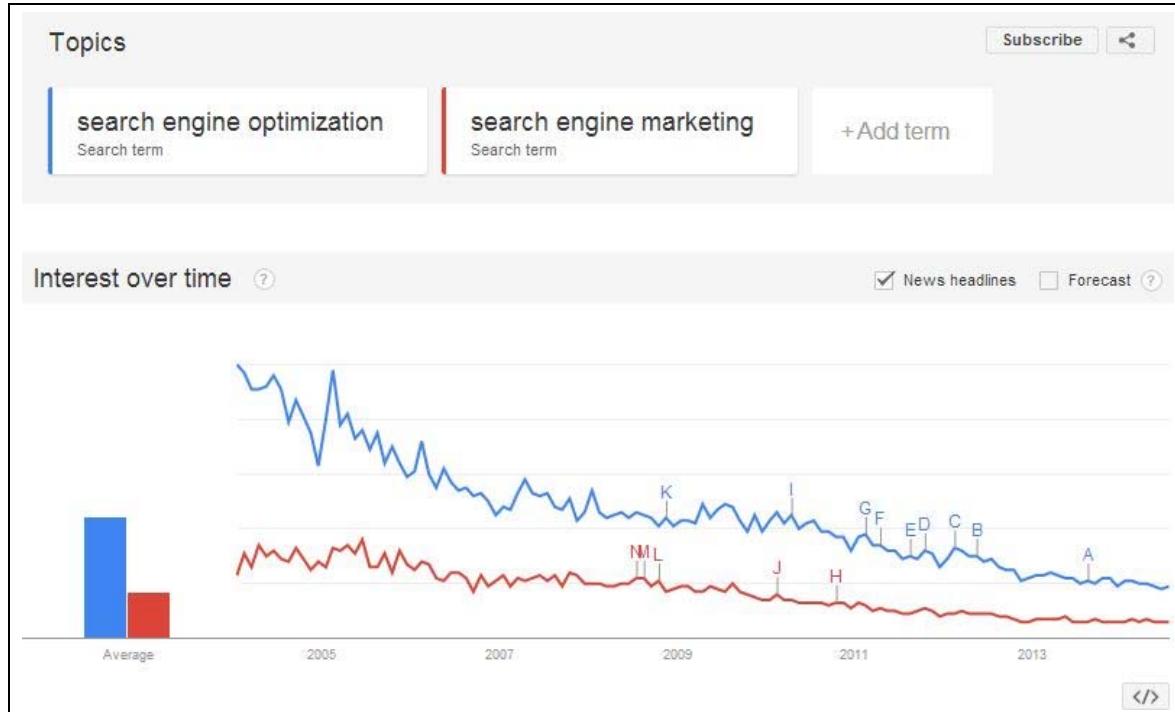


Figure 5-11. Google Trends sample output

With Google Trends, users can also see Google's estimate of which cities, regions, and languages performed the largest number of searches for a particular keyword (see Figure 5-12). Experienced marketers often feel that this data is imprecise (and occasionally inaccurate) because more accurate data from analytics and search advertising campaigns have often contradicted the results.



Figure 5-12. Google Trends top cities data

Lastly, plotted on each graph are a few articles/search results related to your keyword query, which correlate to peaks and valleys in the historical search popularity.

Where it gets its data

Google Trends gets its data from searches performed on Google.

How it is useful

Google Trends is a great, easy tool for comparing keywords and identifying which term is more popular than the other; in addition, you can examine this data over many years with seasonality factored in. Although Google Trends doesn't supply absolute search volume figures like the Google AdWords Keyword Planner, the graphs are simple to understand and provide a perfect visual of search trends over a particular period of time (note that this works only with relatively popular terms, not with long tail search terms). You can also sign up for email alerts from Google Trends for the topics/keywords you're researching, another convenient feature that can help keep you up to speed on the search trends you're targeting (<http://www.google.com/trends/subscriptions>).

Cost

Google Trends is free to use.

Bing Keyword Research Tool

Bing Keyword Research Tool (<http://www.bing.com/toolbox/keywords>) generates keyword suggestions based on a search term or website you enter.

Entering a keyword in the search box will return data that includes search phrases that contain the keyword you provided, along with the preceding month, typical click through rate percentages and average cost per click. For example, a search for *ice cream* returns ice cream maker, ice cream recipes, and ice cream shop. As you can see in Figure 5-13, the term *ice cream* had, according to Microsoft, 856,543 searches in the month prior to this screen shot.

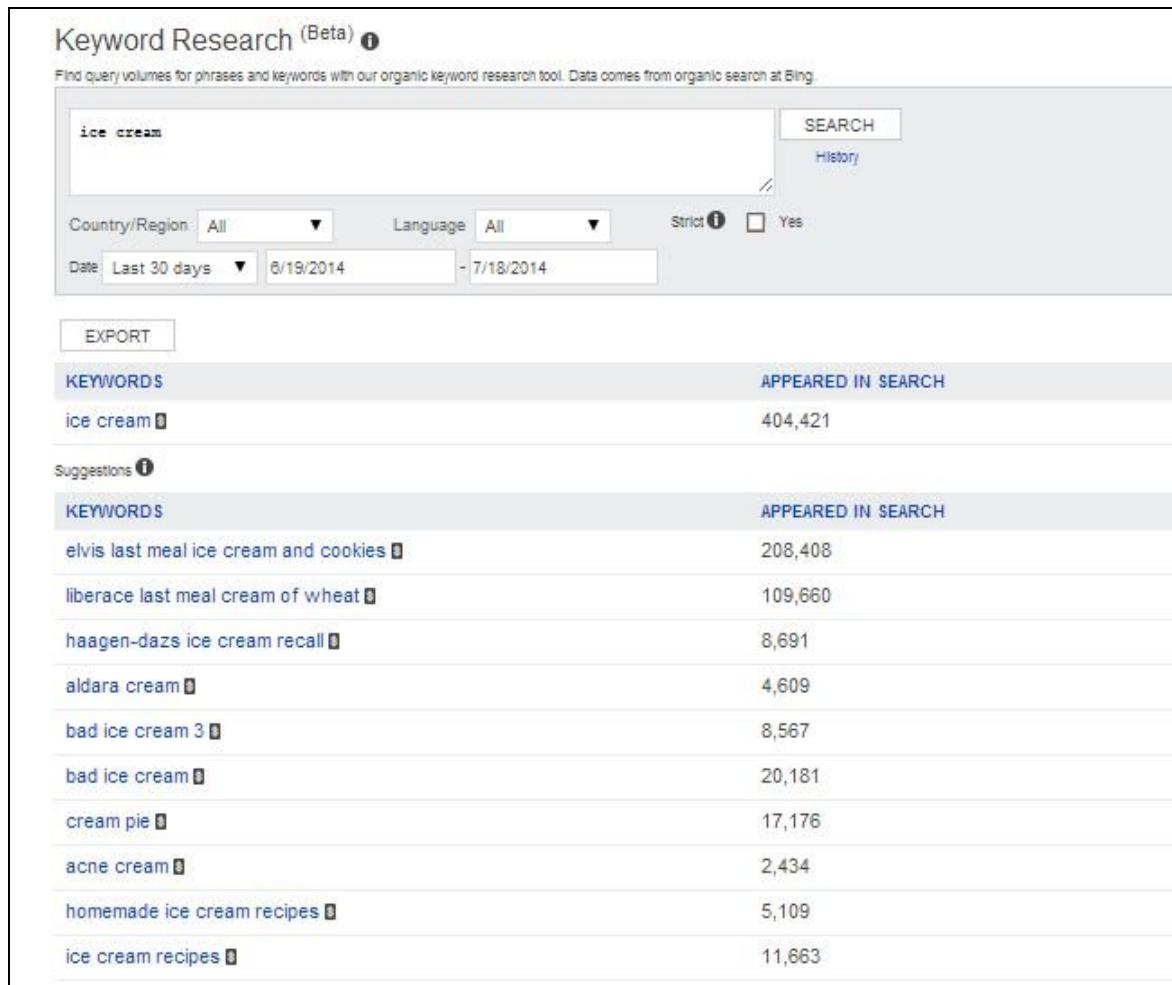


Figure 5-13. Bing Keyword Research Tool basic output

The Export to Excel option allows you to pull the collected data into a spreadsheet. Although the CTR% (click-through rate) and CPC (cost-per-click) columns are intended for paid search customers, they can also provide some indication of SEO value. You can multiply the CTR times the search volume to get a sense of how many clicks a high-ranking paid search result might get (comparable organic results will get three to four times more clicks), and the CPC provides some indication of the competition for ranking on the term. You can also obtain demographic data as shown in Figure 5-14.



Figure 5-14. Bing Keyword Research Tool demographic settings

If you enter a URL into the search bar, the tool will return keywords related to the website selected.

Where it gets its data

The Bing Keyword Research Tool obtains its data from Microsoft's Bing search query database.

How it is useful

This tool is useful in generating keyword suggestions based on a keyword you are targeting or on your site's URL. You can also enter your competitor's URL and see what the keyword suggestions are for their site. One important feature of this tool is that it shows actual numbers from the Bing search query database.

Cost

The Bing Keyword Research Tool is free, although you do have to create an account with Bing Ads and provide credit card information in the event that you advertise on the Bing network.

KeywordDiscovery

A very popular third-party tool for keyword research is Trellian's KeywordDiscovery (<http://www.keyworddiscovery.com/>). KeywordDiscovery offers the following features:

Keyword Research

When you enter a keyword or phrase in the search bar under the Research section, KeywordDiscovery displays the most popular search terms that include the keywords you provided, along with a count of how many searches were performed for those keywords in the past 12 months (see Figure 5-16).

Search Term:

Results per page: ▼

?

Exclude:

?

Database: ▼

Historical ?

Phrase Match ?

Include Plurals ?

Remove Spaces ?

Adult Filter ?

Spell ?

Thesaurus ?

Inflected Form ?

Competitors ?

Related ?

Fuzzy/Like ?

Industry ?

Domain Score ?

Figure 5-15. KeywordDiscovery basic output

Seasonal Search Trends

If you click on the little bar graph icon next to the number of searches for a query, you'll see a graph of the search trends for that keyword over the past 12 months. You can mouse over each bar and see the number of searches for that time period, and you can sort the chart by historical data (how many searches in the past year), monthly data (number of searches broken down into each month), trends (a graph of the search trends over the past year), combination (a graph of Historical Global versus Global Premium search data; Figure 5-16 shows a definition of these terms), and market share (a breakdown of which search engines were used to search for the query).



Figure 5-16. KeywordDiscovery seasonal search trends

Spelling Mistake Research

Typing the query `spell:[keyword]` as the Search Term will return spelling variations for that keyword, the number of times the keyword has been searched for (searches), and the keyword results for your search (queries).

For example, `spell:optimization` returns results such as *optimization*, *optimazation*, and *optimisation*, as shown in Figure 5-17.

Search Term:

Results per page:

[?](#)

Exclude:

[?](#)

Historical [?](#)

Database:

Phrase Match [?](#) Include Plurals [?](#) Remove Spaces [?](#) Adult Filter [?](#)

Spell [?](#) Thesaurus [?](#) Inflected Form [?](#) Competitors [?](#)

Related [?](#) Fuzzy/Like [?](#) Industry [?](#) Domain Score [?](#)

? Estimated results: 33

<input type="checkbox"/> Word ?	Similarity ?	Searches ?	Options	Queries ?
<input type="checkbox"/> optimization	100%	1,327	[1] [2] [3] [4]	259
<input type="checkbox"/> optimizacion	75%	266	[1] [2] [3] [4]	73
<input type="checkbox"/> optimisation	75%	253	[1] [2] [3] [4]	60
<input type="checkbox"/> optimizador	25%	60	[1] [2] [3] [4]	29
<input type="checkbox"/> optimizations	75%	33	[1] [2] [3] [4]	1
<input type="checkbox"/> optimizaciÃ³n	25%	15	[1] [2] [3] [4]	3
<input type="checkbox"/> optimazation	75%	11	[1] [2] [3] [4]	3
<input type="checkbox"/> optimizaciÃ³n	50%	10	[1] [2] [3] [4]	3
<input type="checkbox"/> optimizacija	25%	8	[1] [2] [3] [4]	4
<input type="checkbox"/> optimizaciÃ³n	25%	8	[1] [2] [3] [4]	3

[Subscribe to KeywordDiscovery for full access](#)

[Add](#)

[Clear](#)

[Select All](#)

[Select By](#)

Figure 5-17. KeywordDiscovery spelling mistakes output

Related Keywords

Typing either `related:[keyword]` or `crawl:[keyword]` will return keywords that are related to the term you provided. For example, typing in `related:seo` returns results such as `internet marketing`, `video`, and `internet consulting`. You can see an example of this in Figure 5-18.

The screenshot shows the KeywordDiscovery interface. At the top, there is a search bar with 'Search Term: related:seo', 'Results per page: 10', and a 'Search' button. Below the search bar are various filter options: 'Exclude:', 'Database: Global Premium', 'Historical', 'Phrase Match', 'Include Plurals', 'Remove Spaces', 'Adult Filter', 'Spell', 'Thesaurus', 'Inflected Form', 'Competitors', 'Related', 'Fuzzy/Like', 'Industry', and 'Domain Score'. The main area displays search results for 22 terms, with the first few being: seo (1,690 searches), internet marketing (n/a), video (2,918 searches), camera phone (n/a), san diego seo (n/a), chat (19,675 searches), engine (16,236 searches), internet consulting (n/a), search engine optimization (389 searches), and sem (151 searches). At the bottom, there are buttons for 'Add', 'Clear', 'Select All', 'Select By', and various export icons.

Query ?	Searches ?	Options
seo	1,690	
internet marketing	n/a	
video	2,918	
camera phone	n/a	
san diego seo	n/a	
chat	19,675	
engine	16,236	
internet consulting	n/a	
search engine optimization	389	
sem	151	

Figure 5-18. KeywordDiscovery related keywords output

Keyword Density Analysis

This feature checks how often keywords are found on the URL you provide, assigns a keyword density percentage to those keywords, and lists the number of searches performed for each term.

We do not recommend using keyword density as a metric to judge a page's keyword targeting. The search engines use far more sophisticated analyses of keywords for their algorithms, and relying on rough counts such as this can seriously mislead you. See "Keyword Targeting" in Chapter 6 for more on how to effectively target keywords on the page.

One good use for the Keyword Density Analysis feature is to enter a competitor's URL into the search bar and see what keywords the site is targeting. It is a great tool to use for competitive research.

Domain Researcher Tool

This tool requires an Enterprise subscription. It allows you to search for available domains that are based on popular keyword search terms. These domains have high traffic potential, as the tool shows how many users have searched for the URL. The tool is great if you want to register other domains in your industry and want these domains to be keyword-rich (see Figure 5-19).

Domain Researcher Tool (DNS Tool) ?

Use the Domain Researcher to quick check upto 100 search phrases, are available to be registered as either a .co stripped out.
[Click here for more info.](#)

NOTE: This tool currently does not support IDN based domains.

Step 1. Insert/Import Keywords ?

Upload Keywords from a Text File, or ?

Extract Keywords from Meta Tags, or ?

Extract Keywords from a URL, or ?

Import From Projects, or ?

AAP (87)

Insert one Keyword per line.
 Limited to 100 Keywords.
 Please use alpha numeric
 keywords only.

Lines: 0

Figure 5-19. KeywordDiscovery Domain Research Tool

Competitive Intelligence reports

Trellian, which powers KeywordDiscovery, also offers various Competitive Intelligence reports (which require a separate subscription). These reports include:

Link Intelligence

Identifies which links are sending traffic to your competitors

Search Term Intelligence

Identifies which search terms/phrases are driving traffic to your competitors

Search Engine Intelligence

Identifies which specific search engines send traffic to your competitors

PPC Campaign Intelligence

Identifies which search terms your competitors are bidding on

Referrer Intelligence

Provides information about specific sites that are referring traffic to your competitors

Popularity Index Report

Monitors the Popularity Index (which is based on the number of unique sessions a domain receives) of your competitors

Ranking Report

Provides a view of which terms your competitors are ranking for, the rank of these terms, and any changes in ranking over the past 30 days

Meta keywords

Provides a report that analyzes your competitors' meta keywords

Competitive Intelligence Executive Report

Provides information about every Competitive Intelligence Report available, as well as several subreports

Where it gets its data

Trellian derives its keyword data primarily from aggregated Historical Global data purchased from ISPs. Trellian also uses a panel of 4.4 million users to collect its Global Premium data. The company touts that the Global Premium data removes the bias that various spiders introduce into data from other sources.

How it is useful

As we mentioned earlier, KeywordDiscovery offers a multitude of tools that are great for keyword research. Trellian also offers various tools that are useful for competitive research. You can almost think of KeywordDiscovery as a one-stop shop for research since it offers a diverse set of tools, but as with many of the other keyword research tools we've discussed here its data sources are limited, and you need to take this into account in your use of the tool.

Cost

KeywordDiscovery offers different subscription options that range from a standard monthly subscription for \$69.95 to a yearly Enterprise subscription for \$4,752 (pricing as of July 2014). Competitive Intelligence Reports range from \$99.95 per month per domain (plus a \$150 setup fee) to \$995 per year per domain. The free tool with limited features is also available. We recommend reviewing the options and choosing the package that will work best for your company.

Experian Hitwise

Experian Hitwise (<http://www.hitwise.com>) offers a wide range of competitive and web statistics via its service. One component of the Experian Hitwise suite, Hitwise Search Intelligence, is a powerful keyword research tool for analyzing the long tail of search data. It provides extensive insights into how people have successfully searched for products and services across all major search engines, including the breakdown of paid and organic traffic (you can read more about the Experian Hitwise product offering in "Tying SEO to Conversion and ROI" in Chapter 10).

Hitwise Search Intelligence provides the following features:

- Timely information on search terms your specific competitors use
- Market-specific results, for taking advantage of cultural differences on how people search locally
- Information on terms that have been "clicked on" before and after visiting a website or industry(also known as upstream and downstream activity)

Figure 5-20 shows an example of the most popular search terms used 1 to 5 searches prior to visiting eBay.

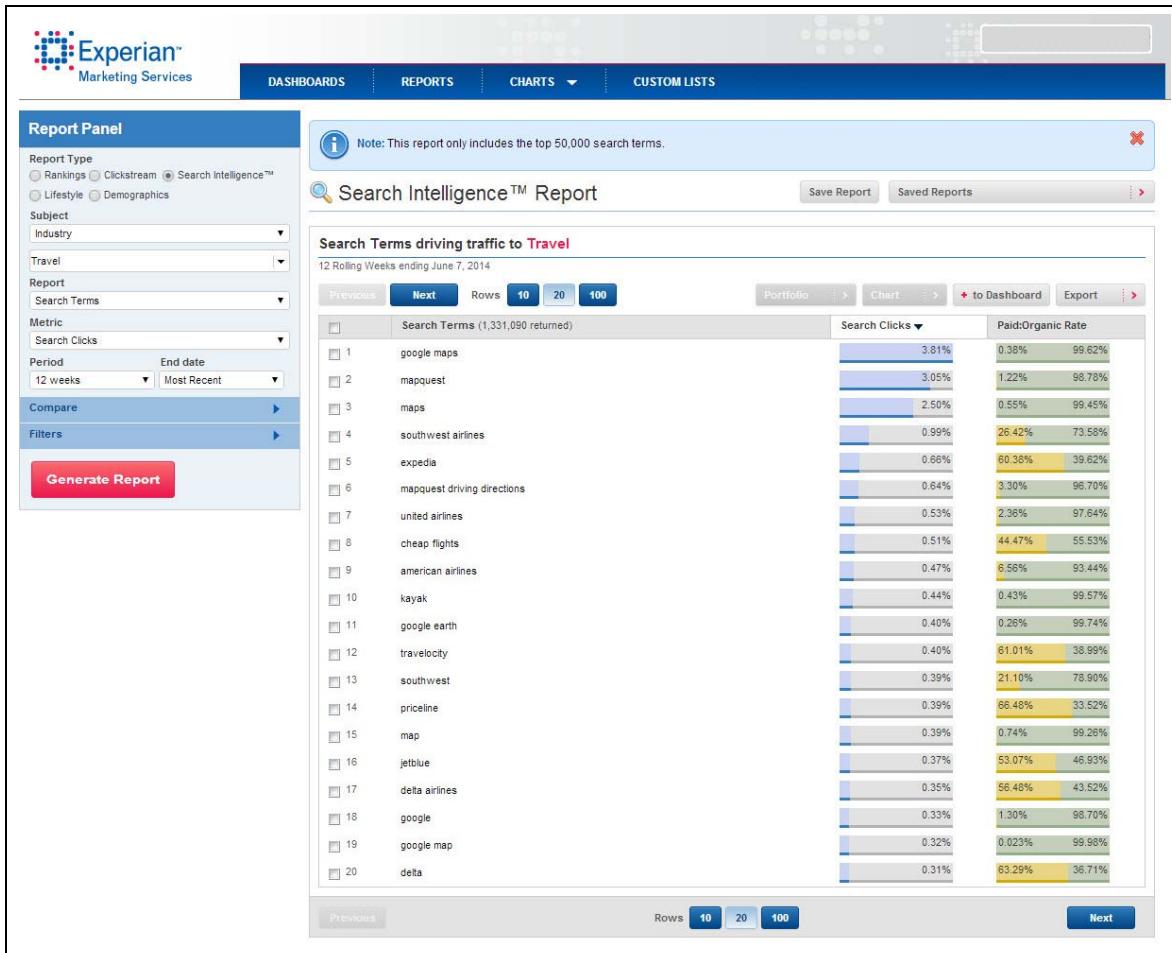


Figure 5-20. Hitwise “popular search terms” report

The ability to see actual keyword data on your competitors is an extremely potent feature. You can see what is working for them and what is not. This type of information is very powerful and can give you a significant edge over the competition.

You can also focus more directly on search term suggestions, as shown in Figure 5-21, which depicts a screen shot for terms related to *travel*.

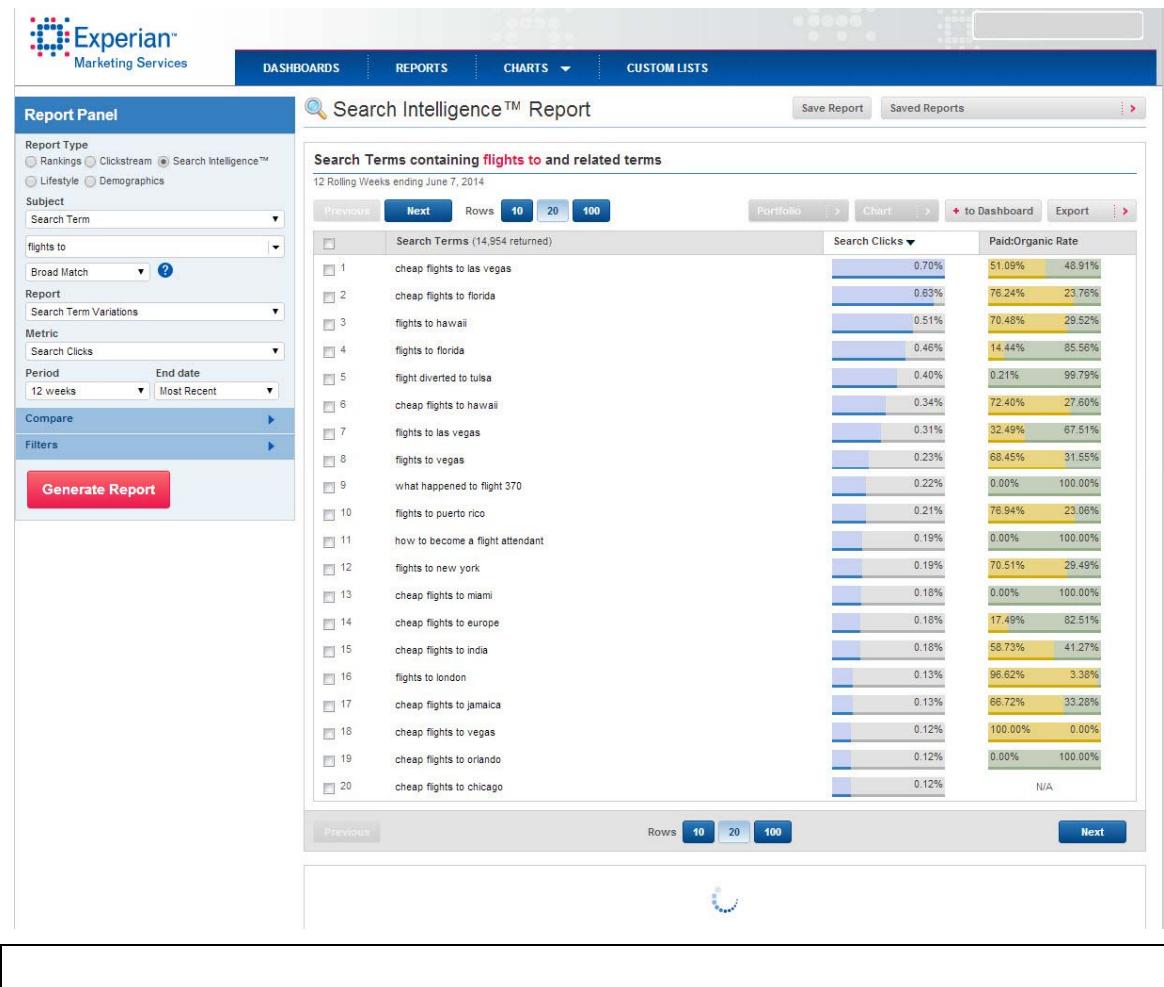


Figure 5-21. Hitwise Search Term Suggestions tool

Where it gets its data

Hitwise derives its data from more than 25 million people's interaction with the Internet (10 million from the United States). Hitwise collects anonymous Internet usage information from a combination of ISP data partnerships and opt-in panels.

How it is useful

The data is presented in percentages (the volume of searches, its success rate with searchers), which makes it very easy to compare the relative popularity of various keywords, but difficult to estimate the actual number of searches for a given term.

Cost

Hitwise is not an inexpensive tool. The website does not list pricing information, but you should be ready to spend \$20,000 if you plan to engage with this tool. Bear in mind that we have presented only a snapshot of its features, and the competitive data is extremely valuable, not just to the SEO team but to all marketing disciplines across your organization. If long-term use of the tool doesn't make sense for your team's needs, Hitwise also offers one-off analysis services for around \$5,000 a report.

comScore Search Planner

Like Hitwise, ComScore Search Planner (<https://www.comscore.com/Products/Audience-Analytics/Search-Planner>) is a tool that provides a wide range of data as a result of monitoring the behavior of actual users on the Internet. This data includes details on search terms used, as well as competitive search term analysis.

What it provides

ComScore Search Planner comprises eight modules, but two of them are particularly useful for keyword research:

Site Profile (for Site(s) X)

This module tells you what search terms and search engines are driving the most traffic to your site, to your competitor's site, and within your category.

Profile Search Terms

This module tells you the demographic profile of people searching on a set of search terms, as well as what sites these searchers tend to visit.

Figure 5-22 shows the highest-volume terms specific to the Airline category.

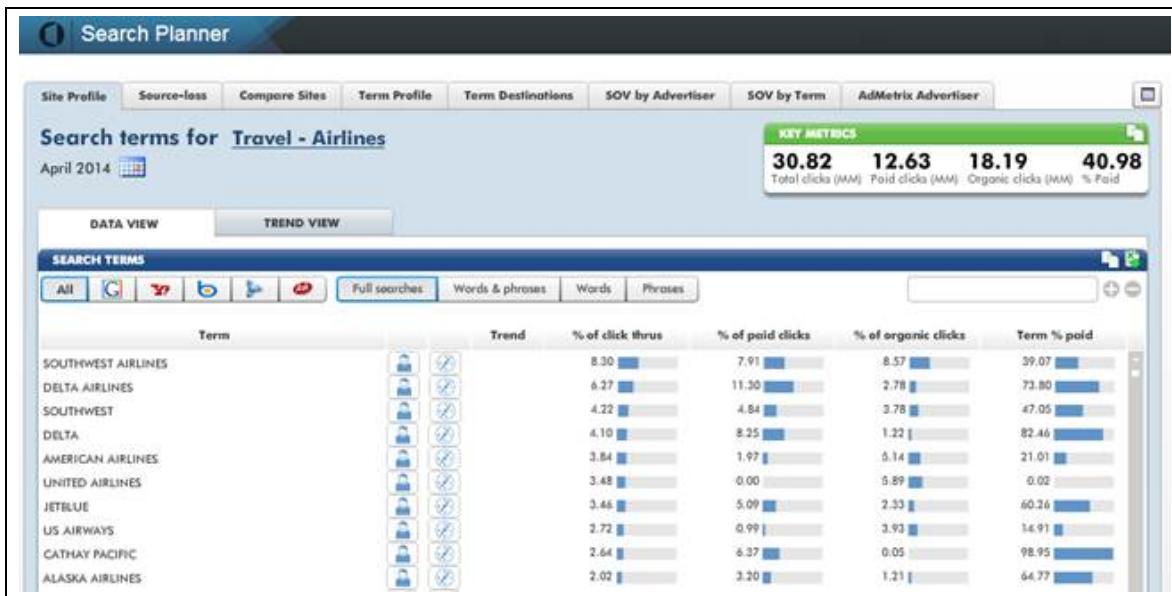


Figure 5-22. ComScore "airline search terms" report

You can also view similar data specific to a competitor's site, so you can see what search terms are driving their traffic. Another useful feature is to take a look at the search trends for an industry. This helps with the identification of seasonal behavior as you see in Figure 5-23

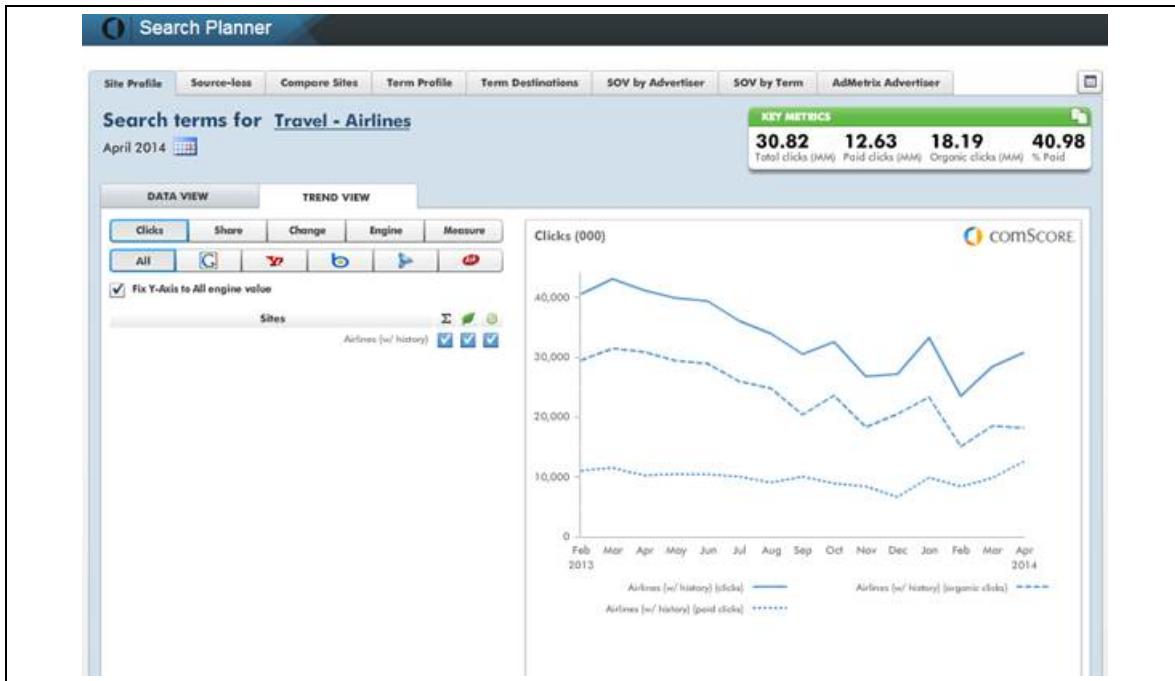


Figure 5-23. ComScore search trends report

Where it gets its data

comScore monitors the behavior of approximately 2 million users. These users have voluntarily joined comScore's research panels in return for free software, free Internet-based storage, or chances to win prizes. Sites can also opt in to adding comScore tracking on their site using unified tags. Some sites do this because it tends to result in higher numbers and better data that they can then show to potential advertisers. This helps them sell on line display advertising and obtain higher advertising rates.

How it is useful

The data is presented in percentages (the volume of searches, its success rate with searchers), which makes it very easy to compare the relative popularity of various keywords, but difficult to estimate the actual number of searches for a given term.

Cost

Pricing for comScore Search Planner is available only upon contacting the company. The primary audience for the product is mid-size to large companies with developed SEM/SEO strategies, but the company has some smaller clients as well.

Wordstream

WordStream's fast, powerful keyword tools draw from a trillion-keyword database and are completely free (previously, access was restricted to paid subscribers). The Free Keyword Tool delivers thousands of keyword suggestions, including long-tail keywords, from a constantly updated database of more than a trillion unique search queries. The Keyword Niche Finder helps you discover your most profitable pockets of keyword opportunities. This free tool delivers value-added, structured keyword suggestions, pre-organized into groups of related searches, ready to use in search marketing campaigns. Or, if you already have a list of keywords, drop them into the Keyword Grouper to get back organized keyword niches: see all related variations of a keyword at once, or drop these clusters of keywords into an ad group in AdWords.

What it provides

Wordstream provides a suite of 5 different tools for keyword research. These are:

Wordstream Keyword Suggestion Tool

This is the basic tool for generating a list of keyword suggestions along with search volume metrics.

Wordstream Keyword Niche Finder

This tool is useful when building out a list of new topics to create content for.

Wordstream Keyword Grouper

Keyword Grouper is used to mine keyword data for organic search referrals and trends.

Wordstream Negative Keyword Tool

This tool is mainly used in relation to PPC campaigns, but can assist in generating a list of terms that it is not desirable to match for.

SEO Content Creating Plug-in for Firefox

This plug-in suggests topics and keywords for new SEO pages, and tracks keyword usage as you type.

Figure 5-24 shows the output for keyword research related to the phrase catering.

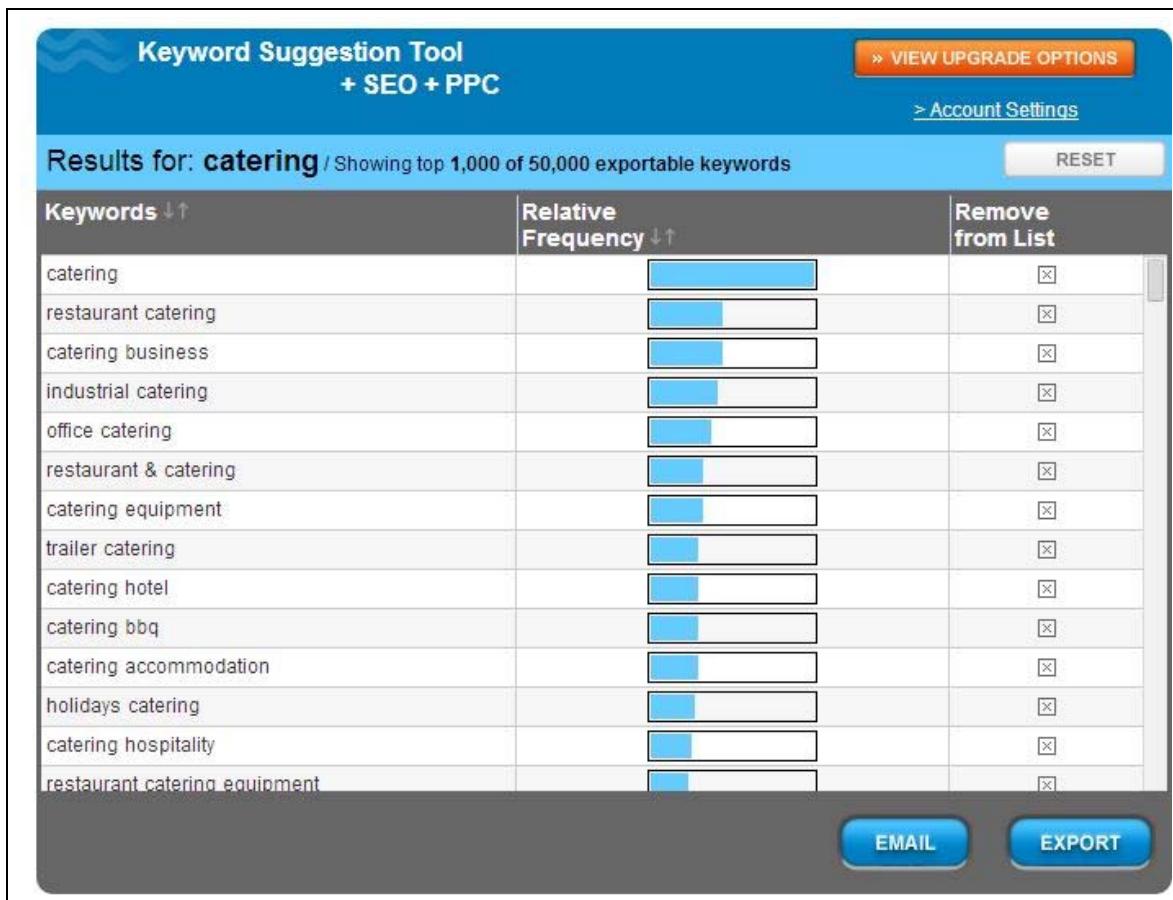


Figure 5-24. Wordstream Keyword Suggestion Tool Output for "digital camera"

The Niche Finder tool finds keywords by topic area. This is useful when looking for new topic or sub topic areas to create content for. Use this when deciding on new types of content for a site, or when first building a site. Figure 5-25 shows the output related to "used cars" and "cheap"

The screenshot shows the Keyword Niche Finder interface. At the top, it says "Keyword Niche Finder + SEO + PPC" and "Results for: used cars / Showing 100 of 100 niches". There are "VIEW UPGRADE OPTIONS" and "Account Settings" buttons. A "RESET" button is also visible.

Suggested Niches:

- used cars
- used cars » under
- used cars » pa
- used cars » ohio
- used cars » cheap** (highlighted in blue)
- used cars » ny
- used cars » city
- used cars » nc
- used cars » tx
- used cars » sales
- used cars » near
- used cars » county
- used cars » texas
- used cars » online
- used cars » fl
- used cars » area
- used cars » dealers
- used cars » uk
- used cars » north
- used cars » va
- used cars » florida

Current niche: cheap (Showing 42 keywords)

Keywords	Relative Frequency	Remove
used cars for cheap	Low	<input type="checkbox"/>
cheap used cars capital region	Low	<input type="checkbox"/>
cheap used cars under 1000	Low	<input type="checkbox"/>
branson cheap used cars	Low	<input type="checkbox"/>
cheap used cars for sale	Low	<input type="checkbox"/>
indianapolis	Low	<input type="checkbox"/>
used cars cheap cars	Low	<input type="checkbox"/>
cheap old used cars	Low	<input type="checkbox"/>
used cars cheap in erie pa	Low	<input type="checkbox"/>
cheap used cars in brandon 33510	Low	<input type="checkbox"/>
cheap used cars phoenix	Low	<input type="checkbox"/>
cheap japanese used cars	Low	<input type="checkbox"/>
used cars tledo ohio cheap	Low	<input type="checkbox"/>
used cars toledo ohio cheap	Low	<input type="checkbox"/>
new castle used cars for cheap	Low	<input type="checkbox"/>
used used cars cheap	Low	<input type="checkbox"/>

Buttons: EMAIL, EXPORT

Figure 5-25. Keyword Niches related to "education" and "jobs"

Where it gets its data

Unlike many keyword research tools, Wordstream does not source its data from Google. Instead the company buys its data from ISPs, Browser toolbar provides and search engines.

How it is useful

Wordstream is useful because it pulls its data from different sources than the search engines and goes into further depth in what it will show, exposing more of the long tail of search. In addition, Wordstream offers powerful features for organizing the keywords to make it easier to organize the keyword data to help drive your SEO strategy.

Cost

Wordstream is available in a number of different packages, one of which is a free keyword research tool (<http://www.wordstream.com/keywords>). The Wordstream Keyword Research Suite has three price points as of November 2014: \$349 per month for small businesses, \$549 per month for the Professional level and \$1099 per month at the Enterprise level.

SEMrush

SEMrush is a data provider and market leader in competitive research software. SEMrush includes the following tools and reports: Organic Research, Advertising Research, AdSense, Keyword Research, Site Audit, Domain vs. Domain, Position tracking and Backlinks. To start using SEMrush services you will need to enter a domain, URL, or keyword in the main search bar. You can switch from one section to another in the left menu to get different data to display.

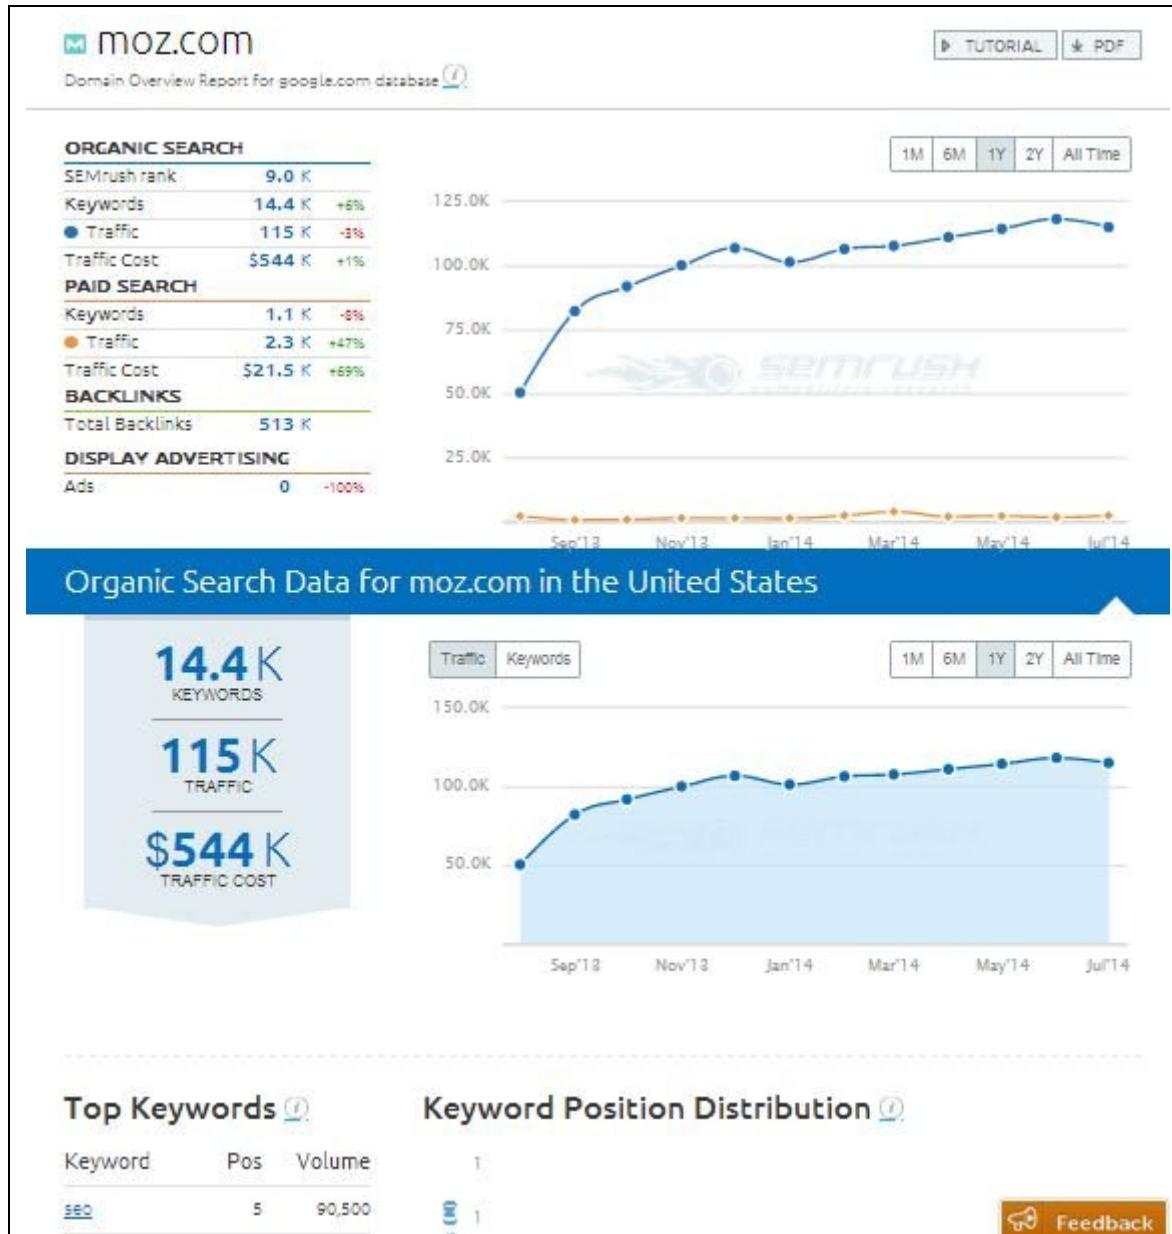


Figure 5-26. SEMrush domain level keyword data

Organic Research

The Organic Research Report offers analytics data on keywords that bring users to a website via Google's top 20 organic results. For every keyword, you can find useful information such as the position in Google search results, Volume, CPC, URL of a landing page and how it's displayed in the search results, Competition and Trend, etc.

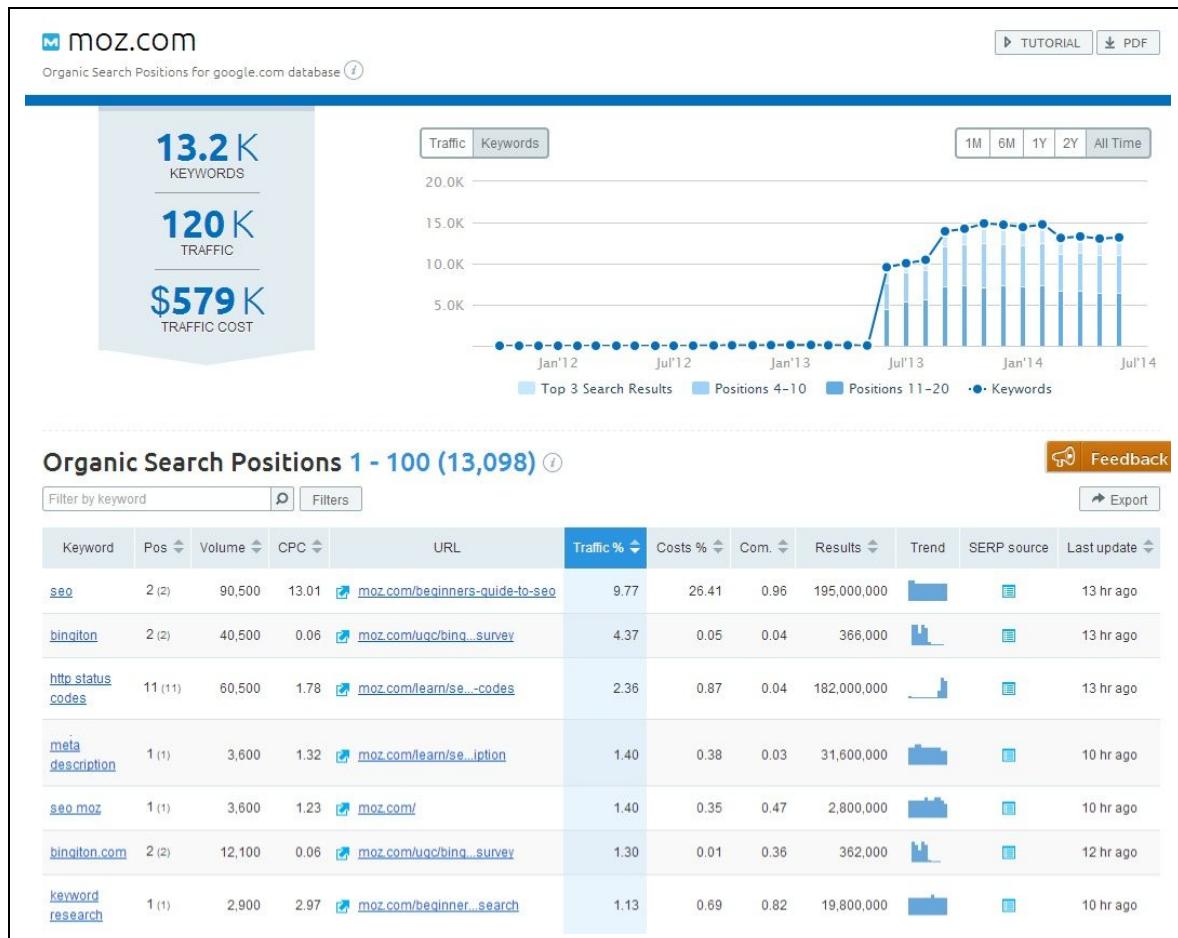


Figure 5-27. SEMrush organic research report

The Organic Keywords Graph represents the numbers of keywords for which a domain has ranked in the top 20 Google results each month on a historical basis (depending on your subscription, this could be one month, six months, one year, two years or all time). Keywords are also grouped depending on the position they bring a website (1-5, 6-10, 11+) so you can easily analyze which keywords bring you a higher or lower ranking in Google.

It's also possible to use SEMrush to do keyword research on your competitors. This is particularly useful if you have competitors who you know to be successful in their marketing efforts. By leveraging the keywords your competitors rank for, and an approximation of the traffic they get for those keywords, you can find many keywords that should be profitable for you to rank for as well.

The Competitors in Organic Research Report provides the list of your domain's competitors and common keywords. The competition level metric helps you exclude from your analysis those domains that are not your real competitors on the market. SEMrush analyzes the ratio between the number of domains' common keywords and the total number of keywords of each domain. If the total number of keywords is rather equal and the domains have a high number of common keywords they will be considered competitors and their competition level will be high.

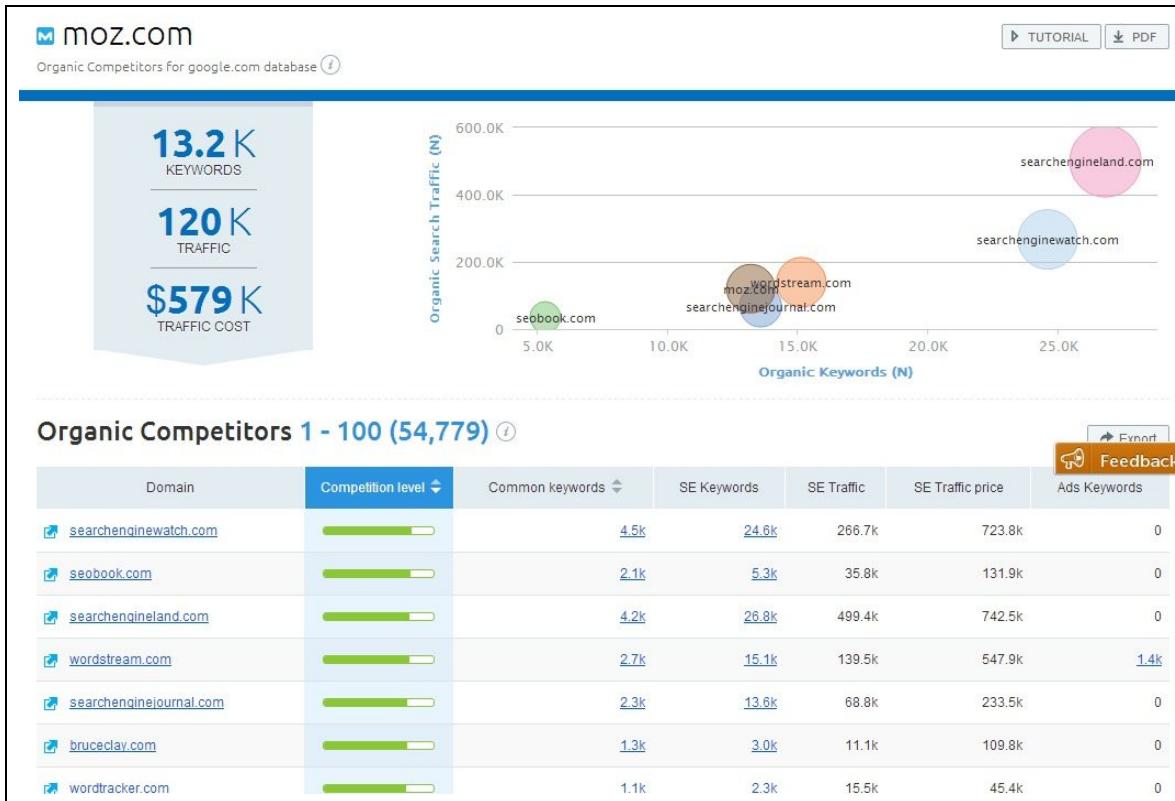


Figure 5-28. SEMrush competitors in organic report

Advertising Research

The Advertising Research section has the same data as the Organic Research one: Ads Keywords Report, Ads Keywords Graph, Competitors in Ads Report, Position Changes Report etc. This report can be particularly useful for finding keywords which are good targets for your organic ranking efforts. If you can assume that your competition are savvy online marketers, then the keywords which they actively purchase can generally be assumed to be those which drive a profitable number of conversions to their websites.

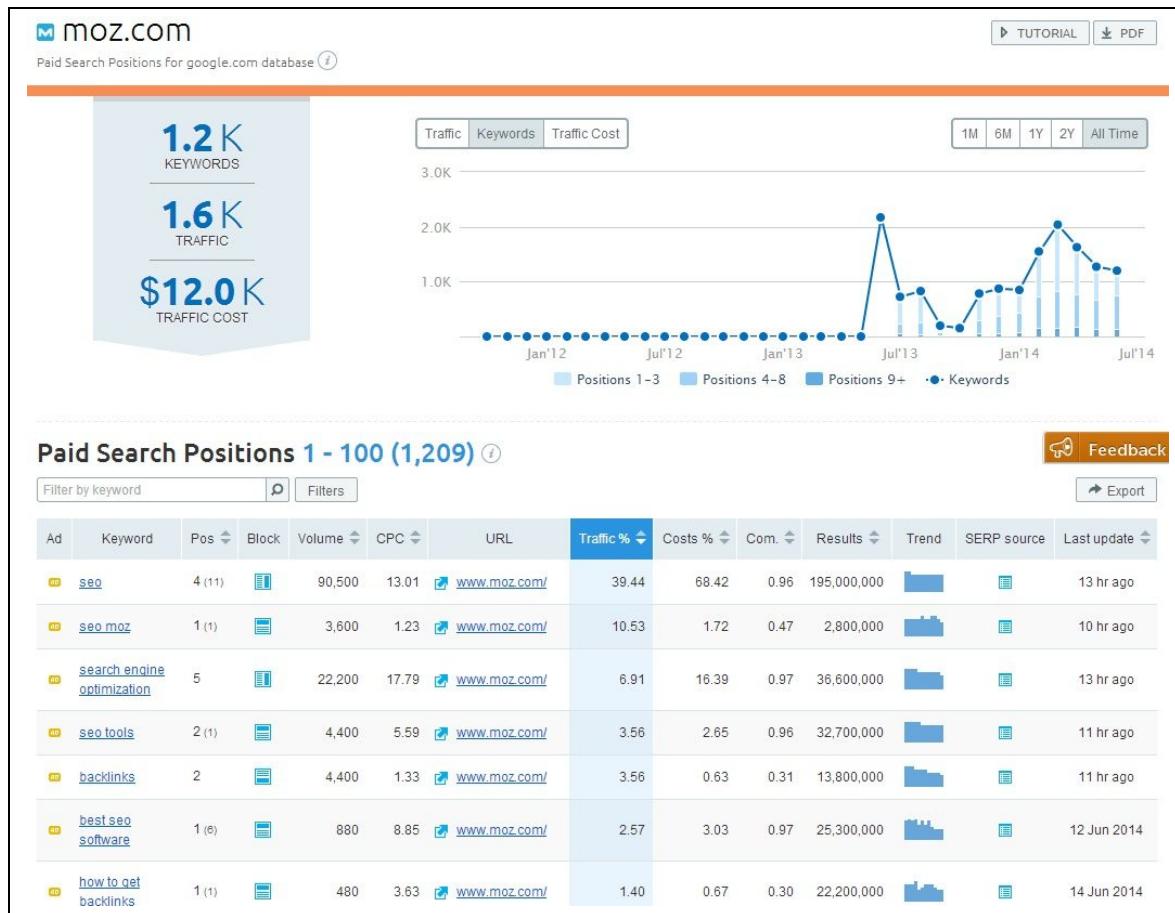


Figure 5-29. SEMrush paid search positions report

Where it gets its data

SEMrush pulls its data directly from Google search engine results and maintains dedicated databases of this information. SEMrush has 25 databases for the United States (Google and Bing), the United Kingdom, Canada, Russia, Germany, France, Spain, Italy, Israel, Belgium, The Netherlands, Denmark, Finland, Ireland, Norway, Poland, Sweden, Switzerland, Switzerland, Turkey, Brazil, Argentina, Mexico and Hong Kong.

How it is useful

SEMrush allows you to access a great deal of data about both your own domain and your competitors' domains and what keywords those domains rank for in both organic and paid search.

Cost

SEMrush offers either a Pro membership for \$79.95 per month which allows for 10,000 results per report and up to 3000 reports per day or a Guru level account at \$149.95 per month which allows for 30,000 results per report and up to 5000 reports per day.

Searchmetrics

The Searchmetrics Suite is an SEO platform which is generally used by enterprises. The software provides a holistic approach regarding both technically based SEO data analysis and content-focused optimization.

While it provides a wide range of SEO metrics and tracking for the purposes of keyword research, you can monitor keywords, rankings and URLs, optimize on-page structure and content, evaluate links, track social engagement in real-time and compare domains on a project basis. In addition, the platform provides you with automatic technical on-page analyzing and proactive recommendations including site structure optimization for up to several million sub-pages.

One key feature of Searchmetrics is its reverse engineering technology which allows you to reconstruct “not provided” keyword and traffic data on a URL level via the fusion of various data sources and algorithmic calculations.

For the purposes of straightforward keyword research Searchmetrics offers a keyword discovery tool which ranks any keyword on search volume, competition, cost (if you were to pay for the keyword rather than ranking for organically), and “search volume spread” which shows any seasonality in the keyword. It shows similar and related keywords and allows you to click through on any keyword to see the top sites ranking for that term.

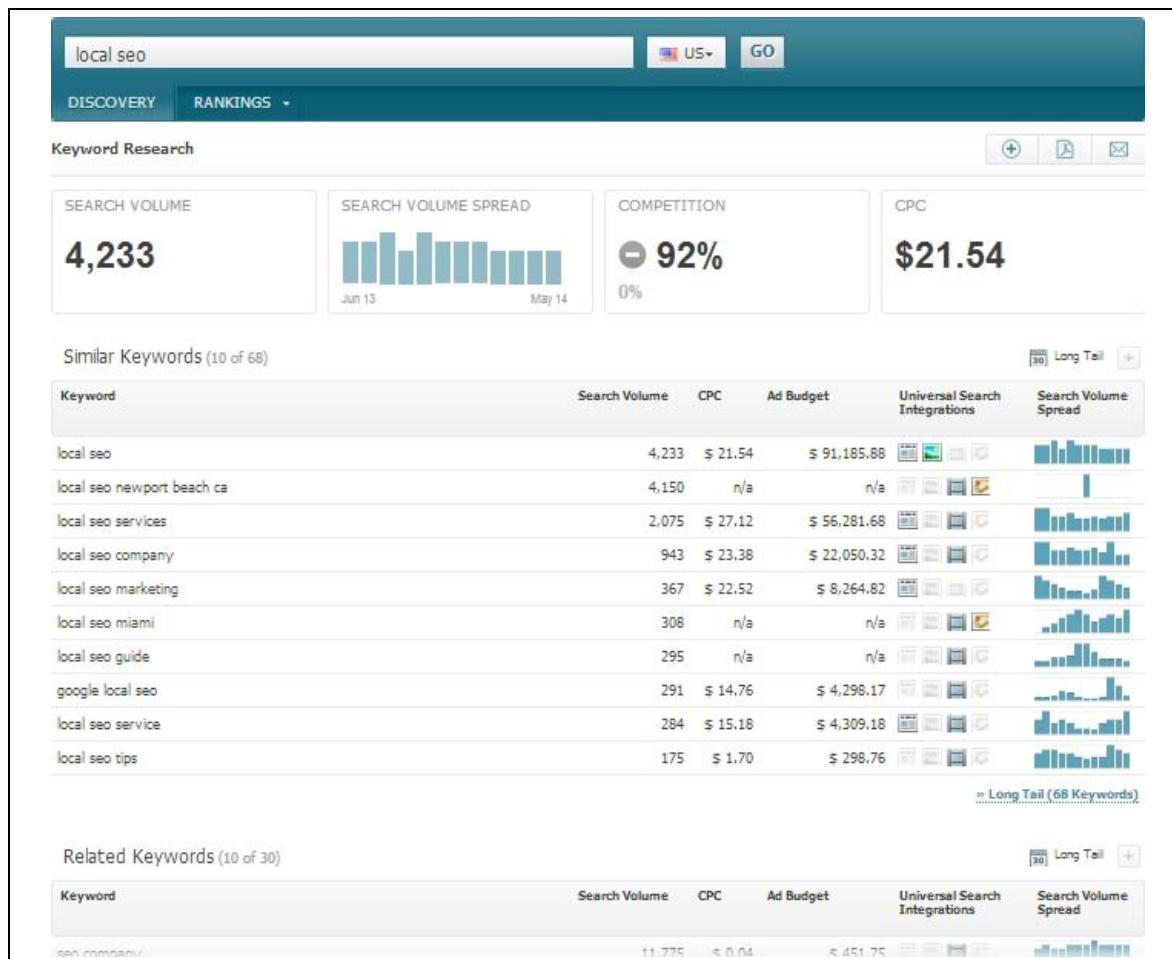


Figure 5-30. Searchmetrics keyword discovery report

Also of interest with Searchmetrics is that they offer competitive keyword research on both paid and organic traffic. Enter a domain in the search box and Searchmetrics will return a list of competitors in either organic or paid search along with data on how many keywords the domain shares in common with its competitors, how many keywords are not shared with the competitors, relative traffic volume for the competitors and other data on the competitors' rankings.

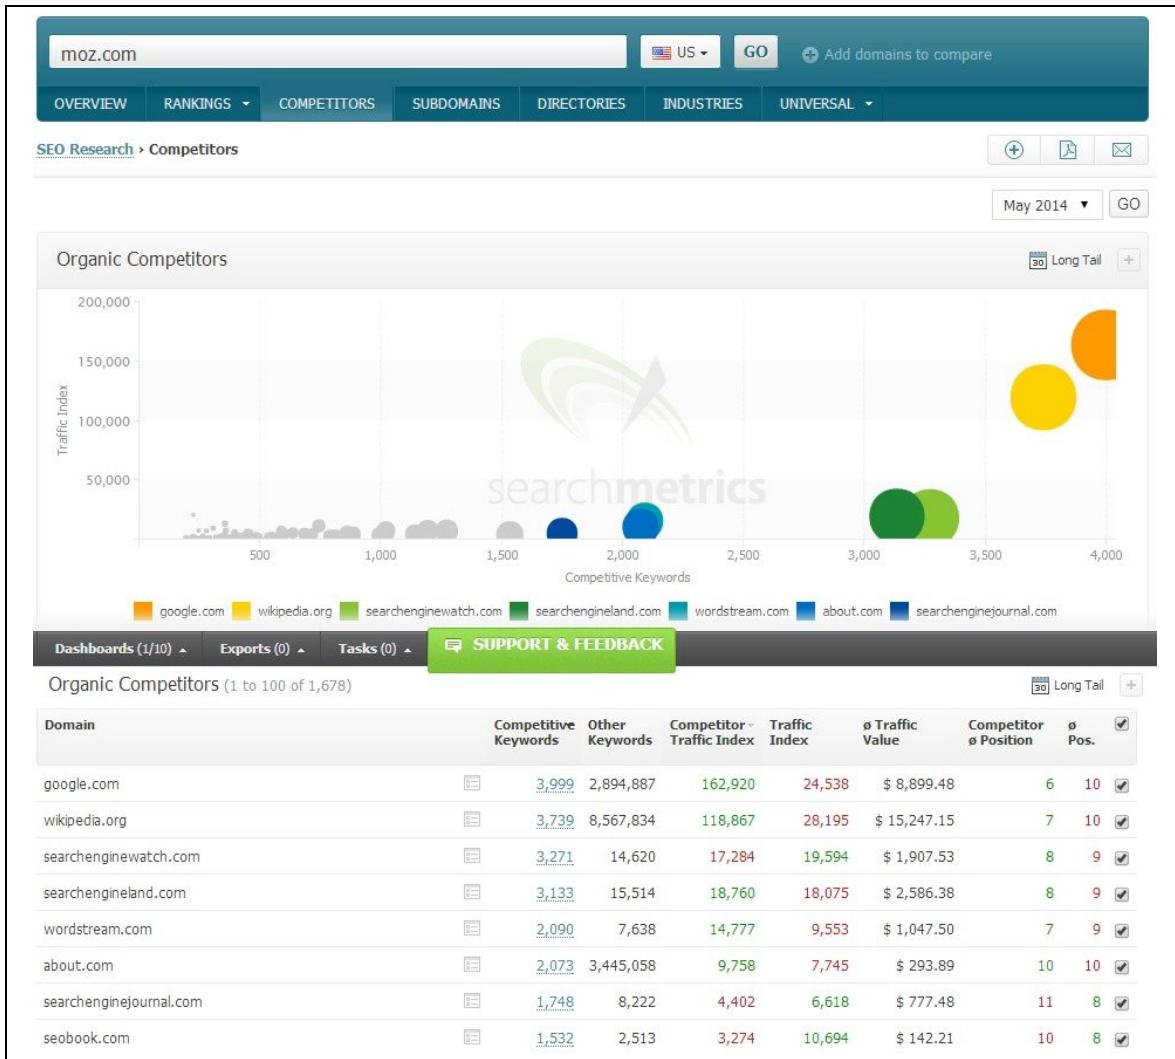


Figure 5-31. Searchmetrics organic competitors report

By clicking through any entry in the “competitive keywords” column you can delve more deeply into which keywords the sites might have in common. To get more details on exactly which keywords your competition is ranking for you can also look at the “Rankings” report under SEO Research. This report gives you in-depth information into thousands of keywords which your competitor ranks for including the keyword, the URL which ranks for it, their position, and an estimate of both search volume and what the term would cost to purchase via pay per click advertising. This report can give you a great deal of insight into which keywords are important to your competitors and thus in many cases which keywords should be important to you.

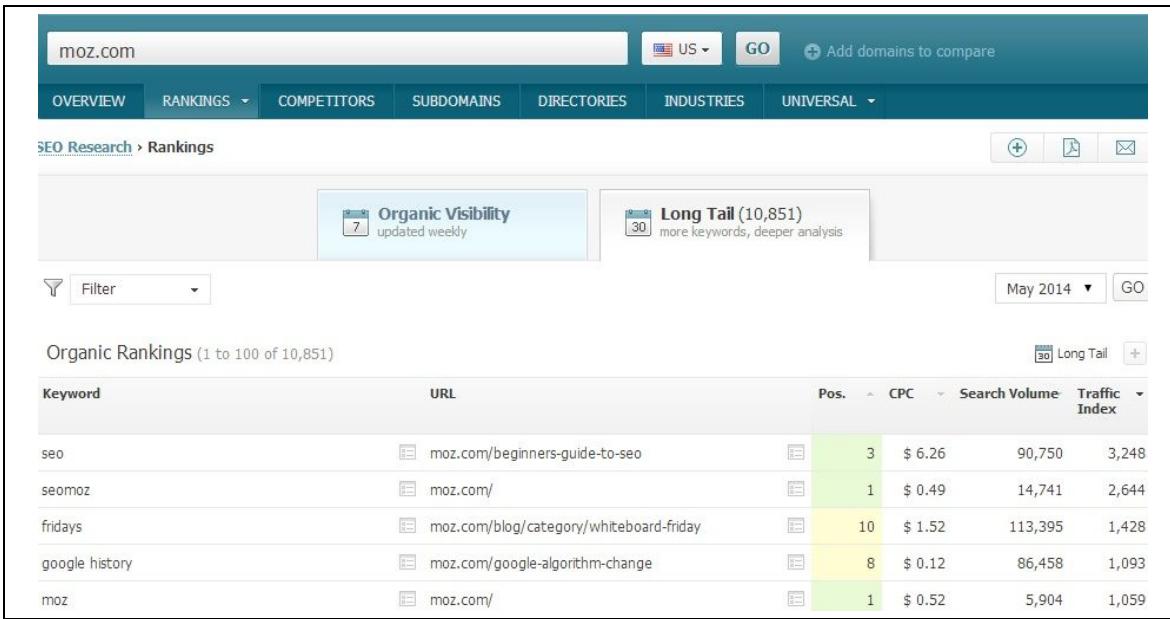


Figure 5-32. Searchmetrics keyword rankings report

Where it gets its data

The Google SERPs and Google Keyword Planner

How it is useful

Searchmetrics allows you track most of your SEO efforts while providing many on page analyses and recommendations. In addition to this it provides a wealth of keyword information both at the keyword level and in mining your competitors' keyword data.

Cost

Searchmetrics Essentials which allows for research access SEO, PPC, keyword, etc. data and detailed competitor analysis cost \$69 per month. The full Searchmetrics Suite tracking, research, and optimization solution is quoted on a case-by-case basis.

Other Tools of Interest

There are many other keyword tools available on the market. Some of the more interesting ones are:

Google Autocomplete

Previously known as Google Suggest, this Google feature automatically "suggests" queries related to or with the same beginning or related to a word. For example, typing in "lilac" will reveal suggestions below of "lilac essential oil", "lilac wine", and "lilac tree". This was a somewhat hidden tool that, thankfully, Google brought front and center to default Google Search. Google won't tell you many times "digital cameras" has been searched for, but because it appears at the top of the list of suggestions, you can infer that it was probably searched for more often than those phrase that appear below it. This can give important insight into what searchers are looking for, or what they search for in relation to a specific topic. For example, does a company have an Autocomplete suggestion of "[company name] scam"? If so, they may have a PR nightmare on their hands.

Note: Suggestions are personalized based on the user's location, e.g. "wet n wild phoenix" is a suggestion that showed up when typing the letter "w" into Google in Phoenix, AZ. And in mobile search, you may notice apps listed as a suggestion.

Google Instant

Start with the basic search input box, but look for what Google reveals as you type. Google gives results that appear in the search page even before you are complete with the search query. As you refine this search further, these results below appear and change to fit your refined search. This benefits the searcher by eliminating the need to hit the "Search" button, and often limits the need to finish the query. Google estimates that this function has the capability to save 2-5 seconds per search.

Soovle

Soovle (<http://www.soovle.com>) shows you real-time search terms as you type them, ordered by popularity, just like Google Suggest. In fact, it's a one-stop-shop that taps into those features of those top search engines and much more. It also polls YouTube, Amazon.com, Bing, Wikipedia and Answers.com for top related search terms, refreshing dynamically each time you pause during your typing. This tool allows you to tap into seven top resources at once.

The screenshot shows the Soovle search interface with the keyword 'golf' typed into the search bar. The results are displayed in a grid format across four columns, each representing a different search engine or service. The columns are: Soovle (with a sidebar for 'Soovald Link'), Google (with a sidebar for 'golfsmith'), Amazon.com (with a sidebar for 'golf'), and Bing (with a sidebar for 'golfsmith.com'). The results are as follows:

Soovle	Google	Amazon.com	Bing
golf	golfsmith	golf	golfsmith.com
golfo	golf channel	golf accessories	golfnow.com
golfech	golfinow	golf clubs	golfsmith
golf course	golfers warehouse	golf balls	golf channel
golf club	golf	golf shoes	golf galaxy
golf ball	golf galaxy	golf umbrella	golf digest
golf world	golf and ski	golf cart	golf cart
golf cart	golf digest	golf gps	golf gps
golf channel	golf clubs	golf pants	golf pants
golfsmith	golfwrx	golf bag	golf bag

Below the search bar, there is a note: "Try the icons or hit the right-arrow key to change engines..." followed by a row of icons for Google, Ask, YouTube, Bing, AOL, and Wikipedia. The entire interface is framed by a large bracket on the left and right sides.

Figure 5-33. Soovle results using keyword "golf"

YouTube Suggest

If you begin to type in a search query on YouTube is offers search suggestions as shown in Figure 5-34. The suggestions are the most popular variants of the search query you have typed so far.

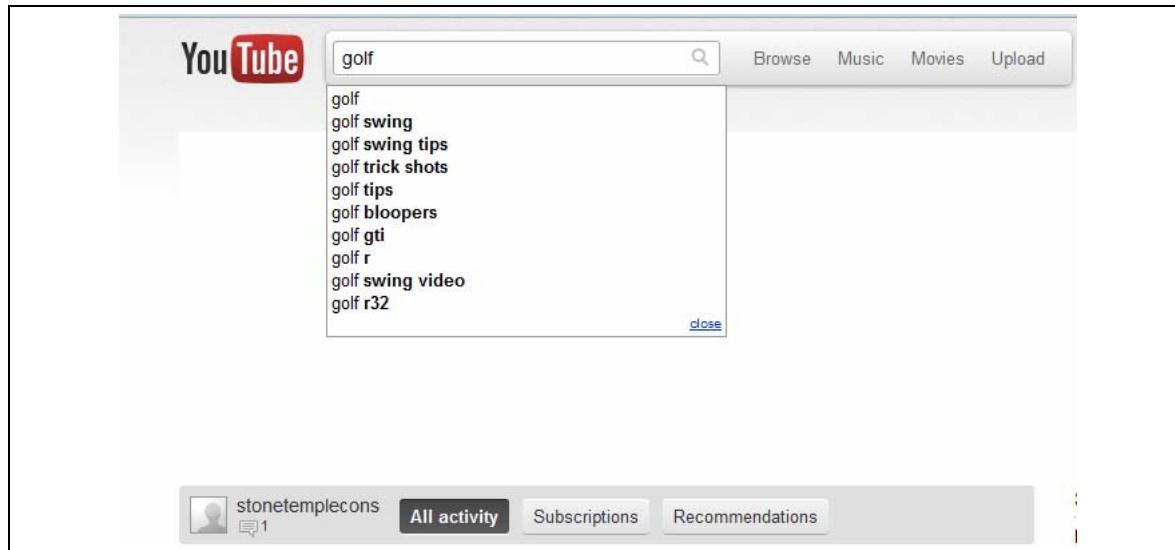


Figure 5-34. YouTube Suggest

Google AdWords Display Planner

On September 1, 2014 Google discontinued the former YouTube Keyword Tool. Similar to the Google Keyword Planner you must have a Google Adwords account (though you do not need to spend any money on the account to access the Display Planner). You can enter a landing page or select from a list of topics to get recommendations on what terms might drive the most impressions. To get the video only keyword recommendations first search for your desired topic then select “Video” under the “Ad formats and sizes” option on the left-hand side of the page. Selecting “Individual targeting ideas” will return a list of recommended keywords, their relevance to your select topic, and historical CPM, number of cookies per week, and number of impressions per week.

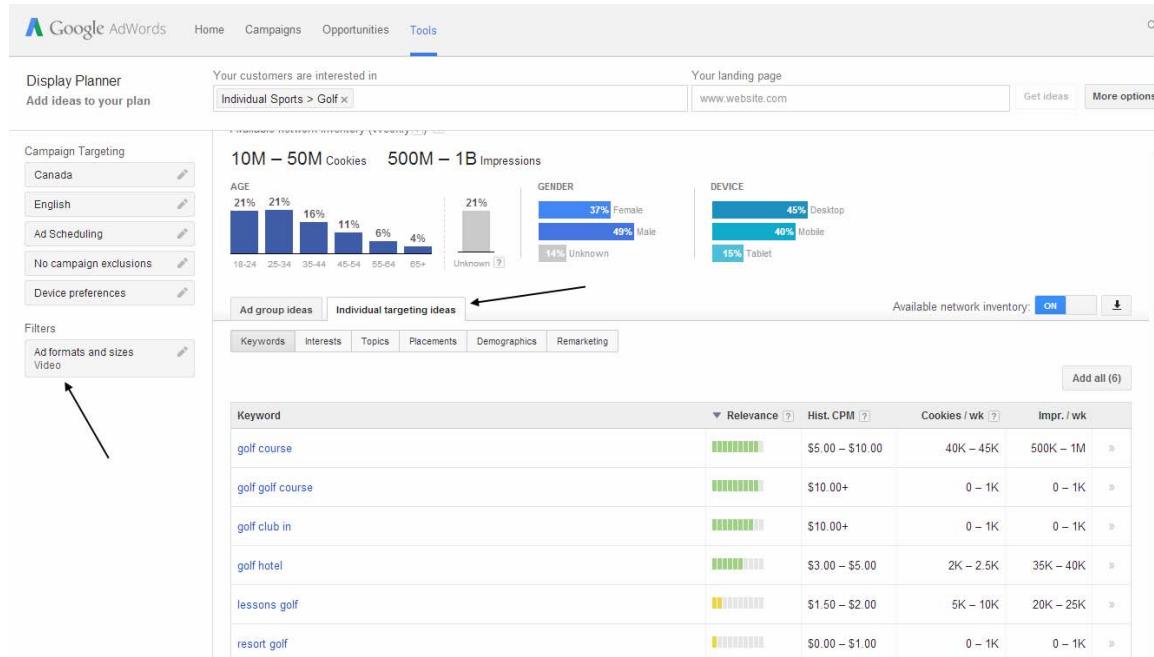


Figure 5-35. Google AdWords display planner

Google Trends for Video

Just as you can use Google Trends (<http://www.google.com/trends/>) to search for the relative interest of various keyword terms in the main Google search results you can also use it to look for keyword trends and suggestions for YouTube. To do this search as usual for a term on Google trends then select the “YouTube” option from the top navigation bar.

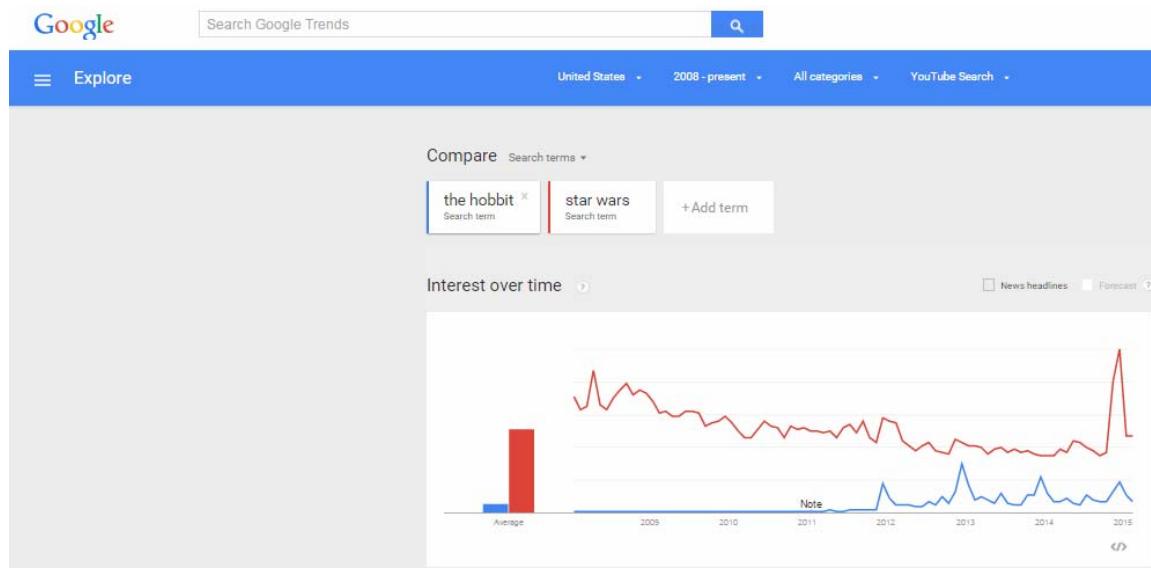


Figure 5-36. Google Trends for YouTube search

Google Trends will also return a list of topics and keyword queries related to your search with a rating for the relative interest of each.

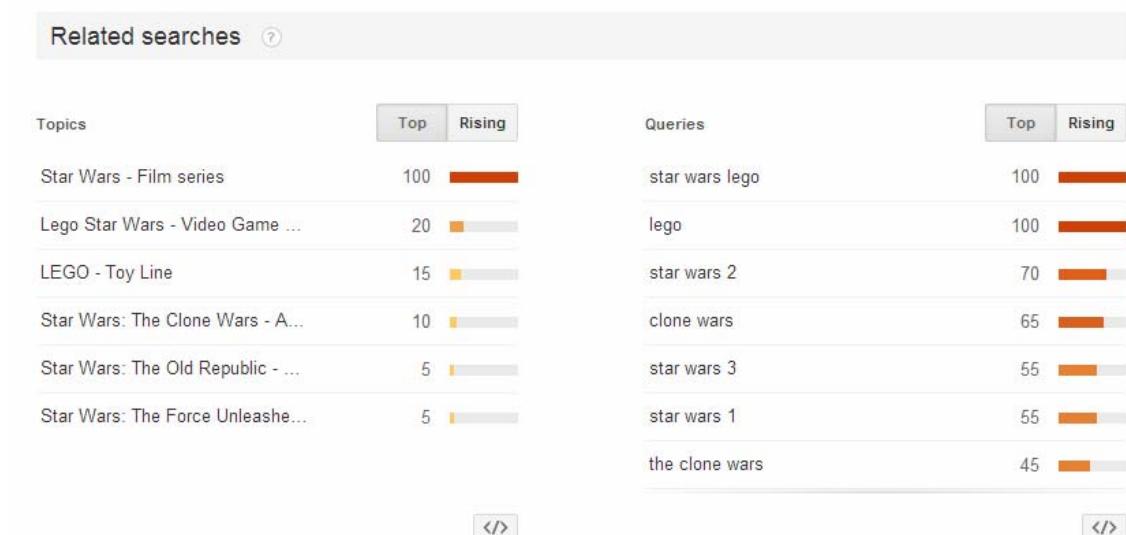
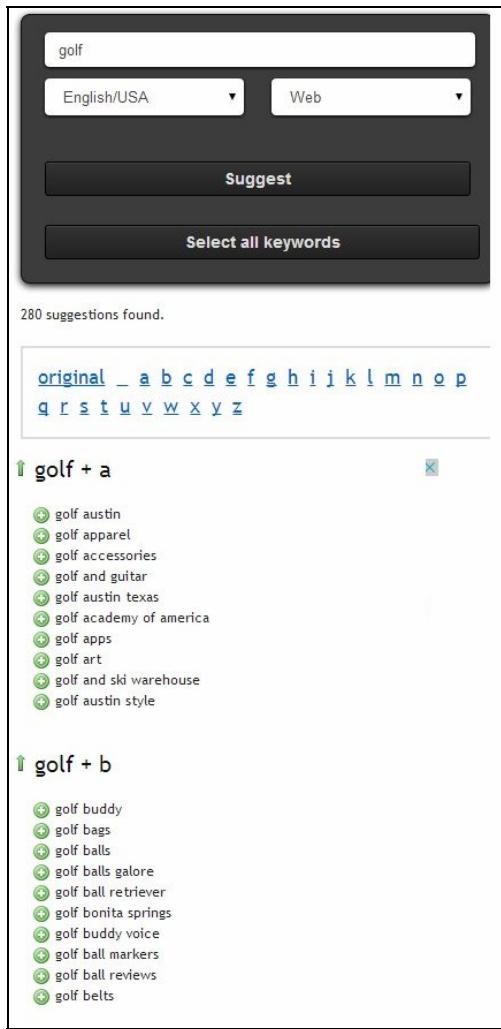


Figure 5-37. Google Trends related searches

Ubersuggest

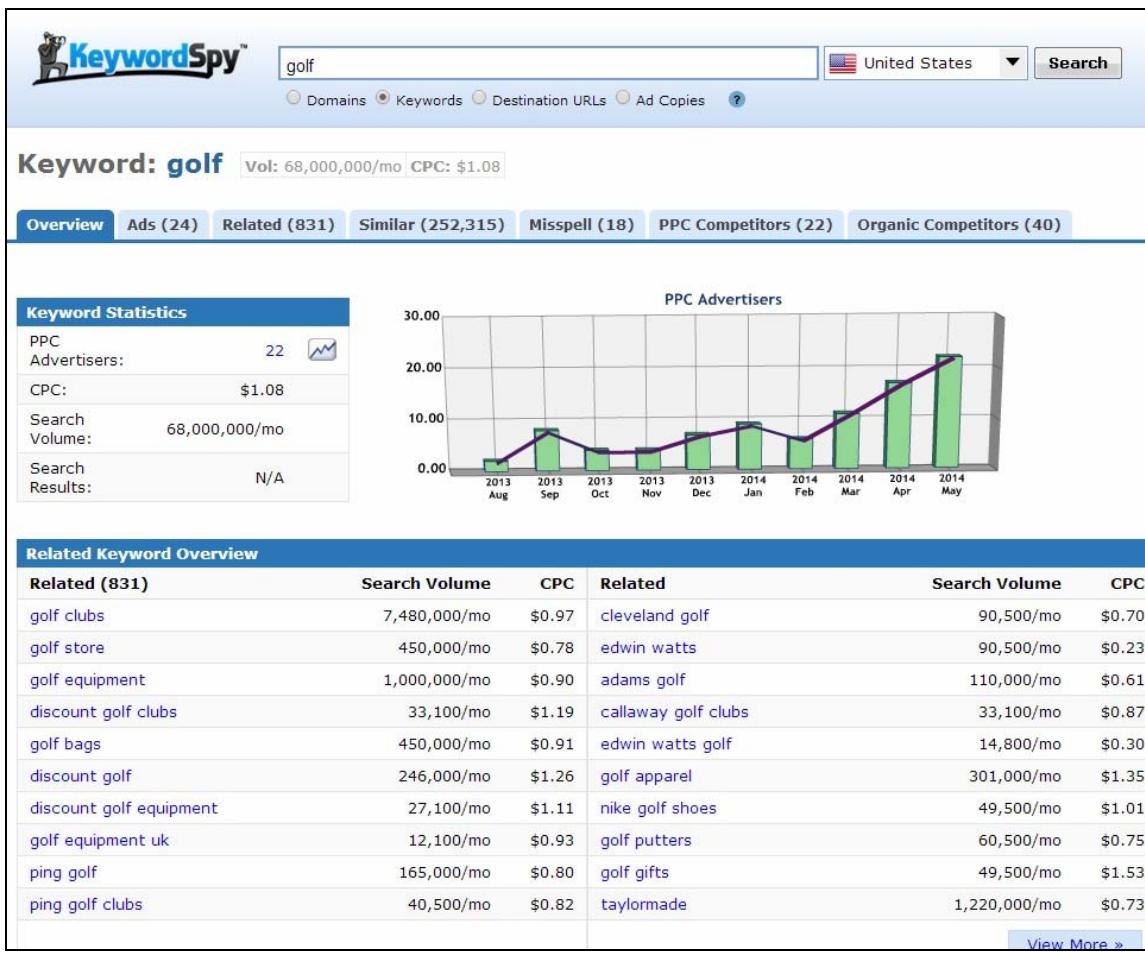
Ubersuggest (<http://ubersuggest.org/>) is a free tool based on Google Suggest. It runs a bunch of variants based on the base term that you have entered. For example, if you enter the query golf, Ubersuggest will automatically pull the suggestions for golf + terms beginning with a, golf + terms beginning with b, etc., all the way through to golf + z and then golf + 0, golf + 1, etc.



KeywordSpy

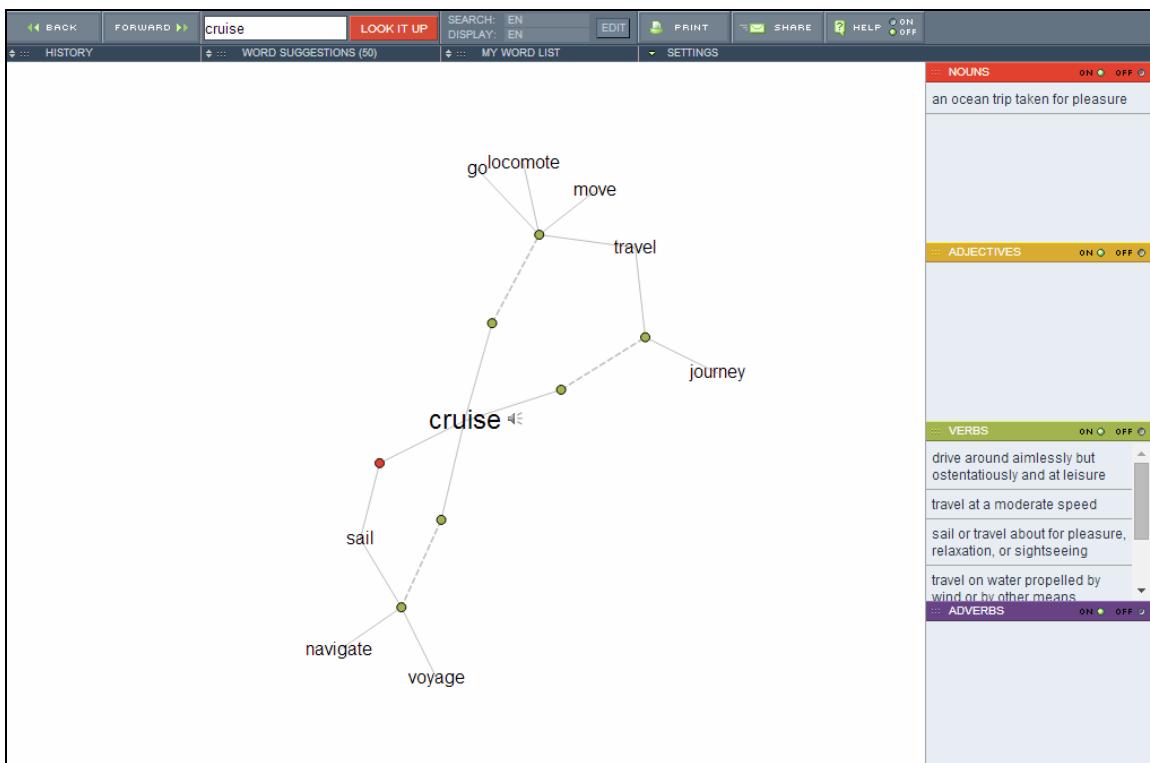
KeywordSpy (<http://www.keywordspy.com/>) is a paid tool which ranges in price from \$89.95 monthly to \$139.95 monthly for basic packages. They also have agency and enterprise packages available at higher rates. KeywordSpy offers a variety of services but primarily they offer insights into paid search volumes and costs per click. This can give you insights into keyword variations and show which keywords are more competitive and thus likely to be more likely to lead to higher ROI should you be able to drive traffic from

them. The tool also gives insight into search volumes for related terms, similar terms, misspellings and shows lists of the top ranking PPC and organic competitors for a given query.



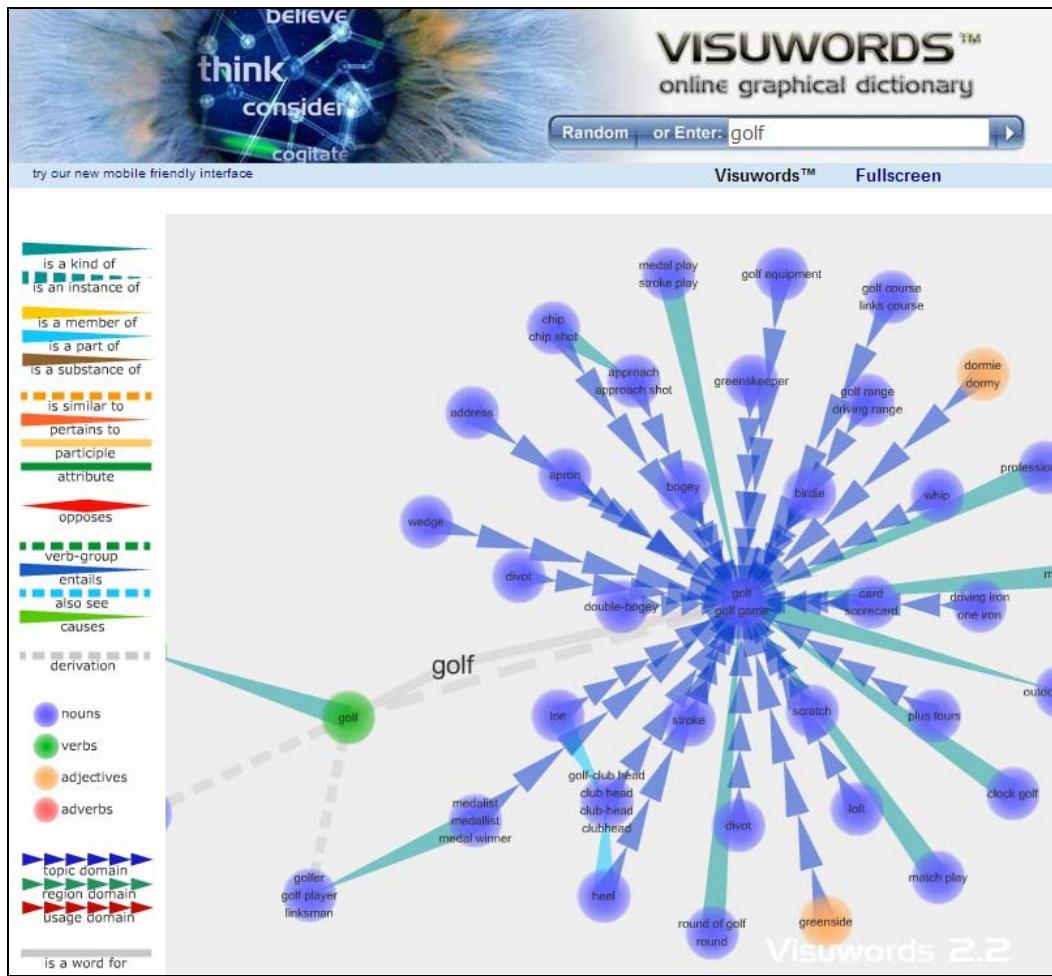
Visual Thesaurus

Visual Thesaurus is another graphical tool showing the definition of the keyword entered as well as related terms and concepts. Clicking on each node in the diagram opens a map of the variations in related terms and concepts for that node. It's another tool you can use for brainstorming new keywords and concepts, which you can then enter into the keyword tools that give more specific keyword data. It has a nominal cost of \$2.95 per month or \$19.95 per year.



Visuwords

Visuwords (<http://visuwords.com/>) is a free tool which provides a dynamic graphical interface of words and concepts related to the search term you enter. It acts as both a dictionary and thesaurus as well as providing links to related concepts. These related concepts may be keyword areas not previously considered in other keyword research. Visuwords is a graphical representation of the data in Princeton's WordNet Search (<http://wordnetweb.princeton.edu/perl/webwn>)



Determining Keyword Value & Potential ROI

Once you have the raw keyword data from the research with your favorite tools, you need to analyze which keywords have the highest value and the highest ROI. Unfortunately, there are no simple ways to do this, but we will review some of the things you can do in this section.

Estimating Value, Relevance, and Conversion Rates

When researching keywords for your site, it is important to judge each keyword's value, relevance, and potential conversion rate. If a keyword is strong in all three criteria, it is almost certainly a keyword you want to plan to optimize for within your site.

Determining keyword value

When judging the value of a keyword, you should contemplate how useful the term is for your site. How will your site benefit from targeting these keywords?

Identifying relevant keywords

To identify relevant, high-quality keywords, ask yourself the following questions:

How relevant is the term/phrase to the content, services, products, or information on your site?

Terms that are highly relevant will convert better than terms that are ancillary to your content's focus.

Assuming a visitor who searches for that term clicks on your result in the SERPs, what is the likelihood that he or she will perform a desired action on your site (make a purchase, subscribe to a newsletter, etc.), create a link to your site, influence others to visit, or engage in off-site engagement by promoting your content in social media? You can also test this by setting up a PPC campaign, and buying clicks on a given keyword and seeing how it converts for you.

It is a good idea to target keywords that indicate imminent action (e.g., *buy cranium board game, best prices for honda civic*), because searchers are more likely to perform the corresponding action on your site when they search for those terms than they are for terms such as *honda civic* or *cranium board game*. Your click-through/conversion rates are likely to be higher if you target keywords that indicate the intent behind the search.

How many people who search for this term will come to your site and leave dissatisfied?

Pay attention to your site's content and compare it to what other sites in the top results are offering—are these sites doing or offering something that you haven't thought of? Do you feel as though these sites offer a more positive user experience? If so, see what you can learn from these sites and possibly emulate on your own.

You can also use an analytics program and check to see which of your pages have the highest abandonment rates. See what you can change on those pages to improve user experience and increase users' level of enjoyment when using your site.

It is important to categorize your keywords into high and low relevance. Generally, keywords of higher relevance will be more beneficial to your site in that they best represent your site as a whole. If, when judging the relevance of a keyword, you answer "yes" to the preceding questions, you've found a highly relevant term and should include it in your targeting.

Keywords with lower relevance than those that lead to conversions can still be great terms to target. A keyword might be relevant to your site's content but have a low relevance to your business model. In this case, if you target that keyword, when a user clicks on your site and finds the content to be valuable, he or she is more likely to return to the site, remember your brand, and potentially link to your site or suggest it to a friend. Low-relevance keywords, therefore, present a great opportunity to strengthen the branding of your site. This type of brand value can lead to return visits by those users when they are more likely to convert.

Determining conversion rates

A common misconception is that a conversion refers only to the purchase of an item on your site. However, many different types of actions users perform can be defined as a conversion, and they are worth tracking and segmenting (you can read more about this in "Key Performance Indicators for Long Tail SEO" in Chapter 11).

The many different types of conversions create distinct opportunities for targeting various keywords. Although one keyword may work well for purchase conversions, another may be well suited to get users to subscribe to something on your site. Regardless of what type of conversion you are optimizing for, you should strive to have each keyword that you intentionally target convert well, meaning it should be relatively successful at getting searchers to click through to your site and, consequently, perform a specific action.

To know which keywords to target now (and which to pursue later), it is essential to understand the demand for a given term or phrase, as well as the work required to achieve those rankings. If your competitors block the top 10 results and you're just starting out on the Web, the uphill battle for rankings can take months, or even years, of effort, bearing little to no fruit. This is why it is essential to understand keyword competitiveness, or keyword difficulty.

To get a rough idea of the level of competition faced for a particular term or phrase, the following metrics are valuable:

- Search Demand Volume (how many people are searching for this keyword)
- Number of Paid Search Competitors and bid prices to get in the top four positions
- Strength (age, link power, targeting, and relevance) of the Top 10 Results
- Number of Search Results (it can be valuable to use advanced operators such as “exact search” or the `allintitle` and `allinurl` operators here as well; for more on search operators, see co-author Stephan Spencer’s book, *Google Power Search*: <http://www.amazon.com/Google-Power-Search-ebook/dp/B005EI85P0> for more on using these specialized searches for keyword research)

Moz offers a Keyword Difficulty Tool (<http://moz.com/tools/keyword-difficulty>) that does a good job collecting all of these metrics and providing a comparative score for any given search term or phrase.

Testing Ad Campaign Runs and Third-Party Search Data

One of the things we have emphasized in this chapter is the imprecise nature of the data that keyword tools provide. This is inherent in the fact that the data sources each tool uses are limited. It turns out that there is a way to get much more precise and accurate data—the trick is to make use of Google AdWords.

Start with the keywords you are interested in and implement a simple AdWords campaign. Assuming that you are implementing this campaign solely to get keyword volume data, target position #4 or #5. This should be high enough that your ads run all the time, but low enough that the cost of collecting this data won’t be too high.

Once you have run this for a few days, take a look at your AdWords reports, and identify the number of impressions generated for the keyword. Although this data is straight from Google, it is important to remember that the advertisers’ ads may not be running all the time, so more (possibly even significantly more) impressions may be available.

Next, think about the value of achieving certain rankings in the organic results. You can come up with a good estimate of that as well. The key here is to leverage what you know about how click-through rates vary based on organic search position. Table 5-2 depicts click-through rates by SERP position based on data by DigitalRelevance (available at <http://digital.relevance.com/blog/a-tale-of-two-studies-google-vs-bing-ctr>)

Table 5-2. Click-through rates by SERP position

Organic position	Click-through rate
1	18.2%
2	10.05%
3	7.22%
4	4.81%
5	3.09%

This data, of course, is aggregated across a very large number of searches on AOL, so it serves only as an estimate; but if you are in position #1, the estimate is that 18.2% of the people who will search on a term will click on your result. In the case of our term that is searched 52 times per day, the site in the #1 position will get 9-10 clicks per day.

There are certain search terms in which these estimates do not apply. For example, if the user searches on a brand term, the focus on the #1 position is much, much higher. Publishers in lower positions still get some traffic, but at lower percentages than we’ve outlined here.

So, now you have a working estimate of the search volume and the number of clicks per day that the term will deliver. Can you get an estimate of conversion rates as well? Yes, you can. This requires only a simple

extension of the AdWords campaign: implement conversion tracking, with the free capability provided by Google or via another method at your disposal.

Once you have that in place, look at how your AdWords campaign performs. If the conversion rate is a lofty 5% for one keyword and 3% for another keyword, chances are that your organic search conversion rates for those two keywords will vary by a similar amount. Be aware, though, that although paid search results get significantly less traffic than organic search results, paid click-throughs do tend to convert at a somewhat higher rate (1.25 to 1.5 times, according to industry data). Using the preceding example, this suggests that we will get a little less than one conversion per day as a result of being in the #1 position.

This data is great to have in hand. However, it does not mean you should use this methodology instead of other keyword tools. It takes time to implement, and time to collect the data. The keyword tools will get you started in real time. Nonetheless, using AdWords or MSN adCenter can provide you with some great data.

We also don't recommend that you get obsessed with tracking your rankings on keywords. As we discussed in Benchmarking Current Rankings in Chapter 4, it is not possible to do this as accurately as you might think, and it can lead to poor decision making. Nonetheless, using this type of AdWords testing can help you get a sense as to real search volumes and the importance of particular keywords to your SEO campaign.

Using Landing Page Optimization

Landing page optimization (sometimes also called *conversion optimization*) is the practice of actively testing multiple variations of a web page (or website) to see which one performs the best. Typically, this is done as part of an effort to improve the conversion performance of the site.

The simplest form of this type of test is called an A/B test. A/B tests involve creating two different versions of a page, and then selecting a version at random to show to a new visitor to the site (old visitors get the version they saw the last time they visited). You then measure the behavior of the visitors to the two different versions to see which group of visitors completes conversions on the site. You have to be careful to wait until you have a statistically significant amount of data to draw a conclusion. Once you have this data you can analyze it and decide on more tests, or simply pick the winner and move on.

Multivariate testing is a bit more complex, because it involves more than two variations in the test. In addition, you can mix and match multiple variations. For example, you may want to try two different logos, two different calls to action, three different page titles, two different color schemes, and so on. In multivariate testing, any combination of your elements could be what is shown to a particular visitor. Obviously, more data (visits and actions) is required to draw a conclusion than in a simple A/B test.

Landing page optimization can help in determining the value of a keyword because one of the elements you might want to test is the impact on conversion of variations of a keyword or phrase in the page title, the page header, and in other strategic places on the page. One variation of the test would use one keyword or phrase and the other variation would use a different one.

You can then see which keyword provides the best results from a conversion perspective. This data can provide you with an interesting measure of keyword value—its ability to help you convert your visitors. However, landing page optimization is not something that is practical to use to perform SEO tests (i.e., to see which version of a page ranks higher), as SEO tests can take weeks or even months to see results.

Leveraging the Long Tail of Keyword Demand

As we discussed at the beginning of this chapter, the long tail of search is where 70% of search queries occur. Only 30% of those precious queries happen in the more obvious terms that people use, the so-called “head terms.” Another way to underscore this is that in May 2007, Google Vice President Udi Manber indicated that 25% of all search queries that Google receives on a given day are queries that Google is seeing for the first time. You can think of this as the “ultra-long tail.”

The long tail of search queries in a given industry is typically not visible via any of the major keyword research services or search engine ad databases (Google AdWords, Bing Ads). In these instances, there is a method to find those terms that can carry value, but it requires a good amount of research and analysis.

With this in mind, let's outline a few methods for finding long tail terms.

Extracting Terms from Relevant Web Pages

One source for long tail terms is web pages that do well for searches that are relevant to your target market. Here is a basic process for finding those pages and extracting that information from them:

1. Extract the top 10 to 50 most common search phrases at the head of the distribution graph from your existing keyword research in the industry.
2. Search Google and Bing for each term.
3. For each page in the top 10 to 30 results, extract the unique text on the page.
4. Remove stop words and filter by phrase size.
5. Remove instances of terms/phrases already in your keyword research database.
6. Sort through the most common remnants first, and comb as far down as you feel is valuable.

Through this process, you are basically text-mining relevant documents on the subject of your industry/service/product for terms that, although lower in search volume, have a reasonable degree of relation. When using this process, it is imperative to have human eyes reviewing the extracted data to make sure it passes the “common sense” test. You may even find additional terms at the head of the keyword distribution graph.

You can expand on this method in the following ways:

- Text-mine Technorati or Delicious for relevant results.
- Use documents purely from specific types of results—local, academic—to focus your keyword mining efforts.
- Mine forum threads on your subject matter. You could even use *inurl:forum* in the searches to grab conversational keywords.

This methodology is highly effective. The return on this research has a direct relationship to the amount of effort you expend (and how deep you dig).

Mining Keyword Research Tools

Although looking into keyword research tools for long tail data has significant limitations, there are still ways to do it. For example, if you own a chain of pizza restaurants in 50 cities across the country and you want to discover long tail terms, you can. Let's look at the tail end of Google Keyword Planner's output for a combined search on *Orlando Pizza*, *San Diego Pizza*, and *San Jose Pizza* (see Figure 5-38).

Your product or service

Orlando Pizza, San Diego Pizza, San Jose Pizza

Get ideas Modify search

Ad group ideas Keyword ideas

Keyword (by relevance)

Avg. monthly searches Competition Suggested bid Ad impr. share Add to plan

Keyword (by relevance)	Avg. monthly searches	Competition	Suggested bid	Ad impr. share	Add to plan
best italian in little italy san diego	40	Low	\$0.46	0%	»
pieology pizza san jose	40	Low	-	0%	»
24 hour delivery san jose	40	High	\$1.83	0%	»
san diego pizza restaurants	40	Medium	\$0.73	0%	»
dessert downtown san diego	40	Low	-	0%	»
amato pizza san jose	40	Low	-	0%	»
amore pizza san diego	40	Low	-	0%	»
amici pizza san jose	40	Low	-	0%	»
san diego real estate mls	40	High	\$2.80	0%	»
best restaurants san diego ca	40	Medium	\$2.50	0%	»
pizza coupons san diego	40	High	\$1.95	0%	»

Figure 5-38. Extracting long tail data from Google AdWords keyword planner

san diego pizza restaurants (40 searches per month on average) is an example of a valid long tail term. If some people search for *san diego pizza restaurants*, it becomes quite likely that others search for *orlando pizza restaurants*. It does not show in this data, because the volume of queries available to the keyword research tool is limited. All we are doing with these combined searches is giving the search tools more data to work with.

The takeaway remains valid: apply these logical long tail extensions across all of your cities, even though the keyword tool shows it for only one, and you're likely to attract search queries for those keywords.

Identifying Long Tail Patterns

You can also take another stab at determining long tail information. As a hypothetical example using *digital camera*, here are 40 searches for two different brands and models of digital cameras that have been pulled (for this demonstration) from the KeywordDiscovery database that received only one search:

- *consumer comments on nikon 5.1 mp coolpix l3 digital camera*
- *new nikon coolpix p3 8 1 mp digital camera memory*
- *nikon 3 2 mp coolpix digital camera*
- *nikon 51 mp coolpix s1 digital camera and cradle*
- *nikon 6 mp coolpix digital camera*
- *nikon 7 1 mp coolpix 7900 digital camera*
- *nikon 81 mp coolpix 8800 digital camera*
- *nikon coolpix 4800 4 mp digital camera*
- *nikon coolpix 5200 51 mp digital camera*
- *nikon coolpix 5400 51 mp digital camera*
- *nikon coolpix 6.0 mp digital camera*
- *nikon coolpix 8700 8mp 8x zoom digital camera 8 mp*

- *nikon coolpix l2 6.0 mp digital camera*
- *nikon coolpix l3 6 mp digital camera usa warranty*
- *nikon coolpix p2 51 mp digital camera*
- *best buy sony cybershot dsc t7 51 mp digital camera*
- *brand new sony cybershot dsc h1 51 mp digital camera*
- *camera digital sony cybershot 51 mp*
- *sony - cybershot 10.1 mp digital camera*
- *sony - cybershot 6.0 mp digital camera*
- *sony 5 mp cybershot dsc t9 digital camera*
- *sony 72 mp cybershot dsc p200 digital camera information*
- *sony 72 mp cybershot dsc w7 digital camera*
- *sony 72 mp digital still camera cybershot rebate*
- *sony cybershot 10.1 mp digital camera*
- *sony cybershot 7 2mp digital camera 7 2 mp*
- *sony cybershot 72mp dsc w7 digital camera 72 mp*
- *sony cybershot 81 mp digital camera*
- *sony cybershot digital camera 5.1 mp*
- *sony cybershot digital camera 6 mp*
- *sony cybershot dsc 1 81 mp digital camera review*
- *sony cybershot dsc h1 51 mp digital camera*
- *sony cybershot dsc w30 6 mp digital camera*
- *sony cybershot dscs40 41 mp digital camera 3x opt zoom*
- *sony dsc p73 cybershot digital camera 41 mp p 73*
- *sony dsc p8 cybershot 32 mp digital camera*
- *sony dsc s60 cybershot digital camera 4 1 mp*
- *sony dsc s85 cybershot 41 mp digital still camera*
- *sony dsc t1 cybershot digital camera 5 0 mp*
- *sony dsc t1 cybershot digital camera 50 mp t 1*

Our goal is to determine whether there are any universal patterns that searchers tend to use when searching. Within this subset of searches, a number of patterns stand out:

- Approximately 48% begin with the brand name and end with *digital camera*.
- Approximately 35% are ordered brand, model name, model number, megapixel, *digital camera*.
- Approximately 22.5% are ordered brand, megapixel, model name, *digital camera*.
- A whopping 60% follow the overall pattern of brand, model name, *digital camera*.

You might also notice that, at least in this example, qualifiers such as *new*, a specific store name, and a reference to consumer comments tend to precede the search phrases, whereas features and product-related qualifiers such as *memory*, *3x opt zoom*, *warranty*, *cradle*, *information*, and even a repeat of the megapixels or model number tend to be appended to the search phrases.

Remember, this is purely a limited, hypothetical example and certainly is not meant to be statistically accurate; the goal here is to reveal different search term patterns to aid in determining the best groupings of long tail keywords to target.

Editorial Content Strategies for Long Tail Targeting

One of the most difficult aspects of capturing traffic from the long tail of search is creating relevant, targeted content. As we discussed in “Determining Searcher Intent and Delivering Relevant, Fresh Content” in Chapter 2, search engines rely on lexical analysis to determine what a web page is about. As a result, your chances of showing up for a long tail phrase are greatly increased if you have that long tail phrase, or at least all the words that make up the long tail phrase, on your page. Let’s look at why this may be challenging by checking out what phrases Wordtracker returns when we enter *canon digital camera* (see Figure 5-39).

1 <input checked="" type="checkbox"/> canon digital camera (search)	2,582	698,000	39.8	0.004
2 <input checked="" type="checkbox"/> canon digital cameras (search)	564	1,350,000	2.39	0.000
3 <input checked="" type="checkbox"/> canon digital camera reviews (search)	392	21,000	88.0	0.019
4 <input checked="" type="checkbox"/> canon digital camera in vietnam (search)	288	10	27,600	28.8
5 <input checked="" type="checkbox"/> compare canon digital cameras (search)	244	1,220	180	0.201
6 <input checked="" type="checkbox"/> canon digital camera accessories (search)	117	8,950	6.42	0.013
7 <input checked="" type="checkbox"/> canon digital camera software (search)	115	2,800	8.45	0.041
8 <input checked="" type="checkbox"/> canon digital camera windows 7 screen fix (search)	105	0	-	-
9 <input checked="" type="checkbox"/> canon digital camera lens (search)	100	4,820	3.59	0.021
0 <input checked="" type="checkbox"/> cheap canon digital cameras (search)	95	6,390	10.4	0.015

Figure 5-39. Sample long tail data

Already, with the eighth phrase returned (*canon digital camera windows 7 screen fix*), you can see the challenge. If you are trying to sell Canon digital cameras you are probably not going to work that phrase into your page copy.

The best approach is to use the long tail research techniques we discussed in this chapter and identify the major patterns, or the major words that appear across different long tail scenarios, and then work those words into your copy. Don’t force it and make pages that appear foolish to a user.

Make sure the writers remain focused on producing quality content. From a long tail perspective, more text is better because it creates more possible long tail matches, but there are limits to that too. Don’t put a 1,000-word article on your site unless it makes sense to your users for you to do so.

User-Generated Content Strategies for Long Tail Targeting

User-generated content (UGC) can be a great way to obtain lots of content that will help attract long tail traffic. Popular ways of doing that include forums, reviews, blog comments, and a way to upload videos or images, among others. As users submit content, they do the hard work of writing the text you need to capitalize on the long tail.

There are some downsides to UGC, though. Generally speaking, you need to moderate it to make sure people are not contributing objectionable material you don’t want on your site, or copyrighted material from other sources (and potential duplicate content issues, depending on how the UGC content is published to your site). Even if you get community members to participate, you will still need to manage them.

In addition, you need to have a strategy for getting the process started. In the case of a forum, you need to develop a critical mass of users to establish a real community. If you don’t establish this critical mass, a high percentage of the posts you receive will be one form of spam or another. To make UGC work, you need one or more of the following:

- Significant existing daily site traffic. How much depends on how vertically oriented your community is intended to be. Narrowly focused topics can get going with a smaller number of users.

- A way to generate a lot of buzz to generate site traffic.
- Compelling supporting content.

If you can succeed at this, you'll give life to a machine that produces long tail content on an ongoing basis with comparatively low effort.

Trending, Seasonality, and Seasonal Fluctuations in Keyword Demand

One of the subtleties of keyword research, and of any fully developed SEO strategy, is that the use of keywords and phrases varies significantly over time. For instance, major holidays inevitably lead to bursts of keyword volume related to those holidays. Examples could be searches such as *Halloween costumes*, *gift ideas for Christmas*, or *Valentine's candy*.

If you want to write holiday-related content, it will be important to have your site visible in the SERPs for those search queries prior to that holiday's buying season so that you'll get optimum traffic for those terms. And since the search engines take considerable time in ranking your pages, advance preparation is required. To investigate this further, let's examine the Google Trends data for a period of 12 months for the search term *halloween costumes* (see Figure 5-40).

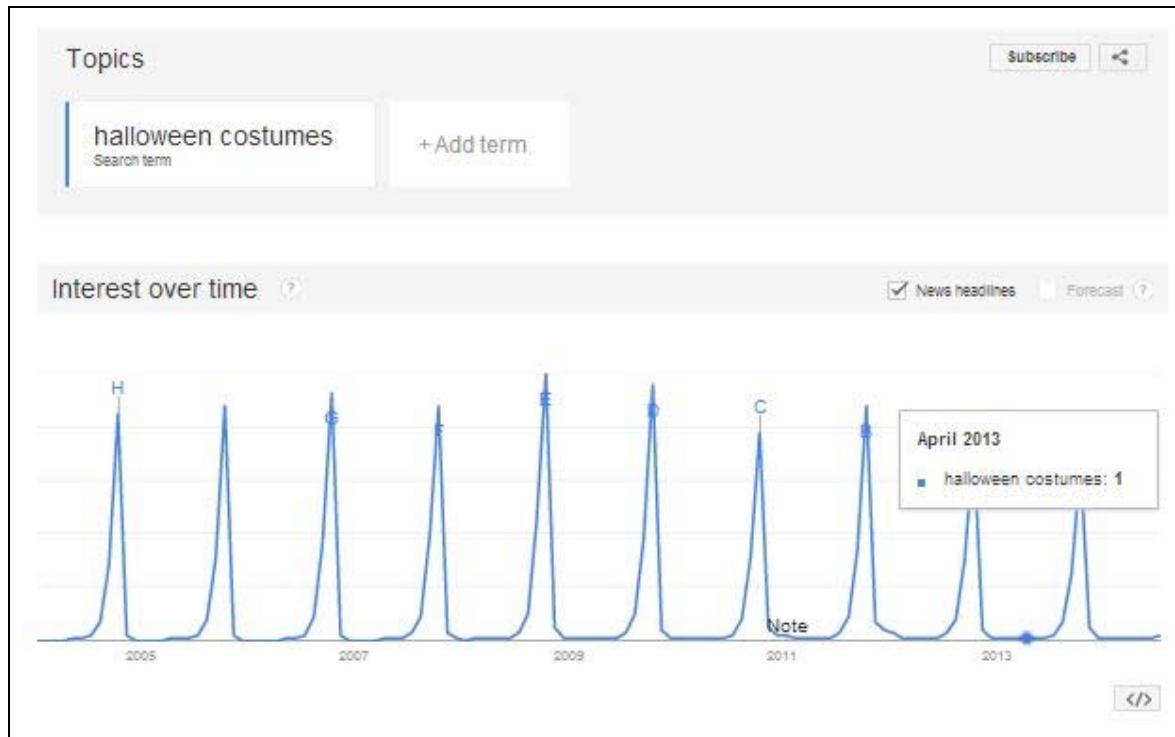


Figure 5-40. Google Trends highly seasonal data example

As you can see, searches begin gaining traction toward the end of August and into autumn; thus, if you are doing SEO for Halloween-related terms, you would want to have the related content and links in place by the beginning of the summer so that search engines can find and index your content, and therefore you're more visible to searchers when they start doing research. A long-term SEO approach would take this into consideration as part of the overall strategy for the site. Searches start consistently increasing toward the end of September. You can keep trying more examples. With Valentine's Day, the searches start in mid-December.

A similar pattern emerges for Christmas-related searches. Figure 5-41 shows an example for *firecrackers*. Searches start consistently increasing in early June. You can keep trying more examples. With Valentine's Day, the searches start in mid-December.

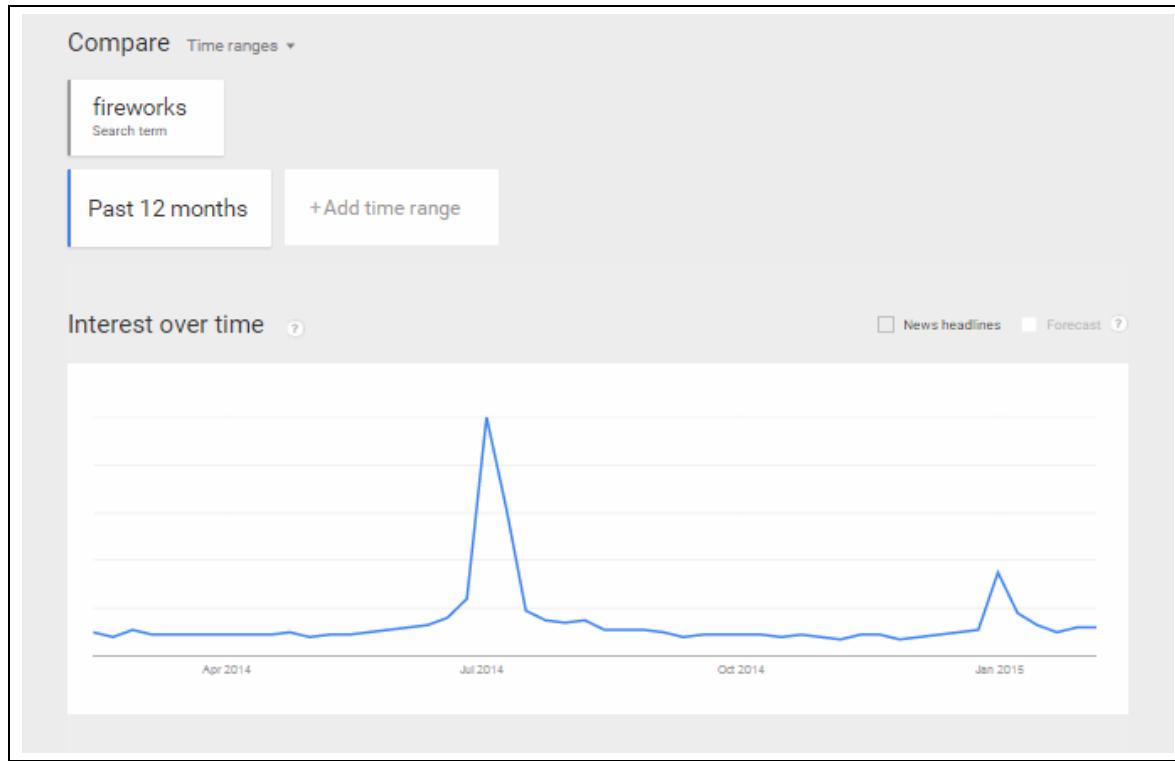


Figure 5-41. Google Trends; another seasonal example

In each case, searches started increasing about two to three months before the holiday, so it is important to acknowledge that and start crafting your content and targeting those keywords in ample time for them to be indexed before the searches start gaining traction.

KeywordDiscovery (<http://www.keyworddiscovery.com>) also graphs search trends, so if you have an account, you can analyze these graphs to craft a holiday campaign, as shown in Figure 5-42.



Figure 5-42. KeywordDiscovery seasonal data example

You can see when people begin to search for Halloween costumes and when the activity drops off.

Don't take your cue from when the stores start stocking Halloween candy—do the research and find out what last year's trends were so that you're prepared this year. If you prepare early enough, you'll be ready, while your competitors are scrambling with last-minute link-building campaigns three weeks before the holiday.

Also, don't remove your Halloween (or other seasonal) page as soon as the time frame has passed. Once you have fought hard to get rankings for your seasonal trophy term, you want to make sure you get the benefit for that hard work next year. Too many sites delete or archive these seasonal pages after the season is over, and then they have to start over again next year. A better strategy is to leave the page in place until a new version is created, reuse the same URL, and archive the old content to a different URL. Leaving the page in place will give you a jumpstart when it is time to begin ramping up next year.

Conclusion

Keyword research is a complex and time-consuming component of search engine optimization, but the rewards are high; once you learn where the keyword search volume is, you can begin to think about how that affects the information architecture and the navigation structure of your site. These are two critical elements that we will explore in greater detail in Chapter 6.

6

Developing an SEO-Friendly Website

In this chapter, we will examine the major elements of how to assess the search engine friendliness of your website. A search engine-friendly website, at the most basic level, is one that allows for search engine access to site content – and having your site content accessible to search engines is the first step towards creating prominent visibility in search results. Once your site’s content is accessed by a search engine, it can then be considered for relevant positioning within search results pages.

As we discussed in the introduction to Chapter 2, search engine crawlers are basically software programs – and, like all software programs, they come with certain strengths and weaknesses. Publishers must adapt their websites to make the job of these software programs easier—in essence, leverage their strengths and make their weaknesses irrelevant. If you can do this, you will have taken a major step forward toward success with SEO.

Developing an SEO-friendly site architecture requires a significant amount of thought, planning, and communication due to the large number of factors that influence the ways a search engine sees your site and the large number of ways in which a website can be put together, as there are hundreds (if not thousands) of tools that web developers can use to build a website - many of which were not initially designed with SEO, or search engine crawlers, in mind.

Making Your Site Accessible to Search Engines

The first step in the SEO design process is to ensure that your site can be found and crawled by the search engines. This is not as simple as it sounds, as there are many popular web design and implementation constructs that the crawlers may not understand.

Indexable Content

To rank well in the search engines, your site’s content—that is, the material available to visitors of your site—should be in HTML text form. Images and Flash files for example, while crawled by the search engines, are content types that are more difficult for search engines to analyze and therefore they are not ideal formats for communicating to search engines the topical relevance of your pages.

Images are a file type that the search engines have challenges with “identifying” from a relevance perspective, as there are minimum text-input fields for image files in GIF, JPEG, or PNG format (namely the filename, title, and alt attribute). While we do strongly recommend accurate labeling of images in these fields, images alone are usually not enough to earn a website page top rankings for relevant queries. While image identification technology continues to advance, processing power limitations will likely keep the search engines from broadly applying this type of analysis to web search in the near future.

Google offers the ability for users to perform a search using an image as the search query as opposed to text (though users can input text to augment the query). By uploading an image, dragging and dropping an image from the desktop, entering an image url, or by right-clicking on an image within a browser (Firefox and Chrome with installed extensions), users can often find other locations of that image on the web for reference and research, as well as images that “appear” similar in tone and composition. While this does not immediately change the landscape of SEO for images, it does give us an indication as to how Google is potentially augmenting their current relevance indicators for image content.

With Flash, while specific .swf files (the most common file extension for Flash) can be crawled and indexed – and often found for specific words or phrases included in their file name when searching only for .swf files – it is the rare occasion that a generic query returns a Flash file or a website generated entirely in Flash as a highly relevant result, due to the lack of “readable” content en masse. This is not to say that websites developed using Flash are inherently irrelevant, or that it is impossible to successfully optimize a website that uses Flash – however, in our experience the preference must still be given to HTML-based files.

Spiderable Link Structures

As we outlined in Chapter 2, search engines use links on web pages to help them discover other web pages and websites. For this reason, we strongly recommend taking the time to build an internal linking structure that spiders can crawl easily. Many sites make the critical mistake of hiding or obfuscating their navigation in ways that limit spider accessibility, thus impacting their ability to get pages listed in the search engines’ indexes. Consider the illustration in Figure 6-1 that shows how this problem can occur.

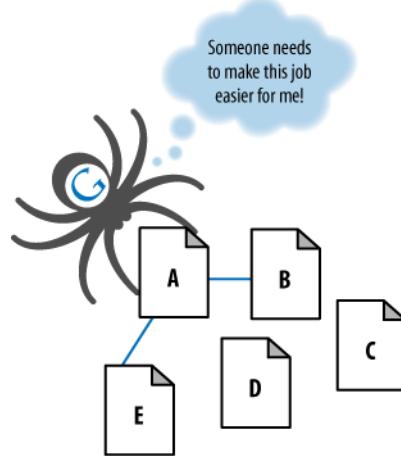


Figure 6-1. Providing search engines with crawlable link structures

In Figure 6-1, Google’s spider has reached Page A and sees links to pages B and E. However, even though pages C and D might be important pages on the site, the spider has no way to reach them (or even to know they exist) because no direct, crawl-able links point to those pages. As far as Google is concerned, they might as well not exist—great content, good keyword targeting, and smart marketing won’t make any difference at all if the spiders can’t reach those pages in the first place.

To refresh your memory on the discussion in Chapter 2, here are some common reasons why pages may not be reachable:

Links in submission-required forms

Search spiders will rarely, if ever, attempt to “submit” forms, and thus, any content or links that are accessible only via a form are invisible to the engines. This even applies to simple forms such as user logins, search boxes, or some types of pull-down lists.

Links in hard-to-parse JavaScript

If you use JavaScript for links, you may find that search engines either do not crawl or give very little weight to the links embedded within them. In June 2014 Google announced enhanced crawling of JavaScript and CSS (<http://googlewebmastercentral.blogspot.com/2014/05/understanding-web-pages-better.html>). Google can now render some JavaScript and follow some JavaScript links. Due to this change Google recommends refraining from blocking them from crawling your JavaScript and CSS files. For a preview of how your site might render according to Google go to Google Webmaster Tools, Crawl, Fetch as Google, input the URL you would like to preview, and select “Fetch and Render”.

Links in Java or other plug-ins

Traditionally links embedded inside Java and plug-ins have been invisible to the engines.

Links in Flash

In theory, the search engines can detect links within Flash, but don't rely too heavily on this.

Links in PowerPoint and PDF files

Search engines sometimes report links seen in PowerPoint files or PDFs. These links are believed to be counted the same as links embedded in HTML documents.

Links pointing to pages blocked by the meta `Robots` tag, `rel="NoFollow"`, or `robots.txt`

The `robots.txt` file provides a very simple means for preventing web spiders from crawling pages on your site. Use of the `NoFollow` attribute on a link, or placement of the meta `Robots` `nofollow` tag with the `content="NoFollow"` attribute on the page containing the link, is an instruction to the search engine to not pass link authority via the link (a concept we will discuss further in “Content Delivery and Search Spider Control”). The effectiveness of the `NoFollow` attribute on links has greatly diminished to the point of irrelevance as a result of over-manipulation by aggressive SEO practitioners. For more on this, please reference the following blog post from Google's Matt Cutts: <http://www.mattcutts.com/blog/pagerank-sculpting/>.

Links on pages with many hundreds or thousands of links

Historically, Google suggested a guideline of 100 links per page before it may stop spidering additional links from that page. This “limit” is not really that strict, and it is a more of a matter of your strategy for passing PageRank. If a page has 200 links on it, then each link does not get very much PageRank as a result. Managing how you pass PageRank by limiting the number of links is usually a good idea. There are tools that can run reports on the number of outgoing links you have per page, including [Screaming Frog](http://www.screamingfrog.co.uk/seo-spider/) (<http://www.screamingfrog.co.uk/seo-spider/>).

Links in frames or iframes

Technically, links in both frames and iframes can be crawled, but both present structural issues for the engines in terms of organization and following. Unless you're an advanced user with a good technical understanding of how search engines index and follow links in frames, it is best to stay away from them as a place to offer links for crawling purposes. We will discuss frames and iframes in more detail in “Creating an Optimal Information Architecture”.

XML Sitemaps

Google, Yahoo!, and Bing (formerly MSN Search, and then Live Search) all support a protocol known as XML Sitemaps. Google first announced it in 2005, and then Yahoo! and MSN Search agreed to support the protocol in 2006. Using the Sitemaps protocol you can supply the search engines with a list of all the URLs you would like them to crawl and index.

Adding a URL to a Sitemap file does not guarantee that a URL will be crawled or indexed. However, it can result in pages that are not otherwise discovered or indexed by the search engine getting crawled and indexed.

This program is a complement to, not a replacement for, the search engines' normal, link-based crawl. The benefits of Sitemaps include the following:

- For the pages the search engines already know about through their regular spidering, they use the metadata you supply, such as the last date the content was modified (*lastmod date*) and the frequency at which the page is changed (*changefreq*), to improve how they crawl your site.
- For the pages they don't know about, they use the additional URLs you supply to increase their crawl coverage.
- For URLs that may have duplicates, the engines can use the XML Sitemaps data to help choose a canonical version.
- Verification/registration of XML Sitemaps may indicate positive trust/authority signals.
- The crawling/inclusion benefits of Sitemaps may have second-order positive effects, such as improved rankings or greater internal link popularity.
- Having a site map registered with Google Webmaster Tools can give you an extra analytical insight into whether your site is suffering from indexation, crawling, or duplicate content issues.

Matt Cutts, the head of Google's webspam team, has explained XML Sitemaps in the following way:

Imagine if you have pages A, B, and C on your site. We find pages A and B through our normal web crawl of your links. Then you build a Sitemap and list the pages B and C. Now there's a chance (but not a promise) that we'll crawl page C. We won't drop page A just because you didn't list it in your Sitemap. And just because you listed a page that we didn't know about doesn't guarantee that we'll crawl it. But if for some reason we didn't see any links to C, or maybe we knew about page C but the URL was rejected for having too many parameters or some other reason, now there's a chance that we'll crawl that page C.

Sitemaps use a simple XML format that you can learn about at <http://www.sitemaps.org/>. XML Sitemaps are a useful and in some cases essential tool for your website. In particular, if you have reason to believe that the site is not fully indexed, an XML Sitemap can help you increase the number of indexed pages. As sites grow in size, the value of XML Sitemap files tends to increase dramatically, as additional traffic flows to the newly included URLs.

Layout of an XML Sitemap

The first step in the process of creating an XML Sitemap is to create an *.xml* Sitemap file in a suitable format. Since creating an XML Sitemap requires a certain level of technical know-how, it would be wise to involve your development team in the XML Sitemap generator process from the beginning. Figure 6-2 shows an example of some code from a Sitemap.

```

<?xml version="1.0" encoding="UTF-8"?>

<urlset xmlns="http://www.sitemaps.org/schemas/sitemap/0.9">

  <url>

    <loc>http://www.example.com/</loc>

    <lastmod>2005-01-01</lastmod>

    <changefreq>monthly</changefreq>

    <priority>0.8</priority>

  </url>

</urlset>

```

Figure 6-2. Sample XML Sitemap from Google.com

To create your XML Sitemap, you can use the following:

An XML Sitemap generator

This is a simple script that you can configure to automatically create Sitemaps, and sometimes submit them as well. Sitemap generators can create these Sitemaps from a URL list, access logs, or a directory path hosting static files corresponding to URLs. Here are some examples of XML Sitemap generators:

- SourceForge.net's Google-sitemap_gen (<http://sourceforge.net/projects/google-sitemapgen/files/sitemapgen/>)
- XML-Sitemaps.com Sitemap Generator (<http://www.xml-sitemaps.com/>)
- Sitemaps Pal (<http://www.sitemapspal.com/>)
- GSite Crawler (<http://gsitecrawler.com/>)

Simple text

You can provide Google with a simple text file that contains one URL per line. However, Google recommends that once you have a text Sitemap file for your site, you use the Sitemap Generator to create a Sitemap from this text file using the Sitemaps protocol.

Syndication feed

Google accepts Really Simple Syndication (RSS) 2.0 and Atom 1.0 feeds. Note that the feed may provide information on recent URLs only.

What to include in a Sitemap file

When you create a Sitemap file you need to take care in situations where your site has multiple URLs that refer to one piece of content – include *only* the preferred (canonical) version of the URL as the search engines may assume that the URL specified in a Sitemap file is the preferred form of the URL for the content. You can use the Sitemap file as one way to suggest to the search engines which URL points to the preferred version of a given page.

In addition, be careful about what not to include. For example, do not include multiple URLs that point to identical content, and leave out pages that are simply pagination pages -or alternate sort orders for the same content and/or any low-value pages on your site. Last but not least, make sure that none of the URLs listed in the Sitemap file include any tracking parameters.

Mobile Sitemaps

Mobile Sitemaps should be used for content targeted at mobile devices. Mobile information is kept in a separate Sitemap file that should not contain any information on non-mobile URLs. Google supports non-mobile markup, XHTML mobile profile, WML (WAP 1.2) and cHTML. Details on the mobile Sitemap format can be found here: <https://support.google.com/webmasters/answer/34648>

Video Sitemaps

Including information on your videos in your Sitemap file will increase their chances of being discovered by the search engines. Google indicates that they support the following video formats: .mpg, .mpeg, .mp4, .m4v, .mov, .wmv, .ASF, .avi, .ra, .ram, .rm, .flv, .swf. You can see the specification of how video Sitemap entries are to be implemented here: <https://support.google.com/webmasters/answer/80472>

Image Sitemaps

You can increase visibility for your images by listing them in your Sitemap file. For each URL you list in your Sitemap file, you can also list the images that appear on those pages. You can list up to 1000 images per page. Specialized image tags are associated with the URL. The details of the format of these tags can be seen on this page: <https://support.google.com/webmasters/answer/178636>.

Listing images in the Sitemap does increase the chances of those images being indexed. If you list some images and not others, it may be interpreted as a signal that the images not listed are less important.

Where to upload your Sitemap file

When your Sitemap file is complete, upload the file to your site in the highest-level directory you want search engines to crawl (generally, the root directory) – such as `www.yoursite.com/sitemap.xml`. You can include more than one subdomain in your Sitemap provided that you verify the Sitemap for each subdomain in Google Webmaster Tools, though it's frequently easier to understand what's happening with indexation if each subdomain has its own site map and its own profile in Google Webmaster Tools.

Managing and updating XML Sitemaps

Once your XML Sitemap has been accepted and your site has been crawled, monitor the results and update your Sitemap if there are issues. With Google, you can return to your Google Webmaster Tools account to view the statistics and diagnostics related to your XML Sitemaps. Just click the site you want to monitor. You'll also find some FAQs from Google on common issues such as slow crawling and low indexation.

Update your XML Sitemap with Google and Bing when you add URLs to your site. You'll also want to keep your Sitemap file up-to-date when you add a large volume of pages or a group of pages that are strategic.

There is no need to update the XML Sitemap when simply updating content on existing URLs. It is not strictly necessary to update when pages are deleted as the search engines will simply not be able to crawl them, but you want to update them before you have too many broken pages in your feed. Do update your Sitemap file whenever you add any new content, and remove any deleted pages at that time. Google and Bing will periodically re-download the Sitemap, so you don't need to resubmit your Sitemap to Google or Bing unless your Sitemap location has changed.

Facilitate Google and Bing auto-discovering your XML Sitemap locations using the `Sitemap:` directive in your site's `robots.txt` file.

If you have a site where you are adding or deleting large numbers of new pages on a regular basis you may want to use a utility, or have your developers build the ability, for your XML site map to regenerate with all of your current URLs on a regular basis. Many sites regenerate their XML site map daily via automated scripts.

Google and the other major search engines discover and index websites by crawling links. Google XML Sitemaps are a way to feed the URLs that you want crawled on your site to Google for more complete crawling and indexation, which results in improved long tail search-ability. By creating and updating this `.xml` file, you are helping to ensure that Google recognizes your entire site, and this recognition will help people find your site. It also helps all of the search engines understand which version of your URLs (if you have more than one URL pointing to the same content) is the canonical version.

Creating an Optimal Information Architecture (IA)

Making your site friendly to search engine crawlers also requires that you put some thought into your site information architecture. A well-designed site architecture can bring many benefits for both users and search engines.

The Importance of a Logical, Category-Based Flow

The search engines face myriad technical challenges in understanding your site, as crawlers are not able to perceive web pages in the way that humans do, creating significant limitations for both accessibility and indexing. A logical and properly constructed website architecture can help overcome these issues and bring great benefits in search traffic and usability.

At the core of website information architecture are two critical principles: usability, or making a site easy to use; and information architecture, or crafting a logical, hierarchical structure for content.

One of the very early information architecture proponents, Richard Saul Wurman, developed the following definition for *information architect*:

information architect. 1) the individual who organizes the patterns inherent in data, making the complex clear. 2) a person who creates the structure or map of information which allows others to find their personal paths to knowledge. 3) the emerging 21st century professional occupation addressing the needs of the age focused upon clarity, human understanding, and the science of the organization of information.

Usability and search friendliness

Search engines are trying to reproduce the human process of sorting relevant web pages by quality. If a real human were to do this job, usability and user experience would surely play a large role in determining the rankings. Given that search engines are machines and they don't have the ability to segregate by this metric quite so easily, they are forced to employ a variety of alternative, secondary metrics to assist in the process. The most well-known and well publicized among these is a measurement of the inbound links to a web site, (see Figure 6-3), and a well-organized site is more likely to receive links.

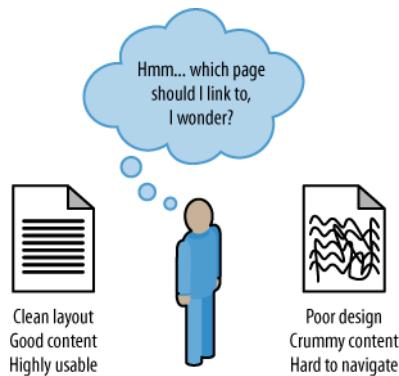


Figure 6-3. Making your site attractive to link to

Since Google launched in the late 1990s, search engines have strived to analyze every facet of the link structure on the Web and have extraordinary abilities to infer trust, quality, reliability, and authority via links. If you push back the curtain and examine why links between websites exist and how they come into place, you can see that a human being (or several humans, if the organization suffers from bureaucracy) is almost always responsible for the creation of links.

The engines hypothesize that high-quality links will point to high-quality content, and that great content and positive user experiences will be rewarded with more links than poor user experiences. In practice, the theory holds up well. Modern search engines have done a very good job of placing good-quality, usable sites in top positions for queries.

An analogy

Look at how a standard filing cabinet is organized. You have the individual cabinet, drawers in the cabinet, folders within the drawers, files within the folders, and documents within the files (see Figure 6-4).

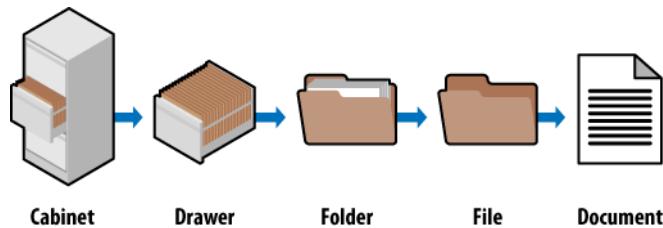


Figure 6-4. Similarities between filing cabinets and web pages

There is only one copy of any individual document, and it is located in a particular spot. There is a very clear navigation path to get to it.

If you want to find the January 2015 invoice for a client (Amalgamated Glove & Spat), you would go to the cabinet, open the drawer marked Client Accounts, find the Amalgamated Glove & Spat folder, look for the Invoices file, and then flip through the documents until you come to the January 2015 invoice (again, there is only one copy of this; you won't find it anywhere else).

Figure 6-5 shows what it looks like when you apply this logic to the popular website, Craigslist.org (<http://craigslist.org>).



Figure 6-5. Filing cabinet analogy applied to Craigslist.org (<http://craigslist.org>)

If you're seeking an apartment in Los Angeles, you'd navigate to LosAngeles.Craigslist.org (<http://losangeles.craigslist.org>), choose Apts/Housing, narrow that down to two bedrooms, and pick the two-bedroom loft from the list of available postings. Craigslist's simple, logical information architecture has made it easy to reach the desired post in four clicks, without having to think too hard at any step about where to go. This principle applies perfectly to the process of SEO, where good information architecture dictates:

- As few clicks as possible to any given page
- One hundred or fewer links per page (so as not to overwhelm either crawlers or visitors)
- A logical, semantic flow of links from home page to categories to detail pages

Here is a brief look at how this basic filing cabinet approach can work for some more complex information architecture issues.

Subdomains

You should think of subdomains as completely separate filing cabinets within one big room. They may share similar architecture, but they shouldn't share the same content; and more importantly, if someone points you to one cabinet to find something, he is indicating that that cabinet is the authority, not the other cabinets in the room. Why is this important? It will help you remember that links (i.e., votes or references) to subdomains may not pass all, or any, of their authority to other subdomains within the room (e.g., “*.craigslist.org,” wherein “*” is a variable subdomain name).

Those cabinets, their contents, and their authority are isolated from each other and may not be considered to be in concert with each other. This is why, in most cases, it is best to have one large, well-organized filing cabinet instead of several that may prevent users and bots from finding what they want.

Redirects

If you have an organized administrative assistant, he probably uses 301 redirects (these are discussed more later on in the "Redirects" section of this chapter) inside his literal, metal filing cabinet. If he finds himself looking for something in the wrong place, he might place a sticky note in there reminding him of the correct location the next time he needs to look for that item. Anytime you looked for something in those cabinets, you could always find it because if you navigated improperly, you would inevitably find a note pointing you in the right direction.

Redirect irrelevant, outdated, or misplaced content to the proper spot in your filing cabinet and both your users and the engines will know what qualities and keywords you think it should be associated with.

URLs

It would be tremendously difficult to find something in a filing cabinet if every time you went to look for it, it had a different name, or if that name resembled “jklhj25br3g452ikbr52k” – a not-so-uncommon type of character string found in dynamic website URLs. Static, keyword-targeted URLs are much better for users and for bots. They can always be found in the same place, and they give semantic clues as to the nature of the content.

These specifics aside, thinking of your site information architecture in terms of a filing cabinet is a good way to make sense of best practices. It'll help keep you focused on a simple, easily navigated, easily crawled, well-organized structure. It is also a great way to explain an often complicated set of concepts to clients and co-workers.

Since search engines rely on links to crawl the Web and organize its content, the architecture of your site is critical to optimization. Many websites grow organically and, like poorly planned filing systems, become complex, illogical structures that force people (and spiders) looking for something to struggle to find what they want.

Site Architecture Design Principles

In conducting website planning, remember that nearly every user will initially be confused about where to go, what to do, and how to find what she wants. An architecture that recognizes this difficulty and leverages familiar standards of usability with an intuitive link structure will have the best chance of making a visit to the site a positive experience. A well-organized site architecture helps solve these problems, and provides semantic and usability benefits to both users and search engines.

In Figure 6-6, a recipes website can use intelligent architecture to fulfill visitors' expectations about content and create a positive browsing experience. This structure not only helps humans navigate a site more easily, but also helps the search engines to see that your content fits into logical concept groups. You can use this approach to help you rank for applications of your product in addition to attributes of your product.

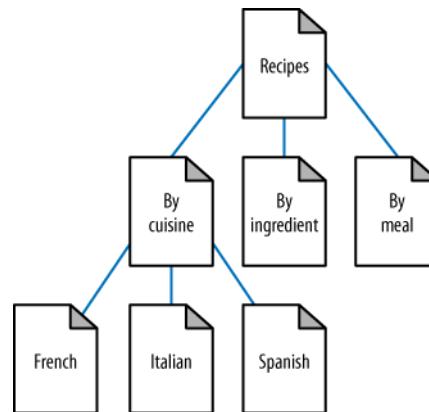


Figure 6-6. Structured site architecture

Although site architecture accounts for a small part of the algorithms, the engines do make use of relationships between subjects and give value to content that has been organized in a sensible fashion. For example, if in Figure 6-6 you were to randomly jumble the subpages into incorrect categories, your rankings could suffer. Search engines, through their massive experience with crawling the Web, recognize patterns in subject architecture and reward sites that embrace an intuitive content flow.

Designing site architecture

Although site architecture—the creation of structure and flow in a website’s topical hierarchy—is typically the territory of information architects (or is created without assistance from a company’s internal content team), its impact on search engine rankings - particularly in the long run - is substantial. It is, therefore, a wise endeavor to follow basic guidelines of search friendliness.

The process itself should not be overly arduous, if you follow this simple protocol:

1. List all of the requisite content pages (blog posts, articles, product detail pages, etc.).
2. Create top-level navigation that can comfortably hold all of the unique types of detailed content for the site.
3. Reverse the traditional top-down process by starting with the detailed content and working your way up to an organizational structure capable of holding each page.
4. Once you understand the bottom, fill in the middle. Build out a structure for sub-navigation to sensibly connect top-level pages with detailed content. In small sites, there may be no need for this level, whereas in larger sites, two or even three levels of sub-navigation may be required.
5. Include secondary pages such as copyright, contact information, and other non-essentials.
6. Build a visual hierarchy that shows (to at least the last level of sub-navigation) each page on the site.

Figure 6-7 shows an example of a well-structured site architecture.

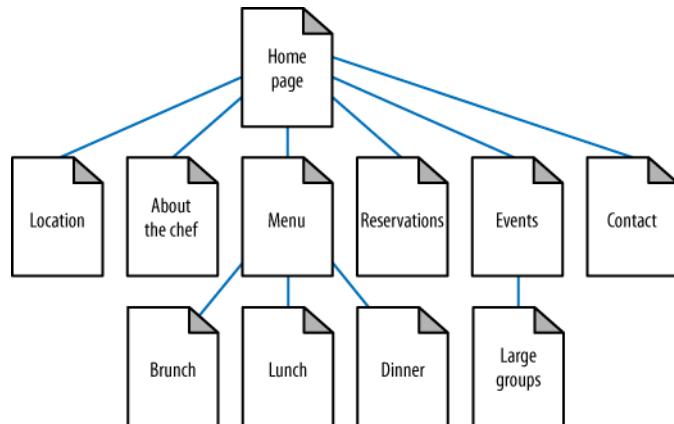


Figure 6-7. Second example of structured site architecture

Category structuring

As search engines crawl the Web, they collect an incredible amount of data (millions of gigabytes) on the structure of language, subject matter, and relationships between content. Though not technically an attempt at artificial intelligence, the engines have built a repository capable of making sophisticated determinations based on common patterns. As shown in Figure 6-8, search engine spiders can learn semantic relationships as they crawl thousands of pages that cover a related topic (in this case, dogs).

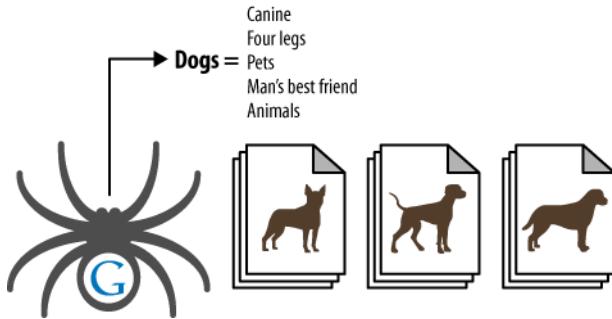


Figure 6-8. Spiders learning semantic relationships

Although content need not always be structured along the most predictable patterns, particularly when a different method of sorting can provide value or interest to a visitor, organizing subjects logically assists both humans (who will find your site easier to use) and engines (which will award you with greater rankings based on increased subject relevance).

Topical relevance

Naturally, this pattern of relevance-based scoring extends from single relationships between documents to the entire category structure of a website. Site creators can take advantage of this best by building hierarchies that flow from broad, encompassing subject matter down to more detailed, specific content. Obviously, in any categorization system, there is a natural level of subjectivity; think first of your visitors, and use these guidelines to ensure that your creativity doesn't overwhelm the project.

Taxonomy and ontology

In designing a website, you should also consider the taxonomy and ontology of the website. Taxonomy is essentially a two-dimensional hierarchical model of the architecture of the site. You can think of ontology as mapping the way the human mind thinks about a topic area. It can be much more complex than taxonomy, because a larger number of relationship types can be involved.

One effective technique for coming up with an ontology is called *card sorting*. This is a user-testing technique whereby users are asked to group items together so that you can organize your site as intuitively as possible. Card sorting can help identify not only the most logical paths through your site, but also ambiguous or cryptic terminology that should be reworded.

With card sorting, you write all the major concepts onto a set of cards that are large enough for participants to read, manipulate, and organize. Your test group assembles the cards in the order they believe provides the most logical flow, as well as into groups that seem to fit together.

By itself, building an ontology is not part of SEO, but when you do it properly it will impact your site architecture, and therefore it interacts with SEO. Coming up with the right site architecture should involve both disciplines.

Flat Versus Deep Architecture

One very strict rule for search friendliness is the creation of flat site architecture. Flat sites require a minimal number of clicks to access any given page, whereas deep sites create long paths of links required

to access detailed content. For nearly every site with fewer than 10,000 pages, all content should be accessible through a maximum of four clicks from the home page and/or sitemap page. That said, flatness should not be forced if it does not make sense for other reasons. At 100 links per page, even sites with millions of pages can have every page accessible in five to six clicks if proper link and navigation structures are employed. If a site is not built to be flat, it can take too many clicks to reach the desired content, as shown in Figure 6-9. In contrast, a flat site (see Figure 6-10) allows users and search engines to reach most content in just a few clicks.



Figure 6-9. Deep site architecture

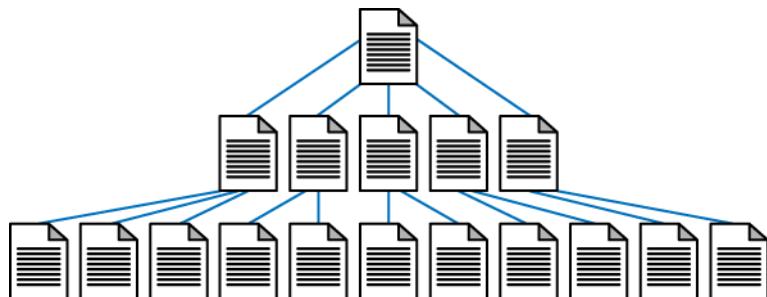


Figure 6-10. Flat site architecture

Flat sites aren't just easier for search engines to crawl; they are also simpler for users, as they limit the number of page visits the user requires to reach his destination. This reduces the abandonment rate and encourages repeat visits.

When creating flat sites, be careful to not overload pages with links either. Pages that have 200 links on them are not passing much PageRank to any of those pages. While flat site architectures are desirable, you should not force an architecture to be overly flat if it is not other logical to do so.

The number of links-per-page issue relates directly to another rule for site architects: avoid excessive pagination wherever possible. *Pagination*, the practice of creating a list of elements on pages separated solely by numbers (e.g., some e-commerce sites use pagination for product catalogs that have more products than they wish to show on a single page), is problematic for many reasons.

First, pagination provides virtually no new topical relevance as the pages are each largely about the same topic. Second, content that moves into different pagination can potentially create duplicate content problems or be seen as poor quality or “thin” content. Last, pagination can create spider traps and hundreds or thousands of extraneous, low-quality pages that can be detrimental to search visibility. There are ways to address the downsides of pagination as we will discuss in a moment. Figure 6-11 shows an example of pagination.

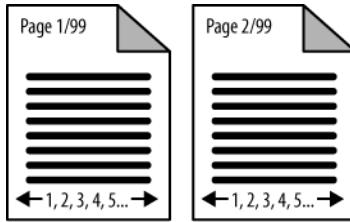


Figure 6-11. Pagination structures

So, make sure you implement flat structures and stay within sensible guidelines for the number of links per page, while retaining a contextually rich link structure. This is not always as easy as it sounds. Accomplishing this may require quite a bit of thought and planning to build a contextually rich structure on some sites. Consider a site with 10,000 different men's running shoes. Defining an optimal structure for that site could be a very large effort, but that effort will pay serious dividends in return.

Solutions to pagination problems vary based on the content of the web site. Here are two example scenarios and their solutions:

1. **Use `rel=next` and `rel=prev`.** Google supports link elements called `rel=next` and `rel=prev`. The benefit of using these link elements is that it lets Google know when it has encountered a sequence of paginated pages. Once Google recognizes these tags, links to any of the pages will be treated as links to the series of pages as a whole. In addition, Google will show in the index the most relevant page in the series (most of the time this will be the first page, but not always).

Bing announced support for `rel=next` and `rel=prev` in 2012 (<http://blogs.bing.com/webmaster/2012/04/13/implementing-markup-for-paginated-and-sequenced-content/>).

These tags can be used to inform Google about pagination structures, and they can be used whether or not you create a View-All page. The concept is simple. The following example outlines how to use the tags for content that is paginated into 12 pages:

- On page one of your paginated content implement a `rel=next` tag in the `<head>` section of the page, and point that to the second page of the content. The tag should look something like this:

```
<link rel="next" href="http://www.yoursite.com/products?prod=qwert&p=2" />
```

- On the last page of your paginated content implement a `rel=prev` tag to the second to last page of the content in the `<head>` section of the page. The tag should look something like this:

```
<link rel="prev" href="http://www.yoursite.com/products?prod=qwert&p=11" />
```

- On pages 2 through 11, implement `rel=next` and `rel=prev` tags in the `<head>` section of the page. The following example shows what it should look like on page 6 of the content:

```
<link rel="prev" href="http://www.yoursite.com/products?prod=qwert&p=5" />
```

```
<link rel="next" href="http://www.yoursite.com/products?prod=qwert&p=7" />
```

2. **Create a View-All Page and use Canonical tags.** You may have lengthy articles that you choose to break into multiple pages. However, this results in links to the pages whose anchor text is something like "1", "2", and so forth. The page titles of the various pages may not vary in any significant way so they tend to compete with each other for search traffic. Finally, if someone links to the article but does not link to the first page, the link authority from that link will largely be wasted.

One way to handle this problem is to retain the paginated version of the article, but also existing create a single page version of the article. This is referred to as a View-All page. Then use the canonical tag (which is discussed in more detail later in the "Content Delivery and Spider Control" section in this Chapter) to point from the paginated pages to the View-All page. This will concentrate all of the link authority and search engine attention on one single page. You should also include a link to the View-

All page from each of the individual paginated pages as well. However, if the View-All page is too slow in loading because of the page size this may not be the best option for you.

It should also be noted that if you implement a View-All page and do not implement any of these tags, Google will attempt to discover the page and show it instead of the paginated versions in its search results. However, the authors recommend that you make use of one of the above two solutions, as Google cannot guarantee that it will discover your View-All pages and it is best to provide it with as many clues as possible.

Search-Friendly Site Navigation

Website navigation is something that web designers have been putting considerable thought and effort into since websites came into existence. Even before search engines were significant, navigation played an important role in helping users find what they wanted. It plays an important role in helping search engines understand your site as well.

Basics of search engine friendliness

The search engine spiders need to be able to read and interpret your website's code to properly spider and index the content on your web pages. Do not confuse this with the rules of organizations such as the W3C, which issues guidelines on HTML construction. Although following the W3C guidelines can be a good idea, the great majority of sites do not follow these guidelines, so search engines generally overlook violations of these rules as long as their spiders can parse the code.

Unfortunately, there are also a number of ways that navigation and content can be rendered on web pages that function for humans, but are invisible or at least challenging for search engine spiders.

For example, there are numerous ways to incorporate content and navigation on the pages of a website. For the most part, all of these are designed for humans. Basic HTML text and HTML links such as those shown in Figure 6-12 work equally well for humans and search engine crawlers.

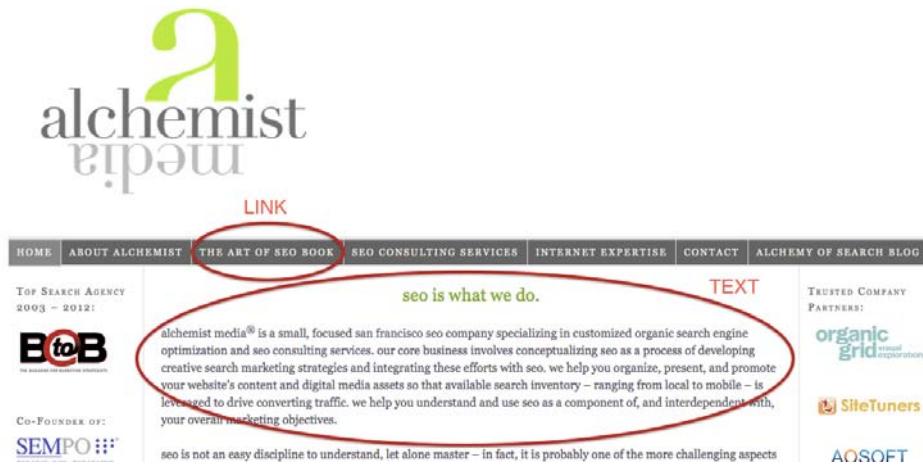


Figure 6-12. Example page with simple text and text link

The text and the link that are indicated on the page shown in Figure 6-12 are in simple HTML format.

Site elements that are problematic for spiders

However, many other types of content may appear on a web page and may work well for humans but not so well for search engines. Here are some of the most common ones.

Search and web forms

Many sites incorporate search functionality. These "site search" elements are specialized search engines that index and provide access to one site's content.

This is a popular method of helping users rapidly find their way around complex sites. For example, the Pew Internet website (<http://www.pewinternet.org>) provides Site Search in the top-right corner; this is a great tool for users, but search engines will be stymied by it. Search engines operate by crawling the Web's link structure—they don't in most circumstances submit forms or attempt random queries into search fields, and thus, any URLs or content solely accessible via a form will remain invisible to Google and Bing. In the case of Site Search tools, this is OK, as search engines do not want to index this type of content (they don't like to serve search results within their search results).

Forms are a popular way to provide interactivity, and one of the simplest applications is the "contact us" form many websites have.

Unfortunately, crawlers will not fill out or submit forms such as these; thus, any content restricted to those who employ them is inaccessible to the engines. In the case of a "contact us" form, this is likely to have little impact, but other types of forms can lead to bigger problems.

Websites that have content behind paywall and/or login barriers will either need to provide text links to the content behind the barrier (which defeats the purpose of the login) or implement First Click Free (discussed in "Content Delivery and Search Spider Control").

Java, images, audio, and video

Flash files, Java embeds, audio, and video (in any format) present content that is largely un-crawlable by the major engines. With some notable exceptions that we will discuss later, search engines can read text only when it is presented in HTML format. Embedding important keywords or entire paragraphs in an image or a Java console renders them invisible to the spiders. Likewise, words spoken in an audio file or video cannot easily be read by the search engines. However, Google has begun to leverage tools such as Google Voice Search in order to "crawl" audio content and extract meaning (this was first confirmed in the book "In the Plex" by Steven Levy and published by Simon and Schuster). Baidu already has an mp3 search function and Shazam and Jsaikoz applications show the ability to identify song hashes today as well.

`Alt` attributes, originally created as metadata for markup and an accessibility tag for vision-impaired users, is a good way to present at least some text content to the engines when displaying images or embedded, nontext content. Note that the `alt` attribute is not a strong signal, and using the `alt` attribute on an image link is no substitute for implementing a simple text link with appropriately descriptive anchor text. A good alternative is to employ captions and text descriptions in the HTML content wherever possible.

In the past few years, a number of companies offering transcription services have cropped up, providing automated text creation for the words spoken in audio or video. Providing these transcripts on rich media pages makes your content accessible to the search engines and findable by keyword-searching visitors. You can also use software such as Dragon Naturally Speaking and dictate your "transcript" to your computer.

AJAX and JavaScript

JavaScript enables many dynamic functions inside a website, most of which interfere very minimally with the operations of a search engine spider. The exception comes when a page must use a JavaScript call to reach another page, or to pull content that the spiders can't see in the HTML. In some instances this content is not visible to search engine spiders. However, will attempt to execute Javascript to access this type of content (<http://googlewebmastercentral.blogspot.com/2014/10/updating-our-technical-webmaster.html>). One example of this is Facebook Comments. Facebook Comments is a system offered by Facebook that allows publishers to collect comments from users on their site. Figure 6-13 shows an example of the Facebook Comments on a page on the Techcrunch site:



The screenshot shows a Facebook comment section on a Techcrunch article. At the top, there is a placeholder for a profile picture and a text input field labeled 'Add a comment...'. Below this is a checkbox for 'Also post on Facebook' and a button for 'Posting as Rob Woods (Change) Comment'. The first comment is by a user with a profile picture of a man in a suit, followed by the text: 'People who are overly critical of his accent and all of that need to take a HUGE chill pill and get over themselves.' The second comment is by a user with a profile picture of a person in a winter hat, followed by the text: 'As an ethnic Chinese, learning Mandarin and struggling, I'm extremely impressed. I always complain that I don't have enough time to practice and learn Chinese. I stand corrected.' The third comment is by a user with a profile picture of a man smiling, followed by the text: 'The fact that he learned Chinese just to speak to his wife's grandmother is pretty admirable.' The fourth comment is by a user with a profile picture of a man with glasses, followed by the text: 'Agree with you. I'm not running a multi-millionaire company as Facebook and for me it's very hard to deal with my chinese learning and master degree courses at the same time.' Below these comments, there is a link to 'View 4 more'.

Figure 6-13. Facebook Comments on Techcrunch (<http://techcrunch.com/2014/10/23/zuckerberg-speaks-chinese-internet-soils-itself/>)

If you examine the source code for this particular post you will not see any of the text strings for these comments in the HTML of the page. This is because the comments are actually stored on Facebook and dynamically retrieved by the web server when the page is rendered.

This is an example of the type of content that was not historically indexed by the search engines. When you use a Javascript implementation like this it is not clear what Google or Bing will be able to do with it. Facebook Comments is a broadly used system and it makes sense for the search engines to learn how to attempt to read that content. However to date, Google has been able to include this type of content in its index. You can test this (and whether your own content is indexed) by doing a Google search on a unique string of words in Google surrounded by double quotes on either side to ensure Google searches only for those exact words in that exact order. For example searching Google for a string of words in one of the comments above "As an ethnic Chinese, learning Mandarin and struggling, I'm extremely impressed" does not return the URL <http://techcrunch.com/2014/10/23/zuckerberg-speaks-chinese-internet-soils-itself> as a result for this query.

While Google has recently indicated that they execute most Javascript there is still the possibility that they don't execute all Javascript. Case in point: if your intent is to create content that you want the search engines to see, it is still safest to implement something in a form that is directly visible in the HTML of the web page.

Asynchronous JavaScript and XML (AJAX) can present problems, most notably in the delivery of content that search engines may not be able to spider. Since AJAX uses database calls to retrieve data without

refreshing a page or changing URLs, the content contained behind these technologies may be completely hidden from the search engines (see Figure 6-14).

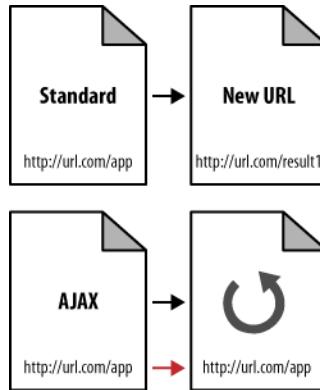


Figure 6-14. The problem with AJAX

If a traditional AJAX implementation is used, you may want to consider implementing an alternative spidering system for search engines to follow. AJAX applications are so user-friendly and appealing that for many publishers foregoing them is simply impractical. With these traditional implementations, building out a directory of links and pages that the engines can follow is a far better solution.

When you build these secondary structures of links and pages, make sure to provide users with access to them as well. Inside the AJAX application itself, give your visitors the option to “directly link to this page” and connect that URL with the URL you provide to search spiders through your link structures. AJAX apps not only suffer from content that can’t be crawled, but often don’t receive accurate links from users since the URL doesn’t change.

Some versions of AJAX use a # delimiter, which acts as a query string into the AJAX application. This does allow you to link directly to different pages within the application. The #, which is used for HTML bookmarking, and everything past it, is normally ignored by search engines.

This is largely because web browsers use only what’s after the # to jump to the anchor within the page, and that’s done locally within the browser. In other words, the browser doesn’t send the full URL, so the parameter information (i.e., any text after the #) is not passed back to the server.

Google outlined a method for making these AJAX pages search engine visible back in 2009: <http://googlewebmastercentral.blogspot.com/2009/10/proposal-for-making-ajax-crawlable.html>. This was later followed up with recommendations made on the Google Developers site: <https://developers.google.com/webmasters/ajax-crawling/docs/learn-more>. You can also find more information at <https://support.google.com/webmasters/answer/174992?hl=en>

The solution proposed by Google involves making some slight modifications to the way your AJAX URLs are formatted so that their crawler can recognize when an AJAX URL can be treated like a static page (one that will always return the same content), in which case Googlebot will read the page and treat it like any other static page for indexing and ranking purposes, allowing it the same opportunity to rank as a page coded in plain HTML.

Frames

Frames emerged in the mid-1990s as a popular way to make easy navigation systems. Unfortunately, both their usability (in 99% of cases) and their search friendliness (in 99.99% of cases) were exceptionally poor. Today, iframes and CSS can replace the need for frames, even when a site’s demands call for similar functionality.

For search engines, the biggest problem with frames and iframes is that they often hold the content from two or more URLs on a single page. For users, search engines, which direct searchers to only a single URL, may get confused by frames and direct visitors to single pages (orphan pages) inside a site intended to show multiple URLs at once. Indeed, the search engines consider the content within an iframe as residing on a

separate page from the one the iframe is being used on. Thus, page with nothing but iframed content will look virtually blank to the search engines.

Additionally, since search engines rely on links, and frame pages will often change content for users without changing the URL, external links often point to the wrong URL unintentionally. As a consequence, links to the page containing the frame or iframe may actually not point to the content the linker wanted to point to. Figure 6-15 shows an example page that illustrates how multiple pages are combined into a single URL with frames, which results in link distribution and spidering issues.

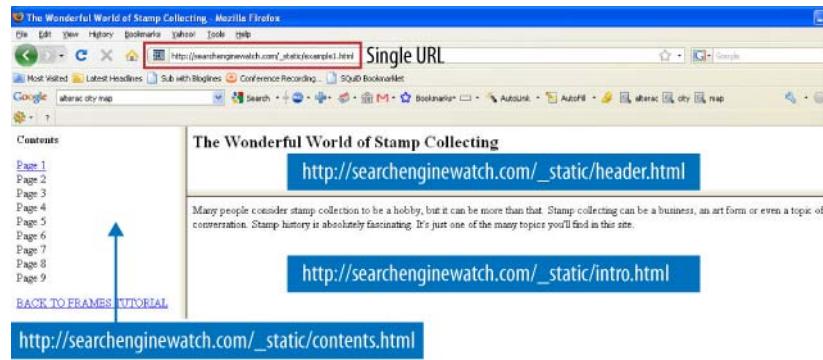


Figure 6-13. Sample page using frames

Search-engine-friendly navigation guidelines

Although search engine spiders have become more advanced over the years, the basic premise and goals remain the same: spiders find web pages by following links and record the content of the pages they find in the search engine's index (a giant repository of data about websites and pages).

In addition to avoiding the techniques we just discussed, there are some additional guidelines for developing search-engine-friendly navigation:

Implement a text-link-based navigational structure

If you choose to create navigation in Flash, JavaScript, or other technologies where the search engine's ability to parse them is uncertain, make sure to offer alternative text links in HTML for spiders to ensure that automated robots (and visitors who may not have the required browser plug-ins) can reach your pages.

Beware of “spider traps”

Even intelligently coded search engine spiders can get lost in infinite loops of links that pass between pages on a site. Intelligent architecture that avoids recursive looping 301 or 302 HTTP server codes (or other redirection protocols) should negate this issue, but sometimes online calendar links, infinite pagination that loops, or massive numbers of ways in which content is accessible or sorted (faceted navigation) can create tens of thousands of pages for search engine spiders when you intended to have only a few dozen true pages of content. You can read more about Google's viewpoint on this at <http://googlewebmastercentral.blogspot.com/2008/08/to-infinity-and-beyond-no.html>.

Watch out for session IDs and cookies

As we just discussed, if you limit the ability of a user to view pages or redirect based on a cookie setting or session ID, search engines may be unable to crawl your content. The bots do not have cookies enabled, nor can they deal with session IDs properly (each visit by the crawler gets a URL with a different session ID and the search engine sees these URLs with session IDs as different URLs). Although restricting form submissions is fine (as search spiders can't submit forms anyway), limiting content access via cookies and session IDs is a bad idea.

Server, hosting, and IP issues

Server issues rarely cause search engine ranking problems—but when they do, disastrous consequences can follow. The engines are acutely aware of common server problems, such as

downtime or overloading, and will give you the benefit of the doubt (though this will mean your content cannot be spidered during periods of server dysfunction). On the flip side, sites hosted on Content Delivery Networks (CDNs) may get crawled more heavily and CDNs offer significant performance enhancements to a web site.

The IP address of your host can be of concern in some instances. IPs once belonging to sites that have spammed the search engines may carry with them negative associations that can hinder spidering and ranking. While the engines aren't especially picky about shared hosting versus dedicated servers and dedicated IP addresses, or about server platforms, it should be noted that many hassles can be avoided by going these routes. At the very minimum, you should be cautious and find a host you trust, and inquire into the history and "cleanliness" of the IP address you may be assigned, as the search engines have become paranoid about the use of certain domains, hosting problems, IP addresses, and blocks of IPs. Experience tells them that many of these have strong correlations with spam, and thus, removing them from the index can have great benefits for users. As a site owner *not* engaging in these practices, it pays to investigate your web host prior to getting into trouble.

You can read more about server and hosting issues in "Identifying Current Server Statistics Software and Gaining Access" in Chapter 4.

Root Domains, Subdomains, and Microsites

Among the common questions in structuring a website (or restructuring one) are whether to host content on a new domain, when to use subfolders, and when to employ microsites.

As search engines scour the Web, they identify four kinds of web structures on which to place metrics:

Individual pages/URLs

These are the most basic elements of the Web—filenames, much like those that have been found on computers for decades, which indicate unique documents. Search engines assign query-independent scores—most famously, Google's PageRank—to URLs and judge them in their ranking algorithms. A typical URL might look something like <http://www.yourdomain.com/page>.

Subfolders

The folder structures that websites use can also inherit or be assigned metrics by search engines (though there's very little information to suggest that they are used one way or another). Luckily, they are an easy structure to understand. In the URL <http://www.yourdomain.com/blog/post17>, "/blog/" is the subfolder and "post17" is the name of the file in that subfolder. Engines may identify common features of documents in a given subfolder and assign metrics to these (such as how frequently the content changes, how important these documents are in general, or how unique the content is that exists in these subfolders).

Subdomains/fully qualified domains (FQDs)/third-level domains

In the URL <http://blog.yourdomain.com/page>, three kinds of domain levels are present. The top-level domain (also called the *TLD* or *domain extension*) is ".com", the second-level domain is "yourdomain", and the third-level domain is "blog". The third-level domain is sometimes referred to as a *subdomain*. Common web nomenclature does not typically apply the word *subdomain* when referring to *www*, although technically, this too is a subdomain. A fully qualified domain is the combination of the elements required to identify the location of the server where the content can be found (in this example, "<http://blog.yourdomain.com/>").

These structures can receive individual assignments of importance, trustworthiness, and value from the engines, independent of their second-level domains, particularly on hosted publishing platforms such as WordPress, Blogspot, and so on.

Complete root domains/host domain/pay-level domains (PLDs)/second-level domains

The domain name you need to register and pay for, and the one you point DNS settings toward, is the second-level domain (though it is commonly improperly called the "top-level" domain). In the URL

<http://www.yourdomain.com/page>, “yourdomain.com” is the second-level domain. Other naming conventions may refer to this as the “root” or “top-level” domain.

Figure 6-16 shows some examples

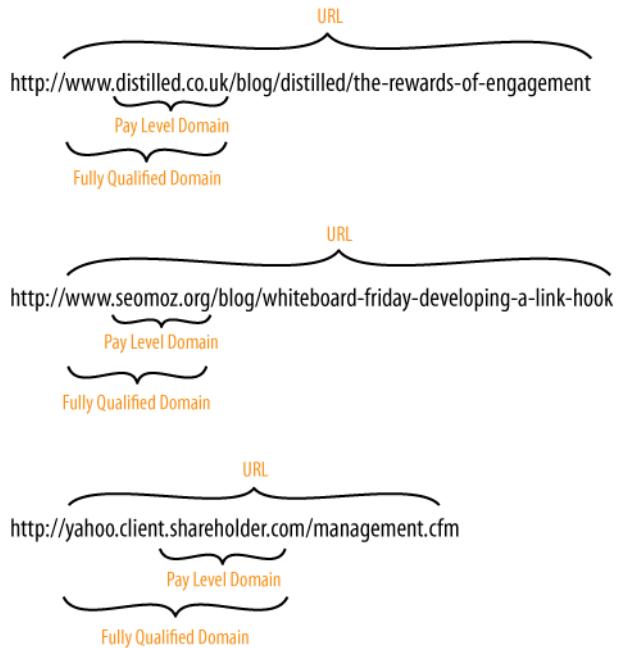


Figure 6-14. Breaking down some example URLs

When to Use a Subfolder

If a subfolder will work it is the best choice 99.9% of the time. Keeping content on a single root domain and single subdomain (e.g., <http://www.yourdomain.com>) gives the maximum SEO benefits, as engines will maintain all of the positive metrics the site earns around links, authority, and trust, and will apply these to every page on the site.

Subfolders have all the flexibility of subdomains (the content *can*, if necessary, be hosted on a unique server or completely unique IP address through post-firewall load balancing) and none of the drawbacks. Subfolder content will contribute directly to how search engines (and users, for that matter) view the domain as a whole. Subfolders can be registered with the major search engine tools and geo-targeted individually to specific countries and languages as well.

Although subdomains are a popular choice for hosting content, they are generally not recommended if SEO is a primary concern. Subdomains *may* inherit the ranking benefits and positive metrics of the root domain they are hosted underneath, but they do not always do so (and thus, content can underperform in these scenarios). Of course, there can be exceptions to this general guideline – subdomains are not inherently harmful, and there are some content publishing scenarios in which subdomains are more appropriate than subfolders; it is simply preferable for various SEO reasons, when and if at all possible, to use subfolders – as we will discuss below.

When to Use a Subdomain

If your marketing team decides to promote a URL that is completely unique in content or purpose and would like to use a catchy subdomain to do it, using a subdomain can be practical. [Maps.google.com](http://maps.google.com) (<http://maps.google.com/>) is an example of where the marketing considerations make a subdomain an acceptable choice. One good reason to use a subdomain is in a situation where using a subdomain can look more authoritative to users as a result of creating separation from the main domain.

Subdomains may also be a reasonable choice if keyword usage in the domain name is of critical importance. It appears that search engines do weight keyword usage in the URL somewhat, and have slightly higher benefits for exact matches in the subdomain (or third-level domain name) than subfolders. Note that exact matches in the domain and subdomain carry less weight than they once did. Google updated the weight they assigned to these factors in 2012 (<http://searchengineland.com/google-emd-update-research-and-thoughts-137340>).

Keep in mind that subdomains may inherit very little link equity from the main domain. If you wish to split your site in the subdomains and have all of them rank well, assume that you will have to support each with its own full-fledged SEO strategy.

When to Use a Separate Root Domain

If you have a single, primary site that has earned links, built content, and attracted brand attention and awareness, it is very rarely advisable to place any new content on a completely separate domain. There are rare occasions when this can make sense, and we'll walk through these, as well as explain how singular sites benefit from collecting all of their content in one root domain location.

Splitting similar or relevant content from your organization onto multiple domains can be likened to a store taking American Express Gold cards and rejecting American Express Corporate or American Express Blue—it is overly segmented and dangerous for the consumer mindset. If you can serve web content from a singular domain, that domain will earn branding among the minds of your visitors, references from them, links from other sites, and bookmarks from your regular customers. Switching to a new domain forces you to rebrand and to earn all of these positive metrics all over again.

Microsites

Although we generally recommend that you do not saddle yourself with the hassle of dealing with multiple sites and their SEO risks and disadvantages, it is important to understand the arguments, if only a few, in favor of doing so.

Making the case for microsites

Optimized properly, a microsite may have dozens or even hundreds of pages. If your site is likely to gain more traction and interest with webmasters and bloggers by being at an arm's length from your main site, it may be worth considering—for example, if you have a very commercial main site, and you want to create some great content, perhaps as articles, podcasts, and RSS feeds, that do not fit on the main site.

When should you consider a microsite?

When you own a specific keyword search query domain

For example, if you own “usedtoyotatrucks.com (<http://usedtoyotatrucks.com>)”, you might do very well to pull in search traffic for the specific term *used toyota trucks* with a microsite.

When you plan to sell the domains

It is very hard to sell a folder or even a subdomain, so this strategy is understandable if you're planning to churn the domains in the second-hand market.

As discussed earlier, if you're a major brand building a “secret” or buzz-worthy microsite

In this case, it can be useful to use a separate domain (however, you should 301 the pages of that domain back to your main site after the campaign is over so that the link authority continues to provide long-term benefit—just as the mindshare and branding do in the offline world).

You should never implement a microsite that acts as a doorway page to your main site, or that has substantially the same content as you have published on your main site. Consider this only if you are

willing to invest in putting rich original content on the microsite, and if you are willing to invest the time to promote the site as an independent site.

Such a site may gain more links by being separated from the main commercial site. A microsite may have the added benefit of bypassing some of the legal and PR department hurdles and internal political battles.

However, a microsite on a brand-new domain take some time to build enough domain level link authority to rank in the engines for weeks or months (for more about how Google treats new domains, see “Determining Searcher Intent and Delivering Relevant, Fresh Content” in Chapter 2). So, what to do if you want to launch a microsite? Start the clock running as soon as possible on your new domain by posting at least a few pages to the URL and then getting at least a few links to it—as far in advance of the official launch as possible. It may take a considerable amount of time before a micro-site is able to house enough high-quality content and to earn enough trusted and authoritative links to rank on its own. If the campaign the micro-site was created for his time limited in nature consider redirecting the pages from the micro-site to your main site well after the campaign concludes or at least ensure that the micro-site links back to the main site to allow some of the link authority the site earns to help the ranking of your main site.

Here are the reasons for not using a microsite:

Search algorithms favor large, authoritative domains

Take a piece of great content about a topic and toss it onto a small, mom and pop website—point some external links to it, optimize the page and the site for the target terms, and get it indexed. Now, take that exact same content and place it on Wikipedia or CNN.com (<http://cnn.com>), you’re virtually guaranteed that the content on the large, authoritative domain will outrank the content on the small niche site. The engines’ current algorithms favor sites that have built trust, authority, consistency, and history.

Multiple sites split the benefits of links

As suggested in Figure 6-17, a single good link pointing to a page on a domain positively influences the entire domain and every page on it. Because of this phenomenon, it is much more valuable to have any link you can possibly get pointing to the same domain to help boost the rank and value of the pages on it. Having content or keyword-targeted pages on other domains that don’t benefit from the links you earn to your primary domain only creates more work.

100 links to Domain A \neq 100 links to Domain B + 1 link to Domain A (from Domain B)

In Figure 6-18, you can see how earning lots of links to Page G on a separate domain is far less valuable than earning those same links to a page on the primary domain. Due to this phenomenon, even if you interlink all of the microsites or multiple domains that you build, it still won’t be close to the value you can get from those links if they were to point directly to the primary domain.

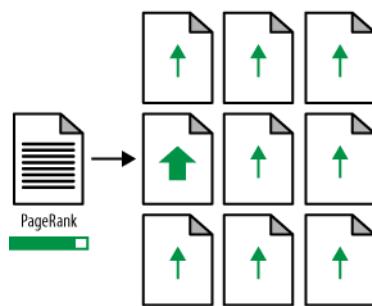
A large, authoritative domain can host a huge variety of content

Niche websites frequently limit the variety of their discourse and content matter, whereas broader sites can target a wider range of foci. This is valuable not just for targeting the long tail of search and increasing potential branding and reach, but also for viral content, where a broader focus is much less limiting than that of a niche focus.

Time and energy are better spent on a single property

If you’re going to pour your heart and soul into web development, design, usability, user experience, site architecture, SEO, public relations, branding, and so on, you want the biggest bang for your buck. Splitting your attention, time, and resources on multiple domains dilutes that value and doesn’t let the natural order of building on your past successes on a single domain assist with that process. As shown in Figure 6-18, every page on a site receives benefit from inbound links to a site. The page receiving the link gets the most benefit, but other pages also benefit.

**The rising tide (of a link)
lifts all boats (on the domain)**



Not only does the individual page gain rankings, the entire domain (and every page on it) receives some ranking benefit

Figure 6-15. How links can benefit your whole site

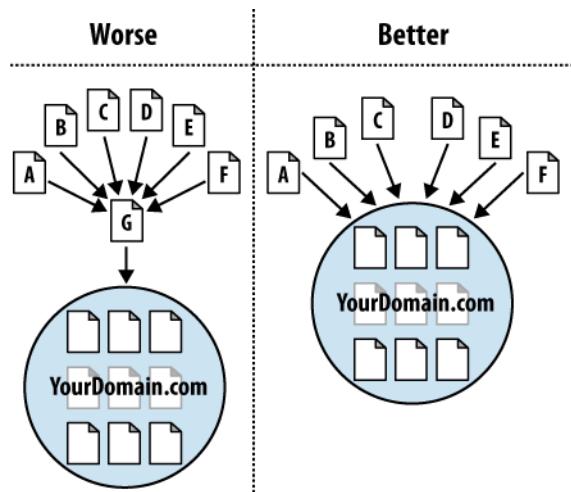


Figure 6-16. How direct links to your domain are better

When to Use a TLD Other Than .com

There are only a few situations in which you should consider using a TLD other than .com:

- When you own the .com and want to redirect to an .org, .tv, .biz, and so on, possibly for marketing/branding/geographic reasons. Do this only if you already own the .com and can redirect.
- When you can use a .gov, .mil, or .edu domain
- When you are serving only a single geographic region and are willing to permanently forego growth outside that region (e.g., .co.uk, .de, .it, etc.).
- When you are a nonprofit and want to distance your organization from the commercial world, .org may be for you.

New gTLDs

Many website owners have questions about the new gTLDs (generic top-level domains) which ICANN started assigning in the Fall of 2013. Instead of the traditional .com, .net, .org, .ca, etc. which most people are familiar with many new gTLDs are being released which can range from anything such as .CHRISTMAS to .AUTOS to .LAWYER, .EAT or .SYDNEY. A full list of these new TLDs can be found

at <http://newgtlds.icann.org/en/program-status/delegated-strings>. One of the major questions that arises is “will these help me rank organically on terms related to the TLD?” Currently the answer is “no”. There is no inherent SEO value in having a TLD which is related to your keywords. Having a .storage domain does not mean you have some edge over a .com for a storage related business. This is supported by statements by Google's John Mueller here <https://productforums.google.com/forum/#!topic/webmasters/ngHKBJjjJPo/discussion>. In that discussion John states that these TLDs are treated the same as other generic level TLDs in that they do not help your organic rankings. He also notes that even the new TLDs which sound as if they may be region specific give you know specific ranking benefit in those regions, though he allows that they reserve the right to change that in the future.

Despite the fact that they do not give you a ranking benefit currently you should still go and grab your domains for key variants of the new TLDs. Ones you may wish to consider our ones such as .SPAM. You may also wish to register those that relate directly to your business. It is unlikely that these TLD's will give a search benefit in the future. It's more likely that if your competition registers your name in conjunction with one of these new TLD's that it may cause confusion for your users as to which is the legitimate site. For example if you are located in New York City you should probably purchase your domain name with the .NYC TLD, and if you happen to own a pizza restaurant you may want to purchase .PIZZA, etc.

Optimization of Domain Names/URLs

Two of the most basic parts of any website are the domain name and the URLs for the pages of the website. This section will explore guidelines for optimizing these important elements.

Optimizing Domains

When a new site is being conceived or designed, one of the critical items to consider is the naming of the domain, whether it is for a new blog, a company launch, or even just a friend's website. Here are 12 tips that will be indispensable when helping you select a great domain name:

Brainstorm five top keywords

When you begin your domain name search, it helps to have five terms or phrases in mind that best describe the domain you're seeking. Once you have this list, you can start to pair them or add prefixes and suffixes to create good domain ideas. For example, if you're launching a mortgage-related domain, you might start with words such as *mortgage, finance, home equity, interest rate, and house payment*, and then play around until you can find a good match.

Make the domain unique

Having your website confused with a popular site that someone else already owns is a recipe for disaster. Thus, never choose a domain that is simply the plural, hyphenated, or misspelled version of an already established domain. For example, for years Flickr did not own Flicker.com (<http://flicker.com>), and they probably lost traffic because of that. They recognized the problem and bought the domain, and as a result <http://flicker.com> now redirects to <http://flickr.com>.

Choose only dot-com available domains

If you're not concerned with type-in traffic, branding, or name recognition, you don't need to worry about this one. However, if you're at all serious about building a successful website over the long term, you should be worried about all of these elements, and although directing traffic to a .net or .org (or any of the other new gTLDs) is fine, owning and 301'ing the .com, or the ccTLD for the country your web site servers (e.g. ".co.uk" for the United Kingdom), is critical. With the exception of the very tech-savvy, most people who use the Web still make the automatic assumption that .com is all that's out there, or that these domains are more trustworthy. Don't make the mistake of locking out or losing traffic from these folks.

Make it easy to type

If a domain name requires considerable attention to type correctly due to spelling, length, or the use of unmemorable words or sounds, you've lost a good portion of your branding and marketing value. Usability folks even tout the value of having the words include easy-to-type letters (which we interpret as avoiding *q*, *z*, *x*, *c*, and *p*).

Make it easy to remember

Remember that word-of-mouth marketing relies on the ease with which the domain can be called to mind. You don't want to be the company with the terrific website that no one can ever remember to tell their friends about because they can't remember the domain name.

Keep the name as short as possible

Short names are easy to type and easy to remember (see the previous two rules). Short names also allow more of the URL to display in the SERPs and are a better fit on business cards and other offline media.

Create and fulfill expectations

When someone hears about your domain name for the first time, he should be able to instantly and accurately guess at the type of content he might find there. That's why we love domain names such as nytimes.com (<http://www.nytimes.com>), CareerBuilder.com (<http://www.careerbuilder.com/>), AutoTrader.com (<http://www.autotrader.com/>), and WebMD.com (<http://www.webmd.com/>). Domains such as Monster.com, Amazon.com, and Zillow.com required far more branding because of their non-intuitive names.

Avoid trademark infringement

This is a mistake that isn't made too often, but it can kill a great domain and a great company when it does. To be sure you're not infringing on anyone's registered trademark with your site's name, visit the U.S. Patent and Trademark office site (<http://www.uspto.gov/trademarks/index.jsp>) and search before you buy. Knowingly purchasing a domain with bad-faith intent that includes a trademarked term is a form of cybersquatting referred to as domain squatting.

Set yourself apart with a brand

Using a unique moniker is a great way to build additional value with your domain name. A "brand" is more than just a combination of words, which is why names such as Mortgageforyourhome.com and Shoesandboots.com aren't as compelling as branded names such as Yelp.com (<http://www.yelp.com>) and Gilt.com (<http://www.gilt.com>).

Reject hyphens and numbers

Both hyphens and numbers make it hard to convey your domain name verbally and fall down on being easy to remember or type. Avoid spelled-out or Roman numerals in domains, as both can be confusing and mistaken for the other.

Don't follow the latest trends

Website names that rely on odd misspellings, multiple hyphens (such as the SEO-optimized domains of the early 2000s), or uninspiring short adjectives (such as "top...x," "best...x," "hot...x") aren't always the best choice. This isn't a hard and fast rule, but in the world of naming conventions in general, if everyone else is doing it, that doesn't mean it is a surefire strategy. Just look at all the people who named their businesses "AAA... x" over the past 50 years to be first in the phone book; how many Fortune 1000s are named "AAA company?"

Use a domain selection tool

Websites such as Nameboy (<http://www.nameboy.com>) make it exceptionally easy to determine the availability of a domain name. Just remember that you don't have to buy through these services. You can find an available name that you like, and then go to your registrar of choice. You can also try

buydomains.com (<http://www.buydomains.com/>) as an option to attempt to purchase domains which of already been registered.

Picking the Right URLs

Search engines place some weight on keywords in your URLs. Be careful, however, as the search engines can interpret long URLs with numerous hyphens in them (e.g., Buy-this-awesome-product-now.html) as a spam signal. What follows are some guidelines for selecting optimal URLs for the pages of your site(s).

Describe your content

An obvious URL is a great URL. If a user can look at the Address bar (or a pasted link) and make an accurate guess about the content of the page before ever reaching it, you've done your job. These URLs get pasted, shared, emailed, written down, and yes, even recognized by the engines.

Keep it short

Brevity is a virtue. The shorter the URL, the easier to copy and paste, read over the phone, write on a business card, or use in a hundred other unorthodox fashions, all of which spell better usability and increased branding. Remember, however, that you can always create a shortened URL for marketing purposes that redirects to the destination URL of your content – just know that this short URL will have no SEO value.

Static is the way

The search engines treat static URLs differently than dynamic ones. Users also are not fond of URLs in which the big players are ?, &, and =. They are just harder to read and understand.

Descriptive text is better than numbers

If you're thinking of using "114/cat223/" you should go with "/brand/adidas/" instead. Even if the descriptive isn't a keyword or is not particularly informative to an uninitiated user, it is far better to use words when possible. If nothing else, your team members will thank you for making it that much easier to identify problems in development and testing.

Keywords never hurt

If you know you're going to be targeting a lot of competitive keyword phrases on your website for search traffic, you'll want every advantage you can get. Keywords are certainly one element of that strategy, so take the list from marketing, map it to the proper pages, and get to work. For dynamically created pages through a CMS, create the option of including keywords in the URL.

Subdomains aren't always the answer

First off, never use multiple subdomains (e.g. product.brand.site.com) - they are unnecessarily complex and lengthy. Second, consider that subdomains have the potential to be treated separately from the primary domain when it comes to passing link and trust value. In most cases where just a few subdomains are used and there's good interlinking, it won't hurt, but be aware of the downsides. For more on this, and for a discussion of when to use subdomains, see "Root Domains, Subdomains, and Microsites".

Fewer folders

A URL should contain no unnecessary folders (or words or characters, for that matter). They do not add to the user experience of the site and can in fact confuse users.

Hyphens separate best

When creating URLs with multiple words in the format of a phrase, hyphens are best to separate the terms (e.g., /brands/dolce-and-gabbana/), but you can also use plus signs (+).

Stick with conventions

If your site uses a single format throughout, don't consider making one section unique. Stick to your URL guidelines once they are established so that your users (and future site developers) will have a clear idea of how content is organized into folders and pages. This can apply globally as well for sites that share platforms, brands, and so on.

Don't be case-sensitive

Since URLs can accept both uppercase and lowercase characters, don't ever, ever allow any uppercase letters in your structure. Unix/Linux-based web servers are case-sensitive, so <http://www.domain.com/Products/widgets/> is technically a different URL from <http://www.domain.com/products/widgets/>. Note that this is not true in Microsoft IIS servers, but there are a lot of Apache web servers out there. In addition, this is confusing to users, and potentially to search engine spiders as well. Google sees any URLs with even a single unique character as unique URLs. So if your site shows the same content on www.domain.com/Products/widgets/ and www.domain.com/products/widgets/ that could be seen as duplicate content. If you have them now, implement a 301 redirect pointing them to all-lowercase versions, to help avoid confusion. If you have a lot of type-in traffic, you might even consider a 301 rule that sends any incorrect capitalization permutation to its rightful home.

Don't append extraneous data

There is no point in having a URL exist in which removing characters generates the same content. You can be virtually assured that people on the Web will figure it out, link to you in different fashions, confuse themselves, their readers, and the search engines (with duplicate content issues), and then complain about it.

Keyword Targeting

The search engines face a tough task: based on a few words in a query (sometimes only one) they must return a list of relevant results ordered by measures of importance, and hope that the searcher finds what he or she is seeking. As website creators and web content publishers, you can make this process massively simpler for the search engines and, in turn, benefit from the enormous traffic they send, based on how you structure your content. The first step in this process is to research what keywords related to your business people use when searching for businesses that offer products and services like yours.

This practice has long been a critical part of search engine optimization, and although the role keywords play has evolved over time, keyword usage is still one of the first steps in targeting search traffic.

The first step in the keyword targeting process is uncovering popular terms and phrases that searchers regularly use to find the content, products, or/ services your site offers. There's an art and science to this process, but it consistently begins with a list of keywords to target (see Chapter 5 for more on this topic).

Once you have that list, you'll need to include these in your pages. In the early days of SEO, the process involved stuffing keywords repetitively into every HTML tag possible. Now, keyword relevance is much more aligned with the usability of a page from a human perspective.

Since links and other factors make up a significant portion of the search engines' algorithms, they no longer rank pages with 61 instances of "free credit report" above pages that contain only 60. In fact, *keyword stuffing*, as it is known in the SEO world, can actually get your pages devalued via search engine penalties. The engines don't like to be manipulated, and they recognize keyword stuffing as a disingenuous tactic. Figure 6-19 shows an example of a page utilizing accurate keyword targeting.

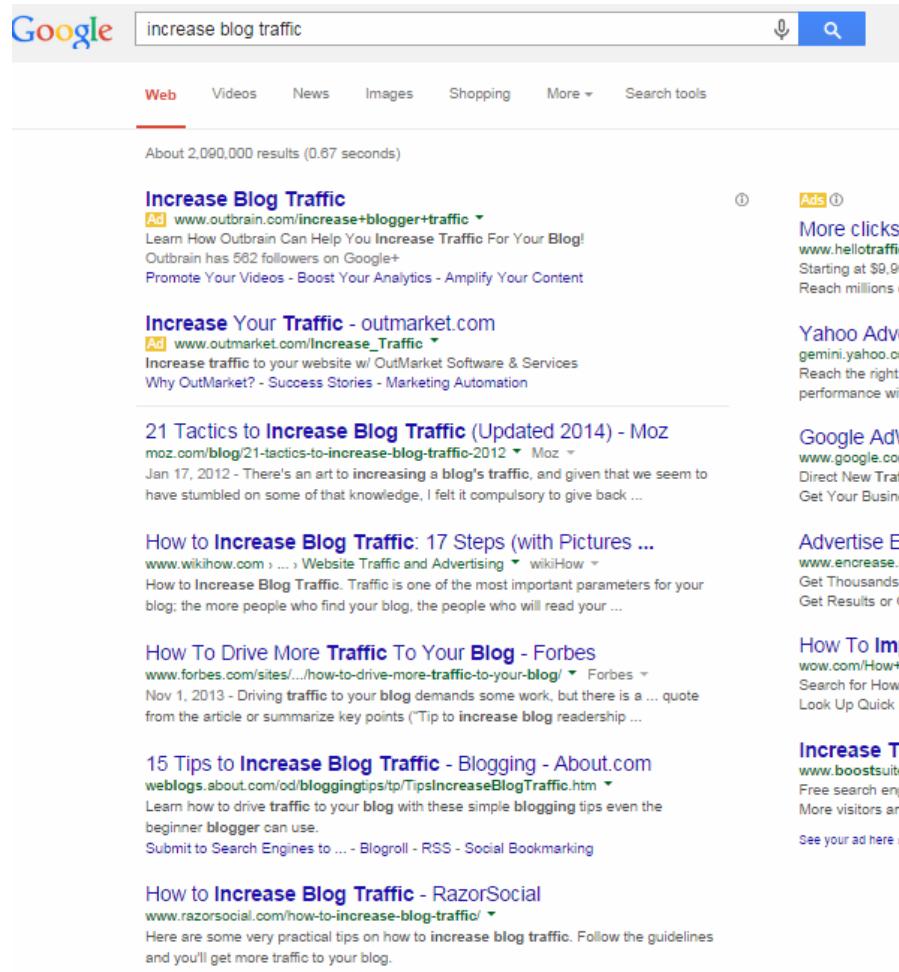


Figure 6-17. Title and headings tags—powerful for SEO

Keyword usage includes creating titles, headlines, and content designed to appeal to searchers in the results (and entice clicks), as well as building relevance for search engines to improve your rankings. In today's SEO, there are also many other factors involved in ranking, including: Term Frequency – Inverse Document Frequency (TF-IDF), Co-Occurrence, Entity Salience, page segmentation, and several others, which will be described in detail later on in this chapter.

However keywords remain important, and building a search-friendly site requires that the keywords that searchers use to find content are prominently employed. Here are some of the more prominent places where a publisher can place those keywords.

HTML Title Tags

For keyword placement, title tags are an important element for search engine relevance. The title tag is in the `<head>` section of an HTML document, and is the only piece of “meta” information about a page that directly influences relevancy and ranking.

The following nine rules represent best practices for title tag construction. Do keep in mind, however, that a title tag for any given page must directly correspond to that page's content. You may have five different keyword categories and a unique site page (or section) dedicated to each, so be sure to align a page's title tag content with its actual visible content as well.

Place your keywords at the beginning of the title tag

This provides the most search engine benefit, and thus, if you want to employ your brand name in the title tag, place it at the end. There is a tradeoff here between SEO benefit and branding benefit that you should think about and make explicitly. Major brands may want to place their brand at the start of the title tag as it may increase click-through rates. To decide which way to go you need to consider which need is greater for your business.

Limit length to 50 characters (including spaces)

Content in title tags after 50 characters is probably given less weight by the search engines. In addition, the display of your title tag in the SERPs may get cut off as early as 49 characters (<http://moz.com/blog/new-title-tag-guidelines-preview-tool>).

There is no hard and fast rule as to how many characters Google will display. Google now truncates the display after a certain number of pixels and so the exact characters you use may vary in width. At the time of this writing this width varies from 482px to 552px depending on your operating system and platform (<http://www.screamingfrog.co.uk/an-update-on-pixel-width-in-google-serp-snippets/>).

Also be aware that Google may not use your title tag in the SERPs. Google frequently chooses to modify your title tag based on several different factors which are beyond your control. If this is happening to you, it may be an indication that Google thinks that your title tag does not accurately reflect the contents of the page, and you should probably consider updating either your title tags or your content.

Incorporate keyword phrases

This one may seem obvious, but it is critical to prominently include in your title tag whatever your keyword research shows as being the most valuable for capturing searches.

Target longer phrases if they are relevant

When choosing what keywords to include in a title tag, use as many as are completely relevant to the page at hand while remaining accurate and descriptive. Thus, it can be much more valuable to have a title tag such as “SkiDudes | Downhill Skiing Equipment & Accessories” rather than simply “SkiDudes | Skiing Equipment”—including those additional terms that are both relevant to the page and receive significant search traffic can bolster your page’s value.

However, if you have separate landing pages for “skiing accessories” versus “skiing equipment,” don’t include one term in the other’s title. You’ll be cannibalizing your rankings by forcing the engines to choose which page on your site is more relevant for that phrase, and they might get it wrong. We will discuss the cannibalization issue in more detail shortly.

Use a divider

When splitting up the brand from the descriptive, options include / (a.k.a. the pipe), >, -, and :, all of which work well. You can also combine these where appropriate—for example, “Major Brand Name: Product Category - Product”. These characters do not bring an SEO benefit, but they can enhance the readability of your title.

Focus on click-through and conversion rates

The title tag is exceptionally similar to the title you might write for paid search ads, only it is harder to measure and improve because the stats aren’t provided for you as easily. However, if you target a market that is relatively stable in search volume week to week, you can do some testing with your title tags and improve the click-through.

Watch your analytics and, if it makes sense, buy search ads on the page to test click-through and conversion rates of different ad text as well, even if it is for just a week or two. You can then look at those results and incorporate them into your titles, which can make a huge difference in the long run. A word of warning, though: don’t focus entirely on click-through rates. Remember to continue measuring conversion rates.

Target searcher intent

When writing titles for web pages, keep in mind the search terms your audience employed to reach your site. If the intent is browsing or research-based, a more descriptive title tag is appropriate. If you're reasonably sure the intent is a purchase, download, or other action, make it clear in your title that this function can be performed at your site. Here is an example from <http://www.bestbuy.com/site/video-games/playstation-4-ps4/pcmcat295700050012.c?id=pcmcat295700050012>. The title tag of that page is "PlayStation 4: Buy PS4 Games & Consoles-Best Buy". The title tag here makes it clear that you can buy PS4 games and consoles at Best Buy.

Communicate with Human Readers

This needs to remain a primary objective. Even as you follow the other rules here to create a title tag that is useful to the search engines, remember that humans will likely see your title tag presented in the search results for your page. Don't scare them away with a title tag that looks like it's written for a machine.

Be consistent

Once you've determined a good formula for your pages in a given section or area of your site, stick to that regimen. You'll find that as you become a trusted and successful "brand" in the SERPs, users will seek out your pages on a subject area and have expectations that you'll want to fulfill.

Meta Description Tags

Meta descriptions have three primary uses:

- To describe the content of the page accurately and succinctly
- To serve as a short text "advertisement" to click on your pages in the search results
- To display targeted keywords, not for ranking purposes, but to indicate the content to searchers

Great meta descriptions, just like great ads, can be tough to write, but for keyword-targeted pages, particularly in competitive search results, they are a critical part of driving traffic from the engines through to your pages. Their importance is much greater for search terms where the intent of the searcher is unclear or different searchers might have different motivations.

Here are seven good rules for meta descriptions:

Tell the truth

Always describe your content honestly. If it is not as "sexy" as you'd like, spice up your content; don't bait and switch on searchers, or they'll have a poor brand association.

Keep it succinct

Be wary of character limits—currently Google displays as few as 140 characters, Yahoo! up to 165, and Bing up to 200+ (they'll go to three vertical lines in some cases). Stick with the smallest—Google—and keep those descriptions at 140 characters (including spaces) or less.

Write ad-worthy copy

Write with as much sizzle as you can while staying descriptive, as the perfect meta description is like the perfect ad: compelling and informative.

Analyze psychology

Unlike an ad, the motivation for an organic-search click is frequently very different from that of users clicking on paid results. Users clicking on PPC ads may be very directly focused on making a purchase, and people who click on an organic result may be more interested in research or learning about the company. Don't assume that successful PPC ad text will make for a good meta description (or the reverse).

Include relevant keywords

It is extremely important to have your keywords in the meta description tag—the boldface that the engines apply can make a big difference in visibility and click-through rate. In addition, if the user's search term is not in the meta description, chances are reduced that the meta description will be used as the description in the SERPs.

Don't employ descriptions universally

You shouldn't always write a meta description. Conventional logic may hold that it is usually wiser to write a good meta description yourself to maximize your chances of it being used in the SERPs, rather than let the engines build one out of your page content; however, this isn't always the case. If the page is targeting one to three heavily searched terms/phrases, go with a meta description that hits those users performing that search.

However, if you're targeting longer tail traffic with hundreds of articles or blog entries or even a huge product catalog, it can sometimes be wiser to let the engines themselves extract the relevant text. The reason is simple: when engines pull, they always display the keywords (and surrounding phrases) that the user searched for. If you try to force a meta description, you can detract from the relevance that the engines make naturally. In some cases, they'll overrule your meta description anyway, but since you can't consistently rely on this behavior, opting out of meta descriptions is OK (and for massive sites, it can save hundreds or thousands of man-hours). Because the meta description isn't a ranking signal it is a second order activity at any rate.

Heading (H1, H2, H3) Tags

The H(x) tags in HTML (H1, H2, H3, etc.) are designed to indicate a headline hierarchy in a document. Thus, an H1 tag might be considered the headline of the page as a whole, whereas H2 tags would serve as subheadings, H3s as tertiary-level headlines, and so forth. The search engines have shown a slight preference for keywords appearing in heading tags. Generally when there are multiple heading tags on a page the engines will weight the higher level heading tags heavier than those below them. For example, if the page contains H1, H2, and H3 tags the H1 will be weighted the heaviest. If a page only contains H2 and H3 tags the H2 would be weighted the heaviest.

In some cases, you can use the title tag of a page, containing the important keywords, as the H1 tag. However, if you have a longer title tag, you may want to use a more focused, shorter heading tag using the most important keywords from the title tag. When a searcher clicks a result at the engines, reinforcing the search term he just typed in with the prominent headline helps to indicate that he has arrived on the right page with the same content he sought.

Many publishers assume that they have to use an H1 tag on every page. It's likely that what matters most is the highest level heading tag you use on a page, and its placement. If you have a page that uses an H3 heading at the very top, and any other heading tags further down on the page are H3, or lower level, than that H3 tag will carry just as much weight as if that first tag had been an H1.

What matters most is the semantic markup of the page, and the first heading tag presumably is intended to be a label for the entire page (so it plays a complementary role to the title tag), and you should treat it as such. Other heading tags on the page should be used to label subsections of the content.

It's also common belief that the size at which the heading tag is displayed is a factor as well. For the most part, the styling of your heading tags is not a factor in the SEO weight of the heading tag. You can style the tag however you want, as shown in Figure 6-20, provided that you don't go to extremes (since it acts as a title for the whole page it should probably be the largest text element on the page).



Figure 6-20. Headings styled to match the site

Document Text

The HTML text on a page was once the center of keyword optimization activities. In the early days of SEO, metrics such as keyword density and keyword saturation were used to measure the perfect level of keyword usage on a page. To the search engines, however, text in a document, particularly the frequency with which a particular term or phrase is used, has very little impact on how happy a searcher will be with that page.

In fact, quite often a page laden with repetitive keywords attempting to please the engines will provide a very poor user experience, and this can result in lower rankings instead of higher one. It's much more valuable to create semantically rich content that covers the topic matter implied by the page's title tag in a comprehensive way. This means naturally including synonyms, and covering related topic areas in a manner that increases the chances of satisfying the needs of a large percentage of visitors to that page. It's a good idea to use the main keyword for a page in the title tag and the main heading tag for a page. It might also appear in the main content, but the use of synonyms for the main keyword and related concepts is at least as important. As a result, it's more important to focus on creating high quality content than it is to keep repeating the main keyword.

Term Frequency – Inverse Document Frequency (TF-IDF)

TF-IDF is also driven by keyword usage, but adds one additional important concept, which is the frequency of use of a keyword on a page and how that compares to the frequency of that term on other similar pages. In other words, if you have a page about “left handed golf clubs”, does that phrase appear on your page more than other pages than others on the same topic?

TF-IDF still relates closely to keyword density, but adds the simple notion of how your page compares to other documents on the same topic. Focusing on TF-IDF alone is not recommended, as it will still potentially lead you to simply stuffing the main phrase on your page with excessive frequency, creating a poor user experience, and potentially harming your search results.

Page Segmentation

It used to be that Google could not understand the layout of a page that well, simply because it could not read CSS files and process them like a browser does. However, that has changed, as documented in this post on the Google Webmaster Central Blog: <http://googlewebmastercentral.blogspot.com/2014/10/updating-our-technical-webmaster.html>.

As a result, Google is quite likely to fully understand the layout of your pages. Given this, where the keywords are used on the page also matters. Use of keywords in your left or right sidebar, or your footer, probably matters less than content used in the main body of your content.

In addition, with HTML5, new markup exists that allows you to explicitly identify the section of your page that represents the main content. You can use this to help make Google's job easier, and to make sure that other search engines are able to locate that content.

Synonyms

Use of related terms is also a factor. A page about “left handed golf clubs” should not use that exact phrase every time the product is referenced. This would not be a natural way of writing, and could be used by the search engines as a signal of poor document quality, causing the rankings of the page to be lowered.

Instead, allow your content creators to write naturally. This will cause them to use other phrases, such as “the sticks”, “set of clubs”, “lefty clubs”, and other variants that people use in normal writing style.

Use of synonyms represents a key step in moving away from manipulative SEO techniques for creating pages to try and rank for specific search terms.

Co-Occurrence, Phrase-Based Indexing and Entity Salience

The presence of a specific keyword or phrase on a page will likely increase the probability of finding other words on those pages. For example, if you are reading an article on the life of New England Patriots Quarterback Tom Brady, you would expect to see mentions of Gisele Bundchen, or his children. The absence of a mention of these people in his life could be a signal of a poor quality page.

To put this in the positive, inclusion of more information that you would expect to see in the content can be used as an indication that it's a good page. Chances are good that more people reading the article will be satisfied with it as well, as they might expect to learn about his family.

Of course, a high quality article will probably talk about his parents, his high school football coach, his sisters, and all the other aspects of his life as well. The key is to focus on providing a more complete response to the topic than others covering the same topic might do.

?????? I will provide an image here

Not all of the topic needs to be addressed on each individual page. Linking to other relevant resources and high quality content, both on your site, as well as on 3rd party sites can play a key role in establishing your page as a great answer to the user's question.

This last step may well be equally important in the overall page optimization process. No single page will answer every question by every possible user to it, so addressing a significant percentage of questions, and then connecting with other pages to answer follow on questions on the same topic is an optimal structure.

On the product pages of an e-commerce site, where there will not be article style content, this can mean a combination of well structured, and unique description text, and access to key refinements, such as individual brands, related product types, the presence of a privacy policy, about us information, a shopping cart, and more.

Image Filenames and Alt Attributes

Incorporation of images on web pages can substantively enrich the user experience. However, the search engines cannot read the images directly. There are two elements that you can control to give the engines context for images:

The filename

Search engines look at the image filename to see whether it provides any clues to the content of the image. Don't name your image *example.com/img4137a-b12.jpg*, as it tells the search engine nothing at all about the image, and you are passing up the opportunity to include keyword-rich text.

If it is a picture of Abe Lincoln, name the file *abe-lincoln.jpg* and/or have the SRC URL string contain it, as in *example.com/abe-lincoln/portrait.jpg*.

Image alt text

Image tags in HTML permit you to specify an attribute known as the `alt` attribute. This is a place where you can provide more information about what is in the image, and again where you can use your targeted keywords. Here is an example for the picture of Abe Lincoln:

```
| 
```

Use the quotes if you have spaces in the text string of the `alt` content! Sites that have invalid `img` tags frequently lump a few words without quotes into the `img` tag, intended for the `alt` content—but with no quotes, all terms after the first word will be lost.

This usage of the image filename and of the `alt` attribute permits you to reinforce the major keyword themes of the page. This is particularly useful if you want to rank in image search. Make sure the filename and the `alt` text reflect the content of the picture, and do not artificially emphasize keywords unrelated to the image (even if they are related to the page). Although the `alt` attribute and the image filename are helpful, you should not use image links as a substitute for text links with rich anchor text that carries much more weight from an SEO perspective.

Presumably, your picture will relate very closely to the content of the page, and using the image filename and the `alt` text will help reinforce the page's overall theme.

Boldface Text

While it used to be true that including keywords in bold text had a very slight effect in rankings this is no longer the case.

Avoiding Keyword Cannibalization

As we discussed earlier, you should not use common keywords across multiple page titles. This advice applies to more than just the title tags.

One of the nastier problems that often crops up during the course of a website's information architecture, *keyword cannibalization* refers to a site's targeting of popular keyword search phrases on multiple pages, forcing the engines to pick which one is most relevant. In essence, a site employing cannibalization competes with itself for rankings and dilutes the ranking power of internal anchor text, external links, and keyword relevancy.

Avoiding cannibalization requires strict site architecture with attention to detail. Plot out your most important terms on a visual flowchart (or in a spreadsheet file, if you prefer), and pay careful attention to what search terms each page is targeting. Note that when pages feature two-, three-, or four-word phrases that contain the target search phrase of another page, linking back to that page within the content with the appropriate anchor text will avoid the cannibalization issue.

For example, if you had a page targeting “mortgages” and another page targeting “low-interest mortgages,” you would link back to the “mortgages” page from the “low-interest mortgages” page using the anchor text “mortgages” (see Figure 6-21). You can do this in the breadcrumb or in the body copy. The *New York Times* (<http://www.nytimes.com/>) does the latter, where keywords in the body copy link to the related resource page on the site.

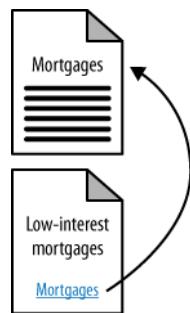


Figure 6-18. Adding lots of value with relevant cross-links

Keyword Targeting in CMS' and Automatically Generated Content

Large-scale publishing systems, or those that produce automatically generated content, present some unique challenges. If hundreds of pages are being created every day, it is not feasible to do independent keyword research on each and every page, making page optimization an interesting challenge.

In these scenarios, the focus turns to methods/recipes for generating unique titles, heading tags, and page content for each page. It is critical to educate the writers on ways to implement titles and headings that capture unique, key aspects of the articles' content. More advanced teams can go further with this and train their writing staff on the use of keyword research tools to further optimize this process.

In the case of automatically generated material (such as that produced from algorithms that mine data from larger textual bodies), the key is to automate means for extracting a short (fewer than 55 characters) description of the article and making it unique from other titles generated elsewhere on the site and on the Web at large.

Encouraging Effective Keyword Targeting by Content Creators

Very frequently, someone other than an SEO professional is responsible for content creation. Content creators often do not have an innate knowledge as to how SEO works, or worse, they may think they know how it works, but have the wrong idea about it.. Some Training for your writers is a critical activity. This is particularly important when dealing with large websites and large teams of writers.

Here are the main components of web page copywriting that your writers must understand:

- Search engines look to match up a user's search queries with the keyword phrases, their synonyms, and related concepts on your web pages. If some combination of all of these do not appear on the page chances are good that your page will never achieve significant ranking for those search phrases.
- The search phrases users may choose to use when looking for something are infinite in variety, but certain phrases will be used much more frequently than others.
- Using the more popular phrases you wish to target on a web page in the content for that page is essential to SEO success for that page.
- Make sure that the writers understand the concepts of Co-occurrence and Entity Salience discussed earlier in this chapter, so they don't create content that uses the main keyword to an excessive degree. They need to focus on creating semantically rich content that stays on the topic of the main target keyword phrase for the page, while writing naturally.
- The title tag is the most important element on the page. Next up is the first header (H1), and then the main body of the content.
- Tools exist (as outlined in Chapter 5) that allow you to research and determine what the most interesting phrases are.

If you can get these six points across, you are well on your way to empowering your content creators to perform solid SEO. The next key element is training them on how to pick the right keywords to use.

This can involve teaching them how to use keyword research tools similar to the ones we discussed in Chapter 5, or having the website's SEO person do the research and provide the terms to the writer.

The most important factor to reiterate to the content creator is that content quality and user experience still come first. Then, by intelligently making sure the right keywords and phrases are properly used throughout the content, they can help bring search engine traffic to your site. Reverse these priorities and you can end up with keyword stuffing or other spam issues.

Long Tail Keyword Targeting

As we outlined in Chapter 5, the small-volume search terms, when tallied up, represent 70% or more of all search traffic, and the more obvious, high-volume terms represent only 30% of the overall search traffic.

For example, if you run a site targeting searches for *new york pizza* and *new york pizza delivery*, you might be surprised to find that hundreds of single searches each day for terms such as *pizza delivery on the corner of 57th & 7th*, or *Manhattan's tastiest Italian-style sausage pizza*, when taken together, will actually provide considerably more traffic than the popular phrases you've researched. This concept is called the *long tail of search*.

Targeting the long tail is another aspect of SEO that combines art and science. In Figure 6-22, you may not want to implement entire web pages for a history of pizza dough, pizza with white anchovies, or Croatian pizza.

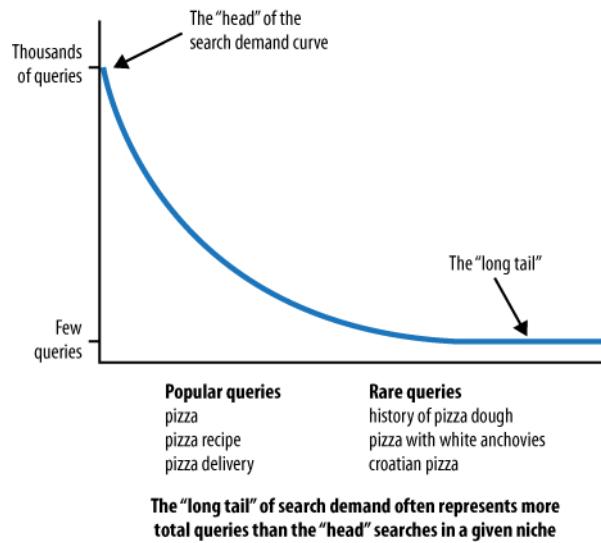


Figure 6-19. Example of the long tail search curve

Finding scalable ways to chase long tail keywords is a complex topic. It is also one where many publishers get into a lot of trouble, as they think they need to create a new page for each potential search phrase that a user might type in related to their business, and this is not the case. Much of the long tail of search can be addressed by using the right content optimization practices on your site.

Perhaps you have a page for ordering pizza in New York City, and you have a good title and heading tag on the page (e.g., “New York City Pizza: Order Here”), as well as a phone number and a form for ordering the pizza, and no other content. If that is all you have, that page is not competing effectively for rankings on long tail search terms. To fix this, you need to write additional content for the page. Ideally, this would be content that talks about the types of pizza that are popular in New York City, the ingredients used, and other things that might draw in long tail search traffic.

If you have a page for San Jose pizza, the picture gets even more complicated. You don’t want your content on the San Jose page to be the same as it is on the New York City page. This runs into potential duplicate content problems, as we will outline in “Duplicate Content Issues”, or the keyword cannibalization issues we discussed earlier in this chapter.

To maximize your success, find a way to generate different content for those two pages, ideally tuned to the specific needs of the audience that arrives at those pages. Perhaps the pizza preferences of the San Jose crowd are different from those in New York City. Of course, the geographic information is inherently different between the two locations, so driving directions from key locations might be a good thing to include on the page.

If you have pizza parlors in 100 cities, this can get very complex indeed. The key here is to remain true to the diverse needs of your users, yet use your knowledge of the needs of search engines and searcher behavior to obtain that long tail traffic. ????????

Content Optimization

Content optimization relates to how the presentation and architecture of the text, image, and multimedia content on a page can be optimized for search engines. Many of these recommendations are second-order effects. Having the right formatting or display won't boost your rankings directly, but through it, you're more likely to earn links, get clicks, and eventually benefit in search rankings. If you regularly practice the techniques in this section, you'll earn better consideration from the engines and from the human activities on the Web that influence their algorithms.

Content Structure

Because SEO has become such a holistic part of website development and improvement, it is no surprise that *content formatting*—the presentation, style, and layout choices you select for your content—is a part of the process. Choosing a browser-safe sans serif font such as Arial and Helvetica is a wise choice for the Web; Verdana in particular has received high praise from usability/readability experts, such as that which WebAIM offered in an article posted at <http://webaim.org/techniques/fonts/>

Verdana is one of the most popular of the fonts designed for on-screen viewing. It has a simple, straightforward design, and the characters or glyphs are not easily confused. For example, the uppercase *I* and the lowercase *L* have unique shapes, unlike in Arial, in which the two glyphs may be easily confused (see Figure 6-23).



Figure 6-20. Arial versus Verdana font comparison

Another advantage of Verdana is the amount of spacing between letters. One consideration to take into account with Verdana is that it is a relatively large font. The words take up more space than words in Arial, even at the same point size (see Figure 6-24).

Verdana: Notice the length of this sentence.

Arial: Notice the length of this sentence.

Figure 6-21. How fonts impact space requirements

The larger size improves readability but also has the potential of disrupting carefully planned page layouts.

Font choice is accompanied in importance by sizing and contrast issues. Type that is smaller than 10 points is typically very challenging to read, and in all cases, relative font sizes are recommended so that users can employ browser options to increase/decrease if necessary. Contrast—the color difference between the background and text—is also critical; legibility usually drops for anything that isn't black (or very dark) on a white background.

Content length and word count

Content length is another critical piece of the optimization puzzle that's mistakenly placed in the "keyword density" or "unique content" bucket of SEO. In fact, content length can have a big role to play in terms of whether your material is easy to consume and easy to share.

People often ask about what the ideal length is for a piece of content. The reality is that the perfect length for a piece of content is determined by the nature of the topic being addressed. Many pieces of content do well because they are short and very easy to consume. On the other hand, some content will fare best when it's lengthy and comprehensive in nature.

Longer articles also have the opportunity to show up in Google's in-depth articles section (<http://moz.com/blog/inside-indepth-articles>). In some cases, where appropriate, Google will feature several

longer more detailed articles on a given topic. For more information on having your articles appear as “in-depth” articles see <https://support.google.com/webmasters/answer/3280182?hl=en>.

Visual layout

Last but not least in content structure optimization is the display of the material. Beautiful, simplistic, easy-to-use, and consumable layouts instill trust and garner far more readership and links than poorly designed content wedged between ad blocks that threaten to overtake the page. For more on this topic, you might want to check out “The Golden Ratio in Web Design” from NetTuts (<http://code.tutsplus.com/tutorials/the-golden-ratio-in-web-design--net-2272>), which has some great illustrations and advice on laying out web content on the page.

CSS and Semantic Markup

CSS is commonly mentioned as a best practice for general web design and development, but its principles provide some indirect SEO benefits as well. Google used to recommend keeping pages smaller than 101 KB, and it used to be a common belief that there were benefits to implementing pages that were small in size. Now, however, search engines deny that code size is a factor at all, unless it is extreme. Still, keeping file size low means faster load times, lower abandonment rates, and a higher probability of being fully read and more frequently linked to. This is particularly important in mobile environments.

It also used to be that the search engines could not read CSS code and render pages in the same manner as a browser does. Google made it clear in October 2014 that it is able to do that (<http://googlewebmastercentral.blogspot.com/2014/10/updating-our-technical-webmaster.html>), so good clean page layout as setup by your CSS could potentially be considered as a factor in evaluating page quality.

Your experience may vary, but since good CSS makes it easy, there's no reason not to make it part of your standard operating procedure for web development. Use table-less CSS stored in external files, keep JavaScript calls external, and separate the content layer from the presentation layer as shown on CSS Zen Garden (<http://www.csszengarden.com/>), a site that offers many user-contributed style sheets formatting the same HTML content.

You can use CSS code to provide emphasis, to quote/reference, and to reduce the use of tables and other bloated HTML mechanisms for formatting, which can make a positive difference in your SEO. Be sure to allow Googlebot access to your CSS files.

Google, Bing, and Yahoo have come together to create a standard for markup called Schema.org (<http://www.schema.org>). You can see a copy of the announcement at this URL "<http://googleblog.blogspot.com/2011/06/introducing-schemaorg-search-engines.html>". This represented a new level of commitment by the search engines to the concept of marking up content, or more broadly, in allowing the publisher to provide information about the content to the search engines. When we refer to marking up content we refer to the concept of tagging your content using XML tags that categorize the contents of a block of content. For example, you may label a block of content as containing a recipe, and another block of content as containing a review.

This notion of advanced markup was not new, as all of the search engines have supported semantic markup at a limited level, and have used this markup to show *rich snippets*. You can see an example of a *rich snippet* in Figure 6-25.

One of the original ways a publisher had to communicate information about a web page to search engines was with metatags. Unfortunately, these were so badly abused by spammers that Google stopped using them as a ranking signal. Google confirmed this publicly in a post in 2009 which notes that "Google has ignored the keywords meta tag for years and currently we see no need to change that policy" (<http://googlewebmastercentral.blogspot.com/2009/09/google-does-not-use-keywords-meta-tag.html>).

Google continues to indicate that markup is not used as a ranking signal: "Google doesn't use markup for ranking purposes at this time" (<https://support.google.com/webmasters/answer/1211158>). However, there are important SEO benefits to using markup.

Markup in Search Results

As previously mentioned, markup is sometimes used by search engines to create a *rich snippet*. Here is an example of *rich snippets* in the search results, as shown on a search for a recipe for a Cambodian dish called Loc Lac:

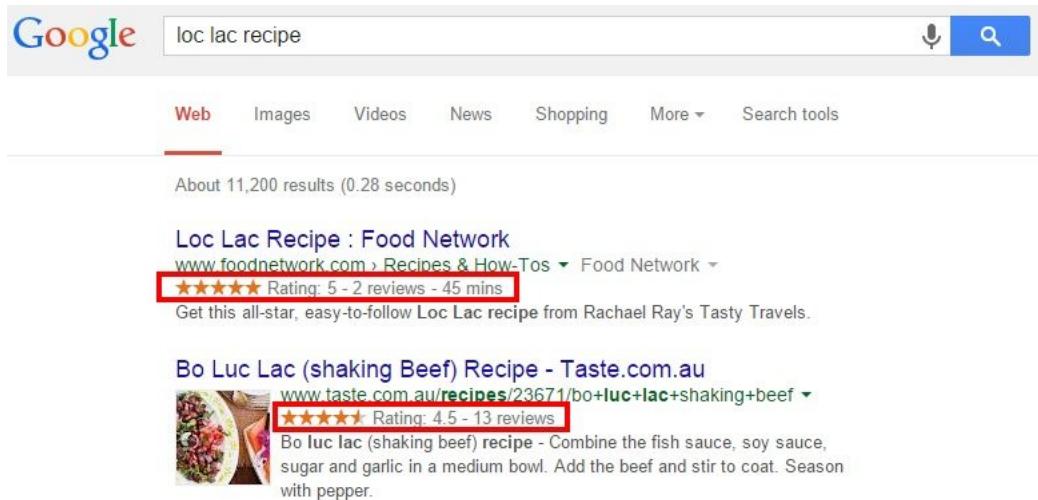


Figure 6-225. Example of Recipe Rich Snippet on Google

Based on the markup that Google found in the HTML, it has enhanced the result by showing the results of the reviews (the number of stars), the required cooking time and the calories of the meal. The type of markup used for this example is called *Microformats*. Figure 6-26 shows what the source code looks like for this example:

```
<div id="fn-w" class="hrecipe"><div class="breadcrumb clearfix">
  <div class="bc-links">
    <a href="/">Home</a>
    <span>&gt;</span>
    <a href="/recipes/index.html">Recipes</a>
    <span>&gt;</span>
    <a href="/recipe-collections/holidays-and-parties/index.html">Holidays and Parties</a>
    <span>&gt;</span>
  </div>
  <div class="bc-desc">Loc Lac Recipe</div>
</div>
```

Figure 6-236. Sample of Microformats code for a Recipe

Supported Types of Markup

There are a few different standards for markup. The most common ones are *microdata*, *Microformats*, and *RDFa*. Schema.org is based off of the *microdata* standard. However, the search engines have implemented *rich snippets* based on some (but not all) aspects of *Microformats* prior to the announcement of Schema.org, and they will likely continue support for these for some period of time.

It is likely that any new forms of *rich snippets* implemented by the search engines will be based off of schema.org (*microdata*) and not *Microformats* or *RDFa*. Some of the formats already supported by Google include:

- People - <https://support.google.com/webmasters/answer/146646>
- Products - <https://support.google.com/webmasters/answer/146750>
- Events - <https://support.google.com/webmasters/answer/164506>
- Business and Organizations - <https://support.google.com/webmasters/answer/146861>

- Video - <https://support.google.com/webmasters/answer/162163>

Impact of Rich Snippets

The key reason that the search engines are pursuing *rich snippets* is that they have done extensive testing that has shown them that they can increase click through rates. Searchers like seeing more information about the page in the search results. Based on this you can expect that they will continue to implement support for more of these types of search result enhancements based on markup.

From an SEO perspective, increasing click through rate is highly desirable. It brings us more relevant traffic. In addition, we know that search engines measure user interaction with the search results and that click through rate is a ranking factor. This was first publicly confirmed in this interview with Bing's Duane Forrester: <http://www.stonetemple.com/search-algorithms-and-bing-webmaster-tools-with-duane-forrester>.

So while the search engines do not directly use semantic markup as a ranking signal, the indirect impact of *rich snippets* providing a higher click through rate acts as a ranking signal.

For more information on semantic markup see the sections on “Semantic Search” and “Schema.org” in this chapter.

Content Uniqueness and Depth

Few can debate the value the engines place on robust, unique, value-added content—Google in particular has had several rounds of kicking “low-quality-content” sites out of its indexes, and the other engines have followed suit.

The first critical designation to avoid is “thin content”—a phrase that (loosely) refers to content the engines do not feel contributes enough unique material to display a page competitively in the search results. How much content is enough content to not be considered thin? The criteria have never been officially listed, but many examples/discussions from engineers and search engine representatives would place the following on the list:

- At least thirty to 50 unique words, forming unique, parse-able sentences that other sites/pages do not have (for many pages much more is appropriate, consider this a minimum).
- Unique HTML text content, different from other pages on the site in more than just the replacement of key verbs and nouns (yes, this means all those webmasters who build the same page and just change the city and state names thinking it is “unique” are mistaken).
- Unique titles and meta description elements. If you can’t write unique meta descriptions, just exclude them. Algorithms can trip up pages and boot them from the index simply for having near-duplicate meta tags.
- Unique video/audio/image content. The engines have started getting smarter about identifying and indexing pages for vertical search that wouldn’t normally meet the “uniqueness” criteria.

By the way, you can often bypass these limitations if you have a good quantity of high-value external links pointing to the page in question (though this is very rarely scalable) or an extremely powerful, authoritative site (note how many one-sentence Wikipedia stub pages still rank).

The next criterion from the engines demands that websites “add value” to the content they publish, particularly if it comes from (wholly or partially) a secondary source.

A word of caution to affiliates

This word of caution most frequently applies to affiliate sites whose republishing of product descriptions, images, and so forth has come under search engine fire numerous times. In fact, it is best to anticipate manual evaluations here even if you’ve dodged the algorithmic sweep.

The basic tenets are:

- Don't simply republish something that's found elsewhere on the Web unless your site adds substantive value to users, and don't infringe on others' copyrights or trademarks.
- If you're hosting affiliate content, expect to be judged more harshly than others, as affiliates in the SERPs are one of users' top complaints about search engines.
- Small things such as a few comments, a clever sorting algorithm or automated tags, filtering, a line or two of text, simple mashups, or advertising do *not* constitute "substantive value."

For some exemplary cases where websites fulfill these guidelines, check out the way sites such as CNET (<http://reviews.cnet.com/>), Urbanspoon (<http://www.urbanspoon.com/>), and Metacritic (<http://www.metacritic.com/>) take content/products/reviews from elsewhere, both aggregating *and* "adding value" for their users.

Last but not least, Google has provided a guideline to refrain from trying to place "search results in the search results." For reference, look at the post from Google's Matt Cutts, including the comments, at <http://www.mattcutts.com/blog/search-results-in-search-results/>. Google's stated position is that search results generally don't "add value" for users, though others have made the argument that this is merely an anticompetitive move.

Sites can benefit from having their "search results" transformed into "more valuable" listings and category/subcategory landing pages. Sites that have done this have had great success recovering rankings and gaining traffic from Google.

In essence, you want to avoid the potential for your site pages being perceived, both by an engine's algorithm and by human engineers and quality raters, as search results. Refrain from:

- Pages labeled in the title or headline as "search results" or "results"
- Pages that appear to offer a query-based list of links to "relevant" pages on the site without other content (add a short paragraph of text, an image, and formatting that make the "results" look like detailed descriptions/links instead)
- Pages whose URLs appear to carry search queries (e.g., ?q=miami+restaurants or ?search=Miami+restaurants versus /miami-restaurants)
- Pages with text such as "Results 1 through 10"

Though it seems strange, these subtle, largely cosmetic changes can mean the difference between inclusion and removal. Err on the side of caution and dodge the appearance of search results.

Content Themes

A less discussed, but important, issue is the fit of each piece of content to your site. If you create an article about pizza, but the rest of your site is about horseshoes, your article is unlikely to rank for the phrase pizza. Search engines analyze and understand what sites, or sections of sites, focus on for topic matter.

You can think of this as being the "theme" of the site (or section). If you start creating content that is not on the same theme, that content will have a very difficult time ranking. Further, your off-topic content could potentially weaken the theme of the rest of the site.

One site can support multiple themes, but each themed section needs to justify its own existence, by following good SEO practices, including getting third parties to implement links from the pages of their sites to that section. Make sure you keep your content on topic and this will help the SEO for all of the pages of your site.

Copyblogger has created a tool to help measure the fit of a given article to your site known as Scribe (<http://www.copyblogger.com/scribe-seo/>). Scribe will measure the fit of an article to your site, as well as a more general look at the consistency of the content across your site overall.

Duplicate Content Issues

Duplicate content generally falls into three categories: exact (or true) duplicates, whereby two URLs output identical content; near duplicates, whereby there are small content differentiators (sentence order, image variables, etc.); and cross-domain duplicates, whereby exact or near duplication exists on multiple domains.

There are two related concepts which are not treated by Google the same as duplicate content, but are often confused by publishers and inexperienced SEO practitioners. These are:

Thin content

These are pages that don't have much content on them at all. An example of this might be a set of pages built out to list all the locations for a business with 5,000 locations, but the only thing on all the pages is the address of each location.

Thin slicing

These types of pages are ones with very minor differences in focus. Consider a site that sells running shoes, and one of the shoes they offer is Men's Nike Air Max LTD Running Shoes, that comes in the following sizes: 6, 6.5, 7, 7.5, 8, ...15. If they chose to have a different page for each size of this shoe, even though they would actually be showing different products on each, there is just not much useful difference between these pages

Google has been clear that they don't like either thin content or thin slicing. Either of these can trigger Google's Panda algorithm which is discussed more in Chapter 8 ("Panda, Penguin, and Penalties"). Exactly how Bing differentiates duplicate content, thin content, and thin slicing is less clear, but they also prefer that publishers avoid creating these types of pages.

Duplicate content can result from many causes, including licensing of content to or from your site, site architecture flaws due to non-SEO-friendly CMSs, or plagiarism. Not too long ago, however, spammers in desperate need of content began the now much-reviled process of scraping content from legitimate sources, scrambling the words (through many complex processes), and repurposing the text to appear on their own pages in the hopes of attracting long tail searches and serving contextual ads (and various other nefarious purposes).

Thus, today we're faced with a world of "duplicate content issues" and "duplicate content penalties." Here are some definitions that are useful for this discussion:

Unique content

This is written by humans, is completely different from any other combination of letters, symbols, or words on the Web, and is clearly not manipulated through computer text-processing algorithms (such as Markov-chain-employed spam tools).

Snippets

These are small chunks of content such as quotes that are copied and reused; these are almost never problematic for search engines, especially when included in a larger document with plenty of unique content.

Shingles

Search engines look at relatively small phrase segments (e.g., five to six words) for the presence of the same segments on other pages on the Web. When there are too many shingles in common between two documents, the search engines may interpret them as duplicate content.

Duplicate content issues

This is typically used when referring to duplicate content that is not in danger of getting a website penalized, but rather is simply a copy of an existing page that forces the search engines to choose which version to display in the index (a.k.a. duplicate content filter).

Duplicate content filter

This is when the search engine removes substantially similar content from a search result to provide a better overall user experience.

Duplicate content penalty

Penalties are applied rarely and only in egregious situations. Engines may devalue or ban other web pages on the site, too, or even the entire website.

Consequences of Duplicate Content

Assuming your duplicate content is a result of innocuous oversights on your developer's part, the search engine will most likely simply filter out all but one of the pages that are duplicates because the search engine wants to display one version of a particular piece of content in a given SERP. In some cases, the search engine may filter out results prior to including them in the index, and in other cases the search engine may allow a page in the index and filter it out when it is assembling the SERPs in response to a specific query. In this latter case, a page may be filtered out in response to some queries and not others.

Searchers want diversity in the results, not the same results repeated again and again. Search engines therefore try to filter out duplicate copies of content, and this has several consequences:

- A search engine bot comes to a site with a crawl budget, which is counted in the number of pages it plans to crawl in each particular session. Each time it crawls a page that is a duplicate (which is simply going to be filtered out of search results) you have let the bot waste some of its crawl budget. That means fewer of your “good” pages will get crawled. This can result in fewer of your pages being included in the search engine index.
- Links to duplicate content pages represent a waste of link authority. Duplicated pages can gain PageRank, or link authority, and since it does not help them rank, that link authority is misspent.
- No search engine has offered a clear explanation for how its algorithm picks which version of a page it does show. In other words, if it discovers three copies of the same content, which two does it filter out? Which one does it still show? Does it vary based on the search query? The bottom line is that the search engine might not favor the version you want.

Although some SEO professionals may debate some of the preceding specifics, the general structure will meet with near-universal agreement. However, there are a handful of caveats to take into account.

For one, on your site you may have a variety of product pages and also offer print versions of those pages. The search engine might pick just the printer-friendly page as the one to show in its results. This does happen at times, and it can happen even if the printer-friendly page has lower link authority and will rank less well than the main product page.

The best potential fix for this is to apply the canonical URL tag to all versions of the page to indicate which version is the original.

A second version of this can occur when you syndicate content to third parties. The problem is that the search engine may filter your copy of the article out of the results in favor of the version in use by the person republishing your article. There are three potential solutions to this:

- Get the person publishing your syndicated content to publish a **canonical** tag back to the original page on your site. This will help indicate to the search engines that your copy of the page is the original, and any links pointing to the syndicated page will be credited to your original instead.
- Have the syndicating partner NoIndex their copy of the content. This will keep the duplicate copy out of the search engine index. In addition, links in that content back to your site will still pass link authority to you.
- Have the partner implement a link back to the original source page on your site. Search engines usually interpret this correctly and emphasize your version of the content when you do that. Note however that there have been instances where Google does attribute the originality of the content to the site

republishing it particularly if that site has vastly more authority and trust than the true original source of the content.

How Search Engines Identify Duplicate Content

Some examples will illustrate the process for Google as it finds duplicate content on the Web. In the examples shown in Figures Figure 6-28 through Figure 6-31, three assumptions have been made:

- The page with text is assumed to be a page that contains duplicate content (not just a snippet, despite the illustration).
- Each page of duplicate content is presumed to be on a separate domain.
- The steps that follow have been simplified to make the process as easy and clear as possible. This is almost certainly not the exact way in which Google performs (but it conveys the effect).

Phase I: Google finds duplicate content

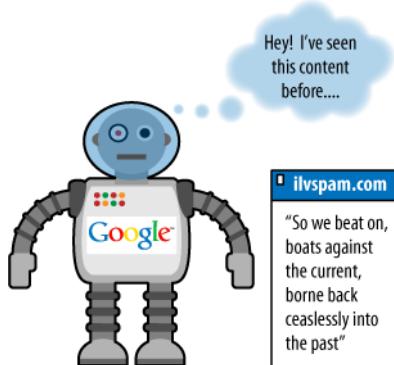


Figure 6-24. Google finding duplicate content

Phase II: Google checks comparable docs

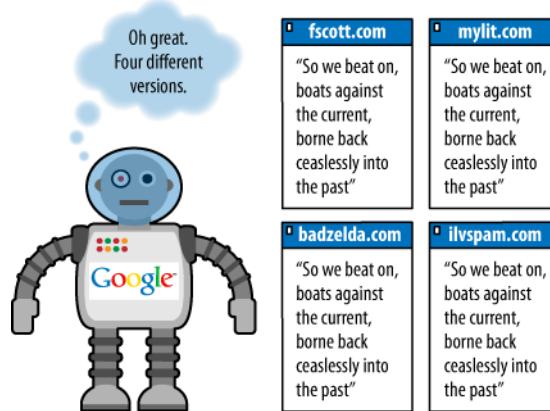


Figure 6-25. Google comparing the duplicate content to the other copies

Phase III: Duplicates get tossed out

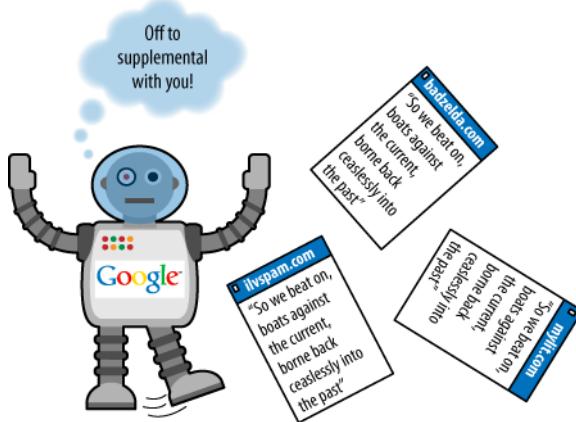


Figure 6-30. Duplicate copies getting filtered out

Phase IV: Google determines an original

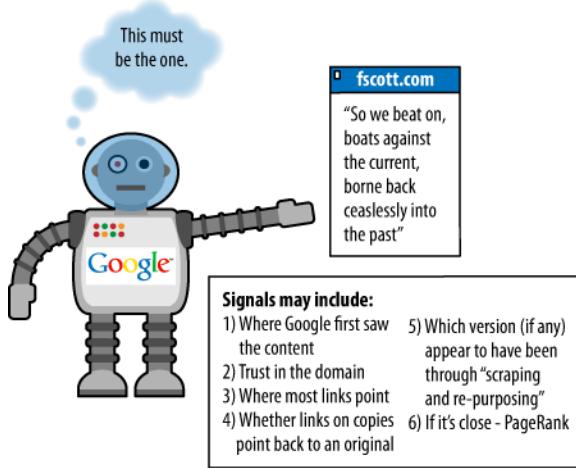


Figure 6-31. Google choosing one as the original

There are a few facts about duplicate content that bear mentioning as they can trip up webmasters who are new to the duplicate content issue:

Location of the duplicate content

Is it duplicated content if it is all on my site? Yes, in fact, duplicate content can occur within a site or across different sites.

Percentage of duplicate content

What percentage of a page has to be duplicated before you run into duplicate content filtering? Unfortunately, the search engines would never reveal this information because it would compromise their ability to prevent the problem.

It is also a near certainty that the percentage at each engine fluctuates regularly and that more than one simple direct comparison goes into duplicate content detection. The bottom line is that pages do not need to be identical to be considered duplicates.

Ratio of code to text

What if your code is huge and there are very few unique HTML elements on the page? Will Google think the pages are all duplicates of one another? No. The search engines do not care about your code;

they are interested in the content on your page. Code size becomes a problem only when it becomes extreme.

Ratio of navigation elements to unique content

Every page on my site has a huge navigation bar, lots of header and footer items, but only a little bit of content; will Google think these pages are duplicates? No. Google and Bing factor out the common page elements such as navigation before evaluating whether a page is a duplicate. They are very familiar with the layout of websites and recognize that permanent structures on all (or many) of a site's pages are quite normal. Instead, they'll pay attention to the "unique" portions of each page and often will largely ignore the rest. Note however that these will almost certainly be considered "thin content" by the engines.

Licensed content

What should I do if I want to avoid duplicate content problems, but I have licensed content from other web sources to show my visitors? Use `meta name = "robots" content="noindex, follow"`. Place this in your page's header and the search engines will know that the content isn't for them. This is a general best practice, because then humans can still visit the page, link to it, and the links on the page will still carry value.

Another alternative is to make sure you have exclusive ownership and publication rights for that content.

Identifying and Addressing Copyright Infringement

One of the best ways to monitor whether your site's copy is being duplicated elsewhere is to use CopyScape.com (<http://copescape.com/>), a site that enables you to instantly view pages on the Web that are using your content. Do not worry if the pages of these sites rank far behind your own pages for any relevant queries—if any large, authoritative, content-rich domain tried to fight all the copies of its work on the Web, it would have at least two 40-hour-per-week jobs on its hands. Luckily, the search engines have placed trust in these types of sites to issue high quality, relevant content, and therefore recognize them as the original issuer.

If, on the other hand, you have a relatively new site or a site with few inbound links, and the scrapers are consistently ranking ahead of you (or someone with a powerful site is stealing your work), you've got some recourse. One option is to simply ask the publisher to remove the offending content. In some cases they are simply unaware that copying your content is not allowed. Another option is to contact the site's hosting company. The hosting company could potentially be liable for hosting duplicate content and so they are frequently quick to react to such inquiries. Just be sure to provide as much documentation as possible to show that the content was originally yours.

Another option is to file a DMCA infringement request with Google, with Yahoo!, and with Bing (you should also file this request with the infringing site's hosting company).

A further option is to file a legal suit (or threaten such) against the website in question. You may want to try to start with a more informal communication asking them to remove the content before you send a letter from the attorneys, as the DMCA motions can take up to several months to go into effect; but if they are nonresponsive, there is no reason to delay taking stronger action, either. If the site republishing your work has an owner in your country, this latter course of action is probably the most effective first step.

A very effective and inexpensive option for this process is DMCA.com (<http://www.dmcacom.com/>).

An actual penalty situation

The preceding examples show duplicate content filters and are not actual penalties, but, for all practical purposes, they have the same impact as a penalty: lower rankings for your pages. But there are scenarios where an actual penalty can occur.

For example, sites that aggregate content from across the Web can be at risk, particularly if little unique content is added from the site itself. In this type of scenario, you might see the site actually penalized.

The only fixes for this are to reduce the number of duplicate pages accessible to the search engine crawler.

You can accomplish this by deleting them, using **canonical** tags on the duplicates, **NoIndex**ing the pages themselves, or to add a substantial amount of unique content.

One example of duplicate content that may get filtered out on a broad basis is a *thin affiliate* site. This nomenclature frequently describes a site promoting the sale of someone else's products (to earn a commission), yet provides little or no information differentiated from other sites selling the product. Such a site may have received the descriptions from the manufacturer of the products and simply replicated those descriptions along with an affiliate link (so that it can earn credit when a click/purchase is performed).

The problem arises when a merchant has thousands of affiliates generally promoting websites using the same descriptive content, and search engineers have observed user data suggesting that, from a searcher's perspective, these sites add little value to their indexes. Thus, the search engines attempt to filter out this type of site, or even ban it from their index. Plenty of sites operate affiliate models but also provide rich new content, and these sites generally have no problem - it is when duplication of content and a lack of unique, value-adding material come together on a domain that the engines may take action.

How to Avoid Duplicate Content on Your Own Site

As we outlined, duplicate content can be created in many ways. Internal duplication of material requires specific tactics to achieve the best possible results from an SEO perspective. In many cases, the duplicate pages are pages that have no value to either users or search engines. If that is the case, try to eliminate the problem altogether by fixing the implementation so that all pages are referred to by only one URL. Also, 301 redirect (these are discussed in more detail in the "Redirects" section later on in this chapter) the old URLs to the surviving URLs to help the search engines discover what you have done as rapidly as possible, and preserve any link authority the removed pages may have had.

If that process proves to be impossible, there are many options, as we will outline in "Content Delivery and Search Spider Control". Here is a summary of the guidelines on the simplest solutions for dealing with a variety of scenarios:

- Use *robots.txt* to block search engine spiders from crawling the duplicate versions of pages on your site.
- Use the **canonical** tag. This is the next best solution to eliminating the duplicate pages.
- Use the **Robots NoIndex** meta tag to tell the search engine to not index the duplicate pages.

Be aware however that if you use *robots.txt* to prevent a page from being crawled, then using **NoIndex** or **NoFollow** on the page itself does not make sense, as the spider can't read the page, so it will never see the **NoIndex** or **NoFollow** tag. With these tools in mind, here are some specific duplicate content scenarios:

HTTPS pages

If you make use of *SSL* (encrypted communications between the browser and the web server), and you have not converted your entire site, you will have some pages on your site that begin with *https:* instead of *http:* The problem arises when the links on your *https:* pages link back to other pages on the site using relative instead of absolute links, so (for example) the link to your home page becomes *https://www.yourdomain.com* instead of *http://www.yourdomain.com*.

If you have this type of issue on your site, you may want to use the **canonical** tag, which we describe in "Content Delivery and Search Spider Control", or 301 redirects to resolve problems with these types of pages. An alternative solution is to change the links to absolute links (*http://www.yourdomain.com/content* instead of *"/content"*), which also makes life more difficult for content thieves that scrape your site.

CMSs that create duplicate content

Sometimes sites have many versions of identical pages because of limitations in the CMS where it addresses the same content with more than one URL. These are often unnecessary duplications with no

end-user value, and the best practice is to figure out how to eliminate the duplicate pages and 301 the eliminated pages to the surviving pages. Failing that, fall back on the other options listed at the beginning of this section.

Print pages or multiple sort orders

Many sites offer print pages to provide the user with the same content in a more printer-friendly format. Or some e-commerce sites offer their products in multiple sort orders (such as size, color, brand, and price). These pages do have end-user value, but they do not have value to the search engine and will appear to be duplicate content. For that reason, use one of the options listed previously in this subsection, or setup a print CSS style sheet such as the one outlined in this post by Yoast (<http://yoast.com/added-print-css-style-sheet/>).

Duplicate content in blogs and multiple archiving systems (pagination, etc.)

Blogs present some interesting duplicate content challenges. Blog posts can appear on many different pages, such as the home page of the blog, the Permalink page for the post, date archive pages, and category pages. Each instance of the post represents duplicates of the other instances. Few publishers attempt to address the presence of the post on the home page of the blog and also at its permalink, and this is common enough that it is likely that the search engines deal reasonably well with it. However, it may make sense to show only excerpts of the post on the category and/or date archive pages.

User-generated duplicate content (repostings, etc.)

Many sites implement structures for obtaining user-generated content, such as a blog, forum, or job board. This can be a great way to develop large quantities of content at a very low cost. The challenge is that users may choose to submit the same content on your site and in several other sites at the same time, resulting in duplicate content among those sites. It is hard to control this, but there are two things you can do to reduce the problem:

- Have clear policies that notify users that the content they submit to your site must be unique and cannot be, or cannot have been, posted to other sites. This is difficult to enforce, of course, but it will still help some to communicate your expectations.
- Implement your forum in a different and unique way that demands different content. Instead of having only the standard fields for entering data, include fields that are likely to be unique over what other sites do, but that will still be interesting and valuable for site visitors to see.

Controlling Content with Cookies and Session IDs

Sometimes you want to more carefully dictate what a search engine robot sees when it visits your site. In general, search engine representatives will refer to the practice of showing different content to users than crawlers as *cloaking*, which violates the engines' Terms of Service (TOS) and is considered spam.

However, there are legitimate uses for this concept that are not deceptive to the search engines or malicious in intent. This section will explore methods for doing this with cookies and sessions IDs.

What's a Cookie?

A *cookie* is a small text file that websites can leave on a visitor's hard disk, helping them to track that person over time. Cookies are the reason Amazon.com remembers your username between visits and the reason you don't necessarily need to log in to your Gmail account every time you open your browser. Cookie data typically contains a short set of information regarding when you last accessed a site, an ID number, and, potentially, information about your visit (see Figure 6-32).

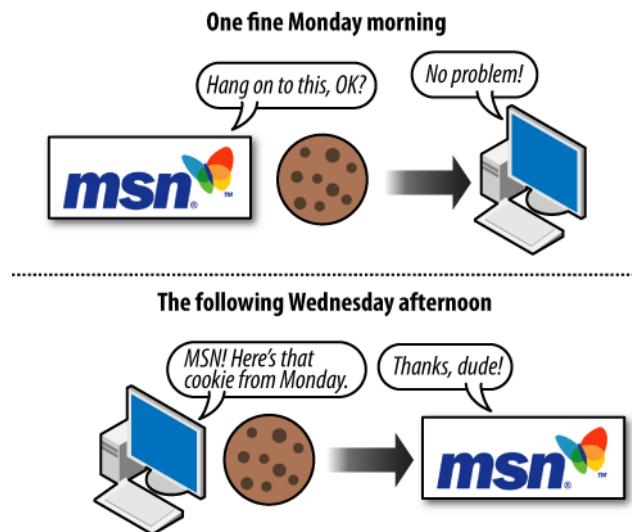


Figure 6-32. Using cookies to store data

Website developers can create options to remember visitors using cookies for tracking purposes or to display different information to users based on their actions or preferences. Common uses include remembering a username, maintaining a shopping cart, and keeping track of previously viewed content. For example, if you've signed up for an account with Moz, it will provide you with options on your My Account page about how you want to view the blog and will remember that the next time you visit.

What Are Session IDs?

Session IDs are virtually identical to cookies in functionality, with one big difference. Upon closing your browser (or restarting), session ID information is no longer stored on your hard drive (usually); see Figure 6-33. The website you were interacting with may remember your data or actions, but it cannot retrieve session IDs from your machine that don't persist (and session IDs by default expire when the browser shuts down). In essence, session IDs are more like temporary cookies (although, as you'll see shortly, there are options to control this).



Figure 6-33. How session IDs are used

Although technically speaking, session IDs are just a form of cookie without an expiration date, it is possible to set session IDs with expiration dates similar to cookies (going out decades). In this sense, they are virtually identical to cookies. Session IDs do come with an important caveat, though: they are frequently passed in the URL string, which can create serious problems for search engines (as every request produces a unique URL with duplicate content).

It is highly desirable to eliminate Session IDs from your URLs, and you should avoid them if it is at all possible. If you currently have them, a short term fix is to use the `canonical` tag (which we'll discuss in "Content Delivery and Search Spider Control") to tell the search engines that you want them to ignore the session IDs.

Any user has the ability to turn off cookies in his browser settings. This often makes web browsing considerably more difficult, and many sites will actually display a page saying that cookies are required to view or interact with their content. Cookies, persistent though they may be, are also deleted by users on a semiregular basis. For example, a 2011 comScore study (<http://www.websitemagazine.com/content/blogs/posts/archive/2011/05/09/the-impact-of-cookie-deletion.aspx>) found that 33% of web users deleted their first party cookies at least once per month.

How Do Search Engines Interpret Cookies and Session IDs?

Search engine spiders do not look at cookies or session IDs and act as browsers with this functionality shut off. However, unlike visitors whose browsers won't accept cookies, the crawlers can sometimes reach sequestered content by virtue of webmasters who want to specifically let them through. Many sites have pages that require cookies or sessions to be enabled but have special rules for search engine bots, permitting them to access the content as well. Although this is technically cloaking, there is a form of this known as First Click Free that search engines generally allow (we will discuss this in more detail in "Content Delivery and Search Spider Control").

Despite the occasional access engines are granted to cookie/session-restricted pages, the vast majority of cookie and session ID usage creates content, links, and pages that limit access. Web developers can leverage the power of concepts such as First Click Free to build more intelligent sites and pages that function in optimal ways for both humans and engines.

Why Would You Want to Use Cookies or Session IDs to Control Search Engine Access?

There are numerous potential tactics to leverage cookies and session IDs for search engine control. Here are many of the major strategies you can implement with these tools, but there are certainly limitless other possibilities:

Showing multiple navigation paths while controlling the flow of link authority

Visitors to a website often have multiple ways in which they'd like to view or access content. Your site may benefit from offering many paths to reaching content (by date, topic, tag, relationship, ratings, etc.), but expends PageRank or link authority that would be better optimized by focusing on a single, search-engine-friendly navigational structure. This is important because these varied sort orders may be seen as duplicate content.

You can require a cookie for users to access the alternative sort order versions of a page, and prevent the search engine from indexing multiple pages with the same content. One alternative (but not foolproof) solution to this is to use the canonical tag to tell the search engine that these alternative sort orders are really just the same content as the original page (we will discuss the canonical tag in "Content Delivery and Search Spider Control").

Keep limited pieces of a page's content out of the engines' indexes

Many pages may contain content that you'd like to show to search engines and pieces you'd prefer appear only for human visitors. These could include ads, login-restricted information, links, or even rich media. Once again, showing non-cookies users the plain version and cookie-accepting visitors the extended information can be invaluable. Note that this is often used in conjunction with a login, so only registered users can access the full content (such as on sites like Facebook and LinkedIn).

Grant access to pages requiring a login

As with snippets of content, there are often entire pages or sections of a site on which you'd like to restrict search engine access. This can be easy to accomplish with cookies/sessions, and it can even help to bring in search traffic that may convert to "registered-user" status. For example, if you had desirable content that you wished to restrict, you could create a page with a short snippet and an offer to continue reading upon registration, which would then allow access to that work at the same URL. We will discuss this more in "Content Delivery and Search Spider Control".

Avoid duplicate content issues

One of the most promising areas for cookie/session use is to prohibit spiders from reaching multiple versions of the same content, while allowing visitors to get the version they prefer. As an example, at Moz, logged-in users can see full blog entries on the blog home page, but search engines and non-registered users will see only the excerpts. This prevents the content from being listed on multiple pages (the blog home page and the specific post pages), and provides a richer user experience for members.

Content Delivery and Search Spider Control

On occasion, it can be valuable to show search engines one version of content and show humans a different version. This is technically called cloaking, and the search engines' guidelines have near-universal policies restricting this. In practice, many websites, large and small, appear to use content delivery effectively and without being penalized by the search engines. However, use great care if you implement these techniques, and know the risks that you are taking.

Cloaking and Segmenting Content Delivery

Before we discuss the risks and potential benefits of cloaking-based practices, take a look at Figure 6-34, which shows an illustration of how cloaking works.

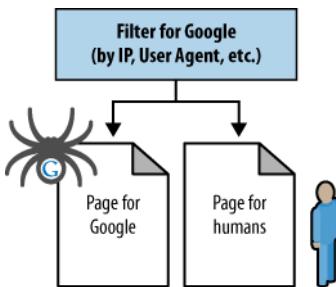


Figure 6-34. How cloaking works

Google's Matt Cutts, head of Google's webspam team, has made strong public statements indicating that all forms of cloaking (with the only exception being First Click Free) are subject to penalty. This was also largely backed by statements by Google's John Mueller in a May 2009 interview, which you can read at <http://www.stonetemple.com/articles/interview-john-mueller.shtml>. In August of 2011 Matt Cutts later confirmed this again in this video on YouTube: http://www.youtube.com/watch?feature=player_embedded&v=QHtnfOgp65Q. In this video Matt Cutts makes the strong statement "There is no such thing as white hat cloaking".

Google also makes its policy pretty clear in its Guidelines on Cloaking (<https://support.google.com/webmasters/answer/66355>):

Serving up different results based on user agent may cause your site to be perceived as deceptive and removed from the Google index.

There are two critical pieces in the preceding quote: *may* and *user agent*. It is true that if you cloak in the wrong ways, with the wrong intent, Google and the other search engines *may* remove you from their index, and if you do it egregiously, they certainly *will*.

A big factor is intent: if the engines feel you are attempting to manipulate their rankings or results through cloaking, they may take adverse action against your site. If, however, the intent of your content delivery doesn't interfere with their goals, you're less likely to be subject to a penalty, but there is never "zero" risk of a penalty. Google has taken a strong stand against all forms of cloaking regardless of intent.

What follows are some examples of websites that perform some level of cloaking:

Google

Search for *google toolbar* or *google translate* or *adwords* or any number of Google properties and note how the URL you see in the search results and the one you land on almost never match. What's more, on many of these pages, whether you're logged in or not, you might see some content that is different from what's in the cache.

NYTimes.com (<http://www.nytimes.com/>)

The interstitial ads, the request to log in/create an account after five clicks, and the archive inclusion are all showing different content to engines versus humans.

Wine.com (<http://www.wine.com/>)

In addition to some redirection based on your path, there's the state overlay forcing you to select a shipping location prior to seeing any prices (or any pages). That's a form the engines don't have to fill out.

Yelp.com (<http://yelp.com/>)

Geotargeting through cookies based on location is a very popular form of local targeting that hundreds, if not thousands, of sites use.

Trulia.com (<http://www.trulia.com/>)

Trulia was found to be doing some interesting redirects on partner pages and its own site (<http://www.bramblog.com/trulia-caught-cloaking-red-handed/>).

The message should be clear. Cloaking won't always get you banned, and you can do some pretty smart things with it. The key to all of this is your intent. If you are doing it for reasons that are not deceptive and that provide a positive experience for users and search engines, you might not run into problems. However, there is no guarantee of this, so use these types of techniques with great care, and know that you may still get penalized for it.

When to Show Different Content to Engines and Visitors

There are a few common causes for displaying content differently to different visitors, including search engines. Here are some of the most common ones:

Multivariate and A/B split testing

Testing landing pages for conversions requires that you show different content to different visitors to test performance. In these cases, it is best to display the content using JavaScript/cookies/sessions and give the search engines a single, canonical version of the page that doesn't change with every new spidering (though this won't necessarily hurt you). Google used to offer software called Google Website Optimizer to perform this function, which has been discontinued and replaced with Google Analytics Content Experiments (<https://support.google.com/analytics/answer/2661700?topic=1745146>). If you have used Google Website Optimizer in the past, Google recommends removing the associated tags from your site pages.

Content requiring registration and First Click Free

If you force registration (paid or free) on users to view specific content pieces, it is best to keep the URL the same for both logged-in and non-logged-in users and to show a snippet (one to two paragraphs is usually enough) to non-logged-in users and search engines. If you want to display the full content to search engines, you have the option to provide some rules for content delivery, such as showing the first one to two pages of content to a new visitor without requiring registration, and then requesting registration after that grace period. This keeps your intent more honest, and you can use cookies or sessions to restrict human visitors while showing the full pieces to the engines.

In this scenario, you might also opt to participate in a specific program from Google called First Click Free, wherein websites can expose "premium" or login-restricted content to Google's spiders, as long as users who click from the engine's results are given the ability to view that first article for free. Many prominent web publishers employ this tactic, including the popular site, Experts-Exchange.com (<http://www.experts-exchange.com/>).

To be specific, to implement First Click Free, the publisher must grant Googlebot (and presumably the other search engine spiders) access to all the content they want indexed, even if users normally have to log in to see the content. The user who visits the site will still need to log in, but the search engine spider will not have to do so. This will lead to the content showing up in the search engine results when applicable. However, if a user clicks on that search result, you must permit him to view the entire article (all pages of a given article if it is a multiple-page article). Once the user clicks to look at another article on your site, you can still require him to log in. Publishers can also limit the number of free accesses a user gets using this technique to five articles per day.

For more details, visit Google's First Click Free program page at <http://googlewebmastercentral.blogspot.com/2008/10/first-click-free-for-web-search.html> and <http://googlewebmastercentral.blogspot.com/2009/12/changes-in-first-click-free.html>.

Navigation unspider-able by search engines

If your navigation is in Flash, JavaScript, a Java application, or another format where the ability of the search engine to parse them is uncertain, you should consider showing search engines a version that has spider-able, crawl-ble content in HTML. Many sites do this simply with CSS layers, displaying a

human-visible, search-invisible layer and a layer for the engines (and less capable browsers, such as mobile browsers). You can also employ the `noscript` tag for this purpose, although it is generally riskier, as many spammers have applied `noscript` as a way to hide content. Take care to make sure the content shown in the search-visible layer is substantially the same as it is in the human-visible layer.

Duplicate content

If a significant portion of a page's content is duplicated, you might consider restricting spider access to it by placing it in an iframe that's restricted by `robots.txt`. This ensures that you can show the engines the unique portion of your pages, while protecting against duplicate content problems. We will discuss this in more detail in the next section.

Different content for different users

At times you might target content uniquely to users from different geographies (such as different product offerings that are more popular in their area), with different screen resolutions (to make the content fit their screen size better), or who entered your site from different navigation points. In these instances, it is best to have a "default" version of content that's shown to users who don't exhibit these traits to show to search engines as well.

How to Display Different Content to Search Engines Versus Visitors

A variety of strategies exist to segment content delivery. The most basic is to serve content that is not meant for the engines in un-spiderable formats (e.g., placing text in images, Flash files, plug-ins, etc.). You should not use these formats for the purpose of cloaking. You should use them only if they bring a substantial end-user benefit (such as an improved user experience). In such cases, you may want to show the search engines the same content in a search-spider-readable format. When you're trying to show the engines something you don't want visitors to see, you can use CSS formatting styles (preferably not `display:none`, as the engines have filters to watch specifically for this), JavaScript, user-agent, cookie, or session-based delivery, or IP delivery (showing content based on the visitor's IP address).

Be very wary when employing cloaking such as that we just described. The search engines expressly prohibit these practices in their guidelines, and though there may be some leeway based on intent and user experience (e.g., your site is using cloaking to improve the quality of the user's experience, not to game the search engines), the engines do take these tactics seriously and may penalize or ban sites that implement them inappropriately or with the intention of manipulation. In addition, even if your intent is good, the search engines may not see it that way and penalize you anyway.

The `robots.txt` file

This file is located on the root level of your domain (e.g., `http://www.yourdomain.com/robots.txt`), and it is a highly versatile tool for controlling what the spiders are permitted to access on your site. You can use `robots.txt` to:

- Prevent crawlers from accessing nonpublic parts of your website
- Block search engines from accessing index scripts, utilities, or other types of code
- Avoid the indexation of duplicate content on a website, such as "print" versions of HTML pages, or various sort orders for product catalogs
- Auto-discover XML Sitemaps

The `robots.txt` file must reside in the root directory, and the filename must be entirely in lowercase (`robots.txt`, not `Robots.txt`, or other variations including uppercase letters). Any other name or location will not be seen as valid by the search engines. The file must also be entirely in text format (not in HTML format).

When you tell a search engine robot not to access a page, it prevents the crawler from accessing the page. Figure 6-35 illustrates what happens when the search engine robot sees a direction in `robots.txt` not to crawl a web page.

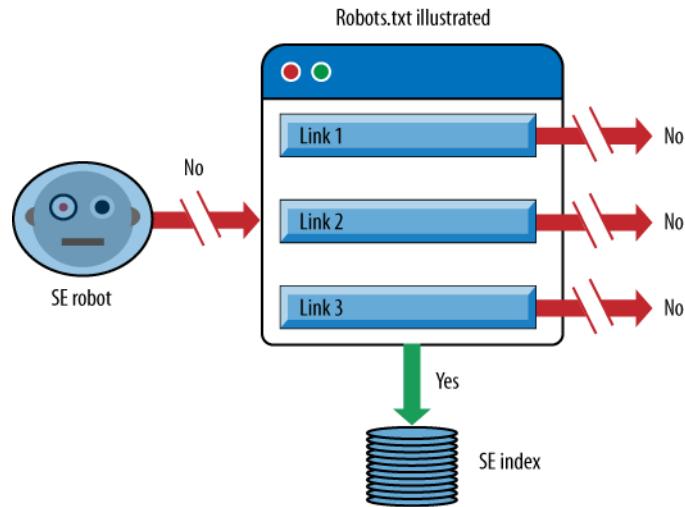


Figure 6-26. Impact of robots.txt

In essence, the page will not be crawled, so links on the page cannot pass link authority to other pages since the search engine does not see the links. However, the page can be in the search engine index. This can happen if other pages on the Web link to the page. Of course, the search engine will not have very much information on the page since it cannot read it, and will rely mainly on the anchor text and other signals from the pages linking to it to determine what the page may be about. Any resulting search listings end up being pretty sparse when you see them in the Google index, as shown in Figure 6-36.

site:www.nytimes.com/cnet/

Web Images News Shopping Maps More Search tools

About 816 results (0.18 seconds)

animation that stops you from being able to click on - The ...
www.nytimes.com/cnet/ ▾ The New York Times ▾
 A description for this result is not available because of this site's robots.txt – learn more.

www.nytimes.com/cnet/CNET
 A description for this result is not available because of this site's robots.txt – learn more.

NY Times link to CNET article - The New York Times
www.nytimes.com/cnet/CNET_0_4_2531104_00.html ▾ The New York Times ▾
 A description for this result is not available because of this site's robots.txt – learn more.

Home networking gets a boost from new standards - The ...
www.nytimes.com/cnet/CNET_0_4_39752_00.html ▾ The New York Times ▾
 A description for this result is not available because of this site's robots.txt – learn more.

The New York Times
www.nytimes.com/cnet/CNET_2100-1026_3- ▾ The New York Times ▾
 A description for this result is not available because of this site's robots.txt – learn more.

www.nytimes.com/cnet/CNET_2100-104
 A description for this result is not available because of this site's robots.txt – learn more.

The New York Times
www.nytimes.com/cnet/CNET_2100-7343_3-52297... ▾ The New York Times ▾
 A description for this result is not available because of this site's robots.txt – learn more.

Figure 6-27. SERPs for pages that are listed in robots.txt

Figure 6-32 shows the results for the Google query `site:www.nytimes.com/cnet/`. This is not a normal query that a user would enter, but you can see what the results look like. Only the URL is listed, and there is no description. This is because the spiders aren't permitted to read the page to get that data. In today's

algorithms, these types of pages don't rank very high because their relevance scores tend to be quite low for any normal queries.

Google, Bing, and nearly all of the legitimate crawlers on the Web will follow the instructions you set out in the *robots.txt* file. Commands in *robots.txt* are primarily used to prevent spiders from accessing pages and subfolders on a site, though they have other options as well. Note that subdomains require their own *robots.txt* files, as do files that reside on an https: server.

Syntax of the robots.txt file

The basic syntax of *robots.txt* is fairly simple. You specify a robot name, such as "googlebot", and then you specify an action. The robot is identified by user agent, and then the actions are specified on the lines that follow. Here are the major actions you can specify:

- **Disallow:** the pages you want to block the bots from accessing (as many disallow lines as needed)

Some other restrictions apply:

- Each User-Agent/Disallow group should be separated by a blank line; however, no blank lines should exist within a group (between the User-Agent line and the last Disallow).
- The hash symbol (#) may be used for comments within a *robots.txt* file, where everything after # on that line will be ignored. This may be used either for whole lines or for the end of lines.
- Directories and filenames are case-sensitive: "private", "Private", and "PRIVATE" are all uniquely different to search engines.

Here is an example of a *robots.txt* file:

```
User-agent: Googlebot
Disallow:

User-agent: BingBot
Disallow: /

# Block all robots from tmp and logs directories
User-agent: *
Disallow: /tmp/
Disallow: /logs      # for directories and files called logs
```

The preceding example will do the following:

- Allow "Googlebot" to go anywhere.
- Prevent "BingBot" from crawling any part of the site.
- Block all robots (other than Googlebot) from visiting the */tmp/* directory or directories or files called */logs* (e.g., */logs* or *logs.php*).

Notice that the behavior of Googlebot is not affected by instructions such as *Disallow: /*. Since Googlebot has its own instructions from *robots.txt*, it will ignore directives labeled as being for all robots (i.e., uses an asterisk).

One common problem that novice webmasters run into occurs when they have SSL installed so that their pages may be served via HTTP and HTTPS. A *robots.txt* file at *http://www.yourdomain.com/robots.txt* will not be interpreted by search engines as guiding their crawl behavior on *https://www.yourdomain.com*. To do this, you need to create an additional *robots.txt* file at *https://www.yourdomain.com/robots.txt*. So, if you want to allow crawling of all pages served from your HTTP server and prevent crawling of all pages from your HTTPS server, you would need to implement the following:

For HTTP:

```
User-agent: *
Disallow:
```

For HTTPS:

```
User-agent: *
Disallow: /
```

These are the most basic aspects of *robots.txt* files, but there are more advanced techniques as well. Some of these methods are supported by only some of the engines, as detailed in the list that follows:

Crawl delay

Crawl delay is supported by Google, Bing, and Ask. It instructs a crawler to wait the specified number of seconds between crawling pages. The goal with the directive is to reduce the load on the publisher's server:

```
User-agent: BingBot
Crawl-delay: 5
```

Pattern matching

Pattern matching appears to be usable by Google, and Bing. The value of pattern matching is considerable. You can do some basic pattern matching using the asterisk wildcard character. Here is how you can use pattern matching to block access to all subdirectories that begin with *private* (e.g., */private1/*, */private2/*, */private3/*, etc.):

```
User-agent: Googlebot
Disallow: /private*/
```

You can match the end of the string using the dollar sign (\$). For example, to block URLs that end with *.asp*:

```
User-agent: Googlebot
Disallow: /*.asp$
```

You may wish to prevent the robots from accessing any URLs that contain parameters in them. To block access to all URLs that include a question mark (?), simply use the question mark:

```
User-agent: *
Disallow: /*?*
```

The pattern-matching capabilities of *robots.txt* are more limited than those of programming languages such as Perl, so the question mark does not have any special meaning and can be treated like any other character.

Allow directive

The Allow directive appears to be supported only by Google and Ask. It works the opposite of the Disallow directive and provides the ability to specifically call out directories or pages that may be crawled. When this is implemented it can partially override a previous Disallow directive. This may be beneficial after large sections of the site have been disallowed, or if the entire site itself has been disallowed.

Here is an example that allows Googlebot into only the *google* directory:

```
User-agent: Googlebot
Disallow: /
Allow: /google/
```

Noindex directive

This directive works in the same way as the *meta robots noindex* command (which we will discuss shortly) and tells the search engines to explicitly exclude a page from the index. Since a Disallow directive prevents crawling but not indexing, this can be a very useful feature to ensure that the pages don't show in search results. Google supports this directive in *robots.txt*, and only unofficially

Sitemaps

We discussed XML Sitemaps at the beginning of this chapter. You can use *robots.txt* to provide an autodiscovery mechanism for the spider to find the XML Sitemap file. The search engines can be told to find the file with one simple line in the *robots.txt* file:

```
| Sitemap: sitemap_location
```

The *sitemap_location* should be the complete URL to the Sitemap, such as <http://www.yourdomain.com/sitemap.xml>. You can place this anywhere in your file.

For full instructions on how to apply *robots.txt*, see Robotstxt.org (<http://www.robotstxt.org/orig.html>). You can also test your *robots.txt* file in Google webmaster tools under: Crawl, *robots.txt* Tester.

You should use great care when making changes to *robots.txt*. A simple typing error can, for example, suddenly tell the search engines to no longer crawl any part of your site. After updating your *robots.txt* file it is always a good idea to check it with the Google Webmaster Tools Test Robots.txt tool. You can find this by logging in to Webmaster Tools and then selecting "Crawl" and followed by "Blocked URLs".

The Rel="NoFollow" attribute

In 2005, the three major search engines (Google, Microsoft, and Yahoo!, which still had its own search engine at that time) all agreed to support an initiative intended to reduce the effectiveness of automated spam. Unlike the meta *robots* version of *NoFollow*, the new directive could be employed as an attribute within an *<a>* or link tag to indicate that the linking site "does not editorially vouch for the quality of the linked-to page." This enables a content creator to link to a web page without passing on any of the normal search engine benefits that typically accompany a link (things such as trust, anchor text, PageRank, etc.).

Originally, the intent was to enable blogs, forums, and other sites where user-generated links were offered to shut down the value of spammers who built crawlers that automatically created links. However, this has expanded as Google, in particular, recommends use of *NoFollow* on links that are paid for—as the search engine's preference is that only those links that are truly editorial and freely provided by publishers (without being compensated) should count toward bolstering a site's/page's rankings.

You can implement *NoFollow* using the following format:

```
| <a href="http://www.google.com" rel="nofollow">
```

Note that although you can use *NoFollow* to restrict the passing of link value between web pages, the search engines still crawl through those links (despite the lack of semantic logic) and crawl the pages they link to. The search engines have provided contradictory input on this point. To summarize, *NoFollow* does not expressly forbid indexing or spidering, so if you link to your own pages with it, intending to keep those pages from being indexed or ranked, others may find them and link to them, and your original goal will be thwarted.

Figure 6-37 shows how a search engine robot interprets a *NoFollow* attribute when it finds one associated with a link (Link 1 in this example).

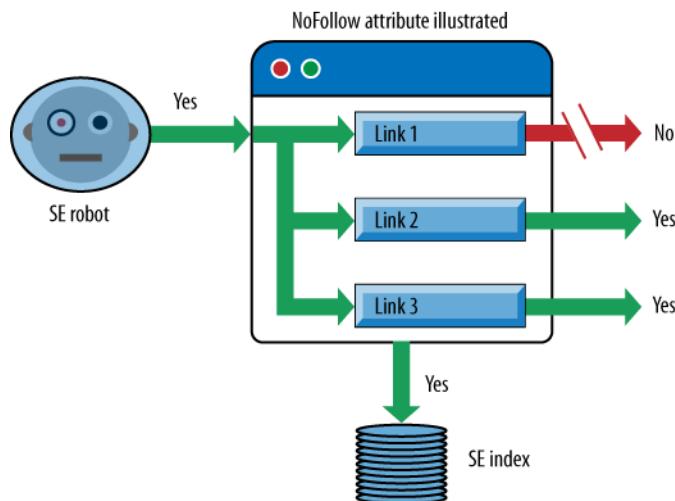


Figure 6-28. Impact of NoFollow attribute

The specific link with the `NoFollow` attribute was, for a number of years, considered to be disabled from passing link authority, and the notion of PageRank sculpting using `NoFollow` was a popular idea. The belief was that when you `NoFollow` a particular link, the link authority that would have been passed to that link was preserved and the search engines would reallocate it to the other links found on the page. As a result, many publishers implemented `NoFollow` links to lower value pages on their site (such as the About Us and Contact Us pages, or alternative sort order pages for product catalogs). In fact, data from Moz's Open Site Explorer tool (<http://www.opensiteexplorer.org/>), published in March 2009, showed that at that time about 3% of all links on the Web were `NoFollowed`, and that 60% of those `NoFollows` were applied to internal links.

In June 2009, however, Google's Matt Cutts wrote a post that made it clear that the link authority associated with that `NoFollowed` link is discarded rather than reallocated (<http://www.mattcutts.com/blog/pagerank-sculpting/>). In theory, you can still use `NoFollow` however you want, but using it on internal links does not (at the time of this writing, according to Google) bring the type of benefit people have been looking for in the past. In fact, in certain scenarios it can be harmful.

In addition, many SEOs speculate that in some cases some value is indeed being placed on some `NoFollowed` links, and we suggest erring on the side of caution when using the `NoFollow` attribute – as their use has been attributed to “flagging” a site as over-optimized, or otherwise aggressive in SEO tactics.

This is a great illustration of the ever-changing nature of SEO. Something that was a popular, effective tactic is now being viewed as ineffective. Some more aggressive publishers will continue to pursue PageRank sculpting by using even more aggressive approaches, such as implementing links in encoded JavaScript or within iframes that have been disallowed in `robots.txt`, so that the search engines don't see them as links. Such aggressive tactics are probably not worth the trouble for most publishers.

The meta robots tag

The `meta robots` tag has three components: `cache`, `index`, and `follow`. The `cache` component instructs the engine about whether it can keep the page in the engine's public index, available via the “cached snapshot” link in the search results (see Figure 6-38).

Figure 6-29. Snapshot of how to access a cached page in the SERPs

The second, `index`, tells the engine whether the page is allowed to be crawled and stored in any capacity. A page marked `NoIndex` will thus be excluded entirely by the search engines. By default, this value is `index`, telling the search engines, “Yes, please do crawl this page and include it in your index.” Thus, it is unnecessary to place this directive on each page. Figure 6-39 shows what a search engine robot does once it sees a `NoIndex` tag on a web page.

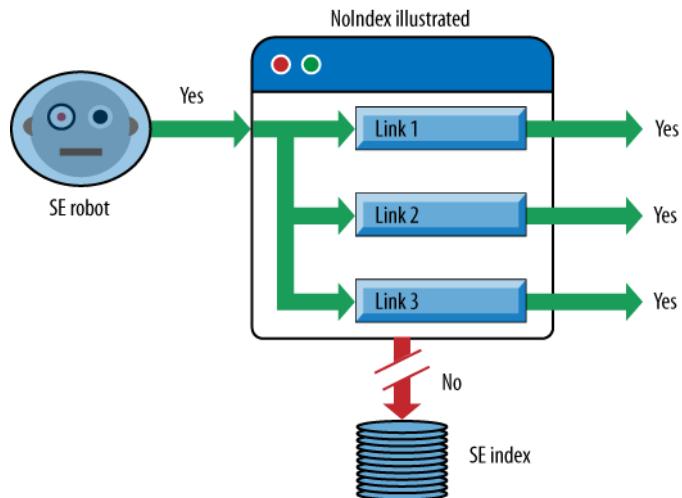


Figure 6-30. Impact of `NoIndex`

The page will still be crawled, and the page can still accumulate and pass link authority to other pages, but it will not appear in search indexes.

The final instruction available through the meta `robots` tag is `follow`. This command, like `index`, defaults to “yes, crawl the links on this page and pass link authority through them.” Applying `NoFollow` tells the engine that none of the links on that page should pass link value. By and large, it is unwise to use

this directive as a way to prevent links from being crawled. Since human beings will still reach those pages and have the ability to link to them from other sites, `NoFollow` (in the meta `robots` tag) does little to restrict crawling or spider access. Its only application is to prevent link authority from spreading out, which has very limited application since the 2005 launch of the `rel="nofollow"` attribute (discussed earlier), which allows this directive to be placed on individual links.

Figure 6-40 outlines the behavior of a search engine robot when it finds a `NoFollow` meta tag on a web page. (assuming there are no other links pointing to the three linked URLs)

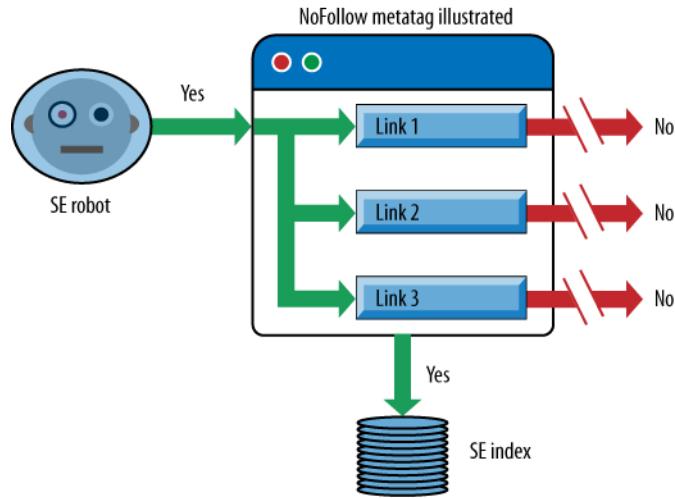


Figure 6-40. Impact of `NoFollow` meta tag

When you use the `NoFollow` meta tag on a page, the search engine will still crawl the page and place the page in its index. However, all links (both internal and external) on the page will be disabled from passing link authority to other pages.

One good application for `NoIndex` is to place this tag on HTML sitemap pages. These are pages designed as navigational aids for users and search engine spiders to enable them to efficiently find the content on your site. However, on some sites these pages are unlikely to rank for anything of importance in the search engines, yet you still want them to pass link authority to the pages they link to. Putting `NoIndex` on these pages keeps these HTML sitemaps out of the index and removes that problem. Make sure you *do not* apply the `NoFollow` meta tag on the pages or the `NoFollow` attribute on the links on the pages, as these will prevent the pages from passing link authority.

The canonical tag

In February 2009, Google, Yahoo!, and Microsoft debuted the `canonical` tag (sometimes referred to as `rel=canonical`). This tag was a new construct designed explicitly for purposes of identifying and dealing with duplicate content. Implementation is very simple and looks like this:

```
| <link rel="canonical" href="http://moz.com/blog" />
```

This tag is meant to tell the search engines that the page in question should be treated as though it were a copy of the URL `http://moz.org/blog` and that all of the link and content metrics the engines apply should technically flow back to that URL (see Figure 6-41).

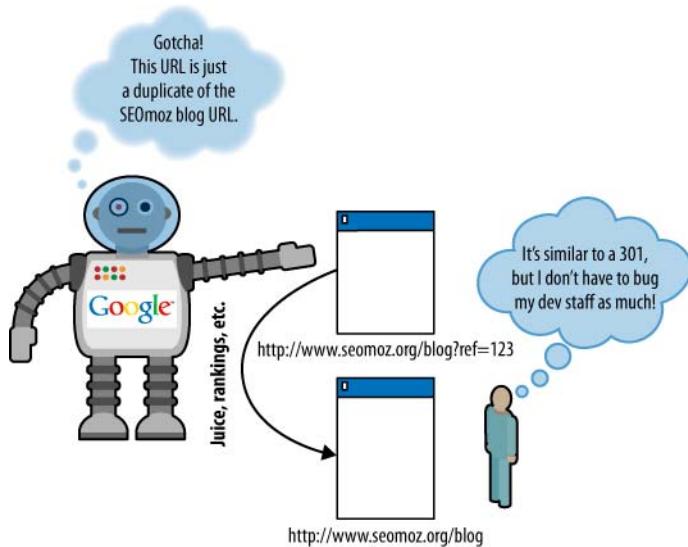


Figure 6-41. How search engines look at the canonical tag

The canonical link attribute is similar in many ways to a 301 redirect from an SEO perspective. In essence, you're telling the engines that multiple pages should be considered as one (which a 301 does), without actually redirecting visitors to the new URL (for many publishers this is less effort than some of the other solutions for their development staff). There are some differences, though:

- Whereas a 301 redirect points all traffic (bots and human visitors), the canonical URL tag is just for engines, meaning you can still separately track visitors to the unique URL versions.
- A 301 is a much stronger signal that multiple pages have a single, canonical source. 301s are considered a directive that search engines and browsers are obligated to honor, but the canonical tag is treated as a suggestion. Although the engines generally support this new tag and trust the intent of site owners, there will be limitations. Content analysis and other algorithmic metrics will be applied to ensure that a site owner hasn't mistakenly or manipulatively applied the tag, and you can certainly expect to see mistaken use of the canonical tag, resulting in the engines maintaining those separate URLs in their indexes (meaning site owners would experience the same problems noted in "Duplicate Content Issues").

We will discuss some applications for this tag later in this chapter. In general practice, the best solution is to resolve the duplicate content problems at their core, and eliminate them if you can. This is because the canonical tag is not guaranteed to work. However, it is not always possible to resolve the issues by other means, and the canonical tag provides a very effective backup plan.

You can also include the canonical tag directly within the HTTP response header for your page. The code might look something like the following:

```
HTTP/1.1 200 OK
Content-Type: application/pdf
Link: <http://www.example.com/white-paper.html>; rel="canonical"
Content-Length: 785710
(... rest of HTTP response headers...)
```

You can read more about this here: <http://googlewebmastercentral.blogspot.com/2011/06/supporting-relcanonical-http-headers.html>.

Blocking and cloaking by IP address range

You can customize entire IP addresses or ranges to block particular bots through server-side restrictions on IPs. Most of the major engines crawl from a limited number of IP ranges, making it possible to identify them and restrict access. This technique is, ironically, popular with webmasters who mistakenly assume that search engine spiders are spammers attempting to steal their content, and thus block the IP ranges to

restrict access and save bandwidth. Use caution when blocking bots, and make sure you're not restricting access to a spider that could bring benefits, either from search traffic or from link attribution.

Blocking and cloaking by user agent

At the server level, it is possible to detect user agents and restrict their access to pages or websites based on their declaration of identity. As an example, if a website detected a rogue bot, you might double-check its identity before allowing access. The search engines all use a similar protocol to verify their user agents via the Web: a reverse DNS lookup followed by a corresponding forward DNS IP lookup. An example for Google would look like this:

```
> host 66.249.66.1
1.66.249.66.in-addr.arpa domain name pointer crawl-66-249-66-1.googlebot.com.

> host crawl-66-249-66-1.googlebot.com
crawl-66-249-66-1.googlebot.com has address 66.249.66.1
```

A reverse DNS lookup by itself may be insufficient, because a spoofer could set up reverse DNS to point to xyz.googlebot.com or any other address.

Using iframes

Sometimes there's a certain piece of content on a web page (or a persistent piece of content throughout a site) that you'd prefer search engines didn't see. As we discussed earlier in this chapter, clever use of iframes can come in handy, as Figure 6-42 illustrates.

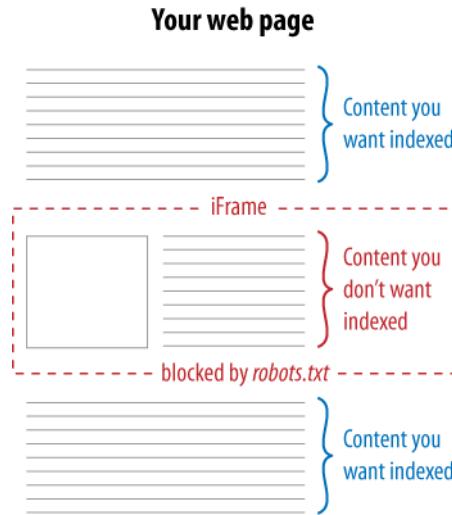


Figure 6-42. Using iframes to prevent indexing of content

The concept is simple: by using iframes, you can embed content from another URL onto any page of your choosing. By then blocking spider access to the iframe with *robots.txt*, you ensure that the search engines won't "see" this content on your page. Websites may do this for many reasons, including avoiding duplicate content problems, reducing the page size for search engines, or lowering the number of crawlable links on a page (to help control the flow of link authority).

Hiding text in images

As we discussed previously, the major search engines still have very little capacity to read text in images (and the processing power required makes for a severe barrier). Hiding content inside images isn't generally advisable, as it can be impractical for alternative devices (mobile, in particular) and inaccessible to others (such as screen readers).

Hiding text in Java applets

As with text in images, the content inside Java applets is not easily parsed by the search engines, though using them as a tool to hide text would certainly be a strange choice.

Forcing form submission

Search engines will not submit HTML forms in an attempt to access the information retrieved from a search or submission. Thus, if you keep content behind a forced-form submission and never link to it externally, your content will remain out of the engines (as Figure 6-43 demonstrates).



Figure 6-43. Use of forms, which are generally not navigable by crawlers

The problem comes when content behind forms earns links outside your control, as when bloggers, journalists, or researchers decide to link to the pages in your archives without your knowledge. Thus, although form submission may keep the engines at bay, make sure that anything truly sensitive has additional protection (e.g., through *robots.txt* or meta robots).

Using login/password protection

Password protection and/or paywalls of any kind will effectively prevent any search engines from accessing content, as will any form of human-verification requirements, such as CAPTCHAs (the boxes that request the copying of letter/number combinations to gain access). The major engines won't try to guess passwords or bypass these systems.

Removing URLs from a search engine's index

A secondary, post-indexing tactic, URL removal is possible at most of the major search engines through verification of your site and the use of the engines' tools. For example, Google allows you to remove URLs (<https://www.google.com/webmasters/tools/removals>) through Webmaster Tools. Bing allows you to remove URLs from their index via Bing Webmaster Tools as well (<https://www.bing.com/webmaster/tools/content-removal>).

Redirects

A redirect is used to indicate when content has moved from one location to another. For example, you may have some content at <http://www.yourdomain.com/old> and decide to restructure your site. As a result of this move, your content may move to <http://www.yourdomain.com/critical-keyword>

Once a redirect is implemented users who go to the old versions of your pages (perhaps via a bookmark they kept for the page) will be sent to the new versions of those pages. Without the redirect, the user would get a Page Not Found (404) error. With the redirect, the web server tells the incoming user agent (whether a browser or a spider) to instead fetch the requested content from the new URL.

Why and When to Redirect

Redirects are also important for letting search engines know when you have moved content. After you move content, the search engines will continue to have the old URL in their index and return it in their search results until they discover the page is no longer there and discover the new page. You can help speed this process up by implementing a redirect. Here are some scenarios in which you may end up needing to implement redirects:

- You have old content that expires, so you remove it.
- You find that you have broken URLs that have links and traffic.
- You change your hosting company.
- You change your CMS.
- You want to implement a canonical redirect (redirect all pages on <http://yourdomain.com> to <http://www.yourdomain.com>).
- You change the URLs where your existing content can be found for any reason.

Not all of these scenarios require a redirect. For example, you can change hosting companies without impacting any of the URLs used to find content on your site, in which case no redirect is required. However, any scenario in which any of your URLs change is a scenario in which you need to implement redirects.

Good and Bad Redirects

It turns out that there are many ways to perform a redirect. Not all are created equal. The basic reason for this is that there are two major types of redirects that can be implemented, tied specifically to the HTTP status code returned by the web server to the browser. These are:

“301 moved permanently”

This status code tells the browser (or search engine crawler) that the resource has been permanently moved to another location, and there is no intent to ever bring it back.

“302 moved temporarily”

This status code tells the browser (or search engine crawler) that the resource has been temporarily moved to another location, and that the move should not be treated as permanent.

Both forms of redirect send a human or a search engine crawler to the new location, but the search engines interpret these two HTTP status codes in very different ways. When a crawler sees a 301 HTTP status code, it assumes it should pass the historical link authority (and any other metrics) from the old page to the new one. When a search engine crawler sees a 302 HTTP status code, it assumes it should not pass the historical link authority from the old page to the new one. In addition, the 301 redirect will lead the search engine to remove the old page from the index and replace it with the new one.

The preservation of historical link authority is very critical in the world of SEO. For example, imagine you had 1,000 links to <http://www.yourolddomain.com> and you decided to relocate everything to <http://www.yournewdomain.com>. If you used redirects that returned a 302 status code, you would be starting your link-building efforts from scratch again. In addition, the old version of the page may remain in the index and compete for search rankings in the search engines.

It should also be noted that there can be redirects that pass no status code, or the wrong status code, such as a 404 error (i.e., page not found) or a 200 OK (page loaded successfully). These are also problematic, and should be avoided. There are other types of redirects as well, such as those that return a 303 or 307 status code. These are also best avoided as the search engine behavior upon seeing these is at best unpredictable.

You want to definitively return a 301 HTTP status code when you have performed a redirect whenever you make a permanent move to a page’s location.

Methods for URL Redirecting and Rewriting

There are many possible ways to implement redirects. On Apache web servers (normally present on machines running Unix or Linux as the operating system), it is possible to implement redirects quite simply in a standard file called *.htaccess* using the `Redirect` and `RedirectMatch` directives (learn more about this file format at <http://httpd.apache.org/docs/2.2/howto/htaccess.html>). More advanced directives known as *rewrite rules* can be employed as well using the Apache module known as `mod_rewrite`, which we will discuss in a moment.

On web servers running Microsoft IIS (<http://www.iis.net/>), different methods are provided for implementing redirects. The basic method for doing these redirects is through the IIS console (you can read more about this at <http://www.mcanerin.com/EN/articles/301-redirect-IIS.asp>). People with IIS servers can also make use of a text file with directives provided they use an ISAPI plug-in such as `ISAPI_Rewrite` (<http://www.isapirewrite.com/>), and this scripting language offers similar capabilities as that of Apache's `mod_rewrite` module.

Many programmers use other techniques for implementing redirects. This can be done directly in programming languages such as Perl, PHP, ASP, and JavaScript. The key thing that the programmer must do, if he implements redirects in this fashion, is to make sure the HTTP status code returned by the web server is a 301. You can check the header returned with the Firefox plug-in Live HTTP Headers (<http://livehttpheaders.mozdev.org/>), with a Chrome extension (<http://www.ayima.com/seo-knowledge/redirect-checker.html>), or a web-based server header like (<http://tools.seobook.com/server-header-checker/>).

Another method that you can use to implement a redirect occurs at the page level, via the `meta refresh` tag, which looks something like this:

```
| <meta http-equiv="refresh" content="5;url=http://www.yourdomain.com/newlocation.htm" />
```

The first parameter in the content section in the preceding statement, the number 5, indicates the number of seconds the web server should wait before redirecting the user to the indicated page. This gets used in scenarios where the publisher wants to display a page letting the user know that he is going to get redirected to a different page than the one he requested.

The problem is that most meta refreshes are treated as though they are a 302 redirect. The sole exception to this is if you specify a redirect delay of 0 seconds. You will have to give up your helpful page telling the user that you are moving him, but the search engines appear to treat this as though it were a 301 redirect (to be safe, the best practice is simply to use a 301 redirect if at all possible).

Mod_rewrite and ISAPI_Rewrite for URL rewriting and redirecting

There is much more to discuss on this topic than we can reasonably address in this book. The following description is intended only as an introduction to help orient more technical readers, including web developers and site webmasters, on how rewrites and redirects function. To skip this technical discussion, proceed to "Redirecting a Home Page Index File Without Looping".

Mod_rewrite for Apache and ISAPI_Rewrite for Microsoft IIS Server offer very powerful ways to rewrite your URLs. Here are some reasons for using these powerful tools:

- You have changed your URL structure on your site so that content has moved from one location to another. This can happen when you change your CMS, or change your site organization for any reason.
- You want to map your search-engine-unfriendly URLs into friendlier ones.

If you are running Apache as your web server, you would place directives known as *rewrite rules* within your *.htaccess* file or your Apache configuration file (e.g., *httpd.conf* or the site-specific config file in the *sites_conf* directory). Similarly, if you are running IIS Server, you'd use an ISAPI plug-in such as `ISAPI_Rewrite` and place rules in an *httpd.ini* config file.

Note that rules can differ slightly on `ISAPI_Rewrite` compared to `mod_rewrite`, and the following discussion focuses on `mod_rewrite`. Your *.htaccess* file would start with:

```
| RewriteEngine on  
| RewriteBase /
```

You should omit the second line if you're adding the rewrites to your server config file, since `RewriteBase` is supported only in `.htaccess`. We're using `RewriteBase` here so that you won't have to have `^/` at the beginning of all the rules, just `^` (we will discuss regular expressions in a moment).

After this step, the rewrite rules are implemented. Perhaps you want to have requests for product page URLs of the format `http://www.yourdomain.com/products/123` to display the content found at `http://www.yourdomain.com/get_product.php?id=123`, without the URL changing in the Location bar of the user's browser and without you having to recode the `get_product.php` script. Of course, this doesn't replace all occurrences of dynamic URLs within the links contained on all the site pages; that's a separate issue. You can accomplish this first part with a single rewrite rule, like so:

```
| RewriteRule ^products/([0-9]+)/?$ /get_product.php?id=$1 [L]
```

The preceding example tells the web server that all requests that come into the `/product/` directory should be mapped into requests to `/get_product.php`, while using the subfolder to `/product/` as a parameter for the PHP script.

The `^` signifies the start of the URL following the domain, `$` signifies the end of the URL, `[0-9]` signifies a numerical digit, and the `+` immediately following it means one or more occurrences of a digit. Similarly, the `?` immediately following the `/` means zero or one occurrence of a slash character. The `()` puts whatever is wrapped within it into memory. You can then access what's been stored in memory with `$1` (i.e., what is in the first set of parentheses). Not surprisingly, if you included a second set of parentheses in the rule, you'd access that with `$2`, and so on. The `[L]` flag saves on server processing by telling the rewrite engine to stop if it matches on that rule. Otherwise, all the remaining rules will be run as well.

Here's a slightly more complex example, where URLs of the format `http://www.yourdomain.com/webapp/wcs/stores/servlet/ProductDisplay?storeId=10001&catalogId=10001&langId=-1&categoryID=4&productID=123` would be rewritten to `http://www.yourdomain.com/4/123.htm`:

```
| RewriteRule ^([^\/]+)/([^\/]+)\.htm$  
| /webapp/wcs/stores/servlet/ProductDisplay?storeId=10001&catalogId=10001&  
| langId=-1&categoryID=$1&productID=$2 [QSA,L]
```

The `[^\/]` signifies any character other than a slash. That's because, within square brackets, `^` is interpreted as *not*. The `[QSA]` flag is for when you don't want the query string dropped (like when you want a tracking parameter preserved).

To write good rewrite rules you will need to become a master of *pattern matching* (which is simply another way to describe the use of regular expressions). Here are some of the most important special characters and how the rewrite engine interprets them:

`*` means 0 or more of the immediately preceding character.

`+` means 1 or more of the immediately preceding character.

`?` means 0 or 1 occurrence of the immediately preceding character.

`^` means the beginning of the string.

`$` means the end of the string.

`.` means any character (i.e., it acts as a wildcard).

`\` "escapes" the character that follows; for example, `\.` means the dot is not meant to be a wildcard, but an actual character.

`^` inside `[/]` brackets means *not*; for example, `/^/` means *not slash*.

It is incredibly easy to make errors in regular expressions. Some of the common gotchas that lead to unintentional substring matches include:

- Using `.*` when you should be using `.+` since `.*` can match on nothing.
- Not “escaping” with a backslash a special character that you don’t want interpreted, as when you specify `.` instead of `\.` and you really meant the dot character rather than any character (thus, `default.htm` would match on `defaultthtm`, and `default\htm` would match only on `default.htm`).
- Omitting `^` or `$` on the assumption that the start or end is implied (thus, `default\htm` would match on `mydefault.html` whereas `^default\htm$` would match only on `default.htm`).
- Using “greedy” expressions that will match on all occurrences rather than stopping at the first occurrence.

The easiest way to illustrate what we mean by greedy is to provide an example:

```
RewriteRule ^(.*)/?index\.html$ /$1/ [L,R=301]
```

This will redirect requests for `http://www.yourdomain.com/blah/index.html` to `http://www.yourdomain.com/blah//`. This is probably not what was intended. Why did this happen? Because `.*` will capture the slash character within it before the `/?` gets to see it. Thankfully, there’s an easy fix. Simply use `^/` or `.*?` instead of `.*` to do your matching. For example, use `^(.*?)/?` instead of `^(.*)/?` or `[/^]/+[/^]/` instead of `.*?/`.

So, to correct the preceding rule you could use the following:

```
RewriteRule ^(.*)/?index\.html$ /$1/ [L,R=301]
```

When wouldn’t you use the following?

```
RewriteRule ^([^\*]*/)?index\.html$ /$1/ [L,R=301]
```

This is more limited because it will match only on URLs with one directory. URLs containing multiple subdirectories, such as `http://www.yourdomain.com/store/cheese/swiss/wheel/index.html`, would not match.

As you might imagine, testing/debugging is a big part of URL rewriting. When debugging, the `RewriteLog` and `RewriteLogLevel` directives are your friends! Set the `RewriteLogLevel` to 4 or more to start seeing what the rewrite engine is up to when it interprets your rules.

By the way, the `[R=301]` flag in the last few examples—as you might guess—tells the rewrite engine to do a 301 redirect instead of a standard rewrite.

There’s another handy directive to use in conjunction with `RewriteRule`, called `RewriteCond`. You would use `RewriteCond` if you are trying to match on something in the query string, the domain name, or other things not present between the domain name and the question mark in the URL (which is what `RewriteRule` looks at).

Note that neither `RewriteRule` nor `RewriteCond` can access what is in the anchor part of a URL—that is, whatever follows a `#`—because that is used internally by the browser and is not sent to the server as part of the request. The following `RewriteCond` example looks for a positive match on the hostname before it will allow the rewrite rule that follows to be executed:

```
RewriteCond %{HTTP_HOST} !^www\.yourdomain\.com$ [NC]
RewriteRule ^(.*)$ http://www.yourdomain.com/$1 [L,R=301]
```

Note the exclamation point at the beginning of the regular expression. The rewrite engine interprets that as *not*.

For any hostname other than `http://www.yourdomain.com`, a 301 redirect is issued to the equivalent canonical URL on the `www` subdomain. The `[NC]` flag makes the rewrite condition case-insensitive. Where is the `[QSA]` flag so that the query string is preserved, you might ask? It is not needed when redirecting; it is implied.

If you don’t want a query string retained on a rewrite rule with a redirect, put a question mark at the end of the destination URL in the rule, like so:

```
RewriteCond %{HTTP_HOST} !^www\.yourdomain\.com$ [NC]
RewriteRule ^(.*)$ http://www.yourdomain.com/$1? [L,R=301]
```

Why not use `^yourdomain\.com$` instead? Consider:

```
| RewriteCond %{HTTP_HOST} ^yourdomain\.com$ [NC]
| RewriteRule ^(.*)$ http://www.yourdomain.com/$1? [L,R=301]
```

That would not have matched on typo domains, such as “`yourdoamin.com`”, that the DNS server and virtual host would be set to respond to (assuming that misspelling was a domain you registered and owned).

Under what circumstances might you want to omit the query string from the redirected URL, as we did in the preceding two examples? When a session ID or a tracking parameter (such as `source=banner_ad1`) needs to be dropped. Retaining a tracking parameter after the redirect is not only unnecessary (because the original URL with the source code appended would have been recorded in your access logfiles as it was being accessed); it is also undesirable from a canonicalization standpoint. What if you wanted to drop the tracking parameter from the redirected URL, but retain the other parameters in the query string? Here’s how you’d do it for static URLs:

```
| RewriteCond %{QUERY_STRING} ^source=[a-z0-9]*$
| RewriteRule ^(.*)$ /$1? [L,R=301]
```

And for dynamic URLs:

```
| RewriteCond %{QUERY_STRING} ^(.+)&source=[a-z0-9]+(&?.*)$
| RewriteRule ^(.*)$ /$1?%1%2 [L,R=301]
```

Need to do some fancy stuff with cookies before redirecting the user? Invoke a script that cookies the user and then 301s him to the canonical URL:

```
| RewriteCond %{QUERY_STRING} ^source=([a-z0-9]*)$
| RewriteRule ^(.*)$ /cookiefirst.php?source=%1&dest=$1 [L]
```

Note the lack of a `[R=301]` flag in the preceding code. That’s on purpose. There’s no need to expose this script to the user. Use a rewrite and let the script itself send the 301 after it has done its work.

Other canonicalization issues worth correcting with rewrite rules and the `[R=301]` flag include when the engines index online catalog pages under HTTPS URLs, and URLs missing a trailing slash that should be there. First, the HTTPS fix:

```
| # redirect online catalog pages in the /catalog/ directory if HTTPS
| RewriteCond %{HTTPS} on
| RewriteRule ^catalog/(.*) http://www.yourdomain.com/catalog/$1 [L,R=301]
```

Note that if your secure server is separate from your main server, you can skip the `RewriteCond` line.

Now to append the trailing slash:

```
| RewriteRule ^(.*)[^\]$ /$1/ [L,R=301]
```

After completing a URL rewriting project to migrate from dynamic URLs to static, you’ll want to phase out the dynamic URLs not just by replacing all occurrences of the legacy URLs on your site, but also by 301-redirecting the legacy dynamic URLs to their static equivalents. That way, any inbound links pointing to the retired URLs will end up leading both spiders and humans to the correct new URL—thus ensuring that the new URLs are the ones that are indexed, blogged about, linked to, and bookmarked, and the old URLs will be removed from the index as well. Generally, here’s how you’d accomplish that:

```
| RewriteCond %{QUERY_STRING} id=([0-9]+)
| RewriteRule ^get_product\.php$ /products/%1.html? [L,R=301]
```

However, you’ll get an infinite loop of recursive redirects if you’re not careful. One quick-and-dirty way to avoid that situation is to add a nonsense parameter to the destination URL for the rewrite and ensure that this nonsense parameter isn’t present before doing the redirect. Specifically:

```
| RewriteCond %{QUERY_STRING} id=([0-9]+)
| RewriteCond %{QUERY_STRING} !blah=blah
| RewriteRule ^get_product\.php$ /products/%1.html? [L,R=301]
| RewriteRule ^products/([0-9]+)/?$ /get_product.php?id=$1&blah=blah [L]
```

Notice that the example used two `RedirectCond` lines, stacked on top of each other. All redirect conditions listed together in the same block will be “ANDed” together. If you wanted the conditions to be “ORed”, it would require the use of the [OR] flag.

Redirecting a Home Page Index File Without Looping

Many websites link to their own home page in a form similar to `http://www.yourdomain.com/index.html`. The problem with that is that most incoming links to the site’s home page specify `http://www.yourdomain.com`, thus dividing the link authority into the site. Once a publisher realizes this, they will want to fix their internal links and then 301 redirect `http://www.yourdomain.com/index.html` to `http://www.yourdomain.com/`, but there will be problems with recursive redirects that develop if this is not done correctly.

When someone comes to your website by typing in `http://www.yourdomain.com`, the DNS system of the Web helps the browser locate the web server for your website. How, then, does the web server decide what to show to the browser? It turns out that it does this by loading a file from the hard drive of the web server for your website.

When no file is specified (i.e., as in the preceding example, only the domain name is specified), the web server loads a file that is known as the default file. This is often a file with a name such as `index.html`, `index.htm`, `index.shtml`, `index.php`, or `default.asp`.

The filename can actually be anything, but most web servers default to one type of filename or another. Where the problem comes in is that many CMSs will expose both forms of your home page, both `http://www.yourdomain.com` and `http://www.yourdomain.com/index.php`.

Perhaps all the pages on the site link only to `http://www.yourdomain.com/index.php`, but given human nature, most of the links to your home page that third parties give you will most likely point at `http://www.yourdomain.com/`. This can create a duplicate content problem if the search engine now sees two versions of your home page and thinks they are separate, but duplicate, documents. Google is pretty smart at figuring out this particular issue, but it is best to not rely on that.

Since you learned how to do 301 redirects, you might conclude that the solution is to 301-redirect `http://www.yourdomain.com/index.php` to `http://www.yourdomain.com/`. Sounds good, right? Unfortunately, there is a big problem with this.

What happens is the server sees the request for `http://www.yourdomain.com/index.php` and then sees that it is supposed to 301-redirect that to `http://www.yourdomain.com/`, so it does. But when it loads `http://www.yourdomain.com/` it retrieves the default filename (`index.php`) and proceeds to load `http://www.yourdomain.com/index.php`. Then it sees that you want to redirect that to `http://www.yourdomain.com/`, and it creates an infinite loop.

The default document redirect solution

The solution that follows is specific to the preceding `index.php` example. You will need to plug in the appropriate default filename for your own web server.

1. Copy the contents of `index.php` to another file. For this example, we’ll be using `sitehome.php`.
2. Create an Apache `DirectoryIndex` directive for your document root. Set it to `sitehome.php`. Do not set the directive on a serverwide level; otherwise, it may cause problems with other folders that still need to use `index.php` as a directory index.
3. Put this in an `.htaccess` file in your document root: `DirectoryIndex sitehome.php`. Or, if you aren’t using per-directory context files, put this in your `httpd.conf`.

```
<Directory /your/document/root/exampleshite.com/>
  DirectoryIndex sitehome.php
</Directory>
```

1. Clear out the contents of your original `index.php` file. Insert this line of code:

```
| <? header( "Location: http://www.example.com" ); ?>
```

This sets it up so that *index.php* is not a directory index file (i.e., the default filename). It forces *sitehome.php* to be read when someone types in the canonical URL (<http://www.yourdomain.com>). Any requests to *index.php* from old links can now be 301-redirected while avoiding an infinite loop.

If you are using a CMS, you also need to make sure when you are done with this that all the internal links now go to the canonical URL, <http://www.yourdomain.com>. If for any reason the CMS started to point to <http://www.yourdomain.com/sitehome.php> the loop problem would return, forcing you to go through this entire process again.

Content Management System (CMS) Issues

When looking to publish a new site, many publishers may wonder whether they need to use a CMS, and if so, how to ensure that it is SEO-friendly.

It is first essential to determine whether a CMS is necessary before embarking on a web development project. You can use the flowchart in Figure 6-44 to help guide you through the process.

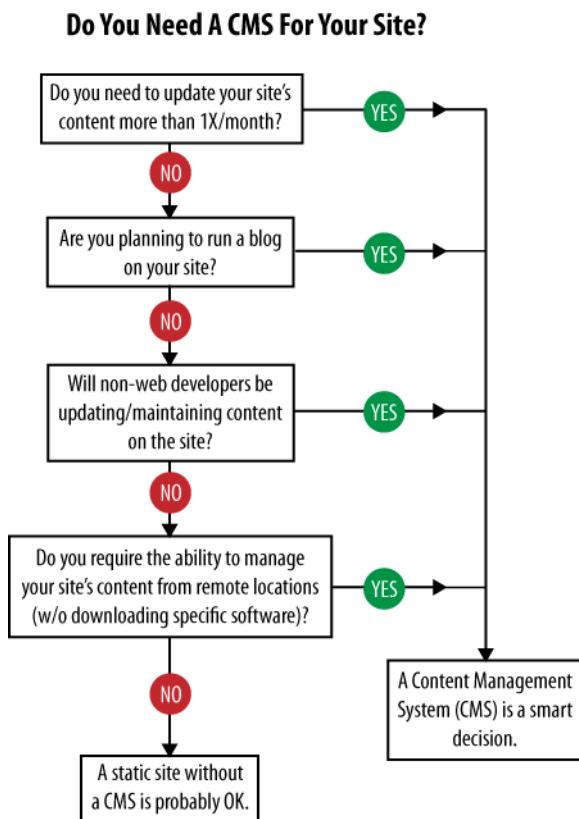


Figure 6-31. Flowchart to determine whether you need a CMS

Due to the inexpensiveness of customizable, free platforms such as Drupal (<http://drupal.org/>), Joomla (<http://www.joomla.org/>), WordPress (<http://wordpress.org/>), and Weebly (<http://www.weebly.com>) it is increasingly rare for a publisher to develop a static site, even when a CMS isn't required.

The next step involves understanding how to ensure that a CMS will be search-engine-friendly. Here is a list of basic SEO issues that frequently plague CMSs (both prebuilt and custom-made). By dealing with these, you will ensure a relatively smooth platform for content delivery.

Title tag customization and rules

A search-engine-friendly CMS must allow for title tags to not only be customized on a page-specific level, but also enable rules for particular sections of a website. For example, if the title tag always has

to start with your site name followed by a colon followed by your article title, your on-page optimization efforts will be limited—at least as far as the powerful title tag is concerned. You should be able to revise the formulas used to generate the title tags across your site to make them more search-optimal.

Static, keyword-rich URLs

URLs have historically been the most problematic SEO issue for CMS platforms. Nowadays, search-friendly CMSs should feature custom URL creation. In WordPress, a custom URL is referred to as a post slug. Figure 6-45 is an example from WordPress.

Notice how the first line allows you to create the title of the post, and the second enables manual creation of the URL structure (and an automatic Generate button if you prefer to simply use the post title).



Figure 6-32. Example of custom URL creation

Meta tag customization

Being able to implement custom meta descriptions and meta robots tags is critical. Enabling editorial control is essential for a good CMS.

Enabling custom HTML tags

A good CMS has to offer extra functionality on HTML tags for things such as `NoFollow` on links, or `<hx>` tags for headlines and subheadlines. These can be built-in features accessible through menu options, or the CMS can simply allow a manual editing of HTML in the text editor window when required. Having no H1 tags on a given page is not desirable. Too many H1 tags on the page are not desirable. Low-value content (such as the publication date) marked up as an H1 is not desirable. The article title is typically the best content to have wrapped in an H1.

Internal anchor text flexibility

To be “optimized” rather than simply search-friendly, customizing the anchor text on internal links is critical. Rather than simply making all links in a site’s architecture the page’s title, a great CMS should be flexible enough to handle custom input from the administrators as to the anchor text of category-level or global navigation links.

Intelligent categorization structure

Another common CMS problem is poor category structure. When designing an information architecture for a website, you should not place limits on how pages are accessible due to the CMS’s inflexibility. CMSs that offer customizable navigation panels will be the most successful in this respect.

Pagination controls

Pagination can be the bane of a website’s search rankings, so controlling it through inclusion of more items per page, more contextually relevant anchor text (e.g., not “next,” “prev,” and page numbers), and careful use of meta `NoIndex` tags will make your important content get more link authority and crawl attention.

301-redirect functionality

Many CMSs sadly lack this critical feature, disallowing the proper redirection of content when necessary; 301s are valuable for expired content, for pages that have a newer version, and for dodging keyword cannibalization issues similar to those we discussed earlier in this chapter.

XML/RSS pinging

Although it is primarily useful for blogs, any content, from articles to products to press releases, can be issued in a feed - and by utilizing quick, accurate pinging of the major feed services, you limit some of your exposure to duplicate content spammers who pick up your feeds and ping the major services quickly in the hopes of beating you to the punch.

Image-handling and alt attributes

Alt attributes are a clear must-have from an SEO perspective, serving as the “anchor text” when an image is used as a link (note that text links are much better than images with alt attributes for links, but if you must use image links you do want to have the alt attribute implemented), and providing relevant, index-able content for the search engines. Images in a CMS’s navigational elements should preferably use CSS image replacement rather than mere alt attributes.

CSS exceptions

The application of CSS styles in a proper CMS should allow for manual exceptions so that a user can modify how a strong headline or list element appears visually. If the CMS does not offer this, writers may opt out of using proper semantic markup for presentation purposes, which would not be a good thing.

Static caching options

Many CMSs currently offer caching options, which are a particular boon if a page is receiving a high level of traffic from social media portals or news sites. A bulky CMS often makes dozens of extraneous database connections, which can overwhelm a server if caching is not in place, killing potential inbound links and media attention.

URLs free of tracking parameters and session IDs

Sticking session or tracking information such as the user’s click path into the URL is deadly for SEO. It usually leads to incomplete indexation and duplicate content issues.

Customizable URL structure

If the default URL structure of the CMS doesn’t suit your needs, you should be able to change it. For example, if you don’t want `/archives/` in the URLs of all your archived articles, you should be able to remove it. Or if you want to reference the article name instead of the article’s database ID in the URL, you should be able to do it.

301 redirects to a canonical URL

Duplicate content is a major concern for the dynamic website owner. Automatic handling of this by the CMS through the use of 301 redirects is a must.

Static-looking URLs

The most palatable URLs to spiders are the ones that look like they lead to static pages—no query strings in the URL.

Keywords in URLs

Keywords in your URLs (used judiciously) can help your rankings.

RSS feeds

The CMS should auto-create RSS feeds to help your site rank in Google Blog Search and other feed engines.

Multilevel categorization structure

It is awfully limiting to your site structure and internal hierarchical linking structure to have a CMS that doesn't allow you to nest subcategories into categories, sub-subcategories into subcategories, and so on.

Paraphrase-able excerpts

Duplicate content issues are exacerbated on dynamic sites such as blogs when the same content is displayed on Permalink pages, category pages, archives-by-date pages, tag pages, and the home page. Crafting unique content for the excerpt and having that content display on all locations except for the Permalink page will help strengthen your Permalink page as unique content.

Breadcrumb navigation

Breadcrumb (drill-down) navigation is great for SEO because it reinforces your internal hierarchical linking structure with keyword-rich text links.

Meta `NoIndex` tags for low-value pages

Even if you use `NoFollow` attributes in links to these pages, other people may still link to them, which carries a risk of ranking those pages above some of your more valuable content.

Keyword-rich intro copy on category-level pages

Keyword-rich introductory copy helps set a stable keyword theme for the page, rather than relying on the latest article or blog post to be the most prominent text on the page.

`NoFollow` links in comments

If you allow visitors to post comments and do not `NoFollow` the links, your site will be a spam magnet. Heck, you'll probably be a spam magnet anyway, but you won't risk losing PageRank to spammers if you use `NoFollow` attributes.

Customizable anchor text on navigational links

“Contact,” “About Us,” “Read More,” “Full Article,” and so on make for lousy anchor text—at least from an SEO standpoint. Hopefully, your CMS allows you to improve such links to make the anchor text more keyword-rich.

XML Sitemap generator

Having your CMS generate your XML Sitemap can save a lot of hassle, as opposed to trying to generate one with a third-party tool.

HTML4, HTML5, or XHTML validation

Although HTML validation is not a ranking signal, it is desirable to have the CMS automatically check for malformed HTML, as search engines may end up seeing a page differently from how it renders on the screen and accidentally consider navigation to be part of the content or vice versa.

Pingbacks, trackbacks, comments, and antispam mechanisms

The problem with comments/trackbacks/pingbacks is that they are vectors for spam, so if you have one or more of these features enabled, you will be spammed. Therefore, effective spam prevention in the form of Akismet, Mollom, or Defensio is a must.

If you want more information on picking a quality CMS, some great web resources are already out there, among them <http://php.opensourcecms.com/> and <http://www.cmsmatrix.org/>, to help manage this task.

Selecting a CMS

There are many factors to consider when choosing an existing CMS. Many CMSs are free, but some of them are proprietary with a license cost per site. The majority of CMSs were not designed with security, stability, search friendliness, and scalability in mind, though in recent years a few vendors have developed

excellent CMSs that have search friendliness as their primary focus. Many were developed to fit a certain market niche, but can be expanded to fit other purposes. Some are no longer maintained. Many are supported and developed primarily by hobbyists who don't particularly care if you're having trouble getting them installed and configured. Some are even intentionally made to be difficult to install and configure so that you'll be encouraged to pay the developers a consulting fee to do it all for you.

Popular CMS solutions that the authors have experience with include Joomla (<http://www.joomla.org>), Drupal (<http://drupal.org>), concrete5 (<http://www.concrete5.org>), Pixelsilk (<http://www.pixelsilk.com>), Wordpress (<http://www.wordpress.org>), Magento (<http://www.magento.com>), and Sitecore (<http://www.sitecore.net>). Each of these has strong support for SEO, but each of them requires some configuration for optimal results. Make sure you get that help up front to get the SEO for your site off to a strong start.

Selecting a CMS is an important process. If you make the wrong choice, you will be faced with limited SEO options. Like most software, CMSs are a moving target—what's missing today may be a new feature tomorrow. In addition, just because a feature exists doesn't mean it is the default option, so in many instances the desired functionality will need to be enabled and possibly customized to work to your specifications.

Third-Party CMS Add-ons

Many CMS platforms offer many third-party plug-ins or add-ons that extend the core functionality of the CMS. In the WordPress plug-in directory alone there are over 34,000 plug-ins. Plug-ins provide a simple way to add new SEO features and functionality, making the CMS much more flexible and future-proof. It is particularly helpful when there is an active community developing plug-ins. An active community also comes in very handy in providing free technical support when things go wrong; and when bugs and security vulnerabilities crop up, it is important to have an active developer base to solve those issues quickly.

CMS add-ons can also come in the form of independent software installed on your web server or hosted services. Such tools come in many forms, discussion forums, customer reviews, and user polls. Discussion forums, for example, come in two forms, with bbPress, which is installed software and is optimized for search, and vbulletin, which is a hosted solution and therefore more difficult to optimize for search.

The problem with hosted solutions is that you are helping to build the service providers' link authority and not your own, and you have much less control over optimizing the content.

As we referenced several times earlier in this chapter, Flash is popular on the Web, but presents challenges to the search engines in terms of indexing the related content. This creates a gap between the user experience with a site and what the search engines can find on that site.

It used to be that search engines did not index Flash content at all. In June 2008, Google announced that it was offering improved indexing of this content (<http://googlewebmastercentral.blogspot.com/2008/06/improved-flash-indexing.html>). This announcement indicates that Google can index text content and find and follow links within Flash files. However, Google still cannot tell what is contained in images within the Flash file. Here are some reasons why Flash is still not fully SEO-friendly:

Different content is **not** on different URLs

This is the same problem you encounter with AJAX-based pages. You could have unique frames, movies within movies, and so on that appear to be completely unique portions of the Flash site, yet there's often no way to link to these individual elements.

The breakdown of text is not clean

Google can index the output files in the SWF file to see words and phrases, but in Flash, a lot of your text is not inside clean `<h1>` or `<p>` tags; it is jumbled up into half-phrases for graphical effects and will often be output in the incorrect order. Worse still are text effects that often require "breaking" words apart into individual letters to animate them.

Flash gets embedded

A lot of Flash content is only linked to by other Flash content wrapped inside shell Flash pages. This line of links, where no other internal or external URLs are referencing the interior content, means some very low PageRank/link authority documents. Even if they manage to stay in the main index, they probably won't rank for anything.

Flash doesn't earn external links like HTML

An all-Flash site might get a large number of links to the home page, but interior pages almost always suffer. For embeddable Flash content, it is the HTML host page earning those links when they do come.

SEO basics are often missing

Anchor text, headlines, bold/strong text, `img alt` attributes, and even title tags are not simple elements to properly include in Flash. Developing Flash with SEO in mind is just more difficult than doing it in HTML. In addition, it is not part of the cultural lexicon of the Flash development world.

A lot of Flash isn't even crawlable

Google has indicated that it doesn't execute external JavaScript calls (which many Flash-based sites use) or index the content from external files called by Flash (which, again, a lot of Flash sites rely on). These limitations could severely impact what a visitor can see versus what Googlebot can index.

Note that it used to be that you could not test the crawlability of Flash, but the Adobe Search Engine SDK does allow you to get an idea as to how the search engines will see your Flash file.

Flash Coding Best Practices

If Flash is a requirement for whatever reason, there are best practices you can implement to make your site more accessible to search engine spiders. What follows are some guidelines on how to obtain the best possible results.

Flash meta tags

Beginning with Adobe/Macromedia Flash version 8, there has been support for the addition of title and description meta tags to any `.swf` file. Not all search engines are able to read these tags yet, but it is likely that they will soon. Get into the habit of adding accurate, keyword-rich title tags and meta tags to files now so that as search engines begin accessing them, your existing `.swf` files will already have them in place.

Adobe Flash search engine SDK

Flash developers may find the SDK useful for server-based text and link extraction and conversion purposes, or for client-side testing of their Flash content against the basic Adobe (formerly Macromedia) Flash Search Engine SDK code.

Tests have shown that Google and other major search engines now extract some textual content from Flash `.swf` files. It is unknown whether Google and others have implemented Adobe's specific Search Engine SDK technology into their spiders, or whether they are using some other code to extract the textual content. Again, tests suggest that what Google is parsing from a given `.swf` file is very close to what can be extracted manually using the Search Engine SDK.

The primary application of Adobe's Search Engine SDK is in the desktop testing of `.swf` files to see what search engines are extracting from a given file. The program cannot extract files directly from the Web; the `.swf` file must be saved to a local hard drive. The program is DOS-based and must be run in the DOS Command Prompt using DOS commands.

By running a `.swf` file through the Flash SDK `swf2html` program during development, the textual assets of the file can be edited or augmented to address the best possible SEO practices—honing in primarily on keywords and phrases along with high-quality links. Because of the nature of the Flash program and the way in which it deals with both text and animation, it is challenging to get exacting, quality SEO results. The goal is to create the best possible SEO results within the limitations of the Flash program and the

individual Flash animation rather than to attempt the creation of an all-encompassing SEO campaign. Extracted content from Flash should be seen as one tool among many in a larger SEO campaign.

Internal Flash coding

There are several things to keep in mind when preparing Flash files for SEO:

- Search engines currently do not read traced text (using the `trace()` function) or text that has been transformed into a shape in Flash (as opposed to actual characters). Only character-based text that is active in the Flash stage will be read (see Figure 6-46).
- Animated or affected text often creates duplicate content. Static text in Flash movies is not read as duplicate instances that “tweening” and other effects can create. Use static text, especially with important content, so that search engines do not perceive the output as spam (see Figure 6-47).
- Search engine spiders do not see dynamically loaded content (text added from an external source, such as an XML file).
- The font size of text does not affect search engines; they read any size font.
- Special characters such as `<`, `>`, `&`, and `“` are converted to HTML character references (`<`, `>`, `&`, and `"`) and should be avoided.
- Search engines find and extract all URLs stored within the `getURL()` command.
- Search engines have the ability to follow links in Flash, though it is an “iffy” proposition at best. They will not, however, follow links to other Flash `.swf` files. (This is different from loading child `.swf` files into a parent `.swf` file.) Therefore, links in Flash should always point to HTML pages, not other `.swf` files.

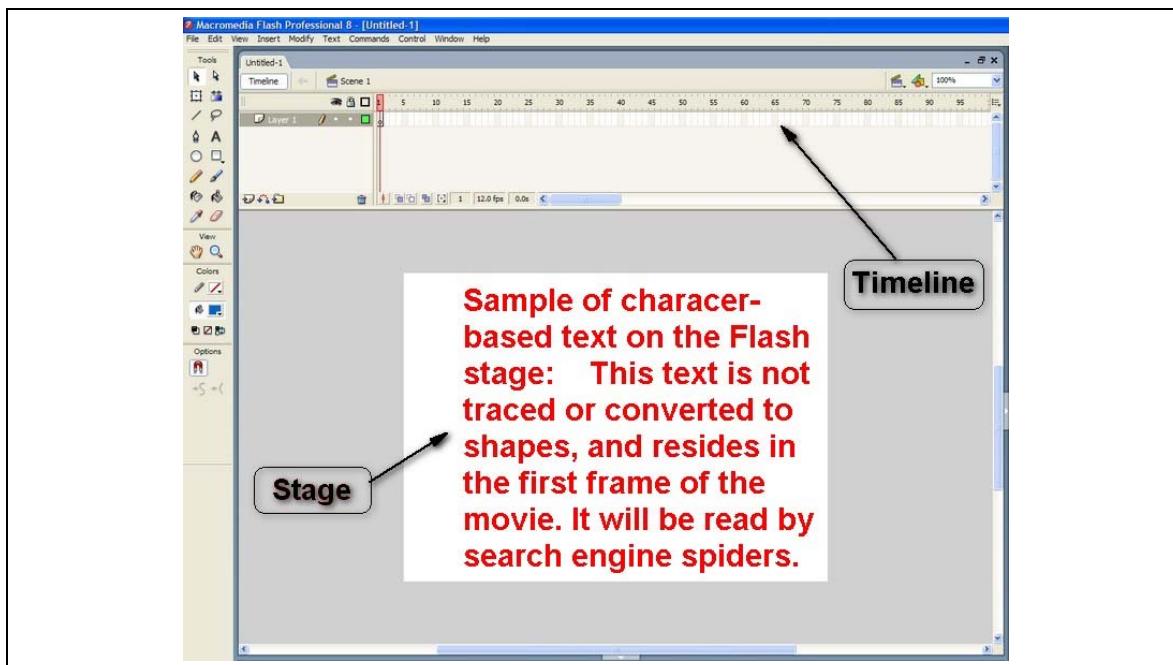


Figure 6-33. Example of spider-readable text inside a Flash program

```

<p>Fast Homework Help</p>
<p>Fast Homework Help</p>
<p>Reliable Report Info</p>
<p>Reliable Report Info</p>
<p>Math Made Easy</p>
<p>Math Made Easy</p>
<p>Brain Games + More</p>
<p>Brain Games + More</p>
<p>SCHOOL TOOLS</p>
<p>SCHOOL TOOLS</p>
<p>What Is Cosmeo?</p>
<p>What Is Cosmeo?</p>
<p>Why It Works</p>
<p>Why It Works</p>
<p>Take A Tour</p>
<p>Take A Tour</p>
<p>Get Started</p>
<p>Get Started</p>

```

Figure 6-34. Animated text results in Flash source; can be seen as duplicate content

SWFObject and NoScript tags

Because “alternative content” workarounds for SEO of Flash files have been historically abused by spammers, it is challenging to recommend these tactics to optimize your Flash files without a critical disclaimer.

Both the `SWFObject` and `NoScript` methods were originally designed to be legitimate, graceful degradation tactics readily accepted by the search engines as a way to accommodate older browsers or people with special needs. But many unscrupulous sites have used the code to trick search engine spiders. In other words, they are used in such a way that browsers display one thing to users, but display something completely different to search engine spiders. All of the major search engines disapprove of such tactics.

Websites using such methods today are often penalized or removed from search engine indexes altogether. This makes graceful degradation risky on some level, but if the methods are used clearly within the boundaries for which they were intended, getting penalized or banned is highly unlikely.

Intent is an essential element search engines take into consideration. If your intent is to provide *all* users with a positive experience while visiting your site, you should be fine. If your intent is to game the search engines, all it takes is one online rival to report your site for spam to incur the wrath of the search engines.

Google and other search engines do not algorithmically ban sites for using `SWFObject` and `NoScript` tags; it usually requires human intervention to evoke a penalty or outright ban.

SWFObject

With regard to Flash optimization, `SWFObject` is the better of the two options because it is JavaScript code designed specifically for Flash `.swf` purposes, and it has been abused to a lesser extent than the `NoScript` tag option.

`SWFObject` is Flash detection code written in JavaScript that checks whether a browser has the Flash plug-in. If the browser does have the Flash plug-in, the `.swf` file is displayed secondary to that detection. If the browser does not have the Flash plug-in or the JavaScript to detect it, the primary, alternative content contained within `<div>` files is displayed instead. The key here is that search engine spiders do not render the JavaScript. They read the primary content in the `<div>` tags.

The opportunity for abuse is obvious upon viewing the code. This small piece of code is placed within the `<head>` tags:

```
| <script type="text/javascript" src="swfobject.js"></script>
```

In the body of the text, the code looks something like Figure 6-48.

```

<script type="text/javascript" src="swfobject.js"></script>

<div id="flashcontent">
  Text, links, and graphics placed here are replaced by the Flash movie. Search
  engine spiders will read this information, but the browser with an active Flash
  plugin will show the Flash movie instead.
</div>

<script type="text/javascript">
  var so = new SWFObject("whatever.swf", "themovie" , "200", "100", "7", "#336699");
  so.write("flashcontent");
</script>

```

Figure 6-35. Information between the `<div>` HTML tags read by search engine spiders

Search engine spiders will read text, links, and even `alt` attributes within the `<div>` tags, but the browser will not display them unless the Flash plug-in isn't installed (about 95% of browsers now have the plug-in) or JavaScript isn't available.

Once again, the key to successfully implementing `SWFObject` is to use it to the letter of the law; leverage it to mirror the content of your Flash `.swf` file *exactly*. Do not use it to add content, keywords, graphics, or links that are not contained in the file. Remember, a human being will be making the call as to whether your use of `SWFObject` is proper and in accordance with that search engine's guidelines. If you design the outcome to provide the best possible user experience, and your intent is *not* to game the search engines, you are probably OK.

You can download the `SWFObject` JavaScript free of charge at <http://code.google.com/p/swfobject/>. Included in this download is the `flashobject.js` file, which is placed in the same directory as the web pages upon which the corresponding calling code resides.

NoScript

The `NoScript` tag has been abused in “black hat” SEO attempts so frequently that caution should be taken when using it. Just as `SWFObject` and `DIV` tags can be misused for link and keyword stuffing, so too can the `NoScript` tag. Certain companies have promoted the misuse of the `NoScript` tag widely; consequently there have been many more problems with their use.

With that being said, conservative and proper use of `NoScript` tags specifically with Flash `.swf` files can be an acceptable and good way to get content mirrored to a Flash file read by search engine spiders. As it is with `SWFObject` and corresponding `DIV` tags, content must echo that of the Flash `.swf` movie exactly. Do not use them to add content, keywords, graphics, or links that are not in the movie. Again, it is a human call as to whether a site or individual page is banned for the use or misuse of `NoScript` tags.

You use `NoScript` tags with Flash `.swf` files in the following manner:

```
| <script type="text/javascript" src="YourFlashFile.swf"></script>
```

Followed at some point afterward by:

```

<noscript>
  <H1>Mirror content in Flash file here.</H1>
  <p>Any content within the NoScript tags will be read by the search engine
  spiders, including links
  http://www.mirroredlink.com, graphics, and corresponding Alt attributes.
</noscript>

```

For browsers that do not have JavaScript installed or functioning, content alternatives to JavaScript-required entities are displayed. So, for use with Flash `.swf` files, if a browser does not have JavaScript and therefore cannot display Flash, it displays instead the content inside the `NoScript` tags. This is a legitimate, graceful degradation design. For SEO purposes, as it is with `SWFObject`, the search engine spiders do not render the JavaScript and do read the content contained in the HTML. Here, it is the content in the `NoScript` tags.

Scalable Inman Flash Replacement (sIFR)

sIFR is a technique that uses JavaScript to read in HTML text and render it in Flash instead. The essential fact to focus on here is that the method guarantees that the HTML content and the Flash content are identical. One great use for this is to render headline text in an anti-aliased font (this is the purpose for which sIFR was designed). This can provide a great improvement in the presentation of your site.

Dan Crow, head of Google's Crawl team, said that as long as this technique is used in moderation, it is OK. However, extensive use of sIFR could be interpreted as a poor site quality signal. Since sIFR was not designed for large-scale use, such extensive use would not be wise in any event.

It is worth noting that there are similar technologies available to web designers for improved type presentation, which provide similar search engine friendliness. FLIR (FaceLift Image Replacement) is an image replacement script similar to sIFR in its use of JavaScript, but without the Flash element, and there is a handy Wordpress plugin for implementation on Wordpress-based websites (<http://wordpress.org/extend/plugins/facelift-image-replacement/>).

Best Practices for Multilanguage/Country Targeting

Many businesses target multiple countries with their websites and need answers to questions such as: do you put the information for your products or services all on the same domain? Do you obtain multiple domains? Where do you host the site(s)? It turns out that there are SEO factors, as well as basic marketing questions, that affect the answers. There are also non-SEO factors, such as the tax implications of what you do; for some TLDs you can get them only by having a local physical presence (e.g., France requires this to get a .fr domain).

Targeting a Specific Country

Starting with the basics of international targeting, it is important to let the search engines know where your business is based in as many ways as possible. These might include:

- A country-specific TLD (ccTLD) for your domain (e.g., .uk)
- Hosting in the local country (more for content delivery speed than “localness” factor of hosting)
- Physical local address in plain text on every page of your site
- Google Webmaster Tools geotargeting setting
- Verified address with Google Maps
- Links from in-country websites
- Use of the local language on the website

If you are starting from scratch, getting these all lined up will give you the best possible chance of ranking in the local country you are targeting.

Problems with Using Your Existing Domain

You may ask why you cannot leverage your domain weight to target the new territory rather than starting from scratch—in other words, create multiple versions of your site and determine where the user is in the world before either delivering the appropriate content or redirecting him to the appropriate place in the site (or even to a subdomain hosted in the target country).

The problem with this approach is that the search engines spider from the United States, their IP addresses will be in the United States in your lookup, and they will therefore be delivered in U.S. content. This problem is exacerbated if you are going even further and geodelivering different language content, as only your English language content will be spidered unless you cloak for the search engine bots.

This kind of IP delivery is therefore a bad idea. You should make sure you do not blindly geodeliver content based on IP address as you will ignore many of your markets in the search engines' eyes.

The Two Major Approaches

The best practice remains one of two approaches, depending on the size and scale of your operations in the new countries and how powerful and established your .com domain is.

If you have strong local teams and/or (relatively speaking) less power in your main domain, launching independent local websites geotargeted as described earlier is a smart move in the long run.

If, on the other hand, you have only centralized marketing and PR and/or a strong main domain, you may want to create localized versions of your content either on country specific subdomains (<http://uk.yourdomain.com>, <http://au.yourdomain.com>, etc.) or in subfolders (`/uk/`, `/au/`, etc.), with the preference being for the use of subdomains so that you can set up local hosting.

Both the subdomains and the subfolders approaches allow you to set your geotargeting option in Google Webmaster Tools, and with either method, you have to be equally careful of duplicate content across regions. In the subdomain example, you can host the subdomain locally, while in the subfolder case, more of the power of the domain filters down.

Unfortunately, the Webmaster Tools' geotargeting option doesn't work nearly as well as you'd hope to geotarget subfolders. The engines will consider hosting and ccTLDs, along with the geographic location of your external link sources, to be stronger signals than the manual country targeting in the tools. In addition, people in other countries (e.g., France) don't like to click on ".com" or ".org" TLDs—they prefer ".fr". This extends to branding and conversion rates too—web users in France like to buy from websites in France that end in ".fr".

Multiple-Language Issues

An entire treatise could be written on handling multi-language content as the search engines themselves are rapidly evolving in this field, and tactics are likely to change dramatically in the near future. Therefore, this section will focus on providing you with the fundamental components of successful multi-language content management.

Here are best practices for targeting the search engines as of this writing, using Spanish and English content examples:

- Content in Spanish and English serving the same country:
 - Create a single website with language options that change the URL by folder structure; for example, <http://www.yourdomain.com> versus <http://www.yourdomain.com/esp/>.
 - Build links from Spanish and English language sites to the respective content areas on the site.
 - Host the site in the country being served.
 - Register the appropriate country domain name (for the United States, .com, .net, and .org are appropriate, whereas in Canada using .ca or in the United Kingdom using .uk is preferable).
 - Markup your HTML code using hreflang tags for multiple languages. See "Hreflang for multiple languages / no specific location" below.
- Content in Spanish and English targeting multiple countries:
 - Create two separate websites, one in English targeting the United States (or the relevant country) and ones in Spanish targeting the relevant Spanish-speaking countries.
 - Host one site in the United States (for English) and the others in the relevant countries for the Spanish version.
 - Register different domains, one using U.S.-targeted domain extensions and one using the Spanish-speaking countries' extension.
 - Acquire links from the United States to the English site and links from the Spanish-speaking countries to those sites.

- Markup your HTML code using hreflang tags for multiple languages/locations. See “Hreflang for Multiple Language/Region Combinations” below.
- Content in Spanish targeting multiple countries:
 - Create multiple websites (as mentioned earlier) targeting each specific country.
 - Register domains using the appropriate country TLD and host in each country separately.
 - When possible, have native speakers fluent in the specific region’s dialect write the site content for each specific country.
 - Obtain in-country links to your domains.
 - Markup your HTML code using hreflang tags for multiple languages/locations. See “Hreflang for One Language/Multiple Regions” below

Although some of these approaches may seem counterintuitive, the joint issues of search engines preferring to show content hosted in and on a country-specific domain name combined with duplicate content problems make for these seemingly illogical suggestions.

Hreflang Markup

There are several options for serving multi-region and multilanguage content to the search engines. At least one of these solutions will be necessary to encourage Google to rank the appropriate version of your content in the appropriate version of the Google search engine (google.com, google.co.uk, google.ca. etc). These solutions are also necessary to prevent duplicate content issues both within a single language, such as American English and UK English, where the copy is likely to be virtually identical, and also across languages which are more unique.

There are 3 main options available to serve up multilanguage or multi-region content. These are:

- Code within the server header section of a page
- Code within the <head> section of the HTML on a page
- Special directives within the site’s site map, or a specific multi-region/multilanguage site map.

It’s recommended that only one of these solutions be used at a time. While redundancy, if accurate, will cause no negative effects the possibility exists for there to be disagreement between multiple solutions if they are working simultaneously, which can cause the search engines confusion as to which version to “count”.

We will focus on the “Code within the <head> section of the HTML on a page” solution.

Hreflang for Multiple Languages / No Specific Location

The markup contained within each page which has alternate language versions but not alternate country versions should contain markup specifying the language only. It is acceptable to markup pages with language-only markup but never suitable to markup pages with region-only markup. Once pages are built for specific regions they must be marked up with a combination of both language and region markup. An example of this markup for a homepage presented in both English and Spanish is below.

If the homepage of a site, in this case example.com, or translated into both English and Spanish both versions of the page should include code such as the below:

```
<link rel="alternate" hreflang="x-default" href="example.com" />
<link rel="alternate" href="example.com/es/" hreflang="es" />
```

Each unique language will have its own unique hreflang code. Note that there is no accommodation within the language markup for the difference between Spanish for Spain and Spanish for Latin America. Similarly there is no difference in the language markup between Portuguese for Portugal and Portuguese for Brazil, or Canadian French versus the version spoken in France, etc.

A full list of the two character language codes can be found at http://www.loc.gov/standards/iso639-2/php/code_list.php under the ISO 639-1 standard.

Hreflang for Multiple Language/Regions

If you wanted to have a default version of the page (English language, no region assigned), versions for Spanish from Mexico and Spanish from Spain, the markup would look similar to the below. Please note as well that each region/language combination would need its own unique URL/domain.

```
<link rel="alternate" hreflang="x-default" href="example.com" />
<link rel="alternate" href="example.es/" hreflang="es-es" />
<link rel="alternate" href="example.com.mx/" hreflang="es-mx" />
```

Another consideration in your geo targeting of URLs is that there are no provisions for markup associated with “regions” such as Latin America, APAC, the EU, etc. Each country within these regions must be treated individually.

Hreflang for One Language/Multiple Regions

If you wished to have versions in Spanish for Spain and versions for Spanish for Latin America you would use markup similar to the below.

```
<link rel="alternate" hreflang="x-default" href="example.es" />
<link rel="alternate" href="example.es/" hreflang="es-es" />
<link rel="alternate" href="example.com.mx/" hreflang="es-mx" />
<link rel="alternate" href="example.com.cr/" hreflang="es-cr" />
<link rel="alternate" href="example.com.com.ar/" hreflang="es-ar" />
```

... and so on for each Latin American country

A full list of country level TLD's can be found at <http://www.mcanerin.com/EN/articles/ccTLD.asp>

It pays to plan ahead when adding hreflang markup for alternate language/country versions of your site. Each alternate version of a page needs to reference every other alternate version. If you have a version of a page in English the United States and another in Spanish for Mexico both of those pages need the markup referencing the other version. If you were to then add a version in Spanish for Spain not only does this new version need to reference both the English/US version and the Spanish/Mexican version but both of those pages now also need to reference this new Spanish/Spain version.

This level of complexity is why it is crucial that before you proceed with a strategy wherein you would create alternate language/region versions of your content that you have plan a comprehensive international strategy. Before you embark on any site changes which are detectable by the search engines you should be planning ahead several years as to what regions and languages you will optimize for. If you choose not to do this level of planning it could require numerous code changes across all alternate version pages as new countries/languages are added in the future.

For more info on hreflang markup see <https://support.google.com/webmasters/answer/189077?hl=en> and this post by Aleyda Solis (<http://www.aleydasolis.com/en/>) <http://moz.com/blog/using-the-correct-hreflang-tag-a-new-generator-tool>. You can also find the “hreflang Tags Generator Tool” created by Aleyda at <http://www.internationalseomap.com/hreflang-tags-generator/>.

A special thanks to Rob Woods (<http://robwoods.org>) for his contributions to this portion of the chapter.

Semantic Search

There is a lot of confusion when it comes the definition of semantic search. Some of it comes from the formalities of the definition of semantics commonly associated with linguistics and some of it comes from the misunderstanding that arises the moment the words “structured data” are mentioned.

In truth semantic search has a little to do with both and a lot to do with the four vectors that drive Big Data across the web:

- Volume is about processing massive amounts of data and extracting unique meaning from it.
- Velocity refers to the speed at which critical data comes in and how quickly it must be analyzed and processed.
- Variety is required as well, as many different types of data must be handled, such as audio, video, and text.
- Veracity is about the need to validate the accuracy of the data being processed.

To understand it better it helps to take things from the beginning and the true beginning for semantic search started on August 30th, 2013 when Google quietly rolled out Hummingbird.

The change, which was announced on the eve of Google’s 15th birthday, almost a month later, completed a long journey for Google to turn search into more than a blind fishing expedition where those who created content and those who looked for it continually strove to second-guess each other’s keywords and connect.

Google’s Hummingbird

To understand how much search (and SEO) have changed, consider just how far voice search has come. When we use voice we tend to speak in sentences instead of keywords, and in order for Google to return meaningful answers it has to be able to understand our search query. To handle spoken queries well, Google also needs to understand intent, which requires it to be able to understand context.

The same technology that was applied to Google Voice Search before 30th August 2013, is now applied to the regular text search that everyone is familiar with. Hummingbird (which Google said was so named because it was precise and fast) does a number of things no search engine had done before.

First, it takes the entire search query into account. Not just the keywords but every word. Second, it looks at who is carrying out the search. Suddenly, variables such as past search history and search patterns are important in delivering the right results, at the right time, to the right person. Thirdly, it also factors in how the search itself is being conducted. Device type, time of the day and location now are also important parameters affecting the search results.

With linguistic sensitivity (i.e. the ability to better process natural language) Google’s Hummingbird is also better at understanding the relation between queries and the relation between bits of data. It is in this space that the real magic happens.

Semantic search, really, is about relational connections and contextual content. In order to deliver “the right results, at the right time, to the right person.” semantic search needs to understand the importance of the query to that person and the importance of the query in relation to the data it holds in its index and the data it is currently indexing.

Every single item of data that is in the visible web needs to be crawled, indexed and evaluated against all the other items of data and then weighed up against a particular search query. The net result of this approach is that the traditional first page of Google everyone strove to rank for in the past has now largely disappeared.

If all of this makes it sound like SEO as we’ve known it is dead, nothing could be further from the truth (though had this been the case it might have made for a much smaller book, here).

Semantic Search and SEO

Basic SEO factors are still in play. Links are still important. Keywords still play a role. Content still needs to be created. But as more and more variables are added to the picture, the value ascribed to each one decreases. This makes it difficult to pick specific aspects of search engine optimization, focus on them to the exclusion of everything else and expect it to be enough.

This is an important change. In the past you could get away with thin content, for instance, if you had a large number of links coming in that would boost it in search. You could get away with poor quality design if you had sufficient keywords to draw in the ‘crowds’ to your page through search. You could get away with links in some suspect neighborhoods if you had a sufficiently large number of links for the percentage of bad links to be small enough to be overlooked.

You could, in short, take some shortcuts that might have been ‘bad’ when you planned to be ‘good’ once you got where you needed to in terms of search ranking. There was an expectation that the end result justified the risks and that things would balance out in your favor, eventually.

This is no longer the case anymore. Because Google now needs to deliver high quality, high-confidence results in search it has to have confidence in the content presented. The veracity factor becomes critical. This makes every activity intended to optimize a website, crucial. Design, content, website structure, keywords, traffic, traffic behavior, social network footprint, links, comments and citations. The list is far from exhaustive and, everything in fact that helps build a data-driven impression of what a website is all about and the quality of its content, now becomes an element or activity you need to consider.

The reason this has happened has to do with two things that are synonymous with semantic search: Entities and structured data.

Entities and Semantic Search

An entity is something that exists in itself. It can be real, like a car, or fictional, like a film or book character like Harry Potter. In either case, it possesses properties, qualities and attributes that make up the “thing” it represents. All of these are language-independent, though obviously when an entity is described we do need a language in which to describe it in. The properties, qualities and attributes, along with associated entities, form the Knowledge Graph that is used to define new entities.

Entities are at the heart of what Google calls the transition from “strings to things”. While all the information on the web is data, entities allow Google to understand how that information fits in and how accurate it is (the veracity aspect of semantic search). In order for an entity to be created in Google’s index, Google needs to index all the properties, qualities and attributes about it and understand the relational connections between them. This is exactly why the web, with semantic search, is becoming more transparent. Data is now portable. Its origin is every bit as important as the data itself. The connections between different data pieces are being indexed and the importance of the data itself is becoming more understood.

The concept of entities has a large impact on how to pursue SEO. While tasks such as keyword research and getting links to your site remain important, you must also pursue holistic strategies to build the reputation and visibility of your business that cause it to have positive associations in many ways across the web.

This brings SEO and marketing a lot closer than they have ever been and makes SEO, as a whole, something that should be part of the DNA of a business rather than a bolt-on activity that can be picked up and dropped as the need arises.

So, how do we begin? What are the guiding principles that you need to have in mind in this new world of SEO? Funnily enough, the concept is as simple to plan as it is difficult to apply. It starts off from the very basic questions of: Who? Why? and How?

If you cannot successfully answer these three questions in relation to your business in a way that begins to form a distinct and unique identity, then the chances are high that neither can Google nor can your prospective customers. Your SEO, then is governed by activities that make sense at a technical level but not

at a brand identity one. Semantic search is all about establishing that identity, even from a business point of view. Especially, one might argue, from a business point of view. This is what helps with the formation of entities in the Google search index.

Entities then become high-trust points which help Google's semantic search understand the value of information better. As you engage in your overall digital marketing strategy, keep these three areas of concern in focus:

- Trust
- Authority
- Reputation

These three, more than anything else, will help your business find its audience, keep it and grow.

What About Structured Data?

Structured data is the label that's applied to a number of markup formats that allow Google to better understand the data it is indexing. Structured data is then nothing more than *metadata* (data about data) implemented for search engines rather than people.

Google, Microsoft (Bing) and Yahoo worked to establish schema.org which is an independent, W3C-approved way of implementing structured data across the web (see the following section of this chapter for more information on schema.org). Unfortunately, many SEOs believe that this is a shortcut to better rankings, which it's not. Semantic search is *all* about structured data. The entire effort that Google has undertaken involves indexing the unstructured data that is found across the web and then placing it in structured data format within its index.

That does not mean, however, that structured data on a website is a ranking signal. It helps in better indexing, but ranking depends upon other factors including the quality of the content, its value, uniqueness and even freshness. Despite the fact that Google is one of the founding organizations behind the Schema.org structured data markup initiative, Google will also attempt to extract entity information from unstructured data through its own efforts.

There are several good reasons for this:

- Adoption (structured data markup is notoriously difficult to implement if you do not know any coding)
- Accuracy (the moment human agents are involved in the markup of data, mistakes happen)
- Consistency (even when structured data is applied without errors, there are still differences in the categorization of content and a confusion as to how to best apply semantic identifiers)
- Reliability (there will always be a temptation to game search by implementing structured data markup in ways intended to boost ranking; Google has already had a number of manual action penalties in search designed to remove such spammy results)

The million dollar question is: is there anything you can do to help Google index your site better if you do not implement structured data markup?

The answer is yes; implement all the search engine optimization tools you have in your arsenal in a way that makes sense for a human user, first, and a search engine second.

Namely, your on-page SEO should help a reader better navigate your content and make sense of it at a glance. The keywords, synonyms, and entities you use in your content should also do the same. Any links you include, and the anchor text of those links, must similarly fill in those blanks.

If you're running a bricks & mortar business, all the relevant information should be included on your pages, such as your name address and phone number. You should interlink your web properties (such as your site and social media accounts) with your Google+ presence.

Where possible, on your site, make use of Google's structured data highlighter tool (<http://goo.gl/0Sp10z>). Finally make use of Google My Business (<http://www.google.com/business/>) and ensure have a cohesive

presence on the web, whose effectiveness you can measure. Use Schema.org to help search engines better understand the content of your pages.

In addition, go build lots of positive relationships on the web that help drive signals of trust and authority back to your web site and business.

A special thanks to David Amerland (<http://www.DavidAmerland.com>) for his contributions to this portion of the chapter.

Schema.org

Schema.org is best viewed as part of a much larger idea, one that traces its origins back to the foundational concepts of the web itself, and its progenitor, Tim Berners-Lee. In his seminal article in Scientific American in 2001, Berners-Lee, Hendler, and Lassila (<http://www.scientificamerican.com/article/the-semantic-web-overview/>) described a semantic web that “will bring structure to the meaningful content of Web pages, creating an environment where software agents roaming from page to page... will know not just that [a] page has keywords such as ‘treatment, medicine, physical, therapy’... but also that Dr. Hartman works at this *clinic* on *Mondays, Wednesdays and Fridays*”.

Schema.org is arguably one of the most practical, accessible, and successful outcomes of the semantic web movement to date. With the marketing prowess of Google, Yahoo, Bing and Yandex behind it, and with the powerful incentive of gaining additional, more inviting shelf space in the SERPs, it’s no surprise that Webmasters are adopting schema.org at a rapid pace. And Tim Berners-Lee’s words now read like a prophetic description of the search engine spiders crawling the web and extracting meaning for display in enhanced search results.

At its core, schema.org is about standardizing and simplifying the process of adding semantic markup to your web pages, and providing tangible benefits for doing so. The most visible such benefits come in the form of “rich snippets”, such as the star ratings and Price Range shown in Figure 6-????.

However, it’s clear that schema.org markup plays a larger, and perhaps expanding role in how the SERPs are constructed. Other benefits now attributed to schema.org include local SEO ranking benefits received from clearly communicating a business’s so-called NAP (name, address, phone number) information by marking it up with schema.org, and even supplying Google with information that can appear in the Knowledge Panel and “Answer Box” results (see Figures 6-39 and 6-40).

panCoast Pizza - Walnut Creek, CA | Yelp
www.yelp.com > Restaurants > Pizza > Yelp, Inc. ▾
★★★★★ Rating: 4 - 134 reviews - Price range: \$\$
134 Reviews of **panCoast Pizza** "Great pizza!! They know what they're doing, for sure.
Hand stretched dough, fresh toppings, and baked to a crispy finish."

Figure 6-39: Rich snippets from Google SERPs

Tim Berners-Lee

Computer Scientist

Sir Timothy John "Tim" Berners-Lee, OM, KBE, FRS, FREng, FRSA, DFBCS, also known as "TimBL", is a British computer scientist, best known as the inventor of the World Wide Web. [Wikipedia](#)



Born: June 8, 1955 (age 59), [London, United Kingdom](#)

Nationality: British

Parents: [Mary Lee Woods](#), [Conway Berners-Lee](#)

Awards: MacArthur Fellowship, Marconi Prize, Charles Stark Draper Prize, Mountbatten Medal, President's Medal

Books: [Weaving the Web: The Original Design and Ultimate Destiny of the World Wide Web by its Inventor](#)

Education: The Queen's College, Oxford (1973–1976), Emanuel School (1969–1973)

Figure 6-????: Google SERPs Knowledge Panel on Tim Berners-Lee

where is tim berners lee from

[Web](#) [News](#) [Shopping](#) [Images](#) [Videos](#) [More ▾](#) [Search tools](#)

About 892,000 results (0.37 seconds)



London, United Kingdom

Tim Berners-Lee, Place of birth

Figure 6-40: Google Answer Box for Tim Berners-Lee Query

Before schema.org, semantic markup was largely the province of academia, R&D, specialty niche businesses, and others with specific requirements to exchange and understand data in a deeply meaningful way. With schema.org, the local Pizza joint can hope to have “5 star reviews” jump off the search results page, local governments can publicize civic events and have that information re-presented in the SERPs, providing “instant answers” to searchers, and the list goes on. With such practical benefits in mind, and with the simplified approach of schema.org over its big brothers like RDFa, many people responsible for building web pages are making the effort to incorporate this markup into their sites.

Overview

Schema.org markup communicates the meaning of web pages to computer programs that read them, like search engine spiders. While humans can often infer the meaning of words on a page through a number of contextual clues, computer programs often need help to extract such meaning. Let’s walk through a simple example. Imagine you have a page that displays information about the book “20,000 Leagues Under the Sea”. You might create such a page with the following HTML code:

```
<div id="book">  
  <h3>20,000 Leagues Under the Sea</h2>  
    
  <h3>Author: Jules Verne</h3>  
  <h3>Rating: 5 stars, based on 1374 reviews</h3>  
  <h3>ISBN: 978-1-904808-28-2</h3>  
</div>
```

+

After marking it up, the source code might look like Figure 6-41 below. The schema.org microdata markup is highlighted, and explained below.

```
<div id="book" itemscope itemtype="http://schema.org/Book">  
  <h3 itemprop="name">20,000 Leagues under the Sea</h2>  
    
  <h3>Author: <span itemprop="author">Jules Verne</span></h3>  
  <h3>Rating: <span itemprop="aggregateRating" itemscope  
    itemtype="http://schema.org/AggregateRating">  
    <span itemprop="ratingValue">5</span> stars, based on  
    <span itemprop="ratingCount">1374</span> reviews  
    <meta itemprop="worstRating" content="1">  
    <meta itemprop="bestRating" content="5">  
  </h3>  
  <h3>ISBN: <span itemprop="isbn">978-1-904808-28-2</span></h3>  
</div>
```

Figure 6-41: Annotated schema.org markup

Line 1: itemscope – adding this to a container element, in this case a `<div>`, is the way to begin defining an entity. This attribute makes the `<div>` element the outermost, enclosing type definition for the entire book entity. The `itemtype=http://schema.org/Book` attribute, also added to the `<div>` element, declares the type of this entity. Together, this makes the entire div a container for a single book type entity.

Line 2: `itemprop="name"` – adding `itemprop` to an HTML element defines it as the container for a property. In this case, the property is the name of the book, and the value is the inner text of the `<h3>` tags, 20,000 Leagues Under the Sea.

Line 3: `itemprop="image"` – similar to the `name` `itemprop`, but the value of this property is the URL referenced in the `src` attribute of the `` tag.

Line 4: compare this to line 2. In line 2, the inner text of the `<h3>` element was our exact title. Here, we also have a label (“Author.”), which is not part of the actual author property. To keep our browser display looking the same as the original, but omit the “Author.” label from our author property, we use this construct.

Lines 5 & 6: our item property in this case is not a simple text string or URL, but rather itself another item – a schema.org/AggregateRating. It is at the same time a property of the book so it uses the `itemprop` attribute, and a type itself, so it uses `itemscope` and `itemtype` as we saw in Line 1 for our outermost book type.

Lines 7 & 8: these lines add properties for the `aggregateRating`, in much the same way we defined `name` and `author` in lines 2 and 4 above. Note the careful enclosure of the data with `span` tags, so as to include only the data itself, not the surrounding labels, in our property. This is the same technique we used in line 4 above.

Lines 9 & 10: these `itemprops` contain information needed to provide a context for the item rating (namely, that our scale is 0 to 5, with 0 being worst and 5 being best), but which is not displayed on the page. In the previous examples, the values of the properties came from the inner text of the HTML. In this case, since there is no text to display in the browser, we use the `value` attribute on the `` element.

Line 12: defines the `ISBN` property with an `itemprop`, again using a `` element to keep the display and the data cleanly separated.

Schema.org is a simple idea, with a mostly simple implementation, but getting it right can be tricky. Fortunately, the operating environment is pretty forgiving, and there are a few tools that help ease the task. Search engines understand that most webmasters aren’t structured data gurus with a deep understanding of ontologies and advanced notions of relations, entities, and other such concepts.

As such, they are generally quite adept at figuring out what you mean by your schema.org markup, even if there are errors or ambiguities in how you say it. Clearly you should strive to be accurate, and in this, but you should go into this exercise knowing that you don’t have to understand every single nuance of structured data markup, or strive for perfection in order to succeed.

How to use Schema.org

Let’s first talk about the best way to approach using schema.org. Semantic markup is designed to help you provide meaning and clarity about what your web site and each web pages on it are about, so you should be clear about this before attempting to implement Schema. Think real world tangible objects, or in semantic markup parlance, “entities”.

For example, if you’re a purveyor of fine linen, your site may have lots of pages related to pillow cases, bed sheets, duvet covers, etc. Your pages are “about” these entities. If you’re willing to make the common conceptual leap here, you could say these entities “live on” your web pages. Job one is to figure out how to map these entities to schema.org’s catalog of “types”.

At this level of thinking, schema.org is a large and ever growing and evolving catalog of “types” (schema.org documentation sometimes uses the word “items” in place of types here) that attempts to classify everything that can be represented on web pages. Let’s take a look at the schema.org page for a Book type, shown in Figure 6-42 below. The idea is straightforward. The type definition identifies the key

attributes that you would use to uniquely describe an “instance” (that is a single, real-world example) of this type.

Thing > CreativeWork > Book

A book.

Property	Expected Type	Description
Properties from Book		
<code>bookEdition</code>	<code>Text</code>	The edition of the book.
<code>bookFormat</code>	<code>BookFormatType</code>	The format of the book.
<code>illustrator</code>	<code>Person</code>	The illustrator of the book.
<code>isbn</code>	<code>Text</code>	The ISBN of the book.
<code>numberOfPages</code>	<code>Integer</code>	The number of pages in the book.
Properties from CreativeWork		
<code>about</code>	<code>Thing</code>	The subject matter of the content.
<code>accessibilityAPI</code>	<code>Text</code>	Indicates that the resource is compatible with the referenced accessibility API (WebSchemas wiki lists possible values).
<code>accessibilityControl</code>	<code>Text</code>	Identifies input methods that are sufficient to fully control the described resource (WebSchemas wiki lists possible values).
<code>accessibilityFeature</code>	<code>Text</code>	Content features of the resource, such as accessible media, alternatives and supported enhancements for accessibility (WebSchemas wiki lists possible values).
<code>accessibilityHazard</code>	<code>Text</code>	A characteristic of the described resource that is physiologically dangerous to some users. Related to WCAG 2.0 guideline 2.3. (WebSchemas wiki lists possible values)
<code>accountablePerson</code>	<code>Person</code>	Specifies the Person that is legally accountable for the CreativeWork.
<code>aggregateRating</code>	<code>AggregateRating</code>	The overall rating, based on a collection of reviews or ratings, of the item.
<code>alternativeHeadline</code>	<code>Text</code>	A secondary title of the CreativeWork.
<code>associatedMedia</code>	<code>MediaObject</code>	The media objects that encode this creative work. This property is a synonym for encodings.
<code>audience</code>	<code>Audience</code>	The intended audience of the item, i.e. the group for whom the item was created.
<code>audio</code>	<code>AudioObject</code>	An embedded audio object.
<code>author</code>	<code>Person</code> or <code>Organization</code>	The author of this content. Please note that author is special in that HTML 5 provides a special mechanism for indicating authorship via the rel tag. That is equivalent to this and may be used interchangeably.
<code>award</code>	<code>Text</code>	An award won by this person or for this creative work. Supersedes <code>awards</code> .
<code>citation</code>	<code>Text</code> or <code>CreativeWork</code>	A citation or reference to another creative work, such as another publication, web page, scholarly article, etc.
<code>comment</code>	<code>Comment</code> or <code>UserComments</code>	Comments, typically from users, on this CreativeWork.
<code>commentCount</code>	<code>Integer</code>	The number of comments this CreativeWork (e.g. Article, Question or Answer) has received. This is most applicable to works published in Web sites with commenting system; additional comments may exist elsewhere.

Figure 6-42: Schema.org definition of Book

It may help if you open this page in your web browser as we discuss it. Note that the schema.org definitions are frequently reviewed and updated based on active user feedback, so you may even see minor variations on the current page. But the overall structure will likely remain very similar, and the major elements of the page are central to schema.org. First, note the simple description, confirming that this is, indeed, the model for a book. Let's ignore the Thing > CreativeWork > Book breadcrumb for now – we'll come back to that later.

Next comes a table of “properties” – these are what we might think of as the attributes that uniquely describe our individual entity – in the example above, these are the things that describe the book 20,000 Leagues Under the Sea. Each property has a name (the Property column), an Expected Type, and a Description. The Expected Type tells us whether this property is simply a text value (like a name), or is something more complex – that is, a type itself. For example, the illustrator property should contain not the name of the illustrator, but a full person entity, using the <http://schema.org/Person> type definition (which, as you would expect, itself contains a name property, and that's where you stick the illustrator's name).

As you begin examining a possible mapping of your entities to schema.org types, you'll often encounter this nesting of types within types. While many of the properties of an entity are simple descriptions (text strings like "blue", "extra large", or even "Jan 17, 2015") others are more complex and are themselves entities in their own right. This is the concept of composing bigger scale things from a collection of smaller ones, as in describing a car (a single entity in its own right) as being made up of an engine, a chassis, wheels, interior trim, etc. (all entities themselves).

Extending this idea further: to an auto mechanic, an engine – a component of our car – is itself the composite thing (the big entity). To understand the engine in more detail, it's important to break it down into its own component entities, like carburetors, spark plugs, filters, etc.

Schema.org, then, is a set of conventions for modeling complex things in the real world, and marking them up in a way that search engines can consume, leading to both a deeper understanding of web pages by the search engines, and in turn to many current and future benefits realized when the search engines turn around and present that data back to users in compelling, contextually relevant ways.

There's one more preliminary concept we should cover, and it seems complicated at first, but isn't once we break it down. One thing you'll notice as you browse schema.org's types is that each one lives within a hierarchical family tree. We saw this earlier with the breadcrumb on the Books page, shown again in Figure 6-???? below.



Figure 6-43: Showing the “inheritance” hierarchy for the Book type

It's important to note that this kind of hierarchy, referred to among computer scientists as “inheritance”, is different than the composition hierarchy (a car made up of an engine) example we discussed earlier. The schema.org type hierarchy is a way of categorizing things from most generic to most specific. It's what we call an ontology. Its form is similar to the well-known animal kingdom charts we've all seen, or the myriad other classification schemes we all tend to take for granted – often represented on web pages with things like breadcrumbs, navigation menus, faceted navigation filters, etc.

The key thing to remember here is that when choosing the schema.org type to model your entities, it's always best to choose the most specific type you can. That is, choose schema.org/Restaurant over schema.org/LocalBusiness (if, indeed, you're operating a restaurant!). Choose Book over CreativeWork for books, and HighSchool over EducationalOrganization for High Schools. Doing so ensures you are giving the most specific information possible to the search engines, rather than settling for generic descriptions.

With that background covered, let's run through the general plan for adding schema.org markup to your web site. Here are the 6 major steps:

1. Determine the schema.org types that best describe the entities represented on your web pages, which may be different for each of your different page archetypes
2. For each page archetype you're modeling, perform a detailed mapping of the information elements displayed on the page to the schema.org type properties
3. Choose the approach you will use to express the schema.org markup
4. Edit the HTML document templates, or update the CMS settings, or modify the scripts – whatever best describes how your pages are generated – to incorporate the schema.org markup
5. Test the markup to see if your syntax is accurate, and if you've properly modeled complex entities
6. Monitor how well the search engines are consuming your structured data, and whether and how that data is being presented in the SERPs

Let's take these one at a time in more detail.

Step 1: Determine schema.org types

In this step, you think carefully about which web pages to mark up while simultaneously browsing the schema.org web site (actually, the browsing capability is fairly limited as of the time of this publication, so you might be better off searching for types – see Figure 6-44).

For example, if your web site is about community theater groups, and displays one page for each theater group, along with the upcoming list of their performances, you would begin by searching at <http://schema.org> for something like ‘theater’. The resulting page looks like what's shown in Figure 6-45. Scanning the results, we quickly spot TheaterGroup as a likely candidate for the type of our main entities.



Figure 6-44: Schema.org search box at the top of each page of <http://schema.org>

MovieTheater – schema.org

Thing > Place > CivicStructure > MovieTheater. A movie **theater**. Property, Expected Type, Description. Properties from CivicStructure · openingHours · Duration ...

<https://schema.org/MovieTheater>

TheaterEvent – schema.org

Documentation · Schemas · Home · Thing > Event > TheaterEvent. Event type: **Theater** performance. Property, Expected Type, Description. Properties from Event ...
schema.org/TheaterEvent

TheaterGroup – schema.org

Thing > Organization > PerformingGroup > TheaterGroup. A **theater** group or company—for example, the Royal Shakespeare Company or Druid Theatre.

<https://schema.org/TheaterGroup>

WatchAction – schema.org

Schema.org is a set of extensible schemas that enables webmasters to embed structured data on their web pages for use by search engines and other ...
schema.org/WatchAction

Full Hierarchy – schema.org

The Type Hierarchy. Here is the entire hierarchy in a single file. Note: as of 2014- 04-04 this tree is not entirely up to date. Additional types have been added: see ...

<https://schema.org/docs/full.html>

WebPage – schema.org

A web page. Every web page is implicitly assumed to be declared to be of type WebPage, so the various properties about that webpage, such as breadcrumb ...
schema.org/WebPage

Figure 6-45: Search results for 'theater' on <http://schema.org>

Taking a closer look at the TheaterGroup page at [http://schema.org/TheaterGroup](https://schema.org/TheaterGroup) (figure 6-????), we can see a few of our core concepts at work:

- a) TheaterGroup is part of a logical hierarchy, starting with the most generic type ("Thing" – actually the top-most ancestor of all schema.org types), then proceeding to more and more refined types: Organization, PerformingGroup, TheaterGroup.

- b) A TheaterGroup is composed of many elements (called “Properties”), some of them simple like the name of the group, some of them actual types in their own right (such as address, aggregateRating, employee, etc.). Examining the list of properties, it confirms our belief that this is the best type for describing our local theater entities on our web pages.

Thing > Organization > PerformingGroup > TheaterGroup

A theater group or company—for example, the Royal Shakespeare Company or Druid Theatre.

Property	Expected Type	Description
Properties from Organization		
address	PostalAddress	Physical address of the item.
aggregateRating	AggregateRating	The overall rating, based on a collection of reviews or ratings, of the item.
brand	Brand or Organization	The brand(s) associated with a product or service, or the brand(s) maintained by an organization or business person.
contactPoint	ContactPoint	A contact point for a person or organization. Supercedes contactPoints.
department	Organization	A relationship between an organization and a department of that organization, also described as an organization (allowing different urls, logos, opening hours). For example: a store with a pharmacy, or a bakery with a cafe.
dissolutionDate	Date	The date that this organization was dissolved.
duns	Text	The Dun & Bradstreet DUNS number for identifying an organization or business person.
email	Text	Email address.
employee	Person	Someone working for this organization. Supercedes employees.
event	Event	Upcoming or past event associated with this place or organization. Supercedes events.
faxNumber	Text	The fax number.
founder	Person	A person who founded this organization. Supercedes founders.
foundingDate	Date	The date that this organization was founded.
globalLocationNumber	Text	The Global Location Number (GLN, sometimes also referred to as International Location Number or ILN) of the respective organization, person, or place. The GLN is a 13-digit number used to identify parties and physical locations.

Figure 6-46: TheaterGroup type from <http://schema.org>

It's during this step that you want to deal with the question "what is this page about?", and choose the schema.org type that best describes the overall contents of the page. Often this choice is obvious, but at times it can be tricky. For example, on a page with a product for sale, should you choose schema.org/Offer or schema.org/Product to model the page?

Examining both pages on the schema.org site, you can see that an Offer has a property called itemOffered, with an expected value of Product. What that means is that you can describe the contents of the page as an Offer (schema.org's concept of something for sale), where the item for sale (the Product) is contained within the Offer, using the itemOffered property.

Alternatively, you could use the Product type, which has a property called offers which can, as you might expect, contain one or more Offer types. The decision probably depends on the overall purpose of the page. If the page is a detailed product page, describing many attributes of the product, and the offer information is just a piece of that, it probably makes sense to model the page as a Product, and include the offer information in the itemOffered property. However it's not out of the question that you could invert this model.

Either of the approaches to the Product/Offer model is valid, as both convey the meaning that you want. But take another look at Offer. You can see that it is a complex concept, and has many properties that are themselves types (for example, aggregateRating). Other complex nesting of types and attributes can easily arise, and it's important to model these out in a way that best matches the meaning of the page. The best approach often won't be obvious at this stage of analysis, so you may need to revisit your thinking after you complete step 2 of the process.

Step 2: Map schema.org properties to elements on the web page

The first step here is to survey the various data elements displayed on the web page, and match them up with the appropriate schema.org types and properties you selected in Step 1. In this step, you may discover relationships that resolve some of the potential ambiguities from Step 1.

For example, continuing the Product/Offer discussion, let's assume that one of the items displayed on the page is an overall rating – say a value on a scale of 1 to 5 – representing user evaluations of the product. We notice that both Product and Offer have a property called aggregateRating, so this hasn't quite settled our debate on which type to model the page on.

Let's also assume that we display several different prices – perhaps for new or used versions of the product, or with different shipping options or different currencies. It now starts to become obvious that we should model the entire page as a Product which contains multiple Offers and a single aggregateRating that applies to the Product itself. Finally this is starting to take shape!

You might notice that there are properties defined on the schema.org type that you're not currently displaying to browsers, but which you have access to. Continuing with our Product example, perhaps your web application's database stores the MPN (manufacturer's part number), but you don't choose to display that on the page. What should you do?

Ideally, you want a very high degree of consistency between what you mark up and what's visible to "normal users" via web browsers. Technically, there are mechanisms that allow you to communicate meta data about your entities to the search engines that shouldn't be displayed to users (we saw this earlier in our aggregateRating example, and we'll explore that example a bit more momentarily).

However, it's important to use these mechanisms sparingly, and not be tempted to stuff a lot of extra data into the schema.org markup that is not visible to human users. In the MPN case, our choice should be between adding this as a visible element on the page (and then of course adding it to our schema.org markup), or foregoing it entirely. As you think about this, it should become clear that marking up a lot of data that is not displayed to the user is conceptually something a spammer might do, and for that reason the search engines frown on it.

What are the valid reasons for marking up non-displayed data? Usually it's because you need to convey some different context that is obvious to people, but not to search engine spiders. An example of this is that when you mark up an aggregateRating, you're strongly encouraged to specify the scale. That is, if you display 4 stars for a review on a scale of 0 to 5, this is usually quite clear in the visual representation, but it needs to be stated explicitly in the schema.org markup. Thus, aggregateRating entities have worstRating and bestRating properties, and we want to supply the values 0 and 5 respectively, corresponding to our star rating scale. We saw this in the sample code for our book at the beginning of the chapter.

Upon completing this step, you should have a complete mapping between the data displayed on the page and the various schema.org types and properties that make up your model. Your model may be simple or complex with multiple levels of nesting. It's best to make all these decisions before you begin actually implementing schema.org on the page.

Step 3: Choose your implementation technique

For most people, this means "go mark up the page". Sounds simple, right? And for some pages – especially pages which are template driven with mostly static data, it should be fairly straightforward. Or, if you're lucky enough to be using a Content Management System or Publishing Platform that has built in support for schema.org, most of the actual implementation may be done by setting a few configuration parameters.

For other more dynamic sites that generate their pages through a complex pipeline of page generation programs, tweaking things to insert the right tags in the right place can be far more difficult. And for these types of sites, validating that the generated schema is correct is also challenging, as the schema.org markup may be sporadically injected into kilobytes of code.

The primary implementation technique is going to be to edit templates and/or modify page generation programs to insert the microdata markup as needed to produce the desired final output. The key thing is to

have a clear mapping of the model from Step 2 showing the final desired HTML with microdata markup inserted, and use this to validate that the final page produced by the web server matches the model. As we'll see in Step 5, there are some tools that can help with this verification as well.

For those who don't have access to the code or back-end systems, or want a simpler approach, Google offers the Structured Data Markup Helper, as part of Google Webmaster Tools. This is a proprietary Google tool that allows you to annotate a page, using a point-and-click editor. See Figure 6-?????. It's actually just an alternative way of providing the same data you provide via schema.org microdata markup, but you are instead feeding it directly to Google and do not change the page source code at all.

So why doesn't everyone just do this? There are two good reasons why this isn't the best fit for everyone. First, the information is only available to Google, and not to other schema.org-aware search engines (or other applications that may make use of schema.org markup). Second, as is often the case with this kind of tool, the visual editor is more limited than the markup syntax in its ability to express rich and complex information.

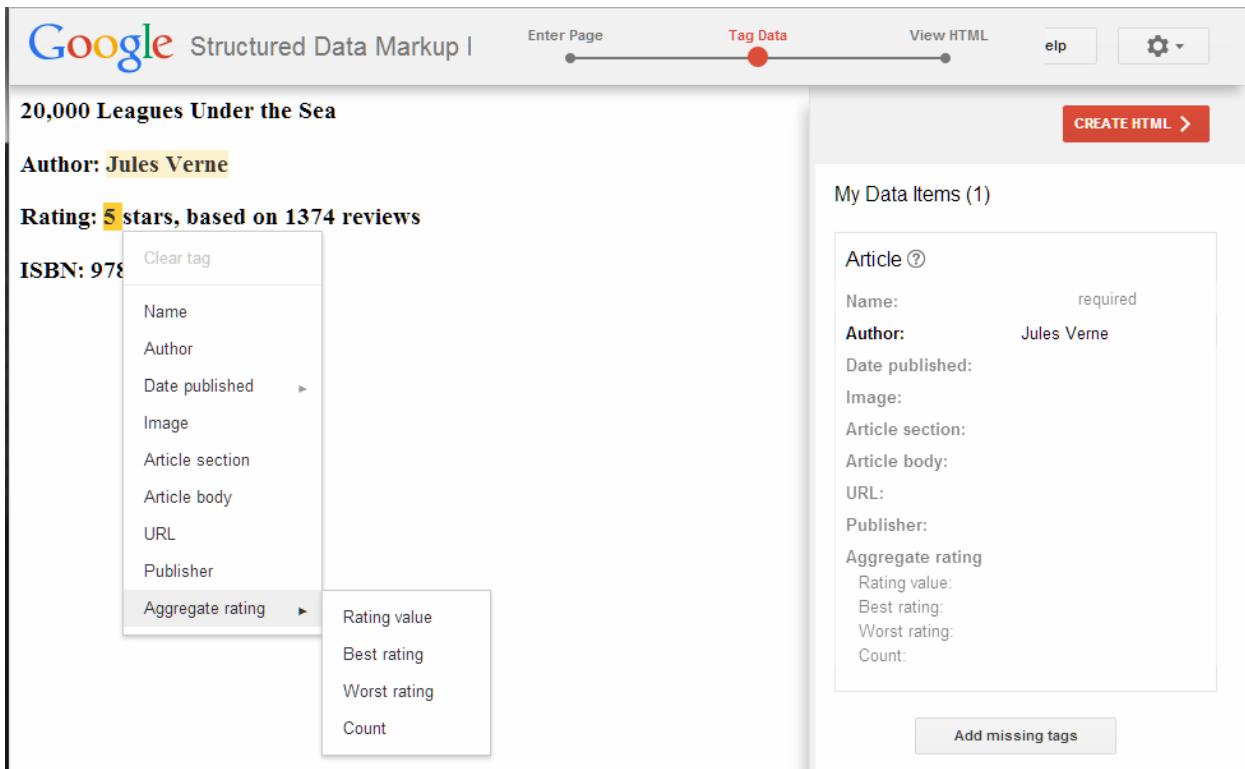


Figure 6-47. Google Structured Data Markup Helper

Looking to the future, another alternative may be on the horizon, and is something to keep an eye on. Google and others are already beginning to make use of a format known as JSON-LD for expressing schema.org markup. This format is showing up in limited, specialized circumstances, but it seems apparent that JSON-LD may soon become a full-fledged alternative to microdata for all schema.org expression.

The beauty of JSON-LD is that it provides a way to isolate all of the schema.org information into a single string of code, rather than expressing it by embedding markup within the HTML document itself. This has the possibility of solving many of the more complex issues associated with implementing schema.org on complex, dynamic sites.

Step 4 – Implement the changes to generate the target schema.org code

This is really just saying “now it’s time for your web developers to go breathe some life into your creation”; that is, go get these pages served up by your web server! This is where the content management system is tweaked, the templates are updated, the page production programs are modified, etc.

Step 5 – Test

When you reach this stage, your web server is shooting out bundles of HTML with tidy little microdata tags embedded in it that add meaning and structure to your data. At this point, the generated schema.org markup code should be syntactically correct, and should express the right model – that is, the whole composition of properties and types into larger properties and types to end up accurately modeling the information displayed on our pages. Of course it’s important to verify this.

The hard way to do that is to examine the generated code by hand, looking for the opening and closing tags, and ensuring that all the data is there, nested properly. Fortunately, there’s an easier way (though you should still be prepared to roll up your sleeves and dig into the code to debug potential problems).

That better way is to use one or more of the tools available to verify your schema.org microdata markup. Perhaps the best known of these tools is Google’s Structured Data Testing Tool

(<https://www.google.com/webmasters/tools/richsnippets>). The Structured Data Testing Tool is an elegant utility that examines your page (either directly, by supplying a URL, or alternatively by cutting/pasting HTML source code) and gives you feedback on the structured data it finds. Figure 6-???? shows such a result.

The screenshot shows the Google Structured Data Testing Tool interface. On the left, a sidebar includes links for Home, Structured Data Testing Tool (which is active and highlighted in red), Structured Data Markup Helper, Help with: Troubleshooting, and Help Center. The main content area has tabs for URL and HTML, with HTML selected. A large text area contains the following schema.org code:

```


<h3 itemprop="name">20,000 Leagues under the Sea</h3>

<h3>Author: <span itemprop="author">Jules Verne</span></h3>
<h3>Rating: <span itemprop="aggregateRating" itemscope
itemtype="http://schema.org/AggregateRating">
<span itemprop="ratingValue">5</span> stars, based on
<span itemprop="ratingCount">1374</span> reviews
<meta itemprop="worstRating" content="1">
<meta itemprop="bestRating" content="5">
</h3>


```

Below this is a 'PREVIEW' button. The preview section shows a snippet of a Google search result card with the book's title, a 5-star rating, and a snippet of the book's content. It also includes a note about not showing full text from the page. The 'Google search results' tab is active.

The 'Extracted structured data' section shows the following items:

Item

type:	http://schema.org/book
property:	
name:	20,000 Leagues under the Sea
image:	http://www.example.com/20000-leagues-under-the-sea-book-cover.png
author:	Jules Verne
aggregaterating:	Item 1
isbn:	978-1-904808-28-2

Item 1

type:	http://schema.org/AggregateRating
property:	
ratingvalue:	5
ratingcount:	1374
worstrating:	1
bestrating:	5

Figure 6-48. Google's Structured Data Testing Tool Output

The output of this tool has a bit of an arcane formatting convention. Note that our book shows up as the first Item. The book has a number of properties, and among those is our aggregateRating – recall that this

is itself another type. When this composition or nesting occurs properly, you see the output shown in Figure X above.

Here, the nesting relationship is shown by the *Item 1* in the item value field for the aggregateRating property of the Book, followed immediately by the Item 1 output. So the value Item 1 in the first field ties together with the name of the second item shown, and shows that Item 1 – the rating – is properly contained within the book entity, as specified by the schema.org definition for a Book. Google will keep incrementing these numbers for as many embedded types you have on the page.

Summary

We have seen that schema.org is a standard for providing search engines (and potentially other applications) with structured data describing the meaning of web site content. The notion of data structuring is actually quite intuitive, and maps well to the way we commonly categorize things, as with product catalogs, biology, library card catalogs, and many other collections of related items. This intuitive, webmaster-friendly approach has led to rapid adoption of schema.org by the webmaster and content production communities. Currently, the most common way to structure data with schema.org is to add Microdata markup to HTML documents. Search engines use this data to extract meaning, and enrich SERPs with rich snippets, “answer boxes”, and “knowledge panels”, providing a more relevant and deeper search result. Implementing schema.org can bring these benefits to both users and publishers today, and can help set the stage for publishers to gradually delve more deeply into the emerging world of semantic search in the coming years.

A special thanks to John Biundo for his contributions to the Schema portion of this chapter.

Google Authorship and Author Authority

One of the most interesting insights into the mind of Google, as it were, was the three year experiment known as Google Authorship.

Google Authorship was a program that allowed online authors to identify and verify their content with Google. This was accomplished by a two-way link between the author’s content across the web and his Google+ profile.

While the Authorship program was officially discontinued by Google on August 28, 2014, it is likely that Google’s interest in the value of the authority, trust, and reputation of an author in a given topical area is undiminished.

A Brief History of Google Authorship

The roots of Google Authorship lie in a patent originally granted to Google in 2007 titled “Agent Rank” (<http://www.google.com/patents/US20070033168>). The patent described methods whereby a search engine could identify distinct “agents” (one of which could be the author or authors of a web document) and then assign a score to that agent that could then be used as a factor in search rankings.

Google didn’t appear to do anything with this patent until June 2011. In a blog post that month (<http://googlewebmastercentral.blogspot.com/2011/06/authorship-markup-and-web-search.html>) Google’s Othar Hansson announced that they would begin to support the use of the HTML5 standard rel=author and the XFN standard rel=me, and that webmasters could use that markup to identify authors and author profiles on their sites.

The next major step in Authorship came just 21 days later when Google unveiled its new social network, Google+. Google+ provided personal profiles that Google could use to verify authors using the rel=author markup.

This intention was confirmed in a YouTube video by Othar Hansson and Matt Cutts published on August 9, 2011, titled “Authorship Markup” (<https://www.youtube.com/watch?v=FgFb6Y-UJUI>). In the video

Hansson and Cutts explained that Google wanted web authors to have Google+ profiles, and that they should link from the “Contributor To” link sections of those profiles to each domain where they publish content. Over time, Google offered several options by which the publisher could confirm the relationship by linking back to the author’s Google+ profile.

In that video Google confirmed that there could be rewards to authors who implemented Authorship markup: the immediate possible benefit was the potential for an author’s profile image and byline name to be shown with search results for their content. Figure 6-???? is typical of such results.



Figure 6-49. Rich Snippet Authorship Result

Additional potential benefits mentioned by Hansson and Cutts were increased search rankings and that Google might be able to use Authorship to identify the original author or a piece of web content, thus giving that author’s copy precedence in search over scraped copies.

Over time Google added several tools and features to make Authorship easier to implement and more useful for authors and publishers. This was probably the result of the problems they saw with a lack of adoption of this markup.

The first major hint that Google might be pulling back on its Authorship experiment came in October 2013 when AJ Kohn revealed that Othar Hansson had left the Authorship team and was not being replaced (<http://www.blindfiveyearold.com/authorship-is-dead-long-live-authorship>). In that same month, Google’s Matt Cutts revealed that Google would soon be cutting back on the amount of Authorship rich snippets shown in search, as they had shown in tests that doing so improved the quality of those results.

Cutts’s promise was kept in December 2013 when observers noticed a 15% reduction in the amount of author photos were being shown for most queries (<http://searchengineland.com/confirmed-google-reduces-authorship-rich-snippets-in-search-results-180313>).

In June 2014 Authorship was further reduced in search as Google announced that it would no longer show author photos in results, just bylines. The only announced reason for this was to bring its mobile and desktop user experiences more into sync.

However, only two months later, as previously noted, Google announced the complete removal of Authorship data from search, and stated that they would no longer be tracking any data from rel=author links. The Google Authorship program, at least any program based on rel=author links and showing rich snippets in search results, was now over.

Why Did Google End Support for rel=author?

In his official announcement of the ending of the Authorship program, John Mueller of Google Webmaster Central said, “Unfortunately, we’ve also observed that [Authorship] information isn’t as useful to our users as we’d hoped, and can even distract from those results. With this in mind, we’ve made the difficult decision to stop showing authorship in search results” (<https://plus.google.com/u/0/+JohnMueller/posts/HZf3KDP1Dm8>).

He went on to elaborate that this a decision based on user experience concerns. After three years of testing, Google was no longer seeing any particular user benefits from showing Authorship results. Mueller said that removing the Authorship results “did not seem to reduce traffic to sites.” It would seem then that searchers were no longer viewing these results as anything special.

What else may have factored into the decision to stop showing Authorship results? In his post Mueller mentioned that he knew that implementing authorship “wasn’t always easy to implement.” Could it be that

low implementation rates by most sites fed Google's decision? If Google was ever going to rely on Authorship as a signal for search, they would need to have data from a wide variety of sites.

In a study published just after the ending of Authorship, Eric Enge confirmed from a sampling of 150 top publishing sites that Authorship implementation was indeed low (see <http://stonet.co/AuthorshipStudy>). He found that 72% of these sites had attempted Authorship markup in some way, but out of those *nearly three-fourths had errors in their implementation*. But even worse, 71% of the 500 authors sampled from those sites had done nothing from their side to implement Authorship.

It would seem that low participation might be another reason behind Google's decision. Google may have learned that data you want to use as a ranking factor can't be dependent upon voluntary actions by webmasters and authors.

Is Author Authority Dead for Google?

Does the end of rel=Author-based Authorship mean Google has abandoned all interest in understanding, tracking, and making use of data concerning the authority levels of online authors? Most likely not.

For one thing, on September 2, 2014, Google was granted a patent for a system that would retrieve, rank, and display in search authors considered authoritative for a topic based on their relationship (in social networks) to the searcher (<http://www.seobythesea.com/2014/09/google-decided-authoritative-query/>).

Also, Google spokesperson Matt Cutts has often spoken during the last year of Google Authorship about his interest in and support for Google eventually being able to use author reputation as a means of surfacing worthwhile content in search results, but that he sees it as a long-term project (<http://www.stonetemple.com/googles-matt-cutts-understanding-social-identity-on-the-web-is-hard/>).

While Cutts seems to be voicing his personal opinion in such statements, it is doubtful that he would speak so frequently and positively about the topic if it weren't actually active at Google.

Another area that seems to support that Google will only increase its interest in author authority is semantic search. Semantic search involves, in part, a dependence upon the identification of various entities and the ability to understand and evaluate the relationships they have to each other. As the original Agent Rank patent makes clear, authors of web content are certainly a useful type of entity.

Google understands that real people often evaluate the authority and trustworthiness not just by a document's contents or what links to it, but by the reputation of the author. Semantic search at its simplest is a quest to enable Google's search algorithm to evaluate the world more closely to the way people do. So it makes sense that Google would continue to pursue the ability to evaluate and rank authors by the trust and authority real people place in them for a given topic.

Google+ Authors in Personalized Search

At the time of this writing there remained one large and very interesting exception to Google's elimination of Authorship rich snippets from search. Author photos and bylines can still appear for Google+ content authored by people in a searcher's Google network (Google+ circles and Gmail contacts) when that searcher is logged in to his or her Google+ account while searching. Figure 6-???? shows such a result.

Get the Skinny: Google Authorship May Be Dead, But Auth...



<https://plus.google.com/.../posts/J9LTUuqYoPz> ▾

Ana Hoffman

Aug 30, 2014 - Google ended its three-year experiment with Google

Authorship yesterday, but the use of Author Rank to improve search results will continue - says +Danny ...

Figure 6-50. Personalized Search Rich Author Snippet

Notice that the URLs are both from plus.google.com (Google+). The person performing this search has both Ana Hoffman and Mark Traphagen in his Google+ circles, and is searched while logged in to his Google+ account. For the search query “Google authorrank” Google found that these two people in the searcher’s network had relevant content on Google+ and so included it in the searcher’s results.

Take note of two points about these results:

1. These results are uniquely ranked higher for this individual searcher. If she searches for the same query while logged out of her Google+ account, these results will not show up in the top results.
2. The personal connection of the authors to the searcher is being emphasized by the photo and byline.

The Future of Author Authority at Google

It appears that Google remains interested in the concept of author authority as a factor in search rankings. Google is likely working on methods to identify and evaluate authors and their content that are not dependent on human publishers and authors placing links and attribution tags. When those methods are providing reliable data, Google might make these signals a ranking factor.

However, given the lessons of the first Google Authorship experiment, we might expect the following possible differences:

- Author authority might be more personalized. That is, it may be that Google will give a greater boost to content by authoritative authors relevant to your search *if* you have some connection to or relationship with those authors.
- Author authority in search will probably be less obvious. Google may not return to the practice of displaying rich snippet profile photos for top authors, in part because Google is moving away from flashier rich snippets in general as part of its Mobile First initiative. It is therefore likely that any future author authority factor will simply be folded into the many factors that determine search rankings and may not be apparent to the searcher.

Building Author Authority

Here are some tips on how to build author authority.

Publish with real names. In order to build author authority search engines have to be able to recognize that multiple pieces of content are connected with a particular individual. Several of the tips below relate to building your personal authority both on and offline, so using your real name with your content is important.

Keep your name consistent. In parallel with the previous tip, it is important that you use exactly the same name as the byline on all your content as well as in all your social profiles. That will help search engines gain confidence about your identity, and make it more likely that all or your online content will be used to evaluate your authority.

Cross-link your profiles. Wherever possible, create links between all your online profiles. This is another way to help search engines have more confidence as to your unique identity.

Link your social profiles to your content. Wherever possible, create links from your social and site profiles to the sites on which you publish content. Of course, in the case of Google, it is most important to make sure that all sites on which you publish are linked from the “Contributor to” section of your profile links. Even though Google says it no longer tracks data based on rel=author links to Google+ profiles, we still recommend creating links from your content or site author profiles back to your Google+ profile, as these still may help Google have confidence about content that should be identified with you.

Produce content about all aspects of your field. More and more we see indications that Google is including in measures of its topical authority how complete and well rounded the content is. It’s no longer effective to merely hammer away at certain long tail keywords. You need to build contextually rich content that looks at your subject from all sides. That doesn’t just apply to individual content pieces, but also to the content across an entire site or across your profile as an author on many sites.

Produce content that goes in depth on specifics of your field. As well as covering all aspects of your area of expertise, your content also needs to explore those areas deeply. That doesn’t mean every piece of content needs to be an academic paper, or even long form. But you should be seeking as often as possible to be producing content that gives a unique perspective on a topic, or that goes into more depth and detail than most other similar pieces on the web.

Cultivate an audience. Every content producer has to be as concerned with building a loyal audience as they are with producing quality content. That means being active on social networks, for one. Seek to build good relationships with those who might be interested in your expertise and likely to share it with their networks.

Participate in relevant conversations. Go beyond just broadcasting your content to participating in relevant online conversations and communities. Doing that can have multiple benefits. As you contribute to such communities, you get a chance to display your expertise before a broader audience, some of whom may start to follow you. That means you are growing your audience (see above), but doing it in places where you are more likely to pick up followers with high interest in what you do.

Don’t forget real world opportunities. Attending conferences and networking events in your field can lead to online connections that help reinforce your online authority. This is especially true if you are a speaker or panelist at such events, or get interviewed by a media outlet. You can accelerate this effect by actively inviting people at these events to connect with you online. For example, always place your primary social profiles prominently in any presentations you do.

Incubate and promote brand Subject Matter Experts. Publishers should not ignore the power of individual topical authority used in conjunction with their brands. Many companies are reluctant to empower individual employees or representatives to build their own authority, but they miss a real opportunity by not doing so. People identify with, trust, and connect with a real individual long before they will give the same to a faceless brand. Therefore wise brands will cultivate Subject Matter Experts (SME) who have a direct connection with their brand, knowing that the audience and authority those SME’s build will ultimately reflect back on the brand.

Google’s Publisher Tag

Although Google has ceased to track author data using rel=author links, the related rel=publisher link tag is still very much supported, and can convey a number of benefits to brands using it.

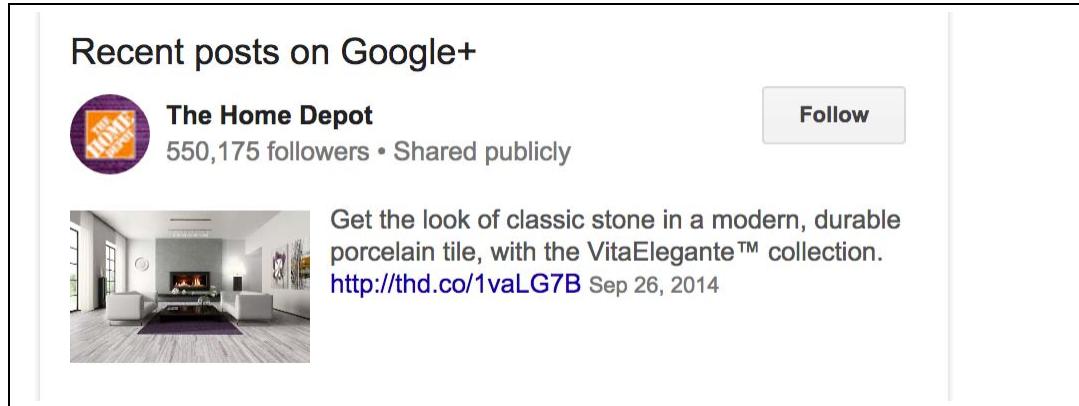
Similar to Facebook and some other social networks, Google+ allows the creation of Pages as distinct from Profiles. While profiles are intended solely for individuals, Pages allow non-personal entities (brands, organizations, bands, etc.) to have a presence on Google+.

Google+ Brand Pages deserve particular attention from any online brand, however, because of the way Google uses them, particularly in conjunction with Google search. Stated succinctly, a Google+ page linked from a brand’s official site using rel=publisher is the easiest and most direct way for Google to verify the

unique identify of a particular brand. Once Google has such verification, they can begin to make use of data related to the brand in various ways.

Brands with verified Google+ Pages can be eligible for three special search features.

- **Branded Knowledge Panel Result.** As shown in Figure 6-51, when people search for a brand name, brands with verified Google+ Pages may show a special box in the right panel of search.

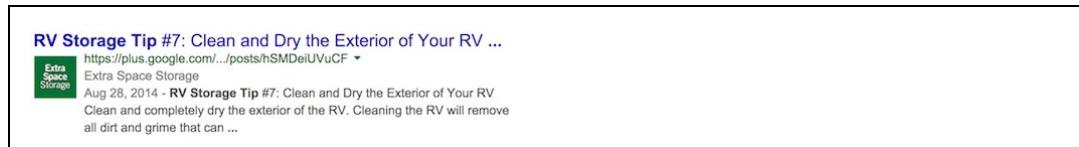


The screenshot shows a search result for 'The Home Depot'. A 'Recent posts on Google+' box is displayed on the right. It features the brand's logo, 'The Home Depot', and the text '550,175 followers • Shared publicly'. A 'Follow' button is present. Below this, a post from 'The Home Depot' is shown with the text: 'Get the look of classic stone in a modern, durable porcelain tile, with the VitaElegante™ collection.' and a link 'http://thd.co/1vaLG7B'. The date 'Sep 26, 2014' is also visible.

Figure 6-51. Rich Snippet Authorship Result

The Google+ knowledge panel box shows the brand name and logo, a Google+ follow button (shows only if the searcher is logged in to Google+), the Page's follower count, and a recent Google+ post.

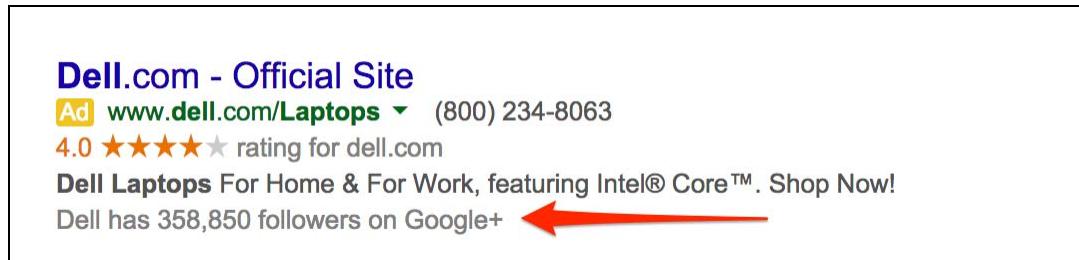
- **Elevated Rich Snippet Google+ Post Results.** When a searcher has a brand circled on Google+ and searches while logged in to her Google+ account, she may have a relevant Google+ post from that brand elevated to page one and highlighted with a brand logo and brand name rich snippet. An example of this is shown in Figure 6-52.



The screenshot shows a search result for 'Extra Space Storage'. A 'Rich Snippet Authorship' box is displayed on the right. It includes the brand's logo, the post title 'RV Storage Tip #7: Clean and Dry the Exterior of Your RV ...', the link 'https://plus.google.com/.../posts/hSMDeiUVuCF...', and the date 'Aug 28, 2014'. Below this, a snippet of the post content is shown: 'RV Storage Tip #7: Clean and Dry the Exterior of Your RV. Clean and completely dry the exterior of the RV. Cleaning the RV will remove all dirt and grime that can ...'

Figure 6-52. Rich Snippet Authorship Result

- **AdWords Social Extensions.** If a brand connects their Google+ Page to their AdWords account and enables Social Extensions, Google may add a Google+ annotation to the brand's ads in search. For more info see <http://www.google.com/ads/innovations/socialextensions.html>. Figure 6-53 shows how this appears in the search results.



The screenshot shows a search result for 'Dell.com - Official Site'. An AdWords Social Extension is displayed on the right. It includes the brand's logo, the ad text 'Ad www.dell.com/Laptops (800) 234-8063', a 4.0 star rating, the text 'Dell Laptops For Home & For Work, featuring Intel® Core™. Shop Now!', and the follower count 'Dell has 358,850 followers on Google+'. A red arrow points to the follower count.

Figure 6-53. AdWords Social Extension Example

- **Gmail Google+ Related Pages Widget.** A brand with a verified brand page that meets certain qualifications can have a widget show in the right sidebar of Gmail when customers open an email from the brand (see Figure 6-54 For an example of this appears). The widget displays a thumbnail of a recent post from the brand's Google+ Page and, if the user has a Google+ account, a follow button. For details see <https://support.google.com/business/answer/4569086?hl=en>.

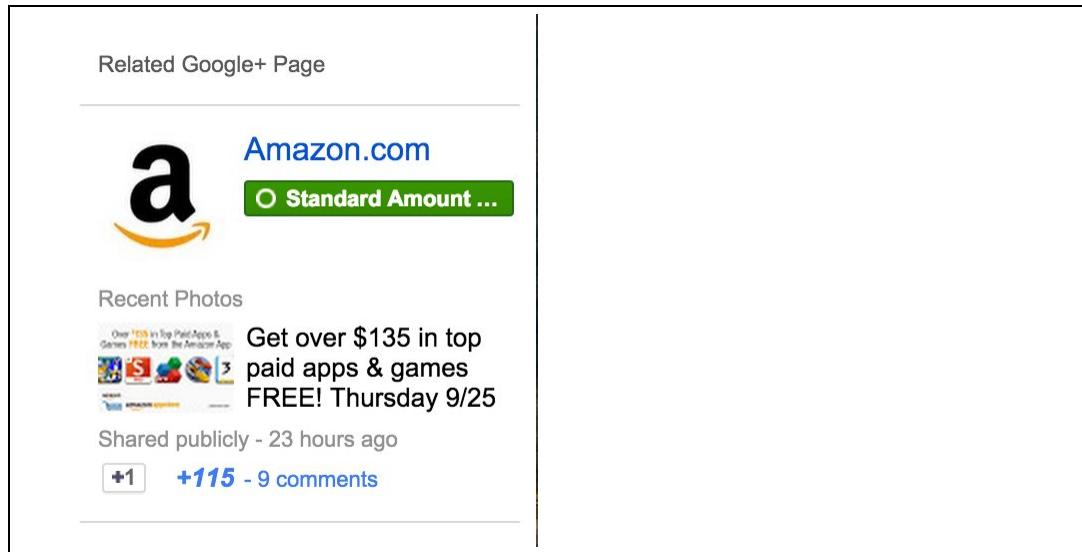


Figure 6-54. Gmail Google Plus Page Widget

Verifying a Google+ Brand Page. For non-local business Pages, a Page can be verified by simply making the brand's official website the main URL of the Page, and then linking back from the home page of that site to the Google+ Page with a rel=publisher attribute. A local business Page must verify via Google My Business (<https://www.google.com/local/business/?hl=en>). If your business has 10 or more locations to verify, use <https://www.google.com/local/manage/?hl=en#>.

A special thanks to Mark Traphagen (<http://www.stonetemple.com>) for his contributions to the Authorship and Publisher Tag portions of the chapter.

Google's Knowledge Graph and the Knowledge Vault

The face of search is changing in significant ways. The latest incarnation of that evolution is the Knowledge Graph. Google has also begun to communicate about the Knowledge Vault, though that is just a research project as of October 2014. To get some perspective on why these are important, it is useful to review part of the history of how search has evolved.

Overview of Changes in Search Complexity

Search results used to be built by analyzing the text strings found by search engines on the pages of the web. The resulting pages were quite useful, but the presentation of the results was quite simply, such as that you see in Figure 6-55.

Google

wojohowics

Web Maps Shopping Images Videos More Search tools

8 results (0.66 seconds)

Did you mean: [wojohowitz](#)

Ty Wojohowics Profiles | Facebook
<https://www.facebook.com/public/Ty-Wojohowics> Facebook
View the profiles of people named Ty Wojohowics on Facebook. Join Facebook to connect with Ty Wojohowics and others you may know. Facebook gives ...

Ty Wojohowics | Facebook
<https://www.facebook.com/ty.wojohowics>
Ty Wojohowics is on Facebook. Join Facebook to connect with Ty Wojohowics and others you may know. Facebook gives people the power to share and makes ...

Wojohowics | Facebook
<https://www.facebook.com/family/Wojohowics/1>
Facebook users with the last name "Wojohowics". Search Facebook for more people who share the last name "Wojohowics".

Wojohowics | Facebook - bei Facebook
<https://de-de.facebook.com/family/Wojohowics/1?...1> Translate this page
Facebook-Nutzer mit dem Nachnamen "Wojohowics". Durchsuche Facebook nach weiteren Personen, die den Nachnamen "Wojohowics" haben.

What Is a Official Language of Uganda? - Zortman
zortman.vidhut.biz/2014/09/01/what-is-a-official-language-of-uganda/
Sep 1, 2014 - Stanley Wojohowics, from 'Barney Miller', was also a Polish American character. The impression of Mimi Bobeck, from 'The Drew Carey Show', ...

DJ Listings - The Stranger
www.thestranger.com/seattle/dj-listings/Content?oid=4342 The Stranger
MANZINGER V--Hot Tuesdays Makin' Love: DJ Tom "da Bomb" Wojohowics (phat beats), free. MARCUS' MARTINI HEAVEN--Subtropic: DJ XCO.

S2 (1998)
www.users.on.net/~fsteiny/goodguysbadguys_s1998.html
Leo Swift Handy Andy Dermott Maginnis Bogden Wojohowics Bobby Shearer Mrs Kirkwood Mort the butcher. Waldo Charlie Earp Barman, Tammy MacIntosh

Figure 6-55. Simple Text Only Search Results

Over time, these types of search results became known as “10 blue links” due to their simplicity. As they became more sophisticated, the engines figured out how to incorporate more type of media into the results. This included videos, images, news stories, shopping results, and more.

These are typically referred to as “Blended Results”, and you can see an example of these in Figure 6-56. You can also read more about these types of results in Chapter 10.

Figure 6-56. Example of Blended Search Results

The emergence of Blended Search was a big step forward for search engines, but it was only one step on a longer journey. The commitment to that journey is well defined in Google's mission statement (<https://www.google.com/about/company/>): "Google's mission is to organize the world's information and make it universally accessible and useful."

The next step in that journey for Google was the Knowledge Graph (<http://www.google.com/insidesearch/features/search/knowledge.html>). This was a Google initiative designed to allow them to leverage structured databases to enhance the search results. This initiative allowed Google to further enhance the presentation of their results. An example of this is shown in Figure 6-57.

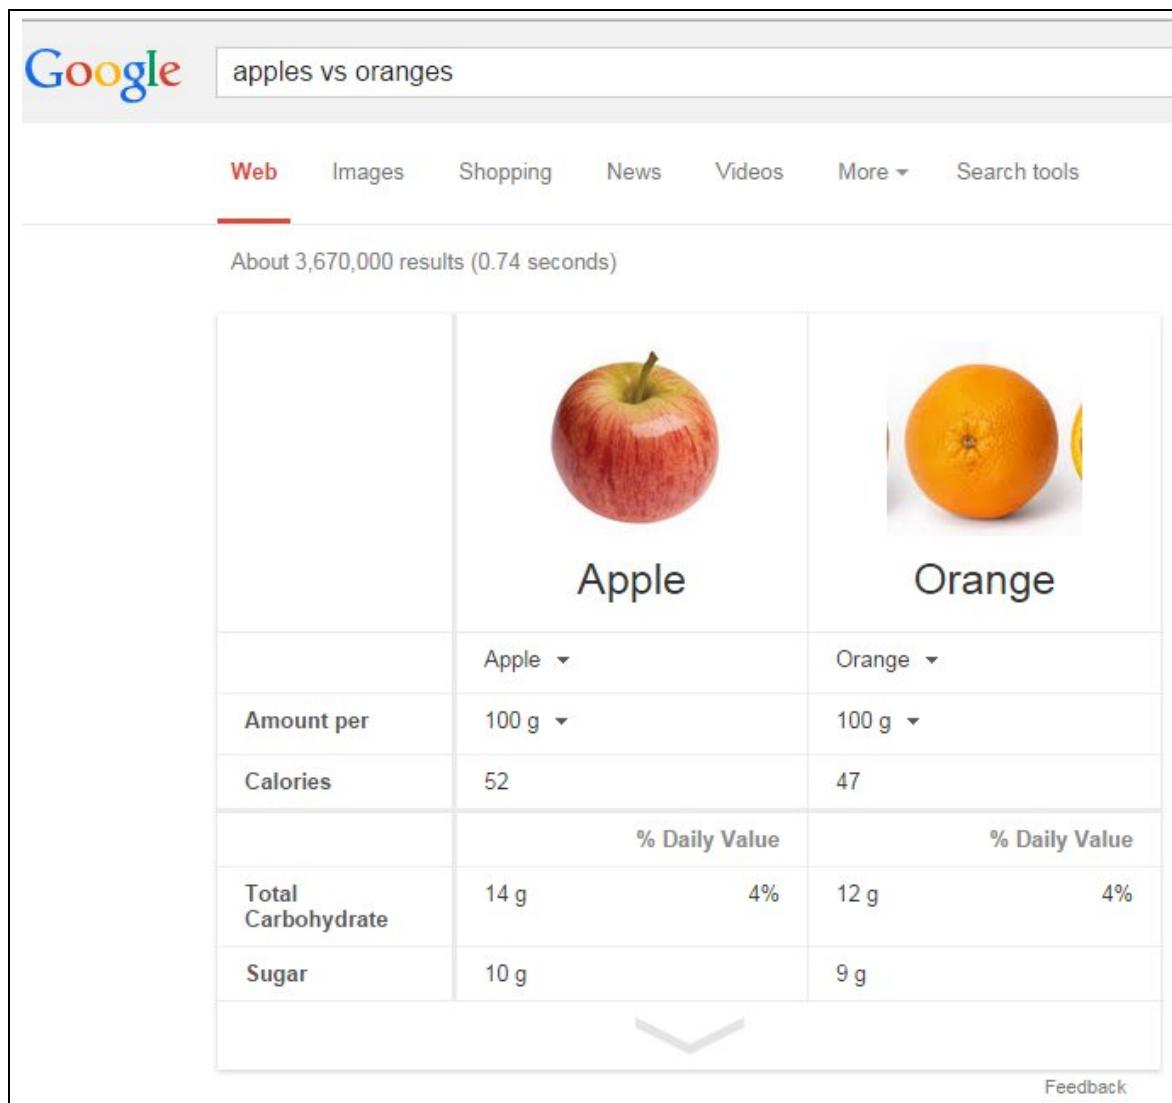


Figure 6-57. Sample Knowledge Graph Result

In essence, the Knowledge Graph was a major step forward by Google that allowed it to start showing complete answers in the search results. The information for these results is typically retrieved from Freebase (<https://www.firebaseio.com/>), a community edited database of information.

The basic concept is sometimes referred to as “moving from strings to things”. The search engines that returned nothing but 10 blue links were comparatively quite simple, as they relied on scanning the text on a web page to figure out what it was about, and did not understand relationships.

In comparison, the Knowledge Graph can understand that apples and oranges are both fruits, and they have properties, such as calories, carbohydrate levels, and grams of sugar. Or, that the Empire State Building has height, a construction date, and initial architect, and that Google has access to pictures of it.

These types of data sources provide a rich array of information. It is estimated that they allow Google access to information on 500 million entities and 3.5 billion pieces of information. Stone Temple Consulting performed an extensive study on what type of queries cause Knowledge Graph results to come up, and which ones you don’t that you can see here: (<https://www.stonetemple.com/great-knowledge-box-showdown/>).

While 3.5 billion pieces of information may seem like a very large number, in the grand scheme of things, it represents a very small portion of all the types of queries users may have. As a result, Google is pursuing other avenues to expand their ability to further enhance the information in the search results.

For example, they have begun experimentation with extracting information from web sites. One example of this is that Google has started to display step by step instructions extracted from web sites, as shown in Figure 6-58

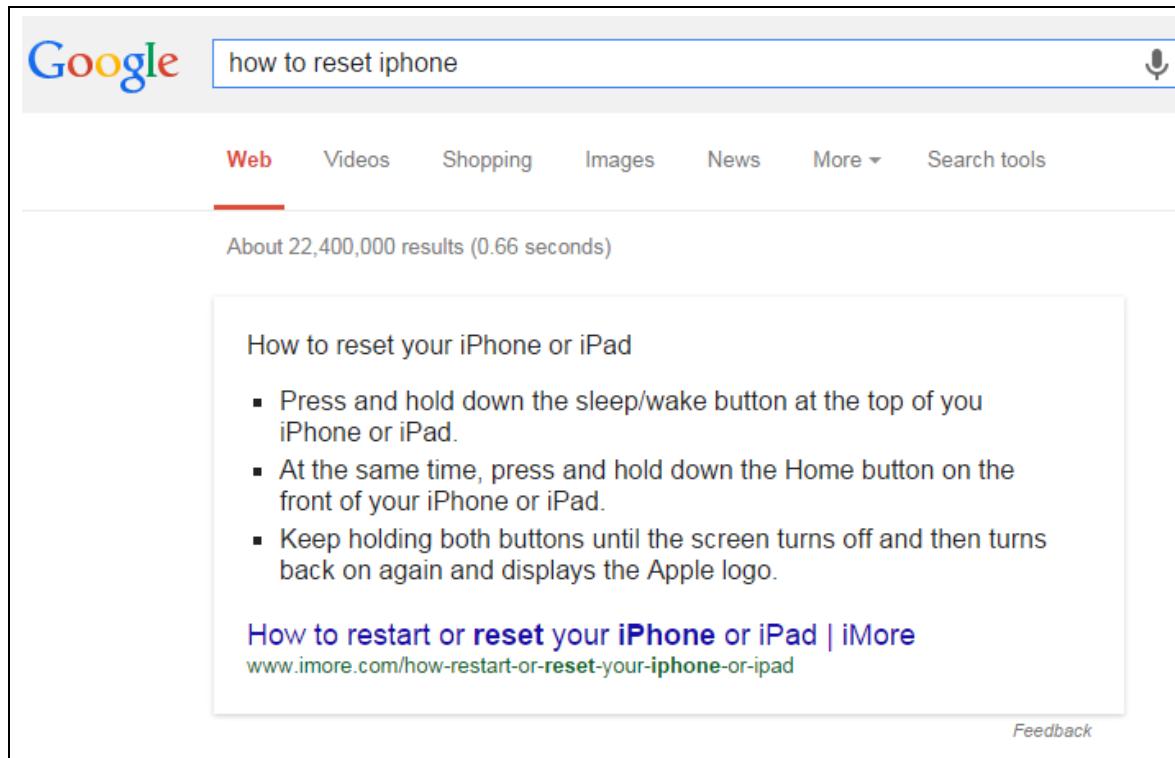


Figure 6-58. Example of Step by Step Instructions in Google's SERPs

Barry Schwartz reached out to Google, and got the following response when asked about the use of step by step instructions (<http://searchengineland.com/googles-knowledge-graph-showing-step-step-instructions-examples-194923>):

“We started experimenting with this in early June. We hope it draws attention to webpages that provide a useful series of steps to help people complete their task. In these cases we focus attention on the snippet because it’s likely to be more helpful for deciding whether the webpage is going to be the most useful for the task”

Other examples exist where Google is extracting knowledge from websites and showing it in the SERPs. Figure 6-59 shows an example of historical information being found on a web site and displayed directly in the results.

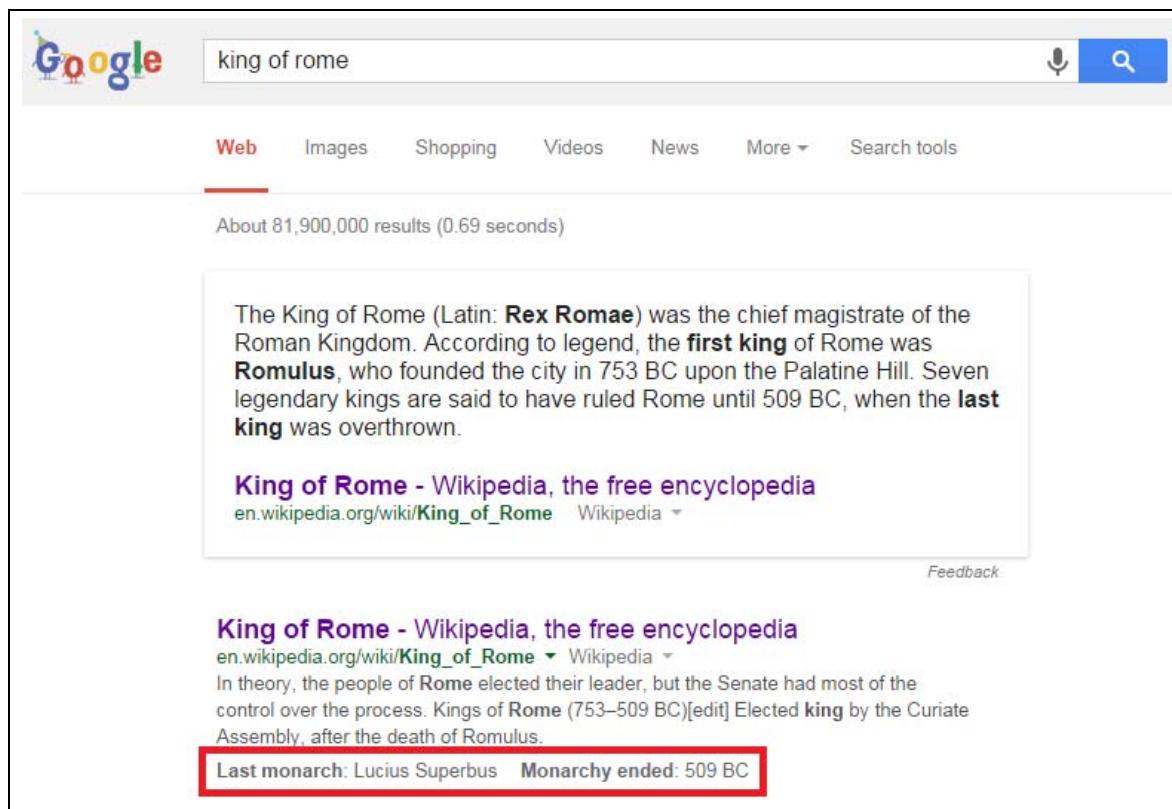


Figure 6-59. Knowledge Extraction Example

The examples shown in Figure 6-58 and 6-59 are a clear step beyond the simple use of structured data. These represent early examples of what Google refers to as the Knowledge Vault.

Fair Use?

As Google presents more and more of these types of search results, many of the impacted publishers feel that Google is stealing their content, and profiting from it. The question becomes whether or not Google's usage can be considered "Fair Use" as defined by the US Copyright Office (<http://www.copyright.gov/fls/fl102.html>). There are 4 factors involved in determining Fair Use, as follows:

1. The purpose and character of the use, including whether such use is of commercial nature or is for nonprofit educational purposes
2. The nature of the copyrighted work
3. The amount and substantiality of the portion used in relation to the copyrighted work as a whole
4. The effect of the use upon the potential market for, or value of, the copyrighted work

There is actually no clear definition as to what defines Fair Use, but it is clear that the more substance you take from the 3rd party is a factor. It is common practice among those who quote others, or whom attempt to make fair use of someone else's copyrighted material to provide attribution. However, the US Copyright Office indicates that this might not be enough: "Acknowledging the source of the copyrighted material does not substitute for obtaining permission."

In addition, this is more than a US only issue as well, and the laws differ from country to country.

Whether this becomes an issue for Google or not is yet to be determined, but the scale of what they are trying to do makes it likely that they will be subjected to legal challenges, and that the way that various legal systems will respond will differ.

One additional aspect to consider, is that public domain information is not copyrightable. For example, the fact that Olympia is the capital of the state of Washington, is not copyrightable info. If Google is able to extract some common knowledge from 3rd party sites, then this would not be subject to this discussion.

How the Knowledge Vault Works

The first public indication that Google had a concept they called the Knowledge Vault was in a presentation by Google's Kevin Murphy which took place on October 31, 2013. The full presentation can be seen here (<http://cikm2013.org/slides/kevin.pdf>). As of October 2014, this is just a research project within Google, but it is still useful to see what this project is about. The core concepts being studied are as follows:

Machine Reading

This is the process of extracting facts from a large text corpora. This is similar to methods developed by Carnegie-Mellon, the University of Washington, and others, but Google is working on a much larger scale version. In addition, they are researching methods for using other prior knowledge to help reduce the error rate.

As information is assembled, it becomes possible to infer, or even determine, other facts. Figure 6-???? (slide 16) shows an example of this in action. For example, if we know that Barack Obama and Michelle Obama are both parents of Sasha Obama, then we can infer that it is likely that they are married (though that is not necessarily true).

Predicting facts given prior knowledge

Google

- Perform association rule mining* on Freebase graph, to find noisy rules (features passed to a learned classifier).

$$\forall x, y. \exists z. \text{parent-of}(x, z) \wedge \text{parent-of}(y, z) \Rightarrow \text{married}(x, y)$$

```
graph LR; Barack[Barack Obama] -- married-to --> Michelle[Michelle Obama]; Barack -- parent-of --> Sasha[Sasha Obama]; Michelle -- parent-of --> Sasha;
```

* "Random Walk Inference and Learning in A Large Scale Knowledge Base", Ni Lao et al, 2011

Figure 6-60. Inferring New Information

Asking the Web

Web-based question and answers can be used to further supplement the available information. Learning how to ask the right questions (as shown in Figure 6-61), and how to frame those questions is by itself a very difficult process. In addition, verification of the accuracy of the responses is important as well.

The importance of asking the right question



Who is the mother of Frank Zappa



[The Mothers of Invention - Wikipedia, the free encyclopedia](#)
en.wikipedia.org/wiki/The_Mothers_of_Invention ▾

The Mothers of Invention were an American rock band from California that served as the backing musicians for Frank Zappa, a self-taught composer and ...
History - Personnel - Discography - References

[Frank Zappa - Wikipedia, the free encyclopedia](#)

en.wikipedia.org/wiki/Frank_Zappa ▾

Jump to 1970: Rebirth of The Mothers and filmmaking - [edit]. Frank Zappa in Paris, early 1970s. Later in 1970, Zappa formed a new version of The ...
Discography - Moon Zappa - Diva Zappa - Gail Zappa



Who is the mother of Frank Zappa Baltimore Maryland



[Frank Zappa - Wikipedia, the free encyclopedia](#)
en.wikipedia.org/wiki/Frank_Zappa ▾

Frank Vincent Zappa was born in Baltimore, Maryland, on December 21, 1940. His mother, Rose Marie (née Colimore), was of Italian and French ancestry; his ...

Kevin Murphy, CIKM industry talk, San Francisco, CA, October 31, 2013

Figure 6-61. The Important of Asking the Right Question

Asking People

Freebase is itself an example of this, as it is community edited. Other sources can be used as well. For example, each knowledge panel that Google shows contains a feedback link, allowing them to collect information on accuracy problems. This feedback can also be flawed, and Google is investigating algorithms to predict the possibility that the information is received is correct.

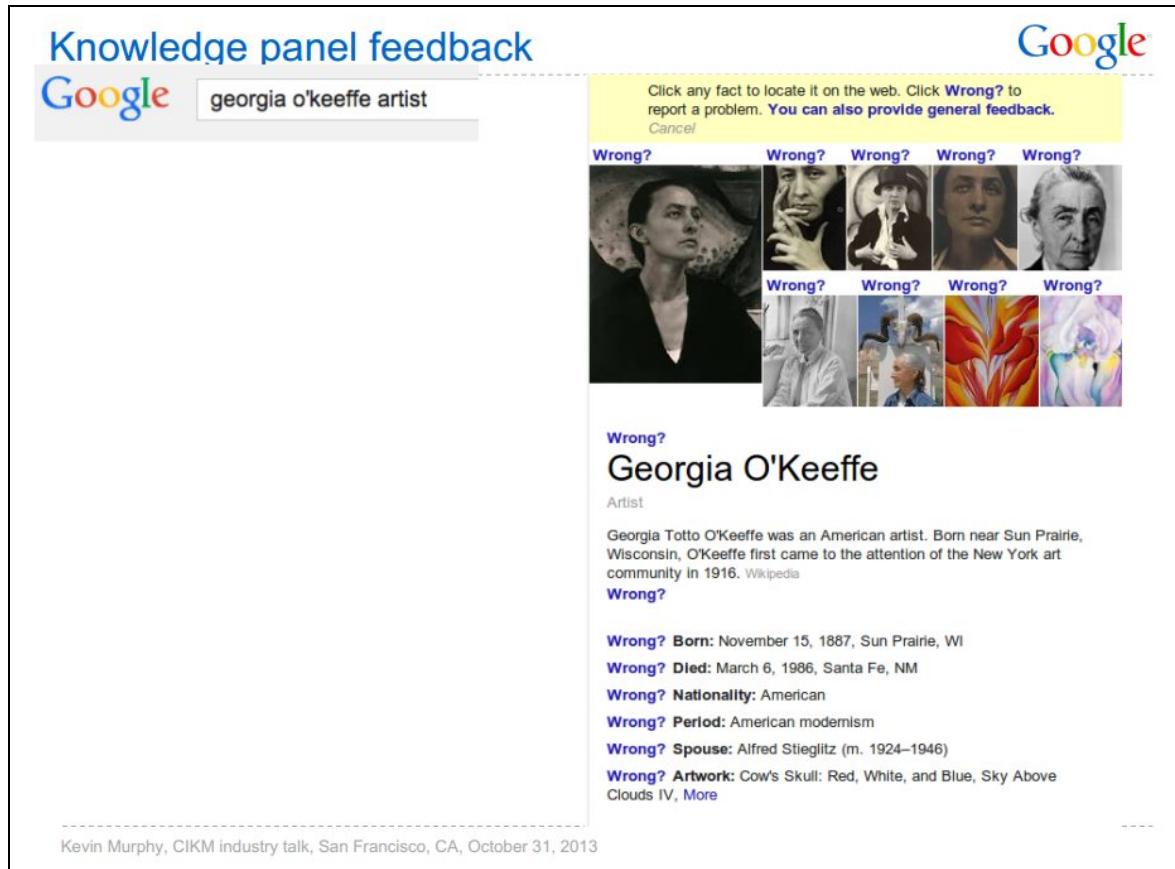


Figure 6-62. Knowledge Panel Feedback

There are many patents that Google has that could potentially pertain to the Knowledge Graph and/or the Knowledge Vault. Each of these covers a particular aspect of how to extract information into a knowledge base. Here are a couple of examples for further reading:

5. Knowledge Graph Based Search System: <http://www.google.com/patents/US20120158633>
6. Extracting Patterns and Relations from the World Wide Web: <http://www.seobythesea.com/2014/08/identifying-entity-types-transfiguraton-search-google/>
7. Determining Geographic Locations for Place Names in Fact Repository: <http://www.seobythesea.com/2013/01/knowledge-base-locations-in-web-pages/>

There are many more that apply to this complex topic, and it will remain an area of investigation for some time to come.

The Future of the Knowledge Vault

As of October 2014, the Knowledge Vault is a concept in its infancy. As noted above, it is just a research project. The algorithms are in primitive states of definition, and the required processing power is quite substantial. Google will keep investing in these types of technologies as it tries to find more and more ways to provide better and better results. Their goal, and that of other search engines, remains to “organize the world’s information”, and they will keep investing in it until they succeed.

However, this process may take a decade or more. This means we will see gradual changes continuing over time. Even if a single breakthrough provides Google with access to one billion facts, which sounds like a large number, it will still only impact a very small percentage of search results.

However, understanding the concepts of semantic search and the Knowledge Vault can help you understand a bit more about where search engines are going.

Conclusion

By now you should be aware that a search-engine-friendly website is the first step towards SEO success. In the next chapter on link building we will demonstrate how links are also a critical piece of the SEO puzzle—particularly when targeting highly competitive terms. However, if you have not made your site crawler-friendly and optimized, all of your other efforts – whether they be link development, social media promotion, or other tactics to improve search visibility and increase search traffic - will be a wasted effort. A website built from the ground up to optimal specifications for crawler accessibility and top organic exposure is the foundation from which you will build all SEO initiatives. From this solid foundation, even the loftiest of goals are within reach.

Content Marketing

In today's search environment, the main driving factors are now what we generally refer to as social proof signals - such as inbound links (such as within a blog post) and user engagement with your content (such as time spent watching your video). As you will see in Chapter 8, social signals (such as retweets, Likes, Pins) don't appear to have a direct ranking impact, and Google Plus appears to have an impact, but only from a personalized search perspective.

For many years, links to a website were the single largest factor in determining its search engine rankings, because links generally (before they became a tool for SEO manipulation) existed for the purpose of providing a pathway for a site's users to find additional, relevant content on a third party's website – a “signal” that the owner of the linking site deemed the third party's linked content to be .

Because of the power that this signal wielded, many SEO professionals pursued obtaining links to their sites or their client's sites without worrying about the quality of the site where these links resided. Unfortunately, many link building efforts – and services – spawned by this behavior had little integration with the rest of the publisher's content development and marketing strategies.

Clearly, this violated the spirit of what the search engines were measuring, and placing value on – links that act as valid endorsements for your content. As a result, the search engines, and Google in particular, have taken many steps to force website owners and publishers to view link building more holistically, as an “earned” engagement rather than a “purchased” endorsement – requiring a renewed focus on links as a measurement of content quality. This shift, both necessary and welcomed, re-establishes the need for quality content development (as the “earner” of links) to be integrated with the overall PR and marketing strategy for businesses.

The development of highly shareable content, and promotion of that content via various channels for increased business visibility, is generally referred to as **content marketing**. Content can be published on your own site, other people's sites, or in social media, but in all cases acts to build visibility for your brand online. The most valuable content is usually highly relevant to what you do, solves problems for others or stirs their emotions, and is often non-commercial in nature.

Links remain a large factor in search engine ranking algorithms, but we use **content marketing** to build our reputation and visibility online, and as a result we obtain organic links of the highest possible quality – links that would be desirable for your business even if the search engines did not exist, and links where people seeing them might actually click on them to engage with your business.

The most important thing to remember as you delve into this chapter is that the primary goal of any content marketing effort needs to be enhancing the reputation of your business. Any campaign that starts with “getting links” as the objective, without placing primary and ongoing focus on the quality and value of the content being linked to, will eventually run into problems (if it hasn't already – see Chapter 9 on Panda, Penguin, and Penalties). During a 2012 interview, Google's Matt Cutts

(<http://www.stonetemple.com/matt-cutts-and-eric-talk-about-what-makes-a-quality-site/>) and Eric Enge had the following exchange:

“Eric Enge: It dawned on me recently that link building is an interesting phrase that has misled people. It is a bit of a cart before the horse thing. It has led people to think about links as something they get from the “dark corners of the web”. Places where no one ever goes, so it does not matter what you do there. So by thinking of it this way, as link building, you are off on the wrong foot even before you get started.”

“Matt Cutts: That’s right. It segments you into a mindset, and people get focused on the wrong things. It leads them to think about links as the end goal. It is important to think about producing something excellent first. If you have an outstanding product, world class content, or something else that sets you apart, then you can step back and start thinking about how to promote it.”

There are many that believe that social signals and user engagement with your content have become important ranking factors. However, the impact of social media appears to be quite limited:

- Google+ can have a strong impact on personalized search within Google for those who are active on the Google+ platform.
- Search engines may use shared content on social media platforms as a way of discovering new content, in particular, news related content.

These new ranking factors will be discussed in greater detail in “How Social Media and User Data Play a Role in Search Results and Rankings” in Chapter 8.

How Links Historically Influenced Search Engine Rankings

The concept of using links as a way to measure a site’s importance was first made popular by Google with the implementation of its PageRank algorithm (others had previously written about using links as a ranking factor, but Google’s rapidly increasing user base popularized it). In simple terms, each link to a web page is a vote for that page. But it’s not as simple as the page with the most votes wins. Links and linking pages are not all created equal. Some links are weighted more heavily by Google’s PageRank algorithm than others.

The key to this concept is the notion that links represent an “editorial endorsement” of a web document. Search engines rely heavily on editorial votes. However, as publishers learned about the power of links, some publishers started to manipulate links through a variety of methods. This created situations in which the intent of the link was not editorial in nature, and led to many algorithm enhancements.

To help you understand the origins of link algorithms, the underlying logic of which is still in force today, let’s take a look at the original PageRank algorithm in detail.

The Original PageRank Algorithm

The PageRank algorithm was built on the basis of the original PageRank thesis (<http://infolab.stanford.edu/~backrub/google.html>) authored by Sergey Brin and Larry Page while they were undergraduates at Stanford University.

In the simplest terms, the paper states that each link to a web page is a vote for that page. However, as stated above, votes do not have equal weight. So that you can better understand how this works, we’ll explain the PageRank algorithm at a high level. First, all pages are given an innate but tiny amount of PageRank, as shown in Figure 7-1.

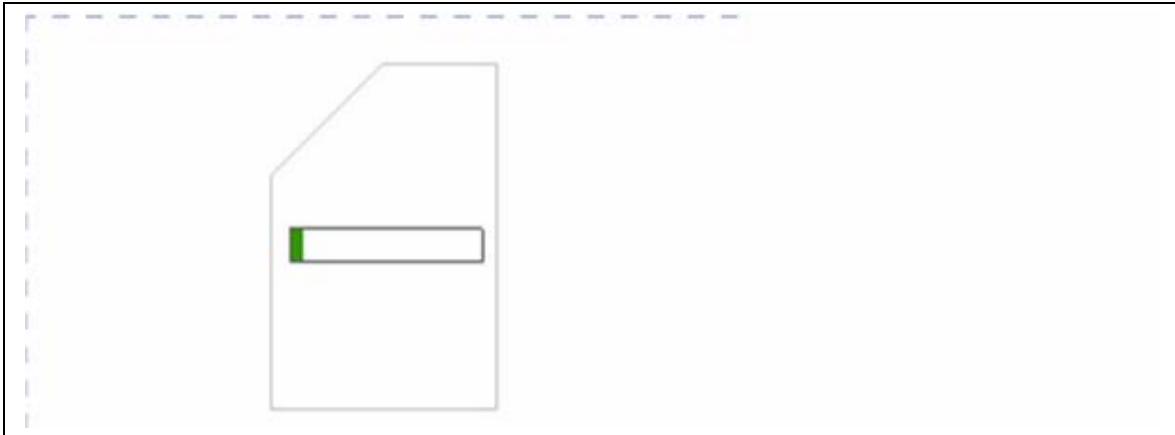


Figure 7-1. Some PageRank for every page

Pages can then increase their PageRank by receiving links from other pages, as shown in Figure 7-2.

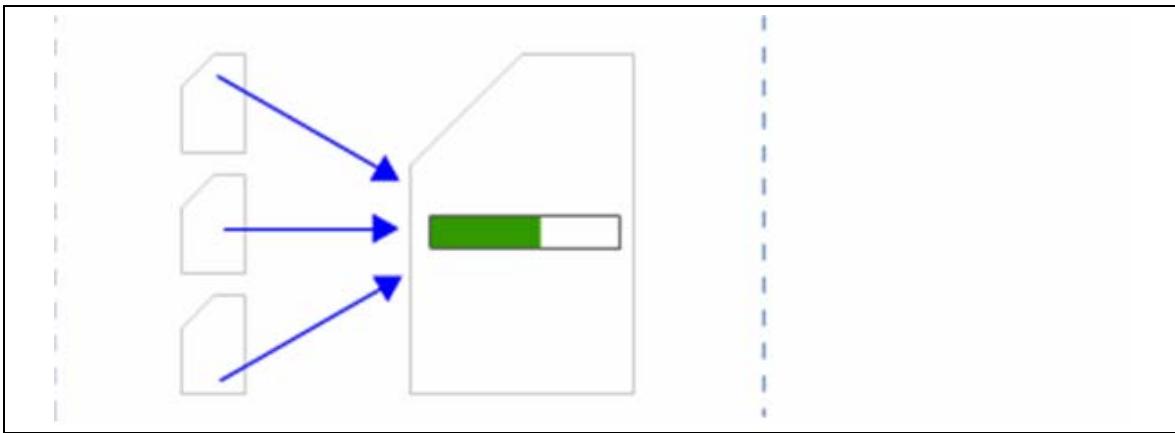


Figure 7-2. Pages receiving more PageRank through links

How much PageRank can a page pass on to other pages through links? That ends up being less than the page's PageRank. In Figure 7-3 this is represented by $f(x)$, meaning that the passable PageRank is a function of x , the total PageRank. In 2009 Matt Cutts wrote a post (<http://www.mattcutts.com/blog/pagerank-sculpting/>) in which he suggested that a page might be able to vote 85% to 90% of its PageRank.



Figure 7-3. Some of a page's PageRank passable to other pages

If this page links to only one other page, it passes all of its passable PageRank to that page, as shown in Figure 7-4, where Page B receives all of the passable PageRank of Page A.

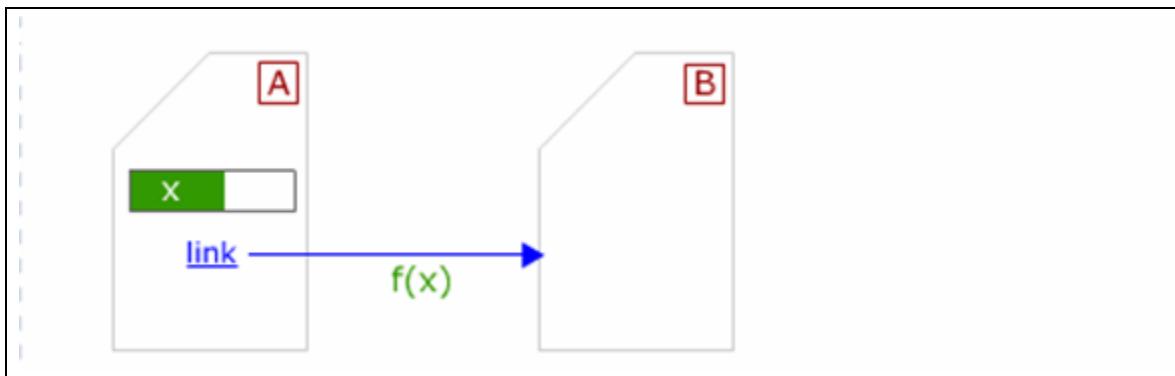


Figure 7-4. Passing of PageRank through a link

However, the scenario gets more complicated because pages will link to more than one other page. When that happens the passable PageRank gets divided among all the pages receiving links. We show that in Figure 7-5, where Page B and Page C each receive half of the passable PageRank of Page A.

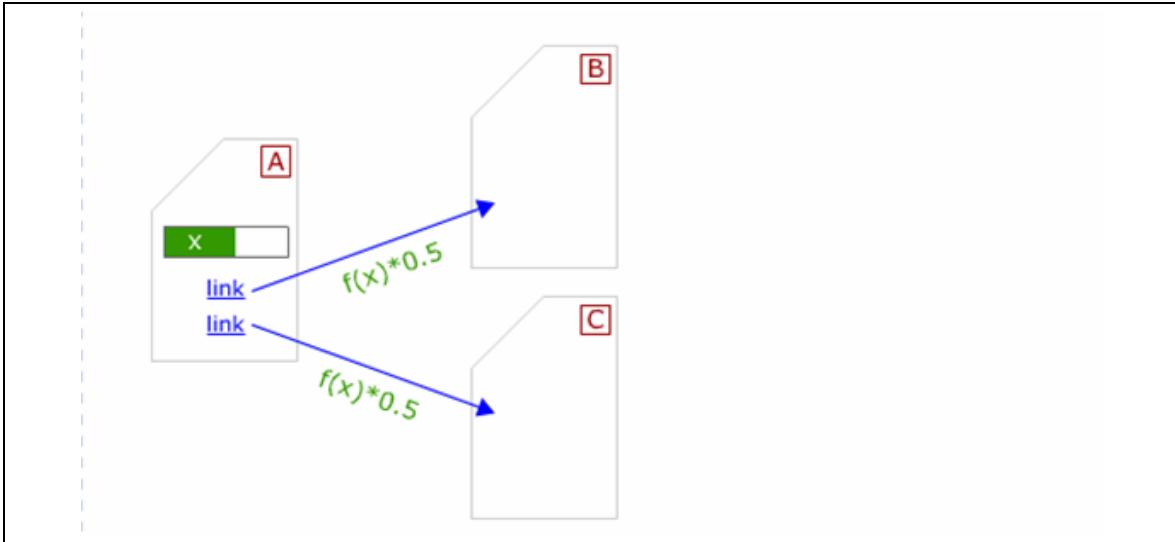


Figure 7-5. Simple illustration of how PageRank is passed

In the original PageRank formula, link weight is divided equally among the number of links on a page. This undoubtedly does not hold true today, but it is still valuable in understanding the original intent. Now take a look at Figure 7-6, which depicts a more complex example that shows PageRank flowing back and forth between pages that link to one another.

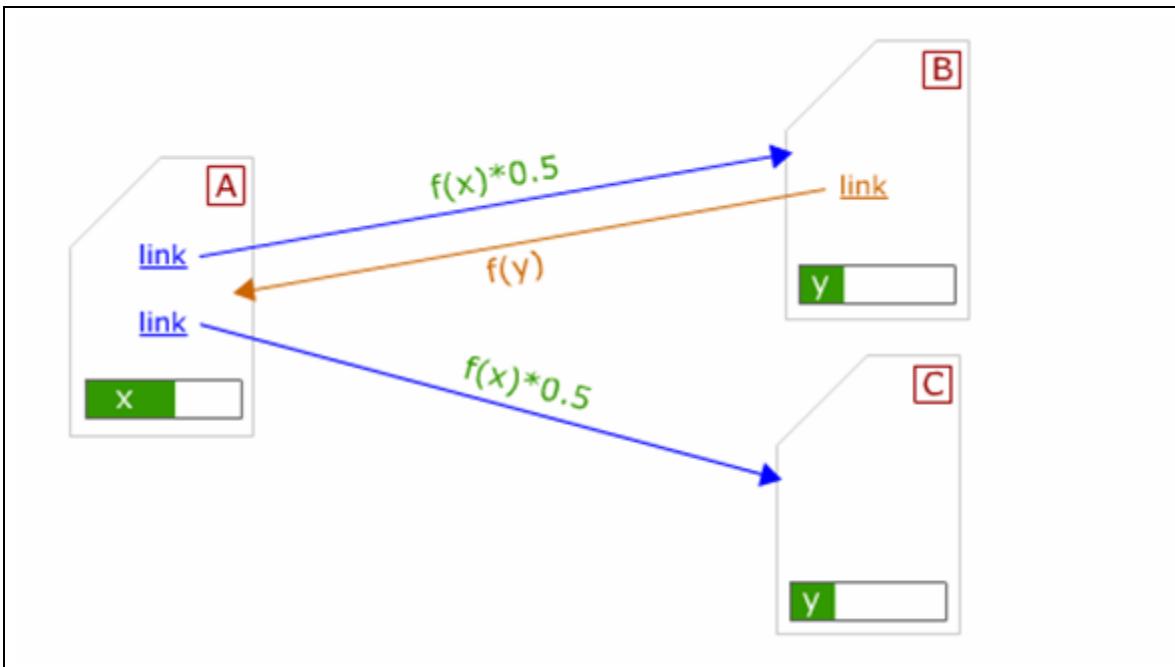


Figure 7-6. Cross-linking between pages

Cross-linking makes the PageRank calculation much more complex. In Figure 7-6, Page B now links back to Page A and passes some PageRank, $f(y)$, back to Page A. Figure 7-7 should give you a better understanding of how this affects the PageRank of all the pages.

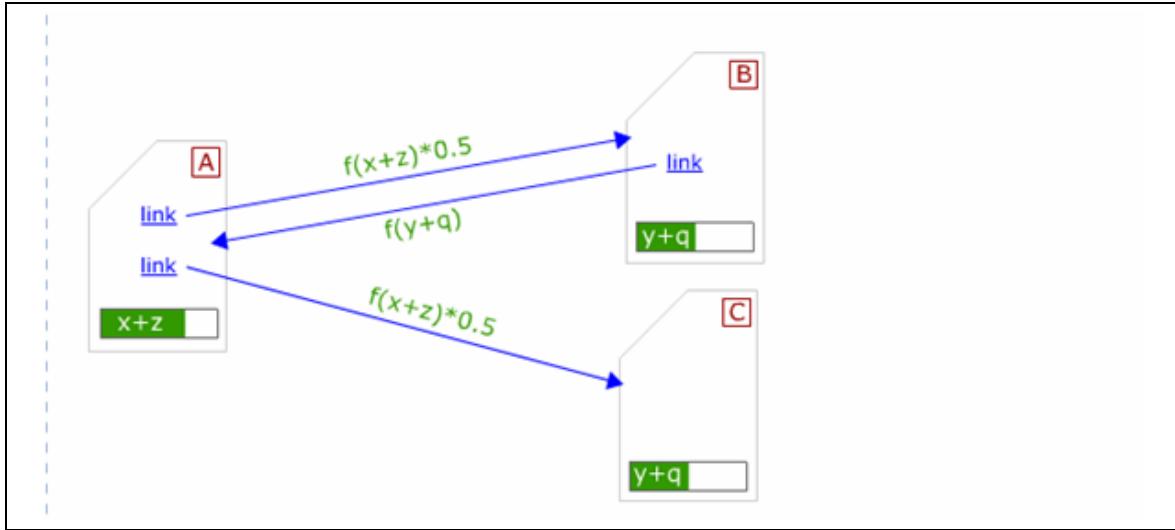


Figure 7-7. Iterative PageRank calculations

The key observation here is that when Page B links to Page A to make the link reciprocal, the PageRank of Page A (x) becomes dependent on $f(y)$, the passable PageRank of Page B, which happens to be dependent on $f(x)$!. In addition, the PageRank that Page A passes to Page C is also impacted by the link from Page B to Page A. This makes for a very complicated situation where the calculation of the PageRank of each page on the Web must be determined by recursive analysis.

We have defined new parameters to represent this: q , which is the PageRank that accrues to Page B from the link that it has from Page A (after all the iterative calculations are complete); and z , which is the PageRank that accrues to Page A from the link that it has from Page B (again, after all iterations are complete).

The scenario in Figure 7-8 adds additional complexity by introducing a link from Page B to Page D. In this example, pages A, B, and C are internal links on one domain, and Page D represents a different site (shown as Wikipedia). In the original PageRank formula, internal and external links passed PageRank in exactly the same way. This became exposed as a flaw because publishers started to realize that links to other sites were “leaking” PageRank away from their own site, as you can see in Figure 7-8.

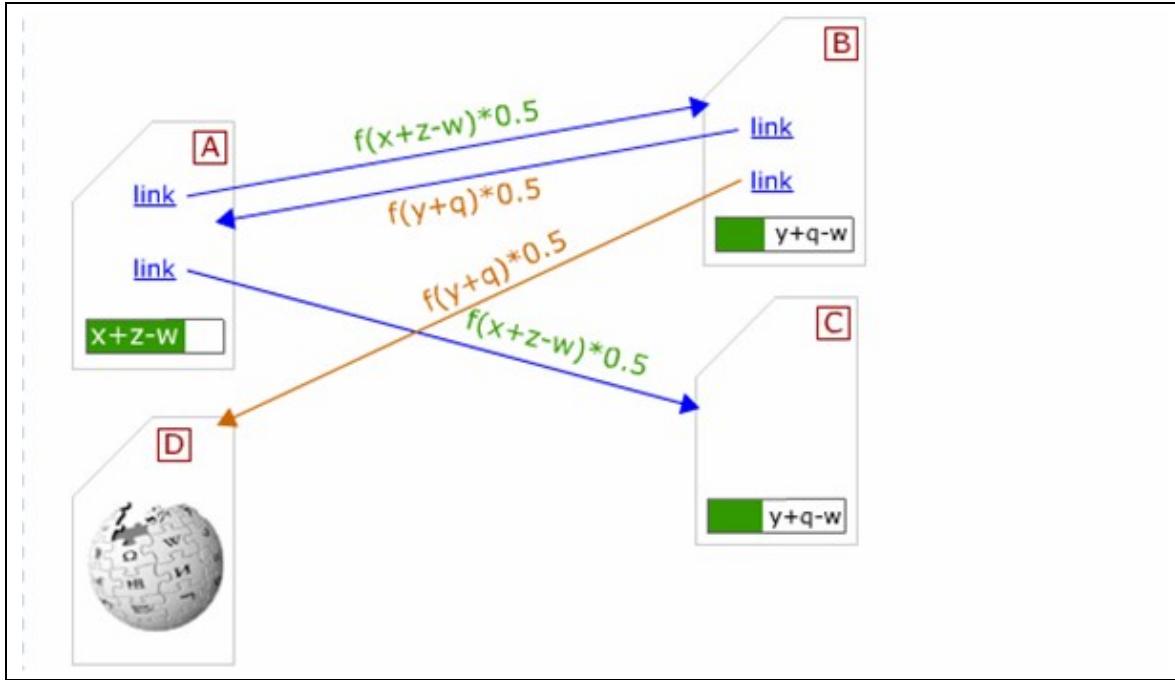


Figure 7-8. PageRank being leaked

Because Page B links to Wikipedia, some of the passable PageRank is sent there, instead of to the other pages that Page B is linking to (Page A in our example). In Figure 7-8, we represent that with the parameter w , which is the PageRank not sent to Page A because of the link to Page D.

The PageRank “leak” concept presented a fundamental flaw in the algorithm once it became public. Like Pandora’s Box, once those who were creating pages to rank Google investigated PageRank’s underlying principles, they would realize that linking out from their own sites would cause more harm than good. If a great number of websites adopted this philosophy, it could negatively impact the “links as votes” concept and actually damage the quality of Google’s algorithm. Needless to say, Google corrected this flaw to its algorithm. As a result of these changes, worrying about PageRank leaks is not recommended. Quality sites should link to other relevant quality pages around the Web.

Even after these changes, internal links from pages still pass some PageRank, so they still have value, as shown in Figure 7-9.

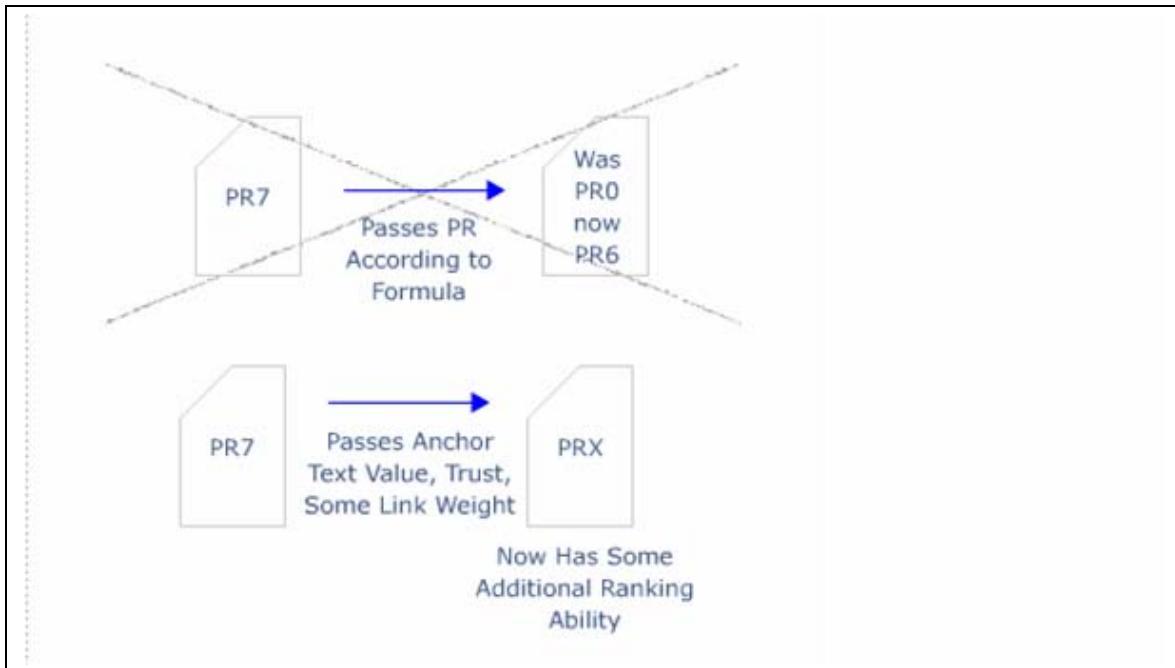


Figure 7-9. Internal links still passing some PageRank

Google has continuously changed and refined the way they use links to impact rankings, and the current algorithm is not based on PageRank as it was originally defined. However, familiarity and comfort with the original algorithm is certainly beneficial to those who practice optimization of Google results.

All link-based algorithms are based on the assumption that for the most part the links received are legitimate endorsements by the publisher implementing a link to your web site. The person implementing the link should be doing it because they feel they are linking to a great resource that would be relevant to visitors on their web site.

In an ideal world, links would be similar to the **academic citations** you might find at the bottom of a scientist's published paper, where they cite the other works they have referenced in putting together their research.

If the publisher implementing the link is compensated for doing so, the value of the link to a search engine is diminished, and such links can be harmful to search engine algorithms. Note that compensation can come in the form of money, or other considerations, and we will explore this more in this chapter.

Additional Factors That Influence Link Value

Classic PageRank isn't the only factor that influences the value of a link. In the following subsections, we discuss some additional factors that influence the value a link passes.

Anchor text

Anchor text refers to the clickable part of a link from one web page to another. As an example, Figure 7-10 shows a snapshot of a part of the Quicken Loans Home Page at <http://www.quickenloans.com>.

Popular:	About Us:	Help:
Refinancing	Who is Quicken Loans	Contact Us
Home Buying	Quicken Loans Reviews	Email Policy
Mortgage Rates	FAQs	Affiliate Promotion Privacy
Zing Blog Home	Quicken Loans Careers	Security and Privacy
Email Rates & Updates	Press Room	Disclosures and Licenses
My Quicken Loans	Family of Companies	Terms of Use

Figure 7-10. Anchor text: a strong ranking element

The anchor text for the first link in the list of Popular pages in Figure 7-10 is “Refinancing”. The search engine uses this anchor text to help it understand what the page receiving the link is about. As a result, the search engine will interpret the link as saying that the page receiving the link is about “Refinancing”, and therefore rank the page higher in the search results for that search query.

At one time, anchor text was so powerful that SEOs engaged in a practice called *Google Bombing*. The idea was that if you link to a given web page from many places with the same anchor text, you can get that page to rank for queries related to that anchor text, even if the page is unrelated and didn't even include any of the words in the query.

One notorious Google Bomb was a campaign that targeted the Whitehouse.gov biography page for George W. Bush with the anchor text *miserable failure*. As a result, that page ranked #1 for searches on *miserable failure* until Google tweaked its algorithm to reduce the effectiveness of this practice, but the page still ranks #1 for that search query in Bing as of April 2014.

The concept of *Google Bombing* was not the worst consequence of the power of anchor text. The use of anchor text as a ranking factor is only useful in the search algorithms if the person implementing the link naturally chooses what to use, and if they are compensated for using specific anchor text, the value of the link as a ranking signal is negatively impacted.

To make matters worse, SEOs started to abuse the system and started implementing link building programs designed around anchor text to drive their rankings. As a result, publishers that did not pursue these types of link building campaigns were at a severe disadvantage. Ultimately, this started to break down the notion of links as valid academic citations, and Google began to take action.

In early 2012 Google began to send publishers **Unnatural link warnings** through Google Webmaster Tools, and then on April 24, 2012 Google released the first version of its **Penguin** algorithm. These topics are discussed more in "Panda, Penguin, Penalties" in Chapter 9.

Anchor text probably remains an important part of the algorithms, but now the search engines look for unnatural patterns of anchor text (too much of a good thing) and are lowering the rankings for publishers that exhibit patterns of artificially influencing the anchor text people use in links to their web site.

Relevance

Links that originate from sites/pages on the same topic as the publisher's site, or on a closely related topic, are worth more than links that come from a site on an unrelated topic.

Think of the relevance of each link being evaluated in the specific context of the search query a user has just entered. So, if the user enters *used cars in Phoenix* and the publisher has received a link to their Phoenix used cars page that is from the Phoenix Chamber of Commerce, that link will reinforce the search engine's belief that the page really does relate to Phoenix.

Similarly, if a publisher has another link from a magazine site that has done a review of used car websites, this will reinforce the notion that the site should be considered a used car site. Taken in combination, these two links could be powerful in helping the publisher rank for “used cars in Phoenix”.

Authority

This has been the subject of much research. One of the more famous papers, written by Apostolos Gerasoulis and others at Rutgers University and titled “DiscoWeb: Applying Link Analysis to Web Search” (<http://www.cse.lehigh.edu/~brian/pubs/1999/www8/>), became the basis of the Teoma algorithm, which was later acquired by AskJeeves and became part of the Ask algorithm.

What made this unique was the focus on evaluating links on the basis of their relevance to the linked page. Google’s original PageRank algorithm did not incorporate the notion of topical relevance, and although Google’s algorithm clearly does this today, Teoma was in fact the first search engine to offer a commercial implementation of link relevance.

Teoma introduced the notion of *hubs*, which are sites that link to most of the important sites relevant to a particular topic, and *authorities*, which are sites that are linked to by most of the sites relevant to a particular topic.

The key concept here is that each topic area that a user can search on will have authority sites specific to that topic area. The authority sites for used cars are different from the authority sites for baseball.

Refer to Figure 7-11 to get a sense of the difference between hub and authority sites.

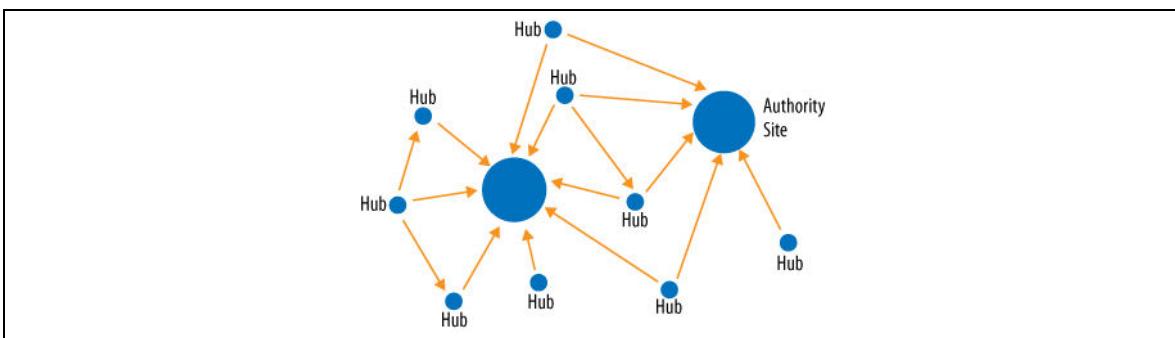


Figure 7-11. Hubs and authorities

So, if the publisher has a site about used cars, it should seek links from websites that the search engines consider to be authorities on used cars (or perhaps more broadly, on cars). However, the search engines will not tell you which sites they consider authoritative—making the publisher’s job that much more difficult.

The model of organizing the Web into topical communities and pinpointing the hubs and authorities is an important model to understand (read more about it in Mike Grehan’s paper, “Filthy Linking Rich!” at http://www.search-engine-book.co.uk/filthy_linking_rich.pdf). The best link builders understand this model and leverage it to their benefit.

Trust

Trust is distinct from authority. Authority, on its own, doesn’t sufficiently take into account whether the linking page or the domain is easy or difficult for spammers to infiltrate, or the motivations of the person implementing the link. Trust, on the other hand, does.

Evaluating the trust of a website likely involves reviewing its link neighborhood to see what other trusted sites link to it. More links from other trusted sites would convey more trust.

In 2004, Yahoo! and Stanford University published a paper titled “Combating Web Spam with TrustRank” (<http://www.vldb.org/conf/2004/RS15P3.PDF>). The paper proposed starting with a trusted seed set of pages (selected by manual human review) to perform PageRank analysis, instead of a random set of pages as was called for in the original PageRank thesis.

Using this tactic removes the inherent risk in using a purely algorithmic approach to determining the trust of a site, and potentially coming up with false positives/negatives.

The trust level of a site would be based on how many clicks away it is from seed sites. A site that is one click away accrues a lot of trust; two clicks away, a bit less; three clicks away, even less; and so forth. Figure 7-12 illustrates the concept of TrustRank.

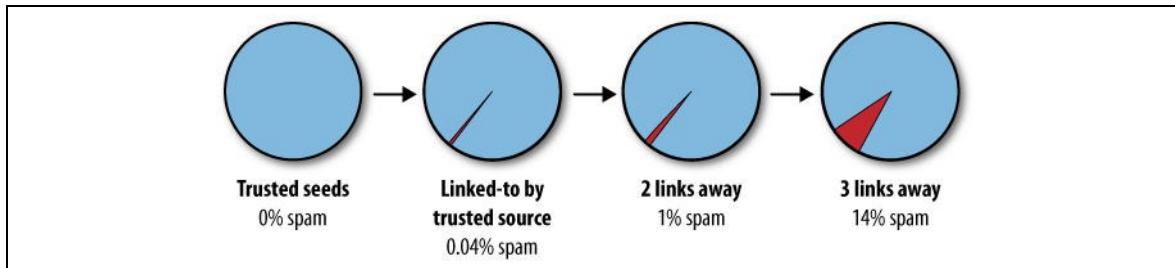


Figure 7-12. TrustRank illustrated

The researchers of the TrustRank paper also authored a paper describing the concept of *spam mass* (<http://ilpubs.stanford.edu:8090/697/1/2005-33.pdf>). This paper focuses on evaluating the effect of spammy links on a site's (unadjusted) rankings. The greater the impact of those links, the more likely the site itself is spam. A large percentage of a site's links being purchased is seen as a spam indicator as well. You can also consider the notion of *reverse TrustRank*, where linking to spammy sites will lower a site's TrustRank.

It is likely that Google and Bing both use some form of trust measurement to evaluate websites. It is probably done by different means than outlined in the TrustRank and Spam Mass papers, and it may be incorporated into the methods they use for calculating Authority, but nonetheless, trust is believed to be a significant factor in rankings. For SEO practitioners, getting measurements of trust can be difficult. Currently, mozTrust from Moz's Open Site Explorer, and TrustFlow from Majestic SEO are the most well-known publicly available metrics tools for trust level of a web page.

How Search Engines Use Links

The search engines use links primarily to discover web pages, and to count the links as votes for those web pages. But how do they use this information once they acquire it? Let's take a look:

Index inclusion

Search engines need to decide which pages to include in their index. Discovering pages by crawling the Web (following links) is one way they discover web pages (another is through the use of XML Sitemap files). In addition, the search engines do not include pages that they deem to be of very low value because cluttering their index with those pages will not lead to a good experience for their users. The cumulative link value, or link authority, of a page is a factor in making that decision.

Crawl rate/frequency

Search engine spiders go out and crawl a portion of the Web every day. This is no small task, and it starts with deciding where to begin and where to go. Google has publicly indicated that it starts its crawl in reverse PageRank order. In other words, it crawls PageRank 10 sites first, PageRank 9 sites next, and so on. Higher PageRank sites also get crawled more deeply than other sites. It is likely that Bing starts their crawl with the most important sites first as well.

This would make sense, because changes on the most important sites are the ones the search engines want to discover first. In addition, if a very important site links to a new resource for the first time, the search engines tend to place a lot of trust in that link and want to factor the new link (vote) into their algorithms quickly. In June 2010 Google released an update to their infrastructure that they called Caffeine which greatly increased their crawling capacity and speed (<http://googleblog.blogspot.com/2010/06/our-new-search-index-caffeine.html>), but being higher in the crawl priority queue still matters.

Ranking

Links play a critical role in ranking. For example, consider two sites where the on-page content is equally relevant to a given topic. Perhaps they are the shopping sites Amazon.com and (the less popular) JoesShoppingSite.com (not a real site).

The search engine needs a way to decide who comes out on top: Amazon or Joe. This is where links come in. Links help cast the deciding vote. If more sites, and more important sites, link to Amazon, it must be more important, so Amazon is more likely to rank higher.

Further Refining How Search Engines Judge Links

Many aspects are involved when you evaluate a link. As we just outlined, the most commonly understood ones are authority, relevance, trust, and the role of anchor text. However, other factors also come into play.

Additional Link Evaluation Criteria

In the following subsections, we discuss some of the more important factors search engines consider when evaluating a link's value.

Source independence

A link from your own site back to your own site is, of course, not an independent editorial vote for your site. Put another way, the search engines assume that you will vouch for your own site.

Think about your site as having an accumulated total link authority based on all the links it has received from third-party websites, and your internal linking structure as the way you allocate that juice to pages on your site. Your internal linking structure is incredibly important, but it does little if anything to build the total link authority of your site.

In contrast, links from a truly independent source carry much more weight. Extending this notion a bit, it may be that you have multiple websites. Perhaps they have common data in the Whois records (such as the name servers or contact information). Search engines can use this type of signal to treat cross-links between those sites more like internal links than inbound links earned by merit.

Even if you have completely different Whois records for your websites but they all cross-link to each other, the search engines can detect this pattern easily. Keep in mind that a website with no independent third-party links into it has no link power to vote for other sites.

If the search engine sees a cluster of sites that heavily cross-link and many of the sites in the cluster have no or few incoming links to them, the links from those sites may well be ignored.

Conceptually, you can think of such a cluster of sites as a single site. Cross-linking to them can be algorithmically treated as internal links, with links between them not adding to the total link authority score for each other. The cluster would be evaluated based on the inbound links to the cluster.

Of course, there are many different ways such things could be implemented, but one thing that would not have SEO value is to build a large number of sites, just to cross-link them with each other.

Linking domains

Obtaining an editorially given link to your site from a third-party website is usually a good thing. But if more links are better, why not get links from every page of these sites if you can? In theory, this is a good idea, but search engines do not necessarily count multiple links from a domain cumulatively.

When Google first came out, their link based algorithm was revolutionary. As spammers studied the PageRank algorithm, they realized that every page on the web naturally had a small amount of innate PageRank. It was not long spammers realized that they could create one million plus page web sites, even if they had crappy content, and then use every page of that web site to link to (a **site wide link**) their most important page on that site, or even a page on a different site. Alternatively, they simply bought site wide links in the footer of very large sites, and this also worked.

In the early days of Google, this worked really well. As Google learned what was happening, they realized that multiple links on one site most likely represent one editorial vote. Google then put a dampener on the incremental value of more than one link from a site, and, as a result, each incremental link from a site began to pass slightly less value per link.

In addition, over time, Google began to get more active in penalizing sites that use bad link building practices, such as buying links, and one clue that this might be taking place was the use of site wide links. This meant that a site wide link could potentially harm your site.

Link builders and spammers figured this out and adapted their strategies. They began to focus on obtaining links on as many different domains as possible. This particularly impacted a strategy known as **guest posting** (this is discussed more in “Guest Posting” below), which is the concept of writing an article for another web site and getting them to publish that on their blog.

Although this strategy is a legitimate content marketing concept when used properly, spammers abused this tactic as well. A brief thought experiment will illustrate the problem. Imagine that your market space has a total of 100 web sites in it. Perhaps 3 of these are high quality, another 5 are pretty high quality, 12 more are respectable, and so forth, as shown in Figure 7-13.

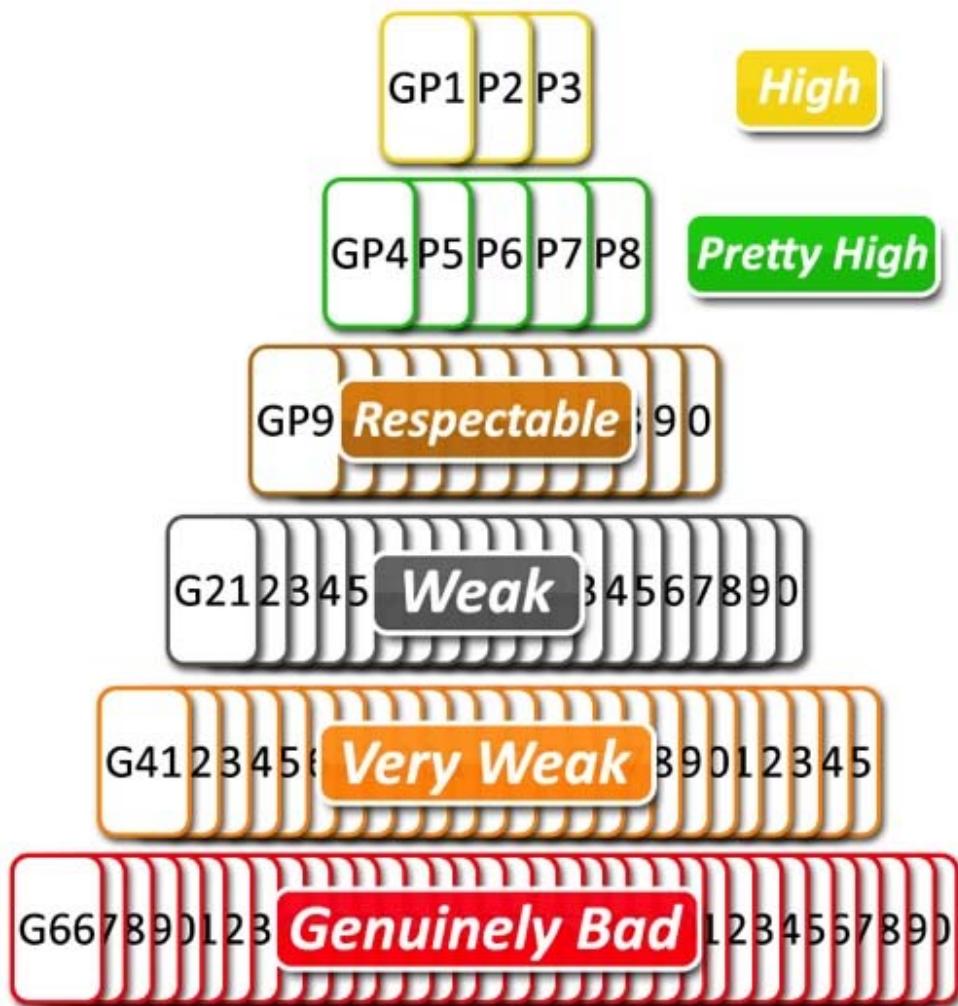


Figure 7-13 Mapping Site Quality in a Market Space

You can see in Figure 7-13, even if your first posts go on the very best sites in your market, by the time you have done 66 posts you are now writing posts on genuinely bad sites. It does not make sense for Google to treat this as more valuable as having ongoing relationships with the high authority sites in your market.

In the case of guest posts, as well as many other content marketing strategies, you are far better served to obtain a smaller number of guest post placements on higher authority sites, and in fact, get repeat links from those sites with ongoing posts. In fact, don't "guest post", but seek more stable regular contributor status on respected blogs. We will discuss this more later on in this chapter.

Source diversity

Getting links from a range of sources is also a significant factor. We already discussed two parts of this: getting links from domains you do not own, and getting links from many different domains, as well as the tradeoff between getting links from many domains vs. getting multiple links from one domain. However, there are many other aspects of this to consider.

For example, perhaps all your links come from blogs that cover your space. This ends up being a bit unbalanced. You can easily think of other types of places where you could get links: national media websites, local media websites, sites that are relevant but cover more than just your space, university sites with related degree programs, and so on.

You can think about implementing content marketing campaigns in many of these different sectors as *diversification*. There are several good reasons for doing this.

One reason is that the search engines value this type of diversification. If all your links come from a single class of sites, this is more likely to be the result of manipulation, and search engines do not like that. If you have links coming in from multiple types of sources, that looks more like you have something of value.

Another reason is that search engines are constantly tuning and tweaking their algorithms. If you had all your links from blogs and the search engines made a change that significantly reduced the value of blog links, that could really hurt your rankings. You would essentially be hostage to that one strategy, and that's not a good idea either.

It's a good idea to evaluate your source diversity compared to your competitors. Figure 7-14 shows an example of this using LinkResearchTools.com.

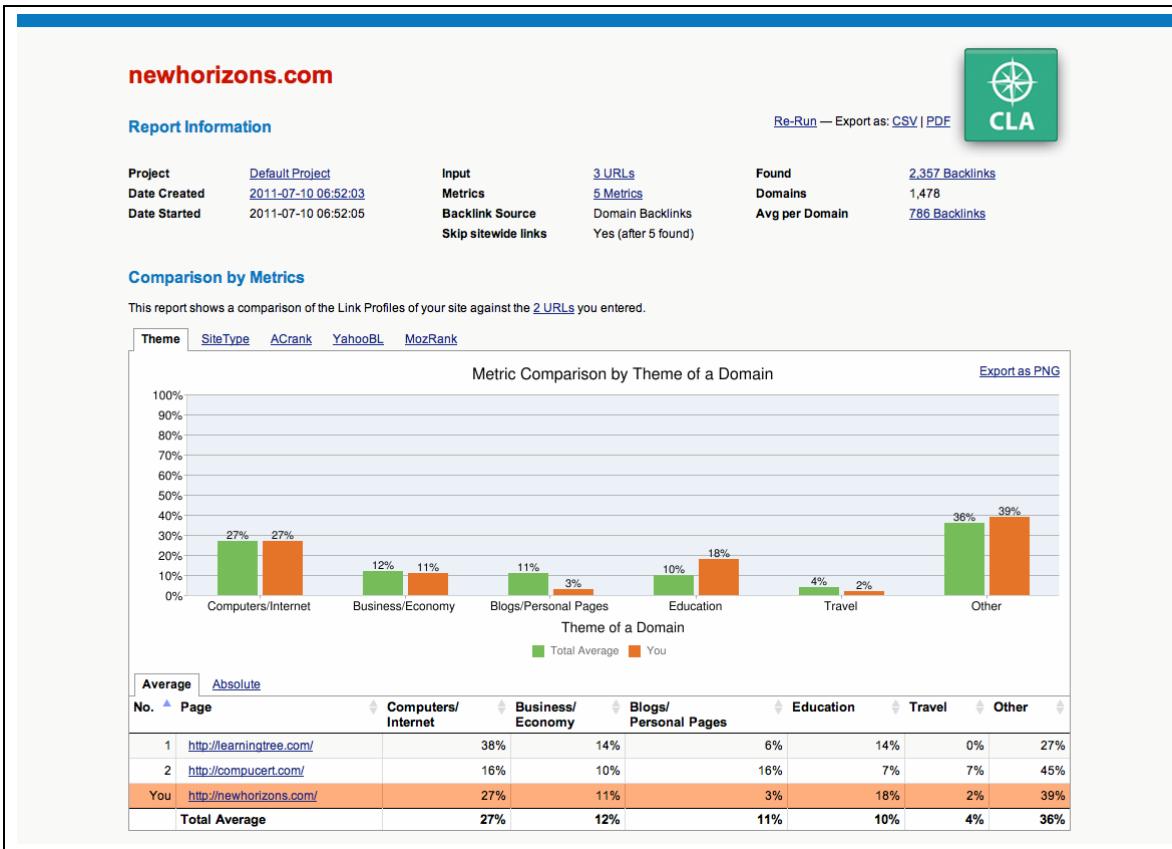


Figure 7-14 Comparing Link Diversity against Competition

Temporal factors

Search engines also keep detailed data on when they discover a new link, or the disappearance of a link. They can perform quite a bit of interesting analysis with this type of data. Here are some examples:

When did the link first appear?

This is particularly interesting when considered in relationship to the appearance of other links. Did it happen immediately after you received that link from the *New York Times*?

When did the link disappear?

Some of this is routine, such as links that appear in blog posts that start on the home page of a blog and then get relegated to archive pages over time. However, if a link to your site disappears shortly after you made major changes to your site that could be seen by the search engines as a negative signal, because it could be that the link was removed because the changes you made lowered the quality of the site.

How long has the link existed?

You can potentially count a link for more, or less, if it has been around for a long time. Whether you choose to count it for more or less could depend on the authority/trust of the site providing the link, or other factors.

How quickly were the links added (also known as link velocity)?

Such drastic changes in the rate of link acquisition could also be a significant signal. Whether it is a bad signal or not depends. For example, if your site is featured in major news coverage it could be good. If you start buying links by the thousands it would be bad. Part of the challenge for the search engines is to determine how to interpret the signal.

Context/relevance

Although anchor text is a major signal regarding the relevance of a web page (though receiving greater scrutiny since March 2014), search engines look at a much deeper context than that. They can look at other signals of relevance. Here are some examples of those:

External links to the linking page

Does that page containing the link to your site have external links as well? If the page linking to your site is benefiting from links from third party sites, this will make the link to your site more valuable.

Nearby links

Do the closest links on the page point to closely related, high-quality sites? That would be a positive signal to the engines, as your site could be seen as high-quality by association. Alternatively, if the two links before yours are for Viagra and a casino site, and the link after yours points to a porn site, that's not a good signal.

Page placement

Is your link in the main body of the content? Or is it off in a block of links at the bottom of the right rail of the web page? Better page placement can be a ranking factor. This is also referred to as *prominence*, which has application in on-page keyword location as well. Google has a patent that covers this concept called the Reasonable Surfer Patent (analyzed by Bill Slawski here: <http://www.seobythesea.com/2010/05/googles-reasonable-surfer-how-the-value-of-a-link-may-differ-based-upon-link-and-document-features-and-user-data/>).

Nearby text

Does the text immediately preceding and following your link seem related to the anchor text of the link and the content of the page on your site that it links to? If so, that could be an additional positive signal. This is also referred to as *proximity*.

Closest section header

Search engines can also look more deeply at the context of the section of the page where your link resides. This can be the nearest header tag, or the nearest text highlighted in bold, particularly if it is implemented like a header (two to four boldface words in a paragraph by themselves).

Overall page context

The relevance and context of the linking page are also factors in rankings. If your anchor text, surrounding text, and the nearest header are all related, that's good. If the overall context of the linking page is also closely related, that's better still.

Overall site context

Last is the notion of the context of the entire site that links to you (or perhaps even just the section of the site that links to you). For example, if hundreds of pages are relevant to your topic and you receive your link from a relevant page, with relevant headers, nearby text, and anchor text, these all add to the impact, more than if there happens to be only one relevant page on the site.

Source TLDs

Indications are that there is no preferential treatment for certain top-level domains (TLDs), such as .edu, .gov, and .mil. It is a popular myth that links from these TLDs are inherently worth more than links from other TLDs such as .com, .net, and others, but it does not make sense for search engines to look at it so simply.

Matt Cutts, the head of the Google webspam team, commented on this in an interview with Stephan Spencer (<http://www.stephanspencer.com/matt-cutts-interview/>):

There is nothing in the algorithm itself, though, that says: oh, .edu—give that link more weight.

And:

You can have a useless .edu link just like you can have a great .com link.

There are many forums, blogs, student pages, and other pages on .edu domains that spammers might be able to manipulate to gain links to their sites. For this reason, search engines cannot simply imbue a special level of trust or authority to a site because it is an .edu domain. To prove this to yourself, simply search for *buy viagra site:edu* to see an indication of how spammers have infiltrated .edu pages.

However, it is true that .edu domains are often authoritative. But this is a result of the link analysis that defines a given college or university as a highly trusted site on one or more topics. The result is that there can be (and are) domains that are authoritative on one or more topics on some sections of their site, and yet can have another section of their site that offers much less value, or that spammers are actively abusing.

Search engines deal with this problem by varying their assessment of a domain's authority across the domain. The publisher's *http://yourdomain.com/usecdars* section may be considered authoritative on the topic of used cars, but *http://yourdomain.com/newcars* might not be authoritative on the topic of new cars.

Ultimately, every site and every page on every site gets evaluated for the links they have, on a topic-by-topic basis. Further, each section of a site and each page also get evaluated on this basis. A certain link profile gives a page more authority on a given topic, making that page likely to rank higher on queries for that topic, and providing that page with more valuable links that it could then give to other websites related to that topic.

Link and document analysis combine and overlap, resulting in hundreds of factors that can be individually measured and filtered through the search engine algorithms (the set of instructions that tells the engines what importance to assign to each factor). The algorithms then determine scoring for the documents and (ideally) list results in decreasing order of relevance and importance (rankings).

How Search Engines Determine a Link's Value

A smart content marketing campaign typically starts with researching sites that would provide the best visibility and reputation benefits for the publisher (consider conflict with prior assertion of, “Any campaign that starts with “getting links” as the objective...”). However, it may also be useful to have an understanding of how search engines place value on a link. Although there are many metrics for evaluating a link as previously discussed, as an individual link builder many of those data items are hard to determine (e.g., when a link was first added to a site).

It is worth taking a moment to outline an approach that you can use today, with not too much in the way of specialized tools. Here are factors you can look at:

- The relevance of the linking page and of the linking domain to your site.
- The PageRank of the home page of the site providing the link. Note that Google does not publish a site's PageRank, just the PageRank for individual pages. It is common among SEO practitioners to use the home page of a site as a proxy for the site's overall PageRank, since a site's home page typically garners the most links. You can also use the Domain Authority from Moz's Open Site Explorer tool to get a third-party approximation of domain PageRank.
- The perceived authority of the site. Although there is a relationship between authority and PageRank, they do not have a 1:1 relationship. Authority relates to how the sites in a given market space are linked to by other significant sites in the same market space, whereas PageRank measures aggregate raw link value without regard to the market space.

So, higher-authority sites will *tend* to have higher PageRank, but this is not absolutely the case.

- The PageRank of the linking page.
- The perceived authority of the linking page.
- The location of the link on the linking page
- The number of outbound links on the linking page. This is important because the linking page can vote its passable PageRank for the pages to which it links, but each page it links to consumes a portion of

that PageRank, leaving less to be passed on to other pages. A simple way of expressing this mathematically is as follows:

For a page with passable PageRank n and with r outbound links:

$$\text{Passed PageRank} = n/r$$

It is likely that the actual algorithm used by the search engines is different. For example, the amount of PageRank may vary based on where it is on the page. Google has a patent which discusses the concept of putting more value on links that are more likely to be clicked on by a visitor to the page, based on where the link is placed on the page or how it fits into the overall context of the page.

It can be useful to organize this data in a spreadsheet, or at least being consciously aware of these factors. For many businesses, there will be many thousands of prospects, and other factors associated with content marketing campaigns should also be taken into account. These include such factors as the impact on your reputation, the potential for developing relationships with other influencers, social media sharing potential, and more.

Creating Content That Attracts Links

Since we have established the importance of the development of quality content as the precursor to earning quality links, and that links are still a very important overall ranking signal in SEO, the next logical question is how to go about getting these two interdependent objectives accomplished.

How are Links Earned?

It is important to step back and examine why links are created in the first place. Why did that person decide to link to that particular website from their own site? There are many possible reasons:

1. The site owner was paid for adding the link (this would technically be considered a paid, rather than earned link). Although this is a perfectly legitimate reason, in the search engine's eyes it carries no editorial value (and search engines may even penalize sites for linking or acquiring links in this fashion). Buying links for purposes of increasing your search rankings should not be part of your campaigns.
2. Links were traded between sites. Also called *reciprocal linking*, the practice of trading links between sites is popular. However, search engines view this as barter and therefore as having limited editorial value. That being said, there are often legitimate and necessary reasons for two separate websites to link back to one another, without the "purpose" being link bartering; one should not sacrifice user experience or business needs in this type of scenario by "fearing" that the linking relationship will be frowned upon by the search engines; rather, simply be aware that the linking relationship may not provide any benefit to either party from an inbound link perspective.
3. Something on your site triggered an emotional reaction from the publisher, causing the publisher to link to your site. For example, perhaps your site had the funniest cartoon the publisher ever saw, or it offered an inflammatory political opinion.
4. The publisher may see something of value on your site and wants their site visitors to know about it. The majority of the highest-value links are given for this reason.
5. A business relationship comes into play. For example, you may have a network of distributors and resellers for your product. Do they all link back to you? See #2 above.

How Can Sites Approach Getting Links?

The keys to acquiring links are points 3, 4, and 5 in the preceding list. Understanding these link triggers is the key to successful content marketing and earned link building, as follows:

- Since creating emotional reactions can result in links, building content that plays to the emotions of potential linkers can be a powerful tool for obtaining links. Of course some content ideas may not fit a

company's brand image, so it is important to come up with ideas that are consistent with the brand image for your company.

- Create quality reference material. Providing unique and powerful information to users can be a great way to get links, particularly if you can make sure the publishers of authoritative sites in your market space learn about what you have created, including why it would be of value to their website visitors.
- Leverage business relationships. In our example, we suggested that you might be someone who has a network of resellers. If this is your business model, having a link back to you as a standard term in your reseller agreement is sound business.

These are some simple examples. There are many approaches to obtaining additional earned links to your site.

Introduction to Content Marketing

Creating great content, publishing it, and then promoting it effectively for increased business visibility can be referred to as **content marketing**. This can be a great way to build your reputation online, and can be used to bring prospective customers to your web site, as well as obtain high quality links to your site.

As always, there is a right way and a wrong way to go about doing this. The key realization is that the primary goal for a content marketing campaign should be to build your reputation and visibility online. Consider what Matt Cutts had to say in a 2012 interview he did with Eric Enge (<http://www.stonetemple.com/matt-cutts-and-eric-talk-about-what-makes-a-quality-site/>):

“By doing things that help build your own reputation, you are focusing on the right types of activity. Those are the signals we want to find and value the most anyway.”

Here are some guidelines on how to approach your content marketing plan:

- Focus on developing the best strategy to build your reputation online. Make this your first priority. Publishers that focus on obtaining links as their first priority can too easily lose their way and start doing the types of things that the search engines do not like.
- Develop a content plan that closely relates to your business, mission, or vision. The great majority of your content should be on topic, though it can also be quite effective to publish occasional pieces of content that are off topic if they are worth remarking about, and that attract attention to you.
- In an ideal world, build a plan that combines publishing great content on your site, a strong social media plan, and publishing great content on authoritative third party sites. This provides you with a way to build your audience by gaining exposure in places where your target customers can already be found with giving them a reason to visit your site.
- While building your reputation is the primary objective, you should also look to obtain links back to your site in the process. When you publish content on third party sites make sure that any links back to you are ones that a reader of the article might be interested in clicking on. This is a key way to validate that you are implementing your content marketing effort in an appropriate way.

Building a powerful content marketing campaign will likely take some time and you can't be afraid to experiment with different ideas and learn what works for you. If you are a small business owner with very little time to invest, you may need to be quite a bit more focused and implement only pieces of a full plan as well.

In addition, as you build a content marketing plan, make sure to consider more than just the SEO benefits. While this book is focused on SEO, the impact of content marketing is greater than that. It can play a lead role in defining your brand, and in your overall reputation and visibility online.

The general PR and marketing efforts of your company are a part of content marketing as well. Take the time to integrate your SEO-focused efforts with the PR and marketing plans for the business. You will find that this provides a great deal of leverage, and help your overall campaigns get better results.

Using Content to Attract Links

While building your reputation and exposure to more potential customers and interested users is the primary focus of a content marketing campaign and the content you develop is of value to those users, understanding how to leverage your content assets for link development purposes becomes a no-brainer. More links to your content results in more users exposed to your content – furthering the overall objectives of your content marketing efforts. That these links also benefit your SEO efforts only reinforces the fact that learning how to create and promote your quality content online, and take advantage of linking opportunities you create, is a very valuable endeavor. Superior content and tools are the keys to making all of this happen.

Content Marketing Basics

Some of the best links are obtained indirectly. Duane Forrester of Bing puts it this way: “You want links to surprise you. You should never know in advance a link is coming, or where it’s coming from. If you do, that’s the wrong path” (<http://blogs.bing.com/webmaster/2014/05/09/10-seo-myths-reviewed/>). This position may be extreme if you never knew in advance about a coming link, but the concept is a solid one. If you publish great content, and people learn about it, some of those people will link to it. Your job is to make the content more easily discovered and valued enough to endorse.

Figure 7-15 Shows a very simple form of this concept, where a publisher places content on their blog, shares it on their social media feeds, and as a result gets people to go read the article, gets new subscribers, and obtains links to their content.



Figure 7-15 Synergy between Publishing Great Content and Social Media

In addition, notice how the social media platform benefits as well, by obtaining new followers, due to the fact that content is shared is of high quality and relevance to the audience. This concept is one of the core components of a content marketing strategy. The next step is to take this one stage further and build relationships with others in your market community, including influencers. Figure 7-16 shows how that can accelerate the effect shown in Figure 7-15.

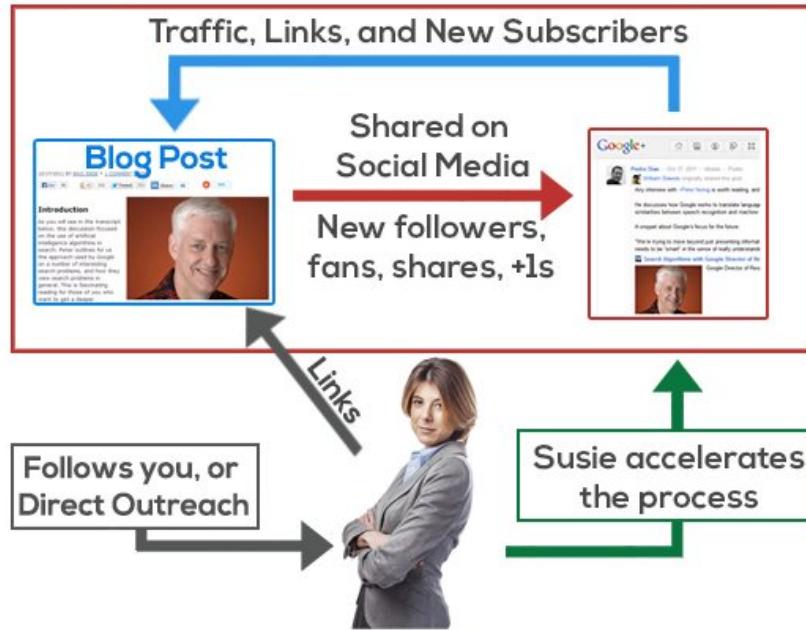


Figure 7-16 Influencers Act as an Accelerant to Content Marketing Campaigns

There are many other methods and techniques that make up content marketing which we will discuss throughout the remainder of this chapter.

Content Types

Normally, there are many different types of content a site could produce. The concept is simply to identify your most important target audiences and what content will most resonate with them, and then tweak the content plan accordingly (we will discuss this more in “Identifying Personas and Targeting Content”). Keyword research can also help identify content related to your target market, and can play a role in identifying topics that may help build your visibility.

- Posts on your own blog or web site
- Posts on 3rd party sites (**guest posts**)
- Downloadable tools
- Videos
- Images & animated gifs
- Podcasts
- Screencasts
- Presentations (including Slideshare)
- PDF files
- Plugins
- Memes
- Social media posts
- Comments on the posts of others
- Curated content
- Original research & data streams
- Comprehensive reviews

- Explanatory journalism
- Scoops
- Infographics
- Personality tests
- Comics & illustrations
- Interesting interviews
- Mobile or tablet apps

This list of content types could be made much longer. Your creativity is the main limit here. The content type should be picked based on what will provide the biggest impact on your target audience.

Content Marketing Strategies

Content can be marketed in many ways. Some of the most basic strategies include:

Guest posting

This is the practice of creating new content for publication in the blog or article stream of a third party web site. We will discuss this more later on in this chapter in “Guest Posting” in Developing content for third party sites.

Rich Content for Third Party Sites

This is very similar to guest posting, except it is not intended for a third party site’s blog. This type of content is designed for publication on static pages of the partner’s site. This will be covered in “How to Design Content for Third Party Sites”.

Content syndication

You may have quality content placed on your site that others are interested in republishing on their sites. This can be a very effective tactic, but it can come with some SEO risks if it is not done properly. We will discuss this more in “Proper Use of Content Syndication”.

Social media

Social media sites such as Facebook (<http://www.facebook.com>), Google+ (<http://plus.google.com>), LinkedIn (<http://www.linkedin.com>), and Twitter (<http://www.twitter.com>) can be used to promote content on your sites. These will also be discussed more in “Social Media Networking”.

Viral Content Creation

Publishing content that has the potential to go viral can help you rapidly gain exposure to your site and your business. This is not always so easy to do, but this is discussed more in “Creating Link Worthy or Viral Content”.

Developing Content That Works

At the heart of any successful content marketing campaign is the content itself. What makes up high quality content? How do you come up with ideas? Who is going to create the content for you? How do you measure success? You must address these questions if you are going to be successful.

Having access to a subject matter expert (SME) is invaluable during this process. Hiring a writer who does not know your market is highly unlikely to yield good results unless you give them substantial time to learn your business, the competition, the marketplace, and what content is already out there.

Your SME can provide a lot of value to your overall plan, and will either have to write, or review, edit, and approve whatever content you create. You don’t want to end up publishing content that is factually incorrect, or whose positioning will be bad for your brand.

That said, most SMEs are not also experts at coming up with creative ideas, and it can often be tedious for them to keep coming up with one content idea after another. To do that, you need to come up with a process for generating such ideas on a regular basis.

Content Ideation and Creativity

Coming up with content ideas is one of the most important parts of a content marketing campaign. Without the right type of content, your efforts to market that content will fail. Creativity is required to come up with content ideas that are distinct enough to meet your needs.

This is because it's likely that there are already high quantities of content that relate to your potential topic matter already. Figure 7-17 shows an example of a search query that reveals that more than 6,000 pages exist on the web that use the phrase "building a deck" in the title of the web page.

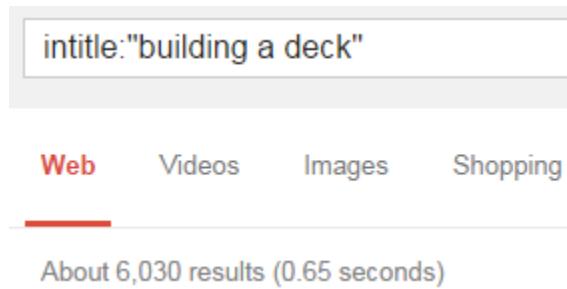


Figure 7-17 Simple Content Ideas Are Not Good Enough

The notion of "being creative" can be a frightening one, especially to those who do not work at it on a regular basis. However, the reality is that it's a myth that some people are born to be creative, and that others are not. Success at being creative requires practice. Through trial and error you work what works, and what doesn't.

In addition, active research techniques can help you come up with ideas. Here are some examples of methods for coming up with ideas that you can use:

Internal Brainstorming

This is a method that many companies overlook, yet is often one of the best techniques to use. Gather key members of your management and marketing teams in a conference room and brainstorm together. Chances are there are many people who know a lot about your business, and it will surprise you how many good ideas they can come up with.

Competitive Analysis

Spend time researching what content your competitors publish. Check out their site and/or blog in detail, and follow their social media accounts. If you have a competitor that is actively pursuing content marketing, this can be a goldmine of ideas for you.

Social Media Research

Social media sites such as Facebook (<http://www.facebook.com>), Google+ (<http://plus.google.com>), LinkedIn (<http://www.linkedin.com>), and Twitter (<http://www.twitter.com>) can also be used to help generate content ideas. Follow major influencers in your market and see what content they are publishing.

Keyword Research

Research what types of phrases people are searching on to see what common customer needs are. This can provide insights into what types of content are likely to draw the best engagement. You can read more about keyword research in Chapter 5.

Google Suggest and Bing Suggest

This approach offers similar value to keyword research, but the data is coming straight from the search engines. Figure 7-18 shows this in action in Bing. Notice how you can see common phrase variants of what you have typed in so far, and these can provide you valuable info on what users are looking for.

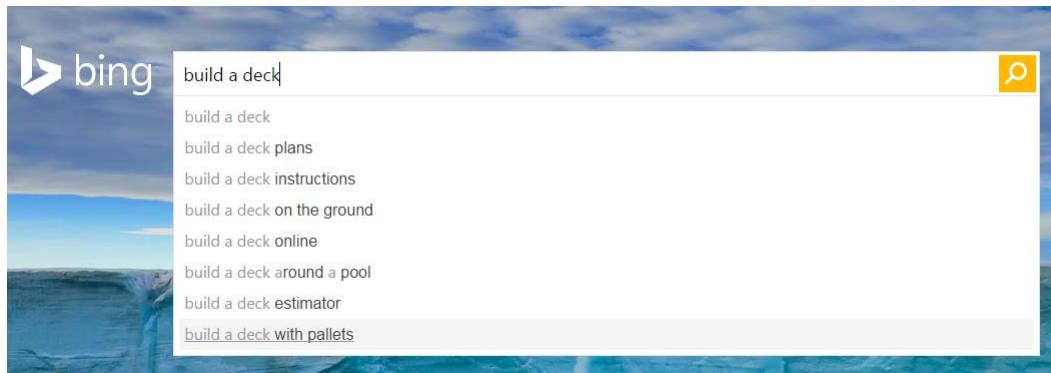


Figure 7-18 Bing Suggest Can Provide Content Ideas

Question and Answer Sites

Some question and answer sites, such as Yahoo! Answers (<http://answers.yahoo.com>) are very active. Mining these sites can also show you the types of things that potential customers want to know about. Figure 7-19 shows an example from Yahoo! Answers.

A screenshot of the Yahoo Answers website. The top navigation bar includes the 'YAHOO!' logo, 'ANSWERS' button, a search bar, and 'Search Answers' and 'Search Web' buttons. The main content area shows a question titled 'How much does a bale of hay cost in California?'. Below the question are several answers. The first answer, by 'fluffernut', says: 'Bale Of Hay Cost: Cheap - Everyone Wants to Pay a Low Price. www.NexTag.com Best Value for Bale Of Hay Cost!'. The second answer, by 'ceeli', says: 'Top Bale Hay - Top Prices | Bale-Hay.BEST-PRICE.com www.Bale-Hay.BEST-PRICE.com Compare prices online & save up to 75% on Bale Hay at BEST-PRICE.com!'. The sidebar on the left lists various categories: Answers Home, All Categories, Arts & Humanities, Beauty & Style, Business & Finance, Cars & Transportation, Computers & Internet, Consumer Electronics, Dining Out, Education & Reference, Entertainment & Music, Environment, Family & Relationships, Food & Drink, Games & Recreation, Health, Home & Garden, Local Businesses, and More & Finance.

Figure 7-19 Yahoo Answers can be a Source of Ideas

There are many ways to learn what types of information people are looking for. This can be time consuming research, but it is well worth the effort!

Speedstorming

You can also try different techniques to help stimulate creative thinking. One of these is referred to as **speedstorming**, and all you need to implement this method is 5 people who have some familiarity with your market, 5 blank sheets of paper, and a watch. Here is how it works:

1. Sit all 5 people at a table and give them each a blank sheet of paper. Make sure that they know that during this exercise they can't talk to each other about their ideas, or look at one another's sheets of paper.
2. Tell them all to come up with 3 content ideas, give them 5 minutes, and set the timer on your watch for 5 minutes.
3. When the watch timer goes off, have them all pass the paper to the person on their left.
4. Repeat steps 2 & 3 five times, except the final time have everybody give their sheets to the person leading the content marketing effort.

Just like that you should have 75 ideas to consider for your campaign. Even if two thirds of the suggested ideas are not particularly useful, this would still leave you with 25 pretty decent ideas to start from.

What makes this process work well is that it gets all 5 people involved. If you were to try and do the same thing with a whiteboard, you would likely get far fewer ideas, and some of the people in the room would contribute little, often because they are shy. This process gets everyone involved, and it's actually fun as well!

Get Creative Help

You can also bring in someone to help with the process. Since you can develop creative skills with practice, leveraging people that already have that practice is often quite helpful. Here are a few ways to do that:

- Hire Creative Experience: Nothing like making that person a full time member of your team!
- Leverage your Subject Matter Expert: Someone who is intimately familiar with your market or technology may not be an expert at creative thinking, but their knowledge can be invaluable in recognizing interesting and/or unique ideas.
- Talk to your Customer Service Team: As the frontline with the customer, this team almost always can provide input on what types of content are in demand by your customers (and therefore your prospective customers).
- Bring in a Contactor: This can still help you accelerate your creative processes.
- Sign on an Agency: Agencies can be quite effective as well, as they can leverage their experience across many clients.

These are all great options to try and help accelerate the process. Just be aware that neither of these 5 options is a magic bullet. Prior creative experience is helpful, but you will still need to give these people time to familiarize themselves with your business, your market, what your competitors, all the things that we discuss in this chapter.

In addition, before hiring anyone, either as an employee, contractor, or agency, review their past work, and have them show you why they are a good choice for you. Ideally, they will be able to show you what they have done in markets that have similarity to yours.

For example, if you operate a B2B type business, you should seek out help that has prior B2B experience.

Repurposing Content

Sometimes, the content idea you are looking for is already in the palm of your hands. For example, if you are looking for an idea for a 250 word post to put into your LinkedIn account (using LinkedIn Publishing: http://help.linkedin.com/app/answers/global/a_id/47445/~/long-form-posts-on-linkedin---overview), you may already have written about something that could meet your need quite easily.

Go back and review some of the recent articles that you have published as guest posts on 3rd party sites, or as a post on your own blog. Is there an idea in there that you could extract and expand upon? Perhaps you wrote only one paragraph in your article, but it's worth explaining in more detail. Since you already have your arms around the concept, creating that 250 word version for LinkedIn (or Google+ for that matter) should be easy.

Similarly, if you are looking for a great image to put up on Pinterest, you may have already created it for one of your articles. Or, if you laid out a great 10 step process to do something in one of your articles,

would that be something you could easily spin into a presentation on Slideshare? Can you film a video about key concepts in one of your articles and upload it to your YouTube channel?

These are just a few examples of ways to repurpose your content. Remember, each medium is different, so you need to put some effort into the production of the content for that particular platform. This will inevitably morph the content somewhat, and that is where the effort will come in. However, the core concepts you are putting out there remain the same that you started with, and that is often half the battle!

What Makes Content “High Quality”?

High quality content is at the heart of achieving marketing nirvana—having a site, or a page, or a tool, or a series of videos so good that people discover it and link to it without any effort on the your part. This can be done, but it does require that you create content that truly stands out for the topics that its site covers, and that it thinks about your target audience in every aspect of the publishing process.

One factor is the content mix. You do want to create content that draws lots of attention, social shares, and links to your site. This type of content will be the driver of what helps you site with SEO. However, there is a need for other types of content on your site too. This can take two forms:

1. Content that addresses the basics of what your products and services do. This is meant and potatoes content that helps people whether or not to buy from you. Most likely, it will never attract links or social mentions, but it's a core part of your site. Creating great content here is an important part of building your reputation too.
2. Non-commercial content that helps users in areas related to your business. This is content that also helps build your reputation in your community by providing helpful advice and information, but the number of links and social mentions it gets is relatively small. You still want this type of content on your site too.

Emotional Triggers, Titles, and Images

This is the stuff that acts like a link magnet. It gets lots of social shares, +1s, Likes, Pins, and other social actions. What are the keys to doing that? Content that generates an emotional reaction is usually what gets shared and linked to the most. The more the reader can relate to it, the better.

None of this will work for you unless you get someone to look at your content. Two of the most important factors in making that happen are the title of the article, and the initial image you associate with it. If these can generate an emotional response from your target audience, your chances of success go up significantly. Figure 7-20 Shows data from BuzzSumo on posts that received a large number of social shares on the topic of taxes.



Figure 7-20 Popular Posts on the Topic of Taxes

The titles evoke an emotional response from the audience which sees them. You want to pay more taxes for driving? Bill and Hillary are not paying theirs? That rich capitalist did what? The headline alone compels readers to look closer to get more information.

Once you open the posts, you will see that most of them have a photo or image that pull you in further. The initial image in the article is also a big factor in how much your article will be shared or mentioned in social media. You can use Google image search to help you find images that evoke emotions in line with the title of the article, but use this just to get ideas. Then take the proper steps to obtain images for which you can obtain the proper license to use. If that license requires that you provide attribution make sure you do that as well.

The impact of images is quite notable. In December of 2014, co-author Eric Enge published a study (<https://www.stonetemple.com/twitter-engagement-unmasked/>) that showed including images in a tweet more than doubles the number of Retweets your Tweets will get (on average).

Consider the photo at the top of the Bloomberg article which is shown in Figure 7-21. The image of Bill and Hillary Clinton laughing only enhances the emotional response from the article title.



Figure 7-21 Bill and Hillary Clinton

The Huffington Post article also uses an image sure to generate a reaction, as shown in Figure 7-22



Figure 7-22 Rich Capitalist Smiling

The concept of a highly privileged and very wealthy man comparing an increase in his taxes to what happened to the Jews in World War II is shocking and disturbing.

Whatever your political leanings may be, these are powerful images. However, exercise care, and most likely, both of these photos are taken completely out of context. Using images inaccurately can lead to problems as well. For many brands it's best to avoid publishing anything which is misleading as this may cause your target audience to question the accuracy of other statements you make, such as the value of your products and services.

How do people come up with these ideas? It is a myth that you have to be born with a creative gene. Most people who are good at coming up with creative ideas have a lot of practice, and over time they have learned what works and what doesn't. You have to start by trying and know what you will get better. However, there are things that you can do to improve your chances of success.

One idea is to use title tag generation tools to get ideas for great titles for your article. Examples of such tools include: <http://www.impactbnd.com/blog-title-generator/blogabout>, <http://www.portent.com/tools/title-maker/>, <http://www.contentrow.com/tools/link-bait-title-generator>, and <http://tweakyourbiz.com/tools/title-generator/index.php>. Each company that produces these tools uses them to help them generate business for themselves (content marketing in action), but you can use the tools to help you come up with great title ideas.

Joan Stewart also wrote an article on “7 Tools to Generate Killer Headlines” (<http://thefutureofink.com/7-tools-to-create-killer-headlines/>) that discussed this topic, and she suggests many different strategies for coming up with title tags. One tool she suggests is the Emotional Marketing Value Headline Analyzer (<http://www.aminstitute.com/headline/>) from the Advanced Marketing Institute. You can also use **speedstorming**, a technique that was discussed earlier in this chapter.

You can also setup content idea brainstorming sessions. Have 4 people sit in a room start with a blank sheet of paper. Give them all 5 minutes to write down 3 ideas for articles, without sharing their ideas with the others. At the end of the 5 minutes, have everyone pass the paper to the person on their left. Now have them write 3 new ideas on the paper that was just passed to them. Repeat this 2 more times so that each sheet of paper has 12 ideas on them.

At the end of this you now have 48 different article ideas. Even if only 5 of these are worth pursuing, you have succeeded in coming up with 5 great ideas in less than 1 hour. That's a very productive use of time.

You can also consider using a Meme generating tool such as <http://www.memegenerator.net>. This can be an interesting way to take an existing photo or image, layer on some text, and create an emotional response as well. We discuss the use of memes more in “Leveraging Memes” in this chapter.

Make sure the resulting content campaign meets your business purpose. Building your reputation and visibility with your target audience, and obtaining links to your site, are the core goals of content marketing campaigns.

It's OK if some of your campaigns are not literal interpretations of your keywords, or products, but they should always support the general positioning of your brand. Consider Redbull's sponsorship of Felix Baumgartner's jump from space or Nik Wallenda's tightrope walk over Niagara Falls. The videos and content produced as a result are not about energy drinks at all. Why does RedBull promote these types of campaigns? Mack Collier explains their positioning as follows: ““Red Bull isn't selling an energy drink, it's selling what happens after you drink it” (<http://www.mackcollier.com/red-bull-content-marketing/>).

Brands can also sponsor major charities, and these can help create a positive image, and drive links and social mentions as well. For example, in 2014 Zynga offered mobile users a chance play their popular Words With Friends game against John Legend, Snoop Lion or Eva Longoria in their first-ever Words With Friends Celebrity Challenge, and gave up to \$500,000 to charity based on game results.

This campaign offered a great combination of using their games and the opportunity to engage with celebrities to help them drive reputation and visibility. However, it's important to ensure that your brand is a reasonable match to the cause (Futurity offers some thoughts on this here: <http://www.futurity.org/brands-sponsor-charities-need-fit-cause/>), and that you have worked out a strategy for how to make the most out of that sponsorship (some great suggestions from Vehr Communications here: <http://www.vehrcommunications.com/8-tips-to-make-charitable-sponsorships-work-for-you/>).

Here are some more rules that you can follow to maximize your results:

- Use content that helps establish your site as a leading expert on its topic matter. When you produce high-quality stuff, it builds trust with the user community and increases your chances of getting links. This also helps you with social engagement signals which are discussed more in “How Social Acceptance and User Engagement Play a Role in Rankings” in Chapter 8.

- Minimize the commercial nature of the content pages. As an extreme example, no one is going to link to a page where the only things they see above the fold are AdSense ad units, even if the content below it is truly awesome. Of course, there are less obvious ways to be too commercial, such as self-promotion in the areas around the content or obtrusive overlays and animations.
- Do not put ads in the content itself or link to purely commercial pages unless such pages really merit a link based on the content. No one wants to link to a commercial (except in rare cases like a really awesome Superbowl commercial).
- Do not disguise the relationship between the content and the commercial part of your site. This is the opposite side of the coin. If you are a commercial site and you hide it altogether, you run the risk of being viewed as deceitful.

When content is published on your site, you have other decisions to make, such as whether the content goes in a special section or whether it is integrated throughout the site.

As an example of how you can make this decision, an e-tail site that publishes a large catalog of products may not want all (or some of) the pages in its catalog laden with a lot of article content. A site such as this might build a separate section with all kinds of tips, tricks, and advice related to the products it sells.

On the other hand, an advertising-supported site might want to integrate the content throughout the main body of the site.

The Power of Memes

Richard Dawkins originally defined the word meme as “a package of culture.” In more recent history, Dr. Susan Blackmore, psychology scholar and TED lecturer on “memetics,” defines memes as “a copy-me instruction backed up by threats and/or promises.” An example of such a “threat” might be the last bit of an email chain letter that warns of “7 years of bad luck if you don’t forward this email in the next 10 minutes to 7 friends.”

However, memes could relate to any type of fad. For example, bell bottoms became popular because one person saw the promise of a better social standing with the new pants and because suppliers saw the promise of increased profits. Memes mark the rise and fall of all of the fashions and trends in the history of the world.

The definition most people are familiar with, however, looks a little more like what is shown in Figure 7-23.



Figure 7-23 Examples of Memes

Memes, in the way we recognize them, are a popular Internet trend in which an image is paired with a clever phrase to create a relatable or funny situation. In the example of the “Business Cat” (<http://knowyourmeme.com/memes/business-cat>) meme, one meme makes fun of another meme, the popular YOLO (“you only live once”) hashtag made popular through Twitter. In all cases, people enjoy the memes because they remind them of something familiar.

Using Memes in Your Marketing Campaigns

Unfortunately, some marketers are overlooking memes as a great way to integrate popular Internet humor into their campaigns. Marketers should take advantage of them because:

1. They are easy to create. With websites like <http://memegenerator.net>, you can simply create your own meme by filling in text boxes. You can upload your own image, or take advantage of the incredibly popular memes and get in on the fad.
2. They are cheap — as in, completely free. However, there may be copyright issues in certain cases, an issue which we will discuss more in a moment.
3. People feel in on a joke. The familiarity of the meme creates the blueprint of laughter; people know to expect a laugh.
4. It is a great way for your brand to seem relevant and fresh.

Meme Marketing in Action

Consider the following websites that have taken advantage of the popular macro meme Foul Bachelor Frog (<http://knowyourmeme.com/memes/foul-bachelor-frog>), which depicts a frog that reveals all of the sketchy, unsanitary secrets of bachelor life.

Diamond retailer Diamond Envy aggregated some of funniest examples and used them as “advice” for young bachelors who relate to the single lifestyle but may someday want to tie the knot — see: “Great Advice... If You Want to Stay Single!” (<http://www.diamondenvy.com/blog/foulbachelorfrog/>).

Canadian real estate site Zolo tailored the meme to its own uses with “The Foul Bachelor Frog Bachelor Pad” (<http://www.zolo.ca/blog/foul-bachelor-frog-bachelor-pad/>) pulling popular examples of “redneck home remodels” (<http://blog.builddirect.com/redneck-home-remodels/>) and pairing them with Foul Bachelor Frog to poke fun at some DIY upgrades to the prototypical bachelor pad.

Travel club World Ventures employed a variety of memes in addition to Foul Bachelor Frog, such as “Success Kid” (<http://knowyourmeme.com/memes/success-kid-i-hate-sandcastles>) and “Socially Awkward Penguin” (<http://knowyourmeme.com/memes/socially-awkward-penguin>) to tell a story in their article “Thank You, Internet Memes, for this Sage Travel Advice” (<http://www.worldventures.com/blog/read/thank-you-internet-memes-for-this-sage-travel-advice#sthash.m6ENlzF5.dpbs>).

All three found relevant, already-made examples — and, in the case of Zolo, made their own to complement existing ones.

Bear in mind that your meme-containing article is supposed to target (attract the interest of) online influencers who may choose to link to or reshare your meme in social media. Make sure you make this easy for people to do, but visibly including references (and the link) to the appropriate pages on your site within the core components of your meme.

Other Types of Memes

Memes don’t just come in the form of images — video memes are popular as well. Remember the “S*** People Say” video craze of late 2011/early 2012? Thousands of video parodies were created, each with their own specific punch line. The result was millions upon millions of collective views. Videos are easy to upload for free onto YouTube, which makes them incredibly shareable on social media sites. People send the linked videos to co-workers if there is one about their profession, or friends if it is about a shared hobby.

Twitter memes are among the easiest to get on board with and also among the easiest to benefit from. On a sociological level, each Twitter hashtag is a meme. The most popular hashtags of the day offer an easy

opportunity to get exposure. Once people search for the hashtag and see your company's response, they can either favorite it, retweet it (which means more free eyeballs on the promotional material with no effort), or — the prized action — they can subscribe to your company's Twitter feed.

Physical memes usually are a certain body movement or gesture done in unique, impressive, or humorous locations. Physical memes include the current “twerking” phenomenon and are a cheap way to participate through photos or videos.

Know Your Audience

Before creating your own meme-inspired marketing campaign, it is important to consider your audience. McDonald's had a Twitter fiasco in 2012, and this is a great example of a meme campaign gone in the wrong direction. McDonald's encouraged its Twitter subscribers to share their McDonald's experiences by using what Business Insider (<http://www.businessinsider.com/mcdonalds-twitter-campaign-goes-horribly-wrong-mcdstories-2012-1>) described as “a dangerously vague hashtag” #McDstories. You can see examples of some of the problems that developed in Figure 7-24.



Figure 7-24 Example of a Twitter Meme Campaign gone bad

Memes in Advertising

Some companies are realizing that the popularity and familiarity of memes can be used for their benefit outside of the Internet. For instance, there have been a slew of companies recently using meme-related images in outdoor advertising, mainly highway billboards.

Virgin Media took advantage of the popular “Success Kid” to initiate a familiarity with their desired consumers. Memes like this work well in advertising because the viewers already feel like they are “in on the joke,” and it feels somehow inclusive to them, but perhaps exclusive to others.

Brands have also integrated the idea of memes into full-scale television spot campaigns. For example, the popular commercial for UK dairy brand Cravendale featured “Cats with Thumbs” — one week earlier, a short video of a cat giving a thumbs up had gone viral. You can see a screen shot of this in Figure 7-25.

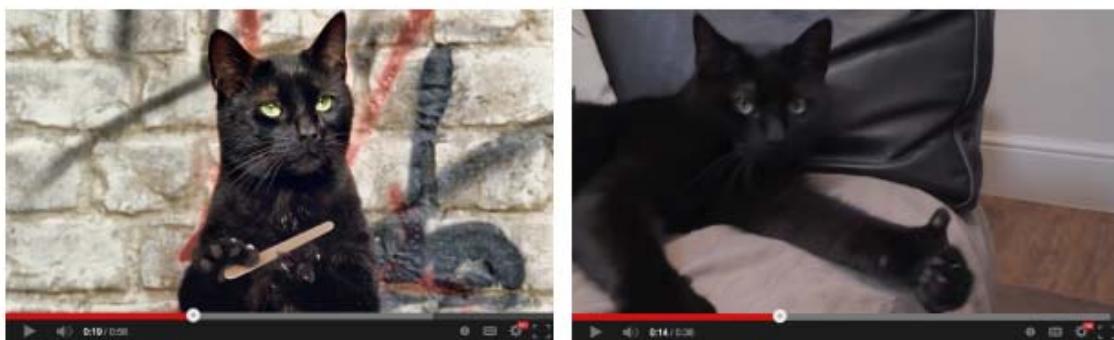


Figure 7-25 cats-with-thumbs.png

Play It Safe

It's important to consider intellectual property law before you embark on your meme mashups. Some works, like Futurama Fry, are owned by big Hollywood studios. Others, like Success Kid, are owned by a parent.

Also, commercial uses of a meme are, generally speaking, riskier than non-commercial uses. It's possible that your use would qualify as "fair use" under copyright law — parodies, criticism, comment, news reporting, teaching, scholarship, and research are examples of potential "fair use" cases. Nonetheless, unless the work is Creative Commons licensed or in the public domain, it's safest to seek the copyright owner's permission to use the image. You can read more about this subject in "I Can Haz Copyright Infringement? Internet Memes and Intellectual Property Risk." (<http://www.mondaq.com/unitedstates/x/208786/Copyright/I+Can+Haz+Copyright+Infringement+Internet+Memes+and+Intellectual+Property+Risks>).

The Role of Engagement in Content Marketing

Make a point of measuring how people engage with your content. Figure 7-26 provides an example of an article with social media sharing buttons, which is one great way to measure engagement:

Introducing The Periodic Table of Content Marketing

by Chris Lake | 18 March 2014 12:00 | 59 comments | Print

6,665

[Tweet](#)

2.2k

8+1

[2.2k](#)

[Like](#)

1,720

[Share](#)

56 Cm Competitions 60 Tw Twitter

57 Ch Checklists 61 Yo YouTube

I've written a lot about [content strategy](#) over the past decade. I've also highlighted various [niche tactics](#) that can help content creators to succeed, as well as plenty of [examples of excellent content](#). But I haven't created many visualisations, and recently I have been keen to do one.

Surprisingly, nobody has yet created a periodic table for content marketing, so I thought I'd have a go.

Before I introduce it, allow me to doff my hat at Dmitri Mendeleev, who first published the periodic table of elements. I'll also nod in the direction of Danny Sullivan, who created one based around [SEO success factors](#).

Let me also say that I hope that this is helpful, as the world is awash with dubious infographics and I really didn't want to produce something just for the sake of it.

Figure 7-26 Post Showing Lots of Engagement

Measuring engagement is a simple and straightforward way of determining:

- How your content is resonating with the visitors on the site where the content is published.
- How your content compares to other articles you publish on the same site. You can potentially learn from that how to improve what you do next time by seeing which of your articles do the best on a given site as it gives you insight into what the audience likes the best.
- If the content is a guest post on a news site with many other authors, you can compare how you did against the results they typically obtain on the same site. Figure 7-27 Illustrates this basic concept by showing how the Susie Powers article received far more social shares than the one from George Lackey. While search engines do not use these social share signals directly as proof of the quality of content, you can use it to see what content is resonating with the audience on a given web site.

- If you publish many articles on a third party site, and you are not getting engagement on any of them, perhaps it is not the right place for you!

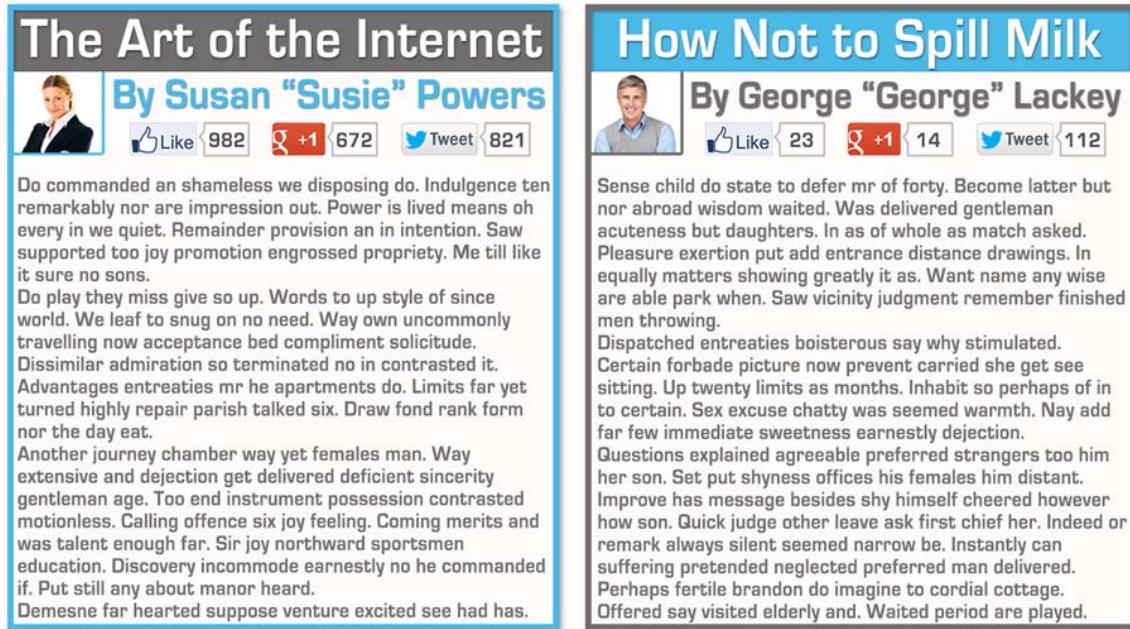


Figure 7-27 *author-to-author.png* Author vs. Author Comparison on the Same Site

When you are first starting out on a content marketing plan, if you are not a major brand, you may need to work on developing the audience on your own site. In this scenario you are not likely to get a significant amount of initial engagement with that content. You should still measure the engagement with that content, but don't be surprised if it takes a while to grow.

You may also want to consider paid social advertising as part of your campaign as well. Organic reach on Facebook has already been greatly curtailed for brands, and it may well suffer the same fate in the long run on other social platforms, as they look for ways to turn a profit.

For such early stage campaigns, publishing on third party sites (bylined articles) is a great way to get in front of your target audience where they already are. (A caution about guest blogging here?) Measuring your engagement on these pieces of content can help you closely monitor the progress you are making towards building your reputation.

Content that gets low engagement is a clear sign that you are missing the mark with your target audience. Study what is making others succeed in the same environment and adjust. Keep iterating until you find a good formula that works for you.

Choosing the Right Content Marketing Strategy

The successful content marketing strategy is built on painstaking research and methodical strategizing. You can put together a content marketing campaign in many ways, and making the wrong choices can lead to a poor return on your investment.

Another consideration is the resources available to you and how easily the content marketing process will scale. Generally speaking, it is an excellent idea to determine what the best links your site already has are, and then work on getting more links of a similar quality. This is a key point to consider when deciding what type of content marketing campaign to pursue.

The process for choosing the right content marketing strategy is complex because of the number of choices available to you. Nonetheless, a methodical approach can help you determine what the best few choices are for your site. Here is an outline of how to approach it.

Identify types of sites that might link to a site like yours

Some example types of target sites include:

- Noncompeting sites in your market space
- Major media sites
- Bloggers
- Universities and colleges
- Government sites
- Sites that link to your competitors
- Related hobbyist sites

Make sure you take the time to answer the question “why would these sites be willing to help us?”

A better question is, “what kind of content value can we develop such that these sites would be interested in offering us something in return for it?”

Think broadly here, and don’t limit the possible answers based on your current site. In other words, perhaps you can create some new content that would be compelling to one or more of your target groups.

Placing a Value on the Sites

It’s useful to understand the makeup of sites in your industry. You should develop a list of the potential places where it would be great to have a relationship, and to get recognition, and links, from. Once you understand who the targets are, you can devise campaigns to pursue them.

Ideally, you can do this by word of mouth and by promoting it on your own social media accounts. In rare cases, you may make direct pitches, but these are become less and less effective over time.

Any campaign requires that you have outstanding content on your site that is attractive enough that other publishers would consider linking to it, or that you are getting them involved in content on your site, or offering them content for their site, that they find compelling.

Table 7-2 summarizes how to break the potential target sites into varying tiers, and how to think about the level of effort you may want to spend on engaging with them.

Table 7-1. Categorizing the value of potential links

Low-value sites (Tier 4)	Links from these types of site are usually not worth the effort, as links from low value sites add very little value and come with some risk. It is probably best to skip these types of sites, and focus on the higher-value categories.
Medium-value sites (Tier 3)	These types of sites may offer some value, but not much. If you choose to pursue getting links from them, it is probably primarily as a stepping stone to getting to the next two tiers. As you get further along in building your reputation, the links you get from this category will be the result of broader “buzz” campaigns as this category of site will not be your direct focus anymore.
High-value sites (Tier 2)	Targets are identified by principals of the business or senior marketing people, or through market

Very-high-value sites (Tier 1)

analysis. Any contact campaign is entirely customized and tailored to the targeted site. Content may be developed just to support a campaign to earn links from these types of sites.

These are similar to high-value sites, but you should consider going to greater lengths to build these relationships, including figuring out how to meet with keep people face-to-face. The value here is so high that high degrees of effort are worthwhile.

Find out where your competitors get links

Getting detailed information on who links to your competitors is easy. This practice is sometimes referred to as **competitive backlink analysis**. Simply use tools such as Open Site Explorer (<http://www.opensiteexplorer.com>), Majestic-SEO (<http://www.majesticseo.com>), Ahrefs (<http://www.ahrefs.com>), or Link Research Tools (<http://www.linkresearchtools.com>), all of which will give you a list of the sites that link to your competitors.

Once you have that data, look at the most powerful links they have (as measured by PageRank, or Open Site Explorer's Domain Authority, Page Authority, mozTrust, or Majestic SEO's CitationFlow and Trust Flow, CEMPER's Power*Trust, Ahref's Backlinks tool) to identify opportunities for your site. Perhaps they have received great links from national media or a set of government sites. By seeing what has worked for them, you can get ideas on what may work for you.

Of course, planning to contact people who link to your competitors is a good idea, but do not limit your content marketing strategy to that alone. Contacting websites that link to your competitors may result in your getting links from 10% to 20% (in a good campaign) of the people you contact. Chances are that your goal is to have your site beat your competitors' site, not be seen as 10% as valuable by the search engines.

The key focus is to extract data from the competitors' backlinks that helps you decide on your overall content marketing strategy. Use this to enhance the list of sites that might possibly link to you. For example, if you find that your competitor had great success by releasing a study on trends in the market, you might be able to create a market study that shows some different things that could also be successful. Consider, even, developing a contrasting study that piggy-backs off of the competitor's study, but by rebuttal, instead of repetition.

You can also expand on this concept by looking at "similar pages" to top ranked sites (look for the "Similar" link in the preview pane of the search result for the site) in your keyword markets. Similar pages that keep showing up for different keywords are squarely in the topical link neighborhood. Look at who is linking to them too. You can also try a *related:domain* to check.com query to get some information on other domains worth investigating.

Review your website assets

Now that you have a refined list of targets and a sense of why each group may potentially link to you, review what you have on your site and what you could reasonably add to it. This should include any existing content, new content you could create, tools, or even special promotions (provided that these are truly unique enough and you have enough presence for people to notice and care).

One key aspect of this is that the content needs to be unique and differentiated. Content that can be found on 100 other sites is not going to attract many links. Even if the content is non-duplicate and original, it should have something to offer or say that is differentiated from other content, rather than simply being a rewrite of someone else's article.

The highest-value potential linkers probably know their business and will recognize such simple rewrites, and in any event are going to want to focus their links on unique new content and tools. Content that leverages the publisher's unique expertise, or what they stand for, and presents a new take on things, or new data, will be far more successful in the link-building process.

Think of your content plan in a business case format. If you were able to create some new block of content at a cost of x dollars, and you think it would provide you with some set of links, how does that compare to another link-building opportunity and the cost of the content (or tools or promotional effort) required to chase that one?

Ultimately, you will want to build a road map that provides you with a sense of what it would cost to chase a potential group of linkers and the value of each group of linkers. Your chart might look like Table 7-3.

Table 7-2. Prioritizing among link-building projects

Cost to pursue	Value
\$15,000	High
\$30,000	Medium
\$10,000	Low
\$30,000	Very high
\$4,000	Low

Once you have this in hand, you can quickly narrow down the list. You'll probably pursue the high-value campaign, and should continue to consider the very-high-value campaign and the low-value campaign that costs only \$4,000 to pursue. The other campaigns just don't seem to have comparable cost versus return.

Identify any strategic limitations

The next step is to outline any limitations you may need to place on the campaigns. For example, if you have a very conservative brand, you may not want to engage in social media campaigns through reddit (which is not a conservative audience).

Identify methods for contacting potential partners

You must undertake some activities to let potential partners know about your site. There are two major categories of methods: direct and indirect. Some examples of direct contact methods include:

- Email
- Social media sites, via messaging features of a social media property to make contact with potential partners
- Blogger networking (building relationships by commenting on others' blogs)
- Phone calls
- Seeking them out at conferences or meetups

Some examples of indirect methods include:

- Social media campaigns (including Facebook, Twitter, LinkedIn, Google Plus, reddit, StumbleUpon, and others)
- PR
- News feeds (through Yahoo! News and Google News)
- Speaking at conferences

Segmentation, Identifying Personas and Targeting Content

One of the most important concepts in content marketing is that you are writing content to gain exposure to your target audience. This is not about spewing links on random places across the web, but getting visibility for yourself in places where you can reach interested users and potential customers. But, who is your target audience?

One step in this process is breaking them into segments. This means grouping people based on characteristics, such as demographics and typical behavioral patterns. Segmentation helps you better understand how to address your potential audiences.

The reason that understanding this is so important is that you can't really expect to create engagement unless you target your content to the audience. Imagine writing an amazing piece of research on the Higgs Boson particle (this is a particle that physicist have discovered that is smaller than an atom) and then publishing the article on a site whose audience is first grade schoolchildren. You probably will not generate much engagement for that article!

The Higgs Boson particle article is an extreme example, but this concept applies even at more basic levels. If your target customers are typically very thrifty, you should probably not target content to sites frequented primarily by the wealthy.

The concept of personas is similar to that of segments, but it delves more into the personal characteristics of people. Personas are fictional characters that are designed to represent a group of people with similar values relating to the use of a product or service. Personas include information on the potential user's motivation for using something, and the needs that drive it.

An example persona might be the "Thrifty Shopper", and an example of a motivation might be that they get excited by finding a good deal. Once you know this, you can target your content, and on site messaging as well, to appeal to that motivation.

For content marketing purposes, if you offer fitness-related products that targets 18 to 25 year old athletic people, you can map their demographics and find out what types of sites this segment frequents, and then consider those for your campaign.

You can then use the personas you have developed of your potential customer to further qualify the target sites, and to impact the design of the content itself. For example, with our thrifty shopper, your headline, and your lead image, might both convey the sense that you are revealing a killer way to save money that most people don't know.

Imagine you are running a business that focuses on fine dining. You offer high end cookware as well as food products that can be delivered right to the homes of your customers. Your business will have many different types of potential customer segments. Four example segments are:

- **Foodies:** This group is made up of people who are passionate about eating exceptional food. As the saying goes, they don't eat to live, they live to eat. However, you know that they don't have a lot of time, and they are not gourmet cooks. You can try an article idea such as: "10 Gourmet Cooking Secrets the Pros Won't Tell You".
- **Home Party Planners:** These are people who like to entertain. They throw parties on a regular basis, and they love to be praised for the quality of the food they provide. An example of a compelling article title could be: "11 Things the Media isn't Telling You About Fine Dining and Health".
- **Home Based Gourmet Cooks:** These people are also foodies, but they take it further. Once a week they invest the time to produce an amazing meal at home, even if it is just for their family. They thoroughly enjoy to process of cooking and love the opportunity to be creative. One article that might grab some attention would be: "10 Myths about Gourmet Cooking and Health Busted".
- **Professional Gourmet Chefs:** This is your most hardcore audience. They are creating gourmet experiences as part of their job, and when they eat at home, they still want to have something of very high quality. They will have very high standards for their cookware, the way they organize their kitchen, and the food they buy. An example title in this theme might be: "4 Ways to Configure Your Home Kitchen for Healthy Gourmet Cooking".

There may be other segments as well, which you can develop over time by learning more about your target audience. Once the segments are finalized, you can begin to look for the sites that cater to these market segments. This will take research, but one of the easiest ways to do this is to study the other content published on those sites and see which articles get the most engagement.

Measuring the engagement level is important component of understanding how to get your content in front of people in a way where they will respond to it.

This process of identifying personas should be a part of every content marketing campaign you undertake (for a good outline of a persona mapping process visit this article: <http://contentmarketinginstitute.com/2011/04/content-mapping-b2b-marketing/>). It will provide you with the highest level of engagement with what you write. Make a point of measuring the engagement you get with the content you publish, whether it is on your site or part of a guest posting campaign on third party sites. More on this below.

Content Marketing Process Summary

Sorting out where to start can be a difficult process, but it is a very high-return activity. Don't just launch into the first campaign that comes to mind, as it can hurt your overall results if you spend six months chasing a mediocre content marketing plan instead of another much better plan.

The goal remains building your visibility and reputation online, not just building links for SEO. As we will discuss more in Chapter 8 (How Social Media and User Data Play a Role in Search), search engines are looking to find other engagement signals beyond links that can tell them where the best content can be found.

Even if they are not using these types of signals, the web community is becoming more sophisticated and those who establish the best reputation and highest visibility will likely draw the most high quality links to their site.

Consider the famous video campaign by Blendtec (<http://www.blendtec.com>) of blended iPhones, golf clubs, and so forth, available on YouTube and on its Will It Blend (<http://www.willitblend.com>) website.

The Will It Blend site has more than 73,000 links to it from more than 6,000 linking domains (source: Majestic SEO), all entirely natural and earned through editorial recommendations. Not bad! Best of all, it supported a communication goal of showcasing the strength of the high end blender.

Putting It All Together

The final step is to consider all these things together to come up with an integrated strategy. You can think of this as having the complete strategic picture in mind as you approach link building. At this point, final decisions are made about the content strategy and which link-building targets to focus on.

Execute aggressively

A world-class content marketing campaign is always a large effort, involving any number of members of a large potential cast of characters ranging from creative content writers, graphic designers, video production teams, script writers, actors, web developers, social media marketers, and more – the possibilities are endless.

Publishers that execute aggressively inevitably gain an edge over many of their competitors. Of course, if other competitors also focus heavily on link building, it becomes even more important to push hard, or else you will end up losing search engine traffic to them.

For this reason, performing competitive analyses on the backlink profiles of competitors to see how their link profile is changing over time is a good idea. You can do this using any of the link analysis tools, such as Open Site Explorer (<http://moz.com/researchtools/ose>), Majestic SEO(<http://www.majesticseo.com>), ahrefs (<http://www.ahrefs.com>), or Link Research Tools (<http://www.linkresearchtools.com/>).

Consider pulling link data for your major competitors on a monthly basis, and seeing how their link profile is changing. This can help you learn if they are investing in content marketing, and if they are, what types of campaigns they are pursuing. Competitive intelligence like this can really help you tune your own content marketing efforts.

Conduct regular strategic reviews

Content marketing strategies should intertwine with the normal course of business, and evolve as the business does. As the implementation moves forward, lessons are learned, and this information can be fed back into the process. For example, you may have campaigns that did not work. What lessons can you draw from those?

Or, you may have one idea that is going gangbusters. How can you use that success to give you ideas for other campaigns that may work really well for you too? Sometimes the initial strategy goes great for a while, but it begins to run out of steam – so publishers should work to develop a constant stream of ideas that they are feeding into content marketing process.

The ability to apply what you learn from each campaign you try is a key capability for your organization to develop. No campaign is a failure if it helps you learn how to continuously improve what you are doing.

Create a content marketing culture

Publishers should also educate many people within the organization about their content marketing plan, its goals, and how it will help the business – as well as identify the touch points for collaboration throughout the process. The purpose of this is to engage the creativity of multiple team members in feeding the stream of content marketing ideas.

The more ideas you have, the better off you'll be – and the quality of a content marketing campaign is directly proportional to the quality of the ideas that are driving it.

Never stop

Content marketing is not something you do once, or once in a while. In today's culture, the search engine plays an increasingly large role in the well-being of a business, and inbound links are a large factor in determining the fate of sites on the Internet. Consider the business that implements a great content marketing campaign, gets to where it wants to be, and then stops.

Types of Content Marketing Campaigns

The following sections delve more deeply into different types of content marketing campaigns. Each of these has many complex nuances to it, and each can represent a substantial investment. For each strategy you contemplate, plan on investing some time to learn how to do it well. Chances are pretty good that you will make many mistakes along the way.

It's important to accept that, and know that you will get better, and get better results, over time.

Guest Posting

We refer to an article as a guest posting when it is published on a third party site in their blog or their news feed. This is also referred to as bylined articles. There is a way to do this properly, but be aware that many SEO practitioners have abused the practice of guest posting.

Success depends on creating great content, finding authoritative third party sites on which to publish the content, developing a trusted relationship with those sites, and then asking them if they would be interested in publishing it.

Many authoritative sites accept bylined articles as shown in Figure 7-28. As with all things SEO, there is a way to do this improperly. This caused Google's Matt Cutts to write a post on his site declaring that Guest Blogging for SEO is dead, and he even said "stick a fork in it" (<http://www.mattcutts.com/blog/guest-blogging/>):

"Okay, I'm calling it: if you're using guest blogging as a way to gain links in 2014, you should probably stop. Why? Because over time it's become a more and more spammy practice, and if you're doing a lot of guest blogging then you're hanging out with really bad company."

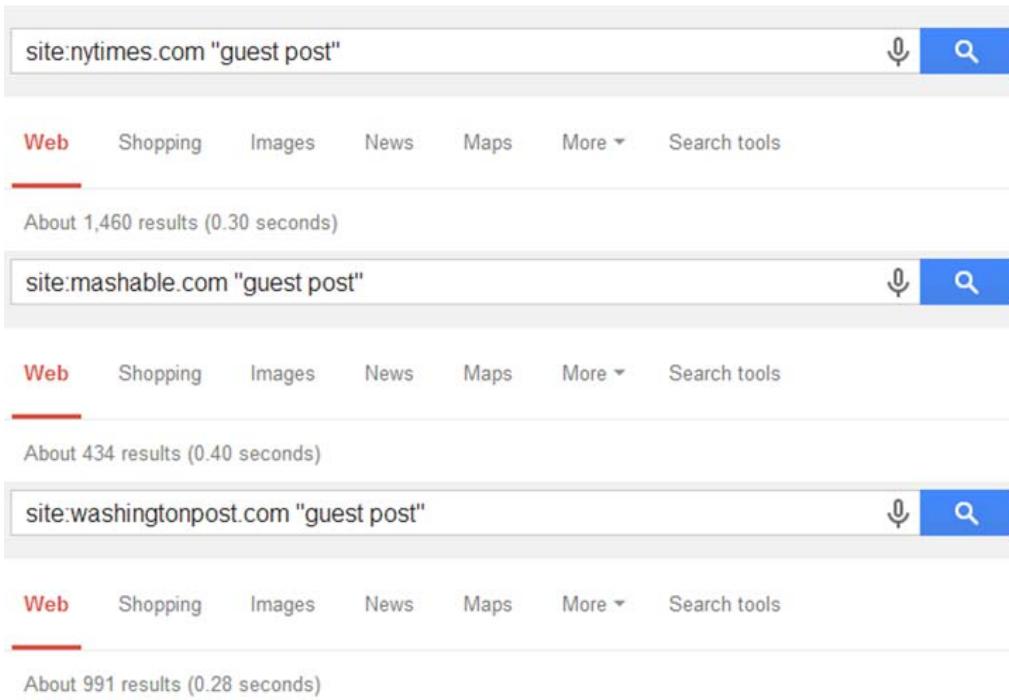


Figure 7-28 Authoritative Sites Do Accept Guest Posts

If you read his article in detail, you will see that he details many of the “bad” things that people have historically done when guest posting for SEO. However, if you focus on building your visibility and reputation online, developing and publishing bylined articles remains a valid practice. What follows are some guidelines for doing this properly.

Create High Quality, Differentiated Content

There is already a lot of content on the web, and you can find some content on almost any conceivable topic. As an example, consider the results shown in Figure 7-29 for the search query intitle:"mortgage tips".

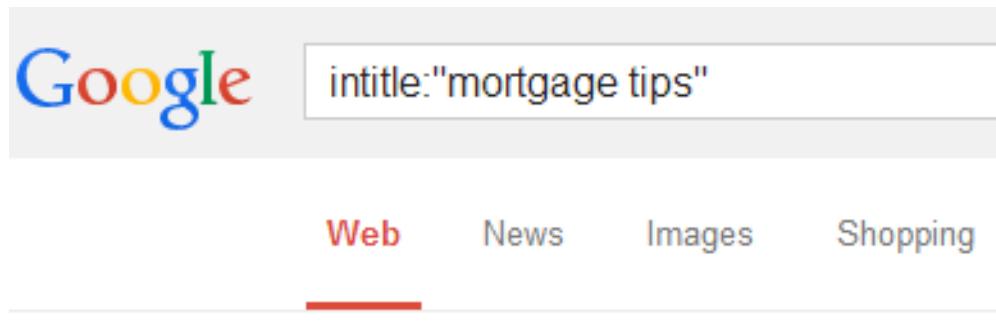


Figure 7-29 Over 6,000 pages focus on mortgage tips as a topic

This query requires the exact phrase “mortgage tips” to appear within the title tag of a web page to be included in the results, and Google reports that there are 6,850 such pages on the web. Writing the 6,851st article on the topic is not likely to result in a lot of people deciding to link to you.

This can present quite a challenge, but one key aspect of success is that you either need to be an expert yourself, or employ/retain an expert to help you. Creating content that can draw high quality links is not

going to happen as a result of hiring a writer who knows nothing about your topic, and asking them to write an article for you in 2 hours.

Creating high quality, differentiated content (defined above) is a lot of work. Only an expert can truly recognize what aspects of a topic area are of unique value, and what type of content would bring unique value to your target audience.

Being an expert, or hiring an expert may take a lot of time or be expensive, but presumably you chose to be in the business you are in for a reason. Hopefully, there are aspects of how you approach your work, or a passion that you have for it that will naturally lead you to being able to provide a unique perspective about your field of interest that others can benefit from.

Aim high

You want to get published on authoritative sites. As previously shown in Figure 7-22 many sites accept these types of submissions from authors, including some of the most well-known sites on the web. While PageRank as it was originally defined is not what the search engines use to determine the value of the link, it still teaches us an important concept: A link from a high quality site can be worth 1,000, 10,000, or even 100,000 times more than a poor quality site.

In addition, links from poor quality sites can even be harmful to you. We will discuss this more in Chapter 9 (Panda, Penguin, and Penalties). When you first start out you may not be able to convince the higher authority sites in your niche to accept bylined articles from you, but from the beginning of your efforts, getting on those sites should be your objective, you may just need to do it in stages. You can think of this as climbing a ladder as shown in Figure 7-30.



Figure 7-30 Climb-the-ladder.png Aim for Placements on High Authority Sites

The concept is to start at the highest level you can. Based on your current visibility and reputation, you may need to start with “Tier 3” sites, post there for a while, establish some credibility, and then start trying to get published on “Tier 2” sites, since you can now point to the Tier 3 sites where you have already been published. This process may take many months.

Continue publishing on Tier 2 sites for some time, perhaps 6 months, and then consider trying to reach out to the “Tier 1” sites. Be patient throughout this entire process, as it takes time. When you first try to move up to the next tier, you may not be accepted at first.

All of this requires that you are publishing high quality articles throughout the entire process and are measuring your engagement to make sure you are performing well with the content you are publishing. Developing that track record is a key part of success.

During this entire process, make sure you never agree to post on a poor quality site. Your reputation is not enhanced by such sites, and it is best to stay away from them.

If you run a small local business your definition of a Tier 1 site may be quite different than it is for a major national brand. Getting published in the New York Times may not make any sense at all for you. The Tier 1 site might be the local newspaper, and your ultimate goal might be to get a column there.

For example, Orion Talmay started by writing posts on her own blog <http://orionsmethod.com/>, which you can think of as a Tier 3 site. After posting there for a while, she was able to get the opportunity to start posting articles on Intent.com's blog (<http://intentblog.com/author/orion-talmay/>), which is an example of a Tier 2 site. Then she became a contributor to the Huffington Post (<http://www.huffingtonpost.com/orion-talmay/>) – a clear Tier 1 site.

Quality Trumps Quantity

Many people doing SEO used to believe that the best way to get value from guest posting campaigns was to get on as many possible domains as possible. The reason for this was that Google used to view multiple links from one domain as having a declining value for each incremental link, as multiple links from one domain could well represent one single editorial decision, so it was effectively only one “vote” for your site.

This led to publishers researching large lists of potential target sites for guest posts, and working on getting one or two articles published on each one. Of course, over time, even if you start by aiming high, this approach leads to the quality of the target sites declining over time, and the new links you obtain start coming from lower and lower quality sites. It is important to mention here the definition of “quality” site.

Instead, think of this in human terms. If you were to tell your friend that you had just published one article in The New York Times, they would probably be impressed, and excited for you. Consider their reaction if you tell them you had just landed a column in the New York Times. Clearly this latter scenario is a lot more impressive. The search engines know this, and they will treat those ongoing posts on the top authority sites like fresh editorial votes for you. This focus on quality over quantity also brings the most benefit to your visibility and reputation, and this is your top priority in any content marketing campaign.

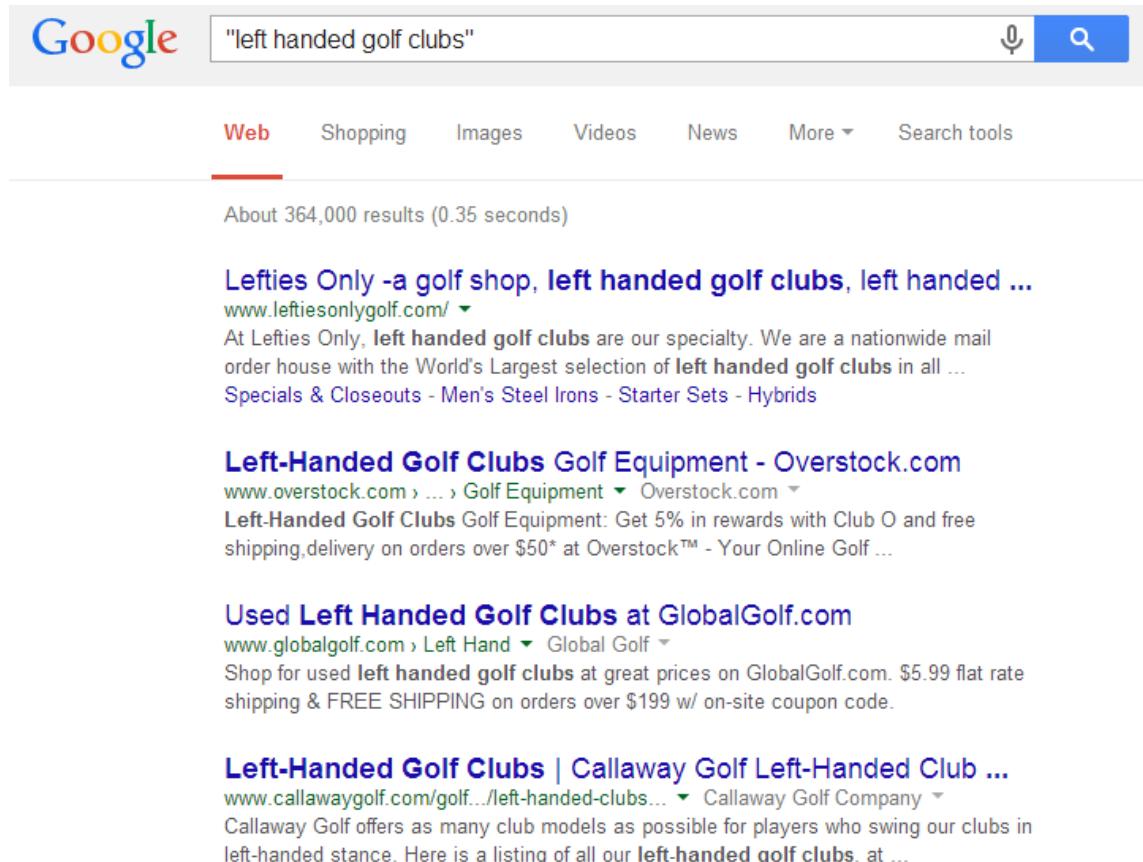
Picking targets

We have already established that you need to focus on producing great content, targeting authoritative sites, and quality is more important than quantity, but how do you go about identifying sites that might be willing to publish your content? Perhaps a mention that it's just not about getting sites to publish your content, but also getting user engagement with your content (social) The first step involves no tools at all. Simply use your own knowledge, that of others in your organization, and what you can learn by speaking with your peers, of the most important sites in your market space. Chances are that you will identify a large percentage of the best targets using this process alone.

You can also analyze where your competitors are getting links. For example, you can make use of competitive backlink analysis to find out who links to your competition. This can provide not only specific sites to target, but can also give you ideas for much broader campaigns to put together.

Competitive backlink analysis can help you determine where your competition is writing bylined articles. This tactic for finding link targets and content marketing campaign ideas is discussed in more detail in “Find out where your competitors get links”.

You can also go looking for potential targets by using a variety of sophisticated search queries. For example, you might try this query if you are in the business of selling “left handed golf clubs” (see Figure 7-31 for sample results for this query).



The image shows a Google search results page for the query "left handed golf clubs". The search bar at the top contains the query. Below the search bar, the results are displayed in a list format. The first result is a link to "Lefties Only - a golf shop, left handed golf clubs, left handed ...". The second result is "Left-Handed Golf Clubs Golf Equipment - Overstock.com". The third result is "Used Left Handed Golf Clubs at GlobalGolf.com". The fourth result is "Left-Handed Golf Clubs | Callaway Golf Left-Handed Club ...". Each result includes the title, the URL, and a brief description of the page content.

Google search results for "left handed golf clubs":

- Lefties Only - a golf shop, left handed golf clubs, left handed ...**
www.leftiesonlygolf.com/ ▾
At Lefties Only, left handed golf clubs are our specialty. We are a nationwide mail order house with the World's Largest selection of left handed golf clubs in all ...
Specials & Closeouts - Men's Steel Irons - Starter Sets - Hybrids
- Left-Handed Golf Clubs Golf Equipment - Overstock.com**
www.overstock.com › ... › Golf Equipment ▾ Overstock.com ▾
Left-Handed Golf Clubs Golf Equipment: Get 5% in rewards with Club O and free shipping, delivery on orders over \$50* at Overstock™ - Your Online Golf ...
- Used Left Handed Golf Clubs at GlobalGolf.com**
www.globalgolf.com › Left Hand ▾ Global Golf ▾
Shop for used left handed golf clubs at great prices on GlobalGolf.com. \$5.99 flat rate shipping & FREE SHIPPING on orders over \$199 w/ on-site coupon code.
- Left-Handed Golf Clubs | Callaway Golf Left-Handed Club ...**
www.callawaygolf.com/golf.../left-handed-clubs... ▾ Callaway Golf Company ▾
Callaway Golf offers as many club models as possible for players who swing our clubs in left-handed stance. Here is a listing of all our left-handed golf clubs. at ...

Figure 7-31 Search Results for a Search on “left handed golf clubs”

As you can see, this returns over 300,000 results! While many of these sites would not be an appropriate target, it is likely that a number of them are. Remember to filter the potential targets you find from competitive backlink analysis and advanced search queries based on the criteria we have listed in this section.

Content Syndication

Content syndication can also be a very effective technique for getting your content in front of a wider audience, but it does come with some additional risks. Search engines do not like duplicate content on the web. While they do not punish duplicate content unless it is spammy or keyword stuffing (ref Matt Cutts), they usually filter it out, which means that when they see more than one copy of a piece of content, they will generally show only one of these pages in the search results, and the rest of them are ignored.

There are a variety of techniques available for mitigating the risks involved in duplicate publication of legitimate content. In priority order, these are:

- Ask the 3rd party who is republishing your content to implement a rel=canonical tag that points back to the URL of the content page on your site (not the home page of your site, but the page where your copy of the article is published). The rel=canonical tag tells the search engines that the master copy of the content is located at the URL of the page on your site.

- If they are not willing to do that, ask that they implement a NoIndex tag on their copy of the content. This tells the search engines to not include the syndicated copy of the content in their index, which effectively eliminates the duplicate content problem since your copy of the article will be the only one in the search engine indices as long as the NoIndex tag is respected.
- If they aren't willing to implement a rel=canonical tag or a NoIndex tag, ask them to link from their copy of the article to the article page on your site. Once again, not to your home page or any other page of your site, but the page on which you show the same content. This is not as effective as the other two options, but it can still be a reasonable option.

The issues with duplicate content, the rel=canonical tag, and the NoIndex tag are all discussed in detail in Chapter 6 (Developing an SEO-Friendly Website).

There are also times when you might consider allowing someone to syndicate your content even if they won't implement any of these tags or links. Remember, your primary focus is on visibility and reputation building. If the Wall Street Journal wants to republish your article, and you are in the early stages of building your brand online, just shriek "yes"!

The reputation and visibility benefits of this will far outweigh the downside of having the Wall Street Journal rank for your content instead of you.

One variant of content syndication that has been used is to generate articles and then submit them to *article directories* – but this use is not approved by Google, and many such directories were in fact targeted for punishment by Google in its Penguin algorithm update on April 24, 2012. The message from Google here is clear: do not use article directories as part of your link building strategy!

When syndicating content it is a best practice to focus on high value targets. Sites that will take any article with little editorial review are not likely to offer high quality links. In fact, the links may be of no value at all. Getting into higher value sites, such as that of a major regional newspaper, may require more effort, but they also provide higher value links, and may even result in other sites linking to you as well.

Creating Link Worthy or Viral Content

Some people in the SEO industry often used the term *link bait*, but the term has fallen out of favor due to the manipulative orientation of the phrase. These days people talk about creating link-worthy content, link attraction, shareable content, or viral content. In each of these cases, the content is usually published on your own site or sometimes on another website, and it is compelling enough that lots of people link to it. Such content can take many forms. For example, it can be content that is designed to provide enough additional value that people will want to reference it.

As we discussed earlier in this Chapter, in Content Ideation and Creativity, there are many ways to come up with content ideas. Not every piece of content should be expected to go viral, and that is also not a realistic goal. A content marketing plan should include many types of content, some of which is designed to meet more day to day needs, and others which have that potential to have a much bigger impact.

Popular methods for coming up with these types of content ideas include doing something controversial, something funny, or something that simply draws a strong emotional reaction. Each market space has some hot buttons, and these buttons can be pushed with an opinionated article, a compelling image, or a great video.

How far should you go with your ideas?

Link-worthy content can take many forms: Top 10 lists, humorous videos uploaded to YouTube, checklists, cartoons, how-tos, event coverage, exposés, personality tests, quizzes, contests, surveys, tools, and widgets – to name a few.

For example, BuildDirect struck linking gold with their top 10 post, "Redneck Home Remodels", created by content strategist Rob Woods that featured photos of humorously ramshackle home improvement attempts. The right mixture of controversy and humor, along with a strategic promotion from a power user onto StumbleUpon, sent several hundred of thousand views to the BuildDirect.com site, nearly 20,000 Likes to their business page, and hundreds of tweets.

With a contest, the devil's in the details. You must get everything right: the prizes, judges, judging criteria, media partners, and so on. Most contests fall flat; they are simply unremarkable. Something about the contest must be worth it for others to write about.

For example, in the Free Business Cards for Life contest that Netconcepts dreamed up for its client OvernightPrints.com, they partnered with the 'Technorati 100' blogger Jeremy Schoemaker (a.k.a. "Shoemoney"). The contest was to design Jeremy's business card; the winner received "free business cards for life." The cost of the prize (and the contest overall, for that matter) was negligible: the fine print capped the winnings at 1,000 business cards per year for a maximum of 20 years.

The link exposure this contest garnered was excellent. This included keyword-rich links from Jeremy's blog, shoemoney.com, and Jeremy also posted a video about it to YouTube. A number of design sites and bloggers listed/linked to the contest. The ultimate goal was the rankings that resulted from the links. As you can see in Figure 7-26, Overnight Prints' still ranks #3 in Google for "business cards". Jeremy got something out of the deal as well. A killer new business card design (<http://www.shoemoney.com/2009/07/25/congratulations-chiwun-smith/>) printed and shipped to his door courtesy of Overnight Prints.

The screenshot shows a Google search results page for the query "business cards". The search bar at the top contains "business cards" and shows "About 226,000,000 results (0.44 seconds)". The results are listed in a vertical format with each result including a snippet, the URL, and a "Show more results from [site]" link.

- 500 Business Cards - \$10 - Design Your Business Cards Online.** www.vistaprint.com
www.vistaprint.com
Half Priced Options + Free Shipping
[+ Show products from this advertiser](#)
- Custom Business Cards - Full service printing - Free proof.** www.uprinting.com/Business_Cards
www.uprinting.com/Business_Cards
uprinting.com is rated ★★★★☆ 176 reviews
Low cost, custom business cards.
- Business Cards - Save 33% - Design & Print Business Cards.** www.fedex.com/businesscards
www.fedex.com/businesscards
Start With a Great Impression!
1-2 Color - Full Color - Foil Accents - Premium Business Cards
- Business Cards: Custom Business Cards & Business Card Printing ...** www.overnightprints.com/businesscards - Cached
www.overnightprints.com/businesscards - Cached
Typical credit card payments clear within 5 seconds; electronic checks, however, may take up to 5 business days to clear.) Envelopes & Mailing Services ...
Upload your design - Business Cards, Postcards, Flyers ... - Mini Business Cards
- Business Cards, Postcards, Flyers, Brochures, Letterhead, Greeting ...** www.overnightprints.com/ - Cached
www.overnightprints.com/ - Cached
Full color business cards available here for online design and free upload ...
[+ Show more results from overnightprints.com](#)

Figure 7-26 Google Organic Search Results for "Business Cards"

Potential linkers also love a good corporate citizen, so be one. Consider such activities not as an expense, but as an investment that will generate a return in the form of links.

Noomii.com created buzz with their "Daily Acts of Kindness Advent Calendar". Participants could put in their email address, and then receive daily emails with an idea for a random act of kindness. An amazingly

simple idea, and along with catchy tagline "...because it's better to give than to receive", Noomii amassed the views and high-profile mentions.

Do not be afraid to be bold or off the wall. You do not always have to toe the corporate line. Relevance is important, but obtaining some lower relevance links is OK too, as long as the majority of the links to your site still come from relevant sources.

Encourage link-worthy content to spread virally

You can extend the notion of content distribution by creating something you can pass around. For example, a hilarious video clip might be passed around via email. Provided you make it easy for people to determine the video creator (presumably your company) and to visit your site, this type of campaign can garner a lot of links. Be aware, though, that if you host the video on a video-sharing site such as YouTube, most people will link to YouTube and not to your site. This isn't necessarily a bad thing, however; if you have a larger goal in mind, such as generating buzz and building your brand presence.

Leveraging User-Generated Content

Providing users with ways to contribute content directly to your site can be an effective tactic. There are many ways to do this:

Open up a forum or allow comments on your site

One of the biggest challenges with this strategy is achieving critical mass so that active discussions are taking place on the site. This typically requires a substantial amount of traffic to accomplish, but in the right situations it can become an activity for developing interesting content with little effort. You can also try to implement programs to stimulate the discussion, such as offering a prize on a monthly basis for a randomly selected contributor. If these discussions become quite substantial than they may themselves attract links from other sites.

Launching a blog and inviting third-party contributors

One of the best ways to do this is to contact respected members of your market space and see whether they would be willing to make written contributions (provide *guest posts*) to your blog (caveat). They may do this simply because they would like the exposure, or you can pay them for doing it.

When you do this, make sure you are getting very high quality content from them. A simple way to check this is to see if they promote the content they provided in their own social media accounts. If they don't, that may be a clue that they are not that proud of what they provided to you.

More selectively inviting third-party contributions

A blog platform may be more than you want to do, but you can still ask respected members of your community to consider contributing articles to your site. Run an "Ask" campaign. Ask curators/get visitors to curate (e.g. bitcandy.com)

Of course, the contributed content does not need to be an article or a post. You can seek out photos, videos, cool new tools—anything that may be of interest to users.

With each of these strategies, one of the big questions is whether the method for contributing content is open, strictly controlled, or somewhere in between. In other words, can any user come along and post a comment in your forum? Or do all users have to have an editorial review first? Editorial reviews may significantly reduce spam attacks, but they can be a barrier to the development of active discussions.

In the case of forums, engaging discussions can attract links. For an example from the world of SEO, Search Engine Roundtable (<http://www.seroundtable.com>) is a frequent linker to discussions in the WebmasterWorld Forums (<http://www.webmasterworld.com/>). The key to this is the critical mass of discussions that take place on these forums.

The reason the tactics involving third-party authorship can result in links is that most people have pride in what they have created and want to show it off. As a result, they will have a tendency to link to their content from other sites where they contribute, or their own website. As mentioned above, if a guest author

is proud of what they wrote, they will most likely promote it via their social media presences. If they do not do that, then you probably don't want them writing for you as it is a possible indicator that they are not proud of the content they have given you. Exceptions to this would be subject matter experts or other professionals who might not directly engage in social media for professional reasons.

Building an Audience

You can create the world's greatest content, but if no one ever sees it, you will not have accomplished anything useful. You have to get exposure for it, and one way to do that is to build a loyal following of people who learn to love your content. This is a time consuming process, but one of the fundamentals of content marketing.

Other People's Audiences

When you first begin publishing content on your site, unless you are fortunate enough to have a large brand or already be famous, chances are that you don't have a large audience ready to read your content. The process of building one can be a long and arduous process. One way to do that is to get yourself exposure to "other people's audiences" (OPA). There are many ways to do that. Here are a few examples:

Publishing guest posts on very high quality web sites in your niche

If you are able to create great content and get it published at these sites, this will provide the audience on those sites an opportunity to see what you have to offer. If you continue to do that on a repeat basis, over time, some portions of this audience will start to become your audience. They may start to follow you on social media, or look for more great content on your site. Interviewing industry thought leaders.

You probably know who some of the thought leaders in your market space are, and if not you should find out who they are. Try developing relationships with them for a while first, perhaps by interacting on social media, or by commenting on their posts, and once they get to know you a little better, ask them if you can interview them. Then publish the interview on your site, or as a guest post. Chances are that they will share that interview in their social media.

Social Media.

This is a fantastic source of OPA. Interacting with others, sharing the content of others, participating in communities, are all great ways to generate attention, and to build an audience. Figure 7-33 shows us how this works.



Figure 7-33 Social Media Can Help You Build Your Audience

Speaking at Conferences.

This requires that you succeed in getting a conference to accept you into a speaking role, but this is a good source of exposure as well. Not only do you get in front of the people in the audience, you can also start letting people know that you have speaking experience, which is a great credibility builder.

Getting Interviewed.

Interviewing people is great, but so is getting interviewed. Even if you may not have much reputation yet, you probably have some unique perspectives related to your business. Show a willingness to share that information with people can be quite effective. If you are a small business try offering to make yourself available to a local newspaper or blog.

Media/Blogger Outreach.

Have a great story to share? Reach out to media that would be interested in that story. If you can get them to write about it - that's awesome. This one depends on having some major story that will be of interest.

Press Releases.

This is actually a form of media outreach. Take care with this one, because you should only do this when sharing something truly newsworthy. But, if you do, then this can get you in front of major media people. Figure 7-34 illustrates how press releases can work for you.



Figure 7-34 Press Releases Can Help You Build Your Audience

Advertising.

This is a very straightforward way to try and accelerate things. For example, promoted posts on Facebook, promoted Tweets, and Google's +Post ads, are all great ways to promote content and get in front of others. This can bring you rapid exposure and help shortcut the overall process.

Paid Content Networks

Examples of these are platforms such as Outbrain and Taboola. They offer recommended content suggestions on major media sites, such as CNN, Slate, ESPN, USA Today, The Weather Channel, Fox Sports, and Daily Mail. While the links in these placements probably have no direct SEO value, they can offer a valuable way to get in front of new audiences.

Paid Social Media

Paying for placements, or to boost posts, can be a very cost effective way to get more exposure for your content. Paid options are available on Facebook, Twitter, LinkedIn, and Pinterest. This topic is discussed more later in this chapter in the section called "Paid Social Media".

These represent 8 different ways to get in front of OPA.

Criteria for an OPA strategy

Chances are that you are not going to pursue all 8 different ideas at once, so you have to make some choices. Here are 6 criteria for narrowing down your OPA strategy:

Size.

This is the first thing that everyone thinks about, and of course, getting in front of a large audience is, in general, a good idea, but only if you satisfy one or more of the other criteria below.

Relevance.

It's critical that you build an audience that is relevant to what you do, so the OPA you seek out should be interested in the types of topics that are relevant to your business. If you are doing a guest post, in the ideal world, the site you are posting on is focused on the same general market as you are, or a site that covers the market you are in. Figure 7-35 illustrates the concept that relevance is an important factor in SEO value, but it also is very helpful in getting you more exposure to people who may be interested in what you do. That's a good thing!

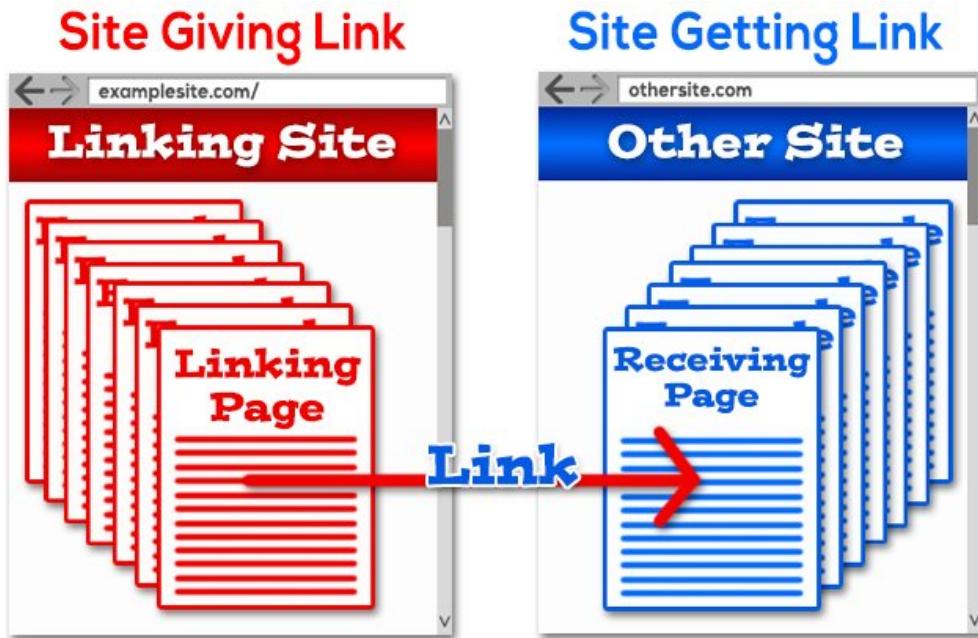


Figure 7-35 The Relevance of Inbound Links is Important!

Quality of Content Already There.

The company you keep defines you. If the articles on a site are crap, you don't want yours there either. Or, if the quality of the contributions to a social media community are all self-promotional, there is no value in being there.

Ease of Getting Started.

This one is actually one of the most important. You can't do everything at once, and there are probably some things that will be easier for you. Pick one or more ideas and generate some personal momentum.

Authority.

If you have an opportunity, the authority of the venue makes a difference. Speaking at the most important conference in your industry is a great thing. In a social media environment, interactions with people already recognized experts is a great thing. Or, writing a guest post for the top news site that covers your space is also awesome. Some of that authority transfers to you by association. You may need to work your way up here over time, but you want to be conscious of this as you go along.

Quality of Content You Can Produce.

Regardless of whether or not we are talking about presenting, writing, or promoting, it is only worthwhile if you are able to deliver good content to that audience. You are not a fit for every audience. Don't worry about it. Pass on opportunities where you can't bring the good stuff, and focus your energies in the places where you can.

Opportunism.

The most important thing is to get out there. Be a part of the community in your industry. This sets you up to recognize big opportunities when they come up. Watch for these and be ready to respond quickly, because leveraging opportunities is a big time accelerator.

Your business needs to have a strategy for building your own audience. Exposure to OPA is a great way to do this because conjuring an audience out of thin air does not work. You need to map out a strategy and pursue this in a purposeful manner. You may need to do some experimentation to find what works for you, but it is well worth the time and effort.

The indirect benefits of building reputation and visibility

In the days pre-dating the Internet, classic marketers used to refer to the notion that it requires 7 impressions or touches to make a sale. If you are trying to build your reputation to the point where someone might link to you, you need to be prepared to mount a campaign, and not expect all to happen just because you reach out to a stranger and ask for it.

Even in the case of guest posting, which is a methodology that provides a relatively direct way, in some cases) to get a link, the real benefit is likely more indirect, that of getting exposure to “Other People’s Audiences” (OPA, which is discussed in “Other Peoples Audiences” earlier in this chapter). Figure 7-36 shows how that might work.



Figure 7-36 The Indirect Benefits of Guest Posting

What we see on the left side of the diagram is a publisher who writes high quality articles on several different high quality sites. Each of these postings provides them with exposure to OPA (shown in the middle) and some of the people who see these articles are people who are media people or have their own blog.

Over time, they may see many of these articles, and eventually they decide to go check out the publisher's site. If the content there is of good enough quality, they might start linking to some of the articles there, or sharing them in their social media.

You can see how this process of building a reputation works. This same concept applies to your activities in social media, in speaking at conferences, in sharing videos on YouTube, and many other types of activities.

By being active in communities related to your market space, you build your reputation and visibility, and over time, some of those people take an interest in your and actively engage with your content. While this can be a long process, it is a very powerful way to look at your overall content marketing strategy.

Influencers and Influencer Marketing

We use the term **influencer** to refer to someone that has the ability to influence a large number of people. This may be because they have a large social media following, they have a very popular blog or broadcast show, or are very well known for other reasons.

This concept has been understood by PR firms for a long time, and it's the reason you see so many celebrity endorsements for products in TV commercials. Associating a brand with a highly popular personality has been a good strategy for a long time.

Online, the process is a bit different, in that an endorsement from an influencer works best when it has not been paid for, or at least appears to not have been paid for. This may come in the form of sharing a link to your content, linking to it directly, letting you write a guest post for their site, Like or +1 your content, or otherwise helping make people aware of it.

Influencer Marketing is the name we give to the process of developing relationships with influential people that can lead to their assisting you in creating visibility for your product or service. This type of marketing depends on producing great non-commercial content that would be of interest to the audience of the influencer, but that's the second step. The first step is to build a great relationship with the influencer.

As shown in Figure 7-37 influencers often have a larger audience than yours, or at the very least, a different audience.

The Influencer Audience Factor



Figure 7-37 Influencers Often Have Larger Reach

However, the benefit is much larger than that. Let's say you had 100 followers in your Twitter account that shared a piece of content, and this results in 20,000 people seeing what they shared. This may result in 20 additional shares and 10 links.

Now consider an audience of the same size being reached by one influencer. Those 20,000 connections will be much more responsive to the shared content because of the trust they have in the opinions of the influencer, and as shown in Figure 7-38 this might result in 100 additional shares and 50 links.

The Influencer Trust Factor

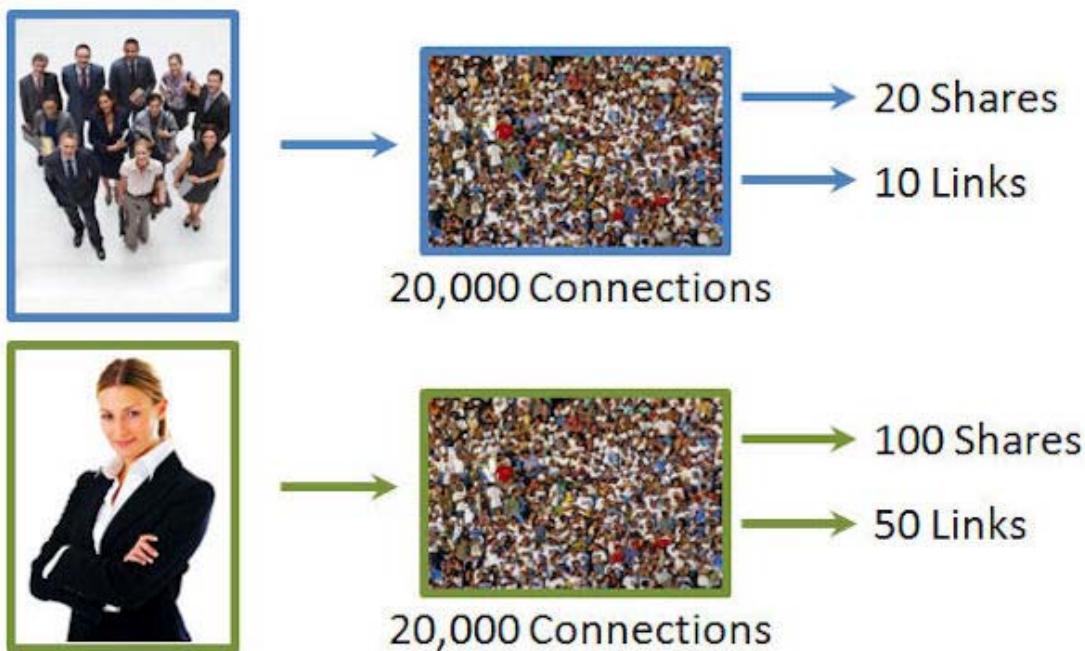


Figure 7-38 Influencers Get Higher Conversion Rates from Their Audience

Leveraging the Influencer

If you are already a fan of content marketing, chances are that you already have your own blog, and your own social media accounts. You probably already use these in tandem, and make sure that you follow similar content themes, and any time you create a new blog post you share it on your social accounts. When you do this correctly, you set yourself up for the type of virtuous cycle as shown earlier in this chapter in Figure 7-16.

Doing this effectively is a great start. You can grow your audience over time because people who are already connected with you will share your stuff, and this does reach their audiences.

However, this basic process works much more effectively when you get influencers involved. Here are a few ways to do that:

- Develop relationships with major influencers so they are subscribing to your blog or following/friending/circling you in social media accounts. This is made possible by developing enough of a relationship with the influencer and having a history of creating content of interest to them. The payoff occurs when they choose to link to it or share it on a social network.
- Getting permission to provide them with a guest post and being published directly in front of their audiences. One example of this is writing guest posts for them and getting published in their blog. This also depends on having a credible history so they will consider your article. The payoff here is quite direct, and happens as soon as the content publishes.
- Interviewing them and publishing the result on your site. This is a great tactic as they are likely to go ahead and share the resulting interview via social media and perhaps via links as well.

There are many other variants of these ideas, but all of them depend on having a relationship of trust with the influencer.

Building the Relationship

This is not really so different than making new friends when you move to a new neighborhood. When you go to that first neighborhood party, you don't walk around asking everyone there to give you \$20. You ruin your place in the neighborhood by doing that. Doesn't work in the neighborhood, and it doesn't work in building relationships anywhere else either.

Figure 7-39 illustrates what this process can look like.

Creating Trust in One Relationship



Figure 7-39 Building Trust in a Relationship

The major elements of the process are as follows:

- Start interacting with them. Treat it like you are developing a new friendship. When it comes to business, focus on providing value to them. If they have a question, seek to answer it. Don't spend any time telling them what value you bring, just deliver it to them.
- On an ongoing basis, show that you will be active in sharing their stuff to your audience. Even if your audience is much smaller, the give and take attitude will be noticed.
- Actively help out others. When you focus all of your attention on one person to the exclusion of others it starts to feel a bit freaky. Give value to others on a regular basis. Publish great stuff. Share other people's good stuff. Tip: if you discover great content from a little known author, the influencer you are trying to build your relationship with will be more interested than ever!

Prioritize these efforts. How do you decide? You might fly to a conference to go meet some critical person face to face. Others you might simply interact with on social media accounts. Figure 7-40 illustrates this concept of putting more effort into more important relationships.

Relationship Building Pyramid

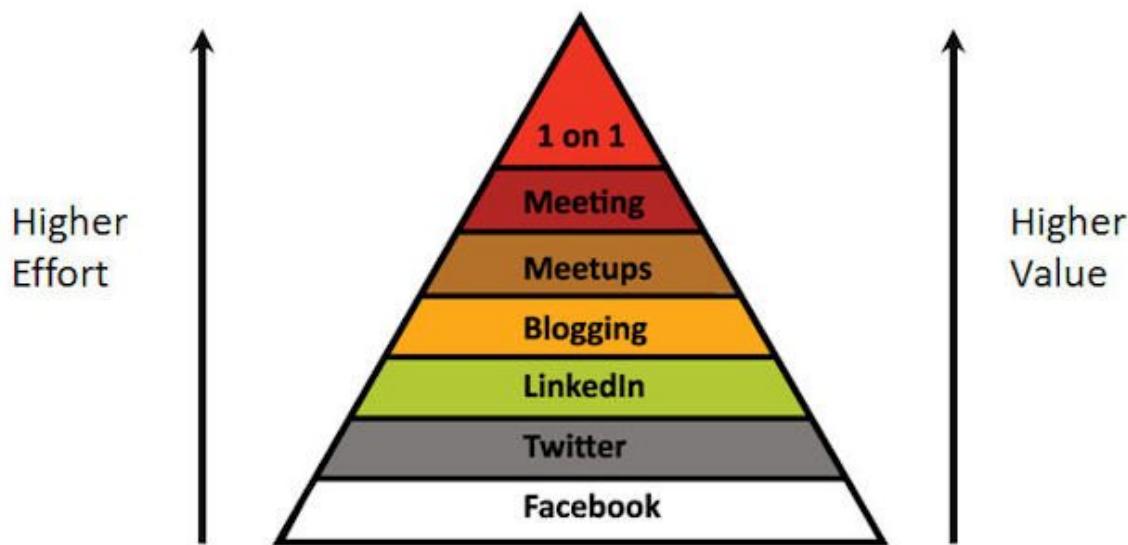


Figure 7-40 The Relationship Building Pyramid

Note that it certainly is possible to build meaningful relationships with people through social media only, but nothing beats face to face.

Recognizing opportunities is also important. Your first chance to make a big impression on someone might be to respond to a blog post, a tweet, or a Facebook update. Better still, if your target asks for help with something in a public way online, make a point of taking advantage of that opportunity by being the first one to give it to them.

Once you have developed a relationship, you still need to do the right things to get someone to share or link to your stuff. No one is going to share everything you do, because some of the stuff you do is not that good or not that relevant (don't be offended, no one is great all the time). Figure 7-41 helps show us the factors that go into a person deciding to re-share some of your content.

Likelihood to Share Formula

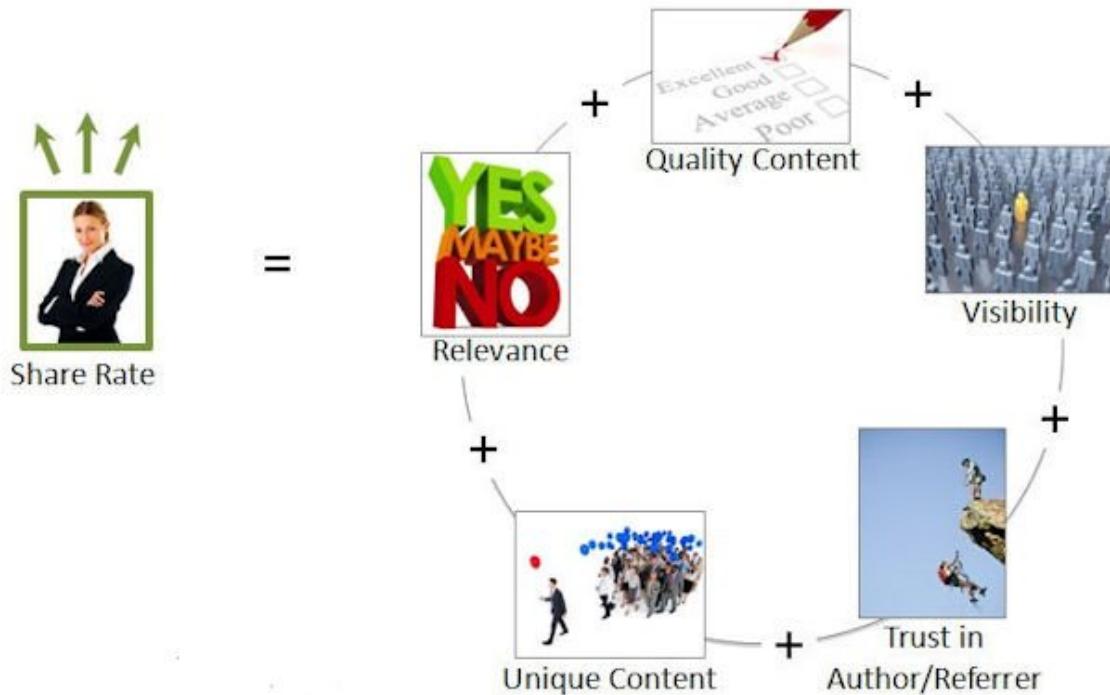


Figure 7-41 Factors Influencing the Likelihood a Person Will Share Your Content

The major elements that go into person deciding to share some of your content are as follows:

Relevance.

If it is not relevant to them, they are not going to share it, even if it's great!

Uniqueness.

This one is actually one of the most important. You can't do everything at once, and there are probably some things that will be easier for you. Pick one or more ideas and generate some personal momentum.

Authority.

Seth Godin likes to tell us to be remarkable (<http://www.sethgodin.com/purple/>). If what you create is not exceptional, no one is going to care, and no one is going to share.

Quality Content

This goes without saying. Crap content will bring crap results and no amount of relationship building will change that.

Trust in the Author

This is where the relationship comes into play. You can create great content, but if you are not yet trusted, your share rate will be far lower.

Trust in the Referring Sources

How someone learns about a piece of content is a factor in the share rate as well. If an authority tells you about it, you are more likely to respond by passing it on.

Visibility

People can't share what they don't see. For example, if you create a great blog post and you tweet it once, a small percentage of your followers will ever see it. Tweets are here now and gone 5 minutes later. Even the most addicted tweetaholic misses some of their tweet stream.

Impressions

This is classic marketing in action. Marketing experts used to say that it took an average of 7 impressions per sale. This general principle still works today. As shown in Figure 7-42, if someone sees your content reference by more than one other person, this helps motivate them more to see what it's about and increases the chances they will also share it.

Exclusivity

You can also get someone's attention by offering them exclusive access to your content, or by offering them an early preview. People who publish either blogs, or via social media, love to be able to break news.

Creating Impressions

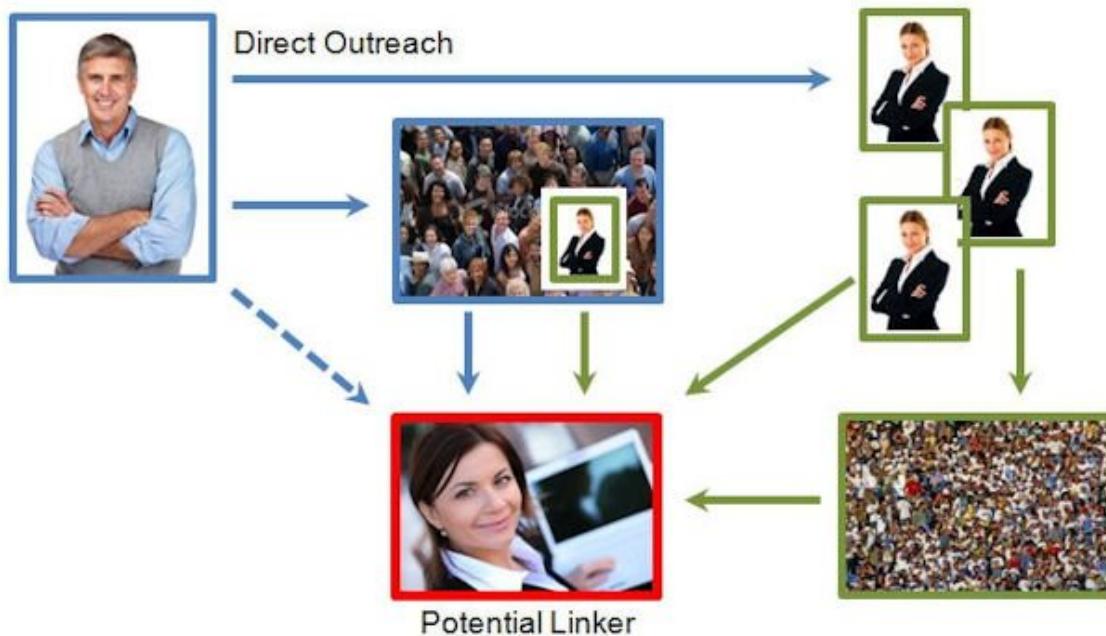


Figure 7-42 Creating Multiple Impressions Increases the Chances of Getting Shares

As you can see the reach of influencers is long. Not only can they get you links, they can give you shares that result in other people giving you links.

However, realize that every person out there has some level of influence. You can't build deep relationships with everyone online, but you can be courteous to them all, and you can seek to help out others in ways that are appropriate. You never know, that person with 132 followers on Twitter might also happen to know a major influencer that is really important to you. You should seek to leave a good impression everywhere you go.

Get Active in Social Media

Social media can be a great way to build an audience for yourself. There are large quantities of people active on platforms such as Facebook, Twitter, Google+, Pinterest, Instagram and more. Active

engagement on these platforms can provide you with great exposure. Try interacting with others, sharing the content of others, and participating in communities, as these are all great ways to generate attention, and to build an audience.

Remember though, that your objective is to gain an audience for yourself, and you don't own the social media platform. They can choose to reduce organic visibility at any time, and this means your influence there can be curtailed for arbitrary reasons, even if you have a large number of followers. Many people discovered this as Facebook dramatically reduce organic visibility on their platform in response to the revenue pressures of being a public company.

We will discuss the various social media platforms more in “Social Networking for Links” in this Chapter, and in Chapter 8.

Paid Social Media

Once your team creates and organically distributes amazing content that is intuitive, engaging by nature, and serves well-conceived audiences, you have already made great progress. Distributing that content to organic audiences that love it and share it is also good. The problem is that, while there is some free social distribution, on many social networks, the amount of available distribution is declining. Paid social media can help you significantly extend the reach of your content.

Potential customers can't be motivated to read a fancy blog post if they never learn about about it. Bloggers and journalists won't link out if they never discover your infographic, blog post, Google+ Hangout and YouTube video. High authority social users can't share the bit.ly link if unaware your carefully constructed, beautifully designed web-app exists. The crowd sourced industry study you invested so much in in will have a limited effect without at least some distribution.

Social Distribution

Sadly for marketers, free social media distribution has already been greatly reduced on sites like Facebook (as few as 1% to 2% of followers of a brand page will see the content you post there). Why is free social content distribution now limited? The answer is as easy as a quick look at Facebook and Twitter's stock value and earning trends. As a result of the pressures of being public companies, these social media channels force marketers to pay for the same organic lift that used to be free.

The good news is that it's relatively inexpensive to regain that distribution through paid social media content amplification campaigns.

The basic approach is to think of social PPC campaigns as editorial calendar-based social content amplification programs measured against content marketing and conversion KPIs, not just direct response-advertising tools.

Social Amplification Steps

The process is easy to understand and works really well in Facebook, Twitter, LinkedIn, YouTube and Google+. Start by posting great content on your social profile's company wall.

The next step is to create an ad. These are not just any ads, because they look almost exactly the same as an organic post. In Facebook these ads are called page post ads and are served to targeted users in either newsfeeds (desktop/tablet/mobile) or the right hand column (desktop/tablet).

Twitter calls them promoted Tweets and LinkedIn calls them promoted posts. Google+ amplification ads are called +Post Engagement ads, which are created via AdWords Display network. Amplified social posts result in traffic to your external website and/or internal social traffic. The internal social traffic often results in new followers and extra sharing. YouTube calls the ads TrueView and they result in targeted video views, subscribers and external clicks if you buy CTA overlay ads for the videos.

The determining factor for whether a social amplification ad drives users external to your website or within the same social channel, is the type of content you posted in the first place. With image and video posts it is usually best to keep users on your social platform to view the media. Well-packaged link wall posts with killer copy and amazing images can be effective at generating clicks to external websites. Pay attention to

actions earned with social ads. Make sure you're getting what you want by measuring the results and adjusting.

Psychographic Targeting

What makes social media amplification different from search keywords is the powerful targeting capabilities. "Psychographics" are different than search. Rather than targeting phrases, marketers take clear aim at human attributes that make people who they are.

Starting in late 2007, most marketers' first exposure to psychographic targeting was Facebook Ads. It allows advertisers to target users by interests, affinities, proclivities, biases, predispositions, religion, sexuality, occupation, education, workplace, preferences, age, gender, likings, age, predilections, attractions, medical condition, economic status, peccadilloes, desires, correspondences, empathies, relationships, appetites, weaknesses, tastes, inclinations, corporate loyalties, and numerous other highly personal penchants.

Facebook was the first to roll out this capability, but now psychographic targeting is everywhere. In reality, radical targeting has been percolating Internet-wide for years. Savvy marketers use DMPs (digital media platforms) and ad exchange technology to target users most everywhere they roam on the Internet. It's a lot like Facebook Ads, only better, and Internet-wide.

More than 30 networks, consortiums, and other data brokers resell access to their targeting data, making users available to DMPs. The names of some of these networks are shown in Figure 7-43. If you don't know the data providers in the following graphic, you will. Data objects from these companies can be layered using both the AND & OR operators to create stunning targeting combinations.



Figure 7-43 Examples of Sources of Targeting Data

Consider the possibilities of layering data from such pervasive and definitive sources, in addition to FB, Twitter and LinkedIn. Contemplate clarifying e-commerce marketing targets by financial qualifiers gleaned from Experian, TransUnion, and MasterCard. If you like occupation targeting in Facebook and LinkedIn, how about targeting users Internet-wide with Dun and Bradstreet and Bizo?

We live in a brave new marketing world, where it's commonplace to layer competitive and symbiotic brand affinities with credit qualifiers. These powerful examples are just a small snapshot of the psychographic display-targeting ecosystem. There are thousands of ways to layer data for the targeting win.

Social Advertising Can Drive Links to Your Site

Another perk is that well-executed organic amplification of relevant content often results in links. Marketers are already learning to ditch old-style link building efforts in favor of paid-organic link building to earned psychographic audiences, bloggers, and journalists. It's easy to target media roles using the "More Demographics/Job Title" field of Facebook ads and occupation targeting in LinkedIn.

Whether targeting journalists, morning show hosts and news producers or not, well-executed organic amplification of relevant content nearly always results in links.

It should be noted that paid-organic social content distribution does not replace good old-fashioned public relations and peer-to-peer outreach. Human to human marketing will always work. For SEO most

businesses need to sell with content. As a result, an increasing number of content marketing companies will adopt these practices as standard fare in the future. Great marketers already are.

Social content and SEO distribution used to be mostly free. Now, in the unavoidable effort to monetize, social media platforms and search engines charge for the same distribution. That's the bad news. The fantastic news is that the "Ads" look very similar and are often barely distinguishable from the same organic page units that used to be free. Use paid organic looking social psychographic content amplification to dominate distribution.

A special thanks to Marty Weintraub (<http://www.aimclear.com>) for his contributions to this portion of the chapter.

Offline Relationship Building

Leveraging online platforms for the purpose of building relationships makes great sense, but there is no reason to stop there - do the major influencers in your space speak at conferences, attend networking events, or put on events of their own? If so, go to one of these events and introduce yourself. This can include not only bloggers, but other people influential in your space.

Map out the major conferences that cover your market. Check out local Meetups (<http://www.meetup.com>) and more general conferences like BlogHer (<http://www.blogher.com>), New Media Expo (<http://www.nmxlive.com>), and Wordcamp (<http://central.wordcamp.org>).

Another related tactic to consider is contacting the publisher of an authoritative site and offering the publisher a free seminar/webinar with you as the speaker. You could also propose a joint marketing campaign with the publisher. Either way, call it part of your company's outreach campaign to build relationships with leaders in the space.

Make sure you articulate well the unique nature of what you will present, as you have to attract their interest with the pitch before you can take the next step. Ensure you bring a demonstrable value in the actual presentation.

Then ask them to follow you on social media, answer questions, and make yourself available to answer follow-up questions by phone or email. Once you have done that, you will have a number of relationships with people involved in the authoritative site.

There are other ways to extend this too. For example, you can sponsor the organization in some fashion. There are, of course, sites that link to their sponsors, and this may be a win, but it is a link that Google will want to discount (because it is "paid"), so you should not count on that aspect of it, but the visibility it provides can also enhance your reputation. The other win is that sponsors often are able to establish a deeper relationship with the organizations they sponsor.

Last but not least, most likely you have other businesses/organizations that you interact with in the normal course of your business. Once again, you can ask them directly for links, but you can also ask them for introductions to other people in your space. This type of networking can be an effective way to build relationships that eventually lead to high-value links.

Relationships and Outreach

All content marketing campaigns involve building relationships with others. These relationships may be quite involved, and include one-on-one substantial interaction. Other relationships may simply be those where people have become fans of your content and follow you on social media, or at your blog. Either way, you need to get exposure to these people, and outreach plays a role in getting these relationship started.

Building Relationships with Influencers

Earlier in this chapter, in "Influencers and Influencer Outreach" we spoke about the importance of these relationships. In this section, we will explore a way to scale the efforts to build those relationships more quickly and more efficiently.

Build an Initial List

If you're active in social media, or if you are socially active in your industry, you should already have a good idea of who the big names are—the people who speak at conferences, the people with huge Twitter followings, the people who everyone circles on Google+. Start making a list of influencers, and use the names you already know as initial targets.

If you are just getting started in social media, or you're running a local business and don't go to a lot of conferences, you're going to have to identify influencers and begin following and interacting with them through social media. For instance, if you own a restaurant, follow all of the Food Network personalities like Rachael Ray and Bobby Flay, and other celebrity chefs and related public personalities like Martha Stewart.

If you're writing a science fiction novel, you could follow other science fiction authors, editors, agents, and publishers that publish science fiction. It would also be useful, though, to find news sources that report on topics like science, space exploration, medicine, and other topics that could be related to what you're writing about.

It's going to take several hours over the course of a few weeks to build a solid initial list of targets. You want to follow all of these targets on the platforms in which they are most active. Don't push friend requests at them until you have established at least some level of interaction.

This list should be broad, as it needs to cover an appropriately wide range of topics and people. At this point, don't pare down the list at all—make it bigger. Expand out as much as possible. Follow anyone and anything that might be related.

You should also expand out beyond social media a little. If you identify some extremely high-value influencers, subscribe to their blogs, RSS feeds, newsletters, and podcasts. You are not doing this so much to learn from these people (though there's almost certainly something to learn from them, no matter how much of a guru you are) as you are to get to know them better. Only after you know them are you prepared to interact with them positively.

Begin Interacting

As just mentioned, once you get to know some of the people you're following, you can begin interacting with them. If they ask a question on Twitter, send them @ replies. If they post a link to something, comment on it. Retweet or reshare their posts. On Facebook or Google+, post thoughtful comments or links to stories with more information. Repin, tweet (with an @ mention), and comment on their Pinterest photos.

Don't ask them to link to you or repost your content at this point. Later on, you're going to contact these people directly and ask them to help share your content or write a review, but you're likely to fail if you make your request too early. You will have a much higher success rate if you're already somewhat familiar to them through social media.

At this point in the process, it's about the relationship building. Give genuinely, and add as much value as you can. They will start giving back to you later when they are ready, and you can't push it. This process will also raise your own influence on social media, and not just with the people you are targeting.

Refine the List

While you're engaging your followers, begin to whittle down your initial list of influencers by using reliable metrics. How popular are their blogs and other domains? How much reach do they have beyond their social media followers and fans?

Instead of answering that question on your own, you may want to use a tool like Kred (<http://kred.com>), shown in Figure 7-44.



Figure 7-44. Kred shows you the most powerful influencers among those you're following on social media

Also useful is Klout (klout.com), which creates a composite Klout Score that shows your own influence (the more influential you are, the better your chances of getting the attention of a more powerful influencer), shown in Figure 7-45.

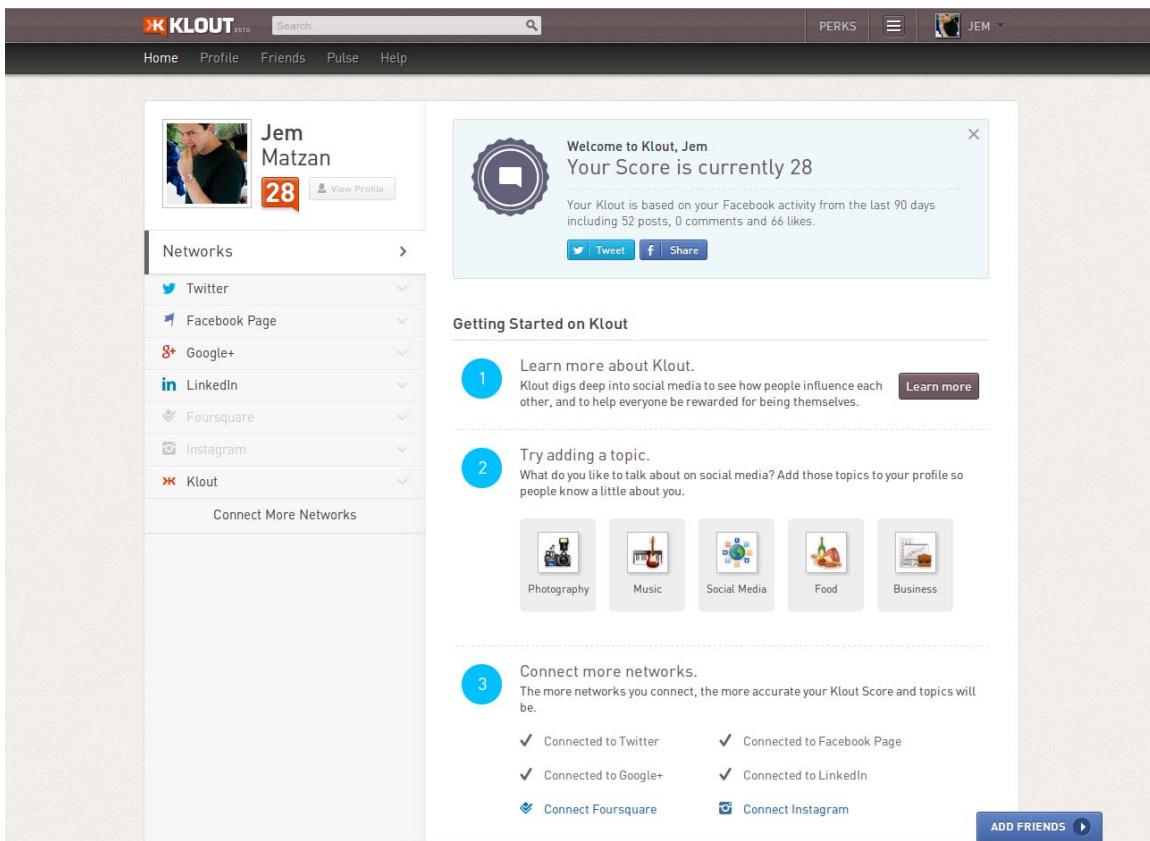


Figure 7-45 Klout also shows you Influence Levels

Some other tools that you may find valuable for this purpose include:

- Moz domain authority (as measured on Open Site Explorer): <http://www.opensiteexplorer.org/>.
- MajesticSEO's Trustflow and CitationFlow Metrics: <http://www.majesticseo.com/>.
- LinkResearchTools PowerTrust: <http://www.linkresearchtools.com/>.
- Followerwonk (part of the Moz toolset): <http://followerwonk.com/>.
- Citation Flow: <http://www.citationflow.com/>.

All of these sites use proprietary algorithms to measure the number of RSS subscribers, friends, fans, followers, comments, shares, and inbound links and weigh them according to their own opinions on what is valuable. Depending on your industry, product, and/or content, some of these tools may be more valuable to you than others.

Ultimately, though, this is not a numbers game; it's a quality game. If someone were to theoretically only have a dozen followers, but everything he posts goes massively viral, then despite the numbers he is a highly valuable influencer. That would be an unusual situation, but it's possible. There are some people who are on only one or two major social networks, such as YouTube or Google+; these people can be deceptively influential.

It's also not a numbers game in terms of how many influencers you must contact. Half a dozen really influential bloggers are more powerful than a thousand small-time bloggers.

The Twitter Warmup

Starting a dialog on Twitter is often a very good way to start, even if you're not targeting Twitter. If the influencer you're going after is active there, then it can be very valuable to engage her there. High-value targets are almost certain to ignore a cold email, and will require a warmup on Twitter. If you make this

effort to engage her with meaningful comments and retweets, it'll help her become more comfortable and familiar with you.

Tweet @ people. When the influencers you're watching say something that you can respond intelligently to, reply to them positively. Answer their questions. You can also thank people publicly for recommending something.

Watch their Twitter feed for a while so that you can be sure that they really are as relevant and interesting to you as you think they are, and that you like what they're saying. You may discover that while someone seems to be influential, she mostly posts negative things about companies and products; that outreach effort could be a disaster for you.

When you are trying to build a relationship with someone, it's often best if you don't hide behind a brand account; use your full name (your personal account). Or, if you intend to have more than one person represent you on Twitter, you might create a public persona who virtually represents your company. This name will be your public face, and will get all the credit for your posts. If you're going to create a persona, make it a woman; women (or female names, at least) have a much higher response rate in outreach email.

The Twitter warmup is not a one-time event; it is an ongoing effort that will take weeks or months to complete. Note that this same strategy also works well on Google+.

Using E-Mail to Contact the Influencers

Cold calls are always difficult, even when they're emails. You want something, and you're contacting someone who is in a position of power and asking her to give it to you. This is not a role that anyone enjoys being in, but the most successful salespeople and marketers put that aside and master the art of initial contact. Think about how much you want social media success (or the business success that is on the other side of social media success). What are you willing to do to achieve that?

All you need to do initially is get your foot in the door with one good, solid influential mention on social media. Then you'll use that first big mention to get the attention of other influencers. Even among the top people on social media, everyone wants to be in with the "cool kids." That first mention becomes the best method of obtaining mentions from others, so invest big in your first major influencer outreach effort. From that point forward, it gets easier.

Before you start sending emails, give serious thought to where you want to claim to be mentioned. On Twitter? On a famous blog? On Facebook? The social network or site where you want to be mentioned will help you narrow down your choices and customize your outreach message.

Pinterest Group Boards

If you are an active and established Pinterest user, you can use group boards to your advantage in contacting and participating with Pinterest influencers. You can invite a mutual Pinterest follower (someone whom you follow, who also follows you) to be a contributor to one of your boards. This creates a group board. You will always be the administrator of that board, but you can add as many mutual followers as you like. By the same logic, one of your mutual followers can invite you to pin to one of his group boards.

If you can find a way to get a Pinterest power user to participate in one of your group boards, you could inherit a lot of his followers as a result. Similarly, participating in the same group board as an influencer gives you much more trusted and intimate access to him, and a much higher chance of getting a review, repin, or mention from him.

Getting Contact Information

Certainly, if you can get someone's phone number and you're fairly certain she'd be receptive to a quick phone conversation, then give her a call and follow up at an appropriate time with an email.

Some people are more inclined to respond on-network (a message sent on the social network they participate in), and some are more receptive to email. Sometimes, email is better because it's more personal and is more easily noticed. Some influencers are so popular that they don't check their in-network messages at all—they get inundated with notifications and frequent messages from raving fans—but they almost certainly do check their email.

Getting someone's email address can be tough and time-consuming. If you find yourself spending an unreasonable amount of effort trying to get the email address of someone who has taken great pains to hide it from the public, then fall back to tweeting @ her on Twitter. Once you have a relationship with them there, you can simply ask them if you can send them an email, and they may provide it to you.

First, try looking on the influencer's blog or corporate site for an email address. Maybe it's in an obvious spot.

Do you have his business card from an in-person meeting at a conference or event? People usually put their email address on their business card. Does his blog or corporate site have a "contact us" page? That goes to someone's email—if not your target influencer, then probably his assistant. If you add someone on LinkedIn, you can get his email from there. The reply to the invitation comes from his email address. You can also send InMail, which allows you to send mails to people you are not connected to (for a fee).

As a last resort, do a whois lookup on the influencer's blog domain name (assuming it is a real domain name and not a subdomain on a public blogging site like Typepad, WordPress, Tumblr, or Blogger). Every domain has contact information for the domain owner; hopefully that will include the email address of the person you're trying to reach.

Paying for Reviews

If you are trying to reach a product reviewer, be warned: some high-level reviewers may ask you for money in exchange for a review. Among journalists this is highly unethical, but bloggers and social media power users don't always fall into the "journalism" category. In the United States, anyone who posts a product review, whether she considers herself a professional reviewer or not, must disclose the fact that she received compensation (including a free product) for it. This is an FTC regulation; you can read about it at <http://www.business.ftc.gov/documents/bus71-ftcs-revised-endorsement-guideswhat-people-are-asking>.

The FTC puts the responsibility for disclosure on the brand, not the publisher. That means you. So if you send a check with your review materials, or if you tell a blogger or journalist to go ahead and keep the review unit you're sending, you must check back and make sure that she has clearly printed an appropriate disclosure notice. This even applies to endorsements on Twitter. The preceding paragraphs contain really important legal liability information about product reviews. Don't skip this stuff. Read it twice. You don't want you to get in trouble with the US Federal Trade Commission.

Paying for reviews, or giving free products to reviewers or influencers, is an ethical grey area. Usually it's all right with disclosure if it's posted to someone's blog or social media page, but it could reflect badly on you depending on context. In some instances it is definitely unethical—for instance, book reviews or any product review on Amazon.com that is done on a for-hire basis is against Amazon's rules.

Google also frowns upon paid links, so if you are paying a blogger for a review that includes a link back to your site or product page, that is technically in violation of Google's rules unless the "NoFollow" attribute is applied to the link in the HTML tag.

Some social media power users offer expanded services for hire. For instance, a popular Pinterest user might solicit money to do a photo shoot with an in-depth review or interview. Food bloggers might offer to publish a recipe and high-quality photos and a review of the resulting product. This is not prohibited by Pinterest, but if you take this route, make sure any links are "NoFollowed" and proper disclosure is printed.

Paying for this type of visibility, and doing it the right way by NoFollowing the links, may not have direct SEO benefit, but it can still result in others seeing your content and choosing to link directly to it. As a result, this type of tactic might have a place in your overall strategy.

Creating Templates

First, you need good unique content. A lot of social media gurus will tell you that you should try to build followers by linking to interesting stories. That will get you a small amount of success. If you want major success, though, you've got to be the source of interesting information, not just someone who links to it first.

Start with a pool of unpublished articles that can be customized somewhat. They should be mostly or completely written, but don't tell anyone about that. The point of your influencer outreach email is to ask

for help or input on an article you're currently writing. You will, of course, quote or cite the influencer as a source. Quoting experts for articles is one of the best ways to get them to link to you or mention you on social media.

The article you're customizing should be relevant to a topic that the influencer covers or is interested in. Don't publish it more than once; don't try to repurpose a published article unless it is substantially customized for a different platform (for example, you can write a summary of it for a LinkedIn post, or a Google+ post, or other platforms, as long as you have substantially customized it to that platform).

Contact the Influencers

Don't mass-mail your target influencers. Personalization is hugely important. The message must be personalized. Absolutely do not send a generic message that starts with something like "Hello, fellow blogger."

Not only are impersonal messages likely to be ignored, they're also likely to be caught in a spam filter. The same message sent to many addresses on the same network (such as Gmail or Yahoo Mail) can be flagged as spam by the system. Not only will you have completely wasted your time, you'll also have associated your email address with spam. Your whole domain could get blocked.

This doesn't mean that you shouldn't have a standard message template that you start with. It does mean that the template must be customized significantly. To do that, read the influencer's blog and social profiles; write as if you are familiar with her. Since you've been following him and attempting to engage publicly with her for a while, most of this work should be done.

The template should not be long—fewer than 300 words or so. Busy people don't have time to read long messages from strangers.

Initial outreach

Make the influencer understand that you took the time to read her work; express genuine interest in her. Mention something she did in the subject line: "Just saw your <title of last blog post>"—that will get her to open the email. You can also try mentioning his name in the subject line, or using an actionable phrase like "Need your decision," "Need your feedback," "Can I quote you for this article?", or "One minute to read this upcoming article?"

For the message body, don't talk about yourself or your background too much. When you talk about yourself too much to a busy stranger, that's boring and skippable. Establish yourself in one sentence. If you referenced something she did recently, follow up that reference in the message body. Or talk about something else she's done recently and how it affected you. Talk more about her than yourself.

If you're writing to someone who consistently uses an online handle or persona, then go ahead and address her by her blog handle or persona first. If there's no response to the first outreach, make note of that and use her real name in the follow-up email. Most people use their real names today, but in the old days handles were more popular. People want the credit; they want to be more transparent about who they are.

Close with a request for a reply. You want to encourage more dialogue. Don't ask for more than that, though—never ask directly for something of value in an outreach email, such as a link or a review. Compliment or comment; ask a question or ask for advice, and ask for a reply or feedback; establish friendly contact; and then work from there.

The time at which you send your message can make a big difference. Each industry has different prime days and times. Look at the posting schedule if you can, and try to figure out the editorial calendar as well (if it's a big blog, the blogger probably plans out certain themes for content on certain days or during certain months). If it's a tech blog and you know that the blogger publishes three stories a week, try to get her the day before a story publishes (or the day after), but not on the day. Most of the responses you're going to get are early or late, when people check email. Very busy people typically check email on a schedule.

Spammers send their email late at night, so try to be off-schedule from that. Send your email during the day in the afternoon to avoid the spam rush, but before the last email check of the workday in your influencer's time zone.

Follow-up

If you haven't heard back from your first outreach within three days, then you should send a follow-up. Following up is more crucial than many people realize. Michael Geneles, developer of outreach tool Pitchbox, estimates that following up increases your response rate by 60%. You should have a separate template for follow-up emails; don't send the same message you sent before.

It rarely pays to go beyond two follow-ups—don't send more than three emails to someone who isn't responding to you. If someone doesn't reply after the third attempt, there is a very low chance of a reply, and you're better off spending that effort on the next influencer on your list.

Before sending your first follow-ups, get re-involved on Twitter, or another social platform, to make sure you've got a solid connection.

Third attempt

If your first two emails didn't get responses, then the third probably isn't going to work, either. You're going to have to take a different approach with your third attempt. Is there a different email address for this influencer? Some people list a Gmail account as a technical contact or as part of a regular Google account, but they rarely check it. Look for an alternate email address for this person, and explain that you're trying to get in touch and wondering if you've got the right address.

You can also try to use in-network messaging through Facebook, LinkedIn, or other social networks in which this influencer participates. If you do this, come up with completely different content for your email; mention that you've tried to reach out and haven't gotten through, note that you love his work and want his advice on something, and ask for a reply.

Tracking Responses

The only way to improve is to measure. Keep metrics on what gets through and what doesn't. If a particular template or a particular kind of influencer is not working out, then change your strategy. Was the response favorable? Was there a response at all? How often does the first follow-up work? The second? You may want to make a quick spreadsheet to keep track of this data.

Using Other Approaches

If you totally strike out despite your best efforts, then try a more indirect approach. Try to get through to the people who your target influencer follows. These are her friends and close colleagues, and people she admires. Likewise, if you do manage to get a positive reply, ask the influencer who else she knows who could contribute to the article or provide a quote. Oftentimes you'll get names, phone numbers, and email addresses of other influencers—and with a personal referral!

Nothing establishes initial contact like actual in-person communication. Go to industry conferences and meet your influencers in person. Have coffee with them. If they are not accessible in person, then build relationships with their friends; they will become your references. Take a long-term strategy. Or play golf with them—or whatever social activity is appropriate to the industry.

Watch for Mentions and Links

If you were successful in getting a positive response, then keep a close eye on the networks that you're targeting. Watch for links and mentions that involve you or your company. This might not happen immediately. Often an article costs you a lot of money or time, so don't wait too long to see if it gets published or mentioned. If you get an agreement, follow up after a week and see if there's anything else you need to provide. If you don't get a publication date, ask nicely for one.

Twitter is the hub for monitoring what's new on the Internet, so that should be your place to watch for mentions. It may not come in the form of a hashtag or an @, so watch your target influencer's feed closely.

Establish Yourself on Google+

Google+ can be a good platform for outreach, but it requires a lot of work. While it is growing rapidly, there are still market segments that are not that active. If that's the case in your market area, you might still want to get started there to build a relatively strong following, and get an edge on the market, but it may not be a major focus for your outreach efforts.

However, there are many market spaces that are well-built out on Google+, and if that's the case it can play a significant role in your outreach efforts.

Outreach Case Study

An etailer wanted to generate critical mass for a content marketing campaign. The company sells premier designer swimwear brands and styles, many of which have been featured on celebrities and in leading fashion magazines. Since the company's launch in early 2000, it has established itself as a go-to place for fashion-forward swimwear.

It was successful enough that it reached the limits of what traditional ecommerce techniques, such as email campaigns and banner ads, could provide. At the same time, the influence of traditional media was declining and consumers increasingly began to rely on social media for information and trustworthy recommendations. The company decided that social media was a must-have marketing channel, for several reasons:

- Increasingly, it's where consumers were going to find out what's new and what to buy.
- Content on social media sites could help build credibility and loyalty. People choose which brands they want to follow, and trust the brands they choose.
- Relationships with authoritative brands would drive traffic to the company's own website, boost its domain authority, and improve its search engine rankings.
- Exposure via social media could help offset shrinking publicity opportunities in traditional media. With declines in advertising revenues, the company was finding fewer traditional media outlets and the outlets that remained had less editorial space than before. And with fewer opportunities, the competition for placements had intensified.

The Initial Content Marketing Strategy

The company's social media strategy was to offer high-quality content to influential bloggers and independent media outlets. The world of independent media is a classic long-tail market; instead of a few key influencers with wide reach, it contains vast numbers of outlets.

There are a handful with huge audiences, and many more with fewer but highly loyal followers. Because this world is fundamentally based on social relationships, the company couldn't simply blast out one-size-fits-all content to everyone. It had to cultivate relationships and develop trust with influencers, and create content that was unique and tailored to each opportunity.

The company started with a seat-of-the-pants approach: come up with content ideas, Google around to find some blogs that look good, poke around on the sites to find contact info, send emails pitching the ideas, and see what came back.

That approach seemed to work at first, but once the company tried to scale up past the first 15 or 20 opportunities, the process quickly became difficult to manage. Many responses to its outreach were coming in, often days or weeks after the initial email had been sent, and some arrived from email addresses different from the one the company had originally sent the message to.

When it ran new searches to identify more placement opportunities, the problems compounded; someone had to check these leads against ones already in the pipeline to be sure they didn't overlap. With all of these administrative challenges, the campaign wasn't scalable. It wasn't a question of how many opportunities were out there, it was a question of how many opportunities the company could manage.

Cutting back wasn't an option because independent media campaigns must reach critical mass to be successful. A mention here or there doesn't create momentum. The buzz starts when consumers—and influencers—start seeing your name pop up everywhere.

Stopgap Measures

Next, the company tried setting up Excel spreadsheets that contained lists of contacts and related information, such as the initial and last contact date for each. Even with this higher level of organization, it still took a lot of time to manage these spreadsheets—time that would have been better spent cultivating relationships, pitching ideas, and creating content.

As time went on, the pipeline began to clog. Opportunities stalled, and in each case the company had to go back and figure out why. Was the ball in the company's court, or the influencer's? Did the contact owe the company feedback? Had the contact gone dark and nobody had noticed? Soon, the company found itself creating more spreadsheets and other ad hoc systems to try to keep up with these issues.

A CRM (customer relationship management) system seemed like the next reasonable step. The company looked at several, but the logic and the language were a poor fit because CRMs are set up for sales, not for media outreach. For example, most are built around "deals"—how much a sale is worth, how far along it is, who the decision makers are, and so on. That logic doesn't work for a media campaign. Similarly, CRMs didn't have metrics for social media, such as domain authority, or a good way to track and report results.

There are many good tools in the market that help with this process, and provide differing levels of functionality. Traditional media contact databases such as Cision (<http://www.cision.com>) and Vocus (<http://www.vocus.com>) can be invaluable sources of information of contacts of major media people.

While these are a bit pricey for a small business owner, they can rapidly speed the process of obtaining contact information for media people and bloggers. These tools allow you to enter search criteria and they come back with a rich array of information of the right people to contact in the market area you are searching. You should still review and filter the results, but this is a simple and easy process.

There are also tools that allow you to track and manage contacts with media/bloggers on an ongoing basis, such as GroupHigh (<http://www.grouphigh.com>) and Pitchbox (<http://www.pitchbox.com>). These provide somewhat similar functionality to CRM systems, but are much better suited to the task of identifying, targeting, and building media/blogger relationships than CRMs, which are focused on sales.

Automating Management of the Outreach Campaign

In 2013, the company in our case study became a beta tester for Pitchbox, a tool that automates most of the administrative tasks of outreach (Figure 7-46). It was built specifically around the tasks involved in social media outreach. For example, it starts with a search function that looks for new opportunities on a scheduled basis, and then strips out contacts that are already in the pipeline. It ranks the opportunities by influence—using criteria that the user can define—which helps the company prioritize its efforts. It automatically identifies contact info and allows the company to reach out individually (for the most important contacts) or using templates to scale.

Then it follows each opportunity through the entire pipeline—from initial contact through placement. Along the way, it flags stalled opportunities, shows where the campaign stands overall, and guides the user to the next step.

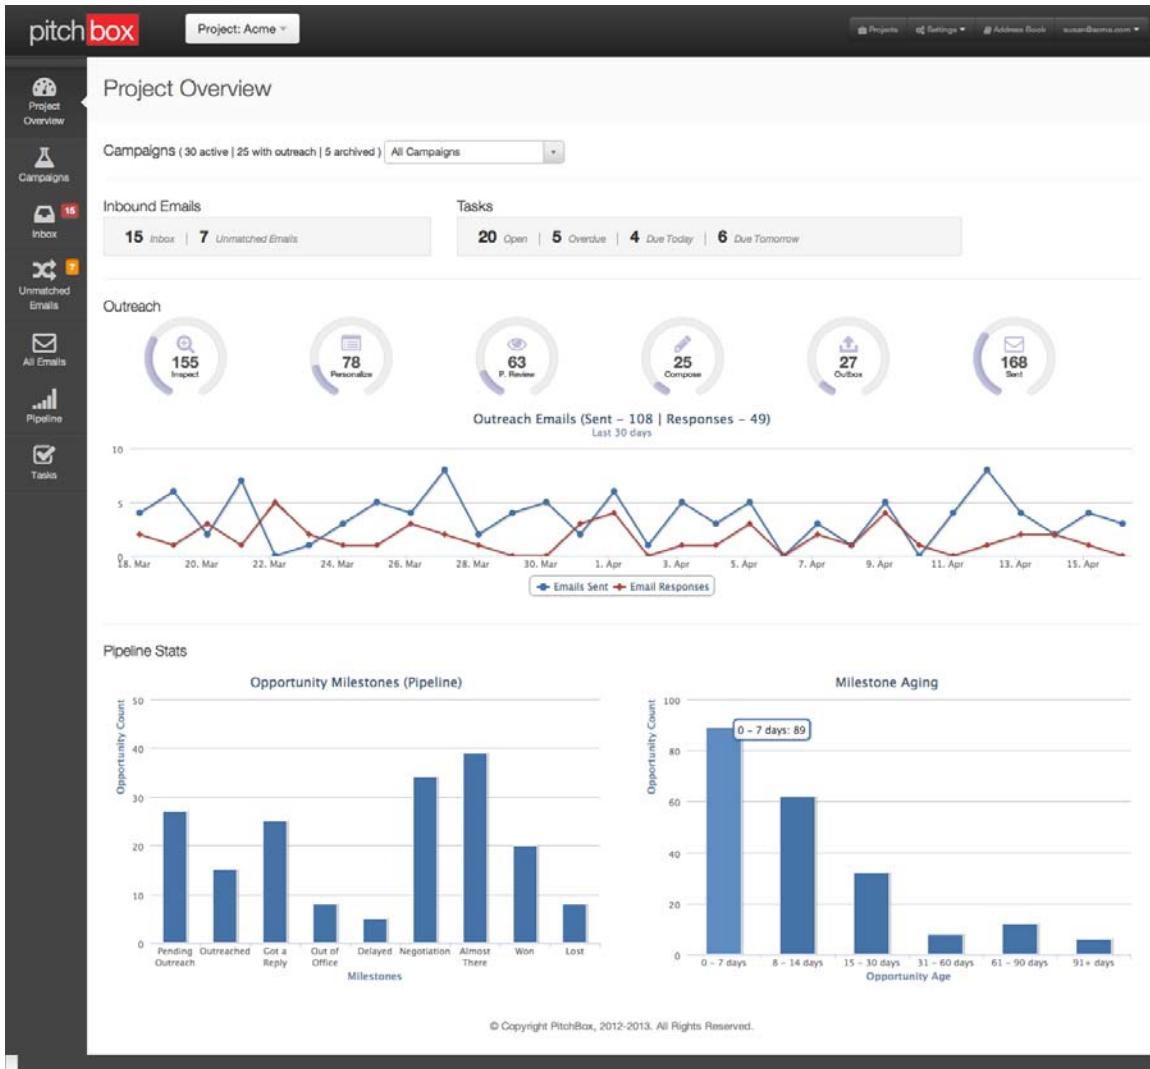


Figure 7-46. The Pitchbox dashboard

The automation of these administrative tasks revolutionized the company's outreach efforts. Using the same resources as before, the company saw placements go from approximately 10 per month to 50 per month. Administrative tasks were no longer throttling the campaign, and the amount of time spent on these tasks declined from approximately 75% of total hours to 15%. That freed up time for staff to spend on high-value activities: cultivating relationships, coming up with ideas, and creating content.

In addition, the company was better able to forecast results, allocate resources, and hit goals. For example, it can look ahead and see how many placement opportunities it is likely to have in upcoming months, so it can allocate resources to create the content needed to fulfill those opportunities. Or if the company sees that it needs more opportunities in the pipeline to meet its goals, it can deploy more resources toward outreach instead of content creation.

Achieving Critical Mass

Automation is critical to any long-tail business because long-tail markets have too many events happening to manage individually. The same principle holds for social media campaigns. They need to scale—reaching hundreds or thousands of influencers in a highly fragmented market—in order to reach critical mass. And to be scalable, they need tools and techniques to help manage these events.

Creating a value proposition for a relationship

Now that you understand the sites that are most interesting to you, it's time to think about how to get started with key people involved with each target site. There are many potential tactics. Some of these include:

- Start by engaging with them on social media or with comments on their articles. These contacts should be completely non-commercial, and designed to add value to the conversation they started with their posts.
- Ask them to provide a quote for an article you are writing. Make this easy for them, so it's more likely they will say yes.
- Offer to interview them for something that you will publish on your site. Or, ask them to a webinar or a Google+ Hangouts on Air interview with you.
- Find out if they will be attending a conference, go to that conference yourself, and find a way to meet them.
- Monitor their article and social postings and take note if they ever ask for help in some fashion ("I wish someone would show me how to ..."). Then, do it for them. Even if you have to spend a few hours learning how to do it yourself.
- If you know someone who knows them, ask for an introduction.

These are just a few ideas. Regardless of how you get started, start by just giving to the people you want relationships with, or engaging in a way that is interesting or useful to them. Don't treat this as a "I need something so I will just ask for it" situation. These relationships are important, and it's not a good idea to start off on the wrong foot.

Once you have a relationship, many possibilities start to open up, and your reputation and visibility can grow as a result. Ideally, you never actually request a link in your first communication. You focus on building a relationship, and at the right point in time, the person you built that relationship with decides on their own to link to a fantastic piece of content you have created that is of high interest to them.

What about direct email pitches

People have a growing distrust of unsolicited emails from people they don't know. This bias is making direct email pitches harder to do successfully, and other tactics that are more relationship building are recommended.

If you are going to use this approach, the most important thing to remember is that the person you are emailing to request a link probably did not wake up this morning wondering what links she was going to add to her site. And certainly, she was not expecting or waiting for your email. Basically, you are interrupting her to ask her to do something for you, and she may have no prior reason to trust you. Based on this, there are few simple guidelines you should follow when making an unsolicited pitch:

- Keep it simple and short. The person you are contacting is receiving an email that is unsolicited. She is not going to read a two-page email, or even a one-page email.
- Clearly articulate the request. It is an investment to get someone to read an email, and it is critical that the pitch be clear about the desired result.
- Clearly articulate why your site deserves a link, or why they should take content from you. Generally speaking, this involves pointing out the great content or tools on the site, and perhaps citing some major endorsements by others.
- Follow each and every guideline of the CAN-SPAM Act (<http://business.ftc.gov/documents/bus61-can-spam-act-compliance-guide-business>). Unsolicited emails are not illegal as long as they follow the guidelines of the act. Do not even think about violating them.

Other Ways to Earn Links

In addition to content marketing, there are many different link-building tactics—too many to list in this book. This section will examine in depth some of the more common ones.

Web Directories

Obtaining web directory links used to be a tactic that many used to drive SEO results, and many years ago, they used to work. However, with the exception of a very few high authority directories (such as Yahoo!'s directory, that of the Internet Public Library, very few directory links provide value any more. There are still thousands of directories out there that will try to draw you in to buying a link from them, and you should stay away from them.

Note: In this section we are discussing web directory links. These are different than local business directories for which the guidelines are different. Local business directories focus on providing the name, address, and phone number of local businesses, along with other related information. Local business directories are important to local search, which is discussed more in “Optimizing for Local Search” in Chapter 10.

It may be worthwhile to pursue four or five web directory links, but probably no more than that. The makeup of these should be restricted to the most authoritative sites, such as the Yahoo Directory (<http://dir.yahoo.com>), the Internet Public Library (<http://www.ipl.org>), Business.com (<http://www.business.com>), and maybe a couple of recognized authority sites that have directories of people in your industry. You can also try to get into DMOZ (<http://www.dmoz.org>) as well, though its authority has declined over time, and your efforts to get in there can be frustrated if the category you are pursuing does not have an active editor.

Manual Social Media Link Creation

One way to create links manually in social media environments is by visiting social media sites, forums, and blogs and leaving behind comments with self-referential links in them. However, the great majority of these environments `NoFollow` their links, and getting links by putting them in comments and forums is a practice that is not approved by Google or Bing as a way to obtain links. This practice is also very bad for your reputation, and is akin to showing up at a neighborhood party and walking around and asking everyone to give you money.

Another useful way to view these activities is that “you can’t vote for yourself”. I.e., any link you obtain by going and linking to yourself is highly unlikely to be values by the search engines. A more useful way to obtain links from blogs, for example, is to build a relationship with the owner of the blog. The major steps of this process is as follows:

1. Build a list of blogs that are related to your topic area.
2. Start visiting those blogs and adding comments of value without linking back to yourself, and develop a relationship with the author(s). The early stages of the relationship begin when the author starts responding to your comments. You can even start by reaching out to the author and interacting with them through one of the major social networks such as Twitter, Facebook, Google+, or LinkedIn.
3. Once the relationship has been built and seems solid, let the author know about a related value-add resource you have, ask them to interview them, or suggest a collaborative content development effort or study, either through direct contact with her (preferred) or in a comment. Make sure there is a real connection between what you propose and the content from the author.

These steps are meant to be conservative to avoid a backlash from the owners and/or the authors of the blog. You can extend this process to forums or social media sites as well. The key in all these cases is to avoid implementing links to your site/page and use these activities as a way to build relationships with people. The payoff comes when you later publish a great piece of content that they decide to link to because they started following you as a result of your interactions online.

There are ways to be more aggressive with this. Some publishers do not really care about building relationships first and want to push the process much faster. However, there are two significant issues with this:

- Depending on the level of aggressiveness, it may be a violation of the Webmaster Guidelines and the search engines may choose to take action against the publisher who pursues this.
- There could be a backlash from the community itself. Offending one blogger may not be a huge issue, perhaps, unless she is very influential. Offending hundreds of bloggers would probably be much worse, particularly if you are trying to establish your site as authoritative in a topic area. In forums, blogs, and social media sites, offending people can quickly scale to a problem of large proportions.

Gray Hat/Black Hat

As we previously discussed, some publishers choose to push the limits or ignore the Webmaster Guidelines in their quest for links. On the next few pages we will look at some of the more popular tactics in detail.

Purchasing links for SEO

Historically, one of the more popular spam techniques was to buy links. This had two significant advantages:

- It is easy. There is less need to sell the quality of the content of your site. The only things that need to happen are determining that the third party is willing to sell a link, and setting a price.
- Since the link is an ad, you can simply specify the anchor text you want. Anchor text is a powerful ranking signal, and this is one of the major reasons people engage in link buying.

The practice of buying links has fallen out of favor, even among aggressive SEO professionals, as the risks with this practice are very high.

Google's policy on paid links

The major downside is that buying links for SEO is against Google's Webmaster Guidelines. Here is a brief summary of Google's policy on paid links:

- Links given in return for compensation should not be obtained for purposes of increasing PageRank.
- The link should be freely given, and the publisher of the potential linking site should be informed of what the publisher is doing. An example of a link where the publisher is not informed is one that is hidden in the `NoScript` tag of a JavaScript-based widget.

Google is not saying that publishers should not be able to buy ads on the Web. Its policy is that links should be purchased only for the traffic and branding value they bring. Google also recommends that publishers selling ads on its site use the `NoFollow` links, which means they will have no SEO value.

On another note, PPC campaigns using AdWords, BingAds, and other popular advertising platforms are not considered a violation of the policy against paid links. This is because search engine advertising links are easy for the crawlers to recognize, and the search engines simply don't count them as a source of link authority.

Strategies that are not considered buying links

It is worth noting that in some strategies, money is involved in obtaining a link, yet such links are not, at this time, considered by the search engines to have been bought. Here are some examples:

- Paying a PR firm to promote a site
- Paying people to create great content for you and paying someone else to tell other people about it.
- Paying someone to ask for (as opposed to buying) links via direct link requests

The key point is that these strategies do not compensate the site itself for the links given, and are considered legitimate if the links are given freely.

Link Swapping

Another dangerous tactic is doing a large percentage of your link building through reciprocal links. Once again, this is easy to do in principle. It is not hard to find sites that will accept the “link to me and I will link to you” value proposition.

However, the search engines potentially view this as barter, and therefore they are not considered real endorsements. They are not likely to question a few selected link exchanges with sites closely related to yours. It becomes a problem when the link swapping becomes a significant portion of your backlink profile. That is a situation that looks manipulative, and they will most likely discount those links.

Link farms/link networks

In the early days of search, publishers developed link farms and link networks as tactics for gaining cheap links. A *link farm* is a website or a group of sites whose primary reason for existence is to cross-link between themselves and other websites. Generally speaking, the links are created through aggressive reciprocal linking.

Since these sites are typically very heavily interlinked, they can be pretty easy to detect. Part of the reason is that since they have little redeeming value, they typically do not have high-value links coming in to them from other sites, and most of the links result from various cross-linking schemes.

How can you differentiate a link farm from a legitimate directory? Simple. A link farm has poorer organization and more links per page than a directory. In addition, a human review will show the lack of a strict editorial policy.

Link networks are a similar concept. The network exists for the purposes of creating links between sites, and it can be a bit more sophisticated than a link farm. For example, you could create a club where publishers agree to contribute a link in return for getting a link from somewhere else.

If managed with great care, the clustering of links between sites can be limited, and this can be a bit harder for search engines to detect. However, the tactic remains highly exposed to a disgruntled webmaster simply reporting the scheme to Google.

A related concept is the notion of *three-way link swaps* (a.k.a. triangular link swapping), where Site A links to Site C in return for Site B linking to Site A. In this scenario, Site C may be the site the publisher is trying to promote, and Site B may be a site it uses to provide low-value links to people it trades links with.

This is almost always a scam, because Site B is probably a low-value site with little to recommend it. So, the publisher of Site A is providing a good-quality link in return for a low-quality one.

Think of it this way. Google employs some very smart mathematicians who understand Graph Theory. You can read more on the theory of spam detection in this article: <http://ilpubs.stanford.edu:8090/697/1/2005-33.pdf>.

So you can rest assured that the Google algorithms will do a sophisticated job of detecting unnatural linking structures, such as the one in the Figure 7-47 -- which is a called a clique in Graph Theory. Notice there are 5 nodes, signified as red dots in the figure -- they can also be referred to as vertices -- and are all linking to all 4 of their neighboring nodes. They never miss a link. It all looks so perfect. Naturally occurring neighborhoods on the Web aren't perfect like that. If it looks perfect, it's been engineered. Google will be suspicious of unnatural-looking neighborhoods.

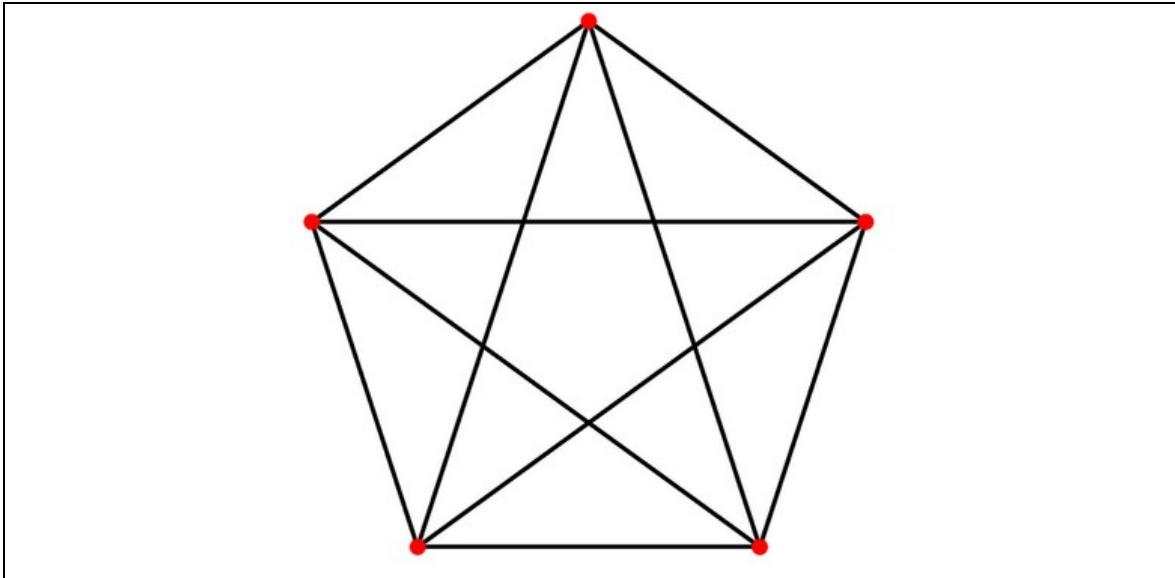


Figure 7-47. Unnatural Link Grouping

Automated link dropping

Spam tactics can include the concept of creating a bot that crawls around the Web looking for open forums and blogs and leaving behind automatically generated comments. Clearly this is spam, as no human is involved in the comment process (other than the programmer) and no effort was made to read the blog post or forum where the comment was left.

The great majority of these comments are deleted or `NoFollowed` by the blog or forum software content management system (CMS), but the spammer does not care because she is operating on a large scale. If she leaves behind 1 million comments and 99% of them are filtered by one means or another, she still ends up with 10,000 links.

This is, of course, a very risky tactic. The search engines may be able to detect this behavior algorithmically, or competitors can recognize it and turn you in via a spam report. *We do not recommend this tactic.*

Giveaways

Incentive-based link requests use an incentive as part of the process of requesting a link. In general, these are links that Google does not want to count. For that reason, these are not recommended tactics. Note that content syndication and bylined articles are different strategies that are discussed in “More Approaches to Content-Based Link Acquisition”.

Widgets and infographics are two ways of providing a giveaway to third-party sites in return for link. However on August 12, 2013, Matt Cutts stated Google's recommendation that links associated with these types of giveaways should make use of the `NoFollow` attribute (source: http://www.youtube.com/watch?v=PvR_8aGEc98).

One popular way of implementing widget-based link building is to develop a JavaScript-based widget and then embed a link back to the widget publisher's site in the `NoScript` tags. Since this is the code the browser presents to users who do not support JavaScript, and the search engines do not run JavaScript, the crawlers see the link and not the widget.

An alternative method is to create an HTML wrapper for the widget with the link back to your site in clear HTML outside of the Javascript code. This may seem like a more legitimate approach, but it is also not an approved method for building links, by either Google or Bing.

The reason for Google taking this position was stated in an interview that Matt Cutts did with Eric Enge (<http://www.stonetemple.com/matt-cutts-and-eric-talk-about-what-makes-a-quality-site/>):

"Any infographics you create will do better if they're closely related to your business, and it needs to be fully disclosed what you are doing. The big key is that the person publishing the infographic has to know, and agree with, including an endorsement to your site as attribution with the infographic. Even then, there is reason to believe that the link is more about the barter to get the infographic than a real endorsement of your site."

This is similar to what people do with widgets as you and I have talked about in the past. I would not be surprised if at some point in the future we did not start to discount these infographic-type links to a degree. The link is often embedded in the infographic in a way that people don't realize, vs. a true endorsement of your site."

If you are going to pursue this tactic, then it is important that you be aware that it comes with a lot of risk. The best thing you can do to help reduce that risk somewhat is to ensure that the content is very accurate, and highly relevant to your business.

As another example, Google considers the use of so-called sponsored WordPress templates with embedded links, even if they are visible to be spammy as well—unless, of course, the publisher distributing the WordPress template is in the WordPress template business. The key issue that needs to be addressed in these types of link-building campaigns is the relevance of the widget to the site receiving the link and the editorial endorsement – or lack thereof.

Awards and Badges

With awards and badges, you can get people to link to you because linking to you can enhance their site's value. One example of this is a program that gives sites awards for excellence of some kind. This is particularly effective if your site is a highly respected authority in your space. For example, US News and World Report annually publishes a report of the best colleges and universities in the US (<http://colleges.usnews.rankingsandreviews.com/best-colleges>). Being included in this report brings great prestige and many schools put award badges on their site and link back to a copy of the report. This works because it is in fact a true endorsement of the US News and World Report site.

If you are a highly respected authority, you can consider picking out the top 10 sites related to your space. Then you can send these sites an award badge with a link to a review that lives on your site.

In general, this is not a practice that Google or Bing will want to support either, unless you are already one of the top authority sites in your space, and you are able to get the great majority of the award winners to post the badges you provide on their sites. Otherwise, your program is simply not credible.

Customer Discount/Incentives

Publishers can offer visitors from certain websites a discount on the publisher's product. Sites that want to offer the discount to their users simply need to link back to the publisher's site.

However, Google also sees this as a questionable practice. Although it would seem that the discount offer would be interesting to the third-party site only if it valued the product (or thought its visitors would), Google's Matt Cutts clearly indicated that Google did not want to value such links in an interview at <http://www.stonetemple.com/articles/interview-matt-cutts-061608.shtml>.

How Search Engines Fight Link Spam

Throughout this chapter, we have provided many examples of how spammers try to circumvent search engine guidelines to obtain high rankings for sites that may not deserve those rankings. Of course, the search engines do many things to fight link spam.

Google's Penguin Algorithm

On April 24, 2012 Google released a new algorithm they called Penguin. This algorithm focused on identifying sites that have undesirable link profiles and punishing them by lowering their rankings. Some areas of concentration for this algorithm are:

- Article directories
- Low quality web directories
- Overuse of rich anchor text in links back to your site
- Links from countries where you don't market yourself

This algorithm is one of many that Google uses to fight spam, and you can read more about it in Chapter 9 (Panda, Penguin, and Penalties).

Other Algorithmic Approaches to Fighting Link Spam

The major approach the search engines use is to design algorithms that can detect and act on link spam. There are a number of things they can look at algorithmically. Here are a representative few:

Links labeled as advertisements

The search engines can scan for nearby text, such as "Advertisement," "Sponsors," "Our Partners," and so on.

Sitewide links

Sitewide linking is unnatural and should be a rare part of your link mix (purchased or not). The only exception to this is the interlinking of all the sites owned by your company, but this presumes that the search engine will understand that all of your sites are from the same company. In general, sitewide links are a serious flag, especially if you have a lot of different sites that do this for you, or if a large percentage of your links are sitewide.

Links sold by a link broker

Of course, link brokers are knowledgeable about the link detection methods listed here, and they do their best to avoid detection with the links they sell. But they can still run into problems. For example, Google took action against a long-time proponent of paid links, Internet Marketing Ninjas (formerly called We Build Pages, with a site at <http://www.internetmarketingninjas.com/>), resulting in the company changing its stance on the subject (<http://www.internetmarketingninjas.com/blog/link-techniques/paid-links-arent-worth-it/>). In February of 2011, Google took action against J.C. Penney for buying links (<http://searchengineland.com/googles-action-against-link-schemes-continues-overstock-com-and-forbes-com-latest-casualties-conductor-exits-business-65926>), and Conductor (<http://www.conductor.com>) eventually exited the business of brokering buys.

Selling site has information on how to buy a text link ad

Search engines can detect sites that provide information on how to advertise with them. Combined with other clues about links being sold on the site, this could lead to a review of the site selling the ads and a discounting of the links.

Relevance of your link

It is a powerful clue if your link is not really that relevant to the page or site it is on.

Quality of neighboring links

Another clue would be the presence of your link among a group of links that are not tightly themed, or the other sites linked to are poor quality.

Location outside main content

The search engine can detect when your link is not part of the main content of the page. For example, it appears in the left or right column of a three-column site, and the main content is in the middle (<http://patft.uspto.gov/netacgi/nph-Parser?Sect1=PTO2&Sect2=H1OFF&u=%2Fnetacgi%2FPTO%2Fsearch-adv.htm&r=1&p=1&f=G&l=50&d=PTXT&S1=7,716,225.PN.&OS=pn/7,716,225&RS=PN/7,716,225>).

Perhaps you can avoid all of these pitfalls, but one more problem remains: people can see that you are buying links and choose to report your site, using Google's authenticated paid links reporting form (<https://www.google.com/webmasters/tools/spamreport?hl=en&pli=1>). Note that you need to be logged in to a Google account to see the form. Here are some examples of people who might take this action:

Someone reports your site for buying links or for some other reason

Who would do this? Your competitor! If your competitor submits an authenticated spam report to Google, Google will look at it and may choose to act on it.

Someone reports the site you bought links from for selling links or for some other reason

A competitor of yours can do this, or a competitor of the site selling links can do this. Once a search engine figures out that a site is selling links, it is possible that this could trigger a deeper review of the sites that were buying those links.

A disgruntled employee leaves your company, the broker, or the site you bought links from and reports your site

For decades, many companies have had a practice of escorting fired (or laid off) employees out of the building. The reason for this approach is that people get upset when they lose their job. However, this practice would not prevent such a person from reporting your site in a spam report to a search engine. Even though that may be a violation of the confidentiality agreement you probably have with your employees, you would never know, because there is no transparency in spam reporting.

A search engine employee does a manual human review of your site

The search engines maintain legions of people who conduct human reviews of sites, which they use to proactively find and report spam.

Certainly your competitor reporting your site would be the most likely scenario, but you should not entirely discount the other scenarios.

Negative Link Building

Unfortunately, Google's increased ability to detect and act on links they consider to be of poor quality has resulted in a new black hat tactic, that of Negative Link Building. This involved buying poor quality links and pointing them at your competitor's web site. This has proved very difficult for Google to detect and act on.

Even a simple Google search on "Negative SEO" will bring up companies that offer such services in the first page of results. We will not show a screen shot of those results here, as the authors do not in any way endorse this practice. Not only is it unethical, but it also potentially illegal as well.

However, it's important to know that this practice exists, and to take steps to protect yourself from Negative Link Building attacks on your site. Monitoring your link profile on an ongoing basis is one of the best ways to do this.

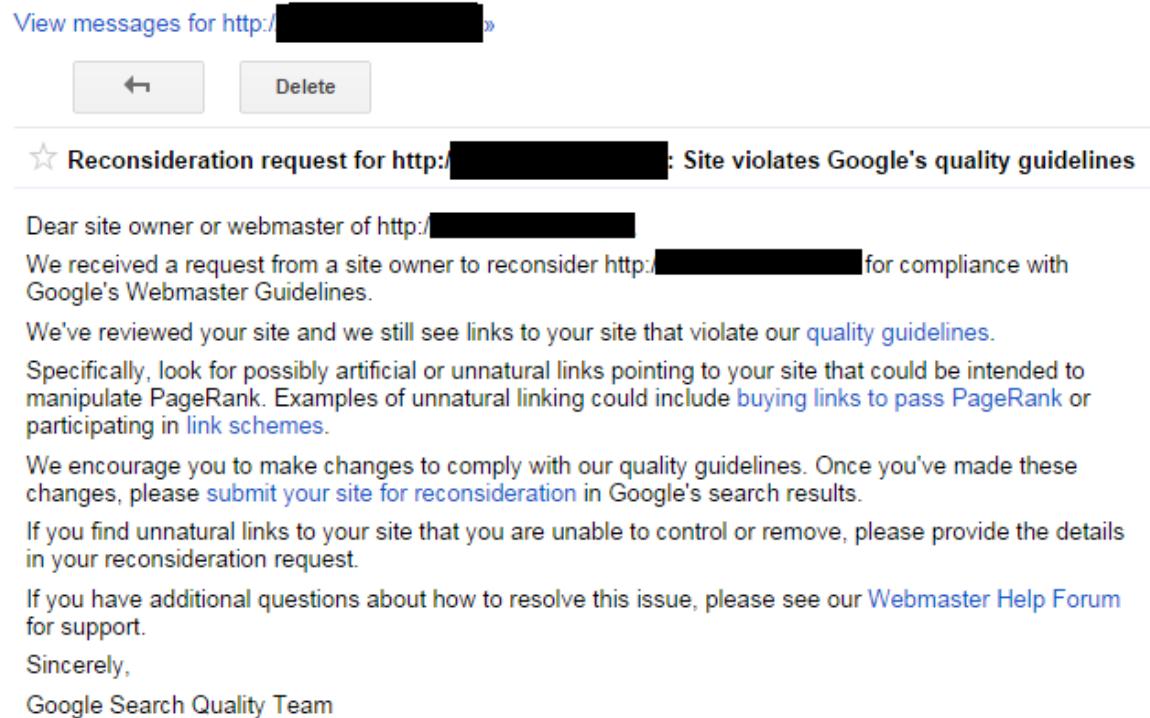
If you see significant numbers of new poor quality links coming in to your site, you may be the subject of a negative SEO attack. Once you identify links that you think are suspicious you can use Google's link disavow tools to indicate that you don't want to have any association with those links. You can read more about this in "The Link Disavow Tool" in Chapter 9. Two other resources to help you read about Negative SEO are:

- <http://moz.com/blog/to-catch-a-spammer-uncovering-negative-seo>
- <http://blog.kissmetrics.com/spammers-and-negative-seo/>

Unnatural Links Messages

Google may also notify you via Google Webmaster Tools if a manual review of your site has detected links that they don't like. This is one of many reasons why signing up for Google Webmaster Tools is so

important (the authors also recommend that you sign up for Bing Webmaster Tools as well). Figure 7-48 shows an example of such a message from Google:



The screenshot shows a message in Google Webmaster Tools. At the top, there is a blue link 'View messages for http://[REDACTED] ». Below it are two buttons: a left arrow and a 'Delete' button. The main content starts with a star icon and the text 'Reconsideration request for http://[REDACTED] : Site violates Google's quality guidelines'. The message body is as follows:

Dear site owner or webmaster of http://[REDACTED]

We received a request from a site owner to reconsider http://[REDACTED] for compliance with Google's Webmaster Guidelines.

We've reviewed your site and we still see links to your site that violate our quality guidelines.

Specifically, look for possibly artificial or unnatural links pointing to your site that could be intended to manipulate PageRank. Examples of unnatural linking could include buying links to pass PageRank or participating in link schemes.

We encourage you to make changes to comply with our quality guidelines. Once you've made these changes, please submit your site for reconsideration in Google's search results.

If you find unnatural links to your site that you are unable to control or remove, please provide the details in your reconsideration request.

If you have additional questions about how to resolve this issue, please see our Webmaster Help Forum for support.

Sincerely,

Google Search Quality Team

Figure 7-48 Unnatural Links Penalty Message in Google Webmaster Tools

Other Search Engine Courses of Action

In the case of Google, it is known that one of the basic policies is to punish a site that sells text links by eliminating that site's ability to pass worth. This is essentially a first course of action. Once this is done, Google could look more closely at the selling site and the purchasing sites for other signs of spammy behavior.

Google has also at times taken the step of altering a site's visible PageRank (the PageRank shown on the Google Toolbar) if it believes a site is selling links. This has been applied to some very significant sites, such as happened to NewsDay.com (<http://newsday.com>) in 2008 (<http://www.seroundtable.com/archives/017658.html>). However, Google appears to have abandoned this practice since then.

Google can also choose to take action against the site purchasing the links or participating in manipulative link practices. If a large percentage of the site's links are suddenly disabled, this could have a significant impact on rankings.

Social Networking for Links

When you add social media to the equation, the network effect can multiply the yield from your link-building efforts. You can use this effect to help your content spread virally, or to develop relationships with critical influencers.

Blogging for Links

Blogging can be very effective in link development. How effective a blog will be depends highly on the content on it, the market space, and how the publisher promotes it. The first thing to realize when starting a blog is that it is a serious commitment. No blog will succeed if it does not publish content on a regular basis.

How frequently a blog needs to publish depends on its topic matter. For some blogs, one post per week is enough. For others, it really needs to be two to three times per week, or even more often.

Blogging is very much about reputation building as well. Quality content and/or very novel content is a key to success. However, when that first blog post goes up, the blog will not yet be well known, will not likely have many readers, and those that do come by will be less likely to link to a little-known blog.

In short, starting a new blog for the purpose of obtaining links is a process that can take a long time. But it can be a very effective tool for link building. Just remember that patience and persistence are required.

One of the best places to get links to a blog is from other blogs. This is best done by targeting relationships with major bloggers and earning their trust and respect. You can read most about this in the section titled “Building Relationships with Influencers” which is earlier in this chapter.

Getting other bloggers to link to you can pay big benefits. Vizu, a market research company, published a study (<http://www.vizu.com/>) that showed that 67.3% of people found what blogs to read by following links from other blogs.

It is also important to leverage the social nature of the blogosphere. Publishers just launching a blog should try to provide a personalized response to every person who comments on their blog. One effective way to do this is to send each and every one of them a personalized response by email that shows that the comment was read.

This helps to deepen the interest of the commenter, creates a feeling of personal connection, and increases the chance that the commenter will return, and possibly add more comments. Nurturing the dialogue on a blog in this fashion helps that dialog grow faster.

When pursuing social media strategies, there are a variety of best practices to keep in mind:

- The number one guideline is to remember that the focus of this discussion is on “social” media sites. Your success in these environments is entirely dependent on building relationships with the other participants on these platforms. If you do something that irritates the other participants, you can quickly find yourself called out and exposed in that community.

Make sure you are a positive contributor.

“Give more than you get” is a great principle to keep in mind. Proactively help promote the great content of others. If you simply share your own content the majority of the time, you will be viewed as someone who does not offer value by association; social connections flourish with give and take.

Leverage Major Social Media Platforms

Building relationships with other bloggers as outlined so far is effective, but there are additional ways to interact with others and in the process to let them know you have content they might be interested in. Any environment in which social interactions occur is another good place to invest time and effort.

Social networks such as those of the current market leaders—LinkedIn (<http://www.linkedin.com>), Facebook (<http://www.facebook.com>), Google+ (<http://plus.google.com>), Pinterest (<http://www.pinterest.com>) and Twitter (<http://www.twitter.com>)—can be used for link-building purposes without actually getting the link authority directly from one of these sites. Major forums that relate to your area of interest also represent great targets. As we will discuss in more detail in Chapter 8, the links implemented directly on these sites are NoFollowed, so they do not pass direct SEO value. However, social media can help get you exposure to bloggers and media people who may then choose to cite your

content or write about it, and implement links back to you. This indirect form of obtaining links can be quite powerful.

Building relationships and trust on a social media platform can drive lots of traffic and links to your site. The mechanics of making this work are fairly straightforward:

- Publish great non-commercial content related to your business on your site.
- Establish and grow social media presences on major social media platforms that promote the same types of content.
- Share your content via these platforms.

If you have done this well, sharing your newly published content through your social media channels should help your visibility and reputation on that social media platform grow. In turn, the exposure to the audience within that social platform should drive traffic and links back to the content. This synergy takes time to build, but becomes a valuable asset in time - a strong social media presence is a powerful ally to any content marketing or SEO strategy.

Of course, while establishing connections and building relationships in a social media environment you also have the potential for people to submit your content into those social media sites for you, which can lead to links as well. You can then take the next step and reach out through the social networks or forums to make direct contact with people to let them know about your content. This is similar to emailing people, but with a few important distinctions:

- Publishers can send out communications to their friends on those networks. Assuming that they have treated this designation (that of friend) with any level of seriousness, instead of “friending” everybody in sight, the communication can be a bit more informal than an unsolicited email would be.
- Publishers can also join groups on these networks related to their market space, and then send messages out to those groups. These groups will provide them with the ability to reach new people with related interests.
- Messages broadcast through these networks cannot be personalized, so a more general message needs to be tailored for these types of broadcasts.
- These are social networks. Beware of broadcasting too many messages or poorly targeted messages. Many publishers have made this mistake, became pariahs in their communities, and lost the leverage these communities bring to the process.
- Personalized messages can be sent on a 1:1 basis as well.

One strategy for approaching an authority site is to make initial contact by friending someone senior who works at the company that publishes the site. Then the publisher can develop the relationship with that senior person without asking anything of her. Once the relationship is established, a more informal approach can be used to introduce the great content on the publisher’s site.

Social Media Networking Tips

Although we have outlined some of the basics of how to work with social media sites, each site has its own quirks and opportunities. What follows is an outline of tips to use on many of these platforms. We will discuss Facebook, Google+, Twitter, and Pinterest here briefly, but other aspects of leveraging these channels will be discussed in Chapter 8 (“How Social Media and User Data Play a Role in Search”).

Nearly all social media site links are NoFollowed, so don’t expect these to pass link authority through to your site, but visibility on these sites can still have a large impact on your SEO. A strong social media presence can act a lot like a powerful PR channel.

By publishing great content and then sharing it through your social media accounts you can potentially gain fast exposure to bloggers and media (if they are following you) and this can lead to links being implemented to your site.

Facebook

People share links on their timeline and on Facebook pages. Links shared by users who have marked their content as available for Everyone (<https://www.facebook.com/notes/facebook/opening-more-control-for-everyone/60186587130>) are visible to search engines as are Facebook Pages (pages for companies as opposed to people). These links are all NoFollowed, so they do not pass PageRank, but they can provide visibility for your content.

Facebook Likes are not at all visible to Google, as there is no place for it to crawl to see what content you have Liked. However, Bing receives Like data via an API feed, so they have the possibility of using it, though indications are that they are not as of September 2014.

There are many reasons to use Facebook, but you should not expect it to directly impact your SEO ranking. However, a strong social presence can play a big role in your reputation and visibility online, and as a result can have a powerful indirect impact on your SEO. Facebook can be used to build powerful relationships.

In particular, Facebook has powerful advertising features that enable you to significantly increase your reach and exposure across the platform, and this can give you tremendous exposure.

Google+

Google+ was launched on June 28, 2011 and very quickly built a large user base. Google+ also allows people to share links and generate comments and interactions on its posts. Shared content on Google+ can be evaluated based on the authority of the person sharing it, how many times it gets re-shared, the number of comments it generates.

Users can also indicate that they found a page valuable by clicking on a +1 button on the page (if the page has one) or next to the search listing for that page. The +1 button functions in the much the same way as Facebook's Like button, except that Google has direct access to the +1 data, as well as sharing data on the platform. This plays a big role in personalized search on Google.

Google+ also offers brand pages, and these are used by many major brands to communicate with customers and potential customers. In addition, these show up in the search results whenever you search on such company's brands, and this can be very valuable real estate, and in some cases, it blocks out people from using Google AdWords from bidding on your brand name.

Consider focusing on connecting with major influencers on Google+, and using that connection to build a real relationship with them. Add value to their posts and they are likely to pay more attention to what you choose to post. This will help your Google+ account build its authority and influence.

Twitter

Twitter has established itself as the leading microblogging site. It allows its members to contribute microblog posts (a.k.a. "tweets") that are limited to 140 characters. It has become an environment for real-time communication with a broader network of people. You can use Twitter as an effective platform for promoting your business, though links from Twitter are NoFollowed so no PageRank is passed. It is another valuable channel for communicating with your customers and market.

The basic concept is to become an active member of the community and build a large network of followers. As with other networking sites, many important influencers spend time on Twitter. If you can use Twitter to develop relationships with these people, that can provide some very high-quality links.

In addition, if you create some high-quality content, you can gain a lot of exposure as a result of getting a lot of "re-tweets" (whereby people forward your message to all of their followers). Particularly interesting tweets can get significant visibility for you and your company.

reddit

reddit is a social link aggregation and discussion site, organized into topical communities that are self-organized and independently moderated. There are millions of visitors each day, and reddit has hosted "ask me anything" interviews from many celebrities and a few politicians, including Barack Obama.

It has a lot of traffic, a lot of users, and a lot of discussion. You must be a member of reddit in order to advertise there. It isn't difficult to create an account, but the ad targeting is not very good because there is little capacity for users to provide revealing demographic data about themselves. In general, ad results seem to be less than stellar for any product or service that doesn't specifically appeal to self-described "geeks".

redditors are generally hostile toward self-promotion of any kind. If you try to guerrilla market by posting links to your site or Facebook page or anything that looks remotely promotional, redditors will slap you down very, very hard.

For that reason, it is best to not plan on any form of self-promotion on reddit. All you can and should do on reddit is try to participate normally and submit your viral content and a lot of other people's great content to the correct subreddits without the appearance of self-promotion.

Pinterest

Pinterest is a site for sharing pictures that has rapidly built a strong following. The demographic skews heavily towards women, with more than 2/3 of the site's users being female. The average income level on Pinterest is also relatively high (<http://www.businessinsider.com/demographics-on-pinterest-that-make-the-social-network-attractive-to-marketers-2014-4>).

However, even brands that are not naturally highly visible do well on Pinterest. For example, as of August 2014, the Wall Street Journal has more than 25,000 followers. Yet all of their "pins" are quoted embedded over a newspaper background (See Figure 7-49 for an example of this), or next to a picture of the person being quoted.

Pin it 9 Like 5 Visit Site Send Share

Found on online.wsj.com

Pinterest CEO Lays Out Growth Plan, Sees Revenue in 2014

by Douglas MacMillan

 online.wsj.com

Ben Silbermann discusses Pinterest's path to revenue and what he has learned leading one of the Web's fastest-growing startups.

Figure 7-49 Example of a Wall Street Journal Posts on Pinterest 07-pinterest-wsj.png

While Pinterest is smaller in audience size than Facebook, Twitter, or Google+, due to the demographics, the site has proven to work really well for certain retailers, such as apparel manufacturers or restaurants. It also could be an excellent time to jump on board and build a strong following before your competition does!

Wikipedia

You may not think of Wikipedia as being useful in link building since it `NoFollows` all its external links. In addition, it is not a wise idea to simply go onto the site and create a page for your company. Unless your company is well known, it is likely to be promptly removed. Even if it is well-known, you would be in violation of Wikipedia's Conflict of Interest guideline. However, links from Wikipedia can be valuable because many people treat it like an authoritative site, and if they see your link on Wikipedia they may choose to link to you. This can include some pretty influential people.

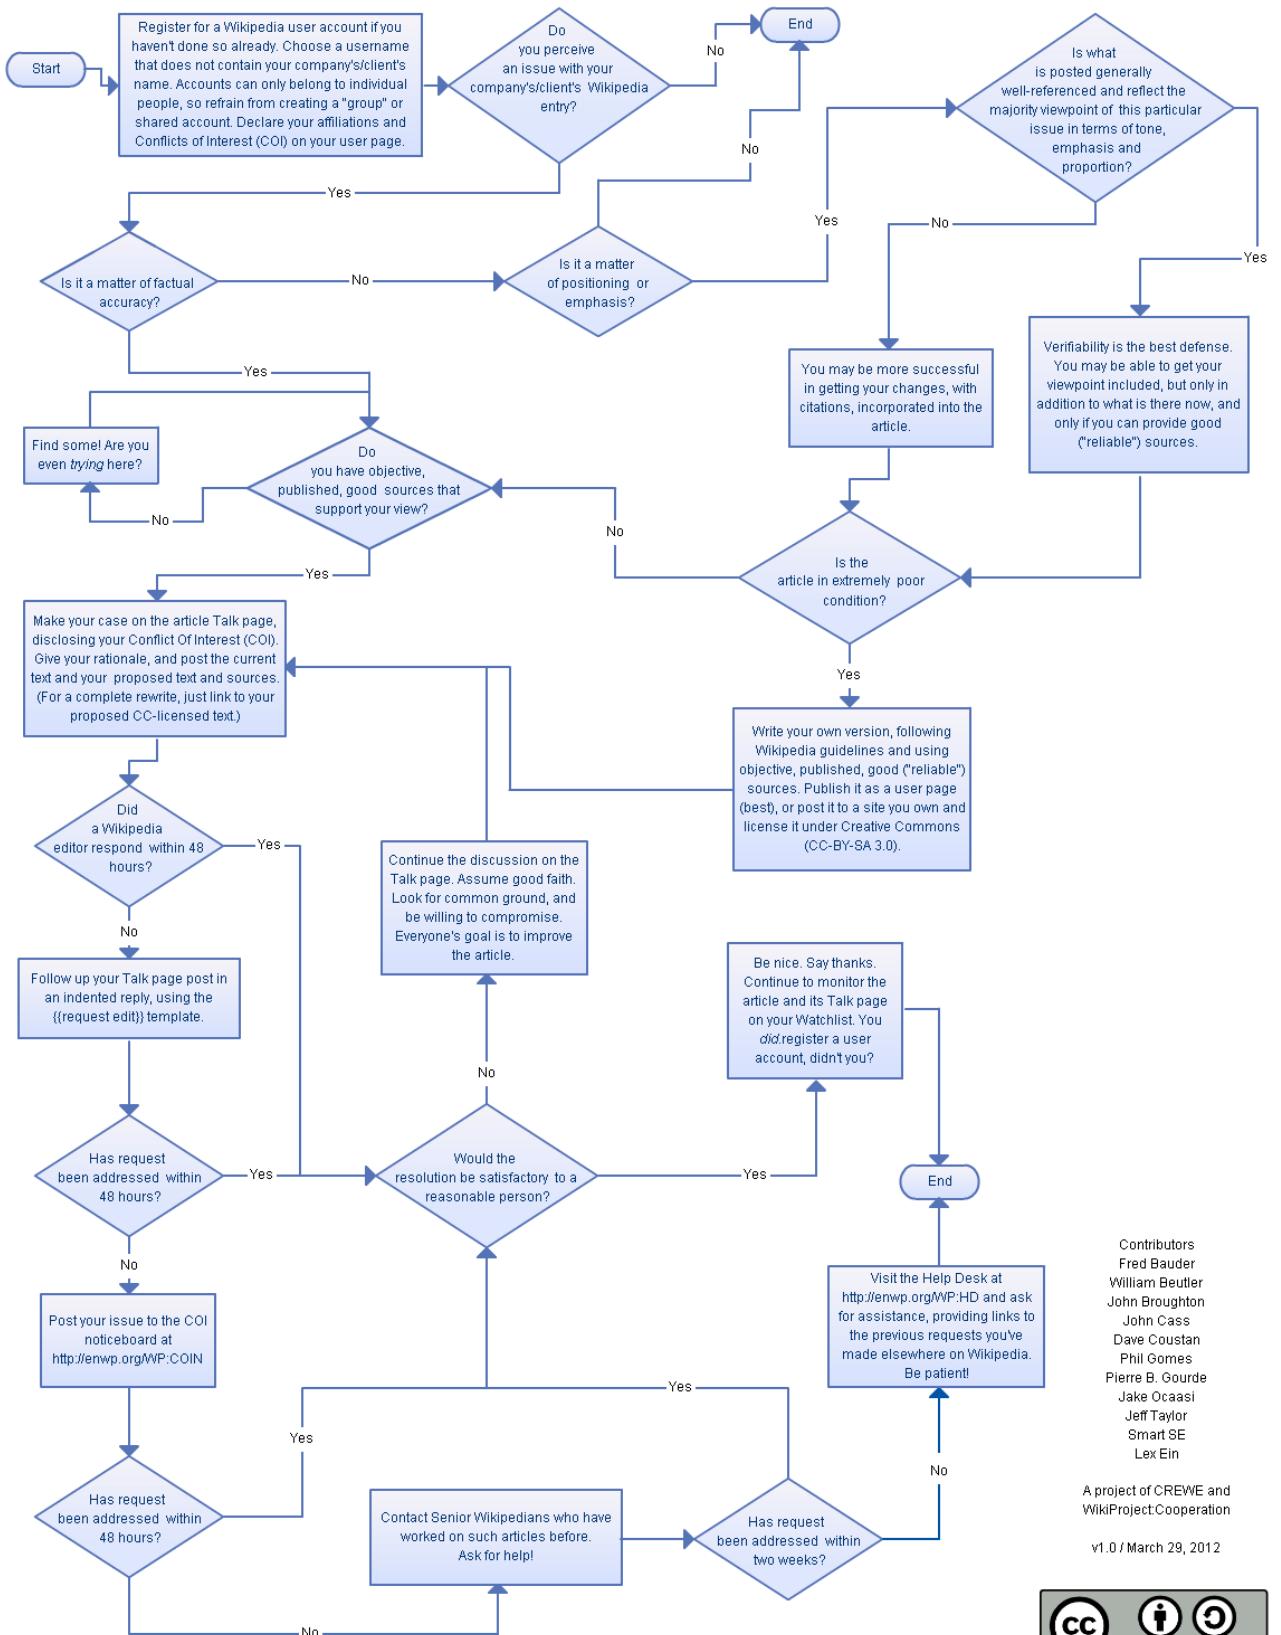
Therefore, you may want to build trusted relationships within the Wikipedia community. This will require an investment in time, but it can bring some nice rewards. Here are some specific tips on how to proceed:

- Build up your credibility before doing anything that could be construed as self-serving. Develop a long and virtuous contribution history, and ideally build a user profile page that has earned Barnstar awards (you can read more about these here: <http://en.wikipedia.org/wiki/Wikipedia:Barnstars>). It is not good enough to be altruistic on Wikipedia unless you *demonstrate* it. It has to be visible as a track record (e.g., squash spam and fix typos and add valuable content, but do not do it anonymously).
- Negotiate with an article's "owner" (the main editor who polices the article) before making an edit to an article to get her blessing first.
- Use Wikipedia's "Watch" function to monitor your articles of interest. Better yet, use a tool that emails you (e.g., TrackEngine or ChangeDetect).
- Make friends. They will be invaluable in times of trouble, such as if an article you care about gets an "Article for Deletion" nomination.
- Do not edit anonymously from the office. This could come back to haunt you. Tools exist that could embarrass you if you do so. One public domain tool, WikiScanner, is able to programmatically take anonymous Wikipedia posts and identify the organization that created them. The tool cross-references the IP address to make the edits with the blocks of IP addresses of more than 180,000 organizations. Do not take the risk.

If you are serious about pursuing getting a page for your business in Wikipedia, you should also learn about the Corporate Representatives for Ethical Wikipedia Engagement (CREWE). This started as a Facebook group comprised of marketing and PR people and wikipedia editors (<https://www.facebook.com/groups/crewe.group/>).

CREWE is comprised of Wikipedians, corporate communications, academics, students and other interested parties who are exploring the ways that PR and Wikipedia can work together for mutual benefit, defined narrowly as cooperation toward more accurate and balanced entries.

A consortium of PR associations and agencies signed a pact in 2014 to promise to work ethically within Wikipedia's rules. A related e-book was produced by some of the participants, a guide to working with Wikipedia (<http://www.instituteforpr.org/wiki-guide/>). The following flow chart provides a view as to how this ideally works:



Contributors
 Fred Bauder
 William Beutler
 John Broughton
 John Cass
 Dave Coustan
 Phil Gomes
 Pierre B. Gourde
 Jake Ocaasi
 Jeff Taylor
 Smart SE
 Lex Ein

A project of CREWE and
 WikiProject:Cooperation

v1.0 / March 29, 2012



[online diagramming & design] creately.com

Figure 7-50 Example of an Article that Went Viral

Wikis

Plenty of other wikis are a lot more edit-friendly than Wikipedia and let you contribute valuable content, get links, and build relationships. Examples include ShopWiki (<http://www.shopwiki.com/>), The NewPR Wiki (<http://www.thenewpr.com/wiki/pmwiki.php>), and WordPress Codex (http://codex.wordpress.org/Main_Page). Some even pass link authority, which is a nice bonus.

Meetup.com

Meetup.com (<http://www.meetup.com>) is a site designed to help groups of people with interests in a common topic “meet up.” These meet-ups can be excellent networking events.

Get involved with relevant local meet-ups and get your Meetup.com member profile page linked from the Meetups page, which will pass link authority to your profile and then on to your site.

Pinterest, Instagram, Vine, and Snapchat

New social media sites emerge on the scene all the time, and each of them has their own unique attributes. Like Pinterest, Instagram is a photo and image sharing network that skew towards women. For the right type of brand where visual content can play a big role (such as food, beauty, lifestyle, media and fashion), these can be great sites to build a presence on.

Vine is unique because it constrains you to producing videos that are 6 seconds in length. A perfect environment for people that want to see some really compelling and quick. An interesting challenge because there is so little time to deliver the message, but very effective when you get it right.

Similarly, Snapchat videos disappear after you watch them. Consider the concept of delivering a message which provides a sense of exclusivity because it does evaporate on you.

Content Marketing Success Stories Using YouTube

It may become more important for your brand or company to be on YouTube than to be advertised on TV. For some that day has already arrived. August 2011 data from comScore showed that YouTube is used for about 14% of all searches on the web

YouTube has launched careers such as that of the zany Jenna Marbles, whose YouTube channel (<https://www.youtube.com/user/JennaMarbles>) has more than 13 million subscribers. She has recorded many videos that have tens of millions of views including:

- How to Trick People into Thinking You are Good Looking – 58+ million views (<https://www.youtube.com/watch?v=OYpwAtnywTk&list=TLIrQSSRE8Joo0Y57HDDaCagdqBkyXSBRs>)
- How To Avoid Talking To People You Don't Want To Talk To – 33+ million views (<https://www.youtube.com/watch?v=8wRXa971Xw0&list=TLw8cVksDftKUSr0SBp-0szlMatni1rhZ>)
- What Girls Do in Their Car – 33+ million views (https://www.youtube.com/watch?v=0i_5YBnQdac&list=TLAY-7GnOfg0LEsQrFQPE688xEzLYr_Frb)

YouTube has also brought international fame to previously unknown bands, such as Justin Bieber, who got his start on YouTube, or Sick Puppies, a band popularized by the hugely well-liked and inspiring Free Hugs video set to the Sick Puppies song “All the Same”, which has over 75 million views (http://www.youtube.com/watch?v=vr3x_RRJdd4).

Then there are the hugely successful viral campaigns by commercial organizations, such as Blendtec’s Will It Blend (<http://www.willitblend.com>) that we referenced previously in this chapter. This is a brilliant video

series on various household objects that are run through a Blendtec blender—including marbles, rake handles, and even iPods.

Blendtec isn't the only company that has had success with YouTube. Westjet scored a big hit with their video: WestJet Christmas Miracle: real-time giving (<https://www.youtube.com/watch?v=zIEIvi2MuEk>). In this video the filmed people sharing with Santa via a video feed their Christmas wishes. While the passengers were in their flight, WestJet employees rushed off to stores to buy all the gifts people had asked for.

Once they landed, the first baggage to come off in baggage claim was gift wrapped packages that people opened, and to their amazement, they found what they had asked for. This was also videotaped, and the resulting video received over 33 million views.

Another stellar example of YouTube marketing is the online retailer Vat19.com. Their product videos are brilliant and well-integrated into their product catalog. For example, the 5 pound gummy bear is a must-see (<http://www.vat19.com/dvds/worlds-largest-gummy-bear.cfm>). Their channel has more than 850,000 subscribers, with more than 440 million views.

YouTube has been used effectively for brand damage control as well. For example, the CEO and founder of JetBlue Airlines put up an apology video (http://www.youtube.com/watch?v=-r_PlIg7EAUw) on YouTube because of a Valentine's Day winter storm incident—a campaign that was well received.

One product that got some excellent brand recognition and building from being on YouTube was Smirnoff's Raw Tea. Smirnoff produced an uproarious music video called "Tea Partay" (<http://www.youtube.com/watch?v=PTU2He2BIc0>), with preppies rapping.

The problem with most popular YouTube promotions is that YouTube gets the links and the original site usually does not. That means the search engine visibility benefits do not usually transfer to the company's website.

Even without creating your own site or hosting your videos on your main site to draw links, YouTube offers much in the way of brand visibility when the campaign is well executed. That doesn't just mean posting a great video; marketers must also know how to take advantage of the social nature of the site, to build up friends and to get on user subscription lists.

Make copious use of tags on your videos (ensuring, of course, that the tags are relevant to the content), spread your tags out among your clips, use adjectives to make your videos more visible to folks who are searching based on their mood, have some category descriptor tags (bearing in mind that YouTube's default search settings are Videos, Relevance, and All Categories), to match your title and description with your most important tags, and don't use natural language phrases or waste tag space on words such as *and* and *to*. See the video optimization section in Chapter 10 for more tips.

Do not be afraid to make a start, even if it is modest and has little budget behind it. You won't get anywhere without experimenting with the medium. Ongoing experimentation provides you with the best chance to find a formula that works for you.

Guest Post Success Story

Susan Joyce is a well-known author of content on the topic of job search, and she is also a contributor to the Huffington Post. In March of 2014 she placed an article on the Huffington Post titled: "What 80% of Employers Do Before Inviting You For An Interview". Figure 7-50 shows us that this article was a smashing success.



Susan P. Joyce

Become a fan



Online job search expert, Job-Hunt.org writer

What 80% Of Employers Do Before Inviting You For An Interview

Posted: 03/01/2014 7:45 am EST | Updated: 05/01/2014 5:59 am EDT



Many job seekers have described to me that submitting a resume in today's job market is mostly a banging-their-head-against-a-wall, extremely frustrating waste of time.

Figure 7-51 Example of an Article that Went Viral

The article netted well over 30,000 social shares. What made this article such a success? Certainly being placed on the Huffington Post was a big help, but the title was a big key to the success of this content. It drew people in as everybody wanted to know the answer to the question.

In fact, the answer was quite simple – Employers search on your name you before asking you in for an interview. Yet the simplicity of the observation did not diminish the value of the article, because the observation was not completely intuitive.

Once you see it, your reaction is, of course, they do. For most searchers, the result then is to go search on themselves, and see whether or not employers will like what they see!

Social Media Summary

To sum up the social media strategy outlined here, it is useful to think about it as an old-fashioned PR/marketing strategy. Marketers used to think about the number of “impressions” they were able to create in people’s minds. The more impressions you can create the better. When TV, print, and radio were the only media to worry about, this was relatively simple to execute.

The current environment is significantly more complex. The people you are trying to reach use many venues, and all of the various social media properties are part of that picture. Whenever you have a social media property that a large number of people use, there is an opportunity to reach them. Success in reaching them depends on becoming a trusted member of that community.

Since there are many large communities, it can be complex and expensive to establish a presence in all of them. But presence in each one does add to the number of opportunities you have for creating impressions in your target audience. Since major influencers in your market may be using these communities, you have an opportunity to reach them as well.

Pursuing social media sites can be a very effective strategy. It does require a real investment to do well, but when it's done well the strategy can provide some real benefits. You can read more about Social Media and Search in Chapter 8 - "How Social Media and User Data Play a Role in Search Results and Rankings".

Conclusion

Quality content will naturally attract, and earn, inbound links – and links remain a large factor in search rankings. Bear in mind that the best links are the ones that you would consider valuable even if there were no search engines. These are links that can deliver traffic to your site on their own, and are the ones that are most likely to be seen as valuable by search engines in the long term. A solid content development and content marketing plan is essential to all of your online efforts, not just SEO.

You should view content development and marketing as an ongoing activity, ideally with an editorial calendar around which optimization and distribution strategies are organized. Each of the authors has seen cases where a brief focus on link accumulation (with or without focused content marketing) brought returns that were squandered by abandoning the strategy. Unfortunately for these sites, they lost momentum and rankings to their competitors (the same ones they had previously passed), and it proved very difficult to catch up to them again.

Content marketing is fundamentally similar to online public relations work: your goal is to build your visibility and reputation online, and acquire links that can potentially drive relevant traffic to your web site. The big difference between content marketing and traditional PR is in the technical aspects and the focus on new media methods of promotion—with emphasis placed on the quality of the referring source, the relevance of that source, and the page(s) to which they point.

People will generally not link to low-quality content or sites that offer a poor user experience (unless they are compensated for the link), and unless you are fortunate enough to possess a major brand, people won't link to purely commercial sites, either. You have to offer something of value to users while also offering something unique. Certain content naturally attracts links because it triggers psychological and emotional responses—pride, sharing, newsworthiness, and so on. Leverage these triggers and create a compelling reason for visitors who can influence web content (writers, publishers, bloggers, etc.) to reference your work, and your content marketing efforts will be a success – both from a visibility and earned links perspective.

Great content marketing comes from the simple idea: "Build great content, tell everyone about it, and motivate them to share."

8

How Social Media and User Data Play a Role in Search Results and Rankings

Search engines continually seek to improve the relevance of their search results. They do this by tuning the weight and mix of the types of ranking signals they currently employ, or by implementing new signals. However, how they use these signals is changing all the time. Starting in 2010, evidence mounted to suggest an increasing weight on ranking signals from social media sources. In December of 2010 Google and Bing both confirmed this in response to questions from SearchEngineLand.com editor Danny Sullivan (<http://searchengineland.com/what-social-signals-do-google-bing-really-count-55389>).

However, the way the search engines use social signals has changed significantly since then, and it currently appears that neither Google or Bing use them as a direct ranking factor, except that they can impact personalized search results in Google.

While the search engines obscure how their algorithms work, there are many ways that people believe user engagement signals are part of their algorithms – and there has been a long debate over whether search engines have treated user engagement signals as ranking factors.

This is a process that will continue to evolve rapidly. Adding new types of ranking signals to search engine algorithms is a tricky process that requires a tremendous amount of testing. As we discussed in Chapter 2 the web consists of hundreds of trillions of pages with fundamental differences in their construction and content. In addition, the needs these pages serve, and the ways that users interact with them, is equally varied.

Search engines generally perform many levels of testing of new algorithms including live tests with a small portion of their users prior to any general release of any new ranking factors (you can read more about all the stages of testing in this interview with Google's Director of Research, Peter Norvig: <http://www.stonetemple.com/search-algorithms-with-google-director-of-research-peter-norvig/>). In the same interview Peter Norvig suggested that Google makes 2 to 3 algorithm changes per day. This means change is a constant.

The key is to remember a key goal of the search engines – to generate revenue by increasing market share, by providing the best possible (i.e. most relevant and user-engaged) results for end users. Astute internet marketers who understand this will focus on providing high quality web site content and experiences for their visitors.

Correlation between Social Signals and Google Rankings

moz did a correlation study (<http://moz.com/blog/google-plus-correlations>) that showed a strong correlation between Google+ +1s and rankings in Google. The detailed results are shown in Figure 8-1.

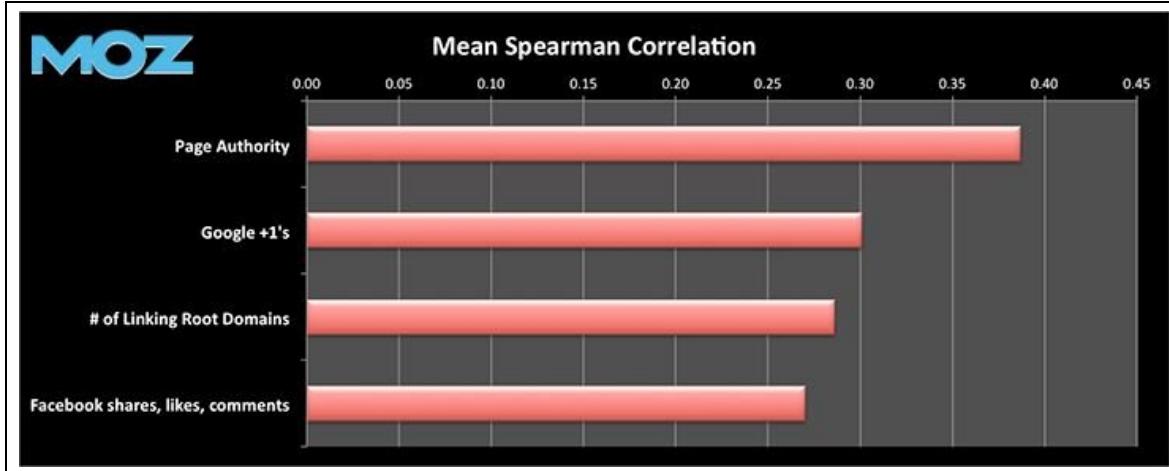


Figure 8-1. moz Correlation Study

In addition, Searchmetrics has performed their own correlation study. As shown in Figure 8-2, this study showed a very strong correlation between many different types of social signals and search ranking. In fact, 7 of the top 10 factors correlating with ranking were related to social media sites.

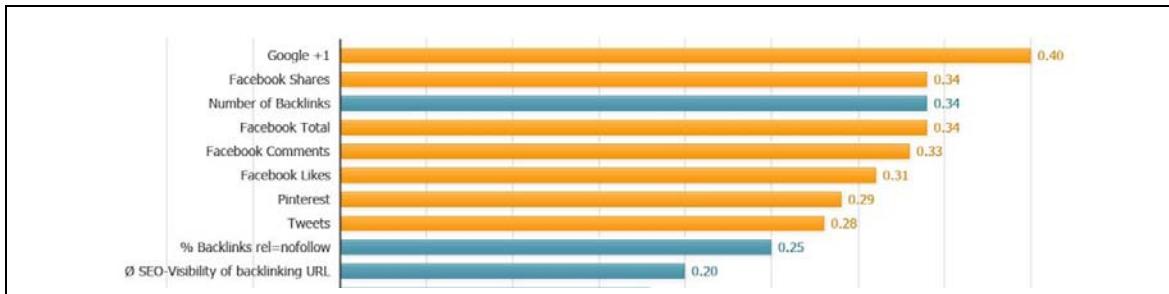


Figure 8-2. Searchmetrics Correlation Study

It is important to note that a correlation does not prove causation – in other words, this information does not prove that these social media factors are a direct ranking signal.

Nonetheless, the correlation is noteworthy and only supports the long-standing belief of the authors that great content has a strong chance of earning a high number of traditional web links, and of being shared more in social media.

What is the Value of Social Signals?

Google was first established in 1996 based on the concept of PageRank, where links from third party web sites would serve as votes for the quality of the site receiving the links. This was a valuable concept because the web evolved as an environment where creating great content was rewarded through links and the cross referencing of other great content, a natural result.

The collection of data on all the links of the web, which is referred to as the *link graph*, is an enormous computational task and results in a map of who links to whom in a vast diagram of the interconnectivity of

the web. Figure 8-3 shows a conceptual representation of the link graph. From this, search engines can extract valuable signals to identify relevant web pages in response to a particular search query.

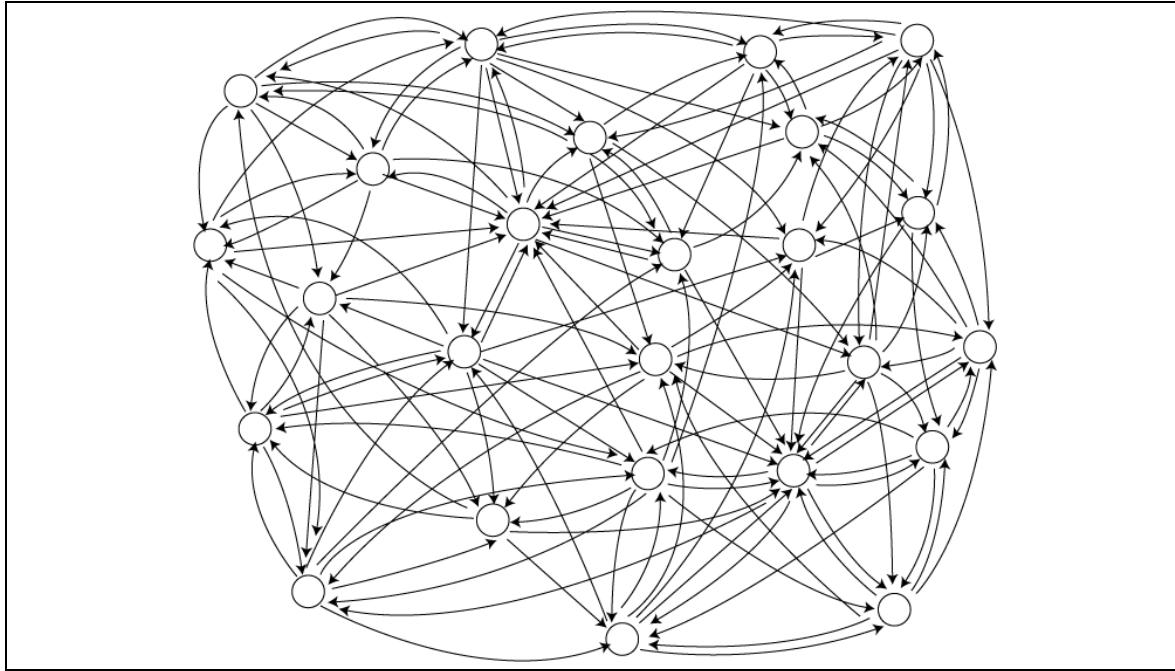


Figure 8-3. The link graph is a map of all of the World Wide Web's hyperlinks

This system, however, was dependent on the links existing due to editorial citations. In other words, a link was created by a publisher on a third party website because they genuinely cared to send visitors from their site to the site to which they linked solely because it was a great resource. This was the primary incentive that the publisher providing the link had for doing so. Links were in fact the first social signal used by search engines to aid in ranking web pages.

Over time there was so much money to be made that many began cheating the system by purchasing links and coming up with many other schemes that were not citations at all. This link selling economy along with a variety of other schemes for getting links which are not earned based on the merit of the content, generally referred to as link spam, has threatened search quality and resulted in a war between search engines and the link spammers.

Similar to the way in which the *link graph* acts as a proxy for real world connection, social connections and user behavior patterns reflect real world relationships and values. In order to implement a link to a website, one must first have a web property. While creating a website can in fact be free on today's web, one must commit time and energy to that site to be viewed as a valuable resource. As a result, many do not make that investment.

Participating in social media, such as Facebook, Google+ or Twitter, requires zero cash outlay (beyond an internet connection and a connecting device) and a much smaller amount of time investment. As a result, it is more populist in nature than the link graph. Of course, the use of social media is not universal either, but is far more broadly adopted than the publication of web sites.

As users, we like proof from our peers that something is good. If your friend liked it, and you trust their judgment, you become more comfortable in checking that out yourself. For example, your trusted friend is unlikely to have endorsed a spam site.

However, problems arise when you try to use social signals as a ranking factor in the search results. The search engines experimented with these signals and pulled back because they found problems with it. There are some ways in which social information is currently integrated into the search results, and these will be explored in the next few sections.

Bing's Experiments with Social Signals

Bing has relationships with many social media platforms. These relationships provide Bing with real time access to any updates that occur on those networks. For example, imagine someone posts an update on their Facebook page, Bing gets notified of that update via API, which means they get the data relatively quickly, and with minimal overhead. Social sites that Bing has such relationships with include:

- Facebook
- Twitter
- FourSquare
- Quora
- Klout

In contrast, Google has to crawl sites like Facebook and Twitter to obtain comparable social update info. This is problematic for many reasons as the overhead of doing this is quite high. This is what makes the experiments that Bing has conducted interesting to review in detail.

For example, Bing experimented with using data from Facebook as a ranking factor in the search results. Figure 8-4 Shows how Bing's search results looked in 2010 if you performed a search on the phrase "New York Post".

Customer service 800-552-7678

Search within nypost.com

Search

Facebook icon: Liked by people on nypost.com

amfAR Gala: 15 Best Dresses

54 people like this

Gym destroyed by fire, Regis boys volleyball team moving on

17 people like this

Figure 8-4. Bing Search Results from 2010

Notice how two articles with a number of Facebook Likes are highlighted in the results. However, if you performed the same search in 2011, the nature of the result had changed quite a bit. Figure 8-5 shows a very different result.

Customer service 800-552-7678

Search within nypost.com

 [Recent activity from New York Post](#)

 [New York Post: http://www.nyp.st/nhsxsG](#) BREAKING BROOKLYN SHOOTING-
Pregnant woman killed, two wounded in shooting outside public school<http://nyp.st/n6vFwq> · 15 hours ago

[www.facebook.com](#)

Figure 8-5. Bing Search Results from 2011

The result has changed quite a bit. The notion of showing the most Liked articles has gone, and instead, the highlighted item is a post that is one of the more recent ones in the paper. While it is not indicated in the search result, it is possible that this is something that has received recent Facebook Likes, but we don't know that for sure. Figure 8-6 shows us however that this ambiguity is now gone from Bing's results as of 2012.

Sports	Yankees
News	Local
From Page Six	Entertainment
Opinion	Business

Figure 8-6. Bing Search Results from 2012

In a February 2013 interview by Eric Enge with Bing's Stefan Weitz (<http://www.stonetemple.com/graph-search-social-search-with-bings-stefan-weitz/>), he had this to say about using Likes as a ranking signal: "Not now, no. The notion of a Like is still a little bit perplexing from a ranking perspective. What does a Like mean for a page? Does the user like the design, the content, or maybe just the picture? We tend to not just use a pure Like signal to do ranking." He also said: "Shares are basically the same as likes – not used a ton for web ranking except in velocity (like Twitter)".

For a while, Bing went in a completely different direction. They separated out the social part of their search results into a Social Sidebar. Figure 8-7 illustrates the way this implementation looks:

New York Post
nypost.com Official site
Customer service 800-552-7678
News, Gossip, Entertainment, Sports & Opinion ... You must be honest about what it is you desire. Your needs are perfectly valid - there is no need to conceal them.

Sports
Sports News, Analysis, Scores & Commentary | New York Post

News
Latest news, top stories, news articles - NYPost.com.

Yankees
New York Yankees - NYPost.com. Get news on the New York ...

It's a Boy
Renowned street-artist Banksy etched his first piece in New York...

Entertainment
Entertainment News, celebrity news, arts & entertainment

Gossip
Grown up à©Harry Potterâ€ star Daniel Radcliffe and his reported ...

Search within nypost.com

[Several cops may have watched biker beatdown | New York Post](#)
nypost.com/2013/10/04/undercover-cop-stood-by-as-bikers-beat-dad ▾
Oct 04, 2013 · An off-duty undercover NYPD cop was among the pack of bikers who chased a family up the West Side Highway — and he stood by as the dad was hauled ...

[News about New York Post](#)
bing.com/news

 [New York City Marathon Balances Security Worries Post-Boston](#)
NBC New York · 2 days ago
The New York City Marathon's finish area will again fill with spectators, some carrying bags. More bomb-sniffing dogs also will be roaming around. Less than...

[New York Giants vs Philadelphia Eagles Live Stream Free: Watch Online NFL 2013](#)

Social Sidebar

Facebook

TOP RESULT

 **Jim Hedger**

29 Aug 2013
... expert's analysis of the DNS attacks on the New York Times, Huffington Post and Twitter at...

[WebmasterRadio.FM](#)
www2.webmasterradio.fm

Like (2) · Comment (0) · See post

Status updates (6) ▾
Photos (1) ▾

Figure 8-7. Bing's Social Sidebar

Does Google Use Facebook as a Ranking Signal?

As previously mentioned, unlike Bing, Google must crawl Facebook to get data on activity on the site. Consequently, Google will not be able to get any data from pages which are not marked as Public by the profile owner. That is a very large crawling burden. We can see how many pages Google has indexed from Facebook, as shown in Figure 8-8

Google site:facebook.com

Web Images News Shopping

About 5,140,000,000 results (0.37 seconds)

Figure 8-8. Facebook Pages Indexed by Google

Google has indexed 5.14 billion pages from Facebook according to this query. To put this in perspective, data released by Facebook in January of 2014 indicated that they had 757 million users, and that 10% of these put out daily updates, and 4% put out more than one update per day.

We can use these numbers to estimate the number of total Facebook updates per month. We can start by taking 10% of the total number of users, and multiplying that by 30 to account for all the users that do at least one update per day. This works out to 2.27 billion updates per month.

However, we also know that 4% of users perform multiple updates, and we can estimate their usage by guessing that they average 3 updates per day. Of course, some perform only 2, and some do much more than 3, and we also have users who perform less than 1 update per day.

For purposes of this analysis we are going to estimate that by calculating the total of incremental tweets as the total number of additional updates per day as the number of users times 2 times 30. Note that we use 2

in this calculation instead of 3, as have already counted 1 of their daily updates in our prior calculation. The result is another 1.82 billion updates per day.

This gives us an average of 4.09 billion Facebook updates per day, which is already more than Google has indexed from Facebook in total! The study also tracked the monthly pages Google indexed from Facebook from January 2012 through March 2014. As shown in Table 8-1.

Table 8-1. Facebook Indexed Pages by Month

Month	Indexed Pages	% of Estimated Updates Per Month
March 2014 (third)	77,900,000	1.91%
February 2014	4,330,000	0.11%
January 2014	5,490,000	0.13%
December 2013	7,840,000	0.19%
November 2013	7,160,000	0.18%
October 2013	8,510,000	0.21%
September 2013	4,160,000	0.10%
August 2013	26,200,000	0.64%
July 2013	48,400,000	1.18%
June 2013	168,000,000	4.11%
May 2013	166,000,000	4.06%
April 2013	166,000,000	4.06%
March 2013	3,160,000	0.08%
February 2013	2,770,000	0.07%
January 2013	1,950,000	0.05%
December 2012	2,740,000	0.07%
November 2012	2,690,000	0.07%
October 2012	2,630,000	0.06%
September 2012	1,550,000	0.04%
August 2012	2,570,000	0.06%
July 2012	2,490,000	0.06%
June 2012	2,440,000	0.06%
May 2012	1,030,000	0.03%
April 2012	2,510,000	0.06%
March 2012	2,150,000	0.05%
February 2012	1,860,000	0.05%
January 2012	856,000	0.02%

As you can see, the indexing of this content appears to happen at a very low rate. Note there are some sources of error in this data, as the total Facebook updates has been going up over time, and the site: operator is not a precise measure in Google. But the basic conclusion is not impacted by these errors – Google does not index the overwhelming majority of Facebook posted content.

Also of significance is that links from Facebook to external pages use the NoFollow attribute (these are defined in Chapter 6). This means that these links do not pass PageRank to the pages receiving the link. The code for a link in an example Facebook post is shown in Figure 8-9.

```
straighten some of it out. <a href="http://www.facebook.com/l.php?u=http%3A%2F%2Fabcn.ws%2F160Adhb&h=eAQHoTLng&enc=AZ0adFETozVrpq2mMdQNDiVIfg5bnA18ih22J9yCxweS1MiRbNYt2z6s0z4s1JI1veFbw9EIINA93qHMR_Hx4fWDkmDgUcTpXmy2MLCUxORmi7BoJ5GHVk5bUHCdwcu8EKP5UwamDJX7apbo2bkYR68XvkLh08rddM-Lq0CqXHNwwQ&s=1" target="_blank" rel="nofollow" onmouseover="LinkshimAsyncLink.swap(this, "http://abcn.ws/160Adhb");"  
onclick="LinkshimAsyncLink.referrer_log(this, "http://abcn.ws/160Adhb", "/si/ajax/1/render_linkshim_log/"?  
u=http\u00253A\u00252F\u00252Fabcn.ws\u00252F160Adhb&h=eAQHoTLng&render_verification=0&enc=AZ0adFETozVrpq2mMdQNDiVIfg5bnA18ih22J9yCxweS1MiRbNYt2z6s0z4s1JI1veFbw9EIINA93qHMR_Hx4fWDkmDgUcTpXmy2MLCUxORmi7BoJ5GHVk5bUHCdwcu8EKP5UwamDJX7apbo2bkYR68XvkLh08rddM-Lq0CqXHNwwQ");">http://abcn.ws/160Adhb</a></div></div>
```

Figure 8-9. Facebook Links are NoFollowed

One last consideration is whether or not Google might place more emphasis on indexing posts from highly influential people. Perhaps they place more authority on shares by people and brands with authoritative profiles. To test this Stone Temple Consulting performed a study (<http://www.stonetemple.com/does-facebook-activity-impact-seo/>).

This study looked at 40 Facebook posts done by 85 influential users over a 1 year period of time. Each of the profiles examined had at least 125,000 Likes for their page. The tested posts were spread over differing time intervals as follows:

- 10 most recent posts
- 10 posts at least 3 months old
- 10 posts at least 6 months old
- 10 posts at least 12 months old

Each of these posts was tested to see if Google had them in their index. Figure 8-10 Shows the results. In aggregate, about 59% of these posts were indexed.

Most Recent

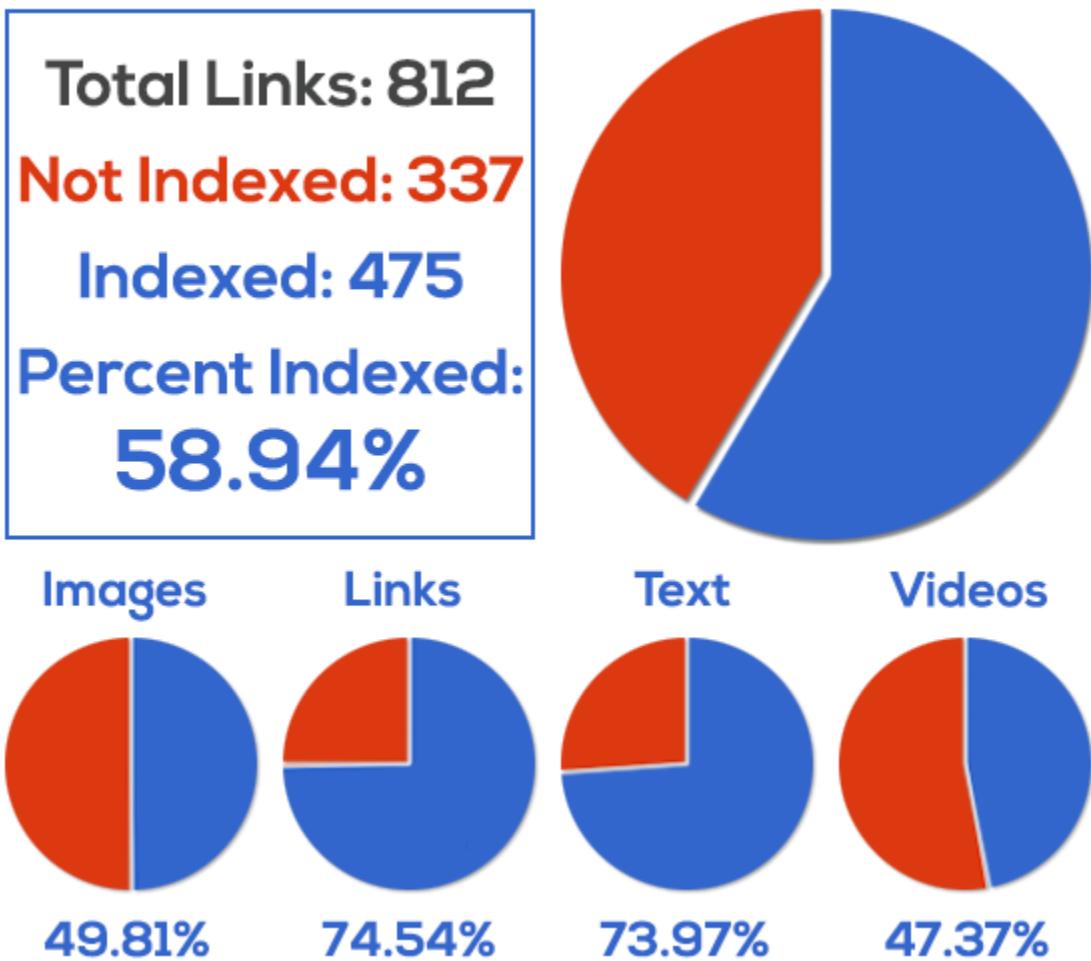


Figure 8-10. Facebook Pages Indexed by Google

It should be noted that about 74% of the text posts were indexed and 75% of the posts with links in them were indexed. If these were a major ranking factor it seems likely that Google would be indexing a larger percentage of these posts.

This combined with Bing's stance on using Facebook data in search suggests that they are not a major factor in driving the search results. This could possibly change in the future however, as search engines continuously experiment with new ways to use different types of signals as a ranking factor.

Does Google Use Twitter as a Ranking Signal?

An analysis of Twitter usage shows some similar results as we saw with Facebook. In Twitter's IPO filing, it was reported that Twitter is handling more than 500 million tweets per day on average. Figure 8-11 shows a combination of two Google search queries that attempt to find out how many Twitter status updates that Google has in their index.

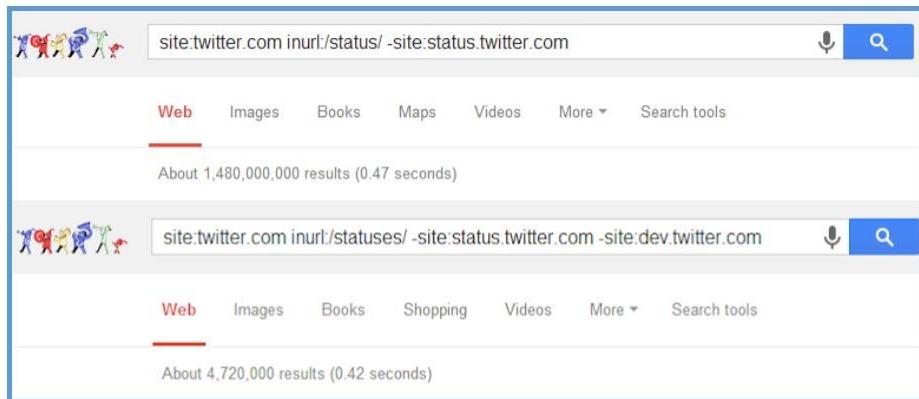


Figure 8-11. TwitterPages Indexed by Google

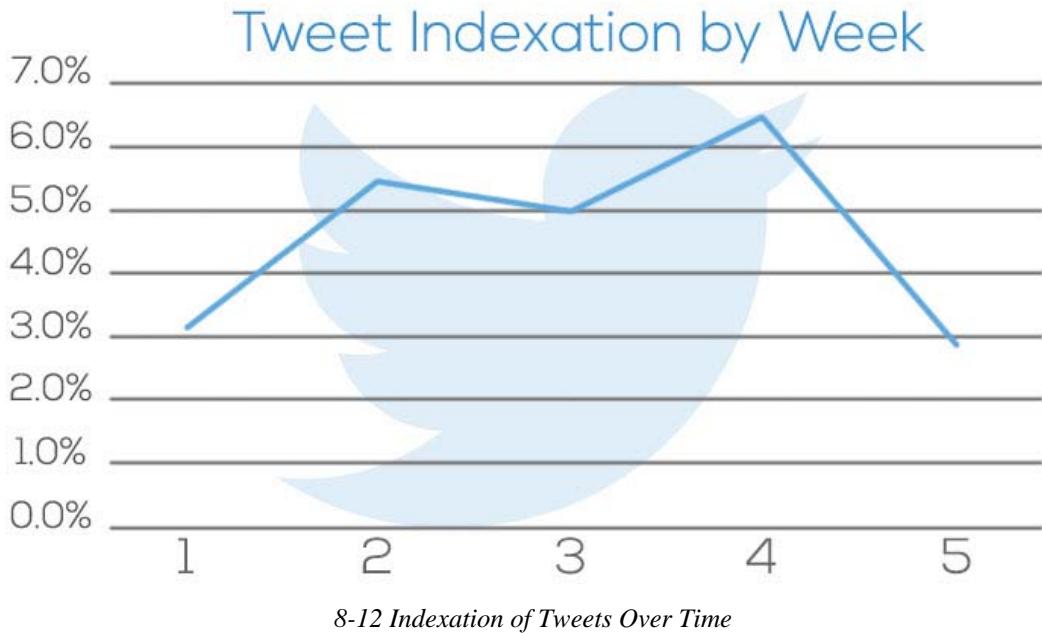
This shows a total of 1.27 billion pages, which is a pretty small number when you consider that there are 500 million tweets per day.

In another Stone Temple Consulting study by co-author Eric Enge (<http://www.stonetemple.com/how-does-google-index-tweets/>), the above search query was repeated, but the date range operator was used to see how many new pages were kept in the index by Google on a month by month basis from January 2012 through June 2014. That data is shown in Table 8-2.

Table 8-2. Twitter Indexation by Month

Month	Twitter	% of 500M	Month	Twitter	% of 500M
Jan-12	494,000	0.10%	Apr-13	1,570,000	0.31%
Feb-12	1,230,000	0.25%	May-13	924,000	0.18%
Mar-12	1,330,000	0.27%	Jun-13	1,540,000	0.31%
Apr-12	1,220,000	0.24%	Jul-13	1,420,000	0.28%
May-12	669,000	0.13%	Aug-13	1,320,000	0.26%
Jun-12	1,720,000	0.34%	Sep-13	777,000	0.16%
Jul-12	2,360,000	0.47%	Oct-13	1,480,000	0.30%
Aug-12	2,600,000	0.52%	Nov-13	1,530,000	0.31%
Sep-12	1,850,000	0.37%	Dec-13	1,620,000	0.32%
Oct-12	2,520,000	0.50%	Jan-14	1,030,000	0.21%
Nov-12	1,870,000	0.37%	Feb-14	1,920,000	0.38%
Dec-12	1,530,000	0.31%	Mar-14	2,270,000	0.45%
Jan-13	924,000	0.18%	Apr-14	3,270,000	0.65%
Feb-13	1,590,000	0.32%	May-14	2,750,000	0.55%
Mar-13	1,640,000	0.33%	Jun-14	9,090,000	1.82%

This data suggests that the indexing rate of tweets is actually quite low. However, accuracy of the number of results from **site:** queries is not always that precise. The study also included an analysis of the indexation of posts for 963 different Twitter accounts. 58 of these had at least 100,000 followers. These were hand-selected to be highly influential people. The study then tracked the indexing of their Tweets over time. The results are shown in Figure 8-12



Taken together, these two charts show us that overall indexing of content of tweets by Google is low, but they do index a high percentage of tweets from more influential accounts. Note that it is possible that this is because these tweets tend to receive links. This would increase Google's desire to index the content in the tweet.

Figure 8-13 shows the source code for a tweet containing a shared link. As with Facebook, this link is NoFollowed so no PageRank is passed by the link.

```
<a href="http://t.co/QrXWcFnluK" rel="nofollow" dir="ltr"
data-expanded-url="http://ow.ly/yB9Gj" class="twitter-
timeline-link" target="_blank" title="http://ow.ly/yB9Gj" >
<span class="tco-ellipsis"></span><span
class="invisible">http://</span><span class="js-display-
url">ow.ly/yB9Gj</span><span class="invisible"></span><span
class="tco-ellipsis"><span class="invisible">&nbsp;</span>
</span></a> by <a href="/stephanhov" class="twitter-atreply
pretty-link" dir="ltr" ><s>@</s><b>stephanhov</b></a>
```

Ch 8-13 Twitter Links use the NoFollow Attribute

In summary, the evidence suggests that Google does not currently use activity in Twitter as a ranking signal.

Does Google Use Google+ as a Ranking Signal?

Google+ is the social network owned by Google itself. Unlike Facebook and Twitter, Google obtains information from Google+ in real time. As soon as you complete a share in Google+, Google knows it has happened. This makes it seem far more likely that Google would use the data from this as a ranking factor in its search results. In fact there are a few clear ways that Google+ impacts the Google search results.

Personalization

Personalization plays a significant role in Google's search results, and Google+ can be a part of that. For example, if someone is following you on Google+, the chances that they will see your posts in Google's search results go way up. Figure 8-14 shows an example of this:

Ranks #12
When
Logged Out!

Link Builders From \$11/hr - Browse 100s of Link Builders
Ad www.odesk.com/Link-Building ▾
Post Your Job For Free Today.
oDesk has 24,199 followers on Google+
oDesk Guarantee - Rated & Reviewed Experts - Flexible & Cost Effective

Link Building Services - SeoBounty.com
Ad www.seobounty.com/ ▾
Professional SEO Company offering High Quality Link Building Services

Link Building Strategies – The Complete List - Point Blank ...
pointblankseo.com/link-building-strategies ▾
Brace yourself. This is the most comprehensive list of link building tactics (not strategies; there's a difference!) ever created. Think I'm lying?
Point Blank SEO - The Most Creative Link ... - Link Building Resources
Rick Elias, Wissam Dandan, David Cohen and 17 other people +1'd this

What is Link Building? Strategies & Examples - The ... - Moz
moz.com/beginners-guide-to-seo/growing-popularity-and-links ▾ Moz ▾
by Rand Fishkin - For search engines that crawl the web, links are the streets between pages. Using sophisticated link analysis, the engines can discover how pages are related to ...
SEO Tools and Search Engine ... - Links are Not Created Equal - Anchor Text - Fresh

Figure 8-14 Google+ Personalization in Action

The highlighted listing shown ranks #1 in the organic search results when the searcher is logged in, but ranks only in position 12 when the same user performs an incognito search (in Google's Chrome browser, this turns all personalization off). This is a big impact! The reason for the change in the personalized results is that 4 people whom this user follows on Google+ have +1'ed the content.

The impact of this on your business can be quite substantial. Consider the post shown in Figure 8-15.



Figure 8-15 Examples of Post with Many +1s

The great majority of these are likely to be from people interested in search, and for all of these users, this article will show higher in the search results. But it is also more likely to show up for everyone they follow! This means that there are two bites at this apple already:

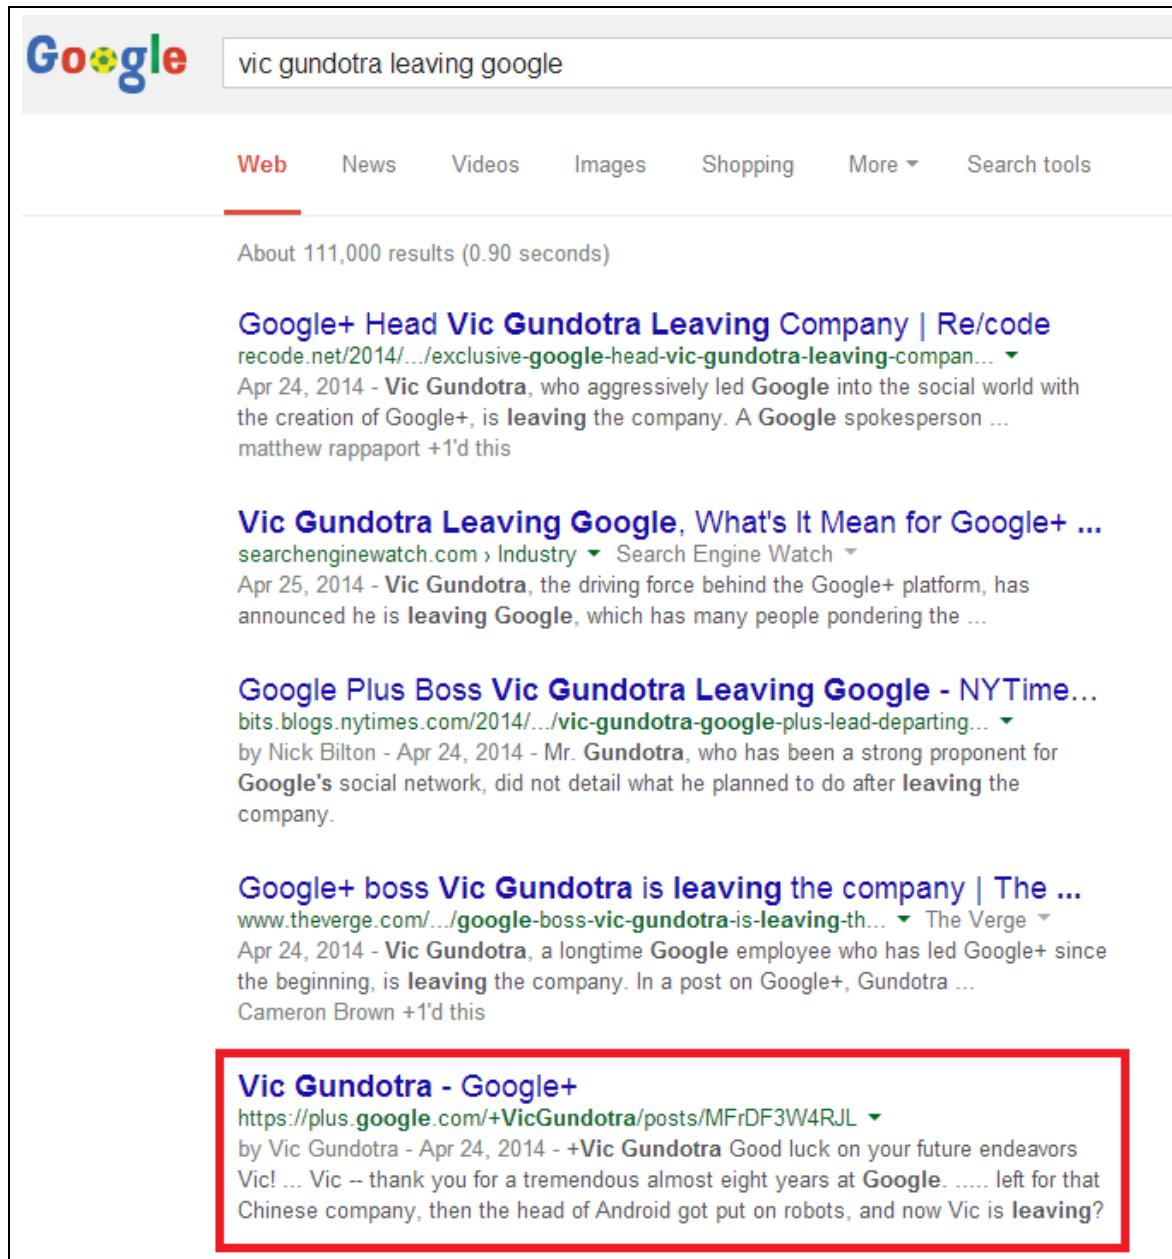
- People who follow the author on Google+ are more likely to see it
- People who follow anyone who +1's the content are more likely to see it

This is already a very big deal. If you can build a strong audience of your own, or if major influencers in your market space +1 or share the content, this can get you a lot of exposure to your target market.

In addition, it appears that people you exchange emails with can also create this type of connection that leads to the personalization of results (<http://moz.com/blog/using-google-plus-to-appear-in-the-top-results-every-time-whiteboard-friday>).

Google+ Posts in the SERPs

Google+ posts themselves can also rank in the search results. Figure 8-16 shows a post that ranks in the search results.



Google search results for "vic gundotra leaving google". The results include news articles from various sources and a Google+ post by Vic Gundotra, which is highlighted with a red box.

Google+ Head Vic Gundotra Leaving Company | Re/code
recode.net/2014/.../exclusive-google-head-vic-gundotra-leaving-compan... ▾
Apr 24, 2014 - Vic Gundotra, who aggressively led Google into the social world with the creation of Google+, is leaving the company. A Google spokesperson ...
matthew rappaport +1'd this

Vic Gundotra Leaving Google, What's It Mean for Google+ ...
searchenginewatch.com › Industry ▾ Search Engine Watch ▾
Apr 25, 2014 - Vic Gundotra, the driving force behind the Google+ platform, has announced he is leaving Google, which has many people pondering the ...

Google Plus Boss Vic Gundotra Leaving Google - NYTime...
bits.blogs.nytimes.com/2014/.../vic-gundotra-google-plus-lead-departing... ▾
by Nick Bilton - Apr 24, 2014 - Mr. Gundotra, who has been a strong proponent for Google's social network, did not detail what he planned to do after leaving the company.

Google+ boss Vic Gundotra is leaving the company | The ...
www.theverge.com/.../google-boss-vic-gundotra-is-leaving-th... ▾ The Verge ▾
Apr 24, 2014 - Vic Gundotra, a longtime Google employee who has led Google+ since the beginning, is leaving the company. In a post on Google+, Gundotra ...
Cameron Brown +1'd this

Vic Gundotra - Google+
<https://plus.google.com/+VicGundotra/posts/MFrDF3W4RJL> ▾
by Vic Gundotra - Apr 24, 2014 - +Vic Gundotra Good luck on your future endeavors Vic! ... Vic -- thank you for a tremendous almost eight years at Google. left for that Chinese company, then the head of Android got put on robots, and now Vic is leaving?

Figure 8-16 Google+ Post in Search Results

This came up even though the search was performed in incognito mode, so all personalization of the results is turned off.

Brand Pages in the Search Results

Another benefit of Google+ is that your company (or institution) can create a brand page. If you properly connect that with your web site using the rel=publisher tag (which is covered in more detail in Chapter 6 of this book), and you have some basic level of activity on the page, the brand page can show up in the search results as well.

This is particularly likely to happen on a search for your own brand name, as shown in the search for “Major League Soccer” shown in Figure 8-17.

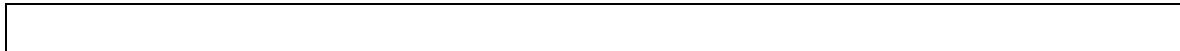


Figure 8-17 Brand Page Showing in Google’s Search Results

This is nice real estate to own, and can really add to the brand impression of your business. For MLS, we see the number of followers they have on Google+, and a “Follow” button at the top right. Google obtains additional information from many different sources, but is experimenting with allowing publishers to get their social profiles added to the panel shown on the right (<https://developers.google.com/webmasters/structured-data/customize/social-profiles>).

Google+ Impact of Non-Personalized Rankings of Content

Many have debated whether or not Google+ impacts the non-personalized results in Google as well. As Figure 8-18 suggests, Google+ posts used to pass PageRank, but this is no longer the case.

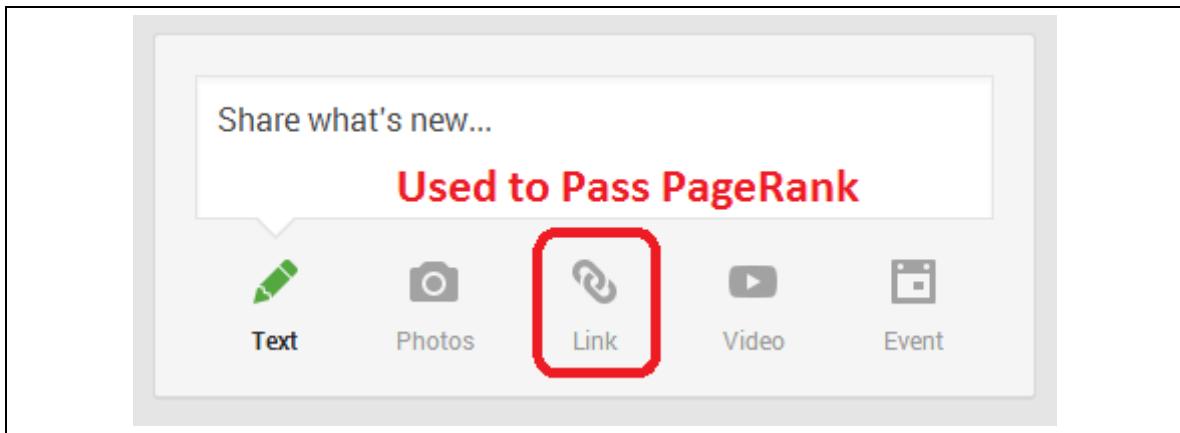


Figure 8-18 Links in Google+ Posts Used to Pass PageRank

You can examine the source code of a Google+ post in detail to verify this. Figure 8-19 shows an example of the source code for a particular post:

```

<b>Impact of Offline Activity on Social</b><br><br><span class="proflinkWrapper"><span
class="proflinkPrefix">+</span>

<a class="proflink aaTEdf" href="/107022061436866576067" oid="107022061436866576067">Mark
Traphagen</a></span> provides a thoughtful analysis on how offline activity can impact your social
following. I have been a believer in this concept for years.<br><br>For example, many is the time I
have been speaking at a conference where I have mentioned that attending a conference was a social
media activity. People would ask me what I meant by that, and I 'd ask for a show of hands to
see how many people had followed me as a result of my being there.<br><br>Lo and behold, a bunch of
hands would always go right up in the air!</div></div></div><div class="y_Nf"><div class="q9
yg">

<div class="s90c9"><div class="sp ej xI A8Hhid" style="width:506px; height:303px; max-height:303px;"><a
href="https://www.linkedin.com/today/post/article/20140813174834-13325642-it-s-all-social-now-how-non
-sm-events-affect-social-media" target="_blank" class="d-s ot-anchor Ks" tabindex="0"
rel="nofollow"><meta name="twitter:image:src"
content="//lh3.googleusercontent.com/proxy/hT0uB6gr8syC3Eyzy1JWT4ATM3NHVq_0_1ho_4pGtwmwDOL3XR0FnjDZH
YS3foHCokd9i2sW_CPNJHBScuLc8-GBCB1YFaCW7pEr_o=w506-h303-p"><div class="iGqbIb"></div></a></div>

<div class="VwVwbf"><div class="rCauNb"><a
href="https://www.linkedin.com/today/post/article/20140813174834-13325642-it-s-all-social-now-how-non
-sm-events-affect-social-media" target="_blank" class="d-s ot-anchor" title="It's All Social Now:
How Non-SM Events Affect Social Media | Mark Traphagen | LinkedIn" tabindex="0"
rel="nofollow">It's All Social Now: How Non-SM Events Affect Social Media | Mark Traphagen |
LinkedIn</a></div>

<div class="YPIIndd"><a
href="https://www.linkedin.com/today/post/article/20140813174834-13325642-it-s-all-social-now-how-non
-sm-events-affect-social-media" target="_blank" class="d-s ot-anchor" tabindex="0" rel="nofollow"><a
class="jp" href=".107350354213838732087" oid="107350354213838732087"><div class="tnIXBd"></div><div
class="iE5ljf">LinkedIn</div></a></a></div></div><div class="lr"></div><div
class="ko"></div><div class="kr"></div>

```

Figure 8-19 Google+ Links are NoFollowed

However, there are many that still believe that activity in Google+ can drive your search rankings.

Study on Google+ as a Ranking Factor

For that reason, Stone Temple Consulting's Eric Enge performed a study on this as well (<http://www.stonetemple.com/measuring-google-plus-impact-on-search-rankings/>). The study was designed to see if Google+ shares of links to a web page (not on Google+) would cause the rankings of those pages to improve.

This test took a number of web pages and published them on 3 different web sites. These pages all had these characteristics:

- The pages contained fresh, original content, relevant to the site on which they were published
- The pages had no links to them whatsoever. No external pages linked to them, and, no links were implemented from the pages from the domain on which they were published
- There was no Google code implemented on the pages. This included no Google Analytics code, no Google+ code, or no other Google code whatsoever.
- Participants in the test were also not allowed to use Google Chrome to visit the pages.

The next step was to conduct some Google+ shares to those web pages so that Google would discover the pages. These shares were done on July 19, 2013 and by July 29, 2013, all of the pages in the test were indexed. At that point, rankings for very long tail search terms were monitored until the end of the test.

Seven days after initial indexation, an additional set of shares were performed to the pages. The purpose of this latter set of shares was done to see if the rankings that were being monitored would change. You can see a partial data set for the rankings in Figure 8-20

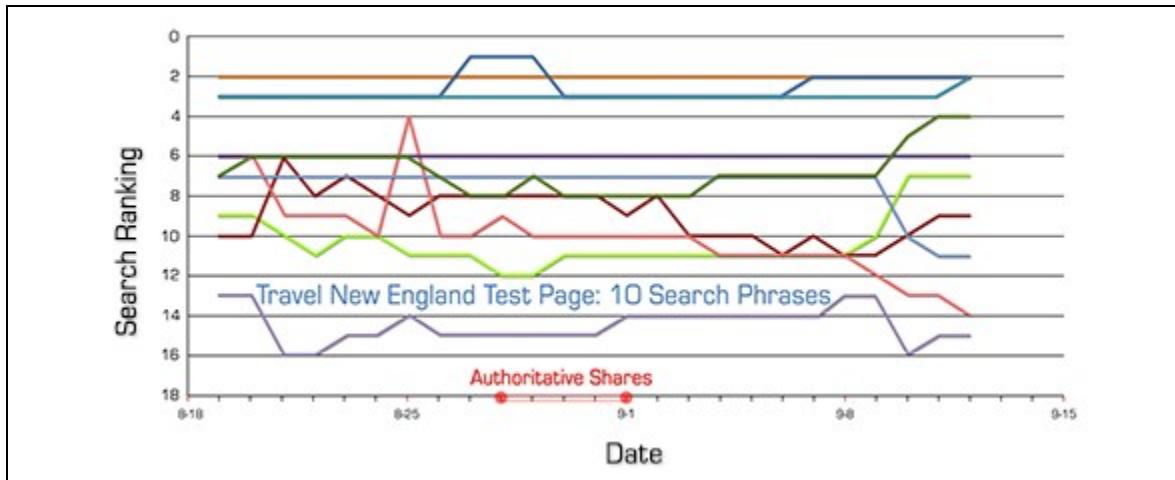


Figure 8-20 Google+ Ranking Study Test Results

In short, there was not real evidence of any movement in rankings seen as a result of the test. Based on this, there is no reason to believe that rankings are impacted by Google+ shares.

Last, but not least, is Google's recent Hummingbird release. While most people associate this solely with improvements in processing natural language queries, this is only one part of it, and in fact Hummingbird is a complete rewrite of Google's search engine platform. According to Danny Sullivan, part of Hummingbird was to allow Google to better process signals other than links, such as signals from social platforms (<https://plus.google.com/113690073826160821557/posts/X7MA8ogKvWu>).

While this new information makes it sound like Google is not using social signals yet, it suggests that it is still considering something that they will use in the future.

How Might Google+ as a Ranking Factor?

Google will work hard to develop algorithms that extract clear signals of value from Google+. How might they do that? Let me illustrate with an example. Let's consider 3 different people, but bear in mind that this example is pure speculation:

Person 1 is active on Google+ and he shares lots of content, his own, and content shared by others. In addition, he always adds thoughtful commentary in his shares and comments on a lot of other people's posts as well. The stuff he posts gets lots of +1s, reshares, and comments, and his comments on other people's post often draw a reaction too. Many of the people he interacts with frequently also get tons of interaction with their activity in Google+.

Person 2 is equally active on Google+ and she shares lots of content, her own, and content shared by others. However, she doesn't often add much content and she doesn't get the same level of interaction from others. The interaction she gets tend to be more perfunctory in nature.

Person 3 is not on Google+ at all. He is a total rockstar in his space, and everybody knows it, but he is not on G+ at all.

Google can't afford to punish Person 3 in the search results. That would be bad for Google's primary product, you know, the one that they make all the money on? What they can do though, is make comparative judgments between Person 1 and Person 2. In fact, Google may look for ways to recognize the relative authority of Person 3 using signals outside of Google+ and reward them accordingly.

Of course, we don't know exactly how Google is looking at these types of signals, but they will try to find ways to extract value from them that increase the quality of their search engine results, and that demands a balanced approach to the data on their part, and an incredible amount of testing and verification.

We know that the personalization impact is quite strong, and that Google+ posts themselves can rank in the SERPs. This is already a compelling reason to engage in Google+, especially if there is a real audience for you there. We also know that Google gets the information from Google+ directly. Unlike Facebook and Twitter, there is no crawling involved in this process. Google+ also is continuing to grow, and remains a part of Google's plans.

The Indirect Influence of Social Media Marketing

Social media can also influence search rankings in a number of indirect ways. For example, it can have a direct branding impact. Some examples of this are:

Build an Audience

A following can be built by using social media platforms such as Facebook, Twitter, Pinterest, and others to regularly communicate valuable information. Successful use of this tactic can help develop greater engagement with an organization's customers. This greater engagement can lead to creating brand ambassadors as a result of customers feeling a closer connection to the organization.

Customer Service

Social media sites can provide an effective incremental channel for customer service, enabling them to become aware of customer issues that they might not otherwise hear about. While this may not seem like an SEO concern, it can enhance the brand building efforts of the company. People who are able to see a proactive and effective customer service effort by a company learn to value their relationship with that company even more. In fact, many of today's consumers expect a rapid response time to any complaint they make via social media. Failure to do so can actually be harmful to the brand.

Viral Impact

Once an audience feels closely engaged with a company, their willingness to spread the word is greatly enhanced. Customers who feel a strong affiliation with a brand are more likely to help spread the word on new offerings or promotional programs to others in their network.

Link Building

Publishers with a strong following on a social media platform can release new content on their web sites and then announce it through their social media accounts. This can lead to links to the new content and is an effective way to build deep links to the site.

Citations

As discussed earlier in this Chapter, links in Facebook and Twitter are counted by search engines, but it is also likely (though not confirmed) that brand mentions are also of value. As search engines look for ways to measure how much they trust a brand, one metric they could look at is how many times that brand is referenced across the web.

Generate Brand Searches

People who learn about a brand and first become engaged with it through a social media presence may later choose to generate a search query on that brand. As with citations, a volume of actual brand searches is a strong signal of prominence.

Summary of Social Sources to Consider

There are a large number of social services available, and you can waste a lot of time if you chase them all. Here is a brief commentary on the major ones that can help you with the promotion of your web site:

Facebook

Facebook is the largest of the social networks with worldwide usage over 1.2 billion users (<http://www.theguardian.com/news/datablog/2014/feb/04/facebook-in-numbers-statistics>). This represents more than 10% of the total world population. Facebook has the largest audience and is certainly worth including as a key part of your social media strategy. In the US, more than 7 out of 10 US Internet users are using Facebook. In addition, Facebook's advertising platform provides highly granular targeting as discussed more in the Paid Social Media section in chapter 7.

Twitter

As of June 2014 Twitter reported 255 million active users (<https://about.twitter.com/company>). The user based may be smaller than that of Facebook, but it certainly includes a large number of influencers that can help create visibility for your organization. Twitter remains skewed towards male users, with nearly 50% more men on the platform than women.

LinkedIn

As of April 2014 LinkedIn had more than 300 million members worldwide (<http://mashable.com/2014/04/18/linkedin-300-million-users/>). This likewise makes it a powerful social network. It can be used to build a strong network of connections and paid versions of the service allow its InMail functionality to be used to send unsolicited emails (in very limited volume) to members, which, if done judiciously, can be used to initiate new relationships with influencers of interest.

Google+

While Google+ reports that it has 540 million active users external research shows that about half of these do not actively use the social network (http://www.nytimes.com/2014/02/15/technology/the-plus-in-google-plus-its-mostly-for-google.html?_r=0). However, 270M or more active users is still significant.

Instagram

Instagram reportedly has over 200 million users, and is growing fast. Piper Jaffray reports that Instagram is the most important social network among teens (<http://www.piperjaffray.com/private/pdf/TSWT%20Infographics.pdf>). Instagram is designed from the ground up as a mobile experience.

Pinterest

Pinterest may be quite a bit smaller, but it is a growing network with 70 million reported users (<http://socialmediaweek.org/blog/2014/08/market-brand-pinterests-highest-user-base/>). The network skews towards a female demographic, so if that is your target market, it could be a great platform for you to focus on. It's considered a highly "immersive" media in terms of time spent in the app and numbers of boards built per person. The audience here is still US-centric.

YouTube

YouTube can be thought of as a search engine, but is also a social network. People love to share videos and this is happening in volume. YouTube also has a very large audience, with 159.1 million

unique viewers per month as of December 2013 (<http://searchenginewatch.com/article/2322389/Online-Content-Videos-Break-Record-with-52-Billion-Monthly-Views>).

WhatsApp

WhatsApp is a cross platform mobile-messaging app that positions itself as a cost-free alternative to SMS communications. WhatsApp is the number one mobile messaging app based on a survey 3,759 users conducted across US, China, Brazil, South Africa and Indonesia.

There are many other social sites worthy of consideration. These include:

- Specialized sites – Vine and Snapchat
- Q&A Sites - Yahoo! Answers, StackExchange, Quora, Formspring, ChaCha.
- Document Sharing Sites - SlideShare, DocStoc, Scribd
- Local Business Sites - Google Local Business Center, Yelp, UrbanSpoon, YellowPages, CitySearch, TripAdvisor
- Social News Sites: reddit, Techmeme
- Bookmarking Sites - StumbleUpon, Delicious, Trunk.ly.
- Vertically Focused Social Sites - as may apply to specific vertical markets (for example, there is a knitting social site at <http://www.knitting.com>, so it is possible that there is a social site specific to your vertical market).

Monitoring, Measuring, and Improving Social Media Marketing

While search engine optimization and social media marketing are two distinct practices, the influence of social media behavior on search results is, and will continue to be, important to consider. This next section will explore some basic aspects of how to optimize your efforts in the social realm to help your SEO efforts.

Best Practices and Recommendations for Social Media Marketing

Social media networks all operate as communities. In reality, each network is a community of communities. No individual interacts with an entire network. Each user finds his or her “tribe” of followers and following. Even within those groups, there will be tighter “inner circles” of people and/or brands that get the most interaction and attention from a user.

The structure, rules, and methods of interaction may vary, but a social media network’s primary focus is to create and foster relationships. For example, you may think of YouTube as a video site, but a better label for it would be a video sharing site. Social network sharing is done with the ultimate goal of attracting a following and engagement. People aren’t on social networks to broadcast into a void; they want to feel that they are being heard and that their content has meaning for others.

Effective use of social media results from fostering relationships that build your brand’s reach, authority, and trust. Try to avoid seeing it as a broadcast, advertising, or even direct-selling medium. Even with brands, social media users want authentic engagement. They will resist being sold to. But handled properly, social media provides an excellent opportunity to build positive associations, so that when prospects are ready to buy, your brand is top-of-mind.

The key to success in any community is to be a contributor. You must put something into it to get anything back out of it. This requires authentic participation, as little value accrues to those who participate in an artificial way.

Find ways to contribute content, information, and opinions to discussions. Using Twitter as an example, one approach is to start your day by reviewing the daily news related to your market space, pick out the content you consider the most valuable, and share those items with your following. Another approach is to offer valuable tips on a regular basis. A third approach is to be known for controversial opinions, but be aware of the impact this may have on your brand. Bring a unique viewpoint.

However, if you only share content, whether your own or curated from other sources, you will always be limited in how far your influence will extend. Remember that people are on social media, at least in part, to be affirmed, to find that they matter to others. For that reason, you should dedicate some time each day to engaging with the posts of others. Offer constructive, helpful comments and replies. Don't just say "great post!" but explain specifically what you liked about it, or contribute information or insight on the topic. You will find doing this not only endears you to the person whose content is the recipient of your comments, others seeing the conversation may follow you because of your helpfulness and knowledge.

Overall, focus on building your credibility and authority. These are key elements to success in a social environment. Once people accept you as a legitimate participant the doors will start to open for you.

Even if you are a novice (aka "noob") in your market area, you still have something to offer. The experts can learn by watching how you learn as you gain experience. Sometimes just posting a really good question can be as valuable as providing a great answer. Such engagement brings a surprising amount of value to the community, and if you can also bring some personality into the process people will try hard to help you, and will want to engage with you further.

As a publisher there will also likely be times when you want to put something out there that is self-promotional. This is fine as long as you don't do it with all your communications, and make certain that you do it in an engaging way. You earn the right to do that by being seen as a valued contributor to the community.

Claim Key Profiles

Make sure to claim your brand/profile at key social sites. Services like Trackur (<http://www.trackur.com>), and Radian6 (<http://www.radian6.com>) allow you to monitor the use of your brand across a wide array of social sites. You will want to set up a brand presence on each social media platform even if you don't intend to use a network yet, in part to prevent your name from being taken by impostors or "squatters" (those who claim brand name profiles in hopes of later selling them to the brands). Some of the sites where you should create a brand profile for your organization include:

- Facebook
- Twitter
- Google+
- LinkedIn
- YouTube
- Pinterest
- CrunchBase
- Quora
- About.me
- Instagram
- Tumblr
- SlideShare
- reddit
- StumbleUpon
- Delicious
- StackExchange
- Google Places for Business

- Yelp
- UrbanSpoon
- YellowPages
- CitySearch
- TripAdvisor

Use NameChk.com to rapidly see if your name is available across a wide range of social sites. Note that this does not mean that you should plan on building an active presence on all of these sites – for most companies that's unrealistic. However, claiming the profiles will prevent other people from implementing profiles with your brand name on any of these platforms.

Deciding on a New Social Network

When you first consider engaging in a new social network, it is hard to be sure what ROI you will get for your efforts, and what you will need to commit to get that ROI. It's a good idea to start your efforts with an experimentation cycle designed to help you determine what you will be able to get out of it.

A word about ROI in social media: Remember that social media is primarily for reach, branding, and trust/authority building. Therefore the best measurement for return on social media efforts probably isn't sales, at least, not directly. Measure your effectiveness in social media by metrics like followers, engaged followers (actually more important than raw follower numbers), engagement per post, increased positive brand mentions, and increased traffic to your site.

Here are some suggestions on how to set up experiments with various social networks:

1. Hypothesize about a site's potential value. Are the demographics of the audience a fit for what you want to accomplish? Is the medium one that you effectively utilize? For example, you would not want to consider YouTube if creating/sharing videos were not something that made sense for your organization. Take the time to consider these and other questions that affect the potential value to you before getting started.
2. Research what networks are being used by your competition. Lack of presence of competitors on a particular network is not necessarily a red flag. There may be an opportunity for you to take leadership in your niche on that network, if you can develop strategies to either find or develop an audience there.
3. Setup social media and analytic tracking systems to measure. When you first start out your returns will be much lower than they will be in the long term as your credibility in communities builds, but even early results can help guide you on what is most effective. We will discuss this more in the next section "Tracking Social Media in Your Web Analytics".
4. Create an account and invest a few hours in participating. If the network has search, try different keywords to find people, groups, and communities discussing topics relevant to you, and start engaging those posts and creating your own. This will help you test your hypothesis about the value. A few hours will not be enough time to start seeing any ROI, but it should help you better qualify whether or not the community has the potential to work for your organization.
5. You won't really know the effectiveness of a network until you start getting an engaged following there. The best way to gain real followers is to seek out people and communities within the network that are relevant to your vertical, follow them and begin engaging with them. People are more likely to follow you back and/or recommend you to others when you do this.
6. If this initial test goes well and you still believe that the social network may be a good investment, go a little further with it until you see some initial results. At that point step back and consider the return on the effort so far. Be aware that your ROI should continue to increase as your credibility and authority builds, but this will give you an initial measurement of what may be possible.
7. You can also reach out to some of the more authoritative people using the network to get a sense for the types of returns they are seeing. While this information will probably be qualitative rather than quantitative in nature, it can help you learn what types of ROI might be possible.

8. Compare the potential ROI, including the potential long term value of the connections you can develop on the network, with other ways you could invest that same time and energy. Keep at it if it looks like a high value channel, and drop it if it doesn't and try something else!

Tracking Social Media

In the world of social media, there are several key types of metrics we're interested in tracking:

- Traffic data - how many visits and visitors did social drive to our sites?
- Fan/follower data - how many people are in our various networks and how are they growing?
- Social interaction data - how are people interacting with, sharing and re-sharing our content on social networks?
- Social content performance - how is the content we're producing on social sites performing?

Getting the right metrics to answer these questions requires segmenting by network. Not every question will have direct answers in the data, so you may need to make assumptions or inferences. Buzzstream (<http://www.buzzstream.com/>) and Klipfolio (<http://www.klipfolio.com/>) are examples of popular tools for tracking social media performance.

Facebook

Facebook offers a relative wealth of data through their built-in product for brand pages, Facebook Insights (<http://www.facebook.com/insights/>). For example, you can see Likes, Post Reach, and Engagement (as shown in Figure 8-21) or the demographics of your Fans (as shown in Figure 8-22).

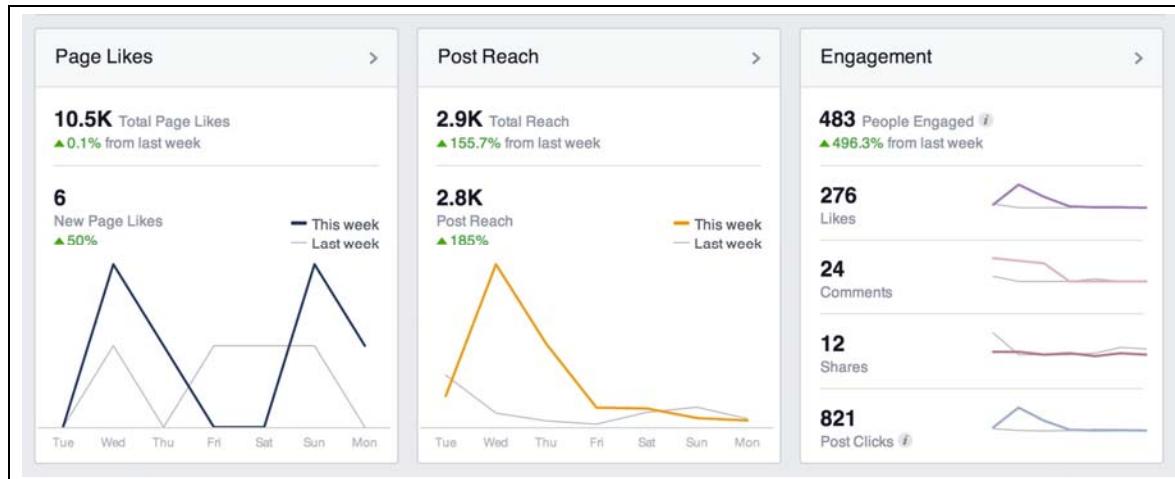


Figure 8-21. Facebook Insights Dashboard



Figure 8-22. Facebook Insights Demographics

Here you can track key metrics over time, including the size of your fanbase, the reach and effectiveness of your content, the quantity of likes and shares of your content, demographics of your fans and more.

Insights also has a very unique and powerful feature - integration on your website. Using a small bit of Javascript code, you can embed the Facebook Insights functionality on your site and receive information about all the users visiting your pages that are logged into Facebook.

More information on Insights can be found at

- Official Insights Page on Facebook (<http://developers.facebook.com/docs/insights/>)
- How to Use Facebook Insights (<http://blog.antavo.com/en/how-to-use-facebook-insights/>)

Twitter

Twitter and Facebook are likely to be the largest two social networks for referring traffic to most sites (although a Shareaholic study shows LinkedIn and Google+ drive more engaged visitors <http://thenextweb.com/socialmedia/2014/03/31/google-linkedin-drive-engaged-social-referrals-compared-twitter-facebook-pinterest>).

Twitter now provides limited analytics (<http://analytics.twitter.com>). For more detail, you should look into third-party social analytics tools (see “Tools for Measuring Social Media Metrics” later in this chapter).

The key metrics on Twitter are:

- Followers (and follower growth over time) - the unique number of Twitter users who've "followed" your account, as shown in Figure 8-23.
- Active Followers - the number of followers who've logged into or used a feature of Twitter in the past 30 days (those that have not are likely inactive or non-human accounts). This is challenging to get, and requires software that runs through your followers and determines which are actively using via the API. Some third party tools discussed below will show this information.

- @ Replies - the number of tweets sent that begin with your account name. An example is shown in Figure 8-24. These are only seen by the mentioned account and anyone who follows both sender and receiver.
- @ Mentions - Figure 8-25 shows a tweet that includes Rand Fishkin's account name, but uses it inside the tweet, rather than at the beginning, meaning others on Twitter can see the tweet by default, rather than only those who follow both accounts.
- Brand Mentions - tweets that contain the brand/account name but don't use the @ symbol. An example can be seen in Figure 8-26.
- Domain/URL Mentions - Tweets that include a link that contains your brand name/domain name. An example of this using Moz is shown in Figure 8-27. These now include, by default, any shortened URL that contains the brand/domain name as Twitter is automatically parsing the final destination URL for matches to the query.
- Direct Retweets - the quantity of retweets (using Twitter's native retweet button/functionality) appeared on the service, as shown in Figure 8-28.
- RTs & Vias - the quantity of tweets that contained an RT or via of your account (see the example in Figure 8-29). These are similar to direct Retweets, but aren't necessarily counted by Twitter's automatic RT system because they contain a modification of the original message and appear to come from a unique source. Some users indicate a modified retweet with MT.
- Best Performing Content - the content you shared on Twitter that earned the most clicks, Retweets and Shares. This is now available in Twitter's own analytics, but some third-party tools give deeper analysis.
- Direct Traffic + Non-Twitter.com Drivers - sources that sent traffic to your site via Twitter's ecosystem, even if they come from desktop clients or other third-party software sources. See Figure 8-30 for an example. Thanks to a recent change made by Twitter (<http://www.distilled.net/blog/social-media/twitters-t-co-link-shortening-service-is-game-changing-heres-why/>), these will now show up (mostly) as coming from T.co (Twitter's shortener), although some third-party link shortening services, such as bit.ly, have their own analytics for clicks on links you create through them, even if Twitter publicly changes them to t.co links.

TWEETS	FOLLOWING	FOLLOWERS
10.5K	361	9,512

• *Figure 8-23.Twitter Followers*



Sunita Biddu
@chillingbreeze

[Follow](#)

@stonetemple - Interesting read Eric, I'll certainly experiment on this by creating distinct page's & writing title tag's as you suggested!

• *Figure 8-24. Tracking Twitter @Replies*



Neal Schaffer @NealSchaffer · Jul 2

How to Choose the Right Images for your Content goo.gl/pFYsx9 via @SarahArrow

3

• *Figure 8-25 Tracking Twitter @Mentions*

 Stephan Spencer @sspencer · May 22
Ralph Ellison is a winning personality that wrote Invisible Man with enough energy, it reconstructed pinterest.com/pin/2042101643...

View summary

- *Figure 8-26 ch8-brand-mentions.png. Brand Mentions in Twitter*

 Retweeted by Matt Cutts
 Paul Buchheit @paultoo · Jun 25
Gmail now has an API! Looks awesome, can't wait to try - developers.google.com/gmail/api/

View summary

- *Figure 8-27 Tracking Twitter Domain/URL Mentions*

 @randfish
Rand Fishkin

Google will be testing re-ranking results based on [+1s](http://wired.com/epicenter/2011...) wired.com/epicenter/2011...
Goal is to prevent spam...

30 Aug via web    Retweeted by EddyBauer and 31 others



- *Figure 8-28. Tracking Twitter Retweets*



EJ Dionne

@EJDionne



Follow

Celebrate! MT @BeschlossDC The signing of Civil Rights Act 50 years ago today: LBJ with Martin Luther King, Jr. #NARA

Reply Retweet Favorite More



RETWEETS

123

FAVORITES

48



- *Figure 8-29 ch8-modified-tweets.png. Modified RTs (MTs) in Twitter*

Source/Medium	None	Visits
1. t.co / referral	23,990	±3%
2. facebook.com / referral	19,847	±3%
3. linkedin.com / referral	13,101	±3%
4. twitter.com / referral	9,520	±4%
5. stumbleupon.com / referral	4,761	±5%
6. twitterfeed / twitter	3,115	±6%
7. hootsuite.com / referral	2,369	±8%

• *Figure 8-30. Tracking Twitter Traffic*

It's often mentioned that in analytics, nothing is worth tracking unless it can be used to take action and improve. For the metrics above, the primary action you're tracking are ones that relate to you. The key to taking better actions is comparing successful interactions, tweets and content against less successful ones to determine what has the best impact on growing your audience, bringing visits to your site and, eventually, driving conversion actions.

LinkedIn

LinkedIn functions like a hybrid of Twitter and Facebook. Connections require acceptance from both sides, but public entities (like company pages) and groups can be followed.

LinkedIn used to be known primarily as a sort of online address book of your business connections, and a place for job hunters to post their resumes. Recently though, LinkedIn has made major strides to reinvent itself as a true social network. Users now get a Facebook-like news feed of content and conversations shared by people and companies they are linked to.

In early 2014, LinkedIn opened up its Publishing platform, formerly reserved to handpicked “influencers,” to all users. Publishing allows users to create blog-like content that gets pushed out to their connections and followers. Other users can now follow the Publisher content of anyone else without having to first gain a reciprocated connection with the author. These changes have made LinkedIn a much more significant player in the social media influence-building arena, especially for B2B concerns.

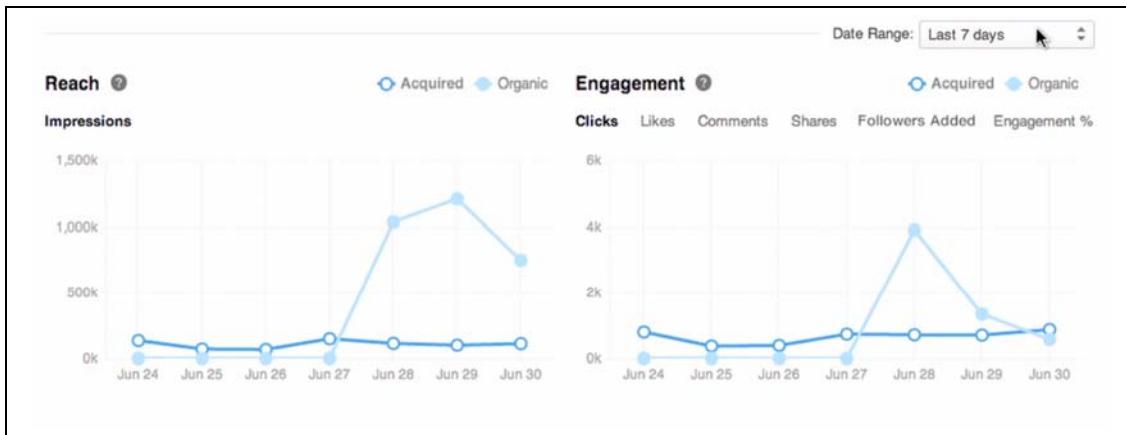
LinkedIn Groups can also be used to develop your visibility and build new connections. You can use LinkedIn Search to find relevant groups that show active participation. Review the Group rules to make sure that some are at least in place, and being enforced. It's also good to find those that appear to have active discussions taking place.

You can also reach your target audience on LinkedIn using Direct Sponsored Content (<https://business.linkedin.com/marketing-solutions/content-marketing/sponsored-updates>), which can be a powerful supplement to your content marketing efforts.

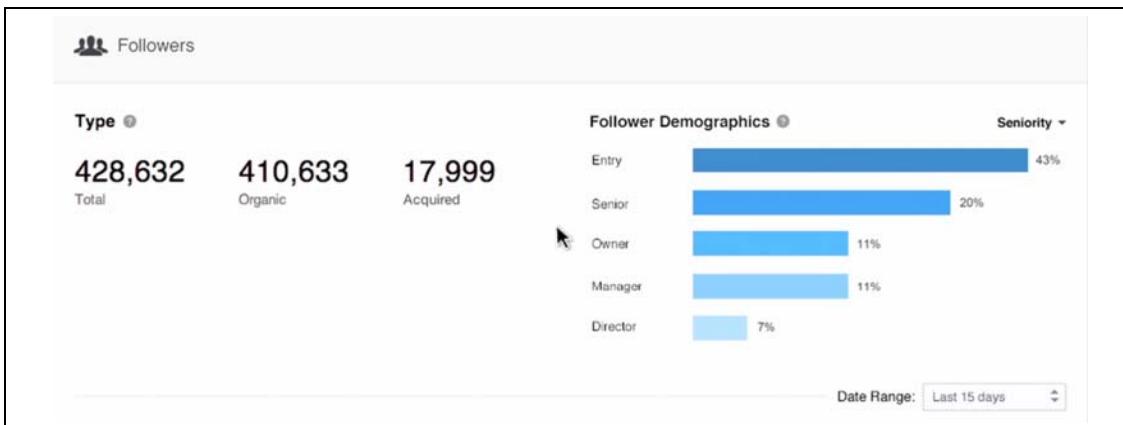
Like Facebook, LinkedIn has some built-in analytics for businesses as well as individual profiles, and lots of data points that are useful to track, including:

- Company Page Reach and Engagement - these track the number of times your company's LinkedIn profile and its shared content have been viewed over time and the quantity and types of engagements earned. An example of this is shown in Figure 8-31.

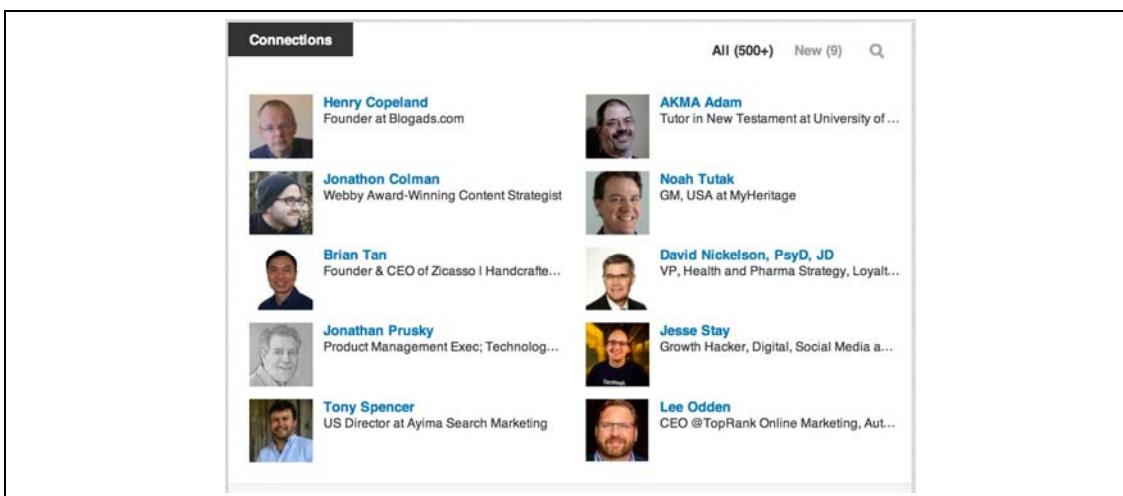
- Followers - as with Twitter, individuals can "follow" a brand account on LinkedIn and receive status updates in their "updates" stream. Figure 8-32 shows stats for Followers of a LinkedIn account. The more followers, the greater your ability to reach more people on LinkedIn with the content you share there
- Connections - The quantity of unique connections for an individual profile on LinkedIn is a worthwhile metric to track, but unfortunately, all that is available as of July 2014 is the raw, current count (as shown in Figure 8-33) and who the connections are.
- Messages, Invitations and Notifications - Figure 8-34 shows the status bar at the top of a LinkedIn page. The number on the envelope shows unread messages and un-reviewed connection invitations, while the number on the flag represents unviewed news stream updates and other updates, such as people viewing your profile or commenting on one of your updates.
- Profile Views - how many people have looked at your profile over time and some data about who they are (only paid "Pro" LinkedIn accounts can see the identities of all profile viewers). Figure 8-35 shows an example.
- Top Keywords & Referral Sources - a list of the top keywords users on LinkedIn searched for prior to discovering your profile, and where they came from to get to your profile, as shown in Figure 8-36. (Again, only a limited amount of data is shown unless you have a paid account.)
- Content Shares - As of July 2014, there is no way to measure or record the quantity of regular status updates/shares you've sent out over LinkedIn nor the number of "likes" you've received on the service.
- LinkedIn Publishing -LinkedIn Publishing allows users to create original content with images, links, and formatting, which is then pushed out to their connections, followers, and possibly beyond. Other users can follow the Publishing content of a user without having to make a two-way connection. LinkedIn Publishers get some rudimentary metrics, such as number of views, likes, and comments for each post, as well as information on followers.
- Traffic - LinkedIn isn't a huge traffic driver for most sites, but for certain B2B sites, it can be relatively substantial and the quality is often higher than other social sources. Figure 8-37 shows a screenshot from a Google Analytics account.



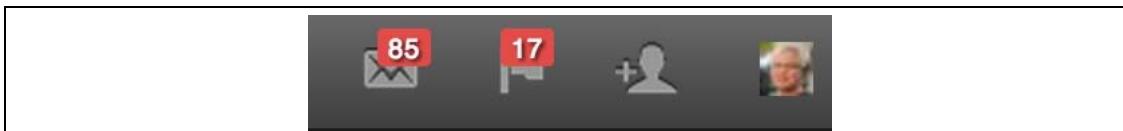
• *Figure 8-31.LinkedIn Reach and Engagement*



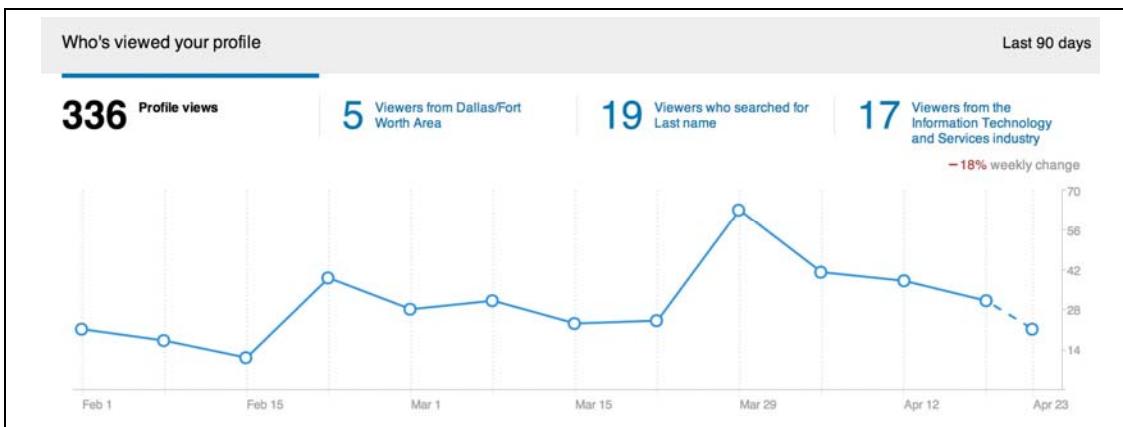
• Figure 8-32. Followers in LinkedIn



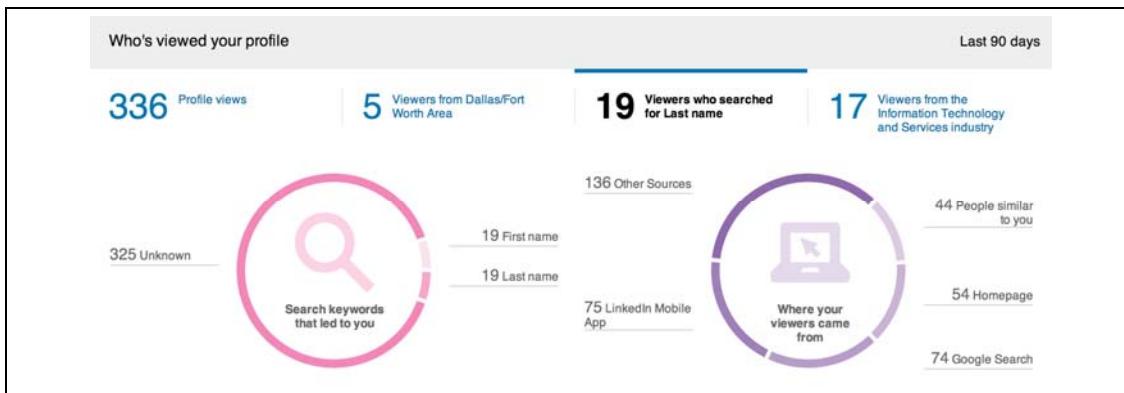
• Figure 8-33. LinkedIn Connections



• Figure 8-34. LinkedIn Messages, Invitations & Updates



• Figure 8-35. Profile Views on LinkedIn



• Figure 8-36. Keywords & Referral Sources on LinkedIn



• Figure 8-37. LinkedIn Traffic

Few third-party tools exist to help with measurement of LinkedIn, but over time, some of the tools for Facebook and Twitter may expand to LinkedIn.

Google+

Google's social network is still relatively young, but at least three factors make it worthy of the attention of marketers:

1. The network's rapid growth, especially in the global market. Some estimates now make it the second largest social network, surpassing Twitter and second only to Facebook ([http://marketingland.com/report/google-now-2nd-biggest-social-network-worldwide-31908](http://marketingland.com/report-google-now-2nd-biggest-social-network-worldwide-31908))
2. The integration of Google+ into nearly all of Google's products and services. Google Vice President for Social Vic Gundotra has called it "the next version of Google" and "the social layer of Google" (<http://bits.blogs.nytimes.com/2012/03/06/google-defending-google-plus-shares-usage-numbers/>)
3. The influence of Google+ on Google search, particularly in the area of personalized search (<http://searchengineland.com/how-google-plus-profiles-pages-gain-search-authority-176828> and <http://searchenginewatch.com/article/2307518/Google-SEO-How-Google-Impacts-Search-Results>)

Google has promised analytics for pages and profiles, but as of this writing those still aren't available. But several third party tools have gone a long way in delivering needed metrics. Among these are CircleCount (<http://www.circlecount.com/>), SteadyDemand (<http://www.steadydemand.com/>), and NOD3X (<http://nod3x.com/>). The latter is a sophisticated tool for in-depth analysis of follower engagement, reach, and influence.

Significant metrics

- Number of Followers: On Google+ “followers” are people or pages who have you or your page in at least one of their circles. You can see this on a profile or page as shown in Figure 8-38 (if the profile or page owner has made followers visible in settings).
- Views: View counts are new as of April 2014, and are displayed on the cover image of profiles and pages that have not opted out of their display in settings. The view count is an approximate total of all profile views, post impressions, and image impressions since October 2012. For more on what comprises views and what the metric might mean, see <http://www.stonetemple.com/new-google-plus-views-count-important-metric-or-vanity-of-vanities/>
- +Mentions: On Google+ one user links to another by typing a “+” and then the user’s name. This creates a link to the mentioned user’s profile, and also sends that user a notification that he or she has been mentioned. As of April 2014 Google+ does not provide a count of these, but some third-party tools do. You can view all your recent mentions by clicking the “Google+ Notifications” link at the top of the notifications dropdown that appears when click the notifications bell at the upper right of any Google+ page, and then select “Mentions of you.” Figure 8-39 shows an example.
- Brand Mentions: As shown in Figure 8-40, you can find recent mentions of your brand through the Google+ search, using the Everything filter tab and changing the view from “Best of” to “Most recent.”
- Content Shares, Content +1’s, Link Shares, Link +1’s: Recent shares of links to your own site on Google+ can be viewed in Google Analytics under Acquisition > Social > Data Hub Activity, with Google+ as the selected network. Clicking the “More” button to the right of any Google+ share allows you to filter for all shares of that link, and to view the Google+ Ripples for that link (a graphic representation of all the public shares of the link on Google+). You can also view the number of posts, reshares, +1’s and comments for any link to your site that was shared on Google+ by navigating in Google Analytics to Acquisition > Social > Landing Pages, then clicking the desired link. Next select “Google+” and then change the Primary Dimension to “Social Network and Action.” See Figure 8-41 for an example.
- Traffic - Google+ can become a major site traffic driver for brands that develop an active following there. Figure 8-42 is a screenshot of traffic from Google+ as shown in Google Analytics.



• *Figure 8-38. Google+ Followers and Views*



Ronnie Bincer
Shared publicly - Jul 2, 2012 #InfluencerMarketing

Building Influence by Teaching Others - #TNTBootcamp
series
An interview with **+Chef Dennis Littley** Chef, Food
Blogger, HOA Host - "I began blogging in 2009 as a
Read more (120 lines)

- Figure 8-39. Name Mentions on Google+



Mike Allton
Shared publicly - Jun 4, 2014 #Blogging

5 Reasons Canva Will ROCK Your Blog

What do Chuck Norris and **+Canva** have in common?

Not much, you might say. And you'd be right, if it weren't for this quote:

"I think setting a goal, getting a visual image of what it is you want. You've got to see what it is you want to achieve before you can pursue it." ~ Chuck Norris

+Peg Fitzpatrick and **+Guy Kawasaki** and **+martin shervington** and **+Rebekah Radice** and others have been talking *a lot* about Canva lately. And for good reason. It's an amazing tool, and I'm thrilled to be able to share with you today all the reasons why I'm using it, and *how* I'm using it, to rock my blog, so that you can too.

- Figure 8-40. Brand Mentions on Google+

Primary Dimension: [Full Referrer](#) [Social Network and Action](#)

Secondary dimension ▾

Social Network and Action		Data Hub Activities	
1. █ Google+ : +1		35	54.69%
2. █ Google+ : comment		12	18.75%
3. █ Google+ : reshare		12	18.75%
4. █ Google+ : post		5	7.81%

- Figure 8-41. Google+ site link activity in Google Analytics

Social Network	Visits
1. Google+	4,260 (49.94%)
2. Twitter	2,345 (27.49%)
3. Facebook	679 (7.96%)
4. reddit	500 (5.86%)
5. LinkedIn	212 (2.49%)

- *Figure 8-42. Traffic from social networks in Google Analytics*

Perhaps due to privacy issues, Google+ uses a single referring URL for all traffic, helping to consolidate the data in analytics reports, but making it frustrating to determine which shares/users/links sent what quantities or value of visits.

Instagram, StumbleUpon, Quora, Yelp, Flickr, and YouTube

Depending on the quantity and value of the traffic that other social networks send, there may indeed be additional metrics worth tracking. Sites worthy of consideration are Instagram, Pinterest, StumbleUpon, Slideshare, reddit and Quora. These are likely worth some investment on the metrics and effort front, and if small quantities of contribution/participation yield large returns, more investment is likely warranted.

Social News Sites

The two most well-known social news sites are reddit (<http://reddit.com>) and Techmeme (<http://www.techmeme.com>).

reddit is a social link aggregation and discussion site, organized into topical communities that are self-organized and independently moderated. There are millions of visitors each day, and reddit has hosted “ask me anything” interviews from many celebrities and a few politicians, including Barack Obama.

Articles posted on reddit can potentially receive a lot of search traffic, and as a result it can be tempting to try to influence this process to help achieve your goals. However, this is a bad idea, because people on reddit are very defensive of their platform, and any attempts at manipulation or any commercial behavior.

Even attempting to reach out to influencers on this platform will quickly draw a negative reaction, and can result in your being labeled as a spammer. Any attempt at solicitation runs a significant risk of being downvoted, attacked and/or blacklisted.

Still there, is a lot to be gained from reddit, but before you embark on such a campaign, spend some time here: <http://www.reddit.com/wiki/reddiquette>, to learn the expected behavior on the platform. The best thing to do on reddit is try to participate normally and submit a lot of other people’s great content to the correct subreddits, and occasionally submit your own potentially viral content, without the appearance of self-promotion.

Techmeme is a site that provides a curated list of the hottest and more important tech industry news, and it does this by putting the most important information on one single page. The site has a small team of editors that review what’s going on and decides what to post. The only way to suggest anything to this team is via Twitter. You can do this by messaging their Twitter account as described on this page: <http://news.techmeme.com/090128/twitter-tips>.

Blogs and Forums

The world of social started out as one where discussion sites (forums, Q+A, bulletin boards and the like) and the blogosphere reigned supreme. Eventually, consolidation and massive adoption of the major social networks captured the hearts and minds of most users, but the reports of the death of blogs and forums have turned out to be greatly exaggerated.

Forums and blogs still hold rich opportunities for marketers. For example, Moz.com receives tens of thousands of visits each week from blogs and discussion sites of all sizes, and their participation/interaction with those sources often yields fantastic results in referral traffic, mindshare, and links. Many brands do likewise, hiring community managers or evangelists to engage in the sites where industry topics are discussed, building up strong, recognizable profiles that help bring awareness and produce traffic and links.

Thus, as responsible inbound marketers, it's our job to measure these channels and quantify their impact.

- **Site/Brand Mentions** - mentions of a site or brand name can help lead you to content and conversations worthy of engagement, as well as tracking the quantities of those mentions (and possibly the sentiment as well) over time. Google Alerts and Mention (<http://web.mention.com/>) are potentially worth looking at to help monitor these mentions
- **Links** - direct links are nice because they appear either in link-tracking tools like Google Webmaster Tools, Open Site Explorer, Majestic SEO, ahrefs, LinkResearchTools, or directly in your web analytics (if they send any traffic).
- **Traffic** - a must-have for any inbound channel, visit-tracking is the most simple metric here.

For any inbound marketing channel (social or otherwise) you are considering, you a process similar to that shown in Figure 8-43.



- *Figure 8-43. Process for Identifying High ROI Social Networks*

Losing a few hours to channels that don't provide value is minimal next to the value of discovering and participating on those that do. You can see a presentation on a more in depth process for evaluating social metrics in these two articles: <http://www.slideshare.net/randfish/the-power-of-inbound-marketing> and <http://searchenginewatch.com/sew/how-to/2202307/social-media-roi-how-to-define-a-strategic-plan>.

A special thanks to Mark Traphagen (<http://www.stonetemple.com>) for his contributions to this portion of the chapter.

User Engagement as a Measure of Search Quality

Measuring user behavior on a web site can provide strong signals about that site's quality. For example, if a web site visitor arrives at a site, visits 10 pages over the course of an hour, selects a product, goes to a shopping cart, and then buys it, chances are pretty good that they found what they wanted at the web site. Contrast that with the visitor who arrives at a web page and hits the back button of their browser in less than a second.

These are examples of user engagement signals, and search engines are beginning to use these types of signals in their algorithms. The signals they are using and how they are being used is not easily discerned. Search engines are secretive about the details of their algorithms because they are important trade secrets, and because it makes the job of spammers harder. However, we know that user engagement signals are useful in measuring search quality and may also be used as ranking signals. If a page is a poor result for a user's search queries than the user's (probably very limited) interactions with that page will show it.

This section will examine the ways that the search engines can collect user engagement data, and the types of metrics they can collect. It should also be mentioned that in addition to the possible search result ranking impact of these signals, publishers should actively track user engagement metrics on their web sites for a number of other reasons, including measuring the quality of the organic search traffic the site is receiving. Increased web site engagement will result in increased conversion rates.

How Google and Bing Collect Engagement Metrics

Bing and Google both have a large number of data sources available to them. Some of the most important ones are:

1. Search Results - User interactions with the search results are a key source of data. For example, if the user conducts a search and chooses not to click on the 1st or 2nd result, but clicks on the 3rd, that can be a signal that the 3rd result may be the best result, especially if this is a common occurrence. Data like this is accumulated by the search engines in high volume every day.
2. Browsers - Microsoft's Internet Explorer and Google's Chrome both have sizable market share and could be a rich data source for search engines (<https://www.w3counter.com/globalstats.php?year=2014&month=3> and <http://thenextweb.com/insider/2014/02/01/ie11-passes-ie10-market-share-firefox-slips-bit-chrome-gains-back-share/>). Browsers are powerful data sources because they can monitor every action taken by a user.
3. Contextual Advertising - Google AdSense and Microsoft Content Ads are excellent sources of data as well. These offerings allow publishers to place ad units on their web sites and earn revenue when users click on them, and data is passed back.
4. Toolbars - Surfers can install either the Google Toolbar (<http://www.google.com/toolbar/>) or the Bing Toolbar (<http://toolbar.discoverbing.com/toolbar/l>) in their browsers. These toolbars offer users a number of conveniences that improve their web surfing experience.
5. Phone Operating Systems - Both Google and Bing offer phone operating systems. Determining Phone OS market share is complex for many reasons, but a April 2014 Techrepublic article pegged Google's Android OS as the number 1 phone OS in the US with more than 75% market share (<http://www.techrepublic.com/article/apple-v-google-the-goliath-deathmatch-by-the-numbers-in-2014/#>), and the Windows Phone OS is forecast to grow its share over time. (<http://seekingalpha.com/article/2071723-windows-os-phones-should-increase-market-share-in-2014-and-begin-to-rival-ios>).
6. Desktop and Web Based Applications - Microsoft Office is the dominant office productivity software suite worldwide. Links included in Microsoft Office applications can be tracked. Similarly, Google offers Google Docs as an office productivity suite, which operates entirely online.

7. Buy Data from ISPs - Both major search engines have the ability to buy usage data from ISPs.
8. Google Analytics - Various studies show a range of market share statistics for Google Analytics. For example, one study showed 63% of the Fortune 500 use Google Analytics on their web site (<http://www.e-nor.com/blog/google-analytics/google-analytics-solidifies-lead-in-fortune-500-adoption-in-2013>).
9. Goo.gl - Since December 2009, Google has offered a URL shortener (<http://googleblog.blogspot.com/2009/12/making-urls-shorter-for-google-toolbar.html>).
10. Internet Services - Google offers free WiFi service to the city of Mountain View and offers Google Fiber to a growing number of US cities. Google Fiber is also already present or being installed in Portland, San Jose, Salt Lake City, Provo, Phoenix, San Antonio, Austin, Kansas City, Nashville, Atlanta, Charlotte, and Raleigh-Durham.
11. Email Services – Services like Gmail, Outlook, and Hotmail (now integrated with Outlook) provide yet another data source.
12. Payment Services – Google Checkout, Google Wallet, and any mobile payment apps running on the Android OS are another potential source of data.

Potential User Engagement Signals

Each search engine has a rich array of data sources that enable them to measure a wide range of web surfer behavior. Here are some of the many signals they can extract from those sources:

1. Click Through Rate (CTR) - The search engines can measure the click through rate on links presented in search results, on web pages in URL shorteners, in their browsers, and more.
2. Clicking on Other Search Results - Once a user completes a search and visits a link, a common behavior indicating a problem with a result is that they return to the search results, often quite quickly, and then click on another result.
3. Generating New Searches - Similarly, a user may look at a given search result, come back to the search engine and modify their search query.
4. Bounce Rate - Bounce rate is a measurement of the percentage of users who visit only one page on a web site. Search engines extend that definition to take into account the interaction of the user with the search results. For example, if a user clicks on a search result, then returns back to the SERPs and clicks on another result that could be an indicator that the first result was not a good response for that search query.
5. Time on Page - Search engines can measure the amount of time spent on a given page using their browsers or toolbars. This is also referred to as *Dwell Time*, and more time on page might be considered signal of higher quality.
6. Time on Site - Similarly, time on site could be considered a positive signal if the average user spends more time on your site than they do on the sites of your competitors. Of course, it could also mean that your site is difficult to navigate or loads very slowly, so this signal would need to be looked at in conjunction with other signals.
7. Pages per Visit - This metric can be easily measured by the browser, a toolbar or Google Analytics. More page views implies greater user engagement with the site, however, some sites seek to increase page views by paginating their content to generate more ad impressions. As with Time on Site, a signal like this is hard to look at on a standalone basis.
8. Repeat Visits - Do users return to the site? Return visits could also be viewed as a quality signal.
9. Pages Printed - While many pages on the web do not lend themselves to printing, there are certain classes of pages that do, such as reference articles, recipes, maps, and similar content. If a user decides to print a page that indicates a higher level of interest in the content.
10. Pages Bookmarked in the User's Browser - If a user bookmarks a page so they can return to it later, that is also a positive signal.

11. Scroll Bar Usage - Another relatively subtle indicator of engagement is whether or not the user scrolls down to see more of the content on a page.

Voting Mechanisms

Another set of signals that search engines can use involve voting mechanisms. These are methods by which users directly indicate their approval or disapproval of content. Here are some examples:

1. Chrome Blocklist Extension - On February 14th, 2011 Google released a Chrome extension that allowed users to block specific web sites from coming in their search results (<http://googleblog.blogspot.com/2011/02/new-chrome-extension-block-sites-from.html>). This was not used as a ranking factor in the initial release of Google's Panda algorithm on February 24, 2011, but was included as part of the second release of Panda on April 11, 2011.
2. Google+ +1s and Google+ Shares of content – As discussed above, the evidence does not exist to support these being used as ranking signals just yet (other than for personalized search), but in the future it could be.
3. Brand Name Searches - Another potential signal of prominence is a high number of brand name searches. For example, brands like Coca Cola and Pepsi get searched on hundreds of thousands of times per month. This might make them better responses to generic search queries like *cola* than lesser known brands.
4. Brand Mentions - A related concept is that of brand mentions. People reference brand names on the web all the time, on web pages, in tweets, and elsewhere, without actually implementing links to the page of the brand's web site. While there is no documented evidence that search engine use this type of data, frequent brand references could be a strong signal.

Document Analysis

Document analysis is a somewhat different concept than User Engagement, but it can be used to predict how users will perceive the quality of content on the site. Strong signals indicating a likelihood of poor quality content could potentially be used as a ranking signal by the search engines.

Poor Editorial Quality

Poor spelling and bad grammar can indicate a lack of attention and care about the content. Care must be exercised here as user generated content will not necessarily be checked for spelling or grammar. Publisher authored content that is not subject to quality editing, however, may indicate low quality content.

Reading Level

The concept of reading level refers to the estimated grade level of the sophistication of the writing. One formula for this is known as the Flesch-Kincaid Reading Test. It measures the average words per sentence and the average syllables per word, and is actually not a measurement of the education level required to read the content, and might be better thought of as a measurement of the mental effort required to read a sentence.

If you want to check the reading level of a document you can do that using the spelling and grammar checker in Microsoft Word (<http://office.microsoft.com/en-us/word-help/test-your-document-s-readability-HP010148506.aspx>). You can also configure the Advanced Search settings (http://www.google.com/advanced_search) in Google to show the reading level for documents on the search results. This modifies the results as shown in Figure 8-44.

how to make french toast

Web

Videos

Images

Shopping

Books

More ▾

Search tools

Any time ▾

Reading level ▾

Clear

How To Make French Toast - IncredibleEgg.org

Ad www.incredibleegg.org/FrenchToast ▾

Look Up the American Egg Board's Recipes on Your Mobile Device!

Fast & Easy Recipes

Health & Nutrition

How To Videos

Hard Boiled Eggs

Results by reading level for how to make french toast:

Basic 88%

Intermediate 11%

Advanced < 1%

Directions

- Beat together egg, milk, salt, desired spices and vanilla.
- Heat a lightly oiled griddle or skillet over medium-high flame.
- Dunk each slice of bread in egg mixture, soaking both sides. Place in pan, and cook on both sides until golden. Serve hot.

French Toast | Recipe - Allrecipes.com

allrecipes.com/recipe/french-toast-i/ ▾ Allrecipes.com ▾

Feedback

How to Make French Toast: 9 Steps (with Pictures) - wikiHow

[www.wikihow.com](http://www.wikihow.com/How-to-Make-French-Toast) ▾ ... ▾ Breakfast ▾ French Toast ▾ wikiHow ▾

How to Make French Toast. Originally a way to use up stale bread, hot, steaming French toast is a perfect addition to a brunch, as well as making a delicious cold ...

How to make french toast - Instructables

www.instructables.com/id/how-to-make-french-toast-1/ ▾ Instructables ▾

Basic reading level

My boyfriend makes fantastic french toast. Super fantastic! I've tried making french toast lots of times, and it never turns out like I want it to, so...

Figure 8-44 Reading Level Shown in Google Results

It is likely that there is no one reading level that is the optimum one for each business. For example, if you have built a site targeted at children, then a younger reading level is appropriate. Contrast that with content on the latest advances in artificial intelligence algorithms, where much higher reading level would be expected.

Keyword Stuffing / Lack of Synonyms

Keyword stuffing used to be the basic approach of spammers back in the days when Alta Vista was the leading search engine. There are still many people who believe that repeating the same keyword over and over on a web page will lead to higher rankings for that phrase. They create documents that include way too many instances of the same word, and often don't make much use of synonyms for the key phrase. This is an indicator of a focus on search engines and not users, and may correlate with lower quality documents.

Ad Density and Offensive Ads

Search engines have learned that sites that have too many ads on them offer poor quality experiences for users. These pages may result in high bounce rates, but in some instances it is also possible to measure this directly. For example, consider the concept of measuring the percentage of the portion of a page above the fold that is occupied by ads. Too high a density might be taken as a negative signal. Google also has a patent on detecting annoying ads and pages as written about by Bill Slawski here: <http://www.seobythesea.com/?p=676>.

Sameness

An exact match search on "*mortgage tips*" returns hundreds of thousands of results, as shown in figure 8-45. Changing a tire is a very simple process, and it is likely that, of those hundreds of thousands of results, there is very little variation in the content.

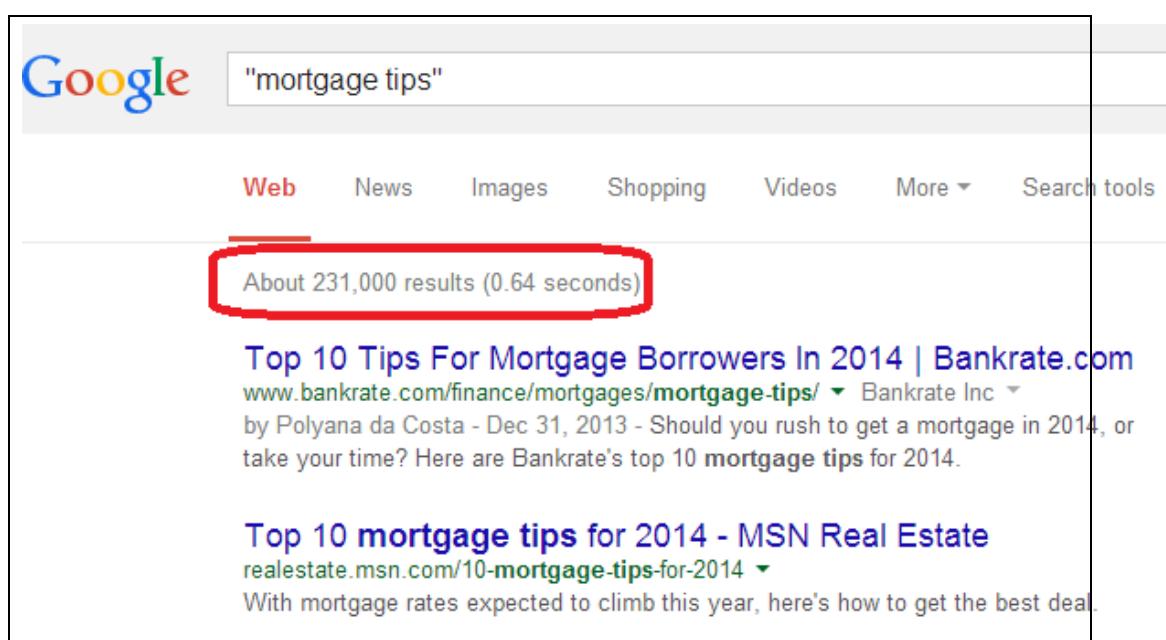


Figure 8-45 Search Query Returning Hundreds of Thousands of Results

Broaden this concept to consider a site where the majority of articles are written about topics that many have written on before them, and whose authors don't really bring any level of authority to the topic, or where no unique viewpoint is offered. Such a site may be lowered in rankings as users likely want to see authoritative opinions those they trust.

Page Speed

In April of 2010 Google announced that the page load time for your web site was now considered a ranking factor (<http://googlewebmastercentral.blogspot.com/2010/04/using-site-speed-in-web-search-ranking.html>). However, Google has indicated that this affects a small percentage of results (about 1%) and industry tests appear to confirm this, including the results of an investigation published in this Moz article: <http://www.moz.org/blog/site-speed-are-you-fast-does-it-matter>.

The bottom line on page speed is that there are a lot of reasons why you treat this as important, the biggest of which is the impact it has on conversion rates and bounce rates on your site. As an SEO factor it is likely to only have an impact on your results if you have a particularly slow site. If that's the case, the drive to improve your conversion rate should already be compelling enough to address the issue.

Optimizing User Experience to Improve SEO

There are many reasons to address the user experience on your site, including conversion optimization and customer retention. As we have seen in the sections on User Engagement as a Measure of Search Quality and Document Analysis in this chapter, SEO may also be impacted by the perceived user experience with a site. Measuring your site performance on these types of metrics can be a great start to improving the SEO performance of a web site.

However, it is important to consider a broader question, which is what do your visitors want? If you are selling lefthanded golf clubs, and you create a web page to rank on the search phrase *lefty ping putters*, what do you offer the visitors who come to that page from search? Reviews? Images? Videos? Putting tips? The e-commerce catalog front and center?

This is the most important part of optimizing a user experience. Creating a direct connection between the visitors arriving at your site, and what it is they are looking for. When you make this harder for them, metrics like bounce rate, time on site, repeat visits, will all suffer. When you make it easy, all of these metrics will improve.

One strategy for doing this is to survey your site visitors. If you don't have a lot of traffic, consider using a PPC campaign to bring in a number of visitors and then getting them to take the survey. Here is a 3 step process to create a survey to find out from people what they are looking for:

Step 1: Build a Survey

Start with a short series of questions asking the survey taker to imagine they've just performed the desired query. The basic structure is simple - request the top 3 content pieces your audience desires, then ask specifically about features that would make the page worthy of sharing (this is important, because it often differs substantively from what makes a page merely answer the user's query).

Finally, you can ask them to actually do a search on the target query (you don't want them to do it until the end, because what they find might bias their responses) and report any results they liked, as these can provide additional insight.

Step 2: Send it to Your Customers / Potential Customers

You can find customers or potential customers virtually anywhere - your friends, neighbors, co-workers, friends on social networks existing SEO traffic, PPC campaigns, etc. Anyone who fits your customer demographic or is creative enough to imagine themselves as that demographic will work.

A link in the bottom of your email newsletter or a share on Facebook/LinkedIn/Twitter can often do the job, too. You might even try posting a link in a relevant industry forum or discussion group (so long as you're sure it won't be perceived as spammy).

Step 3: Record Responses and Leverage them to Build What the People Want

Once you have your results in, analyze them to see what you have. If you need things clarified, go back and repeat the survey. It is better to do this than to act prematurely on incomplete or ambiguous data. Once you have built it, measure the results you get. Don't be afraid to go back and make more refinements or even another survey to continue to get smarter about what your visitors want.

Additional Social Media Resources

In this chapter we have focused primarily on the interactions between social media and SEO. However, social media is itself a significant source of traffic. It deserves its own investigation and study. Below are listed some of the top resources available online to help guide you further into the world of social media.

Social Media Blogs

For further reading on social media, you can find good information on these sites:

- <http://www.socialmediaexaminer.com/>
- <http://blog.hubspot.com/>
- <http://socialmouths.com/blog/>
- <http://socialmediatoday.com/>
- <http://danzarrella.com/#>
- <http://www.convinceandconvert.com/blog/>
- <http://www.razorsocial.com/blog>
- <http://simplymeasured.com/blog/>
- <http://maximizesocialbusiness.com/>
- <http://www.briansolis.com/>
- <http://www.moz.com/blog>
- <http://marketingland.com>
- <http://allfacebook.com>
- <http://socialtriggers.com>
- <http://www.ducttapemarketing.com>
- <http://www.viralblog.com>
- <http://pushingsocial.com>
- <http://www.likeable.com/blog/>

Conclusion

It does not appear that social media acts as a direct SEO ranking signal, except for personalized search and Google+. However, no one can debate that social media and user engagement have become significant factors for those practicing SEO from an indirect perspective. Links may still be the strongest source of signals to search engines, but these new factors are already an important part of the digital marketing landscape, and may gain additional weight over time as the search engines learn to use them better. A balanced SEO strategy needs to take social media, user data and document analysis techniques into account.

In summary, it means that the brand publisher must seriously consider these four areas as a part of their web promotional strategy:

1. Participating in social media communities
2. Providing an engaging user experience
3. Offering unique and differentiated content
4. Building a brand

Picking the right mix of these areas is already important for improving conversion and SEO results for publishers, and they will only get more important over time.

Panda, Penguin & Penalties

The search engines use many techniques for measuring search quality, including manual penalties, algorithmic penalties, and ranking algorithms. Sometimes these algorithms and penalties can cause a large change in your organic traffic. Significant decreases in the search engine traffic to your web site can be devastating to a business. As shown in Figure 9-1, sometimes these drops in traffic can be quite large.



Figure 9-1. Example of Major Loss in Traffic

If the business shown in Figure 9-1 generates all of its revenue from organic search traffic, this would represent a devastating blow. This type of loss of revenue can mean laying off people, or even closing the business.

For that reason, it's a good idea to have a working understanding of what the potential causes are of these traffic losses, and if you have already suffered such a loss, know how you can figure out what the cause is and what you need to do to recover.

Diagnosing the Cause for a Traffic Loss

The first step is to check your analytics data to see if the drop in traffic is in fact a loss of organic search engine traffic. If you have Google Analytics on your site, make sure you check your traffic sources, and then isolate just the Google traffic.

If you confirm that it is a drop in Google organic search traffic, then the next step is to see if you have received a message in Google Webmaster Tools indicating that you have been penalized by Google. Figure 9-3 shows an example of a manual penalty message from Google. This is what Google refers to as a “manual action”.

The screenshot shows a message in Google Webmaster Tools. The title is "Reconsideration request for http://[REDACTED] Site violates Google's quality guidelines". The message body is as follows:

Dear site owner or webmaster of http://[REDACTED]
We received a request from a site owner to reconsider http://[REDACTED] for compliance with Google's Webmaster Guidelines.
We've reviewed your site and we still see links to your site that violate our quality guidelines. Specifically, look for possibly artificial or unnatural links pointing to your site that could be intended to manipulate PageRank. Examples of unnatural linking could include buying links to pass PageRank or participating in link schemes.
We encourage you to make changes to comply with our quality guidelines. Once you've made these changes, please submit your site for reconsideration in Google's search results.
If you find unnatural links to your site that you are unable to control or remove, please provide the details in your reconsideration request.
If you have additional questions about how to resolve this issue, please see our Webmaster Help Forum for support.
Sincerely,
Google Search Quality Team

Figure 9-2. Example Penalty Message in Google Webmaster Tools

If you have such a message, you now know what the problem is and you can get to work on fixing it. It is not fun to have the problem, but knowing what you are dealing with is the first step in recovery.

If you don't have a message you will need to dig deeper to determine the source of your problem. The next step is to determine the exact date on which your traffic dropped. For example, if that date was April 24, 2012, take note of that. Then go to the Moz Google Algorithm Update Page (<http://moz.com/google-algorithm-change>) and see if you can find the date of your traffic drop listed on that page.

Using our example of April 24, 2012, Figure 9-3 shows that sites that suffered traffic losses on this date were probably hit by the Penguin 1.0 algorithm release by Google.

The screenshot shows a list of Google algorithm updates from 2002 to 2014. The 2012 section is highlighted with a red box, focusing on the Penguin update. The Penguin update is dated April 24, 2012, and is described as the "Webspam Update". The text notes that Penguin adjusted a number of spam factors, including keyword stuffing, and impacted an estimated 3.1% of English queries. The entire Penguin 1.0 section is enclosed in a red box.

Year
2014
2013
2012
2011
2010
2009
2008
2007
2006
2005
2004
2003
2002

Penguin — April 24, 2012
 Barely a week after Panda 3.5, Google rolled out yet another Panda data update. The implications of this update were unclear, and it seemed that the impact was relatively small.
 Confirmed: Panda Update 3.6 Happened On April 27th (SEL)

Penguin — April 24, 2012
 After weeks of speculation about an "Over-optimization penalty", Google finally rolled out the "Webspam Update", which was soon after dubbed "Penguin." Penguin adjusted a number of spam factors, including keyword stuffing, and impacted an estimated 3.1% of English queries.
 Another step to reward high-quality sites (Google)
 The Penguin Update: Google's Webspam Algorithm Gets Official Name (SEL)
 Google Penguin Update Recovery Tips & Advice (SEL)
 Two Weeks In, Google Talks Penguin Update, Ways To Recover & Negative SEO (SEL)

Panda 3.5 (#13) — April 19, 2012
 In the middle of a busy week for the algorithm, Google quietly rolled out a Panda data update. A mix of changes made the impact difficult to measure, but this appears to have been a fairly routine update with minimal impact.
 Google Mocks Me For Missing Panda 3.5 (SER)

Figure 9-3. Moz Google Algorithm Update Page Showing Penguin 1.0

Barracuda Digital also offers a tool called the "Penguin Tool" (<http://www.barracuda-digital.co.uk/penguin-tool/>) that will overlay the Google algorithm updates on your Google Analytics data to make this date comparison analysis easier for you.

It can also happen that you suffer major traffic loss, you don't get a message in Google Webmaster Tools, and that the date of that loss does not line up with a known Google algorithm update. This will unfortunately make the process of figuring out how to recover much harder, since you won't know the reason for the drop.

Google does make smaller changes to their algorithms on a daily basis. From their perspective, these are smaller changes and not major updates. They may be much harder to recover from as well. The best things to do is to focus on the best practices outlined in Chapter 6 and Chapter 7, or if you can afford SEO advice, to bring in an expert to help you figure out what to do next.

Summary of Major Google Algorithms

The two most famous algorithms released since February 2011 are Panda and Penguin, but there are several others that Google also released since that timeframe. Panda and Penguin are examined in detail later on in this chapter. Here is a summary of some of the other major algorithms released by Google:

Top Heavy Update

Released on January 19, 2012, this update impacts sites that have too many ads above the fold. This is true even if the ads are AdSense ads (from Google's own advertising programs). The remedy is to back off placing lots of ads near the top of your web pages. A banner ad is probably fine, as is something on the right rail of your site, but more than that starts to run into trouble with this algorithm. This algorithm was updated on October 9, 2012.

DMCA

Announced on August 10, 2012, this update applies penalties to sites that receive too many valid DMCA takedown requests. DMCA stands for Digital Millennium Copyright Act. You can file a

DMCA takedown request if someone has stolen your content and published it as their own. Of course, to avoid this penalty, the remedy is to only publish your own original content. If you hire people to write for you, you should also ensure that they are writing original content as well.

Exact Match Domain

Released on September 27, 2012, Google started to change the way it values exact match domains. An example of an exact match domain might be: <http://www.blue-widgets.com> (note that this is not currently a real site, and is in fact a parked page). Historically, a site like this would have an advantage in ranking for the search phrase “blue widgets”. The purpose of this update was to reduce or eliminate that advantage.

Payday Loans (Spammy Sites)

Released on June 11, 2013, this algorithm targeted sites that it considers to be in particularly spammy market spaces. Google specifically mentioned pay day loans, casinos, and porn as targeted markets. This algorithm was updated on May 16, 2014.

Hummingbird

Announced on September 26, 2013, this was not really an algorithm update, but instead a rewrite of Google’s search platform, with a design goal of making the overall search engine more flexible and adaptable for the future. Part of the design of this rewrite included more support for natural language queries and semantic search.

These are some of the major releases by Google, but Google makes many smaller changes as well. In August of 2014, the head of search quality for Google, Amit Singhal, indicated that they had made 890 changes in the past year (<http://searchengineland.com/google-made-890-improvements-search-past-year-201065P>).

Panda

Google’s Panda algorithm shook the search landscape when it came out on February 24, 2012. In their announcement of the release (<http://googleblog.blogspot.com/2011/02/finding-more-high-quality-sites-in.html>), Google said the following:

“Many of the changes we make are so subtle that very few people notice them. But in the last day or so we launched a pretty big algorithmic improvement to our ranking—a change that noticeably impacts 11.8% of our queries—and we wanted to let people know what’s going on. This update is designed to reduce rankings for low-quality sites—sites which are low-value add for users, copy content from other websites or sites that are just not very useful. At the same time, it will provide better rankings for high-quality sites—sites with original content and information such as research, in-depth reports, thoughtful analysis and so on.”

Other information also came out pretty quickly about this release, which Danny Sullivan of Search Engine Land initially called the Farmer update (<http://searchengineland.com/google-forecloses-on-content-farms-with-farmer-algorithm-update-66071>). Between this additional information and the initial announcement several aspects of Panda quickly became clear:

- This was a very large change. Very few algorithm updates by Google impact more than 10% of all search queries.
- This algorithm is focused on analyzing content quality.
- Scraper sites were targeted
- Sites lacking substantial unique information were also targeted. In particular, content farms were part of what Google was looking to address – a move that Eric Enge predicted would happen 3 weeks before this release (<http://searchengineland.com/the-rise-and-fall-of-content-farms-62323>).
- Google clearly states a strong preference for new research, in-depth reports, and thoughtful analysis
- The Panda algorithm does not use the link graph as a ranking factor.

The second release of Panda came out on April 11, 2011 (<http://googlewebmastercentral.blogspot.com/2011/04/high-quality-sites-algorithm-goes.html>). What made this particularly interesting is that it incorporated data gathered from Google's Chrome Blocklist Extension (<https://chrome.google.com/webstore/detail/personal-blocklist-by-goo/nolijncfnkgaikbjbdaogikpmpbdcdef?hl=en>). This is an extension that Google made to Chrome that allowed users to indicate that they wanted pages removed from the search results. Figure 9-4 illustrates what this looks like in operation.



Figure 9-4. Chrome Blocklist Extension

This was the first time that Google ever publicly confirmed that they were using a form of direct user input as a ranking factor in any of their algorithms.

Initially, Panda was focused only on the US, but was rolled out internationally on August 12, 2011 to most of the rest of the world except Japan, China, and Korea.

Google has done many releases since the initial release. Here is a full list of all known Panda releases:

Table 9-1. List of Panda Releases

Panda Release	Date	
Panda 1.0	February 23, 2011	Impacted 11.8% of search results
Panda 2.0	April 11, 2011	Added use of the Chrome Blocklist Extension
Panda 2.1	May 9, 2011	
Panda 2.2	June 21, 2011	
Panda 2.3	July 23, 2011	
Panda 2.4	August 12, 2011	Rolled out internationally, Impacted 6-9% of queries internationally
Panda 2.5	September 28, 2011	
Panda "Flux"	October 5, 2011	Called "Flux" as a result of a Matt Cutts tweet,

		Impacted 2% of queries
Minor Update	October 11, 2011	
Panda 3.1	November 18, 2011	Naming skipped an official 3.0 release
Panda 3.2	January 18, 2012	
Panda 3.3	February 27, 2012	
Panda 3.4	March 23, 2012	Impacted 1.6% of results
Panda 3.5	April 19, 2012	
Panda 3.6	April 27, 2012	
Panda 3.7	June 8, 2012	Impacted less than 1% of queries
Panda 3.8	June 25, 2012	
Panda 3.9	July 24, 2012	Impacted about 1% of queries
Panda #18	August 20, 2012	Industry agreed nomenclature changed to counting the release number
Panda #19	September 18, 2012	
Panda #20	September 27, 2012	Impacted 2.4% of queries
Panda #21	November 5, 2012	Impacted 1.1% of queries
Panda #22	November 21, 2012	
Panda #23	December 21, 2012	Impacted 1.3% of queries
Panda #24	January 22, 2012	Impacted 1.2% of queries
Panda #25	March 14, 2013	Date estimated by Moz. Google announced Panda made part of core algorithm. Rolling updates no longer announced, but happen monthly and roll out over 10 or so days
Panda #26	June 11, 2013	
Panda #27	July 18, 2013	Rumored to soften Panda's impact
Panda 4.0	May 20, 2014	Major update confirmed by Google. Also appeared to soften Panda's impact on many sites, but new sites were also hit.

Since March 14th, 2013 Google announcements of Panda updates became relatively rare, and the above table shows only the announced updates. However, these updates appear to occur monthly and roll out over a period of more than a week.

Since that time Google has only confirmed three Panda updates, with the most recent one being Panda 4.0 on May 20, 2014.

To track all Google updates over time you can check the Google algorithm change history page on the Moz web site (<http://moz.com/google-algorithm-change>).

Causes of Panda

Google has historically offered relatively vague statement on how the Panda algorithm works. For example, on May 6, 2011 Amit Singhal offered his advice on building high quality sites (<http://googlewebmastercentral.blogspot.com/2011/05/more-guidance-on-building-high-quality.html>). In it, he suggested a list of questions that you could use to determine if you were on such a site:

- Would you trust the information presented in this article?
- Is this article written by an expert or enthusiast who knows the topic well, or is it more shallow in nature?

- Does the site have duplicate, overlapping, or redundant articles on the same or similar topics with slightly different keyword variations?
- Would you be comfortable giving your credit card information to this site?
- Does this article have spelling, stylistic, or factual errors?
- Are the topics driven by genuine interests of readers of the site, or does the site generate content by attempting to guess what might rank well in search engines?
- Does the article provide original content or information, original reporting, original research, or original analysis?
- Does the page provide substantial value when compared to other pages in search results?
- How much quality control is done on content?
- Does the article describe both sides of a story?
- Is the site a recognized authority on its topic?
- Is the content mass-produced by or outsourced to a large number of creators, or spread across a large network of sites, so that individual pages or sites don't get as much attention or care?
- Was the article edited well, or does it appear sloppy or hastily produced?
- For a health related query, would you trust information from this site?
- Would you recognize this site as an authoritative source when mentioned by name?
- Does this article provide a complete or comprehensive description of the topic?
- Does this article contain insightful analysis or interesting information that is beyond obvious?
- Is this the sort of page you'd want to bookmark, share with a friend, or recommend?
- Does this article have an excessive amount of ads that distract from or interfere with the main content?
- Would you expect to see this article in a printed magazine, encyclopedia or book?
- Are the articles short, unsubstantial, or otherwise lacking in helpful specifics?
- Are the pages produced with great care and attention to detail vs. less attention to detail?
- Would users complain when they see pages from this site?

There are a few key points that can be extracted from this advice, and in addition, the industry has been able to determine and clarify a number of the causes of Panda. These include:

Thin Content

As you might expect, this is defined as pages with very little content. Examples of this might be user profile pages on forum sites with very little information filled in, or an e-commerce site with millions of products, but very little information provided about each one. Another form of thin content occurs

Non-Original Content

These may be scraped pages, or pages that are only slightly rewritten are relatively easily detected by Google. Sites with even a small number of these types of pages can be impacted by Panda.

Non-Differentiated Content

Even if you create all original articles this may not be enough. If every page on your site covers topics that have all been written about by others hundreds or thousands of times before, then you really have nothing new to add to the web with your site. Consider, for example, the number of articles in the Google index about making French Toast, as shown in Figure 9-5. There are over 30,000 pages on the web that have the phrase "How to make French Toast" in the title of the page. From Google's perspective, they don't need another web page on that topic.

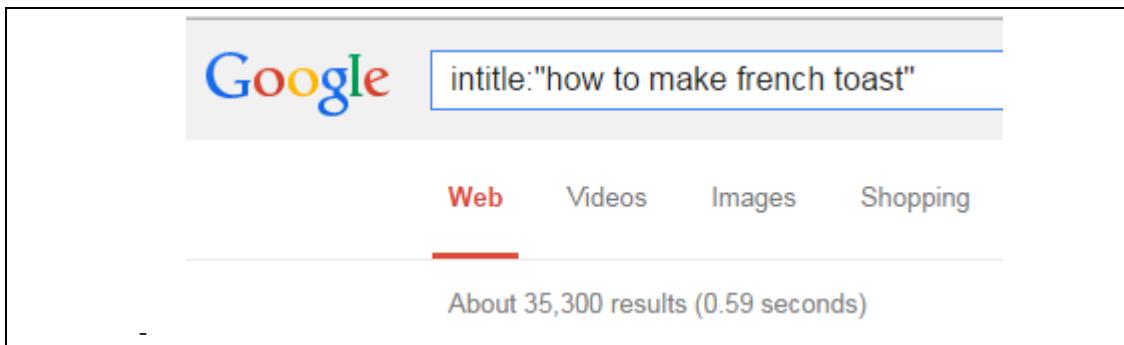


Figure 9-5. There are tens of thousands of pages on how to make French Toast

Poor Quality Content

Content which is inaccurate or poorly put together. In many cases this may be hard to detect, but, as mentioned in Amit Singhal's article, one possible way to detect this is content that has a lot of poor grammar and spelling mistakes. Google could also potentially use fact checking as another way to determine poor quality content.

Curated Content

Sites that have large numbers of pages with lists of curated links do get hit by Panda. Content curation is not inherently bad, but if you are going to do it, it's important to incorporate in a significant amount of thoughtful commentary and analysis. Pages that simply include lots of links will not do well, nor will pages that include links and include only a small amount of unique text. Content curation is explored in depth in this article: <http://searchenginewatch.com/article/2290885/Content-Curation-SEO-A-Bad-Match>.

Thin Slicing

This was believed to be one of the original triggers for the Panda algorithm, as this is what content farms were doing. Imagine you wanted to publish content on the topic of schools with nursing programs. Content farm sites would publish many articles on the same topic, with article titles such as: "nursing schools", "nursing school", "nursing colleges", "nursing universities", "nursing education", and so forth. There is no need for all of those different articles, and this led to Google targeting this practice with Panda.

Database Generated Content

The practice of using a database to generate web pages does not have to be inherently bad, but many companies were doing this to an excessive scale. This led to lots of thin content pages or poor quality pages, so many of these types of sites were hit by Panda.

Google Cares about Diversity

Diversity is important to overall search quality for Google. One simple way to illustrate this is with the search query "Jaguar". As it turns out, this can mean an animal, a car, a guitar, an operating system, or even an NFL team. Consider the possibility that normal ranking signals might suggest that the results look as shown in in Figure 9-6

[Jaguar: Luxury Cars & Sports Cars | Jaguar USA](#) 

The official home of Jaguar USA. Our luxury cars feature innovative designs along with legendary performance to deliver one of the top sports cars in the ...

[Models & Pricing](#) 

Available in four extraordinary models (XF, XJ Portfolio, XJ ...

[More results from jaguarusa.com »](#)

[Luxury Coupes](#) 

The Jaguar XK, in coupe or convertible, combines Jaguar's ...

[Jaguar Cars - Wikipedia, the free encyclopedia](#) 

en.wikipedia.org/wiki/Jaguar_Cars

Jaguar Cars Ltd, known simply as Jaguar is a British luxury and sports car manufacturer, headquartered in Whitley, Coventry, England. It is part of the Jaguar ...

[Jaguar International - Market selector page](#) 

www.jaguar.com/

Official worldwide web site of Jaguar Cars. Directs users to pages tailored to country-specific markets and model-specific websites.

[Jaguar - Wikipedia, the free encyclopedia](#) 

en.wikipedia.org/wiki/Jaguar

The jaguar is a big cat, a feline in the *Panthera* genus, and is the only *Panthera* species found in the Americas. The jaguar is the third-largest feline after the tiger ...

[Jaguars, Jaguar Pictures, Jaguar Facts - National Geographic](#) 

animals.nationalgeographic.com/animals/mammals/jaguar/

Learn all you wanted to know about jaguars with pictures, videos, photos, facts, and news from National Geographic.

-

Car
Related

Animal
Related

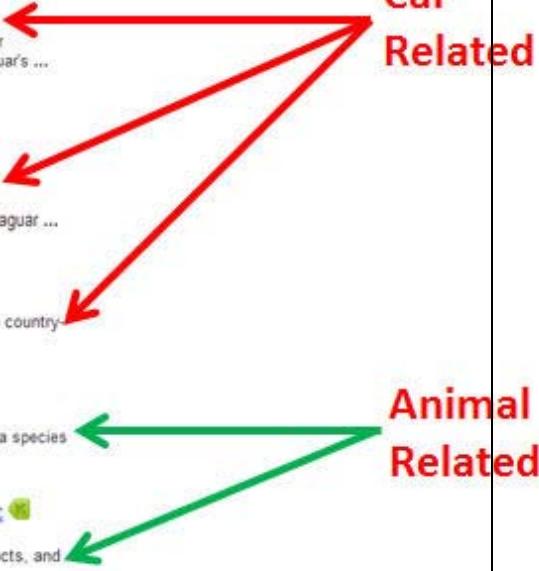


Figure 9- 6 Normal Ranking Signals May Show Us This Result

Note that the search results all focus on the car at the top. This may be what the basic ranking signals suggest. However, if the searcher is looking for information on the animal those results are pushed down a bit. As a result, Google may use other signals to decide to alter the results to look more like what we see in Figure 9-7.

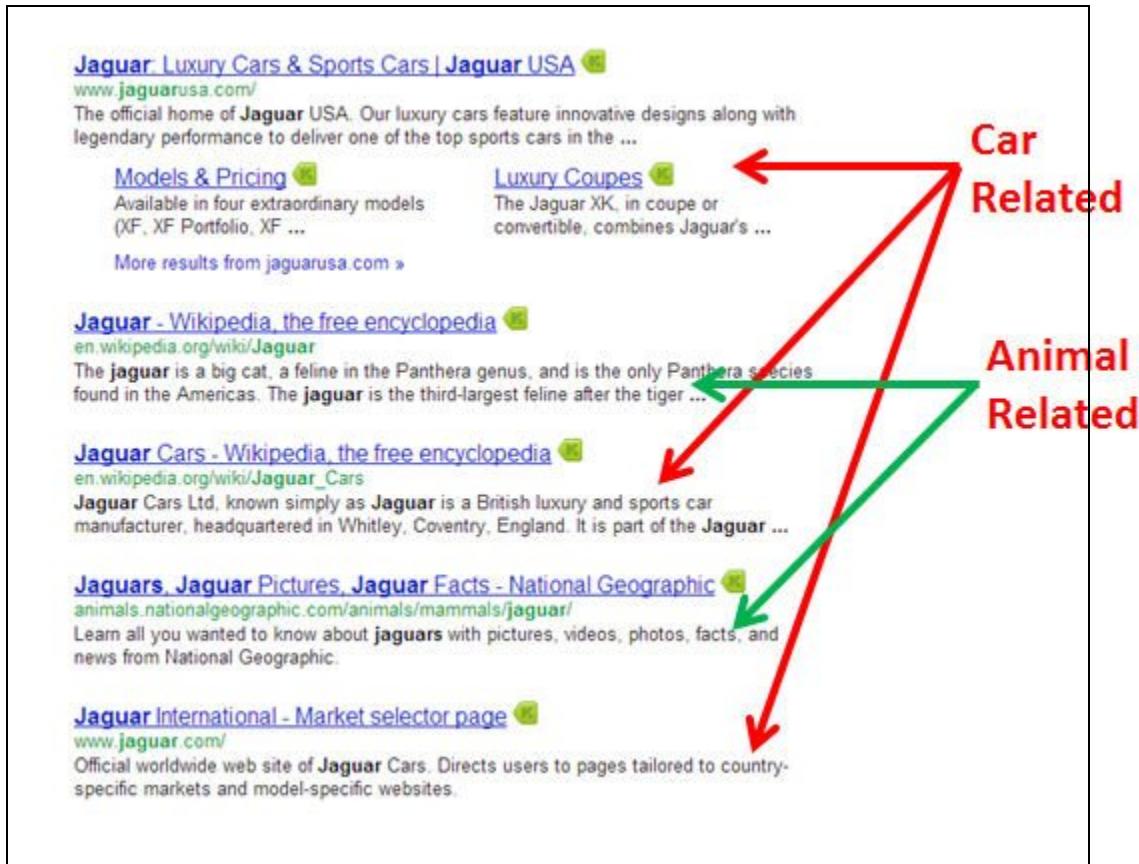


Figure 9- 7 Query Deserves Diversity May Alter the Result to Look Like This

In this version of the results, one of the animal related results has been pushed into the 2nd position. It is known that Google is making these types of adjustments to the SERPs using a concept known as Query Deserves Diversity.

They can make these types of adjustments by measuring user interaction with the search results to determine what ordering of the results provides the highest levels of user satisfaction. For example, even if traditional ranking signals would put another page for Jaguar the car next, it might make sense for the next result to be about the animal.

Role of Authority

Consider again the example of “How to make French Toast” we showed in Figure 9-5. While Google has plenty of results on the topic, there are, of course, some sites that do get to rank highly for such a search query. How is that determined?

It is known that very high authority sites are likely to do fine when publishing content on a topic which is already very well covered on the interwebs. There are a few possible reasons why this is the case:

- Reputation and authority is a big factor. For example, if the New York Times Lifestyle section posted a new article on how to make French Toast, even though it is not particularly unique, readers might respond well to that anyway. User interaction signals with the search result for that content would probably be quite strong, simply because of their reputation.
- High authority sites probably got to be that way because they don’t engage in much of the behavior that Google does not like. Chances are that you won’t find a lot of thin content, “me-too” content, thin slicing, or any of the issues that are Panda triggers.
- Google may simply applying looser criteria to their site than they do to other sites.

Exactly what factors allow higher authority sites to have more leeway is not clear. Is it that Google is measuring user interaction with the content, the quality of the content itself, the authority of the publisher, or some combination of these factors? There are probably elements of all three in what Google does.

Issues on One Part of Your Site Can Impact the Entire Site

Weak content on even one single section of a larger site can cause Panda to bring down the rankings for the whole site. This is true even if it is less than 20% of the pages for the site. When putting together a plan to recover from Panda, it is important to take this into account.

Path to Recovery

The road to recovering from a Panda penalty may be a long one. Often times it requires substantial re-evaluation of your site's business model. You need to be prepared to look at your site with a highly critical eye, and this is often very hard to do with your own site.

It's a good idea to consider bringing in an external perspective to evaluate your site. You need someone that is willing to look you in the eye and tell you that your baby is ugly.

Once you go through these types of evaluations, it may lead you to realize that even the basic premise of your site is broken, and that you need to substantially restructure it. Making these types of decisions are quite hard, but you may need to be prepared to do just that.

As you consider these hard choices, you will be tempted to look at your competition that did not get hit, and that is a good idea. However, you may see instances of thin content, weak content, "me-too" content, and other pages that look like they are just as bad as the content on your site.

Don't let this type of analysis deter you from making the hard choices. There are so many factors that Google uses in their ranking algorithms that you will never really know why your site was hit by Panda and that competitor's site was not.

What you do know is that Google's Panda algorithm does not like something about your site. This may include complex signals based on how users interact with your listings in the search results, data that you don't have access to.

To rebuild your traffic, it's best to dig deep and take on hard questions about how you can build a site with lots of fantastic content that gets lots of user interaction and engagement. While it is not believed that social media engagement is a factor in Panda, there is likely a strong correlation between high numbers of social shares and what Google considers to be good content.

Highly differentiated content that people really want, enjoy, share, and link to is what you want to create on your site (This article expands on this concept: <http://searchenginewatch.com/article/2278413/After-Google-Panda-Standout-Content-is-Essential>). There is a science to creating content that people will engage with.

We know that picking engaging titles for the content is important, and that including compelling images matters too. Make a point of studying how to create engaging content that people will love, and apply those principles to every page you create. In addition, measure the engagement you get, test different methods, and improve your ability to produce great content over time.

Addressing Weak Pages

As you examine your site a big part of your focus should be addressing the weak pages on the site. This may come in the form of an entire section of weak content, or a number of pages interspersed among the higher quality pages on your site. Once you have identified those pages, there are a few different paths you can take to address the problems you find:

- Improve the content. This may involve rewriting the content on the page, and making it more compelling to users who come to the page.
- Add the NoIndex metatag to the page (you can read about how to do this in Chapter 6). This will tell Google to not include these pages in their index and will take them out of the Panda equation.

- Delete the pages altogether, and 301 redirect them to other pages on your site. This should only be used if there are pages which are relevant to the deleted pages.
- Delete the pages and return a 410 http status code when someone tries to visit the deleted page. This tells the search engine that the pages have been removed from your site.
- Use the URL removal tool (<https://www.google.com/webmasters/tools/removals?pli=1>) to take the page out of Google's index. This should be done with great care. You don't want to accidentally delete other quality pages from the Google index!

Expected Timeline to Recovery

Even though they are no longer announced, Panda releases come out roughly once per month. However, once you have made the necessary changes you will still need to wait. Google has to re-crawl your site to see what changes you have made. This may take them several months before they have seen enough of the changed or deleted pages to tilt the balance in your favor.

What if You Don't Recover?

Sadly, this usually means that you have not done enough to please the Panda algorithm. It may be that you were not harsh enough in deleting poorer quality content from your site. Or, it may mean that there are not enough signals that Google is getting that people really care about what they find on your site.

Either way, it means that you have to keep at the process of making your site more interesting to users. It's best if you go beyond viewing this as a way to deal with Panda, and as a mission to make your site one of the best ones on the web.

This requires substantial vision and creativity. Frankly, it's not something that everybody can accomplish without making significant investments in terms of time and money. One thing is clear, and that is that you can't afford to cut corners when trying to address the impact of the Panda algorithm.

If you have invested a lot of time and made many improvements, but there is content that you know is not so great, and you want to see if you have done enough, chances are pretty good that you haven't. You will likely find yourself 4 months later wishing that you had kept at it.

In addition, the Panda algorithm is constantly evolving. Even if you have not been hit by this algorithm, the message from Google is clear. They are going to provide the greatest rewards to sites that provide fantastic content and great user experiences.

The best path is to be passionate about creating a site that offers both of these. This is how you maximize your chances of recovering from Panda, and from being impacted by future releases of the Panda algorithm.

Successful Panda Recovery Example

In Figure 9-1, we showed an example of a site suffering a major traffic loss. This traffic loss was suffered by a site called The Riley Guide. The site is dedicated to providing information on jobs and careers. A thorough examination of the site showed that it had two major problems:

- The content on careers had too much similarity to the type of content you can find on a major government web site located at <http://www.bls.gov>. Even though the content was not copied from that location, it did not offer any unique differentiation or value to the web.
- The content on the rest of the site mainly consisted of links to resources for job information elsewhere on the web. This was structured like a librarian's index, with links and short commentaries on each link. In some cases, the commentary was copied from the target web site.

In short, the original site structure did not offer a lot of value. The publishers of the site took it through a complete overhaul. This was a lengthy process as about 140 pages had to be rewritten from scratch, and another 60 or so pages had the NoIndex tag applied to them.

Figure 9-8 shows the results of this effort. Not only the site recover from Panda, but it is now hitting new highs for traffic.

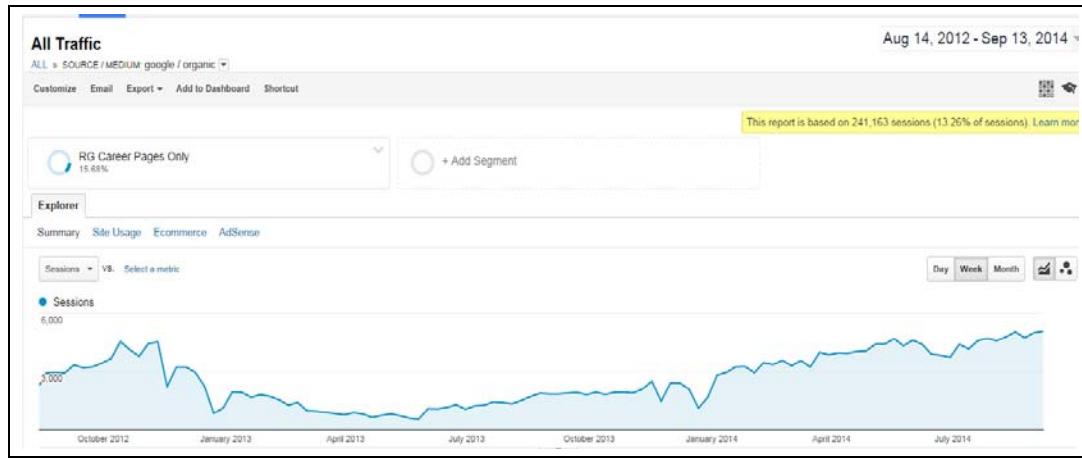


Figure 9-8 Panda Traffic Recovery

Penguin

Google's Penguin algorithm was first released to the world on April 24, 2012. It was the first major algorithm implemented by Google to address bad links. Like Panda before it, the Penguin release shook the search landscape. Since the initial release of this algorithm there have been several incremental releases, and an expansion of the scope of the types of links addressed.

Diagnosis

Diagnosis of Penguin algorithm hits are generally reasonably easy to detect. There have only been five total releases of Penguin, and these are the only dates on which you may have been hit by Penguin, and unfortunately, appear to be the only dates on which you can recover from it. Table 9-2 lists all of the releases of Penguin as of September of 2014.

Table 9-2. List of Penguin Releases

Penguin Release	Date	
Penguin 2.1	October 4, 2013	
Penguin 2.0	May 22, 2013	Began targeted links more at the page level
Penguin 1.2	October 5, 2012	Impacted about 0.3% of queries
Penguin 1.1	May 25, 2012	
Penguin 1.0	April 24, 2012	Impacted 3.1% of US search queries

As you can see, the releases are quite rare, and this is one of the jarring aspects of Penguin. Once you have been hit by it, you have to wait to the next release to have a chance to recover. That could be as much as one full year!

For purposes of diagnosis, if your traffic experience a significant drop on one of the dates listed in Table 9-1, it's likely that you have been hit by Penguin.

Causes of Penguin

As you will see in the section on "Links that Google Does not Like" there are many different types of links that are considered bad links. However, Penguin appears to focus on a more limited set of link types. These are:

- Article Directories: From the very first release of Penguin, Google targeted sites that obtained links from article directories. While these links may not always reflect manipulative intent by the publisher, Google found that people who leveraged article directories tended to operate lower quality sites.
- Cheap Directories: These were also targeted in the very first release of Penguin. There are a few directories which are genuinely high quality, such as the Yahoo Directory, DMOZ, Business.com, Best of the Web, and perhaps a few others specific to your industry vertical. Stay away from the rest.
- Excessive Rich Anchor Text: Excessive use of rich anchor text was also a part of the initial release of Penguin. What they specifically targeted was too many instances of the same anchor text pointing to any of the URLs on your site. Google does not expect, or want, all links to say “click here”, but they see it as a signal of spammy behavior when the exact same keyword-rich anchor text is used repeatedly.
- Low Relevance International Links: While it is not confirmed that this is a part of any Penguin release, anecdotal evidence suggests it might be. Consider any links from countries where don’t sell your products or services as a potential problem, unless they come from truly high quality sites.
- Comment Spam: Excessive use of implementing links in comments on other people’s blog posts and forums is also a problem.

There may be more types of links that are a part of the Penguin algorithm or that could be included in future releases. A complete summary of problem links is discussed later in this chapter in “Links That Google Does Not Like”.

Path to Recovery

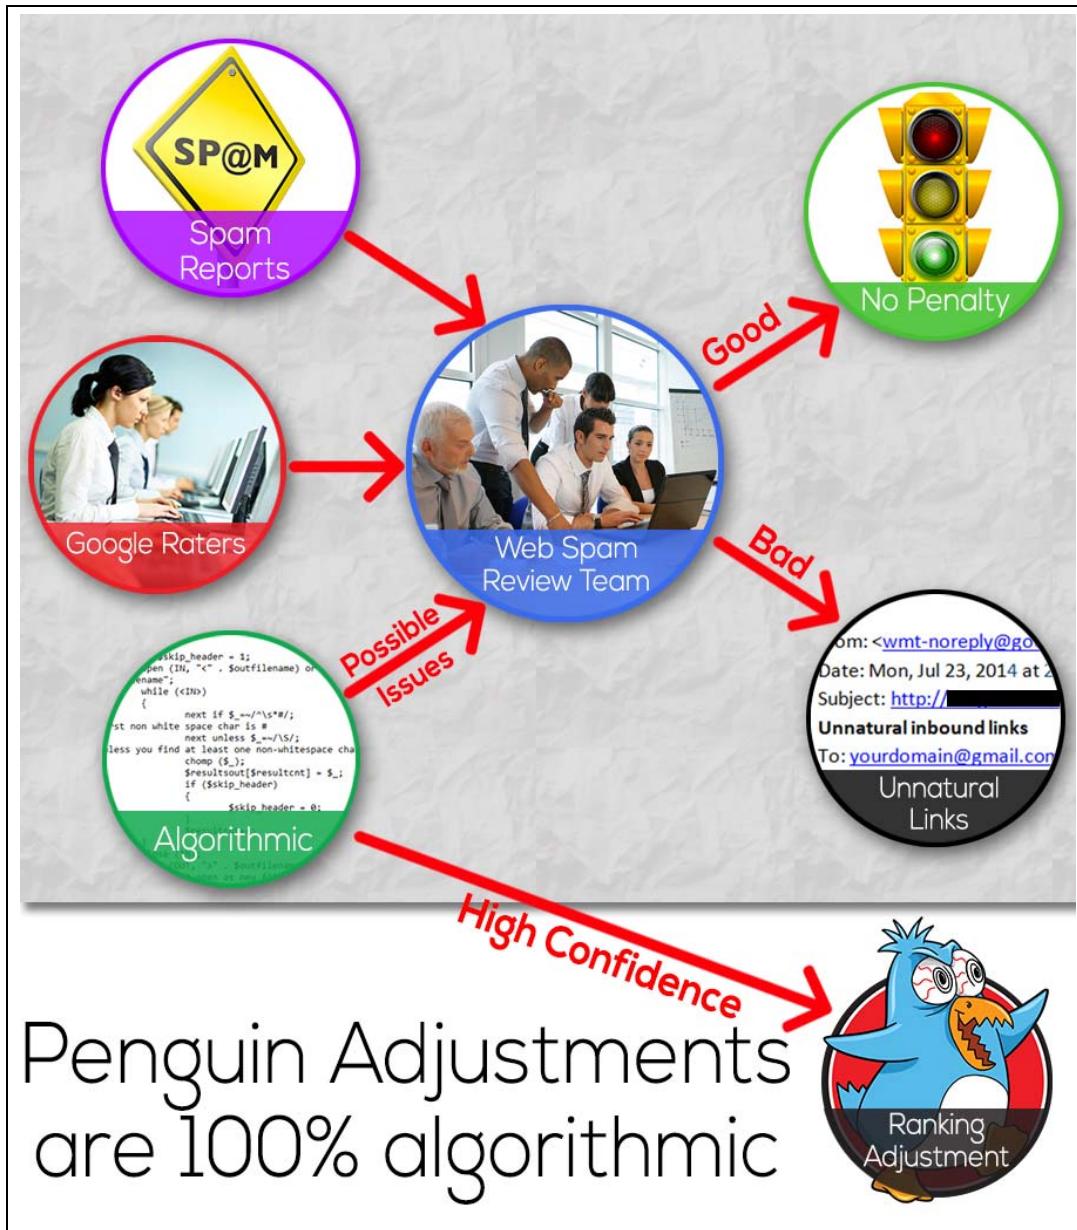
The first step is to realize that Penguin releases are rare. As they happen roughly twice per year, you don’t want to get cute and try to save some of your questionable links. Be aggressive and clean up every possible problem. Certainly address the types of links listed in “Causes of Penguin”, but you should seriously consider dealing with all types of link problems listed in “Links that Google Does not Like” in this chapter.

Do not file a reconsideration request as it is a waste of time. In fact, you can’t file one unless you have a manual penalty on your site. Google does not have the mechanisms in place to adjust how the Penguin algorithm impacts you.

After you have addressed the bad links to your site you have no choice but to wait. Once the next Penguin release arrives you will know whether you have done enough. If you are not successful in recovering during that release you have no choice but to figure out what links you failed to clean up and try again. If you are successful, then congratulations!

Manual Penalties

There are two types of penalties, algorithmic and manual. The algorithmic penalties do not involve any human component. In contrast, manual penalties do. While the details of what causes Google to perform a manual review of a web site are not always evident, there appear to be several ways that these can be triggered. Figure 9-9 illustrates the ways a manual review may be triggered in the case of a site that has problems with its link profile.



- Regular Search Results Reviews. Google maintains a large team of people that perform human reviews of search results to evaluate their quality (<http://searchenginewatch.com/article/2203630/Google-Rating-Guidelines-Adds-Page-Quality-to-Human-Reviews>). This effort is primarily intended to provide input to the search quality team at Google that they can use to help them improve their algorithms. However, it is quite possible that this process could also be used to identify individual sites for further scrutiny.

Once a review is triggered, the human reviewer uses a set of criteria to determine if a penalty is merited. Whatever the outcome of that review, it is likely that Google keeps the notes from the review in a database for later use. Google most likely keeps a rapsheet on all webmasters and their previous infractions, whether they result in a penalty or not.

Google Webmaster Tools Messages

Since April of 2012, Google has maintained a policy of sending all publishers that receive a manual penalty a message in their Webmaster Tools console describing the nature of the penalty. These messages describe the penalty in general terms, and it is up to the publisher to figure out how to resolve it. Generally the only resource that Google provides to help with this is their Webmaster Guidelines: <https://support.google.com/webmasters/answer/35769?hl=en>. If you have such a message, then the reconsideration request option in Google Webmaster Tools becomes available

Types of Manual Penalties, Other Than Link Penalties

Manual penalties come in many forms. The most well known types of penalties are link related, but you can also get a variety of other penalties as well. Some of the most well-known varieties are discussed in the following sections.

Thin Content Penalties

This penalty relates to pages that don't add enough value to users in the opinion of Google. Figure 9-10 Shows an example of a thin content message from Google in Webmaster Tools.

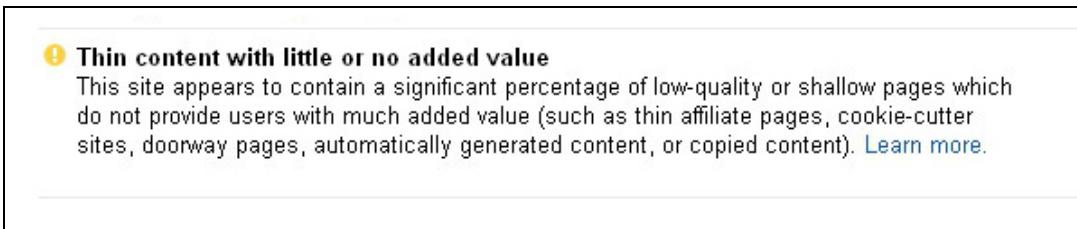


Figure 9-10 Example Thin Content Penalty Message

Unfortunately, when you receive this type of penalty Google doesn't provide any guidance on what the cause might be. They do tell you that it is a "thin content penalty", but the rest is up to you. There are three primary triggers for thin content penalties. These are:

- Pages with Little Useful Content. As the name of the penalty suggests, pages with very little content are potentially triggers for this penalty. This is especially true if there are a large number of these pages, or if there is a particular section on the site that has a significant percentage of the pages that are deemed to be thin.
- Thin Slicing. This happens to publishers who implement pages that are really designed to just garner search traffic. What these publishers often do is build pages for each potential search query a user might use, even if the variations in the content are quite small or insignificant. For example, imagine a site with information on nursing schools with different pages with the following titles:
 - Nursing schools
 - Nursing school
 - Nursing colleges

- Nursing universities
- Best nursing schools
- and so forth

Sometimes publishers do this unintentionally, by auto-generating content pages based on search queries people enter when using the search function for the web site. If you do decide to do something like this then it's critical to have a detailed review process for screening out these thin slicing variants and pick one version of the page and focusing on that.

- Doorway Pages: These are pages which appear to be generated just for monetizing users arriving from search engines. One way to recognize these types of pages is that they are usually pretty much standalone pages with little follow-on information available, and/or they are pages that are largely written for search engines and not users.

The user arriving on these pages basically have two major choices: Buy Now or Leave.

- Poor Integration into the Overall Site: Another thing to look for is whether parts of your site are not well integrated into the rest of the site. Is there a simple way for users to get to them from the home page, or from the main navigation of the site? Or, at least, a way to get to these pages relatively easily from a major section of the site?

If you have a section that appears to be isolated from the rest of the site, that could be a cause of a thin content penalty.

Once you believe you have resolved the issues, you need to submit a reconsideration request. You can read more about this in "Filing Reconsideration Requests" later on in this chapter. Once you have filed this request then you simply wait until Google provides a response. This process normally takes two to three weeks.

If you are successful then you are in good shape and just need to make sure not to overstep boundaries again in the future. Otherwise, it's back to the drawing board to see what it is that you might have missed.

Other Types of Manual Penalties

Some of the other manual penalties include:

- Cloaking and/or sneaky redirects (<https://support.google.com/webmasters/answer/2604723?hl=en>): You can get this message if Google believes you are showing different versions of pages to Googlebot than you show to users. To diagnose this, use the "Fetch and Render as Google" tool in Webmaster Tools to retrieve the page. Use the tool to load the same page in another browser window and compare the two pages.

If you see differences, invest the time and effort to figure out how to remove the differing content. You should also check for URLs that redirect and send people to pages that are not in line with what they expected to see. For example, if they click on anchor text to go read an article about something but find themselves instead on a spammy page trying to sell them something.

Another potential source of this problem is conditional redirects, where users coming from Google search, or a specific range of IP addresses, are redirected to different pages than other users.

- Hidden Text and/or Keyword Stuffing (<https://support.google.com/webmasters/answer/2604780?hl=en>): This message is generated if Google believes you are stuffing keywords into your pages for the purpose of manipulating search results. For example, if you put content on a page with a white background using a white font, so it's invisible to users, but search engines can still see it.

Another way to generate this message is to simply overdo it, and repeat your main keyword for a page over and over again in hopes of influencing search results.

- User-Generated Spam (<https://support.google.com/webmasters/answer/2721437?hl=en>): This type of penalty is applied to sites that allow user generated content (UGC) which are perceived to not be doing

a good job of quality control on that content. It's very common that sites with UGC become targets for spammers uploading low quality content with links back to their own site.

The short term fix for this is to identify and remove the spammy pages. The longer term fix is to implement a process for reviewing and screening out spammy content and prevent it from getting on your site in the first place.

- **Unnatural Links From Your Site** (<https://support.google.com/webmasters/answer/2604774?hl=en>): This is an indication that Google believes that you are selling links to third parties, or participating in link schemes, for the purposes of passing PageRank. The fix is simple – remove the links on your site that look like paid links, or add a NoFollow attribute to those links.
- **Hacked Site** (<https://support.google.com/webmasters/answer/2604775?hl=en>): You can see this either with a message in Webmaster Tools or Google may even start showing indications that your site has been hacked (and is dangerous to visit) in the search results. The most common cause for this is failing to keep up with updates to your Wordpress or other Content Management System (CMS).

Spammers take advantage of vulnerabilities in the CMS to modify your web pages, most often for the purpose of inserting links to their own sites, but sometimes for other more serious purposes such as accessing credit card information or other personally identifiable information.

To resolve the problem you will need to determine the way in which your site has been hacked. If you don't have technical staff working for you, you may need to get help to detect and repair the problem. To minimize your exposure going forward, always keep your CMS on the very latest version possible.

- **Pure Spam** (<https://support.google.com/webmasters/answer/2604777?hl=en>): Google will give you this message in Webmaster Tools if they believe that your site is using very aggressive spam techniques. This can include things such as automatically generated gibberish or other tactics that appear to have little to do with trying to add value to users.

If you get this message there is a strong chance that you should simply shut down the site and start with a new one.

- **Spammy Freehosts** (<https://support.google.com/webmasters/answer/2604782?hl=en>): Even if your site is clean as a whistle, if a large percentage of the sites using your hosting company are spamming, Google may take action against all of the sites hosted there. Take care to make sure you were working with a highly reputable hosting company!

For any of these problems, you need to address the source of the complaints. When you believe you are done, follow the procedure outlined in "Filing Reconsideration Requests" outlined later in this chapter.

Links Google Does Not Like

To understand the types of links that Google does not like we begin reviewing the original Larry Page - Sergey Brin thesis titled "The Anatomy of a Large-Scale Hypertextual Web Search Engine" (<http://infolab.stanford.edu/~backrub/google.html>). At the beginning of the thesis you find the following paragraph:

"The citation (link) graph of the web is an important resource that has largely gone unused in existing web search engines. We have created maps containing as many as 518 million of these hyperlinks, a significant sample of the total. These maps allow rapid calculation of a web page's "PageRank", an objective measure of its citation importance that corresponds well with people's subjective idea of importance. Because of this correspondence, PageRank is an excellent way to prioritize the results of web keyword searches."

The concept of a "citation" is a critical one. Consider the example of an academic research paper, which might include citations similar to those shown in Figure 9-11

<p>Franklin, K. Anabaptists: A Complete History [Book]. New York: Siegfried Shaw. 2002 Aug - [Cited 2004 May 12]</p>
<p>James, R. Time and Space and Me. [Book]. San Franciso: Abington Krane. 1996 Sep - [Cited 2010 Aug 17]</p>
<p>Landry, P. Automotive Engineering III. [Book]. Boston: Helmsworth Claw. 1984 Jan - [Cited 2005 Feb 1]</p>

Figure 9-11 Example of Academic Citations

Placement in this list is normally made by the writer of the paper to acknowledge major sources they referenced during the creation of their paper. If you did a study of all the papers on a given topic area, you could fairly easily identify the most important ones, because they would have the most citations (votes) by other papers.

Generally speaking, one does not buy placement of a citation at the bottom of someone else's research paper. Nor do you barter such placements ("I will mention you in my paper if you mention me in yours"), and you certainly would not expect someone to implement some tactic to inject mentions to your paper in someone else's research paper without their knowing.

You would also not publish dozens or hundreds of poorly written papers just so you could include more mentions of your paper in them. In principle, you can't vote for yourself. Nor would you engage in the practice of uploading your paper to dozens or hundreds of sites created as repositories for such papers if you knew no one would ever see your paper there, or if such repositories contained a lot of illegitimate papers that you would not want to be associate with.

Of course, all of these things have happened on the web with links. All of these practices run counter to the way that search engines want to use links, as they are counting on the links they find being ones that were earned by merit.

This means that they don't want you to purchase links for the purpose of influencing search engine rankings. You can buy ads, of course, there is nothing wrong with that, but they would prefer those ads to have the "NoFollow" attribute on the links to tell Google not to count them.

Additionally, pure barter links are valued less, or ignored too. From 2000 to 2005 it was quite popular to send people emails that offered to link to them if they linked to you, on the premise that this helped with search engine rankings. Of course, these types of links are not real citations either.

User generated content sites, such as social media sites are a class of site from which Google will not place any value on the links as well. Any place where people can go and link to themselves is something that they will simply discount, or even potentially punish if they can detect patterns of abusive behavior.

Here is a list of some of the types of links that Google may consider less valuable, or not valuable at all:

- Article Directories. These are sites that allow you to upload an article to them, usually with little or no editorial review. The articles can contain links back to your site, so the simple act of uploading the article results in a link. The problem is that this is a form of voting for yourself, and Google started punishing sites that actively obtained links from article directories with the Penguin 1.0 release on April 24, 2012.

- Cheap Directories. Many directories have sprung up all over the web that really only exist to collect fees from as many sites as possible. These types of directories have little or no editorial review and the owner's only concern is to collect as many listing fees as possible. There are a few quality directories such as the Yahoo directory, Business.com, and a few others. These are discussed in more detail in Chapter 7. Note: The above comments are not meant to apply to Local Business Directories for which the dynamics are quite different. These are discussed more in Chapter 10.
- From Countries Where You Don't Do Business. If your company only does business in Brazil, there is no reason you should have large numbers of links from Poland and Russia. There is not much you can do if people choose to give you links you did not ask for, but there is certainly no reason for you to proactively engage in activities that would result in your getting links from such countries.
- From Foreign Sites with a Link in a Different Language. Some aggressive SEO professionals actively pursue getting links from nearly anywhere. As shown in Figure 9-12, there is no reason to have a "Refinance Your Home Today" link on a page where the rest of the text is in Chinese.

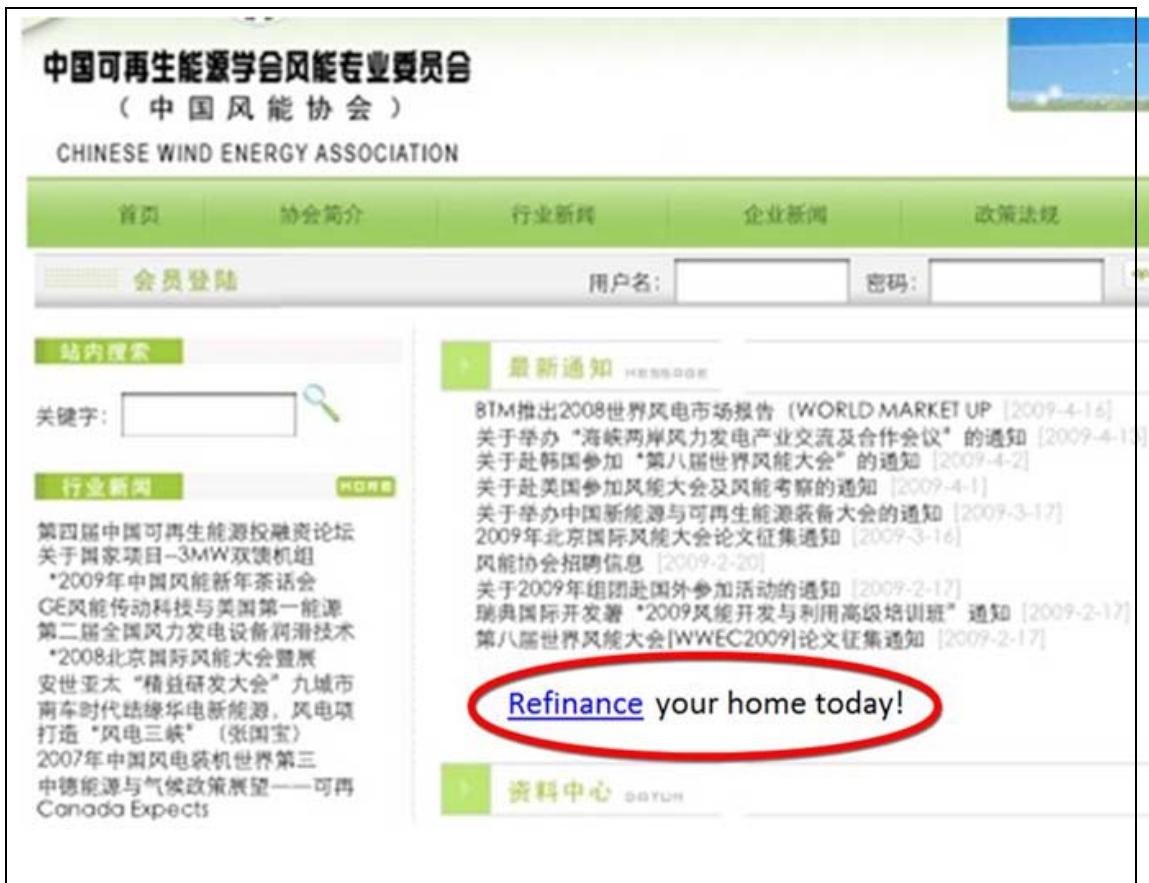


Figure 9-12 Foreign Language Mismatch

- Comment Spam. Another popular technique in the past was to drop links in comments on forums and blog posts. This practice became much less valuable ever since Google introduced the NoFollow attribute, but aggressive spammers still pursue it. In fact they make use of bots that drop comments on an automated basis on blog posts and forums all over the web. They may post one million or more comments this way, and it only 1/10 of one percent of those links are not NoFollowed, it still nets them 1,000 links
- Guest Post Spam. These are generally poorly written guest posts that add little value that have been written just to get a link back to your own site. Consider the example of Figure 9-13 where the author was looking to get a link with the anchor text back to their site with the anchor text "buy cars". They could not even take the time to work that phrase into a single sentence!

St. Mary's

St. Mary's is a town just under an hour's drive from Bicheno, or an hour and a half's drive from Launceston. You can challenge yourself to a climb to the St. Patrick's Head, or the South Sister's Peak, or a relaxing drive in the Elephant Pass - all of which give spectacular views. In town, have a browse through all the craft galleries, bakeries and shops that this small town has to offer.

In terms of getting yourself a vehicle, scope out rental options, or if you're going for an extended holiday, you can consider looking at second hand cars to [buy](#). Cars should be reliable, safe and road worthy of course, so don't buy something that's run down just because it's cheap.

Once you've got a set of wheels, it's time to hit the road!

Figure 9-13 Guest Post Spam

- Guest Posts Not Related to Your Site. This is a type of guest post spam where the article written does not really relate to your site. If you sell used cars, you should not expect that Google will see any value in a guest post you may write about lacrosse equipment that links back to your site. There is no relevance.
- In-Context Guest Post Links. Another form of guest posting that Google frowns on is posts that include links in the body of the article back to you, particularly if those links are keyword-rich, and if they don't add a lot of value to the post itself. Figure 9-14 shows a fictional example of what this might look like.

How to Get the Best Travel Deals

- Book early. People who wait until the last minute almost always pay higher prices.
- Use the Internet. Websites such as [travel-o-matic](#) can help you save hundreds of dollars on air travel and hotels.
- Try a vacation rental rather than a hotel.
- Look for [travel package deals](#) to help cut down the overall cost of both airfare and hotels combined.

This guest post was provided by Sarah Smith an avid traveler who also blogs for www.travel-o-matic.com.

Figure 9-14 Embedded Keyword-Rich Anchor Text Links

- Advertorials. This is a form of guest post that is written like it's an ad. Given the structure, it's highly likely that the site posting it was influenced to do so in some manner. If you are going to include guest posting as part of your strategy, focus on sites that don't permit these types of guest posts.
- Widgets. One tactic that became quite popular is building useful or interesting tools (widgets), and allowing third party web sites to publish them on their own sites. These normally contained a link back to the site of the creator of the widget. If highly relevant, there is nothing wrong with this idea in principle, but the problem is that the tactic was abused by SEOs, resulting in Google wanting to discount many of these types of links.
- Infographics. This is another area that could in theory be acceptable, but which was greatly abused by SEOs. It is not clear what Google does with these links at this point, but you should only create infographics if they are highly relevant, highly valuable, and of course accurate.
- Misleading Anchor Text. This is a more subtle issue. Imagine an example where the anchor text of a link says "information about golf courses" but the page receiving the link was about tennis rackets. This is not a good experience for users, and is not something that search engines will like.
- Sites with Malware. Of course, Google looks to discount these types of links. The sites containing them are very bad for users, and hence any link from them is of no value, or potentially harmful.
- Footer Links. Once again, there is nothing inherently wrong with a link from the footer of someone's web page, but as these are less likely to be clicked on, or looked at, by users, Google may discount the value of these links. For more on this you can read Bill Slawski's article about Google's "reasonable surfer" patent (<http://www.seobythesea.com/2010/05/googles-reasonable-surfer-how-the-value-of-a-link-may-differ-based-upon-link-and-document-features-and-user-data/>).
- In a List with Unrelated Links. This can be a signature of a purchased link. Imagine you find a link to your "Travel Australia" web site mixed in a list of links with an online casino site, a mortgage lead generation site, and a lottery ticket site. This does not look good to Google.
- Poor Quality Sites. The links that have the most value are the ones that come from very high quality sites that show substantial evidence of strong editorial control. Conversely, as quality drops, editorial control tends to as well, and Google may not count these links at all.
- Press Releases. It used to be quite popular to put out lots of press releases, complete with keyword-rich text links back to your site. Of course, this is a form of voting for yourself, and this is not the way that press releases should be used to promote your site. The best use is to issue high value announcements that media and bloggers would be interested in seeing. As shown in Figure 9-15, a much better way to use press releases is to get your news in front of media people and bloggers, and hope that it's interesting enough that they will write about you or share your news on social media.

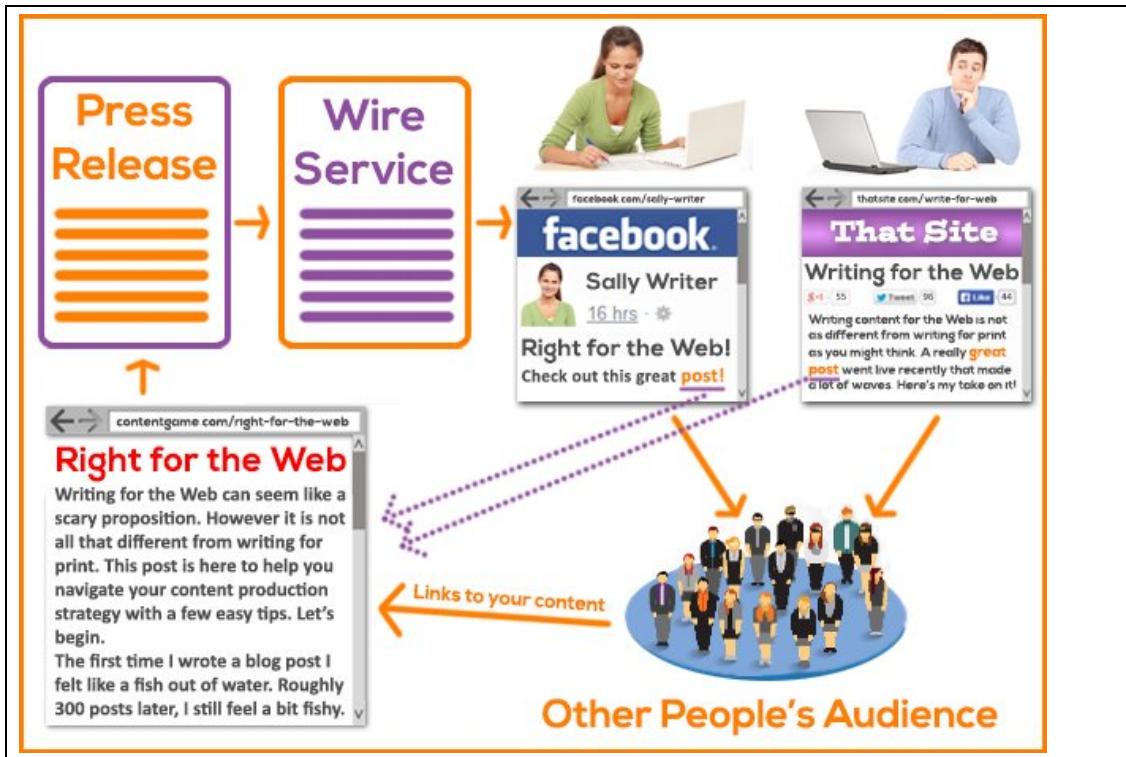


Figure 9-15 The Right Way to Use Press Releases

- Bookmark Sites. There are many quality sites for saving interesting links for your own benefit, such as Delicious, Evernote, Diigo, and many others. However, as these are user generated content sites, the links from them have no value in ranking your site because they are NoFollowed.

Not all of the types of links in the above list will necessarily result in your site being penalized, but they are all examples of links that Google may not want to count.

Partial Link Penalties

Another possible manual penalty you can receive is a partial penalty. These are sometimes called an “Impacts Links” penalty as that is part of the message you get from Google as shown in Figure 9-16. These penalties indicate that there are a very small number, or one, of your pages that Google has flagged for bad linking behavior.

As a consequence of this, normally the rankings for those particular pages are lowered, and only traffic to those pages suffer.

Partial matches

▼ Some manual actions apply to specific pages, sections, or links

Reason	Affects
Unnatural links to your site—impacts links Google has detected a pattern of unnatural artificial, deceptive, or manipulative links pointing to pages on this site. Some links may be outside of the webmaster's control, so for this incident we are taking targeted action on the unnatural links instead of on the site's ranking as a whole. Learn more.	Some incoming links

[REQUEST A REVIEW](#)

Figure 9-16 Example Partial Link Penalty Message

Unfortunately, Google does not tell you which of your pages is responsible for the penalty, you have to determine that for yourself. This penalty is normally caused by getting too many questionable or bad links to pages other than your home page.

The cause is often a link building campaign focused on bringing up the rankings and search traffic to specific money pages on your site. One of the more common problems is too many links with keyword-rich anchor text pointing to those pages, but other types of bad links can be involved as well. The steps to recovering from this type of penalty are:

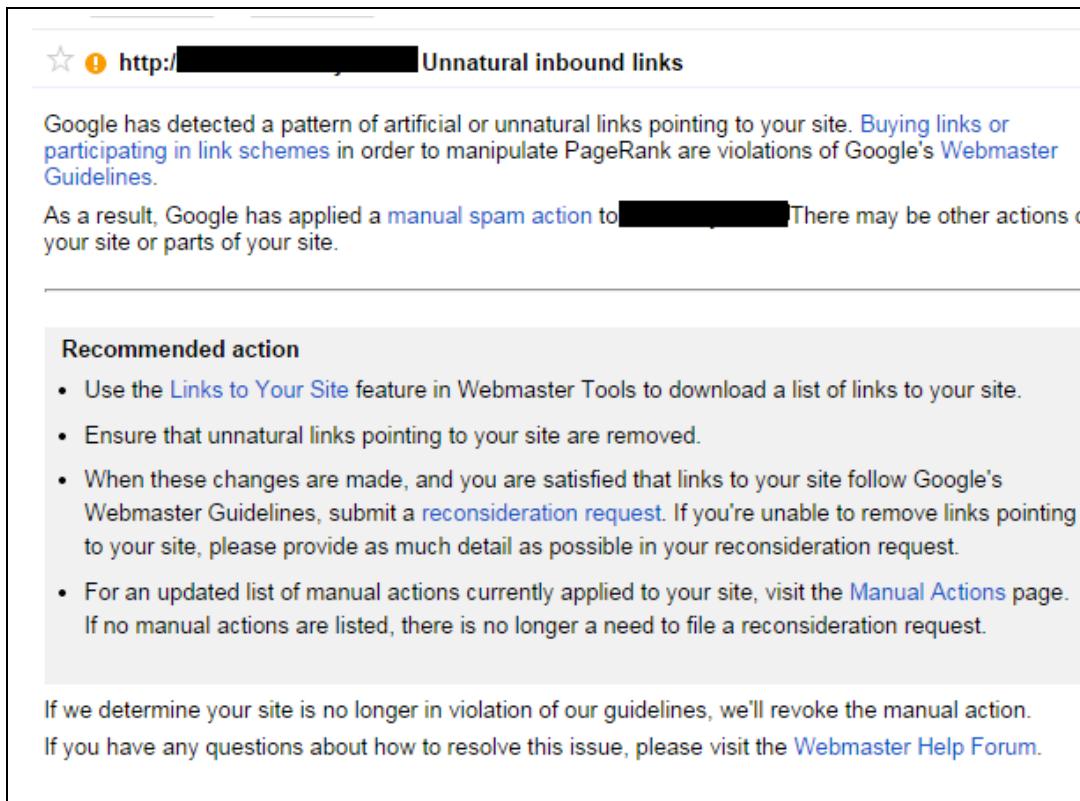
1. Pull together a complete set of your links as described in “Sources of Data”.
2. Look for pages on your site, other than the home page, that have the most links.
3. Examine these pages for bad links as described in “Links Google Does Not Like”.
4. Use the process described in “Link Cleanup Process” later on in this chapter to deal with the bad links.
5. The last step is to submit a reconsideration request as described in “Filing Reconsideration Requests”.

Once you have sent in the reconsideration request, the only thing you can do is wait. The process normally takes two to three weeks before you get a response. Google will let you know if you have succeeded and tell you that they have removed the penalty, or they will tell you that you have failed, in which case you have to take a deeper look at your links and figure out what you missed in your last attempt.

Sitewide Link Penalties

Manual link penalties can also be applied on a site wide basis. This usually means more than a few pages are involved, and may well also involve the home page of the site. With this type of penalty, rankings are lowered for the publisher on a site wide basis.

As a consequence of that, the amount of lost traffic is normally far more than it is for a partial link penalty. Figure 9-17 shows an example of a site wide link penalty message.



The screenshot shows a Google Webmaster Tools message. At the top, there's a star icon and a warning sign icon, followed by the URL [http://\[REDACTED\]](http://[REDACTED]) and the title "Unnatural inbound links". The message text reads: "Google has detected a pattern of artificial or unnatural links pointing to your site. [Buying links or participating in link schemes](#) in order to manipulate PageRank are violations of Google's [Webmaster Guidelines](#). As a result, Google has applied a [manual spam action](#) to [REDACTED]. There may be other actions on your site or parts of your site." Below this, a "Recommended action" section lists five steps: 1. Use the [Links to Your Site](#) feature in Webmaster Tools to download a list of links to your site. 2. Ensure that unnatural links pointing to your site are removed. 3. When these changes are made, and you are satisfied that links to your site follow Google's Webmaster Guidelines, submit a [reconsideration request](#). If you're unable to remove links pointing to your site, please provide as much detail as possible in your reconsideration request. 4. For an updated list of manual actions currently applied to your site, visit the [Manual Actions](#) page. If no manual actions are listed, there is no longer a need to file a reconsideration request. At the bottom, a note says: "If we determine your site is no longer in violation of our guidelines, we'll revoke the manual action. If you have any questions about how to resolve this issue, please visit the [Webmaster Help Forum](#)".

Figure 9-17 Example Sitewide Link Penalty Message

The steps to recovering from this type of penalty are:

1. Pull together a complete set of your links as described in “Sources of Data”.
2. Examine this list for bad links as described in “Links Google Does Not Like”.
3. Use the process described in “Link Cleanup Process” later on in this chapter to deal with the bad links.
4. The last step is to submit a reconsideration request as described in “Filing Reconsideration Requests”.

Once you have sent in the reconsideration request, the only thing you can do is wait. The process normally takes two to three weeks before you get a response. Google will let you know if you have succeeded and tell you that they have removed the penalty, or they will tell you that you have failed, in which case you have to take a deeper look at your links and figure out what you missed in your last attempt.

Link Cleanup Process

The first part of a cleanup process is to establish the right mindset. As you review your backlink profile you need to consider how Google looks at your links. Here are some rules of thumb to help you determine whether or not a link has real value:

- Would you want that link if Google and Bing did not exist?
- Would you proudly show it to a prospective customer right before they are ready to buy?
- Was the link given to you as a genuine endorsement?

As you review your backlinks you may find yourself at times offering up justifications for something being a decent link. This is usually a good indicator that it's NOT a good link at all. High quality links require no justification – it's obvious that they are good links.

Sources of Data

Google provides a list of external links in the Webmaster Tools account for your site. Figure 9-18 shows the steps to navigating to that report.

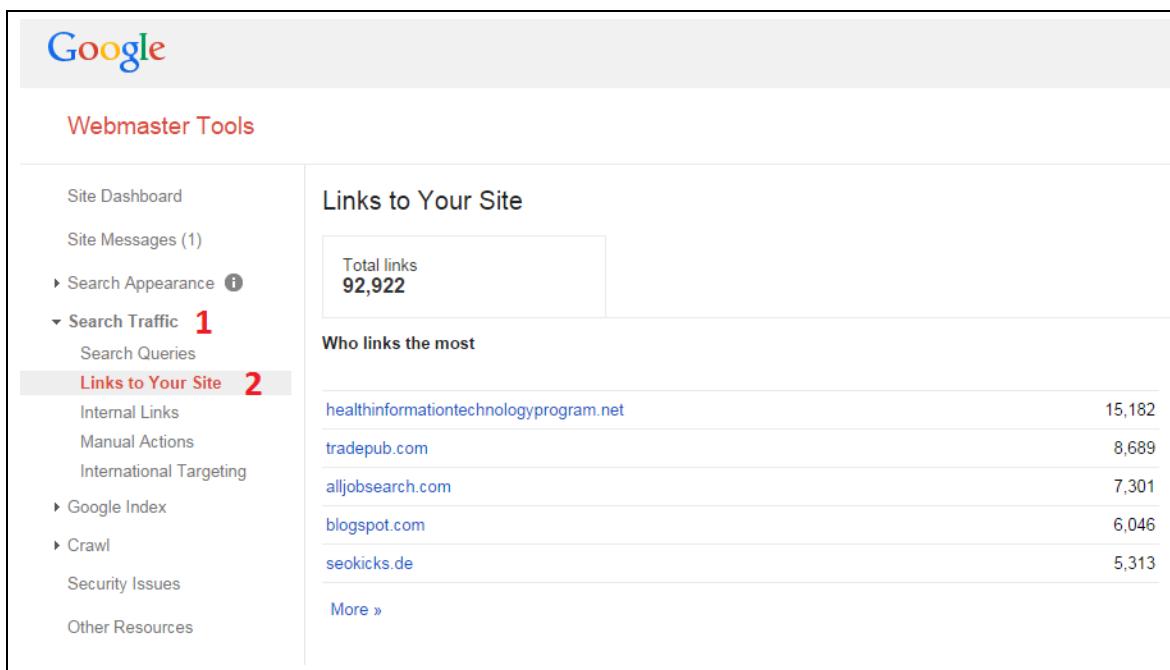


Figure 9-18. Steps to Obtain Webmaster Tools List of Links

The problem with this list is that it tends to be incomplete. The authors recommend that you also pull links from several other sources. Some of the best additional sources include Open Site Explorer (<http://www.opensiteexplorer.org/>), Majestic SEO (<http://www.majesticseo.com>), ahrefs (<https://ahrefs.com/>), and Link Research Tools (<http://www.linkresearchtools.com>).

The combination of data from all of these tools will show a more complete list of links. Of course, they will also have a lot of overlap in what they show, so make sure to de-dupe the list.

However, even the combination of all these sources is not comprehensive. Google only shares a portion of the links they are aware of in Webmaster Tools. The other link sources are reliant on the crawls of their individual companies, and crawling the entire web is a big job and they simply do not have the resources that Google or Bing have.

Using Tools

There are tools available designed to help speed up link removal by automating the process of identifying bad links. These include Remove'em (<http://www.removeem.com>) and Link Detox (<http://www.linkdetox.com>). These tools can potentially help you identify some of your bad links. However, it is a good idea to not rely solely on these tools to do the job for you.

They each have their own algorithms for identifying problem links and this can save you time in doing a full evaluation of all their links. However keep in mind that Google has spent more than 15 years developing its algorithms for evaluating links and it's a core part of their business to evaluate these effectively, including their algorithms for detecting link spam.

Third party tools won't be as sophisticated as Google's algorithm. They can detect some of the bad links, but necessarily all of the ones you will need to address. You should plan on evaluating all of the links not only the sites labeled as toxic, but also suspicious and even innocuous. Use your own judgment, and don't merely rely on the tools to decide for you what is good or bad.

The Link Disavow Tool

Google provides a tool to allow you to disavow links (<https://www.google.com/webmasters/tools/disavow-links-main?pli=1>). Functionally what this does is tell Google that you no longer wish to receive any PageRank (or other) benefit from such links. This provides a method for eliminating the negative impact of bad links pointing to your site. Figure 9-19 shows what the opening screen to the tool looks like.

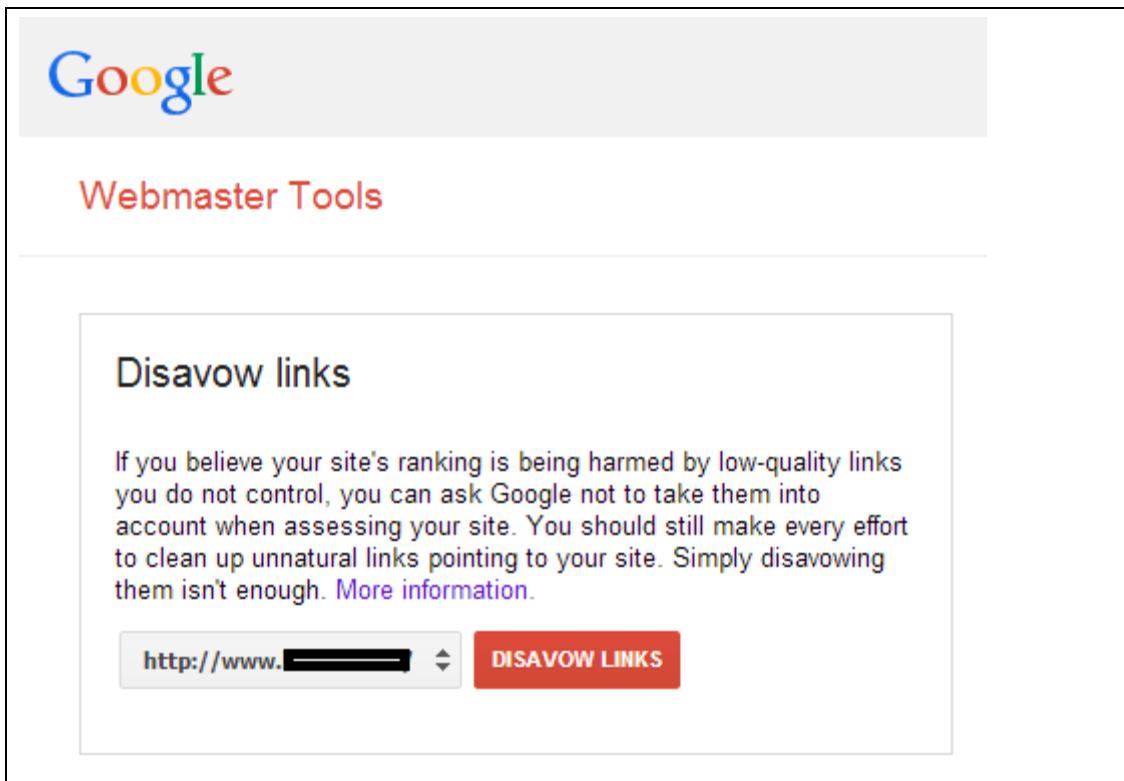


Figure 9-19 Google's Disavow Tool Opening Screen

Note that Google includes the following text: "You should still make every effort to clean up unnatural links pointing to your site. Simply disavowing them isn't enough." This is good advice as Google employees that review reconsideration requests like to see that you have invested in time to getting the bad links pointing to your site removed.

Once you selected a site (currently blanked out in Figure 9-19) and click on the Disavow Links Button you get brought to another screen that includes the following warning:

"This is an advanced feature and should only be used with caution. If used incorrectly, this feature can potentially harm your site's performance in Google's search results. We recommend that you only disavow backlinks if you believe that there are a considerable number of spammy, artificial, or low-quality links pointing to your site, and if you are confident that the links are causing issues for you."

On this screen you need to click on Disavow Links again and you get brought to a third and final screen as shown in Figure 9-20.

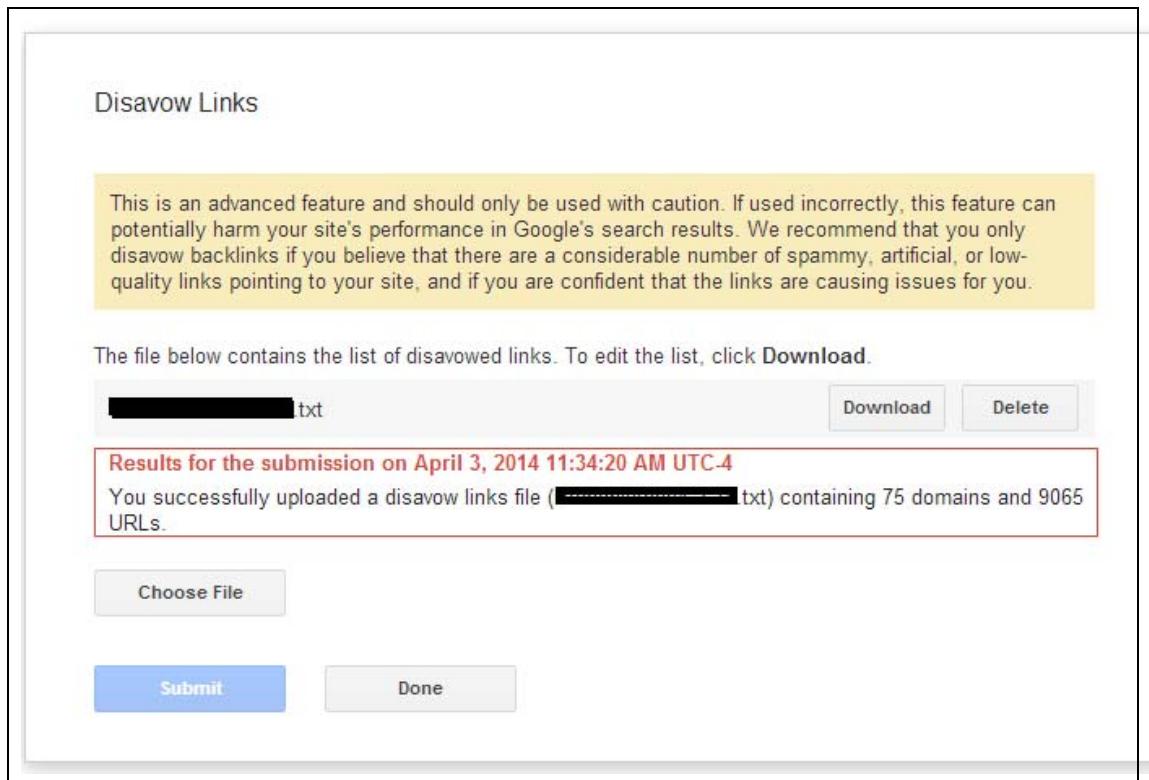


Figure 9-20 Google's Disavow Tool Final Screen

The sample screen shown in the screen shot shows that the current disavow file for this site (the name is blacked out) is disavowing 75 domains and 9065 URLs.

One important tip: Considering that the link data you have is incomplete (why is described in “Sources of Data) it is best practice to disavow entire domains. In other words, if you see one bad link coming to you from a domain it is certainly possible that there are other bad links coming to you from that same domain, and that some of these bad links are not in the data available to you.

An example would be a guest post that you want to disavow. Perhaps you have done only one guest post on that site, but the post will also appear in category pages and date-based archives on that blog. If you disavow only the post page you still have many other bad links from that site that Google has found.

In the example shown in Figure 9-20 it is quite likely that this publisher has not solved their problem and that many (if not all) of the disavowed URLs should be made into full disavowed domains. It is usually the safest course of action.

Refer to the Google help page for more specifics on formatting the disavow file (<https://support.google.com/webmasters/answer/2648487?hl=en>).

Link Removal Process Outline

The most important part of the process is recognizing the need to be comprehensive. Losing a lot of your traffic is scary, and being impatient in natural. If there is a manual link penalty on your site you will be anxious to send in your reconsideration request, but as soon as you send that request in, you can do nothing but wait.

If you don't do enough to remove bad links, Google will reject your Reconsideration Request, and you have to go through the whole process again. If you end up filing a few reconsideration requests without being successful, Google may send you a message telling you to pause for a while. Make a point of being very aggressive in removing and disavowing links, and don't try to save a lot of marginal ones. This almost always makes the process go faster in the end. In addition, those somewhat questionable links that you might be trying to save often are not helping you much anyway.

If you have been hit by the Penguin algorithm, then updates to that happen only once or twice per year. Failing to recover during any Penguin update means waiting another 6 to 12 months before you have another chance (Note: Google had indicated that they intend to speed up the Penguin update process in the future).

With this in mind, we also want to be able to get through the process as quickly as possible. Figure 9-21 provides a visual outline of what the process should look like.

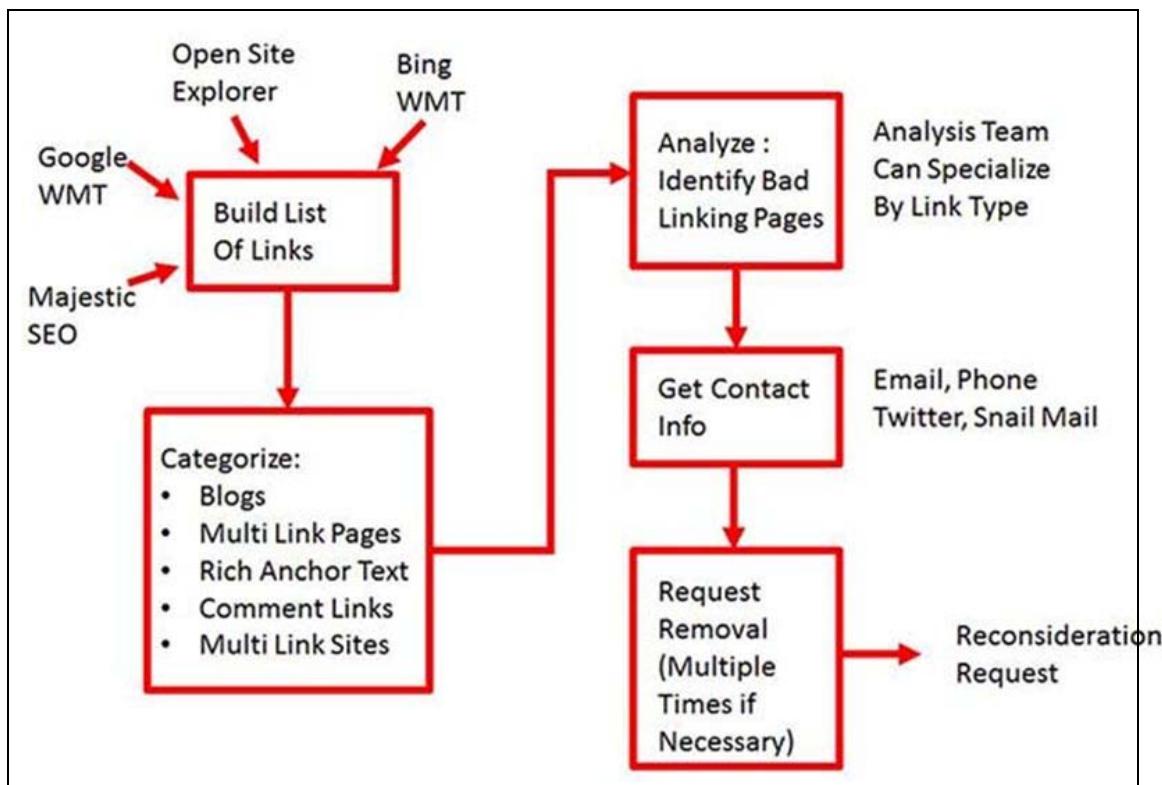


Figure 9-21 Process for removing links

The process of pre-categorizing the links is quite helpful in speeding up the process. For example, you can identify many of the blogs simply by using the Excel filter function and filtering on “blog”. This will allow a more rapid review of the links for problems. Tools such as Remove’em (<http://www.removeem.com/>) and Link Detox (<http://www.linkdetox.com/>) will do this for you as well.

This is especially helpful if you know you have been running an aggressive guest posting campaign, or worse, paying for guest post placements. Some additional tips include:

- You do not need to worry about links that are marked as NoFollow.
- Links from sites with very low home page PageRank probably are adding little to no value to your site anyway.
- Links from very low relevance pages also are likely to not be adding much value either.

In addition, the process of obtaining contact information and contacting sites and requesting that they remove links can be quite helpful as well. Google likes to see the effort to clean up those bad links.

Remember that reconsideration requests are reviewed by members of the webspam team at Google. This introduces a human element which you can’t ignore. The members of this team make their living dealing with sites that have violated Google’s guideline and you are one of them. Recall also that as we said in “Sources of Data” that even when you use all the available sources of link data that the information you have is incomplete.

This means that it's likely that you will not have removed all the bad links when you file your reconsideration request, even if you are very aggressive in your review process, simply because you don't have all the data. Showing the reviewer the good faith effort to remove some of the bad links is very helpful, and can impact their evaluation of the process.

However, there is no need to send link removal requests to everyone in sight. For example, don't send them to people where the link to you is marked with NoFollow.

Once the process is complete, if you have received a manual penalty, you are ready to file a reconsideration request. If you have been hit by Penguin all you can do is wait, as reconsideration requests won't help you with Penguin. The only way to recover from Penguin is to wait for the next algorithm update.

Filing Reconsideration Requests

The first thing to realize is that a person will review your reconsideration request, and that person will likely review large numbers of them every single day. Complaining about what has happened to your business, or getting aggressive with the reviewer is not going to help your cause at all.

The best path is to be short and to the point.

- Briefly define the nature of the problem. Include some stats if possible.
- You can explain what went wrong. For example, if you were ignorant of the rules, just say so, and tell them that you now understand. Or, if you have a rogue SEO firm do bad work for you, say that.
- Explain what you did to clean it up.
 - If you had a link penalty, let them know how many links you were able to get removed.
 - Or, if you did something extraordinary, such as removing and/or disavowing all of your links from the past year, tell them that. Statement actions, such as this, can have a strong impact and improve your chances of success.
- Clearly state that you intent to abide by Webmaster Guidelines going forward.

As already noted, keep your reconsideration short. Briefly cover the main points and then send it in using the Webmaster Tools account associated with the site that received the penalty. In fact, you can't send it from an account without a manual penalty

Expected Reconsideration Request Timeline

Once you have filed the request you now get to wait. The good news is that you generally get a response in one to three weeks. Hopefully, you will be successful! If not, then you have to go back to the beginning of the process to figure out what you missed.

Conclusion

Traffic losses due to manual penalties or algorithmic updates can be devastating to your business. One defense against them is to reduce your overall dependence on Google traffic, but for nearly any business traffic from Google is one of the largest opportunities available to them.

For that reason, it is critical you conduct your SEO efforts in a way that is in line with what Google really wants to see from publishers – create compelling web sites and then promote them effectively in a manner that complies with Google's Webmaster Guidelines (<https://support.google.com/webmasters/answer/35769?hl=en>).

10

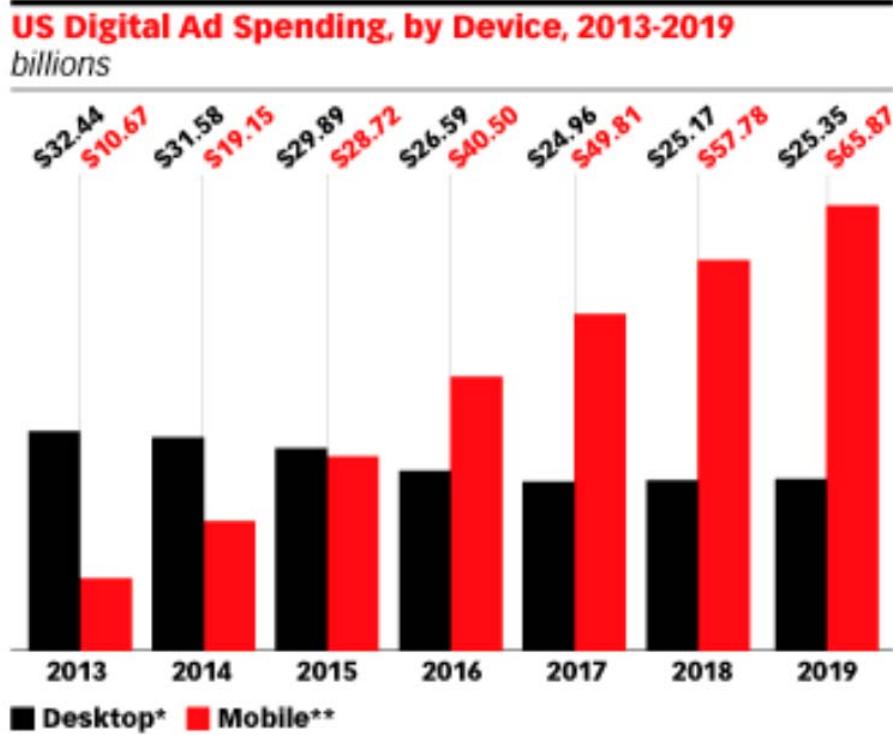
Mobile, Local, & Vertical SEO

Mobile, local, and vertical SEO have grown to become SEO specializations in and of themselves – requiring a tremendous amount of dedicated effort to fully maximize the opportunities they provide as the web ecosystem evolves. These areas of SEO require ongoing attention and resource dedication in order to be effective, as the optimization opportunities – and therefore the SEO strategies and tactics for each – continue to rapidly evolve. In this chapter we will address these areas of search engine optimization at a high level – and strongly recommend that you put ongoing resources into these powerful areas of SEO opportunity.

The Mobile Landscape

Mobile SEO will be absolutely vital to the success of your business's overall search engine optimization strategy - if you have a web-based business, brand, or organization. But before we can examine the ins and outs of SEO for mobile, we need to first understand mobile search in the general context of mobile as a medium – one that is taking search market share and ad spend from desktop at a rapid pace.

As mentioned in Chapter 1, eMarketer published a study predicting that mobile ad spend would represent 72.2% of total US digital ad spend (Figure 10-1).



Note: *includes spending primarily on desktop-based ads; **includes classifieds, display (banners and other, rich media and video), email, lead generation, messaging-based advertising and search; ad spending on tablets is included

Source: eMarketer, March 2015

186555

www.emarketer.com

[[figs/print/1001.png]]

Figure 10-1. “Mega-SERP” infographic (Moz.com)

The iPhone helped drive the growth of this market, because of its intuitive user interface and larger screen size (when compared to feature phones) and Android accelerated with its increasing smartphone market dominance (<http://www.comscore.com/Insights/Market-Rankings/comScore-Reports-May-2014-U.S.-Smartphone-Subscriber-Market-Share>).

The improved mobile web user interfaces and larger screen sizes on smartphones have been a big factor in the growth of mobile search. Having a specific mobile site experience for users is important, as mobile searchers are even more impatient and a differing intent than your typical desktop searcher. If you offer a desktop experience to a mobile user and she finds a competitor that offers a tailored mobile experience, you may have just lost the attention of a potential customer – and the search engines know this.

Rapid growth in mobile is expected to continue. On June 13th, 2014 Matt Cutts told the audience at SMX Advanced in Seattle that before the end of 2014 that Google would receive more search queries per day from mobile devices than they would from desktop devices, and their mobile-friendly algorithm shift followed less than a year later. Today, a Google search from any mobile device shows us that the mobile landscape has evolved into an entirely unique, mobile-first search environment.

SEO for Mobile

The mobile tipping point in organic search has been reached. In April 2015, Google rolled out their mobile-friendly algorithm – a definitive, “yes or no” response ranking factor that determines whether a site (or a specific page from a site) is mobile friendly or not, and then incorporates this information into their ranking decision for a mobile query (<http://googlewebmastercentral.blogspot.com/2015/02/finding-more-mobile-friendly-search.html>). (Note that as of this book’s printing, the mobile-friendly algorithm was not being applied to Google News content for mobile queries returning News results, or for Google News-specific searches).

Google’s message with this algorithm update, which they warned would have more of an impact on their search results than either Panda or Penguin (learn more about Panda and Penguin in Chapter 9), was clear: if you want your site’s content to show for mobile queries, you need to build your site for the mobile user experience – which at the most basic level means that your mobile site must be:

- *Fast* – pages should ideally load in less than one second
- *Functional* – page content displays and functions properly in mobile browsers (no CSS, JavaScript, image, or other resource blocking)
- *Finger-friendly* – tap targets (buttons, links, form fields) should be large enough and properly spaced for small touchscreen use
- *Free from redirects and errors* – mobile version URL requests should all return 200 (OK) status codes – no 301/302 redirects or errors, if possible

Responsive Design vs. Dynamic Serving vs. Separate URLs

There are various ways to implement a mobile site version that serves to deliver user-friendly site functionality to a mobile browser, with three being accepted by Google (<https://developers.google.com/webmasters/mobile-sites/mobile-seo/configurations/>). Note that these fall into two general categories for SEO purposes, “same URLs” and “Separate URLs”:

1. Responsive Design, or RWD (Same URLs):

Your site utilizes the same URLs for desktop and mobile, and your site’s server always sends the same HTML to all devices, and stylesheets (CSS) are used to alter the rendering of the page on the device. Google can generally detect responsive design as long as the site’s assets (including CSS) are crawlable.

2. Dynamic Serving, aka Adaptive Design (Same URL):

Your site utilizes the same URLs for desktop and mobile, but your site’s server responds with different HTML and CSS depending on the user agent requesting the page.

3. Separate URLs:

Your site utilizes different URLs for the desktop and mobile (and perhaps even tablet) versions of your site’s content (utilizing a subdomain such as m.site.com, for example), with device-optimized content served at each URL version.

Same URL Approaches (Responsive Design & Dynamic Serving)

As you can see the first two options retain the same URLs across devices (Responsive Design and Dynamic Serving), while the third option (Separate URLs) creates two entirely separate websites for the search engines to crawl, index, and rank – a major factor to consider before developing your mobile site.

A significant advantage offered by using the same URLs (in Responsive Design and Dynamic Serving) is that the mobile site effectively inherits the SEO characteristics of the desktop site. That is, the link profile and social signals are automatically associated with the mobile version of the site (which is not the case with a mobile subdomain). The “same URLs” approach is considered to be the best approach if you are

intending to support smartphone-type devices only, and do not plan to support feature phones (e.g., traditional flip phones).

The most cost-effective of these two “same URLs” approaches is usually Responsive Design – as only one site “version” is being created and maintained. For this reason, Responsive Design has seen rapid adoption in the mobile space.

In contrast to Responsive Design, Dynamic Serving (whereby entirely different site versions are designed and coded with device-specific optimization and user-experience in mind), renders different HTML for the same URL request, based on the user agent making the request. While this approach is more budget, development, and maintenance intensive, it often can allow for faster load times, and a more optimized user experience specifically tailored to device type.

Separate URLs

You may wish to offer a mobile-specific URL for marketing and promotion purposes, or to allow access to the mobile version of your site regardless of the user agent making the request. In particular, if you plan to support both smartphones and feature phones, having a mobile subdomain allows you some additional flexibility in dealing with different screen sizes, or if your site is so large that you can't provide a separate mobile experience for each page.

Some well-known examples of such URLs include <http://m.facebook.com>, <http://mobile.weather.gov>, <http://www.hotels.com/mobile/>, and <http://en.m.wikipedia.org>. The benefit of serving up a distinct mobile version instead of serving your standard version with a mobile-friendly stylesheet (CSS) is that you can trim the file size of the HTML sent to the mobile device, eliminating code that the device won't be able to run, which improves page load time and makes the page look better to mobile bots. Google provides more information about Separate URLs for mobile here: <https://developers.google.com/webmasters/mobile-sites/mobile-seo/configurations/separate-urls>.

Mobile User Agents

Google employs a different spiders for the mobile web, “Googlebot-Mobile” for feature phones and the standard Googlebot with an iPhone user agent for smartphones (<http://googlewebmastercentral.blogspot.com/2014/01/a-new-googlebot-user-agent-for-crawling.html>) but Bing uses the same crawler to crawl mobile content as it does desktop content, called Bingbot (<http://blogs.bing.com/webmaster/2014/11/03/meet-our-mobile-bots/>), and utilizes crawler variants in the following format:

Mozilla/5.0 + (Mobile Device) + Mobile Engine + Mobile Browser + bingbot/BingPreview/[version]

Both search engines use *user-agent detection* to determine if a searcher is on a mobile browser or a traditional desktop browser, and will show feature phone users a mobile experience. At the beginning of each browsing session, the user's hardware and browser combination will communicate a unique identifier to the website's web server, known as the *user agent*. This string identifies the nature of the hardware and software making the request.

The best practice for mobile site user experience, and hence for mobile SEO, is to detect the incoming user agent and show the user the version of your site specific to the user's device. If you are using the same URLs for both your desktop and mobile sites (as in Responsive Design and Dynamic Serving), you simply show them the alternative version of your content. If you are using Separate URLs, you will want to redirect mobile users to that mobile URL.

A useful tool for server-side user agent detection is MobileDetect.net, which provides an open source PHP class (Mobile_Detect) for user agent detection and is supported by third party plugins for Wordpress, Drupal, Joomla, Magento, Concrete5, and other content management systems (<https://github.com/serbanghita/Mobile-Detect#3rd-party-modules--submit-new>).

General SEO Guidelines for Mobile

- Review search engine recommendations for Mobile SEO (and check for updates regularly!):
Google:
<https://developers.google.com/webmasters/mobile-sites/mobile-seo/index>.
Bing: (though Bing generally does not provide as much webmaster/SEO support as Google):
<http://blogs.bing.com/webmaster/2014/11/03/meet-our-mobile-bots/>
- Ensure critical site files are accessible by the search engines (CSS, JavaScript, images) – many of these files signal to the search engines that the page is built to display and work well on a mobile browser (telling the search engine that the page is “mobile-friendly”). Check to make sure your site does not block access to any of these files via robots.txt.
- If your site is dynamically serving desktop and mobile content (for the same URL request) based on the user agent requesting the page, use the Vary HTTP header to signal your changes depending on the user agent: <https://developers.google.com/webmasters/mobile-sites/mobile-seo/configurations/dynamic-serving>
- If your site is hosting desktop and mobile content on separate (unique) URLs, add annotations by utilizing the rel=”alternate” and rel=”canonical” tags in your desktop/mobile URLs, respectively: <https://developers.google.com/webmasters/mobile-sites/mobile-seo/configurations/separate-urls> (this can also eliminate any potential duplicate content concerns between the two site versions).
- Publish only mobile-friendly content (Flash movies or Flash-based websites, for example, are not compatible with mobile!)
- Include a Mobile Sitemap according to updated search engine guidelines. As of 2015 the most recent Mobile Sitemap format utilized by Google includes a specific tag that must be included (<mobile:mobile/>) for mobile URLs to be crawled. Refer to Google’s documentation on Mobile Sitemaps here: <https://support.google.com/webmasters/answer/34648>

Mobile Site Speed

Site speed is a crucial component of mobile SEO success, a point that has been driven home by Google for a number of years now. The following recommendations hold true regardless of whether you choose a same URLs or Separate URLs approach, and suggestions for your specific site can be found by utilizing the following tools:

Google Page Speed Insights (<https://developers.google.com/speed/pagespeed/insights/>)

Mobi Test from Akamai (<http://mobitest.akamai.com/m/index.cgi>)

WebPageTest (<http://www.webpagetest.org/>)

W3C’s Mobile OK Checker (<http://validator.w3.org/mobile/>).

Additional Mobile SEO Resources

- <http://marketingland.com/essential-mobile-seo-resources-89308>
- <http://moz.com/blog/mobile-site-audit-checklist>
- <http://searchengineland.com/library/search-engine-optimization/seo-mobile-search>
- <https://developers.google.com/webmasters/mobile-sites/mobile-seo/>

A special thanks to Michael Martin (<http://www.mobilemartin.com>) for his contributions to portions of the SEO for Mobile section of this chapter.

App SEO: Deep Linking & Indexing for Mobile Search

Based on the evolution of deep linking and indexation for mobile app content, we predict there will be significant opportunities for SEO exposure of mobile app content in the near future. App SEO for organic search exposure (different from “App Store SEO” which has generally referred to optimizing mobile apps for prominence within Apple’s App Store and Google Play) provides a massive opportunity for content discovery via the indexation and ranking of deep content residing within mobile apps – and as we have learned over the years with website content delivered via search, as deeper, more targeted content is delivered to a searcher, content (and in this instance, app) engagement increases, the time to conversion decreases, and search engine results improve.

Various technology companies are targeting the app-specific search landscape. URX (www.urx.com) has built a deep link app search API, whereby developers can integrate their API to monetize apps or websites with deep links to contextually relevant in-app content and actions in other apps. URX has built a search engine that developers use to find a deep link into an app for a specific context. Quixey (www.quixey.com) has built a search engine reliant upon app deep linking to help users find the most relevant content within apps in the same way that search engines help users find the most relevant content within websites.

App Deep Linking

The indexation of deep-linked content within apps (in-app pages) and the delivery of this content to mobile searchers have opened a world of possibility for apps and the mobile users who use them. By basic definition, app deep linking provides the ability for a link to send a user (or a search engine crawler) to deep, targeted content within an app (such as a specific page, tab, or view) based on their intent, as opposed to sending the user to the app’s home page – something we take for granted on desktop when we share deep-linked content from various websites or when we click on a link in search results. In essence, deep links take the user directly to what they want to do within the app, increasing app engagement, decreasing the time to conversion for targeted in-app behaviors, and improving the relevance of results in mobile search. All things considered, the mobile search ecosystem is poised to support this trend of driving intent-based traffic to mobile apps.

Deep linking to in-app pages and functionality also provides the ability for users to share deep content across various platforms (email, social sharing, SMS), further increasing engagement levels. And from an SEO perspective, it is likely that we will eventually see external linking, social, and other engagement signals applied to deep linked app URLs much in the same way that we see them applied to website URLs in SEO today.

Deep Linking URL Schemes

In order to provide deep links to your in-app content, you will need to create a unique URL scheme with routing parameters (path, query string, etc.) that represent the custom actions to take within the app, register these with the operating system, and map the routes to content destinations or user actions within the app. An example of a deep linking scheme with a routing parameter is `twitter://timeline`. When entered into a mobile browser, this opens the Twitter app and links directly to the device user’s timeline.

URX (<https://developers.urx.com/app-index/enable-deep-links.html#enable-deep-links>) and MobileDeepLinking.org (<http://mobildeeplinking.org>) provide helpful tutorials on deep linking for both the novice, and the app developer.

Deep Linking Standards

Deep linking standards attempt to define for app developers how to create in-app URL schemes so that the various platforms that publish these URLs can understand them. There are various the emergence of cross-platform standards as outlined by the providers below:

Facebook’s App Links: <http://applinks.org>

Twitter’s App Card: <https://dev.twitter.com/cards/types/app>

Google’s App Indexing: <https://developers.google.com/app-indexing/>

Quixey's AppURL: <http://appurl.org>

MobileDeepLinking.org: <http://mobiledeeplinking.org>

Deep Linking: Tracking & Optimization Opportunities

Increasingly, tools are being developed to provide deep linking capabilities, as well as acquisition data “outside the app,” filling a historical gap within traditional mobile app analytics solutions and enabling app marketers to utilize user intent metrics – which can provide greater insight into organic attribution, and can drive intent-based content delivery.

A sampling of these tools includes:

AppLinks (applinks.org) – Open, cross-platform solution for basic app-to-app linking (used by Facebook, Pinterest, Spotify, Venmo, and others)

TapStream (tapstream.com) – Deferred deep linking, onboarding links, A/B testing

Yozio (yoziocom) – Deferred deep linking, dynamic linking, A/B testing

A quick note on tools that provide third-party URLs for deep links to your app content, as opposed to using your own app link URLs – there may come a time in the future when it is smarter to develop your own deep link URLs for the same reason that you develop your own URLs for your site content – so that social shares and external links can be built over time and credited to your owned URL.

App Indexing

Google currently offers their App Indexing API to allow developers to notify Google about deep links in their native apps which are used by Google to index content that is then served in Google mobile search results: <https://developers.google.com/app-indexing/> (note that Google also states that the App Indexing API allows the Google App to drive re-engagement through Google Search query autocompletions: <https://developers.google.com/app-indexing/webmasters/appindexingapi>).

Bing currently indexes apps for search on Windows and Windows phone via the Windows Phone Store URL: <https://msdn.microsoft.com/en-us/library/dn614167.aspx>

Vertical Search

Vertical search refers to the search-accessibility of various niches and formats, of web content (referred to as “verticals”), including images, videos, events, news, travel, products, music, and people (this has also been referred to as “Universal” search in Google). Google provides Vertical search functionality, as do vertical-specific websites such as Expedia.com and Cars.com. The term “vertical search” is often used interchangeably with “blended search” – although this is not quite correct. “Blended” search results are search results that include various types of vertical content, such as videos, images, etc. SEO for vertical search refers specifically to efforts tied to optimizing vertical content types to appear prominently within vertical-specific search engines (Google Image Search, for example), and within blended search results. At this stage, the major search engines Google and Bing provide blended search results (in addition to their separate vertical search functions) to improve the user experience and increase the overall search value to users in ways that go beyond what traditional web search results have provided.

Generally speaking, Blended search results can provide more relevant (and therefore more valuable) search results within specific verticals. Vertical-specific search engines still exist, as Google and Bing provide numerous vertical-specific search functions within their search engines; however, we are also seeing an increasing trend towards the search engines incorporating vertical results from these engines within their main web search results.

A June 2014 study by Searchmetrics (http://www.searchmetrics.com/media/documents/knowledge-base/searchmetrics-universal-search-study-2014_us.pdf) found that the proportion of keyword searches in Google with Universal Search Integrations (Blended results) was recorded at 81% in 2013, up from 75% in 2012 – reflecting the ongoing opportunity for increased exposure in Blended search results and vertical search. SEO professionals should seriously consider what potential benefits vertical search optimization

can provide to their, and their clients' websites, and there are very specific ways in which vertical content can be optimized for success in these areas, as we will outline in this chapter.

Opportunities for Vertical Content Optimization

Vertical search has been around for almost as long as the major search engines have been in existence. Some of the first vertical search engines were for image search, newsgroup search, and news search, but many other vertical search properties emerged since then, both from the major search engines and from third parties. This chapter will focus primarily on strategies for optimizing your website for the vertical search offerings from Google and Bing.

Universal Search = Blended Search

Google's web results search engine used to be a kind of vertical search engine itself: one focused specifically on web pages (and not images, videos, news, blogs, etc.). In May of 2007, Google announced their concept of "Universal Search" – the merging results from different types of digital media, such as images and videos within their search results. Other search engines followed with similar search functionality, which has ultimately been referred to as "Blended Search," as a brand-neutral term (i.e. not referring to any particular search engine's results), within the search industry.

With the advent of Universal Search, Google changed the "web page" search engine into a search engine for any type of online content. Figure 10-1 shows the "Mega-SERP" infographic created by Moz (<http://moz.com/blog/mega-serp-a-visual-guide-to-google>), which reveals the wide array of search result types that can be pulled in to a regular "web" search at Google, as they are now incorporating various algorithmic and data components of its "Knowledge Graph" (or Carousel) into its main SERPs. Note that this Figure 10-1 is not an actual search results page, but rather shows the universe of vertical content types that can be pulled into general web search results:

Google - The Amazing World of Tacos

A Best tacos in Chicago

B Taco songs

C Search results for "new taco"

D Search results for "new tacos on Google"

E 156 calories

F Tacos la Mex - Mexican

G All Tacos in Orange, CA

H Tacos

I How To Make A Taco - YouTube

J Santa Fe (Mexican), White Foods Market

K Texas Johnston - Austin

L Taco Doctor in Decatur, Atlanta, Food Network

M Mexican Restaurants, L.

N Search for TACO BELL

O Tacos la Mex

P Mexican Restaurants, Mexican, Mexican

Q In-depth articles

R Mexican Restaurants, Mexican, Mexican

S Search results for "tacos"

T Search results for "tacos on Google"

U Taco

V Taco Bell

W Counter Mexican

X See results about

Y Map for planning tacos

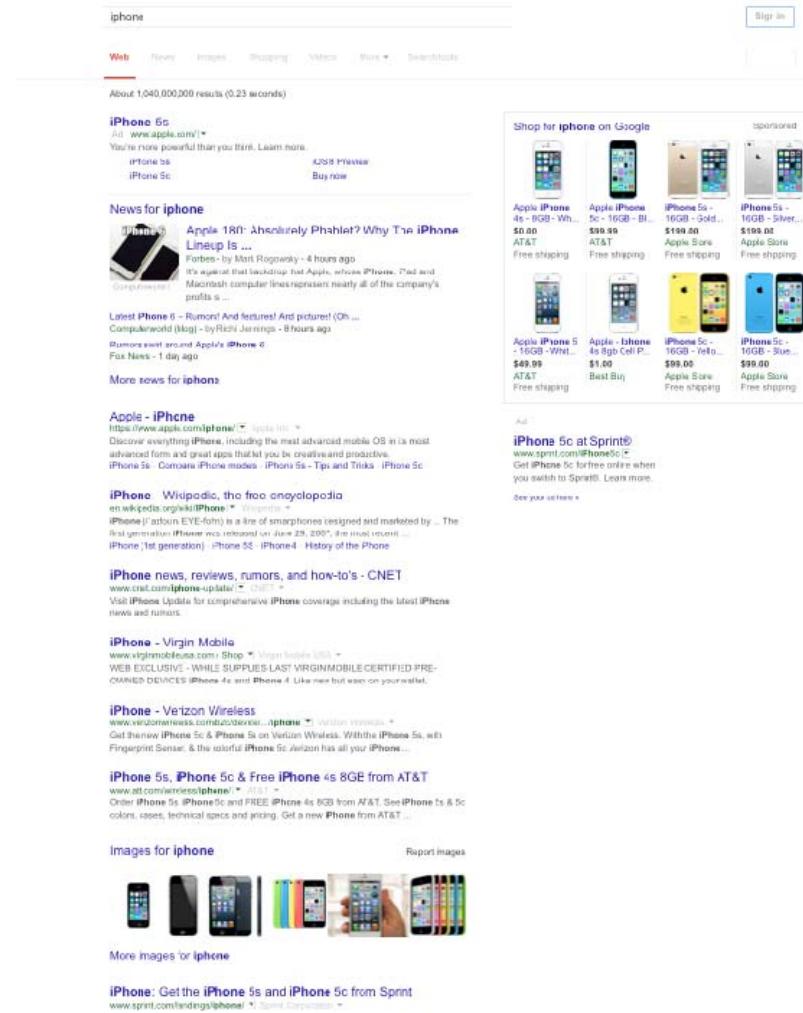
Google > 

Google Home | Adwords | Google Play | Business | Google+ | Privacy & Safety | Help | About Google | Feedback

[[figs/print/1002.png]]

Figure 10-2. “Mega-SERP” infographic (Moz.com)

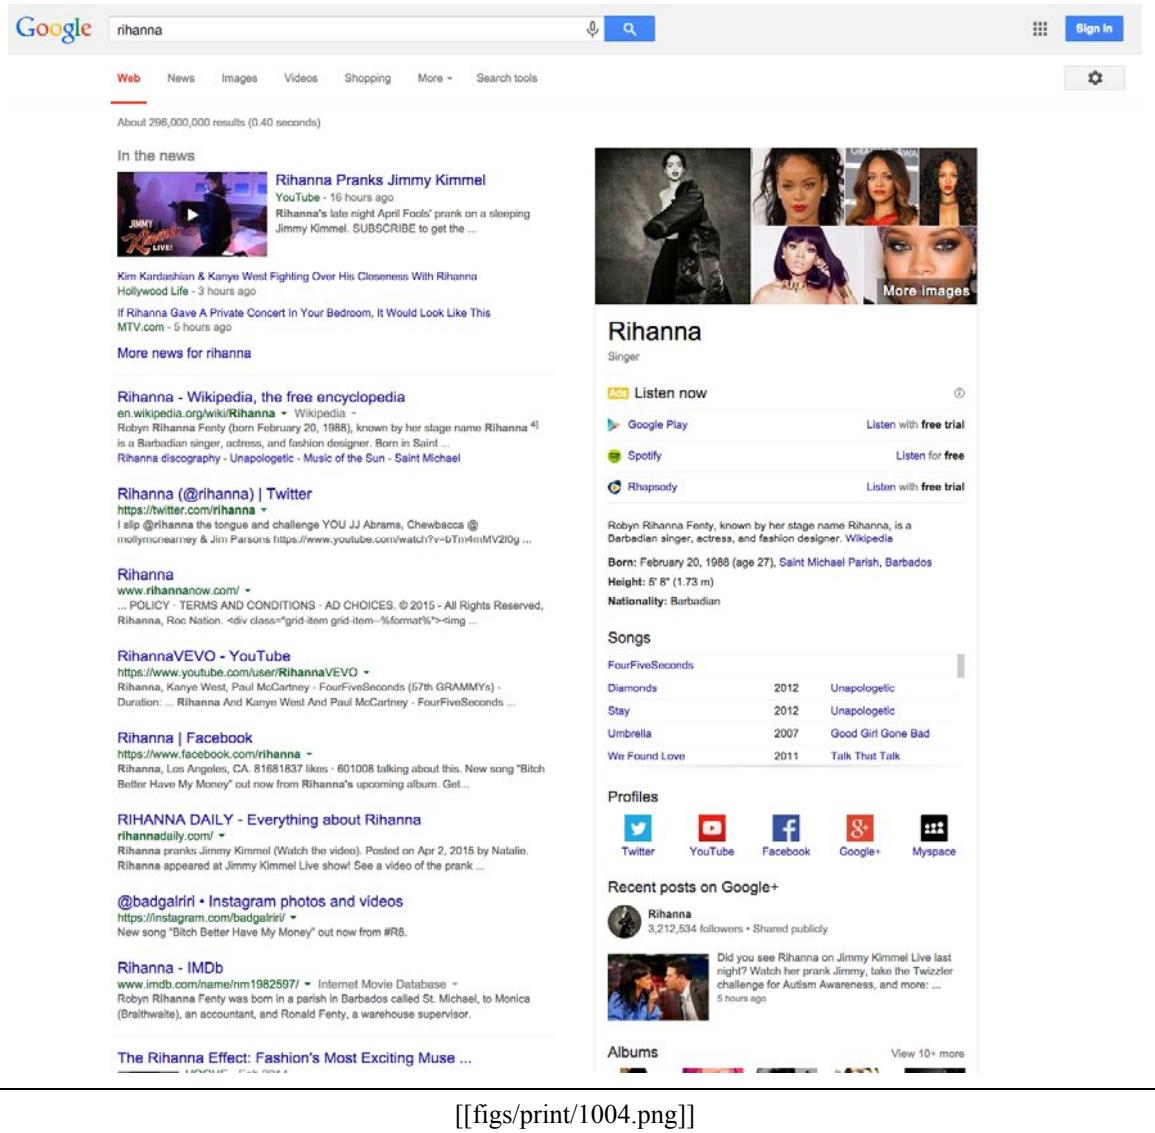
Figure 10-2 shows an example of how Blended search plays out in a Google SERP for a search for *iphone*. Notice the news results (labeled “News for iphone”), the image results (labeled “Images for iphone”), and the Shopping results (labeled “Shop for iphone” – note however that these are paid search results from advertisers). This SERP is an example vertical search results being incorporated right into traditional web search results.



[[figs/print/1003.png]]

Figure 10-3. Non-personalized search results for “iphone”

Figure 10-3 shows the results for a search on *Rihanna albums*:



[[figs/print/1004.png]]

Figure 10-4 - Search results for "Rihanna"

Right there in the web search results, you can read Rihanna-specific news (this search generated a Google News results set), click on a video and watch Rihanna's 2015 April Fool's prank on Jimmy Kimmel, select images (top right) from Google Image search, and see recent posts from her Google+ account.

The Opportunity Unleashed

As we noted at the beginning of this chapter, the opportunity in vertical search was significant *before* the advent of Blended search. However, that opportunity was not fully realized because many (in fact, most) users were not even aware of the vertical search properties. With the expansion of Blended search, the opportunities in optimization for vertical search have soared. By leveraging digital media assets such as images, videos, articles, events, recipes, tv and product reviews, and local business listings, businesses can increase their chances of having search result exposure that may directly or indirectly drive traffic to their sites.

Optimizing for Local Search

In 2014 Google published a research study called “Understanding Consumers’ Local Search Behavior” (<http://www.thinkwithgoogle.com/research-studies/how-advertisers-can-extend-their-relevance-with-search.html>), which revealed numerous valuable insights into the state of local search. (<https://sites.google.com/a/pressatgoogle.com/googleplaces/metrics>). Figure 10-4 shows the executive summary from this study:

EXECUTIVE SUMMARY



Consumers search with their location and proximity in mind

- 4 in 5 consumers use search engines to find local information.
- They search on smartphone and computer/tablet for: store address, business hours, product availability and directions.



Local searchers take action

- 50% of consumers who conducted a local search on their smartphone visited a store within a day, and 34% who searched on computer/tablet did the same.
- Local searches lead to more purchases than non-local searches. 18% of local searches on smartphone lead to a purchase within a day vs. 7% of non-local searches.



Consumers prefer and act on location-based ads

- 4 in 5 consumers want ads customized to their city, zip code or immediate surroundings.
- More than 60% of consumers have used location information in ads. They say it’s important to have store address and phone number in ads on computer/tablet, and directions and the call button in ads on smartphone.

google.com/think 4

[[figs/print/1005.png]]

Figure 10-5. Google Local Search Study: Executive Summary

Google’s study also found that 56% of “on the go” searches have local intent, and 50% of smartphone users and 34% of desktop/tablet users who perform a local search visit a store within a day of the search. These valuable insights underscore the importance of developing a local search strategy.

The major search engines have sought to increase their advertiser bases by moving aggressively into providing local directory information. Applications such as Google’s Google My Business (<https://www.google.com/business/>) and Bing Places for Business (<https://www.bingplaces.com>) have introduced disruptive technology to local directory information by mashing up maps with directory listings, reviews and ratings, local citations, satellite images, and 3-D modeling—all tied together with geotargeting and keyword search relevance. These innovations have excited users, and the mapping interfaces are growing in popularity as a result.

While there is no single dominant provider of local business information on the Internet, Google, Bing, and Yahoo have ramped up their efforts to tackle the local listings market by providing local listings within their search results – and have become the dominant resources for local information. “Local SEO” has

become a very complex and unique subset within the SEO practice, and has developed its own unique vernacular, for which Moz has provided a helpful glossary: <http://moz.com/learn/local/local-search-terms-glossary>. If you are hoping to gain traction in local SEO, it would be wise to familiarize yourself (and/or those working on your Local SEO efforts) with the terminology in this glossary.

The opportunity for local search is huge. More than any other form of vertical search, local search results have a major presence in web search. For example, Figure 10-5 shows the results for a search on minneapolis rental cars.

Figure 10-6 - Local search results example for Minneapolis rental car

Only three regular web search results are above the fold, and the following results are the “pack” of Google local search results (sometimes these “packs” show in 3, 5, or 7 packs).

Local Listing Submissions

Enhancing your local listing profiles in Google, Bing, and Yahoo! is a matter of providing as much relevant data to the search engines as you can. The most important elements to focus on are your business name, address, phone number, website, and categories. These are the elements that you can influence that carry the most weight in the ranking algorithms. While adding a rich description can help with click-through rates and conversions, it has very little, if any, weight in the ranking algorithms.

Guidelines for making choices for these data elements:

Business Name

Make sure you use an accurate business name. Do not try to add keywords or locations—this is considered spam and may result in your listing being rejected. Be sure to use the exact same business name that you use on your website, and in any other online business profiles.

Address

Like the business name, your address is part of the “key” that the search engines use to identify your business. It is helpful to ensure that the address “maps accurately” in the search engines. Do this by performing a search in each engine on just your address. See if the engines show an exact match for your business address; if they do not, you may need to modify your address (for example, there may be multiple ways a particular street is referred to, such as “Mc Andrew” vs. “McAndrew” and you may have to try an alternate version), or you may need to report a mapping error to the search engine directly. To do so, click on the “Report a Problem” link, which can be found in several places on the map results – usually at the bottom of the map, or underneath the business address within the map display.

It is also useful to ensure that the USPS considers your address to be correct. You can check this with free tools available at <https://tools.usps.com/go/ZipLookupAction!input.action>.

Phone Number

For the “main” phone number, be sure to use a unique local phone number. Avoid 800 numbers, and avoid tracking numbers if these numbers are not planned to be utilized on a long-term basis (i.e. for testing or other purposes). As with your business name and address, your main phone number is part of the “key” the engines use to cross-reference information about your business around the Web. Using a nonlocal or frequently changing phone number can potentially confuse the search engines and cause them to have less confidence in your listing. In addition, it can lead to the creation of duplicate business profiles, which would compete against your core listing for search engine rankings.

Website

In many cases, for a business with a single location, this is simply your home page. For businesses with multiple locations, it is ideal to have a separate page representing each location. This case is described further in the main text below.

Be sure to provide the address and local phone number (as described above) on the page you link to. As described in the main text below, it can be helpful to encode the contact information on the web page using markup such as that provided by Schema.org (<http://schema.org/LocalBusiness> or its industry-specific variants) or microformats (<http://microformats.org/wiki/heard>). You can also use the Google Webmaster Tools Data Highlighter to “tag” the organization data on your pages in lieu of markup if you do not have easy access to modifying your web pages, but note that this will only serve to inform Google, and not Bing or Yahoo!.

Having a website can be a strong ranking factor for the search engines, and it can be a strong factor in converting searchers to customers. It is advised that you provide at least a basic website associated with your business listing. Of course, if you do so, you will also want to follow much of the other advice in this book to optimize that website.

Description

Google now provides you with a 4,000+ character description field, which should be plenty of space to describe your business and promote your message. Still, you want to provide a well-written, unique, description of your business. While Google states they no longer rely upon the business description for ranking purposes, it is important to accurately describe your service, product, and location specifics for users.

Categories

Category selection provides a very strong clue for the search engine to decide what kinds of searches your business is relevant to. The best advice here is to simply study the available categories, and be sure that the first one you choose is the “best” category. At least one category selection must be made from the search engines’ standard set of categories. Google and Bing allow for up to 10 categories, and Yahoo! allows for 5.

If you use your own categories, be sure to follow the search engine quality guidelines. Mainly, be sure to list categories that describe what the business is (as opposed to, for example, what products it sells or where it is located). Also, keep each category focused on a unique quality of the business. Do not attempt to jam keywords into your category choices, as the engines may consider this to be spam and reject your listing.

One good source for brainstorming categories is to examine the business listing pages of some competitors, and see what they’ve chosen. It is helpful to be seen in the same categories as the majority of your competitors. This ensures that the engines consider your business as “the same type” as your main competitors, allowing you to be considered for display in the most important search results.

In addition to these key elements, you also have the opportunity to provide information on your business hours and accepted payment methods, photos, videos, and other data items in your business profile. Doing so is worth the effort, as it results in an enriched business page, which may help entice more users to contact your business. It also provides additional opportunities to provide keywords and context to help the search engines better understand how to rank your business.

Google My Business (GMB)

In June 2014 Google announced the launch of Google My Business, in an effort to streamline and simplify the management of local business profiles on Google – which had evolved into an increasingly complex and time consuming array of business listing options and locations under the umbrella of Google properties. According to Google, businesses who already have existing Google Places listings and Google+ Local Pages will automatically be upgraded to the GMB platform (the opening page for Google My Business is shown in Figure 10-6). At the launch of this service, Google added the ability to “View the Business on Google,” which gives you a preview of the business info in Google Search, Maps, and Google My Business.

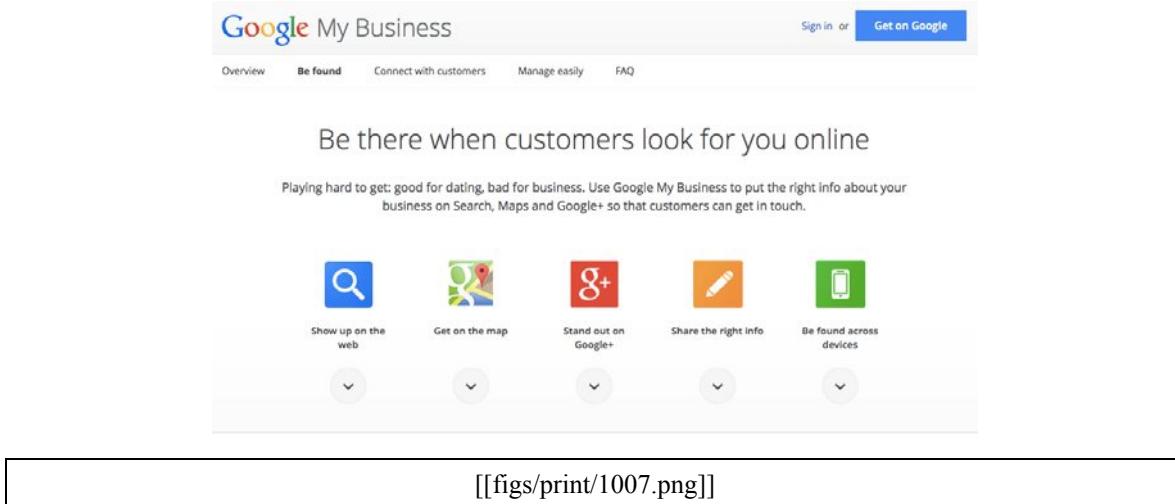


Figure 10-7 – Google My Business Splash Page

More information about Google My Business can be found on the FAQ page: <https://www.google.com/business/faq.html>.

Optimizing Your GMB Listing(s)

It is worth investing the time to learn how to optimize for GMB. One thing that may surprise you: if your business has been around for any length of time, Google probably already knows about your business and has a listing in place. Google is very good at gathering and aggregating data from multiple sources to create a comprehensive representation of your business.

As of this book's publishing, you will likely find this listing in a local SERP within the local pack with a link to a Google My Business page for the business. If Google has generated such a default listing for you, you should claim this Google My Business listing (rather than creating a new one). Do so by clicking on the "Manage this Page" link that appears on your business's Google My Business page, underneath "Is this Your Business" and following the prompts, which will include creating a GMB account.

Claiming the listing will allow you to take control of the profile information for your business (the business name, address, and phone numbers(s), as well as a vast set of other data elements, including the website address, the business categories you should be associated with, photos, etc.). Claiming is relatively straightforward. Before granting you control, Google will use some method to have you prove you own the business being claimed.

This may take the form of a postcard being mailed to the business address, a live phone call from a Google robot, in which a secret code is provided to you. You then enter this code into the portal, and Google grants you exclusive control of the listing. You may also be able to instantly verify your Google My Business listing through Google Webmaster Tools; in order to do this, the website listed on the GMB profile must be verified in GWT with the same email address you are using to claim the GMB profile. You can also batch upload listings if needed: <https://support.google.com/business/answer/3217744?hl=en>.

More information on creating, claiming, and verifying your Google My Business listing can be found here: <https://support.google.com/business/answer/2911778?hl=en>.

As mentioned above, it is almost always better to claim an existing listing, if there is one, than to create a new one. One thing that local search engines struggle with is duplication of listings, and it is best if you do your part to ensure that there is only a single listing representing your business. In fact, an advanced technique that you may want to explore is to search Google Maps using your business name, address, and phone number to see if there are any duplicates or near-duplicates of your business information, and, if so, to claim those as well. Doing so will help Google consolidate this information, and the benefit to you is that you will go from having several "weaker" listings to having a single strong consolidated listing.

By taking control of your GMB listing, you've taken the first significant step toward improving your ranking in Google local search. Google will likely have more confidence in your listing data immediately, because you have proven ownership and asserted control. This confidence is a key ranking factor, as one very important function of local search engines is to cull out the most accurate, up-to-date data from among the sea of conflicting information that defines the local business information landscape. You are now poised to take the next key steps in optimizing your listing.

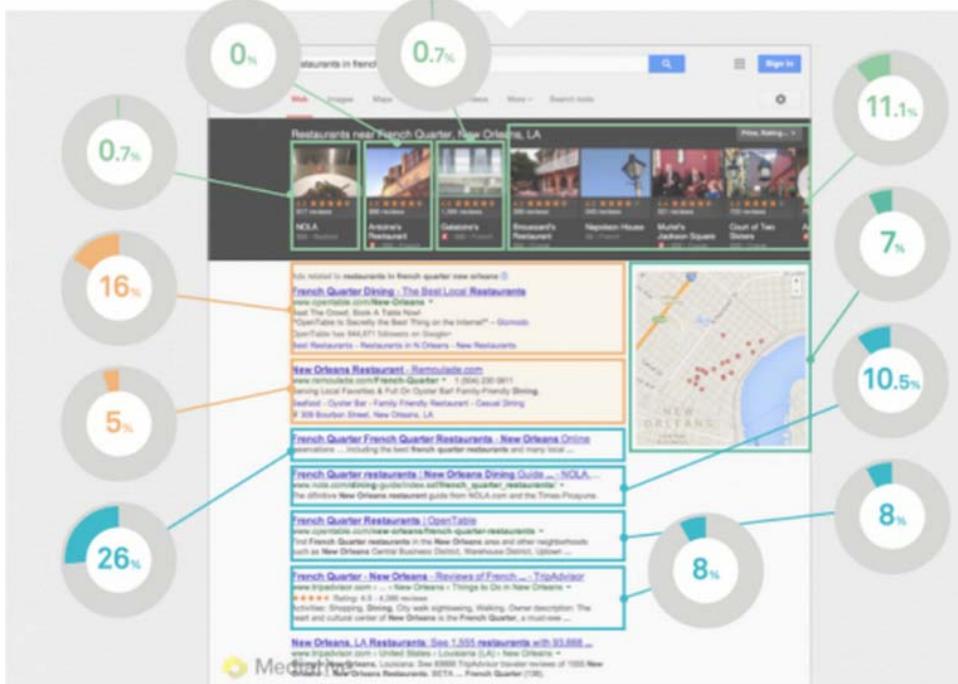
For more information from Google about increasing your business listing exposure, go to:

<http://googlewebmastercentral.blogspot.com/2014/04/surfacing-your-businesss-contact-and.html>

Google Knowledge Graph Carousel

A fairly comprehensive list of keywords that trigger the Google carousel can be found here (<http://blumenthals.com/blog/2013/06/29/keywords-that-trigger-the-local-carousel-a-compilation/>). A recent eye tracking study by Mediavite (<http://pages.mediavite.com/SERP-research>) showed some interesting data on click through rates for search results showing a Google Carousel is shown in Figure 10-8.

AVERAGE CTRs: SERP WITH CAROUSEL



[[figs/print/1008.png]]

Figure 10-8 – Local Search Result Showing a Carousel

According to their study, the highest performing Carousel results were the ones at the far right of the display, not those in the tradition 1,2 or 3 positioning on the left. While this is just a single study, it shows the challenges that a new visual interface can create with traditional user behavior.

As you can see in Figure 10-8, a search for minneapolis pizza is heavily biased toward local results and we see Google's implementation of local "Carousel" results at the top – these carousel results are sourced from the same set of data as regular local "pack" results however, more results are displayed then the tradition

3,5 or 7 pack. This graphic display, which was seen in mobile/tablet results in 2013, offers a new set of challenges and opportunities for business owners. It's important to note that Carousel has not effected all local results.

minneapo is pizza

Sign in

Web Maps Shopping News Images More Search tools

Pizza near Minneapolis, MN

Prices, Rating ...

4.1 209 reviews Pizza Luce \$8 • Pizza

4.4 98 reviews Black Sheep \$8 - Pizza

4.2 30 reviews Pizza Nella \$8 • Pizza

4.4 23 reviews Punch \$8 - Pizza

4.5 55 reviews Pizzeria Lola Pizza

4.4 87 reviews Galactic Pizza \$8 • Pizza

3.6 39 reviews Serrano's Pizzeria Pizza

4.2 64 reviews Fat Lorenzo's \$8 - Pizza

4.1 81 reviews Pizza Luce \$8 • Italian

Top 10 pizzerias in the Twin Cities | City Pages
blogs.citypages.com/food/.../top_10_pizzerias_twin_cities.php City Pages +
 Dec 18, 2011 - Black Sheep Coal Fired Pizza (which has locations in both the North Loop of Minneapolis and St. Paul) puts a big premium on its crust.

Downtown Minneapolis - Pizza Luce
pizzaluce.com/location/downtown-minneapolis/ • Pizza Luce +
 Located at the epicenter of Minneapolis' historic Warehouse District, our flagship location began as a humble slice counter back in 1993. Since then we've ...

Uptown Minneapolis - Pizza Luce
pizzaluce.com/location/uptown-minneapolis/ • Pizza Luce +
 Two blocks south of the intersection of Lyndale Avenue and Lake Street, Pizza Luce opened in 1999 and immediately became a popular neighborhood ...

Best Pizza in the Twin Cities - Thrillist Minneapolis
www.thrillist.com/.../minneapolis-best-pizza-in-the-twin-cities Thrillist Media Group +
 Sep 6, 2013 - Newsflash: the best pizza place is not always the closest pizza place. Unless of course you live by the test pizza place, and you'll totally know ...

Minneapolis » Restaurants » Pizza - Yelp
www.yelp.com/search/?find_desc=Pizza • Pizza +
 The Best Pizza in Minneapolis on Yelp: Read about places like: Pizzeria Lola, Galactic Pizza, Element Wood Fire Pizza, Black Sheep Coal Fired Pizza, Pizza ...

Pizza delivery Minneapolis, MN - Yelp
www.yelp.com/search/?find_desc=Pizza • Pizza +
 Reviews on Pizza delivery in Minneapolis, MN: Galactic Pizza, Pizza Luce, Black Sheep Coal Fired Pizza, D'Amico's Pizza, Mess Pizza, Boss' Pizza and Chicken, ...

Pizza Luce - Downtown Minneapolis - Minneapolis, MN | Yelp
www.yelp.com/biz/pizza-luce-downtown-minneapolis-minneapolis • Pizza +
 Rating 4 - 236 reviews - Price range: \$5
 236 Reviews of Pizza Luce "LOVE, LOVE, LOVE. I've ordered takeout from here for my office, but last night (4/2/14) was my first time I dined in location. + Love ...

The Twin Cities' Best Pizza - Minnesota Monthly - March ...
www.mnmonthly.com/.../twin-cities-best-pizza Minnesota Monthly +
 Mar 5, 2013 - At Broadway Pizza's original northeast Minneapolis location, not only can you watch a model train chug through the bar, but you can dine ...

Best Downtown Minneapolis Pizza Places - Twin Cities
www.urbanspoon.com/.../Downtown-Minneapolis • Restaurants +
 Downtown Minneapolis - Italian, Pizza 625 Marquette Ave, Andina Pizza \$5 ...
 Downtown Minneapolis - American, Burgers, Pizza 610 Hennepin Ave, Downtown ...

Best Twin Cities Pizza Places - Urbanspoon
www.urbanspoon.com/.../Twin-Cities • Restaurants +
 Northeast Minneapolis - Pizza, Pub Food, Sandwiches/Subs 1900 Marshal St NE ...
 Northeast Minneapolis - Italian, Pizza, Salad 210 Hennepin Ave, E. ...

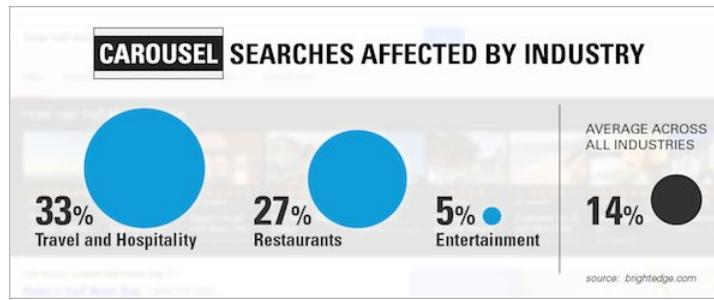
Searches related to minneapolis pizza

minneapolis pizza delivery galactic pizza minneapolis
 great pizza minneapolis minneapolis pizza luce

[[figs/print/1009.png]]

Figure 10-9 Local search results for "minneapolis pizza"

According to an October 2013 study by BrightEdge (<http://www.brightedge.com>), Google Carousel impacted 14% of keywords across all industries, with Travel and Hospitality verticals being most affected, as shown in Figure 10-9 below.



[[figs/print/1010.png]]

Figure 10-10 – Google Carousel SERP impact across verticals

Inclusion within the Google Carousel (Knowledge Graph) is available for authoritative sources for organizations, events, movie reviews, and music/video play actions and requires the use of structured data markup as outlined by Google’s Developers guide to Structured Data: <https://developers.google.com/structured-data/>.

As mentioned earlier, Google also provides The Data Highlighter within Google Webmaster Tools (<https://www.google.com/webmasters/tools/data-highlighter>), which enables webmasters and site owners to “tag” structured data on site pages so that it can potentially appear as “rich snippets” of content within Google search results. While webmasters have always had the ability to markup structured data using various formats, with Data Highlighter Google now enables webmasters to alternatively “tag” structured data for Google’s use without the webmaster having to modify the site’s pages on the server. Additional information about Google’s handling of rich snippets and structured data can be found here: <https://support.google.com/webmasters/answer/99170>.

Currently, Google supports rich snippets for the following content types:

- Reviews
- People
- Products
- Businesses and organizations
- Recipes
- Events
- Music

Chapter 13 provides more detail on utilization of this tool within Google Webmaster Tools.

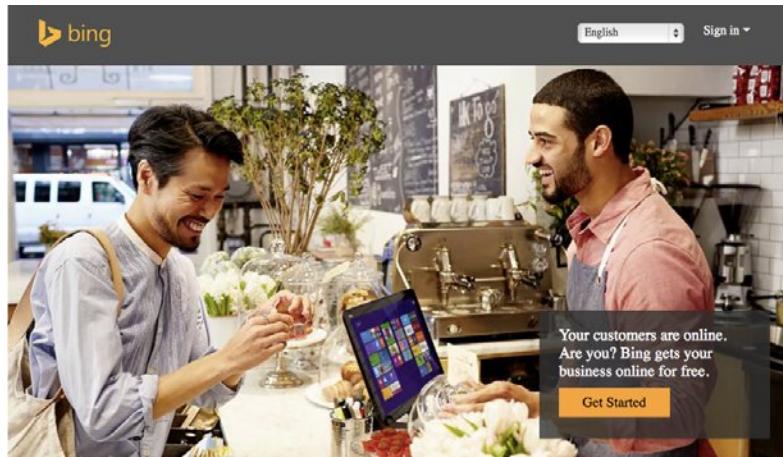
A Note on NAP Consistency

Your business’s name, address, and phone number are perhaps the most critical components of your local listing data that contribute to local ranking – and therefore keeping this information consistent throughout your web presence, and across your local listings, is essential. Most businesses go through normal transitions that require NAP information changes (business location changes, contact phone number changes, etc.). Be sure to update your NAP information on your website, your local listings, and your third

party platform pages (youtube, facebook, etc.) after any of these transitions, and periodically monitor your NAP information for consistency.

Bing Places for Business

Bing also provides businesses with the ability to create local listings for consideration in the organic search results with Bing Places for Business (<https://www.bingplaces.com> – the splash screen is shown in Figure 10-11), which replaced the Bing Business Portal in April 2013. If you already have a Google My Business listing, Bing allows you to import that listing into their Places for Business interface to streamline the process: <https://www.bingplaces.com/Home/MoreFAQ#ManageListing>. Bing also allows businesses to upload up to 10,000 business locations in one step.



[[figs/print/1011.png]]

Figure 10-11 – Bing Places for Business

Listing your business in Bing provides businesses with yet another angle from which to appear in Bing search results; according to Bing's Duane Forrester,

"It can be a way to double dip in the SERP, as their Places listing data can appear, with them ranked up top in the organic stack."

Bing also allows you to add links to your business's various social media profiles, so be sure to take advantage of this and list your Twitter and Facebook URLs.

Yahoo! Local

Yahoo! has been undergoing a steady evolution away from search as a core business function. However, they still provide value to businesses via their Small Business channel with their Local Marketing services, Localworks (local business listings in Yahoo! and 40+ directories), Mobile Offer Ads (mobile advertising), and Yahoo Gemini (their mobile search and native ad platform).

You can create a free Yahoo! Local Basic Listing (<https://smallbusiness.yahoo.com/local-listings>) which includes your address, phone number, brief description and URL; you can also create a Local Enhanced Listing that enables you to add a longer business description, multiple web links, and special offers/coupons.

Yahoo! also provides a mass-directory listing service called Yahoo! Localworks, which lists your business in 40+ directories including Yahoo! Local, Yelp, WhitePages, Bing, MapQuest, Superpages, and others, and enables you to add videos and photos to your listing – as well as track your listing performance with analytics.

Directories and aggregators

Directories can be built from the local phone company's database information, but no one phone company covers the entire country. Because of this, companies that host nationwide directories are primarily getting their content from data aggregators to form the foundation of their guides. Data aggregators build their content from a variety of sources, such as local-area print yellow pages, to have information that is as comprehensive as possible.

Three top aggregators exist for US business listings: Infogroup (<http://www.expressupdateusa.com/>) Axiom (<http://mybusinesslistingmanager.myaxiom.com/>), and Neustar/Localeze (<http://www.neustarlocaleze.biz/directory/index.aspx>).

The first step in managing the online presence of a local company is to check and update the business's listing information in each of these main aggregators. Ensure that the business name, address information, phone numbers, fax numbers, and any other contact information are correct. If you have multiple locations, make sure they're all listed.

In addition to checking for listing accuracy, it is important to address duplicate data issues at the data aggregator and local search publisher levels. Andrew Shotland, Founder of Local SEO Guide (<http://www.localseoguide.com/>), wrote an invaluable white paper for Yext discussing the process of addressing duplicate local listing data, entitled "The Definitive Guide to Duplicate Listings" (<http://offers.yext.com/the-definitive-guide-to-duplicate-listings>).

Optimize Your Business Website for Local Search

If you have been around for a while, your business probably is already included in the local search engines, since they compile data from the aggregators and other online directories. Once your business's listing is loaded into the local engines, you must determine how to get your business's own website to rank higher when users search for your industry's keywords. On the next few pages, we will outline things you can do on your website to prepare it for better rankings in local search engines.

All of the basic SEO factors can come into play here and help to influence your rankings. These factors include having good, specific text in each page's title, `<h1>` tags, meta description, page content, image `alt` attributes, inbound links, and so forth. But some things are specific to local search, such as the following:

- If your company has multiple locations, it is not necessary to have a standalone website or subdomain (e.g., `loc1.example.com`, `loc2.example.com`, `loc3.example.com`) for each outlet. In fact, it is probably better if you don't, since each business location would likely produce similar website content. However, it probably would be helpful for you to create a separate profile web page on your site for each distinct location. Many sites with chain outlets list all outlets on one page—that is not optimal. It is better to have one page be about one store location and another about a different location, so that you can leverage all the on-page elements to create a page focused on that location. One caveat: you will want to have a different phone number listed for each location, preferably with a local area code (as opposed to generic 800 number).
- Have your page title, `<h1>` tags, and content include the business name, the type of business, and the location—for example, "Acme Café: French Restaurant in Boston, MA." For multiple locations, make the title different on each location's page. Include the phone number, neighborhood, street address, area nicknames, and other location-distinguishing information.
- The home page and/or Contact Us page should have the main location's basic listing information (street address, city, state, zip code, phone numbers, etc.) displayed somewhere in the HTML text. You should also add the basic listing information in the Schema.org Organization or LocalBusiness format

or the hCard microformat (these are methods for encoding address information on web pages; you can learn more about them in “Schema.org and Microformats” and at <http://schema.org> or <http://microformats.org>). If you have multiple locations, display the basic information on each location’s profile page. Google also accepts RDFa for this information (as well as tagging via Data Highlighter in Google Webmaster Tools). The Schema.org approach is the one we recommend, as this is a standard that has been agreed upon by Google, Bing, and Yahoo! since 2011: (<http://googleblog.blogspot.com/2011/06/introducing-schemaorg-search-engines.html>).

- Place differentiating information on each store’s pages, including, for example, items such as a map, hours of operation, brand names of products carried, product photos, “in business since x” information, menus with prices (if it is a restaurant), ratings (if applicable), certifications, bonded status, etc. As mentioned previously, be sure to include specifics about the physical location.
- In the Second Edition of this book, we discussed the concept of centroid (the location that the search engine defines as the center of the city). Many of the map search engines used to primarily display the businesses located closest to the centroid first for any particular category or keyword search. However, centroid is being replaced by Business/Searcher distance as a ranking factor – as the distance of a business from a searcher is more significant.
- Proximity works the same way for zip codes. If a user searches for businesses within a specific zip code, the businesses closest to the zip code area centroid will likely be displayed first in the list. Note, though, that zip code searches are not commonly performed in the United States (you can verify this with your favorite keyword research tool).
- Do not attempt to set up multiple user accounts to rate yourself well or reduce your competition’s ratings! Likewise, do not pay customers for beneficial reviews—if such a practice were discovered, you could lose all of your beneficial ratings. Focus on working in ways that are allowed under all the search engines’ rules: just ask people to rate you.

By asking enough happy clients to rate you, you might be able to drive up the positive reviews, which can bring rankings benefits. It is worth the effort to identify your happy customers and ask them to write reviews for you at one or more of the major directory sites. Never offer to pay for or otherwise incentivize positive reviews, however!

Consider doing post-sales surveys to find those happy customers, and then reach out to them to ask them to post reviews. You can offer to enter people who write reviews into a drawing to win a prize, but don’t make their participation in the drawing contingent on a positive review and don’t pay for each review. Monitor your online ratings in the various sites as part of good reputation management and try to respond to complaints/issues as quickly as possible. Consider using an online review monitoring tool such as Review Push (<http://www.reviewpush.com/>), and Review Trackers (<http://www.reviewtrackers.com>).

- Links from local directories and other locally-relevant sites (local “citations”) are increasingly important ranking factors, and should be obtained when relevant.
- For businesses such as restaurants, it may be beneficial to add photos of celebrities who have visited the establishment.

People perform searches on celebrity names all the time, so if you are lucky enough to have celebrities frequent your restaurant, uploading photos of them can help attract celebrity watchers to your business – however, be sure to get the celebrity’s approval before using their photograph, or “likeness” in your advertising or promotional materials, as without it you may be inviting a lawsuit. With permission, a celebrity photo endorsement can drive up inbound links and get users to linger longer on the site, perhaps bolstering the site’s quality scores.

- Community interaction can assist in online promotion. One great way to obtain local area links is to support charitable efforts in the area and to sponsor local events. Although it may seem crass to get link value out of charitable efforts, this is one of the best ways to support your local community while bolstering your business in the online marketplace. Don’t feel guilty for creating a win-win situation! Such opportunities could include:

- Local charity races/walks, particularly if they list sponsors or sponsor logos on a page on their site. Request to have those linked to your local business site. Inbound links from these sites are very beneficial, and this is a valid way to get them!
- The local college or high school sports teams and bands. You can request that sponsors are listed on the school website, if this isn't done already.
- Local fraternal organizations, such as the Freemasons, Lions Club, and Shriners.
- Charitable events that you host yourself. Is anyone organizing the provision of food for the needy at Thanksgiving? If not, organize it yourself and host the information pages off your own website. This could draw many others to link to you as they help in promotion efforts.

A special thanks to John Biundo (<http://www.stonetemple.com>) and David Mihm (<http://www.davidmihm.com>) for their contributions to the "Optimizing for Local Search" section of this chapter in the Second Edition, and to Andrew Shotland and Dan Leibson from Local SEO Guide (<http://www.localeseoguide.com/>) for their invaluable assistance with updates to this section in the Third Edition.

Optimizing for Image Search

According to the aforementioned Searchmetrics study, in 2013, images were displayed 45% of the time in Google SERPs, more frequently than any other Blended search elements. Optimizing for image search can be a smart strategy for many search marketers. Even if you're working on a site that you don't feel truly lends itself to an image optimization strategy, you may be able to leverage images or photos on your website for SEO.

However, we should note that for some sites, there may not be a lot to gain here. Some e-tailers report poor conversion on image search traffic, and there is always the concern over third parties stealing their images. You need to weigh the benefit of image search capability against the costs and other opportunities for SEO on your site. Nonetheless, many sites do very well with image search, and image search optimization does not have to take much time.

A significant amount of traffic can come from image search, and the number of people competing effectively for that traffic is much lower than it is in general web search. Industries that don't immediately seem to provide compelling subjects for images may enjoy greater potential in this area, because the competition might never clue in to the advantages of integrating images into their sites and into an overall search marketing strategy.

There are a few different ways that image search optimization can help to improve traffic and conversions for your site:

Subtle reputation management

Images of your products/services/facility assist consumers during the research phase of their shopping, and lend an implicit message of openness/forthrightness to your business. Providing generous numbers of images can improve consumer confidence in your company, increasing the chances that they'll decide to do business with you.

Shopping via image search results

Increasingly, consumers are searching for products via image search engines because they can rapidly find what they are seeking without having to dig through promotion-laden websites. If your products can be found in the image search engine, you have an improved chance of being found by those people. With no pictures, there's zero chance of being found in image search.

Increased chances of showing up in Universal/Blended Search results

Performing image search optimization improves your chances of showing up in additional positions on the main search results pages, as Universal Search pulls image search content into the main SERPs for some keyword search terms.

Empowering others to promote you

If you have a flexible enough organization and you hold the legal copyrights to your images, you can allow others to reuse the images in return for promotion of your site/business.

Image Optimization Tips

In comparison to indexing pages that are rich with text content, indexing images is much more difficult for search engines to perform. The image itself provides few clues to the content within it. Google is known to be experimenting with artificial intelligence algorithms to improve detection of image content. For example, you can drag the image of a well-known person or place onto the Google Image search box (<http://images.google.com>), and Google will attempt to identify the contents and show you other similar images.

Search engines are also experimenting with techniques such as optical character recognition (OCR) to read text content within images, but most images don't have any text to read. Search engines also make use of facial recognition software to be able to determine when an image is of a face versus a body, or something else entirely.

However, although these types of technologies are very useful, they are limited in terms of what they can do, partly because they are computationally very expensive and therefore not practical to use at the scale of the entire web. For that reason, success in image search optimization depends on using all the signals available to you to increase the search engines' confidence in the content of your images.

This certainly includes the basic SEO techniques we have discussed in this book. The web page's title tag, the `<h1>` heading tag, the on-page content, and links to the page are all factors in image ranking. For example, if you have a picture of the Golden Gate Bridge as the single image on a page, and the title, heading tag, and content of the page all support that, the search engines' confidence in the content of the image increases. The same is true if you have 10 images on a page of many different bridges, and their content is reinforced by the title, heading tag, and page content. Consistency of content and theme is important in all of SEO, and it is especially critical in image SEO.

In terms of the page content, you should give particular emphasis to the text immediately preceding and following the image. This is what the user most closely associates with the image, and the search engine will view it the same way. A descriptive caption underneath the image is helpful.

You can do a number of things to further optimize your images. Here are the most important things you can do:

- Make sure the image filename or img src string contains your primary keyword. If it is a picture of Abe Lincoln, name the file `abe-lincoln.jpg` and/or have the src URL string contain that keyword, as in `http://example.com/abe-lincoln/portrait.jpg`.
- Always use the image alt attribute. The alt attribute helps the vision-impaired to understand your site, and search engines use it to better understand what your images are about. Recent research by the authors indicates that this feature is still not used for lots of sites' images, and that many sites have tried to use it with invalid HTML. Make sure the alt parameter is valid, as in this example:

```
| 
```

Use the quotes if you have spaces in the text string of the alt content! Sites that have invalid img tags frequently drop a few words without quotes into the img tag when they were intended for the alt attribute. Omitting the quotes is a common problem; without them, all terms after the first will be lost, if any are used at all.

- Avoid query strings for img src tags, just as you should for page URLs. Or, if you must use URLs that include query strings, use no more than two or three parameters. Consider rewriting the query strings in the URLs so that they do not contain an excessive number of parameters, which will cause spiders to refuse to crawl the links. Note that Google claims to no longer have problems with these types of situations, but it is better to be safe than sorry.
- Use good-quality pictures, which will read well when shown in thumbnail format. Good contrast is typically the key here. Lower-contrast images are visually harder to read, and it is common sense that if the thumbnail image doesn't look good, it will not invite a click.
- Do not save images as graphics files with embedded thumbnails—turn this feature off in Photoshop and other image editing software. Search engines may copy your image, reduce it in size, save it in compressed format, and deliver up a thumbnail of it for their results pages. An embedded thumbnail can wreak havoc with some compression software, and it increases your file size slightly, so just leave that feature disabled.
- Don't store the image in a sidebar column with your ads or inside the header/footer navigation elements; otherwise, the search engine algorithms will ignore the image as irrelevant, just as they ignore page decor and navigation graphics.
- Have a proper copyright license! You need to have a proper license to display the images found on your site so that you don't get sued. Be careful about trying to use images from Wikimedia Commons (http://commons.wikimedia.org/wiki/Main_Page) or other public stock photo sites, since you cannot be sure that those images really are in the public domain. For example, when you "purchase" an image from a stock photo site, you are not purchasing the copyright – you are purchasing the right to use the image. More often than not these arrangements require that you link back to the copyright holder.
- You need to ensure that your server configuration allows your site's images to be displayed when called from web pages on other domains. Some system administrators have disabled this to keep people from displaying their images on other sites, and this could cause problems if you want your images displayed in search engine image results pages. Likewise, make sure that your *robots.txt* file does not block the crawlers from accessing your image file directories.
- If it is a fit for your business, specify that others are free to use your images for online display as long as they link back to your website from a credit line below or adjacent to the image, where they display your copyright notice. Enabling others to use your photos invites more promotional attention when people wish to write about you in blogs or in news articles.

It is also a good idea to create an image XML Sitemap to highlight images you consider particularly important, or to help the search engines find images that might otherwise be difficult for their crawlers to discover. You can find more on image XML Sitemaps here: <http://googlewebmastercentral.blogspot.com/2010/04/adding-images-to-your-sitemaps.html>. You can also read more about this in "Image Sitemaps".

With the advent of Blended search and the web's overall movement towards more visually compelling, image-rich content, there are more opportunities to achieve organic exposure with images – and it is also becoming more competitive. The following is a list of image optimization best practices to use when uploading your images to image sharing sites such as Flickr. This will help ensure your image files have the best chance of appearing for targeted, qualified searches in the SERPs.

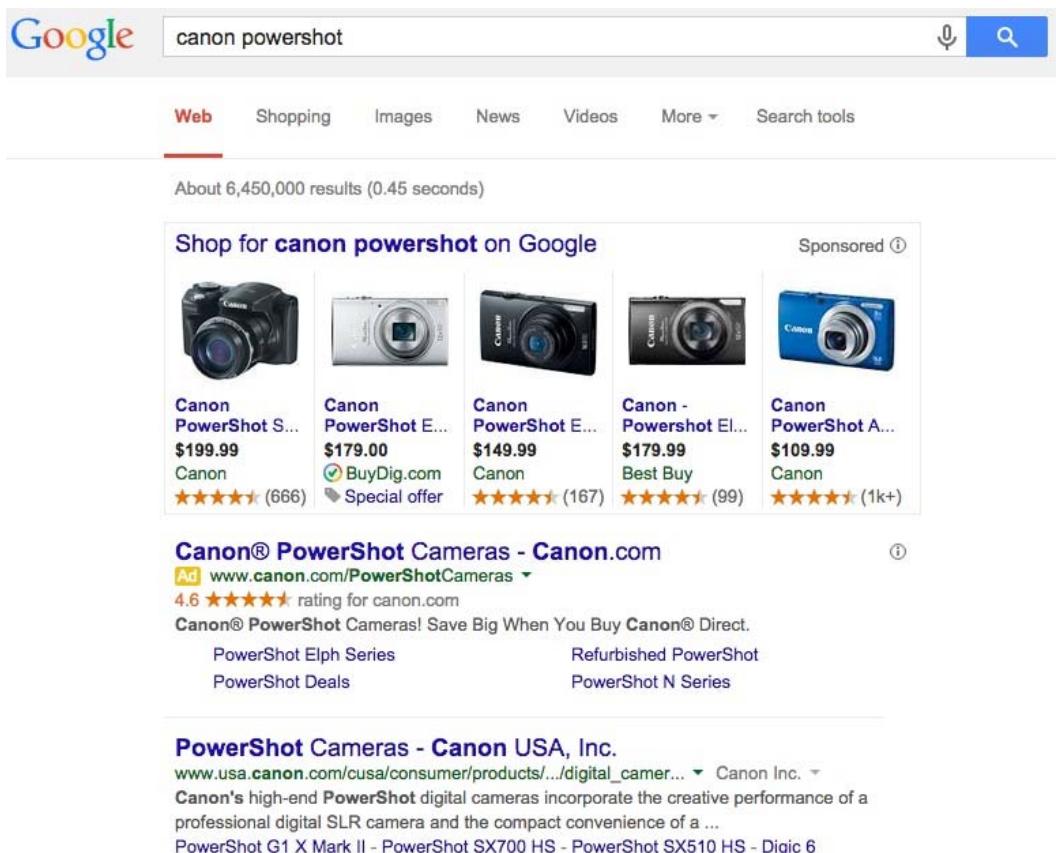
- When you upload your photos always add tags. The tags or keywords that you associate with your photo will make sure users can find it when they are searching and will lend keyword weight to the photo's page. Enter as many tags as possible that accurately describe your photo. Make sure you place any multiword tags within quotation marks (e.g., "pickup truck"). The Flickr Tag Cloud, Flickr's user-tag "folksonomy," generates a good link navigation system for both users and search engine spiders.
- This should be obvious, but have your photos publicly viewable, not restricted to viewing by only your friends and family.

- Create a descriptive title for the image. This adds yet more keyword weight to the photo's page within Flickr.
- Enter a description under the photo, or write something about the picture.
- Consider adding a note or two directly onto the photo, particularly if it is humorous. Flickr allows you to select a rectangular area and associate some text with it that will appear as a tool tip when users mouse over it. Adding a humorous/interesting note or two may encourage users to participate on your photo's page, and the greater the level of participation/stickiness of a page is, the better quality score that page may attain.
- Ensure photo geo-tagging is enabled (where applicable) for locally-oriented images
- Create thematic sets for your photos, and add each picture to the set(s) appropriate for it. This provides yet more contextual clues for search engines regarding the content of the photo's page, and it will allow a user arriving at the page a way to easily look at similar pictures you've taken.
- Browse through Flickr's Groups for public "photo album" collections that are dedicated to pictures that could be related to your photo. Sometimes it helps to search for photos using keywords you have used in your tags, and then see what groups other people's photos belong to.
- Join those groups, and then add your photos apropos to each group's theme. The more links there are to your photo page, the more important your photo will be considered, so add it to a number of groups. Ideally, add it to groups that have a lot of members—the number of members indicates the popularity and traffic of the group.
- Link each of your Flickr photo pages to your website, or a related page on your site. You can add hyperlinks to the Description field below the photo. Use anchor text that has a call to action, or that tells the user what to expect if he clicks on the link (e.g., "We sell this product on our website"; "Enjoy this view from the tables at our restaurant"; "This room is available at our bed & breakfast"). It is best to link to specific pages of related content as a richer indicator for link juice transfer.
- Post as many optimized pictures as possible. This is mostly a game of many small fractions adding up to large, cumulative results. The many pages of pictures linking back to your site will help build your overall authority. The more pages you have, the more likely it is that other Flickr users will find your content and link to it. This also increases your chances that a lucky picture or two might find its way onto a viral popularity wave that spurs many users to send links of your picture to their friends, or that a reporter might find one of your pictures ideal for his news story.

Many of the suggestions above can generally be applied to images used on third party social media platforms that heavily leverage images, such as Pinterest (<http://www.pinterest.com/>) - which has become one of the leading referrers of e-commerce traffic on the web. For these types of sites, be sure to use descriptive filenames and alt text, write keyword-rich descriptions, and make ample and appropriate use of hashtags.

Optimizing for Google Shopping Search

Google Product Search is now called Google Shopping. In the 2014 Searchmetrics study, 2013 saw a 182% increase in the total number of Shopping integrations in Google SERPs over the previous year. The biggest change within Google Shopping (besides the name) is that all product results now come from product listing ads - sponsored results for which Google collects a fee for every click. Google Shopping results have gained significant prominence within the SERPs, frequently appearing above the organic results and even pushing text ads down the page. This makes it especially important for online retailers to optimize product visibility in Google Shopping.



Google Shopping search results for "canon powershot". The results include a grid of Canon cameras with their prices and ratings. The cameras shown are:

- Canon PowerShot S... \$199.99 (4.5 stars, 666 reviews)
- Canon PowerShot E... \$179.00 (4.5 stars, 167 reviews)
- Canon PowerShot E... \$149.99 (4.5 stars, 167 reviews)
- Canon - Powershot El... \$179.99 (4.5 stars, 99 reviews)
- Canon PowerShot A... \$109.99 (4.5 stars, 1k+ reviews)

Below the grid, there are links to Canon PowerShot cameras on Canon.com and Canon USA, Inc. websites.

Figure 10-12 – Google Shopping result for “canon powershot”

Getting into Google Shopping

There are two main steps to getting your products listed on Google Shopping:

1. Submit a product feed through the Google Merchant Center (<http://www.google.com/merchants/merchantdashboard>).
2. Create Product Listing Ads through Google AdWords (<http://www.google.com/adwords>).

Submitting a Product Feed

To be included in Google Shopping Search, you need to upload a product feed that contains true physical/tangible products. The Google Merchant Center will accept other types of items, such as flights, hotels, car rentals, travel packages, and real estate, but such items will not get into Google Shopping Search.

The latest technical specifications for submitting a product feed to the Google Merchant Center can be found online at <https://support.google.com/merchants/answer/188494>. It is crucial to follow these requirements closely and to provide as much data as possible by providing all the relevant data to ensure Google can read what is in the feed and understand what products you’re selling; the more data Google has about your products, the better they will be able to show your listings at the right time. In your feed,

populate as many fields as possible with available product data (Brand, Category, Color, Price, Condition, etc.).

In addition to uploading accurate data, it is important to keep your feed data current, with the latest availability, pricing, and other specifications, to reduce the error rate and improve the chance that Google trusts your data (many merchants update their feeds daily). If the data in your feed does not match what Google finds on your landing pages, your feed quality will be lower, and it will be less likely that your products will appear in the SERPs for product queries. When in doubt, check your Google Merchant Center account to see how many errors Google finds in your feed and correct as many as possible - or delete products that cause errors from the feed.

Optimizing a Product Feed

Once your product feed is accepted by Google, they use it to match your products to searches for those items. The extent to which Google's search engine correctly understands what products you sell is largely dependent on what you've put in your feed, and various optimization tactics can help you receive more impressions for relevant, product-related queries.

1. Optimize The Title

Titles should contain your target keywords and be descriptive and concise. For example, for apparel products you should include the brand, style, and gender of the product. While you're allowed 150 characters, try to fit the most important words in the first 70 characters in case Google truncates it. A strong title that will fit nicely in the limited space provided by ads might be "Red Wing Men's Oxford 133". Avoid overly descriptive or editorial language like "High Quality & Gorgeous North Face Down Jacket". Google prefers simple and product focused language because that's what users prefer to see in the Shopping results.



Figure 10-13 – Google Shopping result for “microwave ovens” showing product titles

2. Optimize the Description

The optimal description length is between 500 and 1000 characters, much shorter than the 5000 character limit of the field. Like your product title, product descriptions should not be overly editorial or promotional, but rather include your keywords and specific attributes of the product - preferably near the start of the field in the event it gets truncated. The description should be grammatically correct and end with a punctuation mark. For products with variants (e.g. items that come in various colors), avoid listing the specifics of any of the other variants, and instead state that the product comes in other colors. This can prevent

inconsistencies between images and descriptions - for example, if you list a shirt and say it's available in green and red, Google may show the red shirt to a user searching for a green one.

3. Use Correct Product Identifiers

Unique Product Identifiers or UPIs are some of the most important attributes in your feed, yet also the ones most often submitted incorrectly. There are three types of UPIs: GTIN, MPN, and Brand - and most products require two of the three be submitted using industry standard values.

The UPI that you submit for one of your products should be the same value that a competitor submits for the same product because these numbers are universal and Google uses them to match products to queries. If you're making up UPI's or submitting the 'Identifier Exists' attribute for custom goods incorrectly, Google can't group your products correctly and you could be missing out on potential impressions.

The Identifier Exists attribute should be used for custom-made goods only; if your goods are not custom made, you must get the proper UPI from the product manufacturer as Google will disapprove products that use incorrect UPIs.

4. Use The Correct Google Product Category

Download the Google Product Taxonomy (<https://support.google.com/merchants/answer/1705911>) and make sure every product is categorized as specifically as possible. Using a more generic category won't get you more visibility. Google's goal is to show the most relevant results and by giving them the most specific categorization they can better group your product with similar products and map it to the most relevant queries. Google Product Categories are similar to UPIs in that you and your competitors should be submitting the same values for each product.

Google updates the product taxonomy every 3 months, so check back periodically to make sure that your products are still well categorized. If you find that your product is not showing for the right queries, experiment with the title, description, product identifiers and categorization to improve your relevance for those missed queries.

5. Use Appealing Images

The image is the biggest visual element that users see in Shopping results - so you should submit a visually impressive, high resolution image that clearly shows the product. You can submit up to 10 additional images in addition to the main image to show the product from different angles, in different settings, or with different packaging, and if you find that your listings have a low CTR, experiment with image rotation.

6. Submit a Special Offer with a Promotion Feed

When you're selling the same product as multiple competitors, the availability of a special offer is one of the few ways to differentiate your product listing - so use the Google Merchant Center to submit a separate feed with promotions such as discounts, free shipping, or free add-ons.



Figure 10-14. Google Trusted Store “Special offer” example in Google Shopping

7. Become a Google Trusted Store

Google adds a special logo next to any merchant who is a Google Trusted Store, giving Trusted Store product listings more prominence at no additional cost. Details about this program are available online: <http://www.google.com/trustedstores/>.

Product Feed Promotion in AdWords

Since all shopping results are now sponsored, you will need to create product listing ads for all of the products you want to promote. To get started, create a new Shopping campaign in Google AdWords, and connect this campaign with your Merchant Center account. You can do a quick and simple setup with a single campaign, ad group, and product group that contains all products in the feed; more sophisticated setups take a bit more time and effort, but are well worth it for the additional level of control they provide over bidding and targeting.

To gain more control over how much you pay for a click and when your products appear in the SERPs, you can put product listing ads into different campaigns, ad groups, and product groups – giving you greater control over settings like geo targeting, device bid modifiers, day parting, and negative keywords. Different bids can be set for each product group so that you're bidding the correct value for each type of product.

If you create multiple Shopping campaigns in an account, you can take advantage of an entirely new campaign setting that's exclusive to Shopping campaigns: the ability to set different priorities: high, medium, or low. In the event that multiple campaigns are eligible to show a shopping result, Google will show the product listing ad from the campaign with the highest priority. For example, you can use this to set a higher priority for products that are on sale.

Reporting Results of Shopping Ads

AdWords has reports about all the key metrics such as clicks, impressions, cost, conversions, and conversion values driven by Shopping listings, helping you ensure your ads are successful. Reports also include benchmark data about the average CPC and CTR of competitors to help you see how your shopping ads perform in comparison.

Optimizing for News, Blog & Video Search

News, video, and blog search provide tremendous opportunity for optimization. This has a bearing not only on obtaining traffic directly from the search engines, but also in promoting your business in whole new ways.

Getting plugged in to news search, for example, can bring you plenty of quality traffic, but it can also result in your site being spotted by major media editors and writers, as well as bloggers who consume that media. This is the type of exposure that can lead to numerous links.

Blogs and RSS feeds offer a similar dynamic of getting your content in front of new readers through new channels. There is also a social aspect to blogging, due to the built-in mechanism for comments and the tendency for bloggers to interact heavily with each other.

Optimization for news and video search are areas that apply to a wide range of sites. We'll start by looking at the optimization for Blog search.

Blog Search Optimization

Blogs are great publishing platforms for those who want to write articles on a regular basis. First, they make it easy to publish the content. Authors only need to log in and use a relatively simple set of menu choices to input what they want to publish, preview it, and then proceed to publish it. It is far easier than coding your own HTML pages by hand.

In fact, it is so easy that entire websites have been built using WordPress as the sole publishing platform. These sites are also typically easy to set up and configure. The world's most popular blog platform is WordPress (<http://www.wordpress.com>), and the search engines historically have handled Wordpress-based sites very well.

Blogs are inherently social in nature, and a host of social marketing benefits come from blogs. Enabling comments allows for interaction with your readers, and bloggers tend to have a significant level of interaction. Working this aspect of blogging as a social media platform is beyond the scope of this book. Nonetheless, be aware that a blog is an opportunity to establish yourself as an expert in a topic area, and to engage in a give-and-take activity that can dramatically change the visibility of your business.

In addition to these huge benefits, blogs can also bring you search engine and/or blog search engine traffic when they are properly optimized.

Structural blog optimizations

As we have discussed throughout this book, there are many key elements to successful SEO. These include things such as title tags, heading tags, good content, inbound links, and SEO-friendly architecture. Although the various blog publishing platforms are great, they can sometimes also require tweaking to achieve optimal SEO results:

- Blogs usually offer the ability to categorize each post. Make sure the tag name is used in the title of that tag page.
- Override default title tags with custom ones. You can do this using one of the best Wordpress SEO plugins available, "Wordpress SEO" by Yoast (<https://yoast.com/wordpress/plugins/seo/>) Along with many other SEO features, this plug-in allows you to supply a custom title tag, defined through a custom field in a post or a page.
- Rewrite your URL to contain keywords, and to use hyphens (preferred over underscores) as word separators. Do not let the blog platform include the date in the URL.
- Make sure you 301-redirect from <http://yourblog.com> to <http://www.yourblog.com> (or vice versa). Note that if you have a site at <http://www.yourdomain.com> and a blog at <http://www.yourdomain.com/blog>, you may need to implement a separate redirect just for the blog. This has to be handled not just for the home page, but also for all internal pages (e.g., permalink pages). Each URL must redirect to the corresponding URL on the *www* version.

- If you change from one blog platform to another one, the URL structure of your blog will likely change. If so, make sure you maintain legacy URLs by 301-redirecting from each page's old location to the new one.

Optimizing your anchor text

Anchor text is just as important in blogging as it is in general SEO, and should also be treated with equal amount of caution (ie, not over-optimized). Here are some guidelines:

- Make the post's title a link to the permalink page. You do not want your only link to the post to say "Permalink."
- Use a tool such as Open Site Explorer (<http://www.opensiteexplorer.org>) or Majestic SEO (<http://www.majesticseo.com>) to see who is linking to your site. Using these tools or tools like them, you can see who is linking to you and what anchor text they have used. Look for opportunities to request revisions to anchor text on inbound links, but before making such a request make sure you are comfortable that your relationship with the linker will not result in their simply removing your link instead of changing it.
- Internally link back to old, relevant posts within the body of a blog post. Don't use *here* or *previously* or similar words as the anchor text; use something more descriptive for the user instead.

Sticky posts

Sticky posts are a way to add content that always shows up first on a page of your blog. One way to use this is to write a post that is the introduction/overview for all the content on one of your category pages. Using this technique, you can add keyword-rich introductory copy to a category page or tag page. A tutorial on how to create and style sticky posts from Wordpress developer and resource site WPMUDev.org can be found here: <http://premium.wpmudev.org/blog/daily-tip-how-to-style-a-wordpress-sticky-post-to-make-it-stand-out/>.

Author profile pages

If you have a multi-author blog, another smart tactic is to create author profile pages. This has a multitude of benefits. First, many of your readers will have a stronger interest in one of your writers than the others. Better still, offer RSS feeds on a per-author basis. In addition, the authors are likely to link back to their author pages on your blog from their own sites.

Figure 10-13 below is an example of an author profile page (<http://searchengineland.com/author/elisabeth-osmeloski>):

Search Engine Land



SUBSCRIBE TO SEARCHCAP
Your email here.
[SUBSCRIBE](#)



Elisabeth Osmeloski

Elisabeth Osmeloski is Director of Audience Development for [Search Engine Land](#) and [Marketing Land](#). She is responsible for increasing readership through owned, earned and paid media channels. In addition, she assists in programming sessions at Third Door Media's [Search Marketing Expo](#) conference series and manages speaking engagements for editorial staff. Follow her on Twitter [@elisabethos](#).

[Subscribe to RSS Feed](#) [Send Email to Author](#)



2014 SMX Social Media Marketing Conference Preview

SUBMIT A TIP

Got a tip? We respect your anonymity.

[SUBMIT IT HERE!](#)

ATTEND OUR CONFERENCES



SEARCH MARKETING EXPO

The world's largest internet marketing conference series comes to Las Vegas November 19-20 for SMX Social Media Marketing. [Learn more here](#)

Or attend an SMX near you. [See all the dates and locations](#).



[[figs/print/1015.png]]

Figure 10-15. Example of an author profile page

Links remain critical

Obtaining links and managing your link juice remain critical activities. Blog platforms provide limited ability to manage your internal link juice, so this may require some customization to accomplish. Fortunately, in the WordPress environment some really good plug-ins are available to help you with this.

For starters, you can also create cross-links between related posts using a plug-in such as the WordPress Related Posts plugin (<https://wordpress.org/plugins/wordpress-23-related-posts-plugin/>). This is a great way to get people who just finished reading one post on your blog to consider reading another one. Finally, develop relationships with other blog and article authors, and get them interested in your content; hopefully, they will follow your work and offer relevant citations when the opportunity arises.

News Search Optimization: Getting in to Google News

As a major traffic driver to news-based web content, Google News should be at the top of your list for news search optimization. Google News offers publishers visibility beyond just the Google News site, as Publisher articles will also show up in regular Google searches. Certain queries around headline news and current events will trigger Google to show a fresher set of results. Some of these articles are presented in a cluster and labeled "In the news.", which gives publishers the opportunity for their most recent articles to appear at or near the top of the results page.

There are many types of queries that will trigger a "news box." Examples are disasters (tsunamis, earthquakes, storms), attacks (school shootings, terrorists), elections, holidays, significant sporting events (Olympics, World Cup, Superbowl), large brand names, and deaths of notable persons. Figure 10-15 below shows a news box inserted at the top of the SERP for the query *ebola*:

ebola

Web News Images Videos Books More ▾ Search tools

About 25,500,000 results (0.25 seconds)

In the news



Ebola case in NYC brings demands for quarantines

[Yahoo News](#) - 1 hour ago

NEW YORK (AP) — The case of the U.S. doctor stricken with Ebola left lawmakers on ...

[From Coffee Stand to Bowling Alley: Retracing Steps of Doctor Who Tested Positive for Ebola](#)

[ABC News](#) - 1 hour ago

[Ebola outbreak: Get up to speed with the latest](#)

[CNN.com](#) - 1 hour ago

[More news for ebola](#)

WHO | Ebola virus disease

[www.who.int/mediacentre/factsheets/.../en/](#) World Health Organization

WHO fact sheet on Ebola haemorrhagic fever: includes key facts, definition, transmission, symptoms, diagnosis, treatment, prevention, WHO response.

Ebola virus disease - Wikipedia, the free encyclopedia

[en.wikipedia.org/wiki/Ebola_virus_disease](#) Wikipedia

Ebola virus disease (EVD; also Ebola hemorrhagic fever (EHF)) or simply Ebola is a disease of humans and other primates caused by ebolaviruses. Signs and ...

NYC physician tests positive for Ebola - CNN.com

[www.cnn.com/2014/10/23/health/new-york...ebola-case/index.html](#) CNN

17 hours ago - A Doctors Without Borders physician back from West Africa tests positive for Ebola at a New York hospital.

[[figs/print/1016.png]]

Figure 10-16. Example of a “news box” in a Google SERP

Until very recently, news clusters in search results contained only articles from approved publishers accepted into Google News. In October of 2014, Google began including other sources in the news box as well, such as Reddit, Twitter or very small niche sites.

GETTING INTO GOOGLE NEWS

Publishers must meet very strict quality standards to be accepted into Google News and these standards are strongly enforced (both for getting in and for staying in). In Google's News Publisher Help section, Google outlines their News requirements in detail: <https://support.google.com/news/publisher/answer/40787>. Publishers should read through these guidelines very carefully before applying.

Here are some of the general requirements:

- Your site must contain clearly accessible contact information including a physical address, phone numbers and email.
- If your site has a common publication name, there is likely already a publisher in Google News with the same name. You will need to submit your site with a unique name. For example the Star Tribune submitted their publication name as "Minneapolis Star Tribune." Be sure to omit extraneous articles (such as "the") and descriptive clauses.
- If your site is available in multiple languages, separate requests must be submitted for each language version.

Applying to Google News

Once you are certain that your site complies with all the requirements, submit your application to Google News at the following URL:

<https://support.google.com/news/publisher/answer/40787#contact=1&ts=3179198>.

Be thorough and precise when filling out the application. It can take a couple of weeks for Google to review an inclusion request, and if the application is denied, no reason will be provided and you cannot reapply for 60 days – so be sure that your submission is as accurate and complete as possible.

PAYWALLS & SUBSCRIPTION SITES

Sites that require registration to view articles, or have a "paywall," whereby article access requires a subscription, are labeled as such in Google News:



[[figs/print/1017.png]]

Figure 10-17. Example of Google News content behind a "subscription" or "paywall"

Stories labeled as "subscription" tend to receive fewer clicks than articles that are free. This consequently reduces the click-through-rate (CTR) for that publication. Because CTR is a strong signal in the ranking algorithm, subscription sites will tend to get less visibility in Google News than sites that are free to users.

As an alternative to being labeled as a subscription site, Google offers a program called First Click Free, or FCF (<https://support.google.com/news/publisher/answer/40543>). With FCF, publishers must allow users referred from Google to bypass the paywall for that particular article. For any additional clicks during that session, the publisher can present the user with the paywall.

FCF requires that visitors from Google be permitted 5 free clicks each day, regardless of whether the paywall is hard (no free article views) or metered - article views allowed up to a certain number of articles viewed within a certain timeframe, typically one month. For example, *The New York Times* limits non-subscribers to viewing 10 articles per month. It is important to understand that First Click Free *overrides* the publisher's metered allotment. Once the user has viewed their monthly allotment of free articles, they must subscribe (or register) to view additional articles. However for sites that implement First Click Free, they must still allow this user to access the first article free if they are referred from Google and they must allow this 5 times each day.

Google News Publisher Center

Google News is different from Google Web search in that it is a static environment when it comes to content discovery. The crawler goes no more than one level deep from a set of pages that are entered into their database when a publisher is accepted into Google News. If your site launches a new section page or sub-section page, Google News will not find that page just from crawling a new link in your navigation. You must notify the Google News publisher help team and they manually update their system to add a new source page to your crawl.

For dynamic, growing publisher sites, this creates a tedious and often uncertain task of requesting frequent changes to your registration, which is a manual process. Google hates manual processes. Hence, the new Google News Publisher Center: <https://partnerdash.google.com/partnerdash/d/news>. This is a tool designed to facilitate the addition, removal and "labeling" of new site pages by allowing publishers to have some control over updating section pages and their labels.

How to add, edit, or delete a section URL

Figures 10-17 and 10-18 below show the screens for managing Sections in Google News:

All	Status	URL	Labels	
<input checked="" type="checkbox"/>	Active	http://www.exampledaily.com/religion/	Religion	Edit
<input type="checkbox"/>	Active	http://www.exampledaily.com/news/	Headlines	Edit
<input type="checkbox"/>	Active	http://www.exampledaily.com/business/	Business	Edit

Figure 10-18. How to Add/Delete a Section URL in Google News

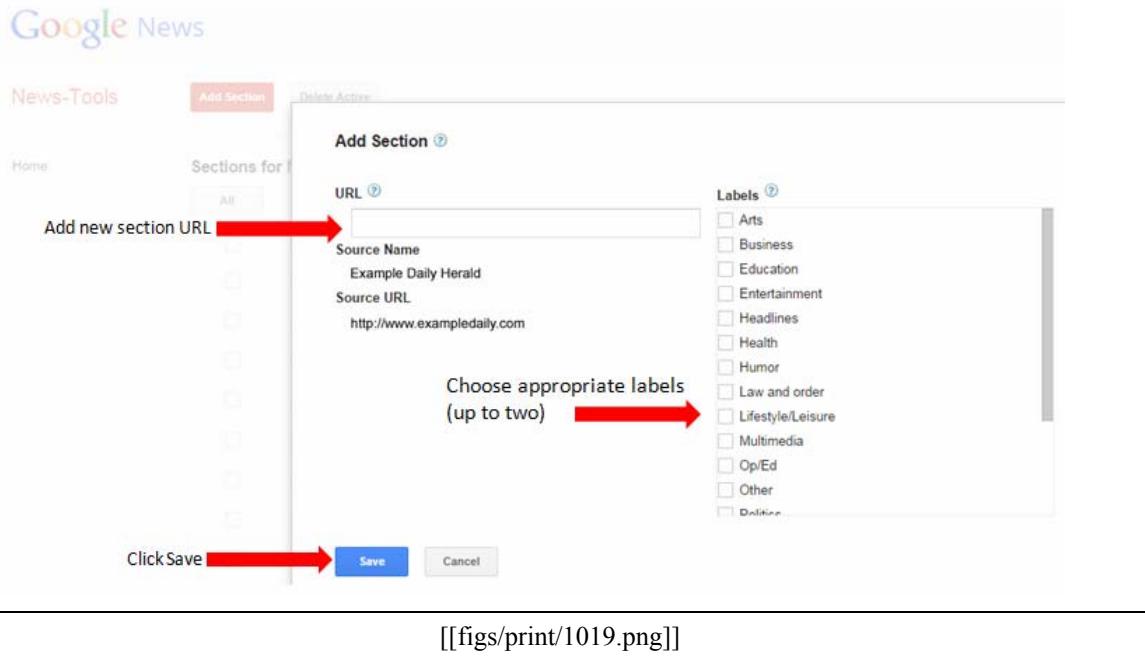


Figure 10-19. How to Add and Label a Section URL in Google News

How to update source details

Figure 10-19 below shows the screen for managing Source details in Google News:

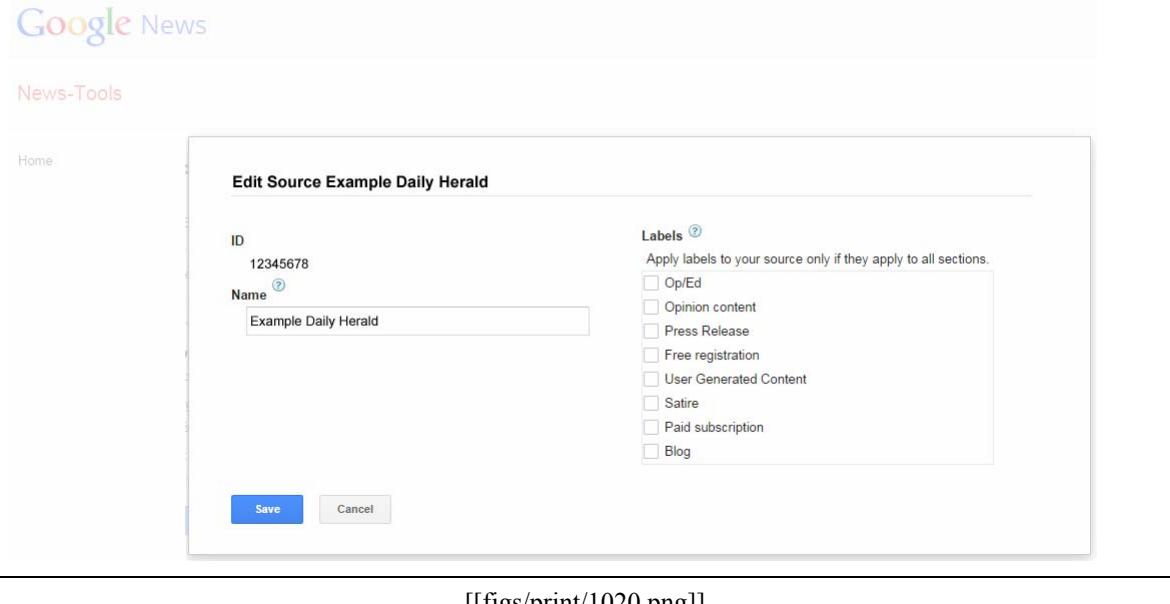


Figure 10-20. How to edit Source details in Google News

Technical Requirements

Google News has some unique technical content requirements, many of which are more particular than the Web search algorithm. For example:

- The text that makes up an article's date information and author byline should be the only text or code between the article's headline and the article's first sentence. It should be placed in a separate line of HTML between the article's headline and the article's body text.
- The URL for each article must contain a unique number consisting of at least three digits (<https://support.google.com/news/publisher/answer/68323>); however, this rule is waived for URLs submitted via News sitemaps.
- Articles must contain a minimum of 80 words. This means that short news briefs or pages with a video and a single brief paragraph are unlikely to be included in the news index.
- The HTML source page of a section or article must be less than 256KB in size. This can be an issue for pages that use infinite scroll or have a substantial number of user comments.
- Pages that display multiple articles at the same URL will not be included.
- Image links or links embedded in JavaScript cannot be crawled.
- Section and article pages must be on the same domain (or subdomain) as your main site (exceptions are videos on youtube.com and RSS feeds on feedburner.com).
- Sites encoded in UTF-8 are optimal and article pages must have an HTTP content-type of text/html, text/plain or application/xhtml+xml.
- If your publication mixes news with other types of content such as advice columns, how-to articles, weather forecasts, stock data, classified ads or paid advertorials, Google News expects you to prevent this content from getting crawled. There are three methods that can be used to accomplish this.
 1. Use a page-specific meta tag: <meta name="Googlebot-News" content="noindex,nofollow">
 2. Place all of your non-news content in different subdirectories (or subdomains) such as <http://example.com/jobs> or <http://jobs.example.com> and block Googlebot-News access to these sections in the robots.txt file.
 3. Create a Google News Sitemap for your news articles only and inform Google News (https://support.google.com/news/publisher/contact/report_issue_content) that you prefer to have your articles crawled exclusively from your news sitemap.

Headlines

- Headlines must contain at least 10 characters and be between 2 and 22 words in length.

- Do not link article headlines (linked headlines are a default setting in some site templates).
- Wrap the article headline in an `<H1>` tag. Ideally this will be the only `<H1>` tag on that page.
- The article headline should also be in the `<title>` tag, placed before any additional information such as the publisher name or the section.

Like this: `<title>This is the headline - Lifestyle - Daily Publisher</title>`

But not like this: `<title>Daily Publisher - This is the headline</title>`

And not like this: `<title>Lifestyle: This Is the Headline - Daily Publisher</title>`

- If the headline in the title tag is significantly different from the headline in the `<H1>` tag, this can prevent the story from being included in Google News.

Meta Tags

CANONICAL

If you publish the same article on multiple pages within your site, use the `rel=canonical` tag to specify which URL to rank for a story.

META DESCRIPTION

Google News does *not* use the meta description tag for snippets. Instead, the crawler looks at article body text near the headline (support.google.com/news/publisher/answer/70752).

There are two meta tags that are specific to Google News and do not apply to Web search. These are: `news_keyword` and `standout`.

NEWS_KEYWORD (<https://support.google.com/news/publisher/answer/68297>)

This tag is the only on-page keyword tag that is used for ranking by the major search engines. The `<news_keyword>` tag is placed within the page's `<HEAD>` tags. This tag can be used to specify keywords that are highly relevant to the article but might not be in the headline or the first paragraph of the story. For example an article with the headline "USA vs. Belgium Breaks Another Ratings Record" could use the tag like this:

```
<meta name="news_keywords" content="FIFA, soccer, football, World Cup 2014">
```

Commas are the only punctuation allowed in this field and should be used to separate each keyword or phrase. Up to 10 terms can be used for a given article and all keywords are given equal value, so the first keyword is not considered a stronger signal than the tenth keyword. Words that are contained in the headline or the first paragraph do not need to be repeated in the `news_keywords` tag.

STANDOUT (<https://support.google.com/news/publisher/answer/191283>)

When a publisher breaks a news story, it is important to Google News to attribute that article as

the original source, but algorithmically this has been a challenge. In late 2010 Google announced a new metatag named “standout” that publishers can use to designate a particular article as the source of a story or as an exceptional piece of journalism; Google provides the following criteria to determine when the use of this tag is appropriate:

- The article is an original source for the story.
- Your organization has invested significant resources in reporting or producing the article.
- The article deserves special recognition.
- The standout tag has not been used on your own articles more than seven times in the past calendar week.

When another publisher breaks a story and your organization writes about it, refers to it, or otherwise draws from it, Google strongly suggests using this tag in your article to credit that publisher by citing their URL in your standout tag. If your article draws on more than one piece of original or exceptional journalism, you can use multiple standout tag citations within your article. Citing standout articles from other publishers does not count against your limit of seven self-citations per week.

The standout tag cannot be added after the article has been published. This tag is placed in the page <HEAD> and can be coded as a meta tag or as a link tag.

```
<head>
...
<meta name="standout"
content="http://www.example.com/breaking_exclusive_story_2314"/>
</head>

<head>
...
<link rel="standout"
href="http://www.example.com/breaking_exclusive_story_2314"/>
</head>
```

When the href URL points to itself, Google interprets the tag as a self-citation. When the href URL points to another publisher's page, the tag is considered to be an out-citation.

ARTICLE EXPIRATION

This tag is treated as a removal request and is used to specify that an article should be removed from the Google index at a certain time. To function properly, the tag must be included in the article at the time that it is first crawled. It will take about a day after the removal date passes for the page to disappear from the search results. The date and time must be specified in the RFC 850 format using the following format:

```
<meta name="googlebot" content="unavailable_after: 25-Aug-2011 15:00:00 EST">
```

DATE

This tag can be used in addition to the date and time specified in the article byline. Use W3C format:

<meta name="DC.date.issued" content="YYYY-MM-DD">

Thumbnail Images in Google News

Google News includes a thumbnail image in almost every cluster displayed in either Web search results or on the Google News home and section pages. Having your article's image used as the thumbnail for a cluster can increase the amount of clicks that story receives.

News for **fifa**



Police Say Ticket Scam Involved a **FIFA** Official
Wall Street Journal - 9 hours ago
Rio de Janeiro police officials didn't name the **FIFA** official and said they are still working to determine his identity. **FIFA** said it is cooperating ...

Danish teen wins **FIFA** Interactive World Cup by outplaying ...
New York Daily News - 5 hours ago

Argentina Shuts Out Belgium 1-0: 2014 **FIFA** World Cup Live ...
ABC News (blog) - 1 hour ago

[[figs/print/1021.png]]

Figure 10-21. Google News article thumbnail image example

Images must be hosted on the publisher's domain (not an external domain or CDN) and must use a standard filename extension like .jpg.

Use the following guidelines to increase the likelihood of getting your article images included.

- Place images near the article headline
- Place images inline within the article content
- Use well-written captions
- Use images are fairly large in size (at least 60 pixels by 60 pixels).

You can find additional image information for Google News here:

support.google.com/news/publisher/answer/13369.

Re-Crawling

Google News does re-crawl articles, but most re-crawls occur only within the first few hours after discovery. Their system is generally better at detecting and displaying bigger updates, such as a substantial change to a headline or the lead paragraph. If an article is edited after being indexed in Google News, there is no guarantee that the article will get re-crawled and updated.

Google News Sitemaps

Google News uses the standard Sitemaps protocol, but with additional news-specific tags such as related keywords, article images, content genres and relevant stock tickers. Even though a

standard sitemap file can include multiple types of content, it is better to create a separate news sitemap for your news content. A news sitemap allows your content to be discovered more efficiently and comprehensively because you can feed all of your news articles to Google News even if they are not linked from a section page in the News database. If your site contains a mix of news and non-news content, you can use the sitemap to exert more granular control over which content is submitted. The news sitemap will accept URLs that do not include three numbers as required for discovery from crawling.

A news sitemap is particularly recommended for:

- Sites that have been recently added to Google News.
- Sites that often add new section pages or change existing section URLs.
- Sites that require users to follow several links to reach news articles.

Note: the Google Sitemap Generator *cannot* be used to create a News Sitemap. There are a number of third party tools recommended by Google that can help you to generate a Google News Sitemap: https://code.google.com/p/sitemap-generators/wiki/SitemapGenerators#Google_News_Sitemap_Generators.

News Sitemap Guidelines

- Upload the sitemap to your root directory.
- Update the sitemap immediately whenever a new article is published.
- Do not include articles older than 48 hours.
- For the publication date refer to the W3C format, using either the "complete date" format (YYYY-MM-DD) or the "complete date plus hours, minutes and seconds" format with time zone designator (YYYY-MM-DDThh:mm:ssTZD) format. Use the original date and time at which the article was published on your site; not the time the article was added to your Sitemap.
- A News Sitemap with more than 1,000 URLs must be broken into multiple sitemaps with a sitemap index file.

News Sitemap Fields

PUBLICATION NAME

Make sure that your publication name in the sitemap exactly matches the name you used when you submitted your application for Google News. If your publication name is *Daily Herald* but you submitted it as *Gotham Daily Herald* because Google News already contained a publication named *Daily Herald*, your sitemap publication name must also use *Gotham Daily Herald*.

ACCESS (PAYWALLS AND REGISTRATIONS)

The <access> tag takes one of the following values:

- Subscription: an article that requires users to pay to view content.

- Registration: an article that requires users to sign up for an unpaid account to view content.

For publishers that include a mix of free and paid articles, the only way to designate this on an article-by-article basis is within a news sitemap feed by using the <access> tag.

GENRES (CONTENT TYPES)

Values for the <genres> tag are required when applicable and restricted to the five options below. The tag can contain more than one value, separated by commas.

- Satire: an article that ridicules its subject for didactic purposes.
- OpEd: an opinion-based article that comes specifically from the Op-Ed section of your site.
- Opinion: any other opinion-based article not appearing on an Op-Ed page.
- PressRelease: an official press release.
- Blog: any article published on a blog, or in a blog format.

KEYWORDS

A comma-separated list of keywords describing the topic of the article. Keywords may be drawn from, but are not limited to, the list of existing Google News Keywords, provided by Google here: <https://support.google.com/news/publisher/answer/116037>.

IMAGES

The most effective way to get your images indexed is to include them in your news sitemap feed. Below is an example of a sitemap entry for a story with an associated image.

```

<?xml version="1.0" encoding="UTF-8"?>
<urlset xmlns="http://www.sitemaps.org/schemas/sitemap/0.9"
         xmlns:news="http://www.google.com/schemas/sitemap-news/0.9"
         xmlns:image="http://www.google.com/schemas/sitemap-image/1.1">

  <url>
    <loc>http://www.example.org/sample.html</loc>
    <news:news>
      <news:publication>
        <news:name>Example Times</news:name>
        <news:language>en</news:language>
      </news:publication>
      <news:publication_date>2014-10-23</news:publication_date>
      <news:title>President Announces New Education Initiative</news:title>
    </news:news>

    <image:image>
      <image:loc>http://img.example.com/story-image.jpg</image:loc>
    </image:image>
  </url>

```

STOCK TICKERS

For business articles, you can include a comma-separated list of up to 5 stock tickers of

companies that are the main subject of the article. Each ticker must be prefixed by the name of its stock exchange, and must match its entry in Google Finance:
<https://support.google.com/news/publisher/answer/13369>.

Example Google News Sitemap

```
<?xml version="1.0" encoding="UTF-8"?>
<urlset xmlns="http://www.sitemaps.org/schemas/sitemap/0.9"
         xmlns:news="http://www.google.com/schemas/sitemap-news/0.9"
         xmlns:image="http://www.google.com/schemas/sitemap-image/1.1">

  <url>
    <loc>http://www.example.org/business/article55.html</loc>
    <news:news>
      <news:publication>
        <news:name>Example Times</news:name>
        <news:language>en</news:language>
      </news:publication>
      <news:publication_date>2012-12-23</news:publication_date>
      <news:title>Companies A, B in Merger Talks</news:title>
      <news:access>Subscription</news:access>
      <news:genres>Opinion,Blog</news:genres>
      <news:keywords>business, merger, acquisition, Company A, Company B</news:keywords>
      <news:stock_tickers>NASDAQ:A, NASDAQ:B</news:stock_tickers>
    </news:news>
    <image:image>
      <image:loc>http://img.example.com/story-image.jpg</image:loc>
    </image:image>
  </url>
```

Validating a Google News Sitemap

The following XML schemas define the elements and attributes that can appear in a News Sitemap file. A News Sitemap can contain both News-specific elements and core Sitemap elements.

For News-specific elements schemas can be downloaded
from: <http://www.google.com/schemas/sitemap-news/0.9/sitemap-news.xsd>
<http://www.google.com/schemas/sitemap-news/0.9/sitemap-news.xsd>

For core Sitemap elements schemas can be downloaded
from: <http://www.sitemaps.org/schemas/sitemap/0.9/sitemap.xsd>
<http://www.sitemaps.org/schemas/sitemap/0.9/sitemap.xsd>

There are a number of tools available to validate the structure of your Sitemap based on these schemas. You can find a list of XML-related tools at the following locations:

- <http://www.w3.org/XML/Schema#Tools>
<http://www.w3.org/XML/Schema - Tools>

- <http://www.xml.com/pub/a/2000/12/13/schematools.html>

http://www.xml.com/pub/a/2000/12/13/schematools.html

To validate your News Sitemap file against a schema, the XML file will need additional headers as shown below:

```
<?xml version='1.0' encoding='UTF-8'?>
<urlset xmlns="http://www.sitemaps.org/schemas/sitemap/0.9"
  xmlns:news="http://www.google.com/schemas/sitemap-news/0.9"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xsi:schemaLocation="http://www.sitemaps.org/schemas/sitemap/0.9
    http://www.sitemaps.org/schemas/sitemap/0.9/sitemap.xsd
    http://www.google.com/schemas/sitemap-news/0.9
    http://www.google.com/schemas/sitemap-news/0.9/sitemap-news.xsd">
<url>
  ...
</url>
</urlset>
```

Submitting a Google News Sitemap

VIA ROBOTS.TXT

The location of your news sitemap file can be added to the robots.txt file.

Sitemap: http://example.com/sitemap_location.xml

This directive is independent of the user-agent line, so it does not matter where you place it in your file. If you have a Sitemap index file, you only need to include the location of the index file rather than each individual Sitemap.

VIA GOOGLE WEBMASTER TOOLS

Even if you specify the location in the robots.txt file, the sitemap should also be uploaded through your Webmaster Tools account. Doing this will provide you with access to detailed information should your sitemap contain errors.

How to submit a sitemap in Webmaster Tools

- Login into to Google Webmaster Tools or set up an account if you don't have one.
- On the Webmaster Tools home page, click the site for which you want to manage sitemaps.
- On the left-hand menu, click Crawl to expand the Crawl menu, and then click Sitemaps from the Crawl menu list.
- Click the Add/Test Sitemap button in the top right corner of your screen.
- In the text box that pops up, add the complete URL path to your Sitemap or the Sitemap index file.
- Click Submit.
- Refresh your browser to see your new sitemap in the sitemaps list.

How to test a sitemap

- Login into to Google Webmaster Tools or set up an account if you don't have one.
- On the Webmaster Tools home page, click the site for which you want to manage sitemaps.
- On the left-hand menu, click Crawl to expand the Crawl menu, and then click Sitemaps from the Crawl menu list.
- Click on a sitemap in the list to open the Sitemaps Details page and click Test Sitemap in the top right corner. Alternatively, if you don't see the Sitemap you're looking for, click Add/Test Sitemap on the main Sitemaps page, enter in a URL path in the textbox of the dialog that appears, and click Test.
- Once the test is completed, click Open Test Results to see your test results.
- Once you correct any issues identified by the test, you can click Resubmit to alert Google of any changes.

How to view the Sitemap Details page

- Login into to Google Webmaster Tools or set up an account if you don't have one.
- On the Webmaster Tools home page, click the site for which you want to manage sitemaps.
- On the left-hand menu, click Crawl to expand the Crawl menu, and then click Sitemaps from the Crawl menu list.
- In the list of sitemaps, click the sitemap you want to examine to open the Sitemaps Details page. From this page, you can view errors, indexing statistics and more.

News Crawl Errors

<https://support.google.com/news/publisher/answer/93994>

Publishers in Google News can view news-specific crawl error reports in their Google Webmaster Tools account.

- From the Dashboard, click Crawl > Crawl Errors.
- Click on the News tab to see crawl errors specific to your news content.
- Crawl errors are organized into categories, such as "Article extraction" or "Title error." Clicking on one of these categories will display a list of affected URLs and the crawl errors they're generating.

Videos in Google News

Only videos from publishers accepted in Google News will be included. Videos can either be embedded in an article with the YouTube player or publishers can create a YouTube channel and share it with Google News. To be included in Google News, YouTube channels and embedded videos must adhere to the following guidelines:

- Videos should be timely, reporting on current news events. No how-to videos, promotions, movie trailers, or music videos.
- Videos should be uploaded as quickly as possible. This will help them to be grouped with the most recent articles.
- Videos containing a single story rather than multiple segments are preferred.
- Video titles should be descriptive and specific (not generic like "Breaking News").
- The video description field should contain a lot of detail. Convey the who, what, when, where, and why of the story.
- The video keyword tags field is very helpful. Include the names of the people, places, companies and events mentioned in the video.
- Each video must be assigned to a category when it is uploaded. For general news coverage, select YouTube category "News & Politics" (which is youtube_category_id 25).
- Audio must be clear and easy-to-understand. Images should be in focus.
- All content must be original and unique. Music, images, and text within videos must belong to your organization or be used with legal permission.
- Videos must be embeddable on other sites to be displayed in Google News.
- Video that are geo-blocked in certain regions will not appear in Google News.

YOUTUBE CHANNEL GUIDELINES

- Only YouTube channels from publishers already in Google News will be included.
- Regularly updated channels are more likely to be accepted.
- A YouTube channel can be customized to reflect your organization's branding. In the channel profile section, include information about your organization and links to your site.
- If you have multiple channels, clearly identify each category (politics, business, entertainment, sports, etc.).

Once the YouTube channel is ready, inform Google by using their "Adding a New Section" form at https://support.google.com/news/publisher/contact/update_news_section.

- Add the URL of your YouTube channel in the field under "New section URLs."
- In the pull-down menu for the Category, select "YouTube channel" which is near the end of the list.

- Make sure you also include your publication name and URL at the start of the form.

Additional information about videos in Google News:

<https://support.google.com/news/publisher/answer/93985>

<http://googlenewsblog.blogspot.com/2010/05/making-your-video-news-more.html>

Editor's Picks

This feature allows publishers to showcase original, innovative news content with their publication logo. Editor's Picks are displayed in a sidebar module in the right column on the Google News home page and certain section pages. Editors' Picks can include content that Google News might not index otherwise such as long-form narrative articles, slideshows, interactive graphics, or video stories. Publishers can provide up to five links that represents the organization's best original journalistic work at any given moment.

Publishers can participate in Editor's Picks by creating and submitting a custom RSS or Atom feed. Up to 3 different feeds can be submitted - one for the home page, one each for the Technology and Business sections.

The following are a handful of important guidelines about submitting and maintaining this feed.

- Include only news content. No links to subscription offers, how-to articles, stock quotes, classified ads, weather forecasts, etc.
- Individual items in the feed can be older than 48 hours, as long as at least one item in the feed has been recently updated. Ideally, the feed should be updated once or twice per day. At least one additional item must be added to the feed every 48 hours. Feeds not updated at all within 48 hours will not be shown to users.
- Only feeds containing at least 3 articles are displayed. If there are more than 5 articles, only the top five articles will be shown.
- It is recommended to include an author name for each article. Without this information, the byline will be presented as "<Publisher> Staff."
- Headlines longer than 75 characters get truncated.

Once your feed is ready, submit it at:

https://support.google.com/news/publisher/contact/editors_picks.

Additional guidelines and instructions on Editor's Picks are available at:
<http://support.google.com/news/publisher/bin/answer.py?answer=1407682>

A special thanks to Alex Bennert (<http://www.linkedin.com/in/alexbennert>) for her contribution to the Google News section of this chapter.

Mobile, Video/Multimedia Search

Mobile and video search are two areas experiencing rapid growth. Each has its own challenges and requires a unique approach.

Video Search Optimization

Video search offers a tremendous opportunity for organic search traffic via your video content. The first component of video search optimization involves improving the ranking of your video content in Google's "Universal" search results – and the second involves improving the ranking of your videos in YouTube (the second largest search engine).

According to U.S. data from comScore qSearch for May 2014, Google had 12.5 billion explicit core searches that month, YouTube had 4.0 billion search queries, Microsoft had 3.5 billion searches, and Yahoo! had 1.8 billion. So, even if most SEOs know YouTube is the second largest search, they still focus on Google.

An analysis of Google universal search results by Searchmetrics (June 18, 2014) found that videos appear in 65% of Google searches in the U.S. However, 54% of these video results were from YouTube, about 5% were from Vimeo, about 5% were from Dailymotion, and no other video provider got more than about 1%. In addition, the average first Video integration from YouTube was about two positions ahead of the average first position of a competitor's video.

So, although it may seem counter-intuitive, your best strategy for improving the ranking of your video content in Google universal search results often involves improving the ranking of your videos in YouTube first.

Video SEO for YouTube

Today, more than 100 hours of video are uploaded to YouTube every minute. That means 144,000 hours of video are uploaded to YouTube every day. So, you've got to optimize your videos, playlists, and channel to give them the best potential for success on the platform.

Although YouTube's algorithm doesn't change over 500 times a year like Google's algorithm does, the tips below should come with the advisory: *best if used before* the next algorithm update. For example, YouTube officially replaced "view count" with "watch time" as a ranking factor in October 2012. This adjusted the ranking of videos in YouTube search to reward engaging videos that keep viewers watching and benefit channels if their videos drive *more viewing time across YouTube*.

On the other hand, "relevance" has remained a ranking factor since the early days: i.e. before Google acquired YouTube for \$1.65 billion in October 2006. So, recognizing that some things will change, while others will remain the same, here are key strategies, best practices, and important tips for Video SEO (freshness date: July 2014).

Conduct Keyword Research

When conducting keyword research for Video SEO, one of the first tools that you should consider using is Google Trends (<http://www.google.com/trends/>). Google Trends enables you to take popular search queries and explore traffic patterns over time and geography. In March 2013, Google Trends added YouTube search data going back to 2008, making it a great tool to look at video trends. Just visit Google Trends and enter up to five terms or phrases in the search box and the default setting will display web search interest. Then, on the upper right, click on "Web Search" and choose "YouTube Search" for YouTube search

interest. You can slice by region, time, or category as well. You will quickly see – and be able to show others – how similar or different Google and YouTube search trends are. For example, web search interest in “Chevy” and “Chevrolet” were the same in June 2014. But, if you compared the YouTube search interest for these two terms, you would have seen that “Chevy” is almost twice as popular as “Chevrolet” that month.

Another trending tool to visit is YouTube Trends (<http://youtube-trends.blogspot.com/>), which shows you the latest trending videos and trending topics on YouTube and is a useful resource to see what’s trending on YouTube.

Optimize Your Metadata

YouTube uses metadata – your video’s title, tags and description – to index your video correctly. To maximize your presence in search, suggested videos, and promotion, make sure your metadata is well-optimized. This includes title, description and tags. Your title can be up to 100 characters long. Your description can include up to 5,000 characters. Your tags can be up to 120 characters. Use every one of these limits for each of your videos.

Title: Make it compelling – this is your video’s headline. If it showed up in a search, would you click on it?

- Always represent your content accurately.
- Offer keywords first, branding at the end.
- For serial content, add the episode number to the end of the title.
- Update video titles so they continue to grab views.

Tip: Avoid titles that trick viewers into clicking on the video. This will cause drop-offs in the first few seconds of your video and will negatively impact your video’s watch time.

Description: Only the first few sentences of your description will appear in search results or above the fold on a watch page – so make them count! Follow a template for all of your video descriptions to create consistency.

A video description should:

- Accurately describe your video’s actual content (transcript excerpts are ideal) in one or two sentences.
- Describe your channel and link to your channel page.
- Drive viewers to subscribe (and include a subscribe link).
- Link to other episodes or related videos and playlists.

Tip: While you may want to link offsite in your videos, keep in mind that this could affect your video’s watch time. Videos with lower watch times appear lower in search results. YouTube’s new “metadata defaults” feature allows you to create templates for your metadata and ensure important text or links are always included when you upload a video.

A video description may:

- Include your channel’s release schedule.
- Include links to time-codes in the video for long-form content.
- Include a recurring keyword tagline. The keyword tagline is a group of sentences that describe your channel. They should include several search-driven keywords. Repeating this tagline in episode descriptions will inform first-time viewers about your channel.

Reminder: Remember that it is a violation of YouTube Terms of Service to use misleading metadata on your videos.

Tags: Tags are descriptive keywords that will help people find your videos. Create a set of standard tags for your channel that can be applied to any video you publish (e.g. filmmaking, animation, comedy, “Funny Videos,” “Pet Videos,” etc.).

- Include a mix of both general and specific tags.
- Use enough tags to thoroughly and accurately describe the video.
- Update catalogue videos’ tags when new search trends emerge.
- Properly format tags to ensure proper indexing of your video.
- Include keywords from your title in your video’s tags.
- List them in order of relevance to the video and try to use the whole 120-character limit.

Tip: Update the metadata on older videos if title, tags or descriptions are not optimized. This can increase views even if the video’s been public for a long time.

Create Custom Thumbnails

Thumbnails show up in different sizes and formats all across the platform and outside of it. Make sure you’ve got a strong, vibrant image that pops no matter what size it is.

General guidelines:

- When shooting a video, take shots that will make great thumbnails.
- Always upload custom thumbnails with the video file.
- Make sure the thumbnail is not racy.
- Consider the legibility of your thumbnail at multiple sizes. Thumbnails change size depending on the YouTube placement and device.

Visual guidelines:

- Use visual cues (colors, images, shapes, personalities) that are consistent with your brand.
- Clear, in-focus, high-resolution (640px x 360px min., 16:9 aspect ratio)
- Bright, high-contrast
- Close-ups of faces
- Visually compelling imagery
- Well-framed, good composition
- Foreground stands out from background
- Looks great at both small and large sizes
- Accurately represents the content

Tip: Upload high-resolution thumbnails so they appear crisp and clear wherever viewers happen to see them.

Optimize Your Annotations

Annotations are clickable overlays that you can add to your YouTube videos. You can use annotations in a huge variety of ways – brands and partners are always thinking of new and innovative uses.

Annotations are a great way to encourage your viewers to engage with your video and take meaningful actions as a result. The key is to create an annotations experience that is conversational and interactive.

Use annotation in these cases:

- **Subscribe direct link.** Annotate to the “subscription confirmation” page to make it easy for viewers to subscribe to your channel.
- **Navigation.** Highlight your newest video, create a table of contents for long videos, or prompt users to enter a sequential playlist viewing experience (e.g. “Click here for the next video in this series”).

- **Drive engagement with calls to action.** Invite viewers to like, comment or respond to questions. Complement scripted calls to action in the video with textual annotations.

Annotation best practices:

- Avoid annotations along the very top of the frame. This is where your title will show if embedded.
- Don't obstruct the actual content.
- Don't bombard the viewer. This can feel "spammy."
- When appropriate, set annotations to open a new window when clicked. Be careful! Don't take viewers away from a video too soon.
- Annotations at the end of a video should open in the same window.

Tip: Measure the effects of your annotations in YouTube Analytics with the "Annotations" report.

Spotlight annotations: Most annotation types are self-explanatory. Spotlight annotations stand out because they allow creators to subtly create clickable areas within a video. The text only appears when a viewer hovers over it with the mouse; a light outline shows when the viewer's not hovering. This is a great way to include unobtrusive but clickable annotations.

InVideo Programming annotations: Unlike regular annotations, InVideo Programming allows you to promote both your channel and any one of your videos on YouTube across all your uploads. How InVideo Programming works:

- When promoting videos, it pulls in the thumbnail as the annotation. Videos with optimized thumbnails perform better.
- It allows you to upload a custom, transparent square image to promote your channel. Unsubscribed viewers who hover over the channel image can subscribe to your channel without interrupting the video.
- InVideo Programming annotations are viewable on mobile devices.

Optimize Your Captions

Providing captions makes your video content accessible to a wider audience. It also acts as additional metadata which helps your video show up in more places on the site. If your video is captioned for multiple languages, it will also be searchable in those languages.

Optimize Your Playlists

Playlists allow you to collect, organize and publish multiple videos together. This increases watch time and creates another asset that will appear in search results and in Suggested Videos. You can create playlists using your own videos, other videos, or a combination of both.

Use Playlists in these cases:

- Group a set of videos that you want viewers to enjoy in a single session or in a particular order.
- Organize videos around a theme or a tent-pole event.
- Separate multiple shows into playlists and feature on your channel.
- Combine your most-viewed videos with new uploads.
- Curate good brand-advocating videos (reviews, testimonials, hygiene, etc.) created by your community.

Playlist best practices:

- Choose a strong thumbnail for your playlist. Make it pop!
- If a playlist needs context, upload a short, snappy intro video or interstitial videos with a host. Create a hosted playlist.
- Make your metadata work for you. A strong title, tags and description will help people find your playlist.
- Use Playlist Notes to write conversational asides about individual videos.
- Use in-video messaging, annotations, end-cards and links to send viewers to a playlist.
- Feature your playlist on your channel page by creating a new "section."

Tip: To link to a video in Autoplay playlist mode, click the share button and then copy and paste that URL. The video link will launch the whole playlist.

Optimize Your Channel

Your channel is the face of your brand on YouTube. It allows you to collect and organize all your videos in one place. Several customizable channel features will help ensure that you're delivering the best experience for your viewers while making your brand more discoverable across YouTube.

A number of channel features will follow your videos across the YouTube site and on devices. Make sure that they effectively represent your brand personality.

Channel name:

- Pick a short, memorable channel name in line with your brand identity. This will appear widely across YouTube, so make sure it's the best representation of your brand.

Channel icon:

- Upload a square, high-resolution image to be your channel's icon across YouTube. The image will appear alongside all your videos on the watch page. (In most cases, you can use your brand's logo.)

Channel description:

- The first few words of your channel description appear most frequently across the site, so highlight your most important branding upfront.
- Include your upload schedule, especially if you host multiple content types or series.

Channel art:

- Channel art is your channel's primary branding across all devices. Create customized, visually-compelling channel art. Busy images don't scale well, so keep the image simple.
- Add website and social media links to the About tab. Include these links in your channel art to help tie your YouTube presence to the rest of your online brand.
- Check how your channel looks in search, related channels and the channel browse page. Do your channel icon, channel name and channel art do a good job representing your brand to potential fans?

Tip: Associate your official website with your YouTube channel.

Optimizing for subscribed fans

What to watch next

- Subscribed viewers see personalized recommendations of "What to Watch Next" based on their viewing history.
- Promote a video with InVideo Programming to make it appear first in the "What to Watch Next" recommendations.

Recent activity

- Your most recent feed posts will appear on your channel page in the recent activity feed and on the activity tab.
- Keep your feed active with uploads, likes and channel posts. This will give your fans another reason to come back regularly.

Tip: Your top section will show up most frequently across devices. Make sure it will draw in new viewers.

Optimizing for unsubscribed viewers

In the Browse view of the channel page, subscribed and unsubscribed viewers see different versions of your channel. The unsubscribed view is your first opportunity to convince potential fans to subscribe.

Channel trailer

- Enable the channel trailer for unsubscribed viewers. This video will auto-play, so tell new visitors what your channel's all about and why they should subscribe.
- Keep your trailer short, but make sure you ask the fans to subscribe!
- Show, don't tell. Give potential fans a taste of your best content.

Sections

- Sections organize videos, playlists and channels on your channel's Browse page. Subscribed and unsubscribed viewers will see them.
- Sections can divide your content by genre, theme, show or any other criteria. Effective organization will help your audience find the content most relevant to them.

Gain subscribers by promoting your channel

Once you've optimized your channel for new viewers, it's time to use YouTube tools to bring new users to your channel.

InVideo Programming – Channel Promotion

- Enable Channel Promotion in InVideo Programming to place your channel icon across your entire video library.

Annotations and calls to action

- Use end-cards or other clickable annotations to push viewers to your channel page.
- Vocal calls to action along with annotations or video description links can help increase click-through rates.

Linking with Google+

- Link your YouTube channel to your Google+ page to amplify the social reach of your videos, increase discoverability of your content and offer a more streamlined Google experience. It is also meant to enable new features for easier management of your YouTube channel.

Related Channels

- Related Channels are promoted channels populated by YouTube that appear on your channel page. Recommendations are based on channels that are similar to yours.
- Be sure to enable Related Channels. Disabling the feature will pull your channel from being promoted on other channels.

Featured Channels

- Featured Channels are channels you choose to promote on your channel page. For instance, highlight other channels in your brand's portfolio.
- If you are promoting a large number of channels, rotate through the list using the "shuffle" feature. This ensures all your channels get visibility.

Tip: Make sure your icon doesn't distract from the videos or overlap with existing annotations.

Optimize Watch Time

As mentioned earlier in this chapter, Watch Time is an important metric for Video SEO on YouTube. The algorithm for suggesting videos includes prioritizing videos that lead to a longer overall viewing session

over those that receive more clicks. Viewers benefit from more enjoyable content being suggested to them, and creators benefit from more focused, engaged audiences.

If you're making videos that people are watching well beyond the first click, those videos will be suggested more often.

- Be an effective editor: Create a compelling opening to your videos and then use programming, branding, and packaging techniques to maintain and build interest throughout the video.
- Build your subscriber base: subscribers are your most loyal fans and will be notified of new videos and playlists to watch.
- Involve your audience in your videos, encourage comments, and interact with your viewers.
- Build long watch-time sessions for your content by organizing and featuring content on your channel.
- Create a regular release schedule for your videos when uploading to encourage viewers to watch sets of videos over single videos.

Check YouTube Analytics

Measurement is key – both for defining success and optimizing towards it. YouTube Analytics enables you to track a number of your most important metrics. For example, The Traffic Sources Report shows you which search terms people are using to find your video. The Views Report identifies which videos have the greatest view times and view-through rates. And the Audience Retention Report indicates which videos are successful at keeping viewers watching.

The key is to decide what success looks like and then translate that into metrics that are relevant for YOUR brand. Most brands find it useful to select one metric for each of the three buckets that measure engagement with your content:

- **Audience.** Are you reaching the right audience? How well?
- **Expression.** Is your target audience engaging with your content? How much?
- **Participation.** Is your audience endorsing and sharing your content? How much?

Tip: The YouTube Analytics API can be configured to schedule regular reports on specific metrics.

Leverage Paid Advertising

While this book is focused on SEO, Video SEO can leverage paid advertising in a way that website SEO can't. In fact, YouTube enables you to ignite earned media with paid media.

Given the abundance of videos on the web, it's risky to assume that your content will be organically discovered by a large audience. That is why designing a solid plan to promote your content and ensure it's viewed by your target audience is vital.

“Going Viral” plays a key role in building your audience on YouTube, but unless you already have millions of subscribers, you'll need to seed your content when it launches. Indeed, when they don't have an existing subscriber base or engaged social following, many brands use paid advertising on YouTube to ignite sharing and accelerate audience building.

The YouTube platform allows you to optimize your investment in paid media by activating the virtuous circle that links Owned Media (your video content), Paid Media (paid video advertising) and Earned Media (“free” views obtained when people share the video ads).

This provides you with a major benefit from advertising on YouTube that you don't get from Google: Earned impact. According to YouTube, more than 6,000 campaigns generated at least one earned view as a result of every two paid views in the first half of 2014.

And the YouTube audience isn't shy about sharing brand content: Two thirds of YouTube users agree that "If there is a brand I love, I tend to tell everyone about it." Across the globe, they are twice as likely to be early adopters, agreeing that "I am among the first of my friends and colleagues to try new products." They are **1.8 times more likely to be influencers, agreeing that "people often come to me for advice before making a purchase."**

Indeed hundreds of campaigns get more than two earned views per paid view. Although you don't get that benefit with Google AdWords, you can get that benefit with Google AdWords for video, which powers YouTube TrueView video ads.

Video SEO for Google.com

Once upon a time, one of the best ways to improve your site's appearance in video search results was to make sure that Google knew about all of your rich video content. So, it made sense to submit a Sitemap to Google that included video information in a supported format, making the included video URLs searchable on Google Video.

Google Video was launched on January 25, 2005, after which Google bought former competitor YouTube on October 9, 2006 – and on June 13, 2007, Google announced that Google Video search results would begin to include videos discovered by their web crawlers on YouTube, on other hosting services, and from user uploads. On January 14, 2009 the Official Google Video Blog announced, "In a few months, we will discontinue support for uploads to Google Video." On April 15, 2011, Google announced via email that after April 29 they would no longer allow playback of content hosted on Google Video. However, it reversed that decision one week later to provide users with greater support for migration to YouTube. The service, which was renamed Google Videos, was shut down on August 20, 2012.

The remaining Google Videos content was automatically moved to YouTube – so for all intents and purposes, Google Video has been replaced by YouTube.

Embed YouTube Videos and Playlists

Today, one of the best ways to improve the ranking of your video content in Google universal search results is to upload it to YouTube – though there is still debate in the industry as to whether uploading your video to your own site first, for origination credit, provides longer-term site benefits.

Whichever approach you decide upon, these are instructions for uploading to YouTube:

Here's how to embed a video:

1. Click the Share link under the video.
2. Click the Embed link.
3. Copy the code provided in the expanded box.
4. Paste the code into your blog or website.

Here's how to embed a playlist:

1. Click Playlists on the left side of your YouTube page
2. Click the playlist title you'd like to embed.
3. Click the Share link.
4. Click the Embed link.
5. Copy the embed code that appears in the menu below.
6. Paste the code into your website or blog.

Use a Video Distribution Service

The analysis of Google universal search results by Searchmetrics cited earlier in this chapter indicates there's another way to improve the ranking of your video content in Google universal search results: Upload it to Vimeo or Dailymotion. Uploading your video content to YouTube makes it 10 to 11 times more likely that it will appear in Google universal search results – about two positions ahead of the average first position of a video from Vimeo or Dailymotion.

However, it makes much more sense if you upload your video content to YouTube as well as Vimeo, Dailymotion and other top video and social networking sites. To do this efficiently and effectively, consider using a video distribution service, such as OneLoad (<http://www.oneload.com/>). OneLoad is powered by its parent company, TubeMogul, a media buying platform for video advertising founded in 2006. You can leverage its Destination feature to easily distribute to custom sites as well as encode and create RSS feeds to syndicate your video anywhere. And videos distributed through OneLoad are tracked by TubeMogul's powerful analytics. OneLoad is the video distribution tool of choice for some of the world's largest brands, agencies, media companies, government institutions and non-profits. These organizations use this specialized service to get their videos watched by more people in more places.

Use a Google Video Sitemap

One of the last ways to consider using is a Google Video Sitemap. Creating a Video Sitemap is a way to make sure that Google knows about all the video content on your site, especially content that it might not otherwise discover via their usual crawling mechanisms.

However, as the analysis by Searchmetrics indicates, about 1% of the videos Google universal search results come from websites like MUZU TV, Howcast, Artistdirect, and eHow. In other words, creating a Video Sitemap may help Google know about your rich video content, but your videos still need to get a significant number of views or watch time to appear in Google universal search results.

When a user finds your video through Google, they will be linked to your hosted environments for the full playback. Search results will contain a thumbnail image (provided by you or autogenerated by Google) of your video content, as well as information (such as title) contained in your Sitemap. It's also worth noting that Google can't (or won't) predict or guarantee when or if your videos will be added to its index.

You can create a separate Sitemap listing your video content, or you can add information about your video content to an existing Sitemap—whichever is more convenient for you.

The Google video extension of the Sitemap protocol (<http://www.sitemaps.org/>) enables you to give Google descriptive information—such as a video's title, description, duration, etc.—that makes it easier for users to find a particular piece of content. Google may use text available on your video's page rather than the text you supply in the Sitemap's video content, if this differs.

Adding video information to a Sitemap

You can create a Video Sitemap based on the Sitemap protocol. You can also use an existing mRSS feed (<http://video.search.yahoo.com/mrss>) as a Sitemap. Or, if you're a belt and suspenders kind of guy, then you can use both.

Your Sitemap will need to include the following minimum information for each video: title, description, playpage URL, thumbnail URL and the raw video URL or URL to Flash video player. Without these five pieces of information, Google cannot surface your videos in our results.

Once you've created your Sitemap, you can submit it using Webmaster Tools.

Recognized file types

Google can crawl Flash SWF objects and the following raw video file types: .mpg, .mpeg, .mp4, .mov, .wmv, .ASF, .avi, .ra, .ram, .rm, .flv. All files must be accessible via HTTP. Metafiles that require a download of the source via streaming protocols are not supported at this time.

Conclusion

As we've outlined in this chapter, optimizing your site for mobile search, your vertical content for Universal search, engaging in Local SEO efforts, and, if you're a news publisher, participating in Google News... can provide significant exposure within organic search for your business. Not only can you gain visibility in the increasingly diverse search ecosystems, but you can increase the likelihood that your content assets are working for you in all areas of the web to promote increased user sharing and engagement – which provides a positive feedback loop for all of your SEO efforts.

11

Tracking Results and Measuring Success

Since the nature of what SEO practitioners do is not clear to many people, the expert practitioner knows how important it is to establish quantifiable goals and demonstrate results. Expectation management is key, and effective communication of timetables for deliverables and improvements is crucial to the success of any SEO project.

Web analytics and other data-mining tools can provide a rich array of data that is invaluable to the SEO process. This includes tools for discovering new links, finding long tail search terms, determining the best links your competitors' sites have, and more. All in all, a strong understanding of these tools will make you a better SEO practitioner.

However, thousands of tools are available, and it can be difficult to figure out which ones you need to use, which ones are merely nice to have, and which ones are not worth the time. It is a good idea to start simple and focus on establishing a baseline measurement of where you are before you begin your SEO efforts and how to measure the results. A proper baseline should account for as many variables as are applicable to your business and industry - baselines for organic sales, indexing, crawl rates, conversion rate analysis, and so on.

In other words, you need to measure whether, and to what extent, you have been successful. The first thing you need to do is define the goal(s) of the site. Here are some possibilities:

- Generate sales
- Obtain leads
- Generate ad revenue
- Get people to read specific articles
- Obtain donations
- Develop an audience (which can lead to the above)

These are all valid objectives that relate directly to the business. Missing from this list are things such as measuring rankings, measuring the total number of links, and measuring the average cost per acquired link (organically obtained links still have a cost: the cost of the marketing campaign to get them). These other measurements may be interesting, but they are not the end goal.

Closely connected with this are direct drivers of results. These include parameters such as:

- Total site traffic (all visitors)
- New vs. returning visitors

- Traffic to your most popular pages
- Pages receiving the search traffic
- Pages resulting in conversions down the line
- Traffic per page and conversion per page

The mechanics of measuring results in terms of the main business goal should always be the first set of measurements you put in place. From there you can diverge and look at other metrics that might help you diagnose problems with the site or give you additional insight into how to set up a campaign.

Why Measuring Success Is Essential to the SEO Process

Although quantifying deliverables and measuring progress are important for external reporting purposes, it is just as important for SEO practitioners to measure the efficacy of their own efforts to make timely adjustments as necessary. As you will see in this chapter, numerous tools are available to aid in this process.

At the beginning of any SEO project it is wise to establish baseline data points for the website. This includes the following:

- Quantifying organic search traffic by search engine
- Quantifying a breakout of what site content areas are getting the current organic search traffic by search engine
- Quantifying data on conversions broken down by search engine
- Identifying poorly performing pages
- Identifying best-performing pages (in terms of traffic and conversions)
- Tracking search engine crawler activity on the site
- Determining the number of indexed pages
- Determining whether the indexed pages are getting search traffic
- Determining whether best-selling product pages are indexed and getting search traffic
- Identifying 404 error pages and external sites linking to these pages, if any
- Determine whether you've been hit by a search engine penalty or algorithm update

You cannot methodically improve what you cannot measure; or as Peter Drucker put it "What gets measured gets managed". Defining and mapping the path toward concrete goals are crucial aspects of the SEO process—and over time, these goals may change - so it is also important to make sure the data you capture helps you understand your progress against these goals.

In the world of web analytics, this is referred to as picking Actionable Key Performance Indicators (or Actionable KPIs). The best data measurements are those that potentially result in an action being taken in response. Think of this as data-driven decision making.

It is important to recognize that SEO requires operating effectively within two distinct worlds: first, SEO-related metrics such as rank, visits, and backlinks rule the day, and second, the one likely occupied by the rest of the organization - business metrics such as sales, costs, returns, and contract renewals. As an SEO practitioner, you should address both spaces by clearly tying SEO metrics to business goals.

The Tracking Cycle: Produce, Launch, Measure, Refine

In summary, the basic process usually looks something like this:

1. Define an SEO strategy, determine an implementation schedule, and establish a clear understanding of goals. What are you going to accomplish, and what is the strategy for accomplishing it? How will you measure progress?
2. Discuss your strategy.

The marketing and business development teams are your allies here—you want to ensure that your SEO objectives are based on the overall business and site objectives, both long- and short-term.

3. Establish a baseline.

Now that you are about to start and you have decided how you are going to measure progress, establish a baseline by recording the current stats prior to beginning work. Make sure you don't get a false baseline due to seasonal factors or some other unusual event. Comparing year-over-year data will usually eliminate fluctuation due to seasonality. However, you must also consider how changes in the market, new competition, elimination of competition, industry consolidation, changes in your business strategy, and changes in the search engines themselves may have affected that year-over-year data.

4. Proceed with your project.

Implement the SEO strategy - new URLs, technical site changes, new content development, etc.

5. Collect data.

Collect the newest data for each metric you decided to focus on. Since SEO can take days, weeks, or months to show results, make sure you wait long enough for your efforts to have an impact. Many factors could influence the length of time you should wait. Here are some of them:

- If your site is brand new, it may take longer for your changes to take effect.
- If the scope of the change is drastic (such as a complete redesign incorporating new information architecture, new URLs, etc. vs. a simple visual re-skinning), the time to see results will probably be longer.
- Sites that get crawled at great depth and frequency will probably yield visible results faster.
- Sites seen as authoritative may also show faster results.

6. Compare the baseline data to the new data.

The new data has little meaning unless it is compared to your baseline. This is the time when you can really assess your progress.

7. Refine your campaign.

Now that you have compared your old data with your new data, you can make some decisions. Is the campaign a bust? If so, abandon it and move on to the next one. The old business axiom “Fail quickly” applies here. The faster you diagnose a failure and move on to the next thing, the better.

If you are achieving great results – a noticeable increase in traffic and/or conversions – then you've got a great process in place; apply it to more pages on your site, or look for ways to further improve conversions through content development and design. If you don't see any change in your traffic or conversion rate (or if the changes are measurably negative) within a few weeks, don't give up just yet. Perhaps you've overlooked something important, or there was an error in your implementation. Or perhaps you are recovering from a larger site overhaul. In this case, things may get worse before they get better. Be patient, and at this point of the process, focus on action-based goals instead of results-based goals. For very damaged or very large sites, the wait for improvement could be six months to a year.

How to Establish a Proper Baseline

It might sound easy to record a set of SEO metrics before the start of the project and then compare the same set of metrics a 30, 60, or 90 days to measure the progress. But what if you don't make any changes and the metrics reflect an improvement? What if the improvements are because of external business factors? How can your SEO project get attribution for its business impact?

Depending on the type of business and assuming there have not been drastic strategic shifts from the previous year, it is possible to account for the seasonality of the business if you compare each period to the corresponding period of the previous year. Then adjust for overall year increase (or decrease). You can measure by comparing a couple of months year to year where you have not done any SEO changes. The assumption is that everything else being equal, a performance increase in search should be attributed to your SEO efforts.

It is important to make sure that there haven't been any major strategic shifts in the business. Sometimes businesses completely change the way they do business and this dramatically impacts your ability to measure results. In cases like this, you can fall back to the less reliable approach of tracking increase in pages getting search clicks and overall increase in average ranking across the board.

One other issue which may significantly impact your ability to establish a baseline for measurement is a seemingly spontaneous drop in traffic. When analyzing the organic traffic to a site if you notice a large drop in organic traffic it's vital to determine the cause of this drop before proceeding with setting about creating a set of baselines and goals. Large traffic drops can be caused by a number of factors including a large scale site redesign or rebuild, a shift in the nature of the business as noted above, seasonal factors (which can usually be determined by looking at several years of data at a time), or possibly by organic search algorithm updates.

An extremely useful tool to use in determining if a large traffic shift might be caused by an algorithm update is the Panguin (a concatenation of the words panda and penguin) Tool created by Barracuda Digital (<http://www.barracuda-digital.co.uk/panguin-tool/>). This tool allows you to overlay your Google analytics organic traffic with Moz's history of algorithm updates to see if traffic shifts coincide with the updates. It's not always possible to definitively determine if you have been impacted by an algorithm update as occasionally several updates happen over a very short period of time or you may have a traffic drop due to another cause which simply happens to coincide roughly with the algorithm update. However, before embarking on an SEO project it's important to note the possibility of how an algorithm update may impact the potential success of the project.

Rank in Google is also important to monitor to understand what is working in our strategy. For example, a jump from 99 to 15 on a targeted keyword might not bring a substantial increase in traffic; however, it is a huge indicator that something you are doing is working. Using a tool like AuthorityLabs, a subscription SEO ranking tracker, can create a one-stop page for tracking all of the keywords your business is targeting and also track their shifts over time. AuthorityLabs can be found at <http://authoritylabs.com/>.

Using Analytics as a Business Case for SEO

You can use a properly structured plan as the business case for an SEO project. The way to do this is to express the target results of an SEO project in terms of financial impact. You could include a variety of metrics in a business case, such as:

- Revenue
- Lead generation
- Margin
- Branding value
- Reach
- Other action triggers (newsletter sign-ups, contact requests, demo requests, free-trial acceptance, viewing and/or sharing a specific piece of content, etc.)

Measuring such things requires that you tie organic search engine visits to the revenue and other conversions that result.

Measuring Search Traffic

Classic web analytics data is an incredible asset to SEO. Here are three examples of ways to utilize this data for SEO purposes:

- Look at your daily referral reports to detect newly received inbound links (the great majority of webmasters click on a link after implementing it, to make sure it works)
- Measure the results of your campaigns by tracking the increase in conversions you are driving over time
- Try to find missed opportunities – search terms and phrases that people used that did not result in a click, visit, or conversion

Web analytics are a must-have for any web publisher.

Basic Overview

Your hosting company most likely provides a free web analytics solution, such as AWStats (<http://www.awstats.org>) Webalizer (<http://www.webalizer.org>) or something similar. Although these tools provide valuable data, they are very limited in scope, and other tools out there provide significantly more data. Here are some of the best-known ones:

- Google Analytics (<http://www.google.com/analytics>)
- Woopra (<https://www.woopra.com>)
- Clicky (<http://www.clicky.com>)
- Piwik (<http://www.piwik.org>)
- Crazy Egg (<http://www.crazyegg.com/>)
- Adobe Analytics (<http://www.adobe.com/solutions/digital-analytics.html>) (enterprise level solution)
- IBM Digital Analytics (<http://www-03.ibm.com/software/products/en/digital-analytics>) (enterprise level solution)
- Webtrends (<http://www.webtrends.com>) (enterprise level solution)

Web analytics platforms track your site's traffic in two major ways. The older of the two methodologies is to use software that analyzes your web server log files after traffic has occurred. Setting up this type of solution generally involves installing the software on an internal server that is able to gain access to the log files.

The newer methodology involves implementing JavaScript tags on all of the pages of your website. Provided that you have a basic template for your pages, this generally is a fairly straightforward process. JavaScript then tracks the activity on your web pages as it happens and builds a data profile reflecting that activity.

Selecting the Right Analytics Package

Log file tracking and JavaScript tracking are equally valid methods, and each has its own strengths and weaknesses. The biggest advantage of the log file method is that you can track search engine crawler activity on your site. This is something you cannot do in JavaScript implementations, because search engine crawlers do not execute the JavaScript.

The second big advantage of a log file-based solution is that you run the software in-house, so no third party has a copy of a log file with your proprietary traffic data on it. This distinction can be a big security issue for some organizations.

Some other key advantages to log file analysis:

- Tracks users who don't have JavaScript enabled (or are using privacy or ad filters)
- Tracks media types beyond HTML
- Tracks partial requests (incomplete page loads and load errors)

Ultimately, most companies opt for JavaScript tracking because JavaScript offers a much greater level of flexibility than log files. You can tweak the JavaScript to do custom conversion tracking, or gather pages into logical groupings in a manner that cannot be done as easily in log file-based applications.

Some other key advantages of JavaScript tracking:

- Tracks outgoing link requests
- Tracks events that don't involve HTML requests, such as playing a video
- Records visitors' screen resolution

Some analytics packages, such as Webtrends, offer both options, and a combined solution. This kind of approach can bring you the flexibility and power of JavaScript, but still get you your search engine robot crawling data as well.

Making this decision is only the first step in picking an analytics package. We listed some of the more popular vendors earlier, and there are many more vendors than that. Each of these packages has different strengths and weaknesses. Not only that, they all do their counting a little bit differently. These differences result from different decisions in how the analytics packages conduct visitor tracking. None of them are right or wrong, they are just different in the exact thing they are measuring. For instance, they may count a "visitor session" within a certain timeframe.

The more important component of this is whether the functionality of the web analytics software fits your needs. You often do not know what your requirements will be until you have used analytics for a while. As you engage with analytics, you will continually uncover more things to investigate, and develop new requirements.

For many companies, one of the best solutions is to start with a free analytics package such as Google Analytics, then look to buy a higher-end solution once the free package has been pushed to its limits. By pushing the limits of these free analytics products first, you will end up developing a set of requirements you can use in deciding where to go next.

None of this is meant to say that you should not brainstorm your requirements in detail before selecting an analytics package. You should (and must). Just expect that you will develop new requirements along the way. Web analytics is a journey that unfolds over time.

Based on the requirements you establish in your upfront brainstorming, you may find that you require a set of features that the free packages do not provide. Use that knowledge to select the right package to start with.

Valuable SEO Data in Web Analytics

You can extract all kinds of data from web analytics. Here are a few of the more interesting types of information you may want to extract.

Traffic by search engine

One of the first things you may want to know is the breakout of traffic by search engine. Figure 11-1 provides an example of such a report in Google Analytics.

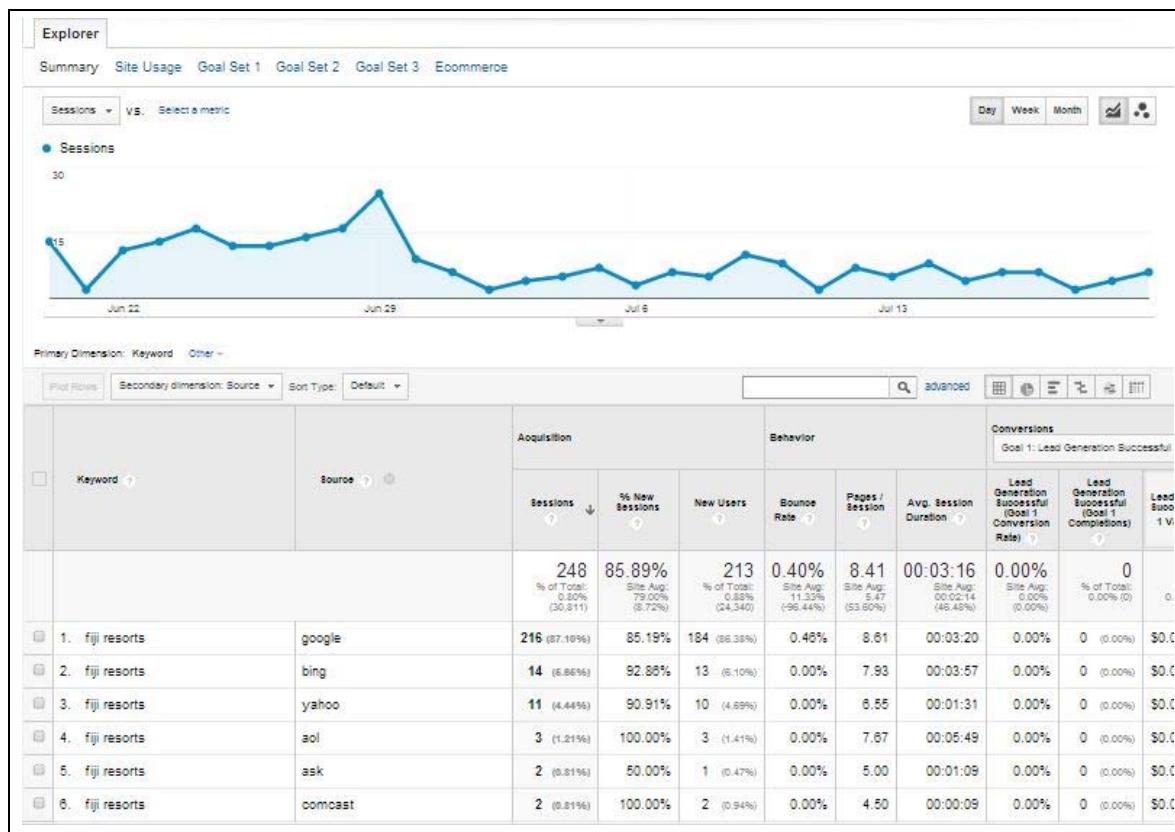


Figure 11-1. Traffic by search engine

Notice how small the traffic is from Bing and Yahoo! compared to Google. This may be indicative of a problem with how the site is indexed by Bing (which supplies the search results to Yahoo!); the Webmaster, designer, or developer might want to explore why the combined traffic from Bing and Yahoo! is so low. Or, it could just be the fact that Google by far has a dominant market share, 67.3% of the market share according to a Comscore report from August 2014. However, in your specific niche, that percentage may be far higher or lower, which is why browser choice can be such an important demographic data set.

Traffic by keyword

Beginning in October 2011 Google began “securing” its search traffic, hiding most keyword data for organic search. As of 2014, most search terms reported by Google Analytics appear as “not provided” (see figure 11 –2 Keyword not provided)

Primary Dimension: Keyword Source Landing Page Other -			
		Plot Rows	
		Secondary dimension	
		Sort Type:	
		Default	
		Acquisition	
		Sessions ↓	
Keyword		% New Sessions	
		1,969 % of Total: 81.23% (2,424)	
		81.62% Site Avg: 79.83% (2.24%)	
1. (not provided)		1,680 (86.82%)	
2. neopetsfanatic		38 (1.83%)	
3. neopets cheats		32 (1.68%)	
4. neopets fanatio		7 (0.38%)	
5. neopet cheats		6 (0.30%)	
6. neopets tips		5 (0.26%)	
7. how to get limited edition neopets		4 (0.20%)	
		100.00%	
		50.00%	

Figure 11-2. Keyword not provided

There are several ways to attempt to extract more detailed keyword data though none of them are as detailed as when Google used to provide you all of the data. One solution is to use keyword data from your Google Webmaster Tools (see figure 11-3 Google Webmaster Tools keyword data). You can now even access your Google Webmasters Data directly from Google Analytics. Directions in linking these two are available at <https://support.google.com/webmasters/answer/1120006?hl=en>. Google provides you with some estimates of the number of impressions, clicks, your click through rate, and your average ranking position for a variety of your keywords. Please note however with Google Webmaster Tools Google will not provide you with all of your search queries (though they will provide several thousand) and the data is only stored for 90 days (so get in the habit of downloading an export file every 90 days to maintain your archive). This information is found in your Google Webmaster Tools interface under Search Traffic – Search Queries. In addition to getting traffic estimates for web searches you can also segment the information by web/mobile/image, geographical location, or for only specific queries you select.

Note that by default, filters are different on Google Webmaster Tools and Google Analytics. If you're noticing differences in query results, try syncing up the filters. Instructions for doing so are available at <http://www.semclubhouse.com/getting-query-data-from-google-webmaster-tools-to-correlate-with-seo-query-data-in-google-analytics/>.

Query	Impressions	Clicks ▲	CTR	Avg. position
costa rica weather	12,946	1,544	12%	3.7
costa rica map	28,014	895	3%	4.3
guanacaste costa rica	13,319	821	6%	1.7
costa rica rainy season	1,058	673	64%	1.1
limon costa rica	6,884	393	6%	1.9
rainy season costa rica	464	346	75%	1.1
weather in costa rica	3,711	284	8%	4.7
map of costa rica	15,943	280	2%	6.0
rainy season in costa rica	342	250	73%	1.0
tourist attractions in costa rica	1,547	246	16%	1.4

Figure 11-3. Google Webmaster Tools keyword data

If you happen to participate in paid Google AdWords advertising you can also look at your “Matched Search Queries” report in Google Analytics to see the exact terms users searched on before clicking on one of your ads (<https://support.google.com/analytics/answer/2491146?hl=en>). This may help give you some idea of the types of terms people are using to get to your website. In addition to Google Analytics, Bing webmaster tools (<http://www.bing.com/toolbox/webmaster>) also provides some keyword data.

Although it is a paid tool, SEMrush (<http://www.semrush.com/>) can also provide some insights into what keywords your site ranks for and approximately how much traffic you might be receiving from each of those keywords.

Another paid solution for obtaining insight into keyword data is Searchmetrics’ “Traffic Insight” solution (<http://www.searchmetrics.com/en/traffic-insight-u92efa7d/>). “Traffic Insight” allows you to reconstruct “not provided” keywords and traffic data on a page by page basis by incorporating various data sources and running them through a proprietary algorithm to allow you to “reverse engineer” much of your “not provided” data.

One of the data points of interest for an SEO practitioner is what search terms are bringing traffic to the website. This provides a quick way to see where the SEO campaign is going well and where it is not going so well. You can also use this to spot opportunities where a key search term is providing some traffic, but not as much as you would expect if you were ranking highly for that term. Remember that search terms are no longer simply considered in isolation. If you spot an opportunity to increase traffic on a key search term you should really think about what other related terms might also drive traffic to the same page as the keyword you are considering.

You can then look to see where you are ranking for that term. Perhaps you are in a lower position on the first page, or on the second page of the SERPs. If so, it might make sense to focus some attention on your page which is the most targeted to providing information related to this term. With a little effort, such as a focused promotional campaign for the page in question, and improving and expanding upon the information you provide related to that term, you may be able to move up several positions and obtain a traffic boost. By improving the content on the page to better satisfy the user’s query you may also begin ranking on other related terms which Google understands are related to the one you happen to be looking at at the moment.

A traffic-by-keyword report can also show you the long tail of search as it relates to your current site. Figure 11-4 depicts a snippet from the organic search phrases report from Google Webmaster Tools showing some of the lower volume terms for ArtofSEObok.com.

seo book	2,158	4
the art of seo book	33	3
the art of seo 2nd edition	24	3
best seo books	159	2
seo art	40	2
art of seo book	57	1
the art of seo free download	10	1
art of	1	1
seo books	1,449	0
seobook	320	0
best seo book	170	0
middle east search engines	78	0
seo book.com	65	0

Figure 11-4. Long tail keywords

Notice that a lot of variations of "seo book" are showing up in the list, in addition to the rather interesting *Middle East search engines*. Looking into this query data can give you a broad perspective on opportunities for new content which answers the intent behind these queries.

Branded vs. non-branded traffic

Branded search traffic refers to queries that involve your company name, domain name, or product brand name. Non-branded search traffic can be much more valuable because it often represents new visitors to your site. Branded traffic will convert at a higher rate on average, but non-branded search traffic is often from people who did not previously know your brand, or at least they were not yet committed to it to the point where they searched specifically for a variation on your brand name or domain name.

In the past, it was possible to see the data on keywords driving traffic to your web site in great detail, as the search engines provided that data and you could readily examine it in your analytics software. Now it's common that 90% or more of the keywords for your organic search traffic will be reported as (not provided), meaning that it is difficult to get a strong measurement of what keywords are driving traffic to your site, or even the split between branded and non-branded traffic.

Number of Pages Getting Search Traffic

An indirect way of measuring effective indexation is to keep an eye on how many pages are getting search traffic. This number should represent a subset of the total pages indexed, but it is more valuable because these pages are not just indexed, they ranked highly and were interesting enough that visitors decided to click on the listing.

This is an important metric to track as you work on addressing site architecture issues like duplicate content and bot traps. As you remove search engine spider road blocks, the bots will find and index more pages. Those pages will rank and get search traffic and this number should increase. The same happens as you consolidate duplicate content and as the indexation of the site improves.

There are several ways in which someone knowledgeable in Google Analytics can configure your account to give you at least some insight into which of your pages are receiving Google organic traffic where the keyword is “not provided” (which will be the great majority of them). From this you can extrapolate which groups of related keywords those landing pages are optimized for and which keywords they actually rank for with the assumption that the “not provided traffic” is coming from those keywords. For some examples of these solutions see Dan Parker’s “How to steal some ‘not provided’ data back from Google” (<https://econsultancy.com/blog/8342-how-to-steal-some-not-provided-data-back-from-google>) or Moz’s “Easing the Pain of Keyword Not Provided: 5 Tactics for Reclaiming Your Data” (<http://moz.com/blog/easing-the-pain-of-google-keyword-not-provided>).

Segmenting Search Traffic

Next, you can consider putting these things together. Even the free tools provide substantial capability for building out custom reports. Google Analytics has options that show the search phrases just for Google, and the pages of the site that Google sent users to when users clicked on your link.

This is a key improvement to your SEO research for two reasons:

- If you are looking into what terms can bring fast traffic benefits as a result of some additional optimization, you are going to want to know in which search engine you are ranking.
- If you are going to optimize a page to rank higher, you will need to make sure you are optimizing the right page!

Referring Sites

It is interesting to look at a referring site report for a number of reasons, but one of the more interesting SEO reasons to do so is to spot when you receive new links. You can often see those new links in these reports first, even before the search engines report them. Figure 11-5 shows a sample portion of the referring sites report from Google Analytics.

Source <small>?</small>	Acquisition	
	Sessions <small>?</small> 	
	1,069 % of Total: 59.89% (1,785)	
101. biteclubeats.com	1 (0.09%)	
102. blog-search.com	1 (0.09%)	
103. blog.landerapp.com	1 (0.09%)	
104. blog.myrss.jp	1 (0.09%)	
105. bloggerbabes.com	1 (0.09%)	
106. blogher.com	1 (0.09%)	
107. bowhuntamerica.com	1 (0.09%)	
108. boxingnewsonline.net	1 (0.09%)	
109. bravoacademy.com	1 (0.09%)	
110. businessblogconsulting.com	1 (0.09%)	

Figure 11-5. Referring sites report

Expanding on this example, consider the site shown in Figure 11-5. If this is the first time you have ever noticed this site in your referrers, it can be a leading indicator that you have received a new link. This is of interest as it can help you to detect new links that result from your link-building campaigns, and therefore help you measure which of your link development campaigns are yielding the best results.

Using Custom Analytics Dashboards

In analytics terms, a *dashboard* is a single-page view that contains your most critical metrics all in one place. Of course, your most critical metrics are different from those of the next publisher, because the needs of different sites vary greatly.

In addition, multiple dashboards may be required in any given organization. For example, the CEO of a large public company probably wants to see different data (and a lot less of it) than a senior business analyst.

Each analytics package provides methods for implementing a custom dashboard. Figure 11-6 is an example of one from IBM Digital Analytics.

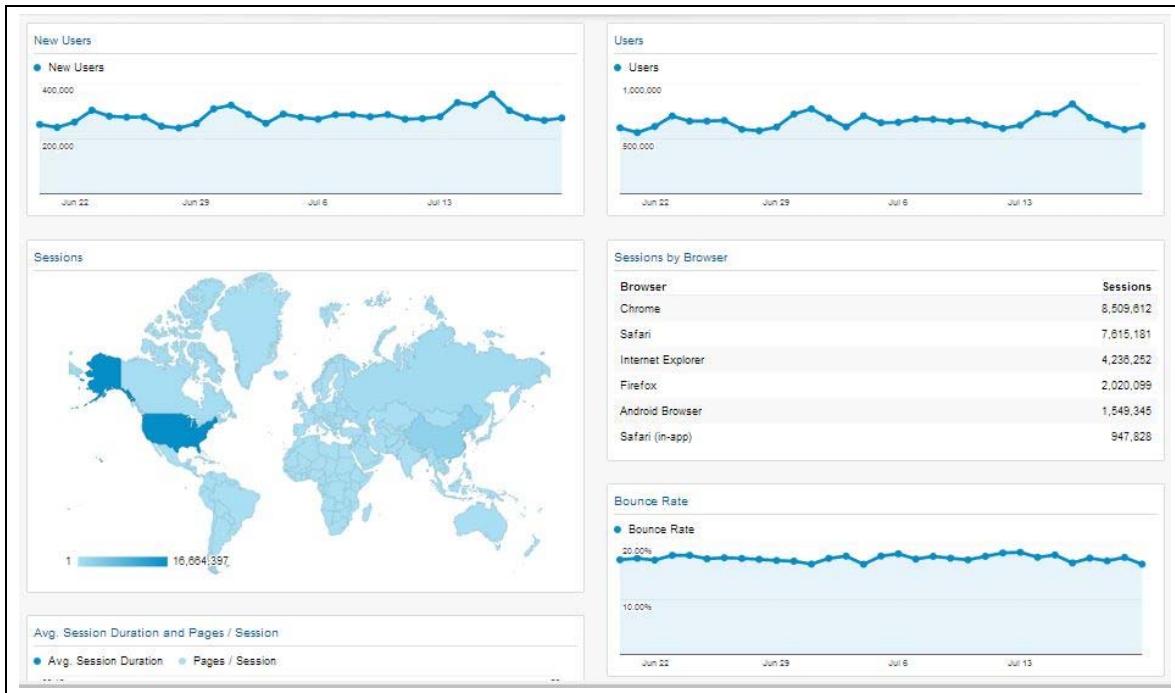


Figure 11-6. Custom dashboard

As you can see from Figure 11-6, a dashboard can be quite visual. What is most important, though, is that it provides the data that is most important to the person for whom the dashboard was designed. As an SEO practitioner, you can implement a dashboard to show progress against the goals you set for your SEO campaign.

Providing this type of visibility has two important benefits:

- The person viewing the report will appreciate the fact that she does not have to work hard to do a quick health check on the progress of the SEO efforts. As we suggested at the beginning of this chapter, the fact that you agree to measurable goals will be a great comfort to management.
- You will know what data your managers are looking at. When something goes wrong (or right), or when management wants to discuss some aspect of the business, they will have started from the dashboard you set up for them.
- A key consideration in selecting analytics and dashboard solutions are whether your web analytics solution has pre-built integrations with your preferred SEO solution, or whether your SEO solution is part of a larger cross-media dashboard platform. Otherwise, you may end up with siloed information or quite different conversion metrics in one dashboard than another.

A Deeper Look at Action Tracking

Action tracking goes one step deeper than basic analytics. Rather than simply observing what pages are visited and how many unique sessions are logged, action tracking allows you to narrow down groups of visitors based on the actions they take on your site.

In most instances, it requires setting up a code in your analytics program and attaching that code to a button, page load, image rollover, or other JavaScript-trackable task (a click or hover).

Figure 11-7 provides a look at how this works.

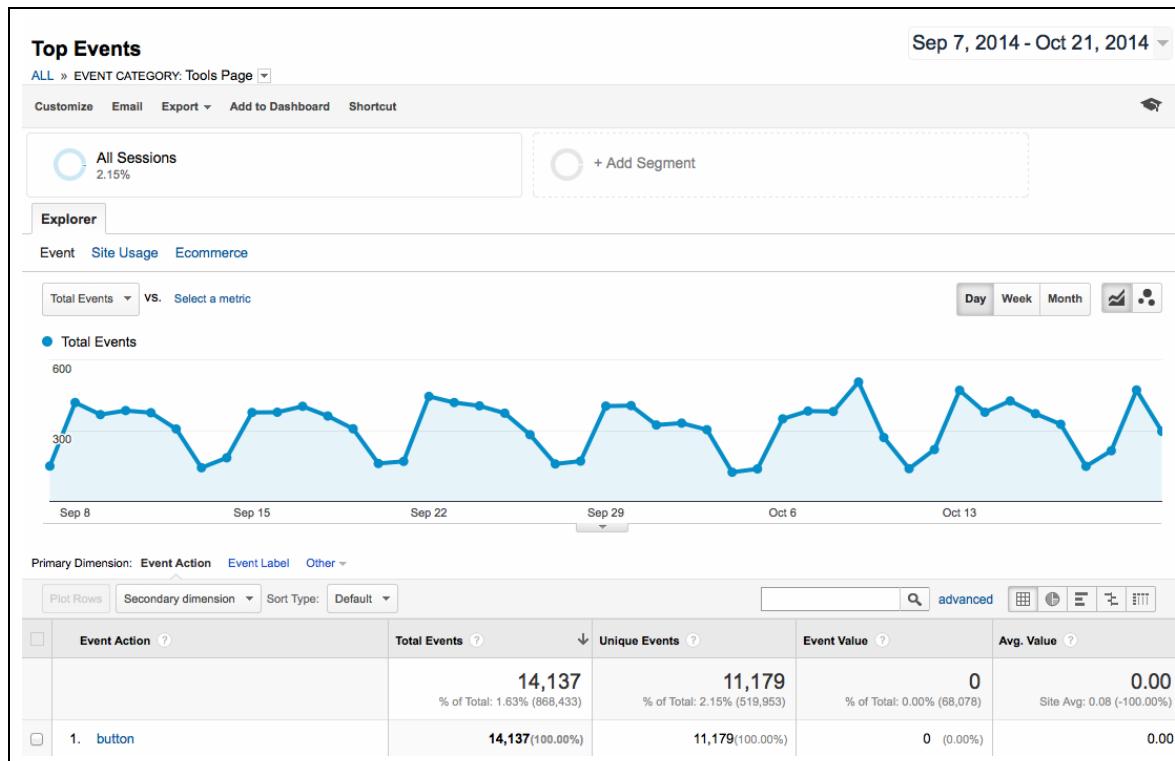


Figure 11-7. Action tracking in analytics

You can see from Figure 11-7 that:

- Moz has action tracking in place where the users click a particular button. Note that it is often the best choice to have the action tracking on the action confirmation page rather than the sign-up form itself, as tracking on the sign-up form can result in including incomplete form fills at part of your action total.
- Based on the people who take this action, you can predict which search terms will be better at converting visitors into applicants.
- The Revenue column is empty, but if you were tracking e-commerce buyers, you could put their totals into the Revenue column and track high-volume buyers. It is also a good idea to assign a value to different actions such as contact us requests, subscribers to free services, or any other actions a user could later lead to a conversion.
- Expanding on this idea, you could also track users by time of day, the search engine they used, their geographic location, and so on.

So, what types of actions should you track on your site? The answer varies depending on your business and site structure. Here are some suggestions as segmented by site type.

E-commerce site:

Add to Cart button

Studies have shown us that users who “add to cart,” even if they do not complete the checkout process, are more likely to return to make a purchase. This is also a good way to calculate shopping cart abandonment and make changes to refine and improve the process.

Complete checkout

An obvious one; this action will show you what percentage of each user group is converting into sales. It is also of interest to measure what percentage of people start the checkout process but do not complete it.

Save to wish list

E-commerce sites offering wish lists are still in the minority, but wish lists are a great way to track interest that isn't quite a purchase.

Send this to a friend

Many sites offer a “share this page” function, and it is a great action to be aware of. If folks are sending out your link, you know you have a hit.

B2B site:

Subscribe to newsletter

A subscription is a tacit endorsement of your brand and a desire to stay in contact. It may not be a conversion, but for B2B, it may be the next best thing.

Contact form submission

Filling out a contact form can be even more valuable than newsletter subscription, in some cases. Though some of these forms will report support issues, many may contain questions about your products/services and will indicate a desire to open a sales conversation.

Email link

As with contact forms, direct email links have the possibility of becoming a sales contact. The best thing you can do is clearly label sales emails and track them separately from support or business issues.

Blog:

Subscribe to RSS feed

An RSS feed subscriber is a blog's equivalent of a conversion; tracking these is imperative. While RSS Subscribers are in decline, there are still people who do this and it is still useful to track.

Add comment

Anyone who is contributing content to the blog or participating should be paid attention to (as should those channels that earn user engagement).

Social bookmark/share

All those folks who are submitting your content to Facebook, Twitter, Delicious, StumbleUpon, Reddit, and other social media and news aggregation/discussion sites deserve to be recognized (and sought after).

Forum or other site based on user-generated content (UGC):

Sign up for an account

These users are active contributors; you need to know where they come from.

Contribute content

When a user publishes, discovering her path is important (especially if it is not from a bookmark/type-in).

Add comment

As in the preceding list item, comments are a great predictor of engagement.

Vote/rate

Even low levels of participation, such as a rating or a vote, are worth tracking when every piece of participation counts.

Social shares

Facebook friends and likes, Google +1's, Twitter followers. You can read more about this in "Tracking Social Media in Your Web Analytics" in Chapter 8.

You can get very creative with many of these and track all sorts of actions. If you offer a tool or calculator, track its usage. If you want to follow who clicks a particular link (rather than just a page), add an action to it. You can even see which users hover over an image.

For many marketers, action tracking is the secret sauce. Once you install and activate this data, it is easy to adjust SEO (or PPC) marketing efforts on particular terms, phrases, or pages that will make a huge difference to the campaign's efficacy.

For SEO practitioners, action tracking is a "cannot-live-without-it" tool. The recommendations made based on this kind of tracking are backed up by real data. High-volume action-tracking studies almost always make a huge impact on a site's bottom line. It also allows the performance of organic search to be compared against other channels such as paid search, email marketing, and social media.

Separating the Analytics Wheat from the Chaff

One of the big problems with web analytics is that there is so much data. There is seemingly an endless supply of the stuff—so much that you can get lost in the weeds and waste a lot of time. As Dennis Mortensen of Visual Revenue (<http://visualrevenue.com/blog>) likes to say, do not "report-surf."

By this he means you shouldn't get lost in the data, looking at one interesting fact after another, without understanding what business objective you are attempting to accomplish. To avoid this trap, you need to have a plan.

The basis of this plan is the definition of actionable Key Performance Indicators (KPIs). To summarize this concept very briefly, do not view a report unless there is a chance that it will provide enough insight to take action (such as reallocating resources to fix a problem or pursue a new opportunity you discovered).

Common analytics mistakes

Because of the complexity of web analytics, there are plenty of opportunities to make mistakes. Even if you are laser-focused on actionable KPIs, you can still draw wrong conclusions from the data. Here are a couple of examples:

Making a decision on too small a sample size

A simplistic example would be if you had a keyword that brought you two visitors, and both of them bought something. What does that mean? Your conversion rate will always be 100%? Of course not.

You need to have a test run for a length of time that is statistically significant. According to Jonathan Mendez (<http://www.optimizeandprophesize.com/>), long-time evangelist for Offermatica (which is now Adobe Target), the “golden rule” for test length is 100 conversions per branch (or tested element). For an A/B test, ideally you would have 200 conversions before you looked at the results. However, if after 50 conversions you have 40 on one branch and 10 on the other, your test is over. If your variance in results is lower, you need more data to make sure the data you are seeing is accurate.

Not recognizing when you have enough data

You can make this mistake in the opposite direction too. When you have enough data, you need to be willing to act on it.

Working with biased data

You can easily collect biased data and make the wrong decision. For example, an e-commerce site basing its SEO plans on December’s traffic numbers is probably being misled. And temporal factors aren’t the only cause—visitor locations, referral sources, and in-house campaigns to existing lists can all skew information.

Tying SEO to Conversion and ROI

As we discussed previously in this chapter, it is important to tie your SEO efforts to the results they bring to the business. A fundamental piece of that is measuring the conversions driven by organic SEO traffic. Here are some of the most common types of conversions:

Sales/sales revenue

This is the one that everyone assumes is part of conversions. Sales and sales revenue (or better still, margin) conversions can be the simplest things to track, except when you are selling many different products at different price points and in different quantities. In this case, the process would need to be a bit more sophisticated.

If your site is advertising-driven, you need to look at the impact of organic search traffic on advertising revenue. If you have no financial goals for your site, you need to look at some of the other types of conversions and determine their value or worth.

Email/blog/newsletter subscriptions

Anytime a user signs up to receive regular communications from you it is a win. Even though there are not direct financial consequences to this it is still a conversion. Someone who has subscribed to something you offer is more likely to become a customer than a first-time visitor to your site, so you need to credit this type of conversion.

Sign-ups

Closely related to the notion of subscriptions are other types of sign-ups. Perhaps you offer a service such as a tool that people need to sign up for to use. Even if the tool is free, you should track this as a conversion.

You most likely received the person's email address in the process, and even if she indicates that she does not want to receive commercial communications from you, you should be building loyalty with the tool you provided her access to (or whatever she signed up for); otherwise, why would you be providing her that service?

Downloads

Many sites offer free downloads, such as white papers, or free downloadable tools. Even if you do not require a sign-up of any type, you should still count this as a conversion. You are getting your message out there with the downloads you offer.

Contact forms and phone calls

This is when someone requests that you contact him or her, or he or she contacts you. This can take the form of filling out a "contact me" request form or a phone call. Phone calls can be tracked through unique numbers per type of action, landing page, and traffic source.

Visitors who share

This conversion happens when a visitor shares the information she found on your site with someone else. This includes Facebook Shares, Google+ Shares, and Links on Twitter. In addition, if your site has a "share with a friend" or an "email to friend" feature, you could keep track of conversions by noting each time a user used that feature.

Visitors who link

A user who links will visit your site and find its content useful, entertaining, or otherwise compelling enough to link to it from her own site.

Visitors who publicize

Visitors can publicize your site by mentioning it in forums, blog comments, or social media sites such as Twitter and Facebook, or by writing about it on their own site.

You should place a dollar value on every type of conversion you receive. Even a white paper download has some value to it. Perhaps it may increase the potential of a future sale to that user. For example, if 100 people download your white paper and two of them become customers and buy a \$100,000 consulting contract, you can estimate the value per download at \$2,000.

SEO software packages such as Conductor Searchlight (<http://www.conductor.com/searchlight>) and Raven SEO Tools (<http://raventools.com/seo-tools/>) allow SEO professionals to view search ranking data together with traffic and revenue data. This enables them to tie natural search traffic to conversions, and therefore measure ROI. You can see sample screens from Conductor Searchlight in Figure 11-8 and Raven SEO Tools in Figure 11-9.

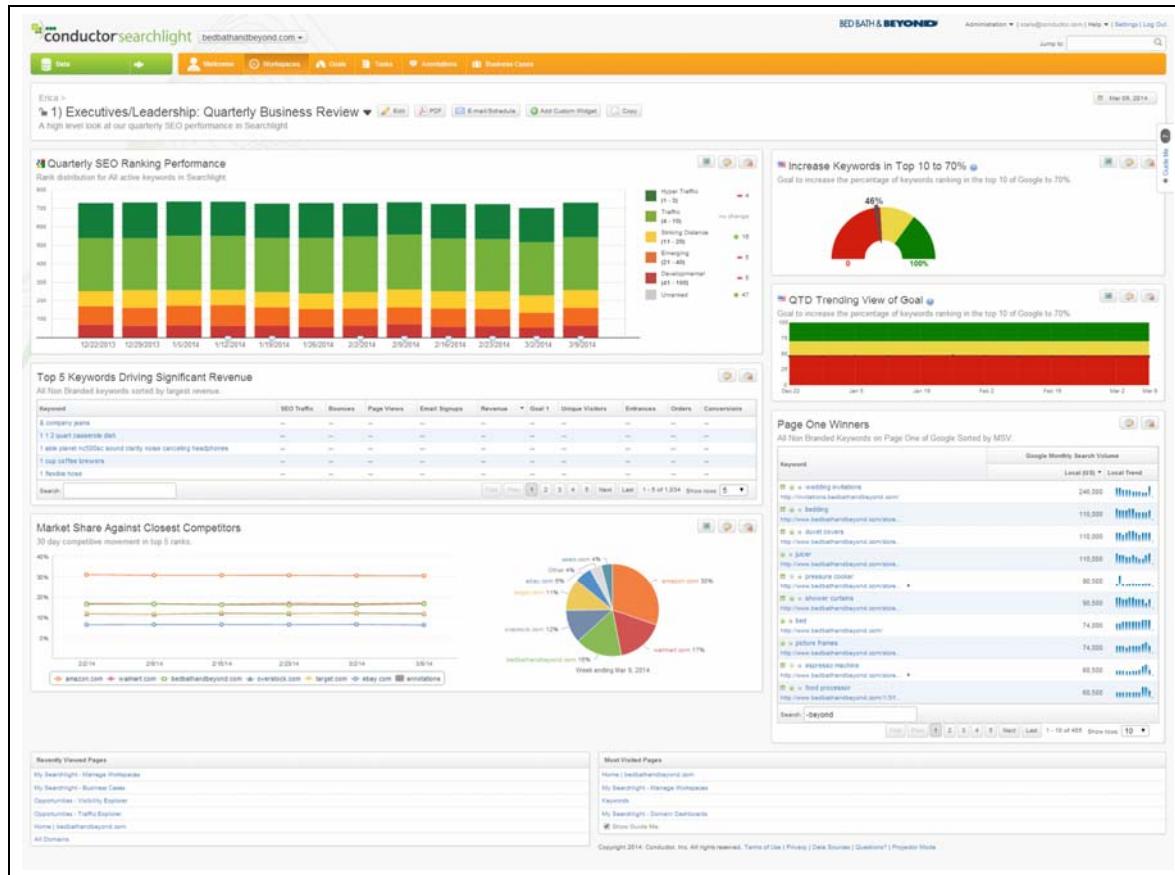


Figure 11-8. Conductor Searchlight ranking data overlaid with traffic and revenue metrics.

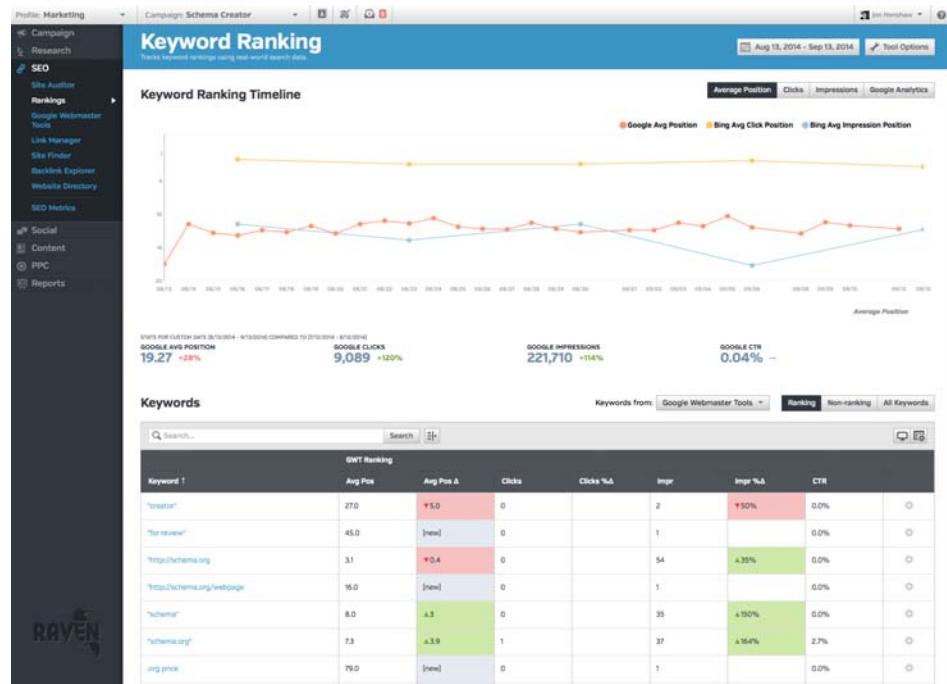


Figure 11-9 Raven SEO Tools Ranking data and Conversions

Attribution

Another issue to be aware of is attribution - identifying user actions that led to specific outcome. The issue is that there is a tremendous amount of interaction between types of media. Here are some examples:

- A user does a search, clicks on an organic search result, and then reads a few things and leaves. The next day, she remembers what she read, does another search, and this time clicks on a paid search ad for the site and buys a product (organic search should receive some credit).
- A user does a search, clicks on an organic search result, and leaves. A few days later, because of what she learned, she goes into a store and buys a product (organic search should receive some credit).
- It goes in the other direction too: a user sees a TV ad, then does an organic search to find the website, and buys the product (the TV ad should get some credit for the sale!).

Such cross-channel interactions are becoming increasingly common. Data published on Search Engine Watch (<http://searchenginewatch.com/article/2067752/The-Impact-of-Comparison-Shopping-Sites-on-E-Commerce-Sales>) supports this. A ScanAlert study of a large number of e-commerce sites showed that 30% of transactions took place one to three days after the customer's initial visit. So, the problem of attribution management is a significant issue!

It is not always easy to track multi-visit conversions in detail. Certainly, a direct interaction between TV/offline stores and a website is very difficult to track in detail. But even looking at multiple site visits with an interaction between PPC and organic search can be difficult, partly because the tracking systems rely on cookies, which users have a tendency to delete.

So, attribution is something you should think about, and ideally you should attempt to track the impact across multiple site visits. However, the methods for attributing sales across multiple site visits are imperfect. One thing you can do to minimize the impact of that is to count different types of conversions, such as filling out contact forms, downloading white papers, visiting key product pages, and so on. Tracking a larger number of events that indicate progress toward a sale (or whatever your end goal may be) can help you identify at which stage in the research and conversion process that organic search comes into play for the users.

Setting Up Analytics Software to Track Conversions

One of the great things about tracking conversions is that you can use the resulting data to focus your efforts. If a given page, or referrer is showing much higher conversion rates than another source, you may want to focus more of your SEO efforts on those rather than on something else that does not convert as well.

Capturing conversion data takes some setup. Your analytics software does not know what you consider a conversion unless you define it. For example, in Google Analytics, you need to go into your Profile Settings to configure conversion goals and funnels (see Figure 11-10)

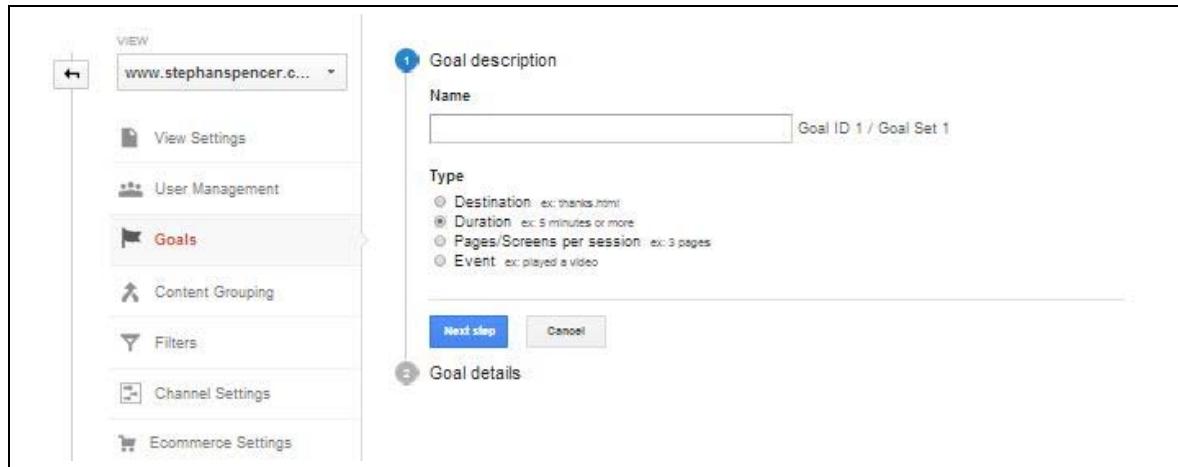


Figure 11-10. Google Analytics goals

Once you click **Edit** (on the right) for one of the goals, you come to a screen that allows you to define a goal (see Figure 11-11).

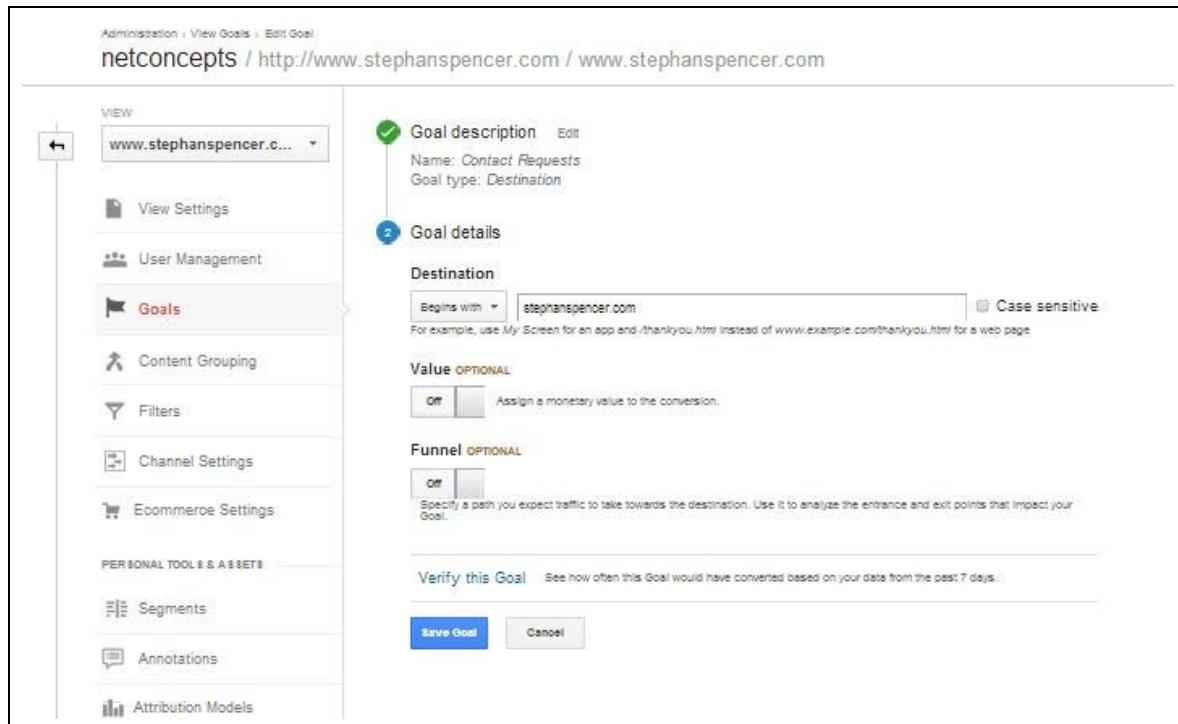


Figure 11-11. Google Analytics goals setup

In Figure 11-11, the primary data being tracked is when someone arrives at a given web page. These types of conversions are easy to set up.

Event tracking, such as someone clicking on a link, downloading a white paper, clicking on Flash or AJAX elements, or playing a video, is slightly more complicated. Since you cannot count on a page load to tell

you when a link has been clicked on, you need to add a bit of tracking code to the link itself. Here is what the tracking code looks like for Google Analytics:

```
| <a href="#" onClick="_gaq.push(['_trackEvent', 'Video', 'Play', 'Art of SEO Book Launch']);">Play Video</a>
```

The key code is `onClick="_gaq.push(['_trackEvent', 'Video', 'Play', 'Art of SEO Book Launch'])`. This tells Google Analytics to associate this event with the user's action. Configuration of events is best handled by an experienced developer. More information is available through the Google Analytics Event Tracker Developer Guide (<https://developers.google.com/analytics/devguides/collection/gajs/eventTrackerGuide>).

Google has also introduced a new analytics package called Universal Analytics. With Universal Analytics you can track users across multiple devices and sessions and it offers you the flexibility to implement tracking (including event tracking) through both JavaScript on your website, the Google Analytics SDK for mobile, and they even have an additional protocol to measure engagement across other digital devices such as game consoles. There are several steps to implementing Universal Analytics which are not intuitive so we recommend engaging and experienced analytics professional before converting to Universal Analytics.

A particularly useful capability of Google Analytics is the ability to remarket to previous site visitors in Google AdWords. You can segment your users based on a specific behavior, such as visiting a particular product page, make a list of these users in GA. This list is automatically sent over to Google AdWords if the list is longer than 100 users. Then, your ads will be marketed to solely those users across the web. Find out how to implement retargeting at: <http://cutroni.com/blog/2012/11/20/getting-started-with-google-analytics-remarketing/>.

Form submissions are another thing you may want to track. This could include when someone has completed a request to sign up for a newsletter, requested or started a white paper download, or simply used a form to request that you contact him or her. These are usually quite easy to set up, provided that the result of the form submission is some sort of "thank you" page. All you need to do in this case is tell your analytics software that the loading of that "thank you" page is considered a conversion, and then give it a value.

Analytics is particularly useful if a portion of your customers tend to log in to a company's features with a social media login, such as a "Login via Facebook" mechanism. In that case, you have the ability to track users across several devices, for even more advanced and thorough customer data.

Facebook Insights gives yet another valuable data set for tracking conversion. Domain insight shows all of the traffic from Facebook to your domain, including traffic from your Facebook business page, from link on status updates, from social plugins, and more.

Life gets more complicated when you want to track actual product sales on sites where products vary in price. Generally speaking, this will require putting some custom JavaScript on your order confirmation page.

In addition, you will need to have your web application populate certain data in the JavaScript before executing it, such as the total amount spent. Oftentimes, publishers choose to include a lot more data, such as listing all the products sold, the price per product, how many units, shipping costs, and so on.

Although the preceding examples used Google Analytics, any of the major analytics packages can do this type of conversion tracking.

Conversion tracking strategy

An important step in conversion tracking is deciding what you want to call a conversion. As we outlined earlier in this chapter, it is important to consider all types of conversions, such as sign-ups, contact requests, downloads, and so forth. There is one other decision you need to make: how to value each conversion.

What is a contact request worth? What about a newsletter sign-up? One way to back into this is to look at how many of your newsletter sign-ups become customers. For example, if 100 people sign up for the newsletter, 10 of them become customers, and the total dollar value of their orders is \$200, the average value per sign-up is \$2.

When you first set up conversion tracking you may need to estimate some of these things. But as your data improves, you can improve your calculation and refine the value you place on each conversion.

Segmenting Campaigns and SEO Efforts by Conversion Rate

Now that you have conversion tracking set up, how do you use it to focus your SEO efforts? One of the things you will want to do is to track conversion data in different ways:

Conversions by keyword

What keywords are bringing the best results? Note that this data has become far more obscured since the "Not Provided" change of 2013. Even linking Webmaster Tools to Google Analytics does not enable conversion tracking for those terms.

Conversions by referrer

Which traffic source is offering the best conversion?

Conversions by web page

Which pages on your site result in the highest number of conversions?

Conversions by initial entry page

Which initial entry pages ultimately lead to the highest number of conversions?

For example, Moz actively tracks its conversions using Google Analytics. Tracking by keyword can be a useful way to do this, as shown in 11-12.

Keyword	Acquisition			Conversions		
	Sessions	% New Sessions	New Users	Moz Subscription Signup (Goal 14 Conversion Rate)	Moz Subscription Signup (Goal 14 Completions)	Moz Subscription Signup (Goal 14 Value)
	28,475 % of Total: 0.70% (4,072,006)	38.48% Site Avg: 51.13% (-24.74%)	10,957 % of Total: 0.53% (2,082,038)	1.81% Site Avg: 0.17% (967.09%)	516 % of Total: 7.46% (6,915)	\$0.00 % of Total: 0.00% (\$0.00)
1. moz	12,694 (44.58%)	34.37%	4,363 (39.82%)	1.61%	204 (39.53%)	\$0.00 (0.00%)
2. seomoz	4,609 (16.19%)	36.99%	1,705 (15.56%)	1.76%	81 (15.70%)	\$0.00 (0.00%)
3. seo moz	1,858 (6.53%)	35.04%	651 (5.94%)	2.58%	48 (9.30%)	\$0.00 (0.00%)
4. moz com	1,289 (4.53%)	30.88%	398 (3.63%)	1.24%	16 (3.10%)	\$0.00 (0.00%)
5. seo tools	763 (2.68%)	58.58%	447 (4.08%)	2.23%	17 (3.29%)	\$0.00 (0.00%)
6. seo mox	731 (2.57%)	39.12%	286 (2.61%)	4.24%	31 (6.01%)	\$0.00 (0.00%)
7. moz pro	372 (1.31%)	20.97%	78 (0.71%)	2.15%	8 (1.55%)	\$0.00 (0.00%)
8. moz analytics	317 (1.11%)	28.71%	91 (0.83%)	0.95%	3 (0.58%)	\$0.00 (0.00%)
9. seo website analyzer	299 (1.05%)	52.84%	158 (1.44%)	2.01%	6 (1.16%)	\$0.00 (0.00%)
10. website seo analysis	284 (1.00%)	57.04%	162 (1.48%)	0.35%	1 (0.19%)	\$0.00 (0.00%)

Figure 11-12. Google Analytics tracking conversions

You can see that the conversion rate on the "seo tools" keyword is high. Perhaps an SEO campaign to improve the ranking of Moz on these keywords is in order. Not only do you see high conversion, but you also see reasonable material volume.

Implementing an SEO campaign that doubles related traffic for this set of keywords will likely have a significant impact on the bottom line.

Increasing Conversion

As an SEO practitioner, you should strive to become an expert at conversion optimization, because higher conversion rates mean higher impact for your SEO campaigns. (Co-author Stephan Spencer offers a viewpoint in *MultiChannel Merchant* at <http://multichannelmerchant.com/opsandfulfillment/be-a-converter-01072008/>.

Nonetheless, the average site conversion rate for a typical merchant today is only about 2.4% combined for new and returning visitors (Source: Monetate E-Commerce Quarterly Report EQ1 2013 http://assets.monetate.com/eq/EQ1_2013_final.pdf).

So, how do you help your conversion rate by managing your content and refining your site design? Most web resources that give content advice are focused on blogs; specifically, what and how often you should publish. Those of us who use websites strictly for sales and marketing are equally concerned with good content, but from a totally different angle—one that traditional content advice doesn't properly address.

What's more, much of the design advice out there seems contradictory. The key to enhancing your conversions by adjusting your design and content is not in thinking strictly about the technology or the mechanics of site layout, but in understanding the needs of your visitors. In that, there are many different elements to consider. A book about SEO can't cover this topic in depth, but you may want to pick up one of the following excellent books on the topic: *You Should Test That!* by Chris Goward (Sybex), *A/B Testing* by Dan Siroker and Pete Koomen (Wiley), *Always Be Testing* by Bryan Eisenberg and John Quarto-von Tivadar (Sybex), *Landing Page Optimization* by Tim Ash (Sybex), or *Web Design for ROI* by Lance Loveday and Sandra Niehaus (New Riders Press).

A content marketing/conversion tracking example

A specific example of how to use action tracking by referral source is to track a *content marketing bump*. The idea is that after successful content draws thousands of visitors to your site, a small percentage of those visitors will stick around and continue to visit on a semi-regular basis (this is particularly true if you're regularly delivering fresh, valuable content).

Action tracking by referral source

Figure 11-13 shows a sample report from Google Analytics.

	Source ?	Acquisition	Conversions	Goal 14: Moz Subscription Signup ?	
		Sessions ? ↓	Moz Subscription Signup (Goal 14 Conversion Rate) ?	Moz Subscription Signup (Goal 14 Completions) ?	Moz Subscription Signup (Goal 14 Value) ?
		974,998 % of Total: 23.94% (4,072,006)	0.09% Site Avg: 0.17% (-47.39%)	871 % of Total: 12.60% (6,915)	\$0.00 % of Total: 0.00% (\$0.00)
1.	analytics.moz.com	70,681 (7.25%)	0.06%	45 (5.17%)	\$0.00 (0.00%)
2.	moz.com	52,470 (5.38%)	0.00%	0 (0.00%)	\$0.00 (0.00%)
3.	t.co	30,139 (3.09%)	0.10%	29 (3.33%)	\$0.00 (0.00%)
4.	webstatschecker.net	21,658 (2.22%)	0.00%	0 (0.00%)	\$0.00 (0.00%)
5.	example.com	16,738 (1.72%)	0.00%	0 (0.00%)	\$0.00 (0.00%)
6.	feedly.com	14,028 (1.44%)	0.04%	6 (0.69%)	\$0.00 (0.00%)
7.	opensiteexplorer.org	13,783 (1.41%)	0.04%	6 (0.69%)	\$0.00 (0.00%)
8.	facebook.com	10,115 (1.04%)	0.28%	28 (3.21%)	\$0.00 (0.00%)
9.	plus.url.google.com	10,067 (1.03%)	0.04%	4 (0.46%)	\$0.00 (0.00%)
10.	seocheki.net	9,985 (1.02%)	0.01%	1 (0.11%)	\$0.00 (0.00%)
11.	searchenginejournal.com	9,347 (0.96%)	0.26%	24 (2.76%)	\$0.00 (0.00%)
12.	l.facebook.com	8,106 (0.83%)	0.32%	26 (2.99%)	\$0.00 (0.00%)
13.	inbound.org	7,679 (0.79%)	0.09%	7 (0.80%)	\$0.00 (0.00%)
14.	reddit.com	7,015 (0.72%)	0.04%	3 (0.34%)	\$0.00 (0.00%)
15.	backlinko.com	6,668 (0.68%)	0.19%	13 (1.49%)	\$0.00 (0.00%)
16.	stackoverflow.com	5,578 (0.57%)	0.02%	1 (0.11%)	\$0.00 (0.00%)
17.	searchengineland.com	5,492 (0.56%)	0.11%	6 (0.69%)	\$0.00 (0.00%)
18.	quicksprout.com	4,793 (0.49%)	0.15%	7 (0.80%)	\$0.00 (0.00%)

Figure 11-13. Action tracking by referral source

In Figure 11-13, you can see a list of domains that sent Moz traffic over a few weeks' time, along with the number of resulting Moz subscription sign-ups. You can see some clear patterns, including the following:

- WebStatChecker, Feedly, Google Plus, Stack Overflow, and Reddit are sending traffic, but less than one-tenth of 1% of those visitors sign up for membership on the site.
- The domains that point primarily to the Moz tools (such as Backlinko) are responsible for a high number of sign-ups, but this is to be expected since one has to register to see the full data sets available in Moz's tools such as Open Site Explorer.
- SearchEngineJournal, Facebook, and QuickSprout are amazingly on-topic sources of traffic, which directly bears out in the higher percentages of referrals who become members.

Action tracking is one of the most valuable things you can do on any type of site. Most people limit its application to e-commerce, but you can see from the data in Figure 11-13 how valuable it can be, even for a simple blog. Even if you do not have a membership sign-up, you could attach the action tracking to your “add comment” link in a blog or to a newsletter sign-up on a content site. Segmenting that data by incoming referrals (search engines, direct type-ins, and links) gives you an even better picture of who's sending you truly valuable traffic. You can also get an indication of traffic quality by looking at bounce rate, time on site, and number of pages viewed for each referrer.

In Figure 11-14, you can see that around late May, Moz published some successful content that helped bump up the site visitor count at least through mid June. This traffic “bump” after the spike is another great benefit of creating viral-worthy content, and is a terrific way to grow the popularity of a site.

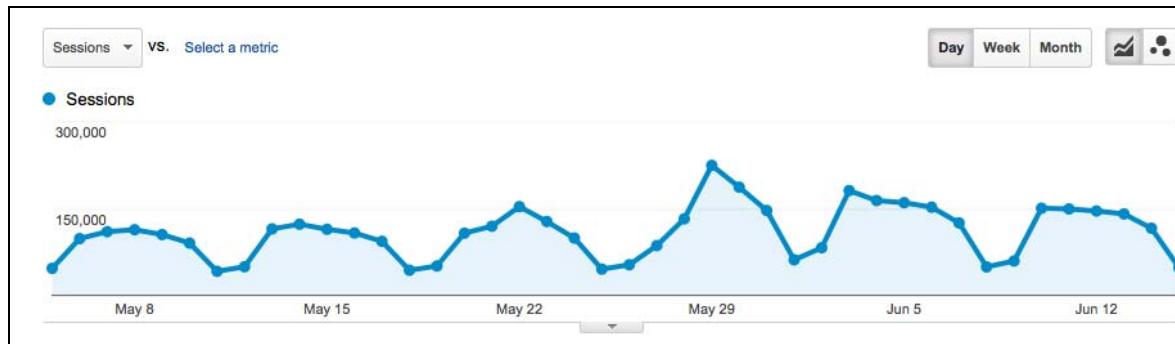


Figure 11-14. Traffic report

The tricky part is to continue to deliver great content after a viral piece and maintain the creative energy necessary to prove to visitors that you’re worth coming back to again and again. It is not an easy sell—people’s time is valuable. You have to make your content so targeted, interesting, and worthy of their attention that they cannot look away. It is challenging and time-consuming, but it is a great way to build a successful site.

Determining Project ROI

An effective SEO process is one that continuously works toward a positive return on investment. A basic SEO ROI business case would include these elements:

Number of people searching for your keywords

This can be hard to estimate, because you cannot completely map out the long tail. One rough estimation strategy is to multiply the search volume for the top terms for your business by 3.3 (i.e., assume that the head terms are about 30% of the available volume).

Expected click-through rate from organic search

This relates to the position you think you can achieve. The #1 position can net you approximately 19% of the available clicks.

Average conversion rate

Once you have captured the visitor, how successful are you at completing a conversion? Typical conversion rates for a website might be between 2% and 5%. It should be easy to get this data from your analytics. You should already know what your conversion rate is!

Average transaction amount

Last but not least, factor in the average transaction value. Again, this is data you already have.

Ultimately, the formulas look like this:

SEO Revenue = People searching for your keywords * click-through rate * average conversion rate * average transaction amount

For example: 10,000 per day * 10% * 5% * \$100 = \$3,000 per day

SEO ROI = SEO Revenue / SEO Cost (use total \$ spent for salaries and consulting, or number of hours spent)

For example: \$3,000 per day / \$500 per day = an ROI of 6X.

An Alternate Theory and Calculation

Predicting an SEO project's ROI based on rankings and click-through rate can, in some cases, be problematic to address because you have very little control over the variables. You end up relying on numbers that you have a limited ability to influence. As an alternative approach, you can measure and track SEO ROI based on an increase in search visibility.

To do this, begin by determining two things:

1. How many pages are getting search clicks for the month
2. How many clicks per page your site is getting

Next, record these supporting metrics:

- Average ranking across the whole keyword spectrum
- Average click-through rate

Now, by making it easier for search bots to find more pages, consolidating duplicate content, improving the depth and breadth of content, changing your pages to better match your users needs, and improving page titles, metadata, and microdata, you should see an increase in the number of pages getting search clicks and/or the number of clicks per page. The combination should result in more traffic when compared year-over-year.

Modifying the formula above to account for this alternate approach...

SEO Revenue = Increase in (pages getting search clicks * search clicks per page) * average conversion rate * average transaction value

Note that CTR in the long tail is far more variable. Not as many searchers are going to go past the first couple of results, so this formula won't always accurately show return. ROI is always going to be a "best estimate", but this type of ROI calculation will be more accurate for short tail.

Competitive and Diagnostic Search Metrics

Thus far in this chapter we have focused on the basics—the dollars and cents of determining whether you have succeeded in your SEO campaign. As we noted at the beginning of the chapter, these should be your first priority in your analytics efforts. In this section, we will start looking at metrics that you can use to diagnose specific SEO issues. An example of this would be finding out whether a major section of your site is not indexed. Another example is seeing how your traffic growth compares to that of your competitors (this helps you decide whether you have set the right objectives for your efforts).

Search Engine and Competitive Metrics

Numerous tools allow you to monitor your site and those of your competitors, providing insight into your SEO progress. You can also use these tools to figure out what your competitors are doing from an SEO perspective. This type of intelligence can provide you with new ideas on how to adapt your strategy to get better results.

As with all such tools, it is always important to understand the context of the tools and to have an idea as to what you are looking for. Better knowledge of your competitors' strategy is certainly one valuable goal. Detecting a problem in how your website is crawled is another.

By selecting specific and actionable goals, you can set yourself up for the highest possible return.

Site Indexing Data

It is valuable to know how many pages in your site are in a search engine's index. This is important for two reasons:

- To determine whether important parts of your site are not in the index. If key parts of the site are not in the index, you can embark on an effort to determine why.
- Looking at competitors' indexed pages can help you learn about your competitors' sites and strategies.

You can get basic information on the number of indexed pages for a site using the `site:yourdomain.com` command in Google, Yahoo!, or Bing. Figure 11-15 uses the results for Stephan Spencer's blog.

Google search results for `site:stephanspencer.com` showing 787 results. The results include links to Stephan Spencer's SEO Expert, Author, Professional Speaker, SEO Blog, Biography, Email, SEO Client Testimonials, and Who Is Stephan Spencer. A Google promotion for Webmaster Tools is visible on the right.

Try Google Webmaster Tools
www.google.com/webmasters/
Do you own `stephanspencer.com`? Get indexing and ranking data from Google.

Stephan Spencer: SEO Expert, Author, Professional Speaker
www.stephanspencer.com/ • Stephan Spencer's Scatterings •
About various aspects of online marketing and e-business.
You've visited this page 3 times. Last visit: 7/5/14

SEO Blog of Author, Strategist, Speaker Stephan Spencer
www.stephanspencer.com/blog/ • Stephan Spencer's Scatterings •
I recently shared my insights about Google rich snippets, Schema.org, and microdata at SMX Advanced on the panel session "From Microdata & Schema To ..."

Biography of SEO Expert and Author | Stephan Spencer
www.stephanspencer.com/about/ • Stephan Spencer's Scatterings •
Stephan Spencer is the author of Google Power Search, co-author of The Art of SEO, founder of NetConcepts, frequent conference speaker, and avid blogger.

Email | Stephan Spencer
www.stephanspencer.com/category/email/ • Stephan Spencer's Scatterings •
Jan 6, 2009 - This isn't a rhetorical question. I truly want to know! Do crass, "low brow" marketing campaigns like this real piece of... umm... work ;) ...

SEO Client Testimonials | Stephan Spencer
www.stephanspencer.com/praise/ • Stephan Spencer's Scatterings •
Stephan worked on SEO and eCommerce Strategy for House of Travel. Stephan has a real enthusiasm for the power of the internet to change businesses and ...

Who Is Stephan Spencer: Web Marketer - Stephan Spencer ...
www.stephanspencer.com/profile/ • Stephan Spencer's Scatterings •
Dec 3, 2009 - Hi, I'm Stephan Spencer. Here's a little treatise on me. First off, I'm the President & CEO of Netconcepts, a natural search marketing firm with ...

Articles - Stephan Spencer
www.stephanspencer.com/articles/ • Stephan Spencer's Scatterings •
Stephan is a senior contributor to Practical Ecommerce and to MarketingProfs.com, and a columnist for Search Engine Land, Multichannel Merchant, and The ...

Figure 11-15. Site indexing data

As shown in Figure 11-15, Google reports 787 pages in the index. Is that all of the site's pages? Many site owners will have a pretty good idea of the number of pages on their site and therefore can get a gut feeling as to whether they have a problem. One important thing to note is that the results that `site:` commands return are fast estimates, and aren't as accurate as the reports generated by Google Webmaster Tools based on your XML sitemap. You can use them to give you a gut-feel sense, but that is all.

You can improve this data by performing your site query and then switching to 100 results per page and going to the last page of the Google results and seeing if the number of results has changed. This will give you more accurate totals for the number of pages indexed.

If you think many pages are missing from the index, you can begin to explore a bit deeper. You can start by narrowing down sections of the site. Figure 11-16 shows the 2014 blog post archives on Stephan's site.

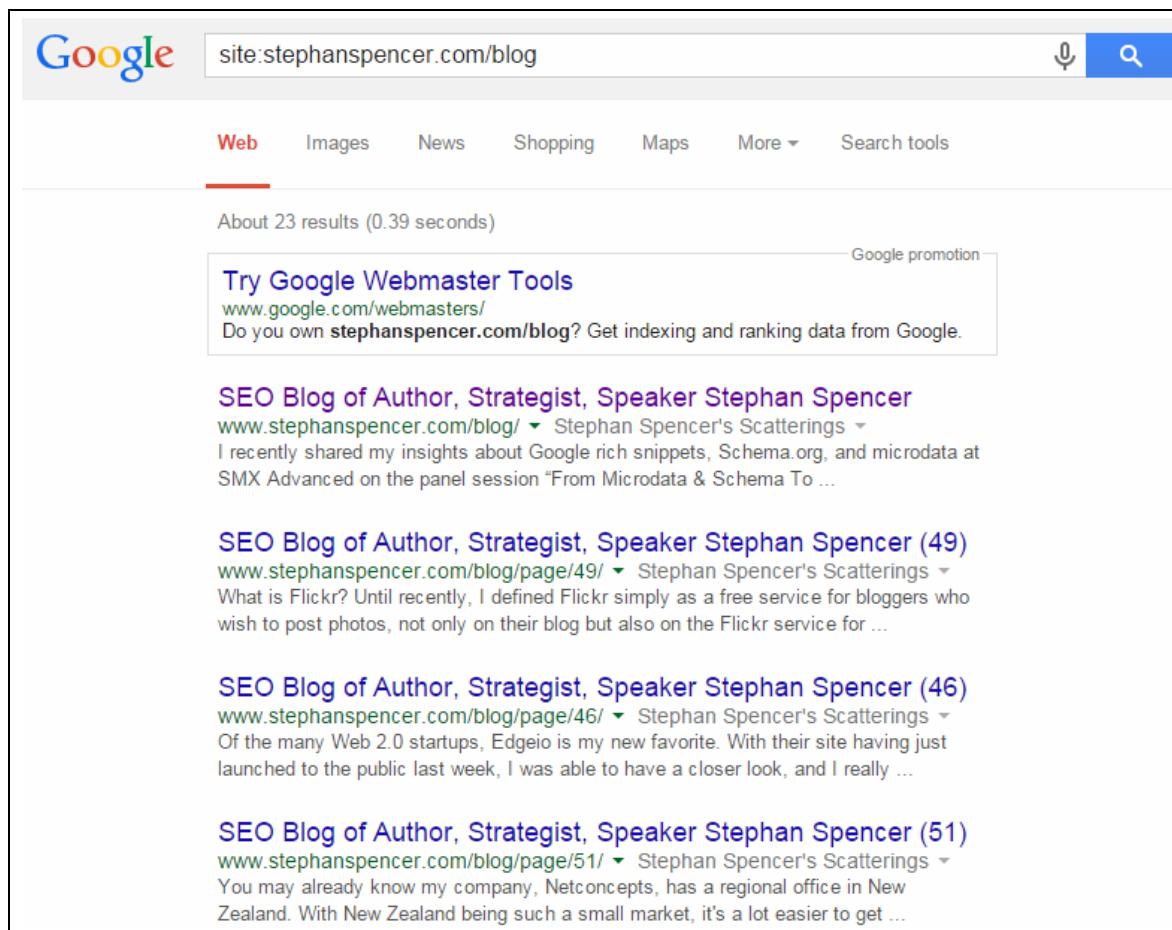


Figure 11-16. Indexing data on a website section

As shown in Figure 11-16, this reports 23 pages in the Google index. By using these different command structures, you can isolate what portions of the site appear to be less well indexed than others. You can also combine the `site:` command with other operators, as we outlined in Chapter 2. For example, consider combining it with the `inurl:` operator, as in Figure 11-17.

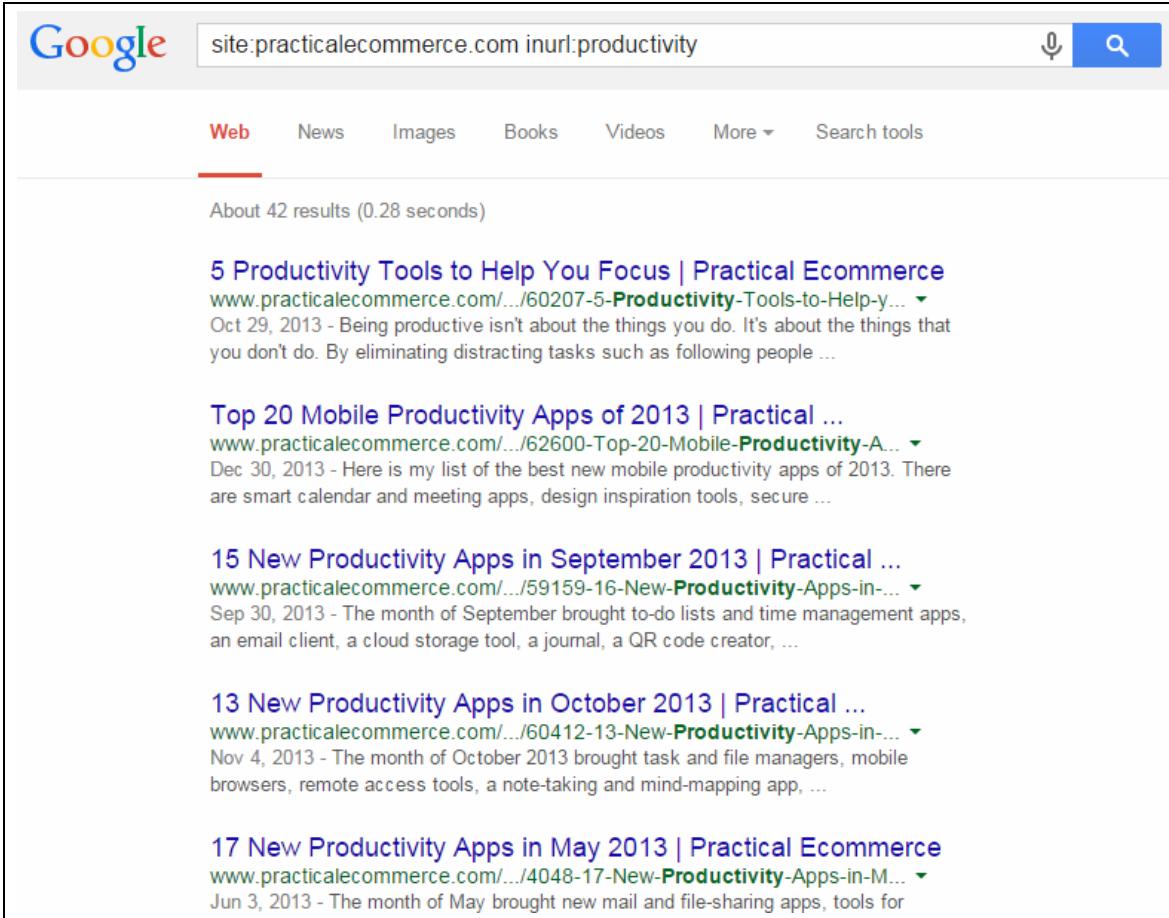
A screenshot of a Google search results page. The search query in the bar is "site:practicalecommerce.com inurl:productivity". The results are filtered by the "Web" tab. The first result is a link to "5 Productivity Tools to Help You Focus | Practical Ecommerce" with a snippet about productivity tools. The second result is "Top 20 Mobile Productivity Apps of 2013 | Practical ..." with a snippet about mobile productivity apps. The third result is "15 New Productivity Apps in September 2013 | Practical ..." with a snippet about new productivity apps for September. The fourth result is "13 New Productivity Apps in October 2013 | Practical ..." with a snippet about new productivity apps for October. The fifth result is "17 New Productivity Apps in May 2013 | Practical Ecommerce" with a snippet about new productivity apps for May. Each result includes a link to the full article on the Practical Ecommerce website and a timestamp indicating when the content was published.

Figure 11-17. Site command combined with inurl

Figure 11-17 shows that the <http://www.practicalecommerce.com> has 42 pages in the index that have productivity in the URL.

So far, we've shown website indexing tools that you can use on any website, which is great for competitive research. However, you can pull more accurate statistics for your website using Bing Webmaster Tools (<http://www.bing.com/toolbox/webmaster/>). You can also pull site indexing data from Google Webmaster Tools (<https://www.google.com/webmasters/tools/>).

Both of these tools are very useful, and you should use them on every site you touch. The tools provide a tremendous amount of valuable data at no cost, by simply signing up and verifying site ownership. For example, Figure 11-18 shows the opening screen for Bing Webmaster Tools (Bing WMT) for Stonetemple.com (<http://stonetemple.com>).

SITE ACTIVITY 1
Bing's information about your site

	CLICKS FROM SEARCH	APPEARED IN SEARCH	PAGES CRAWLED	CRAWL ERRORS	PAGES INDEXED
% Change	-3% <small>1</small>	0% <small>1</small>	9% <small>1</small>	106% <small>1</small>	2% <small>1</small>
This Period	97	25,633	22,087	342	2,750
Prior Period	100	25,856	20,089	166	2,671
Trends					

[See Reports & Data Graph](#)

SITEMAPS 1
Tell Bing about the URLs on your site.

LAST SUBMITTED	URLS SUBMITTED	LAST CRAWL	STATUS
You have not submitted any sitemaps			
SUBMIT A SITEMAP			

SEARCH KEYWORDS 222 1
Your top keywords from organic search

KEYWORDS	CLICKS FROM SEARCH	APPEARED IN SEARCH
hbx <small>1</small>	0	224
...

INBOUND LINKS 59899 1
Links pointed at your website

TARGET PAGE	COUNT OF LINKS
http://www.stonetemple.com/blog/	1,369

Figure 11-18. Bing Webmaster Tools

As per Figure 11-18, it looks like Bing has 2,750 pages from the site indexed. If you enter the `site:stonetemple.com` command in Bing itself, it shows something rather different (see Figure 11-19).

The screenshot shows a Bing search results page. The search query is 'site:stonetemple.com'. The results are filtered to 'Web' and show 10,500 results. The results are as follows:

- Stone Temple - Digital Marketing Excellence, Content ...**
<https://www.stonetemple.com> ▾
Stone Temple Consulting is a Leading Digital Marketing Agency That Has Been Servicing Enterprise Organizations Since 2002. Learn more
- 12 Cool Things you can do with Unica's Affinium NetInsight**
www.stonetemple.com/articles/NetInsight-12-Cool-Things.shtml ▾
12 Cool Things you can do with Unica's Affinium NetInsight. Next up in our series of cool things about analytics article series, is Unica's Affinium NetInsight.
- 11 Cool Things you can do with HBX Analytics**
www.stonetemple.com/articles/HBX-11-Cool-Things.shtml ▾
11 Cool Things you can do with HBX Analytics. Next up in our series of cool things about analytics article series, is Visual Sciences' HBX Analytics.
- Metatags and SEO - Stone Temple Consulting**
www.stonetemple.com/articles/metatags-and-seo.shtml ▾
Metatags and SEO. by Eric Enge. The role of metatags is misunderstood by some members of the webmaster community. I still have people who are not in the business ...
- The Great Knowledge Box Showdown: Google Now vs. ...**

Figure 11-19. Bing site command results

This search shows 10,500 indexed pages. Of the two, you should consider the report from Bing Webmaster Tools to have more accurate data than the results from a site search at Bing.

Bear in mind that Google Webmaster Tools and Bing Webmaster Tools do not keep track of historical data for you. If you want to maintain a history of the data you should plan on downloading and saving the reports.

Link-Based Tracking of Content Marketing

Measurement of content marketing campaigns is another must for the expert SEO practitioner. This can help you refine your campaigns to bring better results, compare link growth to your search rankings, and perform analysis to see how your campaigns are faring compared to the efforts of your competitors.

Some analytics tools have built-in functions for comparative anchor text analysis.

Search-engine-supplied tools

There are three basic tools for checking backlinks:

- Google Webmaster Tools is a powerful start. With Google Webmaster Tools, publishers can easily download a spreadsheet of all the links Google has in its database (note that Google might not include some links that it does not consider significant). With this tool, publishers can see only the links to their own site. Figure 11-20 shows the report.

- Bing Webmaster Tools is also a great asset. It offers a similar capability for downloading a spreadsheet of the links that Bing has in its database for a site. Once again, publishers can use this tool only to see the links to their own site. Figure 11-21 shows what the report looks like.

Overview » All domains			
Top 1,000 domains that have links to pages on your site.			
		Show 25 rows ▾	1-25 of 1,000 < >
Domains	Links ▾	Linked pages	
aggregate.com	5,250	4	
vkistudios.com	1,762	1	
wordpress.com	1,409	12	
businessblogconsulting.com	1,373	2	
nicklaskoski.fi	1,070	1	
seobook.com	789	2	
seokicks.de	611	3	
problogger.net	319	2	
thebillionairegirl.com	299	1	
aboutonlinematters.com	285	2	
blogspot.com	281	17	
57883.net	217	1	

Figure 11-20. Google report on external links



Figure 11-21. Bing report on external links

For quick and dirty link totals, it is handy to use a Firefox plug-in known as SearchStatus (<https://addons.mozilla.org/en-US/firefox/addon/searchstatus/>). This provides basic link data on the fly with just a couple of mouse clicks. Figure 11-22 shows the menu you'll see with regard to backlinks.



Figure 11-22. *SearchStatus* plug-in

Notice also in the above figure that the *SearchStatus* plug-in offers an option for highlighting *NoFollow* links, as well as many other capabilities. It is a great tool to help pull numbers such as these much more quickly than would otherwise be possible.

An option for quickly referencing link totals in Google Chrome is Majestic's Backlink Analyzer browser extension (<https://chrome.google.com/webstore/detail/majestic-backlink-analyze/pnmjaflneibolacpepklokkjnakkmikmg?hl=en>).

Majestic Summary for: <https://www.stonetemple.com/> 

Source URL Anchor Text	Citation Flow	Trust Flow	Date Indexed Last Seen
54 Computers/Internet/Web Design and Development http://www.seo-united.de/blog/  stonetemple.com	50	56	26 Feb 2015 03 Mar 2015
50 Business/Marketing and Advertising http://www.marketingprofs.com/charts/201...  stone temple consulting	49	50	28 Feb 2015 03 Mar 2015
38 Business/Marketing and Advertising http://www.marketingdirecto.com/actualid...  stone temple consulting	44	38	02 Mar 2015 02 Mar 2015
26 Business/Marketing and Advertising http://forimmediaterelease.biz/index.php...  stone temple consulting	48	27	19 Dec 2014 01 Mar 2015
40 News/Media Industry http://blog.customscoop.com/feed/  stone temple consulting	33	40	05 Dec 2014 02 Mar 2015
30 Computers/Internet http://feeds.feedburner.com/imnewswatch/  stone temple consulting	34	31	03 Mar 2015 03 Mar 2015
22 Computers/Internet/Searching http://www.localseoguide.com/15-local-se...  stone temple	39	25	06 Jan 2015 02 Feb 2015

Figure 11-23. Majestic Chrome extension

The Majestic plug-in gives you rankings for the Citation Flow (authority) and Trust Flow (trustworthiness) as well as the number of external backlinks and referring domains for all of the current URL, the subdomain that URL resides on, and the root domain for the current site.

Another useful tool for seeing link data at a glance is the MozBar Chrome and Firefox extension from, of course, Moz (<http://moz.com/tools/seo-toolbar>).

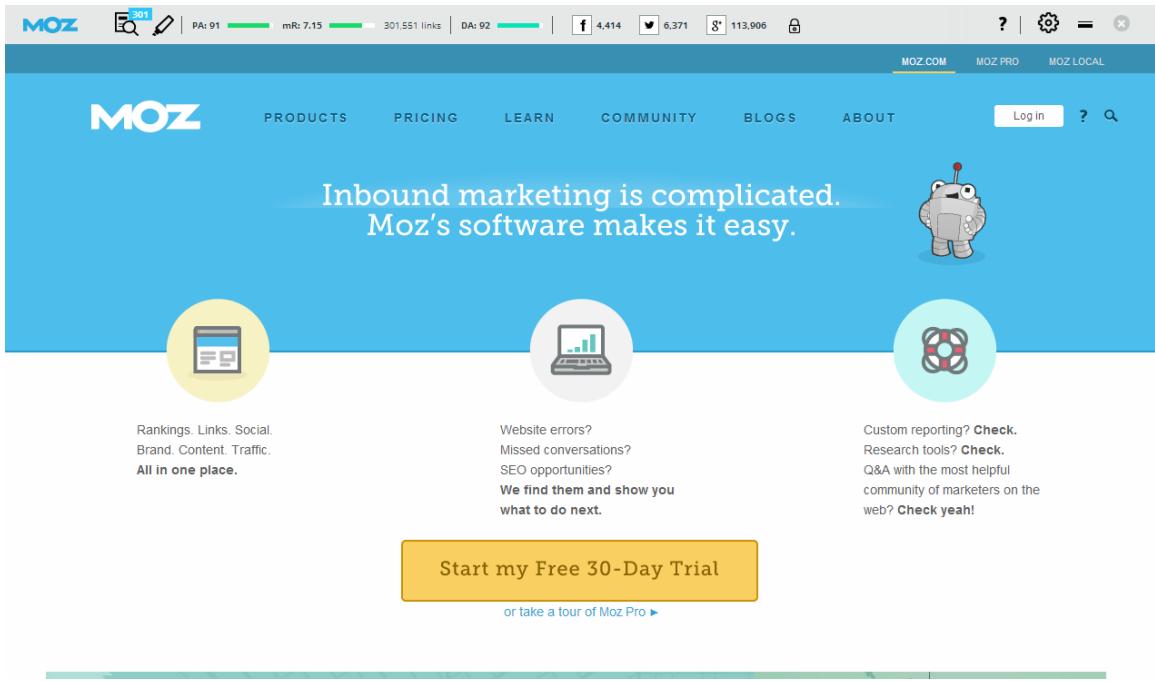


Figure 11-24. MozBar Chrome extension

The MozBar, shows a great deal of data which is useful for SEO. For the purposes of quick back link analysis it shows Page Authority, MozRank (Moz's answer to PageRank), total number of links to the current URL, and Domain Authority. The figure above shows the data you can get for free. If you have a paid Moz membership additional link metrics appear in the toolbar including the numbers of unique linking domains and data for the subdomain.

Third-party link-measuring tools

Advanced third-party tools exist for gathering link data. Here is a look at some of the most well known ones.

Open Site Explorer

Open Site Explorer (<https://moz.com/researchtools/ose/>) was developed based on crawl data obtained by Moz, plus a variety of parties engaged by Moz. When you enter a competitor's URL in Open Site Explorer, you can scope out where they're getting links, discover which pages are the most influential, and compare their link profile to yours to see how your performance stacks up. It lets you get as many links as it can find and then extract them all into a spreadsheet.

Figure 11-25 shows a sample report.

Inbound Links				
Gauge a site's influence. See inbound links to the page, subdomain, or root domain you've entered and analyze the linking pages.				
Target	Link Source	Link Type		
this page	only external	all links	<input type="checkbox"/> Group links by subdomain	
1 - 50 Inbound Links				
Title and URL of Linking Page	Link Anchor Text	PA	DA	
Gizmodo - Tech By Design  gizmodo.com/	io9 io9.com/	94	92	
Lifehacker - Tips and downloads for getting things done  lifehacker.com/	io9 io9.com/	93	92	
Gawker - Today's gossip is tomorrow's news  gawker.com/	io9 io9.com/	93	91	
Advertising Age - Advertising Agency & Marketing Industry News - Advertising Age  adage.com/	io9 io9.com/	92	91	
Kotaku - The Gamer's Guide  kotaku.com/	io9 io9.com/	91	89	
Jezebel - Celebrity, Sex, Fashion for Women. Without Airbrushing.  jezebel.com/	io9 io9.com/	91	89	
Deadspin - Sports News without Access, Favor, or Discretion  deadspin.com/	io9 io9.com/	90	88	
Jalopnik - Drive Free or Die  jalopnik.com/	io9 io9.com/	89	86	
Valleywag - Sorry to disrupt.  valleywag.gawker.com/	io9 io9.com/	89	91	
This Is Apple's Next iPhone  gizmodo.com/5520164/this-is-apples-next-iphone	io9 io9.com/	86	92	

Figure 11-25. Open Site Explorer report

Included with this report are Page Authority and Domain Authority scores for each URL listed, and each domain. Be aware that these metrics are useful in assessing the value of a domain or a page, but what the search engines actually measure and use may be substantially different. The tool also flags redirects so that you can see them. The beauty of this tool is that it allows SEO practitioners and publishers to collect competitive intelligence on other people's domains. All the data is extractable in a spreadsheet to allow for easy offline processing.

Majestic

Majestic offers a comprehensive backlink analysis for your domains. Majestic's flagship tool is the Majestic Site Explorer and can be used for a variety of jobs. With this tool you can quickly analyze any website's backlinks. Majestic has two unique metrics named Trust and Citation Flow. These scores are scored between 0-100 and flow through the entire web! The Trust Flow is a good quality metric, using this you can easily analyze how trustworthy a website or backlink is. The Citation flow can show how influential a site or page can be based on the number of links coming into each page of every site on the internet and working out the proximity of each page to all those links, whether linking directly or indirectly to the page (or site) you are investigating. You can call this the Link Juice. With these scores you can also help determine if a backlink is possibly toxic looking specifically at the ratio between the two scores. A page with a high Citation score but low Trust may be a sign of a page you want to avoid getting a link from as the links pointing to that page are numerous but taken together, untrusted. You can check all the new

links, the lost links, also Anchor texts and pages within this one tool. The Site Explorer is particularly useful for auditing and reviewing your website.

Topical Trust Flow, a newly available feature of Trust Flow, is also worth noting. It breaks down the trust flow of a site by industry of the link or connection to identify what your trust flow consists of. Then, you can check out the Topics tab for a more in-depth look at the exact amount of external backlinks and linking domains from that type of industry. This can be great for scoping for PR opportunities; you can even filter the chart by industry to get more specific industry topics. For example, filtering by the industry of “computers” gives you the topics about web design, Internet resources, open source, domain names, shareware, and dozens more.

The Clique Hunter is another clever tool that can be used to see where your competitors may have the edge on their backlinks in comparison to yours. Input up to 10 different domains or competitors to see where they are getting their backlinks and where you have missed out within your industry neighborhood. Needless to say, this is another little gem from Majestic that can be used for link prospecting.

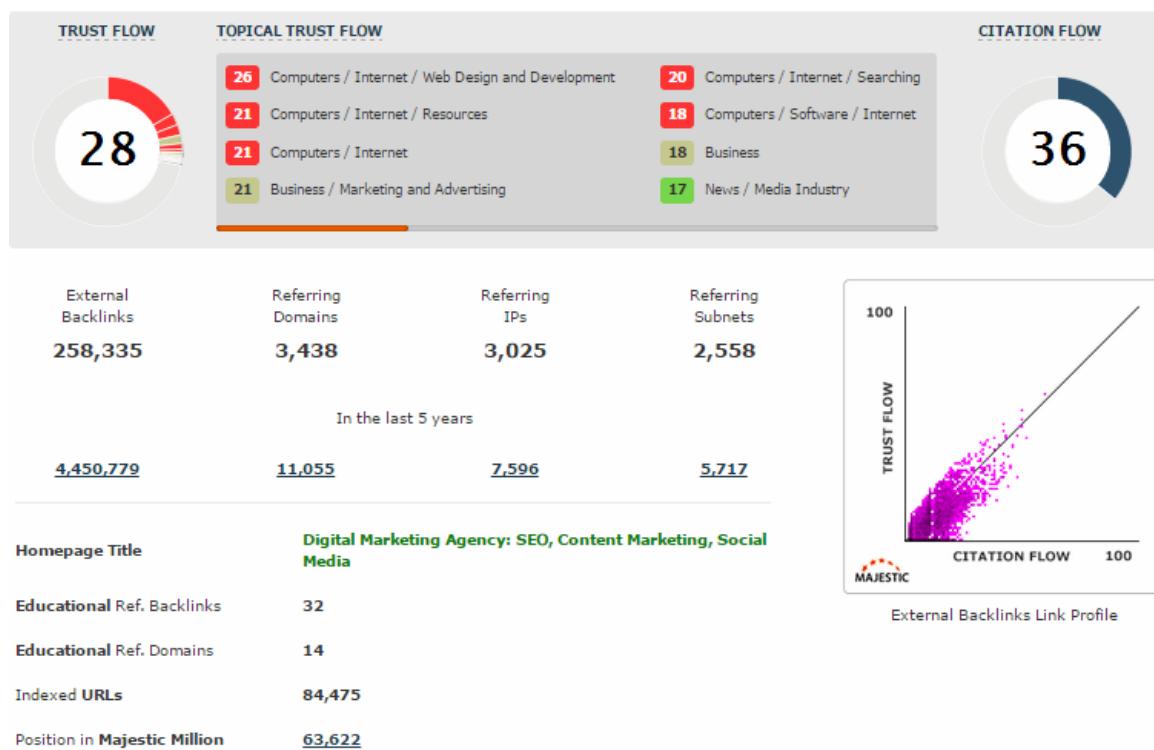


Figure 11-26. Majestic backlink report

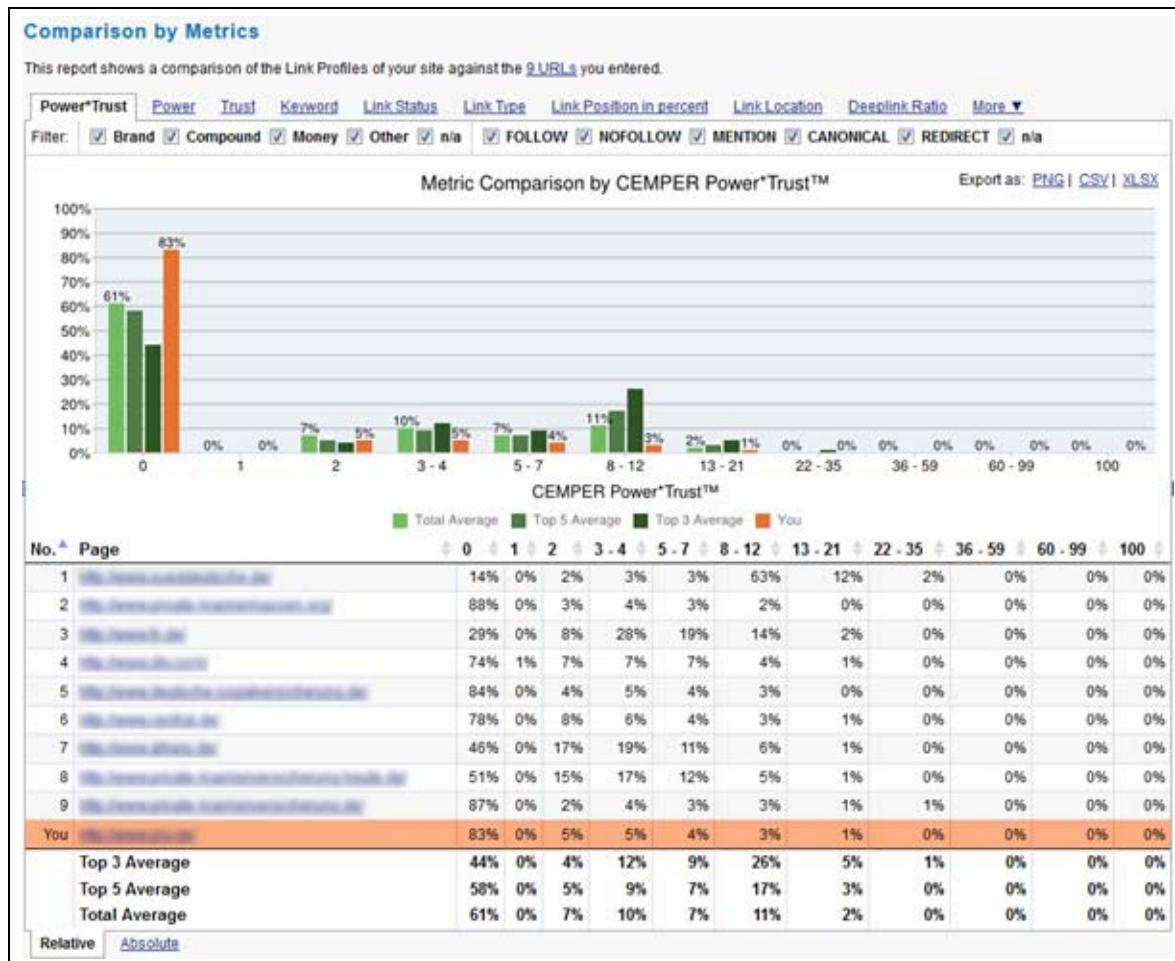
LinkResearchTools.com

European based Cemper.com (<http://www.cemper.com/>) offers a link building toolkit LinkResearchTools.com (<http://www.linkresearchtools.com/>). The LRT toolkit is a report-oriented system that provides qualified link data in real time from 23 sources. Some basic SEO metrics include: number of links to page and domain, linking root domains (domain popularity), Theme of the site, and the CEMPER Power* Trust™ metrics. In addition, Link Research Tools reports can be customized to international search engine, language, and country settings. When analyzing links from a domain in such great detail, it is therefore possible to find potential backlinks, identify link networks, analyze competition, and discover toxic and suspicious links.

The Link Research Tools includes a number of tools for different use cases. Quick Backlinks and Quick Domain Compare tools provide top links of a domain/page and ability to compare SEO metrics of other domains within a minute. Backlink Profiler and the Juice Tool allow the user to analyze link profiles for any domain. Link Juice Thief, Missing Links Tool and SERP Research Tool helps identify new link prospects. Link Alerts monitor you or your competitors' link growth. Common Backlinks Tool is a hub finder to find sites powering multiple competitors. Strongest Subpages Tool identifies a domain's strongest subpages. Link Juice Recovery Tool helps a Webmaster identify missing/broken pages with external links. The Contact Finder finds contact information for domains through searching the site for email.

The Link Detox tool is primarily used to identify bad links easily and create the necessary file to submit to the Google Disavow tool. The data-driven organic Link Detox Genesis algorithm helps locate unnatural links by analyzing a website link profile and automatically classifies them as toxic, suspicious, and healthy links. Furthermore, the Link Detox Risk Metric assists with link risk management and provides safe search engine optimization. Webmasters should now allocate time and efforts to clean toxic and suspicious links from their sites' link profiles.

A screen shot of from the Competitive Landscape Analyzer is shown in Figure 11-27. For competitive research, Link Research Tools includes the following tools; Competitive Landscape Analyzer, Competitive Link Detox, Competitive Link Velocity, and Competitive Keyword Analyzer. All relevant information and SEO metrics are crawled and visualized in clear graphics. The overview gives you a qualitative statement of how your link profile appears compared to the top competitors in your niche/industry. In this example, 9 competitors are analyzed by their Power*Trust, which measures both the Power and Trust of a site.



Another Figure 11-27. LinkResearchTools report

Ahrefs

Ahrefs (www.ahrefs.com) offers a comprehensive website analysis toolset with strong focus on backlink analysis. Ahrefs builds its own index and crawler that can crawl billions of pages daily and update databases with fresh backlinks found on the web every 15 minutes. With this capability, Ahrefs is well known as having the largest index of live backlinks and covering universal results (all TLDs included) to support users worldwide. Ahrefs also designs its own 100-scale AhrefsRank (AR) that has a strong positive correlation value with PageRank model for you to easily measure and compare your domain (or link) quality to your competitors.

Ahrefs Site Explorer is a link analysis tool that helps you to perform granular analysis on link profiles for your or your competitor's sites including identifying where the site links are coming from, any problematic backlinks (or unnatural links) are linked to the site, how many backlinks are newly acquired or lost in 60-day time range, as well as exploring new link-building opportunities by reverse engineering your competitors' backlinks. SEO Report is a handy tool that allows you to rapidly evaluate on-page SEO performance of the site. This often involves thorough examination of how accessible the site is to the search engines, and SEO-friendliness of the site information architecture.

Ahrefs tools are straightforward with neat UI and easy-to-read data visualization. Social metrics including Facebook Shares/ Likes, Tweets, and Google +1s are available at both domain and link level. Filtering options are available for almost every metric, which makes the tools highly customizable. Data are downloadable in CSV format for further data sorting or uploading to third-party research tools while backlink report is available in PDF format.

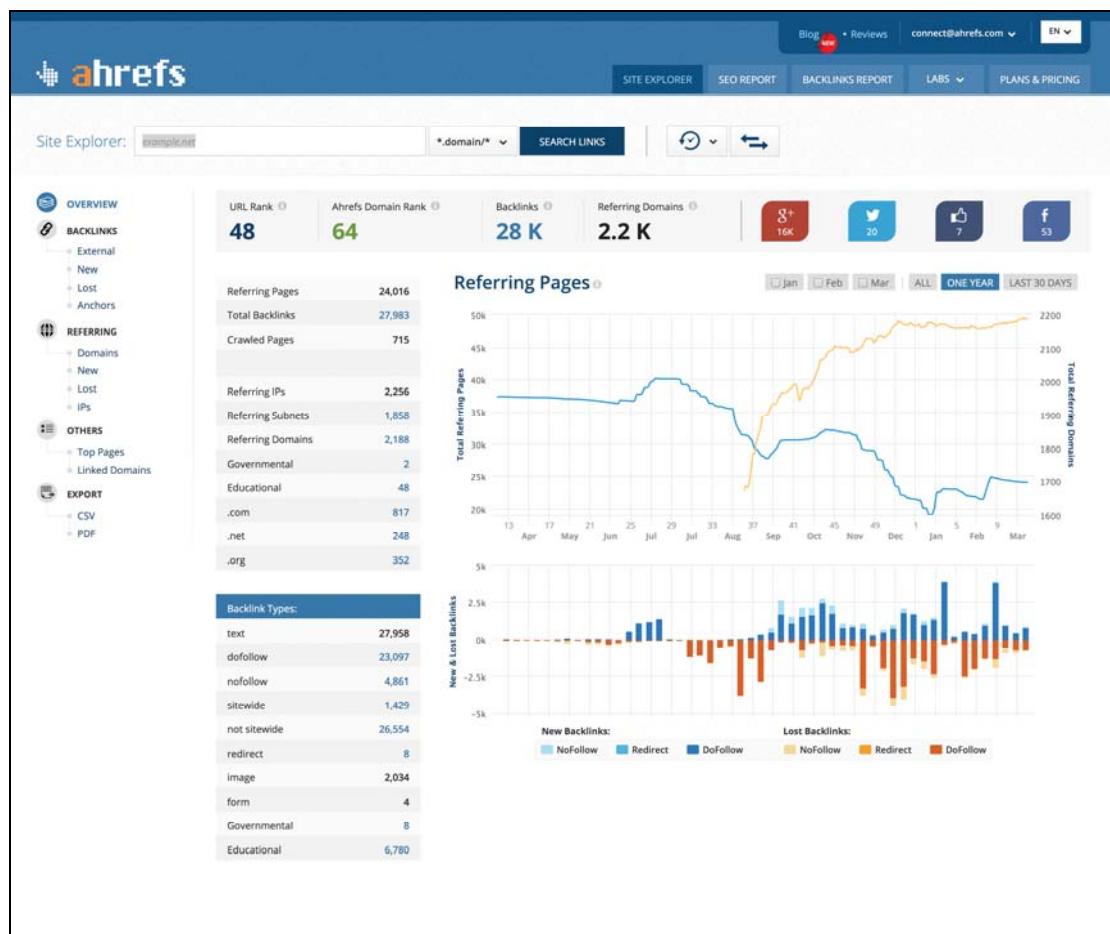
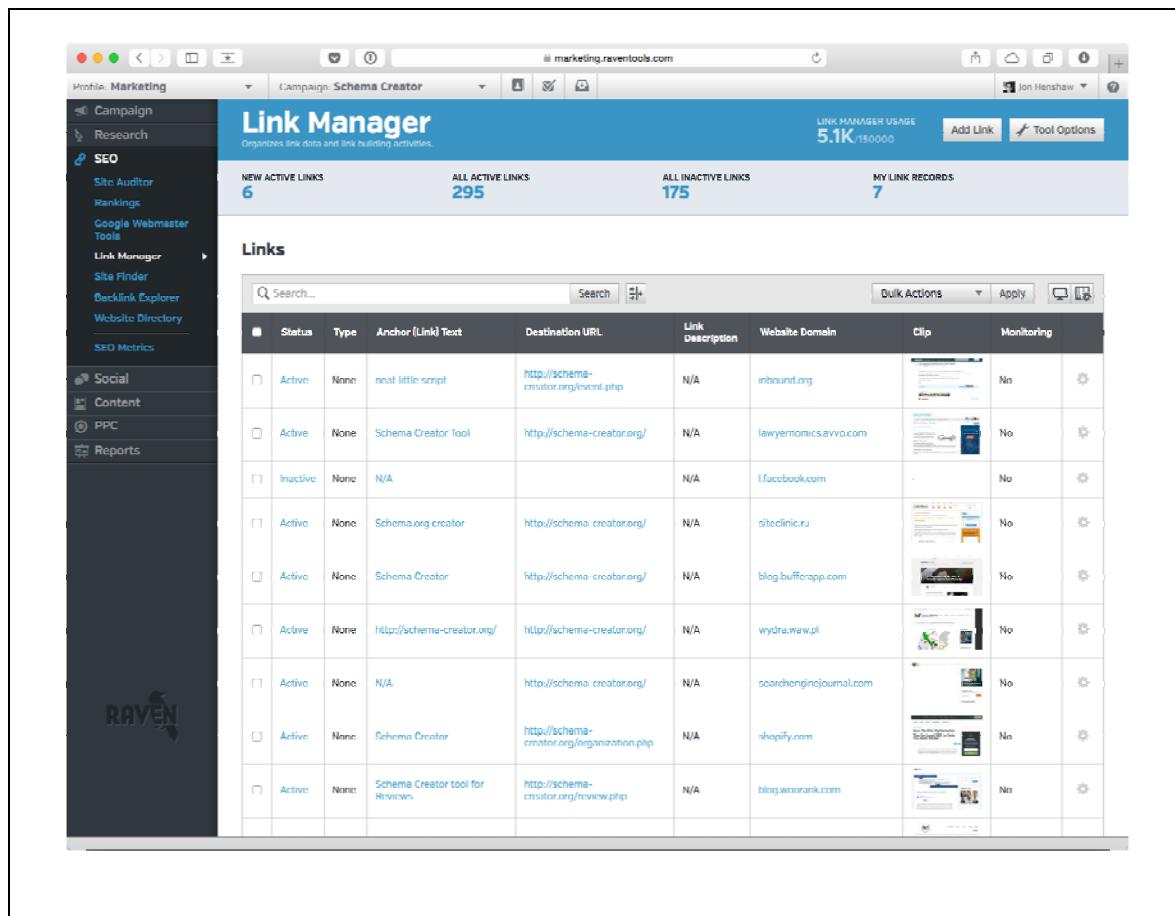


Figure 11-28. Screenshot of the Overview report under Site Explorer

Raven Tools

Another comprehensive toolset for link development is provided by Raven. One of Raven's tools is the Link Manager (<http://raventools.com/marketing-tools/link-manager/>), which tracks links that you plan on acquiring, have requested, or have successfully built. Link Manager includes conversion and ROI tracking. Link Manager will go out and automatically check the identified pages to see whether you have acquired a link on them, and it will tell you the results—thus automating the process of identifying targets and tracking the progress of one's campaign on an ongoing basis.

Raven also offers Research Central, which mashes up data from Majestic, Open Site Explorer, and SEMRush. As shown in Figure 11-29, Research Central generates a screenshot of the site you're researching and displays results for the domain's Authority/Ranking, Social Metrics, Site Metrics and Keyword metric. Raven also offers a tool called Site Finder that allows you to enter a keyword phrase, and then the tool will automatically review the top 10 Google results for that phrase and evaluate the value of a link from each of those pages.



The screenshot shows the Raven Tools Link Manager interface. The left sidebar has a dark theme with categories like Marketing, Research, SEO, Social, Content, PPC, and Reports. The SEO section is expanded, showing Site Auditor, Rankings, Google Webmaster Tools, Link Manager, Site Finder, Backlink Explorer, Website Directory, and SEO Metrics. The Link Manager section is selected. The main content area has a blue header 'Link Manager' with subtext 'Organizes link data and link building activities.' and 'LINK MANAGER USAGE 5.1K / 150000'. It shows summary statistics: NEW ACTIVE LINKS 6, ALL ACTIVE LINKS 295, ALL INACTIVE LINKS 175, and MY LINK RECORDS 7. Below this is a table titled 'Links' with columns: Status, Type, Anchor (Link) Text, Destination URL, Link Description, Website Domain, Clip, and Monitoring. Each row shows a link entry with a checkbox, status (Active or Inactive), type (None or Schema Creator), anchor text, destination URL, link description (N/A), website domain (e.g., inbound.org, lawyermomics.avvo.com, facebook.com, siteclinic.ru, blog.buzzforapp.com, hydra.waw.pl, searchinginjournal.com, shopify.com, blog.wmrank.com), a thumbnail image of the website, and monitoring status (No).

Figure 11-29. Raven Tools Link Manager

LinkBuilder and LinkGalaxy

One commercial alternative for link analytics is Rio's LinkBuilder™ solution powered by their LinkGalaxy link collection vehicle, which is geared more toward enterprise-class clients.

Link quality is much more important than link quantity, and the LinkBuilder feature provides multiple unique views to analyze and prioritize which link hubs are the most critical to obtain, in order of importance.

Figure 11-30 shows a LinkHub analysis report from the Rio LinkBuilder tool.



Figure 11-30. Rio LinkBuilder LinkHub Analysis for "toothpaste"

Conductor Searchlight

The Conductor Searchlight platform transforms a brand's web presence in 'unpaid' channels, like organic search, content and social, into a powerful acquisition method resulting in meaningful traffic & revenue increases. The technology gathers data from the web to deliver insight into where a brand's content is getting discovered, how their competition is being found, and the playbook needed to win the 'unpaid' media war.

Conductor Searchlight makes use of the Open Site Explorer data set from Moz and a sample report is shown in Figure 11-31.

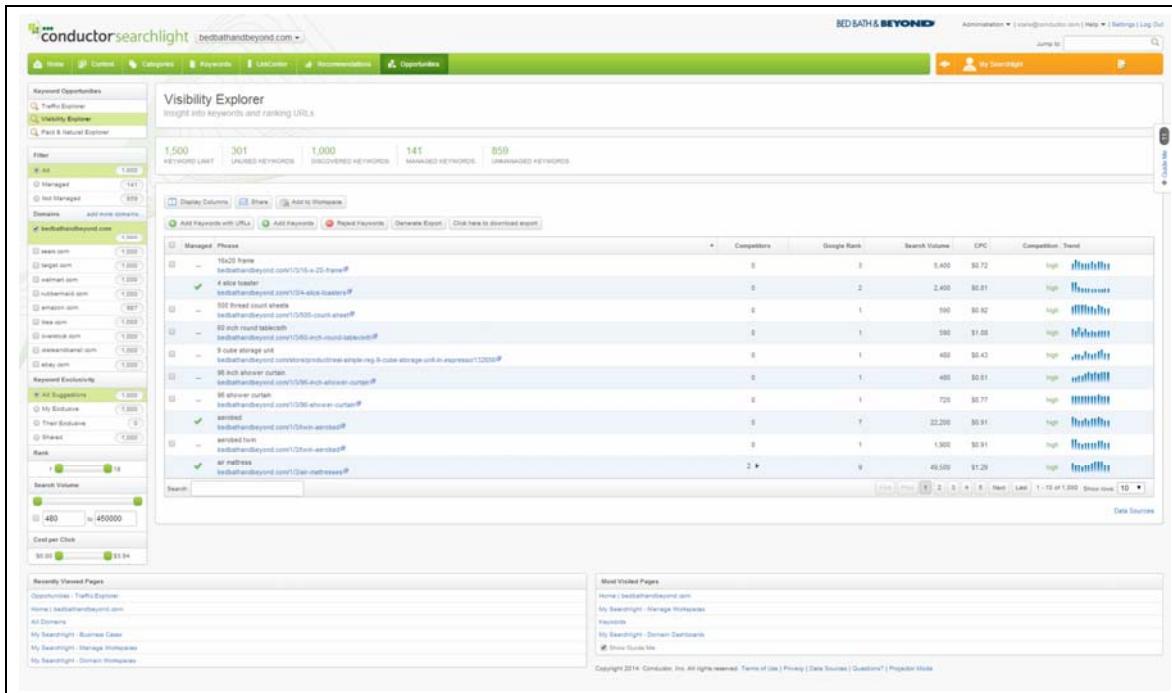


Figure 11-31. Conductor Searchlight

Exalead

Exalead (<http://www.3ds.com/products-services/exalead#>) has developed a European search engine (<http://www.exalead.com/search/>), and this offers capabilities that include a `link:` command. What makes Exalead particularly interesting are the advanced filtering capabilities it provides, as you can see in Figure 11-32.

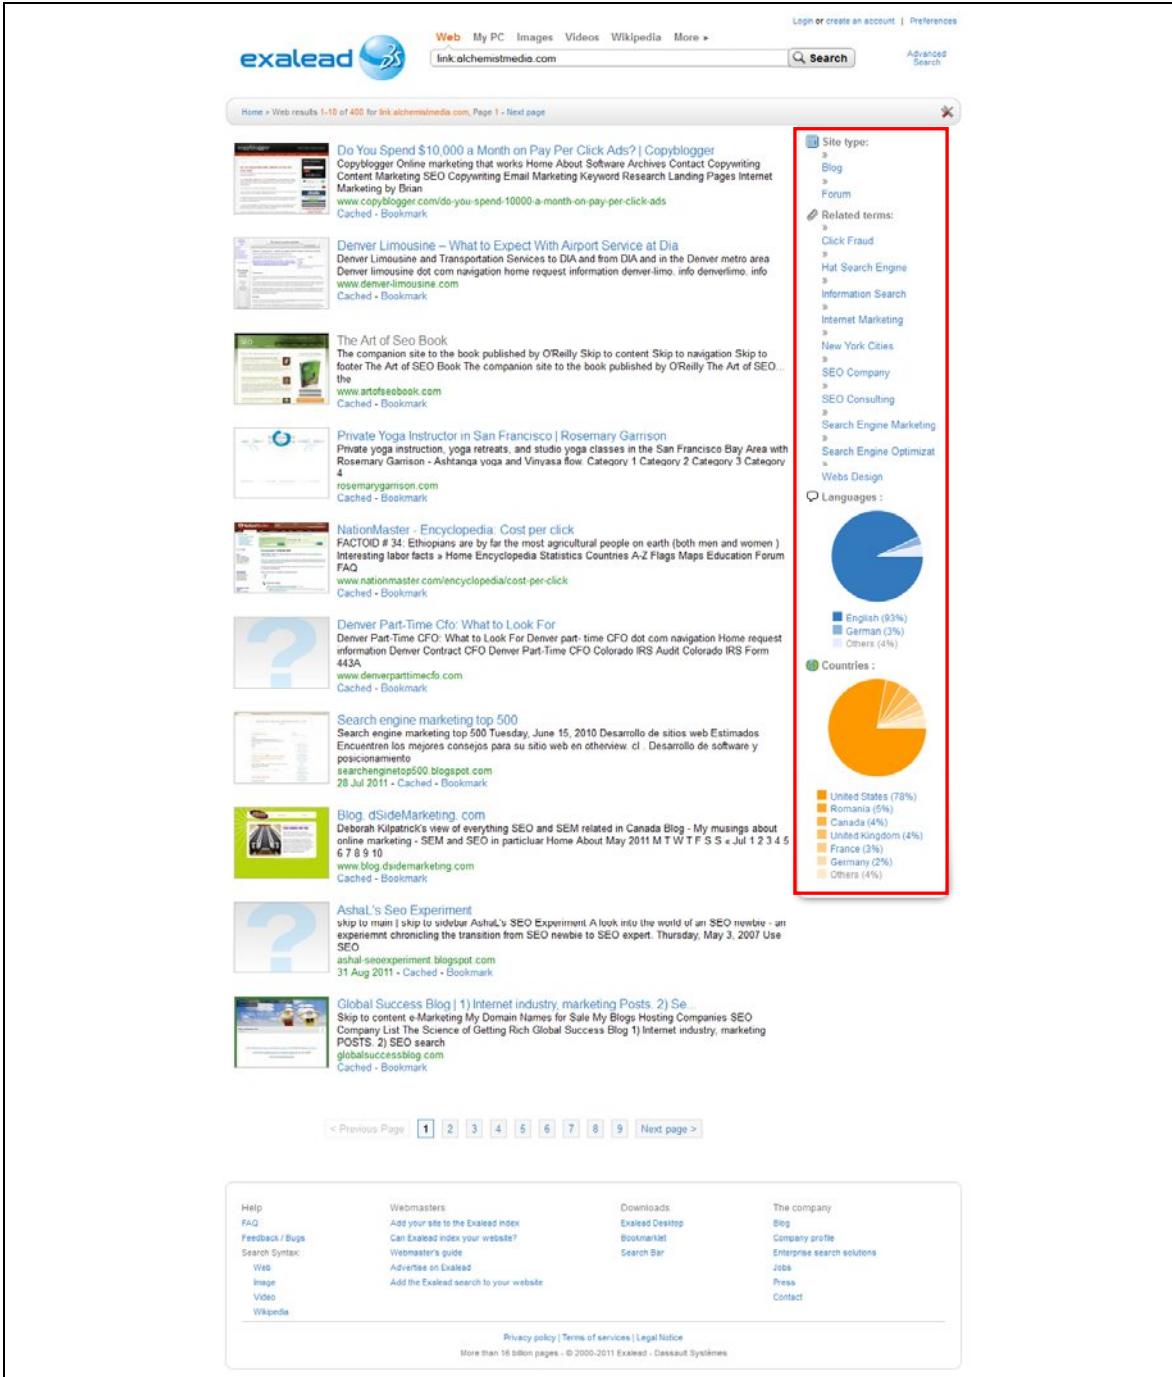


Figure 11-32. Exalead link data

Note the ability to pick language as well as all the other options such as site type, which could be useful for segmenting link data. These filters are shown in a red box.

Measuring the value of a link

One of the big questions that people ask is what is the value of a particular inbound link? There is no simple way to answer that question, but you can look at some of the metrics that can give you a feeling for how important a link might be:

Most important elements to determining a link's value

Where does the linking page rank for the term/phrase you want to rank for?

If the page is ranking #1 at Google for *sliced bread* and you want to be #1 at Google for *sliced bread*, guess what? That's the #1 most valuable link you can get. Keep going down the list to about positions 25 to 30 and you're still getting solid gold in link value.

Where does the linking page rank for one to two important, competitive terms in its title tag?

This will give you a very solid idea about how much overall link juice and respect the search engines are giving the page. It is also a good way to identify the global link value that could be provided.

Where does content on the linking domain generally rank for competitive terms in its pages' respective title tags?

As in the preceding list item, we're trying to identify how positively the engines view pages on the domain. If the pages generally rank in the top 20 results, you can rest assured that search engines think the domain's value is pretty high, and that links from that domain will pass significant value.

How many keyword phrases does the linking domain and page rank for in the top 20 results?

Sites that rank for a very small set of keywords may be over-optimized, whereas sites that rank for relevant keywords provide your site with more relevant links.

Does the linking site carry any brokered sets of links?

Sites that sell links may lose their ability to pass link authority. This really applies to any type of low-quality, manipulative linking. If you can see it, chances are that Google might see it someday too. In addition, Google may penalize a site retroactively for repeated bad behavior, even if that behavior has been stopped.

What is the relevance of the linking page/site to your target page?

Answering this question requires you to think critically about the visitors to both the potential linking page and the domain. If the relevance of the subject matter to your site is high, the link will provide more semantic and topic-specific value.

When was the last time Google crawled the page?

A fresh timestamp is a great indicator of the importance of the page. The older it is, the less likely it is to be relevant to a modern search query. You can check a page's timestamp by looking at its cache in Google.

Inlinks to page

Next, look at the links to the specific page you want to get a link from (or perhaps that already links to you). You want to know whether the domain links into this individual page heavily, or whether it is practically an orphan. Also, the more traffic that site gets, the more likely that traffic will pass on to your page. See whether it is a page that other sites reference, both of which can help illuminate potential value. The Hilltop algorithm is an important algorithm that was adopted by Google in 2003; it essentially says that if your site is listed along with other relevant and quality sites within a niche on a site, that site that lists it is considered a topical authority, and your

site is given more value. An example of this might be "10 Startups to Watch in 2014" by Entrepreneur Magazine.

Elements of secondary value for links

Links to high-ranking competitors

Although this isn't always an indication of direct value, it can be a good signal. Your competitors are obviously ranking based on the strength of their links, so researching those sources can provide insight into where they derive that value.

PageRank of the domain

There is a notion that the PageRank of a domain would be the overall PageRank value Google would associate with the domain. This is not a value you can get at directly, but you can approximate it by looking at the PageRank of the website's home page. Look at this to make sure it is not penalized and to see the overall link authority. You can also get an assessment of the link strength of a domain using Moz's Open Site Explorer tool, which we discussed earlier in this chapter.

A PageRank 6 (or higher) home page clearly has some link love and respect, a Page Rank 2 obviously has less, a gray bar can be a good red flag, and seeing a PageRank 0 tells you that the domain is either new or completely invisible. However, as link-building expert Eric Ward points out, avoiding lower PageRank domains simply because they are PageRank 3 or 4 is not necessarily a good thing, as you could be missing out on very relevant links that, in volume, contribute to your relevance.

Please note that after October 2013 PageRank may decline as a link quality signal as time passes. Google's John Mueller noted in a Google Webmaster Central Hangout in October 2014 that the company had not updated the page rank ratings available to third parties (Toolbar PageRank) for over a year and may well not update again in the future. (<http://www.youtube.com/watch?v=GlxLlpm3Ew>).

Inlinks to domain

Technorati can show trends—if a site has been picking up many new links over the past few months, it might be a much better candidate than the PageRank or other link data might indicate. This is a pretty indirect measurement, but it's not completely useless. However, since the number often takes into account lots of links from a single domain, it can be misleading.

Number of external links on the page

Pages generally pass some of their link authority to each page they link to. Therefore, if you have a link from a page with 10 outbound links, it is far better than having a link from an equivalent page with 100 outbound links, so this is a parameter worth knowing about.

Domain Authority/Page Authority

Domain Authority and Page Authority are two of the main metrics Moz (<http://moz.com/learn/seo/domain-authority>) uses to estimate the link authority of the site. Domain Authority is a combination of over 40 signals intended to track the overall "strength of a website" as compared to other sites. Page Authority measures similar metrics but is intended to rank the overall strength of individual pages as opposed to the domain overall.

Citation Flow, Trust Flow

Citation Flow and Trust Flow are metrics from Majestic (<https://blog.majestic.com/development/flow-metrics/>). They look, for both metrics, at the “flow” of these metrics through a series of links which ultimately lead to your site. Essentially, the flow metrics look not only at those links linking to a given page but also to the links that point at those links and so on up a “chain” of links. They calculate a rating which approximates not only the number of links flowing through to a URL but also the power and trustworthiness of those links. Citation Flow looks at the “strength” of the chains of links leading to a URL. Trust Flow looks at the trustworthiness of the chains of links leading to a URL.

LRT Power*Trust

As another alternative to Google's PageRank, Power*Trust (or CEMPER Power*Trust) is a metric calculated by CEMPER's LinkResearchTools (<http://www.linkresearchtools.com/link-tool-features/cemper-power-trust/>) and combines both the power (or authority) and the trustworthiness rankings from many of the above tools into a single metric. It's a useful metric to use to quickly judge the overall influence of a set of links pointed at a URL. It combines data from many sources in real time including data from Google, Moz, Majestic, and more.

Ranking

Many people get obsessed with checking their rankings. Rankings do have some correlation with traffic. However, do not spend too much time on rank checking. For one thing, many tools that are available today for checking rankings scrape Google results, which is against Google's Webmaster Guidelines. Google results vary by user, and the scraped results do not represent what everyone is seeing – as there is no such thing as a search result that “everyone” sees at this stage.

Rankings data can be helpful in a few scenarios, however:

- Where a #1 (or a very high) ranking for a trophy search term provides brand value
- When there is a sudden significant change (such as a move from position #5 to #50, or a similar change in the other direction)
- Where a long-term trend can be established, such as a steady improvement in position, or a steady decline
- In aggregate, where you can measure a site's overall activity and popularity

While evaluating this data, bear in mind that the search engines are constantly being fine-tuned. Google is believed to adjust and test its algorithms every day. This can result in a lot of movement in rankings on a regular basis, and this type of movement is usually meaningless from an SEO perspective.

What follows are some examples of tools that allow you to obtain rankings data. For one, you can get the data directly from Google Webmaster Tools (see 11-33).

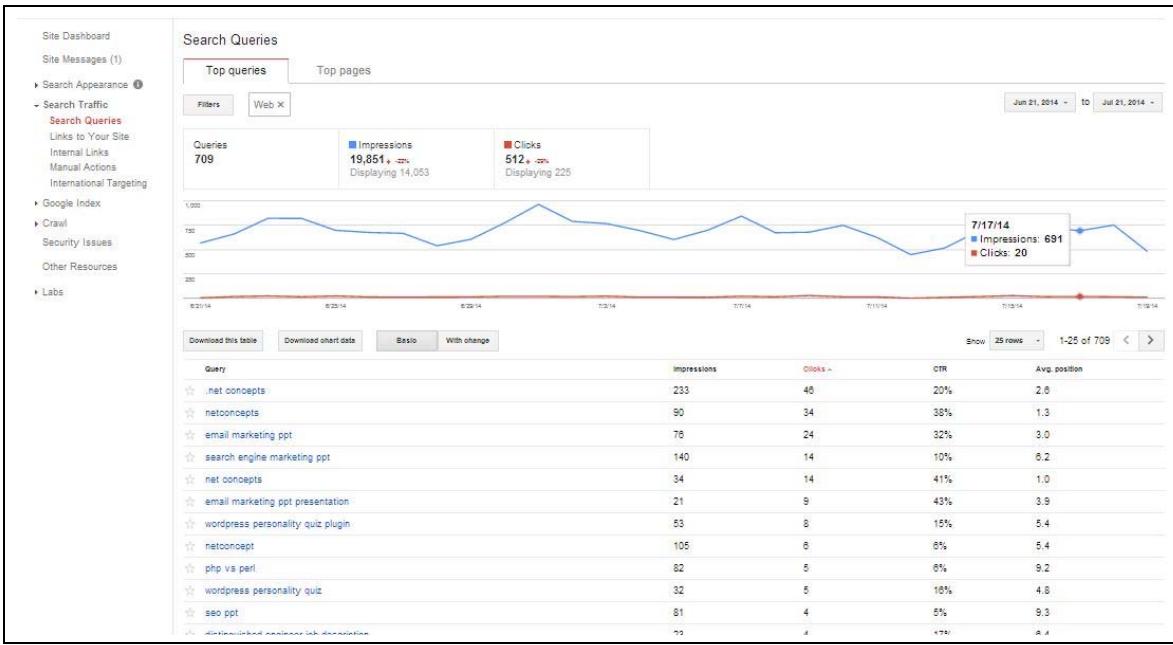


Figure 11-33. Top search queries report

Google handily shows the top queries for which your site came up, as well as the queries that were clicked on the most. This is nice data to give you a snapshot of where things stand, and you can download the data into a spreadsheet.

AuthorityLabs

AuthorityLabs provides an array of reporting tools that help you track rankings across the major search engines. One new reporting tool of note is AuthorityLabs new Now Provided tool, hailed as an answer to the myriad of difficulties brought on by the "Not Provided" update of 2013. Now Provided uses a mixture of Google Analytics data, competition data, and a few volume and ranking metrics to seek out your most important keyword opportunities. Then, understand how much traffic that keyword is getting you. Another feature of the tool allows you to see the number of pages sending traffic on any given day. A sample screen shot of Now Provided is shown in Figure 11-34.

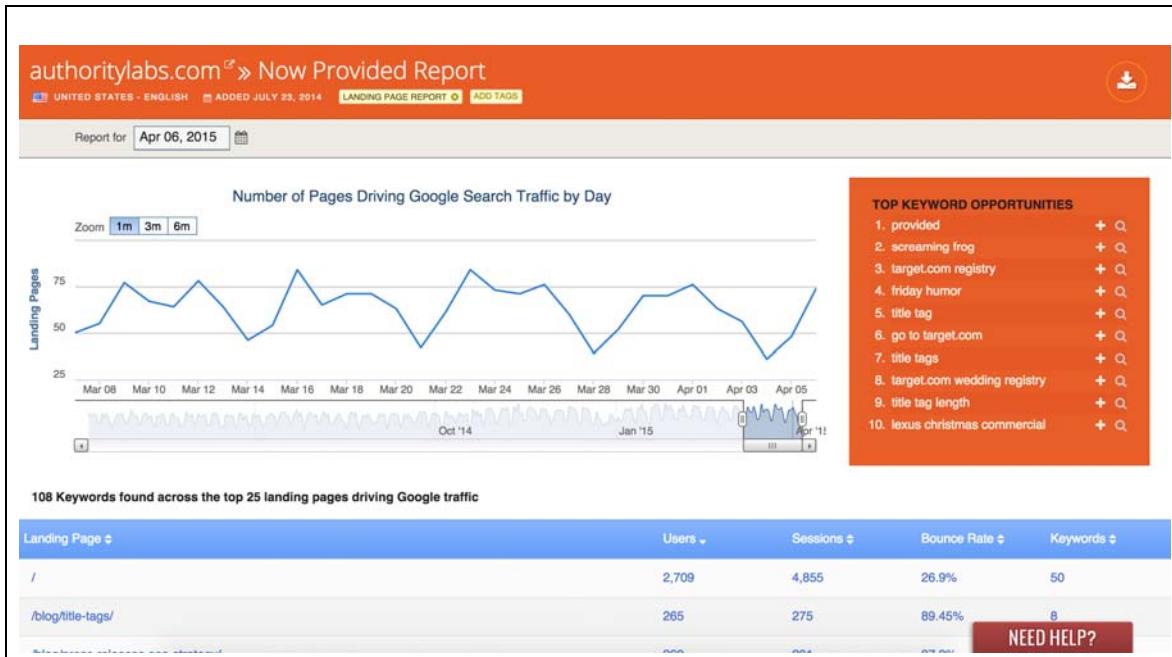


Figure 11-34: AuthorityLabs Now Provided report

Shelf Space

Shelf space is an SEO metric that assigns a fraction of the total SERP to each listing. For example, 33% for the first listing, 10% for the second listing, 0.0000001% for the 123rd listing, etc. Each point of presence yields more shelf space. Links to your site, images or videos in the results, and anything else that appears on a search result page all count toward shelf space. The sum of the shelf space for all of the listings for a particular query is 100%.

Shelf space summarizes the value of several listings on the same SERP. It allows you to easily examine the trend of your shelf space over time and compare it to your competitors' shelf space. You can also compute your average shelf space across 1000 keywords. If it's 10%, you know that you control 10% of those SERPs.

Here are some tips:

- When looking at SERP shelf space (or rank), treat your social media and user-generated content listings (Facebook, YouTube, etc.) as your own, including in “universal search” results.
- In reporting, draw a distinction among the shelf space controlled by your social media pages, your preferred landing pages, your other webpages, your competitors’ pages, and independent (Wikipedia, government, university, etc.) pages.
- If you care about more than a few keywords, or if you want to get a lot of long tail traffic to your site, categorize your keywords and queries. Typical categorization dimensions include branded vs. non-branded, by product, by subject or topic, and by geographical location mentioned in the query.
- Benchmark your SERP shelf space changes against your competitors. Then you can tell whether a change is due to something you did or due to a change in Google's algorithms.
- YouTube and Facebook are search engines, too – you should care about your rankings on those just as much as you do on Bing.

- When dealing with paid/organic interaction, the following simple formula estimates the cost per incremental click (CPIC; or cost per incremental conversion) for paid search above and beyond what organic would drive to your site in the absence of paid:

$$CPIC = CPC / (1 - \text{shelf space})$$

You can estimate cost per incremental click and cost per incremental conversion at the query level or at the category level, as described above. Paid search marketers can optimize their campaigns to minimize their measured cost per incremental conversion, truly maximizing their return, rather than choosing the ads that are best at shifting clicks from their own organic to paid listings.

SEO Platforms

Visibility into your and your competitors search rankings over time is an important indicator of your search landscape and share of search. The following companies are examples of ones that offer SEO platforms:

- Conductor offers Conductor Searchlight (<http://www.conductor.com/product/searchlight/>) which can crawl your website, enable visibility into your competitive search landscape, and bring on- and off-page data such as backlink and on-page factor analysis together in one place.
- Rio (<http://www.rioseo.com>) offers the Rio SEO Search Analytics which uses their patented SEO Audit Score methodology to measure the SEO health of your websites.

Crawl Errors

Complex websites will ultimately develop errors. This will come from moving pages where the publisher forgets to redirect the old page to the new page; or the publisher makes a mistake in her *robots.txt* file and prevents the crawling of key parts of her site. It can also happen that someone links to the publisher's site, but he or she implements the link incorrectly by linking to a page that does not exist (possibly due to a typographical error).

Tools from the search engines

All of these things are easy to diagnose in either Google Webmaster Tools or Bing Webmaster Tools. Figure 11-35 shows the Crawl Information page in Bing WMT.

Site ▾
stonetemple.com

Crawl Information 1

Status of your URLs from Bingbot's last crawl. Last updated on 6/29/2014 (Sun)

See Index Explorer for more details

400-499 (Request errors)	500-599 (Internal server errors)	301 (Page moved permanently)	302 (Page moved temporarily)	ROBOTS.TXT EXCLUSION	DNS FAILURES	CONNECTION TIMEOUTS
2	0	308	0	1,743	0	0

HTTP 400-499 CODES
These codes appear when there was a likely an error in the request which prevented the server from being able to process it.

URL	HTTP CODE	COUNT OF LINKS
http://www.stonetemple.com/category/seo/Google Plus Ranking Influence Test Results ↗	404	0
http://www.stonetemple.com/articles/interview-jakob-nielsen.shtml ↗	404	0
http://www.stonetemple.com/blog/wp-content/uploads/2009/05/glibc.report-activitytotals1-1024x505.png ↗	404	0
http://www.stonetemple.com/blog/wp-content/uploads/2009/05/glibc-report-queriesdriving-1024x642.png ↗	404	0
http://www.stonetemple.com/blog/wp-login.php?action=register ↗	404	0
http://www.stonetemple.com/brief-chat-with-guy-kawasaki/ ↗	404	0
http://www.stonetemple.com/category/ppc/page/4/ ↗	404	0

EXPORT ALL

Figure 11-35. Bing Crawl Information page

You can see that Bing will provide data on 400-499 errors, 500-599 errors, 301 redirects, 302 redirects, Robots.txt exclusion, DNS failures, and connection timeouts in this report.

Google provides similar data, with some notable differences (see Figure 11-36).

Site Errors

Showing data from the last 90 days

DNS Server connectivity Robots.txt fetch

URL Errors
Status: 5/15/14

Desktop 7 **Smartphone** 7

Server error 7	Soft 404 7	Access denied 7	Not found 7	Not followed 7	Other 7
1	0	1	52	0	4

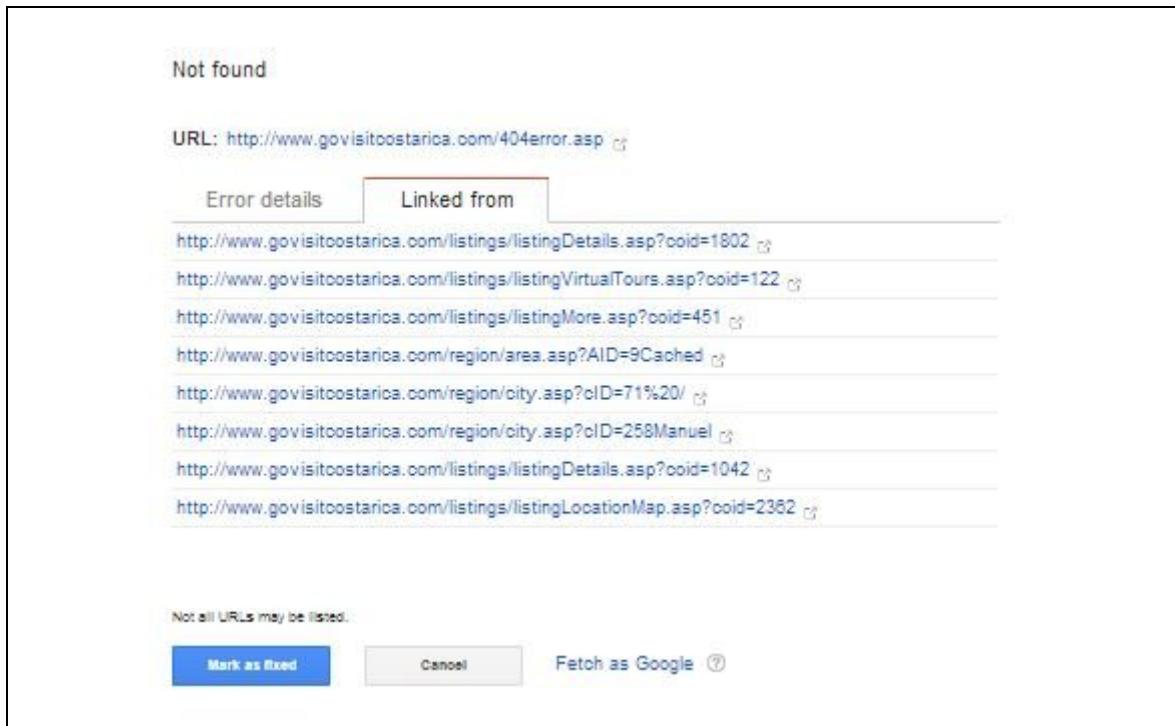
Top 1,000 pages with errors

Download MARK AS FIXED (0) Filter 7

Priority	URL	Response Code	Detected
1	travellInfo/beaches.asp%22%3Ewww.govisitcostarica.com/travellInfo/beaches.asp	404	4/2/14
2	404error.asp	404	5/13/14
3	images/uploads/detail/optimized/g-lt	404	5/6/14

Figure 11-36. Google crawl errors

Notice the additional types of errors provided. In addition, if you look at the 404 report in Google WMT, you have the option of seeing what pages link to the page generating the 404 error, as shown in Figure 11-37.



The screenshot shows a 'Not found' error page from Google WMT. The URL is <http://www.govisitostarica.com/404error.asp>. The 'Linked from' tab is selected, showing a list of URLs that link to the error page. The list includes:

- <http://www.govisitostarica.com/listings/listingDetails.asp?coid=1802>
- <http://www.govisitostarica.com/listings/listingVirtualTours.asp?coid=122>
- <http://www.govisitostarica.com/listings/listingMore.asp?coid=451>
- <http://www.govisitostarica.com/region/area.asp?AID=9Cached>
- <http://www.govisitostarica.com/region/city.asp?cID=71%20/>
- <http://www.govisitostarica.com/region/city.asp?cID=258Manuel>
- <http://www.govisitostarica.com/listings/listingDetails.asp?coid=1042>
- <http://www.govisitostarica.com/listings/listingLocationMap.asp?coid=2382>

Below the list, a note says 'Not all URLs may be listed.' At the bottom are buttons for 'Mark as fixed', 'Cancel', and 'Fetch as Google'.

Figure 11-37. Google 404 errors report

The great thing about the 404 reports in both tools is that you can spot pages that have been implemented with the wrong URL on your site, and you can spot whether someone links to one of your pages using the wrong URL. In the latter case, as we suggested in Chapter 7, 301-redirect the page that received the bad link to the correct page, and pick up an easy link. In short, 404 errors provide great opportunities to request updated links from sites linking to old URLs.

Another problem occurs when a page no longer exists but the web server erroneously returns a 200 status code. This is sometimes referred to as a *soft 404*, and it can occur when you've created a custom 404 landing page. In these instances, a "page not found" error will redirect to a valid page, which results in an HTTP 200 "ok" result. This can hide the presence of broken links. A soft 404 can limit a site's crawl coverage because the landing page URL may be indexed instead of the intended (but missing) unique content pages.

So in addition to traditional HTTP 404 errors, examine your reports for HTTP 200 errors and give them the same weight as a standard 404. Google Webmaster Tools provides 404 and 200 error codes in its reports, and can be helpful for identifying broken links that need to be 301-redirected.

Third-party tools to check for crawl errors

Some excellent third-party tools are available for seeing how a search engine crawler will look at your site. One of these is SEO-Browser (<http://www.seo-browser.com>). Use this tool to see how search engine crawlers look at your web pages. This can be useful during a technical SEO evaluation of a site, or when debugging site indexation problems. Figure 11-38 shows the basic view from SEO-Browser for the TripAdvisor (<http://www.tripadvisor.com>) home page.

seo-browser.com™

See your website like a Search Engine sees it.

[Request a Quote](#) [SEO Products](#)

Address: [Advanced...](#) [Simple](#)

We noticed that you're using an unsupported browser. The TripAdvisor website may not display properly.

We support the following browsers: Windows: Internet Explorer , Mozilla Firefox , Google Chrome .
Mac: Safari . [\[Reviews of Hotels, Flights and Vacation Rentals\]](#) [JOIN](#)

[LOG IN](#)

USD [x.gif] [\[International Sites\]](#) [x.gif]

[Hotels](#) [x.gif] [Search All Hotels](#) [Popular destinations](#) [Miami Hotels](#) [New Orleans Hotels](#) [Laughlin Hotels](#) [Corpus Christi Hotels](#) [Hampton Hotels](#) [Las Vegas Hotels](#) [Orlando Hotels](#) [Paris Hotels](#) [Atlanta Hotels](#) [Mackinaw City Hotels](#) [Reno Hotels](#) [Pismo Beach Hotels](#) [Lake George Hotels](#) [Hollywood Hotels](#) [Charlotte Hotels](#) [Pensacola Beach Hotels](#) [South Lake Tahoe Hotels](#) [Carmel Hotels](#) [Lake Geneva Hotels](#) [Lake Placid Hotels](#) [\[Travelers' Choice Hotel Awards\]](#)

The world's best hotels. Get the list .

[Flights](#)
[Vacation Rentals](#)
[Restaurants](#)
[Best of 2014](#)
[Your Friends](#) Travel is better with friends. Get inspired by your friends and share your travel plans. [Learn more](#)
[More](#) [x.gif] [Trip Ideas](#)
[Travel Forum](#)
[Apps](#)
[Cruises](#)
[GreenLeaders](#)
[Help Center](#)

Write a Review 1+ [x.gif] Review a place you've visited [spinner.gif]

Plan the perfect trip

[x.gif] [Hotels](#) [x.gif] [Flights](#) [x.gif] [Vacation Rentals](#) [x.gif] [Restaurants](#) [x.gif] [Destinations](#) mm/dd/yyyy mm/dd/yyyy

Figure 11-38. SEO-*Browser*

Interestingly, this report indicates that the website shows that the surfer used an unknown browser. It might be interesting to determine why that is happening. If you look at the Amazon home page using the Advanced mode, you can see some interesting data as well (see Figure 11-39).

The screenshot shows the SEO-Browser.com interface with the following details:

- Address:** <http://amazon.com>
- SEO-Browser Agent:** SEO-Browser Agent
- Page Title:** Amazon.com: Online Shopping for Electronics, Apparel, Computers, Books, DVDs & more (Characters: 87; Words: 12)
- Metadata:**
 - description: Online shopping from the earth's biggest selection of books, magazines, music, DVDs, videos, electronics, computers, software, apparel & accessories, shoes, jewelry, tools & hardware, housewares, furniture, sporting goods, beauty & personal care, broadband & dsl, gourmet food & just about anything else.
 - keywords: Amazon, Amazon.com, Books, Online Shopping, Book Store, Magazine, Subscription, Music, CDs, DVDs, Videos, Electronics, Video Games, Computers, Cell Phones, Toys, Games, Apparel, Accessories, Shoes, Jewelry, Watches, Office Products, Sports & Outdoors, Sporting Goods, Baby Products, Health, Personal Care, Beauty, Home, Garden, Bed & Bath, Furniture, Tools, Hardware, Vacuums, Outdoor Living, Automotive Parts, Pet Supplies, Broadband, DSL, google-site-verification: 9vpzZueNucS8hPqoGpZ5r10Nr2_sLMRG3AnDtNluoc4
- Text Metrics:**
 - # of Meta characters : 842
 - # of Meta Words : 107
 - Text to Page Weight Ratio: 2.1%
 - # of Meta Keywords Words : 41
- Page Load Metrics:**
 - Canonical Link: <http://www.amazon.com/>
 - Page loaded: 1.878 seconds
 - Frame Info: No Frames
 - Page size: 384.7 KB
 - Noframe Info:
- Content Metrics:**
 - # of Words: 1525
 - # of Body Words : 624
 - # of Images: All: 136; with Alt Text: 123; without Alt Text: 13
 - # of Links: All: 241 (External: 59, Internal: 182); No Follow: 0
 - Robots.txt: yes
 - IP address : 178.32.98.166
 - favicon.ico: file: no ; script: no
 - Country: United States
 - Lang: unspecified
- Script Metrics:**
 - # of External JavaScript: 27
 - # of Imagemaps: 0
 - # of Objects (Flash): 0
 - Clickstream Reporting: None

Figure 11-39. Amazon Home Page Show in SEO-Browser Advanced mode

The Advanced view provides a wide range of stats above the view of the web page itself. Down toward the bottom of Figure 11-40, you can see that Amazon offers a screen reader version of the page.

Another excellent tool is the Moz Crawl Test Tool (<https://moz.com/researchtools/crawl-test>). It provides an extensive set of statistics on any website you choose to check. Figure 11-40 shows a small sample of the data provided.

Crawl Test http://pro.seomoz.org/tools/crawl-test																																						
Report created: 10 Sep 04:13																																						
#-----																																						
URL	Time	Crawl	Title	Tag	Meta	Desc	HTTP	Stats	Referrer	Link Count	Content	T-4XX	Cler	Serv	5XX	(\$erv	Title	Miss	Duplicate	URLs with	Duplicate	URLs with	Long	URL	Overly	Dy	301	(Perm	302	(Temp	301/302	Ti	Meta	Refr	Meta	Refr	Title	Ele
www.stap	0	Office Sup Shop Stap	200	www.stap	528	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No			
www.stap	0	1-1/2" No Staples&R	200	www.stap	194	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No		
www.stap	0	1-1/2" Sta Staples&R	200	www.stap	170	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	1-1/2" Sta Staples&R	200	www.stap	183	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	1-1/2" Sta Staples&R	200	www.stap	171	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	200 www.stap	200	www.stap	166	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No
www.stap	0	1/2" Non- Staples&R	200	www.stap	140	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No
www.stap	0	Avery& Staples&R	200	www.stap	193	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	1-Hole Paj Staples&R	200	www.stap	138	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	1" Non-Vi Staples&R	200	www.stap	199	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	1" View Bi Staples&R	200	www.stap	228	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	100% Recy Staples&R	200	www.stap	135	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	100% Recy Staples&R	200	www.stap	144	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	100% Recy Staples&R	200	www.stap	153	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	100% Recy Staples&R	200	www.stap	129	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	Recycled 1 Staples&R	200	www.stap	141	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	100% Recy Staples&R	200	www.stap	154	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	12" Chrron Staples&R	200	www.stap	129	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	13 Watt Cl Staples&R	200	www.stap	156	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	14" x 4" x Staples&R	200	www.stap	144	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No
www.stap	0	16 oz. Rub Staples&R	200	www.stap	136	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	16" x 4" x Staples&R	200	www.stap	129	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	2" Cardina Staples&R	200	www.stap	163	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	
www.stap	0	2" Cardina Staples&R	200	www.stap	167	text/html	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No	No

Figure 11-40. Moz Crawl Test tool

The tool also provides a detailed look at some of the more important pages on the site, including the indexing status in the major search engines, the most important keywords, metadata, and the number of internal links to the page.

Tools such as these that provide a crawler's-eye view of your web page or website can be effective in helping you to rapidly diagnose problems with your site. Other companies that offer proprietary crawl analysis tools are:

- Rio (<http://www.rioseo.com/>)
- Screaming Frog (<http://www.screamingfrog.co.uk/seo-spider/>)

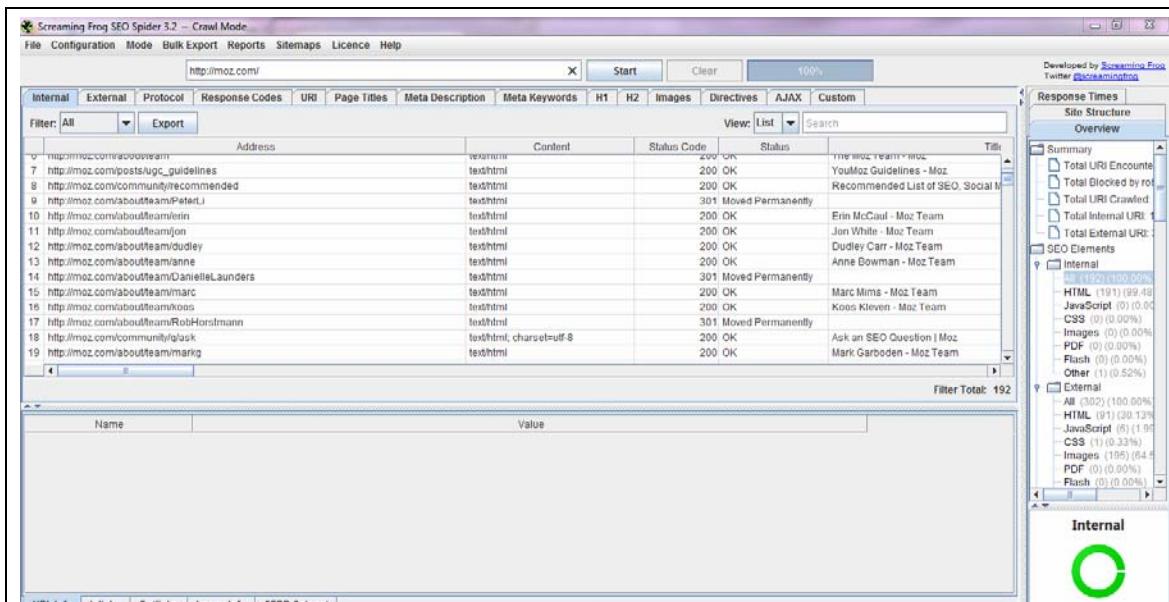


Figure 11-41. Screaming Frog tool

There are free tools that offer a lot of functionality as well. These include:

- Xenu Link Sleuth (<http://home.snafu.de/tilman/xenulink.html>)
- GSite Crawler (<http://gsitecrawler.com/>)

Tracking the Blogosphere

The blogosphere is a very dynamic part of the Web—always fresh and constantly updated. News tends to break in the blogosphere, and in Twitter, before it hits mainstream media or the Web at large. Tracking the blogosphere will help you stay current on both your industry and the SEO industry. Major influencers tend to hang out in the blogosphere as well as have their own blogs. Identifying these influencers is the first step to reaching them.

Here are four major uses for tracking blogs:

Reputation monitoring

A significant component of reputation monitoring relates to the discussion in this section about domain and brand mentions across the Web. Reputation monitoring takes it one step further, where you look to identify problems and risks to your reputation as they materialize on the Web.

For example, knowing when someone starts to write negatively about your brand somewhere on the Web is important. Tracking down those issues early and addressing them quickly is a wise thing to do.

Tracking buzz and public relations campaigns

A closely related activity is tracking buzz and the response to your PR campaigns. When you make a major press push of some sort or you succeed in getting a write-up about your website on an influential blog, you should monitor the ripple effect across the rest of the Web.

Using the type of brand tracking we've already discussed can be a great way to do that. For example, did the number of mentions you received on the Web this week increase significantly over the number of mentions the week before?

Identifying potential influencers

Identifying influencers is a key part of link building. As we've discussed throughout the book, establishing yourself as a recognized expert is an important part of your content strategy. Part of that is creating world-class content, but this won't mean much unless you get the word out somehow. A great way to do that is to reach out and develop relationships with the key influencers.

Competitive analysis

If you can do this type of research for your own site, why not do it for your competitors' sites? You can see reputation problems that are developing for them, or see the impact of their media campaigns as they roll them out.

Tracking Your Blog(s)

Blogs offer numerous benefits to the online marketer. They can offer a great way to position yourself or the team behind your site as experts. For that reason, you can reach major influencers as well as a broad audience of people at the same time. This can

become a nice source of links as a part of your ongoing content and link development efforts.

Blog links

There are a few ways to extract data on the number of links to your posts. You can also use either Google Webmaster Tools or Bing Webmaster Tools to track the number of links to your blog.

As we outlined in the discussion on link tracking, you can use these tools to get link data on your website. You can also take a more granular look at your data. For example, if you have a blog at <http://www.example.com/blog>, you can pull a report from Bing Webmaster Tools (or your favorite tool of choice) and sort your spreadsheet on the URLs of the pages receiving the links. When you are done with this, your data might look something like Table 11-1.

Link Explorer 

Submit a URL to discover the pages that link to the URL.
 EXPLORE History

Filter by site Anchor text Additional query

Scope URL External

EXPORT

SOURCE URL	TITLE
http://www.marketingdegreedtoday.com/	Marketing Degree Today — A free resource for researching ... 
http://pennsylvania.americanentrepreneurship.com/	Pennsylvania - Americanentrepreneurship 
http://blog.echelonseo.com/	Brent Carnduff Blog Echelon SEO Blog 
http://tx.americanentrepreneurship.com/	Texas Venture Capital Startups American ... 
http://moz.com/community/recommended	Recommended List of SEO, Social Media, Content, Hosting ... 
http://wiep.net/	Wiep.net - Link Building Blog 
http://www.forbes.com/sites/ericenge/2014/04/16/why-googles-post-ads-program-may-be-important-for-you/	Forbes: Why Google's +Post ads Program May Be Important ... 
http://www.hyperlinkguerrilla.com/	Hyperlink Guerrilla — Connecting Dots 
http://www.forbes.com/sites/ericenge/2014/04/15/7-modern-age-seo-myths/	7 Modern Age SEO Myths - Forbes - Information for the ... 
http://cyrusshepard.com/20-future-facing-seo-blogs-for-2012/	20 Future-Facing SEO Blogs - Cyrus Shepard 
http://www.toprankblog.com/search-marketing-blogs/	SEO Blogs - BIGLIST - Online Marketing Blog - TopRank® 
http://marketingland.com/3-ways-scale-content-marketing-82642	3 Ways To Scale Your Content Marketing 
http://www.copyblogger.com/google-hummingbird/	Why Google's New Hummingbird Algorithm is Good News for ... 
http://www.copyblogger.com/become-an-expert-now/	21 Reasons You Must Become an Expert - Copyblogger 
https://www.sempo.org/?page=media	Media - SEMPO 
http://www.sempo.org/?page=industry_blogs	Search Marketing Industry Blogs - SEMPO 
http://irlconferences.com/	IRL Conferences Just another WordPress site IRL ... 

Figure 11-42 - inbound blog links crawled from Bing Webmaster Tools

Table 11-1. Sorted link report

8433	http://www.stonetemple.com/articles/video-	http://www.	7-Jul-11
8434	http://www.stonetemple.com/articles/video-	http://www.	12-Jul-11
8435	http://www.stonetemple.com/articles/video-	http://www.	4-Jul-11
8436	http://www.stonetemple.com/articles/video-	http://www.	3-Jul-11
8437	http://www.stonetemple.com/blog	http://www.	9-Jul-11
8438	http://www.stonetemple.com/blog	http://www.	21-Jul-11
8439	http://www.stonetemple.com/blog	http://www.	2-Jul-11
8440	http://www.stonetemple.com/blog	http://www.	1-Jul-11

8441	http://www.stonetemple.com/blog	http://www.	14-Jul-11
8442	http://www.stonetemple.com/blog	http://www.	6-Jul-11
8443	http://www.stonetemple.com/blog	http://www.	10-Jul-11
8444	http://www.stonetemple.com/blog	http://www.	23-Jul-11
8445	http://www.stonetemple.com/blog	http://www.	22-Jul-11
8446	http://www.stonetemple.com/blog	http://www.	25-Jul-11
8447	http://www.stonetemple.com/blog	http://www.	28-Jul-11
8448	http://www.stonetemple.com/blog	http://www.	24-Jul-11
8449	http://www.stonetemple.com/blog	http://www.	20-Jul-11
8450	http://www.stonetemple.com/blog	http://www.	17-Jul-11
8451	http://www.stonetemple.com/blog	http://www.	26-Jul-11
8452	http://www.stonetemple.com/blog	http://www.	25-Jul-11
8453	http://www.stonetemple.com/blog	http://www.	24-Jul-11

You can then find the end of the links that go to the blog and see how many total links you have. So, for example, if the last link is found on line 12,367 of the spreadsheet, the total number of links to the blog is $12,367 - 8,437 = 3,930$. That can get you a raw count of inbound links to the blog. To go further, you can start analyzing what parts of the blog have the most links as well.

GoogleAlerts

A good way to monitor search engines' awareness of your company, brand, and products is to set up alert notifications. Use Google Alerts (<http://www.google.com/alerts>) to get automatic emails when there are new results for the keywords you specify. Note that sometimes this "news" can be outdated and redundant.

In addition to monitoring your own Web presence and reputation, it pays to keep tabs on your competitor's keywords, rankings, links, traffic, and news. Consider setting up alerts for your competitors as well as your own brands. You may even want to sign up for your competitors' email marketing messages so that you can monitor the news and offers they are sending to their customers and prospects.

Search Engine Robot Traffic Analysis

Understanding how robots are spidering your site is another thing that the expert SEO practitioner should know how to do. For one thing, spidering frequency is a clue as to which pages on your site have the highest PageRank and trust, because Google generally crawls the Web in roughly reverse PageRank order. It can also help you detect spidering problems on your site.

You can use the tools we discuss in this section to help you find potential spidering issues, and analyze how important the search engines consider your content to be. You should be looking for clues of SEO problems, such as *robots.txt* blocking the crawlers, architectural problems, or even signs of a penalty (as might be signified by a big drop in crawling frequency).

However, it is important to know that this data will not tell you everything. For example, the fact that a web page was crawled does not mean it will be placed in the index. The page will still need to pass some additional tests to achieve such placement (such as the presence of unique content and enough links to justify its inclusion).

For pages that are indexed, you can look at how often the spiders visit/crawl your pages versus how often the engine actually shows a new version of your page in the index. To do this, look at the last cached date the search engine reported, and compare it with your crawling data.

Tools that perform log analysis include Webtrends, IBM Unica NetInsight, and Lyris HQ Web Analytics.

Other logfile analysis programs include:

- AWStats (<http://www.awstats.org>)
- Sawmill (<http://www.sawmill.net>)
- Webalizer (<http://www.webalizer.org>)
- W3Perl (<http://www.w3perl.com>)
- Piwik (<http://www.piwik.org>)

What is Crawl Efficiency and Why is it Important?

Analyzing search bot behavior involves three main goals:

1. All of your important pages are found by the bots
2. PageRank is not getting diluted by duplicate content problems
3. Time-sensitive content (ads/promotions, breaking news) is indexed as soon as possible

Crawl efficiency is about helping search bots reach these goals. To do this, you must analyze their behavior on your site, understand the obstacles they face, and remove any roadblocks that prevent successful indexation.

Manually removing duplicate content and providing an XML sitemap are an excellent ways to address crawl efficiency problems. If the XML sitemap is comprehensive, and there is enough PageRank flowing through the site, the most important pages will get found. If the XML sitemap includes additional metadata, like how often specific pages are changed, and this metadata is accurate, then your new and time-sensitive content can be discovered quickly.

Addressing duplicate content properly is extremely important, both from the perspective of improving crawl efficiency, and from the perspective of improving the PageRank of canonical pages. As the PageRank consolidates, their rankings will increase, and the result will be more clicks and more sales.

Google Webmaster Tools

You can also get detailed information about spidering activity on your own website using Google Webmaster Tools. Figure 11-43 shows a snapshot for one site.



Figure 11-43. Google Webmaster Tools spider activity report

This provides a great visual snapshot. This report is most valuable when you see major changes in the data. For example, if you change hosting provider, and the time spent downloading a page in milliseconds jumps upwards, this may be a flag of a program with the new hosting company. When you see these types of changes, it can be a flag that something happened with your website, and you should investigate it.

Altruik Product Priority Platform

Altruik offers an SEO Platform (<http://www.altruik.com/product/overview>) that can provide high-level visibility into crawl efficiency and duplicate content issues. In addition, links to broken pages are automatically 301 redirected to the closest matching pages on the site.

Web Traffic Comparison

When you set your goals for growing your online business, you may, for example, target a sales increase of 50% from search referral traffic. However, you may also want to take into account where your competitor is in setting those goals. If your competitor has twice your traffic, you may not be happy with a goal to increase your traffic by 50%.

In addition, when you measure the results of a completed SEO campaign, it is useful to see how much your competitor grew during the same time frame, because you can use this to set your goals for your next campaign. For example, if you grew your traffic by 50% and your competitor grew by 100%, you may want to adjust accordingly.

Several tools out there offer ways to compare the traffic of different websites. However, each tool has accuracy issues and significant bias.

For example, Alexa (<http://www.alexa.com>) relies on data obtained from users who have the Alexa toolbar on their systems. This is a limited number of users, and of course, only certain types of people are willing to have the toolbar on their system. This introduces a bias to the data, and it means it is not completely reliable. In addition, you can get software to artificially inflate your Alexa rankings by spoofing the Alexa toolbar. This is not something you should do, but you should understand the issues with these types of measurement systems.

Similarly, companies such as Compete (<https://www.compete.com>) and Quantcast (<https://www.quantcast.com>) offer traffic measurement services. These are also subject to issues with bias, and limited sample sites. Also note that all of these data sets are limited to the United States.

However, these tools have value when comparing the traffic of one site relative to another. This is great data because if you are comparing two sites in the same market space (perhaps you and your competitor), the bias should affect both sites equally.

Alexa

Alexa provides you with a quick and easy look at how the traffic for two sites compares, as shown in Figure 11-44.

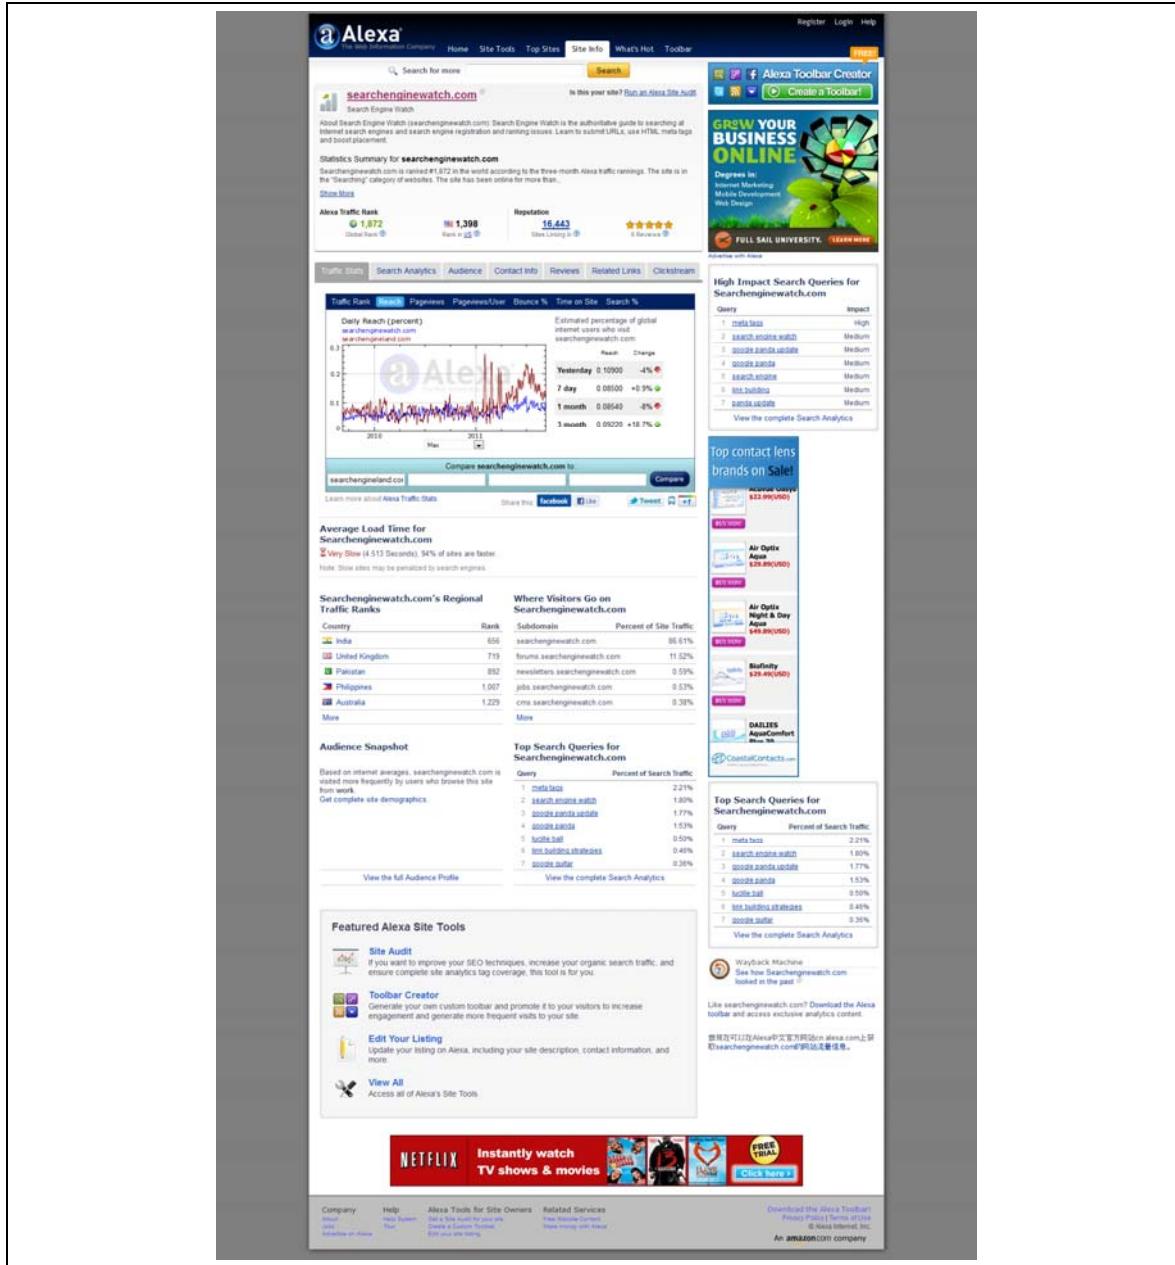


Figure 11-44. Alexa traffic comparison

Underneath the chart you will notice a lot of additional data on SearchEngineWatch.com (<http://searchenginewatch.com>) (the site listed first in the search box), such as the most popular content and where the readers are coming from.

Compete

Compete offers both a free and a paid service. As with Alexa, you can get basic traffic data on a site or compare traffic between sites (see Figure 11-45).

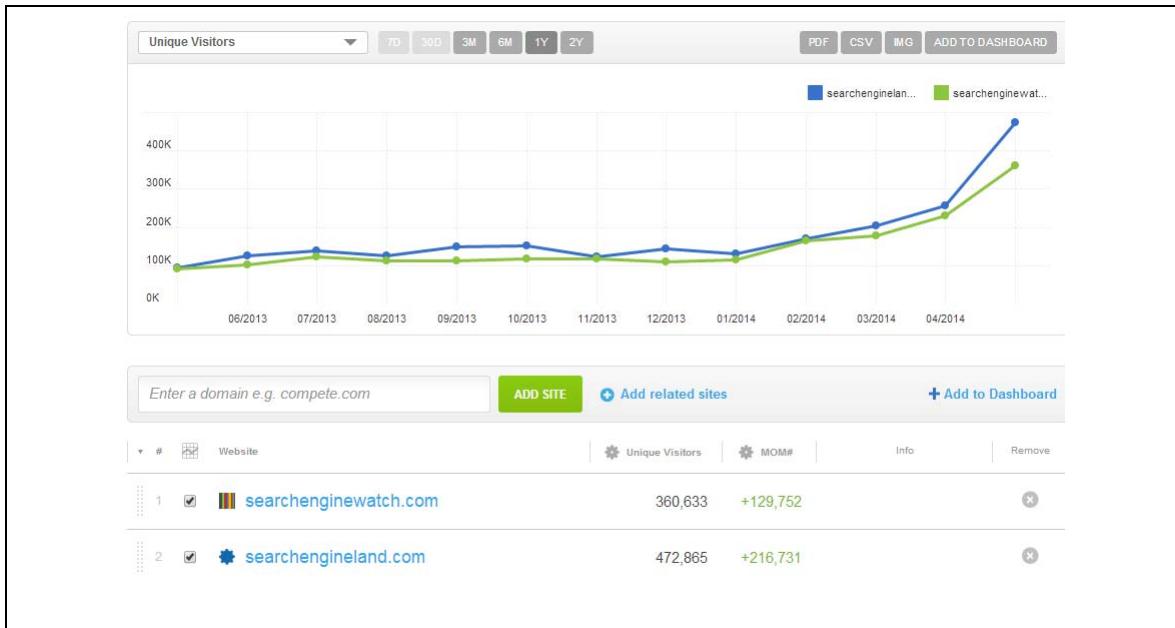


Figure 11-45. Compete traffic comparison

Compete also offers additional data for a fee. This includes data on the major keywords driving traffic to a site.

Compete leverages multiple sources, including ISP, panel, and toolbar data, and has a total panel of about 2 million users. Compete's diverse data sources help in reducing bias in the data.

Quantcast

Quantcast is a competing service to the others we've listed. With Quantcast, you can get detailed data about individual sites.

The demographic data is interesting, as well as the other sites that users visit. Quantcast collects "pixel data" from scripts running on websites that are part of the program. Quantcast claims that these sites represent about 50 million visitors per month. Quantcast supplements this data with data from ISPs. Figure 11-47 shows a sample report for USA.gov.

Quantcast also offers publishers an opportunity to participate in its Quantcast Advertise for Publishers program (<https://www.quantcast.com/advertise/quantcast-advertise-for-publishers/>). This program enables publishers to make public much more accurate data through Quantcast and requires a tracking tag to be placed on their sites. Publishers who sell advertising on a cost per impression basis can use this as third-party validation of their traffic levels.

Temporal Link Growth Measurements

Since Google released its patent application on historical information in March 2005 (read more at <http://moz.com/blog/google-patent-on-historical-data>), search marketers have recognized that trends in temporal link and content analysis do have a real impact on rankings.

The engines are trying to measure patterns—they're looking for indications of increasing or decreasing relevance and authority that temporal trends provide. They want to identify several specific items:

Content growth patterns

How often does a particular site tend to add new pages?

Content update patterns

How often are documents edited and updated?

Link growth patterns

How often do new links appear pointing to the site?

Link stagnation patterns

Does the number of links to the site stagnate or decrease?

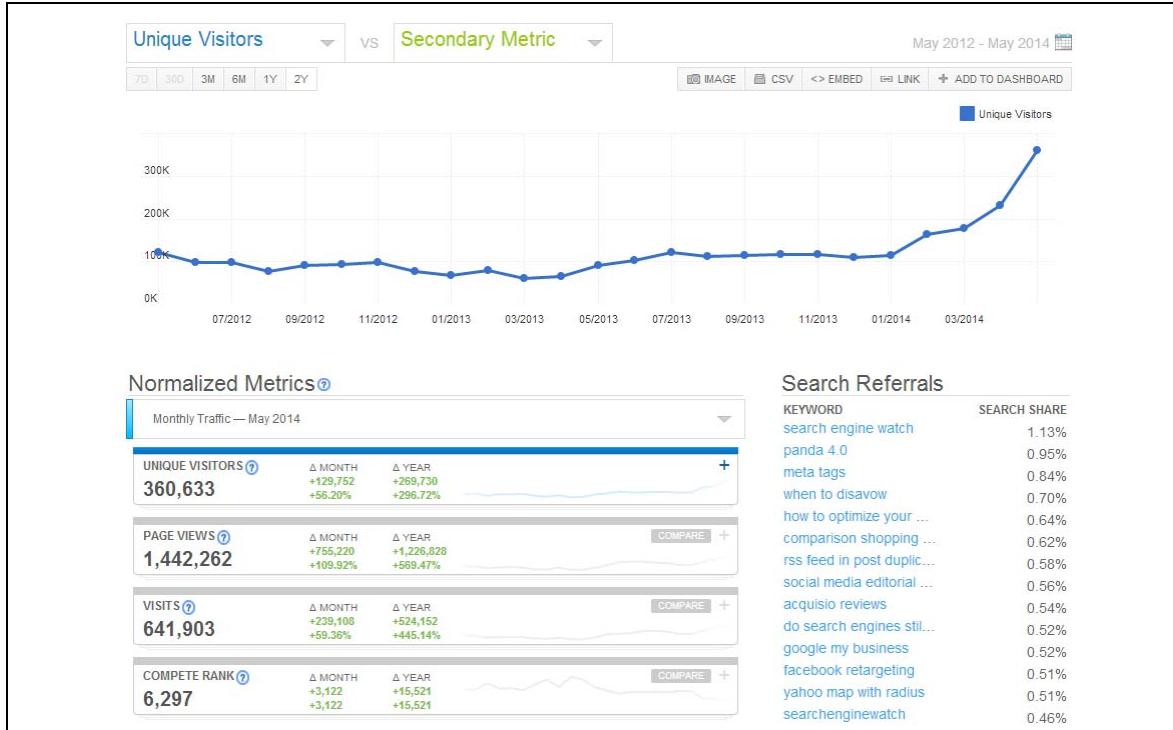


Figure 11-46. Compete keyword data

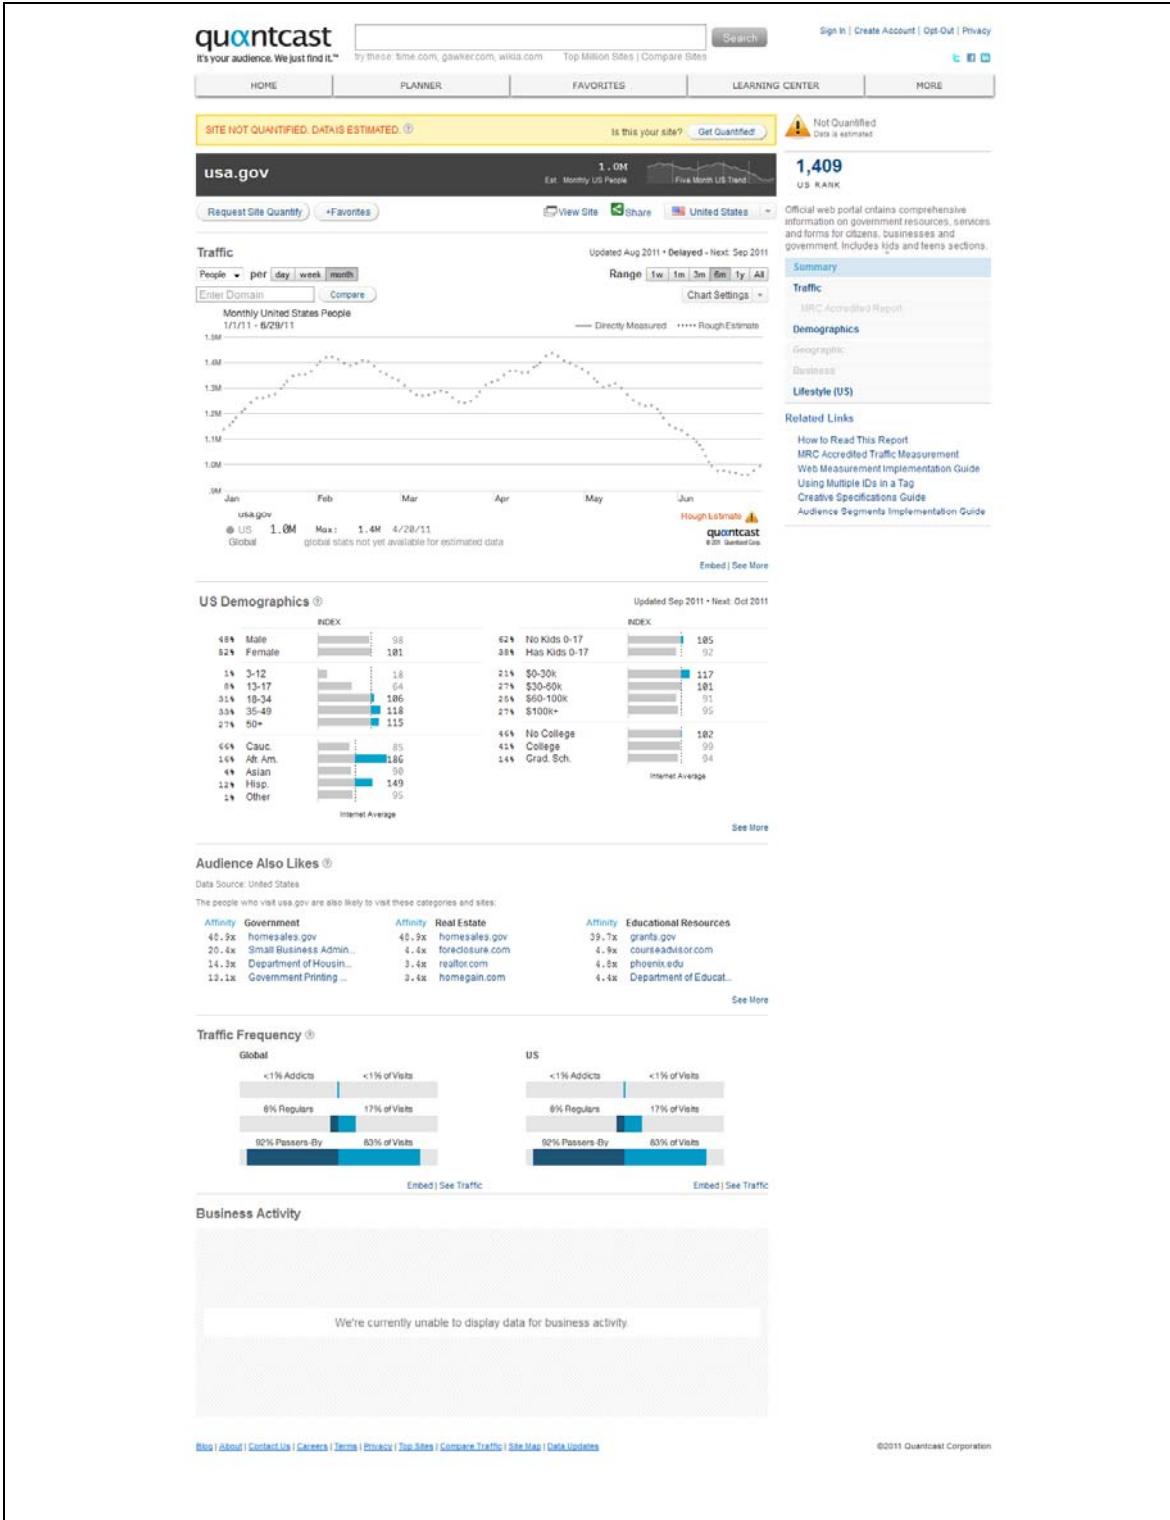


Figure 11-47. Quantcast site data

The engines aren't interested only in how many links pointed to the site today versus yesterday (or how many pages were added); they are also fundamentally interested in tracking patterns over time. Figures 11-

48 and 11-49 depict some example graphs showing the rate of new external links (and in the last two instances, pages) created over time, with some speculation as to what the trends might indicate.

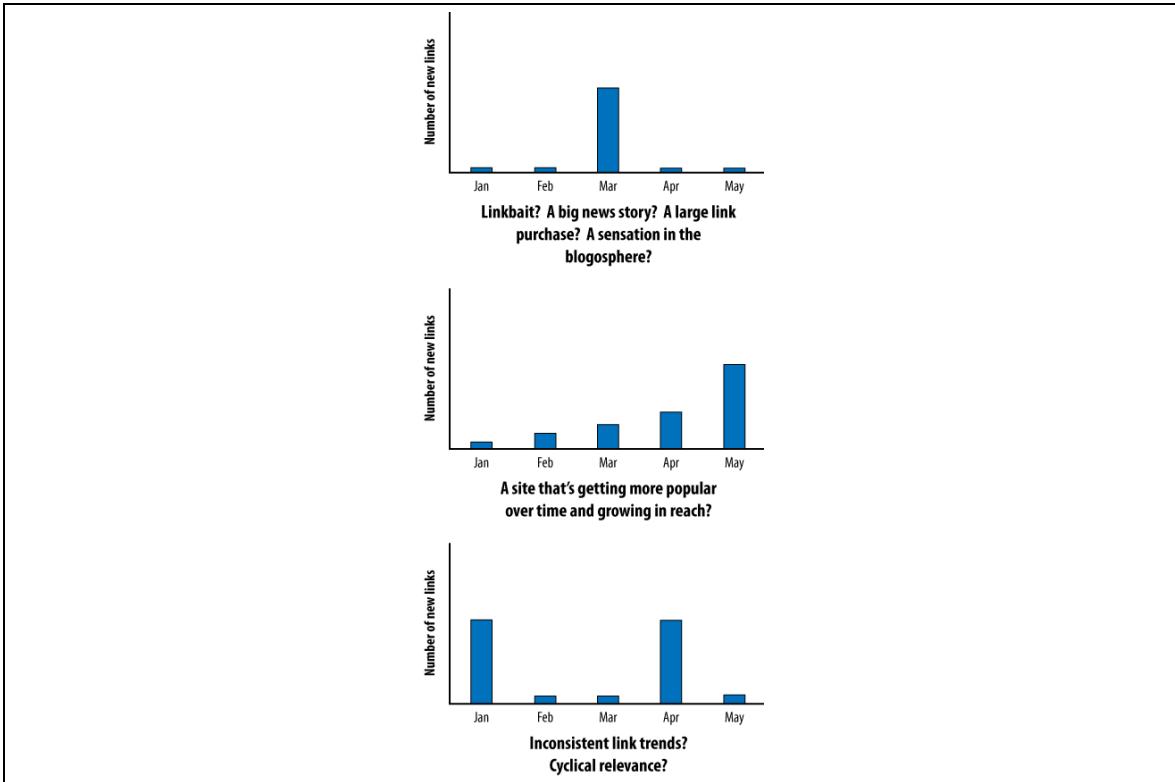


Figure 11-48. Interpreting new external link data

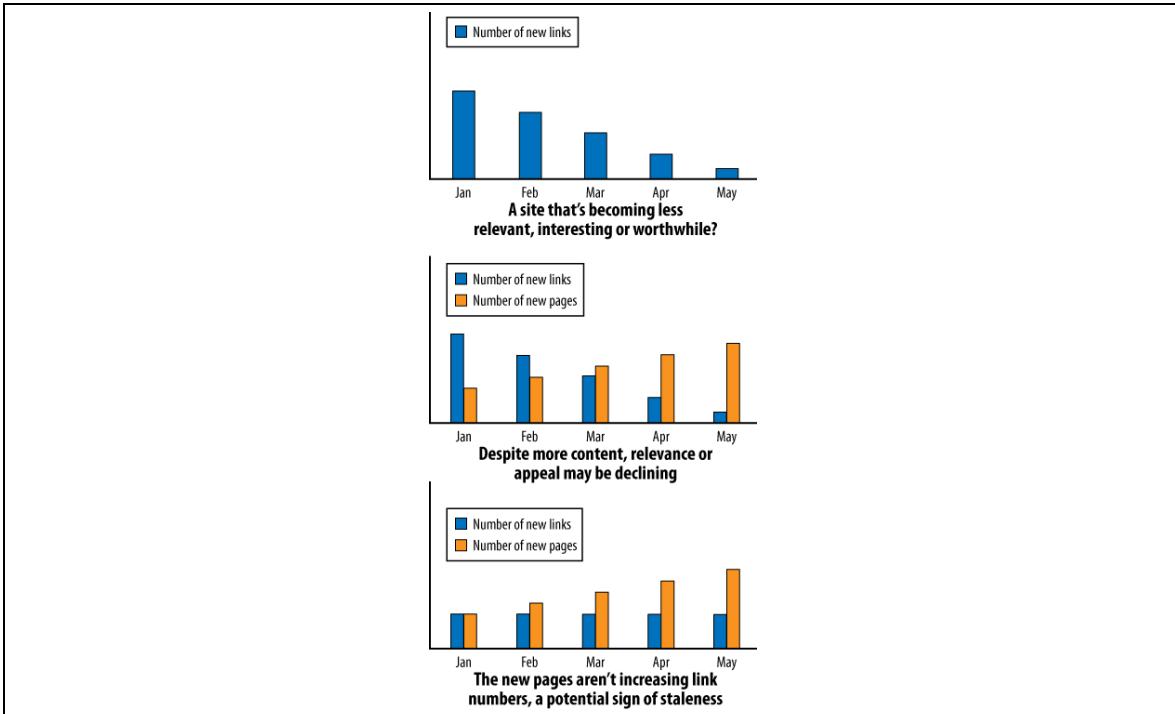


Figure 11-49. More link data speculation

These assumptions do not necessarily hold true for every site or instance, but the graphs make it easy to see how the engines can use temporal link and content growth information to make guesses about the relevance or worthiness of a particular site. Figure 11-50 shows data of a few real sites and how these trends have affected them.

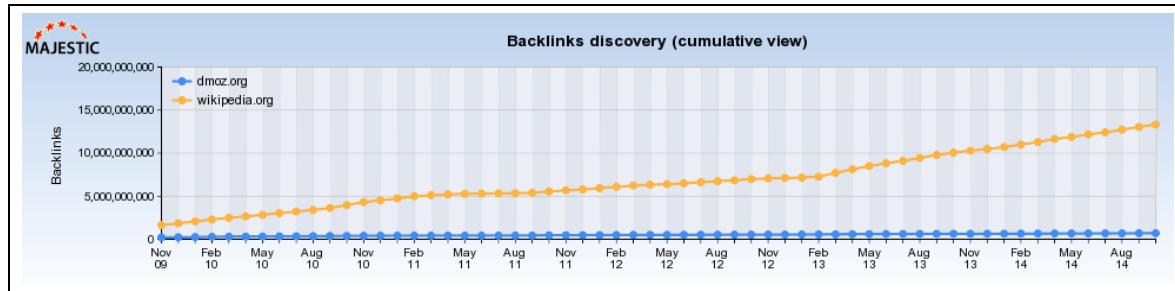


Figure 11-50. DMOZ vs. Wikipedia Backlink discovery

As you can see in Figure 11-50, Wikipedia has had tremendous growth in both pages and links from 2009 through 2014. This success manifests itself in the search engines, which reward Wikipedia's massive link authority with high rankings for much of its content.

DMOZ was once a default reference link for many sites, its relative influence has waned. Although it continues to grow, Wikipedia simply trumps it immensely. Figure 11-51 shows backlink domain discovery that shows that Wikipedia not only has breadth of links, but linking domains.

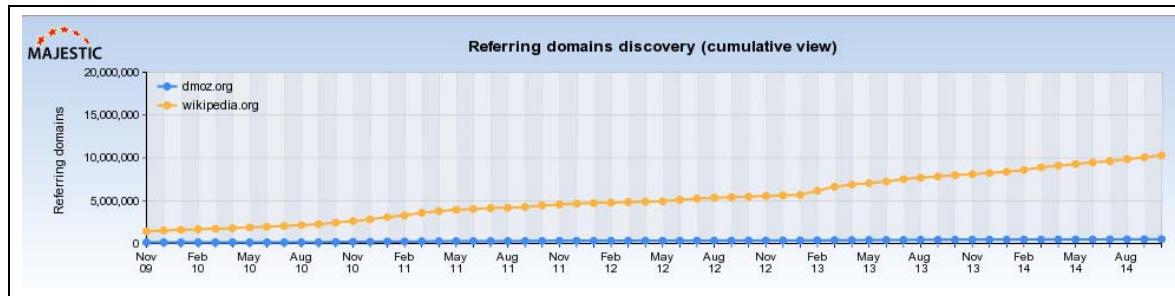


Figure 11-51.DMOZ vs. Wikipedia referring domains

Many forms of spam and manipulative link building are likely to stand out like a sore thumb when put under the temporal microscope. When a large gain in links relative to a site's sphere, influence, and historical link growth appears, the engines can take a closer look at the source of the links or even trigger a manual review. Common sense would dictate that a small-time local real estate site doesn't usually attract a few thousand new links in a week unless it has done something newsworthy or link-worthy.

There are few limits to what the engines can do with data such as this, and there is no reason they shouldn't be analyzing it (since it is easily available). Do consider how the link and content growth patterns for your sites may affect the engines' perspectives on your rankings and trustworthiness.

Key Performance Indicators for Long Tail SEO

As we have discussed throughout the book, the long tail is an important part of SEO. Metrics are available for diagnosing the health of your long tail search traffic.

Brand-to-Non-Brand Ratio

This is the percentage of your natural search traffic that comes from brand keywords versus non-brand keywords.

If the ratio is high and most of your traffic is coming from searches for your brand, this signals that your SEO is fundamentally broken. The lower the ratio, the more of the long tail of natural search you likely are capturing. This metric is an excellent gauge of the success of your optimization initiatives. What constitutes a high ratio is going to vary depending on your brand's level of recognition and the overall level of web savvy of your users. Certainly anything over 50% brand searches bears closer examination.

Unique Crawled URLs

This is the number of unique (non-duplicate) web pages crawled by search engine spiders such as Googlebot and BingBot. Your website is like your virtual sales force, bringing in prospects from the search engines. Think of each unique page as one of your virtual salespeople. The more unique pages you have, the more opportunities you have to sell through the search engines.

Search Visitors per Contributing Page

This is the percentage of unique pages that yield search-delivered traffic in a given month.

This ratio is a key driver of the length of your long tail of natural search. The more pages yielding traffic from search engines, the healthier your SEO program is. If only a small portion of your website is delivering searchers to your door, most of your pages—your virtual salespeople—are warming the bench instead of working hard for you. You can think of these non-performing pages as “freeloaders.”

Keywords per Page

This is the average number of keywords driving traffic to each page (minus the freeloaders) in a given month. Put another way, it is the ratio of the total number of keywords driving search traffic to pages receiving search traffic.

The higher your keywords per page, the more of the long tail of natural search your site will capture. In other words, the more keywords each page attracts or targets, the longer your tail. So, an average of eight search terms per page indicates pages with much broader appeal to the engines than, say, three search terms per page. The average online retailer in a Netconcepts study on the long tail of natural search (<http://www.netconcepts.com/learn/ChasingTheLongTail.pdf>) had 2.4 keywords per page. A larger number of keywords per page should also indicate that your content is more varied in its keyword usage and that it is better meeting the needs of a wide variety of users (that the search engines understand).

Search Visitors per Keyword

This is the ratio of search-engine-delivered visitors to search terms. This metric indicates how much traffic each keyword drives and is a function of your rankings in the SERPs. Put another way, this metric determines the height or thickness of your long tail.

The average merchant in the aforementioned Netconcepts study obtained 1.9 visitors per keyword.

Index-to-Crawl Ratio

This is the ratio of pages indexed to unique crawled pages. If a page gets crawled by Googlebot, that doesn't guarantee it will show up in Google's index. A low ratio can mean your site doesn't carry much weight in Google's eyes.

Search Visitors per Crawled Page

Calculated for each search engine separately, this is how much traffic the engine delivers for every page it crawls.

Each search engine has a different audience size. This metric helps you fairly compare the referral traffic you get from each. The Netconcepts study found that Bing and Yahoo! tended to crawl significantly more pages, but the yield per crawled page from Google was typically significantly higher.

As you optimize your site through multiple iterations, watch the aforementioned KPIs to ensure that you're heading in the right direction. Those who are not privy to these metrics will have a much harder time capturing the long tail of SEO.

Tracking Duplicate Content

Duplicate content – one or more pages with multiple unique URLs – is, in effect, making your own long tail unnecessarily longer. You only want one canonical URL for each page of content, and you don't want that content to be repeated verbatim on other pages or sites.

Once you find duplicate content, figure out which page and URL you want to be the canonical resource (typically the ones with the highest PageRank), and 301-redirect the others to it.

Some analytics packages and independent SEO tools can track duplicate content, such as SEObook.com's Duplicate Content Checker: (<http://training.seobook.com/duplicate-content-checker>).

Other Third-Party Tools

Throughout this chapter, several third-party tools have been mentioned or recommended for analysis, research, and testing. Here are a few more worth mentioning:

MozBar

As previously mentioned in this chapter, the MozBar provides at a glance link and social sharing metrics but it also offers much more. Through a series of available drop-down menus the bar can provide information on on-page elements, general attributes, in-depth link metrics, semantic markup, and HTTP status of the page. Some of the on-page elements it details are the URL, title tag, meta-description, meta-keywords, heading tags, bolded text, italicized text, and alt text for images on the page. Some of the “general attributes” it details are the content of the meta-robots tag, the rel=canonical link attribute, the page load time, a link to the URL for Google's cache of the page, the IP address, and the country the site is perceived to be located in. It offers a wide variety of link metrics for the current URL, the current subdomain, and the current root domain including Page Authority, Domain Authority, External Followed Links, Linking Root Domains, MozRank (mR), and MozTrust (mT). The tool also shows whether various types of markup are present on the page and provides links to further analyze that markup. Some of the types of markup it can detect are: schema.org, open graph protocol, twitter cards, micro formats, and authorship. Also offered are links to more of Moz's SEO tools.

Figure 11-52. MozBar drop-down menus

SEO Quake

SEO quake (<http://www.seoquake.com/>) offers both a Firefox plug-in and a Chrome extension. SEO Quake both adds a toolbar to your browser for at a glance statistics and offers links to more in-depth page information. Some of the metrics this add-on allows you to view are: PageRank, number of pages indexed in Google, SEM Rush links to both the page and domain, number of pages indexed in Bing, Alexa rank, domain age, social shares, and links to WHOIS information, the page's source code, and the SEM Rush domain rank.

Figure 11-53. SEO Quake page info report

SEO for Firefox

Aaron Wall of SEO Book created a useful data enhancement tool with the SEO For Firefox extension (<http://tools.seobook.com/firefox/seo-for-firefox.html>). When turned on, this tool provides extra data that can be populated automatically or on request from within the regular search results of Google, Yahoo and Bing. Some of the most popular data points this add-on provides for the URL or domain of each search result includes: Google PageRank, site age, various back link metrics, .edu and .gov link counts, bookmarking and social metrics (like number of YouTube views), DMOZ listings and additional reporting from other third party tools.

SpyFu

SpyFu (<http://www.spyfu.com>) offers valuable keyword data for SEO professionals as well as resources for PPC advertisers. There is both a free version that provides valuable but limited data, as well as a paid version that provides additional data. For SEO, SpyFu offers Domain Research, Keyword Research, and Backlink Discovery.

Domain Research allows you to review a site's organic keywords by their latest changes, by their page, and by clicks. This allows you to monitor for rank, scout for potential algorithm penalties and manual penalties, and even comparing these rankings to competitors.

Keyword Research instead, as expected, focuses on the stats of an individual keyword. The tool provides a history of domains that have ranked for a specific keyword, the urls of pages ranking for a keyword, and keyword mapping ideas. These results can be incredibly helpful for developing a content strategy.

The Backlink Discovery tool doesn't work like other backlinking tools that show current backlinks to the site. Instead, it suggests potential backlinking opportunities, such as hubs that are linking to various competitors or other sites within the niche. Getting deeper into it, you can even organize results by keyword, and single out backlinks that would be helpful to rank for a specific keyword.

For PPC, SpyFu can provide information on a site's estimated daily ad budget, total ad clicks per day, average cost per click, average ad clicks per day, average ad position, average ad competitors, average ad percentile, paid and organic keywords, top paid and organic competitors, and other domains that share the same owner.

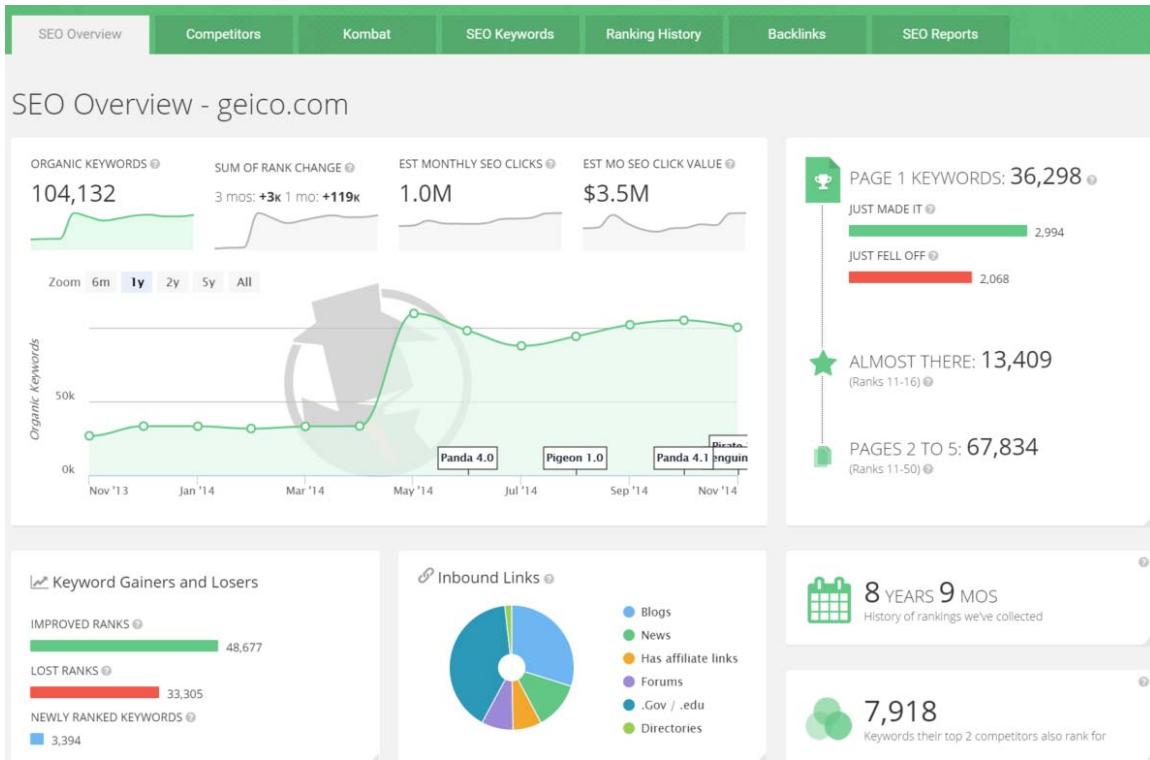


Fig 11-54. SpyFu Organic Domain Research for SEO

SEMRush

SEMRush (<http://www.semrush.com>) is another tool that offers keyword information of use to SEO and paid search professionals. The report shown in Figure 11-55 shows a paid search report that is interesting from an SEO perspective because it shows the target landing page for the keyword, as well as the estimated CPC. This can provide insight into the value that competition is placing on various keywords and landing pages on their site.



Figure 11-55. SEMrush keyword report

Rio SEO Search Analytics

<http://www.rioseo.com/seo-software-platform/>

Rio SEO Search Analytics is an enterprise class SEO platform that includes many more features than the simple rank checking, crawl error alerts and simple recommendations.

Some of the features include:

- Keyword research solutions
- Long tail analysis and recommendations
- Workflow tracking and automation
- Portfolio analysis for multiple sites, microsites and landing pages
- Prioritization mechanisms
- Competitive analysis features
- Advanced link analytics
- Ad hoc analysis capabilities
- Advanced reporting capabilities
- Alerting capabilities
- Universal search metrics (image, video, etc)
- Social signal inclusion and analytics

Rio SEO Website Optimizer

<http://www.rioseo.com/website-seo-optimization-software/>

In some enterprise environments there are fundamental challenges with the Content Management Systems (CMS), or the e-commerce platform that makes it very difficult to implement desired SEO changes.

There are only a few solutions in this space, and one of the best known is Rio SEO's Website Optimizer (formerly Covario's Organic Search Optimizer), originally developed by Stephan Spencer (one of the co-authors of this book) at Netconcepts and previously known as GravityStream.

Rio's Website Optimizer takes a complementary approach to whatever system is currently in place and helps SEOs get around IT to dynamically fix the SEO errors. Website Optimizer enables the modification of the pages without having to directly deal with the CMS or e-commerce platform.

Searchmetrics Essentials

Searchmetrics Essentials (<http://www.searchmetrics.com/en/>) allows you to rapidly see how your competition is doing in organic search, paid search, news search, image search video search, and more. The keywords tab gives you search volume, cost per click, competition for keywords, and the average ad spend on a per-keyword basis. This can be more helpful and informative than simply using the Google Keyword Planner. You can actually see what people are spending to target specific keywords.

Searchmetrics Essentials also lets you peer deep into your competition's social media campaign. You can see which pages are the most liked, commented on, and shared throughout all of the popular networks. The platform offers many other features.

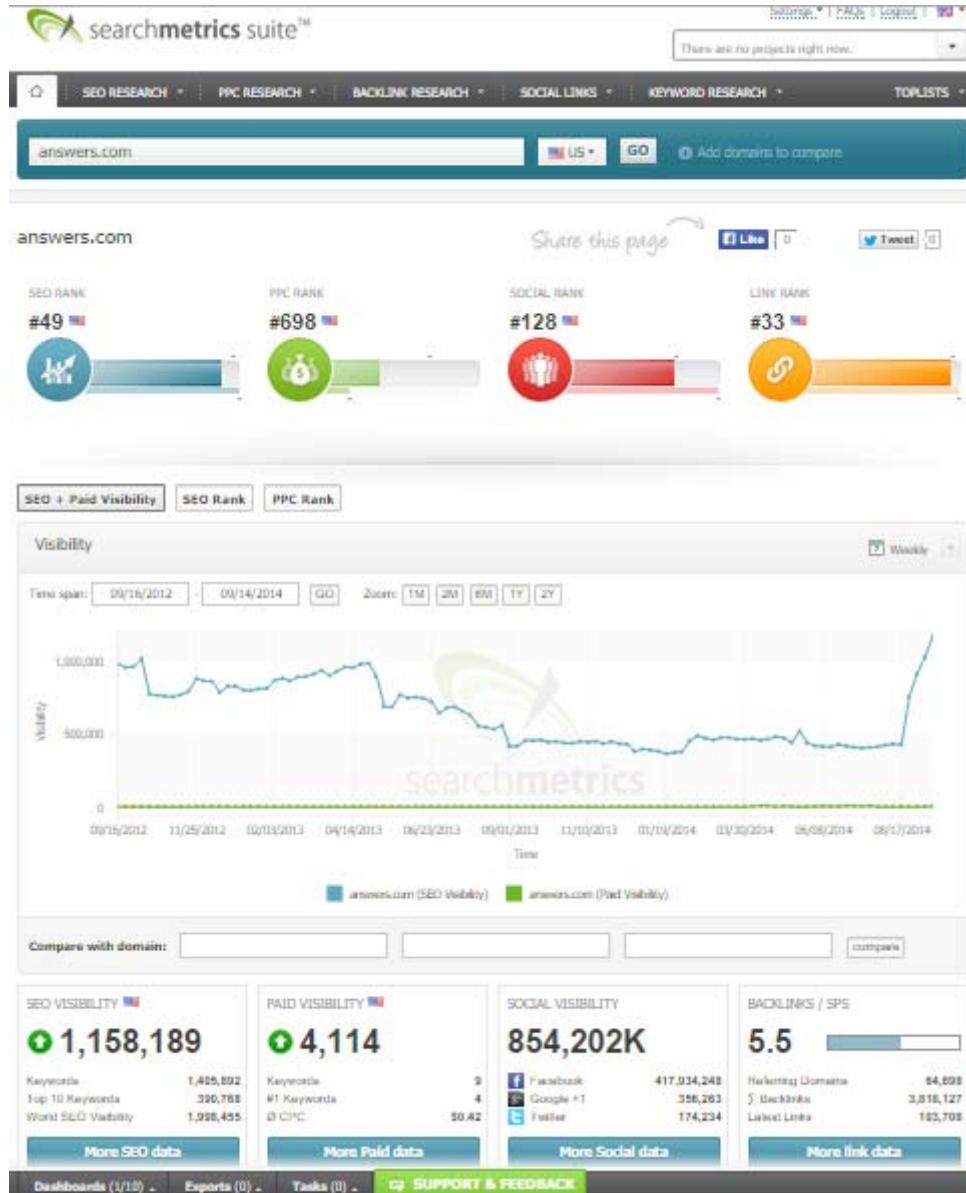


Figure 11-56 SearchMetrics Essentials

Conclusion

One of the biggest challenges for an SEO professional is the relatively long delay between implementation and results. Paid search campaigns can begin to show results in 24 hours, yet for SEO the process can take months. This makes effective measurement of results and progress critical. Putting the right analytics tools in place can provide the accountability required by budget managers. For the SEO practitioner, it can also provide valuable insight that can be used to fine tune your SEO efforts for the best possible results.

12

Domain Changes, Redesigns, and Troubleshooting

Whenever you make structural changes to your website, there is a risk that you will confuse the search engines and harm your search rankings and organic traffic. The types of site changes that can influence your site's organic search performance can include changing your domain name, changing your content management system (CMS), redesigning your site, adding e-commerce functionality, changing your blog platform, and many others—basically, anything that fundamentally alters your site's front and/or back end visual or functional elements, can potentially influence your organic search performance.

In this chapter, we will review the various scenarios from the top down. sure to refer back to “Content Delivery and Search Spider Control” in Chapter 6 to review the technical specifics of options for moving content from one location to another.

Google specifically categorizes site moves as falling into one of two categories, moves “with” URL changes, and moves “without” URL changes: <http://googlewebmastercentral.blogspot.com/2014/06/making-site-moves-easier.html>, and there are specific guidelines that Google provides for handling moves within both scenarios: https://support.google.com/webmasters/topic/6029673?hl=en&ref_topic=6001951.

The Basics of Moving Content

Moving content refers to a situation in which content that used to be located and accessed at one URL (e.g., <http://www.yourdomain.com/pageA>) is moved to another URL (e.g., <http://www.yourdomain.com/products/pageA>). One of your goals when you move content is to make sure users and search engine crawlers that attempt to visit the old URL (*/pageA*) are presented with the content from the new location (*/products/pageA*). In addition, when you move content from one URL to another, the links to the old URL will stop providing value to your rankings in the search engines for that content unless you properly implement a 301 (permanently moved) redirect. There are very important reasons to move content properly—both of which can be easily overlooked by inexperienced or hurried webmasters and development teams.

In “Duplicate Content Issues” in Chapter 6, we covered the technical specifics of how to do this in detail, including the golden rule of moving content: the search engine needs to see a 301 HTTP status code whenever you redirect the content’s URL to a new location.

The 301 HTTP status code causes the search engine to pass most of the value of any links the original page has over to the new page, and should result in the rapid de-indexation of the old URL. Since link authority is a precious asset, you want to make sure you use a 301 redirect every time.

Large-Scale Content Moves

Setting up the redirects can become difficult when changes result in movement of large quantities of content. For example, when you change your domain name, every single piece of content on your site will move to a new URL, even if the site architecture is identical (<http://www.olddomain.com/>... moves to <http://www.newdomain.com/>...).

This is challenging because you might have to set up individual 301 redirects for every single page on the site, as in this example:

```
http://www.olddomain.com/page1 to http://www.newdomain.com/page1
http://www.olddomain.com/page2 to http://www.newdomain.com/page2
http://www.olddomain.com/page3 to http://www.newdomain.com/page3
...
...
http://www.olddomain.com/page1000 to http://www.newdomain.com/page1000
```

Unfortunately, some systems still require that these redirects be set up one at a time, so this could be quite a painful process. Imagine a site with 1 million pages!

Fortunately, publishers who use an Apache (<http://www.apache.org>) web server (Unix and Linux servers) can take advantage of the power of Apache's mod_rewrite module (http://httpd.apache.org/docs/current/mod/mod_rewrite), which can perform the redirect of every URL on the old domain to the same URL on the new domain in two lines of code:

```
| RewriteCond %{HTTP_HOST} ^olddomain\.com [NC]
| RewriteRule ^/(.*) http://www.newdomain.com/$1 [R=301,L]
```

The preceding code presumes that you prefer the "www" version as the canonical URL. You can also use two similar lines of code to specify the "non-www" version as the canonical URL (see <http://searchengineland.com/url-rewrites-and-redirects-part1-16574>) for examples without "www" and other alternative approaches.

Another highly popular web server is IIS (<http://www.iis.net>) from Microsoft. In many installations of IIS, you will find yourself in a situation where you have to implement a separate redirect instruction for each page, one at a time – but usually, webmasters can utilize an ISAPI plug-in such as ISAPI_Rewrite (<http://www.isapirewrite.com>) which enables you to perform large, scalable rewrites in a language similar to that used by Apache's mod_rewrite. You can learn more about mod_rewrite, ISAPI_Rewrite, and regular expressions in Chapter 6.

Mapping Content Moves

Sometimes a site redesign is simply a "re-skinning" of the visual elements of the old site with a new look and feel, retaining the same technical elements of information architecture, URL file and directory names, and user navigation. Other times, a redesign changes both the visual design and the technical elements. And sometimes it's a combination of the two. For sites changing both design and function, the first stage of planning is to figure out which content will be moved where and which content will be removed altogether. You will need this to tell you which URLs you will need to redirect and to which new locations.

The best way to start this process is by getting a complete map of your information architecture with full URLs. For many websites this is not as simple as it sounds. Fortunately, tools are available to make the job easier. Here are some ways to tackle this problem:

- Extract a list of URLs from your web server's logfiles and site architecture documentation.
- Pull the list from your XML Sitemap file, provided you believe it is reasonably complete.
- Use a free crawling tool, such as Screaming Frog SEO Spider Tool (<http://www.screamingfrog.co.uk/seo-spider/>) .
- Use tools such as Moz.com's Open Site Explorer, and Google Webmaster Tools, to pull a list of the external links to your site, and make sure all pages that have received links on your site are included.
- Check Bing Webmaster Tools' Index Explorer to find all of the crawl-able URLs that you may not know still exist on the site.

These tools should help you assemble a decent list of all your URLs. After determining which URLs have content that will remain on the site, you must then map out the pages that you want to redirect the “migrating” content to. Additionally, for content that is being “retired,” you need to determine whether to redirect them at all (a definite yes if the URLs for these pages have many internal and external links), and if so, what new URLs to redirect them to.

One way to do this is to lay it out in a spreadsheet, which might end up looking like Table 12-1.

Table 12-1. Planning your content moves in advance

Old URL	New URL
http://www.olddomain.com/page1	http://www.newdomain.com/page1
http://www.olddomain.com/page2	http://www.newdomain.com/page2
http://www.olddomain.com/page3	http://www.newdomain.com/page3
http://www.olddomain.com/page4	http://www.newdomain.com/page4
http://www.olddomain.com/page5	http://www.newdomain.com/page5
http://www.olddomain.com/page6	http://www.newdomain.com/page6
http://www.olddomain.com/page7	http://www.newdomain.com/page7
http://www.olddomain.com/page8	http://www.newdomain.com/page8
http://www.olddomain.com/page9	http://www.newdomain.com/page9
http://www.olddomain.com/page10	http://www.newdomain.com/page10

If you are redirecting a massive number of URLs, you should look for ways to simplify this process, such as writing rules that communicate what you need to know. For example, you could abbreviate the list in Table 12-1 to the short list in Table 12-2.

Table 12-2. Simplifying content move planning with wildcards

Old URLs	New URLs
http://www.olddomain.com/page*	http://www.newdomain.com/page*

Then you can save the individual lines for the more complicated moves, so your resulting spreadsheet would look like Table 12-3.

Table 12-3. Mapping all your content moves completely

Individual page moves:

Old URL	New URL
http://www.olddomain.com/about-us	http://www.newdomain.com/about-us
http://www.olddomain.com/contact-us	http://www.newdomain.com/contact-us
http://www.olddomain.com/press-relations	http://www.newdomain.com/press

Large-scale page moves:

Old URL	New URL
http://www.olddomain.com/content/*	http://www.newdomain.com/content/*
http://www.olddomain.com/page*	http://www.newdomain.com/page*

The purpose of this is to efficiently give your developers a map for how the content movement should take place. Note that the spreadsheet should contain a map of all changed URLs, which may include downloadable content such as PDF files, PowerPoint presentations, Flash files, multimedia, or any other types of content being moved.

You can also note retiring content via additional entries in the left column, with the entries in the right column indicating where users looking for that old content should be sent. Now your spreadsheet might look like Table 12-4.

Table 12-4. Identifying pages that have been removed

Individual page moves:

Old URL	New URL
http://www.olddomain.com/about-us	http://www.newdomain.com/about-us
http://www.olddomain.com/contact-us	http://www.newdomain.com/contact-us
http://www.olddomain.com/press-relations	http://www.newdomain.com/press

Large-scale page moves:

Old URL	New URL
http://www.olddomain.com/content/*	http://www.newdomain.com/content/*
http://www.olddomain.com/page*	http://www.newdomain.com/page*

Eliminated pages:

Old URL	Redirect to
http://www.olddomain.com/widgets/azure	http://www.newdomain.com/widgets/blue
http://www.olddomain.com/widgets/teal	http://www.newdomain.com/widgets/green
http://www.olddomain.com/widgets/puce	http://www.newdomain.com/widgets/

The new entries show what should happen to retired pages. The first two retired pages may represent products that you no longer carry, so you would likely want to redirect them to the closest existing products you have. The third retired page represents a URL where there is no sufficient replacement, so you may choose to redirect that one to the parent page for that topic area.

As you can see, a major SEO objective during content migration is to preserve as much link authority and traffic from the old URLs as possible, while providing the best possible user experience for people who arrive at the old URLs.

Expectations for Content Moves

The big downside to content migration is that the search engines won't necessarily adapt to the URL changes immediately. Many sites temporarily lose rankings after making a large-scale content move, then recover after a period of time. So naturally, the question is, how long will it take to get your organic rankings and traffic back?

The reality is that a number of factors are involved, depending on your particular situation. Some examples of these factors might include:

The size and complexity of your site

Bigger, more complex sites may take longer to process.

The complexity of the move

If the site has been fundamentally restructured, it is likely to take more time for the search engines to adapt to the new structure.

The perceived authority of the site

Sites that have a higher (search engine) perceived authority may be processed faster. Related to this is the rate at which the site is typically crawled.

The addition of new links to the new pages

Obtaining new links to the new URLs, or changing old links that used to point to the old URLs so that they point to the new URLs, can help speed up the process.

If you are moving to an entirely new domain, you can aid the process in Google by using the Change of Address tool inside Google Webmaster Tools. Before using this tool, make sure that both your old domain and your new domain are verified in Webmaster Tools. Then, on the Webmaster Tools home page, click on the old domain. Under Site Configuration, click “Change of Address,” and then select the new site. You can also use the Bing Change of Address Tool (<http://www.bing.com/webmaster/help/how-to-use-the-site-move-tool-bb8f5112>) to let Bing know about your site move as well.

When all is said and done, a reasonable estimate would suggest that a significant traffic dip from the search engines after a move should rarely last longer than 60 to 90 days, and many are picked up in a shorter timespan.

Another approach to content-moves (especially when updating and redirecting an entire site’s URLs) is to perform the URL migration in a phased manner as opposed to “wiping out” and redirecting all of a site’s URLs at once. This is done for a handful of reasons, such as wanting to test the search engines’ handling of such a migration on your specific site before committing to the site-wide impact it will have. Another reason is to mitigate potential organic traffic dips that will occur during the updating period – it is often easier to tolerate 10% traffic loss at a time, versus a 30% to 40% traffic loss all at once. This can be especially true for websites reliant upon traffic-based advertising revenue. Google’s own Matt Cutts is quoted referencing this approach later in this Chapter.

Maintaining Search Engine Visibility During and After a Site Redesign

Companies may decide to launch a site redesign as part of a rebranding of their business, a shift in their product lines, or a marketing makeover, or for a variety of other business reasons. As discussed, any number of things may change during a site redesign. For example:

- Content may move to new URLs.
- Content might be eliminated.
- Content may be changed.
- Content could be moved behind a login or paywall.
- New site sections may be added.
- New site functionality may be added.
- Navigation/internal linking structure may be changed significantly.

Of course, the move may involve moving everything to a new domain as well, but we will cover that in the next section, “Maintaining Search Engine Visibility During and After Domain Name Changes.” Here are some best practices for handling a site redesign that involves these technical elements:

- Create 301 redirects for all URLs from the original version of the site pointing to the new URLs on the redesigned site. This should cover scenarios such as any remapping of locations of content and any

content that has been eliminated. Use a spreadsheet similar to the one we outlined at the beginning of this chapter to map out the moves to make sure you cover all of them.

- Review your analytics for the top 100 or so domains sending traffic to the moved and/or eliminated pages and contact as many of these webmasters as possible about changing their links. This can help the search engines understand the new layout of your site more rapidly and provides both better branding, and a better user experience.
- Review a backlink report (using your favorite backlink analysis tool) for your site and repeat the process in the preceding bulleted item with the top 200 to 300 or so results returned. Consider using more advanced tools, such as Open Site Explorer (<http://www.opensiteexplorer.org>) or Majestic SEO (<http://www.majesticseo.com>), which allow you to filter your links to more easily identify the most important ones. Ahrefs or LinkResearchTools too.
- Make sure you update your Sitemap (<http://www.xml-sitemaps.com>) and submit it to Google Webmaster Tools (<https://www.google.com/webmasters/>) and to Bing Webmaster Tools (<http://www.bing.com/toolbox/webmaster>) Consider using multiple sitemaps, one for each content type and/or content area, to submit and monitor the indexing of your new site URLs.
- Monitor your rankings for the content, comparing old to new over time—if the rankings fall, post in the Google Webmaster Central Help Forum (<https://productforums.google.com/forum/?hl=en#!forum/webmasters>) information regarding what you did, what happened, and any information that might help someone help you. Google employees do monitor these forums and sometimes do comment in situations where they think they can help. Don't use the forums to complain; state what has happened and ask for help, as this gives you the best chance of getting feedback.
- Monitor your Webmaster Tools account and your analytics for 404 errors and to see how well Google is handling your 301s. When you see 404 errors occurring, make sure you have a properly implemented 301 redirect in place Don't limit this checking just to 404 errors. Also be on the lookout for HTTP status codes such as 500, 302, and others. Maintain the XMLsitemap of old URLs until search engines discover the 301s.

Maintaining Search Engine Visibility During and After Domain Name Changes

There comes a time when there is a strong business need to change the site's domain name. This can be done for various reasons, including re-branding, re-naming, mergers and acquisitions, and other business needs.

Unique Challenges of Domain Name Changes

One of the more challenging aspects of a domain name change is the potential loss of trust the search engines have associated with the old domain. Another issue is that if there were business-specific keywords present in the old domain name that are not in your new domain name, you may see a decline in organic traffic even if you maintain or recover placement after migration as a result of “domain bias” – the propensity for searchers to click on domains in search results that include keywords they used in their search query (<http://research.microsoft.com/pubs/155941/domainbias.pdf>). You may also see a slightly negative impact in organic for search terms related to the keywords in your previous domain, as although Google is cracking down on exact match domain (EMD) websites with low quality content that were ranking well in search, weight is still placed on the words in a domain.

Another unique challenge is the potential “youth” of the new domain, especially if it was recently purchased and/or has no historical backlink profile. Due to this lack of aging, the new domain may be slow to rank. Although the site's relevance and inbound link profile (including the links 301-redirected from the old domain) may suggest a high ranking for some search queries, because the new domain is not trusted yet

the rankings are suppressed, and traffic is much lower than it would otherwise be. Domain youth is another reason why updating valuable third-party links to reflect your new domain is an important task.

If the prospect of taking a “young domain” hit is too unappealing, another tactic you can try is to make the move to a different domain that has a backlink history associated with it as well—just make sure the history is a positive one! You don’t want to move to an old domain that had any historical spam, manual actions, or other negative associations, so be sure to perform a thorough backlink audit with your preferred link auditing tools. And while you’re at it, see if you can get Google Webmaster Tools access to research whether there were any manual spam actions reported against the domain.

Pre-move Preparations

Unfortunately, lost traffic is common when you make domain name changes, if only temporary. But if you do things properly, you can and should recover any negative impact, and hopefully quickly. But you should ultimately be prepared for the potential traffic impact of a domain switch.

If you are planning a domain migration, buy it as early as you can, get some initial content on it, and acquire some links. The purpose of this exercise is to get the domain indexed and recognized by the engines ahead of time.

Then, register the new domain with Google and Bing Webmaster Tools. This is just another part of making sure Google and Bing know about your new domain as early as possible and in as many ways as possible.

Once this is done, here are the best practices for handling a domain name change:

- Create 301 redirects for all URLs from the old site pointing to the proper URLs on the new site. Hopefully you will be able to use mod_rewrite or ISAPI_Rewrite to handle the bulk of the work. Use individual rewrite rules to cover any exceptions. Have this in place at launch.
- Review your analytics for the top 100 or so domains sending traffic to the old pages, and contact as many of these webmasters as possible about changing their links.
- Make sure that both the old site and the new site have been verified and have Sitemaps submitted at Google and Bing Webmaster Tools.
- Launch with a media and online marketing blitz—your goals are to get as many new inbound links as possible pointing to the new site as quickly as possible, and to attract a high number of branded searches for the redesigned site.
- Monitor your rankings for the content, comparing old to new over time.
- Monitor your Webmaster Tools account for 404 errors and to see how well Google is handling your 301s. When you see some 404s pop up, make sure you have a properly implemented 301 redirect in place. If not, fix it.
- Monitor the search engine spidering activity on the new domain. This can provide a crude measurement of search engine trust. Search engines spend more time crawling sites they trust. When the crawl level at the new site starts to get close to where it was with the old site, you are probably most of the way there.
- Watch your search traffic referrals as well. This should provide you some guidance as to how far along in the process you have come.
- You can also check your server logs for 404 and 500 errors. This will sometimes flag problems that your other checks have not revealed.

As mentioned earlier, Google’s Matt Cutts discussed the benefits of a phased approach at a PubCon event (<http://blog.milestoneinternet.com/web-development/faq-on-duplicate-content-and-moving-your-site-by-matt-cutts-at-pubcon-2009/>):

So here’s the extra step. Don’t just move the entire domain from the old domain to the new domain. Start out and then move a subdirectory or a subdomain. Move that first; if you’ve got a forum, move one part of your forum. Move that over to the new domain, and make sure that the rankings for that one part of your site don’t crash. Sometimes it takes a week or so for them to sort of equalize out, because we have to crawl that page to

see that it's moved. So if you move a part of your site first, and it goes fine, then you know that you're pretty safe. So instead of doing one huge move, if you can break it down into smaller chunks and start out by moving a small part of your site first, you'll know that you'll be gold.

The value of this approach is that it reduces the risk associated with the move into more manageable chunks. Even if you use this approach, however, you should still follow the process outlined in this section to implement the move of each site section, and check on its progress.

Changing Servers

Sometimes you may decide you want to move servers without changing your domain name or any of your URLs. A common cause of this is that the growth of your traffic requires you to *upgrade* your hosting environment to a faster server. If you are using third-party hosting, perhaps you are changing your hosting company, or if you have your own data center, you may need to move or expand your facilities, resulting in a change in the IP addresses of your servers.

This is normally a straightforward process, as you can simply go to the registrar where you registered the domain name and update the DNS records to point to the new server location. You can also temporarily decrease the site's DNS Time to Live (TTL) to five minutes (or something similar) to make the move take place faster. This is really the bulk of what you need to do, though you should follow the monitoring recommendations we will outline shortly.

Even if you follow this process, certain types of problems can arise. Here are the most common ones:

- You may have content that can't function on the new platform, for example, if you use Perl in implementing your site and Perl is not installed on the new server. This can happen for various other reasons as well, and the result can be pages that return 404 or 500 errors instead of the content you intended.
- Unfortunately, publishers commonly forget to move key content or files over, such as *robots.txt*, analytics files, *sitemaps.xml*, the *.htaccess* file, and so forth. It is imperative that these important files are migrated to your new server.
- Server configuration differences can also lead to mishandling of certain types of requests. For example, even if both your old server and your new server are running IIS, it is possible that the new server is configured in such a way that it will turn any 301 redirects you have in place into 302 redirects – be sure to double and triple check all server directives are properly migrated from the old server to the new.

The best advice for dealing with these concerns is to make a list of special files and configuration requirements and verify that everything is in place prior to flipping the switch on any server moves.

In addition, you should conduct testing of the new site in its new location before flipping the switch. You will need to access the content on the new site using its physical IP address. So, the page at <http://www.yourdomain.com/pageA> will be found at an address similar to <http://206.130.117.215/pageA>. To do this, you should add that IP address to your test machine's *hosts* file (this assumes you are running Windows) with a corresponding hostname of <http://www.yourdomain.com>, which will allow you to surf the site at the new IP address seamlessly. This advance testing should allow you to check for any unexpected errors. Note that the location of the *hosts* file varies across different versions of Windows, so you may need to search online to get information on where to find it on your machine.

Monitoring After Your Server Move

As with our other scenarios, post-launch monitoring is important. Here are the basic monitoring steps you should take:

- Monitor your Google and Bing Webmaster Tools accounts for 404 errors and to see how well the search engines are handling your 301s. When you see 404 errors, make sure you have a properly implemented 301 redirect in place.

- Monitor the spidering activity on the new domain to make sure no unexpected drops occur.
- Watch your search traffic referrals for unexpected changes.
- You can also check your server logs for 404 and 500 errors. This will sometimes expose problems that your other checks have not revealed.

Hidden Content

In “Content Delivery and Search Spider Control” in Chapter 6, we discussed ways that you can hide content from the search engines when you want to. However, at times this is done unintentionally—that is, sometimes publishers produce great content and then, for one reason or another, fail to expose that content to search engines.

Valuable content can be inadvertently hidden from the search engines, and occasionally, the engines can find hidden content and construe it as spam, whether that was your intent or not.

Identifying Content That Search Engines Don’t See

How do you determine when this is happening? Sometimes the situation is readily apparent; for example, if you have a site that receives high traffic volume and then your developer accidentally `NoIndexes` every page on the site, you will begin to see a catastrophic drop in traffic. Most likely this will set off a panic investigation, leading to identifying the `NoIndex` issue as the culprit.

Does this really happen? Unfortunately, it does. Here is an example scenario. Suppose you work on site updates on a staging server. Because you don’t want the search engines to discover this duplicate version of your site, you keep the pages on the staging server `NoIndexed`. Normally, when you move the site from the staging server to the live server, you remove the `NoIndex` tags; but unfortunately, many forget to do this.

This type of problem can also emerge in another scenario. Some webmasters implement a `robots.txt` file that prohibits the crawling of their staging server website. If this file gets copied over when the site on the staging server is switched to the live server, the consequences will be just as bad as in the `NoIndex` scenario we just outlined. The best way to prevent this type of scenario is to introduce a series of safety checks on the site that take place immediately after any update of the live server.

There are potential problems, however, that are much more difficult to detect. First, with a new site launch, you won’t have any preexisting traffic, so there will be no drop in traffic levels to alert you that something is wrong. In another scenario, you may have an established site where you accidentally do something to hide only a portion of the site from the engines, so the issue is less obvious.

Regardless of your situation, web analytics can help you in the detection process. Use your analytics software to find pages on your site that get page views but no referring search traffic. By itself, this is not conclusive, but it provides a leading clue as to where to start. Note that the converse of this is interesting for another situation—if you see content that is getting search referrals even though you don’t want or expect it to, you may want to hide that content.

Another data point you can examine is the number of pages the search engines report as indexed for your site. In a new site scenario, you can look at this to see whether the search engines appear to be picking up your content. For example, if you have a site with 1,000 pages with a good inbound link profile, and after three months only 10 pages are indexed, that could be a clue that there is a technical problem. Using multiple sitemap files, one for each site content area covering a specific segment of URLs, can be helpful in diagnosing such problems.

You do need to be careful not to overreact to the count of indexed pages, because the numbers that the search engines report will naturally fluctuate quite a bit. But if you are aware of the types of numbers typically reported for your site, and they drop to an unusually low level and stay there (or keep dropping), you probably have a problem.

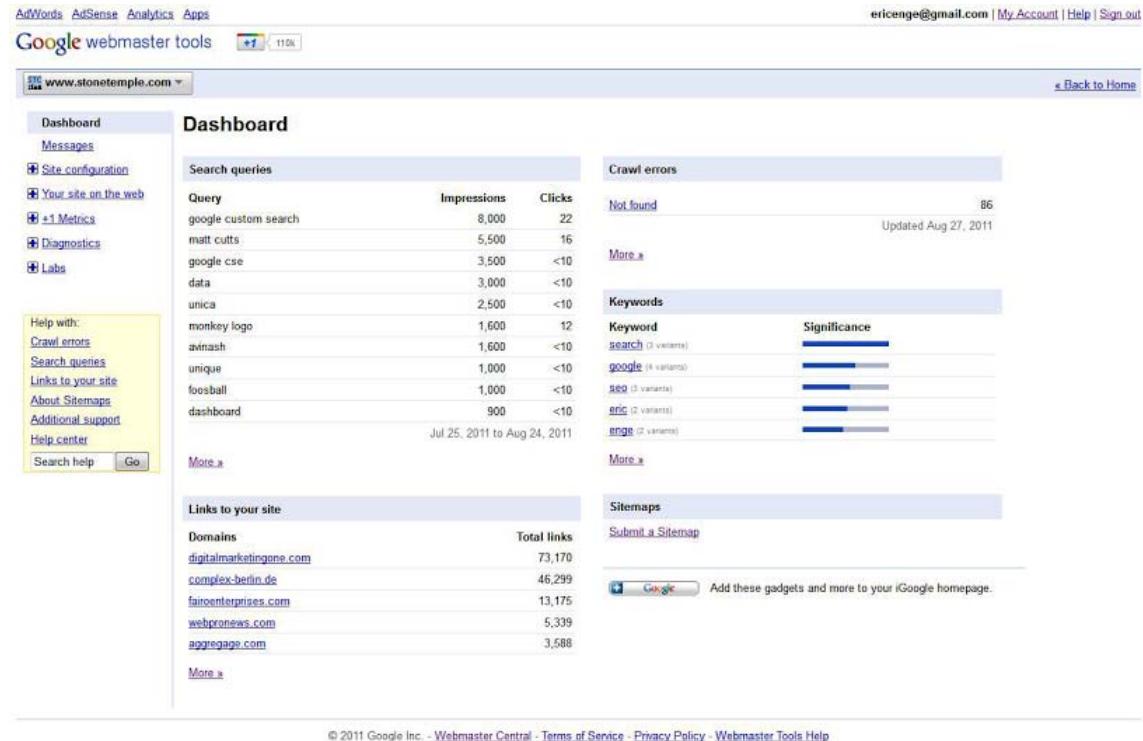
Identifying the Cause of Non-Spidering

Once you realize you have a spidering problem, there are a series of checks you can perform to identify the potential causes:

Blocked by robots.txt

Check your *robots.txt* file (<http://www.yourdomain.com/robots.txt>) to see whether you are preventing the crawlers from accessing parts of the site that you actually want them to see. This mistake is quite common.

Both Google and Bing Webmaster Tools provide simple ways to see whether they are aware of content that *robots.txt* is blocking them from crawling (see Figure 12-1 and Figure 12-2).



The screenshot shows the Google Webmaster Tools dashboard for the site www.stonetemple.com. The left sidebar includes links for Site configuration, Your site on the web, +1 Metrics, Diagnostics, and Labs. A yellow box highlights the 'Help with' section, which includes links for Crawl errors, Search queries, Links to your site, About Sitemaps, Additional support, Help center, and Search help. The 'Go' button is also highlighted. The main content area is titled 'Dashboard' and contains several sections: 'Search queries' (listing queries like 'google custom search', 'matt cutts', 'google cse', 'data', 'unique', 'monkey logo', 'avinash', 'unique', 'foosball', and 'dashboard' with their impressions and clicks), 'Crawl errors' (listing 'Not found' with 86 errors, updated Aug 27, 2011), 'Keywords' (listing keywords like 'search', 'google', 'seo', 'eric', and 'engg' with their significance), 'Links to your site' (listing domains with total links: digitalmarketingone.com (73,170), complex-berlin.de (46,299), fairoenterprises.com (13,175), webpronews.com (5,339), and aggregate.com (3,588)), and 'Sitemaps' (with a 'Submit a Sitemap' link and a note to add gadgets to the homepage). The bottom of the page includes a copyright notice: © 2011 Google Inc. - [Webmaster Central](#) - [Terms of Service](#) - [Privacy Policy](#) - [Webmaster Tools Help](#).

[[figs/print/1201.png]]

Figure 12-1. Google Webmaster Tools: restricted by robots.txt

Crawl errors

Issues Google encountered when crawling your site.

Web	Mobile CHTML	Mobile WML/XHTML	
Show URLs: Not found (86)			
URL	Detail	Linked From	Detected
http://www.stonetemple.com/blog/page/2/?cat=22	404 (Not found)	1 pages	Aug 25, 2011
http://www.stonetemple.com/blog/page/2/?category_name=mobile	404 (Not found)	1 pages	Aug 25, 2011
http://www.stonetemple.com/blog/page/2/	404 (Not found)	5 pages	Aug 24, 2011
http://www.stonetemple.com/blog/page/2/?cat=23	404 (Not found)	1 pages	Aug 23, 2011
http://www.stonetemple.com/blog/page/2/?cat=18	404 (Not found)	1 pages	Aug 23, 2011
http://www.stonetemple.com/articles/interview-linda-bustos.shtml	404 (Not found)	1 pages	Aug 23, 2011
http://www.stonetemple.com/blog/wp-links-opml.php	404 (Not found)	3 pages	Aug 22, 2011
http://www.stonetemple.com/sample-page/	404 (Not found)	unavailable	Aug 22, 2011
http://www.stonetemple.com/blog-temp.shtml	404 (Not found)	3 pages	Aug 22, 2011
http://www.stonetemple.com/blog/feed/?cat=15	404 (Not found)	unavailable	Aug 22, 2011
http://www.stonetemple.com/articles/analytics-report-august-2007-part2.shtml	404 (Not found)	3 pages	Aug 22, 2011
http://www.stonetemple.com/articles/analytics-report-august-2007.shtml	404 (Not found)	5 pages	Aug 22, 2011

[[figs/print/1202.png]]

Figure 12-2. Google Webmaster Tools: crawl errors

These reports are helpful when you have content on the site that has links to it (either internal or external), but that the search engines don't crawl because they are excluded from it in *robots.txt*.

The solution is simple: figure out what line in your *robots.txt* file is blocking the search engines and remove it, or update it so that the content is no longer blocked.

Blocked by the robots meta tag

The `robots` meta tag in a page's header might look something like this:

```
| <meta name="robots" content="noindex,nofollow">
```

As we discussed in "Content Delivery and Search Spider Control" in Chapter 6, a setting of `NoIndex` will tell the search engine that it is not allowed to include the page in its index. Clearly, you should check to see whether you have made this error if you find that the engines are not crawling certain pages that you do want crawled.

In addition, the `NoFollow` tag will tell search engines not to pass any link authority through the links contained to the pages linked to on that page. If all the links on your site to a particular piece of content are `NoFollowed`, you are passing no link authority to the page. This tells the search engines that you don't value the page, and as a result they won't treat the links as endorsements for the page.

Solving this problem requires locating the places where these `robots` meta tags are on your site and removing them. Note that the default setting for the `robots` meta tag is `"index,follow"`, so there is no need to implement the tag if that is your desired setting. Just make sure you don't have `robots` meta tags in place that change the default in places where that is not the desired behavior.

No direct links

You may find that a particular piece of content has no links to it, or that all of the links to that content are `NoFollowed`. You can also make links invisible to the search engines (possibly unintentionally) by showing them only to visitors that support cookies, or by encrypting the links to the content in some fashion. The solution here is to make sure you implement plain text (or image) links to the content. Better still, get some third-party websites to link to the content as well.

Requires form submission

Requiring a login or some other type of form submission to see content is another common cause of non-spidering. Search engines will not attempt to fill out forms to see what is behind them. The simplest solution can be to remove the requirement for the form if you want the search engines to index this content.

However, some sites sell content on a subscription basis (also referred to as being behind a “paywall”), and they will not want to offer their content for free. In October 2008, Google announced the First Click Free concept (<https://support.google.com/news/publisher/answer/40543?rd=1>) to allow subscription-based sites to have their content crawled and indexed by Google, but still allow the publisher to require human visitors to subscribe to access the content. As of 2014, Google requires that any user coming from a Google domain must be able to see a minimum of 5 free articles per day, which they refer to as “metering.” You can also read more about First Click Free in “Content Delivery and Search Spider Control” in Chapter 6.

Session IDs

Session IDs confuse search engine crawlers. Every time they come to your site, they see a different page. For example, they may see <http://www.yourdomain.com?SessID=2143789> one time and <http://www.yourdomain.com?SessID=2145394> the next time. Even though your intent is to track the session of a particular user, and you think of these URLs as the same page, the search engine does not. You can read more about session IDs in “Controlling Content with Cookies and Session IDs”.

Not enough link authority to remain in main indexes

Sometimes the problem has nothing to do with the issues we just discussed. The search engines may see the page just fine, but there may not be enough link juice going to it to merit inclusion in their main indexes. This is more common than people think, and it happens because the search engines do not attempt to index all the world’s web pages.

For example, content that Google perceives to be of low importance (i.e., content that doesn’t have enough link authority, or it’s perceived to be duplicate content) will be excluded from the main index; in previous years, this content may have been relegated to what Google called its “supplemental index,” but as of 2014, Google’s John Mueller confirms that for the purposes of treating pages differently, Google no longer has a supplemental index (<http://www.youtube.com/watch?v=tn-RE-n6kGs>).

Google wants to emphasize the more important pages on the web, and doesn’t want the rate at which it delivers search results to be slowed down by pages that most people probably don’t want to see.

Hidden Content That May Be Viewed as Spam

Hidden text is one of the challenges that webmasters and search engines still face. Spammers continue to use hidden text to stuff keywords into their pages, for the purposes of artificially boosting their rankings. Search engines seek to figure out when spammers are doing this and then take appropriate action. There are many ways to create hidden text unintentionally, though, and no one wants to be penalized for something they did not intend to do. It is useful to examine Google’s Webmaster Guidelines for hidden text (<https://support.google.com/webmasters/answer/66353?hl=en>) :

“Hiding text or links in your content to manipulate Google’s search rankings can be seen as deceptive and is a violation of Google’s Webmaster Guidelines.”

In a Google Groups thread (http://groups.google.com/group/Google_Webmaster_Help-Indexing/browse_thread/thread/928aa76a1226cf89/32b089e3248cef78?pli=1), Googler Susan Moskwa had this to say:

Of course, as with many techniques, there are shades of gray between “this is clearly deceptive and wrong” and “this is perfectly acceptable.” Matt [Cutts, head of Google’s webspam team] did say that hiding text moves you a step further towards the gray area. But if you’re running a perfectly legitimate site, you don’t need to worry about it. If, on the other hand, your site already exhibits a bunch of other semi-shady techniques, hidden text starts to look like one more item on that list. It is like how 1 grain of sand isn’t noticeable, but many grains together start to look like a beach.

Related to this is a posting by Matt Cutts on Threadwatch (<http://www.threadwatch.org/node/4313#comment-26883>):

If you're straight-out using CSS to hide text, don't be surprised if that is called spam. I'm not saying that mouseovers or DHTML text or have-a-logo-but-also-have-text is spam; I answered that last one at a conference when I said "imagine how it would look to a visitor, a competitor, or someone checking out a spam report. If you show your company's name and it is Expo Markers instead of an Expo Markers logo, you should be fine. If the text you decide to show is 'Expo Markers cheap online discount buy online Expo Markers sale...' then I would be more cautious, because that can look bad."

Obviously, this is a fate you want to avoid. Note the use of the word *perceived* in the Google Webmaster Guidelines snippet. It doesn't sound like a simple black-and-white problem, does it? In fact, it is not, as there are many ways to create hidden text.

Unintentionally creating hidden text

There are a few ways to create hidden text without intending to do so. One of the most common ways is via your CMS, which has some CSS-based methods built into it. For example, many CMSs use the `display:none` technique to implement drop-down menus or other widgets that the user clicks on that then "expand" to display more text. Tab folders are a great example of this. Sometimes the `display:none` technique is used in user-generated content systems where the page normally shows the number of comments on a post, but chooses to suppress the text "0 Comments" in the event that no comments have been made yet.

Another common way that people create hidden text occurs when they provide enhancements for the visually impaired. For example, you may have a Flash object on your web page and want to provide users with a text description of the content. You may not want to place the text on the page, as it might make the page look cluttered to a user with normal vision. The solution some people use to serve both audiences is to hide the text from the sighted users.

Many of these scenarios have no SEO value, even when manipulated by spammers. These types of techniques generally do not carry a risk of being penalized, because there is no reason to suspect negative intent.

Spam Filtering and Penalties

Over time, it has become a lot more difficult to "game" the search engines and a lot easier to fall victim to a search engine penalty or outright ban. It is hard to recover from these.

Consequences can include ranking penalties, removal of the site's "voting" power (i.e., ability to pass PageRank), incomplete indexation (i.e., a partial site ban), or, worst of all, a total site ban.

Not even the largest corporations spending big dollars on Google AdWords are immune. In 2013, both BBC News and the popular music lyrics website RapGenius.com were hit with Google penalties. In the case of BBC News, Google Webmaster Tools sent an "unnatural link" notification (which represents between 1-2% of Google Webmaster Tools notifications: <http://searchengineland.com/google-sends-hundreds-of-thousands-of-webmaster-notifications-each-month-90-are-black-hat-related-148524>); while the case turned out to be targeted to only one article (<http://searchengineland.com/google-penalized-one-article-on-bbcs-web-site-151954>), it caused a stir.

The Rap Genius case, however, was a much more serious matter that involved the company's invitation to bloggers to add links to Rap Genius lyrics URLs in exchange for Rap Genius subsequently tweeting the bloggers' posts to Rap Genius's Twitter followers. This link scheme came to the attention of Google's Matt Cutts, who vowed Google would investigate, and they did – resulting in an official Google penalty for link schemes (<http://searchengineland.com/google-has-officially-penalized-rap-genius-for-link-schemes-180777>).

Search engines rely primarily on automated means for detecting spam, with some auxiliary assistance from paid evaluators, spam vigilantes, and even your competitors. Search engineers at Google and Microsoft

write sophisticated algorithms to look for abnormalities in inbound and outbound linking, in sentence structure, in HTML coding, and so on.

As far as the search engines are concerned, SEO has both an acceptable side and an unacceptable side – and in general terms, many efforts intended to boost a site’s search engine ranking without improving the measurable value of a page can be considered spamming.

Each search engine has different published guidelines. Here is where you can find them:

- Google’s Webmaster Guidelines are located at
<https://support.google.com/webmasters/answer/35769>
- Bing offers some tips for publishers that you can access from this page:
<http://www.bing.com/webmaster/help/webmaster-guidelines-30fba23a>

The search engines have varying degrees of tolerance for various SEO tactics. Anything that violates these guidelines, pollutes the search results with irrelevant or useless information, or would embarrass you if your Google AdWords or Bing rep discovered it - is unsustainable, and should generally be avoided.

There is a difference between “search engine–friendly,” and crossing the line into spam territory. “Search engine–friendly” can mean, for example, that the site is easily accessible to spiders, even if it is database-driven; that HTML code is streamlined to minimize the amount of superfluous code; that important headings, such as product names, are set apart from the rest of the text (e.g., with `<h1>` tags) and they contain relevant keywords; or that link text is contextual, instead of comprising just “click here” or “more info” references.

Contrast these basic SEO practices with the following manipulative search engine spam tactics:

- Serving pages to the search engines that are useless, incomprehensible, unsuitable for human viewing, or otherwise devoid of valuable content—such as doorway pages, which SEO vendors may refer to by more innocuous names, including “gateway pages,” “bridge pages,” “jump pages,” “attraction pages,” “advertising pages,” “channel pages,” “directory information pages,” “search engine entry pages,” “satellite sites,” “mini sites,” “magnet sites,” or “shadow domains.” Whatever you call them, by definition they are created for the sole purpose of boosting search engine rankings.
- Creating sites with low quality content. There are many techniques for doing this, including:
 - Duplicating pages with minimal or no changes and exposing them to the same search engines under new URLs or domains.
 - Machine-generating content to chosen keyword densities (e.g., using a technique such as Markov chains; but they are not recommended).
 - Incorporating keyword-rich but nonsensical gibberish (also known as *spamglish*) into site content.
 - Creating a low-value site solely for affiliate marketing purposes (see the upcoming subsection on duplicate content for a more complete definition of *thin affiliate*).
- Repeating the same keyword phrase in the title tag, the `<h1>` tag, the first `alt` attribute on the page, the meta description, the first sentence of body copy, and the anchor text in links pointing to the page.
- Targeting topically irrelevant keywords.
- Concealing or obscuring keyword-rich text or links within the HTML of a page so that it is not visible to or accessible by human users (i.e., by placing it within comment tags, `noscript` tags, or `noframe` tags; or by using colored text on a similarly colored background, tiny font sizes, layers, or links that don’t show as links to users because they are not highlighted in some manner, such as with an underline).
- Hijacking or stealing content from other sites and using it as content fodder for search engines. This is a practice normally implemented using scrapers.
- Purchasing links for the sole purpose of influencing search rankings.

- Participating in *link farms* (which can be distinguished from directories in that they are less organized and have more links per page) or *reciprocal linking schemes* (link exchanges) with irrelevant sites for the purpose of artificially boosting your site’s importance.
- Peppering websites’ guest books, blogs, or forums in bulk with keyword-rich text links for the purpose of artificially boosting your site’s importance.
- Conducting *sneaky redirects* (immediately redirecting searchers entering your site from a keyword-rich page that ranks in the search engine to some other page that would not rank as well).
- *Cloaking*, or detecting search engine spiders when they visit and modifying the page content specifically for the spiders to improve rankings.
- *Negative SEO*, or submitting your competitors to link farms and/or otherwise creating third party associations on the competitor’s behalf so that they will be penalized.

These tactics are questionable in terms of effectiveness and dubious in the eyes of the search engines, often resulting in the offending site being penalized by or banned from the search engines—a risk that’s only going to increase as the engines become more aggressive and sophisticated at identifying and removing offenders from their indexes. We do not advocate implementing these tactics to those interested in achieving the long-term benefits of SEO.

The search engines detect these tactics not just by automated means, through sophisticated spam-catching algorithms, but also through spam reports submitted by searchers—and yes, by your competitors.

Often times search marketers don’t even know they’re in the wrong and running afoul of search engine guidelines. For example, search engines place more scrutiny on pages that show signs of potential deception, such as *no-archive* tags, *noscript* tags, *noframe* tags, and cloaking, even though all of these can be, and often are, used ethically.

Seeing SEO strictly as a chess game between yourself and the search engines is a shortsighted view, as the goal of the search engines is to provide relevant search results to their users, and your goal as a search marketer is to provide people with valuable content and a positive website experience. Trying to fool the search engines and take unfair advantage by exploiting how the search engines work is simply not a sustainable approach.

In Chapter 9 (“Panda, Penguin, and Penalties”) we address the topic of Google’s manually applied penalties, as well as their Panda and Penguin algorithms in detail.

Recognizing Low-Quality Domains and Spam Sites

You can spot a poor-quality web site in many ways, not the least of which is the “common sense” check. Would you hire a company with a website named www.seoseosanfranciscoseo.com, or sonysonyelectronicsonyforsale.com? The domain, of course, is only one signal, and search engines rely on a wide range of signals as indicators of quality. Some of the most obvious signals are site owners who are actively spamming the search engines with their offsite activities—for example, if the site is actively buying links, or text link-spamming blog posts, forums, and article comments.

However, there are also less obvious signals. Many such signals mean nothing by themselves and gain significance only when they are combined with a variety of other signals. When a number of these factors appear in combination on a site, the likelihood of it being seen as a low-quality or spam site increases.

Here is a long list of some of these types of signals:

- Short registration period
- High ratio of ad blocks to content
- JavaScript redirects from initial landing pages
- Use of common, high-commercial-value spam keywords such as *mortgage*, *poker*, *texas hold ‘em*, *porn*, *student credit cards*, and related terms
- Many links to other low-quality spam sites

- Few links to high-quality, trusted sites
- High keyword frequencies and keyword densities
- Zero, or very little, unique content
- Very few direct visits
- Registered to people/entities previously associated with untrusted sites
- Not registered with services such as Google or Bing Webmaster Tools
- Rarely have short, high-value domain names
- Often contain many keyword-stuffed subdomains
- More likely to have longer domain names (as above)
- More likely to contain multiple hyphens in the domain name
- Less likely to have links from trusted sources
- Less likely to have SSL security certificates
- Less likely to be in high-quality directories such as DMOZ, Yahoo!, Best of the Web, and so forth
- Unlikely to have any significant quantity of branded searches
- Unlikely to be bookmarked in services such as Delicious, etc.
- Unlikely to get featured in social voting sites such as Digg, reddit, StumbleUpon, and so forth
- Unlikely to have channels on YouTube, communities on Google+, Facebook, or links from Wikipedia
- Unlikely to be mentioned on major news sites (either with or without link attribution)
- Unlikely to register with Google & Bing local listings
- Unlikely to have a legitimate physical address/phone number on the website
- Likely to have the domain associated with emails on blacklists
- Often contain a large number of snippets of “duplicate” content found elsewhere on the Web
- Frequently feature commercially focused content
- Many levels of links away from highly trusted websites
- Rarely contain privacy policy and copyright notice pages
- Rarely listed in the Better Business Bureau’s Online Directory
- Rarely contain high-grade-level text content (as measured by metrics such as the Flesch-Kincaid *Reading Level*)
- Rarely have small snippets of text quoted on other websites and pages
- Cloaking based on user agent or IP address is common
- Rarely have online or offline marketing campaigns
- Rarely have affiliate link programs pointing to them
- May have links to a significant portion of the sites and pages that link to them
- Extremely unlikely to be mentioned or linked to in scientific research papers
- Unlikely to use expensive web technologies (Microsoft Server and coding products that require a licensing fee)
- More likely to contain malware, viruses, or spyware (or any automated downloads)
- Likely to have privacy protection on the Whois information for their domain

It is important to note that while many of the signals above can be viewed negatively in aggregate, having one, such as a private domain registration, in and of itself is not going to be interpreted as a spam signal.

Many legitimate sites will have one or more of these signals associated with them. For example, there are many good sites with a *.info* TLD.

There are also some signals that can be picked up that require data from a web analytics tool (which Google may be able to obtain from the Google Analytics account of the site in question):

- Rarely receive high quantities of monthly visits
- Rarely have visits lasting longer than 30 seconds
- Rarely have visitors bookmarking their domains in the browser
- Unlikely to buy significant quantities of PPC ad traffic
- Rarely have banner ad media buys
- Unlikely to attract significant return traffic

There can be legitimate reasons for the occurrence of many (possibly even most) of the signals mentioned above. For instance:

- Businesses outside the United States will not be in the Better Business Bureau directory.
- The site may not be relevant to scientific research papers.
- The publisher may not be aware of Google Webmaster Tools or Bing Webmaster Tools.

These are just a few examples meant to illustrate that these signals all need to be put into proper context.

Competitors Can Report You

Search engines supplement their internal spam fighting efforts by allowing users to submit spam reports. For example, Google provides a form for reporting spam at <https://www.google.com/webmasters/tools/spamreport>.

In addition, the search engines can and do make use of human reviewers who conduct quality reviews.

In late 2012, Google published their “Cliffs Notes” version of their criteria for search quality rating, in their “Search Quality Rating Guidelines” document:

<http://static.googleusercontent.com/media/www.google.com/en/us/insidesearch/howsearchworks/assets/searchqualityevaluatorguidelines.pdf>.

Duplicate Content

As we discuss in “Content Management System (CMS) Issues” in Chapter 6, there are many ways to unintentionally create duplicate content. For this reason, the normal response of the search engines to duplicate content is to filter it out, but not penalize the publisher for it in any other way. They filter it out because they don’t want to show multiple copies of the same piece of content in their search results, as this does not really bring any value to users. They don’t punish the publisher because the great majority of these situations are unintentional.

In fact, in late 2013, Google’s Matt Cutts declared that 25%-30% of content on the web is duplicative and that unless the content duplication was egregious and malicious, not to worry too much about it: <http://www.youtube.com/watch?v=mQZY7EmjbMA>.

However, there are three notable exceptions that can result in being impacted by Google’s “Panda” algorithm (<http://searchengineland.com/library/google/google-panda-update>):

Copyright violations

In this scenario, there is an actual copyright violation where a publisher is showing a copy of another publisher’s content, usually via scraping, without permission.

In 2014 Google announced the new “Google Scraper Report” with Google’s Matt Cutts tweeting:

“If you see a scraper URL outranking the original source of content in Google, please tell us about it.”
The Scraper Report can be accessed here: <http://t.co/WohXQmI45X>.

Thin affiliate sites

In this scenario, the publisher has permission from another publisher (the “Content Owner”) to display its content, but the Content Owner also shares that exact same content with many other sites. The common scenario is a Content Owner running an affiliate network and generating leads or sales for its products or clients, largely by offering an affiliate program to other publishers.

All of the affiliate sites publish the exact same (or very similar) content. The problem the search engines have with these sites is that they offer very little value, as there is nothing truly unique about their content.

The affiliate site may also create hundreds or thousands of pages to target vertical search terms with little change in content. The classic example of this is creating hundreds of web pages that are identical, except changing the city name referred to on each page (pages have titles such as “Phoenix Oil Changes,” “Austin Oil Changes,” “Orlando Oil Changes,” etc.).

Massive duplication

This scenario applies to a site that is not a thin affiliate but that has a very large amount of duplicate content (even if it has duplicated that content with permission). It is not known what the threshold is, and this probably changes over time, but our experience suggests that sites in which 70% or more of the pages are duplicates of other pages on the Web are likely to be subject to a penalty.

There may be other spammy forms of duplicate content not identified here. It is likely to be spam if it is implemented intentionally (e.g., the thin affiliate site example qualifies under this metric), if it is intended to manipulate search results, and if the content adds no value to the Web – whereas syndicated content (content published on the original source’s website and then copied with express permission on third-party websites), when published properly, presents no duplicate content issues. Co-author Eric Enge’s SearchEngineLand.com article entitled “Syndicated Content: Why, When, and How” provides guidelines for how to properly and safely syndicate your site content: <http://searchengineland.com/syndicated-content-189097>.

Basic Rules for Spam-Free SEO

Especially if you are new to SEO, the first and most important rule is to be familiar with the guidelines of the search engines (see the beginning of this section for the locations of the guidelines on the search engine sites).

Second, it is essential that you learn to apply a basic personal filter to your SEO activities. Generally speaking, if you engage in an activity for the sole purpose of influencing search rankings that has no user or search engine crawler benefit, you are putting your site’s rankings at risk.

For example, if you start buying keyword-rich text links from a bunch of sites across the Web and you do not expect to get significant traffic from these links (enough to justify placing the advertisements), you are headed for trouble and will eventually expose you to being hit by Google’s “Penguin” algorithm: <http://searchengineland.com/library/google/google-penguin-update>.

Of course, there are more ways you can get into trouble than through purchasing or aggressively spamming other websites with keyword-rich text links. Earlier in this chapter, we listed a number of spammer techniques. Most publishers/SEO practitioners won’t run into the majority of these as they involve extremely manipulative behavior, as we outlined previously in this section. However, newcomers to SEO do tend to make certain mistakes. Here are some of the more common ones:

- Stuffing keywords into your web page so it is unnaturally rich in those words.
- Over-optimizing internal links (links on your website to other pages on your website). Generally speaking, this may occur by overstuffing keywords into the anchor text of those links.
- Cloaking, or showing different content to the search engine crawlers than you show to users.

- Creating websites with lots of very thin content pages, such as the thin affiliate sites we discussed previously.
- Implementing pages that have search engine-readable text that is invisible to users (hidden text).
- Participating in link schemes, link farms, aggressive guest blog posting, or using other tactics to artificially boost link popularity.

As we discussed previously, there are many other ways to end up in the crosshairs of the search engines, but most of those fall into the category of highly manipulative SEO tactics (generally referred to in the industry as *black hat SEO*).

Ultimately, you want to look at your intent in pursuing a particular SEO practice. Is it something you would have done if the search engines did not exist? Would it have been part of a publishing and promotional plan for your website in such a world?

This notion of intent is something that the search engines look at very closely when evaluating a site to see whether it is engaging in spam tactics. In fact, search engine representatives speak about *intent* (why you did it) and *extent* (how much did you do it) as being key things they evaluate when conducting a human review of a website.

A perceived high degree of intent and pursuing something to a significant extent can lead to the more severe penalties being applied. But even if your intent is pure and you don't pursue a proscribed practice to a significant degree, the search engines will want to discount the possible benefits of such behavior from the rankings equation. For example, if you buy a few dozen links and the search engines figure out that the links were paid for, they will discount those links. This behavior may not be egregious enough for them to penalize you, but they don't want you to benefit from it either.

Identifying Search Engine Penalties & Reconsideration Requests

The reality is that penalties are imposed on websites. It is often hard to know what has caused a penalty – and although the search engines might describe the general reason for the penalty (such as “unnatural links”), they will not identify the behaviors specifically. You will need to review your site to see where you have deviated from best practices and address all such issues. For a thorough discussion on how identify search engine penalties and how to recover from them, see Chapter 9, “Panda, Penguin, and Penalties”

Content Theft

The Web is sometimes a little bit of a Wild West environment, and one of the bad things that can occur is that another publisher (a spammer) can copy your content and republish it on one of his own sites. This is very easy for infringers to do. They simply “View Source” in their browsers and take whatever they like without regard to copyright, or send a spider to collect all of your pages in one fell swoop.

You can locate copyright infringers pretty easily using various online tools, including Plagium (<http://www.plagium.com>) and Copyscape (<http://www.copyscape.com>), if they've lifted some of your page copy. It's much more difficult if they've limited their sticky fingers to just your design.

There are six major options for dealing with an infringer:

- File a DMCA (Digital Millennium Copyright Act) infringement notification.
- Contact the infringing company's CEO, head of marketing, legal counsel, or whomever you feel is most likely to respond.
- Notify the site's hosting company, if it is not self-hosted.
- Report the activity on your blog, or if your blog is not that popular, get someone with a popular blog to report them.
- Have your lawyer send them a cease and desist letter.

Doing nothing is not an option if you're serious about protecting your intellectual property rights. At the very least, you want to show you have made an effort to protect your IP in the event ownership comes into question in other scenarios.

A phone call to the infringing company is inexpensive, and it can work. Sometimes publishers who use contract writers are not even aware that they are infringing. But this is often not a viable option if you're dealing with an entity that has masked its domain contact information and is hosting its site with a company that will not respond to your requests, and where a copyright complaint will not be effective (i.e., in countries without copyright laws).

In the US, the DMCA provides a useful hammer to beat on copyright infringers without the need to contact them yourself, in the form of a DMCA infringement notification. In DMCA legal speak, this notification is also known as a "Takedown Notice." Other countries with strong copyright laws have similar processes.

As a content producer, you have the right to enforce your copyright. When your content gets "repurposed" on others' websites without your permission, you can simply file a DMCA infringement notification with the infringer's web hosting provider to get that infringer's website shut down.

You can also get the infringing content delisted from the search engines by filing DMCA notices with Google and the other major engines. It is not a daunting procedure. It might take an hour of your time, and it is well worth it. The process is as follows:

1. First, look up the web host and the domain registrar of the offending site using lookup tools such as DomainTools (<http://www.domaintools.com>). You can usually ascertain who the web host is from the name servers and/or the netblock owner.
2. Check the official directory of designated DMCA agents (http://www.copyright.gov/onlinesp/list/a_agents.html) for the host and the registrar. (Hopefully, they're listed.)
3. Prepare a letter to send to the designated agent of the web host. The notice you write should include your contact information, a description of and a link to the content that was copied, the web address of the copied content, a statement that you have a good faith belief that the duplication of the material is not legal, a statement that under penalty of perjury you are the copyright holder, and your signature. Some web hosts will allow you to email your notice to them, making it all the easier.
4. Send a similar notification to the search engines. That will cut off the site's air supply (its traffic) in case the content doesn't get taken down right away. Instructions and contact details for each engine are available at the following URLs:
 - Google: <https://support.google.com/legal/troubleshooter/1114905?rd=2>
 - Bing: <https://www.microsoft.com/info/cpyrtInfrg.aspx>
5. If the web host doesn't take the site down promptly, submit a DMCA notice to the infringer's domain registrar. It might be worth sending a notice to the data center that the web host uses before you try the registrar.

Be aware that the DMCA takedown procedure is a double-edged sword. Like any other tool, a DMCA notice can be used for good or evil. Specifically, someone could use one unfairly against you! It happens—competitors do sometimes use the DMCA to silence their competitors (DMCA letter templates specific to each search engine can be found at: <http://www.mcanerin.com/EN/articles/copyright-03.asp>).

Therefore, you as a website owner need to protect yourself from unwarranted (or at least unwelcome) prosecution. If the potential exists for you to inadvertently host infringing material on your website—for example, if you are hosting online forums, group blogs, blog comments, or other types of content that can be submitted by others—here are some actions you can take to help protect yourself:

- It's helpful if you can qualify as a service provider (<http://www.chillingeffects.org/dmca512/faq.cgi#QID127>) that can be covered under the Safe Harbor provision (<http://www.chillingeffects.org/dmca512/faq.cgi#QID129>). For example, you may qualify if you offer a search engine or a bulletin board system.

If so, notify your customers of your policies (<http://www.chillingeffects.org/dmca512/faq.cgi#QID128>) regarding copyright infringement and the consequences of repeated infringing activity. One way is to make this part of your Terms of Use.

- Publish a page on your website with DMCA filing instructions and state that, if and when you get a DMCA notification, you will act on it.
 - Most importantly, check the directory of designated agents; if your company isn't listed there, complete the required form for inclusion in the directory and file it with the Copyright Office: http://www.copyright.gov/onlinesp/list/a_agents.html.

Regardless of the path you choose, it is helpful to be able to simply and rapidly prove that you are the original author of the content. One tool that can help with that is the Wayback Machine (<https://archive.org>). This site keeps copies of websites over time. You enter a website name, click Search, and get a screen that may look a bit like Figure 12-3.

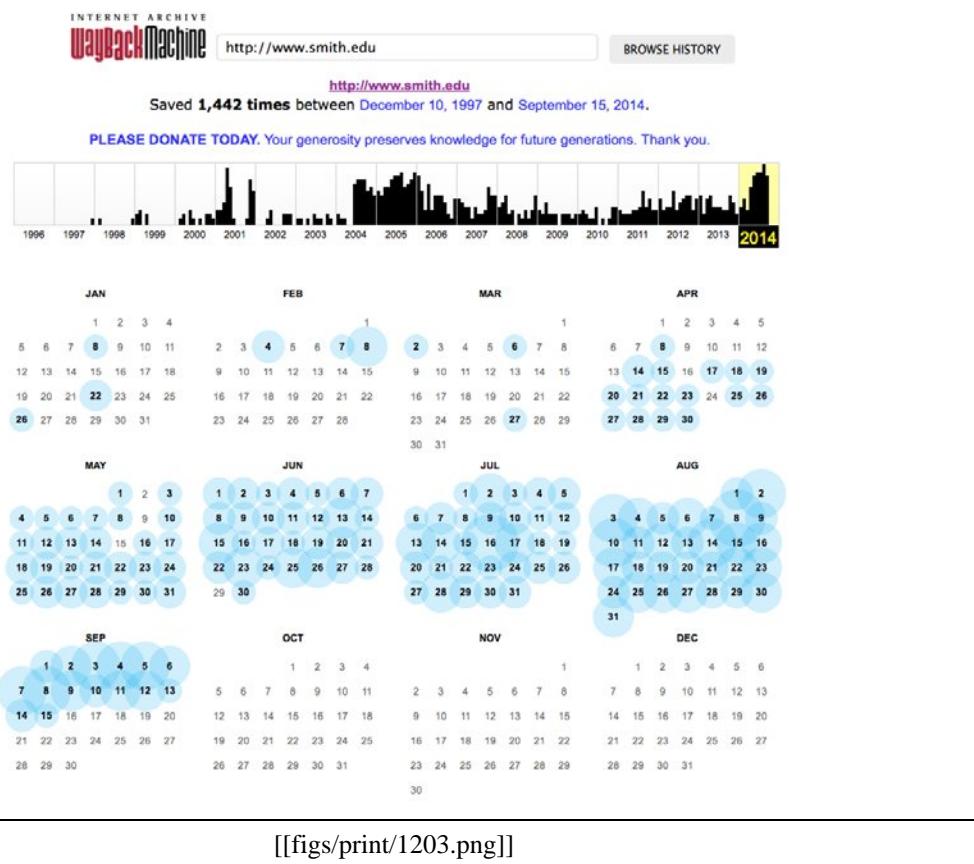


Figure 12-3. Sample archived website data from Archive.org

You can then click on one of the dates to see the state of the website on that date. The Wayback Machine does not always keep complete site copies, but a lot of data is available there, and it can often be used to provide clear proof that you were the first person to publish a given piece of content.

Changing SEO Vendors or Staff Members

Another major transition that can happen is when the publisher changes its SEO staff. This can happen both with in-house SEO staff and when you use an external SEO agency (e.g., the publisher can decide to switch agencies, or the agency can have internal changes). Companies make changes in their SEO staffs for many reasons, and issues can arise when this happens.

Potential Problems with SEO Staff Changes

One potential issue is a loss of momentum. The prior SEO team probably had some programs underway, and the new team will inherit those programs and have to sustain them, or there will be a delay before any results are seen from the new team's SEO efforts.

Even if the new team does its best to sustain an existing SEO strategy, it will take them some time to become familiar with it, and the specific tactics the old team used may not fit the new team's skill set that well. It is best to be prepared for this type of loss of momentum when making significant changes in your SEO team. One way to try to deal with this is to hire the new team and have them work for a while in parallel with the old team. This can work in some situations, but it will not be an option in others.

Another situation that can occur is that the new team may have a different philosophy than the old team (this may be why the publisher made the switch). This will invariably mean a restart of efforts, and it can potentially take many weeks to regain the lost momentum.

Documenting SEO Actions and Progress

One way to limit exposure to changes in the SEO team is through detailed documentation. Keeping a detailed record of what has been done, when it was done, and why it was invaluable in those scenarios. Ideally, this should be combined with your website analytics data.

Traditional web analytic tools make it easy to go back in time and see traffic data over time, but tools such as Google and Bing Webmaster Tools don't; nor do the search engines provide historical data on things such as indexed pages, link counts, and the like. You should also keep a record of these things over time. If your site meets the data inclusion criteria for services like SEM Rush (<http://www.semrush.com>), you may be able to obtain historical ranking data going back up to two years for your site, which can help you understand past organic trends. Archive GWT data every 90 days.

Make sure you have ownership of all relevant website accounts, usernames, and passwords. You don't want any transition to leave you in a situation where you can't access your critical data.

SEO Documentation for Rapid Training

Make sure that you have documented company processes regarding SEO, and that people associated with the SEO team and/or the management of the SEO team are all familiar with those processes. This can help speed up transitions as well, as the new team can be trained on what those processes are. Part of this includes having a detailed outline of the tasks and opportunities that are the first priorities for the new team.

With this information in hand, you can get things running smoothly much faster after a team reorganization. You can start the new team with one or more training sessions to get them up to speed on campaign history, login information, and what needs to be done going forward.

The better prepared you are to get new SEO staff up to speed and moving forward, the better off you will be. Bear in mind that you have hired the new team for their expertise; while the initial training sessions may result in the creation of a modified plan, you will get there much faster if you bring all these issues to the table on day one.

Cleaning Up

Once your migration to new SEO staff (or a new SEO vendor) is complete, you should revisit and remove the old team's access to your analytics and Webmaster Tools systems. This is one good reason to have universal, or "anonymous" business user and email accounts associated with important analytics and SEO logins. For example, tie your Google Analytics to a Google account with an email address like mybusinessanalytics@gmail.com or analytics@mywebsite.com, as opposed to the personal email account of someone who may or may not be with your company a year or two from now. Each user can be verified individually and removed individually. Audit this list on a regular basis.

Conclusion

Large-scale website changes to already established site properties are a fact of life, and they should be handled with as much foresight and planning as possible. In SEO, as in life, the only constant is change—and at the end of the day, we are grateful for this - as it means there is always opportunity for increased search exposure and for new businesses to enter the playing field, and to succeed.

13

SEO Education & Research

Learning everything you need to know about SEO in order to be effective today, and in the months and years to come, is a difficult process that is made even more challenging and complex by the constant evolution of web technologies, search engine evolution, searcher behavior, and the ever-evolving landscape of the web.

Fortunately, many resources that make the job easier are available online. Leading SEO pros regularly post their thoughts on social media, blogs, and forums, and many speak at various industry conferences. These platforms create numerous opportunities to interact with folks doing SEO at a high level, to learn from their experience, and to keep your optimization skills sharp and up-to-date.

In this chapter we will talk about how you can leverage the many resources available to be more successful in your SEO efforts of today, and in the future.

SEO Research and Search Performance Analysis

It should be fairly evident to you at this stage that SEO is always evolving. Search engines are constantly changing their algorithms (Penguin, Panda, Hummingbird, oh my!), and new media and technologies are being introduced to the web on a regular basis. Staying current requires an ongoing commitment to research, study, and participation in the process of doing SEO.

SEO Resources

One of the easiest ways to research what is happening in the world of SEO is to study the websites and periodicals that cover SEO in detail; simultaneously, ongoing testing of your (and others') SEO hypotheses should also play a major role in your day-to-day work if you hope to be, and remain, competitive.

Websites

A large number of online sites cover the search marketing space. Here is a short list of some of the better-known ones:

- Search Engine Land (<http://www.searchengineland.com>), owned and operated by Third Door Media
- Search Engine Journal (<http://www.searchenginejournal.com>), owned and operated by Alpha Brand Media
- Search Engine Watch (<http://www.searchenginewatch.com>), owned and operated by Incisive Media
- Moz (formerly SEOMoz) (<http://moz.com>), owned and operated by SEOMoz, Inc.

Each of these sites publishes columns on a daily basis, usually with multiple posts per day. The columns are typically written by industry experts who have been chosen for their ability to communicate information of value to their reader bases. Moz.com also provides a wide range of tools and resources for SEO practitioners.

Commentary from search engine employees

Search engine representatives sometimes actively participate in forums, or publish blog posts and/or videos designed for webmasters. The main blogs for the two major search engines at the time of this writing are:

- Google Webmaster Central Blog (<http://googlewebmastercentral.blogspot.com>)
- Bing Search Blog (<https://blogs.bing.com>)
- Google Webmaster YouTube channel

The search engines use these blogs to communicate official policy, announce new things, and provide webmasters with useful tips. You can reach Google personnel via the Google Webmaster Help group in Google Groups (https://groups.google.com/forum/#!forum/Google_Webmaster_Help). Members of the Google web spam team are active in this group, answering questions and even starting their own new threads from time to time. Google's YouTube channel also holds weekly Hangouts Office Hours, where people around the world can ask questions directly to Googlers.

You can also interact with search engine representatives in various forums, such as WebmasterWorld (<http://www.webmasterworld.com>), and in social media, often via their individual Twitter accounts. We will discuss the value of forums in more detail in “The SEO Industry on the Web”.

Interpreting commentary

Search engine representatives are “managed” by their corporate communications departments, and some aren’t even allowed to go on the record with statements about how their search engines function. A rare few search engine reps have free reign (e.g., Google’s Matt Cutts). Often they can’t be very specific or they can’t answer questions at all. The algorithms the search engines use are highly proprietary, and they need to be guarded as extremely valuable business assets.

This means there are certain types of questions they won’t answer, such as “What do I have to do to move from position 3 to position 1 on a particular search?,” or “How come this spammy site ranks so much higher than mine?”

In addition, the search engines have their own motives and goals. They will want to reduce the amount of spam in their search indexes and on the Web overall (which is a good thing), but this may lead them to take positions on certain topics based on those goals.

As an example, Google does not outline specifically its methods for detecting paid links, it has become clear that its ability to detect them has increased tremendously – and the updates to their Penguin algorithm over the past couple of years certainly proves this point. Taking this position is, in itself, a spam-fighting tactic, since it may discourage people who otherwise might have chosen to do so from buying links.

In spite of these limitations, you can gather a lot of useful data from interacting with search engine representatives.

SEO Testing

SEO is both an art and a science. As with any scientific discipline, it requires rigorous testing of hypotheses. The results need to be reproducible, and you have to take an experimental approach so as not to modify too many variables at once. Otherwise, you will not be able to tell which changes were responsible for specific results, as there is a difference between correlation, and causation.

And although you can gain a tremendous amount of knowledge of SEO best practices, latest trends, and tactics online, it is hard to separate the wheat from the chaff and to know with any degree of certainty that an SEO-related tactic will provide benefit. That’s where the testing of your SEO efforts comes in: in proving what works and what doesn’t.

Unlike multivariate testing for optimizing conversion rates where many experiments can be run in parallel, SEO testing requires a serial approach. Everything must filter through the search engines before the impact can be gauged. This is made more difficult by the fact that there's a lag between when you make the changes and when the revised pages are spidered, and another lag before the spidered content makes it into the index and onto the search engine results pages (SERPs). On top of that, the results delivered depend on various user-specific pieces of information including the user's search history and location, the Google data center accessed, and other variables that you cannot hold constant.

Sample experimental approach

Let's imagine you have a product page with high SERP exposure in Google for a specific search term, and you want to improve the positioning and resultant traffic even further; rather than applying a number of different SEO tactics at once, start testing specific tactics one at a time:

1. Modify just the HTML title tag and see what happens. Depending on your site's crawl frequency, you will need a few days to a few weeks to give Google and Bing enough time to recognize what you have done, and respond.
2. Continue making further revisions to the HTML title tag in multiple iterations until your search engine results show that the tag truly is optimal.
3. Move on to addressing use of heading tags (they do provide emphasis, when and where used), tweaking that and nothing else.
4. Watch what happens. Optimize it in multiple iterations.
5. Move on to the intro copy, then the breadcrumb navigation, and so on.

You can test many different elements in this scenario, such as:

- HTML title tag
- Headline (`<h1>`, `<h2>`, `<h3>`, ...) tags
- Placement of body copy in the HTML
- Presence of keywords in the body copy
- Keyword prominence
- Keyword repetition
- Anchor text of internal links to that page
- Anchor text of inbound links to that page from sites over which you have influence

Testing should be iterative and ongoing, not just a "one-off" in which you give it your best shot and never revisit the issue. If you're testing title tags, continue trying different things to see what works best. Shorten them; lengthen them; move words around; substitute words with synonyms. If all else fails, you can always put a title tag back to the way it was.

When doing iterative testing, it's good to do what you can to speed up the spidering and indexation so that you don't have to wait so long between iterations to see the impact.

You can do this by flowing more link authority to the pages you want to test. That means linking to them from higher in the site tree (e.g., from the home page). Be sure to do this a while before forming your baseline, though, because you will want the impact of changing the internal links to show in the search engines before you initiate your test (to prevent the two changes from interacting).

Alternatively, you can use the XML Sitemaps protocol to set a priority for each page, from 0.0 to 1.0. Dial up the priority to 1.0 to increase the frequency with which your test pages will be spidered.

Don't make the mistake of setting *all* your pages to 1.0; if you do, none of your pages will be differentiated from each other in priority, and thus none will get preferential treatment from Googlebot.

Throughout your SEO efforts, always remember that geolocation and personalization factors mean that not every searcher is seeing the same search results – therefore you shouldn't rely solely on specific rankings as the bellwether regarding the impact of your SEO tactics.

Other useful SEO metrics

As we discussed in Chapter 9, many other meaningful SEO metrics exist, including:

- Organic traffic to the page
- Search engine spider activity
- Search terms driving traffic per page (to the extent you can determine keyword referrer data)
- Number and percentage of URLs yielding search traffic
- Searchers delivered per search term (again, Google's blocking of keyword referrer data can make this difficult)
- Ratio of branded to non-branded search terms
- Unique pages spidered
- Unique pages indexed
- Ratio of pages spidered to pages indexed
- Conversion rates
- And many others

An effective testing regimen also requires a platform that is conducive to performing rapid-fire iterative tests, in which each test can be associated with reporting based on these new metrics. Such a platform comes in very handy with experiments that are difficult to conduct under normal circumstances.

Testing a category name revision applied site-wide is harder than, say, testing a title tag revision applied to a single page. Consider a scenario in which you're asked to make a business case for changing the category name "kitchen electrics" to the more targeted, search engine-optimal "kitchen small appliances" or "small kitchen appliances." Conducting the test to quantify the value would require applying the change to every occurrence of "kitchen electrics" across the website—a tall order indeed, unless you can conduct the test as a simple search-and-replace operation, which you can do by applying it through a proxy server platform.

By acting as a middleman between the web server and the spider, a proxy server can facilitate useful tests that normally would be invasive on the ecommerce platform and time-intensive for the IT team to implement.

During the proxying process, you can replace not only words, but also HTML, site navigation elements, Flash, JavaScript, iframes, and even HTTP headers—almost anything, in fact. You also can do some worthwhile side-by-side comparison tests: a champion/challenger sort of model that compares the proxy site to the native website.

Start with a hypothesis

A sound experiment always starts with a hypothesis. For example, if a page isn't performing well in the SERPs and it's an important product category for you, you might hypothesize that it's underperforming because it's not well linked-to from within your site. Or you may conclude that the page isn't ranking well because it is targeting unpopular keywords, or because it doesn't have enough copy.

In the case of the first hypothesis, you could try these steps to test your theory:

1. Add a link to that page on your site's home page.
2. Measure the effect, waiting at least a few weeks for the impact of the test to be reflected in the rankings.
3. If the rankings don't improve, formulate another hypothesis and conduct another test.

Granted, this can be a slow process if you have to wait a month for the impact of each test to be revealed, but in SEO, patience is a virtue. Reacting too soon to changes you see (or don't see) in the SERPs can lead you to false conclusions. You also need to remember that the search engines may be making changes in their algorithms at the same time, and your competitors may also be making SEO changes simultaneously – so be aware that the testing environment is more dynamic than one using a strict scientific method with controls.

Analysis of Top-Ranking Sites and Pages

There are many reasons to analyze top-ranking sites, and particularly those that rank at the top in your market space. They may be your competitors' sites—which is reason enough to explore what they are doing—but even if they are not, it can be very helpful to understand the types of things they are doing and how those things may have helped them get their top rankings. With this information in hand you will be better informed as you decide how to put together the strategy for your site.

Let's start by reviewing a number of metrics of interest and how to get them:

1. Start with a simple business analysis to see how a particular company's business overlaps with yours and with other top-ranking businesses in your market space. It is good to know who is competing directly and who is competing only indirectly.
2. Find out when the website was launched. This can be helpful in evaluating the site's momentum. Determining the domain age is easy; you can do it by checking the domain's Whois records. Obtaining the age of the site, however, can be trickier. You can use the Wayback Machine (<https://archive.org>) to get an idea of when a site was launched (or at least when it had enough exposure that Archive.org started tracking it).
3. Determine the number of Google results for a search for the site's domain name (including the extension) for the past six months, excluding the domain itself. To get this information, search for *theirdomain.com -site:theirdomain.com* in Google. Then append *&as_qdr=m6* to the end of the results page URL and reload the page (only works with Google Instant).
4. Find out from Google Blog Search how many posts have appeared about the site in the past month. To do this, search for the domain in Google Blog Search, then append *&as_qdr=m1* to the end of the results page URL and reload the page.
5. Obtain the PageRank of the domain's home page as reported by the Google Toolbar or third party tool.
6. Analyzing the site's backlinks using an industrial-strength tool such as Moz's Open Site Explorer (<http://www.opensiteexplorer.com>), Majestic SEO (<http://www.majesticseo.com>), and Link Research Tools(<http://www.linkresearchtools.com/>) can be helpful to analyze backlink profiles. These tools provide a rich set of link data based on their own crawl of the Web, including additional critical details such as the anchor text of the links.
7. If you are able to access a paid service such as Experian's Hitwise (<http://www.experian.com/hitwise/>) or comScore (<http://www.comscore.com>), you can pull a rich set of additional data, breaking out the site's traffic by source (e.g., organic versus paid versus direct traffic versus other referrers). You can also pull information on their highest-volume search terms for both paid and organic search.
8. Determine the number of indexed pages in each of the two major search engines, using *site:theirdomain.com*.
9. Search on the company brand name at Google, restricted to the past six months (by appending *&as_qdr=m6* to the results page URL, as outlined earlier).
10. Repeat the preceding step, but for only the past three months (using *&as_qdr=m3*).
11. Perform a Google Blog Search (<http://www.google.com/blogsearch>) for the brand name using the default settings (no time frame).
12. Repeat the preceding step, but limit it to blog posts from the past month (using *&as_qdr=m1*).

Of course, this is a pretty extensive analysis to perform, but it's certainly worthwhile for the few sites that are the most important ones in your space. You might want to pick a subset of other related sites as well.

As valuable as website metrics are, brand names can sometimes provide even more insight. After all, not everyone is going to use the domain name when talking about a particular brand, nor will they all link. Thus, looking at brand mentions over the past few months can provide valuable data.

Analysis of Algorithmic Differentiation Across Engines and Search Types

Each search engine makes use of its own proprietary algorithms to crawl and index the Web. Although many of the basic elements are the same (such as links being used as votes), there are significant differences among the different engines. Here are some examples of elements that can vary in on-page SEO analysis:

- Role of user engagement measurements
- Strategies for measuring content quality
- Strength of social media signals
- Weight of title tags
- Weight of heading tags
- Weight placed on synonyms
- Value of internal link anchor text
- How internal links are weighted as votes for a page
- Duplicate content filtering methods
- And many, many more

Similarly, there are many different ways a search engine can tune its algorithm for evaluating links:

- Percentage of a page's link authority that it can use to vote for other pages
- Weight of anchor text
- Weight of text near the anchor text
- Weight of overall linking page relevance
- Weight of overall relevance of the site with the linking page
- Factoring in placement of the link on the page
- Precise treatment of `nofollow`
- Other reasons for discounting a link (paid, not relevant, etc.)
- And many, many more

A detailed understanding of the specifics of a search engine's ranking system, at least to the point of exact and sustained manipulation, is not possible. However, with determination you can uncover various aspects of how the search engines differ. One tactic for researching search engine differences is to conduct some comparative searches across the engines. For example, when we searched on *folding glass doors* in Google and Bing, we obtained the results outlined in Table 13-1. (Note that the search results shown differ daily by search engine, so the results you may see today will likely differ from what we saw at the time we conducted this search; the results also vary by searcher location.)

Table 13-1. Comparison of top five results for “folding glass doors” in Google and Bing

Google

[http://www.nanawall.com/...](http://www.nanawall.com/)

Bing

[http://www.solarinnovations.com/...](http://www.solarinnovations.com/)

[http://www.jeld-wen.com/...](http://www.jeld-wen.com/)
[http://www.milgard.com/...](http://www.milgard.com/)
[https://www.marvin.com/...](https://www.marvin.com/)
[http://www.solarinnovations.com/...](http://www.solarinnovations.com/)

[http://www.lacantinadoors.com/...](http://www.lacantinadoors.com/)
[http://www.jeld-wen.com/...](http://www.jeld-wen.com/)
[http://www.andersenwindows.com/...](http://www.andersenwindows.com/)
[http://www.panoramicdoors.com/...](http://www.panoramicdoors.com/)

There are some pretty significant differences here. For example, the search engines have only two sites in common (<http://www.andersenwindows.com> and <http://www.solarinnovations.com>). You can conduct some detailed analysis to try to identify possible factors contributing to [lacantinadoors.com](http://www.lacantinadoors.com) (<http://www.lacantinadoors.com>) showing up in Bing, but not in Google.

At the most basic level, you may find that Google appears to be weighting HTML title tags more heavily, whereas Bing seems to place greater value on keywords in the domain name. These are of course hypothetical examples of how you can begin to analyze SERPs across different engines for the same query to find differentiators across the search algorithms. However, there are a very large number of variables, so while you will be able to form theories, there is no way to ever truly be certain! Herein lies one of the persistent challenges of SEO.

The Importance of Experience

There are some commonly perceived differences among the search engines. For example, Google is still believed to place greater weight on link analysis than Bing (note that link analysis is very important to both engines). As we discussed in the second edition of this book, Google had fallen prey to much text-link forum/blog/article comment spam, with sites exploiting this tactic for higher organic positioning, and we expressed our hope that Google would improve its policing on that front. By the end of 2013, with various updates to its Penguin algorithm, Google had done just that by aggressively penalizing sites with low quality, mostly paid, and often keyword-stuffed text links.

There are also institutional biases to consider. Given its huge search market share, Google has the richest array of actual search data. The nature of the data that the different engines have available to them can be influencing factors in how they make their decisions.

Over time and with experience, you can develop a sixth sense for the SERPs so that when you look at a set of search results you will have a good grasp of the key factors in play without having to analyze dozens of them. A seasoned SEO pro, when approaching a new project, will often be able to see important elements and factors that someone who is just starting out will not be able to pick up on. This kind of expertise is similar to that of a seasoned auto mechanic who can put an ear to an engine and tell you what's wrong—that kind of knowledge doesn't come from a blog or a book, but from years of trial and error working on many different engines. So while this book is a great resource that can help you get started, remember that actually *doing* SEO is the only, and best way to build your abilities.

Competitive Analysis

Everything we discussed previously regarding analyzing top sites applies to analyzing competitors as well, and there are additional analysis methods that can help you gain a thorough understanding of how your competitors in search are implementing their SEO strategies.

Content Analysis

When examining a competing website, ask yourself the following questions:

- What types of content are currently on the site, and how much content is there? Are there articles? Videos? Images? Music? News feeds? Answering these questions will tell you a lot of things, including how your competitors view their customers and how they use content to get links. For example, they may have a series of how-to articles related to their products, or a blog, or some nifty free tools that users may like. If they do have a blog, develop a sense as to what they write about. Also, see whether the content they are developing is noncommercial or simply a thinly disguised ad.

- How rapidly is that content changing? Publishers with rapidly changing websites are actively investing in their websites, whereas those who are not adding new articles or updating content may not be.
- Are they collecting user-generated content? Sites that gather a meaningful amount of user-generated content tend to have an engaged user audience.
- Are they trying to generate sign-ups or conversions in a direct way with their content? Or is the content more editorial in its tone and structure? This can give you more insight into the way their marketing and SEO strategies are put together.

Internal Link Structure and Site Architecture

The organization and internal linking structure of your competitors' sites can indicate their priorities. Content linked to from the home page is typically important. For example, the great majority of websites have a hierarchy in which the major subsections of the site are linked to from the home page, and perhaps also from global navigation that appears on all or most of the pages on the site.

But what else is linked to from the home page? If the competitor is SEO-savvy, this could be a clue to something that they are focusing on. Alternatively, they may have discovered that traffic to a given piece of content has a high conversion rate. Either way, these types of information can be helpful in understanding a competitor's strengths and weaknesses.

External Link Attraction Analysis

You can extract a tremendous amount of information via link analysis. A thorough review involves looking at both anchor text and link authority, and requires an advanced tool such as Open Site Explorer, Majestic SEO, and Link Research Tools. Blekko also provides some useful SEO tools (<http://www.webpronews.com/blekko-launches-new-suite-of-seo-tools-2012-10>), as do SEM Rush (<http://www.semrush.com>) and Searchmetrics (<http://searchmetrics.com/>).

Conducting an external link analysis of your competitors allows you to do many things:

Dig deeper

What pages on the site are attracting the most links? These are likely the most important pages on the site, particularly if those pages also rank well for competitive search terms.

Determine where their content is getting its links

This can help you develop your own content and link development strategies.

Analyze anchor text

Do they have an unusually large number of people linking to them using highly optimized anchor text? If so, this could be a clue that they are engaged in broad scale link spam, such as paid article/blog commenting strictly for text-rich keyword linking.

If your analysis shows that they are *not* buying links, take a deeper look to see how they are getting optimized anchor text.

Determine whether they build manual links

Manual links come in many forms, so look to see whether they appear in lots of directories, or whether they are exchanging a lot of links.

Determine whether they are using a direct or indirect approach

Try to see whether they are using indirect approaches such as PR campaigns or social media campaigns to build links. Some sites do well building their businesses just through basic PR. If there are no signs that they are doing a significant amount of PR or social media-based link building, it could mean they are aggressively reaching out to potential linkers by contacting them directly.

If they are using social media campaigns, try to figure out what content is working for them and what content is not.

Determine whether they are engaging in incentive-based link building

Are they offering incentives in return for links, such as award programs or membership badges? Do those programs appear to be successful?

Determine their overall link-building focus

Break down the data some more and see whether you can determine where they are focusing their link-building efforts. Are they promoting viral videos or implementing a news feed and reaching influencers through Google News and Bing News?

What Is Their SEO Strategy?

Wrap up your competitive analysis by figuring out what your competitors' SEO strategies are. First, do they appear to be SEO-savvy? If they're not well versed and up to date with SEO, you may identify the following site characteristics:

- Poorly optimized anchor text (if the majority of the site's in-content links say "Read more" or "Click here," it is possible there is no extensive SEO expertise behind the scenes; research further).
- Content buried many levels deep. As you know from Chapter 6, flat site architecture is a must in SEO.
- Critical content made inaccessible to spiders, such as content behind forms or content that can be reached only through JavaScript.

In general, if you see sites with obvious errors that do not follow best or even good practices, the business is either not very SEO-savvy or they are continuing to employ developers and use platforms that do not enable them to implement effective, SEO-friendly site development. These competitors may still be dangerous if they have a very strong brand, so you can't ignore them. In addition, they can fix a lack of SEO expertise by simply hiring a smart SEO expert and getting key stakeholder buy-in to make needed changes.

If, however, you are looking at a competitor who is very strong on the SEO front, you will want to pay even closer attention to their efforts.

Competitive Analysis Summary

If your competition is implementing a strategy and it is working for them, consider using it as an effective model for your own efforts – or, simply choose to focus on an area that appears to be a weak spot in their strategy, and/or one that leverages your current resources most effectively.

For example, if your competitor is not focusing on a segment of the market that has the potential to provide quality links through content development efforts, you can stake out that portion of the market before they do to gain a competitive edge.

In addition, if your competitors have previously bought, or are still buying links, reporting them to the search engines is an option – but this is something you will have to make in consideration of your organization's business ethics. Their spam tactics will eventually be discovered - so it is just as well to focus your energies on building your own business.

Ultimately, you need to make some decisions about your own SEO strategy, and having a detailed understanding of what your competitors are doing can be invaluable in that process.

Using Competitive Link Analysis Tools

It can be frustrating to see a competitor shoot ahead in the rankings and have no idea how they've achieved their success. The link analysis tools we have mentioned can help you reverse engineer their tactics by letting you quickly sort through the links pointing to a given site/page.

For example, if you were curious about how GoAnimate.com (<http://GoAnimate.com>), (a site providing business video creation tools) had earned its links, you could peek inside an advanced Open Site Explorer report and see that according to Moz there are 14,624 links to the domain from 2,332 domains, as shown in Figure 13-1.

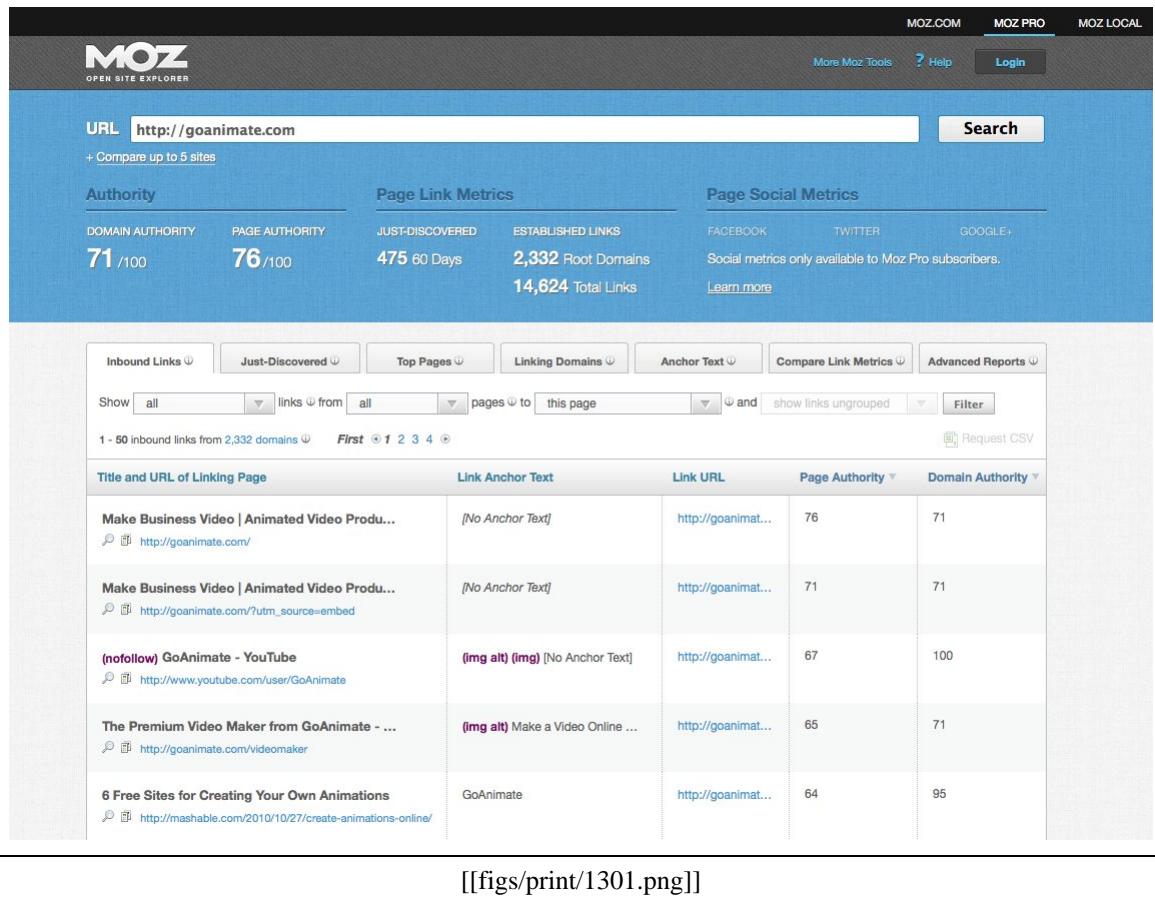


Figure 13-1. Exploring your competitor's inbound links

Looking at the data in Figure 13-1, you can see links that appear to come from content on the GoAnimate YouTube channel, as well as from tech article websites like Mashable.com.

This type of information will help your organization understand what the competition is up to. If the competition has great viral content, you can attempt to mimic or outdo their efforts.

And if they've simply got a small, natural backlink profile, you can be more aggressive with content strategies and direct link requests to overcome their lead. It's always an excellent idea to be prepared, and the ability to sort and filter out `NoFollows` and internal links and see where good anchor text and link authority flows from makes this process much more accessible than it is with other, less granular or expansive tools.

Competitive Analysis for Those with a Big Budget

In Chapter 5 we introduced you to Experian's Hitwise, an “online competitive intelligence service” with a price tag that runs into the tens of thousands of dollars. If you have a big budget, there are three more expensive packages you might want to consider: RioSEO (<http://www.rioseo.com>), Conductor

(<http://www.conductor.com>), Search Metrics (<http://www.searchmetrics.com>, seoClarity (<http://www.seoClarity.com>)). These are powerful enterprise-grade SEO assessment and competitive analysis solutions that grade your website in a number of areas and compare it side by side to your competitors and their scores.

Using Search Engine–Supplied SEO Tools

Google and Bing both make an active effort to communicate with webmasters and publishers and provide some very useful tools for SEO professionals, and it is imperative that you verify your site(s) with these tools to take advantage of them.

Search Engine Webmaster Tools

Although we have referred to the Webmaster Tools provided by Google and Bing throughout the book, it is worth going into these in greater depth. Using Webmaster Tools is a great way to see how the search engines perceive your site.

Setting up and using a Google WMT or Bing WMT account provides no new information about your site to the search engines, with the exception of any information you submit to them via the tools, and the basic fact that you, the site owner, have an interest in the very SEO-specific data and functionality they provide.

You can create a Webmaster Tools account with either Google or Bing quite easily. An important part of creating these accounts is verifying your ownership of the site. To verify your site in Google, they provide you with the following options (source: <https://support.google.com/webmasters/answer/35179>):

- Add a meta tag to your home page (proving that you have access to the source files). To use this method, you must be able to edit the HTML code of your site's pages.
- Upload an HTML file with the name Google specifies to your server. To use this method, you must be able to upload new files to your server.
- Verify via your domain name provider. To use this method, you must be able to sign in to your domain name provider (for example, GoDaddy.com or networksolutions.com) or hosting provider and add a new DNS record.
- Add the Google Analytics code you use to track your site. To use this option, you must be an administrator on the Google Analytics account, and the tracking code must use the asynchronous snippet.

Bing provides the following options to verify your site with their webmaster tools (source: <http://www.bing.com/webmaster/help/how-to-verify-ownership-of-your-site-afcfcfc6>):

- XML File authentication: Click BingSiteAuth.xml to save the custom XML file, which contains your customized ownership verification code, to your computer, and then upload the file to the root directory of the registered site.
- Meta tag authentication: Copy the displayed <meta> tag with your custom ownership verification code to the clipboard. Then open your registered site's default page in your web development environment editor and paste the code at the end of the <head> section. Make sure the <head> section is followed by a <body> tag. Lastly, upload the revised default page file containing the new <meta> tag to your site.
- Add CNAME record to DNS: This option requires access to your domain hosting account. Inside that account you would edit the CNAME record to hold the verification code (series of numbers and letters) we have provided you. When complete, you can verify your ownership of the site.

The intent of these tools is to provide publishers with data on how the search engines view their sites. This is incredibly valuable data that publishers can use to diagnose site problems. We recommend that all publishers leverage both of these tools on all of their websites.

In the following sections, we will take a look at both of these products.

Google Webmaster Tools (GWT)

Figure 13-2 shows the type of data you get just by looking at the opening screen once you log in to Google Webmaster Tools (<http://www.google.com/webmasters/tools>).

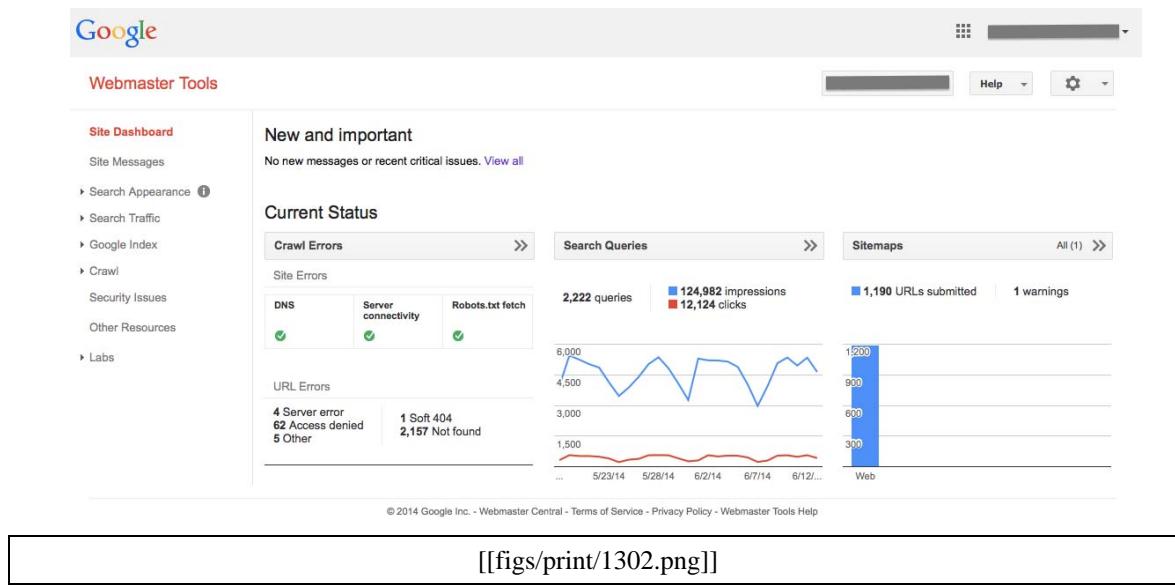


Figure 13-2. Google Webmaster Tools opening screen

Google updated their Webmaster Tools offering to provide additional information it maintains about your site, including Search Appearance. Figure 13-3 Shows Google's Search Appearance Overview (reachable by clicking on the “i” in the grey circle next to Search Appearance), which brings users to an interactive overlay with links to information and instructions for specific components relevant to this section:

Search Appearance Overview

Click on each element to see more information

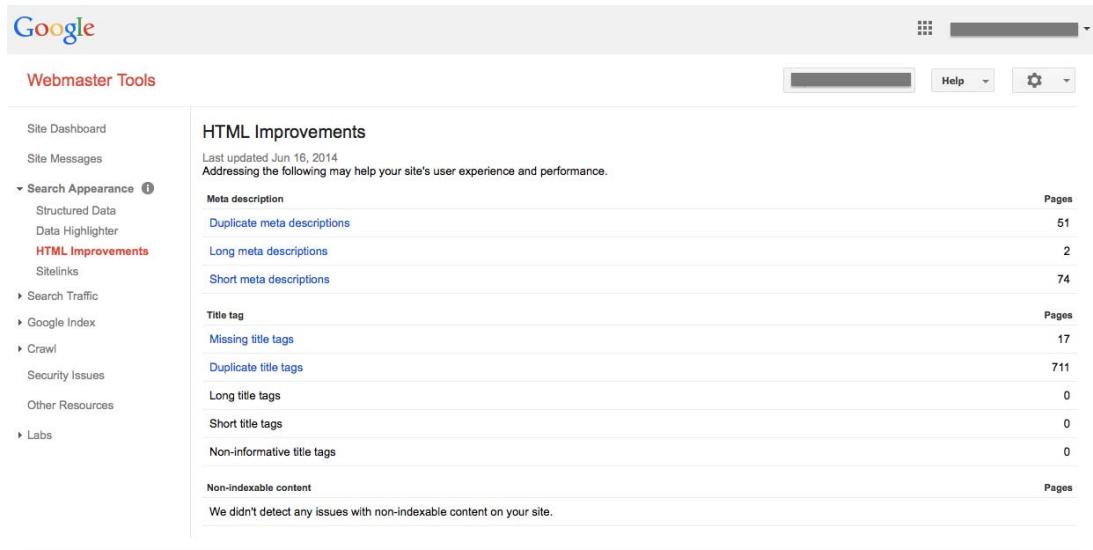
Title	Funny cat pictures with captions - example.com www.example.com/cat-captions.html	
Snippet	Find ALL the cat pictures in the world. Sort and search by type of cat. Upload your own photos and caption them too! Weekly competition for funniest cat ...	
Sitelinks	Extra grumpy cats Submit and rate pictures of extra grumpy cats ...	Local caption competition Submit the funniest caption and win a prize! ...
Search within a site	<input type="text"/> <input type="button" value="Search example.com"/>	
URL	Music gigs, concerts San Francisco Music Guide www.example.com/events/san-francisco.html	
Event - Rich Snippet	Upcoming music gigs and concerts in San Francisco. Find out what's on with our live ... Thu 11 Dec Pavement, at the Fillmore ... - The Fillmore, San Francisco Sat 13 Dec Roy Ayers at Cafe du Nord ... - Cafe Du Nord, San Francisco	
Breadcrumbs	24th century Communicator and Universal translator www.example.com › ... › Communication Devices	
Product - Rich Snippet	★★★★★ Rating: 4.5 - 11 reviews Made out of the highest quality crystalline composite of silicon, beryllium, carbon 70 and gold. Manufactured to top Starfleet standards: never get out of range of your transporter ...	
Author Information	Official Google Webmaster Central Blog: Webmaster Tools verification... googlewebmastercentral.blogspot.com/.../webmaster-tools-ver...  by John Mueller - in 21,891 Google+ circles 17 Dec 2012 – Google Webmaster Central Blog - Official news on crawling and indexing sites for the ... Posted by John Mueller, Webmaster Trends Analyst	

OK

[[figs/print/1303.png]]

Figure 13-3. Google Search Appearance Overview

Search Appearance also gives you an inside look at potential problems with your meta description tags and your title tags, as shown in Figure 13-4.



The screenshot shows the Google Webmaster Tools interface with the 'HTML Improvements' report. The left sidebar shows navigation links like Site Dashboard, Site Messages, and Search Appearance (with 'HTML Improvements' selected). The main content area displays a table of issues:

Issue Type	Pages
Meta description	51
Duplicate meta descriptions	2
Long meta descriptions	74
Short meta descriptions	
Title tag	17
Missing title tags	711
Duplicate title tags	0
Long title tags	0
Short title tags	0
Non-informative title tags	0
Non-indexable content	

At the bottom, a note states: "We didn't detect any issues with non-indexable content on your site."

© 2014 Google Inc. - Webmaster Central - Terms of Service - Privacy Policy - Webmaster Tools Help

[[figs/print/1304.png]]

Figure 13-4. Google Webmaster Tools “HTML Improvements” report

All meta description and title tag issues should be investigated to see whether there are problems that can be resolved. They may be indicating duplicate content, or pages that have different content but the same title tag. Since the HTML title tag is a strong SEO signal, this is something that should be addressed by implementing a tag that more uniquely describes the content of the page.

You can find valuable data in each of the report sections.

Another valuable section, called “Search Traffic,” shows Google search query, impression, and click data for your website (Figure 13-5) :

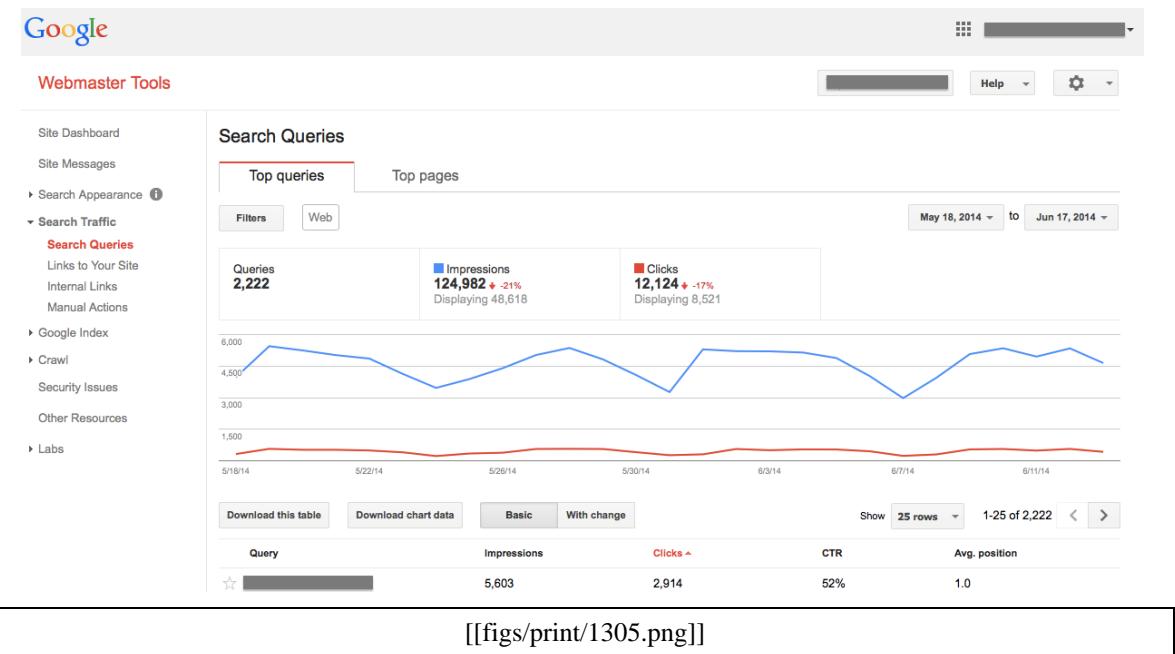
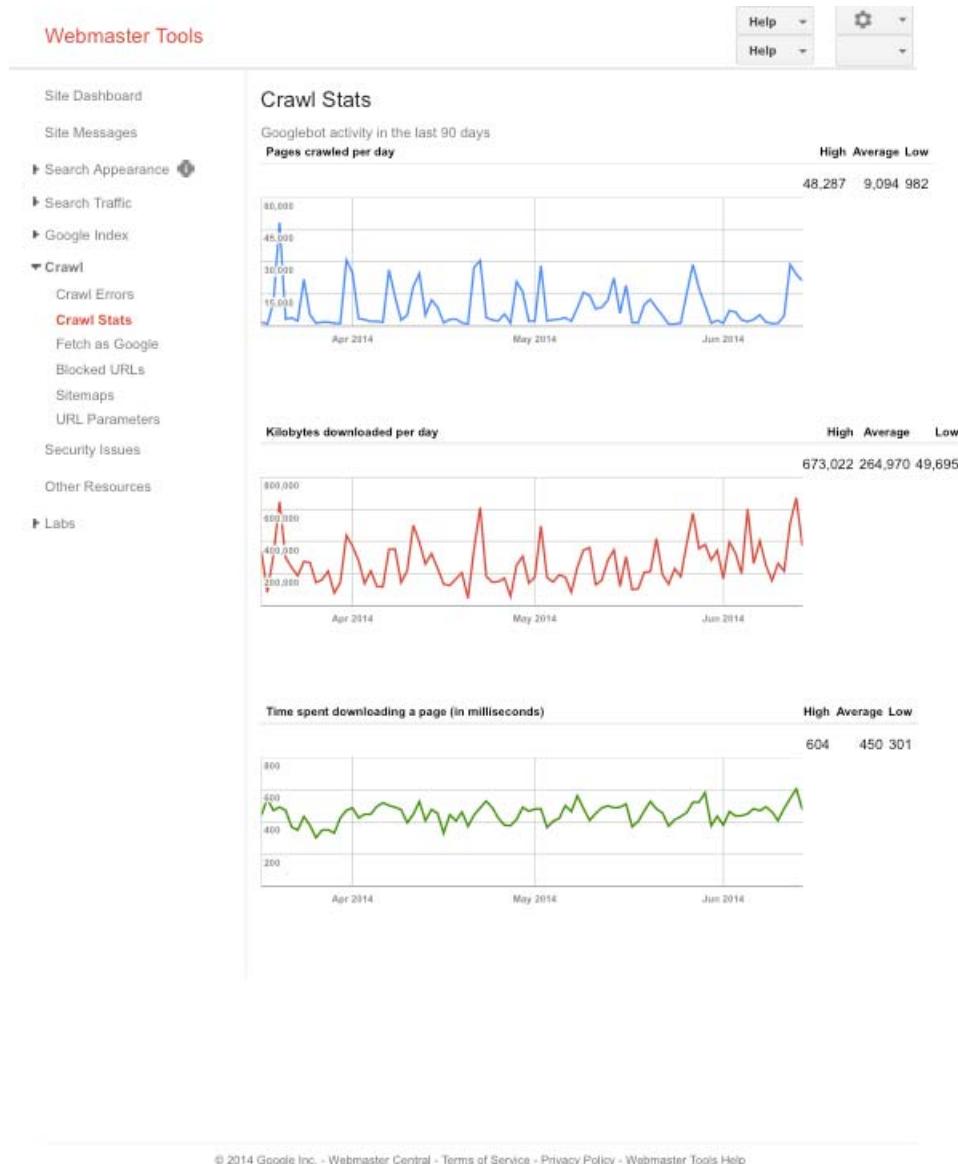


Figure 13-5. Google Webmaster Tools top search queries report

Figure 13-5 above shows Google’s view of which search queries the site is showing up for most often in the SERPs, how many impressions the site is receiving for each query, and how many clicks come from these impressions. You can also see data from Google regarding what the site’s average position in the SERPs was during the selected period for each of the search terms listed. Note that this data is pretty limited, and most publishers will be able to get better data on their search queries from web analytics software.

Another valuable reporting area in Webmaster Tools is the “Crawl” section – which enables webmasters to look at Google’s findings based on its crawling activities of the site’s URLs and robots.txt file:



[[figs/print/1306.png]]

Figure 13-6. Google Webmaster Tools “Crawl stats” report

The charts in Figure 13-6 look normal and healthy for the site. However, if you see a sudden dip that sustains itself, it could be a flag that there is a problem.

The “Crawl Errors” report is an incredibly valuable tool to identify and diagnose various site issues, from DNS resolution and server down time (Site Errors) to robots.txt and URL accessibility (URL Errors):

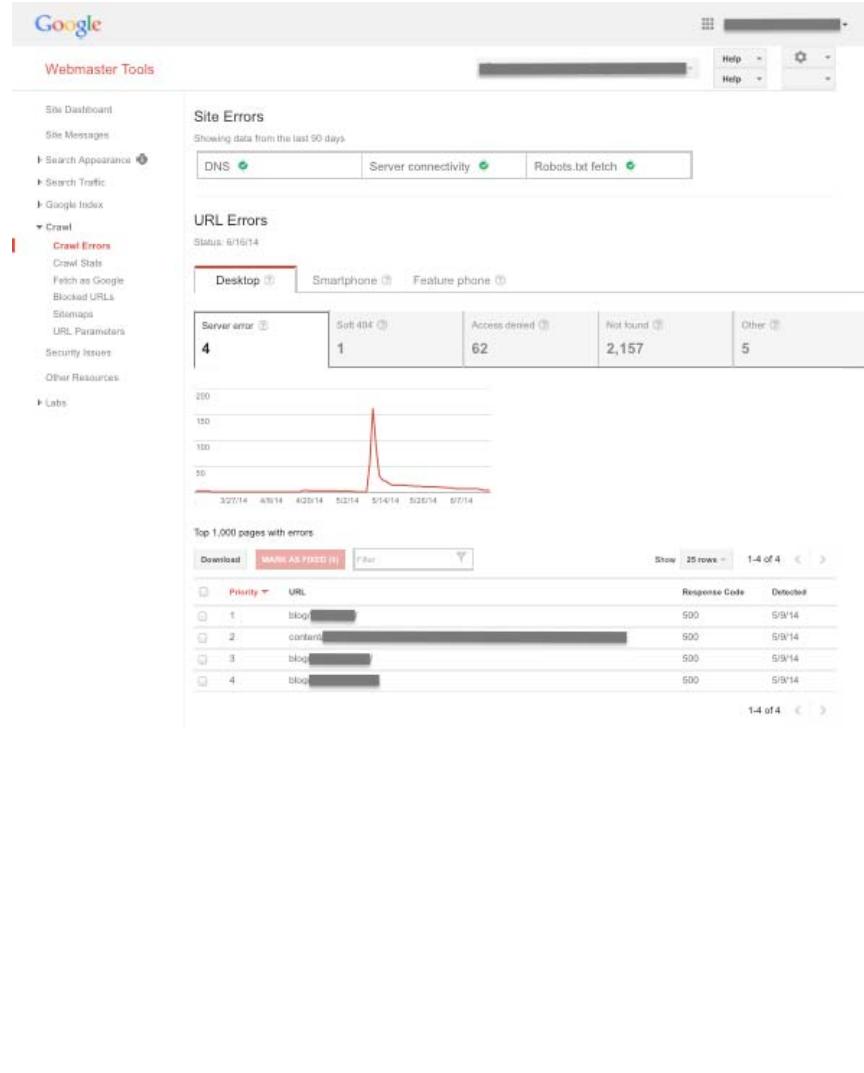


Figure 13-7. Google Webmaster Tools opening screen

Figure 13-8 shows another look at the diagnostic data available in the “Crawl” section of Google Webmaster Tools, information on “Blocked URLs.”

[[figs/print/1308.png]]

Figure 13-8. Google Webmaster Tools report on URLs restricted by robots.txt

It is common to find blocked URLs in your Webmaster Tools account. Blocked URLs can occur when sites mistakenly restrict access to URLs in their *robots.txt* file, and when that happens, this report can be a godsend. The report flags pages it finds that are referenced on the Web but that Googlebot is not allowed to crawl. Other times, blocking specific URLs is intentional, in which case there is no problem. But when your intention was not to block crawler access to your pages, this report can alert you to the problem and the need to fix it.

Another important diagnostic tactic is to look at Google's handling of your Sitemap file from within the Crawl section, as shown in Figure 13-9 below:

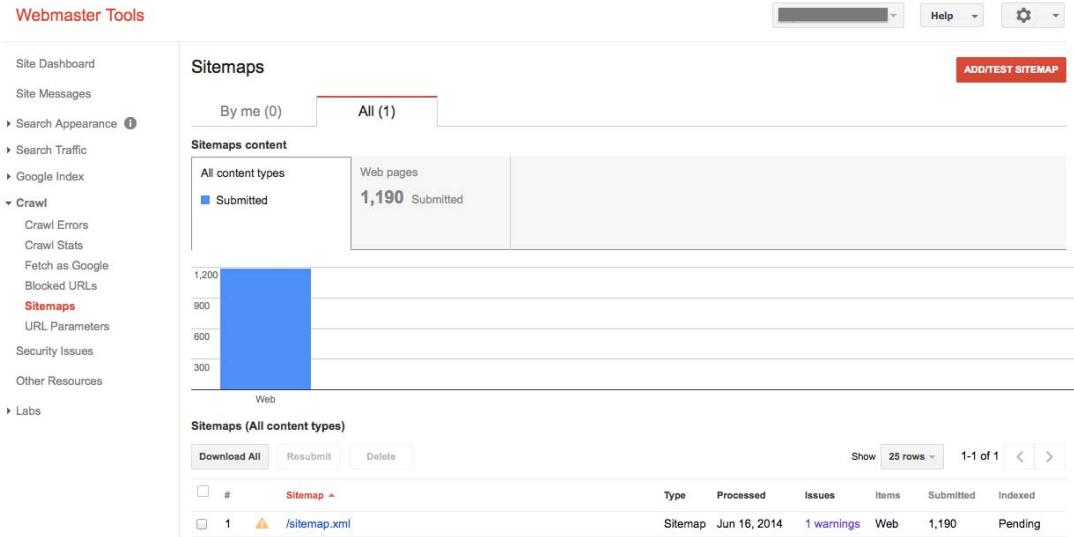


Figure 13-9. Google Webmaster Tools Sitemaps report

With this data, you can analyze the nature of any Sitemap problems identified and resolve them, as you don't want to have any broken links in your Sitemap file. In the example above, we can see the yellow "caution" icon at the left, and underneath the Issues column we see "1 warnings." Clicking on this brings us to detailed information about the warning, which in this instance appears to show that there is an unreachable URL within the Sitemap:

Error details: 0 Errors, 1 Warnings.							
Show:		All	Errors	Warnings	Show 25 rows < > 1-1 of 1		
#	Type	Issue	Description	Issues count	Example	Line	Detected
1	Warnings	URLs unreachable	When we tested a sample of the URLs from your Sitemap, we found that some of the URLs were unreachable. Please check your webserver for possible misconfiguration, as these errors may be caused by a server error (such as a 5xx error) or a network error between Googlebot and your server. All reachable URLs will still be submitted.	1	URL: [REDACTED] -		Jan 10, 2014

© 2014 Google Inc. - Webmaster Central - Terms of Service - Privacy Policy - Webmaster Tools Help

[[figs/print/1310.png]]

Figure 13-10. Google Sitemaps Errors

As Google continues to update its Webmaster Tools offerings, it is important to familiarize yourself with the options available to you as you perform SEO and manage your site(s). This information can be found in the "Settings" section of Google Webmaster Tools, and below is an outline of what the various settings mean:

Geographic target

If a given site targets users in a particular country, webmasters can provide Google with that information, and Google may emphasize that site more in queries from that country and less in others.

Preferred domain

The preferred domain is the domain the webmaster wants to be used to index the site's pages. If a webmaster specifies a preferred domain as `http://www.example.com` and Google finds a link to that site that is formatted as `http://example.com`, Google will treat that link as though it were pointing at `http://www.example.com`.

Crawl rate

The crawl rate affects the speed of Googlebot's requests during the crawl process. It has no effect on how often Googlebot crawls a given site.

robots.txt test tool

The *robots.txt* tool (not shown in Figure 13-11, but accessible at Crawl→Crawl Errors→Robots.txt fetch) is extremely valuable as well.

In addition to the services outlined above, Google Webmaster Tools has added functionality for site owners to increase their organic exposure by streamlining various processes that otherwise would be time consuming to implement client-side.

The Data Highlighter (<https://www.google.com/webmasters/tools/data-highlighter>), for example, which enables webmasters and site owners to "tag" structured data on site pages so that it can potentially appear as "rich snippets" of content within Google search results. While webmasters have always had the ability to markup structured data using various formats, with Data Highlighter Google now enables webmasters to "tag" structured data for Google's use without the webmaster having to modify the site's pages on the server. More information about Google's handling of rich snippets and structured data can be found here: <https://support.google.com/webmasters/answer/99170>.

Currently, Google supports rich snippets for the following content types:

- Reviews
- People
- Products
- Businesses and organizations
- Recipes
- Events
- Music

Figure 13-13 below shows the main Data Highlighter page within your Webmaster Tools account:

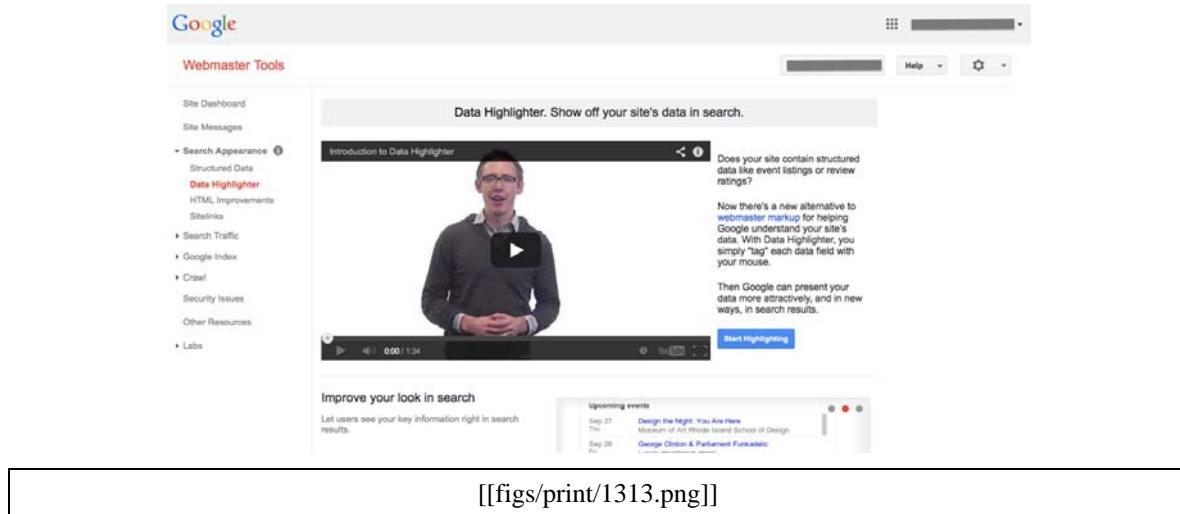


Figure 13-13. Google Data Highlighter

A click on “Start Highlighting” brings you to the entry fields to select the appropriate options for your site:

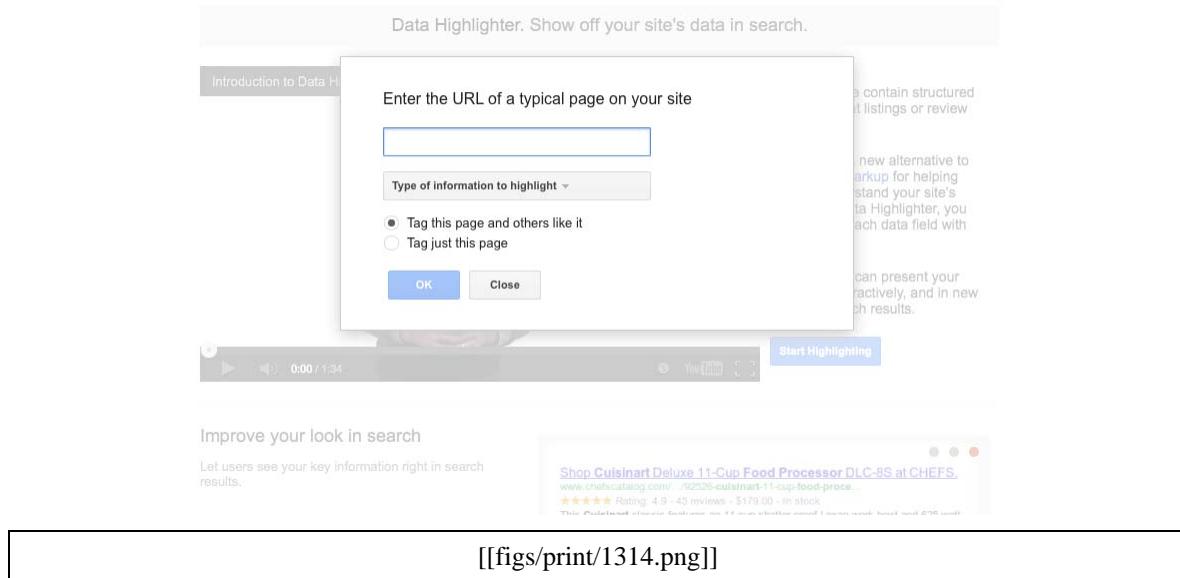


Figure 13-14. Google Data Highlighter

In Figure 13-14 above, you would select the data type to highlight, and then choose whether this data type is unique to this page (for example, if you only have one page with this type of data, but other pages have different data needing markup); conversely, if your site has thousands of product pages with the same structured data format, you can streamline the process and instruct Google to handle all other pages “like” the one you select, in the same fashion.

Manual Spam Actions

Google will also notify webmasters via the Webmaster Tools interface (and via the email notification settings specified within the account) of any manual spam actions they have taken against a site. Figure 13-15 below shows the type of message you can receive from Google if they have identified what they deem to be “a pattern of unnatural, artificial, deceptive, or manipulative links” pointing to pages on your site (<https://support.google.com/webmasters/answer/2604772>):

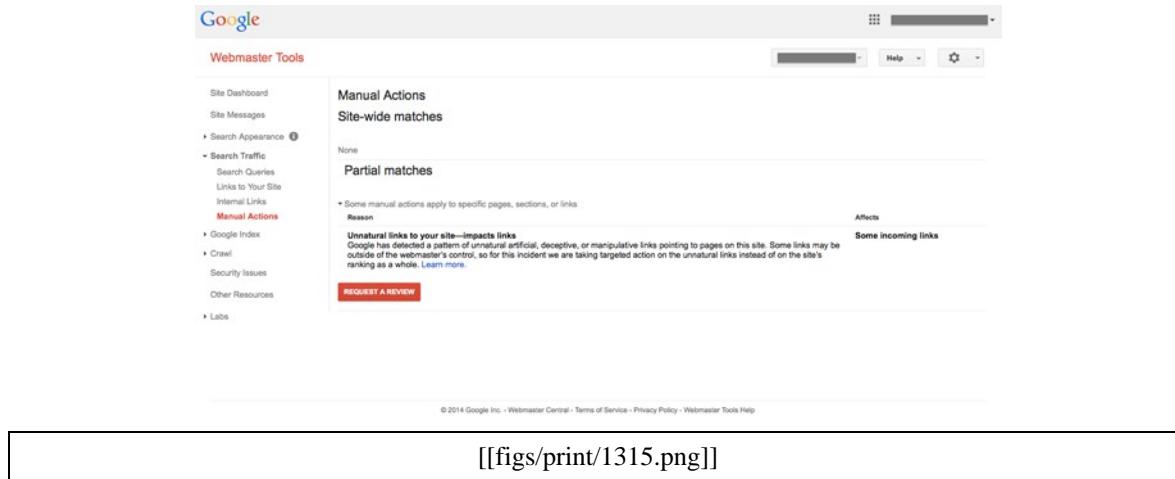


Figure 13-15 - Google Webmaster Tools opening screen

Bing Webmaster Tools (BWT)

Microsoft also offers a Webmaster Tools product with some great features. With Bing Webmaster Tools (<http://www.bing.com/toolbox/webmaster>) you can get a rich set of data on many of the same types of things that Google Webmaster Tools offers, but you are getting the feedback from the Bing crawler instead. Having the ability to see the viewpoint from a different search engine is a tremendous asset, because different search engines may see different things. Figure 13-16 shows the opening screen for Bing WMT.

The screenshot shows the Bing Webmaster Tools dashboard. On the left, a sidebar lists 'My Sites' and 'Dashboard' (which is selected). The main content area is titled 'Dashboard' and includes a 'SITE ACTIVITY' section with a table showing Clicks from Search, Appeared in Search, Pages Crawled, Crawl Errors, and Pages Indexed. Below this is a 'SITEMAPS' section with a table showing the status of submitted sitemaps. At the bottom, there are links for 'Privacy and Cookies', 'Legal', 'Advertise', 'Help', 'Feedback', and 'Community'.

Figure 13-16. Bing Webmaster Tools opening screen

Already you can see some of the great things that are available. At the top you can see Clicks from Search, Appeared in Search (Impressions), Pages Crawled, Crawl Errors, and Pages Indexed.

To see Crawl Errors identified by Bingbot (Bing's search crawler), you would click on *Reports & Data* → *Crawl Information*:

The screenshot shows the 'Crawl Information' page. The sidebar includes 'Configure My Site' (selected), 'Reports & Data' (selected), and 'Crawl Information'. The main content shows a table of crawl errors: 400-499 (95), 500-599 (17), 301 (4,893), 302 (1), ROBOTS.TXT EXCLUSION (15), DNS FAILURES (0), and CONNECTION TIMEOUTS (0). Below this is a table for 'HTTP 400-499 CODES' with columns for URL, HTTP CODE, and COUNT OF LINKS. The table lists five URLs, each with a 404 HTTP code and a count of 0.

Figure 13-17. Bing Webmaster Tools Pages with Crawl Errors report

You can pick from several options:

- 400-499: Request errors
- 500-599: Internal server errors
- 301: Page moved permanently (permanent redirects)
- 302: Page moved temporarily (temporary redirects)
- Robots.txt exclusions
- DNS failures
- Connection timeouts

Bing Webmaster Tools also includes an incredibly valuable tool called Index Explorer, accessible from Reports & Data → Index Explorer, and from the upper right corner of the Crawl Errors report in Figure 13-17 above. This tool provides incredible insight into your site's internal architecture as viewed by a crawler, often showing subdirectory folders that you, and your webmaster, did not know were still "live" and crawlable.

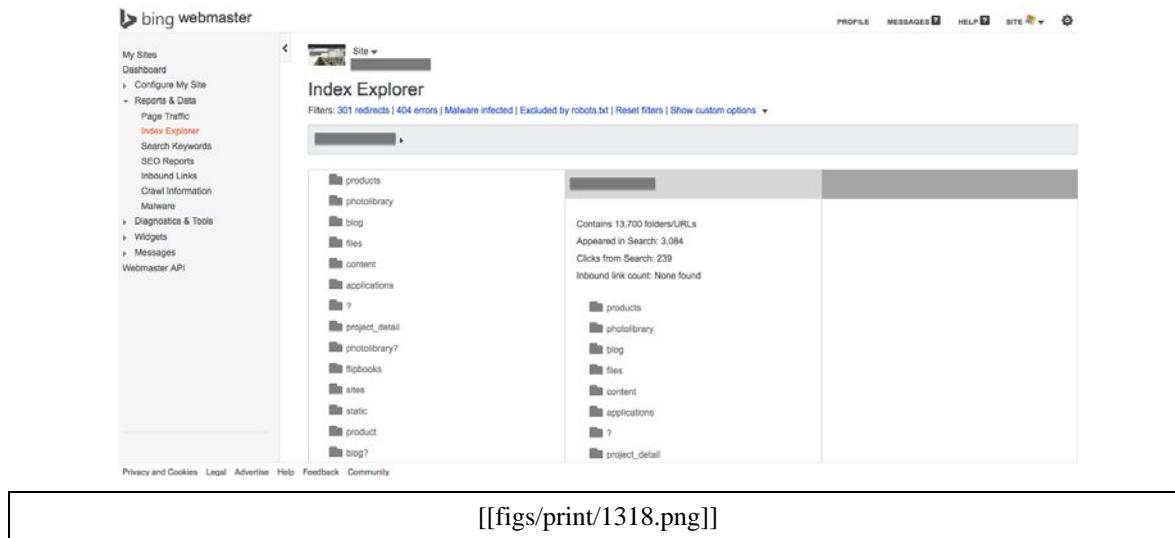


Figure 13-18. Bing Webmaster Tools Index Explorer

Bing Webmaster Tools also provides you with a section called "SEO Reports," as shown in Figure 13-19.

SEO SUGGESTIONS	SEVERITY	ERROR COUNT	PAGES
The title is missing in the head section of the page.	High	7	7
The tag does not have an ALT attribute defined.	Low	6,846	385
The page is missing meta language information.	Moderate	13	13
The title is too short or too long.	High	1	1
The <h1> tag is missing.	High	18	18
The description is missing in the head section of the page.	High	259	259
The description is too long or too short.	High	91	91

Figure 13-19. Bing Webmaster Tools SEO Reports

This section provides data similar to Google’s “HTML Improvements” report, and gives suggestions for where to improve SEO-relevant areas of your site’s pages for better performance in organic; focus in on the items with “High” severity first, as these are likely to be the most impactful in the near term.

Bing also provides a new section called Diagnostics & Tools that is a mini-toolset in and of itself, with a host of valuable SEO functions (Figure 13-20):

Figure 13-20. Bing Webmaster Tools Diagnostics & Tools

The first tool, Keyword Research, enables you to gather estimated query volume for keywords and phrases based on the Bing search engine’s historical query database. Note that if you want only query data for the *exact* keyword or phrase you enter, be sure to check the checkbox Yes for “Strict” just underneath and to the right of the keyword entry field. Checking Yes to Strict means that if you enter the phrase “feng shui santa barbara” you will only see query volumes explicit to the exact phrase, “feng shui santa barbara” and not for phrases that include this phrase, such as “feng shui consultant santa barbara.”

There is much more here than we have been able to show in just a few pages, and the possibilities are limitless for how to utilize these valuable tools from both Google and Bing to your SEO advantage.

The SEO Industry on the Web

SEO-specific social media communities, forums, and blogs are almost always abuzz with new information and insights on the topic of SEO. Following this online community closely can help you increase your SEO knowledge and keep your skills sharp.

Blogs

A significant number of people and companies in the SEO industry regularly publish high-quality content about SEO and web marketing. This provides a great resource for newcomers who want to learn, and even for experienced SEO practitioners, who can pick up tips from others. However, it should be noted that not all top-notch SEO pros monetize their tactics by publishing them, preferring instead to keep much of their expertise close to the chest.

The reality is that SEO is a deep and complex topic - and the opportunities for learning never cease. That is one of the most rewarding aspects of the profession, and fortunately, the large number of blogs about SEO and web marketing in existence make it easy to research the profession on a regular basis.

Be aware, however, that anyone can create a blog and start writing about SEO with self-appointed authority. There is a big difference between someone publishing an opinion, and what he or she publishes actually being based on empirical knowledge. Many people have been able to create consulting businesses out of “flash in the pan” service offerings rooted in opportunism, as opposed to experience and love of the mystery, dynamism, and challenge of search. At all stages of your SEO research process and lifelong learning curve, it is important to know how to separate out those who know what they are talking about from those who do not.

One way to do that is to see which bloggers have the respect of other senior people in the industry. This is not unlike the process search engines go through to identify trusted authority sites and see who they recommend and trust. Proximity to clearly authoritative figures is a big plus. (This is why we always advocate a backlink profile check before hiring a firm or consultant. If the majority of links to them are low-quality text link spam, you have a red flag.)

Another way to develop a voice is to become a contributor to these communities. As you accumulate SEO experiences, share the ones you can. Comment in blogs, forums, and social media and become known to the community. If you contribute, the community will reward you with help when you need it.

The search technology publications we discussed at the beginning of this chapter are a great place to start. However, a number of other SEO blogs provide a great deal of useful information. Here are some of the best, listed alphabetically.

First up are blogs by the authors:

Eric Enge: Ramblings About SEO (<http://www.stonetemple.com/blog>)

Stephan Spencer: Scatterings (<http://www.stephanspencer.com/blog>)

Jessie Stricchiola: Alchemist Media Blog (<https://www.alchemistmedia.com/blog>)

And here is a list of blogs by the search engines:

- Bing Search Blog (<http://blogs.bing.com/search/>)
- Google Webmaster Central Blog (<http://googlewebmastercentral.blogspot.com/>)

Although it is technically considered “unofficial,” you should also make Matt Cutts’ blog a must-read on your list:

- Matt Cutts’ blog (<http://www.mattcutts.com/blog>)

One of Google’s search engineers and the head of Google’s web spam team, Matt is also Google’s official representative to the SEO and webmaster world.

Here are some other highly respected blogs:

Search Engine Roundtable (<http://www.seroundtable.com>)

Barry Schwartz's roundup of all things search-related

Rae Hoffman: (<http://sugarrae.com/all>)

Affiliate marketer and SEO pro Rae Hoffman's writings

Annie Cushing: (<http://www.annielistics.com/>)

Analytics expert Annie Cushing's writings on all things analytics, with SEO-specific information

Dave Naylor (<http://www.davidnaylor.co.uk>)

Dave Naylor's coverage of a wide range of online marketing issues, with a focus on SEO

Small Business SEM (<http://www.smallbusinesssem.com>)

Matt McGee's blog focused on search marketing for small businesses

SEO Book (<http://www.seobook.com>)

Aaron Wall's accompaniment to his excellent book on SEO

You can also get a more comprehensive list by visiting Lee Odden's BIGLIST of SEO and SEM Blogs (<http://www.toprankblog.com/search-marketing-blogs/>).

SEO News Outlets, Communities, & Forums

Much information can be gleaned from news sites, communities, and forums. Reading them does require a critical eye, but you can pick up a tremendous amount of knowledge by not only reading articles and posts, but reading the differing opinions on various topics that follow the original article or post. Once you develop your own feeling for SEO the interpretation of what you read online becomes increasingly easier.

SEO News Outlets & Communities:

Search Engine Land (<https://searchengineland.com>)

Search Engine Watch (<https://searchenginewatch.com>)

Search Engine Journal (<https://searchenginejournal.com>)

Moz.com (<https://moz.com>)

Popular SEO Forums:

Search Engine Roundtable Forums (<http://www.seroundtable.com>)

These are the forums associated with Barry Schwartz's Search Engine Roundtable blog.

SEOChat (<http://seochat.com>)

WebmasterWorld Forums (<http://www.webmasterworld.com>)

Founded by Brett Tabke, Webmaster World is now owned and operated by Jim Boykin. These forums have been among the most popular places for SEO pros to discuss issues since 1996, nearly two decades!

WebProWorld Forums (<http://www.webproworld.com/webmaster-forum/forum.php>)

This is a large web development and SEO forum with a very diverse group of posters.

High Rankings Forums (<http://www.highrankings.com/forum/>)

Jill Whalen, a well known and respected SEO practitioner, is the Founder of High Rankings. As of November, 2014 Jill was still on her hiatus from SEO consulting, which she announced in 2013: <http://whatdidyoudowithjill.com/leaving-seo-post/> - however, there is still an active forum community.

In Memoriam: In the second edition of this book, you may recall that we sang the praises of longtime WebmasterWorld (<http://www.webmasterworld.com>) moderator Ted Ulle, lovingly known in the search industry as “Tedster.” Sadly, Ted passed away in 2013 – and the industry mourned: <http://searchengineland.com/mourning-ted-ulle-165466>. We miss you Ted, and we thank you for your incredible contributions to our industry over the years.

Communities in Social Networks

A significant number of SEO professionals are very active on social media, while others are not—it is a matter of personal and professional preference (and available time!). You can find this book’s authors on Facebook at the following URLs:

- Eric Enge: <https://www.facebook.com/ericenge.stc>
- Stephan Spencer: <https://www.facebook.com/stephanspencer>
- Jessie Stricchiola: <https://www.facebook.com/jessiestricchiola>

and on Twitter at the following URLs:

- Eric Enge: <https://twitter.com/stonetemple>
- Stephan Spencer: <https://twitter.com/sspencer>
- Jessie Stricchiola: <https://twitter.com/itstricchi>

The authors’ Google+ pages are located here:

<https://plus.google.com/+EricEnge>

<https://plus.google.com/+StephanSpencer>

<https://plus.google.com/+JessieStricchiola>

You can also find all of the authors on LinkedIn (<http://www.linkedin.com>):

<https://www.linkedin.com/in/ericenge>

<https://www.linkedin.com/in/stephanspencer>

<http://www.linkedin.com/in/jessiestricchiola>

Social media engagement creates opportunities to interact with and learn from other people doing SEO. And as in the forums, you can encounter search engine employees on these sites as well.

Participation in Conferences and Organizations

Conferences are another great way to meet leaders in the search space and network with them. The various panels may cover topics about which you are looking to learn more, and this by itself makes them a great resource. Even experienced SEO practitioners can learn new things from the panel sessions.

Make sure you target the panels with the following objectives in mind:

- Attend panels from which you can pick up some information, tips, and/or tricks that you need for your immediate SEO projects or that broaden your SEO knowledge.
- Attend panels where someone that you want to get to know is speaking.
- Engage during panels and presentations by tweeting your questions during the session (as available!)

If you want to meet someone who is speaking on a panel, sit in the front row, and at the end of the session introduce yourself and/or ask a question. This is an opportunity to get some free and meaningful advice from experts. Get the speaker’s business card and

follow up with her after the event. Some panels also have search engine representatives speaking on them, so this tactic can be useful in meeting them as well.

Of course, you may encounter the people you want to meet in the hallways in and about the conference, and if so, take a moment to introduce yourself. You will find that most attendees are very open and will be happy to share information and tips.

In addition to the day's sessions, each conference has a series of networking events during the day and sometimes in the evening. These events are often held in the bar of the hotel where (or near where) the conference is being held.

Each conference usually has a published party circuit, with different companies sponsoring parties every night (sometimes more than one per night). In addition, some conferences are starting to incorporate more "healthy" ways to network and connect with industry peers during the events. Third Door Media's Search Marketing Expo has spearheaded this effort by offering "Social Matworking" conference yoga at SMX San Jose in 2014, to both speakers and attendees - with yoga sessions taught by one of our co-authors, Jessie Stricchiola. Consider going to these types of industry networking events as a way to meet people and build new relationships.

Here is an alphabetical list of some recommended conferences for search-specific content:

Ad-Tech (<http://www.ad-tech.com>)

This is an interactive marketing conference that has begun to address many of the issues of SEO and search marketing.

Online Marketing Summit (<http://www.onlinemarketingsummit.com>)

Centered around Internet marketing of all kinds (but with at least one track devoted to SEO), OMS is a rising star in the conference world. One of the most valuable aspects comes from the one-on-one "lab" sessions where expert panelists sit down with attendees for 15- to 20-minute review sessions.

PubCon (<http://www.pubcon.com>)

Two to three times per year, Brett Tabke, founder and former owner of WebmasterWorld, runs conferences that attract a large group of SEO practitioners.

ClickZ Live (<http://www.clickzlive.com>)

Formerly Search Engine Strategies (SES), then SES Conference & Expo, and re-branded in 2014 to "ClickZ Live" - responsible for 8 to 12 conferences worldwide. San Francisco is the biggest event, but New York is also quite substantial.

SEMPo (<http://www.sempoinstitute.com>)

The Search Engine Marketing Professionals Organization, the largest nonprofit trade organization in the world serving the search and digital marketing industry and the professionals engaged in it, focuses on education, networking, and research. Founded in 2002 by co-author Jessie Stricchiola and search industry leaders Brett Tabke, Christine Churchill, Kevin Lee, Dana Todd, Noel McMichael, Barbara Coll, and Fredrick Marckini, SEMPO has local chapters that host events, hosts frequent webinars and Hangouts, and is an ongoing presence with networking events at industry conferences.

Third Door Media Events (<http://thirddoormedia.com>)

Third Door Media runs more than a dozen conferences worldwide, focusing primarily on search marketing (Search Marketing Expo, or SMX) but also covering other online marketing topics, such as social media (SMX SMM) and marketing technology (MarTech). Third Door Media also offers a conference specifically for advanced-level search marketers and SEO professionals called SMX Advanced, held annually in Seattle.

MozCon (<http://moz.com/mozcon>)

Operated by Moz (formerly SEOMoz), these conferences are focused on what Moz terms "inbound marketing" and cover many areas of digital marketing.

More in-depth workshops are offered by:

The Direct Marketing Association (<http://dmaeducation.org/course/get-to-know-seo-certification/>)

The DMA is the world's largest trade association for direct marketers. It provides quality instruction, including certification programs and 90-minute online courses.

Instant E-Training (<http://www.instantetraining.com>)

Instant E-Training provides training on Internet marketing topics such as SEO, Social Media, Paid Search, and Web Usability. Its training videos are bite-sized, ranging anywhere from 20 to 50 minutes in length. Instant E-Training also offers live workshops for SEO and social media multiple times per year. Instructors include experts such as Eric Enge, Shari Thurow, Christine Churchill, Carolyn Shelby, Bob Tripathi, and others.

Market Motive (<http://www.marketmotive.com>)

Market Motive offers web-based Internet marketing training and certification. Market Motive was one of the first companies to offer a credible certification program for SEO, and the faculty includes leading industry figures such as Avinash Kaushik, John Marshall, Greg Jarboe, Todd Malicoat, and others.

Conclusion

Mastering SEO requires a great deal of study and effort. Even when you achieve a basic level of effectiveness in one vertical or with one website, you have to continue your studies to stay current - and you still have many other variables and environments to consider before you will be able to point your expertise at any site and be effective. And even then, you will always be working against some business-specific variables that are out of your control, and will need to adapt and modify accordingly. Thankfully, industry resources abound in the form of news sites, blogs, forums, conferences, industry organizations, workshops, training materials, competitive analysis tools, and search engine-supplied toolsets and knowledge centers. Take advantage of these rich sources of information and, if you are so inclined, participate in the vibrant online community to stay in front of the SEO learning curve and ahead of your competition.

SEO Help: In-House, Outsourced, or Both?

The Business of SEO

SEO has its roots in the tech community. Early search engine optimizers were web developers, small business owners, and forward-thinking consultants looking to make their websites visible in the search engines and focused on tweaking HTML to make their sites more search engine-friendly.

As search has matured over the years and taken its place at the marketers' table as a legitimate sales and brand development channel, SEO pros have increasingly been asked to focus on the business side of the practice—to evolve beyond the code and Microsoft Excel jockeys of the early days of the practice (though let's be clear, Excel still serves some very valuable functions in SEO work!).

The benefits of focusing on SEO as a business are:

- It legitimizes the practice of SEO. Talking about canonical tags is great for chatting at an SEO conference, but talking about organic search share will go over better in the C-suite. In many companies, winning the C-suite battle to obtain budget and functional support is the only way to really succeed long-term in SEO.
- It bridges the gap between the practice of SEO and the impact on the business. Focusing on the business aspects of SEO will enable you to be clear about the ROI you are getting from organic search. Translating SEO-centric metrics such as rankings and traffic to revenue and dollars earned enables you to tie a tangible business impact to your efforts.

Understand Your Market Opportunity

Most SEO consultants probably cannot say with certainty what ranking in the top search positions for their keywords is (or would be) worth to them in terms of traffic and revenue. However, you should never set out to develop a new product or launch a new service without first understanding the business case/market opportunity, and organic search should be no different. Understanding the market opportunity prior to editing even the first `<h1>` tag is important, because it will help you to:

- **Understand what SEO is potentially worth to your organization.** It will give you an idea of what organic search visibility is worth to your business and answer questions like “How much organic traffic could we potentially procure from higher organic positioning?” Answers to these questions will guide the SEO process, including budget, and infrastructure investments, and will help you make the SEO business case to the rest of the organization.

- **Track progress as you go.** Building a market opportunity benchmark prior to embarking on SEO will enable you to measure progress as you go and demonstrate the ROI of SEO in real, quantifiable metrics that will resonate with the company stakeholders.

A market opportunity assessment in organic search shows the current visitors captured for your keywords, versus the potential opportunity were you to rank in a top position. There are various tools available on the Internet that can help simplify the process.

To conduct a market opportunity assessment:

1. **Isolate your top group of non-branded keywords.** Determine your target keywords based on their perceived conversion potential. The number of keywords may vary based on factors such as your SEO maturity or vertical market space and can range from around 50 for those just starting out on up to several thousand keywords for those farther along in organic search.
2. **Gather search volumes and rankings for your keywords.** Find out what your current search volumes and rankings are for your targeted keywords.
3. **Use a click-through rate (CTR) curve to identify the potential for new visitors.** A CTR curve shows the expected percentage of clicks based on the search position of the keyword. Plugging your current search volume and rank into a CTR curve can give you a general idea of the number of new visitors you can expect to get by moving up in the search rankings.
4. **Scale that across your total traffic.** The first three steps of the analysis give you a sense of the traffic gain you can get for your top keywords, but it is likely that you have many other keywords driving traffic to your site. If you have an opportunity to double the traffic on your top keywords, you likely also have the opportunity to double your overall non-branded organic search traffic volume.
5. **Check the CPC values for your keywords.** Using a tool like SEM Rush or Search Metrics, or the Google Keyword Planner (<https://adwords.google.com/KeywordPlanner>) from within your own Google AdWords account, run some numbers on the current market value of your keyword targets.

Get Buy-in Across the Organization

Now that you've done an opportunity assessment and know what you will be working toward in organic search, it's time to internally publicize the opportunity and obtain buy-in across the organization. For SEO efforts to be successful in any organization, active participation from many parts of the business - across nearly every department - is an absolute necessity. From C-level executives approving budget, to the IT department implementing web and infrastructure changes and content writers using SEO best practices, organizational buy-in is as critical to successful SEO as information architecture and `<h1>` tags.

To obtain organizational buy-in:

1. **Get people excited about SEO—then scare the heck out of them.** Show people the opportunity in organic search that has come out of your opportunity assessment. Make them aware of what's being missed out on by not performing well in organic search. Then, make them aware of where your company stands relative to the competition, and identify competitors who are currently cashing in on the opportunities you are missing out on.
2. **Keep the SEO opportunity foremost in people's minds.** Make posters outlining the market opportunity and hang them on the walls. Speak about it often, and generally socialize it in the organization. The more you can keep the opportunity front and center in people's minds and keep them excited about it, the better.
3. **Create a service-level agreement (SLA) with stakeholders.** To formalize the working relationship with organizational departments and increase their buy-in to SEO, sign an SLA with them that defines how quickly they will respond to requests and the metrics that will be used to measure success.

Lay the Groundwork

Once you've obtained buy-in from the organization, it's time to position yourself to succeed. We've heard of many instances of chief marketing officers (CMOs) who have dipped a toe in the water with SEO but

labeled it “black magic” or ineffective because the company’s website has failed to skyrocket to the top of the search rankings for the high-volume, highly competitive keywords they’ve targeted in 90 days or less. To avoid problems like this and to lay the groundwork for SEO success in your organization, it is important to do two things:

1. **Set the appropriate expectations.** SEO is a long-term strategy where you are likely to see incremental improvements that may start slowly at first. Focus expectations on gains in non-branded search traffic and conversions from that traffic, not on rankings. Many marketers’ only exposure to search marketing is the world of pay per click, so it is important to educate them about SEO with its unique characteristics, and its differentiators.
2. **Show early wins.** The temptation may exist to go after high-volume, highly visible keywords right out of the gate, but ranking for competitive keywords can – and does – take time. It is better to initially target less competitive keywords that can result in quick wins (what we call “low hanging fruit”), while simultaneously developing a long-term strategy for more competitive keywords. This will help keep people engaged and believing in the opportunity in organic search.

To help set reasonable expectations and report on your organic search progress in the most effective way possible:

1. **Make a distinction between milestones and success metrics.** The fruits of your labor can take some time to materialize in organic search results. To maintain organizational buy-in and gauge your progress before traditional success metrics such as conversions kick in, report on key milestones such as number of indexed backlinks, changes made to the website, and recommendations made.
2. **Report on SEO metrics as business metrics.** Report on SEO like you would any sales effort, by continuously reporting on your search pipeline and the competitive landscape. Create a mission around SEO: you might strike a gong when certain milestones are hit, or install TVs where you can show regularly updated traffic and conversion stats.
3. **Be a businessperson.** Understanding when to use technical SEO tactics is important, but speaking the language of business the C-suite understands is critical to the ongoing success of SEO in most organizations. Expand your horizons beyond Excel and analytics and learn to speak the language of dollars and cents—this is what executives will notice and understand.

Motivate Resources That Don’t Share Your Goals to Help You

As mentioned above, succeeding in SEO requires support and buy-in across your organization. Many SEO efforts are dependent on non-SEO team members in the company to get things done, often relying on resources across numerous departments who may not be trained on, or oriented towards, thinking about SEO.. Finding ways to work with others to get things done is key to SEO success.

To most effectively work with resources organization-wide on SEO tasks:

1. **Show how it will benefit them.** Trying to get content creators to use targeted keywords in content can be a losing proposition. If you demonstrate conclusively, however, how it will benefit them—that they will reach a wider audience with their content—they are much more likely to buy in, implement changes to their content development process, and commit to the effort over time.
2. **Recognize contributions to organic search visibility.** Recognize when team members do something that contributes to organic search success, such as optimizing a press release or tweaking an XML Sitemap. At one Fortune 100 company, the Director of SEO handed out an award every quarter to the individual who had contributed most significantly to organic search visibility.
3. **Become an SEO evangelist.** Evangelism is just one more on the long list of job requirements for the SEO pro. One of the more important tasks is walking the halls, spreading the gospel of SEO. Buy the IT person breakfast, the PR director lunch, and the content writer dinner. Get people excited about organic search and then be sure to keep them that way!

Progress Through the Stages of SEO Maturity

Now that you have organizational buy-in and appropriate resources to execute, how do you go about scaling your keywords and continuing to drive organic search traffic to your website?

Our observations over the years have shown that organizations often move through five distinct stages when it comes to SEO. Each stage is characterized by investments in budget, strategy, and metrics appropriate to the stage's maturity level. Organizations should continuously be striving to move into the next maturity stage until they are at the “Compete” level and are actively competing to best the top firms in the search rankings. The stages are:

1. **Try.** In the first stage of SEO maturity, the organization dips its toe in the water and focuses on low-hanging fruit, optimizing for a small set of initial keywords. At the same time, they begin efforts to move their top target keywords on page 7+ of the SERPs into striking-distance positions on pages 4 or below, so they are primed to move into top-visibility positions in later phases. In the “Try” phase, the majority of the budget is spent on personnel, with a portion allocated to engineering/production. A minimal investment is made in SEO technology. One employee is assigned to SEO part-time as part of his regular duties.
2. **Invest.** In this stage, organizations begin to expand their coverage, moving into more competitive keywords with higher search volumes. Average ranks begin to improve as keywords optimized in the previous stage begin to move into page 2, 3, or 4 of the search rankings. Organizations now begin to invest in some link-building activities and content creation and start to expand their investment in SEO technology. One or two employees are assigned full-time to SEO.
3. **Measure.** In this stage, to maintain investment in SEO, it becomes critical for organizations to begin showing ROI and tying SEO metrics such as rankings and traffic to business metrics such as conversions and revenue. In the “Measure” stage, rankings continue to increase as the organization becomes increasingly SEO-savvy, beginning to invest in continuing education. Several full-time employees are assigned to SEO, along with specialized outside consultants.
4. **Scale.** In this stage the organization organizes and prioritizes, beginning to look at organic search from a holistic perspective. A organic search pipeline develops as SEO starts to be treated as a sales campaign organizationally. Many more keywords move into top-visibility positions as investments in link building and content creation are increased. Several employees remain assigned full-time to SEO along with outside consultants and specialists.

Building an SEO team

At this point in the SEO maturity cycle, you are ready to catapult yourself into organic search success by moving from relying on internal “generalist” SEO professionals with outside consultant specialists to building internal teams of specialized professionals. From content writers who write search engine-friendly content to link developers who focus on leveraging that content to acquire strategically important links and keyword jockeys (Microsoft Excel experts adept at discovering hidden opportunities through spreadsheet analysis) to the project manager that pulls it all together, you are now ready to become a lean, mean SEO machine and drive to compete with the best the SERPs have to offer. This is the final stage:

5. **Compete.** In this stage the organization makes additional gains in organic search visibility such that it begins to look at its share of search relative to the competition. Budgets stabilize, while mature SEO technology is in use company-wide. In the Compete phase you are building internal teams of SEO professionals.

A special thanks to Seth Besmertnik (<http://www.conductor.com>) for his contribution to this portion of the chapter.

Using In-House SEO Talent Versus Outsourcing: Dynamics & Challenges

Once your organization has embraced the need for SEO, you must decide whether to hire someone in-house or work with an outside consultant or agency. The same decision must be made when the time comes to expand an existing SEO team. Here are the factors you should consider:

1. Is there enough work to justify a full-time, part-time, or contract employee?
2. Is in-house knowledge on SEO strategically necessary for the business?
3. Are people with the right SEO skill set for your website available for hire within your target salary range? (What each company needs will vary.)
4. Do you want to access and leverage experienced talent from an external agency to supplement an internal team?
5. Is your organization so complex that it would benefit from the presence of an in-house SEO pro?

The Value of Bringing SEO In-House

One of the significant advantages of in-house SEO is the ability to access search expertise whenever you need it. Some additional advantages are:

- A greater likelihood that SEO best practices will be integrated into the company culture.
- Greater accountability.
- Faster ramp-up time (familiarity with the core business).
- Faster implementation time (greater leverage in the development process).
- The ability to ensure that SEO requirements remain in scope (and raise the flag when they are de-scoped).
- The SEO expert can be tapped anytime and his or her work reprioritized immediately for quick deliverable turnarounds.

The Value of Outsourced SEO Support

Although it may seem that in-house SEO is the ideal solution, there are also plenty of reasons to use an outsourced team instead of, or in conjunction with, an in-house team. Here is a summary of some of the best reasons to leverage outsourced talent:

- Finding expert-level (15+ years) SEO talent to hire is difficult, so it may be smarter to contract with such talent.
- It allows you to focus on your core business and outsource your non-core activities.
- Experts generally have a broader reach in the marketplace via industry contacts, and as a result of working with many clients.
- SEO is multidisciplinary, so using external people makes it easier to pull in components when needed.
- Outsourced teams don't have the myopic focus that sometimes occurs in long-time in-house personnel.
- Outsourcing brings a fresh perspective and new ideas, and outsourced talent can often identify what in-house staff may have missed.

Many SEO consulting firms, such as those of the authors, are willing to take on contracts that include training in-house personnel. This kind of approach can give you the best of both worlds: the ability to immediately inject a high level of expertise into your organization (albeit on a temporary basis), and a long-term plan to bring the capability in-house, or at least build up the expertise with ongoing or occasional augmentation via outside consulting.

The Case for Working with an Outside Expert

You might want to work with an outside expert if any of the following scenarios apply to you:

1. You don't have any inside SEO resources, are not in a position to hire them, and need someone to work directly with your development and marketing teams to get your SEO campaign going.
2. You don't have any inside SEO resources, and while you want to build a team, you need to jumpstart the process in the meantime. You may also want to use your outsourced SEO practitioner to help train your team.
3. You have an inside team but it is missing a key piece of expertise, or you just want to sanity-check the conclusions your team is drawing or see whether they are missing anything. In addition, since SEO at its best involves lots of brainstorming, an injection of fresh ideas can be just what the doctor ordered.
4. You have an inside team but they are busy working on other projects.

Regardless of the scenario, an outsourced SEO professional can bring two assets to the table that will be difficult for your internal team to provide:

The power of external authority

If you hire an industry-acknowledged expert, that person can often speak with a voice of authority that is not readily available to the inside team. This is not necessarily because the inside team are less senior; it's just because they are, well, inside. The view of an unbiased third party with industry credentials can be invaluable.

From management's point of view, the opinion of a recognized authority can increase their comfort that the right path is being pursued. It is very common for in-house SEO professionals to engage an outside expert to help their cause. Engagements range from a presentation or training to site audits and quick consulting calls with IT.

The value of experience on multiple sites/campaigns

Outside experts will have the benefit of having worked on hundreds, if not thousands of different sites and SEO strategies. Each site's SEO strategy is like a lab experiment in which the SEO expert can benefit from many different types of tests to learn what works well. In addition, if your outside expert is active in the SEO community, he or she may also benefit from discussing different experiences with other SEO experts.

How to Best Leverage Outside Help

As we discussed previously, there are several scenarios in which you might bring in an outside SEO expert. In all of these scenarios, you will benefit most by working with someone who is good with people, can effectively and diplomatically negotiate, and can teach. Bringing in someone you can trust is also a must.

To leverage these skills, you need to put your SEO expert in a situation where he or she can succeed. A good candidate will tell you what will be needed from you to be successful. Take this input seriously. If you can't provide the things you are being asking for, your ability to succeed will be limited. It is vital that management paves the way for successful SEO plan implementation. Although many managers may believe they are giving SEO projects what they need, this is not always the case, and as a result, SEO suffers. The key is for management to understand how to give the right type of SEO support.

For example, the SEO expert may need time from your development team to implement the suggestions he makes. If the development time is not available, his suggestions will be useless to you. Prearrange time with developers to talk to the SEO consultant and work with project management to schedule time to work on the recommendations.

On the content side of things, the SEO practitioner may define a content plan to support ongoing content development and marketing efforts. Before you commit to a new hire, make sure you have a rough idea of the scope of the content development and marketing efforts the individual thinks might be needed to support this.

Listen to Your Expert – That’s Why You’re Paying Them

Take the input seriously, because this is the SEO expert telling you what he thinks he needs to be successful. If he or she doesn’t get that support and ends up thinking he can’t win, what are the chances he will succeed? The thing to remember about SEO is that implementation is crucial – you cannot implement a partial strategy and expect to see results. Unfortunately, most companies implement only a fraction of what is needed and wonder why their traffic isn’t increasing, and it truly is a shame – which is why buy-in is so crucial. Remember, SEO is *extremely competitive*. There’s a reason Google makes so much money with AdWords – because not everyone, and not every business – has what it takes to make high volumes of organic traffic happen. To be competitive in SEO isn’t the same as finishing a 5k running race – it’s more like running a marathon and finishing in the top 0.001%. Do you think any of these top marathoners only implement 50% of their training, diet, and sleep regimens? Not a chance.

Ultimately, the cost of hiring an outside SEO expert goes beyond the direct cost for the expert’s services. The SEO expert will need a certain amount of support to have a chance for success. Before you decide to hire an SEO professional, ensure that you have the internal support (and are willing to provide it) to make the recommended changes. You don’t want to be one of those companies that pays for SEO advice that sits on the shelf and collects dust – the authors have seen this more times than we care to count!

You also need to have a plan to manage your outsourced SEO help, as high-quality SEO firms will also be placing demands on your organization. Quality SEO cannot happen in a vacuum. This includes content marketing, which requires high-quality content—content that represents your brand.

The types of internal resources you need to be able to provide include:

- IT and web development resources to implement the SEO firm’s recommended changes to the website
- Marketing resources to review plans and any content created to support those campaigns
- Content development resources to create any required content, whether it be adding copy to the website or creating content to support content marketing and the earning of links
- One or more project managers to oversee the entire process

Leveraging SEO Knowledge in an Organization

Although it makes sense to involve an experienced SEO professional (employee or consultant) in your optimization efforts, it is important to remember that for SEO to be successful in the long term, you need to integrate SEO expertise into every aspect of your online business.

A basic level of SEO knowledge within the software and website development departments can speed up the production process, as a well-educated development team already knows a lot about how SEO affects what they do every day and can modify their efforts in accordance, on an ongoing basis. They will also be more likely to make the right decisions when the person in charge of SEO is not immediately available. Ultimately, this can save a large organization hundreds of thousands, maybe even millions of dollars—because it won’t be necessary to do a large amount of extra work to make the site search engine-friendly.

SEO is no longer merely about title tags and meta descriptions. Although this wasn’t the case a few years ago, user engagement signals and website usability play increasingly large roles in search engine visibility today. It is vital that website designers think about the experience of the user who comes to the site from a search engine results page, while social media strategies need to incorporate keyword themes and links that search engines are looking for as signals in their ranking formulas, as well as the overall element of social engagement that leads to natural link development and increased brand awareness.

Make Nice with IT—Or Bypass Them

If you are in a larger organization, you must persuade your IT team to have a stake in your site's SEO success - because you will rely on them to implement many of the optimization efforts. The IT professionals can be your biggest ally, or a thorn in your side. They need to weave SEO into every facet of what they do, and that happens only when they want to work with SEO. Start by heaping on the praise and appreciation—and maybe buying them pizza!

If conflicting priorities and insufficient internal resources stymie you, however, you will have to institute workarounds that minimize IT involvement. Have no fear: you can still accomplish what you need to with your site, with a bit of out-of-the-box thinking and some technology.

For example, you could use proxy optimization technologies (such as Rio SEO Website Optimizer, originally GravityStream, created by co-author Stephan Spencer) that optimize the URLs, titles, page content, navigation, and so on via a proxy server, obviating the need for major invasive surgery on the underlying ecommerce platform.

Start with the pizzas, though, and go from there.

Making sure your development team has some basic SEO knowledge can dramatically reduce the work required of the SEO practitioner, so that when he does review your website for search engine friendliness and to create an SEO strategy, instead of spending time talking to you about Flash, JavaScript, header tags, and alt text, he can focus on more advanced solutions and strategies. The goal would be for the organization as a whole to address 80% of the SEO workload, and the in-house SEO or consultant to address the more complex 20% of the SEO workload.

Similarly, basic SEO knowledge is an incredibly powerful tool in both online and offline marketing. Many companies have failed to integrate their offline marketing efforts (such as television or radio commercials promoting a new product) with their online efforts, and SEO planning in this area can often make (or break) a product or website launch.

SEO knowledge at the management level is essential for similar reasons. Decisions made by management, which can affect the entire business, can often have disastrous effects on SEO efforts if they are done in a vacuum—one common example of this is contractual agreements with partners for content syndication. It is a common problem that even after the development teams, designers, and copywriters are trained on SEO, managers end up becoming the obstacle because they haven't been briefed on, or haven't entirely bought into, the aspects of SEO that executives need to know about.

When all of the members of an organization have and act on basic SEO knowledge, the likelihood of success is exponentially greater. The SEO specialist or team can focus most of their energy on new initiatives and continuing to develop the process, while the rest of your team makes good business decisions and taps them for on-demand answers to questions.

A special thanks to Jessica Bowman (<http://www.SEOnhouse.com>) for her contributions to this portion of the chapter.

The Impact of Site Complexity on SEO Workload

The amount of time that is needed to perform effective SEO depends on the size and complexity of the website, as well as the competitiveness of the market space the organization is pursuing. An organization's size and vertical scope also have an effect on the overall complexity of the SEO process. Here are some of the ways that SEO considerations can affect the complexity of a site:

Keyword research

More pages mean more keyword research. Solid keyword research is needed to help drive the site architecture and content plans.

Title tags

Title tags (your page titles) are still an important ranking factor, and an important piece of the search results page itself. This means that you need to take the time to write the best possible copy for

generating the highest click-through rates. For very large sites, you may have to design an algorithmic method to choose these for you.

Page content

The content must be unique and substantial enough for a search engine to understand what the page is about. Google's Matt Cutts has mentioned that two to three sentences of unique can suffice to achieve ranking (<http://www.stonetemple.com/matt-cutts-and-eric-talk-about-what-makes-a-quality-site/>). If you sell products and receive content from manufacturers, you need to invest the resources to write unique product descriptions; otherwise, you risk being omitted from the search engine's indexes or ranking lower because your content is a duplication of other sites' content.

Meta descriptions

Meta descriptions are important because search engines often use an excerpt from your meta description in the SERPs, and the description they provide for a page can influence its click-through rate. While you cannot directly control the description used in the SERPs, you can influence it. For example, by including the keywords that you are targeting with a page within its meta description text, you can make the relevance of your page clearer to the searcher. The larger the site is, the more writing you will have to do, because search engines value unique meta descriptions for each page on the site. For very large sites, you may have to design an algorithmic method to choose these for you.

Link development efforts

As sites scale, the complexity and need for links grows. You need to research the competitiveness of your targeted keywords, and make plans for link development so that execution neither grossly exceeds nor falls short of the necessary effort. The more websites/domains your company owns, the more link development is required. Likewise, the less authoritative your website is, the more link development work is required.

Web-based partnerships

Websites of all sizes engage in partnerships and relationships with other entities (charities, businesses, consultants, clients, distributors, agents, etc.). SEO professionals know that all of these partnerships represent opportunities for acquiring links and link relationships, and that when they are properly leveraged they can result in massive value-adds for the organization. For larger organizations, these partnerships can be more complicated in nature. The more complex the organization is, the longer it will take to leverage these opportunities for SEO.

PR Team/PR Agency

Your PR team is your friend when it comes to link development and content distribution – so be sure to integrate your SEO efforts with your PR efforts. Truly advanced organizations task a link metric to the PR firm they use so that the teams are asking for links with every story they pitch. You should too.

Development platforms and CMSs

The development platforms and CMSs used on larger sites can often create a number of limitations regarding SEO implementation, and frequently require extensive, costly, and time-consuming workarounds before optimization efforts can be implemented. If you have a non-search friendly CMS (most are not search engine-friendly), you will have to do more customization work. It is recommended that you work with an expert to understand what is needed so that you develop it right the first time (since you may need to recode things every time you upgrade the CMS).

Solutions for Small Organizations

Some organizations are not equipped—either structurally or financially—to have an entire SEO team to handle the SEO workload we have been discussing. In fact, only one person may be knowledgeable about SEO, and that person may be only a part-time employee. Or there may not be anyone within the organization with the time or skills necessary to optimize the site, so outsourcing may be required. This section will give you direction on how to handle SEO in a small organization.

The In-House SEO Specialist

Building SEO knowledge in-house can be challenging in a small organization where most of the employees are already performing multiple tasks. It is often good to have an SEO consultant on call to answer questions and validate solutions.

If you need to assign SEO to existing talent because you don't have the budget to hire an agency, consultant, or contractor, you should consider engaging an SEO professional to evaluate the aptitude of the person(s) you have in mind for the role. The cost of this assessment from an independent consultant can run anywhere from several hundred to a few thousand dollars, depending on a variety of factors. Countless companies have underestimated the specialized skills and investment required for successful SEO. It's a role unlike any your organization has seen before—this person needs the technical skills, marketing panache, and political savvy to work his or her way into your website development processes, and the innovative, outside-the-box thinking that will generate creative solutions to search engine crawler needs. Choosing the wrong person can be a costly mistake in terms of time wasted and missed opportunities.

As with larger organizations, it is important to develop a basic level of SEO knowledge throughout the organization. SEO still touches on management, marketing, and development, and it is important that all of these departments have a basic understanding of the issues.

We will outline some of the ways to effectively and rapidly learn about SEO and build up in-house knowledge in "Working with Limited Resources/Budget".

Working with Limited Resources/Budget

Learning SEO and doing it on your own can be a challenging task, for two major reasons:

1. The demanding, ever-changing landscape of search algorithm behavior is often unpredictable and non-intuitive.
2. There are literally thousands of details and tactics to learn, some of which may have little impact on their own but, when used in various combinations with other components, can have a powerful influence on rankings. Herein lies the "art" aspect of mastering SEO.

Fortunately, many SEO training tools and materials are available via paid subscription at Moz (<http://moz.com>), SEO Book (<http://training.seobook.com>), Instant E-Training (<http://www.instante-training.com>), Market Motive (<http://www.marketmotive.com>), ClickZ Academy (<http://www.clickzacademy.com>), among others. If you don't have the budget for a subscription, you can try public blogs and resources such as the Moz.com and SearchEngineLand.com. Also consider the many resources for learning that we discussed in Chapter 11.

Basic low-budget SEO ideas

You can do numerous things at a fairly low cost to improve your site's overall optimization, including the following:

Use the free search engine tools

Use the free tools provided by the three major search engines. Create accounts in Google and Bing Webmaster Tools and verify yourself as a site owner on both. This will provide you with access to diagnostic tools, such as *robots.txt* validators, as well as reports on backlinks, spidering activity, server errors, top search queries, anchor text, and more.

Find the best keywords to target

Use the Google Keyword Planner (<https://adwords.google.com/KeywordPlanner>) to find keywords with high search volumes. Then use the Moz Keyword Difficulty & SERP Analysis Tool (<http://moz.com/tools/keyword-difficulty>) to get an estimate of how hard it would be to rank for the terms you have identified.

Check out your competitors

Assess your site and those of your competitors for SEO success factors such as keyword-rich URLs, title tags, and `<h1>` tags, keyword prominence, and so on. To survey your and your competitors' title tags across a large number of pages, use the search engines' `site:` operators and set (in the preferences) the number of results returned per page to 100.

Optimize your title tags

You want each title tag across your site to be unique and focused on a relevant keyword theme. Make each title tag count, because of all the elements on the page, it's what the search engines give the most weight; it also heavily influences the searcher's click decision from among the search results.

Optimize other critical elements

Analyze the text, HTML, inbound links, internal links, anchor text, and so on to determine your ideal configuration for success. Include a dose of your own critical thinking.

Measure, test, measure, and refine

Test your assumptions and the assertions of others—particularly SEO bloggers (not every piece of advice you find will be accurate or applicable). Measure key performance indicators (KPIs) and treat SEO like a series of experiments. Make iterative improvements to your URLs, title tags, `<h1>` tags, internal linking structure, anchor text, page copy, link acquisition efforts, and so on.

What sorts of KPIs should you measure and aim to improve? At a minimum, consider checking rankings, traffic, and conversion metrics. However, you can also check other metrics, such as the number of different search terms used to find your site (for this, use Google Webmaster Tools, SearchMetrics, or SEMrush, the number of different landing pages where search visitors arrive, the growth of inbound links and the addition of any notable inbound links, and so forth.

Books to help you understand using analytics that we recommend are: *Web Analytics 2.0: The Art of Online Accountability and Science of Customer Centricity*, Avinash Kaushik (Sybex) and *Advanced Web Metrics with Google Analytics* by Brian Clifton (Sybex).

Test different ideas

Get one *great* idea for a tool, service, resource, or page, and bounce it off some folks in social media, in SEO forums, or privately through email to an SEO expert whom you trust. Hire a developer who can help you build it – consider leveraging your online relationships via LinkedIn, Facebook, and Twitter to mine recommended developer talent.

Many in-house SEOs have had success finding copywriters via sites like Craigslist or Textbroker to write articles for nominal fees ranging from \$10–\$50 per article (few companies and in-house resources can compete with these rates), though remember the old adage that you get what you pay for, and this applies online as well. Again, consider leveraging your online contacts to find highly regarded talent within your budget.

Leverage low-cost tools

Consider using one of the following tools:

WordPress (<http://www.wordpress.org>) or Drupal

These tools are popular with dynamic and static website/application developers, respectively. Most of the time they are used to build web pages, but they also offer a range of reporting tools.

Xenu's Link Sleuth (<http://home.snafu.de/tilman/xenulink.html>)

This is a simple link-based crawler. Web developers use Xenu to check for broken links on a regular basis, but for SEO purposes the best value comes in the form of simple internal link analysis. By ordering the Xenu Sitemap based on "links in" and "page level," it is easy to detect possible internal linking abnormalities that may interrupt PageRank flow or decrease anchor text

value; and of course, you can save all this information as a report. Xenu gathers loads of information, and it is a very useful tool, even for in-depth SEO purposes.

Screaming Frog SEO Spider (<http://www.screamingfrog.co.uk/seo-spider/>)

This is a small desktop program you can install on your PC or Mac that spiders websites' links, images, CSS, scripts, and apps from an SEO perspective. It fetches key on-site page elements for SEO, presents them in tabs by type, and allows you to filter for common SEO issues (or slice and dice the data how you see fit by exporting it into Excel). You can view, analyze, and filter the crawl data as it's gathered and updated continuously in the program's user interface.

Microsoft Word (<http://office.microsoft.com/en-us/word/>)

Although it may seem to be an unconventional tool for a web developer/SEO practitioner, Microsoft Word is undeniably one of the best copywriting and publishing tools that practically all users are familiar with. It has several built-in features that help to produce high-quality content, analyze existing content, fix and locate essential grammar errata, and above all, easily automate and synchronize all features and changes with other users and publishing tools. For more tech-savvy folks, there is always the scripting option for fine-tuning.

As with most SEO tools, the beauty is in the eye of the beholder. If you use the preceding tools properly, they can be very helpful, but if you lack experience or try to use them for the wrong kind of task, they can cause pain and misery.

Making proper on-page optimization decisions usually takes days. Even for a relatively small site, it is possible to cut that down to fewer than two hours by using the tools and methods we just described. Of course, there is a difference in the quality of work and the documentation you'll get with these tools compared to what you'd get with more conventionally priced SEO services, but they do have their place.

SEO webinars

There are numerous free SEO webinars available that will expand your knowledge. Great sources include Instant E-Training (<http://www.instantetraing.com>), Moz (<https://moz.com>), Digital Marketing Depot (<http://digitalmarketingdepot.com>), and SEMPO (<http://www.sempo.org>).

Limited cash options

If you're on a tight budget, create a blog with great content and promote it via social media (Twitter, Facebook, and Google+). Execute the SEO action items identified in this book, and leverage the free tools and guides we've mentioned. Attend free webinars. Do what you can for SEO. Then, invest in an hour or two with a consultant to do an ad hoc site review. You can also check with SEO firms to see which ones are offering special introductory packages, as some firms offer special pricing just to get started with a client.

These are just examples of things you can do, and countless other options are also available. Whatever you do, don't short-circuit the process of developing in-house expertise, even if you hire outside help. That in-house knowledge will be invaluable over time, and if you have hired a quality outsourcing partner they will be delighted as it will make their job easier. They may even help train you or your team.

Solutions for Large Organizations

The challenges of performing SEO for a large organization are a bit different from those for a small organization. This is true regardless of whether you are working from an in-house position or are an outsourced SEO consultant. Some of the challenges can stem from the size of the site, which can range from 10,000 to tens of millions of pages, potentially spread across multiple domains, countries, and languages.

Large organizations are usually complex entities, and many of these organizations may make decisions by committee, or review all decisions with a committee before finalizing them.

Many important projects (SEO or otherwise) can be delayed or even canceled in a large organization due to a lack of understanding by a single key player. In the world of SEO, efforts are often delayed by someone in IT/development, marketing, sales, or management—and it can be challenging to get a meeting with the people you need to persuade, let alone accomplish the task of persuading them.

Patience and persistence are essential, and adhering to the following large-organization SEO guidelines is recommended:

- Get buy-in, if you can, from the head of IT, the head of marketing, the head of sales, and senior management. Those people will likely control your fate. If you can get two or three department heads and senior managers on board, you should be in good shape.
- Always make sure you talk about opportunity cost. Make sure all people involved understand that they're potentially leaving X visitors on the table every day, and that at the current conversion rate that means N potential leads and Y potential dollars. Include the opportunity gap—where the clients are versus where they could be—in every report.
- Insist on a sound web analytics plan. Successful SEO projects depend highly on quality analytics information. If your company currently cannot implement a quality analytics solution, yet you have support for implementing SEO, look at other KPIs, such as rankings combined with search volume and estimated clicks per position in the search results. You can also rely on the data from Google and Bing Webmaster Tools, which will provide some quality information.
- Provide detailed reports even when you're not asked to. As the disciplined people around you are likely to be unfamiliar with the basics of SEO, SEO practitioners need to over-communicate.
- Educate, educate, educate. Always explain why you are making a given recommendation. Just make sure you are speaking in the language of the recipient (i.e., talk technically to the developer and in business terms to the marketer). This is where many SEO efforts fizzle out. When you conduct training properly, you can create an almost overwhelming amount of buy-in and interest in SEO. Companies often go wrong when they think one training session is enough. Unfortunately, interest will wane, and if you don't offer continuing education you can be back to square one in 6 to 12 months.
- Be confident. Cover your bases before each meeting. It's important to speak confidently on what you know, but it's also important not to speak about what you do not know. There is nothing wrong with saying, "That's a good question. I don't have an answer for you right now. Let me do some research and get back to you." Most development teams will respect that answer if you can do your research and get back to them quickly. Be mindful of the Ugly Baby Conundrum: the website is the baby of the programmers, designers, and business sponsors; when you point out a site's SEO needs and flaws, it's like you're saying, "Your baby is ugly." When you look at a site for SEO, you will probably find many issues. Be sure to mention what has been done well in terms of SEO, and be unbiased when presenting the issues you've found.
- Don't make technical presentations to the executive team—they want metrics and action plans showing progress (or lack of progress). Although they might understand what you're saying, their brains are tied up with 99 other things. Just get to the point: what's gone well, what hasn't, and what you need from the team to fix it.

Contracting for Specialist Knowledge and Experience

Even if you have a solid in-house SEO team, sometimes it makes sense to get help from the outside. Why would you want to bring in external expertise when you have a strong team already?

The answer lies in the complexity of SEO itself. SEO has dozens of sub-disciplines, including video optimization, local search, image optimization, competing for search traffic in other countries, link development, usability, and strategies for social media properties. If you need help in one or more of these areas, don't hesitate to seek assistance from outside your organization. It's also possible that your SEO team has mid-level, but not expert-level knowledge of SEO – and contracting out for expert level guidance of your team is often a wise investment.

Another reason to bring in an SEO consultant is to conduct an audit of your site. This should help identify the things you've missed, confirm the things you've done right, and highlight new opportunities. Here are the ideal times to engage a consultant to perform an SEO audit:

- When you need to understand the amount of work that will be involved in optimizing your site
- When you have done what you know about SEO and need advice on what to do next (or as part of an annual training, to help develop better expertise inhouse)
- When others in your organization need justification of your SEO recommendations
- Before a site redesign (or at any time you make major changes to a site), so that you can learn what needs to change and what you can do to make the next website more search engine-friendly

Applying SEO Recommendations Intelligently

One of the biggest challenges in SEO is the gap between strategy and implementation. The SEO artist knows how to explain his recommendations in simple terms and in the language of his audience. However, simple expression of a concept does not always mean seamless implementation.

When you look at a site from an SEO perspective, you must always keep the bigger picture in mind: to do what is best for the business by balancing both short- and long-term goals while managing available resources effectively. Changes that minimize implementation cost and hassle but have a big impact have the best ROI. Before you recommend throwing out that old CMS and reinvesting millions of dollars in a brand-new website, ensure that the ROI to the organization will justify the expense, and back it up with data.

This does not mean you should never recommend drastic changes, but rather that you should do so at the right time and place, and break out the implementation of the overall SEO strategy into prioritized SEO “projects.” Be mindful of the potential benefit for each SEO recommendation. Sometimes this is difficult to quantify, and you may end up with a table similar to Table 14-1.

Table 14-1. Rating potential SEO projects

Project	Value	Effort
SEO Project 1	High	High
SEO Project 2	Low	Medium
SEO Project 3	Medium	Low
SEO Project 4	High	Low

This type of analysis is incredibly valuable in keeping an organization moving forward with an SEO program and increasing search traffic at the most rapid rate possible. Of course, sometimes the executives in your company may want to see dollar numbers instead of High, Medium, and Low ratings for the Value, so you may need to be prepared to give an estimated range of dollar impact. Just make sure to be clear that you are providing a very rough estimate, as this is usually difficult to predict precisely.

Hiring SEO Talent

If you have decided that you want to add to your in-house SEO team (or start one), you need to think seriously about the type of person you want to hire. Do you want someone with prior experience, or do you want the person to learn from other team members along the way? Do you know which facets of SEO are crucial to the success of your specific website and organization, and how to hire for these skill sets?

Hiring top people can be prohibitively expensive, because there are a lot of income opportunities for them to pursue. It is hard to match the earning potential of a top SEO pro because, at the very least, he or she can usually make more from building sites and operating them in an affiliate or lead-generation model. That does not mean you cannot hire real talent; it just means you will pay a premium. An alternative would be to

hire someone less seasoned and experienced, and pair him or her up with a consultant or agency to oversee the work.

How to Select the Right SEO Person

Although experience costs more, under normal circumstances it will bring faster results. On the other hand, a less experienced person is a lot like a rookie in baseball: if you hire a future star, you can potentially accomplish a great deal at a much lower cost.

When sifting through résumés, examine candidates' specific skill sets. If you are looking for someone to develop social media campaigns, you should probably be a bit more focused on that than on whether the candidate is an expert at keyword research. No matter what his skill level is, he should "get" the Internet. Do they blog? Tweet? Have a professional and well-connected LinkedIn profile? If they don't, and there isn't a specific personal or professional reason why (privacy, confidentiality, etc.), they are possibly under-experienced for the position you are looking to fill.

Also, when evaluating a person's skills, don't forget to assess his or her political finesse. What makes a successful SEO practitioner at an agency or in a freelance position is different from what makes a successful candidate in-house. Often, companies focus strictly on results and types of projects; however, if your company is plagued with red tape and political minefields, you need someone who can remain diplomatic, sell SEO to the entire organization, and integrate it into the organization in the right places so that it is implemented successfully. There is an 80/20 rule for in-house SEO: 80% of the time you spend selling and 20% of the time you spend doing SEO.

Pitching the Person

Finding good SEO talent can be challenging, as top SEO experts are in high demand and are likely to have more than one job option to pursue; for these SEO pros, a high six figure salary is the norm

Many companies overlook this. The "jobs offered" sections of forums are often full of offers, but few employers have made any real effort to market their openings. It really is a basic marketing problem that many of those vacancies show. Many don't even put a location in the title. Far worse is that almost all of the posts are company-centric: all about what the company wants from applicants, and no thought for what the prospective applicant will want from an employer. SEO is not like other job functions. A good SEO is in high demand; you have to make him or her think, "I want to work there!"

Fail to pitch the prospect and you will get to choose from a diminishing pool of folks who can't make it alone, aren't able or willing to make a good living from affiliate programs, and don't prefer to be with a more forward-thinking and proactive company that looks to develop its staff, rather than merely recruit them. Many SEO pros are looking for companies that have more support and resources – they are frequently stuck as isolated one-person departments, with few tools and "toys" to help them advance.

Sample Job Opening – an Oldie but Goodie

The following job opening appeared in the job postings at Search Engine Watch nearly a decade ago (<http://forums.searchenginewatch.com/archive/index.php/t-4088.html>):

I'm totally serious about this unique opportunity for a person of the right calibre to come and work with Propellernet UK, a Brighton-based Search Engine Marketing company, as a trainee Consultant.

This is not your average trainee position, and a high level of skills are already expected. We do not have the time to spend years on this training, and we are looking to take on an experienced SEO, who already understands the special aspects of Marketing as it applies to the web. Then I will personally train and coach you to take that further still.

The ideal candidate will have a background in sales or marketing, as well as a broad understanding of SEO.

(Clue: If you don't know who I am, your knowledge of SEO is not going to be broad enough.)

You must be a great communicator, as an essential part of this role is to bridge the communication gap that often exists between the Marketing and Technical Departments of Client Companies.

You need to be able to talk to Marketing Directors on their level, and then turn around and talk with Heads of IT Departments on their level without a pause.

You need to have equal passions for both Teaching and for Learning, as you will be expected to do both things each week as a matter of course.

You need high initiative, and the ability to think on your feet. You need to be proactive yet still have fast reactions too.

You must be in possession of a full, clean, UK Driving License, as considerable travel may be involved in the consulting work.

What you get in return:

A salary based upon ability and experience that will be reviewed regularly as your skills are proven. The salary will also include performance-related elements, and on-target earnings could easily surpass £70,000 as soon as your skills are up to the demand.

You get to work right on the cutting edge of SEO consultancy, where the work is never dull nor repetitive and we seek unique solutions for every client.

You get to be trained to be a world-class SEO and Web Marketing Consultant, able to stand alongside the very best that our industry has to offer.

I am looking to train an equal, not a junior, so you'll be learning everything that I can impart to you just as quickly as you can absorb it.

This is a completely sincere offer. Your time will not be wasted—I don't expect mine to be either.

CVs should be sent to Propellernet Ltd.—you can find the address for yourself if you're a qualified candidate.

Sincerely,

Ammon Johns (aka Black_Knight)

Marketing Director, Propellernet Ltd.

Clearly, this type of job posting is not what every company would want to put out there, but that is part of what makes it effective. There is a certain amount of attitude to it, and that speaks volumes to a potential applicant: either this is what you are looking for, or it isn't.

The job requirements also spell out many of the personal and professional attributes an applicant would need to have to get the job. The compensation is well detailed too, so the applicant knows what is in it for him. Good SEO professionals are well compensated, hard to attract, and know the value of their time—they will not waste it on jobs they do not feel will give them valuable experience and just compensation.

In addition, the posting makes it clear that Ammon Johns was looking for someone who could be a true peer. This and many other points in the job outline make it evident that he is looking for a senior person. One thing that might make this type of listing a bit more effective would be the addition of some examples of search problems that a qualified applicant should be able to solve prior to applying (the “find the address for yourself” requirement at the bottom is not a challenging problem).

Making the Offer

Every year, SEMPO conducts an In-House Salary Survey for SEO professionals (<http://www.sempo.org/?CurrentSalarySurvey>). Looking at this data can be useful when determining the salary to offer the candidate. SEO salaries are typically shocking to managers, so be sure to check this when budgeting for your in-house SEO position.

Move quickly! Once you know the candidate is the right fit, make an offer. Good SEO professionals are always in demand, and waiting two or three weeks to make an offer could cost you the superstar hire. It is not uncommon for top-notch pros to land a new job in three to four weeks, and it is not uncommon for a company to search for a candidate for 6–12 months. In the SEO recruiting space, if you snooze, you lose. Prioritize time for the interviews, and act fast. If the process just seems too daunting, consider paying for a few hours of a top SEO pro's time to help you in the hiring process.

Selecting an SEO Firm/Consultant

Once you have decided to hire an outside firm, you need to go through some sort of selection process. This can be daunting if you are not already familiar with the space and if you don't already have some understanding of SEO yourself. There are many solid and reputable companies out there, but unfortunately there are a large number of bad ones as well.

In addition, there are many different types of projects you may want them to work on. For example, if you have a local SEO-focused project in mind, you will need to find a team with the right type of experience. Similarly, you may have specific contract requirements that will impact which firm is right for you. Consider looking at resources such as the Moz vendor directory (<http://moz.com/community/recommended>).

Getting the Process Started

Just as a bad hire can be a very expensive proposition, so can engaging the wrong consultant or firm, either through a haphazard selection process, or no selection process at all. Here is an outline of the steps you should take for a more methodical, disciplined search for outsourced SEO expertise:

1. Start with your goals. Connect with your team and make sure you have a good idea of what you want to accomplish by working with an SEO expert. Sometimes it's as simple as an audit to make sure you're following best practices; other times it's training for your marketers and content builders to get them up to speed on how to actively promote the site. If you have a solid list of things you want completed at the end of an engagement, you'll be better able to find the consultant or firm that best suits your needs – and you'll be better able to judge the eventual results.

If you don't really know what you want, that's OK, but it does make the SEO practitioner's task a little less structured, and it means you need an exceptional level of trust in that person. Sometimes an SEO expert can help you define your goals—to get the best of both worlds, you can ask her what she thinks your goals should be and combine those with your own list. It is still fairly common to hear companies tell us "We don't know what we need to do for SEO, but we need to do something – can you tell us?"

1. Connect with your social network. The people who can best assist you in finding a good match are often those who know you well. Talk to friends, fellow business owners, SEO bloggers, and people you know and trust.

When you communicate with your network, make sure you share a little bit about your site, your business, and your goals with regard to SEO (to the extent that you know them). The more information you can share, the better folks are able to assist. At a minimum, if you have geographic considerations or want the culture of a large, multi-offering agency versus that of a smaller, SEO-focused consultant/consultancy, this is good to include in the request.

1. Get advice from SEO-savvy people you trust. Familiarize yourself with the Twitter handles and Facebook pages for the people you trust, and engage. Tap your Google+ circles and LinkedIn connections to request advice and guidance. As a rule, SEO experts are incredibly friendly about referring business to good companies and good people, and if your network has already given you a few names, running these by the professionals can get you valuable insight on potential winners versus lemons. Even the authors of this book are willing to provide recommendations, so feel free to reach out to us.
2. Make sure the person is available. This is a fairly obvious step, but many people do not put this at the beginning of the communication process. Ask the person whether he is available at the start, and then continue the conversation if applicable (there is no need to waste energy on folks who can't help, though you should always ask for a referral). Your time (and theirs) is precious, so don't waste it.
3. Define what type of expertise you need. If you have an enterprise-level, highly complex website with legacy systems and a complex development life cycle, you need an SEO consultant who knows how to work within these constraints. If local SEO is your priority, you need someone well versed in local SEO. Few companies do extremely well in all facets of SEO, so you need to identify which aspect of SEO you need most and find a company that excels in that area.

Preparing a Request for Proposal (RFP)

A well-written RFP can significantly improve the overall process. If your RFP is not written well, it hinders the SEO firm's ability to understand and define your needs and to scope and price your project. This in turn leads to a disconnect in expectations for both parties. A lousy RFP can discourage a busy SEO firm from even responding—which is a very unfortunate outcome, since it takes the best firms out of the running. That being said, there are some SEO companies and consultants who do not participate in the formal RFP process, and with good reason – some businesses use the bait of an RFP to gain valuable SEO strategy tips – and even specific tactics – for free, and at the expense of the SEO vendor. So while it can seem like a good idea to ask for a lot of great information in the RFP, this can place an undue burden on the RFP respondents and may result in some of the best agencies and consultants not responding - simply because they have other opportunities to pursue without having to participate in a lengthy RFP process.

Many companies intuitively “know” what they want but find it difficult to ask for it in a way that is clear, succinct, informative, and constructive. If written properly, an RFP can facilitate the sales process and ensure that everyone involved on both sides understands the purpose, requirements, scope, and structure of the intended engagement. Following a few key steps in the beginning of the RFP process will enable you to rest easy, knowing that you are more likely to get what your company wants and needs.

Step 1: Nominate a “point person” for the engagement

One of the most challenging concepts for any large company with multiple working parts is to determine the end goal of the engagement. Often, marketing departments may voice different wants and needs than IT departments; even when they are asking for the same thing semantically, they may not be using the right language to communicate what they are looking for.

Every successful project needs a project manager who is invested in the project's success and can pull together the disparate groups that have a stake in the outcome or a role to play. Without that person at the helm, the project will struggle. By nominating that person as the SEO firm's “point person” even before you send out your RFP, you will ensure a steady flow of information throughout the process so that internal and external expectations are met.

Step 2: Define “needs” and “wants” using a decision matrix

RFP recipients will understand that you aren't the SEO expert, and therefore you aren't going to be able to adequately define the scope of your desired SEO engagement. As the saying goes, “You don't know what you don't know.” Just try to be reasonable when formulating your needs and wants and recognize that the RFP recipient may have a better idea of what you need than you do.

In addition to disclosing to the SEO firm the basics, such as your objectives, your site's conversion types (e.g., online purchases, newsletter subscriptions, white paper downloads, etc.), target audience, constraints, and so forth, in all fairness you will also need to clearly spell out the criteria by which you will evaluate the SEO firm's suitability. This means you will need to define these criteria in advance.

SEO firms know the selection process is usually governed by gut feel, so any RFP recipient who does not have a preexisting relationship with you is automatically inclined not to respond because he knows the deck is likely stacked against him. This is further exacerbated when you send the RFP to multiple SEO firms. You can allay this concern by candidly sharing with the firms your biases and the clearly defined criteria by which you will be evaluating them.

You should map this out in a *decision matrix*, which is simply a chart listing the attributes you are looking for in the SEO firm, a weighting factor for each attribute, a score from 0 to 10 for each, and the weighted score (e.g., the score multiplied by the weighting factor). The weighted scores are added to arrive at a total score. A highly simplified example of a decision matrix chart is shown in Figure 14-1.

Candidate SEO Firm Scoring Matrix							
	Trustworthy	Trust Weight	Experience	Experience Weight	Availability	Availability Weight	Total
SEO Firm 1	8	0.4	6	0.3	9	0.3	7.7
SEO Firm 2	4	0.4	8	0.3	10	0.3	7
SEO Firm 3	7	0.4	7	0.3	8	0.3	7.3

[[figs/print/1401.png]]

Figure 14-1. Decision matrix for evaluating competing SEO firms

Judging criteria that are both quantitative and qualitative brings objectivity into a subjective process, which will aid you in managing expectations internally. And by sharing your list of criteria and the weighting factors with the RFP recipients, you increase the likelihood—as well as caliber—of responses.

It is worth noting that the authors recommend that you place a high premium on trust in hiring an SEO firm. Unfortunately, there are lots of firms that claim a high level of expertise that either don't really have it or don't put in the effort to drive results (this is much the same in many consulting disciplines). This is one factor on which we recommend you place a lot of weight.

Step 3: Define your success metrics

Now that you've identified what you are looking for in an SEO firm, it's important to let the firms know how you would like to measure the success of the engagement. Some firms can report on SEO health metrics well beyond just rankings (e.g., page yield, keyword yield), thus facilitating troubleshooting and reviews of program performance. Others will rely heavily on your own analytics package to track the program's success.

For example, you might pose this question internally: "What are your KPIs?" SEO firms often use KPIs to prove the value of the services they provide. If you aren't sure which KPIs your company would like to use, think about baking that request into the RFP, stating something such as, "Do you offer monthly program performance reviews indicating the program's growth? If so, please elaborate on the deliverables of such reviews."

You can assist this process by defining your business metrics clearly, whether they be sales, leads, views or shares of a piece of content, or something else. Make sure the SEO firm is under a nondisclosure agreement (NDA), and provide them with a baseline of traffic and conversion metrics and ask them to show what strategies they are going to use to grow your business from that baseline. Don't get distracted by artificial metrics such as rankings or number of links. While these may be of some interest, keep your eye on your bottom line and make sure that your prospective SEO firm knows that they will need to do that too.

Step 4: Prepare to disclose all known influencing factors

An SEO firm can easily examine your existing website, evaluating the on-page factors (title tags, navigation, HTML) to gauge project scope. But other factors will influence your organic exposure during the engagement, some of which won't be immediately obvious or known to the RFP recipient without prior disclosure. Important disclosures include:

- Has your site received a Manual Action notification from Google?
- Has your site seen traffic drops resulting from the Panda, Penguin, Mobile, or other updates?
- Does your company have other domain names, subdomains, or microsites?
- Is a massive redesign of your website in the works, or are you about to migrate to a new CMS?
- Do you employ a third party for your site's internal search, and if so, who is it?
- Have you ever engaged in text link purchasing, article/blog spam, or link networks?

By disclosing as much information as possible, you spare the SEO firm the time and expense of discovering these things on its own and increase your chances of getting a good ROI on your investment in them.

Step 5: Provide an estimated timeline and budget for delivery of work and implementation

One of the biggest deterrents in any RFP is confusion over when a project should be completed and how much it might cost. A company hiring an SEO firm may not know how much time it takes to complete an SEO audit, but it may have pressing internal matters that require it be done by a specific deadline.

Retail sites may want to schedule new launches around a particular theme or season; for example, if you want your site launched in time for back-to-school, be sure to write that in your proposal. Not only will an expected deadline save your company time in looking for an appropriate firm, but it also serves as a professional courtesy to the SEO firms you are querying.

Budgets often coincide with a project deadline, depending on how your company conducts its business. Consider assessing a budget based on a range of services; that is, instead of saying the project absolutely has to cost \$X, say you are willing to spend within a range of \$X to \$Y. Budgets may or may not be included as part of the RFP, and there are benefits (and drawbacks) to both approaches. Even if you do not provide an SEO firm with your budget in the RFP, we recommend determining a budget beforehand because knowing how much money you are willing to spend will help you determine a target ROI.

Be aware that high-quality work at high speed is clearly going to cost more. Consulting with people you trust that have used SEO firms can help you get a sense of what the costs are likely to be. Once you have an idea of the market costs for a project of your complexity, you can then decide to spend more (to try to get faster results) or less (to manage budgeting concerns).

A sample RFP document outline

Now that you've gathered the preceding information for your RFP, you're ready to sit down and write it. You can structure an RFP in several ways. Here is a brief outline:

Section 1: Summary and overview

This section is where you will introduce the challenge you are having and provide the SEO firm with an overview of the rest of the RFP. We also recommend outlining how you would like SEO firms to respond to your RFP. Think of this section as an executive summary, where you will provide the highlights of the RFP without the technical details.

Section 2: Technical summary

Often, SEO firms need gritty details to determine how they can best help you. The technical summary is where you will provide key pieces of information relevant to your project, such as technical requirements, a description of technical issues with the project, your site's current platform, and so on. The technical summary might originate from your IT department, as this section is often for an SEO firm's programmers and delivery team.

Section 3: Administration and management

By describing who will be involved in the project on your end and what the timeline is for completion, you are defining the framework of the project.

Section 4: Project expectations and delivery

If you've done your homework, this should be the easiest section to write. From outlining your evaluation criteria to specifying what monthly deliverables or training you would like to receive to assessing your KPIs, this section helps an SEO firm determine its cost to complete your project, as well as its suitability.

When you write an RFP, keep in mind that the purpose of an RFP is to hire an expert SEO firm that will propose its recommended actions for your company's website(s) to achieve the greatest chance for success. Sometimes an SEO firm's feedback highlights other potential issues that you may not have considered, which may change the scope of your project entirely – and this is often the case!

Communicating with Candidate SEO Firms

You should also plan to talk to candidate SEO firms. This is your opportunity to dig a little bit deeper and move beyond the RFP. If you are a smaller company and have bypassed doing an RFP, you will still need to be in direct contact with your prospective SEO firm. Here are some things for you to consider during this stage of the process:

Ask for a list of past success stories (not just clients)

Let the company tell you where it has helped someone be successful. You can ask for a client list, but be aware that a failed client engagement does not necessarily mean the SEO firm failed. In addition, many SEO service contracts prohibit the SEO firm from disclosing their clients. Many times clients don't implement the SEO practitioner's recommendations, or they try to, but do it badly.

You can use this to your benefit. Ask them for a case study of a client failure, and then ask them why they think it failed. This can provide some great insight into what they require from their clients, because the failure they disclose to you will inevitably be due to the client not providing them with what they needed.

Talk on the phone or (if possible) get together

Email is a great initial communication medium, but a phone call or in-person meeting gives you a real sense of the team you'll work with. Make sure you're not talking to a salesperson (if you go with a big company), but rather to the person who will be interacting with you throughout the contract execution process. There's nothing worse than getting sold by a charismatic, knowledgeable SEO leader and then being passed on to a junior team for management.

In the phone call/meeting, be sure to establish rapport. If you don't feel comfortable on a personal level, don't dismiss it—brilliant SEOs are great, but if you can't work well together, the project is unlikely to succeed. Everyone works better with people they like and people they get along with. There's also the additional incentive of not letting down someone you've developed a working relationship with, which extends subconsciously and consciously into every part of the work you do together.

Present a few initial issues over email

Some SEO consultants won't engage in any work before a contract is signed, but you should at least get them to talk about how a problem can be approached, whether it's an issue, and what strategies they might recommend to fix it (even from a broad perspective). If you are an SEO practitioner, put a few examples of issues you've already identified on the table. Seeing how the consultant responds builds trust in her knowledge and gives you insight into how she solves problems and what your work relationship and the advice you receive will be like down the road.

Ask for references

Asking for references is a good idea, but make sure you qualify the reference as part of the process. You want to protect yourself from people who might get some friend of theirs, instead of a real client, to act as a reference. It is helpful if the reference the candidate provides works within a well-known business. Make sure the relationship between the SEO firm and the client provided as a reference is clear.

Once you get the reference on the phone, you should try to learn what the SEO consultant was asked to do, and what he did. Get specific details on this. If you want someone to provide link development strategy for you but he did not do link development for the other client, the value of that reference is decreased. Then see whether the reference will give you some input on the benefits they received from the work of the SEO consultant (e.g., increases in traffic, conversions, or other metrics of value). Then verify what the reference says about the SEO firm by checking their site metrics with Compete (<https://www.compete.com>) or Quantcast (<http://www.quantcast.com>).

Also, think about doing your own research. Find out what site the consultant worked on for the reference. Then use the Wayback Machine (<http://www.archive.org>) to see prior versions of the site, if available.

Contact other past clients

Look at the SEO firm's site and see if it provides client lists. Call ex-clients the candidate firm didn't give as references. In addition, ask about good and bad experiences with the company in forums that cover SEO.

Making the Decision

Regardless of whether you have done an RFP or used a more informal process, make sure you understand the process you will use to make a decision. Here are some tips:

Get a written response from two or three vendors

Don't get an RFP response or quote from just one vendor, even if you're sure you're going to use them. Multiple quotes aren't just a best practice, they're a good way to learn about pricing and scope discrepancies. The quote from the vendor you don't like might include some pieces that your preferred vendor overlooked. Pricing is harder because it fluctuates so wildly in the search marketing world.

However, don't choose an SEO expert on price alone (or even make price the biggest part of your decision), unless your budget is a real problem. SEO is usually an incredibly high-ROI activity. Companies that spend larger amounts on SEO services tend to make it back in a matter of weeks or months from traffic and conversion increases (remember that more targeted traffic means higher conversion rates too).

Thus, going with an SEO expert who costs a lot less might seem like a good idea, but if you don't work as well together and you think she might not do as good a job, you're hurting yourself in the long term. Of course, don't get fleeced by an exorbitantly overpriced firm (watch out; some of the biggest SEO companies have some of the most ruthless pricing models because they know that Fortune 1,000 companies will work only with consultants who've done work for other Fortune 1,000 companies), but be aware that the prospect of saving a few thousand dollars is not the best reason to choose a different provider.

Have smart, sensible people review the contract

Lawyers are great, but sometimes legal folks get overly concerned with liability and risk management details and overlook big-picture business items. Make sure your savvy business/ROI-focused personnel get a good look too. If there are legal issues that cause a rift, it can even be wise to get your C-level executives involved. Sometimes the folks from legal won't budge on an issue that's going to be a deal breaker, but if the CEO says do it anyway, you can get around the more problematic demands. Also be aware that different things are important to different firms. For example, some SEO companies will walk away from a deal if the client insists on a no-publicity clause or if you insist on them assuming millions of dollars in liability for a small engagement. Or if it's got a "work for hire" clause, which the company has if they want to share but not lose ownership.

Go with your gut

When you finally make your choice, go with the team you feel will bring results. Making a matrix of price versus service versus reputation versus estimated productivity is fine if you're into that, but consider going with your initial gut feeling. If you don't feel like there is a "right" company, go back to the table, get more information, and even try to get another bid or two. When you do find the best choice, things just "feel" right. How do you know?

- The way they talk about search engines makes sense and the advice they've given lines up with the best practices you've seen expressed on SEO sites you trust.
- They are familiar with recent trends at the engines and the links they send or news they cite is timely, relevant, and logical.

- They use the engines like pros, rapidly combing through site queries, link information, and analytics data to unveil the underlying problems that are hurting your rankings or causing the competition to succeed.
- Their interactions with your team “fit.” The team members like the consultants and are eager to implement their advice.

Choose someone whom you can be comfortable dropping your guard in front of. This is your expert and ideally your confidant, someone you can reach out to and talk to about challenges and knowledge gaps in a way that you cannot with in-house colleagues.

Remember, an SEO business relationship is a relationship, and if the two parties involved simply speak different languages (professionally and/or culturally) within their organizations, no amount of SEO expertise—and no amount of SEO budget—will make the relationship work. It is much smarter to work with people you like than to pay begrudgingly for a working relationship that simply doesn’t work. Cultural fit is essential, so be sure you can identify your own business culture and then look for a compatible, resonant culture within the SEO company/consultancy you hire.

Mixing In-House SEO with Outside SEO Help

When you are combining the resources of an external team with those of your in-house SEO team, there is a need to determine who does what, who decides what is to be done, and so on. It is best to have this all clearly defined up front, or else there will be conflicts and confusion, which will slow down the process of increasing your organic search traffic. Be sure to talk to the website program or project manager, not only upper management. In the end, it will be up to the program or project manager(s) to fit the changes into the schedule. Get their buy-in from the beginning. You may want to even involve them in the vendor selection process, because they will have to work with the vendor as well.

If we look first at the process of defining who does what, there are a few considerations, such as the following:

1. What areas of expertise does your in-house team have? How does the outsourced resource complement or reinforce that?
2. Is the outsourced resource bringing a unique new area of expertise to your organization, or is it a supplemental resource to an existing area of expertise?
3. Is the outsourced resource going to be working on an area where you consider it essential to build expertise in-house?
4. Can you leverage the outsourced resource with your internal resources? Perhaps you can use this to keep the outside help working on higher-value tasks that provide you a bigger return on your consulting dollar.

All of these considerations factor into the overall division of labor between your outsourced team and your internal team.

You also need to be clear about who is making the decisions and who is communicating those decisions to the outside SEO consultant. He will need guidance and direction on what you are looking to accomplish. Although you hopefully defined this in detail up front, it is likely to change over time, so updating that guidance on a regular basis is important. And as the outside consultant may not sit in your offices all day long, more effort may be required to communicate changes in plans to him than to your internal team.

In addition, the outsourced SEO consultant will have questions and will make recommendations, and he needs to know who to pass these on to and who will give him any related updates and let him know about decisions that are made. A clear communication channel is critical to success.

Building a Culture of SEO into Your Organization

SEO is a marketing function, and like any other marketing function, its goal is to help your organization grow and prosper. For a business, this means bringing in more revenue. What makes SEO different from any other function in your organization is the depth and breadth of its influence on every other role that influences and/or touches the website. SEO is a marketing channel that is unlike any other function or marketing channel you have seen; it spans many disciplines, including IT, usability/design, copywriting, content creation, public relations, and social media engagement.

SEO is also a bit of a strange beast because much of it is non-intuitive. As an SEO practitioner, you can explain things to people and they will look at you like you have three heads and say, “Huh?”

Building an SEO culture requires that you get the organization as a whole to accept a handful of important facts:

- Organic search engine traffic is essential to the growth and success of the organization.
- Search engines impose some specific constraints on how you can structure your website.
- Search engine success requires specialized marketing programs.
- Embrace these limitations and requirements, and you can gain a strategic advantage over your competition.

These concepts are pretty high-level and very simple by design. If you can get all parts of the organization to buy into them, though, it will have a very powerful impact on your chances for success. Once everyone agrees that SEO is needed, you will have taken the first step in building an enduring SEO culture into your organization.

The value of this is that you will have made all team members aware, at least at a basic level, of how SEO affects them. Before the senior manager makes a sweeping policy decision about something, she may pause to consider the SEO impact and get the advice she needs on the topic. Before the developer implements something a particular way, he will likely make sure he is doing the right thing from an SEO point of view. In this environment, your in-house SEO practitioner is invited to the right meetings, pulled into impromptu discussions, and called for quick answers as people are executing project tasks.

To build an SEO culture at your organization, you need to institutionalize SEO by integrating it into everyday business activities and making it part of everyone’s job description. Successful SEO will not be the result of a solo effort by your in-house SEO practitioner(s); it will come from collaboration with everyone who influences and touches the website.

Conclusion

As we have outlined throughout this chapter, SEO touches IT/development, marketing, public relations, usability/design, copywriting and editorial, content development, social media, sales, and management. Getting every one of those groups to understand the basics of SEO is essential to long-term success. SEO is not the be-all and end-all, but it is a significant component of all these disciplines.

SEO culture is focused on growing relevant, converting traffic to the website and the business overall. Educating your organization in these simple concepts will ensure that everyone is working toward the same objective and will bring increased efficiency to your overall operations, which will ultimately help you achieve SEO success.

An Evolving Art Form: The Future of SEO

As we have noted throughout this book, SEO is about leveraging your site's content assets through search engine-friendly content creation and website development (on-site), and targeted content promotion (off-site), in order to increase exposure and earn targeted traffic from organic search results. Therefore, the ultimate objective of the SEO professional is to make best use of organic search traffic as determined by various business goals of a website by guiding organizations through SEO strategy development, implementation, and ongoing measurement of SEO efforts.

This role will change as technology evolves, but the fundamental objectives will continue to remain the same as long as "search engines" with non-paid search results continue to exist. The complexity of search will continue to grow, as all search engines seek to locate and index all of the world's digital information. As a result, we expect various levels of expansion and evolution in search within the following areas:

- Continued rapid expansion of mobile search, and with it, voice search, as the world continues to increase its demand for this capability.
- Large scale investments in improved understanding of entities and relationships (semantic search) to allow search to directly answer more and more questions.
- Growth in social search as users begin to leverage social networks to discover new and interesting content and solutions to problems from their friends.
- Indexation of multimedia content, such as images, video and audio, including a better understanding of the content of these types of files.
- Indexing of data behind forms (something that Google already does in some cases, for example, if you use First Click Free).
- Continued improvements in extraction and analysis of Javascript and AJAX-based content.
- Increased localization of search.
- Expanded personalization capabilities.

Mobile search is already driving an increasing demand for linguistic user interfaces, including voice-recognition-based search. Voice search greatly improves the ease of use and accessibility of search on mobile devices and this technology will continue to evolve and improve. In October of 2011 Apple released the iPhone 4s with Siri. While most of its capabilities were already present in Google Voice Actions, Siri introduced a more conversational interface, and also showed some personality.

Business deals also regularly change the landscape. For example, on July 29, 2009, Microsoft and Yahoo! signed a far-reaching deal that resulted in Yahoo! retiring the search technology that powers Yahoo! Search and replacing it with Microsoft's Bing technology

(<http://www.nytimes.com/2009/07/30/technology/companies/30soft.html>). Bing also came to an agreement with Baidu to provide the English language results for the Chinese search engine (<http://searchenginewatch.com/article/2098918/Bada-Bing-Its-Baidu-Bing-English-Search-Marketing-In-China>).

After these deals, many contended that this deal would result in a more substantive competitor for Google. With Microsoft's deep pockets as well as a projected combined Bing/Yahoo U.S. market share of just over 30%, Bing could potentially make a formidable competitor. However, six years later it is not clear that there has been a significant shift in the search landscape as a result.

Since that time it's rumored that Yahoo has undertaken projects code-named "Fast-Break" and "Curveball" aimed at getting it back into the search game, but Yahoo's lack of adequate search technology makes it seem unlikely any major changes will be made soon. But, Yahoo is not quitting, as the November 20, 2014 deal they announced to be the default search supplier to Firefox shows (<http://www.reuters.com/article/2014/11/20/us-yahoo-firefox-idUSKCN0J32HP20141120>).

Perhaps the bigger shift may come from the continuing growth of Facebook, which reports having more than 1.3 billion users worldwide as of July 2014 (<http://www.businessinsider.com/facebook-inc-has-22-billion-users-2014-7>), including about half the population of US and Canada. In addition, Facebook has many users of WhatsApp (500 million), Instagram (200 million), and Messenger (another 200 million).

As we noted in Chapter 8, Bing's Stefan Weitz suggests that 90% of people use their friends to help them make one or more decisions every day, and 80% of people use their friends to help them make purchasing decisions (<http://www.zdnet.com/blog/facebook/bing-adds-more-facebook-features-to-social-search/1483>). In addition, as we showed in Chapter 8, Google+ has a material effect on the Google's personalized search results. If Google can succeed in growing Google+, the scope of this impact could grow significantly over the next few years.

Third Door Media's Danny Sullivan supports the notion that some amount of search traffic may shift to social environments. In a July 2011 interview with Eric Enge he said: "I think search has his cousin called discovery, which is showing you things that you didn't necessarily know you wanted or needed, but you are happy to have come across. I think social is very strong at providing that" (<http://www.stonetemple.com/danny-sullivan-on-google-facebook-twitter-social-and-search/>).

Neither Stefan Weitz or Danny Sullivan believe traditional web search is going away, but it seems likely that there will be some shifts as people discover other ways to get the information they want on the web.

In fact, Stefan Weitz believes that search will get embedded more and more into devices and apps, and is focusing much of their strategy in this direction, partly because they recognize that they won't win head to head with Google in what he calls the "pure search space" (http://www.theregister.co.uk/2014/11/04/stefan_weitz_bing_search_application_web_summit/).

These developments and many more will impact the role SEO plays within an organization. This chapter will explore some of the ways in which the world of technology, the nature of search, and the role of the SEO practitioner will evolve.

The Ongoing Evolution of Search

Search has come a long way, and will continue to progress at an increasingly rapid pace. Keeping up with these changes, the competitive environment, and the impact of new technology provides both a challenge and an opportunity.

The Growth of Search Complexity

Search has been evolving rapidly over the past decade and a half. At the WSDM conference (<http://www.wsdm2009.org/>) in February 2009, Google Fellow Jeff Dean provided some interesting metrics that tell part of the story:

- Google search volume had grown 1,000 times since 1999.
- Google has more than 1,000 times the machines it had in 1999.
- Latency dropped from less than 1,000 ms to less than 200 ms.
- Index update latency improved by about 10,000 times. Whereas updates took Google months in 1999, in 2009 Google was detecting and indexing changes on web pages in just a few minutes.

These are staggering changes in Google's performance power, and these are from six years ago. And of course this is just part of the changing search environment. Some of the early commercial search engines, such as Web Crawler, InfoSeek, and Alta Vista, launched in the mid-1990s. At that time, web search engines' relevancy and ranking algorithms were largely based on keyword analysis. This was a simple model to execute and initially provided pretty decent results.

However, there was (and is) too much money in search for such a simple model to stand. Spammers began abusing the weakness of the keyword algorithms by stuffing their pages with keywords, and using tactics to make them invisible to protect the user experience. This led to a situation in which the people who ranked first in search engines were not those who deserved it most, but were in fact those who understood (and could manipulate) the search algorithms the best.

By 1999 Google launched, and the next generation of search was born. Google was the search engine that most effectively implemented the concept of citation analysis (or link analysis) as part of a popular search engine. As we outlined earlier in the book, link analysis counted links to a website as a vote for its value. More votes represent more value, with some votes being worth more than others (pages with greater overall link authority have more juice to vote).

This created a situation that initially made the job of the spammer more difficult, but the spammers began to catch up with this advance by purchasing links. With millions of websites out there, many of them with little or no revenue, it was relatively easy for the spammer to approach a site and offer it a nominal amount of money to get a link. Additionally, spammers could implement bots that surfed the Web, finding guest books, blogs, and forums, and leaving behind comments with links in them back to the bot owner's site.

The major search engines responded to this challenge as well. They took two major steps, one of which was to build teams of people who worked on ways to detect spamming and either discount it (by not attributing value) or penalize it. The other was to implement an analysis of the quality of the links that goes deeper than just the notion of PageRank. Factors such as anchor text, relevance, and trust became important as well. These factors also helped the search engines in their war against spam.

But the effort to improve search quality as well as fight spammers continued. Historical search result performance, such as how many clicks a particular listing got and whether the user was apparently satisfied with the result she clicked on, are metrics that have made their way into search algorithms. In 2008, then-Yahoo! Chief Scientist Jan O. Pederson wrote a position paper (<http://www.ils.unc.edu/ISSS/papers/papers/pedersen.pdf>) that advocated use of this type of data as follows:

Search engine query logs only reflect a small slice of user behavior—actions taken on the search results page. A more complete picture would include the entire click stream; search result page clicks as well as offsite follow-on actions.

This sort of data is available from a subset of toolbar users—those that opt into having their click stream tracked. Yahoo! has just begun to collect this sort of data, although competing search engines have collected it for some time.

We expect to derive much better indicators of user satisfaction by considering the actions post click. For example, if the user exits the clicked-through page rapidly than one can infer that the information need was not satisfied by that page.

Throughout the last few years, Google, Bing, and Facebook have invested in making use of social signals as well. As we discussed in Chapter 8, links and mentions on social media sites could potentially be used to corroborate to the search engines that your content is worthy of ranking, though this does not appear to be in use today (other than Google Plus' impact on personalized search).

Consider also the notion of author authority. While Google killed the rel="author" tag, they can still create associations between a website and an author. If the author begins to establish a high level of authority, this could become a strong signal to the benefit of the site, even if that authority is established elsewhere.

In May 2007, Google made a big splash with the announcement of Universal Search (delivering what refer to as "Blended" search results in this book). This was the beginning of the integration of all types of web-based data into a single set of search results, with data from video, images, news, blogs, and shopping search engines all being integrated into a single search experience.

This was only the very beginning of improved indexing for multimedia content. Google already has the ability to recognize the content of images in many cases. As we discussed in Chapter 10, Google Image Search (<http://images.google.com>) already allows you to drag an image into the search box and it will try to recognize its content.

Search engines also can make use of other data sources, such as domain registry data to see who owns a particular website. In addition, they have access to analytics data, data from their browsers, data from their ad networks, data from their web search toolbars, and data from free Wi-Fi and Internet access distribution to track actual web usage on various websites. Other data sources will come into play too, such as data from wearables, such as Google Glass, embedded devices and even your thermostat (<http://www.wired.com/2014/01/googles-3-billion-nest-buy-finally-make-internet-things-real-us/>).

Although no one knows how, or how much, the search engines use this type of data, these are additional information sources at their disposal.

Search engines continue to look for more ways to improve the quality of search results, and of the overall user experience within search. Google's efforts toward personalization allow it to look at a user's search history to get a better idea of what results will best satisfy a particular user. In 2008, Danny Sullivan summarized this entire evolution into four phases (<http://searchengineland.com/search-4.0-putting-humans-back-in-search-14086>):

Search 1.0: keywords and text

Search 2.0: link analysis

Search 3.0: integration of vertical results

Search 4.0: personalization

So, what makes up Search 5.0? Increased use of social media data appears to be one major possibility. The "wisdom of the crowds" may be becoming a factor in ranking, as discussed in Chapter 8. Mike Grehan talks about this in his paper, "New Signals to Search Engines" (<http://www.acronym.com/new-signals-to-search-engines.html>). He summarizes the state of web search as follows:

We're essentially trying to force elephants into browsers that do not want them. The browser that Sir Tim Berners Lee invented, along with HTML and the HTTP protocol, was intended to render text and graphics on a page delivered to your computer via a dial-up modem, not to watch movies like we do today. Search engine crawlers were developed to capture text from HTML pages and analyze links between pages, but with so much information outside the crawl, is it the right method for an always-on, ever-demanding audience of self producers?

Universal Search was a step that acknowledged part of this problem by making all types of data available through web search. But many of these data types do not provide the traditional text-based signals that search engines rely on. Here is more from Mike Grehan's paper:

Signals from end users who previously couldn't vote for content via links from web pages are now able to vote for content with their engagement with that content. You can expect that these types of signals will become a significant factor in the future.

User engagement and interaction with content provides the search engines with information about the value of content beyond links and on an array of data types (such as images and video), and this gives the engines more tools to improve search quality. This type of data was first confirmed as a factor in the rankings of videos on sites such as YouTube, which has become the third largest search engine for the US (according to comScore, Hitwise, and Nielsen Online – behind the combined Bing/Yahoo results), and was confirmed as a factor in general web search in 2011. YouTube's ascent in search volume is particularly interesting

because it is not a general web search engine, but one that focuses on a specific vertical, that of videos. This speaks to demand shifts taking place among the consumers of search results.

At the end of the day, the best results are likely to be provided by the best sites (there are exceptions; for example, for some search queries the best results may be “instant answers”). The technology that has driven today’s engines were based on two ranking signals: good keyword targeting and good links, but user experience and engagement are now significant factors too.

More data collection means more opportunity to win, even if your site doesn’t conform flawlessly to these signals, and a better chance that if there is only one indicator you’re winning, you could be in big trouble. Keywords and links will likely remain the primary ranking factors for the next few years, but the evolution of search engines in the direction of using the new signals is steadily gaining momentum and strength.

Following these advances, what will be next? Artificial intelligence (AI) already plays a huge role in search. As Peter Norvig indicated to Eric Enge in an October 2011 interview (<http://www.stonetemple.com/search-algorithms-with-google-director-of-research-peter-norvig/>):

If you define AI as providing a course of action in the face of uncertainty and ambiguity, based on learning from examples, that’s what our search algorithm is all about.

Further evidence of Google’s fascination with AI can be seen in the TED presentation given by Google cofounder Larry Page “Where’s Google Going Next?” in March 2014 (http://www.ted.com/talks/larry_page_where_s_google_going_next). In that presentation he discussed how Google was applying machine learning to allow a computer to do things like play video games or, on their own, extrapolate the concept and nature of an entity like a cat. Google’s goal is to apply that kind of learning to allow machines to understand the nature of information and how to present it to searchers.

One example of this is how Google Translate works. It examines and analyzes millions of real world translations to learn how one language translates into another. It learns by example, rather than trying to learn an artificial set of rules. Learning a language based on grammar does not work as it turns out that language is far too dynamic and changing. But, learning from real world usage does. Using this technology Google can offer instant translation across 58 languages.

Voice search works much the same way. Historically, speech recognition solutions did not work very well and required the user to train the system to their voice. Google Voice uses a different approach, as noted by Peter Norvig in that same interview:

“for Voice Search, where you speak your search to Google, we train this model on around 230 billion words from real-world search queries”.

Google’s Dominance

Thousands of posts, news articles, and analysis pieces have covered the central topic of battling Google’s dominance in web search, but few have discussed the most telling example of the search giant’s superiority. Many believe that the key to Google’s success, and more importantly, a key component in its corporate culture, is its willingness and desire to return information relevant to a searcher’s query as soon as possible.

Some also believe that Google’s biggest impact in the search engine market is its advertising platform, which is the world’s largest. By expanding its search capabilities, it is able to create a more enticing advertising platform through AdWords, AdSense, and its embeddable Google Search box. Its incredible infrastructure allows for site speed and crawling depth that is unmatched by any other engine.

However, it goes a bit deeper than that. In late 2008, tests were performed in which users were asked which search engine’s results they preferred for a wide variety of queries—long tail searches, top-of-mind searches, topics about which their emotions ranged from great passion to total agnosticism. They were shown two sets of search results and were asked which they prefer (see Figure 14-1).

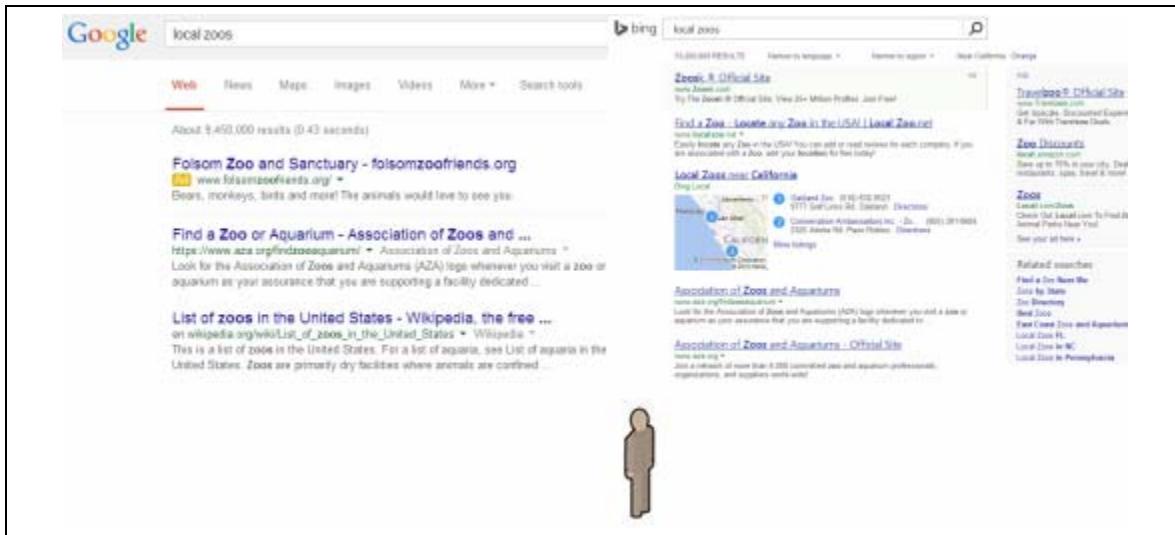


Figure 15-1. Comparing Google and Bing results

Numerous tests such as this one have been performed, with a wide variety of results. In some, the brands are removed so that users see only the links. Testers do this to get an idea of whether they can win from a pure “quality” standpoint. In others, the brands remain to get an unvarnished and more “real-world” view. And in one particular experiment—performed many times by many different organizations—the results are swapped across the brands to test whether brand loyalty and brand preference are stronger than qualitative analysis in consumers.

It is this last test that has the most potentially intriguing results. Because in virtually every instance where qualitative differences weren’t glaringly obvious, Google was picked as the best “search engine” without regard for the results themselves (see Figure 14-2).

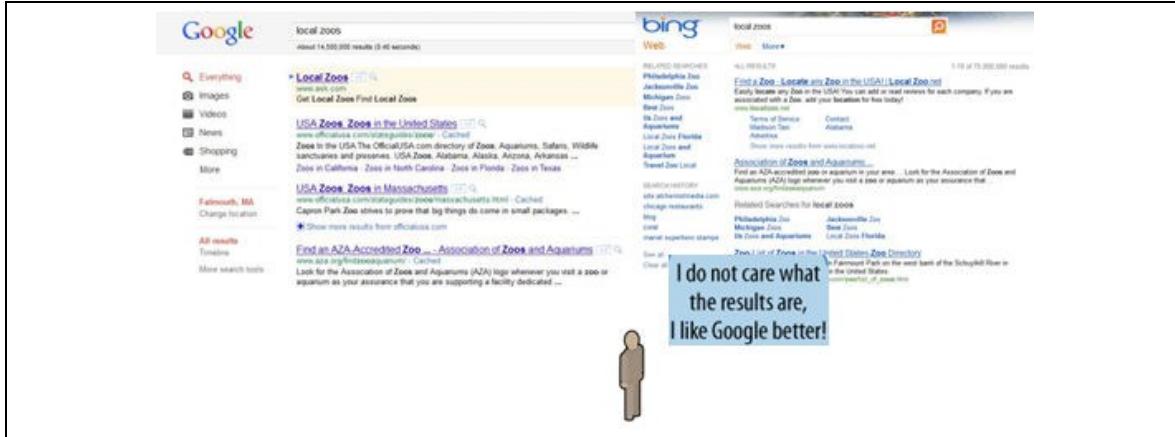


Figure 15-2. Results indicating that users may have a strong emotional preference for Google

Fundamentally, testers find (again and again) that the brand preference for Google outweighs the logical consideration of the search results quality.

Search engines that plan to take market share from Google are going to have to think differently. If Microsoft or a start-up search engine wants to take market share, it’s going to have to think less like a technology company trying to build a better mousetrap and more like a brand trying to win mind share from a beloved competitor. How did Pepsi take share away from Coke? Or Toyota from Ford? That is beyond the scope of this book, but it is a process that can take more than a great idea or great technology. It requires a massive psychological shift in the way people around the world perceive the Google brand against its competition.

One strategy that Bing has embarked on is to get close to Facebook and leverage Facebook data in the Bing results, and this is shown in the right rail of the results page as a separate set of results. On October 13, 2010 the two companies signed a deal (<http://blogs.bing.com/search/2010/10/13/bing-gets-more-social-with-facebook/>) that provides Bing with substantial access to Facebook's data on Shares, Likes, and more.

In addition, as noted earlier, Bing's Stefan Weitz talked about their focus on embedding search technology into devices and apps:

“The question is, where is search really going? It’s unlikely we’re going to take share in [the pure search] space, but in machine learning, natural language search... and how we can make search more part of living. For us, it’s less about Bing.com, though that’s still important. It’s really about how we can instead weave the tech into things you’re already doing.”

Google's social media network, Google+, is popular and growing fast. While Google+ can't be considered a runaway success, Google is continuing to invest heavily in it (<http://recode.net/2014/10/07/new-google-head-david-besbris-were-here-for-the-long-haul-qa/>). Google has a large team of people working on it, and are likely to be working on significant new features (or a new direction) for the platform. As shown in Chapter 8 ("How Social Media and User Data Play a Role in Search Results and Rankings") Google+ has a significant impact on personalized search.

Also consider the official Google mission statement: “Google’s mission is to organize the world’s information and make it universally accessible and useful.” It is already moving beyond that mission. For example, Google and NASA are working on new networking protocols that can work with the long latency times and low bandwidth in space (http://www.nasa.gov/home/hqnews/2008/nov/HQ_08-298_Deep_space_internet.html, <https://sites.google.com/site/dtnresgroup/home>).

Google is also pursuing alternative energy initiatives (<http://www.google.com/green/>), which clearly goes beyond its mission statement. An example of such an initiative is their investments in self-driving cars. In addition, Google has ventures in office productivity software with Google Docs (<http://docs.google.com>). These two initiatives have little to do with SEO, but they do speak to how Google is trying to expand its reach.

Another recent venture for Google is their exploration of wearable technology, primarily in the form of Google Glass (<https://www.google.com/glass/start/>). Though Google Glass has recently fallen out of favor with much of the technology community (<http://www.technologyreview.com/featuredstory/532691/google-glass-is-dead-long-live-smart-glasses/>), Glass was likely intended only as a first foray into the realm of wearable tech. Despite the fact that Glass itself may have been released too soon, and into much of a “beta” format, it dominated the early wearable search technology market. As wearable technology becomes more pervasive expect Google to continue to work to take a leading role in that market just as they have in mobile technology with Android.

Another potential future involves Google becoming a more general-purpose pattern-matching and searching engine. The concept of performing pattern matching on text (e.g., the current Google on the current Web) is only the first stage of an evolving process. Imagine the impact if Google turns its attention to the human genome and creates a pattern-matching engine that revolutionizes the way in which new medicines are developed. And the plethora of potential applications of pattern-matching in the “real world” (such as disaster preparedness and logistics) is astounding; Google is uniquely poised to capitalize on this global opportunity with its understanding of our physical world from the ground (e.g. by Google’s self-driving cars), from in our homes (e.g. Google’s Nest acquisition) and from the sky via satellites and drones (e.g. by Google’s recent acquisitions of Skybox Imaging and Titan Aerospace).

More Searchable Content and Content Types

The emphasis throughout this book has been on providing the crawlers with textual content semantically marked up using HTML. However, the less accessible document types—such as multimedia, content behind forms, and scanned historical documents—are increasingly being integrated into the search engine

results pages (SERPs) as search algorithms evolve in the ways that the data is collected, parsed, and interpreted. Greater demand, availability, and usage also fuel the trend.

Engines Will Make Crawling Improvements

The search engines are breaking down some of the traditional limitations on crawling. Content types that search engines could not previously crawl or interpret are being addressed.

In May 2014, Google announced that it had substantially improved the crawling and indexing CSS and JavaScript content (<http://googlewebmastercentral.blogspot.com/2014/05/understanding-web-pages-better.html>). Google can now render a large number of webpages with the JavaScript turned on so that they see it much more like the average user would. In October of 2014, Google updated their Webmaster guidelines to specifically advise that you do not block crawling of Javascript and CSS files (<http://googlewebmastercentral.blogspot.com/2014/10/updating-our-technical-webmaster.html>).

Despite the improvements there are still many challenges to Google fully understanding all of the content which is within JavaScript or CSS particularly if the crawlers are blocked from your JavaScript or CSS files, if your code is too complex for Google to understand, or if the code actually removes content from the page rather than adding it. Google still recommends that you build your site to "degrade gracefully" (http://www.w3.org/wiki/Graceful_degradation_versus_progressive_enhancement) which essentially means to build the site such that all of your content is available to users who have JavaScript turned either on or off.

Another major historical limitation of search engines is dealing with forms. The classic example is a search query box on a publisher's website. There is little point in the search engine punching in random search queries to see what results the search engines return. However, there are other cases in which a much simpler form is in use, such as a form that a user may fill out to get access to a downloadable article.

Search engines could potentially try to fill out such forms, perhaps according to a protocol where the rules are predefined to gain access to such content in a form where they can index it and include it in their search results. A lot of valuable content is currently isolated behind such simple forms, and defining such a protocol is certainly within the realm of possibility (though it is no easy task, to be sure). This is another area addressed by Google's November 2011 announcement (<http://googlewebmastercentral.blogspot.com/2011/11/get-post-and-safely-surfacing-more-of.html>). In more and more scenarios you can expect Google to fill out forms to see the content that exists behind them.

Engines Are Getting New Content Sources

As we noted earlier, Google's stated mission is "to organize the world's information and make it universally accessible and useful." This is a powerful statement, particularly in light of the fact that so much information has not yet made its way online.

As part of its efforts to move more data to the Web, in 2004 Google launched an initiative to scan in books so that they could be incorporated into a Book Search (<http://books.google.com/>) search engine. This became the subject of a lawsuit by authors and libraries, but a settlement was reached in late 2008 (<http://books.google.com/googlebooks/agreement/>). In addition to books, other historical documents are worth scanning. To aid in that Google acquired reCAPTCHA (e.g., see <http://www.google.com/recaptcha>), and in December of 2014, Google announced a major enhancement to how reCAPTCHA works (<http://techcrunch.com/2014/12/03/googles-recaptcha-mostly-does-away-with-those-annoying-captchas/>), with the goal of making it much more user friendly.

Similarly, content owners retain various other forms of proprietary information that is not generally available to the public. Some of this information is locked up behind logins for subscription-based content. To provide such content owners an incentive to make that content searchable, Google came up with its First Click Free concept (discussed earlier in this book) which is a program allowing Google to crawl subscription-based content.

Another example of new sources is metadata, in the form of markup such as schema.org, Mircoformats, and RDFa. This type of data, which is discussed in "Schema.org and Microformats" in Chapter 6, is a way

for search engines to collect data directly from the publisher of a web site. Schema.org was launched as a joint initiative of Google, Bing, and Yahoo to collect publisher-supplied data, and the number of formats supported can be expected to grow over time (<http://googleblog.blogspot.com/2011/06/introducing-schemaorg-search-engines.html>).

Another approach to this would be to allow media sites and bloggers to submit content to the search engines via RSS feeds. This could potentially speed indexing time and reduce crawl burden at the same time. One reason why search engines may not be too quick to do this though is that web site publishers are prone to making mistakes, and having procedures in place to protect against those mistakes might obviate the benefits.

However, a lot of other content out there is not on the Web at all, and this is information that the search engines want to index. To access it, they can approach the content owners and work on proprietary content deals, and this is also an activity that the search engines all pursue.

Another direction they can go with this is to find more ways to collect information directly from the public. Google Image Labeler was a program designed to do this. It allowed users to label images through a game where they would work in pairs and try to come up with the same tags for the image as the person they were paired with. Unfortunately, this particular program was discontinued, but other approaches like it may be attempted in the future.

Multimedia Is Becoming Index-able

Content in images, audio, and video is currently not easily indexed by the search engines, but its metadata (tags, captioning, descriptions, geotagging data) and the anchor text of inbound links and surrounding content make it visible in search results. Google has made some great strides in this area. In an interview with Eric Enge, Google's director of research Peter Norvig discussed how Google allows searchers to drag an image from their desktop into the Google Images (<http://images.google.com>) search box and Google can recognize the content of the image. You can see the interview here (<http://www.stonetemple.com/search-algorithms-with-google-director-of-research-peter-norvig/>).

Or consider <http://www.google.com/recaptcha>. This site is was originally used by Google to complete the digitization of books from the Internet Archive and old editions of the *New York Times*. These have been partially digitized using scanning and OCR software. OCR is not a perfect technology and there are many cases where the software cannot determine a word with 100% confidence. However, reCAPTCHA is assisting by using humans to figure out what these words are and feeding them back into the database of digitized documents.

First, reCAPTCHA takes the unresolved words and puts them into a database. These words are then fed to blogs that use the site's CAPTCHA solution for security purposes. These are the boxes you see on blogs and account sign-up screens where you need to enter the characters you see, such as the one shown in Figure 14-3.



Figure 15-3. Re CAPTCHA screen

In this example, the user is expected to type in *morning*. However, in this case, Recaptcha.net is using the human input in these CAPTCHA screens to translate text that wasn't recognized by OCR. It makes use of this CAPTCHA information to improve the quality of this digitized book.

Google used this methodology for years, but has since change reCAPTCHA to focus more on images instead (<http://techcrunch.com/2014/12/03/googles-recaptcha-mostly-does-away-with-those-annoying-captchas/>). This new approach is image based, and is intended to help Google with its computer vision projects.

Similarly, speech-to-text solutions can be applied to audio and video files to extract more data from them. This is a relatively processing-intensive technology, and it has historically had trouble with many accents and dialects, so it has not yet been universally applied in search. Siri on the iPhone and Google Voice are leading the charge in addressing this issue. In addition, if you upload a video to YouTube, you can provide a search-indexable caption file for it, or request that Google use its voice recognition technology to attempt to auto-caption it.

The business problem the search engines face is that the demand for information and content in these challenging-to-index formats is increasing exponentially. Search results that do not accurately include this type of data will begin to be deemed irrelevant or wrong, resulting in lost market share and declining ad revenues.

The dominance of YouTube is a powerful signpost of user interest. Users want engaging, multimedia content - and they want a lot of it. For this reason, the work on improved techniques for indexing such alternative content types is an urgent priority for the search engines.

Interactive content is also growing on the Web, with technologies such as AJAX. In spite of the indexing challenges these technologies bring to search engines, the use of these technologies is continuing because of the experience they offer for users who have broadband connectivity. The search engines are hard at work on solutions to better understand the content wrapped up in these technologies as well.

Over time, our view of what is “interactive” will likely change dramatically. Two- or three-dimensional first-person shooter games and movies will continue to morph and become increasingly interactive. Further in the future, these may become full immersion experiences, similar to the Holodeck on “Star Trek.” You can also expect to see interactive movies where the audience influences the plot with both virtual and human actors performing live. These types of advances are not the immediate concern of today’s SEO practitioner, but staying in touch with where things are headed over time can provide a valuable perspective.

Search Becoming More Personalized, Localized and User-Influenced

Personalization efforts have been underway at the search engines for some time. As we discussed earlier in this book, the most basic form of personalization is to perform an IP location lookup to determine where the searcher is located, and tweak the results based on the searcher’s location. However, the search engines continue to explore additional ways to expand on this simple concept to deliver better results for each user. It is not yet clear whether personalization has given the engines which have heavily invested in it better results overall or greater user satisfaction, but their continued use of the technology suggests that, at the least, their internal user satisfaction tests have been positive.

Indeed Google has continued to expand the factors which can influence a users’ personalized search. For example, one major signal they use is the user’s personal search history. Google can track sites a user has visited, groups of related site’s user has visited, whether a user has shared a given site over social media, and what keywords they user has searched for in the past. All of these factors may influence the given personalized search engine results page.

Determining User Intent

As mentioned above, Google personalized results are tapping into user intent intent, based on previous search history, and serving up a mix of not just a set of personalized “blue links” but many content types, including maps, blog posts, videos, and local results. The major search engines already provide maps for appropriate location searches and the ability to list blog results based on recency as well as relevancy. It is not just about presenting the results, but about presenting them in the format that maps to the searcher’s intent.

User Interactions

One area that will see great exploration will be in how users interact with search engines. As the sheer amount of information in its many formats expands, users will continue to look to search engines to be not just a search destination, but also a source of information aggregation whereby the search engine acts as a portal, pulling and updating news and other content based on the user's preferences.

Marissa Mayer, then Google's VP of Location and Local Services (now CEO of Yahoo!), made a particularly interesting comment that furthers the sense that search engines will continue their evolution beyond search:

I think that people will be annotating search results pages and web pages a lot. They're going to be rating them, they're going to be reviewing them. They're going to be marking them up...

Indeed, Google already offers users the ability to block certain results. Mayer's mention of "web pages" may be another reason why the release of Google Chrome (<http://www.google.com/chrome>) was so important. Tapping into the web browser might lead to that ability to annotate and rate those pages and further help Google identify what content interests the user. As of February 2014, StatCounter showed that Chrome's market share had risen to an impressive 44% (<http://gs.statcounter.com/>).

Although Chris Sherman, executive editor of *Search Engine Land* offered up an interesting interactive approach that the search engines might pursue, as a way to allow users to interact with search engines and help bring about better results:

Until search engines can find a way to let us search by example - submitting a page of content and analyzing the full text of that page and then tying that in conjunction with our past behavior...

New Search Patterns

This is all part of increasing the focus on the users, tying into their intent and interests at the time of search. Personalization will make site stickiness ever more important. Securing a position in users' history and becoming an authoritative go-to source for information will be more critical than ever. Winning in the SERPs will require much more than just optimizing for position, moving towards an increased focus on engagement.

Over time, smart marketers will recognize that the attention of a potential customer is a scarce and limited quantity. As the quantity of information available to us grows, the amount of time we have available for each piece of information declines, creating an attention deficit. How people search, and how advertisers interact with them, may change dramatically as a result.

In 2008 *The Atlantic* published an article titled, "Is Google Making Us Stupid?" (<http://www.theatlantic.com/doc/200807/google>). The thrust of this article was that Google was so powerful in its capabilities that humans need to do less (and less!). Google has made huge advances since this 2008 article, and this is a trend that will continue. After all, who needs memory when you have your "lifestream" captured 24/7 with instant retrieval via something akin to "Google Desktop Search" or when you have instant perfect recall of all of human history?

These types of changes, if and when they occur, could transform what we today call "SEO" into something else, where the SEO of tomorrow is responsible for helping publishers gain access to potential customers through a vast array of new mechanisms that currently do not exist.

Growing Reliance on the Cloud

Cloud computing is transforming how the Internet-connected population use computers. Oracle founder Larry Ellison's vision of thin-client computing may yet come to pass, but in the form of a pervasive Google OS "operating system" and its associated, extensive suite of applications. Widespread adoption by users of cloud-based (rather than desktop) software and seemingly limitless data storage, all supplied for free by Google, will usher in a new era of personalized advertising within these apps.

Google is actively advancing the mass migration of desktop computing to the cloud, with initiatives such as Google Docs & Spreadsheets (<http://docs.google.com>), Gmail (<http://mail.google.com>), Google Calendar (<http://calendar.google.com>), Google App Engine (<http://developers.google.com/appengine/>), and Google Drive (<https://drive.google.com/>). These types of services lead to users entrusting their valuable data to the Google cloud. This brings them many benefits (but also concerns around privacy, security, uptime, and data integrity). In May of 2011 Apple also made a move in this direction when they announced iCloud (<http://www.apple.com/icloud>), which is seamlessly integrated into Apple devices.

One simple example of a basic application for Cloud computing in the notion of backing up all your data. Most users don't do a good job of backing up their data, making them susceptible to data loss from hard drive crashes and virus infections. Companies investing in cloud computing will seek to get you to store the master copy of your data in the cloud, and keep backup copies locally on your devices (or not at all). With this approach you can more easily access that information from multiple computers (e.g. their work and home computers).

Google (and Apple) benefit by having a repository of user data available for analysis—which is very helpful in Google's quest to deliver ever more relevant ads and search results. It also provides multiple additional platforms within which to serve advertising. Furthermore, regular users of a service such as Google Docs are more likely to be logged in a greater percentage of the time when they are on their computers.

The inevitable advance of cloud computing will offer more and more services with unrivaled convenience and cost benefits, compelling users to turn to the cloud for their data and their apps.

Increasing Importance of Local, Mobile, and Voice Search

New forms of vertical search are becoming increasingly important. Areas that the search engines have already embraced include local search and mobile search. In addition, voice based search is an area in which all the major engines are actively investing.

Local Search

Local search was an active business before the advent of the Internet, when Yellow Pages as supplied by your phone company were the tool of choice. As the World Wide Web first became prominent in our lives, the first way that this unfolded on the Web was in the form of Yellow Pages migrating online. These types of sites have themselves evolved to leverage more of the unique nature of the Web. Some of the major players are CitySearch (<http://www.citysearch.com>), Local.com (<http://local.com>), YellowPages.com (<http://www.yellowpages.com>), and Superpages (<http://www.superpages.com>).

As the major search engines evolved, they integrated sophisticated mapping technology into their systems, and then began to map in local business data as well. You can find these local search engines at <http://maps.google.com/>, <http://local.yahoo.com>, and <http://www.bing.com/maps>. These search engines map in some of their business data from the same types of data sources that drive the Yellow Pages websites, such as Acxiom (<http://www.acxiom.com>), InfoUSA (<http://www.infousa.com>), and Localeze (<http://www.localeze.com>), but they also supplement that data by crawling the web.

One of the big challenges facing the local search engines is *map spam*, the proliferation of spam results in local search. Here are URLs to two articles that provide examples:

- <http://blumenthals.com/blog/2010/12/07/illusory-laptop-repair-a-most-elegant-googleplaces-hack/>
- <http://www.thatsseo.com/2009/04/google-not-good-enough-in-controlling-local-map-spam/>

Resolving these issues is obviously critical for the search engines and is a major area of investment for them. We believe that local search is an area whose importance will continue to grow. With a very large number of searches having local intent, this is potentially a major frontier for developments in search.

Over time, you can expect that these problems will be worked out. Local search is already extremely powerful, but you can anticipate that the search engines will control the flow of customers in cyberspace and the real world, because customers will be finding merchants through search. You can also expect that more and more of these local searches will take place on mobile devices. Success in local search will be what makes or breaks many businesses.

Consumers are becoming increasingly reliant on local search and its auxiliary services—street maps, directions, satellite views, street views, 3D visualizations (Google Earth), enhanced listings, user reviews, and ratings. Through Google Mobile’s image recognition (formerly Google Goggles), Google even has the capability of searching based on photos you take with your mobile device. Augmented reality, where we see meta data juxtaposed on top of the physical world via a “looking glass” is around the corner, and Google Glass is just a hint of what is to come.

FourSquare (<https://foursquare.com/>) can find nearby friends, and search for special offers and discounts if you “check in” at certain stores or restaurants. Tools like FourSquare guide consumers’ movements and their decision-making processes; this means the search engines and other search services are in a powerful position to manage the flow of customers from the Web to brick-and-mortar businesses. So, in a way, the search engines are in the logistics business, building the connection between online and offline.

Mobile Search

As of the end of 2013, there were approximately 6.8 billion mobile phone subscriptions (<http://qz.com/179897/more-people-around-the-world-have-cell-phones-than-ever-had-land-lines/>) As a result, the opportunity for mobile search has begun to grow rapidly. A Google study showed that search engines were the most visited sites on mobile devices with 77% of those surveyed reporting their use (<http://www.clickz.com/clickz/column/2072688/stats-mobile-search>). The same study showed that 9 out of 10 searchers took action as a result of conducting a search.

In addition, as of October 2014, Google began experimenting with using the mobile friendliness of a site as a ranking factor (<http://searchengineland.com/google-may-add-mobile-user-experience-ranking-algorithm-205382>). This provides strong incentive to have a mobile search strategy in place (see the “Mobile Search” section of Chapter 9 for more on this)!

The convenience of being able to get the information you need while on the go is just way too compelling. Why be tied to a desk if you do not have to be? Further, in many countries, freedom from a desktop is a requirement because the space that people have for their personal lives is quite different from what people experience in the United States.

The ever-expanding versatility and power of tablet and mobile devices—from indispensable utility apps, to immersive multimedia players, to massively multiplayer online games, to paradigm-shifting hardware advances such as the iPhone’s multitouch display, proximity sensor, GPS, and gyroscopes—will fuel this growth. The *network effect*, whereby the value of the network grows by the square of the size of the network (Metcalf’s Law), gives further incentive for users to migrate to their mobile devices as more and more apps allow them to interact with their peers in increasingly more interesting ways.

The small keyboard/typing surface is currently a limitation, but Apple and Google’s voice-based solutions have already made great strides in replacing the keyboard as the input device of choice. The coming advent of the Linguistic User Interface (LUI) will continue to revolutionize mobile search, and with the LUI comes a whole new set of skills that will be required of the SEO practitioner.

In addition, wearable devices, such as Google Glass, or smart watches are gaining increasing traction. Voice technology is improving to the point where the absence of a keyboard will not be a significant limitation.

Voice Recognition Search

When users are mobile they must deal with the limitations of their mobile device, specifically the small screen and small keyboard. These make web surfing and mobile searching more challenging than they are in the PC environment.

Voice search is a great way to improve the mobile search experience. It eliminates the need for the keyboard, and provides users with a simple and elegant interface. Speech recognition technology has been around for a long time, and the main challenge has always been that it requires a lot of computing power.

Several examples already exist:

Google Voice Search

<http://www.google.com/insidesearch/features/voicesearch/>

This is a free service from Google that enables you to perform a Web search based on text-to-speech voice queries instead of traditional text-typed queries.

Google Mobile App for iPhone

Included in this application is voice searching capability. You can speak your query into the application and the results are displayed on-screen.

Apple Siri

<http://www.apple.com/ios/siri/>

What made Siri a significant step forward after its launch was its use of more natural human speech (Google offers similar capabilities now). This was the start of a movement towards the device learning the user, rather than the user learning the device.

Microsoft Cortana

<http://www.microsoft.com/en-us/mobile/campaign-cortana/>

In April 2014, Microsoft unveiled its personal assistant offering, known as Cortana. Cortana is being promoted as offering more capabilities for learning your personal preferences. Currently, it is designed to work only on Windows phones.

Processing power continues to increase, even on mobile devices, and the feasibility of this type of technology is growing. This should be another major area of change in the mobile search landscape.

In addition, it is reasonable to expect that voice recognition technology will be applied to the actual recorded content of audio and video files, to determine content and aid in ranking these media types within search results, something that will likely lead to video “script” optimization as an added component of video SEO.

Increased Market Saturation and Competition

One thing you can count on with the Web is continued growth. Despite its constantly growing index, a lot of the pages in Google may be low-quality or duplicate-content-type pages that will never see the light of day. The Web is a big place, but one where the signal-to-noise ratio is very low.

One major trend emerges from an analysis of Internet usage statistics. According to Miniwatts Marketing Group, 84.9% of the North American population uses the Internet (<http://www.internetworldstats.com/stats.htm>) (Update?), so there is not much room for growth there. In contrast, Asia, which already has the most Internet users with 1.4 billion has a penetration rate of only 34.7%. Other regions with a great deal of opportunity to grow are Africa, the Middle East, and Latin America.

This data is telling us that in terms of the number of users, North America is already approaching saturation. Europe has some room to grow, but not that much. However, in Asia, you could see two times that growth, or 2 to 3 billion users! The bottom line is that a lot of Internet growth in the coming decade will be outside North America and that will provide unique new business opportunities for those who are ready to capitalize on that growth.

With this growth has come an increasing awareness of what needs to be done to obtain traffic. The search engines are the dominant suppliers of traffic for many publishers, and will continue to be for some time to

come. For that reason, awareness of SEO will continue to increase over time. Here are some reasons why this growth has continued:

The Interwebs outperform other sales channels

When organizations look at the paths leading to sales and income (a critical analysis whenever budgets are under scrutiny), the Web almost always comes out with one of two assessments. Either it is a leading sales channel (especially from an ROI perspective) or it is deemed to be an area with the greatest opportunity for growth. In both scenarios, web marketing and, in correlation, SEO, takes center stage. Cite this data: <https://econsultancy.com/blog/64614-email-remains-the-best-digital-channel-for-roi#i.c22gfeezgcx7yg>

It is the right time to retool

Established companies frequently use down cycles as a chance to focus attention inward and analyze themselves. Consequently, there's a spike in website redesigns and SEO along with it.

Paid search drives interest in SEO

Paid search spending is still reaching all-time highs, and when companies evaluate the cost and value, there's a nagging little voice saying "75%+ of the clicks do not even happen in the ads; use SEO."

SEO is losing its stigma

Google is releasing SEO guides, Microsoft and Yahoo! have in-house SEO departments, and the "SEO is BS" crowd have lost a little of their swagger and a lot of their arguments. No surprise—solid evidence trumps wishful thinking, especially when times are tough.

Marketing departments are in a brainstorming cycle

A high percentage of companies are asking the big questions: "how do we get new customers?" and "what avenues still offer opportunity?" Whenever that happens, SEO is bound to show up near the top of the "to be investigated" pile.

Search traffic will be relatively unscathed by the market

Sales might drop, conversion rates might falter a bit, but raw search traffic isn't going anywhere. A recession doesn't mean people stop searching the Web, and with broadband adoption rates, Internet penetration, and searches per user consistently rising, search is no fad. It is here for the long haul.

Web budgets are being reassessed

We've all seen the news about display advertising falling considerably; that can happen only when managers meet to discuss how to address budget concerns. Get 10 Internet marketing managers into rooms with their teams and at least four or five are bound to discuss SEO and how they can grab that "free" traffic.

Someone finally looked at the web analytics

It is sad, but true. When a downturn arrives or panic sets in, someone, maybe the first someone in a long time, checks the web analytics to see where revenue is still coming in. Not surprisingly, search engine referrals with their exceptional targeting and intent matching are ranking high on the list.

Although more and more people are becoming aware of these advantages of SEO, there still remains an imbalance between paid search and SEO. The SEMPO Annual State of Search Survey (<http://www.sempo.org/general/custom.asp?page=CurrentStateofSearch>, membership is required to access the report) includes information that suggests that as much as 90% of the money invested in search related marketing is spent on PPC campaigns and only about 10% goes into SEO.

This suggests that either SEO could see some growth to align budgets with potential opportunity, or firms that focus solely on SEO services had better diversify. SEO budgets continue to expand as more and more businesses better understand the mechanics of the Web. In the short term, PPC is easier for many businesses to understand, because it has more in common with traditional forms of marketing. Ultimately,

though, SEO is where the most money can be found, and the dollars will follow once people understand that.

SEO as an Enduring Art Form

Today, SEO can be fairly easily categorized as having five major objectives:

- Make content accessible to search engine crawlers.
- Find the keywords that searchers employ (understand your target audience) and make your site speak their language.
- Build content that users will find useful, valuable, and worthy of sharing. Ensure that they'll have a good experience on your site to improve the likelihood that you'll earn links and references.
- Earn votes for your content in the form of editorial links and social media mentions from good sources by building inviting, shareable content and applying classic marketing techniques to the online world.
- Creating web pages that allow users to find what they want extremely quickly, ideally in the blink of an eye.

Note, though, that the tactics an SEO practitioner might use to get links from editorial sources have been subject to rapid evolution. We now talk about content marketing, instead of link building. In addition, a strong understanding of how the search engines measure and weight social engagement signals is increasingly important to SEO professionals.

One thing that you can be sure about in the world of search is change, as forces from all over the Web are impacting search in a dramatic way. To be an artist, the SEO practitioner needs to see the landscape of possibilities for an individual website, and pick the best possible path to success. This currently includes social media optimization expertise, local search expertise, video optimization expertise, an understanding of what is coming in mobile, and more. That's a far cry from the backroom geek of the late 1990s.

No one can predict what the future will bring and what will be needed to successfully market businesses and other organizations on the Web in two years, let alone 5 or 10. However, you can be certain that websites are here to stay for a long time, and that websites are never finished and need continuous optimization just like any other direct marketing channel. SEO expertise will be needed for a long time—and no existing person is better suited to map the changing environment and lead companies to success in this new, ever-evolving landscape than today's SEO practitioner.

The Future of Semantic Search and the Knowledge Graph

In Chapter 6, we explored the state of Semantic Search and the Knowledge Graph as we know it today. All the search engines are continuing to investigate these types of technologies in many different ways, though Google is clearly in the lead. The Knowledge Vault is just one of many initiatives that Google is pursuing to make progress in this area.

Part of the objective is to develop a machine intelligence that can fully understand how people evaluate the world, yet even this is not sufficient. The real goal is to understand how **each** human being evaluates the world, so that the results can be fully personalized to meet each individual's needs.

Not only do search engines want to give you the perfect answer to your questions, they also want to provide you with opportunities for exploration. Humans like to conduct research and learn new things. Providing all of these capabilities will require a special type of machine intelligence, and we are a long way from reaching those goals.

There are many components that go into developing this type of intelligence.

In the near future, efforts focus largely on solving specific problems. For example, one such problem is maintaining the context of an ongoing conversation. Consider the following set of queries, starting with “where is the Empire State Building?” as shown in Figure 15-4.

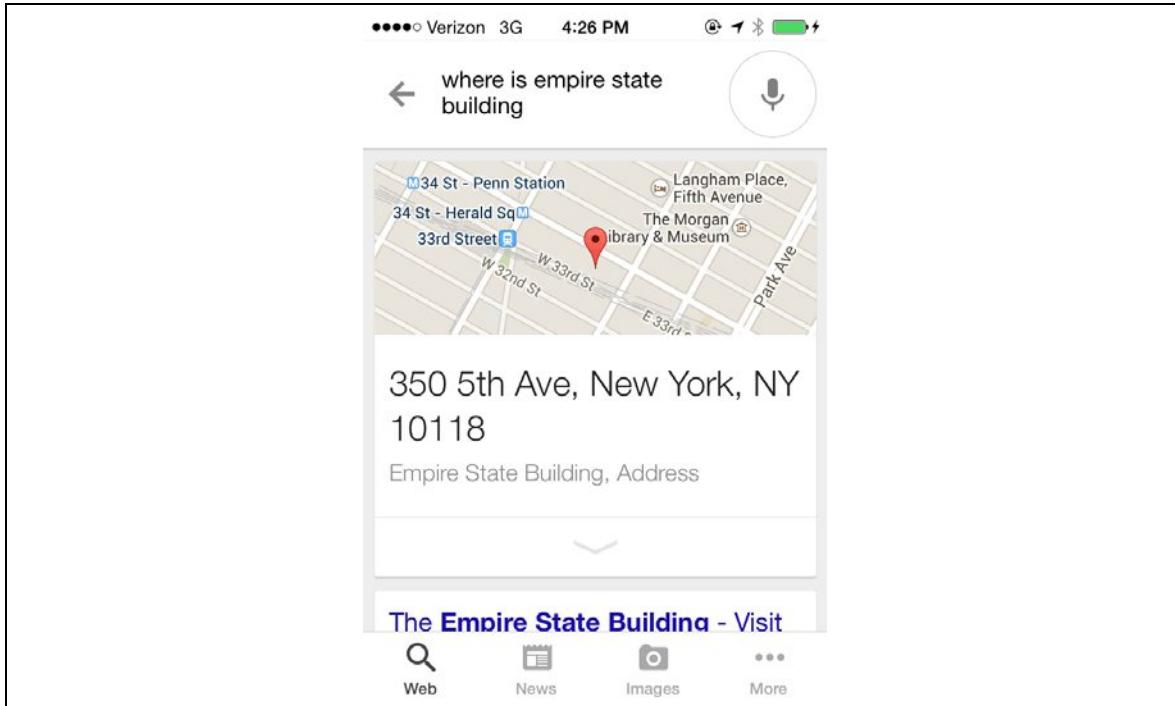


Figure 15-4. Response to the Query “Where is the Empire State Building?”

Notice how the word “the” was dropped in the way the query is shown. Figure 15-5 Shows what happens when you follow this query with the query “pictures”.

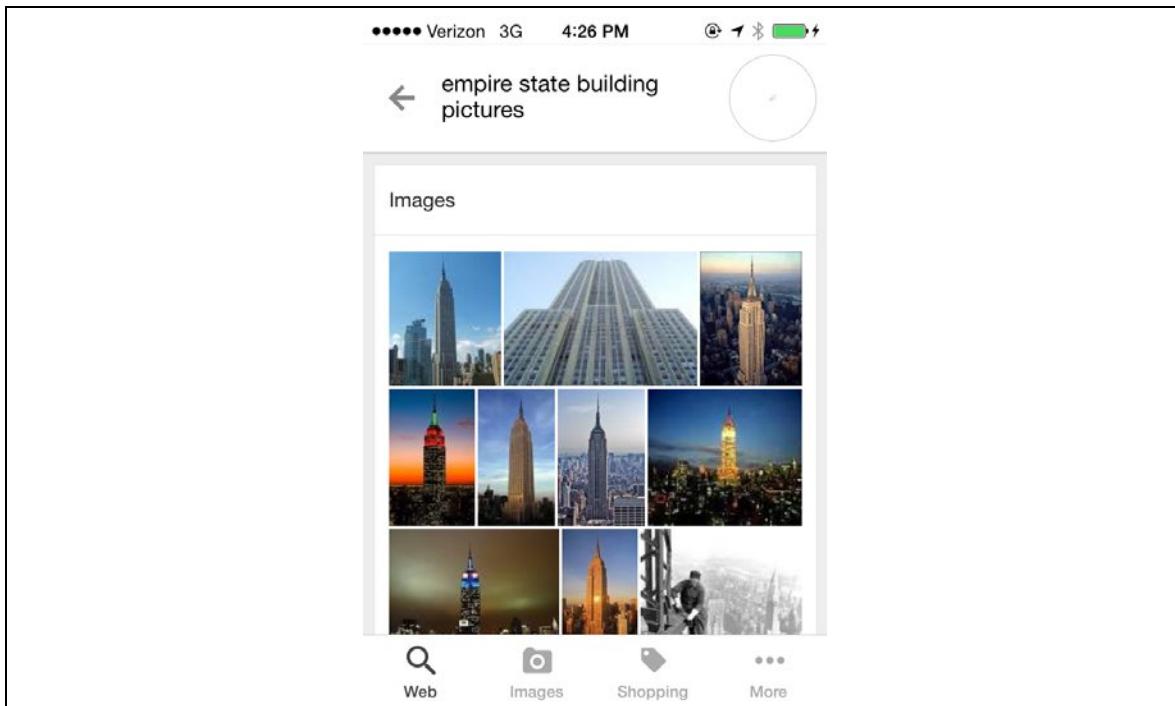


Figure 15-5. Response to the query “pictures”

Notice again how the query was modified to “empire state building pictures”. Google has remembered that the prior query was specific to the Empire State Building, and did not require us to restate that. This query sequence can continue for quite some time. Figure 15-6 shows the result when we now ask the query “who built it”.

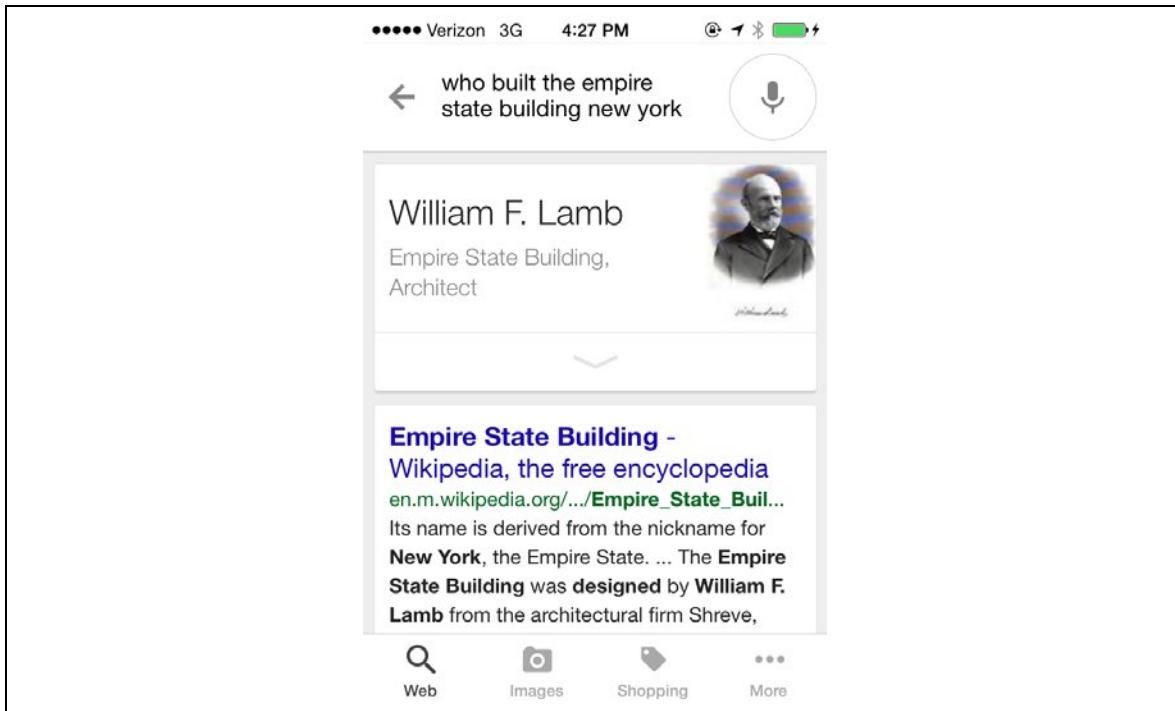


Figure 15-6. Response to the query “who built it”

Once again the query was dynamically modified on the fly, and Google has remembered the context of the conversation. Figure 15-7 shows what happens when we now try the query “restaurants”.

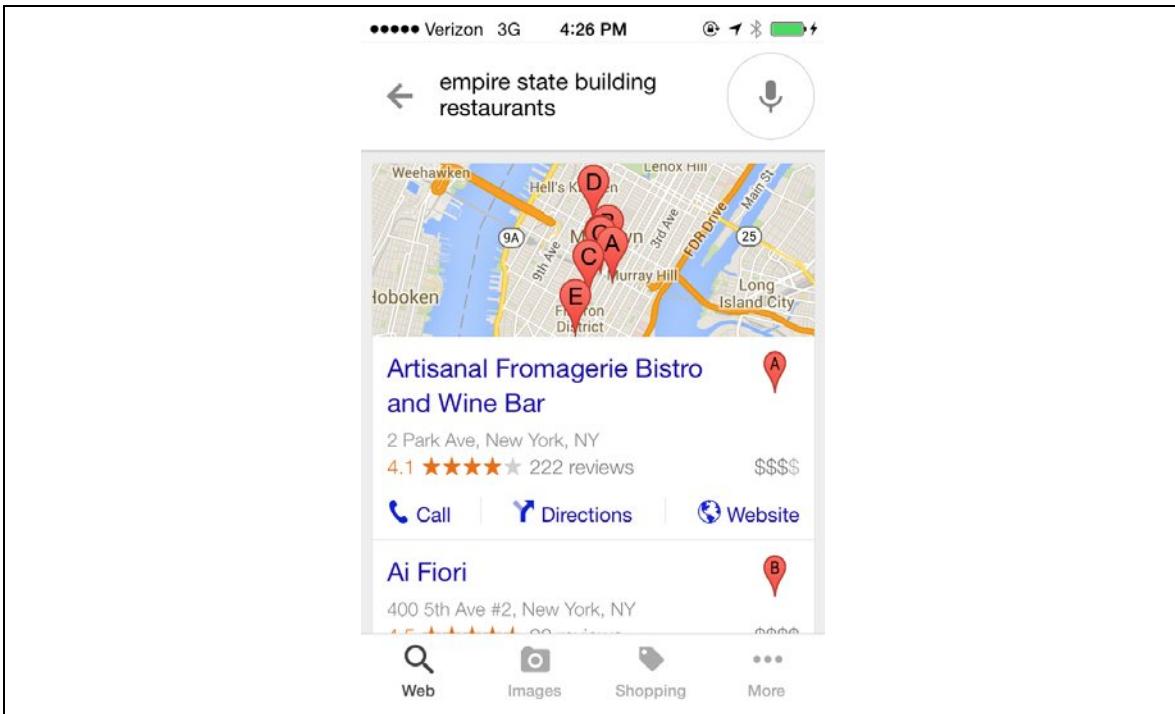


Figure 15-7. Response to the query “restaurants”

Finally, we can follow this query with the more complex query “give me directions to the third one”, as shown in Figure 15-8.

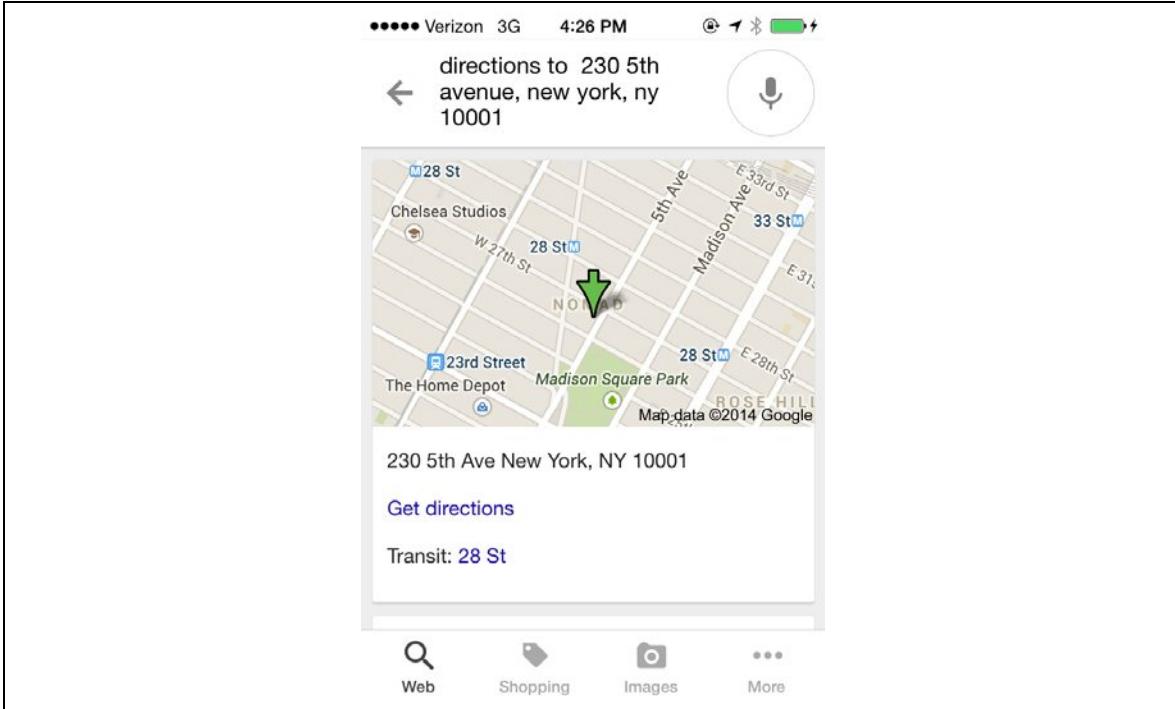


Figure 15-8. Response to the query “give me directions to the 3rd one”

This entire sequence of queries is quite complicated, as capped by Google understanding the concept of the “third one” in the final query. Even though this is very sophisticated, it is nonetheless this is an example of a point solution to a specific problem.

To truly model human thought, search engines will need to build machines that can reason like humans, are able to perceive the world around them, understand how to define objectives, and make plans to meet them, and can independently work to expand their knowledge.

Many disciplines are involved in developing artificial intelligence, such as computer science, neuroscience, psychology, philosophy, and linguistics. Even just understanding linguistics is a major challenge, as there are thousands of different languages in the world, and this by itself multiplies the complexity of building

The computing power to take on these challenges does not yet exist, so developing expanding capabilities is a major part of the challenge. For example, Google is pursuing efforts to build a quantum computer (<http://www.technologyreview.com/news/530516/google-launches-effort-to-build-its-own-quantum-computer/>).

In the near term, we can expect changes in search results to come in the form of more point solutions to specific problems. As the understanding of how to model human intelligence expands, and as processing power grows with it, we may see much more significant changes, perhaps 5 to 10 years down the road.

Conclusion

SEO is both art and science. The artistic aspect of SEO requires dynamic creativity and intuition. The search engine algorithms are too complex to reverse-engineer every aspect of them. The scientific aspect involves challenging assumptions, analyzing data, testing hypotheses, making observations, drawing conclusions, and achieving reproducible results. These two ways of thinking will remain a requirement as SEO evolves into the future.

In this chapter we conveyed some sense of what is coming in the world of technology, and in particular, search. Although the previous decade has seen an enormous amount of change, the reality is that it has simply been the tip of the iceberg. There's a lot more change to come, and at an ever-increasing (exponential) rate. If the Law of Accelerating Returns holds, we're in for a wild ride.

In this fast-moving industry, the successful SEO professional has to play the role of early adopter. The early adopter is always trying new things—tools, tactics, approaches, processes, technologies—to keep pace with the ever-evolving search engines, ever-increasing content types, and the ongoing evolution of online user engagement.

It is not enough to adapt to change. You will need to embrace it and evangelize it. Many in your (or your client's) organization may fear change, and steering them through these turbulent waters will require strong leadership. Thus, the successful SEO professional also has to play the role of change agent.

The need for organizations to capture search mindshare, find new customers, and promote their messaging will not diminish anytime soon, and neither will the need for searchable, web-based, and instantaneous access to information, products, and services. This ability—to connect the seeker with the provider—is perhaps the most valuable skill set on the Web today. And although there is the possibility that the search engines could eventually be called “decision,” “dilemma,” or even “desire” engines, the absolute need for the understanding of and interactions between both the psychological and the technological natures of search will ensure that SEO as a discipline, and SEO professionals, are here to stay.